

1 Two Architectures for Simulating Biomarker-Treatment
2 Interactions: Implications for Statistical Power Under
3 Carryover

4 pmsimstats team

5 2026-07-31

6 **Contents**

7	1 Abstract	2
8	2 1. Introduction	3
9	3 2. Two DGP Architectures	8
10	3.1 2.1 Architecture A: Direct Mean Moderation	8
11	3.2 2.2 Architecture B: Differential Correlation (MVN)	9
12	3.3 2.3 Mathematical Comparison	11
13	4 3. Consequences for Statistical Power Under Carryover	12
14	4.1 3.1 Empirical demonstration	12
15	4.2 3.2 Why the architectures diverge	16
16	4.3 3.3 Robustness check at $N = 70$	18
17	5 4. Biological Assumptions	20
18	5.1 4.1 When Architecture A is appropriate	21
19	5.2 4.2 When Architecture B is appropriate	22
20	5.3 4.3 The prazosin-PTSD case	25
21	6 5. Implications for Simulation Study Design	27
22	6.1 5.1 Choosing the appropriate architecture	27
23	6.2 5.2 Reporting requirements	27
24	6.3 5.3 When the choice is uncertain	29
25	7 6. Alternative Analysis Strategies Under Architecture B	29
26	7.1 6.1 Restrict analysis to on-drug observations	29
27	7.2 6.2 Weighted analysis	30
28	7.3 6.3 Within-subject contrast	31
29	7.4 6.4 Two-stage random slopes	31
30	7.5 6.5 Exclusion of contaminated observations	32
31	7.6 6.6 Design-level solutions	33

32	7.7 6.7 Qualitative comparison	34
33	8 7. Relationship to the Broader Literature	34
34	8.1 7.1 Treatment effect heterogeneity	34
35	8.2 7.2 Prevalence of each architecture	36
36	8.3 7.3 Crossover and N-of-1 trial methodology	38
37	8.4 7.4 Precision medicine trial design	38
38	9 8. Conclusions	41
39	10 Conflict of interest	42
40	11 Funding	42
41	12 Data availability statement	42
42	13 Reproducibility	43
43	14 References	43

44 1 Abstract

45 **Background.** Simulation studies of predictive biomarker-treatment interactions
46 in N-of-1 clinical trials require a data-generating process (DGP) that specifies how
47 the biomarker moderates treatment response. Two fundamentally different DGP
48 architectures are used in practice, and the choice between them determines the
49 extent to which carryover effects reduce statistical power to detect the interaction.
50 The published simulation literature does not distinguish them.

51 **Methods.** We formalize two architectures: direct mean moderation (Architec-
52 ture~A), under which the biomarker scales the treatment effect in the popula-
53 tion mean structure; and differential correlation (Architecture~B, the multivari-
54 ate normal approach), under which the interaction emerges from treatment-state-
55 dependent correlation between biomarker and response in the joint distribution.
56 Power consequences under both architectures are quantified by Monte Carlo sim-
57 ulation on an aggregated N-of-1 design framework with continuous-time AR(1)
58 residual correlation in the analysis model. The biomarker, treatment phase, and
59 carryover decay are parameterized to the prazosin/PTSD reference application.

60 **Results.** Under Architecture~A, power is largely preserved under carryover
61 because the proportional relationship between biomarker and drug response is
62 maintained at every timepoint. Under Architecture~B, carryover erodes the dif-
63 ferential correlation signal and reduces power. At the prazosin-calibrated refer-
64 ence cell ($N = 35$ per randomization path, i.e. effective analysis $N = 70$ for
65 the crossover (CO) and open-label-with-blinded-discontinuation (OL+BDC) de-
66 signs and 140 for Hybrid after path concatenation; $c_{bm} = 0.45$, OL+BDC de-
67 sign) across $t_{1/2} \in \{0, 0.5, 1.0\}$ weeks, Architecture~B power drops from 0.777
68 to 0.539 (a 30.6% relative loss; $z = 7.64$, $p < 10^{-13}$ on the B-versus-A con-

trast), while Architecture~A drops only from 0.751 to 0.730 (a 2.8% relative loss). The B-versus-A absolute power-loss gap of 21.7 percentage points at OL+BDC, $t_{1/2} = 1.0$ is approximately 11 times the matched Architecture~A loss; the Hybrid design produces a smaller but still significant gap ($z = 3.96$, $p < 10^{-4}$); the traditional crossover design produces no detectable gap because its within-subject AR(1) structure dominates the carryover-mediated correlation channel that Architecture~B uses. The 36-cell production factorial (2 architectures $\times$ 3 designs $\times$ 3 $t_{1/2} \times 2 c_{bm}$ levels) was run at 1{,}000 replicates per cell with 100% convergence; the Monte Carlo standard error (MCSE) on cell power is 0.013-0.016. Type~I error sits in $[0.010, 0.074]$, conservative relative to the nominal level and with no architecture-specific inflation. The conservatism reflects the `corCAR1` working-covariance miscalibration characterized in a companion calibration study¹; being common to both architectures, it does not affect the architecture contrast.

Conclusions. The two architectures encode different biological mechanisms, heterogeneity at the mean structure versus heterogeneity at the variance structure, and the choice between them is a substantive scientific question rather than a parameterization convenience. We recommend explicit declaration of the DGP architecture in any biomarker-validation simulation study and provide guidance for selecting the architecture appropriate to a given clinical context.

2 1. Introduction

[1]

The central question is biomarker-treatment interaction; simulation studies need a DGP that encodes it.

- Does a baseline biomarker predict differential treatment response, i.e., does it moderate the treatment effect?
- Closed-form power is available only in simplified cases [`@pmsimstats-paper08`]; otherwise simulation is the principal power evaluation tool for proposed trial designs.
- These simulations require a DGP that explicitly encodes the biomarker-treatment interaction mechanism.

rgt to complete.

A central question in precision medicine trial design is whether a baseline biomarker predicts differential treatment response, that is, whether the treatment effect is moderated by the biomarker value. For simplified special cases, such as a balanced strict repeated-measures ANOVA under sphericity, the power of the interaction test admits a closed-form expression, derived together with its equivalence to a

two-sample t -test in a companion paper². For the continuous-time mixed model and the N-of-1 designs considered here, however, no such closed form is available, so simulation studies are the principal tool for evaluating the statistical power of proposed trial designs to detect such interactions, and these simulations necessarily require a data-generating process (DGP) that encodes the biomarker-treatment interaction mechanism.

[2]

Every simulation framework must choose how to generate treatment effect heterogeneity, and that choice has consequences.

- Heterogeneity of treatment effects (HTE) literature (2005, 2020) distinguishes qualitative interactions (direction reverses) from quantitative ones (magnitude varies).
- Any simulation framework must decide how to encode this heterogeneity in the DGP.
- That decision is substantive, not a mere parameterization detail.

rgt to complete.

The statistical literature on treatment effect heterogeneity, developed notably by Rothwell³ and synthesized more recently by Kent et al.⁴, distinguishes between *qualitative* interactions, in which the direction of the treatment effect reverses across biomarker strata, and *quantitative* interactions, in which the magnitude varies but the direction remains consistent. As we shall argue, any simulation framework intended to evaluate power for biomarker-moderated effects must choose how to generate this heterogeneity, and that choice is not merely a technical convenience.

[3]

Architecture A (direct mean moderation) is the near-universal standard; the Hendrickson (2020) multivariate normal (MVN) approach is apparently unique.

- Every major simulation framework uses direct mean moderation: Simon (2010) for stratified designs; Freidlin and Korn (2014) for adaptive enrichment; Renfro and Sargent (2017) for basket/umbrella trials; all N-of-1 frameworks.
- Wang and Schork (2019) and Schork (2022) also use Architecture A despite sharing authorship with Hendrickson et al. (2020), suggesting the MVN parameterization is a paper-specific choice, not a lab convention.
- To the best of our knowledge, no other clinical trial simulation framework uses differential correlation (MVN) to encode the

148

biomarker-treatment interaction.

149

rgt to complete.

150

151

152

153

154

155

156

157

158

159

160

161

162

163

164

165

166

167

168

169

170

171

172

173

174

175

[4]

176

177

178

Architecture B is principled, not merely idiosyncratic: it is the reduced-form second-moment encoding of a latent responder-class mechanism.

179

180

181

182

183

184

185

186

- If the biomarker is a proxy for an unobserved responder subtype, the faithful DGP is a finite mixture; a single MVN reproduces its first two moments.
- Architecture B is thus one point on a continuum from continuous-heterogeneity to explicit latent-class formulations.
- When to prefer it, and whether the approximation is adequate, is taken up in Section 4.2 and formally in a companion paper [pmsimstats-paper03].

187

rgt to complete.

188

189

Although the differential-correlation parameterization is unusual in the simulation literature, it is not arbitrary. It is the reduced-form,

second-moment representation of a latent responder-class mechanism: if the biomarker does not itself determine the drug’s effect but instead indexes an unobserved subpopulation that responds (a disease subtype, an inflammatory state, a genetic background), the biologically faithful DGP is a finite mixture, and a single multivariate normal reproduces that mixture’s first two moments through a treatment-state-dependent biomarker-response correlation. On this reading Architecture~B is not a competitor to mean moderation on equal footing but the covariance-only end of a continuum that runs through mean-and-covariance and location-scale specifications to an explicit latent-class model. When this reduced form is preferable to direct mean moderation, and when the approximation breaks down, is taken up in Section~4.2 and is the subject of a companion paper¹⁴ that develops this biomarker-moderation spectrum and evaluates class-aware analyses at trial-relevant sample sizes.

In plain language, and without the latent-variable vocabulary: suppose patients come in two kinds, those who respond to the drug and those who do not, but we cannot see which kind any given patient is. The biomarker is only an imperfect clue to that hidden status, so two patients with the same biomarker value can still respond differently. Architecture~A cannot capture this, because it assumes the biomarker fixes the size of each patient’s drug effect exactly. Architecture~B captures it by letting the biomarker and the response track each other only while the drug is acting: a higher biomarker tends to go with a larger response, as a statistical tendency rather than a fixed rule. When the drug washes out that tendency fades, which is the reason carryover costs more power under Architecture~B than under Architecture~A.

[5]

The comparison originates not in an error but in recognizing that Hendrickson’s Architecture B is appropriate for its latent-subtype biomarker, whereas the N-of-1 simulation literature defaults to Architecture A.

- Hendrickson et al. (2020) use Architecture B, the correct encoding for a biomarker hypothesized to proxy a latent responder subtype (resting systolic blood pressure for a noradrenergic PTSD subtype).
- Most N-of-1 and biomarker-trial simulation work instead encodes the interaction through Architecture A.
- The divergence became concrete when two implementations of the same pmsimstats pipeline, one using B and one A, reached different conclusions about carryover.

rgt to complete.

Architecture~B is the appropriate data-generating choice for the trial that motivated Hendrickson et al.¹³. The predictive biomarker there is resting systolic blood pressure¹⁵, which is hypothesized to index a latent noradrenergic subtype of PTSD preferentially responsive to alpha-1 adrenergic blockade¹⁶ rather than to determine prazosin exposure. On that hypothesis the biomarker is a proxy for an unobserved responder class, exactly the latent responder-class mechanism for which the differential-correlation encoding is the reduced form (Section~4.3). What brought the architectural distinction into focus was the realization that this choice is not shared by the wider N-of-1 and biomarker-trial simulation literature, which, as reviewed above, encodes the interaction almost exclusively through direct mean moderation (Architecture~A). The mismatch became concrete during development of the pmsimstats framework: the original Hendrickson-derived code generated data under Architecture~B, while a parallel reimplementation adopted Architecture~A, and the two produced qualitatively different conclusions about how carryover erodes power. The distinction is therefore not a coding artifact to be corrected but a substantive modeling choice. The same biomarker-moderation question, encoded in the first versus the second moment of the outcome distribution, yields materially different power under carryover, and the encoding that is appropriate depends on the biology of the biomarker.

[6]

This paper is the first systematic comparison of the two DGP architectures; we formalize, simulate, and give guidance.

- We formalize the distinction between Architecture A (mean moderation) and Architecture B (differential correlation).
- We demonstrate power consequences through Monte Carlo simulation.
- We offer guidance for investigators on how to choose the appropriate architecture.

rgt to complete.

We present here what we believe to be the first systematic comparison of these two DGP architectures for biomarker-treatment interaction power estimation under carryover. We shall formalize the distinction between them, demonstrate its consequences through Monte Carlo simulation, and offer guidance for investigators designing simulation studies for predictive biomarker-moderated treatment effects.

Related documents in this series. The mathematical derivation, code-level implementation details, four implementation attempts (three failed and one correct), and the empirical evaluation of the original Hendrickson¹³ parameter grid under Architecture~A are documented in

19-mean-moderation-implementation-notes.tex (with a brief markdown
summary in 19-mean-moderation-implementation-notes.md). A condensed
empirical-results summary suitable for collaborator presentation is provided in
21-architecture-comparison-summary.md.

279

3 2. Two DGP Architectures

Notation. Throughout this paper, B_i denotes the biomarker value of participant i , D_{it} the binary on-drug indicator at timepoint t (1 on drug, 0 off drug), and BR_{it} the biological response component of the outcome. These three symbols describe the data-generating-process layer. The analysis-model layer uses a related but distinct continuous drug exposure variable, `Dbc`, defined in Section 3.1. Code-style identifiers in typewriter font (`bm`, `Dbc`, `BR`) refer to specific variables in the `pmsimstats` codebase or to R formula syntax; mathematical symbols are used otherwise.

3.1 2.1 Architecture A: Direct Mean Moderation

In this approach the biomarker enters the mean structure of the outcome directly. The treatment effect for participant i at timepoint t is:

$$Y_{it} = \mu_0 + \beta_1 \cdot D_{it} \cdot (1 + \beta_{bm} \cdot B_i) + \epsilon_{it}$$

where D_{it} is the binary on-drug indicator (or a continuous measure of drug exposure), B_i is the participant's biomarker value (typically centered), and β_{bm} is the biomarker moderation parameter. The interaction is explicit in the mean: participants with higher biomarker values have proportionally larger treatment effects.

Key properties:

- The interaction is **deterministic** given the biomarker value and treatment status.
- The biomarker moderation β_{bm} is a **population-level parameter** that applies uniformly to all participants.
- The signal for the interaction is in the **first moment** (conditional mean) of the outcome distribution.
- Adding noise, random effects, or carryover to D_{it} scales the entire treatment effect (including its biomarker-dependent component) proportionally. The **ratio** of drug effect between high-biomarker and low-biomarker participants is preserved.

The multiplicative and additive forms of Architecture A are interchangeable for power calculations.

- The multiplicative form maps directly to the standard additive regression, the effect-modeling equation of Kent et al. (2020).
- $\beta_3 = \beta_1 \cdot \beta_{bm}$; the additive form additionally allows a prognostic main effect.
- Both forms fall within Architecture A and are interchangeable for power calculations on the interaction coefficient.

rgt to complete.

We note that the multiplicative form above maps directly onto the standard additive regression form for predictive biomarker effects used in the trial methodology literature, specifically the effect modeling equation of Kent et al.⁴:

$$Y_{it} = \beta_0 + \beta_1 D_{it} + \beta_2 B_i + \beta_3 D_{it} B_i + \epsilon_{it}$$

The interaction coefficient β_3 in the additive form corresponds to the product $\beta_1 \cdot \beta_{bm}$ in the multiplicative form. The additive form additionally accommodates a prognostic main effect β_2 for the biomarker, which the multiplicative form sets to zero by construction. For the purpose of power calculations on the interaction coefficient, the two forms are interchangeable up to this prognostic-effect distinction, and both fall within Architecture A.

This architecture is used in the pedagogical simulation (`implementations/simple/simulation.R`) and in the exploratory carryover analyses (`simulation_carryover_spectrum.R`) of the pmsimstats project.

3.2 2.2 Architecture B: Differential Correlation (MVN)

In this approach the biomarker and the biological response (BR) component are jointly drawn from a multivariate normal distribution with treatment-state-dependent correlation:

$$\begin{pmatrix} B_i \\ BR_{it} \end{pmatrix} \sim \text{MVN} \left(\begin{pmatrix} \mu_B \\ \mu_{BR}(t, D_{it}) \end{pmatrix}, \Sigma(D_{it}) \right)$$

where the covariance matrix Σ has the following treatment-state-dependent BM-BR correlation:

$$\text{Cor}(B_i, BR_{it}) = \begin{cases} c_{bm} & \text{if } D_{it} = 1 \text{ (on drug)} \\ c_{bm} \cdot e^{-\lambda \cdot t_{sd}} & \text{if } D_{it} = 0 \text{ (off drug)} \end{cases}$$

where t_{sd} is time since drug discontinuation. The exponential-decay form is the default; the limiting cases $\lambda = 0$ (no decay; correlation persists indefinitely during the off-drug phase) and $\lambda \rightarrow \infty$ (immediate decay; correlation drops to zero the instant the drug is discontinued) recover the two no-carryover model variants.

[8]

In Architecture B the interaction is in the covariance structure, not the mean; it emerges from the conditional distribution given the biomarker value.

- The population mean of BR_{it} is the same for all participants; no mean-structure interaction exists.
- Given biomarker value b , the conditional expectation of BR scales with c_{bm} , creating an interaction-like structure.
- Correlation is c_{bm} on drug and decays as $c_{bm}e^{-\lambda t_{sd}}$ off drug; moderation attenuates during the off-drug period.

rgt to complete.

The interaction is not in the mean structure: the population mean of BR_{it} is the same for all participants with the same treatment history. Instead it emerges from the **conditional distribution**: given a participant's biomarker value $B_i = b$, the conditional expectation of their biological response is:

$$E[BR_{it}|B_i = b, D_{it} = 1] = \mu_{BR}(t) + c_{bm} \cdot \frac{\sigma_{BR}}{\sigma_B} \cdot (b - \mu_B)$$

This conditional expectation has an interaction-like structure: the biomarker value b modulates the expected BR. The effective correlation is c_{bm} on drug and $c_{bm} \cdot e^{-\lambda t_{sd}}$ during the off-drug period; the moderation therefore persists with attenuated strength during off-drug periods and vanishes only in the limit of complete decay.

Key properties:

- The interaction is **probabilistic**: it manifests as a tendency, not a deterministic scaling.
- The biomarker moderation c_{bm} is a **correlation parameter** that governs the strength of the association.
- The signal for the interaction is in the **second moment** (covariance structure) of the joint distribution.
- Carryover effects in the DGP inflate the BR means during off-drug periods, making them resemble on-drug means. This reduces the **differential** in the correlation structure between treatment states, weakening the signal that the analysis model detects.

This architecture is used in the Hendrickson et al.¹³ publication code and the audited R/ package in pmsimstats-ng.

A third possibility is that both channels operate at once: the biomarker could scale the mean drug effect (Architecture~A) while simultaneously altering the biomarker-response correlation during active treatment (Architecture~B). Such a combined architecture, carrying independent weights on the two channels, recovers A and B as its single-channel limits and is the natural framework when the operative mechanism is uncertain. Because it introduces a second free parameter and its power behavior warrants separate treatment, we develop it in a companion paper¹⁷ and restrict the present paper to the A-versus-B contrast.

3.3 2.3 Mathematical Comparison

[9]

Under Architecture A, carryover shrinks the signal amplitude but preserves the ratio between biomarker groups, so identifiability is maintained.

- The expected outcome difference between two participants scales linearly with D .
- Under carryover, D is replaced by a decayed value and the absolute magnitude shrinks.
- But the ratio of drug effect between high- and low-biomarker participants is preserved at every exposure level; the signal contracts but remains identifiable.

rgt to complete.

We begin the comparison by considering the expected outcome difference between two participants with biomarker values b_1 and b_2 at the same drug exposure level D under Architecture A:

$$E[Y \mid b_1, D] - E[Y \mid b_2, D] = \beta_1 \cdot \beta_{bm} \cdot (b_1 - b_2) \cdot D$$

The absolute magnitude of this difference scales linearly with D and therefore shrinks during off-drug periods when D is replaced by an exposure-decayed value $D \cdot \text{carryover}$. We note, however, that the **ratio** of the drug effect between high-biomarker and low-biomarker participants is preserved at every level of exposure. The biomarker-treatment signal contracts in amplitude but does not lose its identifiability under carryover.

[10]

Under Architecture B, the biomarker-dependent component of the off-drug response decays exponentially, eventually eliminating the differential signal.

- Off-drug, the conditional expectation includes a biomarker-dependent term scaled by $c_{bm}e^{-\lambda t_{sd}}$.
- As t_{sd} grows, that term shrinks toward zero.
- In the limit, all participants converge to the same off-drug mean regardless of biomarker value; the interaction signal vanishes.

rgt to complete.

Under Architecture B, the analogous quantity is:

$$E[BR|b_1, D = 1] - E[BR|b_2, D = 1] = c_{bm} \cdot \frac{\sigma_{BR}}{\sigma_B} \cdot (b_1 - b_2)$$

During off-drug periods with carryover:

$$E[BR|b, D = 0, t_{sd}] = \mu_{BR, \text{off}}(t_{sd}) + c_{bm} \cdot e^{-\lambda t_{sd}} \cdot \frac{\sigma_{BR}}{\sigma_B} \cdot (b - \mu_B)$$

Here the biomarker-dependent component decays exponentially with time off drug. As t_{sd} increases, the off-drug conditional expectation converges to $\mu_{BR, \text{off}}$ for all biomarker values, eliminating the differential signal entirely.

4 3. Consequences for Statistical Power Under Carryover

4.1 3.1 Empirical demonstration

[11]

The two DGPs are identical except for how they encode the BM-BR linkage; any power divergence is attributable solely to that difference.

- All non-biologic components (time-varying, placebo-belief, residual noise) are generated identically in both architectures.
- The only difference: Architecture B encodes BM-BR interaction in Σ ; Architecture A adds a post-draw mean shift to on-drug BR.
- The same `nlme::lme` analysis model with `corCAR1` is used for both, so divergence in results is attributable to the DGP alone.

rgt to complete.

We ran identical simulation configurations under both architectures using the `pmsimstats-ng` framework, matching the DGP parameters as

closely as possible between the two approaches. We note that the non-biologic response components (time-varying response, placebo-belief accumulation, and residual noise) are generated from the same multivariate normal draw with identical means, variances, within-factor AR(1) autocorrelation, and cross-factor correlations in both architectures. The two DGPs differ *only* in how the biomarker-to-biological-response (BM-BR) linkage is encoded: Architecture~B embeds it in the off-diagonal c_{bm} entries of Σ , whereas Architecture~A omits those entries and adds a post-draw shift $c_{bm} \cdot B_i^* \cdot \sigma_{BR}$ to on-drug BR columns (see 19-mean-moderation-implementation-notes.tex for the scaling derivation). The divergent carryover behavior reported below is therefore attributable to the BM-BR parameterization, not to any difference in how the placebo, time-varying, or residual-noise components are generated. The analysis model was the same in both cases: an `nlme::lme` fit with a `corCAR1` continuous-time autocorrelation structure and a single biomarker-by-drug interaction term, written `bm:Dbc` in R formula syntax.

[12]

Dbc is a continuous drug exposure variable that captures exponential off-drug decay; the `bm:Dbc` interaction tests whether biomarker moderates that exposure-decayed effect.

- `Dbc` = 1 on drug; = $\exp(-\lambda t_{sd})$ off drug, matching the DGP decay rate.
- Unlike the binary D_{it} , `Dbc` smoothly captures off-drug shrinkage.
- The test of `bm:Dbc` is the test of whether the biomarker moderates the (decay-adjusted) drug effect.

rgt to complete.

The analysis-model drug exposure variable `Dbc` is the *continuous* analogue of the binary DGP indicator D_{it} . It equals 1 during on-drug periods and $\exp(-\lambda t_{sd})$ during off-drug periods, where t_{sd} is time since drug discontinuation and λ is the analysis-model carryover decay rate (typically derived from the drug’s pharmacokinetic half-life as $\lambda = \ln 2 / t_{1/2}$, in which case it matches the DGP decay rate of Section 2.2 by construction). Whereas D_{it} flips abruptly between 0 and 1, `Dbc` captures the smooth shrinkage of residual drug effect during the off-drug period. The interaction coefficient on `bm:Dbc` is therefore the test of whether the biomarker moderates the exposure-decayed drug effect.

We follow the ADEMP framework of Morris, White, and Crowther (2019): Aims, Data-generating mechanisms, Estimands, Methods, Performance measures.

Setup:

- Trial designs: CO (traditional crossover), Hybrid (N-of-1 Hybrid),

- OL+BDC (open-label with blinded discontinuation)
- Sample size: $N = 35$ per randomization path (effective analysis $N = 70$ for CO and OL+BDC, 140 for Hybrid, after the paths are concatenated)
 - Biomarker correlation / moderation: $c_{bm} = 0.45$ / $\beta_{bm} = 0.45$
 - DGP carryover half-life: $t_{1/2} \in \{0, 0.5, 1.0\}$ weeks
 - Analysis: matched Dbc with same half-life
 - Replicates: 1,000 per cell (MCSE on a binomial power estimate near 0.5 is approximately $\sqrt{0.5 \cdot 0.5 / 1,000} \approx 0.016$)
 - Implementation: 8-core `parallel::mclapply`; full factorial 2 architectures $\times$ 3 designs $\times$ 3 $t_{1/2} \times$ 2 c_{bm} levels = 36 cells, 36,000 fits, 100% convergence

[13]

$N = 35$ was chosen because it keeps power off the ceiling, making any carryover-induced loss visible; Type I error is in the expected range.

- $N = 35$ is the smallest prazosin-PTSD-calibrated sample size in the literature.
- Neither architecture saturates power at zero carryover, so losses are mechanically observable.
- Type I error across 18 null cells sat in $[0.010, 0.074]$, consistent with the conservative-to-nominal range expected for this specification.

rgt to complete.

The $N = 35$ choice is the smallest of the prazosin-PTSD- calibrated sample sizes used in the published methodological literature; it is a regime in which neither architecture saturates power at the no-carryover cells, so any carryover-induced erosion is mechanically visible. Type I error under the $c_{bm} = 0$ null cells sat in $[0.010, 0.074]$ across the 18 null cells. These rates below the nominal 0.05 are not architecture-specific; they reflect a known conservatism of the standard `corCAR1` Wald test, whose model-based standard error is overstated by roughly 6 to 10 percent under this working-covariance specification. A companion calibration study¹ localizes this miscalibration to the `corCAR1` residual layer and shows that a CR2 bias-reduced cluster-robust standard error restores the realized Type I error to approximately the nominal level, with a near-uniform power increment that leaves the ranking of analysis specifications unchanged. Because the conservatism is common to both architectures, it is common-mode with respect to the B-versus-A contrast that is the object of this section and does not affect it.

Results under Architecture A (direct mean moderation):

Design	$t_{1/2} = 0$	$t_{1/2} = 0.5$	$t_{1/2} = 1.0$
CO	0.739	0.739	0.725
Hybrid	0.992	0.987	0.978
OL+BDC	0.751	0.748	0.730

[14]

Architecture A shows modest power loss under carryover (1–3%), below preliminary expectations from earlier exploratory runs in this project.

- Relative losses: 1.9% (CO), 1.4% (Hybrid), 2.8% (OL+BDC).
- All three fall below the 8–13% range suggested by earlier exploratory runs in this project.
- Qualitative finding: modest erosion; power essentially intact under carryover.

rgt to complete.

Power declines modestly under Architecture A: relative loss from $t_{1/2} = 0$ to $t_{1/2} = 1$ is 1.9% (CO), 1.4% (Hybrid), and 2.8% (OL+BDC). We note that all three values fall well below the 8–13% relative-loss range suggested by earlier exploratory runs in this project for direct-mean-moderation DGPs, though they reproduce the qualitative finding of those runs, namely a modest erosion that leaves power essentially intact.

Results under Architecture B (differential correlation, MVN):

Design	$t_{1/2} = 0$	$t_{1/2} = 0.5$	$t_{1/2} = 1.0$
CO	0.745	0.754	0.723
Hybrid	0.995	0.980	0.944
OL+BDC	0.777	0.694	0.539

[15]

Architecture B shows sharply larger power loss under carryover, especially OL+BDC (30.6%); the B-vs-A gap is significant for OL+BDC and Hybrid but not CO.

- Relative losses: 3.0% (CO), 5.1% (Hybrid), 30.6% (OL+BDC); the OL+BDC loss is $\approx 11\times$ the matched Architecture A loss.
- B-vs-A gap: $z = 7.64$ ($p < 10^{-13}$) at OL+BDC; $z = 3.96$ ($p < 10^{-4}$) at Hybrid.
- CO gap not significant ($z = 0.29$): the CO design’s within-subject AR(1) structure dominates the correlation-decay channel that Architecture B depends on.

rgt to complete.

Power declines more sharply under Architecture B, and most markedly so in the OL+BDC design: relative loss from $t_{1/2} = 0$ to $t_{1/2} = 1$ is 3.0% (CO; not significantly different from Architecture A’s CO loss), 5.1% (Hybrid), and **30.6% (OL+BDC)**. The CO cell shows a small, non-monotone fluctuation ($0.745 \rightarrow 0.754 \rightarrow 0.723$), which is within one MCSE and is consistent with the null finding of no significant architecture-by-carryover interaction in that design. The OL+BDC loss is approximately 11 times the matched Architecture-A OL+BDC loss (2.8%). It falls below the 40-60% relative-loss range suggested by earlier exploratory runs in this project for the MVN differential-correlation DGP, but it confirms the qualitative finding of those runs, namely a substantial and strongly design-dependent erosion. We compare the carryover-induced power loss between architectures using a difference-of-differences z -statistic. Let $\Delta_X = \hat{\pi}_{X,0} - \hat{\pi}_{X,1}$ denote the within-architecture loss from $t_{1/2} = 0$ to $t_{1/2} = 1.0$ for architecture $X \in \{A, B\}$, estimated from $n = 1,000$ independent replicates at each cell. The B-versus-A gap is $\Delta_B - \Delta_A$, with standard error

$$SE = \sqrt{\frac{\hat{\pi}_{B,0}(1 - \hat{\pi}_{B,0})}{n} + \frac{\hat{\pi}_{B,1}(1 - \hat{\pi}_{B,1})}{n} + \frac{\hat{\pi}_{A,0}(1 - \hat{\pi}_{A,0})}{n} + \frac{\hat{\pi}_{A,1}(1 - \hat{\pi}_{A,1})}{n}}$$

since the four proportions come from independent replicate sets. All tests are two-sided. At OL+BDC this gives $z = 7.64$ ($p < 10^{-13}$); at Hybrid, $z = 3.96$ ($p < 10^{-4}$); the CO contrast is not statistically distinguishable from zero ($z = 0.29$) because the CO design’s within-subject AR(1) structure dominates the carryover-mediated correlation decay channel that Architecture B uses to encode the biomarker-treatment interaction.

4.2 3.2 Why the architectures diverge

We turn now to a mechanistic account of why the two architectures produce such different power losses under carryover.

[16]

Architecture A is subject to only one carryover channel (mean blurring); it escapes the second channel (correlation erosion).

- Carryover compresses Dbc contrast and inflates off-drug BR means; both architectures share this channel.
- Architecture A lacks a second channel because β_{bm} is a fixed population parameter and the noise structure is drug-state-independent.

- Result: reduced variance in the interaction term, but signal-to-noise ratio degrades only modestly.

rgt to complete.

Under Architecture A, carryover reduces the magnitude of the drug exposure variable seen by the analysis model (D_{bc} drops below 1 during the off-drug period), and the mean BR at off-drug timepoints is inflated by the residual drug effect exactly as in Architecture B. Both mechanisms compress the between-state treatment contrast and are shared by the two architectures. What Architecture A lacks is the second channel of signal loss: the biomarker moderation coefficient β_{bm} is a fixed population parameter, the biomarker enters the mean structure directly rather than through a correlation, and the noise structure is independent of drug state. The biomarker-by-drug interaction term therefore has reduced variance (because D_{bc} has less contrast between drug states), but the signal-to-noise ratio degrades only modestly.

Under Architecture B, carryover operates on two levels simultaneously:

1. **Mean blurring:** The BR means during off-drug periods are inflated by residual carryover, making them resemble on-drug BR means. This reduces the treatment contrast that the analysis model uses to distinguish drug effect from no-drug. This channel is shared with Architecture A.
2. **Correlation erosion:** The BM-BR correlation during off-drug periods decays exponentially with t_{sd} . This is not just a statistical modeling choice; it is a structural consequence of the DGP. The correlation between biomarker and BR reflects the degree to which the drug effect is active; as the drug washes out, the correlation weakens because the drug-mediated component of BR variance shrinks relative to the drug-independent component. This channel is *unique* to Architecture B, because only here does the biomarker signal live in the second moment of the joint distribution. Under Architecture A the drug-mediated and drug-independent components of BR variance move in tandem (both are scaled by the same D_{bc}-weighted drug effect), so their ratio, and hence the identifiability of the interaction, is preserved even as amplitudes shrink.

[17]

Under Architecture B the interaction estimate is attacked from both sides simultaneously; Architecture A faces only one of the two attacks.

- The bm:D_{bc} coefficient depends on the covariance between B_i and D_{bc}-weighted outcomes.
- Under Architecture B: (1) D_{bc} contrast is compressed; (2) the BM-BR correlation generating the covariance signal is simultaneously decaying.

629 • Architecture A faces only attack (1), not attack (2); this single
630 difference accounts for the divergence in power.

631 *rgt to complete.*

632 The analysis model’s interaction coefficient (the coefficient on `bm:Dbc`)
633 depends on the covariance between B_i and Dbc-weighted outcomes.
634 Under Architecture B, this covariance is being attacked from both
635 sides: the treatment contrast in Dbc is compressed, and the BM-BR
636 correlation that generates the covariance signal is simultaneously de-
637 caying. That Architecture A escapes the second attack entirely ac-
638 counts for the divergence in power loss documented above.

639 4.3 3.3 Robustness check at $N = 70$

640 [18]

641 **An $N = 70$ confirmation run checks whether the qualitative**
642 **ordering holds when power is near the ceiling.**

- 643 • 800 replicates per cell, 12 cells ($2 \text{ architectures} \times 2 c_{bm} \times 3 t_{1/2}$),
644 OL+BDC design.
- 645 • 100% convergence across 9,600 fits.
- 646 • Purpose: verify the Section 3.1 ordering survives saturation at
647 the larger N .

648 *rgt to complete.*

649 The Section~3.1 production results are at $N = 35$, where neither
650 architecture saturates power and the carryover-induced loss is me-
651 chanically visible. We conducted a confirmation run at $N = 70$ to
652 check whether the qualitative ordering survives the saturation that
653 comes with the larger sample size. An 800-replicate-per-cell Monte
654 Carlo run was executed at $N = 70$ across 12 cells (architecture
655 $\in \{\text{MVN, mean moderation}\} \times c_{bm} \in \{0, 0.45\} \times t_{1/2} \in \{0, 0.5, 1.0\}$
656 weeks), all on the OL+BDC design, using 8-core `parallel::mclapply`.
657 Wall-clock 483~s; convergence 100% across 9{,}600 fits.

658 At $c_{bm} = 0.45$:

Architecture	$t_{1/2} = 0$	$t_{1/2} = 0.5$	$t_{1/2} = 1.0$
MVN (B)	0.978	0.945	0.885
Mean moderation (A)	0.976	0.973	0.959

659 [19]

660 **The $N = 70$ run confirms the qualitative ordering; compressed**
661 **absolute magnitudes reflect ceiling saturation, not a smaller**
662 **underlying effect.**

- Losses: 9.3 pp (Architecture B, $p < 10^{-12}$) vs. 1.7 pp (Architecture A, $p \approx 0.05$); ordering matches $N = 35$.
- At $N = 70$ both architectures are within 2–5 pp of the ceiling at $t_{1/2} = 0$, mechanically truncating observable loss.
- $N = 35$ production results remain the canonical reference for the narrative magnitude claims.

rgt to complete.

Within-architecture losses from $t_{1/2} = 0$ to $t_{1/2} = 1$ are 9.3 percentage points (Architecture~B; two-proportion z -test on the paired cells, $z = 7.4$, $p < 10^{-12}$) and 1.7 percentage points (Architecture~A; $z = 2.0$, $p \approx 0.05$), where each z -statistic uses the formula $(\hat{\pi}_0 - \hat{\pi}_1) / \sqrt{\hat{\pi}_0(1 - \hat{\pi}_0)/n + \hat{\pi}_1(1 - \hat{\pi}_1)/n}$ with $n = 800$ replicates per cell. The qualitative ordering matches Section~3.1's $N = 35$ result. The absolute magnitudes are compressed because $N = 70$ at the prazosin-PTSD calibration places both architectures within 2–5 percentage points of the unit ceiling at $t_{1/2} = 0$, mechanically truncating the loss either architecture's curve can absorb. We note that reading the narrative loss-magnitude claim against the $N = 70$ robustness numbers (3% and 9% relative loss for Architecture~B versus 0.4% and 1.8% for Architecture~A) requires acknowledging this saturation effect; the $N = 35$ production results in Section~3.1 are the canonical reference for the narrative magnitude.

[20]

Type I error is well-controlled across all 18 null cells, with no architecture-specific inflation.

- Type I error range $[0.010, 0.074]$, mean 0.043, median 0.042 across 18 null cells at $N = 35$.
- No architecture-specific inflation detected.
- The slightly conservative Architecture B OL+BDC null cell at $t_{1/2} = 1.0$ ($p = 0.010$) is the only notable departure.

rgt to complete.

Type I error. Across all 18 null cells of the $N = 35$ production run, Type~I error sat in $[0.010, 0.074]$, with mean 0.043 and median 0.042. We observe no architecture-specific inflation. The uniformly sub-nominal rates are consistent with the corCAR1 working-covariance miscalibration characterized in the companion calibration study¹; because it affects both architectures alike, it is common-mode with respect to the architecture contrast. Architecture~B's $t_{1/2} = 1.0$ OL+BDC null ($p = 0.010$, slightly conservative) is the smallest cell; the corresponding alternative ($p = 0.539$) is the most informative cell of the Section~3.1 narrative.

[21]

Estimator bias is small and architecturally interpretable; coefficient drift under Architecture A contrasts with correlation-channel erosion under Architecture B.

- Mean $\hat{\beta}_{bm:D_{bc}}$ at $c_{bm} = 0.45$ ranges -0.231 to -0.281 across alternative cells; Architecture A OL+BDC at $t_{1/2} = 1.0$ produces the largest absolute mean.
- Architecture B cells maintain mean near -0.234 across all $t_{1/2}$ values, consistent with erosion through correlation decay rather than coefficient drift.
- Nominal-95% Wald coverage holds across the grid.

rgt to complete.

Estimator bias. Mean estimated $\hat{\beta}_{bm:D_{bc}}$ at $c_{bm} = 0.45$ ranged from -0.231 to -0.281 across the 18 alternative cells, with the $N = 35$ Architecture-A OL+BDC cell at $t_{1/2} = 1.0$ producing the largest absolute mean (-0.281); the Architecture-B cells maintained mean near -0.234 across all $t_{1/2}$ values, consistent with the Architecture-B information channel being eroded through correlation decay rather than through coefficient drift. Empirical SE of $\hat{\beta}_{bm:D_{bc}}$ runs 0.05 to 0.10 across cells; the within-replicate model SE matches at scale, indicating nominal-95% Wald coverage holds across the grid.

5 4. Biological Assumptions

[22]

Architecture choice encodes a biological claim about how the biomarker moderates treatment response.

- The three architectures reflect distinct assumptions about the mechanism of biomarker-treatment moderation.
- Section 4 takes each architecture in turn (4.1, 4.2, 4.3 for Architectures A, B, and C respectively), then discusses the prazosin-PTSD motivating case.

rgt to complete.

The choice between architectures is not merely a statistical convenience; it reflects different assumptions about the biological mechanism by which the biomarker moderates treatment response. We consider each architecture in turn, then take up the prazosin-PTSD motivating application in detail.

5.1 4.1 When Architecture A is appropriate

[23]

Architecture A is appropriate when the biomarker directly determines the magnitude of the drug effect; the scaling is deterministic at the mean-structure level.

- The drug effect scales with the biomarker value: $\text{effect}_i = f(B_i) \cdot \text{drug}$.
- 'Deterministic' means any two participants with the same B_i have the same expected response; individual outcomes still carry noise.
- Appropriate when the biomarker is a direct pharmacokinetic/pharmacodynamic (PK/PD) determinant (eGFR, receptor density, metabolizer status).

rgt to complete.

Architecture A (direct mean moderation) is appropriate when the biomarker governs the **magnitude of the drug's biological effect** for each individual participant. The drug effect for participant i is scaled by their biomarker value: $\text{effect}_i = f(B_i) \cdot \text{drug}$. The scaling is deterministic in the population-mean sense, in that any two participants sharing the same B_i have the same expected response, but individual outcomes still carry the usual noise, random-effect, and residual-biological variability. The label 'deterministic' refers to the mean structure of the DGP, not to the realized response of any single participant.

Clinical examples:

- **Renal clearance of a renally-excreted drug.** Estimated glomerular filtration rate (eGFR) directly governs the elimination rate of drugs such as digoxin or aminoglycoside antibiotics. Participants with higher eGFR clear the drug more rapidly, achieving lower steady-state concentrations and a proportionally smaller drug effect. The biomarker-drug relationship is mechanistic and deterministic: given the biomarker value, the effective dose at steady state is fully determined (up to random noise).
- **Receptor density moderation.** A biomarker that measures target receptor density (e.g., PET imaging of receptor availability for an antagonist) directly determines the pharmacodynamic response to that antagonist. More receptors means more drug effect, proportionally.
- **Genetic metabolizer status.** CYP2D6 metabolizer phenotype determines the rate of drug activation or clearance. Poor metabolizers of a prodrug receive less active compound, reducing their treatment effect proportionally.
- **Dose-response with biomarker-dependent effective dose.** When the biomarker determines the effective dose (through body composition, protein

binding, or volume of distribution), the treatment effect scales with the biomarker through the dose-response curve.

[24]

In all Architecture A cases, the proportional biomarker-drug relationship is preserved during the off-drug period.

- Residual drug effect during the off-drug period maintains the same proportional relationship with the biomarker.
- If participant A has $2\times$ the drug effect of B on-drug, they also have $\approx 2\times$ the residual effect during the off-drug period.
- This preserved proportionality is what Architecture A encodes and why it loses little power under carryover.

rgt to complete.

In all these cases the biomarker’s moderating effect is preserved under carryover because the residual drug effect during off-drug periods maintains the same proportional relationship with the biomarker. If participant A has twice the drug effect of participant B when on drug, they will also have approximately twice the residual effect during the off-drug period. That proportionality is precisely what Architecture A encodes.

5.2 4.2 When Architecture B is appropriate

[25]

Architecture B is appropriate when the biomarker is a probabilistic predictor of response, mediated by latent factors the biomarker imperfectly indexes.

- Appropriate when the biomarker predicts which participants are likely to respond, not the magnitude of any individual’s response.
- The biomarker-response association is population-level and probabilistic, mediated by latent factors (disease subtype, neurobiological heterogeneity, comorbidities).
- Two participants with identical biomarker values may respond very differently.

rgt to complete.

Architecture B (differential correlation) is appropriate when the biomarker predicts **which participants are likely to respond**, but the magnitude of any individual’s response is not deterministically governed by the biomarker. The biomarker and the drug response are associated at the population level, but the association is mediated by latent factors (disease subtype, neurobiological heterogeneity, comorbidities) that the biomarker imperfectly indexes.

Clinical examples:

- **Blood pressure as a PTSD subtype marker.** Elevated resting blood pressure may index a noradrenergic-predominant PTSD phenotype¹³ that is more likely to respond to prazosin (an alpha-1 adrenergic antagonist). The blood pressure does not *cause* the drug to work better; rather, it correlates with an underlying neurobiological state that determines drug responsiveness. Two participants with identical blood pressure may have very different responses because blood pressure is an imperfect proxy for the noradrenergic state.
- **Baseline disease severity as a treatment predictor.** Trials in Alzheimer’s disease, depression, and chronic pain often find that patients with more severe baseline symptoms show larger treatment effects (the ‘floor effect’ or ‘regression to the mean’ phenomenon). The baseline severity score correlates with treatment response, but the relationship is statistical (population-level tendency), not deterministic.
- **Inflammatory biomarkers in psychiatry.** C-reactive protein (CRP) or interleukin-6 levels may predict response to anti-inflammatory augmentation of antidepressants¹⁸. High-CRP patients are more likely to respond, but the relationship is probabilistic: many high-CRP patients do not respond, and some low-CRP patients do.
- **Genetic risk scores.** Polygenic risk scores predict disease trajectory and treatment response, but explain only a fraction of the variance. The remainder reflects unmeasured genetic, epigenetic, and environmental factors.

[26]

Under Architecture B, carryover weakens the biomarker-response correlation because the biomarker’s predictive power depends on the presence of active drug.

- As the drug washes out, the drug-mediated variance component diminishes, weakening the BM-BR correlation.
- The biomarker’s predictive power is tied to active drug effect; without drug, it predicts nothing about non-drug symptom change.
- This is the ‘correlation erosion’ channel unique to Architecture B.

rgt to complete.

In these cases carryover has a different implication: as the drug washes out, the correlation between biomarker and response weakens because the drug-mediated component of response variance diminishes. The biomarker’s predictive power is inherently tied to the presence of active drug effect. Without drug, the biomarker is just a biomarker; it does not predict the non-drug component of symptom change.

[27]

Architecture B is a tractable second-moment approximation to the biologically faithful finite-mixture DGP; adequate for power calculations but not a mechanistic endorsement.

- If the biomarker indexes a discrete responder class, the faithful DGP is a finite mixture; Architecture B approximates its population-level covariance structure.
- Architecture B captures the covariance implications of latent-class membership without committing to the full generative form.
- Adequate for power calculations; should not be read as endorsing MVN differential correlation as the correct mechanism.

rgt to complete.

The latent-subtype framing raises a legitimate modeling question: if the biomarker is really a noisy indicator of an underlying responder class, the biologically faithful DGP would be an explicit finite mixture in which each participant is drawn from one of two (or more) response distributions with class-membership probability a function of B_i . Under such a mixture DGP the individual-level moderation is discrete, not probabilistic, and the MVN differential-correlation parameterization is best understood as a tractable second-moment approximation to a richer mixture model rather than as the generative mechanism itself. The practical consequence is that Architecture~B as implemented here captures the population-level covariance implications of a latent-class mechanism without committing to its full generative form; this is adequate for power calculations that test for the presence of an interaction but should not be read as an endorsement of MVN differential correlation as the correct mechanistic model. Whether the observed power loss under carryover is better mitigated by analysis strategies (Section 6) or by explicit mixture modeling is taken up in a companion paper¹⁴, which derives the conditions under which a single multivariate normal reproduces the first two moments of a two-component mixture, catalogues the regimes in which that approximation fails, and finds in a preliminary evaluation that at trial-relevant sample sizes class-aware analyses are not competitive with the continuous linear-mixed baseline.

[28]

Senn (2016) cautions that apparent personalized response is often a variance artifact; Architecture B presupposes a real, stable latent subtype, which may not hold.

- Senn (2016) argues that apparent biomarker-treatment interactions are often variance-component artifacts, not genuine effects.
- Architecture B presupposes the latent subtype is real, identifiable,

904 and stable across the trial period.
905 • If that presupposition is shaky, the documented power loss may
906 overstate recoverable signal.

907 *rgt to complete.*

908 The latent-subtype framing of Architecture B should also be adopted
909 with appropriate epistemic caution. Senn¹⁹ has argued that the ap-
910 parent ‘personal element in response to treatment’ is often a variance-
911 component artifact rather than a genuine biomarker-treatment inter-
912 action, and that even modest within-subject variability can produce
913 convincing-looking but spurious differential response. Architecture B
914 presupposes that the latent subtype indexed by the biomarker is real,
915 identifiable, and stable across the trial period. When that presupposi-
916 tion is shaky, the apparent power loss under carryover documented in
917 Section 3 may overstate what could be recovered under any analysis
918 model, because the underlying signal it is attempting to estimate may
919 itself be smaller than assumed.

920 5.3 4.3 The prazosin-PTSD case

921 We turn briefly to the motivating clinical application for pmsimstats, which is the
922 active-duty prazosin trial by Raskind et al.¹⁵, whose SBP-stratified secondary anal-
923 ysis is the subject of ongoing reanalysis by the pmsimstats team. The biomarker
924 is resting systolic blood pressure (SBP) and the outcome is CAPS (Clinician-
925 Administered PTSD Scale) score change.

926 [29]

927 **The prazosin-PTSD evidence base is mixed; the PACT null**
928 **result motivates the search for a predictive biomarker, with**
929 **SBP as the leading candidate.**

- 930 • Earlier single-site Raskind trials (2003, 2007, 2013) reported ben-
931 efit on nightmares and PTSD symptoms.
- 932 • PACT (2018) failed to replicate; this has been interpreted as
933 consistent with effect heterogeneity.
- 934 • SBP, as a noninvasive proxy for central noradrenergic tone, is the
935 leading predictive biomarker candidate.

936 *rgt to complete.*

937 The clinical evidence base for prazosin in PTSD is mixed. The ear-
938 lier single-site trials by Raskind and colleagues^{15,20,21} reported benefit
939 on nightmares and global PTSD symptoms in veteran and active-duty
940 samples. The subsequent multisite Veterans Affairs Cooperative Study
941 Program trial¹⁶, known as PACT, did not replicate the effect: prazosin
942 failed to outperform placebo on the primary endpoints. The negative
943 PACT result has been interpreted by several investigators as consis-

tent with effect heterogeneity, namely that prazosin works in some patients but not others, which motivates the search for a predictive biomarker that would identify the responder subgroup. Resting SBP, as a noninvasive proxy for central noradrenergic tone, is the leading candidate.

[30]

The SBP-prazosin mechanism is an Architecture B scenario: SBP is a proxy for noradrenergic state, not a direct PK determinant.

- Elevated SBP indexes heightened noradrenergic tone, a PTSD subtype preferentially responsive to prazosin.
- SBP is a proxy, not a causal determinant; two patients with identical SBP may have very different noradrenergic states.
- This is precisely the Architecture B scenario: probabilistic association, not deterministic scaling.

rgt to complete.

The biological hypothesis is that elevated SBP indexes heightened central noradrenergic tone, which characterizes a subtype of PTSD that is preferentially responsive to alpha-1 adrenergic blockade by prazosin. This is an Architecture B scenario: SBP is a proxy for an underlying neurobiological state, not a direct determinant of drug pharmacokinetics. Two patients with the same SBP may differ in their actual noradrenergic tone, alpha-1 receptor distribution, and downstream signaling, leading to different treatment responses despite identical biomarker values.

[31]

Hendrickson et al. (2020) correctly uses Architecture B for this context; the carryover power loss is clinically relevant for trial design decisions.

- Architecture B (MVN differential correlation) is appropriate for the SBP-prazosin mechanism.
- The documented carryover sensitivity means designs with longer off-drug periods will have meaningfully less power.
- This effect is larger than a direct mean moderation simulation would predict.

rgt to complete.

The Hendrickson et al.¹³ DGP uses Architecture B (MVN differential correlation) for this clinical context. The finding that carryover substantially reduces power under this architecture is therefore clinically relevant: it suggests that trial designs with longer off-drug periods will

984 have meaningfully less power to detect the SBP-prazosin interaction
 985 than designs with shorter off-drug periods, and this effect is larger
 986 than what a direct mean moderation simulation would predict.

987

988 6 5. Implications for Simulation Study Design

989 6.1 5.1 Choosing the appropriate architecture

990 The choice of DGP architecture should be guided by the biological mechanism
 991 hypothesized for the biomarker-treatment interaction:

Criterion	Architecture A	Architecture B
Biomarker role	Causal mediator	Statistical predictor
Mechanism	Deterministic scaling	Probabilistic association
Signal location	Mean structure	Covariance structure
Free parameters	1 (c_{bm})	1 (c_{bm})
Carryover impact	Modest (1–3%)	Design-dependent, up to 31%
Appropriate when	Biomarker determines dose or PK	Biomarker indexes a latent subtype

992 When the operative mechanism is unknown, Architecture~B is the conservative
 993 default for trial-design decisions: its larger carryover sensitivity keeps power esti-
 994 mates from being optimistic relative to Architecture~A. A combined architecture
 995 that mixes the two channels, together with a corresponding sensitivity sweep, is
 996 developed in the companion paper¹⁷.

997 6.2 5.2 Reporting requirements

998 Simulation studies evaluating biomarker-moderated treatment effects should ex-
 999 plicitly state:

- 1000 1. Whether the biomarker-treatment interaction is generated through mean
 1001 moderation, differential correlation, or both channels simultaneously.
- 1002 2. The biological rationale for the chosen mechanism.
- 1003 3. Whether the carryover model (if present) operates on the interaction signal
 1004 consistently with the chosen architecture.
- 1005 4. Sensitivity analyses under alternative architectures, if the biological mecha-
 1006 nism is uncertain.

1007 [32]

1008 **The four reporting requirements are the ADEMP 'D' element**
 1009 **applied specifically to biomarker-validation simulations.**

- ADEMP (2019) prescribes five elements: Aims, Data-generating mechanisms, Estimands, Methods, Performance measures.
- The DGP-architecture choice concerns the D element; the four requirements above specialize that element for architecture-sensitive simulations.

rgt to complete.

The reporting framework above is consistent with the broader ADEMP template of Morris, White, and Crowther²², which prescribes that simulation studies fix five elements before production runs: the **Aims** of the study, the **Data-generating mechanisms**, the **Estimands** (the inferential targets), the **Methods** (analysis specifications compared), and the **Performance measures** (power, bias, coverage, Type I error, and MCSEs of each). The DGP-architecture choice this paper concerns is the **D** element of an ADEMP specification; the reporting requirements above are the specialization of the broader ADEMP framework to the architecture-sensitive class of biomarker-validation simulations.

[33]

Primary estimand: $\beta_{bm:D}$ in the linear mixed-effects (LME) model; performance measures cover power, Type I error, bias, SE, and convergence; 1000 replicates per cell.

- Estimand: $\beta_{bm:D}$ in $Sx \sim \text{bm} + \text{t} + \text{Dbc} + \text{bm:Dbc}$ with random intercept and AR(1).
- Five performance measures: power, Type I error, bias, empirical SE, convergence rate.
- $n_{\text{reps}} = 1000$ throughout; calibrated to achieve $\text{MCSE} < 0.02$ for power near 0.5.

rgt to complete.

The **primary estimand** for this paper’s simulations is the biomarker-treatment interaction coefficient $\beta_{bm:D}$ in the linear-mixed analysis model $Sx \sim \text{bm} + \text{t} + \text{Dbc} + \text{bm:Dbc}$ with random intercept and AR(1) residual correlation. The reported **performance measures** are: (i) power for $\beta_{bm:D}$ at $\alpha = 0.05$, with $\text{MCSE} \sqrt{\hat{\pi}(1 - \hat{\pi})/n_{\text{reps}}}$; (ii) Type I error under $c_{bm} = 0$, with the same MCSE; (iii) bias of $\hat{\beta}_{bm:D}$ relative to the calibrated true value at each parameter cell; (iv) the empirical MCSE of $\hat{\beta}_{bm:D}$ across replicates; and (v) the convergence rate of the linear-mixed fit. Sample size, biomarker effect size, and carryover half-life vary across the parameter grid; the number of replicates per cell is calibrated to achieve MCSE below 0.02 for power at $\hat{\pi} = 0.5$, which corresponds to $n_{\text{reps}} \geq 625$. The simulations reported here use $n_{\text{reps}} = 1000$ throughout.

6.3 5.3 When the choice is uncertain

If the biological mechanism is ambiguous, as it often is in early biomarker development, we suggest investigators take the following steps:

1. Run the simulation under both architectures and report the range of power estimates.
2. Architecture B produces more conservative (lower) power estimates under carryover, making it the safer default for trial design decisions.
3. Design the trial to minimize carryover (adequate washout periods) to reduce sensitivity to the DGP architecture choice.

6. Alternative Analysis Strategies Under Architecture B

This section provides a conceptual overview of six candidate strategies for mitigating the power loss documented in Section 3. None of the strategies below has been evaluated by simulation within this study; the assessments of expected benefit are qualitative and based on the mechanistic account in Section 3.2. A systematic simulation comparison is warranted before any strategy is recommended for applied use.

The power loss documented in Section 3 arises from two distinct sources: (1) compressed variance of the drug exposure variable Dbc when carryover is present, and (2) erosion of the BM-BR correlation during off-drug periods. Source 1 is a property of the analysis model parameterization and can potentially be addressed by choosing a different predictor. Source 2 is a property of the data itself (the signal is genuinely weaker in off-drug observations), and no analysis model can recover signal that is not present. However, analysis strategies that are selective about *which observations contribute to the test* or *how they are weighted* can mitigate the damage.

6.1 Restrict analysis to on-drug observations

The most direct approach is to discard off-drug timepoints entirely and test the biomarker-treatment interaction using only on-drug observations. Every retained observation has the full BM-BR correlation (c_{bm}), so correlation erosion is eliminated. The costs are twofold: a reduction in effective sample size (fewer observations per participant) and, more importantly, the loss of the off-drug reference that separates the biological response from placebo and time-varying background response. Without off-drug observations the biological, placebo, and time-varying components are confounded, which is why the Open-Label design, despite having the most on-drug timepoints, does not dominate the other designs in the Section 3

results. For designs with many on-drug timepoints (OL has 8, the Hybrid design has 5-6 per path), the net effect may be positive: the per-observation signal-to-noise ratio is maximal, and the analysis model simplifies to $\mathbf{Sx} \sim \mathbf{bm} + \mathbf{t} + \mathbf{bm:t}$ since there is no treatment contrast to model. The interaction is detected through the correlation between biomarker and the time-on-drug trajectory.

[34]

Restricting to on-drug observations eliminates correlation erosion but loses the off-drug reference; best for OL and Hybrid designs, worst for CO.

- Benefit: every retained observation has full c_{bm} ; correlation erosion is eliminated.
- Cost: loss of off-drug reference confounds biological, placebo, and time-varying components.
- Practical: best for OL/Hybrid (many on-drug timepoints); least useful for CO (discards half the observations).

rgt to complete.

This approach is most promising for the Open-Label design (which has no off-drug observations regardless) and the Hybrid design (which has a long initial on-drug phase). It is least useful for the Crossover design, where discarding half the observations substantially reduces statistical power.

7.2 6.2 Weighted analysis

[35]

Weighted analysis downweights off-drug observations by drug exposure without discarding them; information-theoretic rationale; straightforward to implement.

- Weight: $w_t = e^{-\lambda t_{sd}}$ for off-drug observations; 1 for on-drug.
- Rationale: higher exposure = stronger correlation signal = more information about the interaction.
- Implemented via the `weights` argument in `nlme::lme`.

rgt to complete.

Rather than discarding off-drug observations entirely, a weighted analysis retains all observations but weights them by their estimated drug exposure level. On-drug observations receive weight 1. Off-drug observations receive weight proportional to the expected residual drug effect:

$$w_t = e^{-\lambda \cdot t_{sd}}$$

This downweights observations where the BM-BR correlation has decayed (and where they contribute more noise than signal to the interaction estimate) without discarding them entirely. The rationale is information-theoretic: observations with higher drug exposure carry more information about the biomarker-treatment interaction because the correlation signal is stronger. This is straightforward to implement in `nlme::lme` using the `weights` argument.

6.3 Within-subject contrast

For designs with both on-drug and off-drug periods (CO, Hybrid, OL+BDC), a summary-measure approach computes a per-participant contrast: the difference between mean on-drug response and mean off-drug response. The biomarker-treatment interaction is then tested by a simple regression of these contrasts on the biomarker:

$$\bar{Y}_{i,\text{on}} - \bar{Y}_{i,\text{off}} = \alpha + \gamma \cdot B_i + \epsilon_i$$

[36]

Within-subject contrast sidesteps the repeated-measures model; trades efficiency for robustness to carryover misspecification.

- Per-participant contrast $\bar{Y}_{\text{on}} - \bar{Y}_{\text{off}}$ is regressed on the biomarker.
- Carryover compresses the contrast magnitude but the contrast-biomarker correlation may be more robust.
- Trade-off: less efficient than a full longitudinal model, but more robust to carryover and correlation misspecification.

rgt to complete.

This sidesteps the `Dbc` parameterization and the repeated-measures model entirely. Carryover affects the magnitude of the contrast (making off-drug means closer to on-drug means), but the *correlation between the contrast and the biomarker* may be more robust because within-subject averaging reduces noise before the biomarker association is tested. The approach trades the efficiency of a full longitudinal model for robustness to misspecification of the carryover and correlation structure.

6.4 Two-stage random slopes

A two-stage approach first estimates participant-specific treatment effects, then tests their association with the biomarker. Stage 1 fits a random-slopes model:

$$Y_{it} = \beta_0 + \beta_1 t + \beta_2 D_{it} + u_{0i} + u_{1i} D_{it} + \epsilon_{it}$$

where β_2 is the population mean drug effect and u_{1i} is participant i 's random deviation from that mean. The participant-specific drug effect is then $\beta_2 + u_{1i}$. Stage 2 regresses the estimated random deviations (or equivalently, the best linear unbiased predictors (BLUPs) of the participant-specific effects) on the biomarker:

$$\hat{u}_{1i} = \delta_0 + \delta_1 B_i + \eta_i$$

[37]

Two-stage random slopes separates within-subject effect estimation from between-subject biomarker association; may be more robust to carryover.

- Stage 1: random-slopes LME estimates participant-specific drug effects ($\beta_2 + u_{1i}$).
- Stage 2: regress BLUPs on biomarker to test association.
- The random slope u_{1i} integrates over the full trajectory; robustness to carryover is plausible but unverified by simulation.

rgt to complete.

This separates the within-subject treatment effect estimation (which uses the longitudinal data and is affected by carryover) from the between-subject biomarker association (which uses the cross-sectional relationship between estimated effects and biomarker values). The approach may be more robust to carryover because the random slope u_{1i} integrates over the participant's entire trajectory, including both on-drug and off-drug phases.

6.5 Exclusion of contaminated observations

[38]

Excluding the first 1–2 off-drug timepoints removes the most contaminated observations; later off-drug timepoints provide a cleaner contrast.

- Early off-drug timepoints: highest residual drug effect, BM-BR correlation actively decaying.
- Later off-drug timepoints: fully washed out, cleaner contrast with on-drug phase.
- Result: comparison between fully-on and fully-off states, avoiding the strained intermediate zone.

rgt to complete.

The first 1-2 off-drug timepoints carry the most carryover contamination: the residual drug effect is highest, and the BM-BR correlation is in the process of decaying. Later off-drug timepoints, where the

drug has fully washed out, provide a cleaner contrast with the on-drug phase. Excluding the early off-drug observations and retaining only the later ones yields a comparison between fully-on and fully-off states, avoiding the intermediate zone where the analysis model is most strained.

7.6 6.6 Design-level solutions

Rather than adapting the analysis model to accommodate carryover, the trial design itself can be modified to reduce carryover exposure:

- **Longer washout periods** between drug phases eliminate residual drug effect before off-drug measurements begin.
- **More on-drug timepoints** relative to off-drug increase the proportion of high-signal observations.
- **Front-loaded drug exposure** (as in the OL+BDC design) places a long initial on-drug block before the off-drug baseline-during-continued-treatment phase, so that the BM-BR correlation has been fully established before any off-drug observations are collected. The alternative ordering (off-drug first, then on-drug) has the appealing property of placing the placebo and time-varying estimation in an uncontaminated window but sacrifices on-drug exposure duration and forces the BM-BR signal to be re-established after a wash-in; the net power trade-off depends on the ratio of carryover half-life to total trial length and is explored empirically in the Section 3 results.

[39]

Design-level changes address the root cause rather than the symptom; cost is longer trial or reduced ability to estimate carryover dynamics.

- Three levers: longer washout, more on-drug timepoints, front-loaded drug exposure (as in OL+BDC).
- Front-loading establishes BM-BR correlation before any off-drug observations; alternative ordering (off-drug first) trades clean placebo estimation for reduced on-drug exposure.
- Addresses root cause; cost is typically longer trial or reduced carryover estimation ability.

rgt to complete.

Design-level solutions address the root cause (carryover in the data) rather than the symptom (power loss in the analysis). Their cost is typically a longer trial duration or reduced ability to estimate the carryover dynamics themselves.

6.7 Qualitative comparison

Table~1 summarizes the strategies along three dimensions: whether they discard observations, implementation complexity, and the expected direction of benefit based on the mechanistic account in Section 3.2. All entries in the ‘Hypothesized benefit’ column are qualitative; no simulation evidence is presented here, and each strategy should be evaluated empirically before use in trial planning.

Table 1: Qualitative summary of analysis strategies for recovering power under Architecture B carryover. All benefit assessments are mechanistic hypotheses, not simulation-supported estimates.

Strategy	Discards data	Complexity	Hypothesized benefit
On-drug only	Yes	Low	Eliminates correlation erosion; forfeits off-drug reference (best for OL/Hybrid)
Weighted	No	Low	Reduces erosion damage proportionally; depends on weight calibration
Within-subject contrast	Aggregates	Low	Robust to carryover misspecification; less efficient than LME
Two-stage random slopes	No	Moderate	Unknown; separates estimation stages; robustness unverified
Exclude early off-drug	Some	Low	Removes most contaminated observations; improves contrast clarity
Design modification	N/A	N/A	Addresses root cause; cost is longer trial or reduced carryover estimation

A systematic simulation comparison of these strategies under Architecture B, analogous to the factorial comparison of analysis models in Section 3, would identify which approach best recovers the power lost to carryover for each trial design. We regard this comparison as important future work and note that the on-drug-only restriction and the weighted analysis are the most tractable candidates for a first empirical evaluation.

7. Relationship to the Broader Literature

7.1 Treatment effect heterogeneity

[40]

The PATH Statement distinguishes risk modeling from effect modeling for predictive HTE analysis; this is an analysis-model distinction, not a DGP distinction.

- Rothwell (2005) establishes the predictive/prognostic biomarker distinction; Kent et al. (2020) synthesize it into the PATH Statement.
- PATH’s risk modeling approach stratifies on a multivariable risk score; effect modeling adds explicit interaction terms.
- PATH favors risk modeling and cautions against data-driven effect modeling because of low power, multiplicity, and overfitting.

rgt to complete.

The statistical literature on treatment effect heterogeneity is extensive. Rothwell³ sets out the foundational distinction between *predictive* biomarkers (those that modify the treatment effect) and *prognostic* biomarkers (those that predict outcome regardless of treatment). The current consensus statement on predictive HTE analysis is the Predictive Approaches to Treatment effect Heterogeneity (PATH) Statement of Kent et al.⁴, which distinguishes two analytic strategies for identifying predictive biomarker effects in randomized trials: a *risk modeling* approach, in which a multivariable risk model (developed without a treatment term) is used to stratify patients and examine variation in absolute treatment benefit across risk strata, and an *effect modeling* approach, in which the regression includes explicit treatment-by-covariate interaction terms. PATH expresses caution about data-driven effect modeling on the grounds of low statistical power, multiplicity, and overfitting, and favors risk modeling when the application criteria are met.

[41]

PATH’s analysis-model axis (risk vs. effect modeling) is orthogonal to our DGP axis; both PATH strategies implicitly assume Architecture A.

- PATH’s distinction is analysis-model; ours is DGP; they are orthogonal.
- Both PATH strategies assume heterogeneity is in the mean structure (Architecture A); Architecture B does not fit either PATH category.
- Architecture A = parametric interaction model; Architecture B = latent-class/mixture-model perspective.

rgt to complete.

We note that the PATH risk-versus-effect-modeling distinction is an analysis-model axis. The Architecture A versus Architecture B distinc-

tion introduced in this paper is a data-generating-process axis, and the two are orthogonal. Both PATH analytic strategies implicitly assume that the underlying treatment effect heterogeneity is encoded in the mean structure of the outcome (the heterogeneity is identifiable from differences in conditional means), which corresponds to Architecture A as a DGP. Architecture B as a DGP is a separate phenomenon that does not map cleanly onto either of PATH’s analysis-model categories: the interaction signal resides in the covariance structure rather than the mean, and the standard regression-based predictive HTE analysis methods endorsed by PATH may have reduced power under Architecture B precisely because they were developed under an implicit Architecture A assumption. Architecture A can accordingly be characterized as a *parametric interaction model* (in PATH’s effect modeling sense). Architecture B can be characterized as a *latent class* or *mixture model* perspective in which the biomarker is a noisy indicator of class membership.

8.2 7.2 Prevalence of each architecture

[42]

Architecture A is the near-universal standard across the biomarker-moderated trial simulation literature; Architecture B is absent from all major threads except Hendrickson et al.

- Simon (2010), Freidlin and Korn (2014), Renfro and Sargent (2017), Haller and Ulm (2018), and the platform trial literature all use Architecture A.
- N-of-1 simulation studies (1997, 2016, 2013) use hierarchical random-effects models, a variant of Architecture A.
- Wang and Schork (2019) and Schork (2022) share authorship with Hendrickson et al. yet adopt Architecture A, suggesting the differential-correlation choice is paper-specific.

rgt to complete.

A survey of the simulation methodology literature reveals that Architecture A (direct mean moderation) is the near-universal standard. Simon⁵ employs it throughout the simulation frameworks he develops for biomarker-stratified designs; Freidlin and Korn⁶ adopt it for adaptive enrichment trials; Renfro and Sargent⁷ use it for basket and umbrella trials; and it underlies the platform trial literature throughout. Haller and Ulm⁸, in simulation work on biomarker-treatment interaction estimation in randomized trials with prognostic covariates, likewise adopt the standard regression specification, focusing on power and bias of the interaction estimator under different covariate adjustment strategies. N-of-1 simulation studies, including

Zucker et al. (1997), Araujo et al.¹⁰, and Duan et al.²³, use hierarchical random-effects models in which individual treatment effects are drawn from a population distribution, a variant of Architecture A in which the biomarker predicts the random treatment effect through the mean structure. Methodological work on power and design for N-of-1 trials with serial correlation and carryover, from both Wang and Schork¹¹ and Schork¹², adopts the same Architecture A specification, even though this work shares senior authorship with Hendrickson et al.¹³. It appears, then, that the differential-correlation parameterization in the latter is a paper-specific design choice rather than a methodological convention of any laboratory or research program.

[43]

Architecture B appears in latent variable / structural equation modeling (SEM) / copula contexts but is unique in N-of-1 clinical trial simulation.

- Architecture B appears in latent variable / SEM / Bayesian joint modeling / copula contexts (e.g., Rizopoulos 2012).
- In N-of-1 clinical trial simulation, Hendrickson et al. (2020) appears unique in using it.

rgt to complete.

Architecture B (differential correlation via MVN) appears primarily in latent variable and structural equation modeling contexts, including the joint modeling framework discussed by Rizopoulos²⁴ (where treatment-state-dependent covariance relates longitudinal biomarkers to survival outcomes rather than to treatment response directly), Bayesian joint modeling frameworks, and copula-based simulation studies. In the clinical trial simulation literature specifically, the Hendrickson et al. (2020) framework appears to be unique in using differential correlation as the mechanism for biomarker-treatment interaction in an N-of-1 trial context.

[44]

The near-universal use of Architecture A means most published power estimates are optimistic for Architecture B biomarkers under carryover.

- Virtually all published power estimates for biomarker-moderated trials assume Architecture A.
- These estimates may be optimistic for biomarkers that operate through Architecture B (covariance-structure mechanism).
- The additional carryover-induced power loss under Architecture B is not accounted for.

rgt to complete.

This predominance of Architecture A has an important implication: the power estimates reported in the vast majority of biomarker-moderated trial simulation studies may be optimistic for biomarkers that operate through an Architecture B mechanism, because these studies do not account for the additional power loss that carryover produces when the interaction signal resides in the covariance structure rather than the mean structure.

8.3 7.3 Crossover and N-of-1 trial methodology

[45]

The crossover trial literature addresses carryover as a nuisance parameter but ignores the architectural distinction; the standard recommendation applies differently under each.

- Senn (2002), Jones and Kenward (2014), and Sturdevant and Lumley (2021) treat carryover as a nuisance parameter; none engages with DGP architecture.
- Standard recommendation (include carryover in model or use adequate washout) applies to both architectures.
- Under Architecture B, crossover designs with short washout periods may be less suitable than previously recognized.

rgt to complete.

The crossover trial literature, represented most thoroughly by Senn²⁵ and Jones and Kenward²⁶, addresses carryover extensively but does not distinguish between architectures for biomarker-treatment interaction modeling. More recent methodological work by Sturdevant and Lumley²⁷ on testing for carryover effects via mixed-effects models likewise treats carryover as a nuisance parameter to be estimated and adjusted for, without engaging with how carryover interacts with biomarker-treatment effect signals. The standard recommendation, namely to always include carryover in the analysis model or to use adequate washout, applies to both architectures but has different power implications under each. The finding that Architecture B is more sensitive to carryover suggests that crossover designs with short washout periods may be less suitable for detecting correlation-based biomarker interactions than previously recognized.

8.4 7.4 Precision medicine trial design

[46]

Enrichment and biomarker-stratified designs assume Architecture A for power calculations; the architectural choice matters differently across three design families.

- Simon (2010) and Freidlin and Korn (2014) frame power calculations in the Architecture A regression idiom.
- Three design families warrant separate consideration: adaptive enrichment, treatment-switching, and basket/umbrella trials.
- The architectural assumption is rarely made explicit, so the practical implications are easy to overlook.

rgt to complete.

Enrichment and biomarker-stratified designs, as discussed by Simon⁵ and Freidlin and Korn⁶, use Architecture A exclusively for power calculations. Three families of design are particularly worth considering in light of the architectural distinction.

[47]

Adaptive enrichment trials are vulnerable to Architecture B power shortfalls precisely when the biomarker mechanism is most uncertain.

- Interim power calculations driving adaptive decisions use Architecture A; if Architecture B holds and off-drug periods occur, realized power will be lower than predicted.
- The mismatch biases adaptation rules toward premature futility stopping or overly narrow enrichment criteria.
- The problem is worst when the biomarker mechanism is unknown, exactly when adaptive enrichment is most often proposed.

rgt to complete.

Adaptive enrichment trials adjust the eligible population mid-trial based on emerging evidence of a biomarker-treatment interaction. The interim power calculations that drive these adaptive decisions are universally framed in the Architecture A regression idiom. If the true biomarker mechanism is closer to Architecture B and any participants pass through washout intervals between dose-finding stages, the realized power for the interaction test will be lower than the adaptation rule expects, biasing decisions toward premature futility stopping or toward overly narrow enrichment criteria. The discrepancy is most acute when the biomarker mechanism is uncertain at the design stage, which is precisely the situation in which adaptive enrichment is most attractive.

[48]

Treatment-switching designs suffer Architecture B power loss through off-drug correlation erosion; parallel-group enrichment designs do not, because there is no off-drug phase.

- Crossover, basket, platform, and N-of-1 hybrid designs all cycle

participants through drug states; each off-drug phase erodes BM-BR correlation under Architecture B.

- Parallel-group enrichment designs escape this loss because each participant stays on a single treatment throughout.
- The architectural choice matters most in within-subject designs, where it is also most likely to be clinically important.

rgt to complete.

Treatment-switching designs (most crossover, basket, and platform trials, as well as N-of-1 hybrid designs) cycle individual participants through different drug states. Under Architecture B, each off-drug phase erodes the BM-BR correlation, so the biomarker-treatment interaction signal in the second and subsequent treatment periods is structurally weaker than in the first. The power loss documented in Section 3 applies here directly. By contrast, parallel-group enrichment designs, in which each participant remains on a single treatment for the duration of the trial, escape this loss because there is no off-drug phase during which the correlation can decay; for these designs, Architectures A and B yield similar power estimates. The architectural choice therefore matters most in exactly the designs where the biomarker mechanism is most likely to be of clinical interest, namely those that deliberately expose each participant to both treatment states in order to estimate within-subject responsiveness.

[49]

Basket and umbrella trials face an additional risk of optimistic aggregated power estimates when Architecture B holds in even a subset of arms.

- Renfro and Sargent (2017) discuss power calculations that average across biomarker-defined arms; the calculation assumes Architecture A throughout.
- Mixed architectures across baskets produce optimistic trial-level power estimates that are difficult to detect from aggregate results.
- Subgroup-specific power diagnostics under both architectures are warranted when the biological mechanism is heterogeneous.

rgt to complete.

Basket and umbrella trials, as discussed by Renfro and Sargent⁷, introduce an additional layer because their power calculations average across multiple biomarker-defined arms. If even a subset of the arms operate through an Architecture B mechanism while the trial-wide power calculation assumes Architecture A throughout, the aggregated power estimates will be biased optimistically in a way that is difficult to detect from trial-level results alone. Subgroup-specific power diagnostics under both architectures should accompany the primary

planning analysis whenever the underlying biology is heterogeneous across baskets.

[50]

Across all three design families, the practical recommendation is to anchor sample size planning on the Architecture B estimate when the biomarker mechanism is uncertain.

- Run power calculations under both architectures and treat Architecture B as the conservative anchor.
- This recommendation parallels the advice in Section 5.3 and is consistent with the PATH Statement (2020).
- The PATH Statement’s caution about data-driven interaction discovery reinforces the case for conservative planning under architectural uncertainty.

rgt to complete.

In all three settings the practical recommendation parallels that of Section 5.3: when the biological mechanism is uncertain, run the power calculation under both architectures and treat the Architecture B estimate as the more conservative anchor for sample size planning. This recommendation is consistent with the PATH Statement⁴, which itself counsels caution about data-driven discovery of treatment-by-covariate interactions and prefers risk modeling approaches when the application criteria are met.

8. Conclusions

In conclusion, six points are to be emphasized.

1. The choice between direct mean moderation (Architecture A) and differential correlation (Architecture B) for generating biomarker-treatment interactions in clinical trial simulations has substantive consequences for power estimates under carryover.
2. Under Architecture A, carryover reduces power only modestly, by 1 to 3% in relative terms across the three designs examined, because the proportional biomarker-drug relationship is preserved during the off-drug period. Under Architecture B the loss is strongly design-dependent: it remains modest in the crossover and hybrid designs (3 to 5%) but reaches roughly 31% in the open-label OL+BDC design, because the differential-correlation signal erodes as the drug effect wanes.

3. The choice should be guided by the hypothesized biological mechanism: Architecture A for biomarkers that directly determine drug pharmacokinetics or pharmacodynamics; Architecture B for biomarkers that statistically predict treatment response through latent subtype membership.
4. For the prazosin-PTSD application with blood pressure as the predictive biomarker, Architecture B (MVN differential correlation) is the more appropriate model, and the resulting power sensitivity to carryover should inform trial design decisions regarding off-drug period duration.
5. The existing clinical trial simulation literature uses Architecture A almost exclusively. The Hendrickson et al. (2020) MVN approach is, to the best of our knowledge, unique in the N-of-1 trial context. This means that published power estimates for biomarker-moderated crossover and N-of-1 designs may be systematically optimistic for biomarkers that operate through a correlation-based (Architecture B) mechanism, because they do not account for the carryover sensitivity documented here.
6. Simulation studies for biomarker-moderated treatment effects should explicitly report their DGP architecture and consider sensitivity analyses under the alternative when the biological mechanism is uncertain.

10 Conflict of interest

The authors declare no conflicts of interest.

11 Funding

This work received no specific funding.

12 Data availability statement

The data and code that support the findings of this study are openly available in the pmsimstats-ng project repository. Per-replicate Monte Carlo outputs and ADEMP-compliant cell-level summaries are versioned under `analysis/data/`; simulation drivers are at `analysis/scripts/`. Exact paths for this manuscript are listed in the Reproducibility section below.

13 Reproducibility

This manuscript was rendered with R~4.6.0 inside the project’s zzcollab Docker container (`rocker/tidyverse:4.6.0`; image hash recorded in `Dockerfile`; package set captured in `renv.lock`). Principal packages used by the simulation pipeline are `nlme` (continuous-time linear-mixed analysis with `corCAR1` residual correlation), `parallel` with `mclapply` (Section~3.1 and 3.3 production runs), `data.table` (the `original-extended` simulation collection), `purrr` (the `tidyverse` simulation collection), and `rmarkdown` for the manuscript build.

Sample-size convention. The Section~3.1 driver (`01-dgp-prototype.R`) generates $N = 35$ subjects independently per randomization path, then concatenates paths before analysis; for two-path designs (CO, OL+BDC) the effective analysis sample is therefore 70 subjects, and for the four-path Hybrid design it is 140 subjects.

Random-number generators. The Section~3.1 driver sets `set.seed(20260509)` and uses per-replicate seeds derived as `100000 * cell_index + rep_index`.

Data and code availability. Sections 3.1 and 3.3: per-replicate outputs at `analysis/data/quick-sim/01-dgp-replicates.rds`; ADEMP cell-level summaries at `01-dgp-summary.txt`; driver at `analysis/scripts/quick-sim/01-dgp-prototype.R`. The manuscript source, bibliography, and CSL file are at `analysis/report/01-dgp-mean-moderation-vs-mvn/`.

14 References

1. pmsimstats team. Type I calibration of fixed-effect interaction tests in linear mixed models: A staged diagnosis of working-covariance misspecification and a cluster-robust remedy. Published online 2026.
2. pmsimstats team. Test-procedure choice for the biomarker-treatment interaction in aggregated N-of-1 trials: Dichotomization, classical RM-ANOVA, mixed-effects, and generalized estimating equations. Published online 2026.
3. Rothwell PM. Treating individuals 5. Subgroup analysis in randomised controlled trials: Importance, indications, and interpretation. *The Lancet*. 2005;365(9454):176-186.
4. Kent DM, Paulus JK, Klaveren D van, et al. The predictive approaches to treatment effect heterogeneity (PATH) statement. *Annals of Internal Medicine*. 2020;172(1):35-45.
5. Simon R. Clinical trial designs for evaluating the medical utility of prognostic and predictive biomarkers in oncology. *Personalized Medicine*. 2010;7(1):33-47.

- 1584 6. Freidlin B, Korn EL. Biomarker enrichment strategies: Matching trial design to biomarker credentials. *Nature Reviews Clinical Oncology*. 2014;11(2):81-90.
- 1585 7. Renfro LA, Sargent DJ. Statistical controversies in clinical research: Basket trials, umbrella trials, and other master protocols. *Annals of Oncology*. 2017;28(1):34-43.
- 1586 8. Haller B, Ulm K. A simulation study on estimating biomarker-treatment interaction effects in randomized trials with prognostic variables. *Trials*. 2018;19(1):128.
- 1587 9. Zucker DR, Schmid CH, McIntosh MW, D'Agostino RB, Selker HP, Lau J. Combining single patient (N-of-1) trials to estimate population treatment effects and to evaluate individual patient responses to treatment. *Journal of Clinical Epidemiology*. 1997;50(4):401-410.
- 1588 10. Araujo A, Julious S, Senn S. Understanding variation in sets of N-of-1 trials. *PLoS ONE*. 2016;11(12):e0167167.
- 1589 11. Wang Y, Schork NJ. Power and design issues in crossover-based N-of-1 clinical trials with fixed data collection periods. *Healthcare (Basel)*. 2019;7(3):84.
- 1590 12. Schork NJ. Accommodating serial correlation and sequential design elements in personalized studies and aggregated personalized studies. *Harvard Data Science Review*. 2022;4(SI3). doi:10.1162/99608f92.c4b04c64
- 1591 13. Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: Statistical methods and practical considerations. *Frontiers in Digital Health*. 2020;2:13. doi:10.3389/fdgth.2020.00013
- 1592 14. pmsimstats team. Latent-class and mixture-model formulations for biomarker-treatment interaction in N-of-1 trials. Published online 2026.
- 1593 15. Raskind MA, Peterson K, Williams T, et al. A trial of prazosin for combat trauma PTSD with nightmares in active-duty soldiers returned from iraq and afghanistan. *American Journal of Psychiatry*. 2013;170(9):1003-1010.
- 1594 16. Raskind MA, Peskind ER, Chow B, et al. Trial of prazosin for post-traumatic stress disorder in military veterans. *New England Journal of Medicine*. 2018;378(6):507-517.

- 1595 17. pmsimstats team. A combined data-generating architecture for biomarker-treatment interactions: Channel super-additivity and mixing-weight sensitivity in aggregated N-of-1 trials. Published online 2026.
- 1596 18. Raison CL, Rutherford RE, Woolwine BJ, et al. A randomized controlled trial of the tumor necrosis factor antagonist infliximab for treatment-resistant depression: The role of baseline inflammatory biomarkers. *JAMA Psychiatry*. 2013;70(1):31-41.
- 1597 19. Senn S. Mastering variation: Variance components and personalised medicine. *Statistics in Medicine*. 2016;35(7):966-977. doi:10.1002/sim.6739
- 1598 20. Raskind MA, Peskind ER, Kanter ED, et al. Reduction of nightmares and other PTSD symptoms in combat veterans by prazosin: A placebo-controlled study. *American Journal of Psychiatry*. 2003;160(2):371-373.
- 1599 21. Raskind MA, Peskind ER, Hoff DJ, et al. A parallel group placebo-controlled study of prazosin for trauma nightmares and sleep disturbance in combat veterans with post-traumatic stress disorder. *Biological Psychiatry*. 2007;61(8):928-934.
- 1600 22. Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Statistics in Medicine*. 2019;38(11):2074-2102. doi:10.1002/sim.8086
- 1601 23. Duan N, Kravitz RL, Schmid CH. Single-patient (N-of-1) trials: A pragmatic clinical decision methodology. *Journal of Clinical Epidemiology*. 2013;66(8):S21-S28.
- 1602 24. Rizopoulos D. *Joint Models for Longitudinal and Time-to-Event Data: With Applications in R*. CRC Press; 2012.
- 1603 25. Senn S. *Cross-over Trials in Clinical Research*. 2nd ed. John Wiley & Sons; 2002.
- 1604 26. Jones B, Kenward MG. *Design and Analysis of Cross-over Trials*. 3rd ed. CRC Press; 2014.
- 1605 27. Sturdevant SG, Lumley T. Statistical methods for testing carryover effects: A mixed effects model approach. *Contemporary Clinical Trials Communications*. 2021;22:100711. doi:10.1016/j.conctc.2021.100711

1607 *Rendered on 2026-07-31 at 11:48 PDT.*

1608 *Source: ~/prj/res/36-pmsimstats-ng/pmsimstats-ng/analysis/report/01-dgp-mean-moderation-vs-mv*