

Three data-generating architectures for biomarker-treatment interactions in aggregated N-of-1 trials, and why the choice changes power under carryover: a tutorial

pmsimstats team

2026-06-17

1 Abstract

Background. Simulation is the standard tool for sizing aggregated N-of-1 trials – designs in which each participant is repeatedly crossed on and off treatment and many single-patient series are pooled – intended to validate predictive biomarkers. Any such simulation must encode how the biomarker moderates the treatment effect in its data-generating process (DGP), and the encoding is a substantive choice. This tutorial, aimed at applied statisticians who may be new to the N-of-1 setting, formalises *three* architectures and shows that the choice among them, usually treated as an implementation detail, materially changes the estimated power loss under carryover.

Methods. We define three DGP architectures. *Architecture A* (direct mean moderation): the biomarker scales the treatment effect in the conditional mean. *Architecture B* (differential correlation, the multivariate-normal approach): the interaction is encoded as treatment-state-dependent correlation in the joint distribution. *Architecture C* (combined): both channels are active at once, with independent parameters $c_{bm,A}$ (mean) and $c_{bm,B}$ (correlation), so that A and B are its single-channel special cases. Holding the analysis model fixed (a linear mixed model with continuous-time AR(1) residual correlation), we quantify by Monte Carlo simulation the power to detect the interaction as carryover increases, across three within-subject designs, calibrated to a prazosin/PTSD reference application.

Results. Under Architecture A power is nearly preserved because the biomarker-effect proportionality is maintained at every occasion: at the reference cell it falls only from 0.751 to 0.730 (2.8% relative) as the carryover half-life grows to one week. Under Architecture B the same carryover erodes the differential correlation that carries the signal, and power falls from 0.777 to 0.539 (30.6% relative), roughly eleven times the Architecture A loss and far beyond Monte Carlo error ($z = 7.64$, $p < 10^{-13}$). The discrepancy is strongly design-dependent: large in designs with a discontinuation phase, negligible in a classical crossover whose within-subject AR(1) structure already dominates the correlation channel. Architecture C behaves intermediately, its carryover sensitivity rising with the weight $c_{bm,B}$ placed on the correlation channel.

Conclusions. The DGP architecture is a substantive scientific assumption, not a parameterisation convenience: it determines how severely carryover appears to reduce power. We recommend that biomarker-validation simulation studies state which of the three architectures they adopt and, when the mechanism is uncertain, report power across the relevant architectures.

2 1. Introduction

[1]

Power for these designs is estimated by simulation, and the simulation must encode the biomarker interaction in its DGP – in more than one possible way.

- The target is the biomarker-by-treatment interaction.
- Power has no closed form here, so it is obtained by Monte Carlo simulation.
- How the interaction is built into the DGP is a modelling choice, the subject of this paper.

To be completed by rgt.

Precision-medicine trials increasingly ask whether a baseline **biomarker** modifies the treatment effect – a biomarker-by-treatment interaction. For the designs considered here the power of a proposed trial cannot be derived in closed form, so it is estimated by Monte Carlo simulation: many trials are generated, analysed, and the rejection rate recorded. The point that motivates this paper is that the simulation must itself encode the interaction in its data-generating process, and there is more than one way to do so.

[2]

There are three architectures: A (mean moderation), B (differential correlation), and C, which activates both channels at once.

- Architecture A: an interaction term scales the treatment effect in the mean; standard across trial-simulation frameworks.
- Architecture B: the biomarker changes the response-outcome correlation, encoding the interaction in the joint distribution (Hendrickson’s N-of-1 approach).
- Architecture C: both channels act together, with A and B as its single-channel special cases.

To be completed by rgt.

We distinguish three architectures. The common one, **direct mean moderation** (Architecture A), adds an interaction term to the outcome equation so the biomarker scales the size of the treatment effect; essentially every standard trial-simulation framework uses it^{1,2}. The second, **differential correlation** (Architecture B), used by Hendrickson et al.³ for N-of-1 trials, instead makes the biomarker change the correlation between treatment response and outcome, so the interaction is encoded in the joint distribution rather than in the mean. The third, **combined** (Architecture C), activates both channels simultaneously through two independent parameters; A and B are exactly its single-channel limits. We treat all three as co-equal data-generating architectures, defined precisely in Section 2.

[3]

The two encodings, found in two implementations of the same software, gave different carryover stories; this paper explains why.

- The original code used Architecture B; a re-implementation used Architecture A.
- The two disagreed about how much carryover hurts power.
- We formalise both and give guidance on which to use.

To be completed by rgt.

We ran into both while building the pmsimstats software: the original code used Architecture B, while a re-implementation had quietly switched to Architecture A, and the two gave different stories about carryover. This paper pins down why. We define the two architectures precisely, show by simulation that they react very differently to carryover, and give simple guidance on which to use.

[4]

Primer: aggregated N-of-1 trials pool repeated within-patient on/off contrasts, and carryover is their distinguishing nuisance.

- A single N-of-1 trial repeatedly crosses one patient on/off; pooling many gives population inference at small N .
- Carryover is residual drug effect decaying over days to weeks, contaminating off-drug occasions.
- Three within-subject designs (CO, OL+BDC, Hybrid) differ in how on/off occasions are arranged.

To be completed by rgt.

The N-of-1 setting, for readers new to it. In a single N-of-1 trial one participant is repeatedly crossed between active treatment and a comparator over a sequence of measurement occasions, so the within-participant on-versus-off contrast estimates that participant's individual treatment effect; an *aggregated* N-of-1 trial pools many such series in a mixed model to obtain population-level inference at far smaller sample sizes than a parallel-group design. The feature that distinguishes these designs analytically is *carryover*: residual drug effect that decays over days to weeks after discontinuation, contaminating off-drug occasions. We study three within-subject designs that differ in how on- and off-drug occasions are arranged – a traditional crossover (CO), an open-label-plus-blinded-discontinuation design (OL+BDC), and a Hybrid N-of-1 design – defined in Section 3.

3 2. Three DGP Architectures

[5]

Notation: B_i biomarker, D_{it} on-drug indicator, BR_{it} biological response, with Dbc the analysis-layer exposure.

- B_i , D_{it} , BR_{it} describe the data-generating layer.
- The analysis layer uses a continuous exposure Dbc (Section 3.1).
- Typewriter identifiers denote code/formula variables; math symbols are used otherwise.

To be completed by rgt.

Notation. Throughout this paper, B_i denotes the biomarker value of participant i , D_{it} the binary on-drug indicator at timepoint t (1 on drug, 0 off drug), and BR_{it} the

biological response component of the outcome. These three symbols describe the data-generating-process layer. The analysis-model layer uses a related but distinct continuous drug exposure variable, `Dbc`, defined in Section 3.1. Code-style identifiers in typewriter font (`bm`, `Dbc`, `BR`) refer to specific variables in the `pmsimstats` codebase or to R formula syntax; mathematical symbols are used otherwise.

3.1 2.1 Architecture A: Direct Mean Moderation

In this approach the biomarker enters the mean structure of the outcome directly. The treatment effect for participant i at timepoint t is:

$$Y_{it} = \mu_0 + \beta_1 \cdot D_{it} \cdot (1 + \beta_{bm} \cdot B_i) + \epsilon_{it}$$

[6]

Under Architecture A the interaction is explicit in the mean: a higher biomarker gives a proportionally larger treatment effect.

- β_{bm} is the biomarker moderation parameter on the centered biomarker.
- The effect scales with B_i deterministically in the conditional mean.
- The signal lives in the first moment of the outcome.

To be completed by rgt.

where D_{it} is the treatment indicator (or a continuous measure of drug exposure), B_i is the participant's biomarker value (typically centered), and β_{bm} is the biomarker moderation parameter. The interaction is explicit in the mean: participants with higher biomarker values have proportionally larger treatment effects.

Key properties:

- The interaction is **deterministic** given the biomarker value and treatment status.
- The biomarker moderation β_{bm} is a **population-level parameter** that applies uniformly to all participants.
- The signal for the interaction is in the **first moment** (conditional mean) of the outcome distribution.
- Adding noise, random effects, or carryover to D_{it} scales the entire treatment effect (including its biomarker-dependent component) proportionally. The **ratio** of drug effect between high-biomarker and low-biomarker participants is preserved.

[7]

The multiplicative form maps onto the standard additive effect-modeling regression, so both are instances of Architecture A.

- The additive interaction β_3 equals $\beta_1 \cdot \beta_{bm}$ in the multiplicative form.
- The additive form adds a prognostic main effect β_2 the multiplicative form sets to zero.
- For interaction power the two are interchangeable up to that prognostic distinction.

To be completed by rgt.

We note that the multiplicative form above maps directly onto the standard additive regression form for predictive biomarker effects used in the trial methodology literature, specifically the effect modeling equation of Kent et al.⁴:

$$Y_{it} = \beta_0 + \beta_1 D_{it} + \beta_2 B_i + \beta_3 D_{it} B_i + \epsilon_{it}$$

The interaction coefficient β_3 in the additive form corresponds to the product $\beta_1 \cdot \beta_{bm}$ in the multiplicative form. The additive form additionally accommodates a prognostic main effect β_2 for the biomarker, which the multiplicative form sets to zero by construction. For the purpose of power calculations on the interaction coefficient, the two forms are interchangeable up to this prognostic-effect distinction, and both fall within Architecture A.

This architecture is used in the pedagogical simulation (`implementations/simple/simulation.R`) and in the exploratory carryover analyses (`simulation_carryover_spectrum.R`) of the pmsimstats project.

3.2 2.2 Architecture B: Differential Correlation (MVN)

In this approach the biomarker and the biological response (BR) component are jointly drawn from a multivariate normal distribution with treatment-state-dependent correlation:

$$\begin{pmatrix} B_i \\ BR_{it} \end{pmatrix} \sim \text{MVN} \left(\begin{pmatrix} \mu_B \\ \mu_{BR}(t, D_{it}) \end{pmatrix}, \Sigma(D_{it}) \right)$$

where the covariance matrix Σ has the following treatment-state-dependent BM-BR correlation:

$$\text{Cor}(B_i, BR_{it}) = \begin{cases} c_{bm} & \text{if } D_{it} = 1 \text{ (on drug)} \\ c_{bm} \cdot e^{-\lambda \cdot t_{sd}} & \text{if } D_{it} = 0 \text{ (off drug)} \end{cases}$$

where t_{sd} is time since drug discontinuation. The exponential-decay form is the default; the limiting cases $\lambda = 0$ (no decay; correlation persists indefinitely during the off-drug phase) and $\lambda \rightarrow \infty$ (immediate decay; correlation drops to zero the instant the drug is discontinued) recover the two no-carryover model variants.

The interaction is not in the mean structure: the population mean of BR_{it} is the same for all participants with the same treatment history. Instead it emerges from the **conditional distribution**: given a participant's biomarker value $B_i = b$, the conditional expectation of their biological response is:

$$E[BR_{it} | B_i = b, D_{it} = 1] = \mu_{BR}(t) + c_{bm} \cdot \frac{\sigma_{BR}}{\sigma_B} \cdot (b - \mu_B)$$

[8]

Under Architecture B the interaction is in the second moment: the biomarker modulates BR through a correlation that decays during washout.

- The conditional mean of BR rises with b via the correlation, not a mean parameter.
- Effective correlation is c_{bm} on-drug, $c_{bm} e^{-\lambda t_{sd}}$ off-drug.

- Moderation persists with attenuated strength off-drug and vanishes only at complete decay.

To be completed by rgt.

This conditional expectation has an interaction-like structure: the biomarker value b modulates the expected BR. The effective correlation is c_{bm} on drug and $c_{bm} \cdot e^{-\lambda t_{sd}}$ during the washout; the moderation therefore persists with attenuated strength during off-drug periods and vanishes only in the limit of complete decay.

Key properties:

- The interaction is **probabilistic**: it manifests as a tendency, not a deterministic scaling.
- The biomarker moderation c_{bm} is a **correlation parameter** that governs the strength of the association.
- The signal for the interaction is in the **second moment** (covariance structure) of the joint distribution.
- Carryover effects in the DGP inflate the BR means during off-drug periods, making them resemble on-drug means. This reduces the **differential** in the correlation structure between treatment states, weakening the signal that the analysis model detects.

This architecture is used in the Hendrickson et al.³ publication code and the audited R/ package in pmsimstats-ng.

3.3 2.3 Architecture C: Combined Mean Moderation and Differential Correlation

Architecture C (combined) applies both interaction mechanisms in sequence. First, the MVN covariance matrix is constructed with BM-BR differential correlation weighted by `c.bm_b`, exactly as under Architecture B. Second, after the MVN draw, the BR component is additively shifted at on-drug timepoints by

$$\Delta BR_{it} = c_{bm,A} \cdot z_{B,i} \cdot \sigma_{BR} \cdot \mathbf{1}[\text{on-drug}]$$

where $z_{B,i} = (B_i - \mu_B)/\sigma_B$ is the standardized biomarker for participant i , exactly as under Architecture A.

The two parameters `c.bm_a` and `c.bm_b` are independent. Setting $c_{bm,B} = 0$ removes the covariance channel and recovers pure Architecture A; setting $c_{bm,A} = 0$ removes the mean-shift channel and recovers pure Architecture B. The boundary cases are enforced algebraically rather than by convention, so the special cases serve as internal validation gates for any simulation study using Architecture C.

We note that when both parameters are set to their full values ($c_{bm,A} = c_{bm,B} = 0.45$), the total interaction effect on the BR mean is approximately double that of either pure architecture at $c_{bm} = 0.45$. A researcher wishing to hold total effect size constant across architectures may instead parameterize as $c_{bm,A} = (1 - \alpha) \cdot c$ and $c_{bm,B} = \alpha \cdot c$ for a mixing weight $\alpha \in [0, 1]$ and total weight c . We do not adopt this constraint in the simulations below, preferring the factorial parameterization that spans the full $(c_{bm,A}, c_{bm,B})$ space.

[9]

The two channels are mechanistically distinct: A is a pharmacodynamic mean shift, B a co-regulatory coupling in the covariance.

- Channel A: the biomarker governs effective dose/receptor occupancy, amplifying the mean drug effect.
- Channel B: on-drug periods couple biomarker and response, a coupling absent in washout.
- They need not scale in proportion, motivating the combined Architecture C.

To be completed by rgt.

The two channels are mechanistically distinct and need not scale in proportion. The mean-shift channel (Architecture A) reflects a pharmacodynamic relationship: the drug's biological effect on BR is amplified in participants whose biomarker is above the population mean, because the biomarker governs the effective dose or receptor occupancy. This mechanism operates at the level of the mean trajectory for each participant.

The covariance channel (Architecture B) reflects a different process: on-drug periods impose a co-regulatory coupling between biomarker expression and treatment response, perhaps because the drug activates a pathway whose output is gated by the biomarker signal. This coupling is absent during off-drug washout, producing the differential correlation structure that Architecture B encodes in the MVN covariance matrix.

A given drug could activate one, both, or neither channel. The simulation practitioner who is uncertain which channel predominates should regard Architecture C with a range of α values as the natural sensitivity framework.

3.4 2.4 Mathematical Comparison

We begin the comparison by considering the expected outcome difference between two participants with biomarker values b_1 and b_2 at the same drug exposure level D under Architecture A:

$$E[Y | b_1, D] - E[Y | b_2, D] = \beta_1 \cdot \beta_{bm} \cdot (b_1 - b_2) \cdot D$$

The absolute magnitude of this difference scales linearly with D and therefore shrinks during off-drug periods when D is replaced by an exposure-decayed value $D \cdot \text{carryover}$. We note, however, that the **ratio** of the drug effect between high-biomarker and low-biomarker participants is preserved at every level of exposure. The biomarker-by-treatment signal contracts in amplitude but does not lose its identifiability under carryover.

Under Architecture B, the analogous quantity is:

$$E[BR|b_1, D = 1] - E[BR|b_2, D = 1] = c_{bm} \cdot \frac{\sigma_{BR}}{\sigma_B} \cdot (b_1 - b_2)$$

During off-drug periods with carryover:

$$E[BR|b, D = 0, t_{sd}] = \mu_{BR, \text{off}}(t_{sd}) + c_{bm} \cdot e^{-\lambda t_{sd}} \cdot \frac{\sigma_{BR}}{\sigma_B} \cdot (b - \mu_B)$$

Here the biomarker-dependent component decays exponentially with time off drug. As t_{sd} increases, the off-drug conditional expectation converges to $\mu_{BR, \text{off}}$ for all biomarker values, eliminating the differential signal entirely.

Under Architecture C both channels are active. The analysis model’s interaction coefficient depends on the covariance between B_i and D_{bc} -weighted outcomes. The mean-shift channel contributes a term that degrades exactly as under Architecture A: the Dbc contrast is compressed during carryover but the ratio of the drug effect between biomarker groups is preserved at every exposure level. The covariance channel contributes a term that degrades as under Architecture B: the BM-BR correlation generating the second-moment signal decays exponentially off-drug. The two channels are additive at the level of the total biomarker-by-drug covariance, so the net power loss under Architecture C lies strictly between the pure-architecture extremes when $c_{bm,A}, c_{bm,B} > 0$. As either parameter approaches zero the corresponding boundary case is recovered.

4 3. Consequences for Statistical Power Under Carryover

4.1 3.1 Empirical demonstration

We ran identical simulation configurations under both architectures using the `pmsimstats-ng` framework, matching the DGP parameters as closely as possible between the two approaches. We note that the non-biologic response components (time-varying response, placebo-belief accumulation, and residual noise) are generated from the same multivariate normal draw with identical means, variances, within-factor AR(1) autocorrelation, and cross-factor correlations in both architectures. The two DGPs differ *only* in how the biomarker-to-biological-response (BM-BR) linkage is encoded: Architecture~B embeds it in the off-diagonal c_{bm} entries of Σ , whereas Architecture~A omits those entries and adds a post-draw shift $c_{bm} \cdot B_i^* \cdot \sigma_{BR}$ to on-drug BR columns (see `19-mean-moderation-implementation-notes.tex` for the scaling derivation). The divergent carryover behavior reported below is therefore attributable to the BM-BR parameterization, not to any difference in how the placebo, time-varying, or residual-noise components are generated. The analysis model was the same in both cases: an `nlme::lme` fit with a `corCAR1` continuous-time autocorrelation structure and a single biomarker-by-drug interaction term, written `bm:Dbc` in R formula syntax.

The analysis-model drug exposure variable `Dbc` is the *continuous* analogue of the binary DGP indicator D_{it} . It equals 1 during on-drug periods and $\exp(-\lambda t_{sd})$ during off-drug periods, where t_{sd} is time since drug discontinuation and λ is the analysis-model carryover decay rate (typically derived from the drug’s pharmacokinetic half-life as $\lambda = \ln 2 / t_{1/2}$, in which case it matches the DGP decay rate of Section 2.2 by construction). Whereas D_{it} flips abruptly between 0 and 1, `Dbc` captures the smooth shrinkage of residual drug effect during washout. The interaction coefficient on `bm:Dbc` is therefore the test of whether the biomarker moderates the exposure-decayed drug effect.

We follow the ADEMP framework of Morris, White, and Crowther (2019): Aims, Data-generating mechanisms, Estimands, Methods, Performance measures.

Setup:

- Trial designs: CO (traditional crossover), Hybrid (N-of-1 Hybrid), OL+BDC (open-label with blinded discontinuation)
- Sample size: $N = 35$ per randomisation path (effective analysis $N = 70$ for CO and OL+BDC, 140 for Hybrid, after the paths are concatenated)
- Biomarker correlation / moderation: $c_{bm} = 0.45$ / $\beta_{bm} = 0.45$

- DGP carryover half-life: $t_{1/2} \in \{0, 0.5, 1.0\}$ weeks
- Analysis: matched Dbc with same half-life
- Replicates: 1,000 per cell (Monte Carlo standard error on a binomial power estimate near 0.5 is approximately $\sqrt{0.5 \cdot 0.5 / 1,000} \approx 0.016$)
- Implementation: 8-core `parallel::mclapply`; full factorial 2 architectures $\times$ 3 designs $\times$ 3 $t_{1/2} \times 2 c_{bm}$ levels = 36 cells, 36,000 fits, 100% convergence

The $N = 35$ choice is the smallest of the prazosin-PTSD- calibrated sample sizes used in the published methodological literature; it is a regime in which neither architecture saturates power at the no-carryover cells, so any carryover-induced erosion is mechanically visible. Type I error under the $c_{bm} = 0$ null cells sat in $[0.010, 0.074]$ across the 18 null cells, consistent with the conservative-to-nominal regime expected for the continuous D_{bc} specification at this sample size.

Results under Architecture A (direct mean moderation):

Design	$t_{1/2} = 0$	$t_{1/2} = 0.5$	$t_{1/2} = 1.0$
CO	0.739	0.739	0.725
Hybrid	0.992	0.987	0.978
OL+BDC	0.751	0.748	0.730

[10]

Under Architecture A the carryover-induced power loss is modest in every design (1-3% relative).

- Relative loss to $t_{1/2} = 1$: 1.9% (CO), 1.4% (Hybrid), 2.8% (OL+BDC).
- Below the 8-13% range earlier exploratory runs suggested.
- Power is essentially intact under carryover.

To be completed by rgt.

Power declines modestly under Architecture A: relative loss from $t_{1/2} = 0$ to $t_{1/2} = 1$ is 1.9% (CO), 1.4% (Hybrid), and 2.8% (OL+BDC). We note that all three values fall well below the 8-13% relative-loss range suggested by earlier exploratory runs in this project for direct-mean-moderation DGPs, though they reproduce the qualitative finding of those runs, namely a modest erosion that leaves power essentially intact.

Results under Architecture B (differential correlation, MVN):

Design	$t_{1/2} = 0$	$t_{1/2} = 0.5$	$t_{1/2} = 1.0$
CO	0.745	0.754	0.723
Hybrid	0.995	0.980	0.944
OL+BDC	0.777	0.694	0.539

[11]

Under Architecture B the loss is much larger and strongly design-dependent, reaching 30.6% at OL+BDC but absent at CO.

- Relative loss: 3.0% (CO, not significant), 5.1% (Hybrid), 30.6% (OL+BDC) – about 11x the Architecture A loss.

- The CO cell fluctuates within MC error, consistent with no architecture-by-carryover interaction there.
- Below the 40-60% range earlier runs suggested, but confirming substantial, design-dependent erosion.

To be completed by rgt.

Power declines more sharply under Architecture B, and most markedly so in the OL+BDC design: relative loss from $t_{1/2} = 0$ to $t_{1/2} = 1$ is 3.0% (CO; not significantly different from Architecture A’s CO loss), 5.1% (Hybrid), and **30.6% (OL+BDC)**. The CO cell shows a small, non-monotone fluctuation ($0.745 \rightarrow 0.754 \rightarrow 0.723$), which is within one Monte Carlo standard error and is consistent with the null finding of no significant architecture-by-carryover interaction in that design. The OL+BDC loss is approximately 11 times the matched Architecture-A OL+BDC loss (2.8%). It falls below the 40-60% relative-loss range suggested by earlier exploratory runs in this project for the MVN differential-correlation DGP, but it confirms the qualitative finding of those runs, namely a substantial and strongly design-dependent erosion.

We compare the carryover-induced power loss between architectures using a difference-of-differences z -statistic. Let $\Delta_X = \hat{\pi}_{X,0} - \hat{\pi}_{X,1}$ denote the within-architecture loss from $t_{1/2} = 0$ to $t_{1/2} = 1.0$ for architecture $X \in \{A, B\}$, estimated from $n = 1,000$ independent replicates at each cell. The B-versus-A gap is $\Delta_B - \Delta_A$, with standard error

$$\text{SE} = \sqrt{\frac{\hat{\pi}_{B,0}(1 - \hat{\pi}_{B,0})}{n} + \frac{\hat{\pi}_{B,1}(1 - \hat{\pi}_{B,1})}{n} + \frac{\hat{\pi}_{A,0}(1 - \hat{\pi}_{A,0})}{n} + \frac{\hat{\pi}_{A,1}(1 - \hat{\pi}_{A,1})}{n}}$$

since the four proportions come from independent replicate sets. All tests are two-sided. At OL+BDC this gives $z = 7.64$ ($p < 10^{-13}$); at Hybrid, $z = 3.96$ ($p < 10^{-4}$); the CO contrast is not statistically distinguishable from zero ($z = 0.29$) because the CO design’s within-subject AR(1) structure dominates the carryover-mediated correlation decay channel that Architecture B uses to encode the biomarker-treatment interaction.

Worked reading of the contrast. At OL+BDC, the within-architecture carryover loss is $\Delta_A = 0.751 - 0.730 = 0.021$ (2.8% relative) and $\Delta_B = 0.777 - 0.539 = 0.238$ (30.6%). The quantity of interest is the difference-of-differences $\Delta_B - \Delta_A = 0.217$ (about 22 percentage points), and the z -statistic above scales it by its standard error to give $z = 7.64$ ($p < 10^{-13}$): the architectures’ carryover behaviour differs well beyond Monte Carlo error. The same contrast is clearly positive for the Hybrid design and indistinguishable from zero for the crossover design.

4.2 3.2 Why the architectures diverge

We turn now to a mechanistic account of why the two architectures produce such different power losses under carryover.

[12]

Under Architecture A, carryover only compresses the treatment contrast; the biomarker signal-to-noise ratio is nearly preserved.

- Carryover shrinks Dbc and inflates off-drug BR means (shared with Architecture B).
- But β_{bm} is a fixed mean parameter with drug-state- independent noise.
- So the interaction variance falls but the signal-to-noise ratio degrades only modestly.

To be completed by rgt.

Under Architecture A, carryover reduces the magnitude of the drug exposure variable seen by the analysis model (D_{bc} drops below 1 during washout), and the mean BR at off-drug timepoints is inflated by the residual drug effect exactly as in Architecture B. Both mechanisms compress the between-state treatment contrast and are shared by the two architectures. What Architecture A lacks is the second channel of signal loss: the biomarker moderation coefficient β_{bm} is a fixed population parameter, the biomarker enters the mean structure directly rather than through a correlation, and the noise structure is independent of drug state. The biomarker-by-drug interaction term therefore has reduced variance (because D_{bc} has less contrast between drug states), but the signal-to-noise ratio degrades only modestly.

Under Architecture B, carryover operates on two levels simultaneously:

1. **Mean blurring:** The BR means during off-drug periods are inflated by residual carryover, making them resemble on-drug BR means. This reduces the treatment contrast that the analysis model uses to distinguish drug effect from no-drug. This channel is shared with Architecture A.
2. **Correlation erosion:** The BM-BR correlation during off-drug periods decays exponentially with t_{sd} . This is not just a statistical modeling choice; it is a structural consequence of the DGP. The correlation between biomarker and BR reflects the degree to which the drug effect is active; as the drug washes out, the correlation weakens because the drug-mediated component of BR variance shrinks relative to the drug-independent component. This channel is *unique* to Architecture B, because only here does the biomarker signal live in the second moment of the joint distribution. Under Architecture A the drug-mediated and drug-independent components of BR variance move in tandem (both are scaled by the same D_{bc}-weighted drug effect), so their ratio, and hence the identifiability of the interaction, is preserved even as amplitudes shrink.

[13]

Architecture B's interaction signal is attacked from both sides, which Architecture A escapes – hence the divergence.

- The **bm:D_{bc}** coefficient depends on the covariance of B_i with D_{bc}-weighted outcomes.
- Under B both the D_{bc} contrast and the BM-BR correlation decay together.
- Under A the second attack is absent, so power is largely preserved.

To be completed by rgt.

The analysis model's interaction coefficient (the coefficient on **bm:D_{bc}**) depends on the covariance between B_i and D_{bc}-weighted outcomes. Under Architecture B, this covariance is being attacked from both sides: the treatment contrast in D_{bc} is compressed, and the BM-BR correlation that generates the covariance signal is simultaneously decaying. That Architecture A escapes the second attack entirely accounts for the divergence in power loss documented above.

4.3 3.3 Robustness check at $N = 70$

[14]

An $N = 70$ confirmation run checks that the architecture ordering survives larger samples.

- Section 3.1 uses $N = 35$, where the loss is mechanically visible (no ceiling saturation).
- The check runs 12 cells at $N = 70$, 800 replicates each, on the OL+BDC design.
- 100% convergence across 9,600 fits.

To be completed by rgt.

The Section~3.1 production results are at $N = 35$, where neither architecture saturates power and the carryover-induced loss is mechanically visible. We conducted a confirmation run at $N = 70$ to check whether the qualitative ordering survives the saturation that comes with the larger sample size. An 800-replicate-per-cell Monte Carlo run was executed at $N = 70$ across 12 cells (architecture $\in \{\text{MVN, mean moderation}\} \times c_{bm} \in \{0, 0.45\} \times t_{1/2} \in \{0, 0.5, 1.0\}$ weeks), all on the OL+BDC design, using 8-core `parallel::mclapply`. Wall-clock 483~s; convergence 100% across 9{,}600 fits.

At $c_{bm} = 0.45$:

Architecture	$t_{1/2} = 0$	$t_{1/2} = 0.5$	$t_{1/2} = 1.0$
MVN (B)	0.978	0.945	0.885
Mean moderation (A)	0.976	0.973	0.959

[15]

At $N = 70$ the B-versus-A ordering persists, though absolute magnitudes are compressed by ceiling saturation.

- Architecture B loses 9.3 points ($z = 7.4$); Architecture A 1.7 points ($z = 2.0$).
- Both sit within 2-5 points of the ceiling at $t_{1/2} = 0$, truncating the absorbable loss.
- The $N = 35$ results remain the canonical reference for the narrative magnitude.

To be completed by rgt.

Within-architecture losses from $t_{1/2} = 0$ to $t_{1/2} = 1$ are 9.3 percentage points (Architecture~B; two-proportion z -test on the paired cells, $z = 7.4$, $p < 10^{-12}$) and 1.7 percentage points (Architecture~A; $z = 2.0$, $p \approx 0.05$), where each z -statistic uses the formula $(\hat{\pi}_0 - \hat{\pi}_1) / \sqrt{\hat{\pi}_0(1 - \hat{\pi}_0)/n + \hat{\pi}_1(1 - \hat{\pi}_1)/n}$ with $n = 800$ replicates per cell. The qualitative ordering matches Section~3.1's $N = 35$ result. The absolute magnitudes are compressed because $N = 70$ at the prazosin-PTSD calibration places both architectures within 2-5 percentage points of the unit ceiling at $t_{1/2} = 0$, mechanically truncating the loss either architecture's curve can absorb. We note that reading the narrative loss-magnitude claim against the $N = 70$ robustness numbers (3% and 9% relative loss for Architecture~B versus 0.4% and 1.8% for Architecture~A) requires acknowledging this saturation effect; the $N = 35$ production results in Section~3.1 are the canonical reference for the narrative magnitude.

Type I error. Across all 18 null cells of the $N = 35$ production run, type~I error sat in $[0.010, 0.074]$, with mean 0.043 and median 0.042. We observe no architecture-specific inflation. Architecture~B's $t_{1/2} = 1.0$ OL+BDC null ($p = 0.010$, slightly conservative) is the smallest cell; the corresponding alternative ($p = 0.539$) is the most informative cell of the Section~3.1 narrative.

Estimator bias. Mean estimated $\hat{\beta}_{bm:D_{bc}}$ at $c_{bm} = 0.45$ ranged from -0.231 to -0.281 across the 18 alternative cells, with the $N = 35$ Architecture-A OL+BDC cell at $t_{1/2} = 1.0$ producing the largest absolute mean (-0.281); the Architecture~B cells maintained mean near -0.234 across all

$t_{1/2}$ values, consistent with the Architecture-B information channel being eroded through correlation decay rather than through coefficient drift. Empirical SE of $\hat{\beta}_{bm:D_{bc}}$ runs 0.05 to 0.10 across cells; the within-replicate model SE matches at scale, indicating nominal-95% Wald coverage holds across the grid.

4.4 3.4 Architecture C: combined-channel power results

We ran a factorial simulation on the Architecture-C parameter grid. The design factors are CO, Hybrid, and OL+BDC; $N = 35$ total subjects per cell (randomised equally across paths); $t_{1/2} \in \{0, 0.5, 1.0\}$ weeks; 1000 replicates per cell; A2 analysis specification. The parameter panel is $c_{bm,A} \times c_{bm,B} \in \{0, 0.22, 0.45\}^2$ (nine cells per design-carryover combination). By construction, the boundary cells ($c_{bm,A} = 0.45, c_{bm,B} = 0$) reduce the combined DGP to the pure Architecture-A functional form, and ($c_{bm,A} = 0, c_{bm,B} = 0.45$) reduce it to the pure Architecture-B form. These boundary cells confirm that the combined parameterisation recovers each limiting case algebraically. They are not, however, numerical checks against the power values of §3.1: the Architecture-C grid here uses $N = 35$ as the total sample size, whereas §3.1 uses $N = 35$ per randomisation path (effective $N = 70/140$). The boundary-cell power values are therefore lower than, and not directly comparable to, the matched pure-architecture cells of §3.1.

Table 1: Architecture C empirical power: 3×3 parameter panel. $N = 35$, carryover $t_{1/2} = 1.0$ week, A2 analysis specification, 1000 replicates per cell. Here $N = 35$ is the total sample size, unlike §3.1 where $N = 35$ is per path (effective $N = 70/140$), so these power values are not directly comparable to the §3.1 tables. Rows index $c_{bm,A}$ (mean-shift channel); columns index $c_{bm,B}$ (covariance channel). By construction, the boundary column $c_{bm,B} = 0$ corresponds to the pure Architecture A form; the boundary row $c_{bm,A} = 0$ corresponds to the pure Architecture B form.

$c_{bm,A}$ (mean channel)	$c_{bm,B}$ (covariance channel)		
	0.00	0.22	0.45
<i>Crossover (CO)</i>			
0.00	0.046	0.064	0.206
0.22	0.070	0.179	0.361
0.45	0.172	0.354	0.528
<i>Hybrid</i>			
0.00	0.034	0.083	0.273
0.22	0.108	0.299	0.621
0.45	0.389	0.642	0.870
<i>OL+BDC</i>			
0.00	0.017	0.049	0.148
0.22	0.086	0.164	0.362
0.45	0.276	0.473	0.688

Table~1 reports empirical power for Architecture-C across the full 3×3 parameter panel at the most demanding carryover condition ($t_{1/2} = 1.0$ week). Several features of the results merit comment.

First, the monotone structure is exact: power increases strictly with each channel parameter, holding the other fixed, across all three designs. This confirms that the two channels contribute independently and additively to detectability.

Second, interior cells exceed the probability-scale additive benchmark. At ($c_{bm,A} = 0.22, c_{bm,B} = 0.22$) in the Hybrid design, power is 0.299. The additive baseline on the probability scale – the

power expected if the two channels contributed independently – is $p(0.22, 0) + p(0, 0.22) - p(0, 0) = 0.108 + 0.083 - 0.034 = 0.157$. The observed value of 0.299 exceeds this benchmark by 0.142, indicating genuine super-additivity: the joint signal compounds rather than aggregates.

Third, the covariance channel ($c_{bm,B}$) and the mean-shift channel ($c_{bm,A}$) contribute comparably to carryover sensitivity. Comparing the no-carryover ($t_{1/2} = 0$) and one-week ($t_{1/2} = 1.0$) panels for the CO design, increasing $c_{bm,B}$ from 0 to 0.45 (fixing $c_{bm,A} = 0$) degrades power from 0.368 to 0.206 – a 44% relative loss. The corresponding loss along the $c_{bm,A}$ axis (fixing $c_{bm,B} = 0$) is from 0.352 to 0.172 – a 51% loss, in the same range. At the equal-weight interior cell ($c_{bm,A} = c_{bm,B} = 0.45$), the Hybrid design retains 0.870 power under one-week carryover, substantially exceeding either pure architecture alone at $c_{bm} = 0.45$ (Architecture~A: 0.389; Architecture~B: 0.273).

Fourth, design choice remains consequential under Architecture~C. At the full-signal, one-week-carryover cell, Hybrid (0.870) substantially outperforms OL+BDC (0.688) and especially CO (0.528). The design ranking established in Section~3.2 is preserved throughout the Architecture~C grid.

5 4. Which recipe matches which biology?

The choice of architecture is really a choice about *why* the biomarker predicts response, so picking one is a biological statement, not a coding preference.

[16]

Use Architecture A when the biomarker mechanically sets the size of the drug effect.

- The biomarker fixes effective dose or receptor occupancy, so it scales the effect proportionally.
- Doubling the biomarker roughly doubles the effect at every timepoint.
- Carryover keeps the same proportion in washout, so power is barely touched.

To be completed by rgt.

Use Architecture A when the biomarker sets the size of the drug effect. If a patient’s biomarker mechanically determines how much drug effect they get – for example, kidney function (eGFR) controlling how fast a drug is cleared, or receptor density controlling how strongly a drug acts – then doubling the biomarker roughly doubles the effect at every timepoint. Carryover does not break this: the leftover effect during washout keeps the same proportion, so power is barely touched.

Typical examples:

- **Drug clearance.** Kidney function (eGFR) sets how fast a renally-cleared drug leaves the body, so it directly scales the steady-state drug level and hence the effect.
- **Receptor density.** If a biomarker measures how many target receptors a patient has, more receptors means proportionally more drug effect.
- **Metabolizer genotype.** A "poor metabolizer" of a prodrug activates less of it, getting a proportionally smaller effect.

[17]

Use Architecture B when the biomarker is only a proxy for a hidden responder state.

- The biomarker imperfectly indexes the real cause, so identical biomarkers can respond differently.
- It predicts response only while the drug is acting.
- In washout the correlation fades and power drops sharply, as in Section 3.

To be completed by rgt.

Use Architecture B when the biomarker only flags who is likely to respond. Sometimes the biomarker is just a *proxy* for some hidden “responder” state. Two patients with the same biomarker can respond very differently, because the biomarker imperfectly indexes the real cause (a disease subtype, an inflammatory state, a genetic background). Here the biomarker predicts response *only while the drug is acting*; once the drug washes out, the correlation that carried the signal fades, and power drops sharply – exactly the Architecture B behaviour we saw in Section 3.

Typical examples:

- **Blood pressure in PTSD.** High resting blood pressure may flag a noradrenergic PTSD subtype likely to respond to prazosin, but it does not *cause* the response; it is a noisy marker of the real underlying state.
- **Inflammatory markers.** High CRP or IL-6 may predict response to anti-inflammatory augmentation of an antidepressant, but only on average – many high-CRP patients still do not respond.
- **Polygenic risk scores.** These predict treatment response while explaining only a fraction of the variance; the rest is unmeasured biology.

(A fully faithful model of this “hidden responder class” idea would be a finite mixture; Architecture B is a convenient correlation-level approximation, adequate for power calculations but not a claim about the true mechanism. Senn⁵ also cautions that apparent biomarker-by-treatment effects are sometimes just variance artifacts, so Architecture B’s signal may be smaller than assumed.)

[18]

The motivating prazosin/PTSD case is an Architecture B story, so its carry-over power loss is a clinically relevant warning.

- Biomarker is resting blood pressure, a proxy for noradrenergic tone, not a driver of drug levels.
- Designs with long off-drug periods may lose more power than an Architecture A simulation predicts.
- Architecture C is the right model if blood pressure also acts through the mean channel; it needs evidence for both parameters.

To be completed by rgt.

The motivating example. The pmsimstats software was built for a prazosin/PTSD trial^{3,6} where the biomarker is resting blood pressure and the outcome is a PTSD symptom score. Blood pressure is thought to be a *proxy* for noradrenergic tone, not a direct driver of drug levels – an Architecture B story. So the large carryover-driven power loss we find

under Architecture B is the clinically relevant warning: designs with long off-drug periods may lose more power to detect this biomarker than a simpler Architecture A simulation would suggest. Architecture C is the appropriate model if blood pressure turns out to act through the mean channel as well as the correlation channel; because it estimates two independent parameters, it calls for independent evidence on each (see Section 2.3).

6 5. Implications for Simulation Study Design

6.1 5.1 Choosing the appropriate architecture

The choice of DGP architecture should be guided by the biological mechanism hypothesized for the biomarker-treatment interaction:

Criterion	Architecture A	Architecture B	Architecture C
Biomarker role	Causal mediator	Statistical predictor	Both
Mechanism	Deterministic scaling	Probabilistic association	Both channels
Signal location	Mean structure	Covariance structure	Mean + covariance
Free parameters	1 (c_{bm})	1 (c_{bm})	2 ($c_{bm,A}$, $c_{bm,B}$)
Carryover impact	Modest (1–3%)	Design-dependent, up to 31%	Intermediate; scales with $c_{bm,B}$
Appropriate when	Biomarker determines dose or PK	Biomarker indexes a latent subtype	Channel uncertain; sensitivity use

[19]

When the mechanism is ambiguous, sweep Architecture C; use Architecture B as the conservative default.

- A mixing weight α spans the two pure-architecture extremes via $c_{bm,A} = (1 - \alpha)c$, $c_{bm,B} = \alpha c$.
- Architecture B is the safer default for design decisions.
- Its larger carryover sensitivity keeps power estimates from being optimistic.

To be completed by rgt.

If the biological mechanism is ambiguous, Architecture~C with a range of mixing weights $\alpha \in [0, 1]$ (where $c_{bm,A} = (1 - \alpha)c$ and $c_{bm,B} = \alpha c$ for fixed total weight c) provides a natural sensitivity framework spanning the two pure-architecture extremes. Architecture~B should remain the default conservative choice for trial design when the channel structure is unknown: its larger carryover sensitivity ensures that power estimates will not be systematically optimistic relative to either pure-architecture alternative.

6.2 5.2 Reporting requirements

Simulation studies evaluating biomarker-moderated treatment effects should explicitly state:

1. Whether the biomarker-treatment interaction is generated through mean moderation, differential correlation, or both channels simultaneously.
2. The biological rationale for the chosen mechanism and, for Architecture C, independent evidence supporting each channel.
3. Whether the carryover model (if present) operates on the interaction signal consistently with the chosen architecture.
4. Sensitivity analyses under alternative architectures, if the biological mechanism is uncertain.

[20]

These requirements specialise the ADEMP template; the architecture choice is its data-generating (D) element.

- ADEMP fixes Aims, Data-generating mechanisms, Estimands, Methods, and Performance measures.
- The DGP-architecture choice is the D element.
- Our reporting rules are ADEMP specialised to architecture-sensitive biomarker simulations.

To be completed by rgt.

The reporting framework above is consistent with the broader ADEMP template of Morris, White, and Crowther⁷, which prescribes that simulation studies fix five elements before production runs: the **Aims** of the study, the **Data-generating mechanisms**, the **Estimands** (the inferential targets), the **Methods** (analysis specifications compared), and the **Performance measures** (power, bias, coverage, type I error, and Monte Carlo standard errors of each). The DGP-architecture choice this paper concerns is the **D** element of an ADEMP specification; the reporting requirements above are the specialisation of the broader ADEMP framework to the architecture-sensitive class of biomarker-validation simulations.

The **primary estimand** for this paper’s simulations is the biomarker-by-treatment interaction coefficient $\beta_{bm:D}$ in the linear-mixed analysis model $Sx \sim bm + t + Dbc + bm:Dbc$ with random intercept and AR(1) residual correlation. The reported **performance measures** are: (i) power for $\beta_{bm:D}$ at $\alpha = 0.05$, with Monte Carlo standard error $\sqrt{\hat{\pi}(1 - \hat{\pi})/n_{\text{reps}}}$; (ii) type I error under $c_{bm} = 0$, with the same MCSE; (iii) bias of $\hat{\beta}_{bm:D}$ relative to the calibrated true value at each parameter cell; (iv) the empirical Monte Carlo standard error of $\hat{\beta}_{bm:D}$ across replicates; and (v) the convergence rate of the linear-mixed fit. Sample size, biomarker effect size, and carryover half-life vary across the parameter grid; the number of replicates per cell is calibrated to achieve MCSE below 0.02 for power at $\hat{\pi} = 0.5$, which corresponds to $n_{\text{reps}} \geq 625$. The simulations reported here use $n_{\text{reps}} = 1000$ throughout.

6.3 5.3 When the choice is uncertain

If the biological mechanism is ambiguous, as it often is in early biomarker development, we suggest investigators take the following steps:

1. Run the simulation under both architectures and report the range of power estimates.
2. Architecture B produces more conservative (lower) power estimates under carryover, making it the safer default for trial design decisions.

3. Design the trial to minimize carryover (adequate washout periods) to reduce sensitivity to the DGP architecture choice.

7 6. Can a smarter analysis recover the lost power?

None of the strategies below was tested by simulation here; the companion paper on carryover-mitigation strategies does that. This section just lists the options and their intuition.

[21]

Architecture B’s loss has two sources; only the variance-compression part is recoverable, by down-weighting low-signal off-drug points.

- Source 1: the treatment predictor’s variance shrinks (recoverable).
- Source 2: the biomarker-response correlation genuinely fades off-drug (lost information).
- Methods that down-weight or drop off-drug points can help with the first.

To be completed by rgt.

Under Architecture B the power loss has two sources: the treatment predictor’s variance shrinks, and the biomarker-response correlation genuinely fades off-drug. The second is lost information that no model can recover. But because off-drug points carry less signal, methods that *down-weight or drop* them can help.

The natural candidates are:

- **Use on-drug points only.** Every kept point has the full correlation, but you lose the off-drug reference and sample size; best for designs with a long on-drug phase, worst for crossover.
- **Weight by drug exposure.** Keep all points but weight each by its expected residual drug effect, $w_t = e^{-\lambda t_{sd}}$, so high-signal points count more. Easy to do with the `weights` argument in `nlme::lme`.
- **Within-subject contrast.** Reduce each patient to one number, (mean on-drug) minus (mean off-drug), then regress that on the biomarker: $\bar{Y}_{i,\text{on}} - \bar{Y}_{i,\text{off}} = \alpha + \gamma B_i + \epsilon_i$. Simple and robust, at some cost in efficiency.
- **Two-stage random slopes.** First estimate each patient’s drug effect u_{1i} from a random-slopes model, then regress those on the biomarker, $\hat{u}_{1i} = \delta_0 + \delta_1 B_i + \eta_i$.
- **Drop the most contaminated early off-drug points, or fix it in the design** with longer washouts or more on-drug time.

7.1 6.1 Qualitative comparison

Table~2 summarises the strategies along three dimensions: whether they discard observations, implementation complexity, and the expected direction of benefit based on the mechanistic account in Section 3.2. All entries in the ‘Hypothesised benefit’ column are qualitative; no simulation evidence is presented here, and each strategy should be evaluated empirically before use in trial planning.

Table 2: Qualitative summary of analysis strategies for recovering power under Architecture B carryover. All benefit assessments are mechanistic hypotheses, not simulation-supported estimates.

Strategy	Discards data	Complexity	Hypothesised benefit
On-drug only	Yes	Low	Eliminates correlation erosion; forfeits off-drug reference (best for OL/Hybrid)
Weighted	No	Low	Reduces erosion damage proportionally; depends on weight calibration
Within-subject contrast	Aggregates	Low	Robust to carryover misspecification; less efficient than LME
Two-stage random slopes	No	Moderate	Unknown; separates estimation stages; robustness unverified
Exclude early off-drug	Some	Low	Removes most contaminated observations; improves contrast clarity
Design modification	N/A	N/A	Addresses root cause; cost is longer trial or reduced carry-over estimation

A systematic simulation comparison of these strategies under Architecture B, analogous to the factorial comparison of analysis models in Section 3, would identify which approach best recovers the power lost to carryover for each trial design. We regard this comparison as important future work and note that the on-drug-only restriction and the weighted analysis are the most tractable candidates for a first empirical evaluation.

8 7. How this fits the wider literature

[22]

Architecture A is near-universal in the literature, so published power estimates may be optimistic for Architecture B biomarkers.

- Simon, Freidlin-Korn, Renfro-Sargent, and standard N-of-1 studies all put the effect in the mean.
- Even N-of-1 carryover work sharing Hendrickson’s authors uses Architecture A, so B looks like a one-off.
- Most published power figures may be optimistic for genuinely Architecture B biomarkers.

To be completed by rgt.

Almost every trial-simulation framework uses Architecture A: Simon¹, Freidlin and Korn², Renfro and Sargent⁸, and the standard N-of-1 simulation studies all put the biomarker effect in the mean. Even methodological work on N-of-1 carryover that shares authors with Hendrickson et al.³ uses Architecture A, which suggests Hendrickson’s differential-correlation choice was a one-off rather than a convention. The practical worry is that most published power estimates may be *optimistic* for biomarkers that really behave like Architecture B, because they ignore the extra carryover-driven loss we have demonstrated.

[23]

The distinction matters most in treatment-switching designs; when the mechanism is unknown, compute power both ways and use the lower number.

- Treatment-switching designs (crossover, N-of-1, basket) put patients through off-drug phases where B’s correlation decays.
- Parallel-group designs have no off-drug phase, so the architecture choice does not matter.
- This is a DGP question, separate from the analysis-model guidance of the PATH statement.

To be completed by rgt.

The distinction matters most exactly in the designs precision medicine cares about: any *treatment-switching* design (crossover, N-of-1 hybrid, basket/platform) puts patients through off-drug phases where the Architecture B correlation decays. Parallel-group designs, with no off-drug phase, are immune and give similar numbers under either architecture. The simple advice when the mechanism is unknown: run the power calculation *both ways* and use the lower (Architecture B) number for planning. This is a data-generating-process question, separate from – and not answered by – the analysis-model guidance in the PATH statement⁴.

To see where this bites in practice, it helps to sort designs by how much off-drug time each patient experiences:

- **Treatment-switching designs** (crossover, N-of-1 hybrid, basket and platform trials) deliberately cycle each patient on and off drug. Under Architecture B every washout erodes the correlation, so the interaction signal in later periods is weaker than in the first – the power loss of Section 3 applies directly.
- **Parallel-group enrichment designs** keep each patient on one treatment throughout. There is no off-drug phase for the correlation to decay in, so Architectures A and B give similar power. The architecture choice simply does not matter here.
- **Adaptive enrichment trials** decide mid-trial whom to keep enrolling, using interim power calculations written in the Architecture A idiom. If the truth is closer to Architecture B and patients pass through washouts, the real power is lower than the rule assumes, which can trigger premature stopping or an over-narrow eligible population.

The pattern is worth stating plainly: the architecture choice matters *most* in exactly the within-subject designs that precision medicine finds most attractive, and *least* in the parallel-group designs where it is often (wrongly) assumed to be harmless.

9 8. Conclusions

In conclusion, six points are to be emphasized.

1. The choice between direct mean moderation (Architecture A) and differential correlation (Architecture B) for generating biomarker-treatment interactions in clinical trial simulations has substantive consequences for power estimates under carryover.
2. Under Architecture A, carryover reduces power only modestly, by 1 to 3% in relative terms across the three designs examined, because the proportional biomarker-drug relationship is preserved during washout. Under Architecture B the loss is strongly design-dependent: it remains modest in the crossover and hybrid designs (3 to 5%) but reaches roughly 31% in the open-label OL+BDC design, because the differential-correlation signal erodes as the drug effect wanes.
3. The choice should be guided by the hypothesized biological mechanism: Architecture A for biomarkers that directly determine drug pharmacokinetics or pharmacodynamics; Architecture B for biomarkers that statistically predict treatment response through latent subtype membership.
4. For the prazosin-PTSD application with blood pressure as the predictive biomarker, Architecture B (MVN differential correlation) is the more appropriate model, and the resulting power sensitivity to carryover should inform trial design decisions regarding off-drug period duration.
5. The existing clinical trial simulation literature uses Architecture A almost exclusively. The Hendrickson et al. (2020) MVN approach is, to the best of our knowledge, unique in the N-of-1 trial context. This means that published power estimates for biomarker-moderated crossover and N-of-1 designs may be systematically optimistic for biomarkers that operate through a correlation-based (Architecture B) mechanism, because they do not account for the carryover sensitivity documented here.
6. Simulation studies for biomarker-moderated treatment effects should explicitly report their DGP architecture and consider sensitivity analyses under the alternative when the biological mechanism is uncertain.

10 Notes and reproducibility

The authors declare no conflicts of interest and received no specific funding. Code and data are in the pmsimstats-ng repository (drivers under `analysis/scripts/`, outputs under `analysis/data/`), with package versions pinned in `renv.lock`; the document renders under R~4.6.0.

One convention to note. In Sections 3.1 and 3.3, $N = 35$ is the number of subjects *per randomisation path*, and paths are concatenated before analysis, so the effective analysis sample is 70 (two-path CO and OL+BDC) or 140 (four-path Hybrid). The Architecture~C results in Section 3.4 instead use $N = 35$ as the *total*, so their power values are lower and are not directly comparable to the Section 3.1 tables.

11 References

1. Simon R. Clinical trial designs for evaluating the medical utility of prognostic and predictive biomarkers in oncology. *Personalized Medicine*. 2010;7(1):33-47.
2. Freidlin B, Korn EL. Biomarker enrichment strategies: Matching trial design to biomarker credentials. *Nature Reviews Clinical Oncology*. 2014;11(2):81-90.
3. Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: Statistical methods and practical considerations. *Frontiers in Digital Health*. 2020;2:13. doi:10.3389/fdgth.2020.00013
4. Kent DM, Paulus JK, Klaveren D van, et al. The predictive approaches to treatment effect heterogeneity (PATH) statement. *Annals of Internal Medicine*. 2020;172(1):35-45.
5. Senn S. Mastering variation: Variance components and personalised medicine. *Statistics in Medicine*. 2016;35(7):966-977. doi:10.1002/sim.6739
6. Raskind MA, Peterson K, Williams T, et al. A trial of prazosin for combat trauma PTSD with nightmares in active-duty soldiers returned from iraq and afghanistan. *American Journal of Psychiatry*. 2013;170(9):1003-1010.
7. Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Statistics in Medicine*. 2019;38(11):2074-2102. doi:10.1002/sim.8086
8. Renfro LA, Sargent DJ. Statistical controversies in clinical research: Basket trials, umbrella trials, and other master protocols. *Annals of Oncology*. 2017;28(1):34-43.