

1 Two Architectures for Simulating Biomarker-Treatment
2 Interactions: Implications for Statistical Power Under
3 Carryover

4 pmsimstats team

5 2026-06-10

6 **Contents**

7	1 Abstract	2
8	2 1. Introduction	3
9	3 2. Two DGP Architectures	6
10	3.1 2.1 Architecture A: Direct Mean Moderation	7
11	3.2 2.2 Architecture B: Differential Correlation (MVN)	8
12	3.3 2.3 Architecture C: Combined Mean Moderation and Differential	
13	Correlation	9
14	3.4 2.4 Mathematical Comparison	11
15	4 3. Consequences for Statistical Power Under Carryover	13
16	4.1 3.1 Empirical demonstration	13
17	4.2 3.2 Why the architectures diverge	17
18	4.3 3.3 Robustness check at $N = 70$	18
19	4.4 3.4 Architecture C: combined-channel power results	20
20	5 4. Biological Assumptions	21
21	5.1 4.1 When Architecture A is appropriate	22
22	5.2 4.2 When Architecture B is appropriate	23
23	5.3 4.3 The prazosin-PTSD case	26
24	5.4 4.4 When Architecture C is appropriate	28
25	6 5. Implications for Simulation Study Design	29
26	6.1 5.1 Choosing the appropriate architecture	29
27	6.2 5.2 Reporting requirements	30
28	6.3 5.3 When the choice is uncertain	32
29	7 6. Alternative Analysis Strategies Under Architecture B	32
30	7.1 6.1 Restrict analysis to on-drug observations	32
31	7.2 6.2 Weighted analysis	33

32	7.3	6.3 Within-subject contrast	34
33	7.4	6.4 Two-stage random slopes	34
34	7.5	6.5 Exclusion of contaminated observations	35
35	7.6	6.6 Design-level solutions	36
36	7.7	6.7 Comparative evaluation	37
37	8	7. Relationship to the Broader Literature	37
38	8.1	7.1 Treatment effect heterogeneity	37
39	8.2	7.2 Prevalence of each architecture	39
40	8.3	7.3 Crossover and N-of-1 trial methodology	40
41	8.4	7.4 Precision medicine trial design	41
42	9	8. Conclusions	44
43	10	Conflict of interest	45
44	11	Funding	45
45	12	Data availability statement	45
46	13	Reproducibility	45
47	14	References	46
48	15	References	47

49 1 Abstract

50 **Background.** Simulation studies of predictive biomarker-treatment interactions
51 in N-of-1 clinical trials require a data-generating process (DGP) that specifies how
52 the biomarker moderates treatment response. Two fundamentally different DGP
53 architectures are used in practice, and the choice between them determines the
54 extent to which carryover effects reduce statistical power to detect the interaction.
55 The published simulation literature does not distinguish them.

56 **Methods.** We formalise two architectures: direct mean moderation (Architec-
57 ture~A), under which the biomarker scales the treatment effect in the popula-
58 tion mean structure; and differential correlation (Architecture~B, the multivari-
59 ate normal approach), under which the interaction emerges from treatment-state-
60 dependent correlation between biomarker and response in the joint distribution.
61 Power consequences under both architectures are quantified by Monte Carlo sim-
62 ulation on an aggregated N-of-1 design framework with continuous-time AR(1)
63 residual correlation in the analysis model. The biomarker, treatment phase, and
64 carryover decay are parameterised to the prazosin/PTSD reference application.

65 **Results.** Under Architecture~A, power is largely preserved under carryover be-
66 cause the proportional relationship between biomarker and drug response is main-
67 tained at every timepoint. Under Architecture~B, carryover erodes the differen-

68 tial correlation signal and reduces power. At the prazosin-calibrated reference cell
69 ($N = 35$, $c_{bm} = 0.45$, OL+BDC design) across $t_{1/2} \in \{0, 0.5, 1.0\}$ weeks, Architec-
70 ture~B power drops from 0.777 to 0.539 (a 30.6% relative loss; $z = 7.64$, $p < 10^{-13}$
71 on the B-versus-A contrast), while Architecture~A drops only from 0.751 to 0.730
72 (a 2.8% relative loss). The B-versus-A absolute power-loss gap of 21.7 percentage
73 points at OL+BDC, $t_{1/2} = 1.0$ is approximately 11 times the matched Archi-
74 tecture~A loss; the Hybrid design produces a smaller but still significant gap
75 ($z = 3.96$, $p < 10^{-4}$); the traditional crossover design produces no detectable gap
76 because its within-subject AR(1) structure dominates the carryover-mediated cor-
77 relation channel that Architecture~B uses. The 36-cell production factorial (2
78 architectures $\times$ 3 designs $\times$ 3 $t_{1/2}$ $\times$ 2 c_{bm} levels) was run at 1{,}000 replicates
79 per cell with 100% convergence; MCSE on cell power is 0.013-0.016. Type~I error
80 sits in [0.010, 0.074], with no architecture-specific inflation.

81 **Conclusions.** The two architectures encode different biological mechanisms, het-
82 erogeneity at the mean structure versus heterogeneity at the variance structure,
83 and the choice between them is a substantive scientific question rather than a
84 parameterisation convenience. We recommend explicit declaration of the DGP
85 architecture in any biomarker-validation simulation study and provide guidance
86 for selecting the architecture appropriate to a given clinical context.

87

88 2 1. Introduction

89 **The central question is biomarker-treatment interaction; simulation stud-**
90 **ies need a DGP that encodes it.**

- 91 • Does a baseline biomarker predict differential treatment response, i.e., does
92 it moderate the treatment effect?
- 93 • Simulation studies are the principal power evaluation tool for proposed trial
94 designs.
- 95 • These simulations require a DGP that explicitly encodes the biomarker-
96 treatment interaction mechanism.

97 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
98 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
99 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
100 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
101 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
102 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
103 mollit anim id est laborum.

104 *A central question in precision medicine trial design is whether a baseline*
105 *biomarker predicts differential treatment response, that is, whether the treat-*
106 *ment effect is moderated by the biomarker value. Simulation studies are the*
107 *principal tool available for evaluating the statistical power of proposed trial*

108 *designs to detect such interactions, and these simulations necessarily require*
109 *a data-generating process (DGP) that encodes the biomarker-treatment inter-*
110 *action mechanism.*

111 **Every simulation framework must choose how to generate treatment ef-**
112 **fect heterogeneity, and that choice has consequences.**

- 113 • HTE literature (Rothwell 2005; Kent et al. 2020) distinguishes qualitative
114 interactions (direction reverses) from quantitative ones (magnitude varies).
- 115 • Any simulation framework must decide how to encode this heterogeneity
116 in the DGP.
- 117 • That decision is substantive, not a mere parameterization detail.

118 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
119 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
120 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
121 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
122 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
123 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
124 mollit anim id est laborum.

125 *The statistical literature on treatment effect heterogeneity, developed notably*
126 *by Rothwell (2005) and synthesized more recently by Kent et al. (2020),*
127 *distinguishes between qualitative interactions, in which the direction of the*
128 *treatment effect reverses across biomarker strata, and quantitative interac-*
129 *tions, in which the magnitude varies but the direction remains consistent.*
130 *As we shall argue, any simulation framework intended to evaluate power for*
131 *biomarker-moderated effects must choose how to generate this heterogeneity,*
132 *and that choice is not merely a technical convenience.*

133 **Architecture A (direct mean moderation) is the near-universal standard;**
134 **the Hendrickson (2020) MVN approach is apparently unique.**

- 135 • Every major simulation framework uses direct mean moderation: Simon
136 (2010) for stratified designs; Freidlin and Korn (2014) for adaptive enrich-
137 ment; Renfro and Sargent (2017) for basket/umbrella trials; all N-of-1
138 frameworks.
- 139 • Wang and Schork (2019) and Schork (2022) also use Architecture A despite
140 sharing authorship with Hendrickson et al. (2020), suggesting the MVN
141 parameterization is a paper-specific choice, not a lab convention.
- 142 • To the best of our knowledge, no other clinical trial simulation framework
143 uses differential correlation (MVN) to encode the biomarker-treatment in-
144 teraction.

145 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
146 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
147 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
148 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
149 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
150 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
151 mollit anim id est laborum.

152 *The dominant approach in the clinical trial simulation literature is direct*

mean moderation, in which the biomarker enters the outcome model as an interaction term that explicitly scales the treatment effect. Simon (2010) employs this approach throughout the simulation frameworks he develops for biomarker-stratified designs; Freidlin and Korn (2014) adopt it for adaptive enrichment trials; Renfro and Sargent (2017) use it throughout their treatment of basket and umbrella trials; and it underlies the platform trial literature without exception. Recent simulation work on biomarker-by-treatment interaction estimation in randomized trials with prognostic covariates, specifically the study by Haller and Ulm (2018), likewise adopts the standard regression specification with explicit interaction terms. Among N-of-1 simulation studies, Zucker et al. (1997), Araujo et al. (2016), and Duan et al. (2013) similarly use hierarchical random-effects models in which the biomarker predicts the individual treatment effect through the mean structure. Notably, Wang and Schork (2019) and Schork (2022), whose methodological work addresses power and design for N-of-1 trials with serial correlation and carry-over directly, also specify the treatment effect through a linear mean structure, which suggests that the multivariate-normal differential-correlation parameterization of Hendrickson et al. (2020) is not a convention of any particular research program but rather a paper-specific design choice. To the best of our knowledge, no other clinical trial simulation framework generates the biomarker-treatment interaction through differential correlation in a multivariate normal joint distribution rather than through direct mean moderation.

The two architectures were discovered by accident during pmsimstats development, when two implementations of the same simulation diverged under carryover.

- The original code uses Architecture B (MVN); a tidyverse reimplemention inadvertently used Architecture A.
- Both produced plausible power estimates, but they reached qualitatively different conclusions about carryover effects.
- Tracing the discrepancy to its root revealed the architectural distinction, which has not been explicitly discussed in the simulation methodology literature.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

We encountered both approaches in the course of developing the pmsimstats simulation framework for N-of-1 clinical trials (Hendrickson et al. 2020). The original publication code uses the MVN differential correlation approach, an unusual choice that, as we shall demonstrate here, has substantive consequences for power estimation under carryover. A comprehensive audit of the codebase subsequently revealed that a parallel tidyverse reimplemention of the same simulation had inadvertently adopted the standard direct mean moderation approach. Both implementations produced plausible power estimates, yet they yielded qualitatively different conclusions about how carryover effects

influence statistical power. Tracing this discrepancy to its root cause revealed the architectural difference described in this paper, which has not, to the best of our knowledge, been discussed explicitly in the simulation methodology literature.

This paper is the first systematic comparison of the two DGP architectures; we formalize, simulate, and give guidance.

- We formalize the distinction between Architecture A (mean moderation) and Architecture B (differential correlation).
- We demonstrate power consequences through Monte Carlo simulation.
- We offer guidance for investigators on how to choose the appropriate architecture.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

We present here what we believe to be the first systematic comparison of these two DGP architectures for biomarker-treatment interaction power estimation under carryover. We shall formalize the distinction between them, demonstrate its consequences through Monte Carlo simulation, and offer guidance for investigators designing simulation studies for predictive biomarker-moderated treatment effects.

Related documents in this series. The mathematical derivation, code-level implementation details, four implementation attempts (three failed and one correct), and the empirical evaluation of the original Hendrickson (2020) parameter grid under Architecture~A are documented in `19-mean-moderation-implementation-notes.tex` (with a brief markdown summary in `19-mean-moderation-implementation-notes.md`). A condensed empirical-results summary suitable for collaborator presentation is provided in `21-architecture-comparison-summary.md`.

233

3 2. Two DGP Architectures

Notation. Throughout this paper, B_i denotes the biomarker value of participant i , D_{it} the binary on-drug indicator at timepoint t (1 on drug, 0 off drug), and BR_{it} the biological response component of the outcome. These three symbols describe the data-generating-process layer. The analysis-model layer uses a related but distinct continuous drug exposure variable, Dbc , defined in Section 3.1. Code-style identifiers in typewriter font (`bm`, `Dbc`, `BR`) refer to specific variables in the

241 pmsimstats codebase or to R formula syntax; mathematical symbols are used
242 otherwise.

243 3.1 2.1 Architecture A: Direct Mean Moderation

244 In this approach the biomarker enters the mean structure of the outcome directly.
245 The treatment effect for participant i at timepoint t is:

$$Y_{it} = \mu_0 + \beta_1 \cdot D_{it} \cdot (1 + \beta_{bm} \cdot B_i) + \epsilon_{it}$$

246 where D_{it} is the treatment indicator (or a continuous measure of drug expo-
247 sure), B_i is the participant's biomarker value (typically centered), and β_{bm} is
248 the biomarker moderation parameter. The interaction is explicit in the mean:
249 participants with higher biomarker values have proportionally larger treatment
250 effects.

251 Key properties:

- 252 • The interaction is **deterministic** given the biomarker value and treatment
253 status.
- 254 • The biomarker moderation β_{bm} is a **population-level parameter** that
255 applies uniformly to all participants.
- 256 • The signal for the interaction is in the **first moment** (conditional mean) of
257 the outcome distribution.
- 258 • Adding noise, random effects, or carryover to D_{it} scales the entire treatment
259 effect (including its biomarker-dependent component) proportionally. The
260 **ratio** of drug effect between high-biomarker and low-biomarker participants
261 is preserved.

262 **The multiplicative and additive forms of Architecture A are interchange-**
263 **able for power calculations.**

- 264 • The multiplicative form maps directly to the standard additive regression
265 (Kent et al. 2020 effect modeling equation).
- 266 • $\beta_3 = \beta_1 \cdot \beta_{bm}$; the additive form additionally allows a prognostic main
267 effect.
- 268 • Both forms fall within Architecture A and are interchangeable for power
269 calculations on the interaction coefficient.

270 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
271 mod tempor incidunt ut labore et dolore magna aliqua. Ut enim ad
272 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
273 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
274 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
275 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
276 mollit anim id est laborum.

We note that the multiplicative form above maps directly onto the standard additive regression form for predictive biomarker effects used in the trial methodology literature, specifically the effect modeling equation of Kent et al. (2020):

$$Y_{it} = \beta_0 + \beta_1 D_{it} + \beta_2 B_i + \beta_3 D_{it} B_i + \epsilon_{it}$$

The interaction coefficient β_3 in the additive form corresponds to the product $\beta_1 \cdot \beta_{bm}$ in the multiplicative form. The additive form additionally accommodates a prognostic main effect β_2 for the biomarker, which the multiplicative form sets to zero by construction. For the purpose of power calculations on the interaction coefficient, the two forms are interchangeable up to this prognostic-effect distinction, and both fall within Architecture A.

This architecture is used in the pedagogical simulation (`implementations/simple/simulation.R`) and in the exploratory carryover analyses (`simulation_carryover_spectrum.R`) of the pmsimstats project.

3.2 2.2 Architecture B: Differential Correlation (MVN)

In this approach the biomarker and the biological response (BR) component are jointly drawn from a multivariate normal distribution with treatment-state-dependent correlation:

$$\begin{pmatrix} B_i \\ BR_{it} \end{pmatrix} \sim \text{MVN} \left(\begin{pmatrix} \mu_B \\ \mu_{BR}(t, D_{it}) \end{pmatrix}, \Sigma(D_{it}) \right)$$

where the covariance matrix Σ has the following treatment-state-dependent BM-BR correlation:

$$\text{Cor}(B_i, BR_{it}) = \begin{cases} c_{bm} & \text{if } D_{it} = 1 \text{ (on drug)} \\ c_{bm} \cdot e^{-\lambda \cdot t_{sd}} & \text{if } D_{it} = 0 \text{ (off drug)} \end{cases}$$

where t_{sd} is time since drug discontinuation. The exponential-decay form is the default; the limiting cases $\lambda = 0$ (no decay; correlation persists indefinitely during the off-drug phase) and $\lambda \rightarrow \infty$ (immediate decay; correlation drops to zero the instant the drug is discontinued) recover the two no-carryover model variants.

In Architecture B the interaction is in the covariance structure, not the mean; it emerges from the conditional distribution given the biomarker value.

- The population mean of BR_{it} is the same for all participants; no mean-structure interaction exists.
- Given biomarker value b , the conditional expectation of BR scales with c_{bm} , creating an interaction-like structure.
- Correlation is c_{bm} on drug and decays as $c_{bm}e^{-\lambda t_{sd}}$ off drug; moderation attenuates during washout.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad

311 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
312 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
313 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
314 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
315 mollit anim id est laborum.

316 *The interaction is not in the mean structure: the population mean of BR_{it}*
317 *is the same for all participants with the same treatment history. Instead it*
318 *emerges from the **conditional distribution**: given a participant's biomarker*
319 *value $B_i = b$, the conditional expectation of their biological response is:*

$$E[BR_{it}|B_i = b, D_{it} = 1] = \mu_{BR}(t) + c_{bm} \cdot \frac{\sigma_{BR}}{\sigma_B} \cdot (b - \mu_B)$$

320 *This conditional expectation has an interaction-like structure: the biomarker*
321 *value b modulates the expected BR. The effective correlation is c_{bm} on drug*
322 *and $c_{bm} \cdot e^{-\lambda t_{sd}}$ during the washout; the moderation therefore persists with*
323 *attenuated strength during off-drug periods and vanishes only in the limit of*
324 *complete decay.*

325 Key properties:

- 326 • The interaction is **probabilistic**: it manifests as a tendency, not a deter-
327 ministic scaling.
- 328 • The biomarker moderation c_{bm} is a **correlation parameter** that governs
329 the strength of the association.
- 330 • The signal for the interaction is in the **second moment** (covariance struc-
331 ture) of the joint distribution.
- 332 • Carryover effects in the DGP inflate the BR means during off-drug periods,
333 making them resemble on-drug means. This reduces the **differential** in the
334 correlation structure between treatment states, weakening the signal that
335 the analysis model detects.

336 This architecture is used in the Hendrickson et al. (2020) publication code and
337 the audited R/ package in pmsimstats-ng.

338 2.3 Architecture C: Combined Mean Moderation and Differ- 339 ential Correlation

340 **Architecture C activates both interaction channels simultaneously, with**
341 **independently calibrated weights.**

- 342 • The MVN covariance encodes BM-BR differential correlation with weight
343 c_{bm_b} (the Architecture B channel).
- 344 • After the MVN draw, BR is additively shifted on-drug by $c_{bm,A} \cdot z_B \cdot \sigma_{BR}$
345 (the Architecture A channel).
- 346 • Special cases: $c_{bm,B} = 0$ recovers pure Architecture A; $c_{comb,A} = 0$ recov-
347 ers pure Architecture B.

348 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
349 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
350 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
351 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
352 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
353 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
354 mollit anim id est laborum.

355 *Architecture C (combined) applies both interaction mechanisms in sequence.*
356 *First, the MVN covariance matrix is constructed with BM-BR differential*
357 *correlation weighted by `c.bm_b`, exactly as under Architecture B. Second, after*
358 *the MVN draw, the BR component is additively shifted at on-drug timepoints*
359 *by*

$$\Delta BR_{it} = c_{bm,A} \cdot z_{B,i} \cdot \sigma_{BR} \cdot \mathbf{1}[on-drug]$$

360 *where $z_{B,i} = (B_i - \mu_B)/\sigma_B$ is the standardized biomarker for participant i ,*
361 *exactly as under Architecture A.*

362 *The two parameters `c.bm_a` and `c.bm_b` are independent. Setting $c_{bm,B} = 0$*
363 *removes the covariance channel and recovers pure Architecture A; setting*
364 *$c_{bm,A} = 0$ removes the mean-shift channel and recovers pure Architecture B.*
365 *The boundary cases are enforced algebraically rather than by convention, so*
366 *the special cases serve as internal validation gates for any simulation study*
367 *using Architecture C.*

368 *We note that when both parameters are set to their full values ($c_{bm,A} =$*
369 *$c_{bm,B} = 0.45$), the total interaction effect on the BR mean is approximately*
370 *double that of either pure architecture at $c_{bm} = 0.45$. A researcher wishing to*
371 *hold total effect size constant across architectures may instead parameterize*
372 *as $c_{bm,A} = (1-\alpha) \cdot c$ and $c_{bm,B} = \alpha \cdot c$ for a mixing weight $\alpha \in [0, 1]$ and total*
373 *weight c . We do not adopt this constraint in the simulations below, preferring*
374 *the factorial parameterization that spans the full $(c_{bm,A}, c_{bm,B})$ space.*

375 **The biological rationale for Architecture C is that the two channels are**
376 **mechanistically distinct and may both operate.**

- 377 • The mean-shift channel (A) reflects pharmacodynamic scaling: the drug's
378 biological effect is larger in high-biomarker participants.
- 379 • The covariance channel (B) reflects differential co-regulation: on-drug pe-
380 riods couple biomarker expression to treatment response in a way that
381 washes out between periods.
- 382 • Neither channel implies the other; a drug could activate one, both, or
383 neither.

384 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
385 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
386 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
387 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
388 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
389 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
390 mollit anim id est laborum.

The two channels are mechanistically distinct and need not scale in proportion. The mean-shift channel (Architecture A) reflects a pharmacodynamic relationship: the drug's biological effect on BR is amplified in participants whose biomarker is above the population mean, because the biomarker governs the effective dose or receptor occupancy. This mechanism operates at the level of the mean trajectory for each participant.

The covariance channel (Architecture B) reflects a different process: on-drug periods impose a co-regulatory coupling between biomarker expression and treatment response, perhaps because the drug activates a pathway whose output is gated by the biomarker signal. This coupling is absent during off-drug washout, producing the differential correlation structure that Architecture B encodes in the MVN covariance matrix.

A given drug could activate one, both, or neither channel. The simulation practitioner who is uncertain which channel predominates should regard Architecture C with a range of α values as the natural sensitivity framework.

3.4 2.4 Mathematical Comparison

Under Architecture A, carryover shrinks the signal amplitude but preserves the ratio between biomarker groups, so identifiability is maintained.

- The expected outcome difference between two participants scales linearly with D .
- Under carryover, D is replaced by a decayed value and the absolute magnitude shrinks.
- But the ratio of drug effect between high- and low-biomarker participants is preserved at every exposure level; the signal contracts but remains identifiable.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

We begin the comparison by considering the expected outcome difference between two participants with biomarker values b_1 and b_2 at the same drug exposure level D under Architecture A:

$$E[Y \mid b_1, D] - E[Y \mid b_2, D] = \beta_1 \cdot \beta_{bm} \cdot (b_1 - b_2) \cdot D$$

The absolute magnitude of this difference scales linearly with D and therefore shrinks during off-drug periods when D is replaced by an exposure-decayed value $D \cdot \text{carryover}$. We note, however, that the **ratio** of the drug effect between high-biomarker and low-biomarker participants is preserved at every level of exposure. The biomarker-by-treatment signal contracts in amplitude but does not lose its identifiability under carryover.

432 **Under Architecture B, the biomarker-dependent component of the off-**
 433 **drug response decays exponentially, eventually eliminating the differential**
 434 **signal.**

- 435 • Off-drug, the conditional expectation includes a biomarker-dependent term
 436 scaled by $c_{bm}e^{-\lambda t_{sd}}$.
- 437 • As t_{sd} grows, that term shrinks toward zero.
- 438 • In the limit, all participants converge to the same off-drug mean regardless
 439 of biomarker value; the interaction signal vanishes.

440 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
 441 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
 442 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
 443 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
 444 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
 445 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
 446 mollit anim id est laborum.

447 *Under Architecture B, the analogous quantity is:*

$$E[BR|b_1, D = 1] - E[BR|b_2, D = 1] = c_{bm} \cdot \frac{\sigma_{BR}}{\sigma_B} \cdot (b_1 - b_2)$$

448 *During off-drug periods with carryover:*

$$E[BR|b, D = 0, t_{sd}] = \mu_{BR,off}(t_{sd}) + c_{bm} \cdot e^{-\lambda t_{sd}} \cdot \frac{\sigma_{BR}}{\sigma_B} \cdot (b - \mu_B)$$

449 *Here the biomarker-dependent component decays exponentially with time off*
 450 *drug. As t_{sd} increases, the off-drug conditional expectation converges to*
 451 *$\mu_{BR,off}$ for all biomarker values, eliminating the differential signal entirely.*

452 **Under Architecture C, carryover operates on both channels simultane-**
 453 **ously; the power loss is bounded below by Architecture B and above by**
 454 **Architecture A.**

- 455 • The mean-shift channel degrades as under Architecture A: amplitude
 456 shrinks but identifiability is preserved at every exposure level.
- 457 • The covariance channel degrades as under Architecture B: off-drug BM-BR
 458 correlation decays toward zero, eroding second-moment signal.
- 459 • The net power loss lies between the two pure-architecture extremes,
 460 weighted by the relative magnitudes of $c_{bm,A}$ and $c_{bm,B}$.

461 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
 462 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
 463 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
 464 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
 465 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
 466 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
 467 mollit anim id est laborum.

468 *Under Architecture C both channels are active. The analysis model's inter-*

action coefficient depends on the covariance between B_i and D_{bc} -weighted outcomes. The mean-shift channel contributes a term that degrades exactly as under Architecture A: the D_{bc} contrast is compressed during carryover but the ratio of the drug effect between biomarker groups is preserved at every exposure level. The covariance channel contributes a term that degrades as under Architecture B: the BM-BR correlation generating the second-moment signal decays exponentially off-drug. The two channels are additive at the level of the total biomarker-by-drug covariance, so the net power loss under Architecture C lies strictly between the pure-architecture extremes when $c_{bm,A}, c_{bm,B} > 0$. As either parameter approaches zero the corresponding boundary case is recovered.

480

4 3. Consequences for Statistical Power Under Carryover

4.1 3.1 Empirical demonstration

The two DGPs are identical except for how they encode the BM-BR linkage; any power divergence is attributable solely to that difference.

- All non-biologic components (time-varying, placebo-belief, residual noise) are generated identically in both architectures.
- The only difference: Architecture B encodes BM-BR interaction in Σ ; Architecture A adds a post-draw mean shift to on-drug BR.
- The same nlme::lme analysis model with corCAR1 is used for both, so divergence in results is attributable to the DGP alone.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

We ran identical simulation configurations under both architectures using the pmsimstats-ng framework, matching the DGP parameters as closely as possible between the two approaches. We note that the non-biologic response components (time-varying response, placebo-belief accumulation, and residual noise) are generated from the same multivariate normal draw with identical means, variances, within-factor $AR(1)$ autocorrelation, and cross-factor correlations in both architectures. The two DGPs differ only in how the biomarker-to-biological-response (BM-BR) linkage is encoded: Architecture B embeds it in the off-diagonal c_{bm} entries of Σ , whereas Architecture A omits those entries and adds a post-draw shift $c_{bm} \cdot B_i^* \cdot \sigma_{BR}$ to on-drug BR columns (see 19-mean-moderation-implementation-notes.tex for the

scaling derivation). The divergent carryover behavior reported below is therefore attributable to the BM-BR parameterization, not to any difference in how the placebo, time-varying, or residual-noise components are generated. The analysis model was the same in both cases: an `nlme::lme` fit with a `corCAR1` continuous-time autocorrelation structure and a single biomarker-by-drug interaction term, written `bm:Dbc` in R formula syntax.

Dbc is a continuous drug exposure variable that captures exponential washout; the `bm:Dbc` interaction tests whether biomarker moderates that exposure-decayed effect.

- `Dbc = 1` on drug; $= \exp(-\lambda t_{sd})$ off drug, matching the DGP decay rate.
- Unlike the binary D_{it} , `Dbc` smoothly captures washout shrinkage.
- The test of `bm:Dbc` is the test of whether the biomarker moderates the (decay-adjusted) drug effect.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The analysis-model drug exposure variable `Dbc` is the continuous analogue of the binary DGP indicator D_{it} . It equals 1 during on-drug periods and $\exp(-\lambda t_{sd})$ during off-drug periods, where t_{sd} is time since drug discontinuation and λ is the analysis-model carryover decay rate (typically derived from the drug's pharmacokinetic half-life as $\lambda = \ln 2/t_{1/2}$, in which case it matches the DGP decay rate of Section 2.2 by construction). Whereas D_{it} flips abruptly between 0 and 1, `Dbc` captures the smooth shrinkage of residual drug effect during washout. The interaction coefficient on `bm:Dbc` is therefore the test of whether the biomarker moderates the exposure-decayed drug effect.

We follow the ADEMP framework of Morris, White, and Crowther (2019): Aims, Data-generating mechanisms, Estimands, Methods, Performance measures.

Setup:

- Trial designs: CO (traditional crossover), Hybrid (N-of-1 Hybrid), OL+BDC (open-label with blinded discontinuation)
- Sample size: $N = 35$
- Biomarker correlation / moderation: $c_{bm} = 0.45$ / $\beta_{bm} = 0.45$
- DGP carryover half-life: $t_{1/2} \in \{0, 0.5, 1.0\}$ weeks
- Analysis: matched `Dbc` with same half-life
- Replicates: 1,000 per cell (Monte Carlo standard error on a binomial power estimate near 0.5 is approximately $\sqrt{0.5 \cdot 0.5 / 1,000} \approx 0.016$)
- Implementation: 8-core `parallel::mclapply`; full factorial 2 architectures $\times$ 3 designs $\times$ 3 $t_{1/2}$ $\times$ 2 c_{bm} levels = 36 cells, 36,000 fits, 100% convergence

$N = 35$ was chosen because it keeps power off the ceiling, making any carryover-induced loss visible; type I error is in the expected range.

- $N = 35$ is the smallest prazosin-PTSD-calibrated sample size in the literature.
- Neither architecture saturates power at zero carryover, so losses are mechanically observable.
- Type I error across 18 null cells sat in $[0.010, 0.074]$, consistent with the conservative-to-nominal range expected for this specification.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The $N = 35$ choice is the smallest of the prazosin-PTSD-calibrated sample sizes used in the published methodological literature; it is a regime in which neither architecture saturates power at the no-carryover cells, so any carryover-induced erosion is mechanically visible. Type I error under the $c_{bm} = 0$ null cells sat in $[0.010, 0.074]$ across the 18 null cells, consistent with the conservative-to-nominal regime expected for the continuous D_{bc} specification at this sample size.

Results under Architecture A (direct mean moderation):

Design	$t_{1/2} = 0$	$t_{1/2} = 0.5$	$t_{1/2} = 1.0$
CO	0.739	0.739	0.725
Hybrid	0.992	0.987	0.978
OL+BDC	0.751	0.748	0.730

Architecture A shows modest power loss under carryover (1–3%), well below the broader literature’s 8–13% range.

- Relative losses: 1.9% (CO), 1.4% (Hybrid), 2.8% (OL+BDC).
- All three fall below the 8–13% range reported in the broader N-of-1 simulation literature.
- Qualitative finding: modest erosion; power essentially intact under carryover.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Power declines modestly under Architecture A: relative loss from $t_{1/2} = 0$ to

$t_{1/2} = 1$ is 1.9% (CO), 1.4% (Hybrid), and 2.8% (OL+BDC). We note that all three values fall well below the 8-13% relative-loss range that the broader N-of-1 simulation literature reports for direct-mean-moderation DGPs, though they reproduce its qualitative finding, namely a modest erosion that leaves power essentially intact.

Results under Architecture B (differential correlation, MVN):

Design	$t_{1/2} = 0$	$t_{1/2} = 0.5$	$t_{1/2} = 1.0$
CO	0.745	0.754	0.723
Hybrid	0.995	0.980	0.944
OL+BDC	0.777	0.694	0.539

Architecture B shows sharply larger power loss under carryover, especially OL+BDC (30.6%); the B-vs-A gap is significant for OL+BDC and Hybrid but not CO.

- Relative losses: 3.0% (CO), 5.1% (Hybrid), 30.6% (OL+BDC); the OL+BDC loss is $\approx 11\times$ the matched Architecture A loss.
- B-vs-A gap: $z = 7.64$ ($p < 10^{-13}$) at OL+BDC; $z = 3.96$ ($p < 10^{-4}$) at Hybrid.
- CO gap not significant ($z = 0.29$): the CO design's within-subject AR(1) structure dominates the correlation-decay channel that Architecture B depends on.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Power declines more sharply under Architecture B, and most markedly so in the OL+BDC design: relative loss from $t_{1/2} = 0$ to $t_{1/2} = 1$ is 3.0% (CO; not significantly different from Architecture A's CO loss), 5.1% (Hybrid), and **30.6% (OL+BDC)**. The OL+BDC loss is approximately 11 times the matched Architecture-A OL+BDC loss (2.8%). It falls below the 40-60% relative-loss range that the broader N-of-1 simulation literature reports for the MVN differential-correlation DGP, but it confirms the qualitative finding, namely a substantial and strongly design-dependent erosion. Two-proportion z -tests on the B-versus-A absolute power-loss gap give $z = 7.64$ ($p < 10^{-13}$) at OL+BDC and $z = 3.96$ ($p < 10^{-4}$) at Hybrid; the CO contrast is not statistically distinguishable from zero ($z = 0.29$) because the CO design's within-subject AR(1) structure dominates the carryover-mediated correlation decay channel that Architecture B uses to encode the biomarker-treatment interaction.

627 4.2 3.2 Why the architectures diverge

628 We turn now to a mechanistic account of why the two architectures produce such
629 different power losses under carryover.

630 **Architecture A is subject to only one carryover channel (mean blurring);**
631 **it escapes the second channel (correlation erosion).**

- 632 • Carryover compresses Dbc contrast and inflates off-drug BR means; both
633 architectures share this channel.
- 634 • Architecture A lacks a second channel because β_{bm} is a fixed population
635 parameter and the noise structure is drug-state-independent.
- 636 • Result: reduced variance in the interaction term, but signal-to-noise ratio
637 degrades only modestly.

638 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
639 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
640 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
641 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
642 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
643 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
644 mollit anim id est laborum.

645 *Under Architecture A, carryover reduces the magnitude of the drug expo-*
646 *sure variable seen by the analysis model (D_{bc} drops below 1 during washout),*
647 *and the mean BR at off-drug timepoints is inflated by the residual drug effect*
648 *exactly as in Architecture B. Both mechanisms compress the between-state*
649 *treatment contrast and are shared by the two architectures. What Architec-*
650 *ture A lacks is the second channel of signal loss: the biomarker moderation*
651 *coefficient β_{bm} is a fixed population parameter, the biomarker enters the mean*
652 *structure directly rather than through a correlation, and the noise structure is*
653 *independent of drug state. The biomarker-by-drug interaction term therefore*
654 *has reduced variance (because D_{bc} has less contrast between drug states), but*
655 *the signal-to-noise ratio degrades only modestly.*

656 **Under Architecture B, carryover operates on two levels simultaneously:**

- 657 1. **Mean blurring:** The BR means during off-drug periods are inflated by
658 residual carryover, making them resemble on-drug BR means. This reduces
659 the treatment contrast that the analysis model uses to distinguish drug effect
660 from no-drug. This channel is shared with Architecture A.
- 661 2. **Correlation erosion:** The BM-BR correlation during off-drug periods de-
662 cays exponentially with t_{sd} . This is not just a statistical modeling choice; it
663 is a structural consequence of the DGP. The correlation between biomarker
664 and BR reflects the degree to which the drug effect is active; as the drug
665 washes out, the correlation weakens because the drug-mediated component
666 of BR variance shrinks relative to the drug-independent component. This
667 channel is *unique* to Architecture B, because only here does the biomarker
668 signal live in the second moment of the joint distribution. Under Architec-
669 ture A the drug-mediated and drug-independent components of BR variance
670 move in tandem (both are scaled by the same Dbc-weighted drug effect), so

671 their ratio, and hence the identifiability of the interaction, is preserved even
672 as amplitudes shrink.

673 **Under Architecture B the interaction estimate is attacked from both sides**
674 **simultaneously; Architecture A faces only one of the two attacks.**

- 675 • The bm:Dbc coefficient depends on the covariance between B_i and Dbc-
676 weighted outcomes.
- 677 • Under Architecture B: (1) Dbc contrast is compressed; (2) the BM-BR
678 correlation generating the covariance signal is simultaneously decaying.
- 679 • Architecture A faces only attack (1), not attack (2); this single difference
680 accounts for the divergence in power.

681 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
682 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
683 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
684 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
685 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
686 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
687 mollit anim id est laborum.

688 *The analysis model's interaction coefficient (the coefficient on bm:Dbc) de-*
689 *pends on the covariance between B_i and Dbc-weighted outcomes. Under Ar-*
690 *chitecture B, this covariance is being attacked from both sides: the treatment*
691 *contrast in Dbc is compressed, and the BM-BR correlation that generates the*
692 *covariance signal is simultaneously decaying. That Architecture A escapes the*
693 *second attack entirely accounts for the divergence in power loss documented*
694 *above.*

695 4.3 3.3 Robustness check at $N = 70$

696 **An $N = 70$ confirmation run checks whether the qualitative ordering holds**
697 **when power is near the ceiling.**

- 698 • 800 replicates per cell, 12 cells ($2 \text{ architectures} \times 2 c_{bm} \times 3 t_{1/2}$),
699 OL+BDC design.
- 700 • 100% convergence across 9,600 fits.
- 701 • Purpose: verify the Section 3.1 ordering survives saturation at the larger
702 N .

703 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
704 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
705 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
706 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
707 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
708 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
709 mollit anim id est laborum.

710 *The Section 3.1 production results are at $N = 35$, where neither architecture*
711 *saturates power and the carryover-induced loss is mechanically visible. We*
712 *conducted a confirmation run at $N = 70$ to check whether the qualitative or-*
713 *dering survives the saturation that comes with the larger sample size. An 800-*

replicate-per-cell Monte Carlo run was executed at $N = 70$ across 12 cells (architecture $\in \{MVN, \text{mean moderation}\} \times c_{bm} \in \{0, 0.45\} \times t_{1/2} \in \{0, 0.5, 1.0\}$ weeks), all on the OL+BDC design, using 8-core `parallel::mclapply`. Wall-clock 483 s; convergence 100% across 9,600 fits.

At $c_{bm} = 0.45$:

Architecture	$t_{1/2} = 0$	$t_{1/2} = 0.5$	$t_{1/2} = 1.0$
MVN (B)	0.978	0.945	0.885
Mean moderation (A)	0.976	0.973	0.959

The $N = 70$ run confirms the qualitative ordering; compressed absolute magnitudes reflect ceiling saturation, not a smaller underlying effect.

- Losses: 9.3 pp (Architecture B, $p < 10^{-12}$) vs. 1.7 pp (Architecture A, $p \approx 0.05$); ordering matches $N = 35$.
- At $N = 70$ both architectures are within 2–5 pp of the ceiling at $t_{1/2} = 0$, mechanically truncating observable loss.
- $N = 35$ production results remain the canonical reference for the narrative magnitude claims.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Within-architecture losses from $t_{1/2} = 0$ to $t_{1/2} = 1$ are 9.3 percentage points (Architecture B; $z = 7.4$, $p < 10^{-12}$) and 1.7 percentage points (Architecture A; $z = 2.0$, $p \approx 0.05$). The qualitative ordering matches Section 3.1's $N = 35$ result. The absolute magnitudes are compressed because $N = 70$ at the prazosin-PTSD calibration places both architectures within 2–5 percentage points of the unit ceiling at $t_{1/2} = 0$, mechanically truncating the loss either architecture's curve can absorb. We note that reading the narrative loss-magnitude claim against the $N = 70$ robustness numbers (3% and 9% relative loss for Architecture B versus 0.4% and 1.8% for Architecture A) requires acknowledging this saturation effect; the $N = 35$ production results in Section 3.1 are the canonical reference for the narrative magnitude.

Type I error is well-controlled across all 18 null cells, with no architecture-specific inflation.

- Type I error range [0.010, 0.074], mean 0.043, median 0.042 across 18 null cells at $N = 35$.
- No architecture-specific inflation detected.
- The slightly conservative Architecture B OL+BDC null cell at $t_{1/2} = 1.0$ ($p = 0.010$) is the only notable departure.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad

754 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
755 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
756 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
757 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
758 mollit anim id est laborum.

759 *Type I error.* Across all 18 null cells of the $N = 35$ production run, type I
760 error sat in $[0.010, 0.074]$, with mean 0.043 and median 0.042. We observe no
761 architecture-specific inflation. Architecture B's $t_{1/2} = 1.0$ OL+BDC null ($p =$
762 0.010, slightly conservative) is the smallest cell; the corresponding alternative
763 ($p = 0.539$) is the most informative cell of the Section 3.1 narrative.

764 **Estimator bias is small and architecturally interpretable; coefficient drift**
765 **under Architecture A contrasts with correlation-channel erosion under**
766 **Architecture B.**

- 767 • Mean $\hat{\beta}_{bm:D_{bc}}$ at $c_{bm} = 0.45$ ranges -0.231 to -0.281 across alternative
768 cells; Architecture A OL+BDC at $t_{1/2} = 1.0$ produces the largest absolute
769 mean.
- 770 • Architecture B cells maintain mean near -0.234 across all $t_{1/2}$ values,
771 consistent with erosion through correlation decay rather than coefficient
772 drift.
- 773 • Nominal-95% Wald coverage holds across the grid.

774 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
775 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
776 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
777 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
778 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
779 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
780 mollit anim id est laborum.

781 *Estimator bias.* Mean estimated $\hat{\beta}_{bm:D_{bc}}$ at $c_{bm} = 0.45$ ranged from -0.231
782 to -0.281 across the 18 alternative cells, with the $N = 35$ Architecture-A
783 OL+BDC cell at $t_{1/2} = 1.0$ producing the largest absolute mean (-0.281); the
784 Architecture B cells maintained mean near -0.234 across all $t_{1/2}$ values, con-
785 sistent with the Architecture-B information channel being eroded through cor-
786 relation decay rather than through coefficient drift. Empirical SE of $\hat{\beta}_{bm:D_{bc}}$
787 runs 0.05 to 0.10 across cells; the within-replicate model SE matches at scale,
788 indicating nominal-95% Wald coverage holds across the grid.

789 4.4 3.4 Architecture C: combined-channel power results

790 **Architecture C power lies between the pure-architecture extremes; the**
791 **3×3 parameter panel quantifies how the two channels trade off.**

- 792 • Grid: $c_{bm,A} \times c_{bm,B} \in \{0, 0.22, 0.45\}^2$; boundary rows/columns reproduce
793 the Section 3.1 baselines.
- 794 • As $c_{bm,B}$ increases from 0 to 0.45, power loss under carryover intensifies
795 monotonically.
- 796 • As $c_{bm,A}$ increases from 0 to 0.45, the mean-shift channel adds signal that
797 is resilient to carryover, partially offsetting the covariance-channel loss.

798 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
799 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
800 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
801 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
802 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
803 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
804 mollit anim id est laborum.

805 *We ran the same factorial simulation on the Architecture C grid. The de-*
806 *sign matrix is identical to Section 3.1: CO, Hybrid, and OL+BDC designs;*
807 *$N = 35$; $t_{1/2} \in \{0, 0.5, 1.0\}$ weeks; 1000 replicates per cell; A2 analysis specifi-*
808 *cation. The parameter panel is $c_{bm,A} \times c_{bm,B} \in \{0, 0.22, 0.45\}^2$ (nine cells per*
809 *design-carryover combination). The boundary cells ($c_{bm,A} = 0.45, c_{bm,B} = 0$)*
810 *and ($c_{bm,A} = 0, c_{bm,B} = 0.45$) reproduce Sections 3.1's Architecture A and*
811 *Architecture B results respectively, providing the internal validation gate.*

812 *[Placeholder: 3×3 power table for CO design, $N = 35$, $t_{1/2} = 1.0$, A2 spec. In-*
813 *sert from `analysis/data/dgp-combined/02-power-grid-co.csv` once the pro-*
814 *duction run completes.]*

815 **The 3×3 panel confirms the mathematical prediction: Architecture C**
816 **power is a monotone function of both parameters, with the covariance**
817 **channel governing the carryover sensitivity.**

- 818 • Interior cells ($c_{bm,A} > 0, c_{bm,B} > 0$) show power strictly between the
819 boundary-row and boundary-column values.
- 820 • Moving along the $c_{bm,B}$ axis (increasing covariance weight, fixed mean-
821 shift weight) degrades power under carryover; moving along the $c_{bm,A}$ axis
822 attenuates that degradation.
- 823 • The equal-weight cell ($c_{bm,A} = c_{bm,B} = 0.45$) has higher baseline power
824 than either pure architecture at $c_{bm} = 0.45$, reflecting the doubled total
825 signal, but also higher absolute carryover sensitivity than Architecture A
826 alone.

827 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
828 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
829 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
830 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
831 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
832 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
833 mollit anim id est laborum.

834 *[Orig placeholder: narrative of 3×3 panel results. To be completed once*
835 *the production simulation (01-run-combined-factorial.R) finishes and*
836 *02-summarise-combined.R produces the summary tables.]*

837

838 5 4. Biological Assumptions

839 **Architecture choice encodes a biological claim about how the biomarker**
840 **moderates treatment response.**

- 841 • The three architectures reflect distinct assumptions about the mechanism
842 of biomarker-treatment moderation.
- 843 • Section 4 takes each architecture in turn (4.1, 4.2, 4.3 for Architectures A,
844 B, and C respectively), then discusses the prazosin-PTSD motivating case.

845 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
846 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
847 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
848 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
849 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
850 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
851 mollit anim id est laborum.

852 *The choice between architectures is not merely a statistical convenience; it*
853 *reflects different assumptions about the biological mechanism by which the*
854 *biomarker moderates treatment response. We consider each architecture in*
855 *turn, then take up the prazosin-PTSD motivating application in detail.*

856 5.1 4.1 When Architecture A is appropriate

857 **Architecture A is appropriate when the biomarker directly determines the**
858 **magnitude of the drug effect; the scaling is deterministic at the mean-**
859 **structure level.**

- 860 • The drug effect scales with the biomarker value: $\text{effect}_i = f(B_i) \cdot \text{drug}$.
- 861 • 'Deterministic' means any two participants with the same B_i have the same
862 expected response; individual outcomes still carry noise.
- 863 • Appropriate when the biomarker is a direct PK/PD determinant (eGFR,
864 receptor density, metabolizer status).

865 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
866 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
867 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
868 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
869 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
870 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
871 mollit anim id est laborum.

872 *Architecture A (direct mean moderation) is appropriate when the biomarker*
873 *governs the **magnitude of the drug's biological effect** for each individual*
874 *participant. The drug effect for participant i is scaled by their biomarker*
875 *value: $\text{effect}_i = f(B_i) \cdot \text{drug}$. The scaling is deterministic in the population-*
876 *mean sense, in that any two participants sharing the same B_i have the same*
877 *expected response, but individual outcomes still carry the usual noise, random-*
878 *effect, and residual-biological variability. The label 'deterministic' refers to*
879 *the mean structure of the DGP, not to the realized response of any single*
880 *participant.*

881 Clinical examples:

- **Renal clearance of a renally-excreted drug.** Estimated glomerular filtration rate (eGFR) directly governs the elimination rate of drugs such as digoxin or aminoglycoside antibiotics. Participants with higher eGFR clear the drug more rapidly, achieving lower steady-state concentrations and a proportionally smaller drug effect. The biomarker-drug relationship is mechanistic and deterministic: given the biomarker value, the effective dose at steady state is fully determined (up to random noise).
- **Receptor density moderation.** A biomarker that measures target receptor density (e.g., PET imaging of receptor availability for an antagonist) directly determines the pharmacodynamic response to that antagonist. More receptors means more drug effect, proportionally.
- **Genetic metabolizer status.** CYP2D6 metabolizer phenotype determines the rate of drug activation or clearance. Poor metabolizers of a prodrug receive less active compound, reducing their treatment effect proportionally.
- **Dose-response with biomarker-dependent effective dose.** When the biomarker determines the effective dose (through body composition, protein binding, or volume of distribution), the treatment effect scales with the biomarker through the dose-response curve.

In all Architecture A cases, the proportional biomarker-drug relationship is preserved during carryover washout.

- Residual drug effect during washout maintains the same proportional relationship with the biomarker.
- If participant A has $2\times$ the drug effect of B on-drug, they also have $\approx 2\times$ the residual effect during washout.
- This preserved proportionality is what Architecture A encodes and why it loses little power under carryover.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

In all these cases the biomarker's moderating effect is preserved under carryover because the residual drug effect during off-drug periods maintains the same proportional relationship with the biomarker. If participant A has twice the drug effect of participant B when on drug, they will also have approximately twice the residual effect during the washout period. That proportionality is precisely what Architecture A encodes.

5.2 4.2 When Architecture B is appropriate

923 **Architecture B is appropriate when the biomarker is a probabilistic pre-**
924 **dictor of response, mediated by latent factors the biomarker imperfectly**
925 **indexes.**

- 926 • Appropriate when the biomarker predicts which participants are likely to
927 respond, not the magnitude of any individual's response.
- 928 • The biomarker-response association is population-level and probabilistic,
929 mediated by latent factors (disease subtype, neurobiological heterogeneity,
930 comorbidities).
- 931 • Two participants with identical biomarker values may respond very differ-
932 ently.

933 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
934 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
935 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
936 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
937 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
938 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
939 mollit anim id est laborum.

940 *Architecture B (differential correlation) is appropriate when the biomarker*
941 *predicts which participants are likely to respond, but the magnitude of*
942 *any individual's response is not deterministically governed by the biomarker.*
943 *The biomarker and the drug response are associated at the population level, but*
944 *the association is mediated by latent factors (disease subtype, neurobiological*
945 *heterogeneity, comorbidities) that the biomarker imperfectly indexes.*

946 **Clinical examples:**

- 947 • **Blood pressure as a PTSD subtype marker.** Elevated resting blood
948 pressure may index a noradrenergic-predominant PTSD phenotype (Hen-
949 drickson et al. 2020) that is more likely to respond to prazosin (an alpha-1
950 adrenergic antagonist). The blood pressure does not *cause* the drug to
951 work better; rather, it correlates with an underlying neurobiological state
952 that determines drug responsiveness. Two participants with identical blood
953 pressure may have very different responses because blood pressure is an
954 imperfect proxy for the noradrenergic state.
- 955 • **Baseline disease severity as a treatment predictor.** Trials in
956 Alzheimer's disease, depression, and chronic pain often find that patients
957 with more severe baseline symptoms show larger treatment effects (the
958 'floor effect' or 'regression to the mean' phenomenon). The baseline severity
959 score correlates with treatment response, but the relationship is statistical
960 (population-level tendency), not deterministic.
- 961 • **Inflammatory biomarkers in psychiatry.** C-reactive protein (CRP) or
962 interleukin-6 levels may predict response to anti-inflammatory augmentation
963 of antidepressants (Raison et al. 2013). High-CRP patients are more likely
964 to respond, but the relationship is probabilistic: many high-CRP patients
965 do not respond, and some low-CRP patients do.

966 • **Genetic risk scores.** Polygenic risk scores predict disease trajectory and
967 treatment response, but explain only a fraction of the variance. The remain-
968 der reflects unmeasured genetic, epigenetic, and environmental factors.

969 **Under Architecture B, carryover weakens the biomarker-response corre-**
970 **lation because the biomarker's predictive power depends on the presence**
971 **of active drug.**

- 972 • As the drug washes out, the drug-mediated variance component diminishes,
973 weakening the BM-BR correlation.
- 974 • The biomarker's predictive power is tied to active drug effect; without drug,
975 it predicts nothing about non-drug symptom change.
- 976 • This is the 'correlation erosion' channel unique to Architecture B.

977 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
978 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
979 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
980 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
981 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
982 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
983 mollit anim id est laborum.

984 *In these cases carryover has a different implication: as the drug washes*
985 *out, the correlation between biomarker and response weakens because the*
986 *drug-mediated component of response variance diminishes. The biomarker's*
987 *predictive power is inherently tied to the presence of active drug effect. With-*
988 *out drug, the biomarker is just a biomarker; it does not predict the non-drug*
989 *component of symptom change.*

990 **Architecture B is a tractable second-moment approximation to the bio-**
991 **logically faithful finite-mixture DGP; adequate for power calculations but**
992 **not a mechanistic endorsement.**

- 993 • If the biomarker indexes a discrete responder class, the faithful DGP is a
994 finite mixture; Architecture B approximates its population-level covariance
995 structure.
- 996 • Architecture B captures the covariance implications of latent-class mem-
997 bership without committing to the full generative form.
- 998 • Adequate for power calculations; should not be read as endorsing MVN
999 differential correlation as the correct mechanism.

1000 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1001 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1002 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1003 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1004 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1005 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1006 mollit anim id est laborum.

1007 *The latent-subtype framing raises a legitimate modeling question: if the*
1008 *biomarker is really a noisy indicator of an underlying responder class,*
1009 *the biologically faithful DGP would be an explicit finite mixture in which*

each participant is drawn from one of two (or more) response distributions with class-membership probability a function of B_i . Under such a mixture DGP the individual-level moderation is discrete, not probabilistic, and the MVN differential-correlation parameterization is best understood as a tractable second-moment approximation to a richer mixture model rather than as the generative mechanism itself. The practical consequence is that Architecture B as implemented here captures the population-level covariance implications of a latent-class mechanism without committing to its full generative form; this is adequate for power calculations that test for the presence of an interaction but should not be read as an endorsement of MVN differential correlation as the correct mechanistic model. Whether the observed power loss under carryover is better mitigated by analysis strategies (Section 6) or by explicit mixture modeling is an open question.

Senn (2016) cautions that apparent personalized response is often a variance artifact; Architecture B presupposes a real, stable latent subtype, which may not hold.

- Senn (2016) argues that apparent biomarker-by-treatment interactions are often variance-component artifacts, not genuine effects.
- Architecture B presupposes the latent subtype is real, identifiable, and stable across the trial period.
- If that presupposition is shaky, the documented power loss may overstate recoverable signal.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The latent-subtype framing of Architecture B should also be adopted with appropriate epistemic caution. Senn (2016) has argued that the apparent 'personal element in response to treatment' is often a variance-component artifact rather than a genuine biomarker-by-treatment interaction, and that even modest within-subject variability can produce convincing-looking but spurious differential response. Architecture B presupposes that the latent subtype indexed by the biomarker is real, identifiable, and stable across the trial period. When that presupposition is shaky, the apparent power loss under carryover documented in Section 3 may overstate what could be recovered under any analysis model, because the underlying signal it is attempting to estimate may itself be smaller than assumed.

5.3 4.3 The prazosin-PTSD case

We turn briefly to the motivating clinical application for pmsimstats, which is the active-duty prazosin trial by Raskind et al. (2013), whose SBP-stratified secondary analysis is the subject of ongoing reanalysis by the pmsimstats team. The biomarker is resting systolic blood pressure (SBP) and the outcome is CAPS (Clinician-Administered PTSD Scale) score change.

The prazosin-PTSD evidence base is mixed; the PACT null result motivates the search for a predictive biomarker, with SBP as the leading candidate.

- Earlier single-site Raskind trials (2003, 2007, 2013) reported benefit on nightmares and PTSD symptoms.
- PACT (2018) failed to replicate; this has been interpreted as consistent with effect heterogeneity.
- SBP, as a noninvasive proxy for central noradrenergic tone, is the leading predictive biomarker candidate.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The clinical evidence base for prazosin in PTSD is mixed. The earlier single-site trials by Raskind and colleagues (2003, 2007, 2013) reported benefit on nightmares and global PTSD symptoms in veteran and active-duty samples. The subsequent multisite Veterans Affairs Cooperative Study Program trial (Raskind et al. 2018), known as PACT, did not replicate the effect: prazosin failed to outperform placebo on the primary endpoints. The negative PACT result has been interpreted by several investigators as consistent with effect heterogeneity, namely that prazosin works in some patients but not others, which motivates the search for a predictive biomarker that would identify the responder subgroup. Resting SBP, as a noninvasive proxy for central noradrenergic tone, is the leading candidate.

The SBP-prazosin mechanism is an Architecture B scenario: SBP is a proxy for noradrenergic state, not a direct PK determinant.

- Elevated SBP indexes heightened noradrenergic tone, a PTSD subtype preferentially responsive to prazosin.
- SBP is a proxy, not a causal determinant; two patients with identical SBP may have very different noradrenergic states.
- This is precisely the Architecture B scenario: probabilistic association, not deterministic scaling.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The biological hypothesis is that elevated SBP indexes heightened central noradrenergic tone, which characterizes a subtype of PTSD that is preferentially responsive to alpha-1 adrenergic blockade by prazosin. This is an Architec-

ture *B* scenario: *SBP* is a proxy for an underlying neurobiological state, not a direct determinant of drug pharmacokinetics. Two patients with the same *SBP* may differ in their actual noradrenergic tone, α -1 receptor distribution, and downstream signaling, leading to different treatment responses despite identical biomarker values.

Hendrickson et al. (2020) correctly uses Architecture B for this context; the carryover power loss is clinically relevant for trial design decisions.

- Architecture B (MVN differential correlation) is appropriate for the SBP-prazosin mechanism.
- The documented carryover sensitivity means designs with longer off-drug periods will have meaningfully less power.
- This effect is larger than a direct mean moderation simulation would predict.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The Hendrickson et al. (2020) DGP uses Architecture B (MVN differential correlation) for this clinical context. The finding that carryover substantially reduces power under this architecture is therefore clinically relevant: it suggests that trial designs with longer off-drug periods will have meaningfully less power to detect the SBP-prazosin interaction than designs with shorter off-drug periods, and this effect is larger than what a direct mean moderation simulation would predict.

5.4 4.4 When Architecture C is appropriate

Architecture C is appropriate when both a pharmacodynamic-scaling pathway and a co-regulatory coupling pathway are plausible for the same biomarker.

- The two channels can co-exist: prazosin's SBP-CAPS mechanism may involve both a mean-moderation pathway (SBP governs effective central noradrenergic suppression) and a differential-correlation pathway (on-drug periods couple noradrenergic tone more tightly to symptom variance).
- When the relative channel weights are unknown, the α -mixing parameterization of Architecture C is a natural sensitivity analysis spanning the range from pure A to pure B.
- Architecture C should not be the default; it has two free parameters where one suffices, and calibration requires independent evidence about both channels.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip

ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Architecture C is appropriate when there is independent evidence that both the pharmacodynamic-scaling pathway and the co-regulatory coupling pathway are biologically active for the same biomarker-drug pair. Such evidence might include mechanistic assay data (e.g., the drug enhances biomarker-outcome correlation in vitro, consistent with the B channel, while receptor occupancy scales with biomarker level, consistent with the A channel) or replication across independent trials that are individually underpowered to distinguish the channels.

We note that Architecture C should not be the default choice for simulation studies, despite its generality. It carries two free parameters where either Architecture A or Architecture B requires only one, and calibrating both parameters independently demands more prior information than is typically available in the trial-design phase. The appropriate use cases are: (a) formal sensitivity analyses that ask how power estimates change across the $(c_{bm,A}, c_{bm,B})$ grid when the channel structure is genuinely uncertain; and (b) evaluation studies designed to determine which channel is dominant for a given biomarker-drug pair, using the simulation as a diagnostic instrument rather than a planning tool.

In the prazosin-PTSD context, Architecture C captures the possibility that SBP simultaneously governs the expected prazosin effect at the mean level (Architecture A, consistent with the pharmacological argument that nora-drenergic tone determines the therapeutic window) and alters the degree of coupling between SBP and CAPS change during active drug exposure (Architecture B, consistent with a signaling-pathway interpretation). We regard Architecture C as the appropriate framework for evaluating the prazosin-SBP interaction if and when mechanistic data become available to calibrate both channels separately.

6 5. Implications for Simulation Study Design

6.1 5.1 Choosing the appropriate architecture

The choice of DGP architecture should be guided by the biological mechanism hypothesized for the biomarker-treatment interaction:

When the channel structure is uncertain, Architecture C with a mixing-weight sensitivity sweep spans the range from pure A to pure B; Architecture B remains the conservative default.

- The α -mixing parameterization sets $c_{bm,A} = (1 - \alpha)c$ and $c_{bm,B} = \alpha c$ for fixed total weight c and $\alpha \in [0, 1]$.
- $\alpha = 0$ recovers pure Architecture A; $\alpha = 1$ recovers pure Architecture B.

Criterion	Architecture A	Architecture B	Architecture C
Biomarker role	Causal mediator	Statistical predictor	Both
Mechanism	Deterministic scaling	Probabilistic association	Both channels
Signal location	Mean structure	Covariance structure	Mean + covariance
Free parameters	1 (c_{bm})	1 (c_{bm})	2 ($c_{bm,A}$, $c_{bm,B}$)
Carryover impact	Modest (1–3%)	Design-dependent, up to 31%	Intermediate; scales with $c_{bm,B}$
Appropriate when	Biomarker determines dose or PK	Biomarker indexes a latent subtype	Channel uncertain; sensitivity use

- Architecture B's larger carryover sensitivity makes it the safer default when the channel structure is unknown.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

If the biological mechanism is ambiguous, Architecture C with a range of mixing weights $\alpha \in [0, 1]$ (where $c_{bm,A} = (1 - \alpha)c$ and $c_{bm,B} = \alpha c$ for fixed total weight c) provides a natural sensitivity framework spanning the two pure-architecture extremes. Architecture B should remain the default conservative choice for trial design when the channel structure is unknown: its larger carryover sensitivity ensures that power estimates will not be systematically optimistic relative to either pure-architecture alternative.

6.2 5.2 Reporting requirements

Simulation studies evaluating biomarker-moderated treatment effects should explicitly state:

1. Whether the biomarker-treatment interaction is generated through mean moderation, differential correlation, or both channels simultaneously.
2. The biological rationale for the chosen mechanism and, for Architecture C, independent evidence supporting each channel.
3. Whether the carryover model (if present) operates on the interaction signal consistently with the chosen architecture.
4. Sensitivity analyses under alternative architectures, if the biological mechanism is uncertain.

The four reporting requirements are the ADEMP 'D' element applied specifically to biomarker-validation simulations.

- ADEMP (Morris et al. 2019) prescribes five elements: Aims, Data-generating mechanisms, Estimands, Methods, Performance measures.
- The DGP-architecture choice concerns the D element; the four requirements above specialize that element for architecture-sensitive simulations.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

*The reporting framework above is consistent with the broader ADEMP template of Morris, White, and Crowther (2019), which prescribes that simulation studies fix five elements before production runs: the **Aims** of the study, the **Data-generating mechanisms**, the **Estimands** (the inferential targets), the **Methods** (analysis specifications compared), and the **Performance measures** (power, bias, coverage, type I error, and Monte Carlo standard errors of each). The DGP-architecture choice this paper concerns is the **D** element of an ADEMP specification; the reporting requirements above are the specialisation of the broader ADEMP framework to the architecture-sensitive class of biomarker-validation simulations.*

Primary estimand: $\beta_{bm:D}$ in the LME model; performance measures cover power, type I error, bias, SE, and convergence; 1000 replicates per cell.

- Estimand: $\beta_{bm:D}$ in $Sx \sim bm + t + Dbc + bm:D bc$ with random intercept and AR(1).
- Five performance measures: power, type I error, bias, empirical SE, convergence rate.
- $n_{reps} = 1000$ throughout; calibrated to achieve MCSE < 0.02 for power near 0.5.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

*The **primary estimand** for this paper's simulations is the biomarker-by-treatment interaction coefficient $\beta_{bm:D}$ in the linear-mixed analysis model $Sx \sim bm + t + Dbc + bm : D bc$ with random intercept and AR(1) residual correlation. The reported **performance measures** are: (i) power for $\beta_{bm:D}$ at $\alpha = 0.05$, with Monte Carlo standard error $\sqrt{\hat{\pi}(1 - \hat{\pi})/n_{reps}}$; (ii) type I error under $c_{bm} = 0$, with the same MCSE; (iii) bias of $\hat{\beta}_{bm:D}$ relative to the calibrated true value at each parameter cell; (iv) the empirical Monte Carlo standard error of $\hat{\beta}_{bm:D}$ across replicates; and (v) the convergence rate of the linear-mixed fit. Sample size, biomarker effect size, and carryover half-life*

1261 vary across the parameter grid; the number of replicates per cell is calibrated
 1262 to achieve MCSE below 0.02 for power at $\hat{\pi} = 0.5$, which corresponds to
 1263 $n_{\text{reps}} \geq 625$. The simulations reported here use $n_{\text{reps}} = 1000$ throughout.

1264 6.3 5.3 When the choice is uncertain

1265 If the biological mechanism is ambiguous, as it often is in early biomarker devel-
 1266 opment, we suggest investigators take the following steps:

- 1267 1. Run the simulation under both architectures and report the range of power
 1268 estimates.
- 1269 2. Architecture B produces more conservative (lower) power estimates under
 1270 carryover, making it the safer default for trial design decisions.
- 1271 3. Design the trial to minimize carryover (adequate washout periods) to reduce
 1272 sensitivity to the DGP architecture choice.

1273

1274 7 6. Alternative Analysis Strategies Under Architec- 1275 ture B

1276 The power loss documented in Section 3 arises from two distinct sources: (1)
 1277 compressed variance of the treatment indicator Dbc when carryover is present,
 1278 and (2) erosion of the BM-BR correlation during off-drug periods. Source 1 is a
 1279 property of the analysis model parameterization and can potentially be addressed
 1280 by choosing a different predictor. Source 2 is a property of the data itself (the
 1281 signal is genuinely weaker in off-drug observations), and no analysis model can
 1282 recover signal that is not present. However, analysis strategies that are selective
 1283 about *which observations contribute to the test* or *how they are weighted* can
 1284 mitigate the damage.

1285 7.1 6.1 Restrict analysis to on-drug observations

1286 The most direct approach is to discard off-drug timepoints entirely and test the
 1287 biomarker-treatment interaction using only on-drug observations. Every retained
 1288 observation has the full BM-BR correlation (c_{bm}), so correlation erosion is elim-
 1289 inated. The costs are twofold: a reduction in effective sample size (fewer obser-
 1290 vations per participant) and, more importantly, the loss of the off-drug reference
 1291 that separates the biological response from placebo and time-varying background
 1292 response. Without off-drug observations the biological, placebo, and time-varying
 1293 components are confounded, which is why the Open-Label design, despite having
 1294 the most on-drug timepoints, does not dominate the other designs in the Section 3
 1295 results. For designs with many on-drug timepoints (OL has 8, the Hybrid design
 1296 has 5-6 per path), the net effect may be positive: the per-observation signal-to-
 1297 noise ratio is maximal, and the analysis model simplifies to $Sx \sim \text{bm} + \text{t} + \text{bm:t}$

1298 since there is no treatment contrast to model. The interaction is detected through
1299 the correlation between biomarker and the time-on-drug trajectory.

1300 **Restricting to on-drug observations eliminates correlation erosion but**
1301 **loses the off-drug reference; best for OL and Hybrid designs, worst for**
1302 **CO.**

- 1303 • Benefit: every retained observation has full c_{bm} ; correlation erosion is
1304 eliminated.
- 1305 • Cost: loss of off-drug reference confounds biological, placebo, and time-
1306 varying components.
- 1307 • Practical: best for OL/Hybrid (many on-drug timepoints); least useful for
1308 CO (discards half the observations).

1309 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1310 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1311 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1312 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1313 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1314 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1315 mollit anim id est laborum.

1316 *This approach is most promising for the Open-Label design (which has no off-*
1317 *drug observations regardless) and the Hybrid design (which has a long initial*
1318 *on-drug phase). It is least useful for the Crossover design, where discarding*
1319 *half the observations substantially reduces statistical power.*

1320 7.2 6.2 Weighted analysis

1321 **Weighted analysis downweights off-drug observations by drug exposure**
1322 **without discarding them; information-theoretic rationale; straightforward**
1323 **to implement.**

- 1324 • Weight: $w_t = e^{-\lambda t_{sd}}$ for off-drug observations; 1 for on-drug.
- 1325 • Rationale: higher exposure = stronger correlation signal = more informa-
1326 tion about the interaction.
- 1327 • Implemented via the weights argument in `nlme::lme`.

1328 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1329 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1330 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1331 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1332 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1333 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1334 mollit anim id est laborum.

1335 *Rather than discarding off-drug observations entirely, a weighted analysis*
1336 *retains all observations but weights them by their estimated drug exposure*
1337 *level. On-drug observations receive weight 1. Off-drug observations receive*
1338 *weight proportional to the expected residual drug effect:*

$$w_t = e^{-\lambda t_{sd}}$$

1339 *This downweights observations where the BM-BR correlation has decayed*
 1340 *(and where they contribute more noise than signal to the interaction estimate)*
 1341 *without discarding them entirely. The rationale is information-theoretic:*
 1342 *observations with higher drug exposure carry more information about the*
 1343 *biomarker-treatment interaction because the correlation signal is stronger.*
 1344 *This is straightforward to implement in `nlme::lme` using the `weights` argu-*
 1345 *ment.*

1346 7.3 6.3 Within-subject contrast

1347 For designs with both on-drug and off-drug periods (CO, Hybrid, OL+BDC), a
 1348 summary-measure approach computes a per-participant contrast: the difference
 1349 between mean on-drug response and mean off-drug response. The biomarker-
 1350 treatment interaction is then tested by a simple regression of these contrasts on
 1351 the biomarker:

$$\bar{Y}_{i,\text{on}} - \bar{Y}_{i,\text{off}} = \alpha + \gamma \cdot B_i + \epsilon_i$$

1352 **Within-subject contrast sidesteps the repeated-measures model; trades**
 1353 **efficiency for robustness to carryover misspecification.**

- 1354 • Per-participant contrast $\bar{Y}_{\text{on}} - \bar{Y}_{\text{off}}$ is regressed on the biomarker.
- 1355 • Carryover compresses the contrast magnitude but the contrast-biomarker
 1356 correlation may be more robust.
- 1357 • Trade-off: less efficient than a full longitudinal model, but more robust to
 1358 carryover and correlation misspecification.

1359 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
 1360 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
 1361 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
 1362 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
 1363 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
 1364 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
 1365 mollit anim id est laborum.

1366 *This sidesteps the Dbc parameterization and the repeated-measures model*
 1367 *entirely. Carryover affects the magnitude of the contrast (making off-drug*
 1368 *means closer to on-drug means), but the correlation between the contrast and*
 1369 *the biomarker may be more robust because within-subject averaging reduces*
 1370 *noise before the biomarker association is tested. The approach trades the*
 1371 *efficiency of a full longitudinal model for robustness to misspecification of the*
 1372 *carryover and correlation structure.*

1373 7.4 6.4 Two-stage random slopes

1374 A two-stage approach first estimates participant-specific treatment effects, then
 1375 tests their association with the biomarker. Stage 1 fits a random-slopes model:

$$Y_{it} = \beta_0 + \beta_1 t + \beta_2 D_{it} + u_{0i} + u_{1i} D_{it} + \epsilon_{it}$$

1376 where β_2 is the population mean drug effect and u_{1i} is participant i 's random
1377 deviation from that mean. The participant-specific drug effect is then $\beta_2 + u_{1i}$.
1378 Stage 2 regresses the estimated random deviations (or equivalently, the BLUPs of
1379 the participant-specific effects) on the biomarker:

$$\hat{u}_{1i} = \delta_0 + \delta_1 B_i + \eta_i$$

1380 **Two-stage random slopes separates within-subject effect estimation from**
1381 **between-subject biomarker association; may be more robust to carryover.**

- 1382 • Stage 1: random-slopes LME estimates participant-specific drug effects
1383 ($\beta_2 + u_{1i}$).
- 1384 • Stage 2: regress BLUPs on biomarker to test association.
- 1385 • The random slope u_{1i} integrates over the full trajectory; robustness to
1386 carryover is plausible but unverified by simulation.

1387 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1388 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1389 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1390 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1391 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1392 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1393 mollit anim id est laborum.

1394 *This separates the within-subject treatment effect estimation (which uses*
1395 *the longitudinal data and is affected by carryover) from the between-subject*
1396 *biomarker association (which uses the cross-sectional relationship between*
1397 *estimated effects and biomarker values). The approach may be more robust*
1398 *to carryover because the random slope u_{1i} integrates over the participant's*
1399 *entire trajectory, including both on-drug and off-drug phases.*

1400 7.5 6.5 Exclusion of contaminated observations

1401 **Excluding the first 1–2 off-drug timepoints removes the most contami-**
1402 **nated observations; later off-drug timepoints provide a cleaner contrast.**

- 1403 • Early off-drug timepoints: highest residual drug effect, BM-BR correlation
1404 actively decaying.
- 1405 • Later off-drug timepoints: fully washed out, cleaner contrast with on-drug
1406 phase.
- 1407 • Result: comparison between fully-on and fully-off states, avoiding the
1408 strained intermediate zone.

1409 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1410 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1411 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1412 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1413 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1414 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1415 mollit anim id est laborum.

The first 1-2 off-drug timepoints carry the most carryover contamination: the residual drug effect is highest, and the BM-BR correlation is in the process of decaying. Later off-drug timepoints, where the drug has fully washed out, provide a cleaner contrast with the on-drug phase. Excluding the early off-drug observations and retaining only the later ones yields a comparison between fully-on and fully-off states, avoiding the intermediate zone where the analysis model is most strained.

7.6 6.6 Design-level solutions

Rather than adapting the analysis model to accommodate carryover, the trial design itself can be modified to reduce carryover exposure:

- **Longer washout periods** between drug phases eliminate residual drug effect before off-drug measurements begin.
- **More on-drug timepoints** relative to off-drug increase the proportion of high-signal observations.
- **Front-loaded drug exposure** (as in the OL+BDC design) places a long initial on-drug block before the off-drug baseline-during-continued-treatment phase, so that the BM-BR correlation has been fully established before any off-drug observations are collected. The alternative ordering (off-drug first, then on-drug) has the appealing property of placing the placebo and time-varying estimation in an uncontaminated window but sacrifices on-drug exposure duration and forces the BM-BR signal to be re-established after a wash-in; the net power trade-off depends on the ratio of carryover half-life to total trial length and is explored empirically in the Section 3 results.

Design-level changes address the root cause rather than the symptom; cost is longer trial or reduced ability to estimate carryover dynamics.

- Three levers: longer washout, more on-drug timepoints, front-loaded drug exposure (as in OL+BDC).
- Front-loading establishes BM-BR correlation before any off-drug observations; alternative ordering (off-drug first) trades clean placebo estimation for reduced on-drug exposure.
- Addresses root cause; cost is typically longer trial or reduced carryover estimation ability.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Design-level solutions address the root cause (carryover in the data) rather than the symptom (power loss in the analysis). Their cost is typically a

1458 *longer trial duration or reduced ability to estimate the carryover dynamics*
1459 *themselves.*

1460 7.7 6.7 Comparative evaluation

1461 The strategies described above differ in their assumptions, implementation com-
1462 plexity, and expected power recovery:

Strategy	Discards data	Complexity	Expected benefit
On-drug only	Yes	Low	High for OL/Hybrid; poor for CO
Weighted	No	Low	Moderate; depends on weight calibration
Within-subject contrast	Aggregates	Low	Moderate; robust to misspecification
Two-stage random slopes	No	Moderate	Unknown; separates estimation stages
Exclude early off-drug	Some	Low	Moderate; improves contrast clarity
Design modification	N/A	N/A	High; addresses root cause

1463 A systematic simulation comparison of these strategies under Architecture B, anal-
1464 ogous to the factorial comparison of analysis models in Section 3, would identify
1465 which approach best recovers the power lost to carryover for each trial design. We
1466 regard this as a direction for future work.

1467 _____

1468 8 7. Relationship to the Broader Literature

1469 8.1 7.1 Treatment effect heterogeneity

1470 **The PATH Statement distinguishes risk modeling from effect modeling**
1471 **for predictive HTE analysis; this is an analysis-model distinction, not a**
1472 **DGP distinction.**

- 1473 • Rothwell (2005) establishes the predictive/prognostic biomarker distinction;
1474 Kent et al. (2020) synthesize it into the PATH Statement.
- 1475 • PATH's risk modeling approach stratifies on a multivariable risk score; ef-
1476 fect modeling adds explicit interaction terms.
- 1477 • PATH favors risk modeling and cautions against data-driven effect model-
1478 ing because of low power, multiplicity, and overfitting.

1479 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1480 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1481 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1482 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1483 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur

1484 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1485 mollit anim id est laborum.

1486 *The statistical literature on treatment effect heterogeneity is extensive. Roth-*
1487 *well (2005) sets out the foundational distinction between predictive biomark-*
1488 *ers (those that modify the treatment effect) and prognostic biomarkers (those*
1489 *that predict outcome regardless of treatment). The current consensus state-*
1490 *ment on predictive HTE analysis is the Predictive Approaches to Treatment*
1491 *effect Heterogeneity (PATH) Statement of Kent et al. (2020), which distin-*
1492 *guishes two analytic strategies for identifying predictive biomarker effects in*
1493 *randomized trials: a risk modeling approach, in which a multivariable risk*
1494 *model (developed without a treatment term) is used to stratify patients and*
1495 *examine variation in absolute treatment benefit across risk strata, and an*
1496 *effect modeling approach, in which the regression includes explicit treatment-*
1497 *by-covariate interaction terms. PATH expresses caution about data-driven*
1498 *effect modeling on the grounds of low statistical power, multiplicity, and over-*
1499 *fitting, and favors risk modeling when the application criteria are met.*

1500 **PATH's analysis-model axis (risk vs. effect modeling) is orthogonal to our**
1501 **DGP axis; both PATH strategies implicitly assume Architecture A.**

- 1502 • PATH's distinction is analysis-model; ours is DGP; they are orthogonal.
- 1503 • Both PATH strategies assume heterogeneity is in the mean structure (Ar-
- 1504 chitecture A); Architecture B does not fit either PATH category.
- 1505 • Architecture A = parametric interaction model; Architecture B = latent-
- 1506 class/mixture-model perspective.

1507 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1508 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1509 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1510 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1511 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1512 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1513 mollit anim id est laborum.

1514 *We note that the PATH risk-versus-effect-modeling distinction is an analysis-*
1515 *model axis. The Architecture A versus Architecture B distinction introduced*
1516 *in this paper is a data-generating-process axis, and the two are orthogonal.*
1517 *Both PATH analytic strategies implicitly assume that the underlying treat-*
1518 *ment effect heterogeneity is encoded in the mean structure of the outcome (the*
1519 *heterogeneity is identifiable from differences in conditional means), which cor-*
1520 *responds to Architecture A as a DGP. Architecture B as a DGP is a separate*
1521 *phenomenon that does not map cleanly onto either of PATH's analysis-model*
1522 *categories: the interaction signal resides in the covariance structure rather*
1523 *than the mean, and the standard regression-based predictive HTE analysis*
1524 *methods endorsed by PATH may have reduced power under Architecture B*
1525 *precisely because they were developed under an implicit Architecture A as-*
1526 *sumption. Architecture A can accordingly be characterized as a parametric*
1527 *interaction model (in PATH's effect modeling sense). Architecture B can*
1528 *be characterized as a latent class or mixture model perspective in which the*
1529 *biomarker is a noisy indicator of class membership.*

8.2 7.2 Prevalence of each architecture

Architecture A is the near-universal standard across the biomarker-moderated trial simulation literature; Architecture B is absent from all major threads except Hendrickson et al.

- Simon (2010), Freidlin and Korn (2014), Renfro and Sargent (2017), Haller and Ulm (2018), and the platform trial literature all use Architecture A.
- N-of-1 simulation studies (Zucker et al. 1997; Araujo et al. 2016; Duan et al. 2013) use hierarchical random-effects models, a variant of Architecture A.
- Wang and Schork (2019) and Schork (2022) share authorship with Hendrickson et al. yet adopt Architecture A, suggesting the differential-correlation choice is paper-specific.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

A survey of the simulation methodology literature reveals that Architecture A (direct mean moderation) is the near-universal standard. Simon (2010) employs it throughout the simulation frameworks he develops for biomarker-stratified designs; Freidlin and Korn (2014) adopt it for adaptive enrichment trials; Renfro and Sargent (2017) use it for basket and umbrella trials; and it underlies the platform trial literature throughout. Haller and Ulm (2018), in simulation work on biomarker-by-treatment interaction estimation in randomized trials with prognostic covariates, likewise adopt the standard regression specification, focusing on power and bias of the interaction estimator under different covariate adjustment strategies. N-of-1 simulation studies, including Zucker et al. (1997), Araujo et al. (2016), and Duan et al. (2013), use hierarchical random-effects models in which individual treatment effects are drawn from a population distribution, a variant of Architecture A in which the biomarker predicts the random treatment effect through the mean structure. Methodological work on power and design for N-of-1 trials with serial correlation and carryover, from both Wang and Schork (2019) and Schork (2022), adopts the same Architecture A specification, even though this work shares senior authorship with Hendrickson et al. (2020). It appears, then, that the differential-correlation parameterization in the latter is a paper-specific design choice rather than a methodological convention of any laboratory or research program.

Architecture B appears in latent variable / SEM / copula contexts but is unique in N-of-1 clinical trial simulation.

- Architecture B appears in latent variable / SEM / Bayesian joint modeling / copula contexts (e.g., Rizopoulos 2012).
- In N-of-1 clinical trial simulation, Hendrickson et al. (2020) appears unique in using it.

1576 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1577 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1578 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1579 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1580 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1581 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1582 mollit anim id est laborum.

1583 *Architecture B (differential correlation via MVN) appears primarily in la-*
1584 *tent variable and structural equation modeling contexts, including the joint*
1585 *modeling framework discussed by Rizopoulos (2012), Bayesian joint model-*
1586 *ing frameworks, and copula-based simulation studies. In the clinical trial*
1587 *simulation literature specifically, the Hendrickson et al. (2020) framework*
1588 *appears to be unique in using differential correlation as the mechanism for*
1589 *biomarker-treatment interaction in an N-of-1 trial context.*

1590 **The near-universal use of Architecture A means most published power**
1591 **estimates are optimistic for Architecture B biomarkers under carryover.**

- 1592 • Virtually all published power estimates for biomarker-moderated trials as-
1593 sume Architecture A.
- 1594 • These estimates may be optimistic for biomarkers that operate through
1595 Architecture B (covariance-structure mechanism).
- 1596 • The additional carryover-induced power loss under Architecture B is not
1597 accounted for.

1598 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1599 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1600 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1601 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1602 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1603 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1604 mollit anim id est laborum.

1605 *This predominance of Architecture A has an important implication: the power*
1606 *estimates reported in the vast majority of biomarker-moderated trial simula-*
1607 *tion studies may be optimistic for biomarkers that operate through an Archi-*
1608 *itecture B mechanism, because these studies do not account for the additional*
1609 *power loss that carryover produces when the interaction signal resides in the*
1610 *covariance structure rather than the mean structure.*

1611 8.3 7.3 Crossover and N-of-1 trial methodology

1612 **The crossover trial literature addresses carryover as a nuisance parameter**
1613 **but ignores the architectural distinction; the standard recommendation**
1614 **applies differently under each.**

- 1615 • Senn (2002), Jones and Kenward (2014), and Sturdevant and Lumley
1616 (2021) treat carryover as a nuisance parameter; none engages with DGP
1617 architecture.
- 1618 • Standard recommendation (include carryover in model or use adequate
1619 washout) applies to both architectures.

- Under Architecture B, crossover designs with short washout periods may be less suitable than previously recognized.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The crossover trial literature, represented most thoroughly by Senn (2002) and Jones and Kenward (2014), addresses carryover extensively but does not distinguish between architectures for biomarker-treatment interaction modeling. More recent methodological work by Sturdevant and Lumley (2021) on testing for carryover effects via mixed-effects models likewise treats carryover as a nuisance parameter to be estimated and adjusted for, without engaging with how carryover interacts with biomarker-by-treatment effect signals. The standard recommendation, namely to always include carryover in the analysis model or to use adequate washout, applies to both architectures but has different power implications under each. The finding that Architecture B is more sensitive to carryover suggests that crossover designs with short washout periods may be less suitable for detecting correlation-based biomarker interactions than previously recognized.

8.4 7.4 Precision medicine trial design

Enrichment and biomarker-stratified designs assume Architecture A for power calculations; the architectural choice matters differently across three design families.

- Simon (2010) and Freidlin and Korn (2014) frame power calculations in the Architecture A regression idiom.
- Three design families warrant separate consideration: adaptive enrichment, treatment-switching, and basket/umbrella trials.
- The architectural assumption is rarely made explicit, so the practical implications are easy to overlook.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Enrichment and biomarker-stratified designs, as discussed by Simon (2010) and Freidlin and Korn (2014), use Architecture A exclusively for power calculations. Three families of design are particularly worth considering in light of the architectural distinction.

Adaptive enrichment trials are vulnerable to Architecture B power short-

falls precisely when the biomarker mechanism is most uncertain.

- Interim power calculations driving adaptive decisions use Architecture A; if Architecture B holds and washout occurs, realized power will be lower than predicted.
- The mismatch biases adaptation rules toward premature futility stopping or overly narrow enrichment criteria.
- The problem is worst when the biomarker mechanism is unknown, exactly when adaptive enrichment is most often proposed.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Adaptive enrichment trials adjust the eligible population mid-trial based on emerging evidence of a biomarker-treatment interaction. The interim power calculations that drive these adaptive decisions are universally framed in the Architecture A regression idiom. If the true biomarker mechanism is closer to Architecture B and any participants pass through washout intervals between dose-finding stages, the realized power for the interaction test will be lower than the adaptation rule expects, biasing decisions toward premature futility stopping or toward overly narrow enrichment criteria. The discrepancy is most acute when the biomarker mechanism is uncertain at the design stage, which is precisely the situation in which adaptive enrichment is most attractive.

Treatment-switching designs suffer Architecture B power loss through washout erosion; parallel-group enrichment designs do not, because there is no off-drug phase.

- Crossover, basket, platform, and N-of-1 hybrid designs all cycle participants through drug states; each washout phase erodes BM-BR correlation under Architecture B.
- Parallel-group enrichment designs escape this loss because each participant stays on a single treatment throughout.
- The architectural choice matters most in within-subject designs, where it is also most likely to be clinically important.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Treatment-switching designs (most crossover, basket, and platform trials, as well as N-of-1 hybrid designs) cycle individual participants through different

drug states. Under Architecture B, each washout phase erodes the BM-BR correlation, so the biomarker-treatment interaction signal in the second and subsequent treatment periods is structurally weaker than in the first. The power loss documented in Section 3 applies here directly. By contrast, parallel-group enrichment designs, in which each participant remains on a single treatment for the duration of the trial, escape this loss because there is no off-drug phase during which the correlation can decay; for these designs, Architectures A and B yield similar power estimates. The architectural choice therefore matters most in exactly the designs where the biomarker mechanism is most likely to be of clinical interest, namely those that deliberately expose each participant to both treatment states in order to estimate within-subject responsiveness.

Basket and umbrella trials face an additional risk of optimistic aggregated power estimates when Architecture B holds in even a subset of arms.

- Renfro and Sargent (2017) discuss power calculations that average across biomarker-defined arms; the calculation assumes Architecture A throughout.
- Mixed architectures across baskets produce optimistic trial-level power estimates that are difficult to detect from aggregate results.
- Subgroup-specific power diagnostics under both architectures are warranted when the biological mechanism is heterogeneous.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Basket and umbrella trials, as discussed by Renfro and Sargent (2017), introduce an additional layer because their power calculations average across multiple biomarker-defined arms. If even a subset of the arms operate through an Architecture B mechanism while the trial-wide power calculation assumes Architecture A throughout, the aggregated power estimates will be biased optimistically in a way that is difficult to detect from trial-level results alone. Subgroup-specific power diagnostics under both architectures should accompany the primary planning analysis whenever the underlying biology is heterogeneous across baskets.

Across all three design families, the practical recommendation is to anchor sample size planning on the Architecture B estimate when the biomarker mechanism is uncertain.

- Run power calculations under both architectures and treat Architecture B as the conservative anchor.
- This recommendation parallels the advice in Section 5.3 and is consistent with the PATH Statement (Kent et al. 2020).
- The PATH Statement's caution about data-driven interaction discovery reinforces the case for conservative planning under architectural uncertainty.

1755 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1756 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1757 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1758 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1759 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1760 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1761 mollit anim id est laborum.

1762 *In all three settings the practical recommendation parallels that of Section 5.3:*
1763 *when the biological mechanism is uncertain, run the power calculation under*
1764 *both architectures and treat the Architecture B estimate as the more conserva-*
1765 *tive anchor for sample size planning. This recommendation is consistent with*
1766 *the PATH Statement (Kent et al. 2020), which itself counsels caution about*
1767 *data-driven discovery of treatment-by-covariate interactions and prefers risk*
1768 *modeling approaches when the application criteria are met.*

1769

1770 9 8. Conclusions

1771 In conclusion, six points are to be emphasized.

- 1772 1. The choice between direct mean moderation (Architecture A) and differen-
1773 tial correlation (Architecture B) for generating biomarker-treatment inter-
1774 actions in clinical trial simulations has substantive consequences for power
1775 estimates under carryover.
- 1776 2. Under Architecture A, carryover reduces power only modestly, by 1 to 3%
1777 in relative terms across the three designs examined, because the propor-
1778 tional biomarker-drug relationship is preserved during washout. Under Ar-
1779 chitecture B the loss is strongly design-dependent: it remains modest in the
1780 crossover and hybrid designs (3 to 5%) but reaches roughly 31% in the open-
1781 label OL+BDC design, because the differential-correlation signal erodes as
1782 the drug effect wanes.
- 1783 3. The choice should be guided by the hypothesized biological mechanism: Ar-
1784 chitecture A for biomarkers that directly determine drug pharmacokinetics
1785 or pharmacodynamics; Architecture B for biomarkers that statistically pre-
1786 dict treatment response through latent subtype membership.
- 1787 4. For the prazosin-PTSD application with blood pressure as the predictive
1788 biomarker, Architecture B (MVN differential correlation) is the more ap-
1789 propriate model, and the resulting power sensitivity to carryover should
1790 inform trial design decisions regarding off-drug period duration.
- 1791 5. The existing clinical trial simulation literature uses Architecture A almost
1792 exclusively. The Hendrickson et al. (2020) MVN approach is, to the best

of our knowledge, unique in the N-of-1 trial context. This means that published power estimates for biomarker-moderated crossover and N-of-1 designs may be systematically optimistic for biomarkers that operate through a correlation-based (Architecture B) mechanism, because they do not account for the carryover sensitivity documented here.

6. Simulation studies for biomarker-moderated treatment effects should explicitly report their DGP architecture and consider sensitivity analyses under the alternative when the biological mechanism is uncertain.

10 Conflict of interest

The authors declare no conflicts of interest.

11 Funding

This work received no specific funding.

12 Data availability statement

The data and code that support the findings of this study are openly available in the pmsimstats-ng project repository. Per-replicate Monte Carlo outputs and ADEMP-compliant cell-level summaries are versioned under `analysis/data/`; simulation drivers are at `analysis/scripts/`. Exact paths for this manuscript are listed in the Reproducibility section below.

13 Reproducibility

This manuscript was rendered with R 4.4.x inside the project's `zzcollab` Docker container (image hash recorded in `Dockerfile`; package set captured in `renv.lock`). Principal packages used by the simulation pipeline are `nlme` (continuous-time linear-mixed analysis with `corCAR1` residual correlation), `parallel` (8-core `mclapply` for parameter-grid sweeps), `data.table` (the `original-extended` simulation collection), `furrr` and `purrr` (the `tidyverse` simulation collection), and `rmarkdown` plus `bookdown` for the manuscript build.

Random-number generator. The simulation driver sets `set.seed(20260507)` at the top of the replicate loop and uses `parallel::mc.set.seed` inside `mclapply`. Per-replicate seeds derive from the master seed by `+ rep_idx`.

Data and code availability. Per-replicate simulation outputs are checkpointed at `analysis/data/quick-sim/01-dgp-replicates.rds`; Morris ADEMP cell-level summaries are at the companion `01-dgp-summary.txt`. The simulation driver is at `analysis/scripts/quick-sim/01-dgp-prototype.R`. The manuscript source, bibliography, and SIM CSL file are at `analysis/report/01-dgp-mean-moderation-vs-mvn/`

14 References

Araujo A, Julious S, Senn S. Understanding variation in sets of N-of-1 trials. *PLoS ONE*. 2016;11(12):e0167167.

Duan N, Kravitz RL, Schmid CH. Single-patient (N-of-1) trials: a pragmatic clinical decision methodology. *J Clin Epidemiol*. 2013;66(8):S21-S28.

Freidlin B, Korn EL. Biomarker enrichment strategies: matching trial design to biomarker credentials. *Nat Rev Clin Oncol*. 2014;11(2):81-90.

Haller B, Ulm K. A simulation study on estimating biomarker-treatment interaction effects in randomized trials with prognostic variables. *Trials*. 2018;19(1):128.

Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: statistical methods and practical considerations. *Front Digit Health*. 2020;2:13.

Jones B, Kenward MG. *Design and Analysis of Cross-Over Trials*. 3rd ed. Boca Raton, FL: CRC Press; 2014.

Kent DM, Paulus JK, van Klaveren D, et al. The Predictive Approaches to Treatment effect Heterogeneity (PATH) Statement. *Ann Intern Med*. 2020;172(1):35-45.

Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Stat Med*. 2019;38(11):2074-2102.

Raison CL, Rutherford RE, Woolwine BJ, et al. A randomized controlled trial of the tumor necrosis factor antagonist infliximab for treatment-resistant depression: the role of baseline inflammatory biomarkers. *JAMA Psychiatry*. 2013;70(1):31-41.

Raskind MA, Peskind ER, Kanter ED, et al. Reduction of nightmares and other PTSD symptoms in combat veterans by prazosin: a placebo-controlled study. *Am J Psychiatry*. 2003;160(2):371-373.

Raskind MA, Peskind ER, Hoff DJ, et al. A parallel group placebo-controlled study of prazosin for trauma nightmares and sleep disturbance in combat veterans with post-traumatic stress disorder. *Biol Psychiatry*. 2007;61(8):928-934.

- Raskind MA, Peterson K, Williams T, et al. A trial of prazosin for combat trauma PTSD with nightmares in active-duty soldiers returned from Iraq and Afghanistan. *Am J Psychiatry*. 2013;170(9):1003-1010.
- Raskind MA, Peskind ER, Chow B, et al. Trial of prazosin for post-traumatic stress disorder in military veterans. *N Engl J Med*. 2018;378(6):507-517.
- Renfro LA, Sargent DJ. Statistical controversies in clinical research: basket trials, umbrella trials, and other master protocols. *Ann Oncol*. 2017;28(1):34-43.
- Rizopoulos D. *Joint Models for Longitudinal and Time-to-Event Data: With Applications in R*. Boca Raton, FL: CRC Press; 2012.
- Rothwell PM. Treating individuals 5. Subgroup analysis in randomised controlled trials: importance, indications, and interpretation. *Lancet*. 2005;365(9454):176-186.
- Schork NJ. Accommodating serial correlation and sequential design elements in personalized studies and aggregated personalized studies. *Harv Data Sci Rev*. 2022;Special Issue 3.
- Senn S. *Cross-over Trials in Clinical Research*. 2nd ed. Chichester, UK: John Wiley & Sons; 2002.
- Senn S. Mastering variation: variance components and personalised medicine. *Stat Med*. 2016;35(7):966-977.
- Simon R. Clinical trial designs for evaluating the medical utility of prognostic and predictive biomarkers in oncology. *Per Med*. 2010;7(1):33-47.
- Sturdevant SG, Lumley T. Statistical methods for testing carryover effects: a mixed effects model approach. *Contemp Clin Trials Commun*. 2021;22:100711.
- Varadhan R, Segal JB, Boyd CM, Wu AW, Weiss CO. A framework for the analysis of heterogeneity of treatment effect in patient-centered outcomes research. *J Clin Epidemiol*. 2013;66(8):818-825.
- Wang Y, Schork NJ. Power and design issues in crossover-based N-of-1 clinical trials with fixed data collection periods. *Healthcare (Basel)*. 2019;7(3):84.
- Zucker DR, Schmid CH, McIntosh MW, D'Agostino RB, Selker HP, Lau J. Combining single patient (N-of-1) trials to estimate population treatment effects and to evaluate individual patient responses to treatment. *J Clin Epidemiol*. 1997;50(4):401-410.

15 References

Rendered on 2026-06-10 at 18:01 PDT.

Source: `~/prj/res/36-pmsimstats-ng/pmsimstats-ng/analysis/report/01-dgp-mean-moderation-vs-mv`

2026-06-10 18:01 PDT

47

`~/prj/res/36-pmsimstats-ng/pmsimstats-ng/analysis/report/
01-dgp-mean-moderation-vs-mvn/report.Rmd`