

# Exposure-weighted versus lagged-treatment carryover specifications under differential-correlation biomarker interactions: a factorial simulation study

pmsimstats team

2026-07-30

## Abstract

**Background.** Aggregated N-of-1 trials pool repeated within-patient on/off contrasts to validate predictive biomarkers. Under the multivariate-normal (MVN) differential-correlation architecture (Architecture B), in which the biomarker-treatment interaction is encoded in a treatment-state-dependent correlation, a companion study found that carryover – residual drug effect that decays after discontinuation – costs up to roughly 31% of power in relative terms. The analyst must still decide how to represent residual exposure in the linear predictor, and it is not established which representation is most robust to the (typically unknown) pharmacokinetic decay shape.

**Methods.** We compare two analysis-model specifications for the drug-exposure regressor under Architecture B: an exposure-weighted continuous indicator that assumes a parametric decay and half-life (M2, the current default), and the binary on-drug indicator augmented by a lagged just-off-drug term that estimates carryover non-parametrically (M3, the classical crossover device). We embed them in an ADEMP-structured factorial simulation<sup>1</sup> crossing two data-generating process (DGP) decay forms (exponential, Weibull) with the two specifications, under three trial designs, two sample sizes, and three interaction strengths (216 cells, 500 replicates each), plus four single-axis sensitivity blocks and a cluster-robust recalibration. Every performance measure is reported with its Monte Carlo standard error (MCSE).

**Results.** M2 attains the highest power only in the two designs with a blinded-discontinuation phase (Hybrid and OL+BDC), where its advantage over M3 is about 10 percentage points (0.860 vs 0.770 at the reference cell). Under the classical crossover (CO) design M2 is markedly inferior (0.488 vs 0.830 for M3, which tracks the binary indicator closely). Where M2 leads, the lead is robust to decay-shape and half-life mis-specification and to autocorrelation strength, and a paired-McNemar high-precision rerun confirms it is real at the highest-leverage cell (+7.2 points,  $p = 1.2 \times 10^{-7}$ ), though the two specifications are statistically tied at three of four other rerun cells. The model-based interaction test is mildly conservative (mean null rejection 0.039), an artifact of the working covariance that a CR2 cluster-robust standard error removes without disturbing the ranking.

**Conclusions.** Under Architecture B, the exposure-weighted predictor M2, reported with a cluster-robust standard error, is recommended for the Hybrid and OL+BDC designs; the lagged-treatment specification M3 is preferable, or at least not worse, under a crossover design. Across all conditions, anticipated dropout governs power far more than the choice between the two specifications, and should dominate design effort accordingly. This manuscript restricts the grid to Architecture B and to the exponential and Weibull decay forms, and does not evaluate the binary on-drug specification (M1) or Architecture A (mean moderation); the full-scope evaluation, including those comparisons, is in the companion report.Rmd / report-slim.Rmd manuscripts.

# 1 Introduction

## 1.1 A short primer on aggregated N-of-1 trials

[1]

**Aggregated N-of-1 trials pool repeated within-patient on/off contrasts to validate predictive biomarkers at far smaller sample sizes than parallel-group trials.**

- A single N-of-1 trial crosses one patient on and off treatment repeatedly; pooling many such series gives population-level inference.
- The estimand is the biomarker-treatment interaction: whether a baseline covariate  $B_i$  modifies the treatment effect.
- Three designs (crossover (CO); open-label plus blinded-discontinuation (OL+BDC); Hybrid) differ in how on- and off-drug occasions are arranged.

*rgt to complete.*

In a single N-of-1 trial one patient is repeatedly crossed between active treatment and a comparator over a sequence of measurement occasions; the within-patient on-versus-off contrast estimates that patient’s individual treatment effect. An *aggregated* N-of-1 trial pools many such single-patient series in a mixed model, trading some within-patient resolution for population-level inference and far smaller sample sizes than a parallel-group trial<sup>2</sup>. The estimand of interest is the biomarker-treatment interaction, i.e. whether a baseline covariate  $B_i$  modifies the treatment effect. We study three designs that differ in how on- and off-drug occasions are arranged: a classical crossover (CO), an open-label plus blinded-discontinuation design (OL+BDC), and a Hybrid N-of-1 design; these are defined in §3.

[2]

**Carryover – residual drug effect that decays after discontinuation – contaminates off-drug occasions and attenuates the interaction test.**

- The drug effect fades over days to weeks rather than switching off, so off-drug occasions carry residual exposure.
- Treating them as cleanly off-drug mis-specifies the design matrix.
- We ask which of two analyst-side exposure representations is most robust to the resulting attenuation.

*rgt to complete.*

The inferential nuisance specific to this setting is **carryover**. When the drug is discontinued its pharmacological effect does not switch off but decays over days to weeks, so off-drug occasions carry residual exposure. Treating them as cleanly off-drug mis-specifies the design matrix and attenuates the interaction test. The question we take up is which of two analyst-side representations of residual exposure best recovers the lost power.

## 1.2 What is already known, and the gap

[3]

**Under the differential-correlation architecture (B), carryover costs roughly 31% of power in relative terms, and the choice of analysis specification for representing residual exposure has not been evaluated against decay-form uncertainty.**

- Architecture B encodes the biomarker interaction in a treatment-state-dependent correlation, which carryover erodes directly.

- A companion study fixed the decay shape at exponential and used a single analysis specification.
- This manuscript restricts to Architecture B, where the carryover problem is largest, and asks which specification is most robust when the decay shape is uncertain.

*rgt to complete.*

A companion paper<sup>3</sup>, extending Hendrickson<sup>2</sup>, shows that under **Architecture B** (multivariate-normal differential correlation) the interaction lives in the treatment-state-dependent covariance  $\text{Cor}(B_i, BR_{it}) = c_{bm} \cdot \phi(t_{sd})$ , which carryover erodes directly: at a half-life of one week, power falls by up to roughly 31% in relative terms, strongly design-dependent (largest under OL+BDC, negligible under CO). Under the alternative Architecture A (direct mean moderation), the corresponding loss is only 1-3%; that architecture, and the third combined Architecture C, are outside the scope of this manuscript, which restricts to Architecture B, the architecture in which the carryover problem is largest and in which Hendrickson’s original results were obtained. That companion work, however, fixed the decay shape at exponential and evaluated a single analysis-model specification throughout, leaving open which analyst-side representation of residual exposure is most robust when the true decay shape is uncertain.

### 1.3 This paper

[4]

**We compare two analyst-side exposure representations – exposure-weighted (M2) and lagged-treatment (M3) – across decay forms under Architecture B.**

- M2 commits to a parametric decay and half-life; M3 estimates carryover non-parametrically via a lagged term.
- Each is crossed against exponential and Weibull DGP decay forms.
- The question is whether one specification is recommendable or the choice is design-dependent.

*rgt to complete.*

We evaluate two analyst-side representations of drug exposure in the linear predictor: an exposure-weighted continuous indicator that commits to a parametric decay form and half-life (**M2**, the current default), and the binary on-drug indicator augmented by a lagged just-off-drug term that estimates carryover non-parametrically (**M3**, the classical crossover device of Jones and Kenward<sup>4</sup>). We cross these against the exponential and Weibull data-generating decay forms and ask whether a single specification can be recommended or whether the choice is design-dependent. The study follows the ADEMP reporting framework of Morris, White, and Crowther<sup>1</sup>, with every performance measure accompanied by its MCSE. Section 3 gives the simulation design in ADEMP order; Section 4 the results; Section 5 the practical implications for trialists.

## 2 Simulation design

### 2.1 Aims

[5]

**The aim is to identify, under Architecture B, which of M2 and M3 is the more robust carryover-mitigation strategy across DGP decay forms and trial designs.**

- Quantify the power cost of analysis-model carryover mis-specification under Architecture B.
- Identify which of M2 and M3 is more robust across the DGP decay-form grid and across trial designs.
- Verify whether the ranking survives half-life mis-specification and other single-axis perturbations (Tier 2).

*rgt to complete.*

The aim of this study is to quantify, under Architecture B, the cost of analysis-model carryover mis-specification in aggregated N-of-1 and crossover trials, and to identify which of the two carryover-mitigation strategies, M2 or M3, is more robust across the DGP decay-form grid and across trial designs. A secondary aim, addressed in the Tier 2 sensitivity blocks, is to verify that the ranking survives perturbation of autocorrelation strength, dropout, half-life mis-specification, and interaction effect size.

## 2.2 Data-generating mechanism

[6]

**This manuscript retains only Architecture B (MVN differential correlation); the biomarker and biological response are jointly MVN with correlation decaying off-drug.**

- $\text{Cor}(B_i, BR_{it}) = c_{bm} \phi(t_{sd})$ , where  $\phi$  is the decay form.
- Architecture A (mean moderation) and the combined Architecture C are out of scope here; see report.Rmd for the cross-architecture comparison.
- Two decay forms for  $\phi$  are evaluated: exponential and Weibull. The linear form is dropped from this manuscript's scope.

*rgt to complete.*

In Architecture B, the biomarker and biological response are jointly drawn from a multivariate normal with treatment-state-dependent correlation  $\text{Cor}(B_i, BR_{it}) = c_{bm} \cdot \phi(t_{sd})$ , where  $\phi$  is the decay form and  $t_{sd}$  is time since drug discontinuation. Architecture A (direct mean moderation) and the combined Architecture C are outside the scope of this manuscript; the full grid comparison across architectures is reported in the companion report.Rmd / report-slim.Rmd manuscripts. In both architectures the decay function  $\phi : \mathbb{R}_+ \rightarrow [0, 1]$  maps  $t_{sd}$  to a residual-effect multiplier. We evaluate two forms here, each parameterized to match a common half-life  $t_{1/2}$ ; the linear form used in the full-scope manuscripts is dropped from this narrower evaluation.

### 2.2.1 Exponential decay

$$\phi_{\text{exp}}(t_{sd}) = \exp(-\lambda t_{sd}), \quad \lambda = \frac{\ln 2}{t_{1/2}}$$

[7]

**The exponential form is the natural reference decay because it corresponds to first-order elimination kinetics.**

- Used throughout the companion manuscript and implemented in `implementations/tidyverse/R/functions`.
- Corresponds to classical pharmacokinetic first-order elimination.
- Retained as the reference against which the Weibull form is compared.

*rgt to complete.*

This is the form used throughout the companion manuscript and currently implemented in `implementations/tidyverse/R/functions.R`. It corresponds to first-order elimination kinetics and is the default assumption in classical pharmacokinetics; we retain it here as the natural reference against which the Weibull form is compared.

### 2.2.2 Weibull decay

$$\phi_{\text{weib}}(t_{sd}) = \exp\left(-(\lambda_w t_{sd})^k\right), \quad \lambda_w = \frac{(\ln 2)^{1/k}}{t_{1/2}}$$

[8]

**The Weibull shape parameter  $k$  controls whether the elimination tail is heavier or lighter than exponential.**

- $k < 1$ : slower elimination with a prolonged low-concentration tail (evaluated at  $k = 0.7$ ).
- $k > 1$ : accelerated off-drug clearance via saturable clearance or active-transporter kinetics (evaluated at  $k = 1.5$ ).
- $k = 1$  recovers the exponential form exactly.

*rgt to complete.*

A shape parameter  $k$  governs whether the elimination tail is heavier ( $k < 1$ , slow elimination) or lighter ( $k > 1$ , accelerated off-drug clearance) than the exponential, which is recovered exactly at  $k = 1$ . We report results at  $k = 0.7$ , representing a drug with a prolonged low-concentration tail, and at  $k = 1.5$ , which captures drugs with saturable clearance or active-transporter kinetics. Together with  $k = 1.0$  (numerically identical to the exponential form), this yields four decay-shape levels in the grid.

## 2.3 Estimand and analysis specifications

[9]

**The estimand is the on-drug biomarker-response correlation  $c_{bm}$ ; the null estimand  $c_{bm} = 0$  is the Type I error row.**

- Under Architecture B the interaction estimand is  $c_{bm}$ , dimensionless.
- Bias, empirical SE, and MSE are reported on the calibrated effect-size scale defined in `docs/19`.
- The null estimand for Type I verification is the  $c_{bm} = 0$  row of the grid.

*rgt to complete.*

Under Architecture B, the interaction estimand is the on-drug biomarker-response correlation  $c_{bm}$ , dimensionless. Bias and empirical standard error are reported on the calibrated effect-size scale (expected on-drug response difference per SD of biomarker at  $t_{1/2} = 0$ ) defined in `docs/19-mean-moderation-implementation-notes.tex`; power and Type I error are reported on the native test-statistic scale. The null estimand for Type I error verification is  $c_{bm} = 0$ , realized by the  $c_{bm} = 0$  row of the parameter grid.

The analysis model is `nlme::lme(Sx ~ bm * X + t, random = ~1 | ptID, correlation = corCAR1(form = ~ t | ptID))` for both specifications. They differ only in the drug exposure predictor  $X$ .

### 2.3.1 M2: exposure-weighted (continuous predictor)

[10]

**M2 uses a continuous exposure-weighted predictor that commits to a specific decay form and half-life, and is the current pmsimstats-ng default.**

- $X = \text{Dbc}_{it}$ : decays off-drug according to the analyst’s assumed form and half-life.
- Well-matched when the analyst’s assumed form coincides with the true DGP form; mis-specified otherwise.
- Sensitivity of M2 to assumed-versus-true mismatch is evaluated in Tier 2 block S2.

*rgt to complete.*

$X = \text{Dbc}_{it}$ , the continuous exposure-decayed predictor constructed to match the analyst’s assumed decay form and half-life. This is the current pmsimstats-ng default specification. At on-drug timepoints  $D_{bc,it} = 1$  regardless of the assumed half-life; at off-drug timepoints the predictor decays as  $D_{bc,it} = \exp(-(\ln 2/\hat{t}_{1/2}) t_{sd,it}) = (1/2)^{t_{sd,it}/\hat{t}_{1/2}}$ , losing exactly half its value for every  $\hat{t}_{1/2}$  weeks since discontinuation. When the analyst’s assumed form coincides with the true DGP form the predictor is well-matched; when it does not, the predictor is mis-specified in a way that Tier 2 block S2 is designed to probe. The half-life assumption shapes the predictor values handed to lme; the model formula is unchanged for any  $\hat{t}_{1/2}$ , so M2 is effectively a different model for each assumed half-life even though the formula is identical.

### 2.3.2 M3: lagged-treatment (estimated carryover term)

[11]

**M3 augments the binary indicator with a lagged-treatment term that estimates carryover magnitude without assuming a decay form, and is invariant to the analyst’s assumed half-life by construction.**

- $L_{it} = D_{i,t-1}(1 - D_{it})$  marks off-drug timepoints immediately following on-drug timepoints.
- Carryover magnitude is estimated from the data; no functional form or half-life assumption is required.
- Follows the classical crossover-trial framework of Jones and Kenward [-@jones2014crossover].

*rgt to complete.*

$X = D_{it}$  augmented with a lagged-treatment indicator  $L_{it} = D_{i,t-1}(1 - D_{it})$  that marks off-drug timepoints following an on-drug timepoint. The carryover magnitude is estimated from the data rather than assumed a priori, so the specification does not require the analyst to posit a functional decay form. This approach follows the classical crossover-trial framework surveyed in<sup>4</sup>. One structural feature bears on interpretation: the lagged term  $L_{it}$  enters the model only as a main-effect nuisance covariate, and the biomarker interaction tested under M3 remains  $\text{bm}:D_{it}$ , the same coefficient as under the binary on-drug specification (M1, not evaluated in this manuscript). M3 is best understood as M1 with a carryover nuisance control rather than as a fundamentally different interaction model; unlike M2, it uses no half-life parameter and is therefore invariant to the analyst’s assumed half-life by construction. This asymmetry – one specification indexed continuously by  $\hat{t}_{1/2}$ , the other not a member of that family at all – is the basis of the Tier 2 block S2 sensitivity contrast (§4.5).

Table 1: Principal simulation grid (Architecture B only). Cell 1 is the matched reference: exponential DGP analyzed under the current pmsimstats-ng default. Off-diagonal cells quantify the cost of DGP-analysis mis-specification.

| DGP decay             | M2 exposure-weighted | M3 lagged-treatment |
|-----------------------|----------------------|---------------------|
| Exponential           | cell 1 (matched)     | cell 2              |
| Weibull ( $k = 0.7$ ) | cell 3               | cell 4              |
| Weibull ( $k = 1.0$ ) | cell 5               | cell 6              |
| Weibull ( $k = 1.5$ ) | cell 7               | cell 8              |

## 2.4 Factorial grid and parameter settings

The principal comparison is a factorial of DGP decay form (four levels: exponential; Weibull at  $k \in \{0.7, 1.0, 1.5\}$ ) against analysis-model specification (M2, M3), under Architecture B only:

| Parameter                         | Values                                |
|-----------------------------------|---------------------------------------|
| Biomarker moderation ( $c_{bm}$ ) | 0.0, 0.30, 0.45                       |
| Carryover half-life ( $t_{1/2}$ ) | 0.0, 0.5, 1.0 weeks                   |
| Weibull shape ( $k$ )             | 0.7, 1.0, 1.5 (Weibull form only)     |
| Sample size ( $N$ )               | 35, 70                                |
| Trial design                      | CO, N-of-1 (Hybrid), OL+BDC           |
| DGP architecture                  | B (MVN differential correlation) only |
| Replicates per cell               | 500 (target); 50 for development      |

[12]

**Restricting to Architecture B and dropping linear decay and M1 shrinks the grid to 216 cells, one-tenth of the full-scope manuscript’s 540.**

- Four decay-shape levels (exponential; Weibull at  $k \in \{0.7, 1.0, 1.5\}$ ), architecture fixed at B.
- Grid =  $4 \times 3 \times 3 \times 2 \times 3 = 216$  cells (decay-shape  $\times$  half-life  $\times$  design  $\times$  sample size  $\times$  biomarker level).
- $c_{bm} = 0$  is the Type I error row;  $t_{1/2} = 0$  is the no-carryover, best-case power reference.

*rgt to complete.*

The 216-cell count follows the same nesting logic as the full-scope manuscript, with the architecture dimension fixed and the linear decay form dropped: the decay-shape and shape combinations number four (exponential; Weibull at  $k \in \{0.7, 1.0, 1.5\}$ ), and the grid is  $4 \times 3 \times 3 \times 2 \times 3 = 216$  cells (decay-shape  $\times$  half-life  $\times$  design  $\times$  sample size  $\times$  biomarker level, all at Architecture B). The  $c_{bm} = 0$  condition serves as the Type I error check, targeting the nominal 5% level. The  $t_{1/2} = 0$  condition removes carryover entirely and provides a best-case power reference at each  $(N, \text{design}, c_{bm})$  combination, against which the power cost of non-zero carryover can be directly read off.

## 2.5 Performance measures and computing environment

For each cell we report Power, Type I error (at  $c_{bm} = 0$ ), Bias and Empirical SE of the interaction estimate (on the calibrated scale), Coverage of the nominal 95% Wald CI, MSE, the non-convergence rate, and the MCSE of each measure, following Morris et al.<sup>1</sup>, Tables 5-6.

[13]

$n_{sim} = 500$  keeps MCSE on power below 0.022; simulations reuse the tidyverse implementation and RNG strategy of the full-scope manuscript, restricted to Architecture B, non-linear decay, and M2/M3.

- Production:  $n_{sim} = 500$ , MCSE  $\leq 0.022$  at  $p = 0.5$ .
- Common random numbers across M2 and M3 within each cell enable paired contrasts.
- Data are the same production run summarized in `analysis/data/02-grid-summary.rds` (540 cells), filtered here to Architecture B, non-linear decay forms, and M2/M3.

*rgt to complete.*

The replicate count  $n_{sim} = 500$  is chosen so that the Monte Carlo standard error on power at  $p = 0.5$  does not exceed 0.022 ( $\sqrt{0.5 \cdot 0.5 / 500} \approx 0.022$ ). This manuscript does not re-run any simulation; it filters the same production-scale grid summary used by the full-scope `report.Rmd` and `report-slim.Rmd` manuscripts (540 cells, `analysis/data/02-grid-summary.rds`, `implementations/tidyverse/` as the primary implementation, R 4.5.3, 128 minutes on eight cores, commit SHA 7018b59, 2026-04-15) down to the 216-cell slice defined by Architecture B, non-linear decay forms, and the M2/M3 specifications. Within each cell, the same simulated dataset is used to fit both remaining specifications (common random numbers), so the M2–M3 contrast is estimated with reduced Monte Carlo variance relative to independent draws. Tier 2 sensitivity blocks and the S6 cluster-robust recalibration are filtered from the same shared outputs (`analysis/data/02-sensitivity-summary.rds`, `analysis/scripts/carryover-sensitivity/output/diag-s6-cr2.rds`); their designs are already anchored at Architecture B. Filtered figures are produced by `03-render-figures-extra-slim.R`, `04-sensitivity-figures-extra-slim.R`, and `07-type1-cr2-figure-extra-slim.R`.

### 3 Results

[14]

All reported results derive from the same 500-replicate production run as the full-scope manuscripts, filtered to Architecture B, exponential/Weibull decay, and M2/M3.

- MCSE on power  $\leq 0.022$  at  $p = 0.5$ ; differences smaller than the relevant MCSE should not be over-interpreted.
- Coverage of nominal 95% Wald CIs is reported for the reference cell.
- Non-convergence rates are reported per cell in the tables that follow.

*rgt to complete.*

#### 3.1 Headline performance table

[15]

The headline table profiles the reference cell; bias, MSE, and coverage are calibrated to M2’s target and are not directly comparable to M3’s.

- Reference cell: Architecture B, exponential DGP, Hybrid design,  $N = 70$ ,  $c_{bm} = 0.45$ ; non-convergence zero throughout.
- Bias/MSE/coverage are scored against M2’s target  $\theta_{true} = -c_{bm}\sigma_{BR} = -3.6$ ; M3 estimates a different coefficient (`bm:Dit`).

Table 2: Headline performance at the reference cell (Architecture B, exponential DGP, Hybrid design,  $N = 70$ ,  $c_{bm} = 0.45$ ,  $n_{sim} = 500$ ). Power is the rejection rate of  $H_0 : \beta_{bm:X} = 0$  at  $\alpha = 0.05$ . Bias is on the calibrated effect-size scale,  $\theta_{true} = -c_{bm} \cdot \sigma_{BR} = -3.6$  (per docs/19 calibration). MCSEs are reported per Morris et al. (2019, Table 6).

| t1/2 | spec | Power (MCSE)  | Bias (MCSE)    | Empirical SE | Coverage (MCSE) | MSE (MCSE)    | Non-conv |
|------|------|---------------|----------------|--------------|-----------------|---------------|----------|
| 0.0  | M2   | 0.988 (0.005) | -0.029 (0.033) | 0.746        | 0.990 (0.004)   | 0.557 (0.034) | 0.000    |
| 0.0  | M3   | 0.990 (0.004) | -0.025 (0.033) | 0.746        | 0.990 (0.004)   | 0.557 (0.034) | 0.000    |
| 0.5  | M2   | 0.948 (0.010) | +0.072 (0.041) | 0.907        | 0.972 (0.007)   | 0.828 (0.057) | 0.000    |
| 0.5  | M3   | 0.920 (0.012) | +0.583 (0.036) | 0.812        | 0.938 (0.011)   | 1.000 (0.059) | 0.000    |
| 1.0  | M2   | 0.860 (0.016) | -0.081 (0.049) | 1.085        | 0.968 (0.008)   | 1.184 (0.077) | 0.000    |
| 1.0  | M3   | 0.770 (0.019) | +1.099 (0.038) | 0.853        | 0.818 (0.017)   | 1.935 (0.094) | 0.000    |

- M3’s apparent bias and under-coverage reflect an estimand mismatch, not an estimator defect; compare only power and Type I error across specifications.

*rgt to complete.*

### 3.2 Matched-decay baseline

[16]

**The matched-decay baseline shows M2 power declining from 0.988 to 0.860 as carryover grows, reproducing Hendrickson’s 0.82.**

- Matched cell (exp DGP, M2, Architecture B,  $N = 70$ ,  $c_{bm} = 0.45$ , Hybrid):  $0.988 \rightarrow 0.948 \rightarrow 0.860$  across  $t_{1/2} \in \{0, 0.5, 1.0\}$ .
- All 500 replicates converged at every  $t_{1/2}$ .
- The  $t_{1/2} = 1.0$  value matches Hendrickson’s reported 0.82, confirming pipeline fidelity.

*rgt to complete.*

Under the matched cell (exponential DGP, M2 exposure-weighted analysis, Architecture-B,  $N = 70$ ,  $c_{bm} = 0.45$ , Hybrid design), power declines from  $0.988 \pm 0.005$  at  $t_{1/2} = 0$  to  $0.948 \pm 0.010$  at  $t_{1/2} = 0.5$  to  $0.860 \pm 0.016$  at  $t_{1/2} = 1.0$  weeks (point estimate  $\pm$  MCSE,  $n_{sim} = 500$ ). All 500 replicates converged in the matched cell at every  $t_{1/2}$ . The value at  $t_{1/2} = 1.0$  is consistent with Hendrickson’s reported power of 0.82 for this cell<sup>2</sup>, confirming that the simulation pipeline reproduces the baseline result.

### 3.3 Power by analysis specification

[17]

**The two specifications coincide at zero carryover and diverge as it grows, but M2’s lead is confined to the two non-crossover designs.**

- Both are near-identical at  $t_{1/2} = 0$ .
- M2’s advantage emerges at  $t_{1/2} = 0.5$  and grows in the Hybrid and OL+BDC designs.
- Under the CO design M2 is dominated by M3 (Section 5).

*rgt to complete.*

The two analysis specifications yield near-identical power at  $t_{1/2} = 0$ , where there is no carryover to model, and diverge progressively as carryover half-life increases. M2 retains a

## Power vs carryover under matched exponential DGP

Architecture B,  $N = 70$ ,  $c_{bm} = 0.45$ , 500 reps/cell




Figure 1: Power versus carryover half-life under the exponential DGP, Architecture B, stratified by analysis specification. M2 retains a modest advantage over M3 in the Hybrid and OL+BDC designs at  $t_{1/2} = 1.0$  weeks, but is dominated by M3 under the crossover (CO) design.

modest but consistent advantage in the Hybrid and OL+BDC designs, where the advantage emerges at  $t_{1/2} = 0.5$  weeks and grows through  $t_{1/2} = 1.0$  weeks. This advantage does not appear under the crossover (CO) design, where M2 is in fact dominated (Section 5): in a crossover the within-subject correlation already carries most of the signal M2’s decaying predictor is designed to recover, so M2 has nothing left to gain and its down-weighting of off-drug points only hurts.

### 3.4 Decay-form mis-specification

[18]

**The heatmap separates two mis-specification directions: wrong analysis given the DGP, and wrong assumed DGP.**

- Matched reference (exp DGP  $\times$  M2): power  $0.860 \pm 0.016$ .
- Same row, M3 column: penalty of the lagged specification under an exponential DGP.
- M2 column, other rows: cost of assuming exponential when the truth is Weibull.

*rgt to complete.*

The heatmap cell at the intersection of the exponential DGP row and the M2 analysis column is the matched reference (power  $0.860 \pm 0.016$ ). The same row’s M3 column quantifies the power penalty of the lagged specification when the DGP is exponential; the M2 column’s other rows quantify the cost of assuming exponential decay in the analysis when the true DGP is Weibull. Even under the heaviest-tailed Weibull DGP ( $k = 0.7$ ), M2’s advantage over M3 narrows but does not reverse; the Discussion (Section 5) takes up the practical interpretation.

### 3.5 Type I error control and cluster-robust recalibration

[19]




Figure 2: Power at the critical cell ( $t_{1/2} = 1.0$  weeks, Architecture B, Hybrid design,  $N = 70$ ,  $c_{bm} = 0.45$ ) across DGP decay forms (rows) and analysis specifications (columns). Under the heavy-tailed Weibull DGP ( $k = 0.7$ ), M2's advantage over M3 narrows because its assumed exponential decay no longer matches truth. The fill scale is stretched to the observed power range (0.66-0.87), not the full  $[0,1]$  range, so that these cell-to-cell differences are visible; the absolute shading should not be compared to other figures in this manuscript.




Figure 3: Empirical Type I error rates under the null ( $c_{bm} = 0$ ), Architecture B, pooled across DGP decay forms and carryover half-lives, stratified by design and analysis specification. Boxplots cover the distribution across cells at  $N = 70$ ; the dashed reference line marks the nominal 0.05 level.

Table 3: Conservatism of the model-based interaction test and its cluster-robust correction, at the null cell ( $c_{bm} = 0$ , exponential DGP,  $t_{1/2} = 1.0$  weeks, Hybrid design,  $N = 70$ , Architecture B, 3,000 replicates).  $\kappa$  is the ratio of the mean standard error to the empirical standard deviation of the interaction estimate;  $\kappa > 1$  indicates a conservative test. Both specifications are conservative under the model-based test, most so the exposure-weighted M2 ( $\kappa = 1.10$ , Type I = 0.029); the CR2 bias-reduced cluster-robust standard error restores  $\kappa$  to near one and Type I to nominal (0.045 to 0.049) for both specifications.

| Specification | $\kappa$ model | Type I model | $\kappa$ CR2 | Type I CR2 |
|---------------|----------------|--------------|--------------|------------|
| M2            | 1.10           | 0.029        | 1.00         | 0.045      |
| M3            | 1.05           | 0.039        | 0.98         | 0.049      |

Table 4: Cluster-robust recalibration of M2 (Power columns) and of the M2–M3 ranking (last two columns), at the reference cell and four stress cells (500 replicates,  $c_{bm} = 0.45$ , Architecture B).  $\kappa$  is the ratio of the mean standard error to the empirical standard deviation of the interaction estimate for M2; a value above one is a conservative test. The M2–M3 columns are computed directly from the per-specification power at this cell under each inference, so they are a genuine check of this manuscript’s own comparison rather than the M2-vs-M1 contrast reported for the full-scope grid. CR2 df is the mean Satterthwaite degrees of freedom for the cluster-robust test on M2. Computed with `clubSandwich`.

| Cell                        | $\kappa$ model | $\kappa$ CR2 | Power model | Power CR2 | M2–M3 model | M2–M3 CR2 | CR2 df |
|-----------------------------|----------------|--------------|-------------|-----------|-------------|-----------|--------|
| Reference, $N = 70$         | 1.15           | 0.99         | 0.829       | 0.863     | +0.082      | +0.090    | 23     |
| Small $N = 35$              | 1.13           | 0.96         | 0.558       | 0.590     | +0.068      | +0.086    | 12     |
| High autocorr, $\rho = 0.9$ | 1.26           | 0.94         | 0.954       | 0.980     | +0.052      | +0.036    | 23     |
| Half-life mis-spec          | 1.14           | 1.00         | 0.759       | 0.795     | +0.012      | +0.022    | 23     |
| 30% MCAR dropout            | 0.99           | 0.85         | 0.148       | 0.126     | +0.014      | -0.000    | 5      |

**Type I error is conservative, reflecting the corCAR1 working-covariance conservatism rather than inflation.**

- Neither specification inflates the rate above nominal in any design.
- The companion calibration study traces the shortfall to a 6-10% over-estimated SE.
- Block S6’s CR2 (below) restores nominal size without changing the ranking.

*rgt to complete.*

Across the null cells, empirical Type I error is uniformly below the nominal 0.05 level for both specifications and all three designs. That shortfall is not sampling slack but the conservatism of the model-based interaction test that the companion calibration study<sup>5</sup> traces to the corCAR1 working covariance: its standard error is over-estimated by six to ten percent and the realized rejection rate falls correspondingly. The conservatism is uniform across specifications, which is why it leaves the comparison intact; the cluster-robust recalibration reported below (block S6) confirms that a CR2 cluster-robust standard error restores the rejection rate to nominal without disturbing the ranking. Table 3 makes the per-specification conservatism and its cluster-robust correction explicit at the null cell.

[20]

**CR2 restores nominal size and recovers two to five power points; the M2-vs-M3 ranking is held or widened at four of five S6 cells but collapses to a statistical tie under 30% dropout.**

- Model-based null rejection sits near 0.03-0.04 across cells and specs; CR2 restores it to

S6: type-I error at the null ( $c_{bm} = 0$ ), model-based vs CR2  
Dashed line = nominal 0.05; grey band = binomial 95% MC interval at 500 reps




Figure 4: Empirical Type I error at the null ( $c_{bm} = 0$ ) under model-based and CR2 cluster-robust inference, across the five S6 stress cells and the two specifications (Architecture B). The model-based test (open) is conservative (near 0.03); CR2 (filled) is restored to the nominal 0.05 (dashed), within the binomial Monte Carlo band (gray). 500 null replicates per cell.

a mean near nominal.

- $\kappa$  falls from 1.07-1.26 to about 1 under CR2 (computed for M2; M3's calibration ratio is comparable, Table reftab:tab-type1-cr2).
- The M2–M3 gap widens under CR2 at the reference, small- $N$ , and half-life-mis-spec cells, narrows slightly at high autocorrelation, and is effectively zero (within noise) under 30% dropout, consistent with the S3 dropout finding below.
- Validated only at Architecture B / Hybrid /  $N = 70$ ; do not extend the reading to the CO design.

*rgt to complete.*

Figure 4 shows the effect of recalibration on test size directly: the model-based null rejection rate sits below nominal across the five S6 cells and both specifications, while the CR2 test restores it to within the Monte Carlo band around 0.05. This is the precondition for the whole comparison – a test that is not correctly sized cannot be preferred on power alone. In the information-rich cells, the calibration ratio  $\kappa$  runs 1.07-1.26 under the model and is restored to about 1 under CR2, and the correctly-sized test recovers two to five points of power. Table 4 reports the M2–M3 gap – this manuscript's own comparison, computed directly from the per-specification power at each S6 cell rather than reused from the M2-vs-M1 contrast reported in the full-scope grid. The gap is held or slightly widened under CR2 at the reference cell (+0.080  $\rightarrow$  +0.092), the small- $N$  cell (+0.068  $\rightarrow$  +0.086), and the half-life-mis-specification cell (+0.012  $\rightarrow$  +0.022); it narrows somewhat under high autocorrelation (+0.052  $\rightarrow$  +0.036, both favoring M2); and at 30% MCAR dropout it collapses from a small M2 lead under the model-based test (+0.014) to an essentially null difference under CR2 (−0.0002), i.e. a statistical tie. This is not a contradiction of the ranking so much as a restatement of the dropout finding in block S3 below: M2's advantage is driven by information in off-drug observations, and once dropout has removed most of that

information there is little left for either the model-based test or its cluster-robust correction to distinguish. This recalibration was validated only in the Architecture~B, Hybrid,  $N = 70$  cells and should not be extended to the crossover design, where M2 does not lead.

### 3.6 Sensitivity analyses

[21]

**Four Tier 2 blocks perturb one axis at a time around the reference cell.**

- Reference: Architecture B, Hybrid,  $N = 70$ ,  $c_{bm} = 0.45$ , exponential DGP,  $t_{1/2} = 1.0$ .
- Each block varies one axis with the others held fixed.
- Blocks are reported separately, not pooled with Tier 1 or each other.

*rgt to complete.*

#### 3.6.1 S1: within-factor autocorrelation

S1: Sensitivity to within-factor autocorrelation

Architecture B, Hybrid,  $N=70$ ,  $c_{bm}=0.45$ , exp DGP,  $t_{1/2}=1.0$




Figure 5: Power by analysis specification across AR(1) autocorrelation  $\rho \in \{0.5, 0.7, 0.8, 0.9\}$ , Architecture B.

[22]

**S1: autocorrelation lifts both specifications equally, preserving M2's  $\approx 0.06$  advantage at every  $\rho$ .**

- Power rises monotonically with  $\rho$  for both.
- M2's advantage holds near 0.06 across  $\rho \in [0.5, 0.9]$ .
- Autocorrelation and the specification act additively, not multiplicatively.

*rgt to complete.*

Both specifications gain power monotonically with  $\rho$ . At  $\rho = 0.5$ , M2 power is  $0.798 \pm 0.018$  versus M3  $0.738 \pm 0.020$ ; at  $\rho = 0.9$ , M2 reaches  $0.956 \pm 0.009$  versus M3  $0.910 \pm 0.013$  (point estimate  $\pm$  MCSE,  $n_{sim} = 500$ ). The M2 advantage over M3, approximately 0.06 in absolute power throughout, is preserved at every AR(1) level: the interaction between autocorrelation strength and analysis specification is additive rather than multiplicative.

## S2: Cost of analyst-truth half-life mismatch

A3 does not depend on assumed half-life; A2 does




Figure 6: Power of M2 and M3 as the analyst’s assumed half-life departs from the true DGP half-life. M3 is invariant to the assumed half-life by construction; M2’s power degrades with mis-specification.

### 3.6.2 S2: analyst-truth half-life mismatch

[23]

**S2: M2 keeps its advantage even when the assumed half-life is wrong by a factor of two or four.**

- M2 averages 0.868 over eight assumed/true half-life cells (range 0.766-0.964).
- M3 is invariant to the assumed half-life (average 0.840).
- M2’s gentle curvature makes it a defensible default under an approximately-known half-life.

*rgt to complete.*

M2 is markedly robust to analyst-truth half-life mismatch. Across the eight S2 cells (DGP  $t_{1/2} \in \{0.5, 1.0\}$  weeks crossed with analyst-assumed  $t_{1/2} \in \{0.25, 0.5, 1.0, 2.0\}$  weeks), M2 averages 0.868, ranging from 0.766 to 0.964; per-cell MCSE ranges from 0.008 to 0.019. M3 is invariant to the analyst-assumed half-life by construction, as it does not consume the parameter, and its power averages 0.840 across the same eight cells. The exposure-weighted predictor’s curvature is sufficiently gentle that mis-specification of the assumed half-life by a factor of two does not erode M2’s advantage over M3 at either DGP half-life, suggesting M2 remains defensible even when the true pharmacokinetic half-life is only approximately known at the design stage.

### 3.6.3 S3: dropout

[24]

**S3: dropout collapses power for both methods, and narrows M2’s advantage as off-drug observations are dropped.**

- At 30% MCAR both fall to about 0.12-0.15, effectively tied.
- The M2–M3 gap shrinks from  $\approx 0.10$  (no dropout) to  $\approx 0.03$  (30% MCAR).

### S3: Sensitivity to dropout

MCAR and MAR (severity-biased); reference cell as above




Figure 7: Power at dropout rates  $\in \{0, 0.1, 0.2, 0.3\}$  under MCAR and MAR mechanisms, Architecture B.

- M2's signal lives in off-drug points, so dropping them erodes its edge; the advantage survives at dropout  $\leq 0.20$ .

*rgt to complete.*

Both specifications lose power monotonically with dropout rate. The MCAR baseline ( $N = 70$ , dropout = 0) gives M2  $0.866 \pm 0.015$ , M3  $0.764 \pm 0.019$ ; at MCAR rate 0.30 these collapse to M2  $0.148 \pm 0.016$ , M3  $0.124 \pm 0.015$ . M2's advantage narrows under heavy dropout: the absolute M2–M3 gap is approximately 0.10 at zero dropout, 0.07 at rate 0.10, 0.04 at rate 0.20, and 0.03 at rate 0.30 under MCAR, narrowing further under MAR at the highest rate ( $\approx 0.004$  at rate 0.30). The pattern is consistent with M2's signal source residing in the exposure-weighted contrast contributed by off-drug observations: when those observations are preferentially dropped, M2's advantage contracts toward M3's. At low-to-moderate dropout (rate  $\leq 0.20$ ) the M2 advantage is preserved.

#### 3.6.4 S4: biomarker-moderation effect-size curve

[25]

**S4: the power curve is sigmoidal; M2's advantage peaks at intermediate effect sizes and vanishes toward the ceiling.**

- Type I error at  $c_{bm} = 0$  is conservative (0.028-0.030).
- M2–M3 peaks at  $+0.102$  at  $c_{bm} = 0.45$ , shrinking to  $+0.042$  at  $c_{bm} = 0.60$ .
- The advantage matters most for moderate biomarker-treatment interactions.

*rgt to complete.*

The expected sigmoidal power curve is observed for both specifications. At  $c_{bm} = 0$ , empirical Type I error is below nominal at  $0.028 \pm 0.007$  (M2) and  $0.030 \pm 0.008$  (M3); power rises through the mid-range and saturates at  $c_{bm} \geq 0.45$ : at  $c_{bm} = 0.6$ , M2 reaches  $0.958 \pm 0.009$  and M3  $0.916 \pm 0.012$ . The M2 advantage is largest at intermediate effect sizes,  $+0.082$  absolute at  $c_{bm} = 0.30$  and  $+0.102$  at  $c_{bm} = 0.45$ , shrinking to  $+0.042$  at  $c_{bm} = 0.60$  where both specifications approach ceiling. The advantage is therefore most practically consequential

#### S4: Biomarker-moderation effect-size curve

Dashed reference line at nominal alpha = 0.05




Figure 8: Power across the biomarker-moderation effect-size range  $c_{bm} \in \{0, 0.10, 0.20, 0.30, 0.45, 0.60\}$  under Architecture B. The  $c_{bm} = 0$  point is the null for Type I verification.

in designs targeting moderate biomarker-treatment interactions.

### 3.7 A higher-precision check

[26]

**A high-precision paired-McNemar rerun finds M2 strictly ahead of M3 only at the highest-leverage cell (+7.2 points,  $p = 1.2 \times 10^{-7}$ ); elsewhere the two are operationally tied.**

- 24-cell OL+BDC rerun at 600 trials/cell, all converging.
- Paired McNemar (same trials): M2 beats M3 by 7.2 points at  $N = 70$ ,  $t_{1/2} = 1.0$ ; other three cells show  $|\Delta| \leq 1.5$  points,  $p \geq 0.18$ .
- Type I error across the null cells stayed in  $[0.003, 0.053]$ .

*rgt to complete.*

Because some Tier 1 gaps are small, a focused 24-cell grid (the OL+BDC design,  $N \in \{35, 70\}$ , true half-life  $\in \{0.5, 1.0\}$ ) was rerun at 600 trials per cell, all converging; because M2 and M3 see the *same* trials, they can be compared with a paired **McNemar test**, more sensitive than treating the specifications as independent. Type I error across the twelve null cells fell in  $[0.003, 0.053]$ , so both specifications hold nominal alpha at the explored  $N$  and DGP half-life combinations. Paired McNemar on M2 versus M3 finds M2 strictly dominant only at the high-leverage corner ( $N = 70$ ,  $t_{1/2}^{\text{DGP}} = 1.0$  weeks,  $\Delta = +7.2$  percentage points,  $p = 1.2 \times 10^{-7}$ ); in the other three cells the paired difference lies between  $-0.5$  and  $+1.5$  percentage points with McNemar  $p \geq 0.18$ , and the two specifications are operationally tied at this resolution. Resolving the M2-versus-M3 gap at the intermediate cells with statistical clarity would require approximately  $2\{, \}$ 000 replicates per cell at the current MCSE.

## 4 Discussion

### 4.1 The main result, and where it breaks down

[27]

**M2’s advantage over M3 is real but conditional: confined to the Hybrid and OL+BDC designs, and reversed under the crossover design.**

- Best cell (Hybrid): M2 beats M3 by  $\approx 10$  points ( $\approx 5$  MCSEs).
- Under CO, M2 is much worse (0.488 vs 0.830 for M3).
- Mechanism: the crossover’s within-subject correlation already carries the signal M2 targets, so M2 only loses there.

*rgt to complete.*

The headline is that M2’s advantage over M3 is real but **conditional**. At its best cell (Hybrid design,  $N = 70$ ,  $c_{bm} = 0.45$ ,  $t_{1/2} = 1.0$ ) M2 beats M3 by about 10 percentage points of power (0.860 vs 0.770) – about five MCSEs, so not luck. In the crossover design it is much *worse* (0.488 vs 0.830 for M3). There is a clean reason for the crossover failure: in a crossover the within-subject correlation already carries most of the signal that M2’s decaying predictor is designed to recover, so M2 has nothing left to gain and its down-weighting of off-drug points only hurts. The ranking “M2 best” therefore holds only in the two non-crossover designs studied here.

## 4.2 What the robustness checks add

[28]

**Where M2 leads, the lead survives all three robustness checks.**

- Decay shape: M2 never drops below M3 even when the assumed form is wrong, up to the heaviest-tailed Weibull.
- Assumed half-life: a two- to four-fold error moves M2 only from 0.83 to 0.77, never below M3.
- Autocorrelation: lifts both equally, so M2’s 5-6 point lead is  $\rho$ -independent.

*rgt to complete.*

Three checks probe whether M2’s win, where it exists, is fragile. **Decay shape (Fig. 2)**. Even when the analyst wrongly assumes exponential decay but the truth is Weibull, M2 still does not drop below M3, within the decay-shape range tested here. **Assumed half-life (Fig. 6)**. M3 does not use a half-life, so it is unaffected; M2 does, yet getting the half-life wrong by a factor of two or four moves its power only from about 0.83 to 0.77, never below M3. **Autocorrelation (Fig. 5)**. Stronger serial correlation raises both methods by about the same amount, so M2’s lead stays near 5-6 points regardless of correlation strength.

## 4.3 Dropout matters more than the method

[29]

**Dropout collapses power for both methods equally, so it dominates the specification choice.**

- At 30% dropout both fall to 0.12-0.25 and become indistinguishable.
- MAR (severity-related) dropout is slightly worse than MCAR, because the highest-signal patients leave first.
- Keeping patients in the trial buys more power than the choice between M2 and M3.

*rgt to complete.*

The most sobering result is about *missing data*, not carryover. When patients drop out (Fig. 7), power collapses for both methods together: at 30% dropout both fall to 0.12-0.25 and become indistinguishable. M2 cannot protect against the loss; it only organizes whatever data remain. Dropout related to severity (MAR) is slightly worse than purely random dropout (MCAR), because the sickest patients – who carry the most signal – leave first. The practical message is blunt: keeping patients in the trial buys far more power than the choice between M2 and M3.

## 4.4 Scope, related work, and limits

[30]

**This manuscript is a narrowed slice of the full-scope grid; its ranking should not be extrapolated to Architecture A, the binary specification M1, or linear decay without consulting the full-scope manuscript.**

- Full grid (report.Rmd / report-slim.Rmd) shows M2’s advantage does not extend to Architecture A, where M2 is marginally dominated throughout.
- M2 power 0.860 reproduces Hendrickson’s 0.82; the Jones-Kenward expectation that shape-free M3 beats M2 under mis-specification is not supported here.
- Two further limitations: a single parameter set, and one-factor-at-a-time sensitivity checks.

*rgt to complete.*

This manuscript deliberately narrows the full-scope grid to Architecture B, non-linear decay, and the M2/M3 contrast, to bring a 30-page evaluation down to a size a reviewer can read at one sitting. The full-scope companion (report.Rmd / report-slim.Rmd) shows that the M2 advantage documented here does not extend to Architecture A (direct mean moderation), where M2 is marginally dominated in every design, nor is the binary on-drug specification M1 dominated by M2 outside the Hybrid/OL+BDC-under-Architecture-B cells studied here; readers who need the cross-architecture or M1 comparison should consult that manuscript rather than extrapolate from this one. Within the narrower scope, our matched-decay (M2) power at the headline cell,  $0.860 \pm 0.016$ , reproduces Hendrickson’s<sup>2</sup> reported 0.82 for the same cell. The crossover tradition of Jones and Kenward<sup>4</sup> predicted that a shape-free lagged term (M3) should beat a parametric predictor (M2) when the decay shape is guessed wrong; our S2 result does not support this in the range we tested, because M2’s penalty for a wrong half-life is too small for M3 to overtake it. Two further limitations bound these conclusions. First, the numbers come from one parameter set (calibrated to a prazosin/PTSD trajectory), so the *absolute* power values are illustrative; the *ranking* is the portable part, and even it is conditional on the design as shown above. Second, the sensitivity checks vary one factor at a time, so they can miss interactions between factors (for example, high autocorrelation combined with heavy dropout); the one joint check reported in the full-scope manuscript (block S5) showed no surprises.

## 5 Conclusions

1. **M2 helps, but only conditionally.** The exposure-weighted predictor gives the highest power – about 10 points over M3 – *only* in the Hybrid and OL+BDC designs, and the lead survives wrong decay-shape and wrong half-life assumptions there. Under the crossover design M2 has no advantage and M3 is at least as good. Where M2 is used, report it with a CR2 cluster-robust standard error (Block S6).

2. **A guessed half-life is fine (where M2 is used).** Getting the half-life wrong by a factor of two to four barely moves M2's power, so a rough pharmacokinetic estimate suffices.
3. **Dropout beats everything.** Power falls sharply as dropout rises (about half lost at 20% dropout), the same for both methods. Preventing dropout matters more than the choice between M2 and M3.
4. **The test is slightly conservative, and fixable.** Both methods reject a little less than the nominal 5% under the null; the CR2 cluster-robust standard error restores the correct rate while recovering a few points of power. The M2-vs-M3 gap is held or widened under CR2 at four of the five S6 stress cells, but collapses to a statistical tie under 30% dropout, consistent with finding 3 above (Block S6).
5. **This is a narrowed slice.** Architecture A, the binary specification M1, and linear decay are outside this manuscript's scope; the full-scope grid in `report.Rmd` / `report-slim.Rmd` is the authority for those comparisons and should be consulted before generalizing beyond Architecture B.

## Notes and reproducibility

The authors declare no conflicts of interest and received no specific funding. This manuscript is a scope-narrowed derivative of `report-slim.Rmd`: it re-filters, but does not re-run, the same production simulation. Cell-level summaries are at `analysis/data/02-grid-summary.rds` and `02-sensitivity-summary.rds`; filtered figures are produced by `03-render-figures-extra-slim.R`, `04-sensitivity-figures-extra-slim.R`, and `07-type1-cr2-figure-extra-slim.R` under `analysis/scripts/carryover-sensitivity/`. No simulation runs during knitting. The production simulation ran under R 4.5.3 and the document is rendered under R 4.6.0, both pinned via `renv.lock`.

## 6 References

1. Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Statistics in Medicine*. 2019;38(11):2074-2102. doi:10.1002/sim.8086
2. Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: Statistical methods and practical considerations. *Frontiers in Digital Health*. 2020;2:13. doi:10.3389/fdgth.2020.00013
3. pmsimstats team. Two architectures for simulating biomarker-treatment interactions in aggregated N-of-1 trials. Published online 2026.
4. Jones B, Kenward MG. *Design and Analysis of Cross-over Trials*. 3rd ed. CRC Press; 2014.
5. pmsimstats team. Type-I calibration of fixed-effect interaction tests in linear mixed models: A staged diagnosis of working-covariance mis-specification and a cluster-robust remedy. Published online 2026.

---

Rendered on 2026-07-30 at 14:06 PDT.

Source: `~/prj/res/36-pmsimstats-ng/pmsimstats-ng/analysis/report/02-carryover-sensitivity/report-`