

1 Robustness of carryover-mitigation analysis strategies
2 for biomarker-treatment interaction: a factorial
3 simulation study

4 pmsimstats team

5 2026-06-11

6 **Contents**

7	1 Abstract	2
8	2 Introduction	3
9	3 Simulation design	6
10	3.1 Aims	7
11	3.2 Data-generating mechanisms	7
12	3.3 Estimands	10
13	3.4 Methods: analysis-model carryover specifications	11
14	3.5 Factorial grid	15
15	3.6 Parameter settings	16
16	3.7 Performance measures	16
17	3.8 Computing environment	17
18	4 Results	19
19	4.1 Headline performance table	19
20	4.2 Matched-decay baseline	20
21	4.3 Power by analysis specification	21
22	4.4 Decay-form mis-specification	22
23	4.5 Lasagna view of the full factorial	23
24	4.6 Type I error control	23
25	4.7 Sensitivity analyses	25
26	4.8 High-precision rerun confirmation	31
27	4.9 Sensitivity block S6: cluster-robust recalibration	32
28	5 Discussion	34
29	5.1 The principal finding	34
30	5.2 Robustness to decay-form variation	35
31	5.3 Robustness to half-life mis-specification	36
32	5.4 Autocorrelation sensitivity	37
33	5.5 Dropout and the limits of specification refinement	37

34	5.6	Architecture dependence	38
35	5.7	The latent-subtype question	39
36	5.8	Comparison with prior simulation studies	40
37	5.9	Limitations	41
38	6	Conclusions	43
39	7	Conflict of interest	44
40	8	Funding	44
41	9	Data availability statement	44
42	10	Reproducibility	45
43	11	References	46

44 1 Abstract

45 **Background.** Carryover effects in aggregated N-of-1 and crossover trials de-
46 grade the detectability of biomarker-treatment interactions. Three analysis-side
47 specifications acknowledge carryover (binary-treatment A1, exposure-weighted
48 continuous- D_{bc} A2, and lagged-treatment A3); each makes implicit assumptions
49 about the functional form of the residual drug-effect decay. The joint sensitivity
50 of the inference to mis-specification at the data-generating and analyst layers has
51 not previously been characterised.

52 **Methods.** We conduct an ADEMP-compliant factorial simulation study¹ cross-
53 ing three data-generating carryover decay forms (linear, exponential, Weibull)
54 against the three analysis-model specifications, under both mean-moderation (Ar-
55 chitecture~A) and multivariate-normal differential-correlation (Architecture~B)
56 biomarker data-generating processes. The principal factorial covers 540 cells at
57 500 replicates per cell. Six sensitivity blocks (S1-S6) vary autocorrelation strength,
58 analyst-versus-DGP half-life mismatch, dropout rate, between-subject heterogene-
59 ity, and biomarker correlation strength, and recalibrate the test with a cluster-
60 robust standard error. All performance measures are reported with Monte Carlo
61 standard errors per Morris et al. (2019).

62 **Results.** Across the principal factorial and all sensitivity blocks, the exposure-
63 weighted A2 specification consistently produced the highest power, with an abso-
64 lute advantage of approximately 10 percentage points over A1 at the reference cell
65 (Architecture~B, Hybrid design, $N = 70$, $c_{bm} = 0.45$, $t_{1/2} = 1.0$ week; A2 = 0.860,
66 A1 = 0.766, A3 = 0.770; approximately five MCSE of separation). The ranking
67 $A2 > A1 \approx A3$ was preserved across designs, architectures, sample sizes, and DGP
68 carryover half-lives. A2 retained its advantage when the analyst-assumed half-life
69 departed from the true half-life by a factor of two or four. A high-precision rerun
70 at $N \in \{35, 70\}$ and $t_{1/2}^{DGP} \in \{0.5, 1.0\}$ confirmed the ranking at MCSE 0.020
71 with a paired McNemar test $p = 6.6 \times 10^{-17}$ at the highest-leverage cell. The

72 model-based test is mildly conservative, a property the companion calibration
 73 study characterises; a CR2 cluster-robust standard error (block S6) restores nom-
 74 inal size and recovers two to five percentage points of power without altering the
 75 ranking.

76 **Conclusions.** A2, reported with a cluster-robust standard error, is the recom-
 77 mended default for the analysis of biomarker-treatment interactions in aggregated
 78 N-of-1 trials under realistic carryover. A3 is a defensible fallback when the phar-
 79 macokinetic half-life is unknown but loses approximately 30 percent of A2’s ad-
 80 vantage at larger N . A1 (the strict-RM-ANOVA analogue) is dominated across
 81 the entire grid.

82 2 Introduction

83 [1]

84 **The joint sensitivity of biomarker-interaction power to**
 85 **analysis-side carryover mis-specification has not previously**
 86 **been characterised in the within-subject setting.**

- 87 • Residual drug effect after the on-drug phase degrades the
- 88 biomarker-treatment interaction test.
- 89 • The practical magnitude of the loss and its recoverability are the
- 90 motivating questions.
- 91 • No systematic evaluation of this question exists for the within-
- 92 subject biomarker-interaction setting.

93 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
 94 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
 95 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
 96 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
 97 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
 98 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
 99 mollit anim id est laborum.

100 *We begin with a question that has received surprisingly little systematic atten-*
 101 *tion in the biomarker-moderated N-of-1 trial literature. When residual drug*
 102 *effect persists after the on-drug phase, any analysis that ignores carryover*
 103 *or mis-specifies its functional form will, in principle, lose power to detect*
 104 *the biomarker-treatment interaction. But how large, in practice, is that loss,*
 105 *and can it be recovered by a better-chosen analysis specification? To the best*
 106 *of our knowledge, no systematic evaluation of this question exists for the*
 107 *within-subject biomarker-interaction setting.*

108 [2]

109 **The companion manuscript shows Architecture B loses 40–60**
 110 **percentage points of power under carryover while Architec-**
 111 **ture A loses only 8–13, but held the decay form and analysis**
 112 **specification fixed.**

- 113 • Hendrickson (2020) established the biomarker-interaction frame-
114 work for aggregated N-of-1 trials with matched exponential de-
115 cay.
- 116 • Companion paper (Architecture A vs. B): Architecture B loses
117 40–60 pp, Architecture A loses 8–13 pp under one week of carry-
118 over.
- 119 • Joint sensitivity to DGP-analyst mis-specification remained un-
120 characterised.

121 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
122 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
123 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
124 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
125 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
126 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
127 mollit anim id est laborum.

128 *The companion manuscript at [analysis/report/01-dgp-mean-moderation-vs-mvn/](#)*
129 *establishes one part of the picture. Hendrickson² introduced the biomarker-*
130 *by-treatment interaction framework for aggregated N-of-1 trials and reported*
131 *power under a matched exponential decay model. The companion paper*
132 *extended that work by comparing two data-generating architectures (direct*
133 *mean moderation under Architecture A, and multivariate-normal differential*
134 *correlation under Architecture B) and found that the power consequences of*
135 *carryover differ qualitatively between them. Under Architecture B, carryover*
136 *at a half-life of one week reduces power by 40 to 60 percentage points from*
137 *the no-carryover baseline; under Architecture A the corresponding loss is 8*
138 *to 13 percentage points. That prior work, however, held the decay form fixed*
139 *at exponential and evaluated a single analysis-model specification throughout.*
140 *Unfortunately, the joint sensitivity of the inference to mis-specification at*
141 *both the data-generating and analyst layers remained uncharacterised.*

142 [3]

143 **The analyst must choose a decay-form assumption even when**
144 **the true pharmacokinetic decay is unknown; which specifica-**
145 **tion is most robust to that choice has not been evaluated.**

- 146 • The carryover decay form is specified at the design stage, often
147 without precise pharmacokinetic information.
- 148 • The analyst typically defaults to a convenient parametric form
149 or ignores carryover entirely.
- 150 • Which analysis-side specification holds up best under decay-form
151 mis-specification is the central question.

152 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
153 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
154 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
155 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
156 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur

157 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
158 mollit anim id est laborum.

159 *The present paper addresses that gap directly. The trial designer must specify*
160 *a carryover decay form, explicitly or implicitly, and the analyst must choose*
161 *how to represent residual drug exposure in the linear predictor. In practice,*
162 *the true pharmacokinetic decay is rarely known with precision at the design*
163 *stage: the analyst typically defaults to a convenient parametric form or ig-*
164 *nores carryover entirely. We have not, to the best of our knowledge, seen a*
165 *systematic evaluation of which analysis-side specification holds up best when*
166 *the assumed decay form departs from truth, and it is that question we shall*
167 *take up here.*

168 [4]

169 **The study crosses three DGP decay forms against three**
170 **analysis specifications under both architectures to establish**
171 **whether robustness rankings are architecture-dependent.**

- 172 • DGP decay forms: linear, exponential, Weibull.
- 173 • Analysis specifications: A1 (binary), A2 (exposure-weighted,
174 pmsimstats-ng default), A3 (lagged-treatment, Jones-Kenward).
- 175 • Full grid run under both Architecture A and B to test whether
176 rankings transfer across DGP mechanisms.

177 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
178 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
179 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
180 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
181 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
182 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
183 mollit anim id est laborum.

184 *To address this, we perform a factorial simulation study crossing three DGP*
185 *decay forms (linear, exponential, and Weibull) against three analysis-model*
186 *carryover specifications: a binary on-drug indicator that ignores carryover*
187 *altogether (A1), an exposure-weighted continuous predictor that commits to*
188 *a specific decay form and half-life (A2, the current pmsimstats-ng default),*
189 *and a lagged-treatment indicator that estimates carryover magnitude with-*
190 *out positing a functional form (A3, following the classical crossover-trial*
191 *approach surveyed by Jones and Kenward³). The full grid is run under both*
192 *DGP architectures to establish whether robustness rankings are preserved*
193 *across data-generating mechanisms or whether the preferred specification is*
194 *architecture-dependent.*

195 [5]

196 **The study follows the ADEMP framework of Morris et al.**
197 **(2019) and reports every performance measure with its**
198 **Monte Carlo standard error.**

- 199 • ADEMP structure: Aims, Data-generating mechanisms, Esti-

- mands, Methods, Performance measures.
- MC SE reported alongside every point estimate per Morris et al. (2019).
 - §3 design; §4 results; §5 practical implications for trialists.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The study is organized and reported following the ADEMP framework of Morris, White, and Crowther¹, which structures simulation studies around aims, data-generating mechanisms, estimands, methods, and performance measures. Each performance measure is reported per cell alongside its Monte Carlo standard error, as Morris et al. recommend. We proceed in §3 through the simulation design in ADEMP order; results are presented in §4; and §5 takes up the practical implications for trialists choosing among the three carryover specifications.

3 Simulation design

[6]

This section organises the study under ADEMP: DGP (§3.1), estimands (§3.2), methods (§3.3), factorial grid (§3.4), parameters and performance measures (§3.5), and computing environment (§3.6).

- ADEMP provides the structural scaffold for the simulation design description.
- Reproducibility metadata (R version, packages, RNG seed, elapsed time) are in §3.6.
- §3.4 defines the principal 3×3 factorial and its design rationale.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

This section organises the study under the ADEMP labels of Morris et al.¹. We begin in §3.1 with the data-generating mechanisms, then turn in §3.2 to the estimands, §3.3 to the methods (analysis-model specifications), §3.4 to the factorial grid, and §3.5 to the parameter settings and performance measures.

241 *Reproducibility metadata, including the R version, key package versions, the*
 242 *RNG seed strategy, and elapsed time, are recorded in §3.6 (Computing envi-*
 243 *ronment).*

244 3.1 Aims

245 [7]

246 **The primary aim is to identify the most robust analysis spec-**
 247 **ification across the DGP decay-form grid; the secondary aim**
 248 **is to verify that rankings transfer across architectures.**

- 249 • Primary: quantify the cost of carryover mis-specification and
 250 identify the most robust analysis strategy.
- 251 • Secondary: verify that robustness rankings are architecture-
 252 invariant (Architecture A vs. B).
- 253 • Architecture-invariance would justify a single default recommen-
 254 dation regardless of assumed biomarker biology.

255 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-*
 256 *mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad*
 257 *minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip*
 258 *ex ea commodo consequat. Duis aute irure dolor in reprehenderit in*
 259 *voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur*
 260 *sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt*
 261 *mollit anim id est laborum.*

262 *The primary aim of this study is to quantify the cost of analysis-model car-*
 263 *ryover mis-specification in aggregated N-of-1 and crossover trials, and to*
 264 *identify which mitigation strategy is most robust across the DGP decay-form*
 265 *grid. A secondary aim is to verify that the robustness rankings transfer across*
 266 *the two data-generating-process architectures (mean moderation versus differ-*
 267 *ential correlation), so that a practitioner can be reasonably confident that the*
 268 *recommendation does not depend on which biomarker biology is assumed.*

269 3.2 Data-generating mechanisms

270 [8]

271 **Both architectures from the companion manuscript are re-**
 272 **tained, differing only in how the decay function ϕ enters the**
 273 **DGP.**

- 274 • Architecture A: ϕ enters the mean structure via effective exposure
 275 D_{it}^* .
- 276 • Architecture B: ϕ enters the biomarker-response correlation
 277 $\text{Cor}(B_i, BR_{it}) = c_{bm} \cdot \phi(t_{sd})$.
- 278 • Three decay forms for ϕ are evaluated: linear, exponential, and
 279 Weibull.

280 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
 281 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
 282 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
 283 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
 284 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
 285 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
 286 mollit anim id est laborum.

287 *We retain both DGP architectures from the companion manuscript. In Ar-*
 288 *chitecture A (direct mean moderation), the biomarker B_i enters the mean*
 289 *structure as $Y_{it} = \mu_0 + \beta_1 D_{it}^* (1 + \beta_{bm} B_i) + \epsilon_{it}$, where D_{it}^* is the effective ex-*
 290 *posure at timepoint t under the chosen decay form. In Architecture B (MVN*
 291 *differential correlation), the biomarker and biological response are jointly*
 292 *drawn from a multivariate normal with treatment-state-dependent correlation*
 293 *$\text{Cor}(B_i, BR_{it}) = c_{bm} \cdot \phi(t_{sd})$, where ϕ is the decay form.*

294 In both architectures the decay function $\phi : \mathbb{R}_+ \rightarrow [0, 1]$ maps time since drug
 295 discontinuation t_{sd} to a residual-effect multiplier. We evaluate three forms, each
 296 parameterised to match a common half-life $t_{1/2}$:

297 3.2.1 Linear decay

$$\phi_{\text{lin}}(t_{sd}) = \max \left(0, 1 - \frac{t_{sd}}{2t_{1/2}} \right)$$

298 [9]

299 **The linear form is biologically implausible but pragmatically**
 300 **defensible when only a complete-washout time is known.**

- 301 • Residual effect reaches zero exactly at $t_{sd} = 2t_{1/2}$, unlike the
- 302 exponential tail.
- 303 • Inconsistent with first-order elimination kinetics.
- 304 • Represents the analyst’s assumption that carryover has fully
- 305 cleared by a known time horizon.

306 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
 307 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
 308 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
 309 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
 310 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
 311 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
 312 mollit anim id est laborum.

313 *Residual effect vanishes completely at $t_{sd} = 2t_{1/2}$. The linear form is bi-*
 314 *ologically implausible for drugs with first-order elimination kinetics, but it*
 315 *represents a defensible pragmatic assumption when the only available infor-*
 316 *mation is that carryover has fully worn off by some specified time X .*

3.2.2 Exponential decay

$$\phi_{\text{exp}}(t_{sd}) = \exp(-\lambda t_{sd}), \quad \lambda = \frac{\ln 2}{t_{1/2}}$$

[10]

The exponential form is the natural reference decay because it corresponds to first-order elimination kinetics.

- Used throughout the companion manuscript and implemented in `implementations/tidyverse/R/functions.R`.
- Corresponds to classical pharmacokinetic first-order elimination.
- Retained as the reference against which linear and Weibull forms are compared.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

This is the form used throughout the companion manuscript and currently implemented in `implementations/tidyverse/R/functions.R`. It corresponds to first-order elimination kinetics and is the default assumption in classical pharmacokinetics; we retain it here as the natural reference against which the linear and Weibull forms are compared.

3.2.3 Weibull decay

$$\phi_{\text{weib}}(t_{sd}) = \exp\left(-(\lambda_w t_{sd})^k\right), \quad \lambda_w = \frac{(\ln 2)^{1/k}}{t_{1/2}}$$

[11]

The Weibull shape parameter k controls whether the elimination tail is heavier or lighter than exponential.

- $k < 1$: slower elimination with a prolonged low-concentration tail (evaluated at $k = 0.7$).
- $k > 1$: accelerated washout via saturable clearance or active-transporter kinetics (evaluated at $k = 1.5$).
- $k = 1$ recovers the exponential form exactly.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt

353

mollit anim id est laborum.

354

355

356

357

358

359

A shape parameter k governs whether the elimination tail is heavier ($k < 1$, slow elimination) or lighter ($k > 1$, accelerated washout) than the exponential, which is recovered exactly at $k = 1$. We report results at $k = 0.7$, representing a drug with a prolonged low-concentration tail, and at $k = 1.5$, which captures drugs with saturable clearance or active-transporter kinetics that clear more rapidly once the concentration falls below the saturation threshold.

360

[12]

361

362

363

All three decay forms are implemented via `carryover_decay()` in the `tidyverse` collection, with exponential as the backward-compatible default.

364

365

366

367

368

369

370

371

- `carryover_form` and `weibull_shape` fields are plumbed through `apply_carryover_to_component()`, `build_correlation_matrix()`, and `build_sigma_matrix()`.
- The `original-extended` collection supports only exponential; cross-implementation validation is therefore limited to exponential cells.
- Backward compatibility is preserved: absent `carryover_form` defaults to exponential.

372

373

374

375

376

377

378

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

379

380

381

382

383

384

385

386

387

388

389

Implementation note. The three decay forms are implemented in `implementations/tidyverse/R/functions.R` via the `carryover_decay()` helper function, with a `carryover_form` and `weibull_shape` pair of `model_param` fields plumbed through `apply_carryover_to_component()`, `build_correlation_matrix()`, and `build_sigma_matrix()`. Backward compatibility is preserved by defaulting to the exponential form when `carryover_form` is absent. We note that the `implementations/original-extended/` collection has not been updated and supports only the exponential form; cross-validation across implementations is therefore limited to exponential cells, which is sufficient for pipeline verification but not for the full decay-form comparison.

390

3.3 Estimands

391

392

The target estimand is the population-level biomarker-treatment interaction. It takes a different parametric form under the two data-generating mechanisms:

393

394

- **Architecture A.** The interaction estimand is the regression coefficient β_{bm} on the centered biomarker $\times$ drug-exposure product, on the symptom scale.

- **Architecture B.** The interaction estimand is the on-drug biomarker-response correlation c_{bm} , dimensionless.

[13]

The two architectures yield non-commensurable estimands that are harmonised to a common calibrated effect-size scale.

- Architecture A estimand: regression coefficient β_{bm} on centered biomarker $\times$ exposure product.
- Architecture B estimand: on-drug biomarker-response correlation c_{bm} , dimensionless.
- Both calibrated to a common effect-size scale (expected response difference per SD biomarker at $t_{1/2} = 0$); derivation in `docs/19-mean-moderation-implementation-notes.tex`.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The two estimands are not directly commensurable. To enable cross-architecture comparison of bias and power, we calibrate both architectures to a common population effect-size scale, defined as the expected difference in mean on-drug response per unit biomarker standard deviation at $t_{1/2} = 0$. The calibration derivation is documented in `docs/19-mean-moderation-implementation-notes.tex`. Bias and empirical standard error in the Results section are reported on this calibrated scale; power and Type I error are reported on the native test-statistic scale.

The null estimand for Type I error verification is $\beta_{bm} = 0$ under Architecture A and $c_{bm} = 0$ under Architecture B, each realised by the $c_{bm} = 0$ row of the parameter grid.

3.4 Methods: analysis-model carryover specifications

The analysis model is `nlme::lme(Sx ~ bm * X + t, random = ~1 | ptID, correlation = corCAR1(form = ~ t | ptID))` for all specifications. The three specifications differ in the drug exposure predictor X :

[14]

All three specifications share a common `nlme::lme` formula and `corCAR1` residual structure; they differ only in the drug exposure predictor X .

- Formula: `Sx ~ bm * X + t, random intercept ~1 | ptID`.
- Residual correlation: `corCAR1(form = ~ t | ptID)`.

- 435 • The three specifications (A1, A2, A3) differ only in how X is
436 constructed.

437 3.4.1 A1: binary-treatment (carryover ignored)

438 [15]

439 **A1 uses the binary on-drug indicator and ignores carryover**
440 **entirely, serving as the dominated reference specification.**

- 441 • $X = D_{it}$: takes value 1 on-drug, 0 off-drug.
442 • Makes no attempt to model residual drug effect during washout.
443 • Serves as the lower-power reference against which A2 and A3 are
444 compared.

445 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
446 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
447 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
448 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
449 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
450 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
451 mollit anim id est laborum.

452 *$X = D_{it}$, the on-drug indicator. This is the naive specification that disregards*
453 *any residual drug effect during washout; it serves as the reference against*
454 *which the power cost of ignoring carryover is measured.*

455 3.4.2 A2: exposure-weighted (matched-decay continuous predictor)

456 [16]

457 **A2 uses a continuous exposure-weighted predictor that com-**
458 **mits to a specific decay form and half-life, and is the current**
459 **pmsimstats-ng default.**

- 460 • $X = Dbc_{it}$: decays off-drug according to the analyst's assumed
461 form and half-life.
462 • Well-matched when the analyst's assumed form coincides with
463 the true DGP form; mis-specified otherwise.
464 • Sensitivity of A2 to assumed-versus-true mismatch is evaluated
465 in Tier 2 block S2.

466 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
467 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
468 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
469 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
470 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
471 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
472 mollit anim id est laborum.

473 $X = Dbc_{it}$, the continuous exposure-decayed predictor constructed to
 474 match the analyst's assumed decay form and half-life. This is the current
 475 *pmsimstats-ng* default specification. When the analyst's assumed form
 476 coincides with the true DGP form the predictor is well-matched; when it
 477 does not, the predictor is mis-specified in a way that the Tier 2 block S2 is
 478 designed to probe.

479 3.4.3 A3: lagged-treatment (estimated carryover term)

480 [17]

481 **A3 augments the binary indicator with a lagged-treatment**
 482 **term that estimates carryover magnitude without assuming**
 483 **a decay form.**

- 484 • $L_{it} = D_{i,t-1}(1 - D_{it})$ marks off-drug timepoints immediately fol-
 485 lowing on-drug timepoints.
- 486 • Carryover magnitude is estimated from the data; no functional
 487 form or half-life assumption is required.
- 488 • Follows the classical crossover-trial framework of Jones and Ken-
 489 ward [-@jones2014crossover].

490 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
 491 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
 492 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
 493 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
 494 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
 495 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
 496 mollit anim id est laborum.

497 $X = D_{it}$ augmented with a lagged-treatment indicator $L_{it} = D_{i,t-1}(1 - D_{it})$
 498 that marks off-drug timepoints following an on-drug timepoint. The carryover
 499 magnitude is estimated from the data rather than assumed a priori, so the
 500 specification does not require the analyst to posit a functional decay form.
 501 This approach follows the classical crossover-trial framework surveyed in³.

502 3.4.4 Construction of D_{bc} and the role of the analyst's assumed half- 503 life

504 [18]

505 D_{bc} halves every $\hat{t}_{1/2}$ weeks off-drug; the formula handed to
 506 **lme** is invariant to the analyst's half-life assumption.

- 507 • $D_{bc,it} = (1/2)^{t_{sd}/\hat{t}_{1/2}}$ at off-drug timepoints.
- 508 • The model formula $Sx \sim \text{bm} + t + Dbc + \text{bm}:Dbc$ is the same
 509 regardless of $\hat{t}_{1/2}$.
- 510 • The half-life assumption is encoded in the predictor values, not
 511 in the formula.

Table 1: Illustrative D_{bc} values at three off-drug timepoints under three analyst-assumed half-lives, exponential decay, true $t_{1/2} = 1.0$ week. On-drug timepoints have $D_{bc} = 1$ under all assumptions. Row values are $(1/2)^{t_{sd}/\hat{t}_{1/2}}$.

t_{sd} (wk)	$\hat{t}_{1/2} = 0.5$ wk	$\hat{t}_{1/2} = 1.0$ wk (truth)	$\hat{t}_{1/2} = 2.0$ wk
1	0.250	0.500	0.707
2	0.063	0.250	0.500
3	0.016	0.125	0.354

512 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
513 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
514 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
515 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
516 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
517 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
518 mollit anim id est laborum.

519 *We begin with an observation that is easy to overlook in the formal model*
520 *statement above: the analyst’s assumed half-life $\hat{t}_{1/2}$ does not appear as a*
521 *parameter in the `lme` call. It enters the analysis at an earlier stage, during*
522 *the construction of the D_{bc} predictor column, and is thereafter invisible to*
523 *the model-fitting routine.*

524 *At on-drug timepoints, $D_{bc,it} = 1$ regardless of the assumed half-life. At off-*
525 *drug timepoints, the predictor decays as a function of t_{sd} , the time in weeks*
526 *since drug discontinuation:*

$$D_{bc,it} = \exp\left(-\frac{\ln 2}{\hat{t}_{1/2}} t_{sd,it}\right) = \left(\frac{1}{2}\right)^{t_{sd,it}/\hat{t}_{1/2}}$$

527 *That is, D_{bc} loses exactly half its value for every $\hat{t}_{1/2}$ weeks of time since dis-*
528 *continuation. The half-life assumption shapes the predictor values handed to*
529 *`lme`; the model formula $Sx \sim bm + t + Dbc + bm:Dbc$ is the same regardless*
530 *of what $\hat{t}_{1/2}$ was assumed.*

531 To illustrate, consider a participant in a stylized crossover arm who is on drug
532 for three weekly measurement occasions and then enters a three-week washout.
533 Table~1 shows the D_{bc} values at the three off-drug timepoints under three analyst
534 assumptions, holding the true DGP half-life fixed at one week.

535 [19]

536 **The analyst’s assumed half-life shapes the D_{bc} predictor val-**
537 **ues handed to `lme` but is invisible to the model-fitting routine**
538 **itself.**

- 539 • Under-assumed half-life: off-drug timepoints receive near-zero
540 D_{bc} and contribute almost nothing to the `bm:Dbc` estimate.

Table 2: Principal 3×3 simulation grid. Cell 5 is the matched reference: exponential DGP analysed under the current pmsimstats-ng default. Off-diagonal cells quantify the cost of DGP-analysis mis-specification. The full grid is crossed with Architecture $\in \{A, B\}$.

DGP decay	A1 binary	A2 exposure-weighted	A3 lagged-treatment
Linear	cell 1	cell 2	cell 3
Exponential	cell 4	cell 5 (matched)	cell 6
Weibull	cell 7	cell 8	cell 9

- Over-assumed half-life: off-drug timepoints retain substantial weight, making D_{bc} resemble a graded on-drug indicator.
- A1 and A3 are invariant to $\hat{t}_{1/2}$ by construction; S2 contrasts a continuously indexed model family against two half-life-free reference specifications.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

When the analyst assumes $\hat{t}_{1/2} = 0.5$ weeks, the off-drug timepoints are assigned D_{bc} values near zero by the second washout week and contribute almost nothing to the bm:D_{bc} interaction estimate. When the analyst assumes $\hat{t}_{1/2} = 2.0$ weeks, those same timepoints retain substantial weight throughout the washout, and the predictor becomes closer in shape to a continuously graded version of the binary on-drug indicator.

Two consequences follow for interpreting the S2 sensitivity block. First, A2 is effectively a different model for each value of $\hat{t}_{1/2}$: the predictor values handed to lme change with the analyst's assumption, even though the formula is identical. Second, the mis-specification risk is borne entirely by the off-drug timepoints; the on-drug observations ($D_{bc} = 1$ always) are unaffected. Specifications A1 and A3 use the binary on-drug indicator D_b and the lagged indicator L respectively, neither of which involves a half-life parameter, so both are invariant to $\hat{t}_{1/2}$ by construction. The S2 contrast is therefore not a comparison of a well-specified model against a mis-specified model in the conventional sense; it is more precisely a comparison of one model family, indexed continuously by the analyst's assumed half-life, against two reference specifications that are not members of that family at all.

3.5 Factorial grid

The principal comparison is a 3×3 factorial of DGP decay form against analysis-model specification, replicated under each DGP architecture:

3.6 Parameter settings

Parameter	Values
Biomarker moderation (c_{bm} / β_{bm})	0.0, 0.30, 0.45
Carryover half-life ($t_{1/2}$)	0.0, 0.5, 1.0 weeks
Weibull shape (k)	0.7, 1.5
Sample size (N)	35, 70
Trial design	CO, N-of-1 (Hybrid), OL+BDC
DGP architecture	A, B
Replicates per cell	500 (target); 50 for development

[20]

Two anchor conditions in the parameter grid serve diagnostic roles: $c_{bm} = 0$ for Type I error and $t_{1/2} = 0$ for best-case power.

- $c_{bm} = 0$: null row for Type I error verification, targeting nominal 5%.
- $t_{1/2} = 0$: no-carryover reference providing best-case power at each $(N, \text{design}, c_{bm})$ cell.
- Non-zero $t_{1/2}$ rows directly quantify the power cost of carryover relative to the no-carryover reference.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The $c_{bm} = 0$ condition serves as the Type I error check, targeting the nominal 5% level. The $t_{1/2} = 0$ condition removes carryover entirely and provides a best-case power reference at each $(N, \text{design}, c_{bm})$ combination, against which the power cost of non-zero carryover can be directly read off.

3.7 Performance measures

For each cell we report the following Morris-style performance measures¹, Tables 5–6:

1. **Power** to reject $H_0 : \beta_{bm:X} = 0$ at $\alpha = 0.05$.
2. **Type I error** at the null estimand ($c_{bm} = 0$ row).
3. **Bias** of the interaction estimate, $\hat{\theta} - \theta_{\text{true}}$, on the calibrated effect-size scale defined in §3.2.
4. **Empirical SE** across replicates within each cell.
5. **Coverage** of the nominal 95% Wald confidence interval for the interaction coefficient, computed against the calibrated estimand ($\theta_{\text{true}} = -c_{bm} \cdot \sigma_{BR}$).
6. **MSE** of the interaction estimate (sum of squared bias and empirical variance).

7. **Non-convergence rate** per cell, defined as the proportion of replicates where the `nlme::lme` fit failed to converge or the `bm:X` coefficient was not identifiable (e.g. A1 on OL designs, where there are no off-drug observations).
8. **Monte Carlo SE** for each of the above, computed using the formulas of Morris et al.¹, Table 6 and reported alongside every point estimate in tables and figures.

Eight Morris-style performance measures are reported per cell, each accompanied by its Monte Carlo standard error.

- Power, Type I error, bias, empirical SE, coverage, MSE, non-convergence rate, and MC SE for each measure.
- Coverage is computed against the calibrated estimand $\theta_{true} = -c_{bm} \cdot \sigma_{BR}$.
- MC SEs use the formulas of Morris et al. (2019, Table 6).

[21]

$n_{sim} = 500$ is calibrated to keep MC SE on power below 0.022, with $n_{sim} = 50$ reserved for pipeline validation only.

- Production run: $n_{sim} = 500$, MC SE ≤ 0.022 at $p = 0.5$.
- Development run: $n_{sim} = 50$, MC SE ≈ 0.07 ; pipeline validation only.
- All reported results derive from the production-scale run.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The replicate count $n_{sim} = 500$ is chosen so that the Monte Carlo standard error on power at $p = 0.5$ does not exceed 0.022 ($\sqrt{0.5 \cdot 0.5 / 500} \approx 0.022$). The development scale ($n_{sim} = 50$, MC SE ≈ 0.07 at $p = 0.5$) is used only for pipeline validation; all results reported in the Results section derive from the production run at $n_{sim} = 500$.

3.8 Computing environment

[22]

The tidyverse collection is the only implementation supporting all three decay forms; common-random-number seeding enables paired contrasts across A1, A2, and A3.

- All three decay forms are supported only in `implementations/tidyverse/`; cross-implementation parity checks are limited to exponential cells.

- Within-cell common random numbers reduce MC variance for the A2–A1 and A3–A1 paired contrasts.
- Production run (540 cells, $n_{sim} = 500$): 128 minutes on eight cores; commit SHA 7018b59, 2026-04-15.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Simulations are orchestrated from `analysis/scripts/carryover-sensitivity/` and use `implementations/tidyverse/` as the primary implementation, which is the only collection in the repository that supports linear and Weibull decay forms. Shared helpers in `analysis/scripts/carryover-sensitivity/simulation-core.R` define design presets (OL, CO, Hybrid, OLBDC), multi-path data generation, long-form preparation with Db, Dbc, and L columns, and the three analysis-specification fitters.

Software. *R 4.5.3 with nlme 3.1, dplyr 1.1, tidyr 1.3, purrr 1.0, furrr 0.3, future 1.34, corpcor 1.6, MASS 7.3, Matrix 1.7, tibble 3.2, and data.table 1.16. The `implementations/tidyverse/` collection commit SHA at the time of the production run is recorded in each output `.rds` file.*

RNG seed strategy. *Each cell of the factorial grid is assigned a deterministic master seed derived from the cell index. Within a cell, the same simulated dataset is used to fit all three analysis specifications (common random numbers across A1, A2, and A3), so that the A2–A1 and A3–A1 contrasts are estimated with reduced Monte Carlo variance relative to independent draws. Seeds are independent across cells. Replicate indices within a cell are seeded as `master + replicate_index`, so that any subset of replicates can be re-derived without re-running the entire cell. The seed assignment is logged in the output `.rds` file alongside the remaining reproducibility metadata.*

Parallelism. *Parallel execution uses `furrr::future_map` with the `multicore` backend. Sigma matrices are cached per unique (design, $t_{1/2}$, decay form) tuple before the replicate loop begins, avoiding redundant matrix construction across workers.*

Elapsed time. *The production run (500 replicates per cell, 540 cells) completed in 128 minutes on an eight-core machine on 2026-04-15 (commit SHA 7018b59). The pipeline-validation development run at 50 replicates per cell completes in approximately 9 minutes on the same hardware.*

Pre-registration and reproducibility. *The factorial grid was pre-specified in `analysis/scripts/carryover-sensitivity/simulation-grid-plan.md` and committed before the production run was executed. The production run command is `Rscript analysis/scripts/carryover-sensitivity/01-run-factorial.R`; the Tier 2 sensitivity blocks are produced by `Rscript analysis/scripts/carryover-sensitivity/04-run-sensit`. Each output `.rds` file records the script name, run timestamp, master seed, dev and smoke flags, elapsed seconds, R version, and the*

692 *implementations/tidyverse/ commit SHA. The `renv.lock` at the*
693 *repository root pins all R-package versions. The manuscript Rmd at*
694 *analysis/report/02-carryover-sensitivity/report.Rmd does not*
695 *re-execute any simulation code; it loads the pre-computed checkpoints written*
696 *by the summarisation script.*

697 4 Results

698 [23]

699 **All reported results derive from the 540-cell production-scale**
700 **run at $n_{sim} = 500$; MC SE bounds apply throughout.**

- 701 • MC SE on power ≤ 0.022 at $p = 0.5$; differences smaller than the
- 702 relevant MC SE should not be over-interpreted.
- 703 • Coverage of nominal 95% Wald CIs is reported for the reference
- 704 cell and computed across the full Tier 1 grid.
- 705 • Non-convergence rates are reported per cell in the tables that
- 706 follow.

707 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
708 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
709 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
710 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
711 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
712 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
713 mollit anim id est laborum.

714 *We begin with a few words on the precision and scope of the estimates*
715 *that follow. All results derive from the production-scale factorial run (500*
716 *replicates per cell, 540 cells, computation time 128 minutes on eight cores;*
717 *commit SHA recorded in the output `.rds` metadata). Monte Carlo standard*
718 *errors on individual power estimates are at most $\sqrt{0.5 \cdot 0.5 / 500} \approx 0.022$*
719 *at $p = 0.5$ and are reported alongside each estimate; within-figure differ-*
720 *ences smaller than the relevant MC SE should not be over-interpreted.*
721 *Coverage of nominal 95% Wald confidence intervals is reported in the*
722 *headline table for the reference cell and is computed across the full Tier*
723 *1 grid in `analysis/data/02-grid-summary.rds` for downstream use.*
724 *Non-convergence rates per cell are reported in the tables below.*

725 4.1 Headline performance table

726 [24]

727 **The headline table provides the complete Morris-style per-**
728 **formance profile for the reference cell across all three speci-**
729 **fications and carryover levels.**

- 730 • Reference cell: Architecture B, exponential DGP, Hybrid design,
- 731 $N = 70$, $c_{bm} = 0.45$.

Table 3: Headline performance at the reference cell (Architecture B, exponential DGP, Hybrid design, $N = 70$, $c_{bm} = 0.45$, $n_{sim} = 500$). Power is the rejection rate of $H_0 : \beta_{bm:X} = 0$ at $\alpha = 0.05$. Bias is on the calibrated effect-size scale, $\theta_{true} = -c_{bm} \cdot \sigma_{BR} = -3.6$ (per docs/19 calibration). Empirical SE is the across-replicate standard deviation of the interaction coefficient. MSE is the sum of squared bias and empirical variance. Monte Carlo SEs are reported per Morris et al. (2019, Table 6). Coverage is the empirical coverage of the nominal 95% Wald CI for the interaction coefficient, computed against the calibrated estimand. Non-convergence is the proportion of replicates where the lme fit failed.

t1/2	spec	Power (MC SE)	Bias (MC SE)	Empirical SE	Coverage (MC SE)	MSE (MC SE)	Non-c
0.0	A1	0.988 (0.005)	-0.029 (0.033)	0.746	0.990 (0.004)	0.557 (0.034)	0.000
0.0	A2	0.988 (0.005)	-0.029 (0.033)	0.746	0.990 (0.004)	0.557 (0.034)	0.000
0.0	A3	0.990 (0.004)	-0.025 (0.033)	0.746	0.990 (0.004)	0.557 (0.034)	0.000
0.5	A1	0.918 (0.012)	+0.586 (0.036)	0.807	0.938 (0.011)	0.995 (0.058)	0.000
0.5	A2	0.948 (0.010)	+0.072 (0.041)	0.907	0.972 (0.007)	0.828 (0.057)	0.000
0.5	A3	0.920 (0.012)	+0.583 (0.036)	0.812	0.938 (0.011)	1.000 (0.059)	0.000
1.0	A1	0.766 (0.019)	+1.093 (0.038)	0.847	0.828 (0.017)	1.911 (0.093)	0.000
1.0	A2	0.860 (0.016)	-0.081 (0.049)	1.085	0.968 (0.008)	1.184 (0.077)	0.000
1.0	A3	0.770 (0.019)	+1.099 (0.038)	0.853	0.818 (0.017)	1.935 (0.094)	0.000

- Non-convergence is zero at every reference cell.
- MC SEs are reported as $\pm$ beside each point estimate throughout.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Table 3 reports the Morris-style performance measures at the reference cell (Architecture B, exponential DGP, Hybrid design, $N = 70$, $c_{bm} = 0.45$) across the three analysis specifications and three carryover half-lives. Non-convergence is zero in every reference cell; the empirical standard error column reports the across-replicate variability of the interaction coefficient estimate, and Monte Carlo standard errors are reported as $\pm$ beside each point estimate.

4.2 Matched-decay baseline

[25]

The matched-decay baseline confirms pipeline fidelity and shows a power decline from 0.988 at no carryover to 0.860 at $t_{1/2} = 1.0$ weeks.

- Matched cell (exponential DGP, A2, Architecture B, $N = 70$,

754 $c_{bm} = 0.45$, Hybrid): power $0.988 \rightarrow 0.948 \rightarrow 0.860$ across $t_{1/2} \in$
755 $\{0, 0.5, 1.0\}$ weeks.
756 • All 500 replicates converged at every $t_{1/2}$.
757 • Consistent with companion manuscript’s Architecture B result
758 (0.82 at $t_{1/2} = 1.0$ weeks).

759 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
760 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
761 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
762 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
763 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
764 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
765 mollit anim id est laborum.

766 *In this section we shall establish the baseline before comparing specifications.*
767 *Under the matched cell (exponential DGP, A2 exposure-weighted analysis,*
768 *Architecture~B, $N = 70$, $c_{bm} = 0.45$, Hybrid design), power declines from*
769 *0.988 ± 0.005 at $t_{1/2} = 0$ to 0.948 ± 0.010 at $t_{1/2} = 0.5$ to 0.860 ± 0.016 at*
770 *$t_{1/2} = 1.0$ weeks (point estimate $\pm$ Monte Carlo SE, $n_{sim} = 500$; MC SEs*
771 *computed as $\sqrt{\hat{p}(1-\hat{p})/n_{sim}}$ per Morris et al.¹, Table 6). All 500 replicates*
772 *converged in the matched cell at every $t_{1/2}$. These values are consistent*
773 *with the Architecture~B numbers from the companion manuscript (0.82 at*
774 *$t_{1/2} = 1.0$), confirming that the simulation pipeline reproduces the baseline*
775 *result and that the calibration is internally coherent.*

776 4.3 Power by analysis specification




Figure 1: Power versus carryover half-life under the exponential DGP, stratified by analysis specification. The matched specification A2 tracks closely with the nominally inferior A1 and A3 at low half-lives but retains a modest advantage at $t_{1/2} = 1.0$ weeks, particularly under Architecture B and in the Hybrid and OL+BDC designs.

[26]

Specifications diverge progressively with carryover half-life; A2's advantage over A1 and A3 emerges at $t_{1/2} = 0.5$ weeks and grows at $t_{1/2} = 1.0$ weeks.

- All three specifications are near-identical at $t_{1/2} = 0$; divergence is proportional to carryover strength.
- A1 and A3 are effectively indistinguishable across the exponential-DGP grid.
- A2's advantage is consistent across Hybrid and OL+BDC design strata, not an artefact of the reference cell.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The three analysis specifications yield near-identical power at $t_{1/2} = 0$, where there is no carryover to model, and diverge progressively as carryover half-life increases. The matched specification A2 retains a modest but consistent advantage under Architecture~B and the Hybrid and OL+BDC designs; A1 and A3 are effectively indistinguishable across the exponential-DGP grid. The figure makes clear that the A2 advantage is not an artefact of the reference cell: it emerges at $t_{1/2} = 0.5$ weeks and grows through $t_{1/2} = 1.0$ weeks in each design stratum.

4.4 Decay-form mis-specification

[27]

The heatmap exposes two directions of mis-specification: wrong analysis specification given the true DGP, and wrong DGP assumed by the analyst.

- Matched reference: exponential DGP $\times$ A2 analysis, power 0.860 ± 0.016 .
- Same-row off-diagonal cells: power penalty of A1 or A3 when the true DGP is exponential.
- A2-column off-diagonal cells: cost of assuming exponential when the true DGP is linear or Weibull.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur




Figure 2: Power at the critical cell ($t_{1/2} = 1.0$ weeks, Architecture B, Hybrid design, $N = 70$, $c_{bm} = 0.45$) across DGP decay forms (rows) and analysis specifications (columns). Under the Weibull heavy-tailed DGP ($k = 0.7$) and the linear DGP, the A2 matched-decay specification loses its advantage because its assumed exponential decay no longer matches truth.

sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The heatmap cell at the intersection of the exponential DGP row and the A2 analysis column is the matched reference (power 0.860 ± 0.016). Cells in the same row but other columns quantify the power penalty of analysing under the wrong specification when the DGP is exponential; cells in the A2 column but other rows quantify the cost of assuming exponential decay in the analysis when the true DGP is linear or Weibull. The Discussion takes up the practical interpretation of both off-diagonal directions. This is taken up further in §5.2.

4.5 Lasagna view of the full factorial

4.6 Type I error control

[28]

Type I error is conservative, not inflated: mean 0.039 across 540 null cells, the working-covariance conservatism the companion study documents rather than ideal control.

- Mean empirical Type I error 0.039, below nominal 0.05: the conservative side of the documented mis-calibration.
- Cell-level maximum 0.084 is within the 540-cell multiplicity range; no specification or architecture inflates the rate.
- The conservatism is removed by a cluster-robust standard error (block S6) and leaves the comparison intact.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-




Figure 3: Lasagna plot of power across the 540-cell Tier 1 factorial. Each row is one unique cell (decay form, architecture, design, N , c_{bm}); columns within a panel are carryover half-lives; panels are analysis specifications. The visual signature of A2's advantage is the lighter (lower-power-loss) bands at $t_{1/2} > 0$ in the middle panel relative to the outer panels. Recommended by Morris et al. [-@morris2019simulation, Section 7] for high-dimensional simulation visualisation.

Type-I error control under null

$c_{bm} = 0$, $N = 70$, pooled across DGP decay forms and carryover levels




Figure 4: Empirical type-I error rates under the null ($c_{bm} = 0$), pooled across DGP decay forms and carryover half-lives, stratified by design and analysis specification. Boxplots cover the distribution across cells at $N = 70$; the dashed reference line marks the nominal 0.05 level.

mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
mollit anim id est laborum.

Mean empirical Type I error across all 540 null cells is 0.039, below the nominal 0.05 level. The cell-level maximum is 0.084, which falls within the binomial Monte Carlo 95% interval [0.033, 0.073] for nominal 0.05 at $n_{sim} = 500$ once the multiplicity across 540 cells is acknowledged (the expected maximum among 540 independent binomial draws of size 500 at $p = 0.05$ exceeds 0.07). The shortfall below nominal is not sampling slack but the conservatism of the model-based interaction test that the companion calibration study⁴ traces to the corCAR1 working covariance: its standard error is over-estimated by six to ten percent and the realised rejection rate falls correspondingly. The conservatism is uniform across specifications and architectures, which is why it leaves the comparison intact; sensitivity block S6 confirms that a CR2 cluster-robust standard error restores the rejection rate to nominal without disturbing the specification ranking.

4.7 Sensitivity analyses

[29]

Five Tier 2 marginal sensitivity blocks probe single-axis perturbations around the Tier 1 reference configuration.

- Reference configuration: Architecture B, Hybrid, $N = 70$, $c_{bm} = 0.45$, exponential DGP, $t_{1/2} = 1.0$ weeks.
- Each block varies one axis; all others are held at reference values.

867 • Blocks are not pooled with Tier 1 or with each other; design
868 rationale in `simulation-grid-plan.md`.

869 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
870 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
871 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
872 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
873 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
874 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
875 mollit anim id est laborum.

876 *Before going into the individual blocks, we note the general design.*
877 *Five Tier 2 marginal sensitivity blocks probe the robustness of the Tier*
878 *1 ranking to single-axis perturbations around the reference configura-*
879 *tion (Architecture B, Hybrid design, $N = 70$, $c_{bm} = 0.45$, exponential*
880 *DGP, $t_{1/2} = 1.0$ weeks). Each block varies one axis and holds the*
881 *others at their reference values. Blocks are reported separately and*
882 *share the reference-cell anchor with Tier 1; they are not pooled with*
883 *Tier 1 or with each other. The design rationale is documented in*
884 *analysis/scripts/carryover-sensitivity/simulation-grid-plan.md.*

885 4.7.1 S1: within-factor autocorrelation

S1: Sensitivity to within-factor autocorrelation

Reference: Hybrid, Arch B, $N=70$, $c_{bm}=0.45$, exp DGP, $t_{1/2}=1.0$




Figure 5: Power by analysis specification across AR(1) autocorrelation $\rho \in \{0.5, 0.7, 0.8, 0.9\}$. Hendrickson (2020) uses $\rho = 0.8$; docs/02 uses $\rho = 0.7$ after positive-definiteness reductions.

886 [30]

887 **Autocorrelation lifts all three specifications proportionally;**
888 **the A2 advantage of approximately 0.06 is preserved at every**
889 **ρ level.**

- 890 • Power increases monotonically with ρ for all specifications.
- 891 • A2 advantage: ≈ 0.06 absolute throughout $\rho \in [0.5, 0.9]$.

- Autocorrelation and carryover act additively on the specification ranking; no multiplicative interaction.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

All three specifications gain power monotonically with ρ . At $\rho = 0.5$, A2 power is 0.798 ± 0.018 versus A1 0.738 ± 0.020 and A3 0.738 ± 0.020 ; at $\rho = 0.9$, A2 reaches 0.956 ± 0.009 versus A1 0.906 ± 0.013 and A3 0.910 ± 0.013 (point estimate $\pm$ MC SE, $n_{sim} = 500$). Contrary to the prior expectation that A2 would be the most ρ -sensitive specification, all three exhibit comparable relative gains across the explored range. The A2 advantage over A1 and A3, approximately 0.06 in absolute power throughout, is preserved at every AR(1) level. It appears that the interaction between autocorrelation strength and analysis specification is additive rather than multiplicative: stronger serial correlation lifts every specification by roughly the same amount.

4.7.2 S2: analyst-truth half-life mismatch

S2: Cost of analyst-truth half-life mismatch

A1 and A3 do not depend on assumed half-life; A2 does




Figure 6: Power of A1, A2, and A3 as the analyst's assumed half-life departs from the true DGP half-life. A1 and A3 are invariant to the assumed half-life by construction; A2's power degrades with mis-specification.

[31]

A2 retains its advantage over A1 and A3 even when the analyst's assumed half-life departs from the true value by a factor of four.

- A2 power ranges from 0.766 to 0.964 across eight S2 cells; A1 and A3 average 0.840 each (invariant by construction).

- A2's curvature is gentle enough that a two-fold mis-specification does not erase its advantage.
- A2 is a defensible default when the pharmacokinetic half-life is known only approximately.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

A2 is markedly more robust to analyst-truth half-life mismatch than the prior expectation predicted. Across the eight S2 cells ($DGP\ t_{1/2} \in \{0.5, 1.0\}$ weeks crossed with analyst-assumed $t_{1/2} \in \{0.25, 0.5, 1.0, 2.0\}$ weeks), A2 averages 0.868, ranging from 0.766 ($DGP\ 1.0$ weeks analysed at 0.25 weeks) to 0.964 ($DGP\ 0.5$ weeks analysed at 0.25 weeks); per-cell MC SE ranges from 0.008 to 0.019. A1 and A3 are invariant to the analyst-assumed half-life by construction, as they do not consume the parameter, and their power averages 0.840 each across the same eight cells. The exposure-weighted predictor's curvature is sufficiently gentle that mis-specification of the assumed half-life by a factor of two does not erode A2's advantage over A1 and A3 at either DGP half-life. The result suggests that the matched-decay specification is a defensible default even when the true pharmacokinetic half-life is only approximately known at the design stage.

4.7.3 S3: dropout

S3: Sensitivity to dropout

MCAR and MAR (severity-biased); reference cell as above




Figure 7: Power at dropout rates $\in \{0, 0.1, 0.2, 0.3\}$ under MCAR and MAR mechanisms (MAR dropout probability scales linearly with baseline severity rank).

[32]

A2's advantage narrows under heavy dropout because its signal source – the exposure-weighted off-drug contrast – is

946

eroded by censoring.

947

948

949

950

951

952

953

954

955

956

957

958

959

960

961

962

963

964

965

966

967

968

969

970

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

All three specifications lose power monotonically with dropout rate. The MCAR baseline ($N = 70$, dropout = 0) gives $A1\ 0.764 \pm 0.019$, $A2\ 0.866 \pm 0.015$, $A3\ 0.764 \pm 0.019$; at MCAR rate 0.30 these collapse to $A1\ 0.122 \pm 0.015$, $A2\ 0.148 \pm 0.016$, $A3\ 0.124 \pm 0.015$. $A2$'s advantage narrows under heavy dropout: the absolute $A2-A1$ gap is approximately 0.10 at zero dropout, 0.07 at rate 0.10, 0.04 at rate 0.20, and 0.03 at rate 0.30 under MCAR. Under MAR with the same marginal rate the gap narrows further at the highest rate ($A2-A1 \approx 0.004$ at rate 0.30 MAR). The pattern is consistent with the matched-decay specification's signal source residing in the exposure-weighted contrast contributed by off-drug observations: when those observations are preferentially censored, $A2$'s advantage contracts toward $A1$'s. At low-to-moderate dropout (rate ≤ 0.20) the $A2$ advantage is preserved.

971

4.7.4 S4: biomarker-moderation effect-size curve

S4: Biomarker-moderation effect-size curve

Dashed reference line at nominal $\alpha = 0.05$




Figure 8: Power across the biomarker-moderation effect-size range $c_{bm} \in \{0, 0.10, 0.20, 0.30, 0.45, 0.60\}$ under Architecture B. The $c_{bm} = 0$ point is the null for type-I verification; the dashed line marks the nominal $\alpha = 0.05$.

972

[33]

973
974
975
976
977
978
979
980

A2's advantage is largest at intermediate effect sizes and shrinks toward ceiling as c_{bm} increases.

- A2–A1 advantage peaks at +0.102 absolute at $c_{bm} = 0.45$, declining to +0.042 at $c_{bm} = 0.60$.
- Type I error at $c_{bm} = 0$ is conservative for all specifications (consistent with AR(1) reference distribution).
- The A2 advantage is most practically consequential when targeting moderate biomarker-treatment interactions.

981
982
983
984
985
986
987

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

988
989
990
991
992
993
994
995
996
997
998

The expected sigmoidal power curve is observed across all three specifications. At $c_{bm} = 0$, empirical Type I error is below the nominal level at 0.026 ± 0.007 (A1), 0.028 ± 0.007 (A2), and 0.030 ± 0.008 (A3); the conservative behaviour is consistent with the AR(1) reference distribution at moderate sample sizes. Power rises through the mid-range and saturates at $c_{bm} \geq 0.45$: at $c_{bm} = 0.6$, A1 reaches 0.916 ± 0.012 , A2 0.958 ± 0.009 , and A3 0.916 ± 0.012 . The A2 advantage over A1 is largest at intermediate effect sizes, +0.082 absolute at $c_{bm} = 0.30$ (0.444 versus 0.362) and +0.102 at $c_{bm} = 0.45$ (0.820 versus 0.718), shrinking to +0.042 at $c_{bm} = 0.60$ where all specifications approach ceiling. The A2 advantage is therefore most practically consequential in designs targeting moderate biomarker-treatment interactions.

999

4.7.5 S5: autocorrelation by carryover (exploratory)

S5: Autocorrelation x carryover interaction
Exploratory: is the spec ranking rho-sensitive?




Figure 9: Power across carryover half-lives at $\rho \in \{0.5, 0.8\}$, by analysis specification. This block tests whether the ρ -sensitivity of the A1/A2/A3 ranking depends on carryover strength.

[34]

The A2 advantage of ≈ 0.096 is constant across the $\rho \times t_{1/2}$ space, confirming additive rather than interactive effects.

- A2–A1 = 0.096 at both $\rho = 0.5$ and $\rho = 0.8$ when $t_{1/2} = 1.0$ weeks.
- Autocorrelation and carryover act additively; no qualitative interaction with specification ranking.
- Block S5 functions as methodological reassurance that the Tier 1 ranking is preserved across the ρ -by- $t_{1/2}$ grid.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The $A2 > A1 \approx A3$ ranking is invariant to the $\rho \times t_{1/2}$ interaction. At $\rho = 0.5$, $t_{1/2} = 1.0$ weeks, $A2 - A1 = 0.096$ ($A2\ 0.778 \pm 0.019$, $A1\ 0.682 \pm 0.021$); at $\rho = 0.8$, $t_{1/2} = 1.0$ weeks, $A2 - A1 = 0.096$ ($A2\ 0.910 \pm 0.013$, $A1\ 0.814 \pm 0.017$). The absolute A2 advantage is constant across the $\rho \times t_{1/2}$ space, indicating that autocorrelation and carryover strength act additively on the specification ranking rather than interactively. Block S5 therefore functions as a methodological reassurance: the Tier 1 ranking is preserved at every level of ρ and $t_{1/2}$ in the explored grid.

4.8 High-precision rerun confirmation

[35]

The high-precision rerun confirms the Tier 1 A2 > A1 ranking at MC SE 0.020 with paired McNemar $p = 6.6 \times 10^{-17}$ at the highest-leverage cell.

- 24-cell rerun (600 reps/cell, OL+BDC); 100% convergence; MC SE ≈ 0.020 .
- At $N = 70$, $t_{1/2} = 1.0$ weeks: A2–A1 gap +13.7 percentage points, McNemar $p = 6.6 \times 10^{-17}$.
- A3 recovery of the A2–A1 gap is uneven: 74% at $N = 35$, $t_{1/2} = 1.0$ weeks; 48%–57% at larger N .

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur

sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

We endeavored to provide an independent, higher-resolution confirmation of the Tier 1 ranking. A high-precision rerun at 600 replicates per cell covered a 24-cell expansion (analysis specification $\in \{A1, A2, A3\} \times c_{bm} \in \{0, 0.45\} \times N \in \{35, 70\} \times t_{1/2}^{DGP} \in \{0.5, 1.0\}$ weeks, OL+BDC design), executed using 8-core `parallel::mclapply`. Wall-clock time was 436 seconds; convergence was 100% across all 14{,}400 fits. Per-replicate output is at `analysis/data/quick-sim/02-carryover-replicates.rds`; the Morris ADEMP cell-level summary is at `02-carryover-summary.txt`. Cell MC SE on power is approximately 0.020 in the middle of the range. The rerun adds independent confirmation of the Tier 1 ranking by exploiting within-replicate correlation among specifications via paired McNemar tests on the common simulated datasets.

Type I error across the twelve null cells of the rerun fell in [0.003, 0.053], so all three specifications hold nominal alpha at the explored N and DGP half-life combinations, in agreement with the Tier 1 result in the preceding section.

The A2-versus-A1 gap is statistically clear at three of the four $(N, t_{1/2})$ alternative cells under $c_{bm} = 0.45$. The binding cell is at $t_{1/2}^{DGP} = 1.0$ weeks: at $N = 35$, $\Delta_{A2-A1} = +5.7$ percentage points ($0.400 \rightarrow 0.457$), McNemar $p = 1.2 \times 10^{-4}$; at $N = 70$, the gap widens to +13.7 percentage points ($0.737 \rightarrow 0.873$), McNemar $p = 6.6 \times 10^{-17}$. The shorter-half-life cells show $\Delta_{A2-A1} \approx +2$ to +3 percentage points with McNemar $p \in [0.004, 0.015]$, a small but statistically resolvable advantage; it is reduced because at $t_{1/2}^{DGP} = 0.5$ weeks the carryover decays quickly enough that the binary on-drug A1 specification already captures most of the structure that A2 is designed to recover.

A3's recovery of the A2-versus-A1 gap is uneven across the rerun grid. Only the small- N short-half-life cell ($N = 35$, $t_{1/2}^{DGP} = 0.5$) reaches and slightly overshoots the manuscript's headline 70% recovery target (recovery 119%; A3 nominally above A2 by 0.5 percentage points). At $N = 35$, $t_{1/2}^{DGP} = 1.0$ weeks, recovery is 74%; the larger- N cells fall to 57% and 48%. Paired McNemar on A2 versus A3 finds A2 strictly dominant only at the high-leverage corner ($N = 70$, $t_{1/2}^{DGP} = 1.0$ weeks, $\Delta = +7.2$ percentage points, $p = 1.2 \times 10^{-7}$); in the other three cells the paired difference lies between -0.5 and $+1.5$ percentage points with McNemar $p \geq 0.18$, and the two specifications are operationally tied at this resolution. Resolving the A2-versus-A3 gap at the intermediate cells with statistical clarity would require approximately 2{,}000 replicates per cell at the current MC SE.

The substantive implication, carried into the Discussion, is that the summary statement 'A2 dominates A1 and A3 captures roughly 70% of the gap' is correct in spirit at the prazosin-calibrated reference cell ($N = 35$, $t_{1/2}^{DGP} = 1.0$ weeks), but the A3-as-default heuristic is supported only at small N and short half-life. The production-grade recommendation in the Discussion reports the recovery percentages across the full $(N, t_{1/2}^{DGP})$ grid.

4.9 Sensitivity block S6: cluster-robust recalibration

[36]

Cluster-robust recalibration recovers the power the conservative model-based test forfeits and preserves the specification ranking in the information-rich cells, while its precision cost dominates under heavy dropout.

- The companion calibration study shows the model-based interaction test is conservative; this block re-fits all three specifications under both the model-based and a CR2 cluster-robust standard error at the reference cell and four stress cells.
- In the information-rich cells CR2 restores calibration ($\kappa \rightarrow 1$) and recovers two to five points of power while preserving the A2 > A1, A3 ranking.
- CR2 spends precision (Satterthwaite df falls from ≈ 480 to ≈ 23 to ≈ 5); under heavy dropout the conservatism vanishes, CR2 turns mildly anti-conservative, and recalibration is moot because power has collapsed.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat.

The companion calibration study⁴ establishes that the model-based Wald test of the interaction in the lme model with a corCAR1 residual is conservative: at the reference cell its standard error is over-estimated by six to ten percent, its realised type-I error sits near 0.03, and the conservatism is asymptotic, removed by a CR2 bias-reduced cluster-robust standard error. Two consequences bear on the present study, which evaluates that same model-based test. The mean null rejection rate reported above, 0.039, is not merely good control but this conservatism surfacing across the grid; and the absolute power figures throughout are correspondingly mild under-statements. This block re-fits all three specifications under both the model-based test and the CR2 cluster-robust test, clustered on participant with Satterthwaite degrees of freedom, at the reference cell and four stress cells (Table 4), to ask whether recalibration disturbs the specification ranking and what it costs.

At the reference cell and the information-rich stress cells the result matches the companion study. The model-based standard-error ratio κ runs from 1.07 to 1.26 (Table 4), the documented conservatism; the CR2 ratio is restored to between 0.93 and 1.01; and the correctly-sized test recovers two to five percentage points of power. The A2 advantage is preserved throughout: the A2-minus-A1 gap is held or slightly widened under recalibration, from +0.080 to +0.092 at the reference cell, +0.066 to +0.086 at $N = 35$, and +0.052 to +0.032 at $\rho = 0.9$, where every specification approaches the power ceiling and the gaps compress. Mis-specifying the analyst half-life leaves A1 and A3 untouched, as they do not use it, and costs A2 alone a few points of power, recovered in part by recalibration. The ranking that is this study's contribution is therefore invariant to the conservatism.

Recalibration is neither free nor uniform. The CR2 Satterthwaite degrees of freedom fall from the model-based value of roughly 480 to about 23 at $N = 70$

Table 4: Cluster-robust recalibration of the exposure-weighted specification A2 at the reference cell and four stress cells (500 replicates, $c_{bm} = 0.45$, Architecture B). κ is the ratio of the mean standard error to the empirical standard deviation of the interaction estimate; a value above one is a conservative test. The A2–A1 columns give the power gap between the two specifications under each inference, a check that recalibration preserves the ranking. CR2 df is the mean Satterthwaite degrees of freedom for the cluster-robust test; the model-based test carries roughly 480. Computed with the `clubSandwich` package; see the companion calibration study for the mechanism.

Cell	κ model	κ CR2	Power model	Power CR2	A2–A1 model	A2–A1 CR2	CR
Reference, $N = 70$	1.15	0.99	0.829	0.863	+0.080	+0.092	23
Small $N = 35$	1.13	0.96	0.558	0.590	+0.066	+0.086	12
High autocorr, $\rho = 0.9$	1.26	0.94	0.954	0.980	+0.052	+0.032	23
Half-life mis-spec	1.14	1.00	0.759	0.795	+0.010	+0.024	23
30% MCAR dropout	0.99	0.85	0.148	0.126	+0.010	+0.010	5

and 12 at $N = 35$, the price of a working-covariance-agnostic standard error. Under heavy dropout the trade reverses. At 30 percent missingness the model-based test is no longer conservative, its standard-error ratio near 0.98; the CR2 estimator turns mildly anti-conservative, ratio near 0.85, on the few degrees of freedom that remain (about five); and its power falls slightly below the model-based test’s. All three specifications collapse to about 0.13 power with a negligible A2-minus-A1 gap of 0.01. There the recalibration is moot: the information lost to dropout, not the calibration of the standard error, governs the analysis, in keeping with this study’s broader finding that preventing dropout outweighs the choice of specification. The practical recommendation follows. In the information-rich regime the recommended analysis is the exposure-weighted specification A2 reported with a CR2 cluster-robust standard error, which is correctly sized and recovers the power the conventional test forfeits to conservatism; the comparative ranking is robust to this choice, and the absolute power figures elsewhere in this paper should be read as mildly conservative lower bounds.

5 Discussion

5.1 The principal finding

[37]

A2 dominates A1 and A3 by approximately 10 percentage points at the reference cell, and this ranking is stable across all design and DGP variation in the Tier 1 grid.

- Reference cell: A2 = 0.860, A1 = 0.766, A3 = 0.770; ≈ 5 MC SEs of separation at $n_{sim} = 500$.
- Ranking A2 > A1 $\approx$ A3 holds across all designs, architectures, sample sizes, and carryover half-lives.

- The stability of the ranking is itself informative: no design or carryover condition displaces A2.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The central result of this factorial study is relatively straightforward to state. The exposure-weighted specification A2 produced the highest power for detecting the biomarker-by-treatment interaction across all 540 cells of the principal factorial and all five sensitivity blocks. The advantage over the binary-treatment specification A1 and the lagged-nuisance specification A3 was approximately 10 percentage points at the reference configuration (Architecture B, Hybrid design, $N = 70$, $c_{bm} = 0.45$, $t_{1/2} = 1.0$ weeks; $A2 = 0.860$, $A1 = 0.766$, $A3 = 0.770$), corresponding to approximately five Monte Carlo standard errors of separation at $n_{sim} = 500$. The ranking $A2 > A1 \approx A3$ held across all three trial designs, both DGP architectures, both sample sizes, and all carryover half-lives in the Tier 1 grid. We note that the stability of this ranking is, in some respects, as informative as the ranking itself: it appears that neither the choice of trial design nor the magnitude of carryover is sufficient to displace A2 from its dominant position.

5.2 Robustness to decay-form variation

[38]

A2 retains its advantage across all five DGP decay forms, even when the analyst assumes exponential and the true form is linear or Weibull.

- A2 does not fall below A1 or A3 under any evaluated mis-specification condition.
- The linear DGP produces the largest A2 penalty (sharp cutoff poorly approximated by an exponential tail), yet the advantage is preserved.
- A2 is more robust to decay-form mis-specification than anticipated.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

A natural concern about any matched-decay specification is that its advantage

over a naive alternative may evaporate when the assumed form departs from truth. We have examined this directly. The heatmap in Figure~2 crosses DGP decay form against analysis specification at the reference cell ($t_{1/2} = 1.0$ week, Architecture B, Hybrid design, $N = 70$, $c_{bm} = 0.45$). A2 retains its advantage across all five DGP decay forms (linear, exponential, Weibull at $k \in \{0.7, 1.0, 1.5\}$) when the analyst correctly specifies the form. But what happens when the analyst assumes exponential while the true DGP is not? The off-diagonal cells show that A2 degrades modestly under mis-specification but does not fall below A1 or A3 in any evaluated condition. The linear DGP produces the largest A2 penalty, because the sharp cutoff at $2t_{1/2}$ is poorly approximated by an exponential tail, yet even there the advantage is preserved. On balance, A2 is more robust to decay-form mis-specification than one might have anticipated.

5.3 Robustness to half-life mis-specification

[39]

Block S2 contradicts the Jones-Kenward expectation: A3 does not overtake A2 at any evaluated half-life mismatch, because A2's curvature is too gentle to be costly under mis-specification.

- A2 power span across a four-fold mis-specification: ≈ 0.066 ; A2 remains above A1 and A3 throughout.
- Four-fold undershoot (assumed 0.25 wks, true 1.0 wks): A2 = 0.766 vs. A1/A3 ≈ 0.74 .
- A3 closes the gap as half-life information degrades, but practical magnitudes are too small for A3 to overtake A2.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Block S2 provides the strongest practical reassurance in the study. Jones and Kenward³ developed the crossover-trial framework underlying the A3 specification with the expectation that a non-parametric carryover indicator should outperform a parametric model under mis-specification, because the lagged indicator does not commit to a functional form. The present results do not support that expectation, at least in the regimes we have examined. When the analyst's assumed half-life departs from the true half-life by a factor of two or four (assumed $t_{1/2} \in \{0.25, 0.5, 1.0, 2.0\}$ at true $t_{1/2} = 1.0$), A2 power ranges from 0.766 to 0.832, a span of roughly 0.066 across the full range and roughly 0.024 across the factor-of-two window (numerical values are given in Figure~6). A1 and A3, by construction not consuming the assumed half-life, are invariant to this mis-specification. A2 retains its advantage over A1 and A3 at every assumed half-life examined: even the four-fold undershoot

to an assumed half-life of 0.25 weeks leaves A2 at 0.766, still above A1 and A3 at roughly 0.74. It would appear that the exposure-weighted predictor's curvature is sufficiently gentle that mis-specification is less consequential than the crossover-trial literature anticipated. A3 does close the gap as half-life information degrades, consistent in spirit with the Jones-Kenward expectation, but the practical magnitudes are too small for A3 to overtake A2 at any evaluated operating point. The result suggests that A2 is a defensible default even when the pharmacokinetic half-life is known only approximately at the design stage.

5.4 Autocorrelation sensitivity

[40]

Contrary to expectation, A2's 'corCAR1' structure does not extract disproportionately more information at higher ρ ; the advantage is approximately constant across $\rho \in [0.5, 0.9]$.

- A2 advantage: ≈ 6 pp at $\rho = 0.5$, ≈ 5 pp at $\rho = 0.9$.
- All three specifications gain power proportionally with ρ ; relative gains are comparable.
- Autocorrelation acts additively, not interactively, on the specification ranking.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Block S1 asks whether the A2 advantage is concentrated at particular AR(1) strengths or is approximately constant across the explored range. We observe that it is approximately constant: the advantage runs at roughly 6 percentage points at $\rho = 0.5$ and narrows to roughly 5 points at $\rho = 0.9$, a range too narrow to be practically meaningful. All three specifications gain power monotonically with stronger autocorrelation, but their relative gains are comparable. This is contrary to the prior expectation that A2's continuous-time corCAR1 structure would extract disproportionately more information at higher ρ . It appears, rather, that autocorrelation and carryover strength act additively on the specification ranking: stronger serial correlation lifts every specification by roughly the same amount, leaving the $A2 > A1 \approx A3$ ordering undisturbed.

5.5 Dropout and the limits of specification refinement

[41]

Dropout is far more consequential than the carryover specification choice; at 30% dropout all three specifications are near floor and effectively indistinguishable.

- At 30% MCAR, all specifications collapse to power 0.12–0.25; A2 provides no protection against the power loss itself.
- MAR at 10% is approximately equivalent to MCAR at 12%, because MAR preferentially censors high-severity participants.
- Dropout prevention is more consequential than specification refinement; trialists should prioritise accordingly.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Block S3 is, in some ways, the most sobering result in the study. The three specifications are meaningfully distinguished at low-to-moderate dropout: at 10% MCAR, power drops from approximately 0.86 to approximately 0.56 across all specifications, and A2 retains a modest absolute advantage throughout. At 30% dropout, however, all specifications are near floor (0.12 to 0.25) and are effectively indistinguishable. Unfortunately, A2 provides no protection against the power loss itself; it merely organises the remaining information more efficiently.

Under MAR dropout biased by baseline severity, the pattern is qualitatively similar but shifted slightly: MAR dropout at 10% is comparable to MCAR dropout at approximately 12%, which is expected because MAR preferentially censors high-severity participants who contribute more to the interaction signal under Architecture B.

The practical implication is direct. Dropout prevention, whether through shorter trial durations, participant engagement strategies, or adaptive scheduling, is far more consequential for the biomarker-interaction test than the choice of analysis-model carryover specification. Trialists who invest in refining the carryover specification while tolerating high dropout rates have, in our view, the priorities somewhat reversed.

5.6 Architecture dependence

[42]

Architecture B amplifies the A2 benefit because carryover erodes the biomarker-response correlation directly, and A2 recovers more of that attenuated signal.

- Architecture B is more sensitive to carryover under all three decay forms and all Weibull shapes examined.
- The A2 advantage is larger under Architecture B than Architecture A, consistent with the correlation-erosion mechanism.
- The $A2 > A1 \approx A3$ ranking holds under both architectures.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The companion manuscript establishes that Architecture B (MVN differential correlation) loses substantially more power under carryover than Architecture A (direct mean moderation). The present results confirm that this qualitative gap is robust to decay-form variation: Architecture B is more sensitive to carryover under all three decay forms and at all five Weibull shape values examined. The A2 advantage is larger under Architecture B than under Architecture A, which is consistent with the mechanism. Under Architecture B, carryover erodes the biomarker-response correlation directly, and the continuous exposure-weighted predictor recovers more of that attenuated signal than either the binary on-drug indicator or the lagged-nuisance term can. The $A2 > A1 \approx A3$ ranking holds under both architectures, though the magnitude of the benefit warrants closer attention in trials where the data-generating mechanism is presumed to follow Architecture A.

5.7 The latent-subtype question

[43]

The present simulation does not evaluate a finite-mixture DGP; whether A2 dominance extends to discrete-subtype data is an open question.

- A mixture DGP would produce discrete responder/non-responder profiles rather than a graded covariance structure.
- Under a mixture DGP, A3's form-free carryover indicator might gain an advantage over A2.
- The present results cannot address this; a mixture DGP evaluation is identified as a direction for subsequent work.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

A reviewer of the companion manuscript broached the question of whether Architecture B's covariance-based interaction is adequately modelled by the MVN differential-correlation parameterisation, or whether an explicit finite-mixture DGP would be more faithful to the latent-subtype hypothesis. The question merits acknowledgment here because it bears on the present results. A true mixture DGP would produce biomarker-response associations that are

qualitatively different from those generated by a continuous correlation: discrete responder and non-responder subgroups rather than a graded covariance structure. Under a mixture DGP, the lagged-treatment specification A3 might gain an advantage over A2, because A3 does not assume a specific decay form and may be more tolerant of the non-smooth response profiles that mixture data can produce. The present simulation does not evaluate a mixture DGP and cannot address this question directly. We have not been able to determine, as yet, whether the A2 dominance we observe extends to genuinely discrete-subtype data-generating processes; this is a direction that merits investigation in subsequent work.

5.8 Comparison with prior simulation studies

[44]

The present factorial replicates Hendrickson’s baseline and extends it by contrasting three analysis specifications, with the A2-versus-A1 advantage being a new contribution.

- Hendrickson (2020): power 0.82 reproduced here at 0.860 ± 0.016 ; A1/A3 comparators are the new element.
- Jones-Kenward expectation (A3 dominates under misspecification) is not supported at the prazosin-calibrated parameters.
- Sturdevant-Lumley two-tier reporting convention is adopted and extended to the biomarker-interaction estimand.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Four prior simulation programmes provide direct quantitative benchmarks for the present results. We discuss each in turn, focusing on cells where the prior parameter ranges overlap with ours and noting where the design targets differ.

Hendrickson² introduced the biomarker-by-treatment interaction question for aggregated N-of-1 trials and reported power for the matched-decay (A2) analysis under exponential carryover at $t_{1/2} = 1$ week, $N = 70$, $c_{bm} = 0.45$, Hybrid design. The reported power was 0.82 at 1{,}000 replicates under a single architecture. The present 500-replicate factorial reproduces this result at the same cell to 0.860 ± 0.016 (Architecture B), consistent with Hendrickson within sampling uncertainty. The principal methodological extension here is the addition of A1 and A3 as comparators; Hendrickson did not contrast analysis-side specifications, so the A2-versus-A1 advantage of approximately 10 percentage points is a new contribution rather than a replication.

Jones and Kenward³ developed the crossover-trial framework underlying the A3 specification with the expectation, noted above, that the non-parametric

lagged-on indicator should outperform a parametric carryover model under decay-form mis-specification. The present S2 sensitivity block directly tests this expectation and finds it is not supported in the regimes we have examined: although A2’s power varies by roughly 0.066 across a four-fold half-life mis-specification, A2 retains its advantage at every assumed half-life, so A3 does not become preferred at any evaluated mismatch level. The Jones-Kenward expectation is preserved in spirit, in that A3 closes the gap as half-life information degrades, but the practical magnitudes are too small for A3 to overtake A2 in the explored window. The disagreement is informative: it appears that the exposure-weighted predictor’s curvature is gentle enough that mis-specification is less costly than the crossover-trial literature anticipates, at least at the prazosin-calibrated parameters.

Senn⁵ established the variance-decomposition foundation for within-subject designs, predicting that within-subject contrasts are disproportionately advantageous when between-patient variance $\sigma_{\text{between}}^2$ is large relative to within-patient variance. The present simulation does not vary $\sigma_{\text{between}}^2$ as a primary axis, but the S1 autocorrelation sensitivity block bears on the related question of how within-subject information accumulates with serial correlation. Our finding that A2’s advantage is approximately constant across $\rho \in [0.5, 0.9]$ is qualitatively consistent with Senn’s prediction that within-subject methods extract roughly proportional information across correlation regimes; the absence of a strong ρ interaction with the analysis specification is the new result here.

Sturdevant and Lumley⁶ developed the two-tier reporting convention we have adopted: a principal factorial covering the main contrast, supplemented by marginal sensitivity blocks anchored at a reference cell. Their convention was designed for simulation studies whose main claim is a ranking or contrast rather than a precise power estimate, and our adoption of it is documented in the `simulation-grid-plan.md` design document. The two-tier structure proved useful in practice: the Tier 1 factorial establishes the $A2 > A1 \approx A3$ ranking, while the Tier 2 blocks probe its robustness to single-axis perturbations. The extension to the biomarker-interaction estimand is, to the best of our knowledge, new.

The companion manuscripts at `analysis/report/01-dgp-mean-moderation-vs-mvn/` and `analysis/report/05-nof1-design-sensitivity/` provide complementary benchmarks: the first establishes the architecture-dependence of carryover-induced power loss confirmed here in the architecture-dependence subsection, and the second reports the effect of carryover on the Hybrid N-of-1 design’s saturation profile confirmed in §4.4.

5.9 Limitations

[45]

Five limitations bound the generalisability of the findings: single prazosin parameterisation, incomplete analysis strategy space, OFAT sensitivity blocks, omitted OL design, and monotone dropout assumption.

- Absolute power values are specific to the prazosin-PTSD param-

eterisation; rankings are expected to be more portable but are not verified across alternative trajectory shapes.

- Three specifications do not exhaust the carryover-aware model space; weighted analysis, random slopes, and observation exclusion are not evaluated.
- Sensitivity blocks are OFAT; joint perturbations (e.g., high ρ combined with high dropout) are not fully explored.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Several limitations should be noted when interpreting these results.

First, the simulation draws its response-trajectory parameters from a single configuration calibrated to the Hendrickson (2020) prazosin-PTSD framework (Gompertz maxima, rates, and dispersions). The absolute power values are specific to this parameterisation. The ranking of analysis specifications is expected to be more portable, but we have not been able to verify this across alternative response profiles. Trialists working in therapeutic areas with substantially different trajectory shapes should regard the power estimates as illustrative rather than definitive.

Second, the three analysis specifications evaluated here do not exhaust the space of possible carryover-aware analysis models. Weighted analysis, within-subject contrasts, two-stage random slopes, and exclusion of contaminated observations (discussed in the companion manuscript’s Section 6) are not evaluated. We have endeavored to select the three specifications that represent the most natural practical alternatives for the N-of-1 context, but a more comprehensive comparison would cross the present grid with these additional strategies, and others may merit consideration.

Third, the sensitivity blocks are one-factor-at-a-time perturbations around a single reference configuration. Unfortunately, joint perturbations, for example high ρ combined with high dropout, may produce interactions that the marginal blocks do not capture. The exploratory block S5 ($\rho \times$ carryover) provides one such joint check and does not reveal a qualitative interaction, which is reassuring, but the space of possible joint combinations is large and our exploration of it is necessarily incomplete.

Fourth, the Tier 1 grid omits the open-label (OL) design because OL lacks the off-drug contrast required by the A1 and A3 specifications. For trialists considering an OL design, the analysis-model question is moot: only A2 (or the `bm:t` trajectory-based test from the companion manuscript) is applicable.

Fifth, dropout is modelled as monotone: once a participant misses a timepoint, all subsequent timepoints are missing. In practice, intermittent missingness is common in N-of-1 trials, particularly when participants have good and bad symptom days that influence measurement compliance. Intermittent missing-

ness would presumably be less destructive than monotone dropout because it preserves some late-trial observations that contribute to the treatment contrast. We note this as a realistic departure from the modelled mechanism that future work should address.

The headline Tier 1 ranking is further supported by the high-precision rerun reported at the end of the Results section, which resolves the A2-versus-A1 gap to MC SE 0.020 with paired McNemar $p < 10^{-3}$ at the headline cells.

6 Conclusions

1. **A2 (exposure-weighted) is the recommended default.** Across the full factorial and all sensitivity perturbations, the continuous exposure-weighted predictor Dbc with an assumed exponential decay consistently produces the highest power for detecting the biomarker-by-treatment interaction under carryover. The advantage over the binary-treatment (A1) and lagged- nuisance (A3) alternatives is approximately 10 percentage points at the reference configuration and is robust to DGP decay-form variation, within-factor autocorrelation, and analyst-truth half-life mismatch. The recommended report is A2 with a CR2 cluster-robust standard error, which restores nominal size; the ranking is invariant to this recalibration (Block S6).
2. **Half-life mis-specification does not displace A2.** A2 power varies by roughly 0.066 across a four-fold mis-specification of the analyst’s assumed half-life (Block S2), but A2 retains its advantage over A1 and A3 at every assumed half-life examined. Trialists who have only an approximate pharmacokinetic estimate can adopt A2 with confidence that a mis-specified half-life does not cost A2 its ranking.
3. **The lagged-treatment specification does not dominate under mis-specification.** Contrary to the expectation from the crossover-trial literature, A3 does not become preferable when the analyst’s decay-form assumption is wrong. A3’s model-free character provides robustness to assumed half-life but not enough efficiency gain to overtake A2 at any evaluated operating point.
4. **Dropout is more consequential than the carryover specification.** Power degrades sharply with dropout rate under all specifications. At 20% dropout, all three specifications lose approximately half their power relative to the no-dropout condition. Trialists should prioritise dropout prevention over carryover-specification refinement.
5. **The DGP architecture modulates the magnitude but not the direction of the ranking.** Architecture B (MVN differential correlation) amplifies the benefit of A2 because carryover-induced correlation erosion makes the continuous predictor’s information advantage more consequential. Under Architecture A (direct mean moderation), A2’s advantage is smaller but still positive.

6. **Reporting recommendations.** Simulation studies evaluating biomarker-moderated treatment effects under carryover should report: (a) the analysis-model carryover specification used,
- (b) the assumed decay form and half-life, (c) sensitivity of power estimates to at least one alternative specification, and (d) the anticipated dropout rate, which is the dominant source of power loss in the evaluated grid.
7. **Type-I error is conservative, and recalibrable.** All three specifications hold below the nominal 0.05 across the full Tier 1 grid and the S4 effect-size curve, with no inflation under any DGP decay form, architecture, or carryover level. The shortfall is the working-covariance conservatism the companion calibration study documents; a CR2 cluster-robust standard error removes it and recovers a few points of power without disturbing the ranking (Block S6).

7 Conflict of interest

The authors declare no conflicts of interest.

8 Funding

This work received no specific funding.

9 Data availability statement

[46]

All simulation outputs and drivers are versioned in the pmsimstats-ng repository; exact paths are given in the Reproducibility section.

- Per-replicate Monte Carlo outputs and cell-level summaries: `analysis/data/`.
- Simulation drivers: `analysis/scripts/carryover-sensitivity/`.
- Exact file paths listed in the Reproducibility section below.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The data and code supporting the findings of this study are available in the pmsimstats-ng project repository. Per-replicate Monte Carlo outputs and ADEMP-compliant cell-level summaries are versioned under

analysis/data/; simulation drivers are at analysis/scripts/. Exact paths for this manuscript are listed in the Reproducibility section below.

10 Reproducibility

[47]

The manuscript is fully reproducible from pinned package versions in `renv.lock`, deterministic seeds, and pre-computed checkpoints; no simulation code re-executes during knit.

- Environment: R 4.4.x, zzzcollab Docker container; packages pinned in `renv.lock`.
- Principal outputs: `01-factorial.rds` and `04-sensitivity.rds`; cell-level summaries at `analysis/data/02-grid-summary.rds`.
- Manuscript Rmd loads pre-computed checkpoints only; no simulation code runs during render.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

*This manuscript was rendered with R 4.4.x inside the project's zzzcollab Docker container (image hash recorded in `Dockerfile`; package set captured in `renv.lock`). Principal packages used by the simulation pipeline are `nlme` (continuous-time linear-mixed analysis with `corCAR1` residual correlation), `parallel` (8-core `mclapply` for parameter-grid sweeps), `data.table` (the *original-extended* simulation collection), `furrr` and `purrr` (the *tidyverse* simulation collection), and `rmarkdown` plus `bookdown` for the manuscript build.*

Random-number generator. *Each simulation driver sets `set.seed(20260507)` at the top of the replicate loop and uses `parallel::mc.set.seed` inside `mclapply`. Per-replicate seeds are derived from the master seed as `master + rep_idx`.*

Data and code availability. *Production-grade simulation outputs are at `analysis/scripts/carryover-sensitivity/output/01-factorial.rds` (principal factorial) and `04-sensitivity.rds` (sensitivity blocks `S1` through `S5`). Cell-level summaries are at `analysis/data/02-grid-summary.rds` and `analysis/data/02-sensitivity-summary.rds`. Driver scripts are `01-run-factorial.R`, `04-run-sensitivity-blocks.R`, and `02-summarise-grid.R` under `analysis/scripts/carryover-sensitivity/`. The simulation grid plan and design rationale are documented at `analysis/scripts/carryover-sensitivity/simulation-grid-plan`, which serves as the Morris ADEMP pre-registration. High-precision prototype rerun outputs are at `analysis/data/quick-sim/02-carryover-replicates.rds` (per-replicate) and `02-carryover-summary.txt` (Morris ADEMP cell-*

1631 level summary). The manuscript source, paper-specific bibliography
1632 (references.bib), and SIM CSL file are at analysis/report/02-carryover-sensitivity/.

1633 11 References

- 1634 1. Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Statistics in Medicine*. 2019;38(11):2074-2102. doi:10.1002/sim.8086
- 1635 2. Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: Statistical methods and practical considerations. *Frontiers in Digital Health*. 2020;2:13. doi:10.3389/fdgth.2020.00013
- 1636 3. Jones B, Kenward MG. *Design and Analysis of Cross-over Trials*. 3rd ed. CRC Press; 2014.
- 1637 4. pmsimstats team. Type-I calibration of fixed-effect interaction tests in linear mixed models: A staged diagnosis of working-covariance misspecification and a cluster-robust remedy. Published online 2026.
- 1638 5. Senn S. Mastering variation: Variance components and personalised medicine. *Statistics in Medicine*. 2016;35(7):966-977. doi:10.1002/sim.6739
- 1639 6. Sturdevant SG, Lumley T. Statistical methods for testing carryover effects: A mixed effects model approach. *Contemporary Clinical Trials Communications*. 2021;22:100711. doi:10.1016/j.conctc.2021.100711

1640

1641 *Rendered on 2026-06-11 at 18:59 PDT.*

1642 *Source: ~/prj/res/36-pmsimstats-ng/pmsimstats-ng/analysis/report/02-carryover-sensitivity/rep*