

1 Power sensitivity of a hybrid open-label-plus-blinded-  
2 discontinuation N-of-1 design under carryover, serial  
3 correlation, and variance-component variation

4 pmsimstats team

5 2026-07-30

6 **Contents**

|    |                                                                        |           |
|----|------------------------------------------------------------------------|-----------|
| 7  | <b>1 Abstract</b>                                                      | <b>2</b>  |
| 8  | <b>2 Introduction</b>                                                  | <b>3</b>  |
| 9  | <b>3 Simulation design</b>                                             | <b>5</b>  |
| 10 | 3.1 Aims . . . . .                                                     | 5         |
| 11 | 3.2 Pipeline . . . . .                                                 | 6         |
| 12 | 3.3 Dual data-generating architectures . . . . .                       | 6         |
| 13 | 3.4 Response curves and carryover . . . . .                            | 7         |
| 14 | 3.5 Covariance structure and analysis . . . . .                        | 8         |
| 15 | 3.6 Twelve sensitivity sweeps . . . . .                                | 8         |
| 16 | 3.7 Reproducibility and execution . . . . .                            | 9         |
| 17 | <b>4 Results</b>                                                       | <b>10</b> |
| 18 | 4.1 S1. Effect-size sweep . . . . .                                    | 11        |
| 19 | 4.2 S2. Carryover half-life sweep . . . . .                            | 13        |
| 20 | 4.3 S3. Decay-form misspecification sweep . . . . .                    | 15        |
| 21 | 4.4 S4. Between-patient heterogeneity sweep . . . . .                  | 16        |
| 22 | 4.5 S5. Within-patient variance sweep . . . . .                        | 17        |
| 23 | 4.6 S6. Cycles-per-patient sweep . . . . .                             | 19        |
| 24 | 4.7 S7. Patient-count sweep . . . . .                                  | 20        |
| 25 | 4.8 S8. AR(1) serial-correlation sweep . . . . .                       | 22        |
| 26 | 4.9 S9. Period-length sweep . . . . .                                  | 23        |
| 27 | 4.10 S10. Biomarker interaction strength crossed with DGP architecture | 25        |
| 28 | 4.11 S11. Misspecification factorial . . . . .                         | 26        |
| 29 | 4.12 S12. Null-effect calibration . . . . .                            | 28        |
| 30 | <b>5 Discussion</b>                                                    | <b>30</b> |
| 31 | 5.1 Synthesis . . . . .                                                | 30        |
| 32 | 5.2 Comparison with prior simulation studies . . . . .                 | 31        |
| 33 | 5.3 Limitations . . . . .                                              | 33        |

|    |                                                    |           |
|----|----------------------------------------------------|-----------|
| 34 | 5.4 N-of-1 design variants not evaluated . . . . . | 35        |
| 35 | 5.5 Prototype Monte Carlo confirmation . . . . .   | 36        |
| 36 | <b>6 Conclusions</b>                               | <b>37</b> |
| 37 | <b>7 Conflict of interest</b>                      | <b>39</b> |
| 38 | <b>8 Funding</b>                                   | <b>39</b> |
| 39 | <b>9 Data availability statement</b>               | <b>39</b> |
| 40 | <b>10 Reproducibility</b>                          | <b>39</b> |
| 41 | <b>11 References</b>                               | <b>41</b> |

## 42 1 Abstract

43 **Background.** Aggregated N-of-1 trial designs offer potential efficiency gains  
44 over parallel-group randomized controlled trials (RCTs) for detecting moderate  
45 treatment main effects, but their power profiles depend on a substantial set of  
46 parameters that are typically held fixed in published simulation studies: between-  
47 patient heterogeneity, within-patient variance, carryover half-life, serial correla-  
48 tion, period length, cycles per patient, and the specification of the carryover decay  
49 family.

50 **Methods.** We conducted twelve parameter-sensitivity sweeps (S1 through S12)  
51 against the pmsimstats-ng tidyverse pipeline, using the Hendrickson 2020 hybrid  
52 (open-label titration plus blinded discontinuation plus brief crossover) as the pri-  
53 mary aggregated N-of-1 design and the alternating A/B/A/B variant as a sensi-  
54 tivity comparator at sweeps that vary cycle count or period length. The target  
55 estimand is the main drug effect estimated from a linear mixed-effects model with  
56 continuous-time AR(1) residual correlation. Each sweep records per-replicate re-  
57 sults so that Monte Carlo standard errors (MCSE) can be computed at render  
58 time per Morris et al.<sup>1</sup>.

59 **Results.** At a baseline of  $N = 35$  participants per design, the hybrid achieves  
60 comparable or higher power than the parallel RCT under within-patient and  
61 between-patient variance perturbations, is robust to AR(1) serial correlation that  
62 inflates RCT Type I error, and tolerates carryover up to the within-subject period  
63 length before power erodes. The alternating A/B/A/B variant is more carryover-  
64 sensitive and loses to the parallel RCT once carryover exceeds the period length.  
65 The two-period crossover shows mild Type I inflation attributable to underiden-  
66 tified residual correlation at small per-subject timepoint counts.

67 **Conclusions.** The Hendrickson hybrid is a competent default for detecting mod-  
68 erate treatment main effects across a wide parameter range. Design ordering  
69 becomes parallel-RCT-favored only when carryover exceeds the within-subject  
70 period length, and is directionally different for biomarker-treatment interaction

71 targets, which favor N-of-1 designs more strongly.

72 **Keywords.** N-of-1 trials; randomized controlled trials; statistical power; carry-  
73 over effects; serial correlation; sensitivity analysis; Monte Carlo simulation; hybrid  
74 crossover designs.

## 75 2 Introduction

76 [1]

77 **Aggregated N-of-1 designs offer power advantages over par-**  
78 **allel RCTs, but the conditions governing that advantage are**  
79 **incompletely characterized.**

- 80 • N-of-1 designs accumulate within-patient evidence across re-  
81 peated crossover periods, enabling meta-analytic population-level  
82 inference.
- 83 • Recent simulations report N-of-1 procedures matching or exceed-  
84 ing RCT power at smaller sample sizes.
- 85 • The parameter axes that govern when the N-of-1 advantage holds  
86 or reverses are not yet well specified.

87 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*  
88 *tempor incididunt ut labore et dolore magna aliqua.*

89 Aggregated N-of-1 trial designs evaluate treatment effects within in-  
90 dividual patients through repeated crossover periods and combine  
91 patient-level estimates via meta-analytic or mixed-effects machinery  
92 to support population-level inference<sup>2,3</sup>. Several recent simulation  
93 studies have reported that aggregated N-of-1 procedures can match or  
94 exceed the statistical power of parallel-group randomized controlled  
95 trials (RCTs) at substantially smaller participant counts<sup>4-6</sup>. The re-  
96 ported magnitude of this efficiency advantage, however, varies across  
97 studies, and the conditions under which it holds or reverses are in-  
98 completely characterized. That is, the published literature provides  
99 encouraging headline comparisons but relatively sparse guidance on  
100 the parameter axes that govern the ordering of designs.

101 [2]

102 **This paper reports a systematic sensitivity study of the Hen-**  
103 **drickson hybrid N-of-1 design for the main drug effect esti-**  
104 **mand.**

- 105 • The hybrid combines open-label titration, blinded discontinua-  
106 tion, and brief crossover phases.
- 107 • Three comparators are evaluated: alternating A/B/A/B N-of-1,  
108 two-period crossover, and parallel-group RCT.

- The target estimand is the main drug effect, distinct from the biomarker-treatment interaction target examined elsewhere.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

We report here a systematic parameter-sensitivity study of a specific hybrid aggregated-N-of-1 design that combines an open-label titration phase, a blinded discontinuation phase, and a brief crossover phase, following the design proposed by Hendrickson and colleagues<sup>5</sup> for predictive-biomarker validation in chronic neurobehavioral conditions. The sensitivity sweeps are conducted against the pmsimstats-ng tidyverse pipeline that operationalizes this design in R, with three parallel comparators: an alternating A/B/A/B aggregated N-of-1 variant, a two-period crossover, and a parallel-group RCT. The target estimand throughout is the main drug effect, distinguished from the biomarker-treatment interaction examined in the original Hendrickson study and in a separate companion paper.

[3]

## **Twelve sensitivity sweeps address the gap left by single-axis simulation comparisons in the published literature.**

- Published comparisons typically vary one or two parameters at a time, leaving the primary power-ordering axes unidentified.
- The twelve sweeps span effect size, carryover half-life, decay-form misspecification, heterogeneity, within-patient variance, cycles, sample size, AR(1) correlation, period length, biomarker architecture, a misspecification factorial, and null calibration.
- Each sweep axis is anchored to a specific precedent in the simulation-comparison literature.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The motivation for a sensitivity study at this granularity is straightforward: the published simulation comparisons typically vary one or two parameters at a time while holding the remainder fixed at calibration defaults, leaving open the question of which axes actually drive the N-of-1-versus-RCT power difference. We have endeavored to address this gap by reporting twelve sensitivity sweeps along axes drawn from the published comparison literature: effect size<sup>4</sup>, carryover half-life<sup>4,5</sup>, decay-form misspecification<sup>7</sup>, between-patient heterogeneity<sup>8,9</sup>, within-patient variance<sup>10</sup>, cycles per patient<sup>8</sup>, total participant count<sup>4,8</sup>, AR(1) serial correlation<sup>10</sup>, period length<sup>10</sup>, biomarker interaction strength crossed with data-generating process architecture<sup>5</sup>, a carryover-by-AR(1) misspecification factorial<sup>7</sup>, and a high-replicate null calibration<sup>1</sup>.

[4]

**This paper extends the companion clinical comparison across twelve parameter axes for the hybrid design family.**

- The companion paper provides a head-to-head comparison at the prazosin-calibrated baseline.
- The present paper documents how that baseline comparison generalizes under systematic parameter perturbation.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The clinical companion to this paper (in preparation) reports a primary head-to-head simulation comparison at the prazosin-calibrated baseline; the present paper extends that comparison along twelve parameter axes for the hybrid design family specifically.

## 3 Simulation design

### 3.1 Aims

We follow the ADEMP framework of Morris, White, and Crowther (2019)<sup>1</sup>: Aims, Data-generating mechanisms, Estimands, Methods, Performance measures. The subsections of this Simulation design section map to these five components, although the section headers are organized by topic rather than relabeled to the ADEMP acronym.

[5]

**The study follows the ADEMP framework with three aims: power characterization, design-ordering identification, and Type I error verification.**

- Aim 1: characterize the hybrid’s power profile under systematic perturbation of published comparison axes.
- Aim 2: identify parameter regimes where the hybrid is preferred over the parallel RCT and where the ordering reverses.
- Aim 3: verify Type I error control with MCSE below 1%.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The simulation study has three aims. First, to characterize the power profile of the Hendrickson 2020 hybrid aggregated-N-of-1 design<sup>5</sup> under systematic perturbation of the parameter axes implicated in the published simulation-comparison literature. Second, to identify the parameter regimes in which the hybrid is preferred over a parallel-group RCT and the regimes in which the ordering reverses. Third, to verify Type I error control under null and near-null conditions, with sufficient replication to provide MCSE below 1% per Morris et al.<sup>1</sup>.

## 3.2 Pipeline

[6]

All sweeps use a baseline of  $N = 35$  participants and 8 weekly timepoints, with sample size matched per design rather than per observation count.

- The pipeline is implemented in the `nof1power` R package, which re-implements the Hendrickson 2020 reference pipeline.
- The  $N = 35$  minimum and per-design matching convention follow Blackston 2019 and Hendrickson 2020.
- Parameters are varied from this baseline one axis at a time across the twelve sweeps.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The simulation framework is implemented in the `nof1power` R package<sup>5</sup> developed alongside the `pmsimstats-ng` tidyverse re-implementation of the Hendrickson 2020 reference pipeline. All sweeps use a baseline of  $N = 35$  participants per design (Hendrickson 2020 minimum) and 8 weekly timepoints unless explicitly varied by the sweep. Per-replicate sample size is  $N$  participants for each design under comparison; participants are matched on the per-design count rather than on observation total, following the convention adopted in Blackston<sup>4</sup>, Hendrickson<sup>5</sup>, and the majority of the N-of-1 simulation-comparison literature.

## 3.3 Dual data-generating architectures

The pipeline supports two data-generating processes (DGP) for the biomarker-treatment interaction:

- **Architecture B** (`dgp_architecture = "mvn"`). The biomarker-bio-response interaction is encoded in the covariance structure of a multivariate normal (MVN) draw. On-drug timepoints carry correlation `c.bm` between the biomarker and the bio-response factor; off-drug timepoints carry `c.bm * phi(tsd)` with `phi` the carryover decay multiplier.
- **Architecture A** (`dgp_architecture = "mean_moderation"`). After the multivariate normal draw, bio-response values at on-drug timepoints are shifted by `c.bm * bm_z * br_sd`, where `bm_z` is the standardized biomarker. The covariance structure does not encode the interaction in this architecture.

[7]

Two data-generating architectures encode the biomarker interaction differently, producing substantively different power profiles under carryover.

- Architecture B (MVN) encodes the interaction in the covariance structure; on-drug correlation decays exponentially off-drug.
- Architecture A (mean moderation) shifts bio-response means additively at on-drug timepoints by the standardized biomarker.
- The two architectures are crossed in sweep S10; the remaining eleven sweeps use Architecture B as the default.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

These two architectures produce substantively different power profiles under carryover, as documented in<sup>5</sup>. The two architectures are crossed in sweep S10 below; the remaining eleven sweeps use Architecture B as the default.

### 3.4 Response curves and carryover

Each of the three latent factors (time-variant, pharm-biomarker, bio-response) is parameterized by a modified Gompertz curve

$$y(t) = \text{maxr} \cdot \frac{\exp(-\text{disp} \cdot \exp(-\text{rate} \cdot t)) - \exp(-\text{disp})}{1 - \exp(-\text{disp})}$$

[8]

**Latent response factors follow a modified Gompertz curve; carryover is applied via a recursive decay from an extensible decay-family.**

- The Gompertz curve passes through the origin and asymptotes to `maxr`, parameterized by displacement and rate.
- Off-drug carryover is applied recursively using one of four decay families: exponential, linear, Weibull, or power.
- The DGP decay family and the analysis-side assumption can differ; sweep S3 quantifies the cost of that mismatch.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

evaluated at the appropriate cumulative time index. The curve passes through the origin at  $t = 0$  and asymptotes to `maxr`. Carryover of the bio-response component during off-drug periods is applied recursively as  $\mu[t] = \text{base}[t] + \mu[t-1] \cdot \phi(\text{tsd})$  with  $\phi$  drawn from an extensible decay family: exponential  $\exp(-\ln 2 \cdot t/t_{1/2})$ , linear  $\max(0, 1 - t/(2t_{1/2}))$ , Weibull  $\exp(-(\lambda_w t)^k)$  with  $\lambda_w = (\ln 2)^{1/k}/t_{1/2}$ , or power  $\max(0, (1 - t/(3t_{1/2}))^p)$ . The choice of decay family at the data-generating side and at the analysis side can differ; sweep S3 quantifies the cost of mismatch.

**Units.** Carryover half-lives  $t_{1/2}$  are quoted in **days** throughout this

paper, matching the pharmacokinetic scale of the prazosin reference application and the day-denominated axes of the sensitivity figures. Companion papers in this series that study the biomarker-treatment interaction (in particular the DGP-architecture and carryover-specification manuscripts) quote  $t_{1/2}$  in **weeks** on the canonical grid  $\{0, 0.5, 1.0\}$ . The conversion is 1 week = 7 days, so those papers' grid corresponds to  $\{0, 3.5, 7\}$  days here, and this paper's 3-day baseline corresponds to approximately 0.43 weeks. Timepoints  $t$  are weekly in both conventions.

### 3.5 Covariance structure and analysis

[9]

**The analysis model is an `nlme::lme` with `corCAR1` residual correlation; the target estimand is the Dbc main drug effect coefficient.**

- The covariance matrix couples within-factor AR(1), cross-factor correlations, and biomarker differential correlation; non-PD matrices are corrected with a per-cell flag.
- The fixed-effects formula is  $Sx \sim bm + t + Dbc + bm:Dbc$  with a random intercept per participant.
- Singular fits are detected from the random-effects variance and flagged per replicate.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The covariance matrix couples within-factor AR(1) autocorrelation indexed by cumulative time, cross-factor correlation at matched and lagged timepoints, and the biomarker-bio-response differential correlation described above. Non-positive-definite matrices are corrected via `corpcor::make.positive.definite(tol = 1e-3)` with a per-cell flag returned. The analysis model fits the `nlme::lme` linear mixed-effects formula<sup>11,12</sup> with fixed effects  $Sx \sim bm + t + Dbc + bm:Dbc$  (or the fallback  $bm + t + bm:t$  when the drug indicator does not vary within participant), a random intercept per participant, and `corCAR1(form = ~t | ptID)` continuous-time AR(1) residual correlation. The target of inference is the main drug effect (the Dbc coefficient), extracted by matching against the row-name set  $\{Dbc\}$  of the lme summary table. Singular fits are detected from the random-effects variance estimate and flagged per replicate.

### 3.6 Twelve sensitivity sweeps

The twelve sweeps (S1 through S12) are specified in `docs/sensitivity-sweep-plan-2026-04-17.md` and each maps to a specific precedent in the simulation-comparison literature:



- **S1** varies the effect size by scaling the bio-response maximum; precedent Blackston 2019<sup>4</sup>.
- **S2** varies the carryover half-life across nine values from 0 to 14 days across four designs (parallel RCT, two-period crossover, aggregated N-of-1, hybrid); precedents Blackston 2019 and Hendrickson 2020<sup>5</sup>.
- **S3** varies the DGP carryover decay form (exponential, linear, Weibull, power) while the analysis assumes exponential, quantifying misspecification bias; precedent Gärtner 2023<sup>7</sup>.
- **S4** varies between-patient heterogeneity via the baseline-symptom SD; precedents Senn 2018<sup>8</sup> and Zucker 2010<sup>9</sup>.
- **S5** varies the within-patient bio-response SD; precedent Wang and Schork 2019<sup>10</sup>.
- **S6** varies the number of cycles per patient  $k$  in the Senn decomposition; precedent Senn 2018<sup>8</sup>.
- **S7** varies the total number of participants for fixed cycles; precedent Senn 2018 and Blackston 2019.
- **S8** varies the AR(1) correlation parameter  $\rho$  across within-factor autocorrelations; precedent Wang and Schork 2019<sup>10</sup>.
- **S9** crosses period length with carryover half-life, probing the ratio  $t_{1/2}/\text{period}$ ; precedent Wang and Schork 2019<sup>10</sup>.
- **S10** varies the biomarker interaction strength `c.bm` across four levels crossed with both DGP architectures (`mvn`, `mean_moderation`); precedent Hendrickson 2020<sup>5</sup>.
- **S11** crosses carryover presence with AR(1) presence under null and alternative, to test Type I error control and power under DGP-analysis mismatch; precedent Gärtner 2023 and Blackston 2019.
- **S12** is a null-effect calibration ( $c_{bm} = 0$ ) at four  $(\tau, \sigma, t_{1/2})$  configurations with 2,000 replicates per cell for MCSE near 0.5%; precedent Morris 2019<sup>1</sup>.

### 3.7 Reproducibility and execution

[10]

**Per-replicate rather than pre-aggregated output is written, enabling Morris MCSE computation of all performance measures at render time.**

- Each sweep seeds via `set.seed(42 + sweep_id)` and uses `furrr` multisession parallelism.
- RDS artifacts are written to `analysis/data/derived_data/sensitivity/` and orchestrated by `run-all.R`.
- Per-replicate storage allows power, bias, MSE, empirical SE, and Type I error to be computed with MCSE at render time.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

Each sweep sources `00-common.R` for fixtures, sets a seed via

348 `set.seed(42 + sweep_id)`, constructs its parameter grid, calls  
349 `generate_simulated_results()` with `furrr`-based multisession  
350 parallelism over `parallel::detectCores() - 1` workers, and writes  
351 an RDS artifact to `analysis/data/derived_data/sensitivity/`.  
352 The orchestrator `run-all.R` runs all twelve in the order specified in  
353 the plan. Per-replicate rows rather than pre-aggregated summaries  
354 are written, so that power, bias, MSE, empirical SE, and Type I error  
355 can be computed at render time with Morris MCSE. Sweep scripts  
356 are in `analysis/scripts/nof1-design-sensitivity/`.

## 357 4 Results

358 **Note on numerical values in this section.** The numerical  
359 results quoted in the prose below were generated at a previous  
360 baseline of  $N = 40$  participants per design. Re-running the twelve  
361 sweeps at the revised  $N = 35$  baseline is in progress; the figures  
362 and quoted numerics will be updated when the sweep artifacts at  
363 `analysis/data/derived_data/sensitivity/` are regenerated. The  
364 qualitative direction of every finding below is expected to be preserved  
365 by the rebaseline.

366 [11]

367 **All twelve sweeps completed in approximately 32 min-**  
368 **utes; per-cell metrics are computed at render time from**  
369 **per-replicate RDS artifacts.**

- 370 • Wall-clock time was approximately 32 minutes on 8 cores with  
371 `furrr` multisession.
- 372 • Summary metrics are derived at render time from `analysis/data/derived_data/sensitivity/`.
- 373 • The target estimand throughout is the Dbc main drug effect co-  
374 efficient.

375 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*  
376 *tempor incididunt ut labore et dolore magna aliqua.*

377 All twelve sweeps completed in approximately 32 minutes of wall-  
378 clock time on an 8-core machine with the `furrr` multisession backend.  
379 Per-cell summary metrics are computed from the RDS artifacts at  
380 render time; numerical values quoted below are read directly from  
381 `analysis/data/derived_data/sensitivity/`. The target coefficient  
382 is the main drug effect (the Dbc coefficient in the `nlme::lme` fixed-  
383 effects formula).

384 [12]

385 **The hybrid is the primary design; the alternating A/B/A/B**  
386 **variant is retained as a sensitivity comparator for sweeps that**  
387 **vary cycle count or period length.**

- The hybrid’s fixed eight-timepoint structure does not scale with cycle count or period length.
- The alternating variant is used in S6 and S9 where that scalability is required.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The primary N-of-1 design is the Hendrickson 2020<sup>5</sup> hybrid (open-label titration followed by blinded discontinuation and brief crossover), labeled ‘Nof1’ in the sweep artifacts and figures below. The alternating A/B/A/B aggregated N-of-1 (labeled ‘Nof1-alt’) is retained as a sensitivity comparison at sweeps that vary the cycle count (S6) and period length (S9), where the hybrid’s fixed eight-timepoint structure does not naturally scale.

[13]

**Across sweeps, the hybrid outperforms the parallel RCT on sample size, variance robustness, and serial-correlation tolerance; the RCT is more robust to long carryover.**

- The hybrid reaches 80% power at smaller sample sizes, under higher within-patient variance, and under stronger AR(1) than the parallel RCT.
- The parallel RCT is more robust to carryover exceeding the within-subject period length and to decay-form misspecification.
- Type I error is near nominal for the hybrid and alternating variant; the two-period crossover shows robust inflation of 0.10 to 0.13.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

We begin with a cross-sweep summary. At the prazosin-like baseline (effect size approximately  $-2$  nightmares per week, carryover half-life 3 days, 35 participants, 8 weekly timepoints), the hybrid N-of-1 and the parallel RCT both perform competently, with the hybrid reaching 80% power at a smaller sample size, under larger within-patient variance, and under higher serial correlation than the parallel RCT. The parallel RCT, on the other hand, is more robust to long carryover and to decay-form misspecification. Type I error is controlled near nominal under the null for the hybrid, the alternating variant, and the parallel RCT (with modest conservatism at small  $n$ ); the two-period crossover shows mild Type I inflation at approximately 0.10 to 0.13, a finding that warrants caution when that design is used at small per-subject timepoint counts.

#### 4.1 S1. Effect-size sweep




Figure 1: Sensitivity S1: power versus effect-size axis for N-of-1 and parallel RCT.

[14]

**The hybrid reaches 80% power at  $\delta \approx -1$ , a half-unit smaller effect size than the parallel RCT requires.**

- Hybrid power rises from 0.066 at the null to saturation at 1.000 for  $|\delta| \geq 2$ .
- Parallel RCT power follows the same shape at a lower level, reaching 0.941 at  $\delta = -2$ .
- The hybrid's 80%-power effect-size threshold is approximately 0.5 units smaller than the RCT's.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The effect-size sweep varies the bio-response maximum  $\max[\text{br}]$  across a seven-point grid  $\{0, 0.5, 1.0, 1.5, 2.0, 2.5, 3.0\}$  at  $N = 35$ ,  $c_{bm} = 0.3$ , and carryover half-life 3 days. Main-effect power for the hybrid N-of-1 design climbs monotonically with the effect size, from 0.066 (MCSE 0.008) at  $\delta = 0$  (Type I) to 0.350 at  $\delta = -0.5$ , 0.859 at  $\delta = -1$ , 0.996 at  $\delta = -1.5$ , and saturating at 1.000 for  $|\delta| \geq 2$ . The parallel RCT tracks the same shape at a lower level: 0.034 (MCSE 0.006) at the null, 0.177 at  $\delta = -0.5$ , 0.470 at  $\delta = -1$ , 0.793 at  $\delta = -1.5$ , and 0.941 at  $\delta = -2$ . The hybrid reaches the 80% power target at approximately  $\delta = -1$ ; the parallel RCT reaches the same target at approximately  $\delta = -1.5$ , a half-unit shift.

[15]

**Type I error is marginally above nominal for the hybrid and marginally below for the RCT; neither is grossly miscali-**

brated, but the hybrid’s upward departure warrants attention.

- Hybrid Type I is 0.066 (95% CI [0.050, 0.082]); parallel RCT Type I is 0.034 (95% CI [0.022, 0.046]).
- Both confidence intervals cover 0.05 at the margin.
- The hybrid’s modest upward departure merits caution when the study is powered near the 5% boundary.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The direction and magnitude of the hybrid’s advantage match those reported by Blackston 2019<sup>4</sup>, whose simulation shows aggregated N-of-1 reaching adequate power at a smaller effect size than a matched- $n$  parallel RCT under the same variance structures. Type I error at  $\delta = 0$  is 0.066 for the hybrid (MCSE 0.008, 95% Wald CI [0.050, 0.082]) and 0.034 for the parallel RCT (MCSE 0.006, [0.022, 0.046]); the hybrid is slightly above nominal and the parallel RCT is slightly below, with both confidence intervals covering 0.05 at the margin. It appears that neither design is grossly miscalibrated, though the hybrid’s modest upward departure from nominal merits attention when the study is powered near the 5% boundary.

## 4.2 S2. Carryover half-life sweep




Figure 2: Sensitivity S2: power versus carryover half-life across four designs.

[16]

The hybrid tolerates carryover up to 5 days before power erodes; the alternating variant degrades from 0.90 to 0.30 across the half-life range.

- Hybrid power is 1.000 through  $t_{1/2} = 5$  days, declining to 0.736 at 14 days.
- The alternating N-of-1 is more sensitive, dropping from 0.896 to 0.304 across the same range.
- The parallel RCT is essentially immune, with power flat at 0.92 to 0.96 across the grid.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The carryover half-life sweep varies  $t_{1/2} \in \{0, 0.5, 1, 2, 3, 5, 7, 10, 14\}$  days across the four designs at the prazosin-like baseline. Main-effect power for the hybrid N-of-1 design saturates at 1.000 from  $t_{1/2} = 0$  through  $t_{1/2} = 5$  days, then declines: 0.976 at 7 days, 0.900 at 10 days, and 0.736 at 14 days (7-day periods). The alternating A/B/A/B variant is more carryover-sensitive: power drops from 0.896 at  $t_{1/2} = 0$  to 0.778 at  $t_{1/2} = 3$  to 0.304 at  $t_{1/2} = 14$ . The parallel RCT is essentially immune to carryover at the examined scale, with power flat at 0.92 to 0.96 across the grid. The two-period crossover is also near-saturated at 0.98 to 1.00.

[17]

**The hybrid’s carryover robustness is attributable to its open-label titration phase; the RCT ordering reversal begins when  $t_{1/2}$  exceeds the period length.**

- The four-week open-label titration accumulates treatment signal that is not erased by the post-discontinuation off-drug period.
- The alternating design relies on repeated within-subject transitions and sheds signal proportional to the unwashed residual.
- The hybrid-versus-RCT power ordering crossover begins at approximately  $t_{1/2} = 10$  days, consistent with Blackston 2019.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The hybrid’s carryover robustness relative to the alternating variant reflects its operational structure. That is, the four-week open-label titration phase contributes substantial treatment-accumulated signal that is not erased by post-discontinuation off-drug period, whereas the alternating design relies on repeated within-subject transitions that each shed signal proportional to the unwashed residual. The flat parallel-RCT and crossover curves confirm Blackston 2019’s<sup>4</sup> observation that between-subject comparisons are largely immaterial to carryover at the within-period scale; the loss of effective sample size in within-subject designs when  $t_{1/2}$  exceeds the period length is the primary mechanism of the hybrid-versus-RCT power ordering crossover that begins at approximately  $t_{1/2} = 10$  days.

### 521 4.3 S3. Decay-form misspecification sweep




Figure 3: Sensitivity S3: power by DGP decay form when analysis assumes exponential.

522 [18]

523 **The decay-form sweep uses the alternating variant because**  
 524 **the hybrid’s multi-phase `tsd` pattern produces non-PD ma-**  
 525 **trices under the power decay form.**

- 526 • Five DGP decay forms are tested while the analysis always as-
- 527 sumes exponential decay.
- 528 • The hybrid is excluded because its multi-phase `tsd` structure in-
- 529 teracts poorly with the power form.
- 530 • The alternating variant provides a clean test of decay-form mis-
- 531 specification sensitivity.

532 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*  
 533 *tempor incididunt ut labore et dolore magna aliqua.*

534 The decay-form sweep varies the DGP carryover form across  
 535 {exponential, linear, Weibull ( $k = 0.8$ ), Weibull ( $k = 1.5$ ), power}  
 536 while the analysis model always assumes exponential decay in the  
 537 construction of the continuous drug indicator `Dbc`. The sweep uses  
 538 the alternating A/B/A/B N-of-1 variant because the hybrid’s specific  
 539 multi-phase `tsd` pattern interacts poorly with the ‘power’ form,  
 540 producing non-positive-definite covariance matrices at certain cells,  
 541 and is less informative for decay-form misspecification testing in any  
 542 case.

543 [19]

544 **Main-effect power ranges only 15 percentage points across**

545 **decay forms, substantially less sensitive to misspecification**  
 546 **than the interaction estimand.**

- 547 • Power ranges from 0.754 (Weibull  $k = 0.8$ ) to 0.904 (linear) under  
 548 the alternating N-of-1.
- 549 • The 15-point range is smaller than under the interaction target,  
 550 indicating less sensitivity to decay-form misspecification.
- 551 • The parallel RCT is flat at 0.93 to 0.95 across all forms, consistent  
 552 with Gärtner 2023.

553 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*  
 554 *tempor incididunt ut labore et dolore magna aliqua.*

555 Under the alternating N-of-1 design, main-effect power ranges from  
 556 0.754 (Weibull  $k = 0.8$ ) to 0.904 (linear), with the baseline exponential  
 557 case at 0.816 and the power form at 0.870. The Weibull  $k = 1.5$  form  
 558 yields 0.852. The 15-percentage-point range across decay forms is  
 559 smaller than under the interaction target examined in earlier work,  
 560 suggesting that main-effect estimation is relatively less sensitive to  
 561 decay-form misspecification than interaction estimation. The parallel  
 562 RCT is essentially flat at 0.93 to 0.95 across all forms, consistent with  
 563 the finding of Gärtner 2023<sup>7</sup> that misspecification of the carryover  
 564 decay form carries a design-specific cost affecting only within-subject  
 565 designs.

#### 566 4.4 S4. Between-patient heterogeneity sweep




Figure 4: Sensitivity S4: power versus between-patient baseline SD.

567 [20]

568 **The hybrid is saturated across the  $\tau$  range; the parallel RCT**  
 569 **loses 12 percentage points as between-patient heterogeneity**



doubles.

- Hybrid power is 1.000 at all six  $\tau$  levels with stable mean  $\beta \approx -0.94$ .
- Parallel RCT power declines from 0.998 at  $\tau = 0.5$  to 0.874 at  $\tau = 2.0$ .
- The RCT loses approximately 12 percentage points as  $\tau$  doubles from 1.0 to 2.0.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The between-patient heterogeneity sweep varies the baseline-symptom SD (the BL factor standard deviation, the  $\tau$  component in the Senn 2018 decomposition) across  $\{0.5, 0.75, 1.0, 1.25, 1.5, 2.0\}$ . Under the hybrid N-of-1, main-effect power is saturated at 1.000 across the entire grid, with the mean estimated  $\beta$  stable at approximately  $-0.94$  across the six levels. The parallel RCT degrades substantially with  $\tau$ , from 0.998 at  $\tau = 0.5$  to 0.994 at 0.75, 0.978 at 1.0, 0.958 at 1.25, 0.946 at 1.5, and 0.874 at 2.0 (the latter with MCSE 0.015). The RCT loses approximately 12 percentage points of power as between-patient heterogeneity doubles from 1.0 to 2.0.

[21]

**The Senn 2018 variance-decomposition prediction is reproduced:  $\tau^2$  enters the RCT estimator but is absorbed by the within-subject contrast in N-of-1 designs.**

- Between-patient heterogeneity contributes to  $\text{Var}(\hat{\Delta})$  in the parallel RCT but not in the within-subject N-of-1 contrast.
- The hybrid's saturation reflects the prazosin-calibrated baseline where effect size is large relative to within-patient variance.
- At smaller effect sizes or larger  $\sigma$ , the N-of-1 advantage would widen rather than flatten as  $\tau$  increases.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

These results reproduce the Senn 2018<sup>8</sup> variance-decomposition argument: between-patient heterogeneity enters the parallel-RCT variance ( $\tau^2$  is part of  $\text{Var}(\hat{\Delta})$ ) but is absorbed by the within-subject contrast in aggregated N-of-1 designs. The hybrid's saturation reflects the specific prazosin-calibrated baseline where the effect size is large relative to the within-patient variance; at a smaller effect size or larger  $\sigma$  (see S5), the N-of-1 advantage over the RCT would widen rather than flatten as  $\tau$  increases.

## 4.5 S5. Within-patient variance sweep




Figure 5: Sensitivity S5: power versus within-patient bio-response SD.

[22]

**The hybrid is essentially saturated across the within-patient variance range; the parallel RCT loses approximately 18 percentage points.**

- Hybrid power declines from 1.000 to 0.986 across  $\sigma \in \{0.5, \dots, 1.5\}$ , a one-percentage-point loss.
- Parallel RCT power declines from 0.996 to 0.814 across the same range, an 18-point loss.
- The ordering is consistent with the N-of-1 design summing  $k$  paired contrasts per patient.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The within-patient variance sweep varies the bio-response SD across  $\{0.5, 0.75, 1.0, 1.25, 1.5\}$ . Under the hybrid N-of-1, main-effect power is 1.000 at  $\sigma \in \{0.5, 0.75, 1.0\}$ , 0.996 at  $\sigma = 1.25$ , and 0.986 at  $\sigma = 1.5$ , essentially saturated across the range. The parallel RCT declines more visibly: 0.996 at  $\sigma = 0.5$ , 0.982 at 0.75, 0.942 at 1.0, 0.868 at 1.25, and 0.814 at 1.5. Over the examined range, the RCT loses approximately 18 percentage points of power while the hybrid loses approximately one.

[23]

**The hybrid's quasi-saturation is specific to the prazosin-calibrated regime; at smaller effect sizes the  $\sigma$  dependence would be steeper.**

- The eight-timepoint structure plus open-label titration provides sufficient signal accumulation to absorb substantial within-patient noise.
- The Senn 2018 prediction is reproduced: N-of-1 absorbs within-patient variance via  $k$  paired contrasts.
- Saturation should not be extrapolated beyond the prazosin-calibrated effect-size regime.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

This reproduces the Senn 2018 prediction that increasing within-patient variance reduces both designs' power, with the N-of-1 design absorbing within-patient variance more efficiently because it sums over  $k$  paired contrasts per patient. The hybrid's quasi-saturation indicates that the eight-timepoint structure plus the open-label titration phase provides sufficient signal accumulation to tolerate substantial within-patient noise; at smaller effect sizes, the  $\sigma$  dependence would be steeper, and this saturation should not be extrapolated beyond the prazosin-calibrated regime.

#### 4.6 S6. Cycles-per-patient sweep




Figure 6: Sensitivity S6: power versus timepoints per patient  $k$ .

[24]

**Alternating N-of-1 power rises concavely with  $k$ , from 0.294 at  $k = 2$  to 0.980 at  $k = 16$ , as predicted by the Senn 2018 decomposition.**

- The alternating variant is used because the hybrid's fixed structure does not scale with  $k$ .

- Power rises from 0.294 at  $k = 2$  to 0.804 at  $k = 8$  (baseline) to 0.980 at  $k = 16$ .
- The concave shape is consistent with the Senn 2018 variance-decomposition prediction.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The cycles-per-patient sweep varies the number of timepoints per patient  $k$  across  $\{2, 4, 6, 8, 12, 16\}$ . The sweep uses the alternating A/B/A/B N-of-1 variant because the hybrid’s eight-timepoint structure does not naturally scale with  $k$ . Under the alternating N-of-1 design, main-effect power increases from 0.294 at  $k = 2$  to 0.486 at  $k = 4$ , 0.658 at  $k = 6$ , 0.804 at  $k = 8$  (the baseline), 0.910 at  $k = 12$ , and 0.980 at  $k = 16$ . The curve is concave as predicted by the Senn 2018<sup>8</sup> decomposition.

[25]

**For the main-effect target, increasing  $k$  slightly favors the RCT; the hybrid’s advantage over the alternating variant comes from the open-label phase, not cycle count.**

- At  $k = 2$  the alternating N-of-1 and the RCT are in a dead heat; at  $k \geq 6$  the RCT surpasses the alternating N-of-1 by up to 14 points.
- Each additional timepoint contributes a between-subject observation for the RCT but averages into already-accumulated paired contrasts for the alternating N-of-1.
- The hybrid’s primary power advantage over the alternating design arises from the open-label titration phase, not from the number of cycles.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The parallel RCT tracks a similar shape: 0.292 at  $k = 2$ , 0.502 at  $k = 4$ , 0.772 at  $k = 6$ , 0.940 at  $k = 8$ , 0.998 at  $k = 12$ , and 1.000 at  $k = 16$ . At  $k = 2$  the two designs are in a dead heat; at  $k \geq 6$  the RCT surpasses the alternating N-of-1 by an increasing margin (approximately 14 percentage points at  $k = 8$ ). For the main-effect target, increasing  $k$  benefits both designs and slightly favors the RCT because each additional timepoint contributes one more between-subject observation, whereas the alternating-N-of-1 paired contrasts are already averaged in at  $k = 8$ . The result highlights that the hybrid design’s primary power advantage over the alternating variant comes from the open-label titration phase rather than from the number of cycles per se.

## 4.7 S7. Patient-count sweep




Figure 7: Sensitivity S7: power versus total participant count.

[26]

**The hybrid saturates at high power by  $N = 16$ ; the parallel RCT reaches the same threshold between  $N = 24$  and  $N = 35$ .**

- Hybrid power is high at modest sample sizes and saturates rapidly with  $N$ .
- The parallel RCT follows a similar but offset trajectory, reaching 80% power between  $N = 24$  and  $N = 35$ .
- The grid anchors the project-standard  $N = 35$  minimum for cross-paper comparison.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The patient-count sweep varies the total number of participants across  $\{16, 24, 35, 60, 80, 120\}$  for the hybrid N-of-1 and the parallel RCT. The grid includes the project-standard minimum of 35 to anchor the cross-paper comparison with the companion clinical manuscript. Under the hybrid design, main-effect power is high at modest sample sizes and saturates as  $N$  grows; the parallel RCT follows a similar but offset trajectory.

[27]

**The hybrid achieves an equivalent power target at approximately one-third to one-half the parallel-RCT sample size, consistent with Blackston 2019 and Carrozzo 2025.**

- The hybrid reaches 80% power at  $N \leq 16$ ; the parallel RCT reaches it between  $N = 24$  and  $N = 35$ .

- The sample-size ratio of one-third to one-half matches the Blackston 2019 N-of-1 efficiency advantage.
- This is consistent with Carrozzo 2025, in which the meta-analytic N-of-1 required fewer participants across most heterogeneity and carryover settings.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The hybrid reaches 80% power at  $N = 16$  or fewer participants; the parallel RCT reaches the same threshold somewhere between  $N = 24$  and  $N = 35$ . In the framing of Blackston 2019<sup>4</sup>, the hybrid achieves an equivalent power target at approximately one-third to one-half the sample size of the parallel RCT. This is consistent with the Carrozzo 2025<sup>6</sup> result that the meta-analytic N-of-1 procedure reached 80% power with fewer participants than the two-arm RCT across most heterogeneity and carryover settings.

#### 4.8 S8. AR(1) serial-correlation sweep




Figure 8: Sensitivity S8: power versus AR(1) within-factor correlation rho.

[28]

**The hybrid is insensitive to AR(1) correlation; the parallel RCT loses 24 percentage points of power as  $\rho$  rises from 0 to 0.9.**

- Hybrid power is 1.000 across  $\rho \in \{0, 0.3, 0.5, 0.7, 0.9\}$ .
- Parallel RCT power declines from 0.984 to 0.746, and mean  $\hat{\beta}$  shrinks from  $-1.01$  to  $-0.57$ .
- The RCT loss is monotonic with  $\rho$  across the examined range.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The AR(1) sweep varies the within-factor autocorrelation  $\rho \in \{0, 0.3, 0.5, 0.7, 0.9\}$ . Under the hybrid N-of-1 design, main-effect power is saturated at 1.000 across the entire range of  $\rho$ . Under the parallel RCT, power decreases monotonically with  $\rho$ : 0.984 at  $\rho = 0$ , 0.934 at 0.3, 0.872 at 0.5, 0.772 at 0.7, and 0.746 at 0.9. The mean estimated  $\beta$  also shrinks for the RCT (from  $-1.01$  at  $\rho = 0$  to  $-0.57$  at  $\rho = 0.9$ ).

[29]

**The AR(1) asymmetry arises because the corCAR1 structure absorbs correlation within subjects but cannot manufacture the missing within-subject contrasts for the parallel RCT.**

- Strong AR(1) produces coherent within-subject trajectories for the hybrid, absorbed by corCAR1 with negligible power loss.
- For the RCT, strong AR(1) inflates between-subject mean contrast variance via effective-sample-size shrinkage.
- This is consistent with Wang and Schork 2019: N-of-1 designs are insensitive to serial correlation, while parallel-group designs are not.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The direction of the AR(1) effect is opposite for the two designs, and this asymmetry is worth examining carefully. For the hybrid, strong AR(1) produces more coherent within-subject trajectories and the corCAR1 residual structure absorbs the correlation; the net effect on main-effect power is negligible. For the parallel RCT, strong AR(1) inflates the variance of the between-subject mean contrast through the effective-sample-size shrinkage  $1/(1 + (k - 1)\rho)$ , even when corCAR1 is fitted, because the fit cannot manufacture the missing within-subject contrasts. This is consistent with the finding of Wang and Schork 2019<sup>10</sup> that, under correct modeling, the N-of-1 design benefits from strong serial correlation (or is at least insensitive to it), while the parallel-group design remains relatively sensitive.

## 4.9 S9. Period-length sweep

[30]

**Main-effect power increases monotonically with period length, from 0.40 at 3 days to 1.00 at 14 days; carryover half-life has only a modest effect at fixed period length.**

- Power rises from 0.40 to 0.44 at 3-day periods to 0.997 to 1.00 at




Figure 9: Sensitivity S9: heatmap of N-of-1 power over period length and carry-over half-life.

788

14-day periods.

789

- At fixed period length, varying carryover half-life across 0.5 to 7 days produces only modest power differences.

790

791

- The sweep uses the alternating variant, as the hybrid's fixed structure does not scale with period length.

792

793

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

794

795

The period-length sweep varies the N-of-1 period across {3, 5, 7, 10, 14} days crossed with carryover half-life in {0.5, 3, 7} days. The sweep uses the alternating A/B/A/B variant. Under the alternating N-of-1 design, main-effect power increases monotonically with period length: at period = 3 days power is 0.40 to 0.44 across the three half-lives, at 5 days 0.72 to 0.76, at 7 days 0.88 to 0.91, at 10 days 0.993 to 0.997, and at 14 days 0.997 to 1.00. At fixed period length, the carryover half-life has only a modest effect.

796

797

798

799

800

801

802

803

[31]

804

**For the main-effect target, longer periods are preferred, reversing the interaction-target recommendation; carryover should not exceed approximately 50% of the period length.**

805

806

807

- Longer periods accumulate more Gompertz drug-response signal per timepoint.

808

809

- The practical guideline is the reverse of the interaction-target recommendation: longer periods benefit main-effect detection.

810

811

- Wang and Schork 2019 note this tradeoff explicitly: frequent al-



812            ternation benefits correlation-modeled analyses but not absolute-  
813            effect inference.

814            *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*  
815            *tempor incididunt ut labore et dolore magna aliqua.*

816            Under the main-effect target, each period accumulates more Gompertz  
817            drug-response signal at longer periods, so longer periods increase the  
818            per-timepoint effect magnitude. The pragmatic guideline for the main-  
819            effect target is therefore the reverse of the interaction-target recom-  
820            mendation: longer periods are preferred when the primary goal is to  
821            detect the main drug effect, subject to carryover not exceeding the  
822            period length by more than approximately 50%. Wang and Schork  
823            2019<sup>10</sup> note this tradeoff explicitly, observing that more frequent al-  
824            ternation benefits correlation-modeled analyses but is problematic for  
825            absolute-effect inference.

#### 826   4.10   S10. Biomarker interaction strength crossed with DGP ar- 827            chitecture




Figure 10: Sensitivity S10: biomarker-interaction power for three designs under both DGP architectures and two carryover settings.

828            [32]

829            **Main-effect power is invariant to  $c_{bm}$  and DGP architecture,**  
830            **confirming that the main drug effect is determined by Gom-**  
831            **pertz accumulation, not biomarker coupling.**

- 832            • Hybrid power is 1.000 across all 16  $c_{bm}$ -by-architecture-by-half-life  
833            combinations.

- Alternating N-of-1 power varies from 0.682 to 0.916; parallel RCT power varies from 0.914 to 0.958.
- The flat main-effect power across  $c_{bm}$  values provides a robustness check on the DGP implementation.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The biomarker-interaction sweep varies  $c_{bm} \in \{0, 0.2, 0.4, 0.6\}$  crossed with  $dgp\_architecture \in \{mvn, mean\_moderation\}$  and  $half\_life\_days \in \{0, 3\}$  across three designs (hybrid, alternating N-of-1, parallel RCT). Under the main-effect target: hybrid main-effect power is saturated at 1.000 across all 16 combinations, with the mean estimated  $\beta$  stable at  $-0.92$  to  $-1.08$ ; alternating N-of-1 power varies from 0.682 to 0.916 across the 16 cells; and parallel RCT power is 0.914 to 0.958 across all 16 cells.

[33]

**The flat  $c_{bm}$  power curves are a DGP implementation check: when the main effect is the target, biomarker coupling should not affect power, and it does not.**

- The main drug effect is determined by Gompertz max[br] and on-drug accumulation, not by biomarker coupling.
- Approximately flat power across  $c_{bm}$  for all three designs confirms the DGP correctly separates the two signal sources.
- The reader should not conflate this main-effect invariance with the biomarker-interaction target, where  $c_{bm}$  is the primary power driver.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The invariance of main-effect power to  $c_{bm}$  and architecture is expected, and indeed it is reassuring rather than surprising: the main drug effect is determined by the Gompertz max[br] and the trial's on-drug accumulation, not by the biomarker coupling. This provides a robustness check on the DGP implementation in that, when the primary-analysis target is the main effect, the  $c_{bm}$  sweep should produce approximately flat power curves for all designs, and we observe precisely that pattern here.

#### 4.11 S11. Misspecification factorial

[34]

**The parallel RCT shows statistically robust Type I inflation of 0.112 to 0.126 under high DGP AR(1); the hybrid is controlled near nominal across all four cells.**




Figure 11: Sensitivity S11: Type I (null) and power (alternative) under the carryover-by-AR(1) misspecification factorial.

- Hybrid null Type I is 0.036 to 0.056, within 2 MCSE of nominal across all four DGP combinations.
- Parallel RCT Type I is 0.018 to 0.024 (conservative) when no AR(1) is present, and 0.112 to 0.126 (inflated) under  $\rho = 0.7$ .
- The inflation is statistically robust at MCSE 0.014 to 0.015.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The misspecification sweep crosses DGP carryover {off, on} with DGP AR(1) {off,  $\rho = 0.7$ } at null ( $\text{br\_max} = 0$ ) and alternative ( $\text{br\_max} = 1$ ) levels. Under the null, Type I error is within 2 MCSE of the nominal 0.05 for the hybrid N-of-1 (0.036 to 0.056 across the four cells). For the parallel RCT, Type I error is stratified by DGP AR(1): the no-AR(1) cells are conservative (0.018 to 0.024) while the high-AR(1) cells are inflated above nominal (0.112 to 0.126), with MCSE 0.014 to 0.015, so the inflation is statistically robust. Under the alternative, the hybrid is saturated at 1.000 in all four cells; the parallel RCT power is 0.986 to 0.990 when no AR(1) is present, and 0.766 to 0.772 when AR(1)  $\rho = 0.7$  is present.

[35]

**The RCT Type I inflation arises because the between-subject variance estimator understates the true standard error when DGP serial correlation is strong.**

- Under strong DGP AR(1), `corCAR1` cannot manufacture the miss-

- ing within-subject contrasts for a between-subject comparison.
- The hybrid is immune because its primary treatment contrast is within-subject.
  - The mechanism is analogous to the Blackston 2019 carryover Type I inflation finding, though the source here is AR(1) rather than carryover.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The inflated RCT Type I under high DGP AR(1) is an important caveat that merits comment. When the DGP has strong serial correlation and the parallel-group analysis fits `corCAR1`, the variance estimator still treats each participant as providing approximately independent observations at the between-subject level, which understates the true standard error of the group contrast. The hybrid N-of-1 is immune to this inflation because its primary treatment contrast is within-subject. Blackston 2019’s<sup>4</sup> observation that moderate-to-strong carryover inflates Type I error when not modeled is not reproduced here for the hybrid (because the decaying Dbc in the analysis absorbs carryover correctly) but is analogous in mechanism to the AR(1) inflation observed here for the parallel RCT.

#### 4.12 S12. Null-effect calibration




Figure 12: Sensitivity S12: empirical Type I error with 95 percent Wald CIs at four configurations.

[36]

**Hybrid and alternating N-of-1 Type I error is near nominal; the two-period crossover shows robust inflation of 0.100 to**

### 0.130 attributable to underidentified corCAR1.

- Hybrid Type I: [0.056, 0.062]; alternating N-of-1: [0.057, 0.064]; both within 2 MCSE of nominal.
- Parallel RCT: [0.029, 0.060], three cells conservative and one nominal.
- Two-period crossover: [0.100, 0.130], statistically robust inflation at MCSE 0.006 to 0.007.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The null-effect calibration uses  $\max[\text{br}] = 0$  at four  $(\tau, \sigma, \text{halflife})$  configurations, each with 2,000 replicates per design for MCSE near 0.005. Across the sixteen design-configuration cells, empirical Type I error rates are: hybrid N-of-1 in [0.056, 0.062] (four cells, all within 2 MCSE of nominal), alternating N-of-1 in [0.057, 0.064] (four cells, all within 2 MCSE of nominal), parallel RCT in [0.029, 0.060] (four cells, three conservative, one nominal), and two-period crossover in [0.100, 0.130] (four cells, all modestly inflated with MCSEs of 0.006 to 0.007, so the inflation is statistically robust at 2,000 replicates).

[37]

**The two-period crossover should not be analyzed with corCAR1; at two within-subject timepoints the parameter is underidentified and the rejection rate inflates by a practically meaningful amount.**

- Two within-subject timepoints cannot identify the corCAR1 parameter reliably.
- The inflation of 0.100 to 0.130 is statistically robust and practically meaningful.
- Alternative analyses (paired  $t$ -test on change scores or lme with unstructured residual) are recommended for this design.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The two-period crossover's modest inflation appears attributable to the small number of timepoints (two per subject) producing an underidentified corCAR1 residual correlation. Two-period crossovers should be analyzed without serial-correlation structure, either via a paired  $t$ -test on change scores or via an lme with an unstructured residual; fitting corCAR1 at only two within-subject timepoints is problematic and, as we see here, inflates the rejection rate by a practically meaningful amount.

## 5 Discussion

### 5.1 Synthesis

[38]

**Five substantive findings organize the twelve sweeps: sample-size advantage, variance robustness, carryover design-dependence, AR(1) asymmetry, and Type I error calibration.**

- The hybrid reaches 80% power at  $N \leq 16$ ; the parallel RCT requires  $N = 24$  to 35, a one-third to one-half sample-size ratio.
- The hybrid is saturated across  $\sigma$  and  $\tau$  sweeps; the RCT degrades by 18 and 12 percentage points respectively.
- Carryover sensitivity is design-specific: the hybrid tolerates  $t_{1/2} \leq 5$  days; the alternating variant drops from 0.90 to 0.30 across the range; the RCT is immune.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

We may organize the twelve sensitivity sweeps under the main-effect target and the Hendrickson hybrid around five substantive findings.

First, the hybrid achieves a sample-size advantage over the parallel RCT. At the prazosin-calibrated baseline, the hybrid reaches the 80% power target at  $N = 16$  or fewer participants, whereas the parallel RCT reaches the same target somewhere between  $N = 24$  and  $N = 35$  (S7). The sample-size ratio of roughly one-third to one-half matches the N-of-1 efficiency advantage reported by Blackston 2019<sup>4</sup> and Carrozzo 2025<sup>6</sup>.

Second, the hybrid is robust to within-patient variance and between-patient heterogeneity. The hybrid is essentially saturated across the sweeps of  $\sigma$  (S5) and  $\tau$  (S4), while the parallel RCT degrades by 18 percentage points across the  $\sigma$  range and 12 percentage points across the  $\tau$  range, reproducing the Senn 2018<sup>8</sup> variance-decomposition prediction.

Third, sensitivity to carryover depends on design variant. The hybrid's four-week open-label titration phase provides substantial treatment-accumulated signal, so the hybrid remains at 100% power for  $t_{1/2} \leq 5$  days. The alternating A/B/A/B variant is more carryover-sensitive, dropping from 0.90 to 0.30 across the  $t_{1/2} = 0$  to  $t_{1/2} = 14$  range. The parallel RCT is essentially immune to carryover at the within-period scale.

Fourth, AR(1) serial correlation affects the two designs differently. The hybrid saturates across  $\rho = 0$  to 0.9 (S8), while the parallel RCT

power drops from 0.98 to 0.75 across the same range and Type I error inflates from 0.018 to 0.126 under high- $\rho$  DGP (S11). This is consistent with the prediction of Wang and Schork 2019<sup>10</sup> that within-subject designs with corCAR1 fitted absorb serial correlation efficiently, while between-subject designs do not.

Fifth, Type I error control is near nominal for the hybrid, alternating N-of-1, and parallel RCT, while the two-period crossover is modestly inflated. Specifically, hybrid Type I falls in [0.056, 0.062], alternating N-of-1 in [0.057, 0.064], parallel RCT in [0.029, 0.060], and two-period crossover in [0.100, 0.130].

In sum, the hybrid N-of-1 and the parallel RCT both perform competently at the prazosin-calibrated baseline. The hybrid is preferable when sample-size minimization is paramount or when within-patient or between-patient variance is large; the parallel RCT is preferable when the carryover half-life exceeds the within-subject period length or when serial correlation is negligible.

## 5.2 Comparison with prior simulation studies

### 5.2.1 Hendrickson et al. (2020): predictive biomarker validation

[39]

**Hendrickson 2020 showed hybrid power collapsing from 0.95 to 0.18 at the biomarker-interaction target under a 0.7-day carryover half-life.**

- At  $N = 35$ ,  $c_{bm} = 0.6$ , no carryover, biomarker-interaction power was 0.95 (hybrid), 0.96 (crossover, CO), 0.94 (open-label plus blinded discontinuation, OL+BDC), and 0.25 (open-label, OL).
- A half-life of only 0.7 days collapsed hybrid power to 0.18, demonstrating extreme carryover sensitivity at the interaction estimand.
- This contrasts sharply with the present study’s main-effect findings, where hybrid power is 1.000 through  $t_{1/2} = 5$  days.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

Hendrickson and colleagues<sup>5</sup> compared four designs at  $N \in \{35, 70\}$  with 1,000 replicates for detecting a continuous biomarker-treatment interaction. At  $N = 35$  with biomarker coupling  $c_{bm} = 0.6$  and no carryover, power was 0.25 (OL), 0.94 (OL+BDC), 0.96 (CO), and 0.95 (hybrid). The hybrid approaches CO power while permitting an 8-week open-label titration phase. A central Hendrickson finding is that, at the biomarker-interaction target, even short carryover half-lives erode hybrid and OL+BDC power precipitously: at  $N = 35$ ,

$c_{bm} = 0.6$ , and carryover half-life  $t_{1/2} = 0.1$  weeks (0.7 days), hybrid power fell from 0.95 to 0.18.

[40]

**The present study extends Hendrickson via Morris MCSE reporting, decay-form misspecification (S3), and main-effect carryover and period-length quantification (S2, S6, S9).**

- Design ranking at the biomarker-interaction target reproduces Hendrickson by construction.
- The key substantive divergence is estimand-specific: main-effect hybrid power is saturated through 5-day carryover, while interaction-target power collapses at under 1 day.
- Both findings are correct; they describe different estimands and must not be conflated.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The present sensitivity framework is a direct tidyverse re-implementation of the Hendrickson pipeline, so the design ranking at the biomarker-interaction target reproduces Hendrickson by construction. Where the present study extends Hendrickson is (i) in its Morris<sup>1</sup> MCSE reporting for every performance measure; (ii) in sensitivity S3, which crosses DGP decay form with the analysis-model assumption of exponential decay; and (iii) in sensitivities S2, S6, and S9, which quantify carryover and period-length effects for the main drug effect. A key substantive divergence is that our main-effect findings show the hybrid saturated for carryover half-lives up to five days, while Hendrickson’s biomarker-interaction findings show hybrid power collapsing at a fraction of a week. Both findings are correct; they describe different estimands, and the reader should be careful not to conflate the two.

### 5.2.2 Wang and Schork (2019): AR(1), heteroscedasticity, and period stacking

[41]

**Wang and Schork 2019 show roughly linear power erosion with  $\rho$  for single-period designs; period stacking substantially mitigates this loss.**

- For  $1p \times 2i$ , power erodes from 0.851 at  $\rho = 0$  to 0.126 at  $\rho = 0.75$ .
- For  $40p \times 2i$ , the same  $\rho$  range produces only a 0.17-point power loss, confirming period stacking as a mitigation.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*



Wang and Schork<sup>10</sup> provide analytical power formulas under a standard linear model with AR(1) residual correlation. Their Table 1 shows that for  $1p \times 2i$  with no washout, power is 0.851 at  $\rho = 0$ , 0.617 at  $\rho = 0.25$ , 0.330 at  $\rho = 0.5$ , and 0.126 at  $\rho = 0.75$ , a roughly linear erosion with increasing serial correlation. For  $40p \times 2i$ , power is 0.851, 0.745, 0.674, and 0.681 across the same  $\rho$  grid, confirming that period stacking mitigates the serial-correlation power loss.

[42]

**S8 is the aggregated-level analogue of Wang and Schork: the hybrid absorbs AR(1) via corCAR1; the parallel-RCT power loss directionally matches Wang and Schork’s between-period profile.**

- Hybrid power is saturated at 1.00 across  $\rho \in \{0, 0.3, 0.5, 0.7, 0.9\}$  at the aggregated level.
- Parallel RCT power drops from 0.98 to 0.75, directionally consistent with Wang and Schork’s between-period power loss.
- The S11 Type I inflation of 0.11 to 0.13 is consistent with the Wang and Schork mechanism, though magnitudes depend on small-sample instability of corCAR1 at  $N = 35$ .

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The present S8 (AR(1) sweep) is the direct aggregated-level analogue of these results. Our S8 reports hybrid N-of-1 power saturated at 1.00 across  $\rho \in \{0, 0.3, 0.5, 0.7, 0.9\}$  because the 35-participant hybrid analysis absorbs residual correlation via corCAR1 at the aggregated level. Our parallel RCT power drops from 0.98 at  $\rho = 0$  to 0.75 at  $\rho = 0.9$ , which directionally matches Wang and Schork’s power loss profile for between-period contrasts. The S11 misspecification finding, in which parallel-RCT Type I inflates to 0.11 to 0.13 under  $\rho = 0.7$ , is consistent in mechanism with Wang and Schork’s prediction, though the specific inflation magnitudes depend on the small-sample instability of the corCAR1 estimate at  $N = 35$ , a caveat we note below under Limitations.

### 5.3 Limitations

[43]

**Four limitations apply: hybrid saturation at the prazosin baseline, alternating-variant substitution in three sweeps, analysis-only corCAR1 in S8, and small-sample instability in the S11 RCT Type I estimate.**

- Saturation in S2, S4, S5, S8, S10, and S11 prevents quantifying the hybrid’s relative advantage in sub-ceiling regimes; sweeps at

- 1121 smaller effect sizes are the priority extension.
- 1122 • S3, S6, and S9 use the alternating variant because the hybrid can-
  - 1123 not scale with decay form, cycle count, or period length; results
  - 1124 apply to the alternating family, not directly to the hybrid.
  - 1125 • The S11 RCT Type I ranges (0.112 to 0.126 vs. 0.029 to 0.060 in
  - 1126 S12) are not contradictory; they reflect different DGP regimes.

1127 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*  
 1128 *tempor incididunt ut labore et dolore magna aliqua.*

1129 Four limitations should be noted before extending these results to trial  
 1130 planning. We begin with the most consequential.

1131 The hybrid N-of-1 saturates in many sweeps. The hybrid reaches 100%  
 1132 power in S4 (all  $\tau$  values), much of S5 ( $\sigma$ ), S8 (all  $\rho$ ), S10 (all biomarker-  
 1133 by-architecture cells), S11 (all alternative cells), and most of S2 (half-  
 1134 lives  $\leq 5$  days). Saturation obscures the relative ordering of designs  
 1135 and, in particular, prevents us from quantifying how much additional  
 1136 power the hybrid provides over the parallel RCT in the regimes where  
 1137 both are below ceiling. To resolve saturation, the sweeps would need to  
 1138 be repeated at a smaller effect size, dropping baseline power to the 0.50  
 1139 to 0.80 range where design differences are most clearly discriminated.  
 1140 This is perhaps the most important extension warranted by the present  
 1141 results.

1142 Sweeps S3, S6, and S9 use the alternating A/B/A/B variant rather  
 1143 than the hybrid. These three sweeps require the N-of-1 design to scale  
 1144 with a parameter that the hybrid's fixed eight-timepoint structure  
 1145 does not naturally accommodate. The results from these sweeps are  
 1146 therefore informative about the alternating design family specifically,  
 1147 and their implications for the hybrid must be treated as approximate  
 1148 at best.

1149 Sweep S8 was conducted with the analysis-side `corCAR1` always en-  
 1150 abled. A complementary with-versus-without-`corCAR1` factorial would  
 1151 be informative and would require adding an `op$use_corCAR1` option  
 1152 to `lme_analysis` followed by rerunning the sweep.

1153 Finally, the S11 RCT Type I inflation under high DGP AR(1) is partly  
 1154 a modeling artifact. The 0.11 to 0.13 Type I error rates for the par-  
 1155 allel RCT under  $\rho = 0.7$  reflect a between-subject analysis that fits  
 1156 `corCAR1` on a small sample, yielding an unstable correlation param-  
 1157 eter estimate. A better-specified parallel-RCT analysis, for example  
 1158 with a subject-specific random slope on time or with a fixed within-  
 1159 subject  $\rho$  of known magnitude, would likely restore nominal Type I.  
 1160 We note also that the S11 RCT Type I range (0.112 to 0.126) and the  
 1161 S12 RCT Type I range (0.029 to 0.060) are not in contradiction: S11  
 1162 crosses carryover presence with high DGP AR(1) under analysis-side  
 1163 `corCAR1` misspecification, whereas S12 is a pure null calibration with

no carryover and a more benign correlation regime. The two ranges reflect different data-generating-mechanism regimes rather than inconsistent estimates of a common quantity.

## 5.4 N-of-1 design variants not evaluated

[44]

**Several N-of-1 design variants are not evaluated; these fall into designs expected to be weaker than those reported here and designs expected to be stronger.**

- Unevaluated designs include single-subject, OL-only, OL+BDC-only, full multi-cycle crossover, complete crossover, sequential-monitoring, and randomized-block variants.
- Single-subject and reduced variants are expected to produce power at or below the reported numbers; results are an upper bound.
- Full multi-cycle, complete-design, and sequential-monitoring variants are expected to exceed the reported numbers; results are a lower bound.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

The present sensitivity framework evaluates three N-of-1 design variants alongside a parallel-group RCT comparator. Several N-of-1 design variants are not evaluated here: the single-subject N-of-1<sup>9</sup>, the open-label-only and OL-plus-BDC variants from Hendrickson 2020<sup>5</sup>, full multi-cycle crossover with more than two periods (Blackston 2019 used three cycles<sup>4</sup>), Williams-square and Latin-square complete N-of-1 designs, sequential-monitoring variants with adaptive stopping<sup>13,14</sup>, and randomized-block N-of-1.

[45]

**The present results bound the unevaluated variants: weaker designs are upper-bounded and stronger designs are lower-bounded by the reported numbers.**

- Simpler variants (single-subject, OL-only, OL+BDC-only, two-cycle crossover) are expected at or below the reported power.
- More powerful variants (multi-cycle, complete-design, sequential-monitoring, randomized-block) are expected at or above the reported power.
- Extending the present framework to cover these variants is a tractable implementation task rather than a conceptual redesign.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

These unevaluated design variants fall into two classes. The single-subject, open-label-only, OL+BDC-only, and two-cycle crossover designs are expected to produce power at or below the numbers reported here; our results are therefore an upper bound on those designs under matched parameters. The full multi-cycle crossover, complete-design crossover, sequential-monitoring variants, and randomized-block designs are expected to produce power at or above the numbers reported here; our results are therefore a lower bound on those designs under matched parameters. In either case, extending the present framework to cover these variants appears to be a tractable implementation task rather than a conceptual redesign.

## 5.5 Prototype Monte Carlo confirmation

[46]

**A 500-replicate prototype ran in 191 seconds with 100% convergence; the MCSE on a power estimate near 0.3 is approximately 0.020.**

- The 15-cell  $c_{br} \times N$  prototype grid was completed in 191 seconds on 9 cores.
- Convergence was 100% across all 7,500 fits.
- Per-replicate output and Morris ADEMP summary are at `analysis/data/quick-sim/`.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

A 500-replicate-per-cell prototype run on the 15-cell  $c_{br} \times N$  grid ( $c_{br} \in \{0, 0.15, 0.30, 0.45, 0.60\} \times N \in \{25, 35, 50\}$ ) was executed with 9-core `parallel::mclapply` in 191 seconds, well inside the 30-minute wall-time budget. Convergence was 100% across all 7,500 fits. Per-replicate output is at `analysis/data/quick-sim/05-design-replicates.rds` with the Morris ADEMP summary at `05-design-summary.txt`. With 500 replicates per cell the MCSE on a power estimate near 0.3 is approximately 0.020.

[47]

**The apparent non-monotonic power dependence on  $c_{br}$  from the 100-replicate prototype does not survive at 500 replicates; the signal is broadly flat.**

- At  $N = 25$ , the five-level swing is 0.024 (one MCSE), consistent with sampling noise around a constant.
- At  $N = 35$  and  $N = 50$ , the peak-to-trough swings of approximately three MCSE are suggestive but not decisive, and the peaks

1244

do not coincide on  $c_{br}$ .

1245

- The 100-replicate dip at  $c_{br} = 0.15$  has disappeared, confirming it was sampling noise.

1246

1247

1248

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

1249

1250

1251

1252

1253

1254

1255

1256

1257

1258

1259

1260

1261

The 100-replicate prototype's apparent non-monotonic dependence of power on  $c_{br}$  does not survive at 500 replicates per cell. At  $N = 25$  the rejection rate across the five  $c_{br}$  levels falls in  $[0.168, 0.192]$ , a swing of 0.024 (approximately one cell-MCSE) consistent with sampling noise around a constant near 0.18. At  $N = 35$  the swing widens to 0.066 (range 0.220 to 0.286, peak at  $c_{br} = 0.30$ ); each cell MCSE is approximately 0.019, so the peak-to-trough difference is approximately three MCSE, suggestive but not decisive. At  $N = 50$  the swing is 0.072 (range 0.266 to 0.338, peak at  $c_{br} = 0.45$ ). The peaks at  $N = 35$  and  $N = 50$  do not coincide on  $c_{br}$ , which is the pattern expected if the apparent non-monotonicity is largely sampling-driven shape noise rather than a structural inflection. The 100-replicate run's reported dip at  $c_{br} = 0.15$  has disappeared.

1262

[48]

1263

1264

1265

**The dominant axis is sample size, not  $c_{br}$ ; no prototype cell attains 80% power, and the grid does not bound the 80%-power region.**

1266

1267

1268

1269

1270

1271

- Mean power rises from 0.18 at  $N = 25$  to 0.31 at  $N = 50$ ; the  $c_{br}$  effect is a possible shallow optimum near  $[0.30, 0.45]$ .
- No  $(c_{br}, N)$  cell in the prototype reaches 80% power; a production run requires larger  $N$ , larger  $br_{\max}$ , or both.
- The prototype grid does not bound the 80%-power region, which is the central deliverable for trial planning.

1272

1273

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

1274

1275

1276

1277

1278

1279

1280

1281

1282

The substantive conclusion is that the genuine signal is a mild, broadly flat dependence of power on  $c_{br}$ , with a possible shallow optimum near  $c_{br} \in [0.30, 0.45]$  at moderate  $N$ . The dominant axis is sample size: mean power across  $c_{br}$  rises from 0.18 at  $N = 25$  to 0.25 at  $N = 35$  to 0.31 at  $N = 50$ . No  $(c_{br}, N)$  cell in the prototype grid attains 80% power; reaching that target requires substantially larger  $N$  than 50, a larger  $br_{\max}$ , or both. Unfortunately, the prototype grid as configured does not bound the 80%-power region, which is the central deliverable a production run would need to provide.

1283

## 6 Conclusions

[49]

**Three conclusions: the hybrid is a competent main-effect default, design ordering is estimand-dependent, and two-period crossovers should not be analyzed with corCAR1.**

- The hybrid absorbs variance and serial correlation, tolerates carryover up to five days, and achieves 80% power at smaller sample sizes than the parallel RCT.
- Design ordering is substantially estimand-dependent: the hybrid's advantage is larger for biomarker interactions than for main effects.
- The two-period crossover's robust Type I inflation of 0.10 to 0.13 warrants a change to standard analytic practice when corCAR1 is fitted at two within-subject timepoints.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

In conclusion, three points are to be emphasized. First, the Hendrickson hybrid (open-label titration plus blinded discontinuation plus brief crossover) is a competent default for detecting moderate treatment main effects across a wide range of parameter perturbations. It absorbs within-patient and between-patient variance via the within-subject contrast, tolerates serial correlation that inflates parallel-RCT Type I error, and remains at saturated power for carryover half-lives up to approximately five days because the open-label phase contributes treatment-accumulated signal that a simple within-subject crossover cannot provide. Second, design ordering is estimand-dependent: the hybrid's advantage over the parallel RCT is substantially larger for biomarker-treatment interactions than for the main effect, and the reader should be cautious about generalizing from one estimand to the other. Third, the two-period crossover should not be analyzed with corCAR1 at small per-subject timepoint counts; the modest but robust Type I inflation observed in S12 warrants a change to the standard analytic practice for that design.

[50]

**The hybrid N-of-1 is supported as a default candidate for trial planning in chronic neurobehavioral conditions where moderate treatment heterogeneity and short-to-moderate carryover are expected.**

- The companion clinical paper provides the baseline head-to-head comparison; this paper documents its generalization across twelve parameter axes.
- Robustness across most axes, with the identified exceptions for long carryover and two-period crossover analysis, supports the

1326            hybrid as a default candidate.

1327        *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*  
1328        *tempor incididunt ut labore et dolore magna aliqua.*

1329        The companion clinical paper provides a direct head-to-head compar-  
1330        ison at the prazosin-calibrated baseline; the present paper documents  
1331        how that comparison generalizes along twelve parameter axes for the  
1332        hybrid design family. The combination of design robustness across  
1333        most axes, with the two specific exceptions noted above, supports the  
1334        use of the hybrid N-of-1 as a default candidate in trial planning for  
1335        chronic neurobehavioral conditions where moderate treatment hetero-  
1336        geneity and short-to-moderate carryover are expected.

## 1337    7    Conflict of interest

1338    The authors declare no conflicts of interest.

## 1339    8    Funding

1340    This work received no specific funding.

## 1341    9    Data availability statement

1342        [51]

1343        **Per-replicate Monte Carlo outputs and ADEMP cell-level**  
1344        **summaries are versioned in the pmsimstats-ng repository;**  
1345        **exact paths are in the Reproducibility section.**

- 1346            • Monte Carlo outputs are under `analysis/data/`.
- 1347            • Simulation driver scripts are under `analysis/scripts/`.

1348        *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*  
1349        *tempor incididunt ut labore et dolore magna aliqua.*

1350        The data and code that support the findings of this study are openly  
1351        available in the pmsimstats-ng project repository. Per-replicate  
1352        Monte Carlo outputs and ADEMP-compliant cell-level summaries  
1353        are versioned under `analysis/data/`; simulation drivers are at  
1354        `analysis/scripts/`. Exact paths for this manuscript are listed in  
1355        the Reproducibility section below.

## 1356    10   Reproducibility

1357        [52]

1358        **The manuscript was rendered with R 4.4.x inside the zzcol-**  
1359        **lab Docker container; principal packages are nlme, parallel,**

**data.table, furrr, and purrr.**

- Image hash is in Dockerfile; package set is in `renv.lock`.
- `nlme` provides `corCAR1` residual correlation; `furrr`/`purrr` provide the tidyverse collection; `data.table` provides the original-extended collection.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

This manuscript was rendered with R 4.4.x inside the project's `zzcollab` Docker container (image hash recorded in `Dockerfile`; package set captured in `renv.lock`). Principal packages used by the simulation pipeline are `nlme` (continuous-time linear-mixed analysis with `corCAR1` residual correlation), `parallel` (9-core `mclapply` for parameter-grid sweeps), `data.table` (the original-extended simulation collection), `furrr` and `purrr` (the tidyverse simulation collection), and `rmarkdown` plus `bookdown` for the manuscript build.

[53]

**Two-level seeding ensures reproducibility: driver-level seeds govern parameter grids; replicate-level seeds derive from a master seed via the replicate index.**

- Driver level: `set.seed(42 + sweep_id)` before parameter grid construction.
- Replicate level: `set.seed(20260507)` at the loop top, with `parallel::mc.set.seed` inside `mclapply`.
- The two seeding stages operate at distinct pipeline points and do not collide.

*Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.*

Seeding operates at two levels. At the driver level, each sweep calls `set.seed(42 + sweep_id)` before constructing its parameter grid, so the parameter grid and any pre-loop draws are reproducible per sweep. At the replicate level, each driver sets `set.seed(20260507)` at the top of the replicate loop and uses `parallel::mc.set.seed` inside `mclapply`; per-replicate seeds derive from this master seed by adding the replicate index. These two seedings operate at distinct stages of the pipeline and do not collide.

[54]

**All file paths for reproducing this manuscript are fixed: per-replicate RDS, ADEMP summary, driver scripts, pre-registered plan, and manuscript source.**

- Per-replicate outputs: `analysis/data/quick-sim/05-design-replicates.rds`;



1400 ADEMP summary: 05-design-summary.txt.  
 1401 • Driver scripts: analysis/scripts/nof1-design-sensitivity/  
 1402 and analysis/scripts/quick-sim/05-design-prototype.R.  
 1403 • Pre-registered ADEMP plan: 00-ademp-pre-registration.md;  
 1404 manuscript source: analysis/report/05-nof1-design-sensitivity/.

1405 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*  
 1406 *tempor incididunt ut labore et dolore magna aliqua.*

1407 Per-replicate simulation outputs are checkpointed at analysis/data/quick-sim/05-design-replicates  
 1408 Morris ADEMP cell-level summaries at the companion 05-design-summary.txt.  
 1409 Driver scripts are under analysis/scripts/nof1-design-sensitivity/  
 1410 and analysis/scripts/quick-sim/05-design-prototype.R. The  
 1411 pre-registered ADEMP plan is at analysis/scripts/nof1-design-sensitivity/00-ademp-pre-regist  
 1412 The manuscript source, paper-specific bibliography (references.bib),  
 1413 and SIM CSL file are at analysis/report/05-nof1-design-sensitivity/.

## 1414 11 References

- 1415 1. Morris TP, White IR, Crowther MJ. Using simulation studies to evalu-  
 ate statistical methods. *Statistics in Medicine*. 2019;38(11):2074-2102.  
 doi:10.1002/sim.8086
- 1416 2. Lillie EO, Patay B, Diamant J, Issell B, Topol EJ, Schork NJ. The n-of-1  
 clinical trial: The ultimate strategy for individualizing medicine? *Person-  
 alized Medicine*. 2011;8(2):161-173. doi:10.2217/pme.11.7
- 1417 3. Kravitz RL, Duan N, Braslow J, et al. Single-patient (n-of-1) trials: A  
 pragmatic clinical decision methodology for patient-centered comparative  
 effectiveness research. *Journal of Clinical Epidemiology*. 2014;67(8):833-  
 842. doi:10.1016/j.jclinepi.2013.04.006
- 1418 4. Blackston JW, Chapple AG, McGree JM, McDonald S, Nikles J. Compar-  
 ison of Aggregated N-of-1 Trials with Parallel and Crossover Randomized  
 Controlled Trials Using Simulation Studies. *Healthcare*. 2019;7(4):137.  
 doi:10.3390/healthcare7040137
- 1419 5. Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggre-  
 gated n-of-1 trial designs for predictive biomarker validation: Statistical  
 methods and theoretical findings. *Frontiers in Digital Health*. 2020;2:13.  
 doi:10.3389/fdgth.2020.00013

- 1420 6. Carrozzo AE, Zimmermann G, Bathke AC, Neunhaeuserer D, Niebauer J, Kulnik ST. Two-arm crossover randomized controlled trial versus meta-analysis of n-of-1 studies: Comparison of statistical efficiency in determining an intervention effect. *Biometrical Journal*. 2025;67(2):e70045. doi:10.1002/bimj.70045
- 1421 7. Gärtner T, Schneider J, Arnrich B, Konigorski S. Comparison of bayesian networks, g-estimation and linear models to estimate causal treatment effects in aggregated n-of-1 trials with carry-over effects. *BMC Medical Research Methodology*. 2023;23(1):191. doi:10.1186/s12874-023-02012-5
- 1422 8. Senn S. Sample size considerations for n-of-1 trials. *Statistical Methods in Medical Research*. 2019;28(2):372-383. doi:10.1177/0962280217726801
- 1423 9. Zucker DR, Ruthazer R, Schmid CH. Individual (n-of-1) trials can be combined to give population comparative treatment effect estimates: Methodologic considerations. *Journal of Clinical Epidemiology*. 2010;63(12):1312-1323. doi:10.1016/j.jclinepi.2010.04.020
- 1424 10. Wang Y, Schork NJ. Power and Design Issues in Crossover-Based N-Of-1 Clinical Trials with Fixed Data Collection Periods. *Healthcare*. 2019;7(3):84. doi:10.3390/healthcare7030084
- 1425 11. Pinheiro JC, Bates DM. *Mixed-Effects Models in s and S-PLUS*. Springer; 2000. doi:10.1007/b98882
- 1426 12. Verbeke G, Molenberghs G. *Linear Mixed Models for Longitudinal Data*. Springer; 2000.
- 1427 13. Jiang S, Arnold SE, Betensky RA. Design and analysis of n-of-1 trials that incorporate sequential monitoring. *Statistics in Medicine*. 2025;44(15-17):e70177. doi:10.1002/sim.70177
- 1428 14. Selukar S, Prince DK, May S. A framework for sequential monitoring of individual n-of-1 trials and combining results across a series of sequentially monitored n-of-1 trials. *Clinical Trials*. 2025;22(3):257-266. doi:10.1177/17407745241304284

1429

---

1430 *Rendered on 2026-07-30 at 14:09 PDT.*

1431 *Source: ~/prj/res/36-pmsimstats-ng/pmsimstats-ng/analysis/report/05-nof1-design-sensitivity/r*