

1 Three-component decomposition of treatment response 2 in aggregated N-of-1 trials: pharmacological, 3 expectation-driven, and natural-history components

4 pmsimstats team

5 2026-06-10

6 Contents

7 Abstract

8 **Background.** Treatment response in aggregated N-of-1 clinical trials reflects
9 three causally distinct mechanisms: pharmacological action of the active com-
10 pound, the patient’s belief that the treatment will work, and natural-history drift
11 in the underlying condition. Standard analysis models lump these into a single
12 ‘treatment effect’ plus residual error, which conflates clinically distinct quantities
13 and compromises both regulatory inference and predictive-biomarker validation.

14 **Methods.** We formalise a three-component decomposition of the within-
15 participant outcome trajectory into a biological response component (BR), a
16 placebo-belief response component (PB), and a time-variant natural-history
17 component (TV), each modelled as a modified Gompertz function with
18 participant-specific random effects. We characterise the trial-design contrasts
19 (open-label, blinded discontinuation, crossover) that supply identifiability of each
20 component, derive the linear-mixed approximations to the full nonlinear model
21 that retain the pharmacological-component target estimand, and quantify the
22 inferential cost of failing to decompose through worked examples and a Monte
23 Carlo simulation study calibrated to the prazosin-PTSD reference parameters.

24 **Results.** The framework is developed and pre-registered for empirical evalua-
25 tion; a 100-replicate pilot of the analysis machinery is reported, and a production
26 simulation programme at 1,000 replicates per cell across $N \in \{35, 70, 100, 150\}$ is
27 pre-registered for execution.

28 **Conclusions.** The three-component framework is the principled basis for
29 predictive-biomarker validation in aggregated N-of-1 trials and clarifies the
30 design contrasts that are mandatory versus optional. A pragmatic analysis
31 specification consistent with the framework but tractable at trial-relevant sample
32 sizes is $Sx \sim \text{bm} + \text{t} + \text{Dbc} + \text{phase} + \text{Dbc:phase} + \text{bm:Dbc} + \text{bm:Dbc:phase}$
33 with random intercept and AR(1) residual correlation.

1 Introduction

1.1 A claim about how N-of-1 trials should be analysed

This paper argues that treatment response in N-of-1 trials comprises three causally distinct components that must be analysed jointly rather than lumped.

- The three components are BR (pharmacological), PB (placebo-belief), and TV (natural history).
- Each component is causally distinct, clinically interpretable, and identifiable from a well-designed N-of-1 trial.
- The standard single-indicator mixed-effects model conflates all three and renders biomarker validation unanswerable.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

This paper advances an opinionated methodological claim: treatment response in aggregated N-of-1 clinical trials is the sum of three causally distinct components and should routinely be analysed as such, rather than estimated as a single lumped quantity. The three components are pharmacological action (BR , the biological response), placebo-belief response (PB , the patient's belief that they are receiving treatment), and time-variant natural history (TV , the trajectory that would have occurred without intervention). Each is causally distinct, each is independently of clinical and regulatory interest, and each is identifiable from a properly designed N-of-1 trial. We shall argue that the standard practice of fitting a single mixed-effects model with one treatment indicator, prevalent across the published N-of-1 literature, conflates these three components, biases the interpretation of treatment effect, and renders predictive-biomarker validation unanswerable as a scientific question.

The novelty lies not in the individual ingredients but in the explicit joint decomposition using trial-design contrasts.

- Simultaneous estimation of all three components from a single trajectory is the key methodological advance.
- Prior N-of-1 literature has not adopted this joint-decomposition framing despite possessing the components.
- The proposed methodology addresses feasibility at trial-relevant sample sizes via a phase-augmented linear-mixed approximation.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur

78 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
79 mollit anim id est laborum.

80 *The framework we advocate is not new in its mathematical ingredients. Each*
81 *individual piece (mixed-effects models with trial-phase indicators, placebo-*
82 *arm contrasts, regression-to-mean adjustment, growth-curve parameterisa-*
83 *tion) has appeared across decades of clinical-trial methodology. What is*
84 *novel, and what we argue should become routine, is the explicit joint de-*
85 *composition: simultaneous estimation of all three components from a sin-*
86 *gle N-of-1 trajectory using trial-design contrasts that uniquely identify each*
87 *one. The published N-of-1 literature has, with rare exceptions, not adopted*
88 *this joint-decomposition framing. We shall document the gap, explain why*
89 *the gap has persisted (parsimony, identifiability concerns, sample-size con-*
90 *straints, methodological inertia), and propose a methodology that is tractable*
91 *at trial-relevant sample sizes through a phase-augmented linear-mixed-effects*
92 *approximation to the full decomposition.*

93 1.2 Why this matters: three failure modes of the lumped analysis

94 **The lumped one-component analysis has three documented failure modes**
95 **that grow with unmodelled component heterogeneity.**

- 96 • The three failure modes are placebo-attribution bias, regression-to-the-
97 mean, and biomarker-validation failure.
- 98 • Each failure mode is documented in the chronic-condition trial literature,
99 not merely hypothetical.
- 100 • The failures are present regardless of sample size.

101 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
102 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
103 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
104 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
105 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
106 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
107 mollit anim id est laborum.

108 *We begin with the failure modes of the standard approach. A one-component*
109 *analysis ('the treatment effect is what we observe minus baseline') has three*
110 *identifiable failure modes that arise when the population displays meanin-*
111 *ful heterogeneity in any of the unmodelled components. None of these is*
112 *hypothetical: each has a documented presence in the chronic-condition trial*
113 *literature.*

114 **Placebo-attribution bias is sample-size-invariant and documented across**
115 **antidepressant, anxiety, and IBS literatures.**

- 116 • Lumped treatment-effect estimates overstate pharmacological efficacy in
117 proportion to *PB* magnitude and variance.
- 118 • Kirsch et al. and Locher et al. demonstrated that antidepressant apparent
119 efficacy is largely explained by placebo response and natural fluctuation.
- 120 • Placebo responsiveness has measurable biological correlates (e.g. COMT
121 genotype), confirming it is a structured phenomenon, not noise.

122 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
123 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
124 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
125 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
126 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
127 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
128 mollit anim id est laborum.

129 *The first is the placebo-attribution failure. When the population contains*
130 *both high and low placebo responders, the lumped treatment effect overes-*
131 *timates the pharmacological component. The bias scales with the magni-*
132 *tude and variance of PB in the study population, and it is invariant under*
133 *sample-size increases: a larger one-component trial converges on the same*
134 *biased estimate. The most consequential example in modern clinical phar-*
135 *macology is antidepressant efficacy. @kirsch2008 analysed FDA-submitted*
136 *data on four widely- prescribed antidepressants and showed that, except at*
137 *the most severe baseline-severity stratum, mean drug-placebo differences fell*
138 *below clinical-significance thresholds, with the apparent 'efficacy' largely ex-*
139 *plained by placebo response and natural fluctuation; @locher2015 reproduced*
140 *this pattern in late-life depression and explicitly invoked regression to the*
141 *mean. @sugarman2014 generalised the finding to anxiety. Patient popula-*
142 *tions differ systematically in placebo responsiveness, and the difference has*
143 *measurable biological correlates: COMT genotype predicts placebo response in*
144 *irritable bowel syndrome (@wang2022comt), indicating that 'placebo' is not a*
145 *homogeneous nuisance but a structured biological phenomenon that interacts*
146 *with the patient's genome and treatment history. A trial whose population*
147 *happens to be placebo-responsive will report a larger 'drug effect' than one*
148 *whose population is not, even when the underlying pharmacology is identical.*

149 **Regression-to-the-mean contributes several points of apparent improve-**
150 **ment in early trial weeks and is absorbed without trace by the lumped**
151 **analysis.**

- 152 • Enrolment at peak severity guarantees a natural-history regression compo-
153 nent irrespective of pharmacology.
- 154 • Hrobjartsson et al. systematic reviews across 200+ trials showed much of
155 conventional placebo response was indistinguishable from natural-history
156 drift.
- 157 • The lumped analysis has no formal mechanism to recover the pharmaco-
158 logical effect once regression is absorbed.

159 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
160 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
161 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
162 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
163 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
164 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
165 mollit anim id est laborum.

166 *The second failure mode is regression-to-the-mean. When patients enrol at*
167 *moments of clinical engagement that coincide with peak symptom severity,*
168 *natural-history regression toward typical severity contributes several points*

of apparent improvement in the early weeks of any trial, without any pharmacological contribution at all. The systematic reviews of @hrobjartsson2001, @hrobjartsson2004, and @hrobjartsson2010 across more than two hundred trials and sixty conditions quantified this directly: in trials with both placebo and no-treatment arms, much of what is conventionally attributed to the placebo response was indistinguishable from regression to the mean and natural-history drift. The lumped analysis absorbs this regression into the treatment-effect estimate and the analyst has no formal mechanism to recover the true pharmacological effect.

Biomarker validation is invalidated when a biomarker correlates with PB or TV, rendering precision-medicine stratification wrong regardless of trial size.

- A biomarker correlated with placebo responsiveness or natural-history trajectory is indistinguishable from a true pharmacological predictor in a lumped analysis.
- The PATH statement defers mechanistic attribution to domain knowledge; the BR-PB-TV decomposition supplies that substrate.
- Rigorously validated pharmacological biomarkers require the decomposition as a precondition, independent of sample size.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The third failure mode is the biomarker-validation failure. The precision-medicine question that motivates many recent N-of-1 trials is whether a baseline biomarker predicts treatment response, formalised in the PATH statement [kent2020path; kent2020pathee]. A biomarker that correlates with placebo responsiveness or with a favourable natural-history trajectory will appear in a one-component analysis as a predictor of the lumped 'treatment effect', indistinguishable from a true pharmacological predictor. The biomarker-stratified prescribing decision that follows from such an analysis is wrong, and the trial's regulatory utility is degraded. The PATH framework explicitly defers the mechanism of treatment-effect heterogeneity to domain knowledge [kent2018bmj; rekkas2020scoping]; the BR-PB-TV decomposition supplies the mechanistic substrate that PATH leaves blank. We argue that predictive-biomarker validation cannot be conducted rigorously in a lumped-analysis framework regardless of how large the trial is.

1.3 The current state of the N-of-1 literature

The N-of-1 literature has matured substantially but motivates three observations that the present paper addresses.

- The design has evolved from individual-patient consultation to a framework supporting population-level inference.

- 214 • CONSORT reporting standards and quality audits document progressive
215 improvement in analytic transparency.
- 216 • Three specific gaps in the literature motivate the decomposition framework
217 proposed here.

218 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
219 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
220 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
221 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
222 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
223 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
224 mollit anim id est laborum.

225 *The aggregated N-of-1 trial design has matured rapidly over the last fifteen*
226 *years. The foundational programmatic statement is @lillie2011; the method-*
227 *ological reviews of @duan2013, the historical-practice audit of @gabler2011,*
228 *and the more recent syntheses of @schork2017nutrition, @schork2018rct,*
229 *@schork2023precision, and @porcino2022 trace the trajectory from individual-*
230 *patient consultations to the aggregated-meta-analytic framework that supports*
231 *population-level inference. Reporting standards are codified in the CONSORT*
232 *extension @vohra2015cent (the headline statement) and @shamseer2015cent*
233 *(the explanation-and-elaboration companion); the reporting-quality audits of*
234 *@li2016cent and its successors document slow but progressive improvement*
235 *in analytic transparency. Three observations about this literature motivate*
236 *the present paper.*

237 **No published N-of-1 trial employs a structural decomposition of treatment**
238 **response despite three decades of methodological literature.**

- 239 • Audit data show the field has relied on paired t-tests, binary-indicator LME,
240 and hierarchical Bayesian regressions.
- 241 • Recent methodological advances (Liao 2023, Blackston 2019, Carrozzo
242 2025) retain the lumped-response framing.
- 243 • The decomposition gap persists even as carryover and identifiability prob-
244 lems are increasingly recognised.

245 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
246 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
247 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
248 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
249 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
250 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
251 mollit anim id est laborum.

252 *The first observation is that the dominant analytic models are simple. The*
253 *audit work of @li2016cent documents that the published N-of-1 trials of the*
254 *preceding three decades used paired t-tests, simple linear mixed-effects mod-*
255 *els with a binary treatment indicator, or hierarchical Bayesian regressions,*
256 *with no published trial employing a structural decomposition of treatment*
257 *response into pharmacological, expectancy, and natural-history components.*
258 *The most recent methodological advances retain this character: @liao2023 de-*
259 *velops Bayesian distributed-lag models with autocorrelated errors as a current*

260 *carryover-aware analysis, but the model still treats the response as a single*
261 *quantity to be estimated. @blackston2019 and @carrozzo2025 compare aggre-*
262 *gated N-of-1 trials with parallel and crossover RCT alternatives by simulation*
263 *and document carryover-induced type-I-error inflation as a key failure mode,*
264 *but again at the level of the lumped treatment effect. None of these papers*
265 *decompose the response.*

266 **Placebo-component separation is widely recognised but almost entirely**
267 **unaddressed at the protocol level in N-of-1 methodology.**

- 268 • The Hendrickson hybrid blinded-discontinuation phase is among the few
269 protocol mechanisms that supply *PB*-versus-*BR* identifiability.
- 270 • Placebo run-in, open-label-placebo, and mediator-analysis approaches have
271 not been imported into N-of-1 designs.
- 272 • The gap between awareness of the problem and protocol-level solutions
273 motivates the present framework.

274 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
275 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
276 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
277 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
278 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
279 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
280 mollit anim id est laborum.

281 *The second observation is that the placebo-component separation problem is*
282 *widely known but largely unaddressed in N-of-1 specifically. The blinded dis-*
283 *continuation phase used in the Hendrickson hybrid design [@hendrickson2020]*
284 *is one of the few protocol-level mechanisms in routine N-of-1 use that supplies*
285 *identifiability of PB versus BR. Other relevant approaches, including placebo*
286 *run-in designs, open-label-placebo trials, and mediator-analysis applications,*
287 *have been developed in parallel but rarely imported into N-of-1 methodology.*

288 **Kaptchuk's within-placebo decomposition is the closest structural ana-**
289 **logue; no published paper applies the corresponding full-response decom-**
290 **position to N-of-1 trials.**

- 291 • Kaptchuk et al. (2008) empirically decomposed the placebo response into
292 observation, ritual, and relationship components.
- 293 • Open-label placebo trials confirm that clinically meaningful PB-only re-
294 sponse is possible even when patients know they receive placebo.
- 295 • The BR-PB-TV framework is structurally analogous but operates at the
296 level of the full treatment response.

297 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
298 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
299 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
300 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
301 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
302 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
303 mollit anim id est laborum.

304 *The third observation is that the closest precedent in the literature is*

Kaptchuk's three-component decomposition of the placebo response itself. In a randomised trial of acupuncture-style placebo for irritable bowel syndrome, @kaptchuk2008components decomposed the placebo response into observation, ritual, and patient-practitioner relationship components, demonstrating empirically that the placebo response is itself a sum of identifiable parts rather than a single 'expectancy' lump. The follow-up open-label placebo trials of @kaptchuk2010olp, @lembo2021olp, and @nurko2022 in adolescent functional pain establish that even disclosed-placebo intervention can produce clinically meaningful PB-only response. The framework we propose is structurally analogous to Kaptchuk's decomposition but applied at the level of the full treatment response (drug, placebo, natural history) rather than within the placebo response alone. To the best of our knowledge, no published paper proposes the corresponding decomposition for the aggregated N-of-1 setting, even though the methodological ingredients, including mixed-effects modelling, blinded-phase contrasts, and trajectory parameterisation, are all in routine use.

The decomposition gap is real, the tools to fill it are mature, and the failure modes are large enough to make it a methodological priority.

- The gap between available methodology and current practice is documented and not attributable to missing tools.
- Failure modes in chronic-condition pharmacology are consequential, not academic.
- The remainder of the paper develops, applies, and evaluates the decomposition framework.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The case we develop in the rest of this paper is therefore that the gap is real, that the methodological tools to fill it are mature, and that the failure modes of the lumped analysis are large enough in chronic-condition pharmacology to make the decomposition a methodological priority rather than an academic refinement.

1.4 The neurobiological and epidemiological reality of the components

Each component of the decomposition reflects a real causal mechanism documented in independent literatures.

- *BR* corresponds to pharmacological action; *PB* to neurobiologically real belief-driven mechanisms; *TV* to natural-history drift.
- The components are not statistical constructs but distinct causal pathways with separate literatures.

- 349 • The section reviews the evidence for each component before establishing
350 the formal model.

351 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
352 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
353 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
354 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
355 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
356 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
357 mollit anim id est laborum.

358 *The decomposition is not a statistical convenience. Each component reflects*
359 *a real causal mechanism documented in independent literatures, and we shall*
360 *briefly recall each in turn before turning to the formal model.*

361 ***BR is the pharmacological target of regulation; PB is neurobiologically***
362 ***real and phase-modulated, which is precisely the design contrast that***
363 ***makes the decomposition identifiable.***

- 364 • *PB* is mediated by endogenous opioid/dopamine release, HPA-axis regula-
365 tion, and top-down attentional shifts.
- 366 • Wager et al. neurologic pain signature provides a clean fMRI dissociation
367 between pharmacological and expectancy contributions.
- 368 • Phase-dependent belief variation (open-label vs. blinded) is the design con-
369 trast that supplies *PB* identifiability.

370 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
371 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
372 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
373 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
374 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
375 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
376 mollit anim id est laborum.

377 *The biological response component represents the pharmacological action of*
378 *the active compound: the chemical binding, downstream signalling, and phys-*
379 *iological response that a regulatory authority cares about when it grants*
380 *approval. The placebo-belief response reflects neurobiologically real mecha-*
381 *nisms documented across pain, depression, Parkinson's, and PTSD research*
382 *[@benedetti2014; @frisaldi2020; @carlino2014; @carlino2016benedetti], me-*
383 *diated by belief-driven release of endogenous opioids and dopamine, antici-*
384 *patory regulation of the hypothalamic-pituitary-adrenal axis, and top-down*
385 *attentional shifts in symptom perception. Direct neuroimaging evidence is*
386 *instructive here: @wager2013nps developed an fMRI-based 'neurologic pain*
387 *signature' that responds to remifentanyl but not to expectancy alone, providing*
388 *a clean neural dissociation between the pharmacological and expectancy con-*
389 *tributions to a clinical pain response. The placebo response is heterogeneous*
390 *across patients with biological correlates [@wang2022comt] and is modulated*
391 *by trial phase: when belief varies, as it does between open-label and blinded-*
392 *discontinuation conditions, the placebo response varies, and the modulation*
393 *is exactly the design contrast that makes the decomposition identifiable.*

394 ***TV* is the untreated counterfactual trajectory; its heterogeneity is the**
395 **central object of precision-medicine inference and must be modelled, not**
396 **absorbed into noise.**

- 397 • *TV* sign depends on disease class: positive for acute recovery, near-zero
398 for stable chronic, negative for progressive conditions.
- 399 • Selection-driven enrolment at peak severity guarantees positive *TV* in early
400 trial weeks regardless of disease class.
- 401 • The PATH statement identifies treatment-effect heterogeneity modelling
402 as mandatory for precision-medicine inference.

403 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
404 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
405 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
406 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
407 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
408 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
409 mollit anim id est laborum.

410 *The time-variant natural-history component reflects whatever trajectory the*
411 *participant would have followed without treatment. For acute conditions this*
412 *trends positive (recovery on intrinsic timescale); for chronic stable condi-*
413 *tions it is near zero on the population mean but with substantial participant-*
414 *level variance; for chronic progressive conditions it trends negative; and*
415 *for selection-driven enrolment processes it is positive in the early weeks*
416 *of the trial regardless of underlying disease class, because of regression to*
417 *the mean [@hrobjartsson2001; @hrobjartsson2010]. The PATH statement*
418 *[@kent2020path; @kent2020pathee] highlights that this heterogeneity is itself*
419 *the central object of precision-medicine inference: heterogeneity of treatment*
420 *effect must be modelled, not absorbed into noise, if the precision-medicine*
421 *question is to be answered.*

422 1.5 What this paper contributes

423 **The paper makes four contributions spanning formalisation, identifiability**
424 **mapping, tractable methodology, and quantification of the inferential cost**
425 **of lumping.**

- 426 • Sections 2–3 formalise the decomposition; section 4 maps it to design
427 contrasts; section 5 develops tractable analysis methods.
- 428 • Sections 6–7 provide worked examples and a Monte Carlo simulation study;
429 section 8 applies the framework to biomarker validation.
- 430 • A companion pedagogical document provides extended worked examples
431 for first-time readers.

432 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
433 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
434 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
435 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
436 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
437 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt

438

mollit anim id est laborum.

439

440

441

442

443

444

445

446

447

448

449

450

451

452

We make four contributions, each addressed in a section below. We begin in section 2 with the formal model and notation. Section 3 develops each component in turn, with biological motivation, parametric form, and identifiability conditions. Section 4 maps the design-contrast structure of the standard hybrid open-label-plus-blinded-discontinuation design onto the identifiability requirements of the decomposition. Section 5 develops the analysis-side methodology, including the linear-mixed-effects approximations and the phase-by-treatment indicator extension that retains BR-PB sensitivity at low parametric cost. Section 6 presents the worked examples, and section 7 presents the simulation study. Section 8 discusses applications to predictive-biomarker validation. The companion pedagogical document at docs/24-component-decomposition-pedagogy.md provides an extended worked-example treatment for readers approaching the material for the first time.

453

454

455

The first contribution is a formal specification of the three-component decomposition as a sum of Gompertz trajectories with participant-specific random effects.

456

457

458

459

460

461

- Each component follows a modified Gompertz function; the joint trajectory is their sum plus observation-level noise.
- Cross-component covariances are permitted, allowing participant-level heterogeneity to co-vary across components.
- The formal specification is the substrate on which identifiability and analysis-side arguments rest.

462

463

464

465

466

467

468

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

469

470

471

472

473

474

Our first contribution is to formalise the three-component decomposition for the N-of-1 trial setting (sections 2 and 3). Each component is modelled as a modified Gompertz function with participant-specific random effects; the joint trajectory is the sum plus observation-level noise; cross-component covariances are permitted. The formal specification is the substrate the framework builds on.

475

476

The second contribution maps identifiability of each component onto specific trial-design contrasts and characterises what is lost without them.

477

478

479

480

481

482

- The hybrid open-label-plus-blinded-discontinuation design supplies identifiability of all three components.
- Designs lacking specific phases lose identifiability of the corresponding component.
- The mapping clarifies which design features are mandatory versus optional for each inferential goal.

483 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
 484 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
 485 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
 486 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
 487 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
 488 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
 489 mollit anim id est laborum.

490 *Our second contribution is to map identifiability onto trial-design contrasts*
 491 *(section 4). The standard hybrid open-label-plus-blinded-discontinuation de-*
 492 *sign, together with phase-by-allocation contrasts, supplies identifiability of all*
 493 *three components. We characterise which designs do and do not supply this*
 494 *identifiability and what is lost in each case.*

495 **The third contribution is a phase-augmented linear-mixed-effects model**
 496 **that captures the BR-versus-PB attribution check at modest parametric**
 497 **cost.**

- 498 • The full nonlinear DGP-matched fit is computationally fragile and degree-
 499 of-freedom-expensive at $N \in [30, 150]$.
- 500 • The phase-augmented LME uses `Dbc:phase` and `bm:Dbc:phase` to test
 501 whether drug and biomarker-by-drug effects attenuate under blinding.
- 502 • This model is the recommended default for trial-relevant sample sizes.

503 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
 504 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
 505 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
 506 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
 507 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
 508 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
 509 mollit anim id est laborum.

510 *Our third contribution is to develop a tractable analysis methodology at trial-*
 511 *relevant sample sizes (section 5). The full nonlinear mixed-effects fit to mirror*
 512 *the DGP is computationally fragile at $N \in [30, 150]$ and consumes more*
 513 *degrees of freedom than trial-relevant sample sizes can spare. We propose a*
 514 *phase-augmented linear-mixed-effects model*

```
515 Sx ~ bm + t + Dbc + phase + Dbc:phase
516       + bm:Dbc + bm:Dbc:phase
517       + (1 | ptID), correlation = corCAR1(...)
```

518 *that captures the practically important part of the decomposition, specifically*
 519 *the BR-versus-PB attribution check via `Dbc : phase` and `bm : Dbc : phase`,*
 520 *at modest parametric cost.*

521 **The fourth contribution quantifies the inferential cost of lumping through**
 522 **worked examples and a calibrated Monte Carlo simulation study.**

- 523 • Two participants with identical lumped responses but different
 524 (*BR, PB, TV*) compositions are indistinguishable under the one-
 525 component analysis.

- The simulation study quantifies bias, type I error, and power across analysis strategies and parameter regimes.
- The calibration uses prazosin-PTSD reference parameters from the Hendrickson et al. dataset.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Our fourth contribution is to quantify the inferential cost of failing to decompose, through worked numerical examples (section 6) and a Monte Carlo simulation study (section 7) calibrated to the prazosin-PTSD reference parameters. The worked example shows how two participants with identical lumped responses but different underlying (BR, PB, TV) compositions appear identical to a one-component analysis and qualitatively different under the decomposition. The simulation study quantifies bias, type I error, and power across the analysis strategies and across the parameter regimes most relevant to chronic-condition trials.

The decomposition is not optional for predictive-biomarker validation: it is the formal mechanism by which the pharmacological precision-medicine question can be asked at all.

- A biomarker predicting the lumped total is not, in general, a pharmacological predictor.
- The decomposition separates the biomarker-by-BR question from biomarker-by-PB and biomarker-by-TV confounding.
- The decomposition should be the default, not an optional sensitivity analysis, in trials with biomarker validation aims.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The framework's central application is to predictive-biomarker validation (section 8). A biomarker that predicts a lumped 'treatment effect' is not, in general, a pharmacological predictor. The decomposition is the formal mechanism that lets the analyst ask the precision-medicine question (does the biomarker predict BR?) rather than the lumped surrogate question of whether the biomarker predicts the total response, including placebo and natural history. We argue that the decomposition should be the default, not an optional sensitivity analysis, in any N-of-1 trial whose primary aim is biomarker validation.

2 Notation and the formal model

This section fixes compact notation and states the formal model; a fuller worked-example treatment is in the companion pedagogical document.

- Definitions are precise enough to support the identifiability and analysis-side arguments in later sections.
- The presentation is kept compact; extended derivations are deferred to the companion document.
- The companion pedagogical document provides numerical worked examples at each formal step.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

In this section we shall fix notation and state the formal model. The three components are defined precisely enough to make the identifiability and analysis-side arguments in later sections rigorous, but we have endeavored to keep the presentation compact. A fuller treatment, with worked numerical examples at each step, is available in the companion pedagogical document at docs/24-component-decomposition-pedagogy.md.

where Y_{it} is the observed symptom score, BL_i is participant i 's baseline score, the three components BR , PB , TV are non-negative reductions in symptom severity attributable to the three causal mechanisms identified in section 1, and ε_{it} is observation-level noise. The sign convention is that an increase in any component reduces Y ; all three components are zero at $t = 0$ by construction (no exposure yet to drug, no belief activation, no elapsed natural history) and grow over the duration of the trial.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

For participant i at trial timepoint t :

$$Y_{it} = BL_i - [BR_{it} + PB_{it} + TV_{it}] + \varepsilon_{it},$$

Each component is driven by its own time-on-state variable, making the components structurally distinct even when they overlap in calendar time.

- BR_{it} accumulates with time-on-drug; PB_{it} with time-on-belief scaled by expectancy $\eta(\text{phase})$; TV_{it} with calendar time from trial entry.

- The distinct time indices are the structural source of identifiability from phase-contrast data.
- All components are zero at $t = 0$ by construction and grow over the trial duration.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Each component is a participant-specific function of its own time-on-state variable. BR_{it} depends on time-on-drug, the cumulative exposure to active medication since drug initiation. PB_{it} depends on time-on-belief, modulated by a phase-specific expectancy factor $\eta(\text{phase})$ that captures the patient's confidence that they are receiving active treatment. TV_{it} depends on time-since-trial-entry, calendar time elapsed regardless of treatment status.

rescaled so that $C(0) = 0$ and $C(t) \rightarrow m_C$ as $t \rightarrow \infty$, for $C \in \{BR, PB, TV\}$. The three Gompertz parameters (m_C, d_C, r_C) are participant-specific random effects drawn from population distributions, and the participant-level heterogeneity in these distributions is the principal source of between-person response variability that the framework is designed to characterise.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Each component follows a modified Gompertz curve:

$$C(t) = m_C \exp(-d_C \exp(-r_C t)) - m_C \exp(-d_C),$$

The residual term has AR(1) structure and the three components are permitted to co-vary through cross-component covariances.

- Within-participant residuals follow $\text{Cov}(\varepsilon_{it}, \varepsilon_{i,t+h}) = \sigma_\varepsilon^2 \rho^{|h|}$.
- Within-time and across-time cross-component covariances permit participant-level heterogeneity in one component to co-vary with the others.
- This structure accommodates biologically plausible correlations, such as high- BR patients also tending toward high- PB .

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip

Table 1: Principal psychometric and pharmacological terms used in the three-component decomposition.

Term	Meaning
BR	Biological response: pharmacological reduction in symptom score attributable to active drug.
PB	Placebo-belief response: reduction attributable to the patient’s belief that they are receiving treatment.
TV	Time-variant component: reduction (or increase) attributable to the natural-history trajectory of the disease.
$\eta(\text{phase})$	Expectancy multiplier for <i>PB</i> : 1 in open-label on-drug, ≈ 0.5 in blinded discontinuation, ≈ 0 in off-drug.
BL_i	Participant i ’s baseline (pre-treatment) symptom score.
D_{it}	Continuous-decay drug indicator on the analysis side: 1 on-drug, $\exp(-\lambda t_{sd,it})$ off-drug.
$t_{sd,it}$	Time since drug discontinuation (off-drug only).
λ	Carryover decay rate, with half-life $t_{1/2} = \ln 2/\lambda$.

ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

ε_{it} is a within-participant residual term with $AR(1)$ serial-correlation structure $\text{Cov}(\varepsilon_{it}, \varepsilon_{i,t+h}) = \sigma_\varepsilon^2 \rho^{|h|}$. The three components are permitted to be jointly correlated through within-time and across-time cross-component covariances c_{cf1t} and c_{cfct} respectively, so that participant-level heterogeneity in any one component can co-vary with heterogeneity in the others.

The principal psychometric and pharmacological terms used below are summarised in Table ??.

3 The three components

This section presents each component with its biological motivation, Gompertz parameterisation, and identifiability conditions to establish the model as causally interpretable.

- The three components are presented in turn: *BR* (biological response), *PB* (placebo-belief), *TV* (natural history).
- Each component is grounded in independent biological literature before its parametric form is stated.
- The section’s purpose is to demonstrate that the model sum reflects causal structure, not mathematical convenience.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

In this section we shall present the three components in turn, each with its

biological motivation, the modified Gompertz parameterisation, and the design contrasts that identify it. The treatment is organised so that the reader can see, by the end of this section, that the formal model is a sum of three causally interpretable trajectories rather than a mathematical convenience.

3.1 Biological response (BR)

BR is the pharmacological target of regulatory approval and biomarker stratification, independent of belief or natural history.

- *BR* arises from chemical action on the body and is zero for inert pills regardless of patient expectation.
- It is the estimand that approval decisions, dose selection, and biomarker-guided stratification require.
- Separating *BR* from *PB* and *TV* is the central precision-medicine problem the framework addresses.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

BR is the physiological reduction in symptom score attributable to the active medication's chemical action on the body. It is independent of belief, suggestion, and natural history: a participant who somehow received active drug without knowing it would still develop a measurable BR, and a participant who took an inert sugar pill while believing it was active drug would have $BR = 0$ no matter how strong their expectation. Methodologically, BR is the quantity that drug regulators and prescribers want: approval decisions, dose selection, and biomarker-guided patient stratification are all framed around BR.

PK/PD considerations produce a saturating *BR* trajectory with a slow initial rise, a peak-growth phase, and a sharp drop to zero on discontinuation.

- Plasma steady state is reached within a few half-lives; pharmacodynamic integration (e.g. receptor adaptation) adds additional lag.
- For prazosin-PTSD, Raskind et al. (2013) calibrate the onset to several weeks of sustained exposure.
- The saturating shape motivates the Gompertz over linear or simple exponential alternatives.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt

726

mollit anim id est laborum.

727

728

729

730

731

732

733

734

735

736

737

The qualitative shape of the BR trajectory follows from pharmacokinetic and pharmacodynamic considerations. Drug concentration in plasma reaches steady state within a small multiple of the elimination half-life; brain concentration follows plasma with a short delay; the downstream pharmacodynamic response (reduction in symptom-relevant neural activity, in the case of prazosin's effect on noradrenergic tone during REM sleep) typically requires several days to weeks of sustained exposure to integrate fully [raskind2013prazosin]. The result is a saturating curve: slow rise in the first week, steepest growth in the second to third week, plateau as receptor adaptation completes, and a sharp drop to zero on discontinuation as the drug clears.

738

739

740

The three Gompertz parameters (m_{BR} , r_{BR} , d_{BR}) capture ceiling, onset rate, and climb shape respectively, calibrated to prazosin-PTSD reference values.

741

742

743

744

745

- m_{BR} is the biological ceiling (maximum reduction at infinite exposure); r_{BR} the onset rate; d_{BR} the displacement shaping the early climb.
- Prazosin-PTSD calibration yields $m_{BR} \approx 11$, $r_{BR} \approx 0.42/\text{week}$, $d_{BR} \approx 5$.
- These are participant-specific random effects whose population distributions encode between-person response heterogeneity.

746

747

748

749

750

751

752

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

753

754

755

756

757

758

759

760

761

762

763

764

The modified Gompertz form $BR(t) = m_{BR} \exp(-d_{BR} \exp(-r_{BR}t)) - m_{BR} \exp(-d_{BR})$ captures this profile compactly. The maximum m_{BR} is the biological ceiling, the largest reduction the drug can produce at infinite exposure. The onset rate r_{BR} controls how rapidly the participant approaches the ceiling: high r_{BR} means a fast onset (effect plateaus in one to two weeks), low r_{BR} a slow climb (six to eight weeks). The displacement d_{BR} shapes the climb: high d_{BR} gives an early-onset curve that rises sharply at first, low d_{BR} a more gradual climb. For PTSD prazosin, calibrated values from the @hendrickson2020 reference dataset are $m_{BR} \approx 11$ (eleven points of nightmare-score reduction at saturation), $r_{BR} \approx 0.42$ per week (half-maximum at roughly four to five weeks of sustained exposure), and $d_{BR} \approx 5$ (moderately steep early climb).

765

766

767

The Gompertz is preferred over linear and simple-exponential alternatives because it captures the asymmetric slow-start saturation profile of chronic-condition pharmacodynamics.

768

769

770

771

- Linear $BR = \beta t$ is biologically implausible past the first few weeks due to unbounded growth.
- Exponential $BR = m(1 - e^{-rt})$ lacks the slow-start phase from receptor adaptation and gene-expression changes.

772 • The Gompertz interpolates between both limits and its three-parameter
773 form avoids overfitting; sensitivity analysis is reported separately in
774 07-gompertz-evaluation.

775 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
776 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
777 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
778 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
779 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
780 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
781 mollit anim id est laborum.

782 *Two alternative families that do not capture the asymmetric saturation profile*
783 *are worth noting for contrast. A linear function $BR = \beta t$ climbs without*
784 *bound and is biologically implausible past a few weeks. An exponential decay*
785 *$BR = m_{BR}(1 - \exp(-rt))$ is acceptable for fast-onset drugs but lacks the*
786 *slow-start phase that characterises most chronic-condition responses, where*
787 *receptor adaptation, feedback regulation, and gene-expression changes produce*
788 *a measurable lag. The Gompertz interpolates between exponential decay (high*
789 *d limit) and linear-in-t growth (low d limit), and its three-parameter form*
790 *is flexible enough to fit a wide range of pharmacodynamic profiles without*
791 *overfitting. Doc 25 of the project documentation provides the full historical*
792 *and biological context for the family choice and quantifies the sensitivity of*
793 *pmsimstats inference to alternative families; that work is reported separately*
794 *as 07-gompertz-evaluation.*

795 ***BR is identified by the blinded-discontinuation contrast, where drug re-***
796 ***moval collapses BR to zero while PB remains intact.***

- 797 • Within-participant on-drug versus off-drug contrasts are the primary iden-
798 tification mechanism.
- 799 • The blinded-discontinuation phase provides the cleanest contrast: *BR*
800 drops, *PB* is preserved.
- 801 • This contrast does not require the analyst to model *PB* or *TV* explicitly
802 to recover *BR*.

803 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
804 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
805 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
806 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
807 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
808 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
809 mollit anim id est laborum.

810 *The BR component is identified by within-participant on-drug versus off-*
811 *drug contrasts, formalised in section 4 below. The most direct contrast is*
812 *the blinded-discontinuation phase: when a participant is silently switched to*
813 *placebo while still believing they are on active drug, BR collapses to zero*
814 *while PB remains roughly intact, and the difference between on-drug and*
815 *off-drug response in that window is a relatively clean estimate of BR that*
816 *does not require the analyst to model PB or TV explicitly.*

3.2 Placebo-belief response (PB)

***PB* is the strict placebo response: improvement that persists when the active drug is silently replaced, provided the patient's belief is unchanged.**

- *PB* is belief-dependent, not drug-dependent; it would survive a silent switch to placebo under sustained belief.
- It is distinct from *BR* in causal origin: *BR* requires chemical action; *PB* requires only belief.
- The definition is strict and operational, not vague or colloquial.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

PB is the reduction in symptom score that arises from the patient's belief that they are receiving active treatment, independent of what they are actually receiving. That is, it is the placebo response in its strict sense: the part of the improvement that would persist if the active drug were silently replaced with an inert pill, provided the patient continued to believe they were on the active drug.

***PB* reflects neurobiologically real mechanisms triggered by belief, not pharmacology; neuroimaging and genetic evidence confirm it is a structured biological phenomenon.**

- Belief-driven endogenous opioid/dopamine release, HPA-axis modulation, and top-down attentional shifts are all documented mechanisms.
- Wager et al. neurologic pain signature dissociates pharmacological from expectancy contributions via fMRI.
- COMT genotype predicts placebo responsiveness, confirming biological structure rather than measurement artefact.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

It is worth being precise about what PB is and is not. PB is not 'fake' improvement, and it is not 'all in the patient's head' in any dismissive sense. The placebo response is mediated by well-characterised neurobiological pathways [@finniss2010; @benedetti2014; @frisaldi2020; @carlino2014; @carlino2016benedetti]. Belief-driven release of endogenous opioids and dopamine produces measurable analgesia and mood elevation; anticipatory regulation of the hypothalamic-pituitary-adrenal axis modulates cortisol and inflammatory cytokines; top-down attentional shifts change how symptoms are perceived

and reported. Each of these mechanisms is a real biological response. What makes them 'placebo' rather than 'drug' is that they are triggered by belief rather than by pharmacology. The neuroimaging evidence is instructive: @wager2013nps developed an fMRI-based 'neurologic pain signature' that responds to remifentanyl but not to expectancy alone, providing a clean neural dissociation between the pharmacological and expectancy contributions to a clinical pain response. Genetic correlates of placebo responsiveness [@wang2022comt] further establish that placebo response is a structured biological phenomenon, not a measurement artefact.

PB is identifiable from BR only when belief and treatment allocation differ, as in the blinded-discontinuation phase.

- The blinded-discontinuation phase is the canonical contrast: patients cannot tell whether they remain on active drug.
- Belief drops to a partial level ($\eta \approx 0.5$) because the patient's expectation is hedged, not absent.
- Designs without a phase that decouples belief from treatment cannot separately identify PB from BR.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

For trial design, the relevant point is that PB is separately identifiable from BR if and only if the trial includes a phase in which the patient's belief about treatment differs from the actual treatment they are receiving. The blinded-discontinuation phase is the canonical such phase: the patient is told that, at some unannounced point during the window, they may be silently switched to placebo, and they cannot reliably tell whether they are still receiving active drug. Belief in treatment correspondingly drops to a partial level (multiplier roughly 0.5) because the patient's expectation is hedged.

The $\eta = 0.5$ blinded-phase expectancy is a modelling assumption supported by literature evidence of 30–60% residual expectations, but requires sensitivity analysis.

- Published placebo literature reports residual expectations of 30 to 60 percent of the open-label level across pain, depression, and PTSD trials.
- Open-label placebo trials confirm clinically meaningful PB-only responses even under explicit disclosure.
- Sensitivity analyses varying η are mandatory before clinical conclusions depend on this single parameter.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in

907 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
908 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
909 mollit anim id est laborum.

910 *PB also follows a Gompertz curve, scaled by the expectancy factor:*

$$PB(t \mid \text{phase}) = \eta(\text{phase}) \cdot m_{PB} \exp(-d_{PB} \exp(-r_{PB}t)) + (\text{offset to start at zero}).$$

911 *Typical expectancy values are $\eta = 1.0$ in open-label on-drug, $\eta \approx 0.5$ in*
912 *blinded discontinuation, and η between 0 and 0.5 in open-label placebo (the*
913 *residual expectation a patient retains when told the previous active phase was*
914 *beneficial). The choice of $\eta = 0.5$ for the blinded phase is a modelling assump-*
915 *tion. The empirical question is how much expectation a patient carries when*
916 *uncertain about allocation, and the published placebo literature suggests it is*
917 *substantial: residual expectations of 30 to 60 percent of the open-label level are*
918 *reported across pain, depression, and PTSD trials [hrobjartsson2001; hrob-*
919 *jartsson2010]. Open-label placebo trials [kaptchuk2010olp; lembo2021olp;*
920 *nurko2022] demonstrate that the residual response can be clinically mean-*
921 *ingful even when patients are explicitly told they are receiving an inert inter-*
922 *vention. Sensitivity analyses should vary η to confirm that conclusions are*
923 *not driven by this single number.*

924 **Kaptchuk et al. (2008) provide the structural precedent; the BR-PB-TV**
925 **framework extends the same decomposition logic to the full treatment**
926 **response.**

- 927 • Kaptchuk decomposed the placebo response itself into observation, ritual,
928 and relationship components empirically.
- 929 • The present framework applies analogous decomposition logic at the level
930 of drug, placebo, and natural history.
- 931 • Both frameworks reject the 'expectancy lump' model in favour of identifi-
932 able additive parts.

933 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
934 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
935 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
936 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
937 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
938 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
939 mollit anim id est laborum.

940 *The decomposition framework as proposed here has a structural analogue in*
941 *the placebo-research literature. @kaptchuk2008components decomposed the*
942 *placebo response itself into observation, ritual, and patient-practitioner rela-*
943 *tionship components in a randomised trial of acupuncture-style placebo for*
944 *irritable bowel syndrome, demonstrating empirically that the placebo response*
945 *is not a single 'expectancy' lump but a sum of identifiable parts. The present*
946 *BR-PB-TV framework is analogous in spirit but operates at the level of the*
947 *full treatment response (drug, placebo, natural history) rather than within the*
948 *placebo response alone.*

949 ***PB* variance scales with η , not just its mean; ignoring this heteroscedastic**
950 **structure miscalibrates standard errors and downstream biomarker tests.**

- 951 • Phase-dependent variance follows $SD(PB \mid \text{phase}) = \eta(\text{phase}) \cdot \sigma_{PB}^{\text{baseline}}$.
- 952 • Open-label phases have high between-patient *PB* variance; blinded phases
- 953 have low variance.
- 954 • Ignoring heteroscedasticity inflates standard errors in open-label and de-
- 955 flates them in blinded phases, with consequences for CI coverage and
- 956 biomarker test calibration.

957 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
958 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
959 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
960 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
961 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
962 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
963 mollit anim id est laborum.

964 *A subtle property of *PB* is that its variance also scales with η , not just*
965 *its mean. When η is high, there is substantial variability across patients*
966 *in how strongly the placebo response is expressed; high-suggestibility patients*
967 *show large *PB*, low-suggestibility patients show essentially none. When η is*
968 *low (blinded), the placebo response is uniformly attenuated and the patient-*
969 *to-patient variance shrinks accordingly. The model therefore parameterises*
970 *$SD(PB \mid \text{phase}) = \eta(\text{phase}) \cdot \sigma_{PB}^{\text{baseline}}$. Ignoring this heteroscedastic structure*
971 *inflates standard errors in the open-label phase and deflates them in the*
972 *blinded phase, with downstream consequences for confidence-interval coverage*
973 *and for the calibration of any biomarker test built on top of the analysis.*

974 3.2.1 Additivity, subadditivity, and the BR-PB interaction

975 **Additivity is a modelling assumption, not biological fact; the section ex-**
976 **amines whether a $BR \times PB$ interaction term is warranted.**

- 977 • The formal model writes total response as a sum; this section asks whether
- 978 the *BR-PB* interaction is zero.
- 979 • The question has inferential consequences for biomarker validation and for
- 980 what the phase-augmented analysis captures.
- 981 • The treatment is preparatory for the generalised interaction model intro-
- 982 duced later in this section.

983 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
984 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
985 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
986 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
987 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
988 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
989 mollit anim id est laborum.

990 *In this section we shall broach a question that the formal model leaves open:*
991 *whether the three components contribute additively to symptom reduction,*

992 or whether there is a meaningful interaction between BR and PB that the
993 additive specification misses.

994 **The additive model may be wrong: published literature documents sub-**
995 **additivity between pharmacological and placebo-belief effects in some**
996 **clinical contexts.**

- 997 • Subadditivity means the combined $BR + PB$ effect is less than the sum
998 of each component measured separately.
- 999 • Evidence from Kaptchuk et al. and Benedetti et al. shows the phenomenon
1000 is context-dependent, not universal.
- 1001 • Formal test: $\partial Y / \partial (BR \cdot PB) < 0$ in a saturated mean-structure model.

1002 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1003 mod tempor incidunt ut labore et dolore magna aliqua. Ut enim ad
1004 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1005 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1006 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1007 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1008 mollit anim id est laborum.

1009 *The decomposition as written in section 2 treats the three components as*
1010 *additive contributors to symptom reduction:*

$$Y_{it} = BL_i - (BR_{it} + PB_{it} + TV_{it}) + \varepsilon_{it}.$$

1011 *The additivity is a modelling assumption, not a fact about biology, and it*
1012 *is worth stating that the assumption may be wrong in directions the frame-*
1013 *work can absorb and in directions it cannot. The published literature on*
1014 *placebo-drug interactions has accumulated evidence that the combined effect*
1015 *of receiving an active drug and believing one is receiving an active drug is,*
1016 *in some clinical contexts, less than the sum of the two components measured*
1017 *separately [kaptchuk2008components; benedetti2014]. The phenomenon is*
1018 *typically termed subadditivity: $\partial Y / \partial (BR \cdot PB) < 0$ when both components are*
1019 *present and their interaction term is included in a saturated mean-structure*
1020 *model.*

1021 **Three candidate mechanisms for subadditivity are shared neural substrate,**
1022 **scale ceiling effects, and belief-dynamics adaptation.**

- 1023 • Shared pain-modulation circuits produce saturating occupancy when both
1024 BR and PB activate them simultaneously.
- 1025 • Scale ceiling effects prevent registration of additional symptom reduction
1026 once $BR + PB$ saturates the measurement range.
- 1027 • Belief-dynamics adaptation: unexpectedly large BR can recalibrate patient
1028 expectation downward, attenuating ongoing PB .

1029 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1030 mod tempor incidunt ut labore et dolore magna aliqua. Ut enim ad
1031 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1032 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1033 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur

sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Three mechanisms have been proposed in the placebo-research literature to explain subadditivity. First, a shared neural substrate: if the descending pain-modulation circuits that mediate placebo analgesia overlap with the circuits modulated by the active drug, simultaneous activation by both mechanisms produces less marginal benefit than activation by either alone, in a saturating-occupancy sense. Second, ceiling effects on the symptom scale: a clinical scale that caps at zero (no symptoms) cannot register additional reduction once a participant has reached that cap, so any participant whose $BR + PB$ already saturates the scale will register no further reduction even if BR or PB would, in isolation, predict more. Third, belief-dynamics adaptation: the patient's expectation may adapt downward when the active drug produces a strong response that exceeds the expected effect, reducing the ongoing PB contribution; equivalently, unexpectedly large BR may attenuate PB via expectation recalibration.

where $\gamma > 0$ encodes subadditivity, $\gamma = 0$ recovers the additive form, and $\gamma < 0$ encodes superadditivity, is the natural extension. The interaction coefficient γ is identifiable from the same trial-design contrasts that identify the marginal components: open-label phases supply $BR \cdot PB$ products at high η , blinded phases supply $BR \cdot PB$ products at attenuated η , and the γ coefficient is estimated from the difference between the two regimes once the marginal components are accounted for.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Whichever mechanism dominates in a given clinical context, the subadditivity question is empirical and design-dependent. Within the BR - PB - TV framework, subadditivity manifests as a non-zero coefficient on a $BR \times PB$ interaction term in the conditional-mean specification. The additive form treats this interaction as zero by construction; a generalised form

$$Y_{it} = \text{BL}_i - [BR_{it} + PB_{it} + TV_{it} + \gamma \cdot BR_{it} \cdot PB_{it}] + \varepsilon_{it},$$

Subadditivity biases biomarker predictive accuracy estimates and is partly absorbed by the phase-augmented analysis even without explicit γ estimation.

- Under subadditivity, the biomarker-by- BR association is partly cancelled, inflating apparent predictive accuracy in the additive model.
- Phase designs attenuating PB (blinded, crossover) push the effective biomarker-treatment estimate closer to the true biomarker-by- BR value.
- The phase-augmented analysis in section 5 absorbs subadditivity as a side benefit.

1078 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1079 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1080 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1081 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1082 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1083 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1084 mollit anim id est laborum.

1085 *For trial design, two consequences follow. First, a biomarker that predicts*
1086 *BR in an additive-decomposition analysis may overstate or understate pre-*
1087 *dictive accuracy if the true mechanism is subadditive: under subadditivity,*
1088 *the biomarker-driven increase in BR is partly cancelled by the $\gamma \cdot BR \cdot PB$ in-*
1089 *teraction, and the marginal biomarker-treatment association is smaller than*
1090 *the biomarker-by-BR association. Second, trial designs that attenuate PB in*
1091 *some phases (blinded discontinuation, crossover) implicitly down-weight the*
1092 *subadditivity contribution to the response in those phases, which means that*
1093 *the effective biomarker-treatment association estimated from those phases is*
1094 *closer to the biomarker-by-BR association than is the open-label estimate.*
1095 *The phase-augmented analysis recommended in section 5 below therefore has*
1096 *the side benefit of partly absorbing subadditivity, even when γ is not estimated*
1097 *explicitly.*

1098 **A forthcoming simulation study will quantify how bias in the biomarker-**
1099 **treatment interaction scales with γ ; the pre-registration predicts linear**
1100 **scaling.**

- 1101 • Bias in $\hat{\beta}_{bm:D}$ is expected to scale linearly with γ at small $|\gamma|$.
- 1102 • Under subadditivity, bias sign is opposite to the true biomarker-by-BR
- 1103 effect; under superadditivity it is aligned.
- 1104 • Empirical quantification is deferred to the production simulation
- 1105 programme in section 7.

1106 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1107 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1108 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1109 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1110 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1111 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1112 mollit anim id est laborum.

1113 *A simulation study quantifying the magnitude of bias in the biomarker-*
1114 *treatment interaction estimate when the additive analysis is fitted to data*
1115 *generated under non-zero γ is forthcoming as part of the simulation pro-*
1116 *gramme described in section 7. The pre-registration is that bias scales linearly*
1117 *with γ at small $|\gamma|$, with the sign opposite to the true biomarker-by-BR effect*
1118 *under subadditivity and aligned with it under superadditivity.*

1119 **Non-zero $\hat{\gamma}$ does not establish true subadditive interaction; latent-class**
1120 **structure with class-correlated components produces the same marginal**
1121 **pattern.**

- 1122 • If class membership correlates with both BR and PB strength, the
- 1123 marginal covariance of $BR \cdot PB$ mimics a direct interaction term.

- This confounding is invisible from the lumped-analysis perspective but is identifiable in a class-aware analysis.
- Apparent subadditivity can arise from latent population heterogeneity with no direct *BR-PB* interaction in the class-conditional DGP.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

*A subtle but inferentially important point merits attention here: a non-zero $\hat{\gamma}$ recovered from an additive-extension analysis does not, on its own, establish that the underlying mechanism is subadditive interaction in the strict sense. The same marginal pattern can arise from latent population structure that the analyst is not modelling. If the trial population is a mixture of treatment-responders and non-responders (in the latent-class sense developed at length in [analysis/report/03-latent-class-mixture-application/](#)), and class membership correlates with both *BR* strength and *PB* strength (responder-class participants have both higher *BR* and higher *PB* on average), then the marginal expected value of *BR*·*PB* over the unobserved class label has a non-zero covariance term whose magnitude depends on the gating function $\pi(B)$. From the lumped-analysis perspective this looks exactly like a γ -weighted *BR*×*PB* interaction even though no such direct interaction was specified in the class-conditional DGP. That is, the latent-class structure is one mechanism that produces apparent subadditivity in a one-component analysis, regardless of whether the underlying biology has a true direct *BR-PB* interaction.*

Distinguishing true subadditivity from latent-class confounding requires either a class-aware analysis or the phase contrasts the *BR-PB-TV* framework provides.

- The two mechanisms produce qualitatively similar marginal effects but are inferentially distinct.
- $\hat{\gamma}$ should be reported alongside a class-aware sensitivity analysis, not as standalone evidence for mechanistic subadditivity.
- The latent-class paper's `pb_class_correlation` simulation axis is the direct test; cross-paper synthesis is the cleanest disentanglement.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The two mechanisms, namely true subadditive interaction and latent-class structure with class-correlated components, produce qualitatively similar marginal effects but are inferentially distinct. Distinguishing them requires

either a class-aware analysis that explicitly models the latent population structure (as the latent-class paper develops at length), or a phase contrast that separates the mechanisms by exploiting the differential expectancy modulation of PB that the BR - PB - TV framework provides. Under the additive-extension analysis fitted to mixture-DGP data with class-correlated components, $\hat{\gamma}$ is identified up to confounding with the class structure, and the analyst should report $\hat{\gamma}$ alongside a class-aware sensitivity analysis rather than as standalone evidence for mechanistic subadditivity. The Study 1 simulation in the latent-class paper includes a `pb_class_correlation` axis that varies $\text{Cor}(BR, PB)$ across class labels; that axis is the direct simulation-side test of the latent-class generation of apparent subadditivity, and the cross-paper synthesis between the two papers is the cleanest way to disentangle the two mechanisms.

3.3 Time-variant natural-history component (TV)

TV is the untreated counterfactual trajectory; it is participant-specific, structured, and can be positive, negative, or near-zero depending on disease class.

- TV is the symptom change that would have occurred without treatment, integrated over the trial duration.
- It varies in sign across disease classes and across patients within a class.
- It is independent of treatment status and belief, distinguishing it structurally from PB .

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

TV is the reduction in symptom score that would have happened even if the patient had received no treatment at all. It is the participant's natural-history trajectory, integrated over the duration of the trial. For some participants TV is positive (symptoms naturally improve over the trial because of life events, therapy, time-in-condition, or regression to the mean); for others TV is negative (symptoms naturally worsen because of trauma anniversaries, accumulating stress, or progressive underlying disease).

TV is a structured participant-specific signal, not noise; its distinguishing feature from PB is independence from treatment status and belief.

- Unlike ε_{it} , TV is participant-level (slow trajectory) and requires explicit modelling, not absorption into noise.
- Off-drug baselines and follow-up windows are the primary data sources for estimating TV .
- TV is modulated by neither treatment assignment nor patient belief, whereas PB is modulated by both.

1214 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1215 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1216 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1217 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1218 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1219 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1220 mollit anim id est laborum.

1221 *Like PB, TV is not noise. It is a structured participant-specific signal*
1222 *that the model can in principle estimate from data, given sufficient repeated*
1223 *measurements and a trial design that includes a long enough off-drug baseline*
1224 *or follow-up window. What distinguishes TV from PB is its origin: TV is*
1225 *independent of treatment status and of belief about treatment, whereas PB is*
1226 *modulated by both.*

1227 3.3.1 When does natural history actually exist?

1228 **The shape of TV depends on disease class, trial timescale, and selec-**
1229 **tion process; none of these can be assumed and all must be recovered**
1230 **empirically.**

- 1231 • The assumption that subjects improve over time is true for some conditions
1232 and systematically wrong for others.
- 1233 • Disease class, trial timescale, and enrolment selection jointly determine the
1234 sign and magnitude of TV.
- 1235 • Recovering the empirical shape of TV is one of the decomposition's primary
1236 objectives.

1237 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1238 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1239 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1240 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1241 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1242 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1243 mollit anim id est laborum.

1244 *A common simplifying assumption in trial analysis is that subjects will im-*
1245 *prove over time, with or without treatment. This is true for some conditions*
1246 *and problematic for others. The shape of TV depends on the underlying dis-*
1247 *ease class, the trial timescale, and the selection process by which participants*
1248 *entered the trial. None of these can be assumed without thought, and the*
1249 *empirical shape of TV is one of the things the decomposition is designed to*
1250 *recover.*

1251 **TV sign varies systematically by disease class: positive for acute recovery,**
1252 **near-zero for stable chronic, negative for progressive, and structured for**
1253 **cyclical conditions.**

- 1254 • Acute conditions (post-operative pain, acute back pain) resolve intrinsically;
1255 substantial TV-improvement appears in the placebo arm.
- 1256 • Progressive conditions (Alzheimer's, ALS) require explicit negative-
1257 trajectory modelling; apparent treatment effect is slowing of progression.

- 1258 • Cyclical/triggered conditions (PTSD with trauma anniversaries, relapsing-
1259 remitting MS) require covariates beyond the Gompertz form.

1260 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1261 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1262 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1263 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1264 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1265 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1266 mollit anim id est laborum.

1267 *Acute conditions tend toward TV-positive: a sprained ankle, a post-operative*
1268 *pain syndrome, an episode of acute back pain all resolve on a timescale of days*
1269 *to weeks even without intervention, and a trial of an analgesic in such a setting*
1270 *will see substantial TV-improvement in the placebo arm. Chronic stable con-*
1271 *ditions (treated hypertension, controlled chronic asthma, long-standing PTSD*
1272 *without active triggers) tend toward TV-zero on the population mean but with*
1273 *substantial individual variance. Chronic progressive conditions (Alzheimer's*
1274 *disease, ALS, late-stage congestive heart failure, many cancers) tend toward*
1275 *TV-negative: a trial of a neuroprotective agent in early Alzheimer's must ex-*
1276 *PLICITLY model the negative natural-history trajectory, since the apparent treat-*
1277 *ment effect is the slowing of progression, not absolute improvement. Cyclical*
1278 *and triggered conditions (PTSD with trauma anniversaries, seasonal-pattern*
1279 *major depression, relapsing-remitting multiple sclerosis) show structured TV*
1280 *that is neither monotone nor stationary, and may require explicit covariates*
1281 *beyond the Gompertz form.*

1282 **Selection-induced TV is real and rarely zero even for stable chronic con-**
1283 **ditions, because enrolment coincides with peak severity triggering help-**
1284 **seeking.**

- 1285 • Regression to the mean alone contributes several points of apparent im-
1286 provement in early trial weeks.
- 1287 • Hrobjartsson et al. systematic reviews across 200+ trials show much of con-
1288 ventional placebo response is indistinguishable from natural-history drift.
- 1289 • Kirsch (2008) and Locher (2015) demonstrate the clinical consequences in
1290 antidepressant and late-life depression trials.

1291 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1292 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1293 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1294 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1295 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1296 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1297 mollit anim id est laborum.

1298 *Selection-induced TV is real and rarely zero, even when the underlying con-*
1299 *dition is stable. Patients enrol in trials at moments of clinical engagement,*
1300 *which often coincide with recent worsening that triggered help-seeking. Regres-*
1301 *sion to the mean alone produces several points of apparent TV-improvement*
1302 *in the early weeks of a trial, even in a population whose underlying condition*
1303 *is stable. The systematic-review evidence of @hrobjartsson2001, @hrobjarts-*

son2004, and @hrobjartsson2010 across more than two hundred trials and sixty conditions quantifies this directly: in trials with both placebo and no-treatment arms, much of what is conventionally attributed to the placebo response was indistinguishable from regression to the mean and natural-history drift. The depression literature provides the most consequential modern example: @kirsch2008, analysing FDA-submitted antidepressant trials, showed that mean drug-placebo differences fell below clinical-significance thresholds except at the most severe baseline-severity stratum, with the apparent 'efficacy' largely explained by placebo response and natural fluctuation; @locher2015 explicitly invoked regression to the mean to explain the severity-by-treatment interaction in late-life depression.

Trial participation itself induces TV -improvement through Hawthorne, structured-visit, and diary-completion effects, independent of pharmacology or placebo.

- Hawthorne, structured-visit, and diary-completion effects each contribute symptom improvements attributable to participation, not treatment.
- These contributions are absorbed into TV in the decomposition, not into BR or PB .
- Ignoring these participation-induced effects inflates the apparent pharmacological and placebo components.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Trial participation itself also induces TV : the Hawthorne effect, the structured-visit effect, and the diary-completion effect appear in the data as TV -improvement that is independent of any pharmacological or placebo response.

3.3.2 Can we assume no TV effect?

Assuming $TV = 0$ is rarely safe in chronic psychiatric or pain trials; the conservative default is to estimate TV from the data.

- Zero- TV assumption is valid only when all four conditions hold: stable disease, short trial, non-severity-peaked recruitment, no ancillary care.
- For most chronic conditions at least one condition fails; the assumption should be tested, not asserted.
- Statistical efficiency lost by estimating TV is smaller than the bias introduced by assuming it zero.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in

1348 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1349 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1350 mollit anim id est laborum.

1351 *But how safe is it to assume TV is zero? The short answer is: rarely safe,*
1352 *and never safe without empirical support. The longer answer is that TV*
1353 *may be assumed zero only if the disease is in a stable chronic phase, the trial*
1354 *is short relative to the disease timescale, the recruitment process does not*
1355 *favour participants at peak severity, and the trial protocol does not introduce*
1356 *ancillary care that itself improves symptoms. For most chronic trials in*
1357 *psychiatric or pain conditions, at least one of these conditions fails. The*
1358 *conservative default is to estimate TV from the data, even at the cost of*
1359 *some statistical efficiency, rather than assume it is zero.*

1360 **The decomposition makes the zero-TV assumption testable and shows**
1361 **that lumping TV into residuals produces autocorrelated, heteroscedastic**
1362 **residuals that bias BR and PB standard errors.**

- 1363 • $\hat{m}_{TV}$ with a variance estimate from the model provides an empirical test
1364 of the zero-TV assumption.
- 1365 • Participant-level TV is structured and correlated with the on/off-drug
1366 schedule; it is not independent noise.
- 1367 • Lumping TV into ε produces autocorrelated within-participant residuals
1368 and heteroscedastic cross-participant residuals.

1369 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1370 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1371 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1372 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1373 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1374 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1375 mollit anim id est laborum.

1376 *The decomposition makes this assumption testable. A model that includes a*
1377 *TV component will return $\hat{m}_{TV}$ with a variance estimate; if both are small*
1378 *relative to $\hat{m}_{BR}$ and $\hat{m}_{PB}$, the analyst has empirical support for the simpli-*
1379 *fying assumption. If $\hat{m}_{TV}$ is meaningfully large, the assumption was wrong,*
1380 *and the trial needs the TV component to produce unbiased inference about*
1381 *the others. Setting TV to zero and lumping it into ε is worse than estimating*
1382 *it: the participant-level TV is structured (each participant has their own slow*
1383 *trajectory) and not independent of the on-drug-versus-off-drug schedule, so*
1384 *lumping it into noise produces residuals that are autocorrelated within partic-*
1385 *ipants and heteroscedastic across them. The resulting standard errors on the*
1386 *BR and PB estimates are wrong in both directions: too small for participants*
1387 *whose TV is small, too large for participants whose TV is large. Estimating*
1388 *TV explicitly fixes both problems simultaneously.*

1389 3.3.3 Why a Gompertz, not a linear time covariate?

1390 **A linear week covariate is insufficient for TV because it enforces linearity,**
1391 **ignores participant heterogeneity, and cannot distinguish TV from BR**
1392 **without the explicit decomposition.**

- A week covariate constrains each additional week to contribute a fixed natural-history change, misrepresenting curved or saturating disease trajectories.
- A single coefficient estimates a population-mean trajectory, ignoring participant-level heterogeneity in natural history.
- Without the decomposition and design contrasts, week absorbs both BR and TV , making them inseparable.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

A natural objection is that a model with week as a covariate already controls for time. But how adequate is such a control? Three reasons suggest it is not enough. First, a week covariate enforces a linear trajectory, constraining each additional week to contribute a fixed amount of natural-history change; real disease trajectories curve, plateau, or accelerate, and the Gompertz form captures the S-shaped or saturating patterns linear time cannot. Second, a single week coefficient estimates a population-mean trajectory that ignores the participant-level heterogeneity in natural history. Third, BR and TV both evolve smoothly over the same trial timescale, and a week covariate cannot tell them apart from observational data alone; the design contrasts discussed in section 4 are what supply identifiability, and without the explicit decomposition the week coefficient absorbs both.

4 Identifiability and trial-design contrasts

Identifiability is supplied by trial design; each phase of the hybrid design contributes a different combination of components to the observed trajectory.

- The decomposition is mathematically specified but requires trial-design contrasts to be statistically identified.
- This section characterises the hybrid open-label-plus-blinded-discontinuation design as the primary identifiability vehicle.
- Each of the three phases supplies a distinct mixture of BR , PB , and TV .

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The mathematical decomposition above is meaningless unless the data contain enough information to pin down each component separately. Trial design

1437 *is what supplies that information. In this section we shall characterise the*
1438 *hybrid open-label-plus-blinded-discontinuation design, used as the worked ex-*
1439 *ample throughout this paper, and show how each of its three phases contributes*
1440 *a different combination of components to the observed trajectory.*

Phase	Drug?	Belief?	Components present
Open-label on drug	Yes	Full	$BR + PB(\eta = 1) + TV$
Blinded discontinuation	Hidden	Hedged	$BR \cdot \mathbf{1}_{\text{on drug}} + PB(\eta = 0.5) + TV$
Open-label crossover	Either	Knows	$BR \cdot \mathbf{1}_{\text{on drug}} + PB(\eta = 1) + TV$

1441 **Three phase-by-allocation contrasts isolate PB , BR , and TV respectively;**
1442 **the LME model combines them but the inference is driven by design.**

- 1443 • Open-label versus blinded contrast yields approximately $0.5 \cdot PB$, isolating
1444 the belief component.
- 1445 • Blinded on-drug versus blinded placebo yields BR alone (since PB and
1446 TV match across arms within the blinded phase).
- 1447 • Off-drug phases yield TV alone, with no BR and PB at the off-drug
1448 expectation level.

1449 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1450 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1451 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1452 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1453 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1454 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1455 mollit anim id est laborum.

1456 *Three phase-by-allocation contrasts then identify the three components. The*
1457 *open-label versus blinded contrast (within drug) yields $PB(\eta = 1) - PB(\eta =$*
1458 *$0.5) \approx 0.5 \cdot PB$, which isolates the belief component. The blinded on-drug*
1459 *versus blinded placebo contrast yields BR alone, since PB and TV match*
1460 *across allocation arms within the blinded phase. All phases off drug yield TV*
1461 *alone, with no BR contribution and PB at the off-drug expectation level.*

1462 **The LME model formally combines all three contrasts, but inferential**
1463 **precision is determined by the design, not model specification.**

- 1464 • Each contrast targets a specific component; the model integrates all three
1465 into a coherent parameter set.
- 1466 • Model specification cannot compensate for missing design contrasts; the
1467 identifying information must exist in the data.
- 1468 • The analyst should evaluate any proposed trial design by verifying each
1469 component has at least one identifying contrast.

1470 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1471 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad

1472 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1473 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1474 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1475 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1476 mollit anim id est laborum.

1477 *Each contrast targets a different component, and the analyst combines them*
1478 *through the linear-mixed-effects model fitted to the full trajectory. The model*
1479 *is the formal mechanism, but the inference is driven by the design contrasts.*

1480 **Designs lacking specific phases lose identifiability of the corresponding**
1481 **component; the decomposition machinery is only as useful as the design**
1482 **allows.**

- 1483 • A pure open-label design cannot separate BR from PB because both are
1484 present at full strength throughout.
- 1485 • A pure blinded design cannot estimate $\eta = 1$ PB because no phase achieves
1486 full-strength belief.
- 1487 • Trial designers should verify each component has at least one phase in
1488 which it is uniquely present before committing to the design.

1489 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1490 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1491 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1492 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1493 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1494 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1495 mollit anim id est laborum.

1496 *A trial that lacks any of these phases loses the corresponding identifiability. A*
1497 *pure open-label design (no blinded discontinuation) cannot separate BR from*
1498 *PB , because both are present at full strength throughout. A pure blinded*
1499 *design (no open-label phase) cannot estimate $\eta = 1$ PB , because no phase*
1500 *has belief at full strength. The decomposition machinery is useful only to*
1501 *the extent that the design supplies identifiability, and trial designers should*
1502 *evaluate any proposed design by checking that each component has at least*
1503 *one phase or contrast in which it is uniquely present.*

1504 5 Analysis-side methodology

1505 **Four constraints make a component-matched analysis impractical at trial-**
1506 **relevant sample sizes; a phase-augmented LME captures the important**
1507 **part of the gap at minimal cost.**

- 1508 • The question of whether to mirror the DGP in the analysis model is rea-
1509 sonable; the section addresses it directly.
- 1510 • The four constraints are identifiability, degrees of freedom, convergence
1511 fragility, and marginal estimand robustness.
- 1512 • A phase-by-treatment indicator is the recommended modest extension that
1513 captures BR -versus- PB sensitivity without the full parametric cost.

1514 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1515 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1516 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1517 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1518 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1519 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1520 mollit anim id est laborum.

1521 *In this section we shall take up the practical question of what analysis model*
1522 *to use. A reasonable reaction to the formal model in section 2 is the fol-*
1523 *lowing: if we are simulating data with three components plus noise, why is*
1524 *the analysis model written as $Sx \sim \mathbf{b}m + \mathbf{t} + D\mathbf{b}c + \mathbf{b}m:D\mathbf{b}c$ with a single*
1525 *random intercept and AR(1) residual correlation? Should we not include*
1526 *three parametric Gompertz components in the analysis, mirroring the DGP,*
1527 *so that the analyst sees the same structure that generated the data? The*
1528 *question is reasonable. Four constraints make a component-matched analysis*
1529 *impractical for most aggregated N-of-1 trial sizes, and a modest extension (a*
1530 *phase-by-treatment indicator) captures the practically important part of the*
1531 *gap at minimal cost.*

1532 5.1 Four reasons not to match the DGP

1533 **The first constraint is identifiability: forcing a three-component model on**
1534 **a design with insufficient contrasts produces unidentified parameters, not**
1535 **component estimates.**

- 1536 • A pure open-label trial cannot distinguish *BR* from *PB*; both rise from
1537 zero together on treatment initiation.
- 1538 • Fitting three components to a one-contrast design produces implausibly
1539 tight standard errors that misrepresent inferential precision.
- 1540 • The analyst must audit whether the design supplies the necessary contrasts
1541 before specifying a component-matched model.

1542 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1543 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1544 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1545 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1546 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1547 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1548 mollit anim id est laborum.

1549 *The first reason is identifiability. The trial design must supply the contrasts*
1550 *that identify each component, and many designs supply only part of it. The*
1551 *decomposition is identified by phase contrasts (open-label versus blinded for*
1552 **PB*, on-drug versus blinded-placebo for *BR*, off-drug timepoints for *TV*). A*
1553 *pure open-label trial has none of these contrasts; in such a trial *BR* and *PB**
1554 *both rise from zero on treatment initiation and are statistically indistinguish-*
1555 *able from each other, regardless of how the analyst writes the model. Forcing*
1556 *an analysis to fit three components when the design supplies contrasts for*
1557 *only one or two of them does not recover the missing components; it produces*
1558 *unidentified or weakly identified parameter estimates with implausibly tight*

standard errors that misrepresent the analyst's actual inferential precision.

The second constraint is degrees of freedom: a component-matched analysis adds approximately a dozen non-linear parameters, inflating moderation estimand variance at $N \in [30, 150]$.

- Gompertz *BR* (3), *PB* (3 + expectancy multiplier), and *TV* (3) each require participant-specific random effects on the maxima.
- At trial-relevant N , the variance of $\hat{\beta}_{bm:D}$ is dominated by participant and timepoint counts.
- Additional parametric structure typically inflates moderation estimand variance more than it reduces bias.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

*The second reason is degrees of freedom. Each component term costs degrees of freedom that come out of the moderation estimand. The biomarker-moderation question (the core scientific aim of the predictive-biomarker validation programme) is asked through the `bm:Dbc` term. A component-matched analysis adds a Gompertz *BR* (three parameters), a Gompertz *PB* (three parameters plus a phase-specific expectancy multiplier), and a Gompertz *TV* (three parameters), each with participant-specific random effects on the maxima. That is on the order of a dozen new parameters, half of them on a non-linear scale. At trial sample sizes of $N \in [30, 150]$, the variance of $\hat{\beta}_{bm:D}$ is dominated by participant and timepoint counts, and adding parametric structure to absorb residual confounding usually inflates the moderation estimand's variance by more than it reduces bias.*

The third constraint is convergence fragility: component-matched nonlinear mixed models converge below 90% at trial-relevant N , contaminating aggregated analyses.

- Simultaneous Gompertz *BR*, *PB*, and *TV* fitting requires `nlme::nlme`, Bayesian hierarchical, or structural-equation approximations.
- Each is more sensitive to starting values, prone to local optima, and likely to fail on individual trial datasets than the LME analogue.
- Non-converging cells at below-90% convergence rates contaminate power and bias estimates in simulation studies.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The third reason is convergence fragility. Component-matched analyses are non-linear in the parameters and their convergence is problematic in practice. Fitting Gompertz BR, PB, and TV simultaneously requires a non-linear mixed model (`nlme::nlme` with self-starting Gompertz forms), a Bayesian hierarchical fit, or a structural-equation approximation. Each is more sensitive to starting values, more prone to local optima, and more likely to fail to converge on a single trial dataset than the linear-mixed analogue. At trial-relevant N , the convergence rate of the component-matched fit is typically below 90%, and the non-converging cells contaminate any aggregated analysis.

The fourth constraint is estimand robustness: the `bm:Dbc` marginal interaction is the right estimand for the biomarker-validation question regardless of the BR-PB split.

- Regulators want to know whether the biomarker predicts treatment response under trial conditions, not a mechanistically decomposed prediction.
- The marginal `bm:Dbc` interaction captures this question whether the moderation is biological or expectation-mediated.
- Component decomposition adds mechanistic precision but is not required for the core biomarker-validation decision.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The fourth reason is that the moderation estimand is sufficiently robust to the BR-versus-PB split for the most common scientific question. For predictive-biomarker validation, the regulator wants to know whether the biomarker predicts treatment response under the trial conditions, not a clean pharmacological-versus-placebo decomposition of that prediction. The simple `bm:Dbc` interaction is the right marginal estimand for the validation question whether the moderation is biological or expectation-mediated.

5.2 A small extension that usually pays off: phase indicators

The phase-augmented LME uses `Dbc:phase` and `bm:Dbc:phase` to test whether drug and biomarker-by-drug effects attenuate under blinding.

- `phase` absorbs systematic differences in *PB* expectancy across trial phases.
- `Dbc:phase` detects whether the drug effect shrinks under blinding, a signature of *PB*-mediation.
- `bm:Dbc:phase` detects whether the biomarker-by-treatment interaction shrinks under blinding, a signature that it is partly biomarker-by-*PB*.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad

minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The smallest useful step toward component-aware analysis, without paying the costs above, is to add a phase indicator and its interaction with the drug indicator. Schematically,

“ $r_{Sx} = \beta_m + t + Dbc + phase + Dbc:phase + \beta_m:Dbc + \beta_m:Dbc:phase + (1 | ptID)$, correlation = corCAR1(...)”

Here `phase` is a categorical covariate (open-label, blinded, post-crossover) that absorbs the systematic differences in PB expectancy across phases. `Dbc:phase` tests whether the apparent drug effect shrinks under blinding, a signature that the open-label drug effect was partly PB-mediated. `$\beta_m:Dbc:phase$` tests whether the moderation itself shrinks under blinding, a signature that the biomarker-by-treatment interaction was partly biomarker-by-PB rather than biomarker-by-BR.

The phase-augmented extension adds at most four parameters yet captures the practically relevant belief-attenuation diagnostic without requiring the Gompertz parametric form.

- Three to four parameters is well within budget at any realistic trial sample size.
- The extension does not separately estimate *BR* and *PB* components but identifies the change in apparent drug effect when belief is hedged.
- For most analyses this is the practically relevant question and is answerable without committing to the Gompertz form.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

*This extension adds three to four parameters at most, well within budget at any realistic N . It does not separately estimate *BR* and *PB* as components, but it does identify the change in apparent drug effect that occurs when belief is hedged. For most analyses, this is the practically relevant question, and it can be answered without committing to the Gompertz parametric form.*

5.3 When to fit the full decomposition

The full decomposition merits the effort only when N is in the hundreds, the design is fully phase-balanced, and the scientific question requires the BR-PB mechanistic split.

- Feasibility conditions: several-hundred N , dense timepoints in every phase,

- 1691 mechanistic BR-PB split required, and informative Bayesian priors avail-
1692 able.
- 1693 • In that regime, `lcmm`, `nlme::nlme` with self-starting Gompertz forms, or
1694 `brms` can return participant-specific component estimates.
 - 1695 • Outside that regime the phase-augmented LME is the better tool.

1696 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1697 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1698 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1699 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1700 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1701 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1702 mollit anim id est laborum.

1703 *A component-matched analysis merits the effort when sample size is in the*
1704 *several hundreds (a phase-3 trial fielded as aggregated N-of-1, or a meta-*
1705 *analysis pooling multiple smaller trials); the trial design provides timepoints*
1706 *in every phase combination, with enough density per cell to constrain the*
1707 *Gompertz curves; the scientific question requires the BR-PB split (a phar-*
1708 *macology paper claiming a specific mechanism, where ruling out a placebo*
1709 *contribution is the central point); and a Bayesian hierarchical fit with infor-*
1710 *mative priors on the Gompertz parameters is feasible. In that regime, `lcmm`,*
1711 *`nlme::nlme` with self-starting Gompertz forms, or a `brms` hierarchical formu-*
1712 *lation can return participant-specific BR, PB, and TV component estimates*
1713 *that mirror the DGP. Outside that regime, the linear-mixed analysis with*
1714 *phase indicators is the better tool.*

1715 **The DGP-analysis asymmetry is deliberate: simulation uses the rich DGP**
1716 **to characterise how a simpler analysis behaves; the analysis need not**
1717 **mirror the DGP to produce useful inference.**

- 1718 • Simulation controls the DGP and can specify any structure; inference must
1719 pay the identifiability cost of every parameter.
- 1720 • The DGP is a characterisation tool; the analysis is a finite-data inference
1721 tool; they have different optimality criteria.
- 1722 • DGP-matching is not required and is often counterproductive at trial-
1723 relevant sample sizes.

1724 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1725 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1726 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1727 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1728 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1729 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1730 mollit anim id est laborum.

1731 *The asymmetry between DGP and analysis is not a flaw of the framework; it*
1732 *is a deliberate consequence of the asymmetry between simulation (where the*
1733 *analyst controls the DGP and can specify any structure, however elaborate)*
1734 *and inference (where the analyst must fit a model to a finite dataset and pay*
1735 *the identifiability cost of every additional parameter). Simulation studies use*
1736 *the rich DGP to characterise how a simpler analysis behaves under realistic*

1737 data generation; they do not require the analysis to match the DGP in order
1738 to produce useful inference.

1739 6 Worked example: information loss without decom- 1740 position

1741 The worked example uses two participants with identical lumped re-
1742 sponses but opposite BR-PB profiles to show that the one-component
1743 analysis cannot distinguish them.

- 1744 • Participant A has high PB ($m_{PB} = 6$) and moderate BR ($m_{BR} = 4$);
1745 Participant B has low PB ($m_{PB} = 1$) and high BR ($m_{BR} = 8$).
- 1746 • Both participants reach approximately 10–11 points of total improvement
1747 by week 8, making them indistinguishable to a one-component analysis.
- 1748 • The blinded-discontinuation phase creates the diagnostic contrast that ex-
1749 poses the underlying difference.

1750 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1751 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1752 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1753 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1754 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1755 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1756 mollit anim id est laborum.

1757 To illustrate the inferential cost of failing to decompose, we shall work through
1758 a simple numerical example. Consider two participants, A and B, enrolled in
1759 a sixteen-week N-of-1 trial of prazosin for PTSD. Both are otherwise similar:
1760 same age, same baseline severity, same trial design. They differ in two latent
1761 traits that the trial does not measure directly. Participant A is a high placebo
1762 responder and a moderate pharmacological responder, with true component
1763 values $m_{BR}^A = 4$, $m_{PB}^A = 6$, $m_{TV}^A = 1$. Participant B is a low placebo
1764 responder and a strong pharmacological responder, with $m_{BR}^B = 8$, $m_{PB}^B = 1$,
1765 $m_{TV}^B = 1$.

1766 Open-label totals are clinically indistinguishable; the one-component anal-
1767 ysis estimates a population drug effect of approximately 10 points with
1768 no differentiation between A and B.

- 1769 • Both participants show approximately 10–11 points of improvement and
1770 present as clinical successes.
- 1771 • A t-test or simple LME on open-label change cannot detect the hetero-
1772 geneity in BR versus PB composition.
- 1773 • The population drug-effect estimate conflates both pharmacological and
1774 placebo contributions.

1775 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1776 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1777 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1778 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1779 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur

1780 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1781 mollit anim id est laborum.

1782 *In the open-label phase (weeks 1 to 8), both BR and PB run at full strength*
1783 *and TV accumulates. By week 8 both participants have reached close to their*
1784 *saturation values:*

1785 *|| BR | PB($\eta = 1$) | TV | Total improvement | |---|---|---| | Participant*
1786 *A | 4.0 | 6.0 | 1.0 | **11.0** | | Participant B | 8.0 | 1.0 | 1.0 | **10.0** |*

1787 *A clinician examining only the totals would conclude that the two participants*
1788 *are responding similarly: both improved by about ten points, both look like*
1789 *clinical successes. A simple t-test on the open-label change would estimate a*
1790 *population drug effect of about ten points and would not distinguish the two*
1791 *participants.*

1792 **The blinded-discontinuation phase reveals the BR-PB heterogeneity: Par-**
1793 **ticipant A retains 4 points while B collapses to 1.5, exposing the diag-**
1794 **nostic contrast.**

- 1795 • Participant A's high *PB* at $\eta = 0.5$ (3 points) sustains improvement even
1796 without active drug.
- 1797 • Participant B's low *PB* (0.5 points) and absent *BR* produce near-complete
1798 return to baseline.
- 1799 • The 4.0 versus 1.5 contrast is exactly what the decomposition uses to
1800 separate *BR* from *PB*.

1801 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1802 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1803 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1804 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1805 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1806 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1807 mollit anim id est laborum.

1808 *In the blinded discontinuation phase (weeks 9 to 12), both participants are*
1809 *silently switched to placebo at the start of week 9. The expectancy multiplier*
1810 *drops from 1.0 to 0.5. By the end of week 12, the components are:*

1811 *|| BR | PB($\eta = 0.5$) | TV | Total improvement | |---|---|---| | Participant*
1812 *A | 0.0 | 3.0 | 1.0 | **4.0** | | Participant B | 0.0 | 0.5 | 1.0 | **1.5** |*

1813 *The two participants now look quite different. Participant A still shows a*
1814 *sizeable improvement (four points) under blinded placebo; participant B has*
1815 *nearly returned to baseline (one and a half points). This is the diagnostic*
1816 *contrast that the decomposition exploits.*

1817 **The one-component model estimates drug at 7.75 points with no princi-**
1818 **pled decomposition, obscuring that A's true *BR* is 4 and B's is 8.**

- 1819 • The 7.75 point estimate is a conflated average of true *BR* values (4.0 and
1820 8.0) plus the placebo wash-out.
- 1821 • The model cannot determine whether A and B differ, nor estimate the
1822 pharmacological drug effect separately.

- 1823 • Placebo and natural-history effects are not identifiable from this contrast
1824 alone.

1825 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1826 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1827 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1828 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1829 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1830 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1831 mollit anim id est laborum.

1832 *A one-component fixed-effects regression of total symptom change on phase-*
1833 *by-allocation indicators sees:*

1834 - Open-label means: 11.0 (A), 10.0 (B). - Blinded-placebo means: 4.0 (A),
1835 1.5 (B). - Estimated drug effect: $11.0 - 4.0 = 7.0$ (A), $10.0 - 1.5 = 8.5$ (B).
1836 Population mean: 7.75. - Estimated placebo effect: not identifiable from this
1837 contrast. - Estimated natural history: not identifiable from this contrast.

1838 *The one-component model has produced a single number conflated with the*
1839 *placebo response that washes out under blinding. It cannot say whether the*
1840 *drug effect is the same for participants A and B; it estimates 'drug' at 7.75*
1841 *points on average, which is an average of the true 4.0 (A) and 8.0 (B) mixed*
1842 *with the placebo wash-out, with no principled way to decompose further.*

1843 **The full decomposition recovers all six component estimates by exploiting**
1844 **phase differential expectancy and the blinded-placebo contrast.**

- 1845 • The LME returns $\hat{m}_{BR}^A = 4.0$, $\hat{m}_{BR}^B = 8.0$, $\hat{m}_{PB}^A = 6.0$, $\hat{m}_{PB}^B = 1.0$ at
1846 the population point estimate.
1847 • Participant-level estimates have meaningful uncertainty at trial-relevant N ;
1848 the simulation study characterises this.
1849 • The phase differential ($\eta = 1$ vs. $\eta = 0.5$) is the identification mechanism
1850 for the PB parameter.

1851 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1852 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1853 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1854 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1855 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1856 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1857 mollit anim id est laborum.

1858 *A linear-mixed-effects analysis with the full BR-PB-TV decomposition, fitted*
1859 *to the same data, can recover the individual components by exploiting the*
1860 *differential expectancy between phases and the on-drug-versus-blinded-placebo*
1861 *contrast. After fitting, the model returns $\hat{m}_{BR}^A = 4.0$, $\hat{m}_{BR}^B = 8.0$, $\hat{m}_{PB}^A = 6.0$,*
1862 *$\hat{m}_{PB}^B = 1.0$, $\hat{m}_{TV}^A = \hat{m}_{TV}^B = 1.0$ (at the population point estimate; participant-*
1863 *level estimates have meaningful uncertainty at N in the trial-relevant range,*
1864 *which the simulation study below characterises).*

1865 **The one-component model produces one biased average; the three-**
1866 **component model produces six clinically actionable quantities, making**

biomarker validation possible.

- The lumped estimate of 7.75 points is a biased average of $BR_A = 4$ and $BR_B = 8$ confounded with placebo wash-out.
- The decomposition supplies individual BR , PB , and TV estimates enabling per-participant clinical decisions.
- Predictive-biomarker validation requires the decomposition because biomarker-by- BR is the target estimand.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The information loss from lumping the components is not abstract. We may compare the inferences directly:

*/ Question / One-component answer / Three-component answer / |—|—|—|
/ Does the drug work for A? / 7.75 points (lumped). / $BR_A = 4.0$, distinct
from her larger placebo response. / Does the drug work for B? / 7.75 points
(lumped). / $BR_B = 8.0$, twice the magnitude of A. / Are A and B different
responders? / No discernible difference. / B has 2x A's BR but 1/6 of A's
 PB . / Is a biomarker predicting BR useful here? / Cannot answer. / Yes
if the biomarker correlates with $\hat{m}_{BR}$. / Should A be continued long-term?
/ Cannot say. / Probably yes ($BR = 4$ is real and pharmacological), but the
larger placebo component will not persist if expectations decay. / Population
mean drug effect? / 7.75 points (biased). / $6.0 = (4.0 + 8.0)/2$ (correct mean
of BR). /*

The one-component model has produced one biased average where the three-component model has produced six clinically actionable quantities. Predictive-biomarker validation, which is a central scientific aim of the trial design family considered here, becomes impossible without the decomposition.

A population-level example shows that cohort differences in open-label totals driven entirely by PB produce apparent non-replicability that the decomposition resolves.

- Cohort 1 (naive, $m_{PB} = 7$) reports 13-point improvements; Cohort 2 (sceptical, $m_{PB} = 2$) reports 8 points; true $m_{BR} = 5$ in both.
- A one-component analysis would falsely conclude the drug works less well in Cohort 2 and trigger searches for confounders.
- The decomposition returns $\hat{m}_{BR} = 5$ in both cohorts, attributing the difference correctly to PB heterogeneity.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur

sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

A second example at the population level illustrates the same problem. Consider two cohorts of N -of-1 participants enrolled in methodologically identical trials of the same drug. Cohort 1 (early trial, naive patients) has mean $m_{BR} = 5$, mean $m_{PB} = 7$, mean $m_{TV} = 1$. Cohort 2 (later trial, patients who have had multiple prior failed treatments and are sceptical) has mean $m_{BR} = 5$, mean $m_{PB} = 2$, mean $m_{TV} = 1$. By construction the drug works equally well in both cohorts. But the open-label totals differ substantially: cohort 1 reports thirteen-point improvements on average, cohort 2 reports eight-point improvements. A one-component analysis comparing the two cohorts would conclude that the drug works less well in cohort 2 and might trigger a search for confounders or sub-population effects. None of this is real; the difference is entirely in PB, specifically in the population baseline level of placebo responsiveness. The three-component decomposition, applied to either cohort, returns $\hat{m}_{BR} = 5$ in both. Unfortunately, this is the kind of apparent non-replicability that pharmacology has spent the last decade struggling to navigate; the decomposition is one of the formal tools that makes it tractable.

7 Simulation study

The simulation study follows ADEMP and quantifies bias and identifiability properties of three analysis strategies across a prazosin-PTSD-calibrated parameter grid.

- The Morris (2019) ADEMP framework structures the study: Aims, DGP, Estimands, Methods, Performance measures.
- The pre-registration is at `analysis/scripts/component-decomposition/00-ademp-pre-registration.md`.
- Three analysis strategies (one-component, phase-augmented, full decomposition) are compared across the parameter grid.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The simulation study reported and pre-registered in this section is structured under the @morris2019 ADEMP framework, specifying Aims, Data-generating mechanisms, Estimands, Methods, and Performance measures. The pre-registered plan is at `analysis/scripts/component-decomposition/00-ademp-pre-registration.md`. A Monte Carlo simulation calibrated to the prazosin/PTSD reference parameters quantifies the bias and identifiability properties of the three analysis strategies described in section 5 across a parameter grid.

7.1 Pilot validation of analysis machinery (100 replicates, 6 cells)

1955
1956

1957
1958
1959
1960
1961

1962
1963
1964
1965
1966
1967
1968

1969
1970
1971

1972
1973

1974
1975
1976
1977

1978
1979
1980
1981
1982
1983
1984

1985
1986
1987
1988
1989
1990
1991
1992

1993
1994
1995

1996
1997

The pilot is a structural check of the analysis pipeline, not a substantive empirical result.

- The production simulation in section 7.2 constitutes the empirical core of the manuscript.
- The pilot confirms only that data generation, model fitting, and summary output run without error.
- No bias or power claims are derived from the 100-replicate pilot.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The production simulation programme described in the following section is the empirical core of this manuscript and has not yet been executed. The pilot reported here confirms only that the analysis machinery runs end-to-end.

The pilot executed a 6-cell, 100-replicate run with perfect convergence, confirming the pipeline is operational.

- The grid crossed three *pb_strength* levels against two analysis strategies.
- Convergence was 1.0 across all six cells under 8-core `mclapply`.
- Output is checkpointed at `analysis/data/quick-sim/`; no inferential claims are drawn.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

A 100-replicate-per-cell pilot was run at the reference Hybrid OL+BDC design, $N = 35$, on the 6-cell slice $pb_strength \in \{0, 6.5, 13\} \times analysis \in \{one-component, phase-augmented\}$, using 8-core `parallel::mclapply`. The convergence rate was 1.0 across all six cells. Per-replicate output is checkpointed at `analysis/data/quick-sim/06-decomp-replicates.rds` and the Morris-format cell-level summary at `analysis/data/quick-sim/06-decomp-summary.txt`. No derived inference is reported from the pilot; substantive bias and power claims are deferred to the production programme below.

Two pilot findings revise the production specification: rank-deficiency forces a formula change, and bias must be read from coefficient means, not rejection rates.

- At $N = 35$, rank-deficiency dropped the D_{bc} :phase term, eliminating the key biomarker-by-treatment-by-phase contrast.

- The predicted bias manifests in the mean of $\hat{\beta}_{bm:D_{bc}}$, not in the Wald-test rejection rate.
- The production programme addresses both by using larger N , a fuller design, and coefficient-mean reporting.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Two design-level observations from the pilot inform the production specification. First, at $N = 35$ the rank-deficiency of the full phase-augmented formula required dropping the D_{bc} :phase interaction term, removing the precise contrast that would have isolated the biomarker-by-treatment-by-phase signal the framework is designed to detect. Second, the predicted bias of the one-component analysis arises in the regression coefficient $\hat{\beta}_{bm:D_{bc}}$, not in the indicator-level rejection rate, so power on the biomarker-by-treatment Wald test is not the right summary statistic to detect the framework's predicted bias. The production programme below addresses both observations through larger N , a richer trial design that admits the full phase-augmented formula, and reporting of the coefficient mean against the implied true value rather than the rejection rate alone.

7.2 Production-simulation specification

The production simulation is organised around four deliverables under the ADEMP framework.

- Driver scripts are located at `analysis/scripts/component-decomposition/`.
- The ADEMP pre-registration is committed prior to production runs.
- The four deliverables address bias, identifiability, type I/power, and participant-level recovery.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The production programme follows the @morris2019 ADEMP framework. The driver scripts will live at `analysis/scripts/component-decomposition/`. We shall describe the four deliverables in turn.

The first deliverable quantifies how much each analysis strategy biases the biomarker-treatment interaction estimate as placebo strength increases.

- The estimand is mean $\hat{\beta}_{bm:D_{bc}}$ with Monte Carlo SE across 1,000 replicates per cell.

- 2042 • Sample sizes $N \in \{35, 70, 100, 150\}$ are crossed with a *pb_strength* grid.
- 2043 • The headline finding is the bias-versus-power trade-off invisible to the 100-
- 2044 replicate prototype.

2045 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-

2046 mod tempor incidunt ut labore et dolore magna aliqua. Ut enim ad

2047 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip

2048 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in

2049 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur

2050 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt

2051 mollit anim id est laborum.

2052 *The first deliverable is the bias of the biomarker-treatment interaction*

2053 *estimate: mean and Monte Carlo standard error of $\hat{\beta}_{bm:D_{bc}}$ for the*

2054 *one-component, phase-augmented, and full-decomposition analyses, across*

2055 *sample sizes $N \in \{35, 70, 100, 150\}$ and a grid of *pb_strength* values, at*

2056 *$n_{reps} = 1,000$ per cell. The headline measure is the bias-versus-power*

2057 *trade-off that the prototype was unable to detect at $N = 35$.*

2058 **The second deliverable assesses whether the full-decomposition analysis**

2059 **is identifiable at trial-relevant sample sizes.**

- 2060 • Identifiability is measured via convergence rates of the nonlinear-mixed-
- 2061 effects fit.
- 2062 • Variance-inflation factors on the Gompertz parameters diagnose collinearity
- 2063 across components.
- 2064 • Sensitivity of the *BR* estimate to $\eta(\text{phase})$ multipliers is reported explicitly.

2065 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-

2066 mod tempor incidunt ut labore et dolore magna aliqua. Ut enim ad

2067 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip

2068 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in

2069 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur

2070 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt

2071 mollit anim id est laborum.

2072 *The second deliverable is a set of identifiability metrics for the full-*

2073 *decomposition analysis: convergence rates of the nonlinear-mixed-effects fit,*

2074 *variance-inflation factors on the component-specific Gompertz parameters,*

2075 *and the sensitivity of the *BR* estimate to the chosen $\eta(\text{phase})$ multipliers.*

2076 **The third deliverable characterises type I error and power across analysis**

2077 **strategies under a full sensitivity grid.**

- 2078 • The DGP is the three-component reference; sensitivity axes are η mismatch,
- 2079 *TV* magnitude, and dropout.
- 2080 • Type I error cells use 5,000 replicates to achieve MCSE 0.003.
- 2081 • All three analysis strategies (one-component, phase-augmented, full-
- 2082 decomposition) are compared.

2083 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-

2084 mod tempor incidunt ut labore et dolore magna aliqua. Ut enim ad

2085 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip

ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The third deliverable is type I error and power for the biomarker-treatment interaction test under each analysis strategy, with the DGP held at the three-component reference and varied across a sensitivity grid (analyst-versus-truth η mismatch, TV-component magnitude, and dropout). Type I cells are run at $n_{reps} = 5,000$ for MCSE 0.003.

The fourth deliverable illustrates which analysis strategies preserve participant-level biomarker stratification.

- Participant-level interaction signals are mapped as a function of (m_{BR}, m_{PB}, m_{TV}) .
- The comparison reveals whether each strategy retains the predictive-biomarker signal.
- The lumped-total regression serves as the baseline that fails to stratify correctly.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The fourth deliverable is the participant-level interaction signal recovered by each analysis, as a function of the underlying participant-level (m_{BR}, m_{PB}, m_{TV}) values, illustrating which analysis strategies preserve the predictive-biomarker stratification that the lumped-total regression does not.

Production runs use the pmsimstats-ng tidyverse collection with checkpointed outputs and a committed pre-registration.

- The tidyverse implementation collection drives the simulation grid.
- Cell-level summaries are checkpointed under analysis/data/; per-replicate files under output/.
- The ADEMP pre-registration is committed at 00-ademp-pre-registration.md before any production run.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The simulation will use the existing pmsimstats-ng tidyverse implementation collection. Cell-level summaries will be checkpointed under