

1 Test-procedure choice for the biomarker-treatment
2 interaction in aggregated N-of-1 trials: dichotomization,
3 classical RM-ANOVA, mixed-effects, and generalized
4 estimating equations

5 pmsimstats team

6 2026-06-19

7 **Contents**

8	1 Abstract	2
9	2 Introduction	3
10	2.1 A claim about how design and test choices interact	3
11	2.2 Why the question matters for chronic-condition pharmacology . . .	4
12	2.3 What the published N-of-1 methodology literature does and does	
13	not do	6
14	2.4 Run-in design and the placebo-selection controversy	9
15	2.5 What this paper contributes	10
16	3 Closed-form strict RM-ANOVA test of the biomarker-treatment	
17	interaction	12
18	3.1 Notation and design	13
19	3.2 Sums of squares and the interaction F -statistic	15
20	3.3 Relation to the linear-mixed Wald test	20
21	4 Cycle-by-period design grid	26
22	5 Simulation design	27
23	6 Results	30
24	6.1 Test-procedure comparison at fixed design	31
25	7 Design-and-test recommendation	36
26	8 Discussion	38
27	8.1 What the Monte Carlo study settles	38
28	9 Conclusions	41

29	10 Conflict of interest	43
30	11 Funding	43
31	12 Data availability statement	44
32	13 Reproducibility	44
33	14 References	45

34 1 Abstract

35 **Background.** Inference about the biomarker-treatment interaction in aggregated
36 N-of-1 trials depends on the analyst’s test procedure: classical repeated-measures
37 ANOVA (RM-ANOVA), the linear mixed-effects model (LME), or generalized esti-
38 mating equations (GEE). A common but rarely examined feature of the classical
39 analysis is that it requires the continuous biomarker to be dichotomized. We
40 quantify the consequences of test-procedure choice for the interaction, separating
41 the cost of dichotomization from the cost of the test procedure proper.

42 **Methods.** We derive the closed-form strict RM-ANOVA F -statistic for the
43 biomarker-stratified treatment-by-time interaction in a balanced split-plot design
44 with categorical predictors and the sphericity assumption, and relate it to the
45 Wald t -test for the corresponding fixed-effects coefficient in a linear mixed model
46 with continuous-time AR(1) residual correlation. We then conduct a Monte Carlo
47 simulation study, designed and reported under the¹ ADEMP framework, that
48 crosses five test procedures (strict RM-ANOVA on the dichotomized biomarker;
49 a continuous-biomarker within-subject contrast; the standard pmsimstats linear-
50 mixed analysis with `corCAR1`; and GEE with naive and with Mancl-DeRouen
51 small-sample sandwich variance) against the biomarker-moderation coefficient
52 $c_{bm} \in \{0, 0.45\}$ and sample size $N \in \{25, 35, 70, 100\}$, a 40-cell factorial at the
53 prazosin-PTSD reference design with 1000 replicates per power cell and 2000 per
54 Type I cell. The continuous-biomarker contrast differs from strict RM-ANOVA
55 only in retaining the continuous biomarker, isolating the cost of dichotomization.
56 The cycle-by-period design grid is specified here but its empirical sweep is deferred
57 to a companion paper.

58 **Results.** The closed-form strict RM-ANOVA test agrees with the linear-mixed
59 Wald test under sphericity and balanced design. In the simulation, most of the
60 power gap between strict RM-ANOVA and the mixed model is the cost of di-
61 chotomization rather than of the test procedure: at $N = 25$, replacing the di-
62 chotomized biomarker with the continuous one raises power by 0.180, while the
63 further step to the full mixed model adds only 0.027, a pattern stable across
64 N (dichotomization 0.08–0.21, test procedure at most 0.04). Among the four
65 continuous-biomarker procedures, the linear-mixed analysis attains the highest
66 power within nominal Type I error; GEE with the naive sandwich inflates Type
67 I error to between $1.5\times$ and $1.9\times$ nominal across all four sample sizes, which the

68 Mancl-DeRouen correction removes at a moderate (about 10-percentage-point)
69 power cost while restoring 93–95% interval coverage. The inline Mancl-DeRouen
70 standard error agrees with the `geesmv` reference to within 2.5%.

71 **Conclusions.** The analyst’s most consequential decision is to keep the biomarker
72 continuous rather than dichotomize it; among continuous-biomarker procedures
73 the linear-mixed analysis is the preferred default and a dichotomization-free
74 within-subject contrast is a close, transparent alternative. Naive GEE should
75 not be used at the sample sizes examined without a small-sample correction.
76 This paper resolves the test-procedure axis at a fixed reference design; the joint
77 cycle-by-period design grid is specified here and deferred to a companion paper.

78 2 Introduction

79 2.1 A claim about how design and test choices interact

80 [1]

81 **Joint optimization of test procedure and design parameters**
82 **is the paper’s central argument.**

- 83 • The two principal choices – test procedure and trial-design pa-
84 rameters – are mutually dependent: validity and power of each
85 test depend on the carryover and autocorrelation structure the
86 design imposes.
- 87 • Treating the choices as independent optimizations, as the
88 published literature typically does, produces locally optimal but
89 jointly suboptimal combinations.
- 90 • The argument is presented as consequential at the sample sizes
91 characteristic of precision-medicine pilot studies.

92 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
93 *tempor incididunt ut labore et dolore magna aliqua.*

94 The present paper develops the methodological case for treating the
95 analyst’s two principal choices in an aggregated N-of-1 trial jointly
96 rather than separately. The two choices are the test procedure for the
97 biomarker-treatment interaction (strict repeated-measures ANOVA,
98 the linear mixed-effects model with continuous-time AR(1) residual
99 correlation, the mixed-effects model for repeated measures, general-
100 ized estimating equations) and the trial-design parameters that gov-
101 ern within-subject treatment cycling (number of treatment cycles, on-
102 drug and off-drug period lengths, on/off symmetry, optional run-in
103 and run-out phases). The validity and power of any given test de-
104 pend on what the design assumes about carryover, autocorrelation,
105 and within-subject variance structure; conversely, the optimal design
106 parameters depend on which test the analyst plans to use. Treating
107 the two as independent optimizations, as the published methodological

literature typically does, produces locally optimal but jointly suboptimal combinations, and we shall argue that this gap is consequential at the sample sizes characteristic of precision-medicine pilot studies.

[2]

The paper’s two-part argument establishes a gap in the literature and then fills it at the test-procedure axis.

- No joint design-by-test sensitivity study exists for the aggregated N-of-1 design class.
- The Monte Carlo experiment shows that linear-mixed with AR(1) dominates strict RM-ANOVA wherever residuals depart from compound symmetry.
- Strict RM-ANOVA is retained as a conservative stress-test baseline rather than a primary recommendation.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The case we develop is two-part. We begin by establishing that a joint design-by-test sensitivity study has not, to the best of our knowledge, been performed for the aggregated N-of-1 design class, and that the absence of one limits the field’s ability to make defensible joint recommendations. We then demonstrate, through a focused Monte Carlo experiment, that the linear mixed-effects test with continuous-time AR(1) residual correlation, used in conjunction with the Hendrickson hybrid open-label-plus-blinded-discontinuation design², dominates the strict classical RM-ANOVA test in essentially every parameter cell where within-participant residuals depart from compound symmetry. That is, strict RM-ANOVA remains useful as a stress test — the most stringent baseline that compound-symmetry machinery can deliver — but not as a recommendation in its own right.

2.2 Why the question matters for chronic-condition pharmacology

In this section we shall set out three reasons why the joint treatment of test procedure and design parameters is warranted specifically for the aggregated N-of-1 design class.

[3]

Physiological autocorrelation in chronic-condition trials violates the sphericity assumption that strict RM-ANOVA requires.

- Chronic-condition response trajectories exhibit serial correlation driven by drug pharmacokinetics, natural-history drift, and measurement reactivity.

- Huynh-Feldt and Greenhouse-Geisser corrections recover nominal Type I only conditionally; the conventional threshold favours Greenhouse-Geisser when $\epsilon < 0.75$ (Blanca et al., 2023).
- The mixed-effects model with `corCAR1` avoids these corrections by modeling AR(1) structure directly.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

We begin with the physiological motivation. Within-participant residuals in chronic-condition trials are autocorrelated not as an analytic convenience but because most chronic-condition response trajectories exhibit serial correlation on a timescale set by drug pharmacokinetics, condition-specific natural-history drift, and measurement reactivity (Hawthorne, structured-visit, and diary-completion effects). Continuous-time AR(1) residual correlation is, in our view, the appropriate structural prior for this physiology, but it does not satisfy the sphericity assumption required by strict RM-ANOVA.³ show through simulation that uncorrected RM-ANOVA inflates Type I error under sphericity violations; they also report, importantly, that the Greenhouse-Geisser-corrected RM-ANOVA and the multilevel linear model perform comparably once the correction is applied, and that the multilevel model is not uniformly safer (it can be the more liberal of the two in some conditions). The degrees-of-freedom corrections^{4,5} recover nominal levels only conditionally.⁶ extend this evidence with finer epsilon resolution, with the conventional rule favouring Greenhouse-Geisser when $\epsilon < 0.75$. The advantage we claim for the mixed-effects model with explicit `corCAR1` residual structure is thus not that it is uniformly more powerful than a well-corrected RM-ANOVA, but that it models the AR(1) structure directly, accommodates the continuous biomarker and unequal time spacing, and so avoids both the dichotomization and the degrees-of-freedom penalty that the classical route incurs.

[4]

Detecting the biomarker-treatment interaction is harder than detecting the average treatment effect (ATE) and is more sensitive to design choices.

- The PATH framework distinguishes effect-modeling from risk-modeling; the former requires more design-aware machinery.
- Within-participant contrasts feeding the interaction estimator are sensitive to trial-design parameters in ways the ATE analysis is not.
- Underpowered test procedures for the interaction will miss predictive-biomarker signal even at nominally adequate sample sizes.

191 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
192 *tempor incididunt ut labore et dolore magna aliqua.*

193 The biomarker-treatment interaction also merits separate attention
194 because it is harder to detect than the average treatment effect (ATE).
195 The PATH framework^{7,8} makes the methodological distinction precise:
196 effect-modeling (the predictive heterogeneity-of-treatment-effects,
197 or HTE, question) requires more design-aware machinery than
198 risk-modeling (the ATE question). The within-participant contrasts
199 that contribute to the interaction estimator are sensitive to the
200 trial-design parameters in ways the ATE analysis is not, and a test
201 procedure that underpowers the interaction term will systematically
202 miss predictive-biomarker signal even at sample sizes that appear
203 adequate for the ATE test.

204 [5]

205 **Cycle count and period length interact with carryover and**
206 **test-procedure choice, so their joint optimum cannot be**
207 **found by optimizing each separately.**

- 208 • Additional cycles increase power for the biomarker-treatment in-
209 teraction but with diminishing returns beyond the fourth or fifth
210 cycle.
- 211 • Shorter periods produce more cycles but worsen on-drug/off-drug
212 separation when carryover is non-trivial.
- 213 • The joint optimum on the (cycles, period length) plane depends
214 on carryover half-life and test procedure together.

215 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
216 *tempor incididunt ut labore et dolore magna aliqua.*

217 Finally, cycle count and period length jointly determine the
218 within-subject contrast structure in ways that cannot be optimized
219 independently.⁹ and¹⁰ establish, in the broader cross-over literature,
220 that adding within-subject treatment cycles contributes informa-
221 tion about the patient-by-treatment interaction variance that no
222 cross-sectional comparator can supply. The N-of-1 specialization of
223 this principle is that more cycles increase power for the biomarker-
224 treatment interaction, but with diminishing returns: cycles beyond
225 the fourth or fifth contribute marginal information at the cost of
226 trial duration. Period length interacts with carryover — shorter
227 periods produce more cycles but worse separation between on-drug
228 and off-drug states once carryover is non-trivial. The joint optimum
229 on the (cycles, period length) plane depends on the carryover half-life
230 and the chosen test procedure together, not on each factor in isolation.

231 **2.3 What the published N-of-1 methodology literature does and**
232 **does not do**

[6]

The majority of published N-of-1 analyses are conducted under a misspecified residual covariance assumption.

- Systematic reviews document that 83.8% of published N-of-1 studies ignore autocorrelation entirely.
- CONSORT extension standards codify reporting quality but do not address the analytic-model decision.
- The analytic-method gap has been identified as the highest-priority methodological problem in the field.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

We begin with the state of the analytic-model literature. The systematic review of¹¹ documents that 83.8% of the published N-of-1 studies they catalogue ignore autocorrelation entirely;¹² and¹³ reach similar conclusions in neurology and behavioral-medicine subdomains. CONSORT extension reporting standards^{14,15} codify reporting quality but do not address the analytic-model decision in any detail, and¹⁶ identifies the analytic-method gap as the highest-priority methodological problem in the field. It appears, then, that the majority of N-of-1 analyses are conducted under a misspecified residual covariance assumption, and that the field has not yet converged on remedies.

[7]

Existing test-procedure comparisons address parallel-group or cluster designs but not the aggregated N-of-1 setting.

- The mixed model for repeated measures (MMRM) consensus applies to parallel-group longitudinal trials with monotone missingness, not N-of-1 designs.
- Only 6.6% of stepped-wedge cluster trials with fewer than six units use appropriate small-sample corrections.
- No existing test-procedure comparison addresses the cross-over or N-of-1 setting directly.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Test-procedure comparisons for longitudinal data are not new. The generic contrast of RM-ANOVA against GEE and mixed-effects models is well established:¹⁷ for instance, report through simulation that RM-ANOVA is appreciably less powerful than GEE and mixed models for continuous longitudinal endpoints, the same ordering we recover. What that literature does not do is treat the biomarker-by-treatment interaction estimand, the aggregated N-of-1 / cross-over design, or the continuous-time AR(1) structure together, nor does

it isolate the dichotomization cost that the classical route imposes specifically when the moderator is continuous. The MMRM consensus established by Mallinckrodt and colleagues^{18–21} applies to parallel-group longitudinal trials with monotone missingness, not to aggregated N-of-1 designs. The GEE methodology of²² with sandwich variance estimation has been adapted to small-cluster settings via bias-corrected estimators^{23,24}, but the small-sample concerns are acute at trial-relevant N-of-1 cluster counts:²⁵ report that in stepped-wedge cluster trials with fewer than six independent units only 6.6% of analyses use appropriate small-sample corrections, illustrating the gap between method availability and practice. None of these papers address the biomarker-interaction estimand in the cross-over or N-of-1 setting directly.

[8]

Design-parameter optimization has been studied in isolation from test-procedure choice, leaving the joint question undressed.

- Hendrickson et al. (2020) is the only paper comparing full design families for biomarker-interaction power, but holds period length fixed and uses a single test procedure.
- Power calculations, autocorrelation effects, and cross-over design literature each address one axis without crossing them.
- No published study crosses test procedure against the design-parameter grid under simulation for the aggregated N-of-1 class.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Design-parameter optimization has been studied, but separately from test-procedure choice.²⁶ develop power calculations for $(AB)^k$ designs that match the N-of-1 within-subject contrast structure;²⁷ document strong autocorrelation effects on permutation-test power; and¹⁰ demonstrates how replicate cross-over designs allow direct estimation of patient-by-treatment interaction variance, the conceptual ancestor of N-of-1 aggregation. The Hendrickson et al. framework² is, to the best of our knowledge, the only paper that compares full design families head-to-head for the biomarker-interaction power question, but it holds period length fixed and uses a single test procedure.²⁸ compare aggregated N-of-1 trials against parallel and cross-over designs by simulation and find the aggregated design more efficient, but analyse a single mixed model and do not vary the test procedure.²⁹ and³⁰ synthesize the broader N-of-1 design literature without addressing the joint design-by-test sensitivity question.

[9]

No published study crosses test procedure against a design-parameter grid for the aggregated N-of-1 design class.

- The gap spans test-procedure comparison (strict RM-ANOVA, MMRM, GEE, mixed-effects with `corCAR1`) crossed against cycles, period lengths, asymmetry, run-in, and run-out.
- The present paper addresses the gap on the test-procedure axis.
- The design-parameter axis is deferred to a companion paper.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The gap, then, is this. Across the recent systematic reviews^{11,12,31} and the methodological literature surveyed above, no published study crosses test procedure (strict RM-ANOVA, MMRM, GEE, mixed-effects with `corCAR1`) against a design-parameter grid (cycles, period lengths, asymmetry, run-in, run-out) under simulation for the aggregated N-of-1 design class. The present paper endeavors to address that gap, at least on the test-procedure axis.

2.4 Run-in design and the placebo-selection controversy

[10]

The Hendrickson hybrid design avoids placebo-run-in selection bias by using open-label stabilization, at some power cost.

- Response-based exclusion during a placebo run-in introduces selection bias that damages external and potentially internal validity (Berger et al., 2003).
- The Hendrickson open-label stabilization phase avoids this bias but forgoes the power that aggressive response-based selection would supply.
- The run-in choice is held fixed in the present paper; varying it is deferred to the companion design-sensitivity paper.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A separate methodological controversy that the joint design-by-test framing intersects concerns the use of placebo run-in phases to exclude placebo responders before randomization.³² show formally that response-based exclusion during run-in introduces a selection bias that damages external validity and may affect internal validity as well. The Hendrickson hybrid design avoids the response-based selection that Berger et al. flag by using an open-label stabilization phase rather than a placebo run-in; the design choice is therefore conservative on the run-in question, though it forgoes the power that aggressive response-based selection would otherwise supply. The cycle-by-period

sweeps reported here take the Hendrickson design’s run-in choice as fixed; varying it is left to the companion design-sensitivity paper at [analysis/report/05-nof1-design-sensitivity/](#).

2.5 What this paper contributes

We make four contributions, each addressed in a section below.

[11]

The first contribution is a closed-form strict RM-ANOVA F -statistic for the biomarker-treatment interaction and its relation to the linear-mixed Wald test.

- The derivation specializes the balanced split-plot model to the biomarker-stratified case under the sphericity assumption.
- The conditions under which the two tests agree and diverge are identified explicitly.
- The derivation extends standard cross-over methodology to the biomarker-stratified case.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

First, we derive the closed-form strict RM-ANOVA F -statistic for the biomarker-stratified treatment-by-time interaction in a balanced split-plot design with categorical predictors and the sphericity assumption (§2). We relate the strict-RM-ANOVA test to the Wald t -test on the biomarker-treatment fixed-effects coefficient in a linear mixed model with continuous-time AR(1) residual correlation, identifying the conditions on residual covariance under which the two tests agree and the regimes in which they diverge. The derivation extends standard cross-over methodology^{9,33} to the biomarker-stratified case and clarifies the connection to the contemporary sphericity-violation literature^{3,6}.

[12]

The second contribution is specification of the cycle-by-period design grid calibrated to the prazosin/PTSD reference application.

- The grid varies cycle count, on-drug and off-drug period lengths, and on/off symmetry around the Hendrickson hybrid reference.
- Calibration is to the prazosin/PTSD application that motivates the pmsimstats program.
- The empirical sweep across this grid is deferred to a companion paper.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Second, we specify a cycle-by-period design grid (§3) that varies the number of treatment cycles, on-drug and off-drug period lengths, and on/off symmetry around the Hendrickson hybrid reference. The grid is calibrated to the prazosin/PTSD reference application that motivates the pmsimstats program.

[13]

The third contribution is a Monte Carlo simulation study crossing test procedures against the design grid under the ADEMP framework.

- The simulation runs on the pmsimstats package data-generating process under Architecture A (mean moderation), exercised through `generateData()` and `lme_analysis()`.
- Five test procedures, including a continuous-biomarker contrast that isolates the dichotomization cost, are crossed against biomarker-moderation coefficient and sample size in a 40-cell factorial.
- Results are reported following the Morris et al. (2019) ADEMP discipline including Monte Carlo standard errors (MCSE), bias, and interval coverage.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Third, we conduct a Monte Carlo simulation study (§4–§5) that crosses five test procedures (strict RM-ANOVA on the dichotomized biomarker; a continuous-biomarker within-subject contrast; the standard pmsimstats linear-mixed analysis with `corCAR1`; and GEE with naive and with Mancl-DeRouen³⁴ small-sample sandwich variance) against the biomarker-moderation coefficient and sample size at the fixed reference design. The continuous-biomarker contrast is the dichotomization-free analogue of the strict RM-ANOVA test and lets us separate the cost of dichotomizing the biomarker from the cost of the test procedure proper. The simulation runs on the pmsimstats package data-generating process under Architecture A (mean moderation) via `generateData()` and `lme_analysis()`, holding fixed the response-curve and covariance-structure axes addressed in `analysis/report/07-gompertz-evaluation/` and `analysis/report/01-dgp-mean-moderation-vs-mvn/`. Results are reported in the¹ ADEMP framework adopted as the project-wide simulation-reporting standard.

[14]

The fourth contribution is a design-and-test recommendation table indexed by carryover half-life and target effect size.

- The recommendation is presented as preferred (cycle count, pe-

riod length, test procedure) combinations rather than a single best design.

- The recommendation is narrowed to the test-procedure axis at the reference design in the present paper.
- The joint cycle-by-period optimization is reserved for the companion paper.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Fourth, we synthesize a design-and-test recommendation table (§6) for typical chronic-condition N-of-1 trial scenarios. The recommendation is presented as a table of preferred (cycle count, period length, test procedure) combinations indexed by the carryover half-life and the target effect size, rather than as a single ‘best’ design.

[15]

The paper’s organization maps each section to one of the four contributions.

- Sections 2 and 3 cover the closed-form derivation and the design-grid specification respectively.
- Sections 4 and 5 present the simulation design and results.
- Sections 6 and 7 synthesize the recommendation and discuss implications.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The remainder of the paper is organized as follows. Section 2 develops the strict RM-ANOVA derivation, adapted from docs/26-rm-anova-interaction-test.md. Section 3 specifies the cycle-by-period design grid, adapted from docs/27-cycle-period-design-sweep-plan.md. Section 4 lays out the simulation design. Section 5 reports the results. Section 6 synthesizes the joint design-and-test recommendation. Section 7 discusses the implications for the methodology of N-of-1 biomarker-validation trials and for the operating characteristics of the pmsimstats simulation program.

3 Closed-form strict RM-ANOVA test of the biomarker-treatment interaction

[16]

The section derives the strict RM-ANOVA closed form and identifies three findings that anchor the rest of the paper.

- The derivation is confined to the strict classical compound-symmetric form; modern extensions appear only for context.

- The simplest design supporting the test has two biomarker strata, two treatment phases, and T within-phase timepoints.
- Three findings from the derivation – equivalence to a two-sample t -test, a closed-form power formula, and five data-generating process (DGP) violations – anchor the simulation study.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In this section we shall derive the closed-form strict RM-ANOVA F -statistic for the biomarker-treatment interaction in a balanced split-plot design with categorical predictors and the sphericity assumption, and relate it to the Wald t -statistic on the corresponding linear-mixed coefficient. The derivation specializes classical split-plot machinery³⁵ to the biomarker-stratified case in the simplest design that supports the test: two biomarker strata, two treatment phases, and T within-phase timepoints. The presentation is deliberately confined to strict RM-ANOVA. Modern extensions — the mixed-effects model for repeated measures³⁶, generalized estimating equations²², and ANCOVA with random subject effects — appear only for context; the closed-form expression we report is the textbook compound-symmetric form. Three findings anchor the rest of the paper:

1. Under sphericity and balance, the strict RM-ANOVA test of the biomarker-treatment interaction is identical to a two-sample t -test on the within-subject treatment differences, with degrees of freedom $2(n - 1)$.
2. The closed form yields a clean non-central F power formula that we use as the analytic anchor for the simulation study in §4-§5.
3. Five assumptions of the closed form are violated by the pmsimstats data-generating process (DGP; continuous biomarker, AR(1) within-subject covariance, exponential carryover, missing data, unequal time spacing). Each violation favors the linear-mixed analysis over the strict RM-ANOVA test in operating characteristics, motivating the primary-and-sensitivity reporting strategy of §6.

3.1 Notation and design

[17]

The notation section establishes a restricted design that is deliberately simpler than the pmsimstats DGP.

- The design is the simplest that supports a biomarker-treatment interaction test in a within-subject setting.
- The gap between this restricted design and the pmsimstats DGP is addressed explicitly in Section 2.4.
- All departures from the DGP are flagged as sources of conservative bias in the strict RM-ANOVA operating characteristics.

518 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
519 *tempor incididunt ut labore et dolore magna aliqua.*

520 We begin with the simplest design that supports a biomarker-
521 treatment interaction test in a within-subject setting. The reader will
522 note that the design structure described here is more restrictive than
523 the pmsimstats DGP; we return to that gap explicitly in §2.4.

524 Specifically, we consider the following structural elements:

- 525 • Two biomarker strata $i \in \{1, 2\}$, formed by dichotomizing the continuous
526 biomarker at a pre-specified cut-point. Strict RM-ANOVA requires cat-
527 egorical predictors; the dichotomization is the price of admission to the
528 framework.
- 529 • n subjects per stratum, balanced. Index $l \in \{1, \dots, n\}$.
- 530 • Two treatment phases $j \in \{1, 2\}$, active drug and placebo. In an N-of-1
531 design these are two phases through which every subject passes; in a parallel-
532 group design they are between-subject and the design is not a split-plot.
- 533 • T within-phase timepoints, indexed $k \in \{1, \dots, T\}$. In the simplest case
534 $T = 1$, the analyst pre-averages within-phase observations to a single mean
535 per subject per phase.

536 [18]

537 **The split-plot structure arises because the biomarker stratum**
538 **is between-subjects while treatment phase and timepoints are**
539 **within-subjects.**

- 540 • Strict RM-ANOVA requires categorical predictors; dichotomizing
541 the continuous biomarker at a cut-point is the price of admission.
- 542 • Each subject is nested in exactly one biomarker stratum but
543 passes through both treatment phases.
- 544 • The design is therefore a canonical split-plot repeated-measures
545 design.

546 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
547 *tempor incididunt ut labore et dolore magna aliqua.*

548 The biomarker stratum is between-subjects (each subject sits in ex-
549 actly one stratum); the treatment phase and the timepoints are within-
550 subjects (each subject sees both phases at every timepoint). This is
551 the canonical split-plot repeated-measures design³⁵.

552 [19]

553 **The classical split-plot model encodes the biomarker-**
554 **treatment interaction as the $(\alpha\beta)_{ij}$ fixed effect under**
555 **compound symmetry.**

- 556 • Sum-to-zero constraints are imposed on all fixed effects; the ran-
557 dom subject effect is nested in biomarker stratum.

- The sphericity (compound symmetry) condition requires equal within-subject covariance across all observation pairs.
- The biomarker-treatment interaction $(\alpha\beta)_{ij}$ and its time-resolved form $(\alpha\beta\tau)_{ijk}$ are the primary targets.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

We write Y_{ijkl} for the symptom outcome at biomarker stratum i , treatment phase j , subject l within stratum i , and timepoint k . The strict classical split-plot model is

$$Y_{ijkl} = \mu + \alpha_i + \beta_j + (\alpha\beta)_{ij} + \pi_{l(i)} + \tau_k + (\alpha\tau)_{ik} + (\beta\tau)_{jk} + (\alpha\beta\tau)_{ijk} + e_{ijkl}, \quad (1)$$

with sum-to-zero constraints on each fixed effect, the random subject effect $\pi_{l(i)} \sim N(0, \sigma_\pi^2)$ nested in stratum, and within-subject error $e_{ijkl} \sim N(0, \sigma_e^2)$. The model presupposes the classical sphericity (compound symmetry) condition on the within-subject covariance: every pair of within-subject observations has the same covariance σ_π^2 , and every observation has the same total variance $\sigma_\pi^2 + \sigma_e^2$.

3.2 Sums of squares and the interaction F -statistic

[20]

The total sum of squares partitions into between-subjects and within-subjects components that are orthogonal under balance.

- The between-subjects component captures biomarker-stratum effects and subject-within-stratum variation.
- The within-subjects component captures treatment-phase effects, the biomarker-treatment interaction, and within-subjects error.
- Orthogonality under balance ensures the interaction F -ratio uses the correct within-subjects error term.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

We begin the derivation by partitioning the total sum of squares. The total sum of squares partitions into between-subjects and within-subjects components:

$$SS_{\text{tot}} = SS_{\text{between}} + SS_{\text{within}}. \quad (2)$$

The between-subjects component decomposes further into SS_A (biomarker stra-

590 tum) and $SS_{S(A)}$ (subject within stratum):

$$SS_A = bnT \sum_{i=1}^a (\bar{Y}_{i...} - \bar{Y}_{...})^2, \quad SS_{S(A)} = bT \sum_{i=1}^a \sum_{l=1}^n (\bar{Y}_{i.l.} - \bar{Y}_{i...})^2. \quad (3)$$

591 The within-subjects component decomposes into SS_B (treatment phase), SS_{AB}
 592 (the biomarker-treatment interaction, the primary target of the test), and a within-
 593 subjects error term $SS_{B \times S(A)}$:

$$SS_{AB} = nT \sum_{i,j} (\bar{Y}_{ij..} - \bar{Y}_{i...} - \bar{Y}_{.j..} + \bar{Y}_{...})^2, \quad (4)$$

$$SS_{B \times S(A)} = T \sum_{i,j,l} (\bar{Y}_{ijl.} - \bar{Y}_{i.l.} - \bar{Y}_{.jl.} + \bar{Y}_{i...})^2. \quad (5)$$

594 [21]

595 **The interaction F -ratio uses the within-subjects error**
 596 **$SS_{B \times S(A)}$ as its denominator, with degrees of freedom**
 597 **$a(n-1)(b-1)$.**

- 598 • SS_{AB} has $df_{AB} = (a-1)(b-1)$ degrees of freedom; the within-
 599 subjects error has $df_{B \times S(A)} = a(n-1)(b-1)$.
- 600 • The within-subjects error term is the appropriate denominator for
 601 the treatment effect and its interaction with the between-subjects
 602 factor.
- 603 • The F -ratio $F_{AB} = MS_{AB}/MS_{B \times S(A)}$ follows a central F under
 604 H_0 and sphericity.

605 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
 606 *tempor incididunt ut labore et dolore magna aliqua.*

607 Corresponding degrees of freedom are $df_A = a-1$, $df_{S(A)} = a(n-1)$,
 608 $df_B = b-1$, $df_{AB} = (a-1)(b-1)$, and $df_{B \times S(A)} = a(n-1)(b-1)$,
 609 the last being the within-subjects error degrees of freedom appropriate
 610 for the treatment effect and its interaction with the between-subjects
 611 factor. The biomarker-treatment interaction is then tested by the F -
 612 ratio

$$F_{AB} = \frac{MS_{AB}}{MS_{B \times S(A)}} = \frac{SS_{AB}/df_{AB}}{SS_{B \times S(A)}/df_{B \times S(A)}}. \quad (6)$$

613 Under the null $H_0 : (\alpha\beta)_{ij} \equiv 0$ and the sphericity assumption, $F_{AB} \sim$
 614 $F(df_{AB}, df_{B \times S(A)})$.

615 3.2.1 The 2×2 split-plot case

616 [22]

Specializing to $a = b = 2$ with $T = 1$ reduces the interaction to the difference of within-subject treatment differences across biomarker strata.

- Setting $T = 1$ corresponds to pre-averaging within-phase observations before computing the interaction.
- The sample interaction contrast $\hat{L}$ is the difference of within-subject treatment differences across strata.
- This specialization yields the cleanest closed-form expression for the interaction F -statistic.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

We now specialize to $a = b = 2$ (two biomarker strata, two phases), setting $T = 1$ for cleanest exposition; that is, the analyst pre-averages within-phase observations to a single value per subject per phase before computing the interaction. The sample interaction contrast is

$$\hat{L} = (\bar{Y}_{11\cdot} - \bar{Y}_{12\cdot}) - (\bar{Y}_{21\cdot} - \bar{Y}_{22\cdot}), \quad (7)$$

[23]

Under sphericity and balance in the 2×2 case, the interaction F -statistic reduces to $nT\hat{L}^2/(4\hat{\sigma}_e^2)$ with degrees of freedom $(1, 2(n-1))$.

- The interaction sum of squares is $SS_{AB} = nT\hat{L}^2/4$; with $df_{AB} = 1$ this equals the interaction mean square.
- Under sphericity, $MS_{B \times S(A)} = \hat{\sigma}_e^2$, yielding the closed form directly.
- The error degrees of freedom are $2(n-1)$, reflecting one degree of freedom per subject in each of the two strata.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

That is, $\hat{L}$ is the difference of within-subject treatment differences across biomarker strata, and is the natural biomarker-treatment interaction statistic. The interaction parameter under the sum-to-zero ANOVA parameterization satisfies $\widehat{(\alpha\beta)}_{11} = \hat{L}/4$ (and the other three cells take values $\pm\hat{L}/4$), so the interaction sum of squares is $SS_{AB} = nT\hat{L}^2/4$, and with $df_{AB} = 1$ this equals the interaction mean square. Under sphericity and balance, $MS_{B \times S(A)} = \hat{\sigma}_e^2$. Substituting into equation (6) yields

$$F_{AB} = \frac{nT\hat{L}^2}{4\hat{\sigma}_e^2}, \quad df_1 = 1, \quad df_2 = 2(n-1). \quad (8)$$

[24]

Equation 8 is the paper’s strict RM-ANOVA F reference expression and is named to distinguish it from the mixed-effects and GEE Wald statistics.

- The expression applies to the $2 \times 2 \times T$ split-plot design and is the textbook compound-symmetric form.
- It serves as the analytic anchor for the power formula in Section 2.3.1 and for comparison with simulation results in Section 5.
- The name ‘strict RM-ANOVA F ’ is used throughout to mark the distinction from Wald statistics that do not require sphericity.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Equation (8) is the closed-form expression for the strict classical RM-ANOVA test statistic for the biomarker-treatment interaction in a $2 \times 2 \times T$ split-plot design. We shall refer to it as the ‘strict RM-ANOVA F ’ throughout the rest of the paper, to distinguish it from the mixed-effects and GEE Wald statistics derived from the same data.

3.2.2 Equivalence to a within-subject t -test

[25]

The strict RM-ANOVA F -test for the biomarker-treatment interaction is equivalent to a two-sample t -test on within-subject treatment differences.

- The within-subject treatment difference Δ_{il} cancels the random subject effect, leaving a pure within-subject contrast.
- $\hat{L} = \bar{\Delta}_1 - \bar{\Delta}_2$ is the difference of stratum-mean treatment differences.
- The equivalence $t_{AB}^2 = F_{AB}$ is well-known in the split-plot literature and makes the test’s assumptions unusually transparent.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

We note that equation (8) is equivalent to a two-sample t -test on the within-subject differences. Define $\Delta_{il} = \bar{Y}_{i1l} - \bar{Y}_{i2l}$ as the within-subject treatment difference for subject l in biomarker stratum i , averaged over T within-phase timepoints. Then

$$\hat{L} = \bar{\Delta}_1 - \bar{\Delta}_2, \quad \bar{\Delta}_i = \frac{1}{n} \sum_{l=1}^n \Delta_{il}. \quad (9)$$

Under sphericity, Δ_{il} are i.i.d. within stratum with $\text{Var}(\Delta_{il}) = 2\sigma_e^2/T$, since the random subject effect cancels in the within-subject difference. The variance of $\hat{L}$ is therefore

$$\text{Var}(\hat{L}) = \frac{2 \cdot 2\sigma_e^2}{nT} = \frac{4\sigma_e^2}{nT}, \quad (10)$$

and the t -statistic for the contrast is

$$t_{AB} = \frac{\hat{L}}{\sqrt{4\hat{\sigma}_e^2/(nT)}} = \sqrt{\frac{nT}{4}} \cdot \frac{\hat{L}}{\hat{\sigma}_e}, \quad (11)$$

[26]

$t_{AB}^2 = F_{AB}$ recovers the interaction F -statistic, confirming that the strict RM-ANOVA test is a two-sample comparison of within-subject treatment effects.

- The identity holds under sphericity because the random subject effect cancels in within-subject differences.
- The test reduces to asking whether the within-subject treatment effect differs between biomarker strata.
- Transparency about this reduction is analytically useful: anything that biases within-subject differences distorts the test in a predictable direction.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

We see that $t_{AB}^2 = F_{AB}$, recovering equation (8). The strict-RM-ANOVA F -test for the biomarker-treatment interaction is therefore identical to a two-sample t -test on the within-subject treatment differences, with degrees of freedom $2(n-1)$. This identity is well-known in the split-plot literature³⁵ and is the analytical anchor for the statement that the strict-RM-ANOVA test of a biomarker-treatment interaction reduces to a simple two-sample comparison of treatment-effect estimates between biomarker strata.

[27]

The transparent reduction to a two-sample t -test will be used in Section 5 to interpret the simulation results.

- The reduction is exploited in Section 5 to connect strict RM-ANOVA results to the mixed-effects findings.
- Biases and inflations in within-subject differences produce predictable distortions in the test's operating characteristics.
- The intuitive appeal of the reduction makes the test's limitations equally transparent.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

That is, under sphericity and balance, the biomarker-treatment inter-

action F -test asks simply whether the within-subject treatment effect differs between biomarker strata, and the answer is given by a two-sample t -test on those within-subject differences. This reduction has intuitive appeal and makes the test's assumptions unusually transparent: anything that biases or inflates the within-subject differences will distort the test in a predictable direction. We shall exploit this transparency when relating the strict-RM-ANOVA results to the mixed-effects results in §5.

[28]

The general $a \times b$ case and the time-resolved three-way interaction extend equation 6 with degrees of freedom scaled by $(a - 1)(b - 1)$.

- For arbitrary a and b , the interaction F -statistic uses $df_1 = (a - 1)(b - 1)$ and $df_2 = a(n - 1)(b - 1)$.
- The three-way interaction $(\alpha\beta\tau)_{ijk}$ tests whether the biomarker modifies the treatment effect across time.
- The corresponding degrees of freedom are $(a - 1)(b - 1)(T - 1)$ and $a(n - 1)(b - 1)(T - 1)$.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

For a biomarker strata, b treatment phases, n subjects per stratum, and T within-phase timepoints, equation (6) generalizes directly with degrees of freedom $df_1 = (a - 1)(b - 1)$ and $df_2 = a(n - 1)(b - 1)$. Including the time factor τ_k explicitly, the test of whether the biomarker modifies the treatment effect across time is the three-way interaction $(\alpha\beta\tau)_{ijk}$; the corresponding F -statistic is $F_{AB\tau} = MS_{AB\tau} / MS_{B\tau \times S(A)}$ with $df_1 = (a - 1)(b - 1)(T - 1)$ and $df_2 = a(n - 1)(b - 1)(T - 1)$.

3.3 Relation to the linear-mixed Wald test

[29]

The strict RM-ANOVA F and the linear-mixed Wald t are asymptotically equivalent only under three conditions that the pmsimstats DGP does not satisfy simultaneously.

- Equivalence requires dichotomized biomarker, compound-symmetric within-subject covariance, and binary Dbc.
- When any condition fails, the mixed-model test extracts more information via continuous biomarker, AR(1) modeling, or continuous exposure-decay.
- In the pmsimstats setting, all three conditions fail, so the mixed-model test dominates strictly.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

We shall now relate the strict-RM-ANOVA F -statistic of equation (8) to the `pmsimstats nlme::lme` Wald t -statistic on the `bm:Dbc` coefficient. The two are asymptotically equivalent under three conditions: (i) the biomarker is dichotomized at the cut-point that defines the strata, (ii) the within-subject covariance is exactly compound-symmetric, and (iii) `Dbc` is binary, the A1 specification of the carryover-sensitivity manuscript rather than the matched-decay A2 specification. When all three conditions hold, the two tests reach the same coefficient and the same standard error up to small-sample correction factors. When any one fails, the mixed-model test extracts more information than the strict RM-ANOVA: the continuous biomarker contributes within-stratum variation, the AR(1) covariance is modeled directly, and the continuous exposure-decay predictor draws on the off-drug trajectory.

[30]

Sphericity violations inflate Type I error; Greenhouse-Geisser and Huynh-Feldt corrections recover nominal control by penalising degrees of freedom.

- Mauchly’s test detects compound-symmetry failure but is not informative about direction; Box’s ϵ is a scalar summary with $\epsilon = 1$ at perfect sphericity.
- The Greenhouse-Geisser correction adjusts both numerator and denominator degrees of freedom by $\hat{\epsilon}$; Huynh-Feldt is less conservative and is recommended when $\hat{\epsilon} > 0.75$.
- The conventional rule recommends Greenhouse-Geisser when $\hat{\epsilon} < 0.75$ to avoid substantial Type I error inflation.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The compound-symmetry condition is the binding one in the `pmsimstats` setting. Equation (8) presupposes sphericity: the within-subject covariance matrix is compound-symmetric. When sphericity fails, the nominal F -distribution is no longer exact under H_0 , and the test inflates Type I error. Two diagnostics and two corrections are standard. Mauchly’s test³⁷ is a likelihood-ratio test for compound symmetry against an unstructured alternative; it often rejects in moderate samples but is not informative about the direction of departure. Box’s ϵ , with $\epsilon = 1$ at compound symmetry and $\epsilon \geq 1/(T - 1)$ at maximum departure, is a scalar measure of departure from sphericity. The Greenhouse-Geisser correction⁵ estimates $\hat{\epsilon}$ from the observed within-subject covariance and adjusts the degrees of freedom: $F_{AB} \sim F(\hat{\epsilon} \cdot df_1, \hat{\epsilon} \cdot df_2)$ under H_0 . The Huynh-Feldt correction⁴ is a

less-conservative variant recommended when $\hat{\epsilon} > 0.75$. The contemporary sphericity-violation literature^{3,6} documents substantial Type I error inflation when sphericity violations go uncorrected, with the conventional rule favouring Greenhouse-Geisser when $\hat{\epsilon} < 0.75$.

[31]

Under the pmsimstats AR(1) DGP, a mixed model with no residual-correlation structure inflates Type I error to 13–17%; corCAR1 restores nominal control, and the pre-averaged strict RM-ANOVA arm is sphericity-robust.

- AR(1) within-factor correlation is not compound-symmetric; Box’s ϵ is strictly below 1.
- Greenhouse-Geisser correction recovers nominal Type I error but substantially attenuates power by penalising degrees of freedom.
- The corCAR1 mixed-model analysis models AR(1) structure directly and does not pay the degrees-of-freedom penalty.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The pmsimstats DGP imposes AR(1) within-factor correlation, which is not compound-symmetric. Under AR(1), Box’s ϵ is strictly below 1, and treating the full multi-timepoint series as repeated within-subject levels under classical RM-ANOVA inflates Type I error unless a degrees-of-freedom correction is applied. The relevant magnitude is recorded in the project’s own audit (docs/06-ar1-residual-correlation.tex): at $T = 8$ timepoints, fitting a mixed model with no residual-correlation structure (the naive analogue of treating AR(1) data as exchangeable) produces Type I error of 13 to 17 percent for the crossover (CO) and open-label (OL) designs, which switching to nlme::lme with corCAR1 returns to nominal. We stress that this 13-to-17-percent figure is for the unstructured/naive fit, not for the strict RM-ANOVA arm evaluated in this paper: that arm pre-averages each phase to a single value (the 2×2 form of §2.2), so its within-subject covariance is 2×2 , sphericity holds by construction, and its Type I error is at nominal (§5). The sphericity penalty is therefore a property of the un-pre-averaged classical analysis; for the pre-averaged strict RM-ANOVA the operative cost is dichotomization, not sphericity. Either way the linear-mixed analysis with corCAR1 avoids both costs, modelling the AR(1) structure directly rather than absorbing it via a degrees-of-freedom penalty.

[32]

The strict RM-ANOVA test is a conservative reduction rather than a competing alternative; option (c) – reporting

both with explicit power-loss tabulation – is operationalized in the simulation study.

- Three defensible responses to a reviewer requesting RM-ANOVA results are identified, ranging from reframing the LME analysis to reporting both with power-loss tables.
- Option (c), reporting both procedures with explicit power-loss tabulation, is the strategy adopted in Sections 4–5.
- This strategy is the most transparent and the most informative for the reviewer.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

For the founding pmsimstats interaction question², the strict-RM-ANOVA test is therefore a conservative reduction of the available analysis rather than a competing alternative. Where a reviewer requests ‘RM-ANOVA results’, three responses are defensible: (a) point at the existing linear-mixed analysis as a continuous-time variant of the RM-ANOVA family, (b) report the strict-RM-ANOVA test from equation (8) as a sensitivity comparator under dichotomized biomarker and binary phase, or (c) report both with explicit power-loss tabulation. The simulation study in §4–§5 operationalizes option (c).

3.3.1 Power calculation under the closed form

[33]

The non-centrality parameter $\lambda = nT\delta^2/(4\sigma_e^2)$ provides the analytic power anchor for the simulation study.

- Under the alternative $H_1 : \hat{L} = \delta$, the F -statistic follows a non-central F distribution.
- The non-centrality parameter scales linearly with n and T and quadratically with δ .
- The formula inverts to a sample-size formula $n = 4\sigma_e^2\lambda/(T\delta^2)$ for power planning.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

For the 2×2 case, equation (8) yields a closed-form power formula that we shall use as the analytic anchor for the simulation study in §4–§5. Under the alternative $H_1 : \hat{L} = \delta$, the statistic $F_{AB} = nT\hat{L}^2/(4\hat{\sigma}_e^2)$ follows a non-central F distribution with non-centrality parameter

$$\lambda = \frac{nT\delta^2}{4\sigma_e^2}, \quad (12)$$

and power at level α is

$$1 - \beta = \Pr \left[F_{AB} > F_{1,2(n-1),1-\alpha} \mid \lambda \right], \quad (13)$$

[34]

All five assumptions of the closed-form power formula are violated to varying degrees in the pmsimstats DGP, so the formula provides an upper bound rather than a calibrated prediction.

- The five assumptions are: sphericity, categorical biomarker with equal stratum sizes, no carryover, no within-treatment relapse, and i.i.d. within-phase observations.
- Violations of sphericity and i.i.d. within-phase observations are two faces of the same AR(1) departure from compound symmetry.
- The analytic power formula is best understood as an upper bound on power available from strict RM-ANOVA under the pmsimstats DGP.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

computable via standard non-central F functions (`stats::pf` in R with the `ncp` argument; `PROC POWER` in SAS). For sample-size calculation, equation (12) inverts to $n = 4\sigma_e^2\lambda/(T\delta^2)$, with λ chosen to deliver the target power at the design α . The formula assumes (i) sphericity, (ii) categorical biomarker at the dichotomized cut-point with equal stratum sizes, (iii) no carryover, (iv) no within-treatment relapse, and (v) i.i.d.

within-phase observations after within-subject differencing. All five assumptions are violated to varying degrees in the pmsimstats DGP, as we shall develop in the next subsection.

[35]

The five assumptions are not independent: sphericity and i.i.d. within-phase observations are two faces of the same AR(1) departure from compound symmetry.

- Under AR(1), compound-symmetry failure drives both Type I error inflation and within-phase autocorrelation simultaneously.
- The analytic power formula is therefore a joint upper bound, not a calibrated predictor for either failure mode in isolation.
- This dependence motivates the Monte Carlo validation in Sections 4–5 rather than reliance on the closed form alone.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

We note that these five assumptions are not independent: violations of sphericity (i) and of the i.i.d. within-phase assumption (v) are, in

the AR(1) setting, two faces of the same departure from compound symmetry. The analytic power formula should therefore be understood as an upper bound on the power available from the strict-RM-ANOVA test under the pmsimstats DGP, not a calibrated prediction.

3.3.2 Limitations of the strict classical formulation

[36]

The five assumptions of the strict classical formulation are empirically problematic in the pmsimstats setting; their severity is presented in order.

- Closed-form simplicity is bought at the price of categorical biomarker, sphericity, no carryover model, balanced design, and equal time spacing.
- Each assumption is violated by the pmsimstats DGP to a degree that favors the linear-mixed analysis.
- The most severe violations – categorical biomarker and no carryover – together account for a 30–50% attenuation of the interaction estimator.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Unfortunately, the closed-form simplicity of equation (8) is bought at the price of five strong assumptions, each of which is empirically problematic in the pmsimstats setting. We present them in approximate order of severity.

1. *Categorical biomarker.* The pmsimstats biomarker (resting SBP in the prazosin-PTSD application) is continuous and has no defensible cut-point. Dichotomizing at the median or at a clinical threshold loses the within-stratum biomarker variation, attenuates the interaction estimator, and inflates Type II error. In the carryover-sensitivity production data the median split gives a $\hat{L}$ standard error roughly 1.4 to 1.7 times that of the continuous-biomarker linear-mixed coefficient, translating to a 30 to 50 percent power penalty.
2. *Sphericity.* Compound-symmetric within-subject covariance is empirically rejected for time-series symptom data: nearby timepoints are more correlated than distant ones. The AR(1) structure pmsimstats implements explicitly violates sphericity. After Greenhouse-Geisser correction, the test recovers nominal Type I error but loses degrees of freedom and power.
3. *No carryover model.* Strict RM-ANOVA has no provision for residual drug effect during the off-drug period; the model treats off-drug observations as if no drug had been administered. Under the pmsimstats DGP with $t_{1/2} = 1$ week, this misspecification biases $\hat{L}$ toward zero (attenuation by 30

to 35 percent in the production data, comparable to the A1 specification evaluated in the carryover-sensitivity manuscript).

4. *Balanced design required.* Equations (6) and (8) assume equal stratum sizes and equal numbers of within-phase timepoints. Real trials have missing data; classical RM-ANOVA handles missingness only through complete-case deletion. The MMRM framework’s likelihood-based handling of missing-at-random (MAR) data³⁶ is genuinely better-behaved.
5. *Equal time spacing.* Strict RM-ANOVA treats time as a categorical factor with implicitly equal-spaced levels. Real N-of-1 trials have unequally-spaced visits; pmsimstats implements continuous-time AR(1) (`corCAR1`), which the strict-RM-ANOVA framework cannot accommodate.

[37]

The five limitations motivate the primary-and-sensitivity reporting strategy; the categorical-biomarker and no-carryover violations are the most consequential.

- The categorical-biomarker and no-carryover violations together account for a 30–50% attenuation of the interaction estimator relative to the linear-mixed analysis.
- The linear-mixed analysis with `corCAR1` is retained as primary; the strict RM-ANOVA result is reported as a sensitivity comparator.
- The Section 6 recommendation operationalizes this strategy for typical chronic-condition N-of-1 scenarios.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

These limitations motivate retaining the linear-mixed analysis with continuous-time AR(1) residual correlation as primary, and reporting the strict-RM-ANOVA closed-form F -statistic as a sensitivity comparator. The five violations are not equally consequential: the categorical-biomarker and no-carryover assumptions are the most problematic in the pmsimstats setting, together accounting for a 30 to 50 percent attenuation of the interaction estimator relative to the linear-mixed analysis. The recommendation operationalized in §6 follows this strategy.

4 Cycle-by-period design grid

[38]

The cycle-by-period grid is specified here but its empirical sweep is deferred; the Section 6 recommendation is correspondingly narrowed to the test-procedure axis.

- The design grid varies cycle count, on-drug and off-drug period lengths, and on/off symmetry around the Hendrickson hybrid reference.
- The present paper’s scope is narrowed to the test-procedure axis at the prazosin/PTSD reference design.
- The joint (cycle count, period length, test procedure) optimum is deferred to the follow-up paper and recorded as a portfolio item.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The cycle-by-period design grid that varies the number of treatment cycles, on-drug and off-drug period lengths, and on/off symmetry around the² hybrid reference is specified in docs/27-cycle-period-design-sweep-plan.md and constitutes the second axis of the joint design-by-test sensitivity question framed in §1. The empirical sweep across this grid is deferred to a follow-up paper. We shall note here, in the interest of transparency, that the present paper narrows its scope to the test-procedure axis at the prazosin/PTSD reference design (Hendrickson hybrid open-label plus blinded discontinuation, OL+BDC; $N \in \{25, 35, 70, 100\}$, $c_{bm} \in \{0, 0.45\}$), which the simulation infrastructure does support at production rep counts. The §6 recommendation is correspondingly scoped to test-procedure choice rather than the joint (cycle count, period length, test procedure) optimum that the design-by-test sensitivity question requires. The follow-up paper executing the cycle-by-period sweep is the natural extension and is recorded as a portfolio item.

5 Simulation design

[39]

The ADEMP framework structures the simulation program around the biomarker-treatment interaction estimand under five test procedures.

- The five procedures are strict RM-ANOVA (dichotomised biomarker), a continuous-biomarker within-subject contrast, linear-mixed with corCAR1, GEE with naive sandwich, and GEE with Mancl-DeRouen small-sample correction.
- The continuous-biomarker contrast separates the cost of dichotomisation from the cost of the test procedure proper; the four continuous-biomarker arms share a common interaction estimand.
- Primary outputs are power and Type I error at $\alpha = 0.05$; secondary estimands (bias, empirical and model SE, 95% CI cover-

age) are reported for the four coefficient-bearing arms.

- All estimates are reported with MCSE following the Morris et al. (2019) discipline.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The simulation program follows the ADEMP framework of¹. The primary estimand is the biomarker-treatment interaction inference under each of the five test procedures. Four operate on the continuous biomarker and share the same interaction estimand $\beta_{bm:D}$ on the symptom scale: a within-subject treatment-effect contrast regressed on the continuous biomarker (the dichotomisation-free analogue of the two-sample t -test of equation (11)); the linear-mixed-effects Wald t on $\hat{\beta}_{bm:D}$; the GEE Wald z on $\hat{\beta}_{bm:D}$ with naive sandwich variance; and the GEE Wald z with Mancl-DeRouen small-sample sandwich variance per³⁴. The fifth, the strict classical RM-ANOVA F -statistic, operates on the median-dichotomised biomarker and so targets a dichotomised estimand on a different scale. The continuous-biomarker contrast is included specifically to separate the cost of dichotomisation (strict RM-ANOVA versus the contrast, both pre-averaged to one value per phase) from the cost of the test procedure proper (the contrast versus the longitudinal linear-mixed and GEE analyses). Primary inferential outputs are power $\pi = \Pr(p < 0.05)$ for $H_0 : \beta_{bm:D} = 0$ at $\alpha = 0.05$, and Type I error under $\beta_{bm:D} = 0$. Secondary estimands, reported for the four coefficient-bearing arms, are bias of $\hat{\beta}_{bm:D}$ relative to a high-precision Monte-Carlo reference value of the estimand, the empirical and the mean model-based standard error, and nominal-95% CI coverage. Strict RM-ANOVA is excluded from these coefficient-based measures because its estimand is on the dichotomised scale.

[40]

MCSE for power and Type I error follow the binomial-proportion form and are reported alongside every estimate.

- Power and Type I error MCSE use $\sqrt{\hat{\pi}(1 - \hat{\pi})/n_{\text{reps}}}$.
- Bias and empirical SE of $\hat{\beta}_{bm:D}$ use the standard within-replicate mean estimator.
- Reporting MCSE alongside every estimate is the Morris et al. discipline adopted project-wide.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Performance measures are reported with MCSE alongside every estimate, following Morris et al.'s recommended discipline. Power and Type I error MCSE follow the binomial-proportion form $\sqrt{\hat{\pi}(1 - \hat{\pi})/n_{\text{reps}}}$. Bias and empirical SE of $\hat{\beta}_{bm:D}$ are reported with

the standard MCSE of the within-replicate mean estimator.

[41]

The full ADEMP pre-registration fixes three studies before production runs commit, providing a pre-specified analytic plan.

- Study 1 is the test-procedure comparison at fixed design (27 cells); Studies 2 and 3 cover the cycle-by-period grid and joint heatmap (240 cells each).
- Pre-registration at 00-ademp-pre-registration.md locks the factorial before production runs commit.
- The present paper reports only Study 1; Studies 2 and 3 are deferred to the companion paper.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The full Aims, Data-generating mechanisms, Estimands, Methods, and Performance measures are pre-registered at analysis/scripts/test-procedure-design-sensitivity which fixes the test-procedure times design-grid factorial across three studies (Study 1: test-procedure comparison at fixed design; Study 2: cycle-by-period grid at fixed test, 240 cells; Study 3: joint design-by-test heatmap, 240 cells) before the production runs commit. The present paper executes a revised Study 1; the executed design departs from the pre-registered Study 1 on several factor levels, and we document every departure in Table 1 below in the interest of full ADEMP transparency.

Table 1: Pre-registered versus executed Study 1. Departures are recorded for transparency; the executed superset of sample sizes (which restores the pre-registered N levels and adds $N = 25$) and the added continuous-biomarker contrast arm respond to referee comments on the first internal draft.

Factor	Pre-registered	Executed (this paper)
Test procedures	3 (RM-ANOVA, LME, GEE-MD)	5 (adds naive GEE and the continuous-biomarker contrast)
Biomarker effect c_{bm}	$\{0, 0.3, 0.6\}$	$\{0, 0.45\}$ (the project-wide reference effect size)
Sample size N	$\{35, 70, 100\}$	$\{25, 35, 70, 100\}$ (superset; adds $N = 25$)
Replicates per cell	1000 (5000 for Type I cells)	1000 power, 2000 Type I
Carryover half-life	$t_{1/2} = 3$ days	$t_{1/2} = 0$ (no carryover at the fixed reference)

[42]

The factorial crosses five procedures by two c_{bm} values by four N values; the inline Mancl-DeRouen SE is benchmarked against `geesmv::GEE.var.md`.

- The simulation crosses five procedures by $c_{bm} \in \{0, 0.45\}$ by $N \in \{25, 35, 70, 100\}$, with 1000 replicates per power cell and 2000 per Type I cell.
- The Mancl-DeRouen correction inflates each cluster's residual via $(I - H_{ii})^{-1}r_i$ in the sandwich meat.
- The inline SE agrees with `geesmv::GEE.var.md` to within roughly 2.5% (mean ratio 1.02 over 25 probe replicates), so the correction magnitude is externally validated.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In the scope of the present paper (test-procedure axis only at the reference design), the simulation runs on the `pmsimstats` package data-generating process under Architecture A (mean moderation), exercised through `generateData()` and the standard `lme_analysis()` fit, rather than on the reduced `implementations/simple` substrate. The factorial comprises $n_{\text{reps}} = 1000$ per power cell and 2000 per Type I cell, across the 40 cells defined by procedure $\in \{\text{strict RM-ANOVA, continuous-biomarker contrast, linear-mixed with corCAR1, GEE with naive sandwich, GEE with Mancl-DeRouen small-sample correction}\} \times c_{bm} \in \{0, 0.45\} \times N \in \{25, 35, 70, 100\}$. Execution uses `parallel::mclapply` with `RNGkind("L'Ecuyer-CMRG")` and `mc.set.seed = TRUE`, so each replicate draws from an independent, reproducible random-number stream. The Mancl-DeRouen correction was implemented inline (no `geesmv`-style external dependency in the production path) by inflating each cluster's residual via $(I - H_{ii})^{-1}r_i$ in the sandwich meat. It produces identical point estimates and larger standard errors than the naive sandwich, and its SE for the interaction coefficient agrees with `geesmv::GEE.var.md` to within approximately 2.5% (mean inline/`geesmv` ratio 1.02, standard deviation 0.02, over 25 probe replicates at $N = 35$, $c_{bm} = 0.45$), so the correction magnitude is externally validated rather than merely plausible.

6 Results

[43]

The 40-cell factorial completed in 587 s with 100% convergence across 60,000 fits; MCSE is approximately 0.016 for power and 0.005 for Type I error.

- Execution used 7-core `parallel::mclapply` with L'Ecuyer-CMRG streams; convergence was 100% in every cell.
- Per-replicate output and the ADEMP summary are archived at `analysis/data/quick-sim/`.
- MCSE on power near 0.5 is 0.016 (1000 reps); MCSE on Type I error near 0.05 is 0.005 (2000 reps), providing adequate resolution to distinguish the procedures.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The simulation program was executed across the 40 cells of the five-procedure-by-effect-by- N factorial in 587 s using 7-core `parallel::mclapply` with `RNGkind("L'Ecuyer-CMRG")` per-stream seeding. Convergence was 100% across all 60{,}000 fits. Per-replicate output is at `analysis/data/quick-sim/08-test-replicates.rds`; the Morris ADEMP cell-level summary is at `analysis/data/quick-sim/08-test-summary.txt`. With 1000 replicates per power cell the MCSE on a power estimate near 0.5 is approximately 0.016, and with 2000 replicates per Type I cell the MCSE on a rate near 0.05 is approximately 0.005. We begin the results with the test-procedure comparison at the fixed reference design.

6.1 Test-procedure comparison at fixed design

[44]

Strict RM-ANOVA targets a dichotomized estimand; the continuous-biomarker contrast isolates the cost of dichotomization from the cost of the test procedure proper.

- Strict RM-ANOVA operates on a median-dichotomized biomarker, pre-averaged to one value per phase; its estimand is on the dichotomized scale.
- The continuous-biomarker contrast uses the same pre-averaging but the continuous biomarker, so it shares the estimand of the linear-mixed and GEE arms.
- The three-rung ladder (RM-ANOVA, contrast, linear-mixed) decomposes the gap between classical RM-ANOVA and the mixed model into a dichotomization component and a test-procedure component.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

We note at the outset that strict RM-ANOVA targets a different estimand from the other four procedures. It operates on a median-dichotomized biomarker, pre-averaged to one value per phase, and so estimates a dichotomized-stratum interaction; the continuous-

biomarker contrast, the linear-mixed model, and both GEE variants estimate the interaction coefficient $\beta_{bm:D}$ on the continuous symptom scale. An earlier version of this comparison contrasted strict RM-ANOVA directly with the mixed model and attributed the resulting power gap to the test procedure; that attribution conflated two distinct costs. We therefore introduce the continuous-biomarker contrast of §4, which differs from strict RM-ANOVA only in retaining the continuous biomarker (it shares the pre-averaging and the absence of any AR(1) machinery) and differs from the linear-mixed model only in discarding the within-phase longitudinal structure. The RM-ANOVA $\rightarrow$ contrast step therefore measures the cost of dichotomization in isolation, and the contrast $\rightarrow$ linear-mixed step measures the cost of the test procedure proper. Reading the two steps separately is the honest decomposition the earlier direct comparison lacked.

Type I error at the null cell ($c_{bm} = 0$), percent $\pm$ MCSE, is shown in Table 2; power at $c_{bm} = 0.45$ in Table 3.

Table 2: Type I error rate (percent $\pm$ MCSE) under $H_0 : \beta_{bm:D} = 0$, 2000 replicates per cell. Nominal level 5%.

Procedure	$N = 25$	$N = 35$	$N = 70$	$N = 100$
Strict RM-ANOVA (dich.)	4.55 $\pm$ 0.47	4.15 $\pm$ 0.45	5.95 $\pm$ 0.53	5.20 $\pm$ 0.50
Contrast (continuous bm)	5.25 $\pm$ 0.50	4.50 $\pm$ 0.46	6.45 $\pm$ 0.55	4.60 $\pm$ 0.47
Linear-mixed (corCAR1)	2.75 $\pm$ 0.37	2.05 $\pm$ 0.32	3.35 $\pm$ 0.40	2.90 $\pm$ 0.38
GEE (naive sandwich)	9.70 $\pm$ 0.66	8.55 $\pm$ 0.63	7.90 $\pm$ 0.60	7.40 $\pm$ 0.59
GEE (Mancl-DeRouen)	6.35 $\pm$ 0.55	5.30 $\pm$ 0.50	6.25 $\pm$ 0.54	5.90 $\pm$ 0.53

Table 3: Power (percent $\pm$ MCSE) at $c_{bm} = 0.45$, 1000 replicates per cell.

Procedure	$N = 25$	$N = 35$	$N = 70$	$N = 100$
Strict RM-ANOVA (dich.)	35.8 $\pm$ 1.52	46.8 $\pm$ 1.58	75.3 $\pm$ 1.36	91.3 $\pm$ 0.89
Contrast (continuous bm)	53.8 $\pm$ 1.58	67.7 $\pm$ 1.48	91.5 $\pm$ 0.88	98.8 $\pm$ 0.34
Linear-mixed (corCAR1)	56.5 $\pm$ 1.57	71.2 $\pm$ 1.43	94.5 $\pm$ 0.72	99.2 $\pm$ 0.28

Procedure	$N = 25$	$N = 35$	$N = 70$	$N = 100$
GEE (naive sandwich)	67.2 ± 1.48	75.9 ± 1.35	94.5 ± 0.72	99.2 ± 0.28
GEE (Mancl-DeRouen)	57.3 ± 1.56	69.1 ± 1.46	92.7 ± 0.82	99.1 ± 0.30

[45]

Most of the gap between strict RM-ANOVA and the mixed model is the cost of dichotomization, not of the test procedure: the continuous-biomarker contrast recovers four-fifths of it.

- At $N = 25$, replacing the dichotomized biomarker with the continuous one (RM-ANOVA $\rightarrow$ contrast) raises power from 0.358 to 0.538, a gain of 0.180.
- Moving from the pre-averaged contrast to the full linear-mixed model adds only 0.027 (0.538 $\rightarrow$ 0.565); the test-procedure component is small.
- The same pattern holds across N : dichotomization costs 0.16–0.21 of power, while the longitudinal test procedure adds at most 0.04.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Four substantive findings emerge from the tables. First, and most consequentially for how the comparison should be read, the power gap between strict RM-ANOVA and the linear-mixed model is predominantly a cost of dichotomizing the biomarker rather than a property of the test procedure. At $N = 25$ the strict RM-ANOVA power is 0.358; the continuous-biomarker contrast, which differs only in retaining the continuous biomarker, reaches 0.538, so dichotomization alone accounts for 0.180 of lost power. The further step to the full linear-mixed model raises power only to 0.565, an increment of 0.027 attributable to the test procedure proper. The decomposition is stable across the grid: the dichotomization component is 0.209 at $N = 35$ (0.468 $\rightarrow$ 0.677), 0.162 at $N = 70$ (0.753 $\rightarrow$ 0.915), and 0.075 at $N = 100$ (0.913 $\rightarrow$ 0.988), while the contrast-to-linear-mixed increment never exceeds 0.04. The earlier framing, which compared strict RM-ANOVA against the mixed model directly and read the difference as a test-procedure effect, therefore overstated the role of the test procedure by roughly a factor of five; the honest statement is that a classical, sphericity-free, AR(1)-free within-subject contrast on the continuous biomarker captures most of the available power, and the mixed model adds a small further increment.

[46]

Naive GEE inflates Type I error across the whole sample-size range, not only at small N ; the Mancl-DeRouen correction restores nominal control at a moderate power cost.

- Naive GEE rejects under the null at 0.097, 0.086, 0.079, and 0.074 for $N = 25, 35, 70, 100$ – between $1.5\times$ and $1.9\times$ nominal throughout.
- The Mancl-DeRouen correction brings the rate to 0.064, 0.053, 0.063, and 0.059, within or close to MCSE of the nominal 0.05 at every N .
- The correction’s power cost is moderate (about 0.10 at $N = 25$, 0.07 at $N = 35$) and the corrected procedure no longer over-rejects.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Second, GEE with the naive sandwich is anti-conservative across the entire sample-size range examined. Its null rejection rate is 0.097 at $N = 25$ and declines only slowly with N to 0.086, 0.079, and 0.074 at $N = 35, 70, 100$ – still roughly $1.5\times$ nominal at $N = 100$. The inflation is therefore not a small-cluster curiosity that disappears at trial-relevant sizes; it persists. The Mancl-DeRouen 2001 small-sample sandwich correction brings the rate within or close to MCSE of nominal at every N (0.064, 0.053, 0.063, 0.059), at a power cost of about ten percentage points at $N = 25$ ($0.672 \rightarrow 0.573$) and seven at $N = 35$ ($0.759 \rightarrow 0.691$). The corrected procedure’s apparent power advantage over the linear-mixed model at small N (Table 3) is the residue of its remaining slight over-rejection, not a genuine efficiency gain.

[47]

Standard-error calibration separates the four coefficient-bearing arms: the contrast and Mancl-DeRouen GEE are well-calibrated, naive GEE under-covers, and the linear-mixed model over-covers.

- The continuous-biomarker contrast and Mancl-DeRouen GEE attain 93–95% coverage with model/empirical SE ratios near 1.
- Naive GEE under-covers at small N (90.0% at $N = 25$) because its model SE is 9–12% too small.
- The linear-mixed model over-covers (97.0–97.7%) because its model SE is 11–16% larger than the empirical SE, the calibration signature of its conservatism.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Third, the secondary estimands (Table 4) locate the procedures' behaviour in the standard-error calibration that the rejection rates only summarize. At $c_{bm} = 0.45$, the continuous-biomarker contrast and the Mancl-DeRouen GEE are well-calibrated: their nominal-95% CI coverage is 93–95% and their mean model-based SE is within a few percent of the empirical SE. Naive GEE under-covers at small N (90.0% at $N = 25$, 91.5% at $N = 35$) because its model-based SE is 9 to 12% smaller than the empirical SE, the calibration counterpart of its Type I inflation. The linear-mixed model over-covers (97.0 to 97.7%) because its model-based SE exceeds the empirical SE by 11 to 16%; this is the same conservatism visible in its sub-nominal Type I error.

Table 4: Secondary estimands at $c_{bm} = 0.45$ for the four coefficient-bearing arms (the interaction estimand is $\beta_{bm:D} = -2.21$, the Monte-Carlo reference value). Coverage is of the nominal-95% Wald interval; SE ratio is mean model-based SE divided by empirical SE.

Arm	Coverage $N = 25$	Coverage $N = 100$	SE ratio $N = 25$	SE ratio $N = 100$	Max rel. bias
Contrast (continuous bm)	94.9	94.8	0.99	0.98	3.0%
Linear-mixed (corCAR1)	97.7	97.0	1.16	1.11	2.6%
GEE (naive sandwich)	90.0	93.3	0.91	0.98	3.8%
GEE (Mancl- DeRouen)	94.3	94.0	1.07	1.02	3.8%

[48]

All four coefficient-bearing arms are approximately unbiased; the linear-mixed conservatism is a variance-calibration effect, not a point-estimate effect, and a Kenward-Roger correction would address it.

- Across all cells, the relative bias of the interaction coefficient is below 4% of the estimand for every coefficient-bearing arm.
- The linear-mixed Type I error is 0.021 to 0.034, below nominal, consistent with corCAR1 degrees-of-freedom accounting under $T = 8$ timepoints.
- A Kenward-Roger or Satterthwaite small-sample correction would address the conservatism and is recorded as a follow-up item.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod

tempor incididunt ut labore et dolore magna aliqua.

Fourth, the comparison is not contaminated by point-estimate bias. The relative bias of $\hat{\beta}_{bm:D}$ is below 4% of the estimand $|\beta_{bm:D}| = 2.21$ for the contrast, the linear-mixed model, and both GEE variants in every cell, so the differences in power and coverage reflect variance estimation and information use rather than systematic mis-estimation. The linear-mixed procedure's Type I error sits below nominal throughout (0.021 to 0.034), the conservatism already visible in its over-coverage and its model/empirical SE ratio above one; it is consistent with `corCAR1` degrees-of-freedom accounting under the short ($T = 8$) serial series. A Kenward-Roger³⁸ or Satterthwaite³⁹ small-sample degrees-of-freedom correction would address the conservatism and recover the small power increment it forgoes; this extension is recorded as a follow-up item.

7 Design-and-test recommendation

[49]

The recommendation is presented as a table of preferred (procedure, scenario) combinations at the reference design rather than as a single best procedure.

- Cycle count and period length are held at reference values; joint optimization over those axes is deferred to the follow-up paper.
- The preferred procedure depends on features of the application – sample size, tolerance for Type I inflation, and reviewer requirements – that the analyst must determine.
- The first-order recommendation is to keep the biomarker continuous; the choice among continuous-biomarker procedures is second-order.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The recommendation table below maps test procedure choice to typical chronic-condition N-of-1 trial scenarios at the prazosin/PTSD reference design (Hendrickson hybrid OL+BDC, biomarker-treatment interaction estimand). The cycle count and period length are held at the reference values defined in §3 because the joint optimization over those axes is deferred to the follow-up paper. We present the recommendation as a table of preferred (procedure, scenario) combinations rather than as a single ‘best’ procedure, because the preferred choice depends on features of the application that the analyst must determine. The overriding recommendation, however, is design-agnostic: retain the biomarker as a continuous predictor, since dichotomization forfeits the large majority of the available power (§5, Table 3).

Scenario	Preferred procedure	Rationale
Default (continuous biomarker, any N)	linear-mixed (corCAR1)	Highest power within nominal Type I; corCAR1 matches AR(1) physiology; well-behaved across N
Lightweight / transparent classical analysis	continuous-biomarker contrast	Within 0.03 of the mixed model on power, nominal Type I, no covariance modelling required
$N \leq 35$, concern about LME conservatism	linear-mixed (corCAR1) with Kenward-Roger	LME sub-nominal and over-covers; KR addresses this without inflating Type I
Robust-SE-required setting (e.g. unmodeled clustering)	GEE Mancl-DeRouen	Restores nominal Type I and 93–95% coverage; modest power cost
Reviewer-mandated ‘classical’ analysis	strict RM-ANOVA	Dichotomized sensitivity comparator; report alongside a continuous-biomarker analysis with the power-loss decomposition
Naive GEE (no small-sample correction)	not recommended	Type I error $1.5\text{--}1.9\times$ nominal across all N ; under-covers

[50]

Keep the biomarker continuous; among continuous-biomarker procedures the linear-mixed model is the preferred default, and naive GEE should not be used without a small-sample correction at any N examined.

- The first-order decision is to retain the continuous biomarker; dichotomization forfeits 0.08–0.21 of power across the grid.
- Among continuous-biomarker procedures, linear-mixed with **corCAR1** has the highest power within nominal Type I; the contrast is a close, transparent alternative.
- Naive GEE is anti-conservative at every N examined and should be replaced by the Mancl-DeRouen correction where a sandwich estimator is required.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The dominant message is that the biomarker should be analysed as a continuous predictor: dichotomizing it for a strict RM-ANOVA forfeits the large majority of the power gap that the present study iso-

lates, while modelling choices among the continuous-biomarker procedures move power only at the margin. Within that class, the linear-mixed analysis with continuous-time AR(1) residual correlation is the preferred default for the biomarker-treatment interaction in aggregated N-of-1 trials at the parameter scale evaluated here, with the dichotomization-free within-subject contrast a close and transparent alternative that requires no covariance modelling. Strict RM-ANOVA remains useful as a dichotomized sensitivity comparator, as the §2 closed-form derivation makes explicit, but it should be reported alongside a continuous-biomarker analysis and the power-loss decomposition rather than on its own. Naive GEE should not be used at the sample sizes examined without a small-sample correction; the Mancl-DeRouen correction restores nominal Type I error and 93–95% coverage at a moderate power cost.

8 Discussion

8.1 What the Monte Carlo study settles

[51]

The simulation settles four substantive questions; the inline Mancl-DeRouen SE is validated against `geesmv::GEE.var.md` to within 2.5%.

- The 40-cell factorial completed in 587 s with 100% convergence across 60,000 fits; MCSE is 0.016 for power near 0.5 and 0.005 for Type I error near 0.05.
- The inline Mancl-DeRouen implementation produces identical point estimates and larger standard errors than the naive sandwich.
- Its SE for the interaction coefficient matches `geesmv::GEE.var.md` to within roughly 2.5% over 25 probe replicates, so the correction magnitude is externally validated.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

We begin the discussion by summarizing what the Monte Carlo study settles at the resolution it was run. The run on the 40-cell five-procedure factorial ($\text{procedure} \in \{\text{strict RM-ANOVA, continuous-biomarker contrast, linear-mixed with } \text{corCAR1}, \text{ GEE with naive sandwich, GEE with Mancl-DeRouen small-sample correction}\} \times c_{bm} \in \{0, 0.45\} \times N \in \{25, 35, 70, 100\}$, with 1000 replicates per power cell and 2000 per Type I cell) was executed in 587 s with 100% convergence across all 60{,}000 fits. Per-replicate output is at `analysis/data/quick-sim/08-test-replicates.rds`; Morris ADEMP summary at `08-test-summary.txt`. With 1000

power replicates the MCSE on a power estimate near 0.5 is approximately 0.016 and with 2000 Type I replicates the MCSE on a rate near 0.05 is approximately 0.005. The Mancl-DeRouen 2001³⁴ correction was implemented inline by inflating each cluster’s residual via $(I - H_{ii})^{-1}r_i$ in the sandwich meat; it produces identical point estimates and larger standard errors than the naive sandwich, and its SE for the interaction coefficient agrees with `geesmv::GEE.var.md` to within approximately 2.5% (mean inline/geesmv ratio 1.02 over 25 probe replicates), so the correction magnitude is externally validated.

[52]

The first resolved question is that the classical-versus-mixed power gap is mostly the cost of dichotomization: the continuous-biomarker contrast recovers about four-fifths of it.

- At $N = 25$, dichotomization (RM-ANOVA $\rightarrow$ contrast) costs 0.180 of power; the test procedure (contrast $\rightarrow$ linear-mixed) adds 0.027.
- The decomposition is stable across N : dichotomization costs 0.08–0.21, the test procedure at most 0.04.
- A classical, sphericity-free, AR(1)-free contrast on the continuous biomarker therefore captures most of the available power.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The study resolves four substantive questions at this resolution. First, and correcting the emphasis of an earlier draft, the power gap between strict RM-ANOVA and the linear-mixed model is predominantly the cost of dichotomizing the biomarker, not a property of the test procedure. The continuous-biomarker contrast, which differs from strict RM-ANOVA only in retaining the continuous biomarker, recovers 0.180 of the 0.207 gap at $N = 25$ ($0.358 \rightarrow 0.538$ of $0.358 \rightarrow 0.565$); the further step to the full mixed model contributes only 0.027. The same decomposition holds across the grid (dichotomization 0.08 to 0.21, test procedure at most 0.04). The earlier direct comparison of RM-ANOVA against the mixed model therefore overstated the test-procedure effect by roughly fivefold; the corrected statement is that keeping the biomarker continuous, not adopting the mixed model per se, is what recovers the power.

[53]

The second resolved question is that naive GEE inflates Type I error across the entire sample-size range, and Mancl-DeRouen restores nominal control.

- Naive GEE rejects at 0.097, 0.086, 0.079, 0.074 for $N =$

1450
1451
1452
1453
1454

1455
1456

1457
1458
1459
1460
1461
1462
1463
1464
1465
1466
1467

1468

1469
1470
1471

1472
1473
1474
1475
1476
1477

1478
1479

1480
1481
1482
1483
1484
1485
1486
1487
1488
1489
1490
1491

25, 35, 70, 100 – still about $1.5\times$ nominal at $N = 100$.

- The Mancl-DeRouen correction brings the rate to 0.064, 0.053, 0.063, 0.059, within or near MCSE of nominal at every N .
- The power cost is moderate (about 0.10 at $N = 25$, 0.07 at $N = 35$).

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Second, naive GEE is anti-conservative across the entire sample-size range, not merely at the smallest N . Its null rejection rate is 0.097 at $N = 25$ and falls only to 0.086, 0.079, and 0.074 at $N = 35, 70, 100$, remaining roughly $1.5\times$ nominal even at $N = 100$. The Mancl-DeRouen correction returns the rate to within or near MCSE of nominal at every N (0.064, 0.053, 0.063, 0.059), at a power cost of about ten percentage points at $N = 25$ ($0.672 \rightarrow 0.573$) and seven at $N = 35$ ($0.759 \rightarrow 0.691$). Naive GEE’s nominally higher power than the linear-mixed model at small N is an artefact of its residual over-rejection rather than a real gain, and disappears once the size distortion is corrected.

[54]

The third resolved question is standard-error calibration: the contrast and Mancl-DeRouen GEE are well-calibrated, naive GEE under-covers, and the linear-mixed model over-covers.

- Contrast and Mancl-DeRouen GEE attain 93–95% coverage with model/empirical SE ratios near 1.
- Naive GEE under-covers (90.0% at $N = 25$) because its model SE is 9–12% too small.
- The linear-mixed model over-covers (97.0–97.7%) because its model SE exceeds the empirical SE by 11–16%.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Third, the standard-error calibration of the four coefficient-bearing arms is what underlies the rejection-rate differences. The continuous-biomarker contrast and the Mancl-DeRouen GEE are well-calibrated, with 93 to 95% nominal-95% coverage and model-based standard errors within a few percent of the empirical standard error. Naive GEE under-covers at small N (90.0% at $N = 25$) because its model-based SE is 9 to 12% too small, the same defect that drives its Type I inflation. The linear-mixed model over-covers (97.0 to 97.7%) because its model-based SE exceeds the empirical SE by 11 to 16%, the calibration counterpart of its sub-nominal Type I error. All four arms are approximately unbiased (relative bias below 4%), so these are pure variance- estimation effects.

[55]

The fourth resolved question is that the linear-mixed conservatism is a degrees-of-freedom artefact a Kenward-Roger correction would remove.

- LME Type I error is 0.021 to 0.034, conservative rather than nominal, consistent with `corCAR1` degrees-of-freedom accounting under $T = 8$ timepoints.
- The conservatism shows up as over-coverage and a model/empirical SE ratio above one, not as point-estimate bias.
- A Kenward-Roger or Satterthwaite correction would recover the small power increment the conservatism forgoes.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Fourth, the linear-mixed procedure's Type I error sits below nominal throughout (0.021 to 0.034), conservative rather than nominal and consistent with `corCAR1` degrees-of-freedom accounting under the short ($T = 8$) serial series. The conservatism is a variance-calibration effect, visible in the over-coverage and the model/empirical SE ratio above one of Table 4, not a point-estimate effect. A Kenward-Roger³⁸ or Satterthwaite³⁹ small-sample degrees-of-freedom correction would address it and recover the small power increment it forgoes, and we record this as a follow-up item; the recommendation table in §6 already lists the Kenward-Roger variant for small- N use.

9 Conclusions

[56]

The biomarker should be kept continuous: dichotomization, not the test procedure, drives most of the classical-versus-mixed power gap; naive GEE is disqualified without a small-sample correction.

- Dichotomizing the biomarker for strict RM-ANOVA forfeits 0.08–0.21 of power across N ; the test procedure proper adds at most 0.04.
- Among continuous-biomarker procedures the linear-mixed model has the highest power within nominal Type I; the contrast is a close, transparent alternative.
- Naive GEE inflates Type I error to 1.5–1.9× nominal across all N ; the Mancl-DeRouen correction restores nominal control at a moderate (about 10-percentage-point) power cost.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In conclusion, three points are to be emphasized. First, the analyst's most consequential choice for the biomarker-treatment interaction in aggregated N-of-1 trials is to keep the biomarker continuous rather than to dichotomize it for a classical RM-ANOVA. The continuous-biomarker contrast recovers about four-fifths of the apparent RM-ANOVA-to-mixed-model power gap (dichotomization costs 0.08 to 0.21 of power across the sample sizes evaluated, while the test procedure proper adds at most 0.04); the linear-mixed analysis with continuous-time AR(1) residual correlation has the highest power within nominal Type I among the continuous-biomarker procedures, with the dichotomization-free within-subject contrast a close alternative. Naive GEE with the standard sandwich estimator inflates Type I error to between $1.5\times$ and $1.9\times$ nominal across the whole $N \in \{25, 35, 70, 100\}$ range, disqualifying its rejection rate as a power summary; the Mancl-DeRouen 2001 small-sample sandwich correction restores nominal Type I control and 93–95% coverage at a moderate (about 10-percentage-point) power cost, making it the operational GEE comparator rather than the naive sandwich.

[57]

The closed-form reduction to a two-sample t -test makes the strict RM-ANOVA transparent; in this study its limitation is dichotomization, with sphericity relevant only to the unpre-averaged classical analysis.

- Under sphericity and balance, the biomarker-treatment interaction F -test reduces to a two-sample t -test on within-subject treatment differences.
- The pre-averaged-to-one-value-per-phase form evaluated here is sphericity-robust by construction and holds nominal Type I error, so its power loss is attributable to dichotomization, not to sphericity.
- Sphericity violation under AR(1) bites only when the full multi-timepoint series is analysed by classical RM-ANOVA without a degrees-of-freedom correction; the linear-mixed model with `corCAR1` sidesteps both issues.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Second, the closed-form strict-RM-ANOVA test of the biomarker-treatment interaction (§2) reduces to a two-sample t -test on the within-subject treatment differences under sphericity and balance. The reduction makes the strict-RM-ANOVA result analytically transparent and connects the simulation findings to the classical split-plot literature. It also pins down what does and does not limit the test here. The strict RM-ANOVA arm we evaluate pre-averages

each phase to a single value, so its within-subject covariance is 2×2 and sphericity holds by construction; consistent with this, its Type I error is at nominal (Table 2), and its power deficit is wholly attributable to the median dichotomization of the biomarker, as the continuous-biomarker contrast demonstrates. The sphericity penalty that motivates much of §2 applies instead to the alternative classical analysis that treats the full $T = 8$ series as repeated within-subject levels: there the AR(1) residual correlation violates sphericity and a Greenhouse-Geisser correction recovers nominal Type I only at a degrees-of-freedom cost, which the linear-mixed analysis with `corCAR1` avoids by modelling the AR(1) structure directly. Both considerations point the same way – away from a dichotomized classical analysis – but through different mechanisms, and we are careful to keep them distinct.

[58]

The cycle-by-period design grid is deferred to a follow-up paper; infrastructure and design specification are in place and the scope limitation is explicitly acknowledged.

- The simulation infrastructure for the cycle-by-period sweep is committed at `analysis/scripts/test-procedure-design-sensitivity/`.
- The design specification is at `docs/27-cycle-period-design-sweep-plan.md`.
- The follow-up paper will deliver the joint recommendation table that the present paper’s Section 6 narrows to the test-procedure-only axis.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Third, the cycle-by-period design grid that constitutes the second axis of the joint design-by-test sensitivity question framed in §1 is deferred to a follow-up paper. The infrastructure for the sweep is in place at `analysis/scripts/test-procedure-design-sensitivity/`, and the design specification is committed at `docs/27-cycle-period-design-sweep-plan.md`. The follow-up paper will deliver the joint (cycle count, period length, test procedure) recommendation table that the present paper’s §6 narrows to the test-procedure-only axis.

10 Conflict of interest

The authors declare no conflicts of interest.

11 Funding

This work received no specific funding.

12 Data availability statement

[59]

Data and code are openly available in the pmsimstats-ng repository at paths listed in the Reproducibility section.

- Per-replicate Monte Carlo outputs and ADEMP cell-level summaries are versioned under `analysis/data/`.
- Simulation drivers are at `analysis/scripts/`.
- Exact paths relevant to this manuscript are listed in the Reproducibility section.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The data and code that support the findings of this study are openly available in the pmsimstats-ng project repository. Per-replicate Monte Carlo outputs and ADEMP-compliant cell-level summaries are versioned under `analysis/data/`; simulation drivers are at `analysis/scripts/`. Exact paths relevant to this manuscript are listed in the Reproducibility section below.

13 Reproducibility

This manuscript’s simulation was executed with R 4.6.x; the project also ships a `zzcollab` Docker container (image hash recorded in `Dockerfile`; package set captured in `renv.lock`) for container-based reproduction. Principal packages used by the Study 1 pipeline are `nlme` (continuous-time linear-mixed analysis with `corCAR1` residual correlation), `geepack` (GEE fitting) plus an inline Mancl-DeRouen 2001 small-sample sandwich correction, `geesmv` (independent benchmark of that correction), `data.table`, and `parallel` (`mclapply` for the parameter-grid sweep), with `rmarkdown` plus `bookdown` for the manuscript build. The Study 1 simulation runs on the pmsimstats package data-generating process under Architecture A via `generateData()` and `lme_analysis()`, not on the reduced implementations/simple substrate.

Random-number generator. The driver sets `RNGkind("L'Ecuyer-CMRG")` and `set.seed(20260507)`, then runs the cell-by-replicate jobs through `parallel::mclapply` with `mc.set.seed = TRUE`, so each replicate draws from an independent, reproducible random-number stream.

Data and code availability. Per-replicate simulation outputs are checkpointed at `analysis/data/quick-sim/08-test-replicates.rds`; the extended Morris ADEMP cell-level summary (rejection rate, MCSE, convergence, bias, empirical and model SE, coverage) at the companion `08-test-summary.txt`, with the Monte-Carlo reference estimand at `08-test-betatruet.txt` and the Mancl-DeRouen benchmark at `08-test-mancl-benchmark.txt`. The driver is `analysis/scripts/test-procedure-design-sensitivity/01-study1-test-procedure.R`

(superseding the earlier `analysis/scripts/quick-sim/08-test-procedure-prototype.R`).
 The pre-registered ADEMP plan, and the deviations recorded in Table 1, are at
`analysis/scripts/test-procedure-design-sensitivity/00-ademp-pre-registration.md`.
 The manuscript source, paper-specific bibliography (`references.bib`), and SIM
 CSL file are at `analysis/report/08-test-procedure-design-sensitivity/`.

14 References

1. Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Statistics in Medicine*. 2019;38(11):2074-2102.
2. Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: Statistical methods and practical considerations. *Frontiers in Digital Health*. 2020;2:13. doi:10.3389/fdgth.2020.00013
3. Haverkamp N, Beauducel A. Violation of the sphericity assumption and its effect on type-I error rates in repeated measures ANOVA and multi-level linear models. *Frontiers in Psychology*. 2017;8:1841. doi:10.3389/fpsyg.2017.01841
4. Huynh H, Feldt LS. Estimation of the Box correction for degrees of freedom from sample data in randomized block and split-plot designs. *Journal of Educational Statistics*. 1976;1(1):69-82. doi:10.3102/10769986001001069
5. Greenhouse SW, Geisser S. On methods in the analysis of profile data. *Psychometrika*. 1959;24(2):95-112. doi:10.1007/BF02289823
6. Blanca MJ, Arnau J, García-Castro FJ, Alarcón R, Bono R. Repeated measures ANOVA and adjusted tests when sphericity is violated: Which procedure is best? *Frontiers in Psychology*. 2023;14:1192453. doi:10.3389/fpsyg.2023.1192453
7. Kent DM, Paulus JK, Klaveren D van, et al. The Predictive Approaches to Treatment effect Heterogeneity (PATH) Statement. *Annals of Internal Medicine*. 2020;172(1):35-45. doi:10.7326/M18-3667
8. Kent DM, Klaveren D van, Paulus JK, et al. The PATH statement: Explanation and elaboration. *Annals of Internal Medicine*. 2020;172(1):W1-W25. doi:10.7326/M18-3668
9. Senn S. The AB/BA crossover: Past, present and future? *Statistical Methods in Medical Research*. 1994;3(4):303-324. doi:10.1177/096228029400300402

- 1668 10. Senn S, Rolfe K, Julious SA. Investigating variability in patient response to treatment: A case study from a replicate cross-over study. *Statistical Methods in Medical Research*. 2011;20(6):657-666. doi:10.1177/0962280210379174
- 1669 11. Natesan Batley P, McClure EB, Brewer B, et al. Evidence and reporting standards in N-of-1 medical studies: A systematic review. *Translational Psychiatry*. 2023;13(1):263. doi:10.1038/s41398-023-02562-8
- 1670 12. Stunnenberg BC, Berends J, Griggs RC, et al. N-of-1 trials in neurology: A systematic review. *Neurology*. 2022;98(2):e174-e185. doi:10.1212/WNL.00000000000012998
- 1671 13. Shaffer JA, Kronish IM, Falzon L, Cheung YK, Davidson KW. N-of-1 randomized intervention trials in health psychology: A systematic review and methodology critique. *Annals of Behavioral Medicine*. 2018;52(9):731-742. doi:10.1093/abm/kax026
- 1672 14. Vohra S, Shamseer L, Sampson M, et al. CONSORT extension for reporting N-of-1 trials (CENT) 2015 statement. *Journal of Clinical Epidemiology*. 2015;76:9-17. doi:10.1016/j.jclinepi.2015.05.004
- 1673 15. Shamseer L, Sampson M, Bukutu C, Schmid CH, et al. CENT 2015: Explanation and elaboration. *Journal of Clinical Epidemiology*. 2015;76:18-46. doi:10.1016/j.jclinepi.2015.05.018
- 1674 16. Kwasnicka D, Inauen J, Nieuwenboom W, et al. Challenges and solutions for N-of-1 design studies in health psychology. *Health Psychology Review*. 2019;13(2):163-178. doi:10.1080/17437199.2018.1564627
- 1675 17. Ma Y, Mazumdar M, Memtsoudis SG. Beyond repeated-measures analysis of variance: Advanced statistical methods for the analysis of longitudinal data in anesthesia research. *Regional Anesthesia and Pain Medicine*. 2012;37(1):99-105. doi:10.1097/AAP.0b013e31823ebc74
- 1676 18. Mallinckrodt CH, Clark SWS, Carroll RJ, Molenberghs G. Assessing response profiles from incomplete longitudinal clinical trial data under regulatory considerations. *Journal of Biopharmaceutical Statistics*. 2003;13(2):179-190. doi:10.1081/BIP-120019265
- 1677 19. Mallinckrodt CH, Clark WS, David SR. Accounting for dropout bias using mixed-effects models. *Journal of Biopharmaceutical Statistics*. 2001;11(1-2):9-21. doi:10.1081/BIP-100104194

- 1678 20. Mallinckrodt CH, Raskin J, Wohlreich MM, Watkin JG, Detke MJ. The efficacy of duloxetine: A comprehensive summary of results from MMRM and LOCF_ANCOVA in eight clinical trials. *BMC Psychiatry*. 2004;4:26. doi:10.1186/1471-244X-4-26
- 1679 21. Prakash A, Risser RC, Mallinckrodt CH. The impact of analytic method on interpretation of outcomes in longitudinal clinical trials. *International Journal of Clinical Practice*. 2008;62(8):1147-1158. doi:10.1111/j.1742-1241.2008.01808.x
- 1680 22. Liang KY, Zeger SL. Longitudinal data analysis using generalized linear models. *Biometrika*. 1986;73(1):13-22. doi:10.1093/biomet/73.1.13
- 1681 23. Wang M, Kong L, Li Z, Zhang L. Covariance estimators for generalized estimating equations (GEE) in longitudinal analysis with small samples. *Statistics in Medicine*. 2016;35(10):1706-1721. doi:10.1002/sim.6817
- 1682 24. Pan W. Sample size and power calculations with correlated binary data. *Controlled Clinical Trials*. 2001;22(3):211-227. doi:10.1016/s0197-2456(01)00131-3
- 1683 25. Tong G, Nevins P, Ryan M, others, Taljaard M. A review of current practice in the design and analysis of extremely small stepped-wedge cluster randomized trials. *Clinical Trials*. 2024;22(1):45-56. doi:10.1177/17407745241276137
- 1684 26. Hedges LV, Shadish WR, Natesan Batley P. Power analysis for single-case designs: Computations for $(AB)^k$ designs. *Behavior Research Methods*. 2023;55(7):3494-3503. doi:10.3758/s13428-022-01971-9
- 1685 27. Bouwmeester S, Jongerling J. Power of a randomization test in a single case multiple baseline AB design. *PLoS ONE*. 2020;15(2):e0228355. doi:10.1371/journal.pone.0228355
- 1686 28. Blackston JW, Chapple AG, McGree JM, McDonald S, Nikles J. Comparison of aggregated N-of-1 trials with parallel and crossover randomized controlled trials using simulation studies. *Healthcare*. 2019;7(4):137. doi:10.3390/healthcare7040137
- 1687 29. Schork NJ, Beaulieu-Jones B, Liang WS, Smalley S, Goetz LH. Exploring human biology with N-of-1 clinical trials. *Cambridge Prisms: Precision Medicine*. 2023;1:e12. doi:10.1017/pcm.2022.15

- 1688 30. Zucker DR, Ruthazer R, Schmid CH. Individual (N-of-1) trials can be combined to give population comparative treatment effect estimates: Methodologic considerations. *Journal of Clinical Epidemiology*. 2010;63(12):1312-1323. doi:10.1016/j.jclinepi.2010.04.020
- 1689 31. Defelippe VM, Thiel GJMW van, Otte WM, Schutgens REG, Stunnenberg B, et al. Toward responsible clinical n-of-1 strategies for rare diseases. *Drug Discovery Today*. 2023;28(10):103688. doi:10.1016/j.drudis.2023.103688
- 1690 32. Berger VW, Rezvani A, Makarewicz VA. Direct effect on validity of response run-in selection in clinical trials. *Controlled Clinical Trials*. 2003;24(2):156-166. doi:10.1016/s0197-2456(02)00316-1
- 1691 33. Senn S. *Cross-over Trials in Clinical Research*. 2nd ed. Wiley; 2002.
- 1692 34. Mancl LA, DeRouen TA. A covariance estimator for GEE with improved small-sample properties. *Biometrics*. 2001;57(1):126-134.
- 1693 35. Maxwell SE, Delaney HD, Kelley K. *Designing Experiments and Analyzing Data: A Model Comparison Perspective*. 3rd ed. Routledge; 2018.
- 1694 36. Mallinckrodt CH, Lane PW, Schnell D, Peng Y, Mancuso JP. Recommendations for the primary analysis of continuous endpoints in longitudinal clinical trials. *Drug Information Journal*. 2008;42(4):303-319. doi:10.1177/009286150804200402
- 1695 37. Mauchly JW. Significance test for sphericity of a normal n-variate distribution. *Annals of Mathematical Statistics*. 1940;11(2):204-209. doi:10.1214/aoms/1177731915
- 1696 38. Kenward MG, Roger JH. Small sample inference for fixed effects from restricted maximum likelihood. *Biometrics*. 1997;53(3):983-997. doi:10.2307/2533558
- 1697 39. Satterthwaite FE. An approximate distribution of estimates of variance components. *Biometrics Bulletin*. 1946;2(6):110-114. doi:10.2307/3002019

1698

1699 *Rendered on 2026-06-19 at 09:12 PDT.*

1700 *Source: ~/prj/res/36-pmsimstats-ng/pmsimstats-ng/analysis/report/08-test-procedure-design-sen.*