

Type-I calibration of fixed-effect interaction tests in linear mixed models: a staged diagnosis of working-covariance mis-specification and a cluster-robust remedy

pmsimstats team

2026-05-20

Contents

7	Abstract	2
8	1 Introduction	2
9	2 A staged diagnostic protocol	3
10	2.1 The four-channel decomposition	4
11	2.2 The working-covariance model ladder	4
12	2.3 The sample-size gradient	4
13	3 Worked example: an interaction test under carryover	5
14	3.1 The analysis model	5
15	3.2 The null configuration	5
16	4 Results	6
17	4.1 The conservatism is standard-error over-estimation	6
18	4.2 The corCAR1 residual layer is the cause	8
19	4.3 The conservatism is asymptotic, not a small-sample artefact	9
20	4.4 A cluster-robust standard error restores calibration	10
21	5 Discussion	11
22	5.1 The result is not about autoregressive correlation	11
23	5.2 Two features that are not obvious in advance	12
24	5.3 When the problem should be expected	12
25	5.4 Limitations	13
26	6 Conclusions	13
27	7 Reproducibility	14
28	References	14

Abstract

Background. The linear mixed model is the standard tool for testing a fixed-effect interaction in longitudinal and clustered data. The model-based Wald test of that interaction is exact only when the working within-cluster covariance is correctly specified. In practice an analyst chooses a convenient parametric residual structure, a continuous-time autoregressive correlation, compound symmetry, or an unstructured form, and then trusts the model-based standard error. The direction and the magnitude of the mis-calibration that follows when the working covariance is wrong are rarely checked.

Methods. We set out a staged diagnostic protocol, a four-channel decomposition of the test statistic, a working-covariance model ladder, and a sample-size gradient, and apply it to a worked example: the test of a biomarker-by-treatment interaction in an aggregated N-of-1 carryover simulation, analysed with an `lme` model carrying a random intercept and a `corCAR1` residual correlation.

Results. At the null the model-based standard error of the interaction is over-estimated by 6 to 10 percent, and the realised type-I error is approximately 0.03 against a nominal 0.05. The degrees-of-freedom rule, the positive-definiteness correction, and point-estimate bias are excluded as causes. The model ladder localises the mis-calibration to the `corCAR1` residual layer: a random intercept alone is well calibrated. The sample-size gradient shows the bias is asymptotic, not a moderate-sample artefact: it does not attenuate from $N = 35$ to $N = 280$. A CR2 bias-reduced cluster-robust standard error restores the standard-error ratio to approximately one and the type-I error to approximately 0.05.

Conclusions. A mis-specified parametric working covariance can bias the model-based standard error of a mixed-model interaction test asymptotically, in either direction. Cluster-robust standard errors are a general remedy, and the staged diagnostic is a reusable instrument for detecting the problem.

1 Introduction

The linear mixed model is, by a wide margin, the most common vehicle for testing whether an effect is moderated by a covariate in longitudinal or clustered data^{1,2}. The quantity of interest is a fixed-effect interaction coefficient, and the inference almost always proceeds through a Wald test: the coefficient estimate is divided by a model-based standard error and referred to a t distribution. That procedure is exact only under a condition that is easy to state and easy to forget. The working covariance, the random-effects structure together with the parametric residual correlation the analyst has assumed, must be correct.

It has been understood since the work of Eicker, Huber, and White that this condition is more demanding than it appears³⁻⁵. When the working covariance is mis-specified, the model-based standard error is no longer a consistent estimator of the true sampling variability of the fixed-effect estimate. Liang and Zeger⁶ carried this insight into longitudinal data analysis: their generalized estimating equations decouple the point estimate, computed under a convenient working correlation, from the variance estimate, computed by an empirical sandwich that re-

68 mains consistent whatever the working correlation. The sandwich, or cluster-robust, variance
 69 estimator has since become a standard instrument, and a substantial literature has refined it
 70 for the small-sample case^{7–10}.

71 Yet in routine mixed-model practice the working covariance is still trusted. An analyst fitting
 72 longitudinal data with `nlme` or a comparable package will choose a residual correlation, very
 73 often a continuous-time first-order autoregressive form, because it is pharmacokinetically or
 74 biologically plausible¹¹, and will then report the model-based standard error of the interaction
 75 without asking what happens if the chosen form is not exactly the form the data follow.
 76 This is not negligence so much as a reasonable working assumption that is seldom audited.
 77 But how large is the error when the assumption fails? Is the resulting test conservative or
 78 anti-conservative? Does the mis-calibration diminish as the sample grows, so that it can
 79 be dismissed as a small-sample nuisance, or does it persist? These questions do not have a
 80 single answer, because the answer depends on the interaction being tested and on the way the
 81 working covariance departs from the truth. They have to be asked of a particular analysis.

82 This paper does exactly that for one analysis, and in doing so offers two contributions that we
 83 believe are more general than the single case. The first is a diagnostic protocol. We decompose
 84 the Wald statistic into four channels through which mis-calibration can enter, and we add two
 85 further instruments, a working-covariance model ladder that localises the cause and a sample-
 86 size gradient that separates a finite-sample artefact from an asymptotic one. The protocol
 87 is not specific to the example and can be applied to any mixed-model interaction test whose
 88 calibration is in doubt. The second contribution is the finding the protocol produced. In our
 89 worked example the model-based interaction test is conservative, not anti-conservative as the
 90 sandwich literature’s emphasis on under-estimated standard errors might lead one to expect;
 91 the conservatism is material, on the order of ten percent in standard-error terms; and it is
 92 asymptotic, persisting undiminished across an eightfold range of sample size. We trace it to
 93 a single component of the working covariance, and we show that a cluster-robust standard
 94 error removes it.

95 The worked example is drawn from the `pmsimstats-ng` project, a simulation study of power
 96 to detect biomarker-treatment interactions in aggregated N-of-1 trials under carryover¹². The
 97 example is incidental to the argument. The phenomenon, a parametric working covariance
 98 that biases a fixed-effect interaction test asymptotically, and the remedy are not properties of
 99 that simulation; they are properties of the mixed-model Wald test, and the example serves only
 100 to make them concrete and quantified. We begin in Section 2 with the diagnostic protocol.
 101 Section 3 describes the worked example. Section 4 reports the staged diagnosis and the
 102 cluster-robust remedy. Section 5 discusses the generality of the result and its limits.

103 2 A staged diagnostic protocol

104 Consider a fixed-effect interaction coefficient θ estimated by $\hat{\theta}$, with a model-based standard
 105 error $\widehat{\text{SE}}$ and an assigned degrees of freedom ν . The Wald test rejects the null $H_0 : \theta = 0$ at
 106 level α when the statistic $T = \hat{\theta}/\widehat{\text{SE}}$ exceeds the critical value $t_{\nu, 1-\alpha/2}$ in absolute value. The
 107 test is correctly calibrated when the realised rejection rate under H_0 equals α . When it does

not, the discrepancy must enter through one of exactly four channels, and the channels are separable by simulation.

2.1 The four-channel decomposition

1. **Point-estimate bias.** If $E(\hat{\theta}) \neq 0$ under H_0 , the statistic is decentred. We measure this by the mean of $\hat{\theta}$ across null replicates.
2. **Standard-error calibration.** If the model-based $\widehat{SE}$ does not estimate the true sampling standard deviation of $\hat{\theta}$, the statistic is rescaled. We measure this by the ratio

$$r = \frac{\overline{\widehat{SE}}}{SD(\hat{\theta})},$$

the mean model-reported standard error over the empirical standard deviation of the estimate across replicates. A ratio above one indicates an over-stated standard error and a conservative test; a ratio below one the reverse.

3. **Degrees of freedom.** If ν is too small, the t reference has tails too heavy and the test is conservative; if too large, the reverse. We measure this by recomputing the p-values against a normal reference and comparing the realised type-I error.
4. **Reference-distribution shape.** Even with $\hat{\theta}$ unbiased, $\widehat{SE}$ calibrated, and ν correct, the statistic may not follow a t distribution. We measure this by the standard deviation of T under H_0 , which should be approximately one, and by the empirical distribution function of the p-values, which should be $\text{Uniform}(0, 1)$.

The four channels are not merely a checklist. They are close to exhaustive: a Wald test is a centred statistic, a scale, a reference distribution, and a degrees-of-freedom choice, and nothing else. Locating the discrepancy in one channel rather than another is the first and most informative step of a diagnosis, because the channels have different causes and different remedies.

2.2 The working-covariance model ladder

The four-channel decomposition says *how* a test is mis-calibrated. It does not say *why*. When the decomposition points to the standard-error channel, the natural follow-up is to ask which part of the fitted model is responsible. We answer this with a model ladder: the same datasets are refitted under a graded sequence of working covariances, from the simplest to the production model, and the standard-error ratio is tracked along the sequence. The rung at which the ratio departs from one identifies the component that introduces the mis-calibration. Because every rung is fitted to the same simulated data, the comparison is free of Monte Carlo noise between models.

2.3 The sample-size gradient

A standard-error mis-calibration may be one of two qualitatively different things. It may be a finite-sample artefact, in which case it attenuates as the sample grows and the working-covariance parameters are estimated more sharply; the small-sample corrections of Kenward

and Roger^{13,14} address mis-calibration of this kind. Or it may be asymptotic, a consequence of a working covariance that does not match the data-generating covariance no matter how much data is collected, in which case it persists. The two are distinguished by running the decomposition across a range of sample sizes. Attenuation toward nominal calibration marks the finite-sample case; a flat profile marks the asymptotic case. The distinction matters, because a finite-sample artefact at a realistic sample size may be tolerable, whereas an asymptotic bias is a property of the method and will not be outgrown.

3 Worked example: an interaction test under carryover

The example is the test of a biomarker-by-treatment interaction in the carryover-sensitivity simulation study of the pmsimstats-ng project. Aggregated N-of-1 trials repeatedly cross each participant between an active treatment and a control condition, and the scientific question is whether a baseline biomarker moderates the treatment effect. Carryover, the persistence of drug effect into a washout period, complicates the analysis, and the study compares analysis-model strategies for accommodating it¹⁵. The reporting follows the ADEMP structure of Morris, White, and Crowther¹⁵.

3.1 The analysis model

For a symptom score $S_{x,it}$ on participant i at timepoint t , the analysis model is the linear mixed model

$$S_{x,it} = \beta_0 + \beta_1 \text{bm}_i + \beta_2 t + \beta_3 X_{it} + \theta (\text{bm}_i \times X_{it}) + b_i + \varepsilon_{it},$$

fitted by `nlme::lme` with a participant random intercept $b_i \sim \text{N}(0, \sigma_b^2)$ and a continuous-time first-order autoregressive residual correlation, `corCAR1(form = ~ t | ptID)`. Here bm_i is the standardised biomarker, X_{it} a drug-exposure predictor, and θ the interaction coefficient of interest. The continuous-time autoregressive residual was adopted in the project for its clinical realism, nearby measurements being more strongly correlated than distant ones, and for the numerical range it affords¹¹. The Wald test of θ is the object of this paper.

3.2 The null configuration

Calibration is a null-hypothesis property, so the diagnosis is run where the biomarker exerts no moderation, $c_{bm} = 0$, and the true interaction is therefore $\theta = 0$. The primary cell is the reference configuration of the parent study: the Hybrid N-of-1 design, $N = 70$ participants, an exponential carryover form with a one-week half-life, and an autoregressive parameter of 0.7. A preliminary check established that at $c_{bm} = 0$ the study's two data-generating architectures collapse to a bit-identical process, so the diagnosis need not stratify by architecture.

The feature of the example that matters for the argument is the relationship between the working covariance and the truth. The simulated symptom score is a sum of three latent factors, each following its own autoregressive process, with timepoint-dependent variances, plus a participant-level baseline term. The working covariance of the analysis model, one ran-

Table 1: Four-channel decomposition at the primary null cell (Hybrid design, $N = 70$, $c_{bm} = 0$, 5000 replicates), for the three analysis specifications A1, A2, A3. The interaction estimate is effectively unbiased; the SE ratio exceeds one; the statistic is correspondingly compressed; the t and z references coincide because the degrees of freedom are large.

Spec	Mean estimate	SE ratio r	SD(T)	Type I (t)	Type I (z)	Mean DF
A1	+0.032	1.069	0.932	0.033	0.034	487
A2	+0.039	1.097	0.909	0.029	0.029	487
A3	+0.033	1.061	0.939	0.035	0.036	486

dom intercept and one single-parameter corCAR1 residual, cannot reproduce that structure exactly. The example is therefore an instance of the general situation the protocol is designed for: a parametric working covariance that is plausible, conventional, and not exactly correct.

4 Results

The diagnosis was carried out at 5000 null replicates for the four-channel decomposition and the model ladder, and at 1500 replicates per point for the sample-size gradient. The cluster-robust remedy was verified at 3000 replicates. All replicate-level outputs are retained for inspection.

4.1 The conservatism is standard-error over-estimation

Table 1 reports the decomposition for the three analysis specifications carried by the parent study, which differ only in the drug-exposure predictor X_{it} and are treated here as three instances of the same interaction test. The reading is consistent across all three. The interaction estimate is effectively unbiased: the null mean is approximately +0.03 on a coefficient whose sampling standard deviation is near unity, a displacement of about 0.04 of a standard deviation. A displacement of that size is statistically detectable at 5000 replicates but is negligible in magnitude, and in any case it points toward rejection and so cannot produce a conservative test. The degrees-of-freedom channel is excluded directly: the model assigns approximately 487 degrees of freedom to the interaction term, and recomputing the p-values against a normal reference moves the realised type-I error by about 0.001, as the near-identical t - and z -reference columns show.

The discrepancy is in the standard-error channel. The model-based standard error exceeds the empirical sampling standard deviation of the estimate by 6 to 10 percent. This single fact accounts for the whole of the conservatism: the standard deviation of the test statistic equals the reciprocal of the standard-error ratio to three significant figures, so the statistic is compressed by exactly the factor by which the standard error is inflated, and no residual reference-distribution anomaly remains to be explained. The realised type-I error at $\alpha = 0.05$ is 0.029 to 0.035 against the nominal 0.05.

Figure 1 confirms the nature of the mis-calibration. The empirical distribution function of




Figure 1: Empirical distribution function of the null p-values at the primary cell, against the Uniform(0,1) diagonal. All three specifications bow uniformly below the diagonal, the signature of a global scale compression rather than a tail anomaly.

Table 2: Working-covariance model ladder at the primary null cell (5000 replicates), for specifications A1 and A2. The same datasets are refitted under each structure. M1, a random intercept alone, is well calibrated; M3, the production model, adds a corCAR1 residual and is conservative.

Model	Structure	r (A1)	r (A2)	Type I (A1)	Type I (A2)
M0	OLS, no within-cluster correlation	1.323	1.289	0.010	0.012
M1	Random intercept only	0.989	0.972	0.049	0.055
M2	corCAR1 residual only	1.023	1.070	0.043	0.036
M3	Random intercept + corCAR1 (production)	1.068	1.097	0.033	0.029

the null p-values lies below the diagonal across the whole unit interval and for all three specifications. A test that failed only in its tails would depart from the diagonal only near zero; the uniform bowing seen here is the signature of a global scale compression, which is what an inflated standard error produces. The positive-definiteness correction available to the data generator was not triggered in any of the cell’s paths and is excluded as a contributor.

4.2 The corCAR1 residual layer is the cause

The decomposition says the standard error is over-estimated. The model ladder, Table 2, says which part of the model does the over-estimating. Four working covariances were fitted to the identical null datasets: ordinary least squares with no within-cluster term (M0), a random intercept alone (M1), a corCAR1 residual alone (M2), and the production combination of a random intercept with a corCAR1 residual (M3).

The result is unambiguous. M1, the random intercept with no residual correlation, is well calibrated: its standard-error ratio sits between 0.97 and 0.99 and its type-I error between 0.049 and 0.055. M3, which adds the corCAR1 residual on top of that random intercept, is conservative: the ratio rises to between 1.07 and 1.10 and the type-I error falls to between 0.029 and 0.033. The corCAR1 residual-correlation layer is therefore the proximate cause of the conservatism. It is worth being precise about what this does and does not say. It does not say that the corCAR1 residual is wrong in any absolute sense; the simulated factors are genuinely autocorrelated, so a residual correlation term is capturing a real feature of the data. It says that adding that term to a model that already carries a random intercept degrades the calibration of the fixed-effect interaction test.

Two further features of the ladder, visible in Figure 2, are worth recording. The first is that M0, ordinary least squares, is not anti-conservative, as the familiar warning that ignoring positive within-cluster correlation under-states standard errors would suggest. It is strongly conservative, with a standard-error ratio near 1.3. The reason is instructive: the interaction here is a within-participant contrast, and ordinary least squares charges it the full residual variance, between-participant and within-participant combined, when the error relevant to a within-participant contrast is only the within-participant part. The familiar warning concerns between-cluster effects; for a within-cluster effect the sign of the OLS error is reversed. The




Figure 2: Interaction standard-error calibration across the model ladder. Left, the ratio of model-reported to empirical standard error; right, the realised type-I error. The dashed lines mark perfect calibration. The random intercept alone (M1) sits on both reference lines; the production model (M3) does not.

Table 3: Sample-size gradient under the production model M3 (1500 replicates per point). The standard-error ratio r and the type-I error (T1) are shown for each specification. Neither trends toward its nominal value as N increases.

N	r A1	r A2	r A3	T1 A1	T1 A2	T1 A3
35	1.072	1.103	1.064	0.038	0.028	0.039
70	1.056	1.083	1.050	0.043	0.031	0.045
140	1.086	1.098	1.079	0.041	0.033	0.042
280	1.052	1.070	1.046	0.037	0.028	0.035

second feature is that the corCAR1 residual without a random intercept (M2) is intermediate. Taken together the ladder shows that calibration is recovered by the random intercept alone and removed by the corCAR1 addition, and that the magnitude and even the sign of the mis-calibration is a property of the working covariance, not a universal constant.

4.3 The conservatism is asymptotic, not a small-sample artefact

If the conservatism were a finite-sample artefact, it would diminish as the sample grew and the working-covariance parameters were estimated more precisely. Table 3 and Figure 3 show that it does not. Across an eightfold range of sample size, from $N = 35$ to $N = 280$ participants, the standard-error ratio holds between 1.05 and 1.10 with no trend toward one, and the type-I error holds between 0.028 and 0.045 with no trend toward 0.05. At the largest sample size the production specification still rejects the true null at 0.028. The conservatism is asymptotic. It is a property of the M3 working covariance and will not be outgrown.




Figure 3: Type-I error and standard-error ratio against sample size under the production model M3. Neither panel trends toward its dashed reference line as the number of participants increases from 35 to 280.

Table 4: Cluster-robust remedy at the primary null cell (3000 replicates). The model-based standard error of the production model is replaced by a CR2 bias-reduced cluster-robust standard error clustered on participant, with Satterthwaite degrees of freedom. The CR2 standard-error ratio is approximately one and the type-I error approximately 0.05.

Spec	r model-based	r CR2	Type I model	Type I CR2	CR2 df
A1	1.062	0.978	0.036	0.048	23.7
A2	1.096	1.000	0.029	0.045	23.7
A3	1.053	0.976	0.039	0.049	23.7

247 The explanation that is consistent with all three results is a working-covariance mismatch
248 rather than a small-sample effect. The simulated symptom score is a mixture of autoregres-
249 sive factors with unequal, timepoint-dependent variances, plus a participant baseline term.
250 The M3 working covariance, a single random intercept and a single-parameter corCAR1 resid-
251 ual, cannot match a mixture of autoregressive processes exactly. Increasing the number of
252 participants estimates the parameters of that working covariance ever more sharply, but it
253 does not change the fact that the covariance being estimated is the wrong one, and so the
254 model-based standard error converges to a biased target. The bias persists because it is a
255 bias of the model, not of the estimate of the model.

256 4.4 A cluster-robust standard error restores calibration

257 The diagnosis points to a remedy that does not require abandoning the corCAR1 model. If the
258 model-based standard error is biased because the working covariance is wrong, the natural re-
259 placement is a standard error that is consistent whatever the working covariance^{5,6}. We replace

the model-based standard error of the interaction with a cluster-robust sandwich standard error, clustered on participant, in the bias-reduced CR2 form of Bell and McCaffrey⁸, with the Satterthwaite degrees of freedom of Pustejovsky and Tipton^{10,16}. The CR2 adjustment and its associated small-sample test were designed for precisely the regime here, a moderate number of clusters, and are implemented for `lme` objects in the `clubSandwich` package^{17,18}.

Table 4 reports the verification. The model-based standard-error ratio reproduces the conservatism, between 1.05 and 1.10. The CR2 cluster-robust ratio is between 0.98 and 1.00, and the realised type-I error rises from approximately 0.03 under the model-based test to between 0.045 and 0.049 under the CR2 test, within Monte Carlo error of the nominal 0.05. The Satterthwaite degrees of freedom for the cluster-robust test are approximately 24, far below the model-based 487; the small-sample reference is doing real work, and a naive normal reference for the sandwich statistic would itself be mildly anti-conservative^{9,19}.




Figure 4: Null p-value distribution functions under the model-based and CR2 cluster-robust tests, by specification. The model-based curves bow below the diagonal; the CR2 curves lie on it.

Figure 4 shows the same result as a distribution function. The model-based p-values bow below the diagonal, as in Figure 1; the cluster-robust p-values lie on it. The point estimate and the corCAR1 model are unchanged; only the standard error of the interaction, and the reference against which it is tested, are replaced.

5 Discussion

5.1 The result is not about autoregressive correlation

It would be a mistake to read this paper as a caution against the continuous-time autoregressive residual in particular. The corCAR1 term is the component the model ladder happened to implicate in this example, but the mechanism is general. The model ladder is itself the evidence: it fitted four working covariances and obtained four different standard-error ratios, from 1.3 under ordinary least squares to 1.0 under a random intercept to near 1.1 under the

283 production model. The mis-calibration is a property of the match between the working co-
 284 variance and the truth, and any parametric residual or random-effects structure, compound
 285 symmetry, an unstructured form, a random slope, a spatial correlation, is capable of produc-
 286 ing it. That a mis-specified working covariance leaves the model-based fixed-effect standard
 287 error inconsistent is, moreover, not a new claim; it is the Eicker-Huber-White result carried
 288 into the mixed-model setting^{3-5,20}. What the worked example contributes is a quantified, de-
 289 composed instance in a realistic longitudinal interaction test, with two features that are not
 290 obvious in advance.

291 5.2 Two features that are not obvious in advance

292 The first is the direction. The sandwich literature is usually motivated by anti-conservative
 293 tests: ignored positive correlation under-states standard errors, and the robust estimator
 294 corrects an inflated rejection rate downward. Here the model-based test is conservative, and
 295 the cluster-robust estimator corrects the rejection rate upward. The ladder explains why the
 296 intuition can mislead. The effect under test is a within-cluster contrast, and for a within-
 297 cluster contrast the relationship between a mis-specified covariance and the standard-error
 298 error need not have the sign that the between-cluster intuition supplies. An analyst who
 299 reasons that “my standard errors are conservative because I have modelled the correlation”
 300 has the logic backward in this case: modelling the correlation, through the `corCAR1` term, is
 301 what introduced the conservatism, and the less elaborate model was the calibrated one.

302 The second is that the bias is asymptotic. It would have been comfortable to attribute the
 303 conservatism to moderate sample size, and an earlier informal reading of the parent study did
 304 so. The sample-size gradient refutes that reading directly. This matters for how the problem
 305 is classified. The small-sample corrections of Kenward and Roger^{13,14} adjust the fixed-effect
 306 standard error and its degrees of freedom for the finite-sample uncertainty in the estimated
 307 variance parameters; they are the right instrument for a finite-sample artefact. They are
 308 not the right instrument here, because the bias is not a finite-sample artefact: it does not
 309 vanish as the variance parameters are estimated with unlimited precision. A Kenward-Roger
 310 correction to a mis-specified working covariance corrects the finite-sample term and leaves the
 311 asymptotic mismatch untouched. The cluster-robust standard error addresses the asymptotic
 312 mismatch directly, which is why it is the appropriate remedy here.

313 5.3 When the problem should be expected

314 The conditions for the problem are not exotic. They are a fixed-effect interaction tested in a
 315 mixed model; a parametric working covariance that is plausible but not known to be exactly
 316 correct, which is to say almost every parametric working covariance in applied use; and an
 317 interaction that is partly or wholly a within-cluster contrast. Wherever those conditions hold,
 318 the calibration of the model-based Wald test is an open question, and the staged protocol of
 319 Section 2 answers it at the cost of a simulation. We would suggest, as a matter of routine
 320 practice, that an interaction test reported from a mixed model with a non-trivial working
 321 covariance be accompanied either by a calibration check of this kind or by a cluster-robust
 322 standard error, which removes the question by being agnostic to the working covariance.

5.4 Limitations

Several limitations bound the scope of the evidence. The diagnosis was carried out for one data-generating family, one estimand, and one production working covariance; the asymptotic conservatism is demonstrated for the corCAR1-plus-random-intercept model, and the generality beyond that rests on the model ladder and on established theory rather than on a sweep of every covariance structure. A broader empirical base, fitting an unstructured or compound-symmetry analysis to the same data, would strengthen the general claim, and is a natural extension. The cluster-robust remedy is not without cost: the Satterthwaite degrees of freedom for the CR2 test are approximately 24 against a model-based 487, so the robust test spends precision to buy validity, and at a genuinely small number of clusters that price would rise. The diagnosis was conducted at the null; the consequence for power, namely that the power estimates of the parent study are mildly pessimistic rather than optimistic, follows from the direction of the standard-error error but was not measured directly here. Finally, the small positive point-estimate bias noted in Section 4.1, while negligible in magnitude and irrelevant to the conservatism, is a genuine finite-sample feature of the estimator that a separate examination might characterise.

6 Conclusions

1. The model-based Wald test of a fixed-effect interaction in a linear mixed model can be materially mis-calibrated when the parametric working covariance does not match the data-generating covariance. In the worked example the mis-calibration is a conservative one: the model-based standard error of the interaction is over-estimated by 6 to 10 percent and the realised type-I error is approximately 0.03 against a nominal 0.05.
2. The four-channel decomposition localises the discrepancy to the standard-error channel and excludes point-estimate bias, the degrees-of-freedom rule, and the positive-definiteness correction.
3. The working-covariance model ladder localises the cause to the corCAR1 residual layer. A random intercept alone is well calibrated; adding the corCAR1 residual removes the calibration.
4. The sample-size gradient shows the mis-calibration is asymptotic. It does not attenuate from $N = 35$ to $N = 280$, and so it is a property of the working covariance and not a small-sample artefact amenable to a Kenward-Roger correction.
5. A CR2 bias-reduced cluster-robust standard error, clustered on participant and referred to a Satterthwaite distribution, restores the standard-error ratio to approximately one and the type-I error to approximately 0.05, without altering the point estimate or the analysis model.
6. The staged protocol, the four-channel decomposition, the model ladder, and the sample-size gradient, is a reusable instrument for any mixed-model interaction test whose calibration is in doubt. We recommend that interaction tests reported from mixed models with a non-trivial working covariance be accompanied by a calibration check of this kind, or by a cluster-robust standard error that is agnostic to the working covariance.

363 7 Reproducibility

364 The diagnosis is driven by four scripts under `analysis/scripts/carryover-sensitivity/`:
 365 `diagnose-null-type1.R` for the four-channel decomposition, `diagnose-null-type1-stages46.R`
 366 for the model ladder and the sample-size gradient, and `diagnose-null-type1-cr2.R` for the
 367 cluster-robust verification, with `check-null-architecture-equivalence.R` establishing the
 368 architecture-invariance of the null. Each writes a replicate-level `.rds` file into the `output/`
 369 subdirectory, and this manuscript reads those files and regenerates every table and figure
 370 inline. The cluster-robust standard errors use the `clubSandwich` package. The examination
 371 plan that preceded the manuscript is retained alongside it as `type1-examination-plan.md`.

372 References

- 373 1. Laird NM, Ware JH. Random-effects models for longitudinal data. *Biometrics*. 1982;38(4):963-974. doi:10.2307/2529876
- 374 2. Diggle PJ, Heagerty P, Liang KY, Zeger SL. *Analysis of Longitudinal Data*. 2nd ed. Oxford University Press; 2002.
- 375 3. Eicker F. Limit theorems for regressions with unequal and dependent errors. In: *Proceedings of the Fifth Berkeley Symposium on Mathematical Statistics and Probability, Volume 1*. University of California Press; 1967:59-82.
- 376 4. Huber PJ. The behavior of maximum likelihood estimates under nonstandard conditions. In: *Proceedings of the Fifth Berkeley Symposium on Mathematical Statistics and Probability, Volume 1*. University of California Press; 1967:221-233.
- 377 5. White H. A heteroskedasticity-consistent covariance matrix estimator and a direct test for heteroskedasticity. *Econometrica*. 1980;48(4):817-838. doi:10.2307/1912934
- 378 6. Liang KY, Zeger SL. Longitudinal data analysis using generalized linear models. *Biometrika*. 1986;73(1):13-22. doi:10.1093/biomet/73.1.13
- 379 7. Mancl LA, DeRouen TA. A covariance estimator for GEE with improved small-sample properties. *Biometrics*. 2001;57(1):126-134. doi:10.1111/j.0006-341X.2001.00126.x
- 380 8. Bell RM, McCaffrey DF. Bias reduction in standard errors for linear regression with multi-stage samples. *Survey Methodology*. 2002;28(2):169-181.
- 381 9. Cameron AC, Miller DL. A practitioner's guide to cluster-robust inference. *Journal of Human Resources*. 2015;50(2):317-372. doi:10.3368/jhr.50.2.317

- 382 10. Pustejovsky JE, Tipton E. Small-sample methods for cluster-robust variance estimation and hypothesis testing in fixed effects models. *Journal of Business & Economic Statistics*. 2018;36(4):672-683. doi:10.1080/07350015.2016.1247004
- 383 11. Pinheiro JC, Bates DM. *Mixed-Effects Models in S and S-PLUS*. Springer; 2000.
- 384 12. Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: Statistical methods and practical considerations. *Frontiers in Digital Health*. 2020;2:13. doi:10.3389/fdgth.2020.00013
- 385 13. Kenward MG, Roger JH. Small sample inference for fixed effects from restricted maximum likelihood. *Biometrics*. 1997;53(3):983-997. doi:10.2307/2533558
- 386 14. Kenward MG, Roger JH. An improved approximation to the precision of fixed effects from restricted maximum likelihood. *Computational Statistics & Data Analysis*. 2009;53(7):2583-2595. doi:10.1016/j.csda.2008.12.013
- 387 15. Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Statistics in Medicine*. 2019;38(11):2074-2102. doi:10.1002/sim.8086
- 388 16. Satterthwaite FE. An approximate distribution of estimates of variance components. *Biometrics Bulletin*. 1946;2(6):110-114. doi:10.2307/3002019
- 389 17. Tipton E. Small sample adjustments for robust variance estimation with meta-regression. *Psychological Methods*. 2015;20(3):375-393. doi:10.1037/met0000011
- 390 18. Tipton E, Pustejovsky JE. Small-sample adjustments for tests of moderators and model fit using robust variance estimation in meta-regression. *Journal of Educational and Behavioral Statistics*. 2015;40(6):604-634. doi:10.3102/1076998615606099
- 391 19. Long JS, Ervin LH. Using heteroscedasticity consistent standard errors in the linear regression model. *The American Statistician*. 2000;54(3):217-224. doi:10.1080/00031305.2000.10474549
- 392 20. Freedman DA. On the so-called Huber sandwich estimator and robust standard errors. *The American Statistician*. 2006;60(4):299-302. doi:10.1198/000313006X152207