

1 Two Architectures for Simulating Biomarker-Treatment
2 Interactions: Implications for Statistical Power Under
3 Carryover

4 pmsimstats team

5 2026-06-10

6 **Contents**

7	1 Abstract	2
8	2 1. Introduction	3
9	3 2. Two DGP Architectures	6
10	3.1 2.1 Architecture A: Direct Mean Moderation	6
11	3.2 2.2 Architecture B: Differential Correlation (MVN)	7
12	3.3 2.3 Mathematical Comparison	8
13	4 3. Consequences for Statistical Power Under Carryover	9
14	4.1 3.1 Empirical demonstration	9
15	4.2 3.2 Why the architectures diverge	12
16	4.3 3.3 Robustness check at $N = 70$	13
17	5 4. Biological Assumptions	15
18	5.1 4.1 When Architecture A is appropriate	16
19	5.2 4.2 When Architecture B is appropriate	17
20	5.3 4.3 The prazosin-PTSD case	19
21	6 5. Implications for Simulation Study Design	20
22	6.1 5.1 Choosing the appropriate architecture	20
23	6.2 5.2 Reporting requirements	21
24	6.3 5.3 When the choice is uncertain	22
25	7 6. Alternative Analysis Strategies Under Architecture B	22
26	7.1 6.1 Restrict analysis to on-drug observations	23
27	7.2 6.2 Weighted analysis	23
28	7.3 6.3 Within-subject contrast	24
29	7.4 6.4 Two-stage random slopes	24
30	7.5 6.5 Exclusion of contaminated observations	25
31	7.6 6.6 Design-level solutions	25

32	7.7	6.7 Comparative evaluation	26
33	8	7. Relationship to the Broader Literature	27
34	8.1	7.1 Treatment effect heterogeneity	27
35	8.2	7.2 Prevalence of each architecture	28
36	8.3	7.3 Crossover and N-of-1 trial methodology	29
37	8.4	7.4 Precision medicine trial design	30
38	9	8. Conclusions	32
39	10	Conflict of interest	33
40	11	Funding	33
41	12	Data availability statement	33
42	13	Reproducibility	33
43	14	References	34
44	15	References	35

45 1 Abstract

46 **Background.** Simulation studies of predictive biomarker-treatment interactions
47 in N-of-1 clinical trials require a data-generating process (DGP) that specifies how
48 the biomarker moderates treatment response. Two fundamentally different DGP
49 architectures are used in practice, and the choice between them determines the
50 extent to which carryover effects reduce statistical power to detect the interaction.
51 The published simulation literature does not distinguish them.

52 **Methods.** We formalise two architectures: direct mean moderation (Architec-
53 ture~A), under which the biomarker scales the treatment effect in the popula-
54 tion mean structure; and differential correlation (Architecture~B, the multivari-
55 ate normal approach), under which the interaction emerges from treatment-state-
56 dependent correlation between biomarker and response in the joint distribution.
57 Power consequences under both architectures are quantified by Monte Carlo sim-
58 ulation on an aggregated N-of-1 design framework with continuous-time AR(1)
59 residual correlation in the analysis model. The biomarker, treatment phase, and
60 carryover decay are parameterised to the prazosin/PTSD reference application.

61 **Results.** Under Architecture~A, power is largely preserved under carryover be-
62 cause the proportional relationship between biomarker and drug response is main-
63 tained at every timepoint. Under Architecture~B, carryover erodes the differen-
64 tial correlation signal and reduces power. At the prazosin-calibrated reference cell
65 ($N = 35$, $c_{bm} = 0.45$, OL+BDC design) across $t_{1/2} \in \{0, 0.5, 1.0\}$ weeks, Architec-
66 ture~B power drops from 0.777 to 0.539 (a 30.6% relative loss; $z = 7.64$, $p < 10^{-13}$
67 on the B-versus-A contrast), while Architecture~A drops only from 0.751 to 0.730

(a 2.8% relative loss). The B-versus-A absolute power-loss gap of 21.7 percentage points at OL+BDC, $t_{1/2} = 1.0$ is approximately 11 times the matched Architecture~A loss; the Hybrid design produces a smaller but still significant gap ($z = 3.96$, $p < 10^{-4}$); the traditional crossover design produces no detectable gap because its within-subject AR(1) structure dominates the carryover-mediated correlation channel that Architecture~B uses. The 36-cell production factorial (2 architectures $\times$ 3 designs $\times$ 3 $t_{1/2}$ $\times$ 2 c_{bm} levels) was run at 1{,}000 replicates per cell with 100% convergence; MCSE on cell power is 0.013-0.016. Type-I error sits in [0.010, 0.074], with no architecture-specific inflation.

Conclusions. The two architectures encode different biological mechanisms, heterogeneity at the mean structure versus heterogeneity at the variance structure, and the choice between them is a substantive scientific question rather than a parameterisation convenience. We recommend explicit declaration of the DGP architecture in any biomarker-validation simulation study and provide guidance for selecting the architecture appropriate to a given clinical context.

2 1. Introduction

A central question in precision medicine trial design is whether a baseline biomarker predicts differential treatment response, that is, whether the treatment effect is moderated by the biomarker value. Simulation studies are the principal tool available for evaluating the statistical power of proposed trial designs to detect such interactions, and these simulations necessarily require a data-generating process (DGP) that encodes the biomarker-treatment interaction mechanism.

The central question is biomarker-treatment interaction; simulation studies need a DGP that encodes it.

- Does a baseline biomarker predict differential treatment response, i.e., does it moderate the treatment effect?
- Simulation studies are the principal power evaluation tool for proposed trial designs.
- These simulations require a DGP that explicitly encodes the biomarker-treatment interaction mechanism.

The statistical literature on treatment effect heterogeneity, developed notably by Rothwell (2005) and synthesized more recently by Kent et al. (2020), distinguishes between *qualitative* interactions, in which the direction of the treatment effect reverses across biomarker strata, and *quantitative* interactions, in which the magnitude varies but the direction remains consistent. As we shall argue, any simulation framework intended to evaluate power for biomarker-moderated effects must choose how to generate this heterogeneity, and that choice is not merely a technical convenience.

Every simulation framework must choose how to generate treatment effect heterogeneity, and that choice has consequences.

- HTE literature (Rothwell 2005; Kent et al. 2020) distinguishes qualitative interactions (direction reverses) from quantitative ones (magnitude varies).
- Any simulation framework must decide how to encode this heterogeneity in the DGP.
- That decision is substantive, not a mere parameterization detail.

The dominant approach in the clinical trial simulation literature is direct mean moderation, in which the biomarker enters the outcome model as an interaction term that explicitly scales the treatment effect. Simon (2010) employs this approach throughout the simulation frameworks he develops for biomarker-stratified designs; Freidlin and Korn (2014) adopt it for adaptive enrichment trials; Renfro and Sargent (2017) use it throughout their treatment of basket and umbrella trials; and it underlies the platform trial literature without exception. Recent simulation work on biomarker-by-treatment interaction estimation in randomized trials with prognostic covariates, specifically the study by Haller and Ulm (2018), likewise adopts the standard regression specification with explicit interaction terms. Among N-of-1 simulation studies, Zucker et al. (1997), Araujo et al. (2016), and Duan et al. (2013) similarly use hierarchical random-effects models in which the biomarker predicts the individual treatment effect through the mean structure. Notably, Wang and Schork (2019) and Schork (2022), whose methodological work addresses power and design for N-of-1 trials with serial correlation and carryover directly, also specify the treatment effect through a linear mean structure, which suggests that the multivariate-normal differential-correlation parameterization of Hendrickson et al. (2020) is not a convention of any particular research program but rather a paper-specific design choice. To the best of our knowledge, no other clinical trial simulation framework generates the biomarker-treatment interaction through differential correlation in a multivariate normal joint distribution rather than through direct mean moderation.

Architecture A (direct mean moderation) is the near-universal standard; the Hendrickson (2020) MVN approach is apparently unique.

- Every major simulation framework uses direct mean moderation: Simon (2010) for stratified designs; Freidlin and Korn (2014) for adaptive enrichment; Renfro and Sargent (2017) for basket/umbrella trials; all N-of-1 frameworks.
- Wang and Schork (2019) and Schork (2022) also use Architecture A despite sharing authorship with Hendrickson et al. (2020), suggesting the MVN parameterization is a paper-specific choice, not a lab convention.
- To the best of our knowledge, no other clinical trial simulation framework uses differential correlation (MVN) to encode the biomarker-treatment interaction.

We encountered both approaches in the course of developing the pmsimstats simulation framework for N-of-1 clinical trials (Hendrickson et al. 2020). The origi-

153 nal publication code uses the MVN differential correlation approach, an unusual
154 choice that, as we shall demonstrate here, has substantive consequences for power
155 estimation under carryover. A comprehensive audit of the codebase subsequently
156 revealed that a parallel tidyverse reimplementation of the same simulation had in-
157 advertently adopted the standard direct mean moderation approach. Both imple-
158 mentations produced plausible power estimates, yet they yielded qualitatively dif-
159 ferent conclusions about how carryover effects influence statistical power. Tracing
160 this discrepancy to its root cause revealed the architectural difference described in
161 this paper, which has not, to the best of our knowledge, been discussed explicitly
162 in the simulation methodology literature.

163 **The two architectures were discovered by accident during pmsim-**
164 **stats development, when two implementations of the same simula-**
165 **tion diverged under carryover.**

- 166 • The original code uses Architecture B (MVN); a tidyverse reimplemen-
167 tation inadvertently used Architecture A.
- 168 • Both produced plausible power estimates, but they reached qualita-
169 tively different conclusions about carryover effects.
- 170 • Tracing the discrepancy to its root revealed the architectural distinction,
171 which has not been explicitly discussed in the simulation methodology
172 literature.

173 We present here what we believe to be the first systematic comparison of these
174 two DGP architectures for biomarker-treatment interaction power estimation un-
175 der carryover. We shall formalize the distinction between them, demonstrate its
176 consequences through Monte Carlo simulation, and offer guidance for investiga-
177 tors designing simulation studies for predictive biomarker-moderated treatment
178 effects.

179 **This paper is the first systematic comparison of the two DGP ar-**
180 **chitectures; we formalize, simulate, and give guidance.**

- 181 • We formalize the distinction between Architecture A (mean modera-
182 tion) and Architecture B (differential correlation).
- 183 • We demonstrate power consequences through Monte Carlo simulation.
- 184 • We offer guidance for investigators on how to choose the appropriate
185 architecture.

186 **Related documents in this series.** The mathematical derivation,
187 code-level implementation details, four implementation attempts (three
188 failed and one correct), and the empirical evaluation of the original Hen-
189 drickson (2020) parameter grid under Architecture~A are documented in
190 **19-mean-moderation-implementation-notes.tex** (with a brief markdown
191 summary in **19-mean-moderation-implementation-notes.md**). A condensed
192 empirical-results summary suitable for collaborator presentation is provided in
193 **21-architecture-comparison-summary.md**.

194

3 2. Two DGP Architectures

Notation. Throughout this paper, B_i denotes the biomarker value of participant i , D_{it} the binary on-drug indicator at timepoint t (1 on drug, 0 off drug), and BR_{it} the biological response component of the outcome. These three symbols describe the data-generating-process layer. The analysis-model layer uses a related but distinct continuous drug exposure variable, D_{bc} , defined in Section 3.1. Code-style identifiers in typewriter font (**bm**, **Dbc**, **BR**) refer to specific variables in the pmsimstats codebase or to R formula syntax; mathematical symbols are used otherwise.

3.1 2.1 Architecture A: Direct Mean Moderation

In this approach the biomarker enters the mean structure of the outcome directly. The treatment effect for participant i at timepoint t is:

$$Y_{it} = \mu_0 + \beta_1 \cdot D_{it} \cdot (1 + \beta_{bm} \cdot B_i) + \epsilon_{it}$$

where D_{it} is the treatment indicator (or a continuous measure of drug exposure), B_i is the participant's biomarker value (typically centered), and β_{bm} is the biomarker moderation parameter. The interaction is explicit in the mean: participants with higher biomarker values have proportionally larger treatment effects.

Key properties:

- The interaction is **deterministic** given the biomarker value and treatment status.
- The biomarker moderation β_{bm} is a **population-level parameter** that applies uniformly to all participants.
- The signal for the interaction is in the **first moment** (conditional mean) of the outcome distribution.
- Adding noise, random effects, or carryover to D_{it} scales the entire treatment effect (including its biomarker-dependent component) proportionally. The **ratio** of drug effect between high-biomarker and low-biomarker participants is preserved.

We note that the multiplicative form above maps directly onto the standard additive regression form for predictive biomarker effects used in the trial methodology literature, specifically the effect modeling equation of Kent et al. (2020):

$$Y_{it} = \beta_0 + \beta_1 D_{it} + \beta_2 B_i + \beta_3 D_{it} B_i + \epsilon_{it}$$

The interaction coefficient β_3 in the additive form corresponds to the product $\beta_1 \cdot \beta_{bm}$ in the multiplicative form. The additive form additionally accommodates a prognostic main effect β_2 for the biomarker, which the multiplicative form sets

to zero by construction. For the purpose of power calculations on the interaction coefficient, the two forms are interchangeable up to this prognostic-effect distinction, and both fall within Architecture A.

The multiplicative and additive forms of Architecture A are interchangeable for power calculations.

- The multiplicative form maps directly to the standard additive regression (Kent et al. 2020 effect modeling equation).
- $\beta_3 = \beta_1 \cdot \beta_{bm}$; the additive form additionally allows a prognostic main effect.
- Both forms fall within Architecture A and are interchangeable for power calculations on the interaction coefficient.

This architecture is used in the pedagogical simulation (`implementations/simple/simulation.R`) and in the exploratory carryover analyses (`simulation_carryover_spectrum.R`) of the pmsimstats project.

3.2 2.2 Architecture B: Differential Correlation (MVN)

In this approach the biomarker and the biological response (BR) component are jointly drawn from a multivariate normal distribution with treatment-state-dependent correlation:

$$\begin{pmatrix} B_i \\ BR_{it} \end{pmatrix} \sim \text{MVN} \left(\begin{pmatrix} \mu_B \\ \mu_{BR}(t, D_{it}) \end{pmatrix}, \Sigma(D_{it}) \right)$$

where the covariance matrix Σ has the following treatment-state-dependent BM-BR correlation:

$$\text{Cor}(B_i, BR_{it}) = \begin{cases} c_{bm} & \text{if } D_{it} = 1 \text{ (on drug)} \\ c_{bm} \cdot e^{-\lambda \cdot t_{sd}} & \text{if } D_{it} = 0 \text{ (off drug)} \end{cases}$$

where t_{sd} is time since drug discontinuation. The exponential-decay form is the default; the limiting cases $\lambda = 0$ (no decay; correlation persists indefinitely during the off-drug phase) and $\lambda \rightarrow \infty$ (immediate decay; correlation drops to zero the instant the drug is discontinued) recover the two no-carryover model variants.

The interaction is not in the mean structure: the population mean of BR_{it} is the same for all participants with the same treatment history. Instead it emerges from the **conditional distribution**: given a participant's biomarker value $B_i = b$, the conditional expectation of their biological response is:

$$E[BR_{it} | B_i = b, D_{it} = 1] = \mu_{BR}(t) + c_{bm} \cdot \frac{\sigma_{BR}}{\sigma_B} \cdot (b - \mu_B)$$

This conditional expectation has an interaction-like structure: the biomarker value b modulates the expected BR. The effective correlation is c_{bm} on drug and $c_{bm} \cdot$

259 $e^{-\lambda t_{sd}}$ during the washout; the moderation therefore persists with attenuated
260 strength during off-drug periods and vanishes only in the limit of complete decay.

261 **In Architecture B the interaction is in the covariance structure,**
262 **not the mean; it emerges from the conditional distribution given**
263 **the biomarker value.**

- 264 • The population mean of BR_{it} is the same for all participants; no mean-
265 structure interaction exists.
- 266 • Given biomarker value b , the conditional expectation of BR scales with
267 c_{bm} , creating an interaction-like structure.
- 268 • Correlation is c_{bm} on drug and decays as $c_{bm}e^{-\lambda t_{sd}}$ off drug; moderation
269 attenuates during washout.

270 Key properties:

- 271 • The interaction is **probabilistic**: it manifests as a tendency, not a deter-
272 ministic scaling.
- 273 • The biomarker moderation c_{bm} is a **correlation parameter** that governs
274 the strength of the association.
- 275 • The signal for the interaction is in the **second moment** (covariance struc-
276 ture) of the joint distribution.
- 277 • Carryover effects in the DGP inflate the BR means during off-drug periods,
278 making them resemble on-drug means. This reduces the **differential** in the
279 correlation structure between treatment states, weakening the signal that
280 the analysis model detects.

281 This architecture is used in the Hendrickson et al. (2020) publication code and
282 the audited R/ package in pmsimstats-ng.

283 3.3 2.3 Mathematical Comparison

284 We begin the comparison by considering the expected outcome difference between
285 two participants with biomarker values b_1 and b_2 at the same drug exposure level
286 D under Architecture A:

$$E[Y \mid b_1, D] - E[Y \mid b_2, D] = \beta_1 \cdot \beta_{bm} \cdot (b_1 - b_2) \cdot D$$

287 The absolute magnitude of this difference scales linearly with D and therefore
288 shrinks during off-drug periods when D is replaced by an exposure-decayed value
289 $D \cdot \text{carryover}$. We note, however, that the **ratio** of the drug effect between high-
290 biomarker and low-biomarker participants is preserved at every level of exposure.
291 The biomarker-by-treatment signal contracts in amplitude but does not lose its
292 identifiability under carryover.

293 **Under Architecture A, carryover shrinks the signal amplitude but**
294 **preserves the ratio between biomarker groups, so identifiability is**
295 **maintained.**

- The expected outcome difference between two participants scales linearly with D .
- Under carryover, D is replaced by a decayed value and the absolute magnitude shrinks.
- But the ratio of drug effect between high- and low-biomarker participants is preserved at every exposure level; the signal contracts but remains identifiable.

Under Architecture B, the analogous quantity is:

$$E[BR|b_1, D = 1] - E[BR|b_2, D = 1] = c_{bm} \cdot \frac{\sigma_{BR}}{\sigma_B} \cdot (b_1 - b_2)$$

During off-drug periods with carryover:

$$E[BR|b, D = 0, t_{sd}] = \mu_{BR, \text{off}}(t_{sd}) + c_{bm} \cdot e^{-\lambda t_{sd}} \cdot \frac{\sigma_{BR}}{\sigma_B} \cdot (b - \mu_B)$$

Here the biomarker-dependent component decays exponentially with time off drug. As t_{sd} increases, the off-drug conditional expectation converges to $\mu_{BR, \text{off}}$ for all biomarker values, eliminating the differential signal entirely.

Under Architecture B, the biomarker-dependent component of the off-drug response decays exponentially, eventually eliminating the differential signal.

- Off-drug, the conditional expectation includes a biomarker-dependent term scaled by $c_{bm}e^{-\lambda t_{sd}}$.
- As t_{sd} grows, that term shrinks toward zero.
- In the limit, all participants converge to the same off-drug mean regardless of biomarker value; the interaction signal vanishes.

4 3. Consequences for Statistical Power Under Carryover

4.1 3.1 Empirical demonstration

We ran identical simulation configurations under both architectures using the pmsimstats-ng framework, matching the DGP parameters as closely as possible between the two approaches. We note that the non-biologic response components (time-varying response, placebo-belief accumulation, and residual noise) are generated from the same multivariate normal draw with identical means, variances, within-factor AR(1) autocorrelation, and cross-factor correlations in both architectures. The two DGPs differ *only* in how the biomarker-to-biological-response (BM-BR) linkage is encoded: Architecture~B embeds it

in the off-diagonal c_{bm} entries of Σ , whereas Architecture~A omits those entries and adds a post-draw shift $c_{bm} \cdot B_i^* \cdot \sigma_{BR}$ to on-drug BR columns (see `19-mean-moderation-implementation-notes.tex` for the scaling derivation). The divergent carryover behavior reported below is therefore attributable to the BM-BR parameterization, not to any difference in how the placebo, time-varying, or residual-noise components are generated. The analysis model was the same in both cases: an `nlme::lme` fit with a `corCAR1` continuous-time autocorrelation structure and a single biomarker-by-drug interaction term, written `bm:Dbc` in R formula syntax.

The two DGPs are identical except for how they encode the BM-BR linkage; any power divergence is attributable solely to that difference.

- All non-biologic components (time-varying, placebo-belief, residual noise) are generated identically in both architectures.
- The only difference: Architecture B encodes BM-BR interaction in Σ ; Architecture A adds a post-draw mean shift to on-drug BR.
- The same `nlme::lme` analysis model with `corCAR1` is used for both, so divergence in results is attributable to the DGP alone.

The analysis-model drug exposure variable `Dbc` is the *continuous* analogue of the binary DGP indicator D_{it} . It equals 1 during on-drug periods and $\exp(-\lambda t_{sd})$ during off-drug periods, where t_{sd} is time since drug discontinuation and λ is the analysis-model carryover decay rate (typically derived from the drug's pharmacokinetic half-life as $\lambda = \ln 2 / t_{1/2}$, in which case it matches the DGP decay rate of Section 2.2 by construction). Whereas D_{it} flips abruptly between 0 and 1, `Dbc` captures the smooth shrinkage of residual drug effect during washout. The interaction coefficient on `bm:Dbc` is therefore the test of whether the biomarker moderates the exposure-decayed drug effect.

`Dbc` is a continuous drug exposure variable that captures exponential washout; the `bm:Dbc` interaction tests whether biomarker moderates that exposure-decayed effect.

- `Dbc` = 1 on drug; = $\exp(-\lambda t_{sd})$ off drug, matching the DGP decay rate.
- Unlike the binary D_{it} , `Dbc` smoothly captures washout shrinkage.
- The test of `bm:Dbc` is the test of whether the biomarker moderates the (decay-adjusted) drug effect.

We follow the ADEMP framework of Morris, White, and Crowther (2019): Aims, Data-generating mechanisms, Estimands, Methods, Performance measures.

Setup:

- Trial designs: CO (traditional crossover), Hybrid (N-of-1 Hybrid), OL+BDC (open-label with blinded discontinuation)
- Sample size: $N = 35$
- Biomarker correlation / moderation: $c_{bm} = 0.45$ / $\beta_{bm} = 0.45$
- DGP carryover half-life: $t_{1/2} \in \{0, 0.5, 1.0\}$ weeks

- Analysis: matched Dbc with same half-life
- Replicates: 1,000 per cell (Monte Carlo standard error on a binomial power estimate near 0.5 is approximately $\sqrt{0.5 \cdot 0.5 / 1,000} \approx 0.016$)
- Implementation: 8-core `parallel::mclapply`; full factorial 2 architectures $\times 3$ designs $\times 3$ $t_{1/2} \times 2$ c_{bm} levels = 36 cells, 36,000 fits, 100% convergence

The $N = 35$ choice is the smallest of the prazosin-PTSD- calibrated sample sizes used in the published methodological literature; it is a regime in which neither architecture saturates power at the no-carryover cells, so any carryover-induced erosion is mechanically visible. Type I error under the $c_{bm} = 0$ null cells sat in $[0.010, 0.074]$ across the 18 null cells, consistent with the conservative-to- nominal regime expected for the continuous D_{bc} specification at this sample size.

$N = 35$ was chosen because it keeps power off the ceiling, making any carryover-induced loss visible; type I error is in the expected range.

- $N = 35$ is the smallest prazosin-PTSD-calibrated sample size in the literature.
- Neither architecture saturates power at zero carryover, so losses are mechanically observable.
- Type I error across 18 null cells sat in $[0.010, 0.074]$, consistent with the conservative-to-nominal range expected for this specification.

Results under Architecture A (direct mean moderation):

Design	$t_{1/2} = 0$	$t_{1/2} = 0.5$	$t_{1/2} = 1.0$
CO	0.739	0.739	0.725
Hybrid	0.992	0.987	0.978
OL+BDC	0.751	0.748	0.730

Power declines modestly under Architecture A: relative loss from $t_{1/2} = 0$ to $t_{1/2} = 1$ is 1.9% (CO), 1.4% (Hybrid), and 2.8% (OL+BDC). We note that all three values fall well below the 8-13% relative-loss range that the broader N-of-1 simulation literature reports for direct-mean-moderation DGPs, though they reproduce its qualitative finding, namely a modest erosion that leaves power essentially intact.

Architecture A shows modest power loss under carryover (1–3%), well below the broader literature’s 8–13% range.

- Relative losses: 1.9% (CO), 1.4% (Hybrid), 2.8% (OL+BDC).
- All three fall below the 8–13% range reported in the broader N-of-1 simulation literature.
- Qualitative finding: modest erosion; power essentially intact under carryover.

Results under Architecture B (differential correlation, MVN):

Power declines more sharply under Architecture B, and most markedly so in

Design	$t_{1/2} = 0$	$t_{1/2} = 0.5$	$t_{1/2} = 1.0$
CO	0.745	0.754	0.723
Hybrid	0.995	0.980	0.944
OL+BDC	0.777	0.694	0.539

the OL+BDC design: relative loss from $t_{1/2} = 0$ to $t_{1/2} = 1$ is 3.0% (CO; not significantly different from Architecture A's CO loss), 5.1% (Hybrid), and **30.6% (OL+BDC)**. The OL+BDC loss is approximately 11 times the matched Architecture-A OL+BDC loss (2.8%). It falls below the 40-60% relative-loss range that the broader N-of-1 simulation literature reports for the MVN differential-correlation DGP, but it confirms the qualitative finding, namely a substantial and strongly design-dependent erosion. Two-proportion z -tests on the B-versus-A absolute power-loss gap give $z = 7.64$ ($p < 10^{-13}$) at OL+BDC and $z = 3.96$ ($p < 10^{-4}$) at Hybrid; the CO contrast is not statistically distinguishable from zero ($z = 0.29$) because the CO design's within-subject AR(1) structure dominates the carryover-mediated correlation decay channel that Architecture B uses to encode the biomarker-treatment interaction.

Architecture B shows sharply larger power loss under carryover, especially OL+BDC (30.6%); the B-vs-A gap is significant for OL+BDC and Hybrid but not CO.

- Relative losses: 3.0% (CO), 5.1% (Hybrid), 30.6% (OL+BDC); the OL+BDC loss is $\approx 11 \times$ the matched Architecture A loss.
- B-vs-A gap: $z = 7.64$ ($p < 10^{-13}$) at OL+BDC; $z = 3.96$ ($p < 10^{-4}$) at Hybrid.
- CO gap not significant ($z = 0.29$): the CO design's within-subject AR(1) structure dominates the correlation-decay channel that Architecture B depends on.

4.2 3.2 Why the architectures diverge

We turn now to a mechanistic account of why the two architectures produce such different power losses under carryover.

Under Architecture A, carryover reduces the magnitude of the drug exposure variable seen by the analysis model (Dbc drops below 1 during washout), and the mean BR at off-drug timepoints is inflated by the residual drug effect exactly as in Architecture B. Both mechanisms compress the between-state treatment contrast and are shared by the two architectures. What Architecture A lacks is the second channel of signal loss: the biomarker moderation coefficient β_{bm} is a fixed population parameter, the biomarker enters the mean structure directly rather than through a correlation, and the noise structure is independent of drug state. The biomarker-by-drug interaction term therefore has reduced variance (because Dbc has less contrast between drug states), but the signal-to-noise ratio degrades only modestly.

Architecture A is subject to only one carryover channel (mean

blurring); it escapes the second channel (correlation erosion).

- Carryover compresses Dbc contrast and inflates off-drug BR means; both architectures share this channel.
- Architecture A lacks a second channel because β_{bm} is a fixed population parameter and the noise structure is drug-state-independent.
- Result: reduced variance in the interaction term, but signal-to-noise ratio degrades only modestly.

Under Architecture B, carryover operates on two levels simultaneously:

1. **Mean blurring:** The BR means during off-drug periods are inflated by residual carryover, making them resemble on-drug BR means. This reduces the treatment contrast that the analysis model uses to distinguish drug effect from no-drug. This channel is shared with Architecture A.
2. **Correlation erosion:** The BM-BR correlation during off-drug periods decays exponentially with t_{sd} . This is not just a statistical modeling choice; it is a structural consequence of the DGP. The correlation between biomarker and BR reflects the degree to which the drug effect is active; as the drug washes out, the correlation weakens because the drug-mediated component of BR variance shrinks relative to the drug-independent component. This channel is *unique* to Architecture B, because only here does the biomarker signal live in the second moment of the joint distribution. Under Architecture A the drug-mediated and drug-independent components of BR variance move in tandem (both are scaled by the same Dbc-weighted drug effect), so their ratio, and hence the identifiability of the interaction, is preserved even as amplitudes shrink.

The analysis model's interaction coefficient (the coefficient on **bm:Dbc**) depends on the covariance between B_i and Dbc-weighted outcomes. Under Architecture B, this covariance is being attacked from both sides: the treatment contrast in Dbc is compressed, and the BM-BR correlation that generates the covariance signal is simultaneously decaying. That Architecture A escapes the second attack entirely accounts for the divergence in power loss documented above.

Under Architecture B the interaction estimate is attacked from both sides simultaneously; Architecture A faces only one of the two attacks.

- The **bm:Dbc** coefficient depends on the covariance between B_i and Dbc-weighted outcomes.
- Under Architecture B: (1) Dbc contrast is compressed; (2) the BM-BR correlation generating the covariance signal is simultaneously decaying.
- Architecture A faces only attack (1), not attack (2); this single difference accounts for the divergence in power.

4.3 3.3 Robustness check at $N = 70$

The Section~3.1 production results are at $N = 35$, where neither architecture saturates power and the carryover-induced loss is mechanically visible. We con-

ducted a confirmation run at $N = 70$ to check whether the qualitative ordering survives the saturation that comes with the larger sample size. An 800-replicate-per-cell Monte Carlo run was executed at $N = 70$ across 12 cells (architecture $\in \{\text{MVN}, \text{mean moderation}\} \times c_{bm} \in \{0, 0.45\} \times t_{1/2} \in \{0, 0.5, 1.0\}$ weeks), all on the OL+BDC design, using 8-core `parallel::mclapply`. Wall-clock 483-s; convergence 100% across 9{,}600 fits.

An $N = 70$ confirmation run checks whether the qualitative ordering holds when power is near the ceiling.

- 800 replicates per cell, 12 cells (2 architectures $\times$ 2 c_{bm} $\times$ 3 $t_{1/2}$), OL+BDC design.
- 100% convergence across 9,600 fits.
- Purpose: verify the Section 3.1 ordering survives saturation at the larger N .

At $c_{bm} = 0.45$:

Architecture	$t_{1/2} = 0$	$t_{1/2} = 0.5$	$t_{1/2} = 1.0$
MVN (B)	0.978	0.945	0.885
Mean moderation (A)	0.976	0.973	0.959

Within-architecture losses from $t_{1/2} = 0$ to $t_{1/2} = 1$ are 9.3 percentage points (Architecture~B; $z = 7.4$, $p < 10^{-12}$) and 1.7 percentage points (Architecture~A; $z = 2.0$, $p \approx 0.05$). The qualitative ordering matches Section~3.1's $N = 35$ result. The absolute magnitudes are compressed because $N = 70$ at the prazosin-PTSD calibration places both architectures within 2-5 percentage points of the unit ceiling at $t_{1/2} = 0$, mechanically truncating the loss either architecture's curve can absorb. We note that reading the narrative loss-magnitude claim against the $N = 70$ robustness numbers (3% and 9% relative loss for Architecture~B versus 0.4% and 1.8% for Architecture~A) requires acknowledging this saturation effect; the $N = 35$ production results in Section~3.1 are the canonical reference for the narrative magnitude.

The $N = 70$ run confirms the qualitative ordering; compressed absolute magnitudes reflect ceiling saturation, not a smaller underlying effect.

- Losses: 9.3 pp (Architecture B, $p < 10^{-12}$) vs. 1.7 pp (Architecture A, $p \approx 0.05$); ordering matches $N = 35$.
- At $N = 70$ both architectures are within 2-5 pp of the ceiling at $t_{1/2} = 0$, mechanically truncating observable loss.
- $N = 35$ production results remain the canonical reference for the narrative magnitude claims.

Type I error. Across all 18 null cells of the $N = 35$ production run, type~I error sat in $[0.010, 0.074]$, with mean 0.043 and median 0.042. We observe no architecture-specific inflation. Architecture~B's $t_{1/2} = 1.0$ OL+BDC null ($p = 0.010$, slightly conservative) is the smallest cell; the corresponding alternative ($p = 0.539$) is the most informative cell of the Section~3.1 narrative.

524 **Type I error is well-controlled across all 18 null cells, with no**
525 **architecture-specific inflation.**

- 526 • Type I error range $[0.010, 0.074]$, mean 0.043, median 0.042 across 18
527 null cells at $N = 35$.
- 528 • No architecture-specific inflation detected.
- 529 • The slightly conservative Architecture B OL+BDC null cell at $t_{1/2} =$
530 1.0 ($p = 0.010$) is the only notable departure.

531 **Estimator bias.** Mean estimated $\hat{\beta}_{bm:D_{bc}}$ at $c_{bm} = 0.45$ ranged from -0.231 to
532 -0.281 across the 18 alternative cells, with the $N = 35$ Architecture-A OL+BDC
533 cell at $t_{1/2} = 1.0$ producing the largest absolute mean (-0.281); the Architec-
534 ture-B cells maintained mean near -0.234 across all $t_{1/2}$ values, consistent with
535 the Architecture-B information channel being eroded through correlation decay
536 rather than through coefficient drift. Empirical SE of $\hat{\beta}_{bm:D_{bc}}$ runs 0.05 to 0.10
537 across cells; the within-replicate model SE matches at scale, indicating nominal-
538 95% Wald coverage holds across the grid.

539 **Estimator bias is small and architecturally interpretable; coeffi-**
540 **cient drift under Architecture A contrasts with correlation-channel**
541 **erosion under Architecture B.**

- 542 • Mean $\hat{\beta}_{bm:D_{bc}}$ at $c_{bm} = 0.45$ ranges -0.231 to -0.281 across alterna-
543 tive cells; Architecture A OL+BDC at $t_{1/2} = 1.0$ produces the largest
544 absolute mean.
- 545 • Architecture B cells maintain mean near -0.234 across all $t_{1/2}$ values,
546 consistent with erosion through correlation decay rather than coefficient
547 drift.
- 548 • Nominal-95% Wald coverage holds across the grid.

549

550 5 4. Biological Assumptions

551 The choice between architectures is not merely a statistical convenience; it reflects
552 different assumptions about the biological mechanism by which the biomarker
553 moderates treatment response. We consider each architecture in turn, then take
554 up the prazosin-PTSD motivating application in detail.

555 **Architecture choice encodes a biological claim about how the**
556 **biomarker moderates treatment response.**

- 557 • The two architectures reflect different assumptions about the mecha-
558 nism of biomarker-treatment moderation.
- 559 • Section 4 takes each architecture in turn, then discusses the prazosin-
560 PTSD case.

5.1 4.1 When Architecture A is appropriate

Architecture A (direct mean moderation) is appropriate when the biomarker governs the **magnitude of the drug's biological effect** for each individual participant. The drug effect for participant i is scaled by their biomarker value: $\text{effect}_i = f(B_i) \cdot \text{drug}$. The scaling is deterministic in the population-mean sense, in that any two participants sharing the same B_i have the same expected response, but individual outcomes still carry the usual noise, random-effect, and residual-biological variability. The label 'deterministic' refers to the mean structure of the DGP, not to the realized response of any single participant.

Architecture A is appropriate when the biomarker directly determines the magnitude of the drug effect; the scaling is deterministic at the mean-structure level.

- The drug effect scales with the biomarker value: $\text{effect}_i = f(B_i) \cdot \text{drug}$.
- 'Deterministic' means any two participants with the same B_i have the same expected response; individual outcomes still carry noise.
- Appropriate when the biomarker is a direct PK/PD determinant (eGFR, receptor density, metabolizer status).

Clinical examples:

- **Renal clearance of a renally-excreted drug.** Estimated glomerular filtration rate (eGFR) directly governs the elimination rate of drugs such as digoxin or aminoglycoside antibiotics. Participants with higher eGFR clear the drug more rapidly, achieving lower steady-state concentrations and a proportionally smaller drug effect. The biomarker-drug relationship is mechanistic and deterministic: given the biomarker value, the effective dose at steady state is fully determined (up to random noise).
- **Receptor density moderation.** A biomarker that measures target receptor density (e.g., PET imaging of receptor availability for an antagonist) directly determines the pharmacodynamic response to that antagonist. More receptors means more drug effect, proportionally.
- **Genetic metabolizer status.** CYP2D6 metabolizer phenotype determines the rate of drug activation or clearance. Poor metabolizers of a prodrug receive less active compound, reducing their treatment effect proportionally.
- **Dose-response with biomarker-dependent effective dose.** When the biomarker determines the effective dose (through body composition, protein binding, or volume of distribution), the treatment effect scales with the biomarker through the dose-response curve.

In all these cases the biomarker's moderating effect is preserved under carryover because the residual drug effect during off-drug periods maintains the same proportional relationship with the biomarker. If participant A has twice the drug effect of participant B when on drug, they will also have approximately twice the residual effect during the washout period. That proportionality is precisely what

603 Architecture A encodes.

604 **In all Architecture A cases, the proportional biomarker-drug rela-**
605 **tionship is preserved during carryover washout.**

- 606 • Residual drug effect during washout maintains the same proportional
607 relationship with the biomarker.
- 608 • If participant A has $2\times$ the drug effect of B on-drug, they also have
609 $\approx 2\times$ the residual effect during washout.
- 610 • This preserved proportionality is what Architecture A encodes and why
611 it loses little power under carryover.

612 5.2 4.2 When Architecture B is appropriate

613 Architecture B (differential correlation) is appropriate when the biomarker pre-
614 dicts **which participants are likely to respond**, but the magnitude of any
615 individual's response is not deterministically governed by the biomarker. The
616 biomarker and the drug response are associated at the population level, but the
617 association is mediated by latent factors (disease subtype, neurobiological hetero-
618 geneity, comorbidities) that the biomarker imperfectly indexes.

619 **Architecture B is appropriate when the biomarker is a probabilistic**
620 **predictor of response, mediated by latent factors the biomarker**
621 **imperfectly indexes.**

- 622 • Appropriate when the biomarker predicts which participants are likely
623 to respond, not the magnitude of any individual's response.
- 624 • The biomarker-response association is population-level and probabilistic,
625 mediated by latent factors (disease subtype, neurobiological hetero-
626 geneity, comorbidities).
- 627 • Two participants with identical biomarker values may respond very dif-
628 ferently.

629 Clinical examples:

- 630 • **Blood pressure as a PTSD subtype marker.** Elevated resting blood
631 pressure may index a noradrenergic-predominant PTSD phenotype (Hen-
632 drickson et al. 2020) that is more likely to respond to prazosin (an alpha-1
633 adrenergic antagonist). The blood pressure does not *cause* the drug to
634 work better; rather, it correlates with an underlying neurobiological state
635 that determines drug responsiveness. Two participants with identical blood
636 pressure may have very different responses because blood pressure is an
637 imperfect proxy for the noradrenergic state.
- 638 • **Baseline disease severity as a treatment predictor.** Trials in
639 Alzheimer's disease, depression, and chronic pain often find that patients
640 with more severe baseline symptoms show larger treatment effects (the
641 'floor effect' or 'regression to the mean' phenomenon). The baseline severity
642 score correlates with treatment response, but the relationship is statistical
643 (population-level tendency), not deterministic.
- 644 • **Inflammatory biomarkers in psychiatry.** C-reactive protein (CRP) or

interleukin-6 levels may predict response to anti-inflammatory augmentation of antidepressants (Raison et al. 2013). High-CRP patients are more likely to respond, but the relationship is probabilistic: many high-CRP patients do not respond, and some low-CRP patients do.

- **Genetic risk scores.** Polygenic risk scores predict disease trajectory and treatment response, but explain only a fraction of the variance. The remainder reflects unmeasured genetic, epigenetic, and environmental factors.

In these cases carryover has a different implication: as the drug washes out, the correlation between biomarker and response weakens because the drug-mediated component of response variance diminishes. The biomarker's predictive power is inherently tied to the presence of active drug effect. Without drug, the biomarker is just a biomarker; it does not predict the non-drug component of symptom change.

Under Architecture B, carryover weakens the biomarker-response correlation because the biomarker's predictive power depends on the presence of active drug.

- As the drug washes out, the drug-mediated variance component diminishes, weakening the BM-BR correlation.
- The biomarker's predictive power is tied to active drug effect; without drug, it predicts nothing about non-drug symptom change.
- This is the 'correlation erosion' channel unique to Architecture B.

The latent-subtype framing raises a legitimate modeling question: if the biomarker is really a noisy indicator of an underlying responder class, the biologically faithful DGP would be an explicit finite mixture in which each participant is drawn from one of two (or more) response distributions with class-membership probability a function of B_i . Under such a mixture DGP the individual-level moderation is discrete, not probabilistic, and the MVN differential-correlation parameterization is best understood as a tractable second-moment approximation to a richer mixture model rather than as the generative mechanism itself. The practical consequence is that Architecture~B as implemented here captures the population-level covariance implications of a latent-class mechanism without committing to its full generative form; this is adequate for power calculations that test for the presence of an interaction but should not be read as an endorsement of MVN differential correlation as the correct mechanistic model. Whether the observed power loss under carryover is better mitigated by analysis strategies (Section 6) or by explicit mixture modeling is an open question.

Architecture B is a tractable second-moment approximation to the biologically faithful finite-mixture DGP; adequate for power calculations but not a mechanistic endorsement.

- If the biomarker indexes a discrete responder class, the faithful DGP is a finite mixture; Architecture B approximates its population-level covariance structure.
- Architecture B captures the covariance implications of latent-class membership without committing to the full generative form.

- Adequate for power calculations; should not be read as endorsing MVN differential correlation as the correct mechanism.

The latent-subtype framing of Architecture B should also be adopted with appropriate epistemic caution. Senn (2016) has argued that the apparent ‘personal element in response to treatment’ is often a variance-component artifact rather than a genuine biomarker-by-treatment interaction, and that even modest within-subject variability can produce convincing-looking but spurious differential response. Architecture B presupposes that the latent subtype indexed by the biomarker is real, identifiable, and stable across the trial period. When that presupposition is shaky, the apparent power loss under carryover documented in Section 3 may overstate what could be recovered under any analysis model, because the underlying signal it is attempting to estimate may itself be smaller than assumed.

Senn (2016) cautions that apparent personalized response is often a variance artifact; Architecture B presupposes a real, stable latent subtype, which may not hold.

- Senn (2016) argues that apparent biomarker-by-treatment interactions are often variance-component artifacts, not genuine effects.
- Architecture B presupposes the latent subtype is real, identifiable, and stable across the trial period.
- If that presupposition is shaky, the documented power loss may overstate recoverable signal.

5.3 4.3 The prazosin-PTSD case

We turn briefly to the motivating clinical application for pmsimstats, which is the active-duty prazosin trial by Raskind et al. (2013), whose SBP-stratified secondary analysis is the subject of ongoing reanalysis by the pmsimstats team. The biomarker is resting systolic blood pressure (SBP) and the outcome is CAPS (Clinician-Administered PTSD Scale) score change.

The clinical evidence base for prazosin in PTSD is mixed. The earlier single-site trials by Raskind and colleagues (2003, 2007, 2013) reported benefit on nightmares and global PTSD symptoms in veteran and active-duty samples. The subsequent multisite Veterans Affairs Cooperative Study Program trial (Raskind et al. 2018), known as PACT, did not replicate the effect: prazosin failed to outperform placebo on the primary endpoints. The negative PACT result has been interpreted by several investigators as consistent with effect heterogeneity, namely that prazosin works in some patients but not others, which motivates the search for a predictive biomarker that would identify the responder subgroup. Resting SBP, as a noninvasive proxy for central noradrenergic tone, is the leading candidate.

The prazosin-PTSD evidence base is mixed; the PACT null result motivates the search for a predictive biomarker, with SBP as the leading candidate.

- Earlier single-site Raskind trials (2003, 2007, 2013) reported benefit on nightmares and PTSD symptoms.

- PACT (2018) failed to replicate; this has been interpreted as consistent with effect heterogeneity.
- SBP, as a noninvasive proxy for central noradrenergic tone, is the leading predictive biomarker candidate.

The biological hypothesis is that elevated SBP indexes heightened central noradrenergic tone, which characterizes a subtype of PTSD that is preferentially responsive to alpha-1 adrenergic blockade by prazosin. This is an Architecture B scenario: SBP is a proxy for an underlying neurobiological state, not a direct determinant of drug pharmacokinetics. Two patients with the same SBP may differ in their actual noradrenergic tone, alpha-1 receptor distribution, and downstream signaling, leading to different treatment responses despite identical biomarker values.

The SBP-prazosin mechanism is an Architecture B scenario: SBP is a proxy for noradrenergic state, not a direct PK determinant.

- Elevated SBP indexes heightened noradrenergic tone, a PTSD subtype preferentially responsive to prazosin.
- SBP is a proxy, not a causal determinant; two patients with identical SBP may have very different noradrenergic states.
- This is precisely the Architecture B scenario: probabilistic association, not deterministic scaling.

The Hendrickson et al. (2020) DGP uses Architecture B (MVN differential correlation) for this clinical context. The finding that carryover substantially reduces power under this architecture is therefore clinically relevant: it suggests that trial designs with longer off-drug periods will have meaningfully less power to detect the SBP-prazosin interaction than designs with shorter off-drug periods, and this effect is larger than what a direct mean moderation simulation would predict.

Hendrickson et al. (2020) correctly uses Architecture B for this context; the carryover power loss is clinically relevant for trial design decisions.

- Architecture B (MVN differential correlation) is appropriate for the SBP-prazosin mechanism.
- The documented carryover sensitivity means designs with longer off-drug periods will have meaningfully less power.
- This effect is larger than a direct mean moderation simulation would predict.

6 5. Implications for Simulation Study Design

6.1 5.1 Choosing the appropriate architecture

The choice of DGP architecture should be guided by the biological mechanism hypothesized for the biomarker-treatment interaction:

Criterion	Architecture A	Architecture B
Biomarker role	Causal mediator	Statistical predictor
Mechanism	Deterministic scaling	Probabilistic association
Signal location	Mean structure	Covariance structure
Carryover impact on power	Modest (1-3%)	Design-dependent, up to 31%
Appropriate when	Biomarker determines effective dose or PK	Biomarker indexes a latent subtype

6.2 5.2 Reporting requirements

Simulation studies evaluating biomarker-moderated treatment effects should explicitly state:

1. Whether the biomarker-treatment interaction is generated through mean moderation or differential correlation.
2. The biological rationale for the chosen mechanism.
3. Whether the carryover model (if present) operates on the interaction signal consistently with the chosen architecture.
4. Sensitivity analyses under the alternative architecture, if the biological mechanism is uncertain.

The reporting framework above is consistent with the broader ADEMP template of Morris, White, and Crowther (2019), which prescribes that simulation studies fix five elements before production runs: the **Aims** of the study, the **Data-generating mechanisms**, the **Estimands** (the inferential targets), the **Methods** (analysis specifications compared), and the **Performance measures** (power, bias, coverage, type I error, and Monte Carlo standard errors of each). The DGP-architecture choice this paper concerns is the **D** element of an ADEMP specification; the reporting requirements above are the specialisation of the broader ADEMP framework to the architecture-sensitive class of biomarker-validation simulations.

The four reporting requirements are the ADEMP 'D' element applied specifically to biomarker-validation simulations.

- ADEMP (Morris et al. 2019) prescribes five elements: Aims, Data-generating mechanisms, Estimands, Methods, Performance measures.
- The DGP-architecture choice concerns the D element; the four requirements above specialize that element for architecture-sensitive simulations.

The **primary estimand** for this paper's simulations is the biomarker-by-treatment interaction coefficient $\beta_{bm:D}$ in the linear-mixed analysis model $S_x \sim bm + t + Dbc + bm : Dbc$ with random intercept and AR(1) residual correlation. The reported **performance measures** are: (i) power for $\beta_{bm:D}$ at $\alpha = 0.05$, with Monte Carlo standard error $\sqrt{\hat{\pi}(1 - \hat{\pi})/n_{\text{reps}}}$; (ii) type I

error under $c_{bm} = 0$, with the same MCSE; (iii) bias of $\hat{\beta}_{bm:D}$ relative to the calibrated true value at each parameter cell; (iv) the empirical Monte Carlo standard error of $\hat{\beta}_{bm:D}$ across replicates; and (v) the convergence rate of the linear-mixed fit. Sample size, biomarker effect size, and carryover half-life vary across the parameter grid; the number of replicates per cell is calibrated to achieve MCSE below 0.02 for power at $\hat{\pi} = 0.5$, which corresponds to $n_{\text{reps}} \geq 625$. The simulations reported here use $n_{\text{reps}} = 1000$ throughout.

Primary estimand: $\beta_{bm:D}$ in the LME model; performance measures cover power, type I error, bias, SE, and convergence; 1000 replicates per cell.

- Estimand: $\beta_{bm:D}$ in $Sx \sim bm + t + Dbc + bm:D bc$ with random intercept and AR(1).
- Five performance measures: power, type I error, bias, empirical SE, convergence rate.
- $n_{\text{reps}} = 1000$ throughout; calibrated to achieve MCSE < 0.02 for power near 0.5.

6.3 5.3 When the choice is uncertain

If the biological mechanism is ambiguous, as it often is in early biomarker development, we suggest investigators take the following steps:

1. Run the simulation under both architectures and report the range of power estimates.
2. Architecture B produces more conservative (lower) power estimates under carryover, making it the safer default for trial design decisions.
3. Design the trial to minimize carryover (adequate washout periods) to reduce sensitivity to the DGP architecture choice.

7 6. Alternative Analysis Strategies Under Architecture B

The power loss documented in Section 3 arises from two distinct sources: (1) compressed variance of the treatment indicator Dbc when carryover is present, and (2) erosion of the BM-BR correlation during off-drug periods. Source 1 is a property of the analysis model parameterization and can potentially be addressed by choosing a different predictor. Source 2 is a property of the data itself (the signal is genuinely weaker in off-drug observations), and no analysis model can recover signal that is not present. However, analysis strategies that are selective about *which observations contribute to the test* or *how they are weighted* can mitigate the damage.

7.1 6.1 Restrict analysis to on-drug observations

The most direct approach is to discard off-drug timepoints entirely and test the biomarker-treatment interaction using only on-drug observations. Every retained observation has the full BM-BR correlation (c_{bm}), so correlation erosion is eliminated. The costs are twofold: a reduction in effective sample size (fewer observations per participant) and, more importantly, the loss of the off-drug reference that separates the biological response from placebo and time-varying background response. Without off-drug observations the biological, placebo, and time-varying components are confounded, which is why the Open-Label design, despite having the most on-drug timepoints, does not dominate the other designs in the Section 3 results. For designs with many on-drug timepoints (OL has 8, the Hybrid design has 5-6 per path), the net effect may be positive: the per-observation signal-to-noise ratio is maximal, and the analysis model simplifies to $Sx \sim \mathbf{bm} + \mathbf{t} + \mathbf{bm:t}$ since there is no treatment contrast to model. The interaction is detected through the correlation between biomarker and the time-on-drug trajectory.

This approach is most promising for the Open-Label design (which has no off-drug observations regardless) and the Hybrid design (which has a long initial on-drug phase). It is least useful for the Crossover design, where discarding half the observations substantially reduces statistical power.

Restricting to on-drug observations eliminates correlation erosion but loses the off-drug reference; best for OL and Hybrid designs, worst for CO.

- Benefit: every retained observation has full c_{bm} ; correlation erosion is eliminated.
- Cost: loss of off-drug reference confounds biological, placebo, and time-varying components.
- Practical: best for OL/Hybrid (many on-drug timepoints); least useful for CO (discards half the observations).

7.2 6.2 Weighted analysis

Rather than discarding off-drug observations entirely, a weighted analysis retains all observations but weights them by their estimated drug exposure level. On-drug observations receive weight 1. Off-drug observations receive weight proportional to the expected residual drug effect:

$$w_t = e^{-\lambda \cdot t_{sd}}$$

This downweights observations where the BM-BR correlation has decayed (and where they contribute more noise than signal to the interaction estimate) without discarding them entirely. The rationale is information-theoretic: observations with higher drug exposure carry more information about the biomarker-treatment interaction because the correlation signal is stronger. This is straightforward to implement in `nlme::lme` using the `weights` argument.

880 Weighted analysis downweights off-drug observations by drug ex-
 881 posure without discarding them; information-theoretic rationale;
 882 straightforward to implement.

- 883 • Weight: $w_t = e^{-\lambda t_{sd}}$ for off-drug observations; 1 for on-drug.
- 884 • Rationale: higher exposure = stronger correlation signal = more infor-
 885 mation about the interaction.
- 886 • Implemented via the `weights` argument in `nlme::lme`.

887 7.3 6.3 Within-subject contrast

888 For designs with both on-drug and off-drug periods (CO, Hybrid, OL+BDC), a
 889 summary-measure approach computes a per-participant contrast: the difference
 890 between mean on-drug response and mean off-drug response. The biomarker-
 891 treatment interaction is then tested by a simple regression of these contrasts on
 892 the biomarker:

$$\bar{Y}_{i,\text{on}} - \bar{Y}_{i,\text{off}} = \alpha + \gamma \cdot B_i + \epsilon_i$$

893 This sidesteps the Dbc parameterization and the repeated-measures model entirely.
 894 Carryover affects the magnitude of the contrast (making off-drug means closer
 895 to on-drug means), but the *correlation between the contrast and the biomarker*
 896 may be more robust because within-subject averaging reduces noise before the
 897 biomarker association is tested. The approach trades the efficiency of a full longi-
 898 tudinal model for robustness to misspecification of the carryover and correlation
 899 structure.

900 Within-subject contrast sidesteps the repeated-measures model;
 901 trades efficiency for robustness to carryover misspecification.

- 902 • Per-participant contrast $\bar{Y}_{\text{on}} - \bar{Y}_{\text{off}}$ is regressed on the biomarker.
- 903 • Carryover compresses the contrast magnitude but the contrast-
 904 biomarker correlation may be more robust.
- 905 • Trade-off: less efficient than a full longitudinal model, but more robust
 906 to carryover and correlation misspecification.

907 7.4 6.4 Two-stage random slopes

908 A two-stage approach first estimates participant-specific treatment effects, then
 909 tests their association with the biomarker. Stage 1 fits a random-slopes model:

$$Y_{it} = \beta_0 + \beta_1 t + \beta_2 D_{it} + u_{0i} + u_{1i} D_{it} + \epsilon_{it}$$

910 where β_2 is the population mean drug effect and u_{1i} is participant i 's random
 911 deviation from that mean. The participant-specific drug effect is then $\beta_2 + u_{1i}$.
 912 Stage 2 regresses the estimated random deviations (or equivalently, the BLUPs of
 913 the participant-specific effects) on the biomarker:

$$\hat{u}_{1i} = \delta_0 + \delta_1 B_i + \eta_i$$

This separates the within-subject treatment effect estimation (which uses the longitudinal data and is affected by carryover) from the between-subject biomarker association (which uses the cross-sectional relationship between estimated effects and biomarker values). The approach may be more robust to carryover because the random slope u_{1i} integrates over the participant's entire trajectory, including both on-drug and off-drug phases.

Two-stage random slopes separates within-subject effect estimation from between-subject biomarker association; may be more robust to carryover.

- Stage 1: random-slopes LME estimates participant-specific drug effects ($\beta_2 + u_{1i}$).
- Stage 2: regress BLUPs on biomarker to test association.
- The random slope u_{1i} integrates over the full trajectory; robustness to carryover is plausible but unverified by simulation.

7.5 6.5 Exclusion of contaminated observations

The first 1-2 off-drug timepoints carry the most carryover contamination: the residual drug effect is highest, and the BM-BR correlation is in the process of decaying. Later off-drug timepoints, where the drug has fully washed out, provide a cleaner contrast with the on-drug phase. Excluding the early off-drug observations and retaining only the later ones yields a comparison between fully-on and fully-off states, avoiding the intermediate zone where the analysis model is most strained.

Excluding the first 1–2 off-drug timepoints removes the most contaminated observations; later off-drug timepoints provide a cleaner contrast.

- Early off-drug timepoints: highest residual drug effect, BM-BR correlation actively decaying.
- Later off-drug timepoints: fully washed out, cleaner contrast with on-drug phase.
- Result: comparison between fully-on and fully-off states, avoiding the strained intermediate zone.

7.6 6.6 Design-level solutions

Rather than adapting the analysis model to accommodate carryover, the trial design itself can be modified to reduce carryover exposure:

- **Longer washout periods** between drug phases eliminate residual drug effect before off-drug measurements begin.
- **More on-drug timepoints** relative to off-drug increase the proportion of high-signal observations.

952 • **Front-loaded drug exposure** (as in the OL+BDC design) places a
 953 long initial on-drug block before the off-drug baseline-during-continued-
 954 treatment phase, so that the BM-BR correlation has been fully established
 955 before any off-drug observations are collected. The alternative ordering
 956 (off-drug first, then on-drug) has the appealing property of placing the
 957 placebo and time-varying estimation in an uncontaminated window but
 958 sacrifices on-drug exposure duration and forces the BM-BR signal to be
 959 re-established after a wash-in; the net power trade-off depends on the ratio
 960 of carryover half-life to total trial length and is explored empirically in the
 961 Section 3 results.

962 Design-level solutions address the root cause (carryover in the data) rather than
 963 the symptom (power loss in the analysis). Their cost is typically a longer trial
 964 duration or reduced ability to estimate the carryover dynamics themselves.

965 Design-level changes address the root cause rather than the symp-
 966 tom; cost is longer trial or reduced ability to estimate carryover
 967 dynamics.

- 968 • Three levers: longer washout, more on-drug timepoints, front-loaded
 969 drug exposure (as in OL+BDC).
- 970 • Front-loading establishes BM-BR correlation before any off-drug obser-
 971 vations; alternative ordering (off-drug first) trades clean placebo esti-
 972 mation for reduced on-drug exposure.
- 973 • Addresses root cause; cost is typically longer trial or reduced carryover
 974 estimation ability.

975 7.7 6.7 Comparative evaluation

976 The strategies described above differ in their assumptions, implementation com-
 977 plexity, and expected power recovery:

Strategy	Discards data	Complexity	Expected benefit
On-drug only	Yes	Low	High for OL/Hybrid; poor for CO
Weighted	No	Low	Moderate; depends on weight calibration
Within-subject contrast	Aggregates	Low	Moderate; robust to misspecification
Two-stage random slopes	No	Moderate	Unknown; separates estimation stages
Exclude early off-drug	Some	Low	Moderate; improves contrast clarity
Design modification	N/A	N/A	High; addresses root cause

978 A systematic simulation comparison of these strategies under Architecture B, anal-
 979 ogous to the factorial comparison of analysis models in Section 3, would identify
 980 which approach best recovers the power lost to carryover for each trial design. We
 981 regard this as a direction for future work.

8 7. Relationship to the Broader Literature

8.1 7.1 Treatment effect heterogeneity

The statistical literature on treatment effect heterogeneity is extensive. Rothwell (2005) sets out the foundational distinction between *predictive* biomarkers (those that modify the treatment effect) and *prognostic* biomarkers (those that predict outcome regardless of treatment). The current consensus statement on predictive HTE analysis is the Predictive Approaches to Treatment effect Heterogeneity (PATH) Statement of Kent et al. (2020), which distinguishes two analytic strategies for identifying predictive biomarker effects in randomized trials: a *risk modeling* approach, in which a multivariable risk model (developed without a treatment term) is used to stratify patients and examine variation in absolute treatment benefit across risk strata, and an *effect modeling* approach, in which the regression includes explicit treatment-by-covariate interaction terms. PATH expresses caution about data-driven effect modeling on the grounds of low statistical power, multiplicity, and overfitting, and favors risk modeling when the application criteria are met.

The PATH Statement distinguishes risk modeling from effect modeling for predictive HTE analysis; this is an analysis-model distinction, not a DGP distinction.

- Rothwell (2005) establishes the predictive/prognostic biomarker distinction; Kent et al. (2020) synthesize it into the PATH Statement.
- PATH's risk modeling approach stratifies on a multivariable risk score; effect modeling adds explicit interaction terms.
- PATH favors risk modeling and cautions against data-driven effect modeling because of low power, multiplicity, and overfitting.

We note that the PATH risk-versus-effect-modeling distinction is an analysis-model axis. The Architecture A versus Architecture B distinction introduced in this paper is a data-generating-process axis, and the two are orthogonal. Both PATH analytic strategies implicitly assume that the underlying treatment effect heterogeneity is encoded in the mean structure of the outcome (the heterogeneity is identifiable from differences in conditional means), which corresponds to Architecture A as a DGP. Architecture B as a DGP is a separate phenomenon that does not map cleanly onto either of PATH's analysis-model categories: the interaction signal resides in the covariance structure rather than the mean, and the standard regression-based predictive HTE analysis methods endorsed by PATH may have reduced power under Architecture B precisely because they were developed under an implicit Architecture A assumption. Architecture A can accordingly be characterized as a *parametric interaction model* (in PATH's effect modeling sense).

Architecture B can be characterized as a *latent class* or *mixture model* perspective in which the biomarker is a noisy indicator of class membership.

PATH's analysis-model axis (risk vs. effect modeling) is orthogonal to our DGP axis; both PATH strategies implicitly assume Architecture A.

- PATH's distinction is analysis-model; ours is DGP; they are orthogonal.
- Both PATH strategies assume heterogeneity is in the mean structure (Architecture A); Architecture B does not fit either PATH category.
- Architecture A = parametric interaction model; Architecture B = latent-class/mixture-model perspective.

8.2 7.2 Prevalence of each architecture

A survey of the simulation methodology literature reveals that Architecture A (direct mean moderation) is the near-universal standard. Simon (2010) employs it throughout the simulation frameworks he develops for biomarker-stratified designs; Freidlin and Korn (2014) adopt it for adaptive enrichment trials; Renfro and Sargent (2017) use it for basket and umbrella trials; and it underlies the platform trial literature throughout. Haller and Ulm (2018), in simulation work on biomarker-by-treatment interaction estimation in randomized trials with prognostic covariates, likewise adopt the standard regression specification, focusing on power and bias of the interaction estimator under different covariate adjustment strategies. N-of-1 simulation studies, including Zucker et al. (1997), Araujo et al. (2016), and Duan et al. (2013), use hierarchical random-effects models in which individual treatment effects are drawn from a population distribution, a variant of Architecture A in which the biomarker predicts the random treatment effect through the mean structure. Methodological work on power and design for N-of-1 trials with serial correlation and carryover, from both Wang and Schork (2019) and Schork (2022), adopts the same Architecture A specification, even though this work shares senior authorship with Hendrickson et al. (2020). It appears, then, that the differential- correlation parameterization in the latter is a paper-specific design choice rather than a methodological convention of any laboratory or research program.

Architecture A is the near-universal standard across the biomarker-moderated trial simulation literature; Architecture B is absent from all major threads except Hendrickson et al.

- Simon (2010), Freidlin and Korn (2014), Renfro and Sargent (2017), Haller and Ulm (2018), and the platform trial literature all use Architecture A.
- N-of-1 simulation studies (Zucker et al. 1997; Araujo et al. 2016; Duan et al. 2013) use hierarchical random-effects models, a variant of Architecture A.
- Wang and Schork (2019) and Schork (2022) share authorship with Hendrickson et al. yet adopt Architecture A, suggesting the differential-correlation choice is paper-specific.

Architecture B (differential correlation via MVN) appears primarily in latent variable and structural equation modeling contexts, including the joint modeling framework discussed by Rizopoulos (2012), Bayesian joint modeling frameworks, and copula-based simulation studies. In the clinical trial simulation literature specifically, the Hendrickson et al. (2020) framework appears to be unique in using differential correlation as the mechanism for biomarker-treatment interaction in an N-of-1 trial context.

Architecture B appears in latent variable / SEM / copula contexts but is unique in N-of-1 clinical trial simulation.

- Architecture B appears in latent variable / SEM / Bayesian joint modeling / copula contexts (e.g., Rizopoulos 2012).
- In N-of-1 clinical trial simulation, Hendrickson et al. (2020) appears unique in using it.

This predominance of Architecture A has an important implication: the power estimates reported in the vast majority of biomarker-moderated trial simulation studies may be optimistic for biomarkers that operate through an Architecture B mechanism, because these studies do not account for the additional power loss that carryover produces when the interaction signal resides in the covariance structure rather than the mean structure.

The near-universal use of Architecture A means most published power estimates are optimistic for Architecture B biomarkers under carryover.

- Virtually all published power estimates for biomarker-moderated trials assume Architecture A.
- These estimates may be optimistic for biomarkers that operate through Architecture B (covariance-structure mechanism).
- The additional carryover-induced power loss under Architecture B is not accounted for.

8.3 7.3 Crossover and N-of-1 trial methodology

The crossover trial literature, represented most thoroughly by Senn (2002) and Jones and Kenward (2014), addresses carryover extensively but does not distinguish between architectures for biomarker-treatment interaction modeling. More recent methodological work by Sturdevant and Lumley (2021) on testing for carryover effects via mixed-effects models likewise treats carryover as a nuisance parameter to be estimated and adjusted for, without engaging with how carryover interacts with biomarker-by-treatment effect signals. The standard recommendation, namely to always include carryover in the analysis model or to use adequate washout, applies to both architectures but has different power implications under each. The finding that Architecture B is more sensitive to carryover suggests that crossover designs with short washout periods may be less suitable for detecting correlation-based biomarker interactions than previously recognized.

The crossover trial literature addresses carryover as a nuisance parameter but ignores the architectural distinction; the standard

recommendation applies differently under each.

- Senn (2002), Jones and Kenward (2014), and Sturdevant and Lumley (2021) treat carryover as a nuisance parameter; none engages with DGP architecture.
- Standard recommendation (include carryover in model or use adequate washout) applies to both architectures.
- Under Architecture B, crossover designs with short washout periods may be less suitable than previously recognized.

8.4 7.4 Precision medicine trial design

Enrichment and biomarker-stratified designs, as discussed by Simon (2010) and Freidlin and Korn (2014), use Architecture A exclusively for power calculations. Three families of design are particularly worth considering in light of the architectural distinction.

Enrichment and biomarker-stratified designs assume Architecture A for power calculations; the architectural choice matters differently across three design families.

- Simon (2010) and Freidlin and Korn (2014) frame power calculations in the Architecture A regression idiom.
- Three design families warrant separate consideration: adaptive enrichment, treatment-switching, and basket/umbrella trials.
- The architectural assumption is rarely made explicit, so the practical implications are easy to overlook.

Adaptive enrichment trials adjust the eligible population mid-trial based on emerging evidence of a biomarker-treatment interaction. The interim power calculations that drive these adaptive decisions are universally framed in the Architecture A regression idiom. If the true biomarker mechanism is closer to Architecture B and any participants pass through washout intervals between dose-finding stages, the realized power for the interaction test will be lower than the adaptation rule expects, biasing decisions toward premature futility stopping or toward overly narrow enrichment criteria. The discrepancy is most acute when the biomarker mechanism is uncertain at the design stage, which is precisely the situation in which adaptive enrichment is most attractive.

Adaptive enrichment trials are vulnerable to Architecture B power shortfalls precisely when the biomarker mechanism is most uncertain.

- Interim power calculations driving adaptive decisions use Architecture A; if Architecture B holds and washout occurs, realized power will be lower than predicted.
- The mismatch biases adaptation rules toward premature futility stopping or overly narrow enrichment criteria.
- The problem is worst when the biomarker mechanism is unknown, exactly when adaptive enrichment is most often proposed.

Treatment-switching designs (most crossover, basket, and platform trials, as

well as N-of-1 hybrid designs) cycle individual participants through different drug states. Under Architecture B, each washout phase erodes the BM-BR correlation, so the biomarker-treatment interaction signal in the second and subsequent treatment periods is structurally weaker than in the first. The power loss documented in Section 3 applies here directly. By contrast, parallel-group enrichment designs, in which each participant remains on a single treatment for the duration of the trial, escape this loss because there is no off-drug phase during which the correlation can decay; for these designs, Architectures A and B yield similar power estimates. The architectural choice therefore matters most in exactly the designs where the biomarker mechanism is most likely to be of clinical interest, namely those that deliberately expose each participant to both treatment states in order to estimate within-subject responsiveness.

Treatment-switching designs suffer Architecture B power loss through washout erosion; parallel-group enrichment designs do not, because there is no off-drug phase.

- Crossover, basket, platform, and N-of-1 hybrid designs all cycle participants through drug states; each washout phase erodes BM-BR correlation under Architecture B.
- Parallel-group enrichment designs escape this loss because each participant stays on a single treatment throughout.
- The architectural choice matters most in within-subject designs, where it is also most likely to be clinically important.

Basket and umbrella trials, as discussed by Renfro and Sargent (2017), introduce an additional layer because their power calculations average across multiple biomarker-defined arms. If even a subset of the arms operate through an Architecture B mechanism while the trial-wide power calculation assumes Architecture A throughout, the aggregated power estimates will be biased optimistically in a way that is difficult to detect from trial-level results alone. Subgroup-specific power diagnostics under both architectures should accompany the primary planning analysis whenever the underlying biology is heterogeneous across baskets.

Basket and umbrella trials face an additional risk of optimistic aggregated power estimates when Architecture B holds in even a subset of arms.

- Renfro and Sargent (2017) discuss power calculations that average across biomarker-defined arms; the calculation assumes Architecture A throughout.
- Mixed architectures across baskets produce optimistic trial-level power estimates that are difficult to detect from aggregate results.
- Subgroup-specific power diagnostics under both architectures are warranted when the biological mechanism is heterogeneous.

In all three settings the practical recommendation parallels that of Section 5.3: when the biological mechanism is uncertain, run the power calculation under both architectures and treat the Architecture B estimate as the more conservative anchor for sample size planning. This recommendation is consistent with the

PATH Statement (Kent et al. 2020), which itself counsels caution about data-driven discovery of treatment-by-covariate interactions and prefers risk modeling approaches when the application criteria are met.

Across all three design families, the practical recommendation is to anchor sample size planning on the Architecture B estimate when the biomarker mechanism is uncertain.

- Run power calculations under both architectures and treat Architecture B as the conservative anchor.
- This recommendation parallels the advice in Section 5.3 and is consistent with the PATH Statement (Kent et al. 2020).
- The PATH Statement’s caution about data-driven interaction discovery reinforces the case for conservative planning under architectural uncertainty.

8. Conclusions

In conclusion, six points are to be emphasized.

1. The choice between direct mean moderation (Architecture A) and differential correlation (Architecture B) for generating biomarker-treatment interactions in clinical trial simulations has substantive consequences for power estimates under carryover.
2. Under Architecture A, carryover reduces power only modestly, by 1 to 3% in relative terms across the three designs examined, because the proportional biomarker-drug relationship is preserved during washout. Under Architecture B the loss is strongly design-dependent: it remains modest in the crossover and hybrid designs (3 to 5%) but reaches roughly 31% in the open-label OL+BDC design, because the differential-correlation signal erodes as the drug effect wanes.
3. The choice should be guided by the hypothesized biological mechanism: Architecture A for biomarkers that directly determine drug pharmacokinetics or pharmacodynamics; Architecture B for biomarkers that statistically predict treatment response through latent subtype membership.
4. For the prazosin-PTSD application with blood pressure as the predictive biomarker, Architecture B (MVN differential correlation) is the more appropriate model, and the resulting power sensitivity to carryover should inform trial design decisions regarding off-drug period duration.
5. The existing clinical trial simulation literature uses Architecture A almost exclusively. The Hendrickson et al. (2020) MVN approach is, to the best

of our knowledge, unique in the N-of-1 trial context. This means that published power estimates for biomarker-moderated crossover and N-of-1 designs may be systematically optimistic for biomarkers that operate through a correlation-based (Architecture B) mechanism, because they do not account for the carryover sensitivity documented here.

6. Simulation studies for biomarker-moderated treatment effects should explicitly report their DGP architecture and consider sensitivity analyses under the alternative when the biological mechanism is uncertain.

10 Conflict of interest

The authors declare no conflicts of interest.

11 Funding

This work received no specific funding.

12 Data availability statement

The data and code that support the findings of this study are openly available in the pmsimstats-ng project repository. Per-replicate Monte Carlo outputs and ADEMP-compliant cell-level summaries are versioned under `analysis/data/`; simulation drivers are at `analysis/scripts/`. Exact paths for this manuscript are listed in the Reproducibility section below.

13 Reproducibility

This manuscript was rendered with R 4.4.x inside the project's `zzcollab` Docker container (image hash recorded in `Dockerfile`; package set captured in `renv.lock`). Principal packages used by the simulation pipeline are `nlme` (continuous-time linear-mixed analysis with `corCAR1` residual correlation), `parallel` (8-core `mclapply` for parameter-grid sweeps), `data.table` (the original-extended simulation collection), `furrr` and `purrr` (the `tidyverse` simulation collection), and `rmarkdown` plus `bookdown` for the manuscript build.

Random-number generator. The simulation driver sets `set.seed(20260507)` at the top of the replicate loop and uses `parallel::mc.set.seed` inside `mclapply`. Per-replicate seeds derive from the master seed by `+ rep_idx`.

Data and code availability. Per-replicate simulation outputs are checkpointed at `analysis/data/quick-sim/01-dgp-replicates.rds`; Morris ADEMP cell-level summaries are at the companion `01-dgp-summary.txt`. The simulation driver is at `analysis/scripts/quick-sim/01-dgp-prototype.R`. The manuscript source, bibliography, and SIM CSL file are at `analysis/report/01-dgp-mean-moderation-vs-mvn/`

14 References

- Araujo A, Julious S, Senn S. Understanding variation in sets of N-of-1 trials. *PLoS ONE*. 2016;11(12):e0167167.
- Duan N, Kravitz RL, Schmid CH. Single-patient (N-of-1) trials: a pragmatic clinical decision methodology. *J Clin Epidemiol*. 2013;66(8):S21-S28.
- Freidlin B, Korn EL. Biomarker enrichment strategies: matching trial design to biomarker credentials. *Nat Rev Clin Oncol*. 2014;11(2):81-90.
- Haller B, Ulm K. A simulation study on estimating biomarker-treatment interaction effects in randomized trials with prognostic variables. *Trials*. 2018;19(1):128.
- Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: statistical methods and practical considerations. *Front Digit Health*. 2020;2:13.
- Jones B, Kenward MG. *Design and Analysis of Cross-Over Trials*. 3rd ed. Boca Raton, FL: CRC Press; 2014.
- Kent DM, Paulus JK, van Klaveren D, et al. The Predictive Approaches to Treatment effect Heterogeneity (PATH) Statement. *Ann Intern Med*. 2020;172(1):35-45.
- Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Stat Med*. 2019;38(11):2074-2102.
- Raison CL, Rutherford RE, Woolwine BJ, et al. A randomized controlled trial of the tumor necrosis factor antagonist infliximab for treatment-resistant depression: the role of baseline inflammatory biomarkers. *JAMA Psychiatry*. 2013;70(1):31-41.
- Raskind MA, Peskind ER, Kanter ED, et al. Reduction of nightmares and other PTSD symptoms in combat veterans by prazosin: a placebo-controlled study. *Am J Psychiatry*. 2003;160(2):371-373.
- Raskind MA, Peskind ER, Hoff DJ, et al. A parallel group placebo-controlled study of prazosin for trauma nightmares and sleep disturbance in combat veterans with post-traumatic stress disorder. *Biol Psychiatry*. 2007;61(8):928-934.

1298 Raskind MA, Peterson K, Williams T, et al. A trial of prazosin for combat trauma
1299 PTSD with nightmares in active-duty soldiers returned from Iraq and Afghanistan.
1300 *Am J Psychiatry*. 2013;170(9):1003-1010.

1301 Raskind MA, Peskind ER, Chow B, et al. Trial of prazosin for post-traumatic
1302 stress disorder in military veterans. *N Engl J Med*. 2018;378(6):507-517.

1303 Renfro LA, Sargent DJ. Statistical controversies in clinical research: basket trials,
1304 umbrella trials, and other master protocols. *Ann Oncol*. 2017;28(1):34-43.

1305 Rizopoulos D. *Joint Models for Longitudinal and Time-to-Event Data: With Ap-*
1306 *plications in R*. Boca Raton, FL: CRC Press; 2012.

1307 Rothwell PM. Treating individuals 5. Subgroup analysis in randomised controlled
1308 trials: importance, indications, and interpretation. *Lancet*. 2005;365(9454):176-
1309 186.

1310 Schork NJ. Accommodating serial correlation and sequential design elements in
1311 personalized studies and aggregated personalized studies. *Harv Data Sci Rev*.
1312 2022;Special Issue 3.

1313 Senn S. *Cross-over Trials in Clinical Research*. 2nd ed. Chichester, UK: John
1314 Wiley & Sons; 2002.

1315 Senn S. Mastering variation: variance components and personalised medicine.
1316 *Stat Med*. 2016;35(7):966-977.

1317 Simon R. Clinical trial designs for evaluating the medical utility of prognostic and
1318 predictive biomarkers in oncology. *Per Med*. 2010;7(1):33-47.

1319 Sturdevant SG, Lumley T. Statistical methods for testing carryover effects: a
1320 mixed effects model approach. *Contemp Clin Trials Commun*. 2021;22:100711.

1321 Varadhan R, Segal JB, Boyd CM, Wu AW, Weiss CO. A framework for the analysis
1322 of heterogeneity of treatment effect in patient-centered outcomes research. *J Clin*
1323 *Epidemiol*. 2013;66(8):818-825.

1324 Wang Y, Schork NJ. Power and design issues in crossover-based N-of-1 clinical
1325 trials with fixed data collection periods. *Healthcare (Basel)*. 2019;7(3):84.

1326 Zucker DR, Schmid CH, McIntosh MW, D'Agostino RB, Selker HP, Lau J.
1327 Combining single patient (N-of-1) trials to estimate population treatment effects
1328 and to evaluate individual patient responses to treatment. *J Clin Epidemiol*.
1329 1997;50(4):401-410.

1330 15 References

1331

1332 *Rendered on 2026-06-10 at 08:50 PDT.*

1333 *Source: ~/prj/res/36-pmsimstats-ng/pmsimstats-ng/analysis/report/01-dgp-mean-moderation-vs-mv*

2026-06-10 08:50 PDT

35

~/prj/res/36-pmsimstats-ng/pmsimstats-ng/analysis/report/
01-dgp-mean-moderation-vs-mvn/report.Rmd