

1 Two Architectures for Simulating Biomarker-Treatment
2 Interactions: Implications for Statistical Power Under
3 Carryover

4 pmsimstats team

5 2026-06-10

6 **Contents**

7	1 Abstract	2
8	2 1. Introduction	3
9	3 2. Two DGP Architectures	6
10	3.1 2.1 Architecture A: Direct Mean Moderation	7
11	3.2 2.2 Architecture B: Differential Correlation (MVN)	8
12	3.3 2.3 Mathematical Comparison	9
13	4 3. Consequences for Statistical Power Under Carryover	11
14	4.1 3.1 Empirical demonstration	11
15	4.2 3.2 Why the architectures diverge	14
16	4.3 3.3 Robustness check at $N = 70$	16
17	5 4. Biological Assumptions	18
18	5.1 4.1 When Architecture A is appropriate	18
19	5.2 4.2 When Architecture B is appropriate	20
20	5.3 4.3 The prazosin-PTSD case	23
21	6 5. Implications for Simulation Study Design	25
22	6.1 5.1 Choosing the appropriate architecture	25
23	6.2 5.2 Reporting requirements	25
24	6.3 5.3 When the choice is uncertain	27
25	7 6. Alternative Analysis Strategies Under Architecture B	27
26	7.1 6.1 Restrict analysis to on-drug observations	27
27	7.2 6.2 Weighted analysis	28
28	7.3 6.3 Within-subject contrast	29
29	7.4 6.4 Two-stage random slopes	29
30	7.5 6.5 Exclusion of contaminated observations	30
31	7.6 6.6 Design-level solutions	31

32	7.7 6.7 Comparative evaluation	32
33	8 7. Relationship to the Broader Literature	32
34	8.1 7.1 Treatment effect heterogeneity	32
35	8.2 7.2 Prevalence of each architecture	33
36	8.3 7.3 Crossover and N-of-1 trial methodology	35
37	8.4 7.4 Precision medicine trial design	36
38	9 8. Conclusions	39
39	10 Conflict of interest	40
40	11 Funding	40
41	12 Data availability statement	40
42	13 Reproducibility	40
43	14 References	41
44	15 References	42

45 1 Abstract

46 **Background.** Simulation studies of predictive biomarker-treatment interactions
47 in N-of-1 clinical trials require a data-generating process (DGP) that specifies how
48 the biomarker moderates treatment response. Two fundamentally different DGP
49 architectures are used in practice, and the choice between them determines the
50 extent to which carryover effects reduce statistical power to detect the interaction.
51 The published simulation literature does not distinguish them.

52 **Methods.** We formalise two architectures: direct mean moderation (Architec-
53 ture~A), under which the biomarker scales the treatment effect in the popula-
54 tion mean structure; and differential correlation (Architecture~B, the multivari-
55 ate normal approach), under which the interaction emerges from treatment-state-
56 dependent correlation between biomarker and response in the joint distribution.
57 Power consequences under both architectures are quantified by Monte Carlo sim-
58 ulation on an aggregated N-of-1 design framework with continuous-time AR(1)
59 residual correlation in the analysis model. The biomarker, treatment phase, and
60 carryover decay are parameterised to the prazosin/PTSD reference application.

61 **Results.** Under Architecture~A, power is largely preserved under carryover be-
62 cause the proportional relationship between biomarker and drug response is main-
63 tained at every timepoint. Under Architecture~B, carryover erodes the differen-
64 tial correlation signal and reduces power. At the prazosin-calibrated reference cell
65 ($N = 35$, $c_{bm} = 0.45$, OL+BDC design) across $t_{1/2} \in \{0, 0.5, 1.0\}$ weeks, Architec-
66 ture~B power drops from 0.777 to 0.539 (a 30.6% relative loss; $z = 7.64$, $p < 10^{-13}$
67 on the B-versus-A contrast), while Architecture~A drops only from 0.751 to 0.730

(a 2.8% relative loss). The B-versus-A absolute power-loss gap of 21.7 percentage points at OL+BDC, $t_{1/2} = 1.0$ is approximately 11 times the matched Architecture~A loss; the Hybrid design produces a smaller but still significant gap ($z = 3.96$, $p < 10^{-4}$); the traditional crossover design produces no detectable gap because its within-subject AR(1) structure dominates the carryover-mediated correlation channel that Architecture~B uses. The 36-cell production factorial (2 architectures $\times$ 3 designs $\times$ 3 $t_{1/2}$ $\times$ 2 c_{bm} levels) was run at 1{,}000 replicates per cell with 100% convergence; MCSE on cell power is 0.013-0.016. Type-I error sits in [0.010, 0.074], with no architecture-specific inflation.

Conclusions. The two architectures encode different biological mechanisms, heterogeneity at the mean structure versus heterogeneity at the variance structure, and the choice between them is a substantive scientific question rather than a parameterisation convenience. We recommend explicit declaration of the DGP architecture in any biomarker-validation simulation study and provide guidance for selecting the architecture appropriate to a given clinical context.

2 1. Introduction

The central question is biomarker-treatment interaction; simulation studies need a DGP that encodes it.

- Does a baseline biomarker predict differential treatment response, i.e., does it moderate the treatment effect?
- Simulation studies are the principal power evaluation tool for proposed trial designs.
- These simulations require a DGP that explicitly encodes the biomarker-treatment interaction mechanism.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

A central question in precision medicine trial design is whether a baseline biomarker predicts differential treatment response, that is, whether the treatment effect is moderated by the biomarker value. Simulation studies are the principal tool available for evaluating the statistical power of proposed trial designs to detect such interactions, and these simulations necessarily require a data-generating process (DGP) that encodes the biomarker-treatment interaction mechanism.

Every simulation framework must choose how to generate treatment effect heterogeneity, and that choice has consequences.

- HTE literature (Rothwell 2005; Kent et al. 2020) distinguishes qualitative interactions (direction reverses) from quantitative ones (magnitude varies).
- Any simulation framework must decide how to encode this heterogeneity in the DGP.
- That decision is substantive, not a mere parameterization detail.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The statistical literature on treatment effect heterogeneity, developed notably by Rothwell (2005) and synthesized more recently by Kent et al. (2020), distinguishes between qualitative interactions, in which the direction of the treatment effect reverses across biomarker strata, and quantitative interactions, in which the magnitude varies but the direction remains consistent. As we shall argue, any simulation framework intended to evaluate power for biomarker-moderated effects must choose how to generate this heterogeneity, and that choice is not merely a technical convenience.

Architecture A (direct mean moderation) is the near-universal standard; the Hendrickson (2020) MVN approach is apparently unique.

- Every major simulation framework uses direct mean moderation: Simon (2010) for stratified designs; Freidlin and Korn (2014) for adaptive enrichment; Renfro and Sargent (2017) for basket/umbrella trials; all N-of-1 frameworks.
- Wang and Schork (2019) and Schork (2022) also use Architecture A despite sharing authorship with Hendrickson et al. (2020), suggesting the MVN parameterization is a paper-specific choice, not a lab convention.
- To the best of our knowledge, no other clinical trial simulation framework uses differential correlation (MVN) to encode the biomarker-treatment interaction.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The dominant approach in the clinical trial simulation literature is direct mean moderation, in which the biomarker enters the outcome model as an interaction term that explicitly scales the treatment effect. Simon (2010) employs this approach throughout the simulation frameworks he develops for biomarker-stratified designs; Freidlin and Korn (2014) adopt it for adaptive enrichment trials; Renfro and Sargent (2017) use it throughout their treatment of basket and umbrella trials; and it underlies the platform trial liter-

155 *ature without exception. Recent simulation work on biomarker-by-treatment*
156 *interaction estimation in randomized trials with prognostic covariates, specif-*
157 *ically the study by Haller and Ulm (2018), likewise adopts the standard regres-*
158 *sion specification with explicit interaction terms. Among N-of-1 simulation*
159 *studies, Zucker et al. (1997), Araujo et al. (2016), and Duan et al. (2013)*
160 *similarly use hierarchical random-effects models in which the biomarker pre-*
161 *dicts the individual treatment effect through the mean structure. Notably,*
162 *Wang and Schork (2019) and Schork (2022), whose methodological work ad-*
163 *resses power and design for N-of-1 trials with serial correlation and carry-*
164 *over directly, also specify the treatment effect through a linear mean structure,*
165 *which suggests that the multivariate-normal differential-correlation parame-*
166 *terization of Hendrickson et al. (2020) is not a convention of any particular*
167 *research program but rather a paper-specific design choice. To the best of*
168 *our knowledge, no other clinical trial simulation framework generates the*
169 *biomarker-treatment interaction through differential correlation in a multi-*
170 *variate normal joint distribution rather than through direct mean moderation.*

171 **The two architectures were discovered by accident during pmsimstats**
172 **development, when two implementations of the same simulation diverged**
173 **under carryover.**

- 174 • The original code uses Architecture B (MVN); a tidyverse reimplementa-
175 tion inadvertently used Architecture A.
- 176 • Both produced plausible power estimates, but they reached qualitatively
177 different conclusions about carryover effects.
- 178 • Tracing the discrepancy to its root revealed the architectural distinction,
179 which has not been explicitly discussed in the simulation methodology lit-
180 erature.

181 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
182 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
183 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
184 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
185 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
186 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
187 mollit anim id est laborum.

188 *We encountered both approaches in the course of developing the pmsimstats*
189 *simulation framework for N-of-1 clinical trials (Hendrickson et al. 2020).*
190 *The original publication code uses the MVN differential correlation approach,*
191 *an unusual choice that, as we shall demonstrate here, has substantive conse-*
192 *quences for power estimation under carryover. A comprehensive audit of the*
193 *codebase subsequently revealed that a parallel tidyverse reimplementa-*
194 *tion of the same simulation had inadvertently adopted the standard direct mean*
195 *moderation approach. Both implementations produced plausible power estimates,*
196 *yet they yielded qualitatively different conclusions about how carryover effects*
197 *influence statistical power. Tracing this discrepancy to its root cause revealed*
198 *the architectural difference described in this paper, which has not, to the*
199 *best of our knowledge, been discussed explicitly in the simulation methodology*
200 *literature.*

201 **This paper is the first systematic comparison of the two DGP architec-**

202 **tures; we formalize, simulate, and give guidance.**

- 203 • We formalize the distinction between Architecture A (mean moderation)
- 204 and Architecture B (differential correlation).
- 205 • We demonstrate power consequences through Monte Carlo simulation.
- 206 • We offer guidance for investigators on how to choose the appropriate ar-
- 207 chitecture.

208 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-

209 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad

210 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip

211 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in

212 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur

213 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt

214 mollit anim id est laborum.

215 *We present here what we believe to be the first systematic comparison of*

216 *these two DGP architectures for biomarker-treatment interaction power es-*

217 *timation under carryover. We shall formalize the distinction between them,*

218 *demonstrate its consequences through Monte Carlo simulation, and offer guid-*

219 *ance for investigators designing simulation studies for predictive biomarker-*

220 *moderated treatment effects.*

221 **Related documents in this series.** The mathematical derivation,

222 code-level implementation details, four implementation attempts (three

223 failed and one correct), and the empirical evaluation of the original Hen-

224 drickson (2020) parameter grid under Architecture~A are documented in

225 19-mean-moderation-implementation-notes.tex (with a brief markdown

226 summary in 19-mean-moderation-implementation-notes.md). A condensed

227 empirical-results summary suitable for collaborator presentation is provided in

228 21-architecture-comparison-summary.md.

229

230 3 2. Two DGP Architectures

231 **Notation.** Throughout this paper, B_i denotes the biomarker value of participant

232 i , D_{it} the binary on-drug indicator at timepoint t (1 on drug, 0 off drug), and BR_{it}

233 the biological response component of the outcome. These three symbols describe

234 the data-generating-process layer. The analysis-model layer uses a related but

235 distinct continuous drug exposure variable, Dbc , defined in Section 3.1. Code-

236 style identifiers in typewriter font (**bm**, **Dbc**, **BR**) refer to specific variables in the

237 pmsimstats codebase or to R formula syntax; mathematical symbols are used

238 otherwise.

239 3.1 2.1 Architecture A: Direct Mean Moderation

240 In this approach the biomarker enters the mean structure of the outcome directly.
241 The treatment effect for participant i at timepoint t is:

$$Y_{it} = \mu_0 + \beta_1 \cdot D_{it} \cdot (1 + \beta_{bm} \cdot B_i) + \epsilon_{it}$$

242 where D_{it} is the treatment indicator (or a continuous measure of drug expo-
243 sure), B_i is the participant's biomarker value (typically centered), and β_{bm} is
244 the biomarker moderation parameter. The interaction is explicit in the mean:
245 participants with higher biomarker values have proportionally larger treatment
246 effects.

247 Key properties:

- 248 • The interaction is **deterministic** given the biomarker value and treatment
249 status.
- 250 • The biomarker moderation β_{bm} is a **population-level parameter** that
251 applies uniformly to all participants.
- 252 • The signal for the interaction is in the **first moment** (conditional mean) of
253 the outcome distribution.
- 254 • Adding noise, random effects, or carryover to D_{it} scales the entire treatment
255 effect (including its biomarker-dependent component) proportionally. The
256 **ratio** of drug effect between high-biomarker and low-biomarker participants
257 is preserved.

258 **The multiplicative and additive forms of Architecture A are interchange-**
259 **able for power calculations.**

- 260 • The multiplicative form maps directly to the standard additive regression
261 (Kent et al. 2020 effect modeling equation).
- 262 • $\beta_3 = \beta_1 \cdot \beta_{bm}$; the additive form additionally allows a prognostic main
263 effect.
- 264 • Both forms fall within Architecture A and are interchangeable for power
265 calculations on the interaction coefficient.

266 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
267 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
268 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
269 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
270 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
271 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
272 mollit anim id est laborum.

273 *We note that the multiplicative form above maps directly onto the standard*
274 *additive regression form for predictive biomarker effects used in the trial*
275 *methodology literature, specifically the effect modeling equation of Kent et al.*
276 *(2020):*

$$Y_{it} = \beta_0 + \beta_1 D_{it} + \beta_2 B_i + \beta_3 D_{it} B_i + \epsilon_{it}$$

277 The interaction coefficient β_3 in the additive form corresponds to the product
 278 $\beta_1 \cdot \beta_{bm}$ in the multiplicative form. The additive form additionally accommo-
 279 dates a prognostic main effect β_2 for the biomarker, which the multiplicative
 280 form sets to zero by construction. For the purpose of power calculations
 281 on the interaction coefficient, the two forms are interchangeable up to this
 282 prognostic-effect distinction, and both fall within Architecture A.

283 This architecture is used in the pedagogical simulation (`implementations/simple/simulation.R`)
 284 and in the exploratory carryover analyses (`simulation_carryover_spectrum.R`)
 285 of the pmsimstats project.

286 3.2 2.2 Architecture B: Differential Correlation (MVN)

287 In this approach the biomarker and the biological response (BR) component
 288 are jointly drawn from a multivariate normal distribution with treatment-state-
 289 dependent correlation:

$$\begin{pmatrix} B_i \\ BR_{it} \end{pmatrix} \sim \text{MVN} \left(\begin{pmatrix} \mu_B \\ \mu_{BR}(t, D_{it}) \end{pmatrix}, \Sigma(D_{it}) \right)$$

290 where the covariance matrix Σ has the following treatment-state-dependent BM-
 291 BR correlation:

$$\text{Cor}(B_i, BR_{it}) = \begin{cases} c_{bm} & \text{if } D_{it} = 1 \text{ (on drug)} \\ c_{bm} \cdot e^{-\lambda \cdot t_{sd}} & \text{if } D_{it} = 0 \text{ (off drug)} \end{cases}$$

292 where t_{sd} is time since drug discontinuation. The exponential-decay form is the
 293 default; the limiting cases $\lambda = 0$ (no decay; correlation persists indefinitely during
 294 the off-drug phase) and $\lambda \rightarrow \infty$ (immediate decay; correlation drops to zero the
 295 instant the drug is discontinued) recover the two no-carryover model variants.

296 **In Architecture B the interaction is in the covariance structure, not the**
 297 **mean; it emerges from the conditional distribution given the biomarker**
 298 **value.**

- 299 • The population mean of BR_{it} is the same for all participants; no mean-
 300 structure interaction exists.
- 301 • Given biomarker value b , the conditional expectation of BR scales with c_{bm} ,
 302 creating an interaction-like structure.
- 303 • Correlation is c_{bm} on drug and decays as $c_{bm}e^{-\lambda t_{sd}}$ off drug; moderation
 304 attenuates during washout.

305 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
 306 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
 307 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
 308 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
 309 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
 310 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
 311 mollit anim id est laborum.

312 *The interaction is not in the mean structure: the population mean of BR_{it}*
 313 *is the same for all participants with the same treatment history. Instead it*
 314 *emerges from the **conditional distribution**: given a participant's biomarker*
 315 *value $B_i = b$, the conditional expectation of their biological response is:*

$$E[BR_{it}|B_i = b, D_{it} = 1] = \mu_{BR}(t) + c_{bm} \cdot \frac{\sigma_{BR}}{\sigma_B} \cdot (b - \mu_B)$$

316 *This conditional expectation has an interaction-like structure: the biomarker*
 317 *value b modulates the expected BR. The effective correlation is c_{bm} on drug*
 318 *and $c_{bm} \cdot e^{-\lambda t_{sd}}$ during the washout; the moderation therefore persists with*
 319 *attenuated strength during off-drug periods and vanishes only in the limit of*
 320 *complete decay.*

321 **Key properties:**

- 322 • The interaction is **probabilistic**: it manifests as a tendency, not a deter-
 323 ministic scaling.
- 324 • The biomarker moderation c_{bm} is a **correlation parameter** that governs
 325 the strength of the association.
- 326 • The signal for the interaction is in the **second moment** (covariance struc-
 327 ture) of the joint distribution.
- 328 • Carryover effects in the DGP inflate the BR means during off-drug periods,
 329 making them resemble on-drug means. This reduces the **differential** in the
 330 correlation structure between treatment states, weakening the signal that
 331 the analysis model detects.

332 This architecture is used in the Hendrickson et al. (2020) publication code and
 333 the audited R/ package in pmsimstats-ng.

334 3.3 2.3 Mathematical Comparison

335 **Under Architecture A, carryover shrinks the signal amplitude but preserves**
 336 **the ratio between biomarker groups, so identifiability is maintained.**

- 337 • The expected outcome difference between two participants scales linearly
 338 with D .
- 339 • Under carryover, D is replaced by a decayed value and the absolute mag-
 340 nitude shrinks.
- 341 • But the ratio of drug effect between high- and low-biomarker participants
 342 is preserved at every exposure level; the signal contracts but remains iden-
 343 tifiable.

344 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
 345 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
 346 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
 347 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
 348 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
 349 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
 350 mollit anim id est laborum.

We begin the comparison by considering the expected outcome difference between two participants with biomarker values b_1 and b_2 at the same drug exposure level D under Architecture A:

$$E[Y \mid b_1, D] - E[Y \mid b_2, D] = \beta_1 \cdot \beta_{bm} \cdot (b_1 - b_2) \cdot D$$

The absolute magnitude of this difference scales linearly with D and therefore shrinks during off-drug periods when D is replaced by an exposure-decayed value $D \cdot \text{carryover}$. We note, however, that the **ratio** of the drug effect between high-biomarker and low-biomarker participants is preserved at every level of exposure. The biomarker-by-treatment signal contracts in amplitude but does not lose its identifiability under carryover.

Under Architecture B, the biomarker-dependent component of the off-drug response decays exponentially, eventually eliminating the differential signal.

- Off-drug, the conditional expectation includes a biomarker-dependent term scaled by $c_{bm} e^{-\lambda t_{sd}}$.
- As t_{sd} grows, that term shrinks toward zero.
- In the limit, all participants converge to the same off-drug mean regardless of biomarker value; the interaction signal vanishes.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Under Architecture B, the analogous quantity is:

$$E[BR|b_1, D = 1] - E[BR|b_2, D = 1] = c_{bm} \cdot \frac{\sigma_{BR}}{\sigma_B} \cdot (b_1 - b_2)$$

During off-drug periods with carryover:

$$E[BR|b, D = 0, t_{sd}] = \mu_{BR, \text{off}}(t_{sd}) + c_{bm} \cdot e^{-\lambda t_{sd}} \cdot \frac{\sigma_{BR}}{\sigma_B} \cdot (b - \mu_B)$$

Here the biomarker-dependent component decays exponentially with time off drug. As t_{sd} increases, the off-drug conditional expectation converges to $\mu_{BR, \text{off}}$ for all biomarker values, eliminating the differential signal entirely.

380

3. Consequences for Statistical Power Under Carryover

3.1 Empirical demonstration

The two DGPs are identical except for how they encode the BM-BR linkage; any power divergence is attributable solely to that difference.

- All non-biologic components (time-varying, placebo-belief, residual noise) are generated identically in both architectures.
- The only difference: Architecture B encodes BM-BR interaction in Σ ; Architecture A adds a post-draw mean shift to on-drug BR.
- The same `nlme::lme` analysis model with `corCAR1` is used for both, so divergence in results is attributable to the DGP alone.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

We ran identical simulation configurations under both architectures using the pmsimstats-ng framework, matching the DGP parameters as closely as possible between the two approaches. We note that the non-biologic response components (time-varying response, placebo-belief accumulation, and residual noise) are generated from the same multivariate normal draw with identical means, variances, within-factor AR(1) autocorrelation, and cross-factor correlations in both architectures. The two DGPs differ only in how the biomarker-to-biological-response (BM-BR) linkage is encoded: Architecture B embeds it in the off-diagonal c_{bm} entries of Σ , whereas Architecture A omits those entries and adds a post-draw shift $c_{bm} \cdot B_i^ \cdot \sigma_{BR}$ to on-drug BR columns (see 19-mean-moderation-implementation-notes.tex for the scaling derivation). The divergent carryover behavior reported below is therefore attributable to the BM-BR parameterization, not to any difference in how the placebo, time-varying, or residual-noise components are generated. The analysis model was the same in both cases: an `nlme::lme` fit with a `corCAR1` continuous-time autocorrelation structure and a single biomarker-by-drug interaction term, written `bm:Dbc` in R formula syntax.*

Dbc is a continuous drug exposure variable that captures exponential washout; the `bm:Dbc` interaction tests whether biomarker moderates that exposure-decayed effect.

- $D_{bc} = 1$ on drug; $= \exp(-\lambda t_{sd})$ off drug, matching the DGP decay rate.
- Unlike the binary D_{it} , D_{bc} smoothly captures washout shrinkage.
- The test of `bm:Dbc` is the test of whether the biomarker moderates the (decay-adjusted) drug effect.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad

425 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
426 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
427 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
428 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
429 mollit anim id est laborum.

430 *The analysis-model drug exposure variable D_{bc} is the continuous analogue*
431 *of the binary DGP indicator D_{it} . It equals 1 during on-drug periods and*
432 *$\exp(-\lambda t_{sd})$ during off-drug periods, where t_{sd} is time since drug discontin-*
433 *uation and λ is the analysis-model carryover decay rate (typically derived*
434 *from the drug's pharmacokinetic half-life as $\lambda = \ln 2/t_{1/2}$, in which case it*
435 *matches the DGP decay rate of Section 2.2 by construction). Whereas D_{it}*
436 *flips abruptly between 0 and 1, D_{bc} captures the smooth shrinkage of residual*
437 *drug effect during washout. The interaction coefficient on $bm:D_{bc}$ is therefore*
438 *the test of whether the biomarker moderates the exposure-decayed drug effect.*

439 We follow the ADEMP framework of Morris, White, and Crowther (2019): Aims,
440 Data-generating mechanisms, Estimands, Methods, Performance measures.

441 Setup:

- 442 • Trial designs: CO (traditional crossover), Hybrid (N-of-1 Hybrid),
443 OL+BDC (open-label with blinded discontinuation)
- 444 • Sample size: $N = 35$
- 445 • Biomarker correlation / moderation: $c_{bm} = 0.45$ / $\beta_{bm} = 0.45$
- 446 • DGP carryover half-life: $t_{1/2} \in \{0, 0.5, 1.0\}$ weeks
- 447 • Analysis: matched D_{bc} with same half-life
- 448 • Replicates: 1,000 per cell (Monte Carlo standard error on a binomial power
449 estimate near 0.5 is approximately $\sqrt{0.5 \cdot 0.5/1,000} \approx 0.016$)
- 450 • Implementation: 8-core `parallel::mclapply`; full factorial 2 architectures
451 $\times 3$ designs $\times 3 t_{1/2} \times 2 c_{bm}$ levels = 36 cells, 36,000 fits, 100% convergence
- 452 **$N = 35$ was chosen because it keeps power off the ceiling, making any**
453 **carryover-induced loss visible; type I error is in the expected range.**
 - 454 • $N = 35$ is the smallest prazosin-PTSD-calibrated sample size in the litera-
455 ture.
 - 456 • Neither architecture saturates power at zero carryover, so losses are me-
457 chanically observable.
 - 458 • Type I error across 18 null cells sat in $[0.010, 0.074]$, consistent with the
459 conservative-to-nominal range expected for this specification.

460 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
461 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
462 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
463 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
464 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur

sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The $N = 35$ choice is the smallest of the prazosin-PTSD- calibrated sample sizes used in the published methodological literature; it is a regime in which neither architecture saturates power at the no-carryover cells, so any carryover-induced erosion is mechanically visible. Type I error under the $c_{bm} = 0$ null cells sat in $[0.010, 0.074]$ across the 18 null cells, consistent with the conservative-to- nominal regime expected for the continuous D_{bc} specification at this sample size.

Results under Architecture A (direct mean moderation):

Design	$t_{1/2} = 0$	$t_{1/2} = 0.5$	$t_{1/2} = 1.0$
CO	0.739	0.739	0.725
Hybrid	0.992	0.987	0.978
OL+BDC	0.751	0.748	0.730

Architecture A shows modest power loss under carryover (1–3%), well below the broader literature’s 8–13% range.

- Relative losses: 1.9% (CO), 1.4% (Hybrid), 2.8% (OL+BDC).
- All three fall below the 8–13% range reported in the broader N-of-1 simulation literature.
- Qualitative finding: modest erosion; power essentially intact under carryover.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Power declines modestly under Architecture A: relative loss from $t_{1/2} = 0$ to $t_{1/2} = 1$ is 1.9% (CO), 1.4% (Hybrid), and 2.8% (OL+BDC). We note that all three values fall well below the 8-13% relative-loss range that the broader N-of-1 simulation literature reports for direct-mean-moderation DGPs, though they reproduce its qualitative finding, namely a modest erosion that leaves power essentially intact.

Results under Architecture B (differential correlation, MVN):

Design	$t_{1/2} = 0$	$t_{1/2} = 0.5$	$t_{1/2} = 1.0$
CO	0.745	0.754	0.723
Hybrid	0.995	0.980	0.944
OL+BDC	0.777	0.694	0.539

Architecture B shows sharply larger power loss under carryover, especially OL+BDC (30.6%); the B-vs-A gap is significant for OL+BDC and Hybrid

498

but not CO.

499

500

501

502

503

504

505

506

507

508

509

510

511

512

513

514

515

516

517

518

519

520

521

522

523

524

525

526

- Relative losses: 3.0% (CO), 5.1% (Hybrid), 30.6% (OL+BDC); the OL+BDC loss is $\approx 11\times$ the matched Architecture A loss.
- B-vs-A gap: $z = 7.64$ ($p < 10^{-13}$) at OL+BDC; $z = 3.96$ ($p < 10^{-4}$) at Hybrid.
- CO gap not significant ($z = 0.29$): the CO design's within-subject AR(1) structure dominates the correlation-decay channel that Architecture B depends on.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Power declines more sharply under Architecture B, and most markedly so in the OL+BDC design: relative loss from $t_{1/2} = 0$ to $t_{1/2} = 1$ is 3.0% (CO; not significantly different from Architecture A's CO loss), 5.1% (Hybrid), and 30.6% (OL+BDC). The OL+BDC loss is approximately 11 times the matched Architecture-A OL+BDC loss (2.8%). It falls below the 40-60% relative-loss range that the broader N-of-1 simulation literature reports for the MVN differential-correlation DGP, but it confirms the qualitative finding, namely a substantial and strongly design-dependent erosion. Two-proportion z -tests on the B-versus-A absolute power-loss gap give $z = 7.64$ ($p < 10^{-13}$) at OL+BDC and $z = 3.96$ ($p < 10^{-4}$) at Hybrid; the CO contrast is not statistically distinguishable from zero ($z = 0.29$) because the CO design's within-subject AR(1) structure dominates the carryover-mediated correlation decay channel that Architecture B uses to encode the biomarker-treatment interaction.

527

3.2 Why the architectures diverge

528

529

We turn now to a mechanistic account of why the two architectures produce such different power losses under carryover.

530

531

Architecture A is subject to only one carryover channel (mean blurring); it escapes the second channel (correlation erosion).

532

533

534

535

536

537

- Carryover compresses Dbc contrast and inflates off-drug BR means; both architectures share this channel.
- Architecture A lacks a second channel because β_{bm} is a fixed population parameter and the noise structure is drug-state-independent.
- Result: reduced variance in the interaction term, but signal-to-noise ratio degrades only modestly.

538

539

540

541

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in

542 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
543 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
544 mollit anim id est laborum.

545 *Under Architecture A, carryover reduces the magnitude of the drug expo-*
546 *sure variable seen by the analysis model (D_{bc} drops below 1 during washout),*
547 *and the mean BR at off-drug timepoints is inflated by the residual drug effect*
548 *exactly as in Architecture B. Both mechanisms compress the between-state*
549 *treatment contrast and are shared by the two architectures. What Architec-*
550 *ture A lacks is the second channel of signal loss: the biomarker moderation*
551 *coefficient β_{bm} is a fixed population parameter, the biomarker enters the mean*
552 *structure directly rather than through a correlation, and the noise structure is*
553 *independent of drug state. The biomarker-by-drug interaction term therefore*
554 *has reduced variance (because D_{bc} has less contrast between drug states), but*
555 *the signal-to-noise ratio degrades only modestly.*

556 **Under Architecture B**, carryover operates on two levels simultaneously:

- 557 1. **Mean blurring:** The BR means during off-drug periods are inflated by
558 residual carryover, making them resemble on-drug BR means. This reduces
559 the treatment contrast that the analysis model uses to distinguish drug effect
560 from no-drug. This channel is shared with Architecture A.
- 561 2. **Correlation erosion:** The BM-BR correlation during off-drug periods de-
562 cays exponentially with t_{sd} . This is not just a statistical modeling choice; it
563 is a structural consequence of the DGP. The correlation between biomarker
564 and BR reflects the degree to which the drug effect is active; as the drug
565 washes out, the correlation weakens because the drug-mediated component
566 of BR variance shrinks relative to the drug-independent component. This
567 channel is *unique* to Architecture B, because only here does the biomarker
568 signal live in the second moment of the joint distribution. Under Architec-
569 ture A the drug-mediated and drug-independent components of BR variance
570 move in tandem (both are scaled by the same D_{bc}-weighted drug effect), so
571 their ratio, and hence the identifiability of the interaction, is preserved even
572 as amplitudes shrink.

573 **Under Architecture B the interaction estimate is attacked from both sides**
574 **simultaneously; Architecture A faces only one of the two attacks.**

- 575 • The $bm:D_{bc}$ coefficient depends on the covariance between B_i and D_{bc}-
576 weighted outcomes.
- 577 • Under Architecture B: (1) D_{bc} contrast is compressed; (2) the BM-BR
578 correlation generating the covariance signal is simultaneously decaying.
- 579 • Architecture A faces only attack (1), not attack (2); this single difference
580 accounts for the divergence in power.

581 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
582 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
583 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
584 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
585 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur

586 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
587 mollit anim id est laborum.

588 *The analysis model's interaction coefficient (the coefficient on `bm:Dbc`) de-*
589 *pends on the covariance between B_i and Dbc-weighted outcomes. Under Ar-*
590 *chitecture B, this covariance is being attacked from both sides: the treatment*
591 *contrast in Dbc is compressed, and the BM-BR correlation that generates the*
592 *covariance signal is simultaneously decaying. That Architecture A escapes the*
593 *second attack entirely accounts for the divergence in power loss documented*
594 *above.*

595 4.3 3.3 Robustness check at $N = 70$

596 **An $N = 70$ confirmation run checks whether the qualitative ordering holds**
597 **when power is near the ceiling.**

- 598 • 800 replicates per cell, 12 cells ($2 \text{ architectures} \times 2 \text{ } c_{bm} \times 3 \text{ } t_{1/2}$),
599 OL+BDC design.
- 600 • 100% convergence across 9,600 fits.
- 601 • Purpose: verify the Section 3.1 ordering survives saturation at the larger
602 N .

603 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
604 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
605 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
606 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
607 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
608 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
609 mollit anim id est laborum.

610 *The Section 3.1 production results are at $N = 35$, where neither architecture*
611 *saturates power and the carryover-induced loss is mechanically visible. We*
612 *conducted a confirmation run at $N = 70$ to check whether the qualitative or-*
613 *dering survives the saturation that comes with the larger sample size. An 800-*
614 *replicate-per-cell Monte Carlo run was executed at $N = 70$ across 12 cells (ar-*
615 *chitecture $\in \{MVN, \text{mean moderation}\} \times c_{bm} \in \{0, 0.45\} \times t_{1/2} \in \{0, 0.5, 1.0\}$*
616 *weeks), all on the OL+BDC design, using 8-core `parallel::mclapply`.*
617 *Wall-clock 483 s; convergence 100% across 9,600 fits.*

618 At $c_{bm} = 0.45$:

Architecture	$t_{1/2} = 0$	$t_{1/2} = 0.5$	$t_{1/2} = 1.0$
MVN (B)	0.978	0.945	0.885
Mean moderation (A)	0.976	0.973	0.959

619 **The $N = 70$ run confirms the qualitative ordering; compressed absolute**
620 **magnitudes reflect ceiling saturation, not a smaller underlying effect.**

- 621 • Losses: 9.3 pp (Architecture B, $p < 10^{-12}$) vs. 1.7 pp (Architecture A,
622 $p \approx 0.05$); ordering matches $N = 35$.
- 623 • At $N = 70$ both architectures are within 2–5 pp of the ceiling at $t_{1/2} = 0$,
624 mechanically truncating observable loss.

- $N = 35$ production results remain the canonical reference for the narrative magnitude claims.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Within-architecture losses from $t_{1/2} = 0$ to $t_{1/2} = 1$ are 9.3 percentage points (Architecture B; $z = 7.4$, $p < 10^{-12}$) and 1.7 percentage points (Architecture A; $z = 2.0$, $p \approx 0.05$). The qualitative ordering matches Section 3.1's $N = 35$ result. The absolute magnitudes are compressed because $N = 70$ at the prazosin-PTSD calibration places both architectures within 2-5 percentage points of the unit ceiling at $t_{1/2} = 0$, mechanically truncating the loss either architecture's curve can absorb. We note that reading the narrative loss-magnitude claim against the $N = 70$ robustness numbers (3% and 9% relative loss for Architecture B versus 0.4% and 1.8% for Architecture A) requires acknowledging this saturation effect; the $N = 35$ production results in Section 3.1 are the canonical reference for the narrative magnitude.

Type I error is well-controlled across all 18 null cells, with no architecture-specific inflation.

- Type I error range $[0.010, 0.074]$, mean 0.043, median 0.042 across 18 null cells at $N = 35$.
- No architecture-specific inflation detected.
- The slightly conservative Architecture B OL+BDC null cell at $t_{1/2} = 1.0$ ($p = 0.010$) is the only notable departure.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

***Type I error.** Across all 18 null cells of the $N = 35$ production run, type I error sat in $[0.010, 0.074]$, with mean 0.043 and median 0.042. We observe no architecture-specific inflation. Architecture B's $t_{1/2} = 1.0$ OL+BDC null ($p = 0.010$, slightly conservative) is the smallest cell; the corresponding alternative ($p = 0.539$) is the most informative cell of the Section 3.1 narrative.*

Estimator bias is small and architecturally interpretable; coefficient drift under Architecture A contrasts with correlation-channel erosion under Architecture B.

- Mean $\hat{\beta}_{bm:D_{bc}}$ at $c_{bm} = 0.45$ ranges -0.231 to -0.281 across alternative cells; Architecture A OL+BDC at $t_{1/2} = 1.0$ produces the largest absolute mean.

- Architecture B cells maintain mean near -0.234 across all $t_{1/2}$ values, consistent with erosion through correlation decay rather than coefficient drift.
- Nominal-95% Wald coverage holds across the grid.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

***Estimator bias.** Mean estimated $\hat{\beta}_{bm:D_{bc}}$ at $c_{bm} = 0.45$ ranged from -0.231 to -0.281 across the 18 alternative cells, with the $N = 35$ Architecture-A OL+BDC cell at $t_{1/2} = 1.0$ producing the largest absolute mean (-0.281); the Architecture B cells maintained mean near -0.234 across all $t_{1/2}$ values, consistent with the Architecture-B information channel being eroded through correlation decay rather than through coefficient drift. Empirical SE of $\hat{\beta}_{bm:D_{bc}}$ runs 0.05 to 0.10 across cells; the within-replicate model SE matches at scale, indicating nominal-95% Wald coverage holds across the grid.*

689

5 4. Biological Assumptions

Architecture choice encodes a biological claim about how the biomarker moderates treatment response.

- The two architectures reflect different assumptions about the mechanism of biomarker-treatment moderation.
- Section 4 takes each architecture in turn, then discusses the prazosin-PTSD case.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The choice between architectures is not merely a statistical convenience; it reflects different assumptions about the biological mechanism by which the biomarker moderates treatment response. We consider each architecture in turn, then take up the prazosin-PTSD motivating application in detail.

5.1 4.1 When Architecture A is appropriate

709 **Architecture A is appropriate when the biomarker directly determines the**
710 **magnitude of the drug effect; the scaling is deterministic at the mean-**
711 **structure level.**

- 712 • The drug effect scales with the biomarker value: $\text{effect}_i = f(B_i) \cdot \text{drug}$.
- 713 • 'Deterministic' means any two participants with the same B_i have the same
714 expected response; individual outcomes still carry noise.
- 715 • Appropriate when the biomarker is a direct PK/PD determinant (eGFR,
716 receptor density, metabolizer status).

717 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
718 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
719 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
720 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
721 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
722 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
723 mollit anim id est laborum.

724 *Architecture A (direct mean moderation) is appropriate when the biomarker*
725 *governs the **magnitude of the drug's biological effect** for each individual*
726 *participant. The drug effect for participant i is scaled by their biomarker*
727 *value: $\text{effect}_i = f(B_i) \cdot \text{drug}$. The scaling is deterministic in the population-*
728 *mean sense, in that any two participants sharing the same B_i have the same*
729 *expected response, but individual outcomes still carry the usual noise, random-*
730 *effect, and residual-biological variability. The label 'deterministic' refers to*
731 *the mean structure of the DGP, not to the realized response of any single*
732 *participant.*

733 Clinical examples:

- 734 • **Renal clearance of a renally-excreted drug.** Estimated glomerular
735 filtration rate (eGFR) directly governs the elimination rate of drugs such
736 as digoxin or aminoglycoside antibiotics. Participants with higher eGFR
737 clear the drug more rapidly, achieving lower steady-state concentrations
738 and a proportionally smaller drug effect. The biomarker-drug relationship
739 is mechanistic and deterministic: given the biomarker value, the effective
740 dose at steady state is fully determined (up to random noise).
- 741 • **Receptor density moderation.** A biomarker that measures target recep-
742 tor density (e.g., PET imaging of receptor availability for an antagonist) di-
743 rectly determines the pharmacodynamic response to that antagonist. More
744 receptors means more drug effect, proportionally.
- 745 • **Genetic metabolizer status.** CYP2D6 metabolizer phenotype deter-
746 mines the rate of drug activation or clearance. Poor metabolizers of a
747 prodrug receive less active compound, reducing their treatment effect pro-
748 portionally.
- 749 • **Dose-response with biomarker-dependent effective dose.** When the
750 biomarker determines the effective dose (through body composition, protein
751 binding, or volume of distribution), the treatment effect scales with the

752 biomarker through the dose-response curve.

753 **In all Architecture A cases, the proportional biomarker-drug relationship**
754 **is preserved during carryover washout.**

- 755 • Residual drug effect during washout maintains the same proportional rela-
756 tionship with the biomarker.
- 757 • If participant A has $2\times$ the drug effect of B on-drug, they also have $\approx 2\times$
758 the residual effect during washout.
- 759 • This preserved proportionality is what Architecture A encodes and why it
760 loses little power under carryover.

761 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
762 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
763 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
764 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
765 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
766 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
767 mollit anim id est laborum.

768 *In all these cases the biomarker's moderating effect is preserved under car-*
769 *ryover because the residual drug effect during off-drug periods maintains the*
770 *same proportional relationship with the biomarker. If participant A has twice*
771 *the drug effect of participant B when on drug, they will also have approxi-*
772 *mately twice the residual effect during the washout period. That proportion-*
773 *ality is precisely what Architecture A encodes.*

774 5.2 4.2 When Architecture B is appropriate

775 **Architecture B is appropriate when the biomarker is a probabilistic pre-**
776 **dictor of response, mediated by latent factors the biomarker imperfectly**
777 **indexes.**

- 778 • Appropriate when the biomarker predicts which participants are likely to
779 respond, not the magnitude of any individual's response.
- 780 • The biomarker-response association is population-level and probabilistic,
781 mediated by latent factors (disease subtype, neurobiological heterogeneity,
782 comorbidities).
- 783 • Two participants with identical biomarker values may respond very differ-
784 ently.

785 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
786 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
787 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
788 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
789 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
790 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
791 mollit anim id est laborum.

792 *Architecture B (differential correlation) is appropriate when the biomarker*
793 *predicts **which participants are likely to respond**, but the magnitude of*
794 *any individual's response is not deterministically governed by the biomarker.*

795 *The biomarker and the drug response are associated at the population level, but*
796 *the association is mediated by latent factors (disease subtype, neurobiological*
797 *heterogeneity, comorbidities) that the biomarker imperfectly indexes.*

798 **Clinical examples:**

- 799 • **Blood pressure as a PTSD subtype marker.** Elevated resting blood
800 pressure may index a noradrenergic-predominant PTSD phenotype (Hen-
801 drickson et al. 2020) that is more likely to respond to prazosin (an alpha-1
802 adrenergic antagonist). The blood pressure does not *cause* the drug to
803 work better; rather, it correlates with an underlying neurobiological state
804 that determines drug responsiveness. Two participants with identical blood
805 pressure may have very different responses because blood pressure is an
806 imperfect proxy for the noradrenergic state.
- 807 • **Baseline disease severity as a treatment predictor.** Trials in
808 Alzheimer's disease, depression, and chronic pain often find that patients
809 with more severe baseline symptoms show larger treatment effects (the
810 'floor effect' or 'regression to the mean' phenomenon). The baseline severity
811 score correlates with treatment response, but the relationship is statistical
812 (population-level tendency), not deterministic.
- 813 • **Inflammatory biomarkers in psychiatry.** C-reactive protein (CRP) or
814 interleukin-6 levels may predict response to anti-inflammatory augmentation
815 of antidepressants (Raison et al. 2013). High-CRP patients are more likely
816 to respond, but the relationship is probabilistic: many high-CRP patients
817 do not respond, and some low-CRP patients do.
- 818 • **Genetic risk scores.** Polygenic risk scores predict disease trajectory and
819 treatment response, but explain only a fraction of the variance. The remain-
820 der reflects unmeasured genetic, epigenetic, and environmental factors.

821 **Under Architecture B, carryover weakens the biomarker-response corre-**
822 **lation because the biomarker's predictive power depends on the presence**
823 **of active drug.**

- 824 • As the drug washes out, the drug-mediated variance component diminishes,
825 weakening the BM-BR correlation.
- 826 • The biomarker's predictive power is tied to active drug effect; without drug,
827 it predicts nothing about non-drug symptom change.
- 828 • This is the 'correlation erosion' channel unique to Architecture B.

829 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
830 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
831 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
832 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
833 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
834 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
835 mollit anim id est laborum.

836 *In these cases carryover has a different implication: as the drug washes*

837 out, the correlation between biomarker and response weakens because the
838 drug-mediated component of response variance diminishes. The biomarker's
839 predictive power is inherently tied to the presence of active drug effect. With-
840 out drug, the biomarker is just a biomarker; it does not predict the non-drug
841 component of symptom change.

842 **Architecture B is a tractable second-moment approximation to the bio-**
843 **logically faithful finite-mixture DGP; adequate for power calculations but**
844 **not a mechanistic endorsement.**

- 845 • If the biomarker indexes a discrete responder class, the faithful DGP is a
846 finite mixture; Architecture B approximates its population-level covariance
847 structure.
- 848 • Architecture B captures the covariance implications of latent-class mem-
849 bership without committing to the full generative form.
- 850 • Adequate for power calculations; should not be read as endorsing MVN
851 differential correlation as the correct mechanism.

852 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
853 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
854 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
855 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
856 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
857 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
858 mollit anim id est laborum.

859 *The latent-subtype framing raises a legitimate modeling question: if the*
860 *biomarker is really a noisy indicator of an underlying responder class,*
861 *the biologically faithful DGP would be an explicit finite mixture in which*
862 *each participant is drawn from one of two (or more) response distributions*
863 *with class-membership probability a function of B_i . Under such a mixture*
864 *DGP the individual-level moderation is discrete, not probabilistic, and*
865 *the MVN differential-correlation parameterization is best understood as a*
866 *tractable second-moment approximation to a richer mixture model rather*
867 *than as the generative mechanism itself. The practical consequence is that*
868 *Architecture B as implemented here captures the population-level covariance*
869 *implications of a latent-class mechanism without committing to its full*
870 *generative form; this is adequate for power calculations that test for the*
871 *presence of an interaction but should not be read as an endorsement of*
872 *MVN differential correlation as the correct mechanistic model. Whether the*
873 *observed power loss under carryover is better mitigated by analysis strategies*
874 *(Section 6) or by explicit mixture modeling is an open question.*

875 **Senn (2016) cautions that apparent personalized response is often a vari-**
876 **ance artifact; Architecture B presupposes a real, stable latent subtype,**
877 **which may not hold.**

- 878 • Senn (2016) argues that apparent biomarker-by-treatment interactions are
879 often variance-component artifacts, not genuine effects.
- 880 • Architecture B presupposes the latent subtype is real, identifiable, and
881 stable across the trial period.
- 882 • If that presupposition is shaky, the documented power loss may overstate
883 recoverable signal.

884 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
885 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
886 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
887 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
888 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
889 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
890 mollit anim id est laborum.

891 *The latent-subtype framing of Architecture B should also be adopted with*
892 *appropriate epistemic caution. Senn (2016) has argued that the apparent*
893 *'personal element in response to treatment' is often a variance-component*
894 *artifact rather than a genuine biomarker-by-treatment interaction, and that*
895 *even modest within-subject variability can produce convincing-looking but spu-*
896 *rious differential response. Architecture B presupposes that the latent subtype*
897 *indexed by the biomarker is real, identifiable, and stable across the trial period.*
898 *When that presupposition is shaky, the apparent power loss under carryover*
899 *documented in Section 3 may overstate what could be recovered under any*
900 *analysis model, because the underlying signal it is attempting to estimate may*
901 *itself be smaller than assumed.*

902 5.3 4.3 The prazosin-PTSD case

903 We turn briefly to the motivating clinical application for pmsimstats, which is
904 the active-duty prazosin trial by Raskind et al. (2013), whose SBP-stratified sec-
905 ondary analysis is the subject of ongoing reanalysis by the pmsimstats team. The
906 biomarker is resting systolic blood pressure (SBP) and the outcome is CAPS
907 (Clinician-Administered PTSD Scale) score change.

908 **The prazosin-PTSD evidence base is mixed; the PACT null result mo-**
909 **tivates the search for a predictive biomarker, with SBP as the leading**
910 **candidate.**

- 911 • Earlier single-site Raskind trials (2003, 2007, 2013) reported benefit on
912 nightmares and PTSD symptoms.
- 913 • PACT (2018) failed to replicate; this has been interpreted as consistent
914 with effect heterogeneity.
- 915 • SBP, as a noninvasive proxy for central noradrenergic tone, is the leading
916 predictive biomarker candidate.

917 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
918 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
919 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
920 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
921 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
922 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
923 mollit anim id est laborum.

924 *The clinical evidence base for prazosin in PTSD is mixed. The earlier single-*
925 *site trials by Raskind and colleagues (2003, 2007, 2013) reported benefit on*
926 *nightmares and global PTSD symptoms in veteran and active-duty samples.*
927 *The subsequent multisite Veterans Affairs Cooperative Study Program trial*

(Raskind et al. 2018), known as PACT, did not replicate the effect: prazosin failed to outperform placebo on the primary endpoints. The negative PACT result has been interpreted by several investigators as consistent with effect heterogeneity, namely that prazosin works in some patients but not others, which motivates the search for a predictive biomarker that would identify the responder subgroup. Resting SBP, as a noninvasive proxy for central noradrenergic tone, is the leading candidate.

The SBP-prazosin mechanism is an Architecture B scenario: SBP is a proxy for noradrenergic state, not a direct PK determinant.

- Elevated SBP indexes heightened noradrenergic tone, a PTSD subtype preferentially responsive to prazosin.
- SBP is a proxy, not a causal determinant; two patients with identical SBP may have very different noradrenergic states.
- This is precisely the Architecture B scenario: probabilistic association, not deterministic scaling.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The biological hypothesis is that elevated SBP indexes heightened central noradrenergic tone, which characterizes a subtype of PTSD that is preferentially responsive to alpha-1 adrenergic blockade by prazosin. This is an Architecture B scenario: SBP is a proxy for an underlying neurobiological state, not a direct determinant of drug pharmacokinetics. Two patients with the same SBP may differ in their actual noradrenergic tone, alpha-1 receptor distribution, and downstream signaling, leading to different treatment responses despite identical biomarker values.

Hendrickson et al. (2020) correctly uses Architecture B for this context; the carryover power loss is clinically relevant for trial design decisions.

- Architecture B (MVN differential correlation) is appropriate for the SBP-prazosin mechanism.
- The documented carryover sensitivity means designs with longer off-drug periods will have meaningfully less power.
- This effect is larger than a direct mean moderation simulation would predict.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

973 *The Hendrickson et al. (2020) DGP uses Architecture B (MVN differential*
 974 *correlation) for this clinical context. The finding that carryover substantially*
 975 *reduces power under this architecture is therefore clinically relevant: it sug-*
 976 *gests that trial designs with longer off-drug periods will have meaningfully*
 977 *less power to detect the SBP-prazosin interaction than designs with shorter*
 978 *off-drug periods, and this effect is larger than what a direct mean moderation*
 979 *simulation would predict.*

980

981 6 5. Implications for Simulation Study Design

982 6.1 5.1 Choosing the appropriate architecture

983 The choice of DGP architecture should be guided by the biological mechanism
 984 hypothesized for the biomarker-treatment interaction:

Criterion	Architecture A	Architecture B
Biomarker role	Causal mediator	Statistical predictor
Mechanism	Deterministic scaling	Probabilistic association
Signal location	Mean structure	Covariance structure
Carryover impact on power	Modest (1-3%)	Design-dependent, up to 31%
Appropriate when	Biomarker deter- mines effective dose or PK	Biomarker indexes a latent subtype

985 6.2 5.2 Reporting requirements

986 Simulation studies evaluating biomarker-moderated treatment effects should ex-
 987 plicitly state:

- 988 1. Whether the biomarker-treatment interaction is generated through mean
 989 moderation or differential correlation.
- 990 2. The biological rationale for the chosen mechanism.
- 991 3. Whether the carryover model (if present) operates on the interaction signal
 992 consistently with the chosen architecture.
- 993 4. Sensitivity analyses under the alternative architecture, if the biological mech-
 994 anism is uncertain.

995 **The four reporting requirements are the ADEMP 'D' element applied**
 996 **specifically to biomarker-validation simulations.**

- 997 • ADEMP (Morris et al. 2019) prescribes five elements: Aims, Data-
 998 generating mechanisms, Estimands, Methods, Performance measures.
- 999 • The DGP-architecture choice concerns the D element; the four require-
 1000 ments above specialize that element for architecture-sensitive simulations.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

*The reporting framework above is consistent with the broader ADEMP template of Morris, White, and Crowther (2019), which prescribes that simulation studies fix five elements before production runs: the **Aims** of the study, the **Data-generating mechanisms**, the **Estimands** (the inferential targets), the **Methods** (analysis specifications compared), and the **Performance measures** (power, bias, coverage, type I error, and Monte Carlo standard errors of each). The DGP-architecture choice this paper concerns is the **D** element of an ADEMP specification; the reporting requirements above are the specialisation of the broader ADEMP framework to the architecture-sensitive class of biomarker-validation simulations.*

Primary estimand: $\beta_{bm:D}$ in the LME model; performance measures cover power, type I error, bias, SE, and convergence; 1000 replicates per cell.

- Estimand: $\beta_{bm:D}$ in $Sx \sim bm + t + Dbc + bm:D bc$ with random intercept and AR(1).
- Five performance measures: power, type I error, bias, empirical SE, convergence rate.
- $n_{\text{reps}} = 1000$ throughout; calibrated to achieve MCSE < 0.02 for power near 0.5.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

*The **primary estimand** for this paper's simulations is the biomarker-by-treatment interaction coefficient $\beta_{bm:D}$ in the linear-mixed analysis model $Sx \sim bm + t + Dbc + bm:D bc$ with random intercept and AR(1) residual correlation. The reported **performance measures** are: (i) power for $\beta_{bm:D}$ at $\alpha = 0.05$, with Monte Carlo standard error $\sqrt{\hat{\pi}(1 - \hat{\pi})/n_{\text{reps}}}$; (ii) type I error under $c_{bm} = 0$, with the same MCSE; (iii) bias of $\hat{\beta}_{bm:D}$ relative to the calibrated true value at each parameter cell; (iv) the empirical Monte Carlo standard error of $\hat{\beta}_{bm:D}$ across replicates; and (v) the convergence rate of the linear-mixed fit. Sample size, biomarker effect size, and carryover half-life vary across the parameter grid; the number of replicates per cell is calibrated to achieve MCSE below 0.02 for power at $\hat{\pi} = 0.5$, which corresponds to $n_{\text{reps}} \geq 625$. The simulations reported here use $n_{\text{reps}} = 1000$ throughout.*

6.3 5.3 When the choice is uncertain

If the biological mechanism is ambiguous, as it often is in early biomarker development, we suggest investigators take the following steps:

1. Run the simulation under both architectures and report the range of power estimates.
2. Architecture B produces more conservative (lower) power estimates under carryover, making it the safer default for trial design decisions.
3. Design the trial to minimize carryover (adequate washout periods) to reduce sensitivity to the DGP architecture choice.

7 6. Alternative Analysis Strategies Under Architecture B

The power loss documented in Section 3 arises from two distinct sources: (1) compressed variance of the treatment indicator D_{bc} when carryover is present, and (2) erosion of the BM-BR correlation during off-drug periods. Source 1 is a property of the analysis model parameterization and can potentially be addressed by choosing a different predictor. Source 2 is a property of the data itself (the signal is genuinely weaker in off-drug observations), and no analysis model can recover signal that is not present. However, analysis strategies that are selective about *which observations contribute to the test* or *how they are weighted* can mitigate the damage.

7.1 6.1 Restrict analysis to on-drug observations

The most direct approach is to discard off-drug timepoints entirely and test the biomarker-treatment interaction using only on-drug observations. Every retained observation has the full BM-BR correlation (c_{bm}), so correlation erosion is eliminated. The costs are twofold: a reduction in effective sample size (fewer observations per participant) and, more importantly, the loss of the off-drug reference that separates the biological response from placebo and time-varying background response. Without off-drug observations the biological, placebo, and time-varying components are confounded, which is why the Open-Label design, despite having the most on-drug timepoints, does not dominate the other designs in the Section 3 results. For designs with many on-drug timepoints (OL has 8, the Hybrid design has 5-6 per path), the net effect may be positive: the per-observation signal-to-noise ratio is maximal, and the analysis model simplifies to $Sx \sim \mathbf{bm} + \mathbf{t} + \mathbf{bm:t}$ since there is no treatment contrast to model. The interaction is detected through the correlation between biomarker and the time-on-drug trajectory.

Restricting to on-drug observations eliminates correlation erosion but loses the off-drug reference; best for OL and Hybrid designs, worst for CO.

- Benefit: every retained observation has full c_{bm} ; correlation erosion is eliminated.
- Cost: loss of off-drug reference confounds biological, placebo, and time-varying components.
- Practical: best for OL/Hybrid (many on-drug timepoints); least useful for CO (discards half the observations).

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

This approach is most promising for the Open-Label design (which has no off-drug observations regardless) and the Hybrid design (which has a long initial on-drug phase). It is least useful for the Crossover design, where discarding half the observations substantially reduces statistical power.

7.2 6.2 Weighted analysis

Weighted analysis downweights off-drug observations by drug exposure without discarding them; information-theoretic rationale; straightforward to implement.

- Weight: $w_t = e^{-\lambda t_{sd}}$ for off-drug observations; 1 for on-drug.
- Rationale: higher exposure = stronger correlation signal = more information about the interaction.
- Implemented via the `weights` argument in `nlme::lme`.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Rather than discarding off-drug observations entirely, a weighted analysis retains all observations but weights them by their estimated drug exposure level. On-drug observations receive weight 1. Off-drug observations receive weight proportional to the expected residual drug effect:

$$w_t = e^{-\lambda \cdot t_{sd}}$$

This downweights observations where the BM-BR correlation has decayed

(and where they contribute more noise than signal to the interaction estimate) without discarding them entirely. The rationale is information-theoretic: observations with higher drug exposure carry more information about the biomarker-treatment interaction because the correlation signal is stronger. This is straightforward to implement in `nlme::lme` using the `weights` argument.

7.3 6.3 Within-subject contrast

For designs with both on-drug and off-drug periods (CO, Hybrid, OL+BDC), a summary-measure approach computes a per-participant contrast: the difference between mean on-drug response and mean off-drug response. The biomarker-treatment interaction is then tested by a simple regression of these contrasts on the biomarker:

$$\bar{Y}_{i,\text{on}} - \bar{Y}_{i,\text{off}} = \alpha + \gamma \cdot B_i + \epsilon_i$$

Within-subject contrast sidesteps the repeated-measures model; trades efficiency for robustness to carryover misspecification.

- Per-participant contrast $\bar{Y}_{\text{on}} - \bar{Y}_{\text{off}}$ is regressed on the biomarker.
- Carryover compresses the contrast magnitude but the contrast-biomarker correlation may be more robust.
- Trade-off: less efficient than a full longitudinal model, but more robust to carryover and correlation misspecification.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

This sidesteps the Dbc parameterization and the repeated-measures model entirely. Carryover affects the magnitude of the contrast (making off-drug means closer to on-drug means), but the correlation between the contrast and the biomarker may be more robust because within-subject averaging reduces noise before the biomarker association is tested. The approach trades the efficiency of a full longitudinal model for robustness to misspecification of the carryover and correlation structure.

7.4 6.4 Two-stage random slopes

A two-stage approach first estimates participant-specific treatment effects, then tests their association with the biomarker. Stage 1 fits a random-slopes model:

$$Y_{it} = \beta_0 + \beta_1 t + \beta_2 D_{it} + u_{0i} + u_{1i} D_{it} + \epsilon_{it}$$

where β_2 is the population mean drug effect and u_{1i} is participant i 's random

deviation from that mean. The participant-specific drug effect is then $\beta_2 + u_{1i}$. Stage 2 regresses the estimated random deviations (or equivalently, the BLUPs of the participant-specific effects) on the biomarker:

$$\hat{u}_{1i} = \delta_0 + \delta_1 B_i + \eta_i$$

Two-stage random slopes separates within-subject effect estimation from between-subject biomarker association; may be more robust to carryover.

- Stage 1: random-slopes LME estimates participant-specific drug effects ($\beta_2 + u_{1i}$).
- Stage 2: regress BLUPs on biomarker to test association.
- The random slope u_{1i} integrates over the full trajectory; robustness to carryover is plausible but unverified by simulation.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

This separates the within-subject treatment effect estimation (which uses the longitudinal data and is affected by carryover) from the between-subject biomarker association (which uses the cross-sectional relationship between estimated effects and biomarker values). The approach may be more robust to carryover because the random slope u_{1i} integrates over the participant's entire trajectory, including both on-drug and off-drug phases.

7.5 6.5 Exclusion of contaminated observations

Excluding the first 1–2 off-drug timepoints removes the most contaminated observations; later off-drug timepoints provide a cleaner contrast.

- Early off-drug timepoints: highest residual drug effect, BM-BR correlation actively decaying.
- Later off-drug timepoints: fully washed out, cleaner contrast with on-drug phase.
- Result: comparison between fully-on and fully-off states, avoiding the strained intermediate zone.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The first 1-2 off-drug timepoints carry the most carryover contamination: the

1198 *residual drug effect is highest, and the BM-BR correlation is in the process*
1199 *of decaying. Later off-drug timepoints, where the drug has fully washed*
1200 *out, provide a cleaner contrast with the on-drug phase. Excluding the early*
1201 *off-drug observations and retaining only the later ones yields a comparison*
1202 *between fully-on and fully-off states, avoiding the intermediate zone where*
1203 *the analysis model is most strained.*

1204 7.6 6.6 Design-level solutions

1205 Rather than adapting the analysis model to accommodate carryover, the trial
1206 design itself can be modified to reduce carryover exposure:

- 1207 • **Longer washout periods** between drug phases eliminate residual drug
1208 effect before off-drug measurements begin.
- 1209 • **More on-drug timepoints** relative to off-drug increase the proportion of
1210 high-signal observations.
- 1211 • **Front-loaded drug exposure** (as in the OL+BDC design) places a
1212 long initial on-drug block before the off-drug baseline-during-continued-
1213 treatment phase, so that the BM-BR correlation has been fully established
1214 before any off-drug observations are collected. The alternative ordering
1215 (off-drug first, then on-drug) has the appealing property of placing the
1216 placebo and time-varying estimation in an uncontaminated window but
1217 sacrifices on-drug exposure duration and forces the BM-BR signal to be
1218 re-established after a wash-in; the net power trade-off depends on the ratio
1219 of carryover half-life to total trial length and is explored empirically in the
1220 Section 3 results.

1221 **Design-level changes address the root cause rather than the symptom;**
1222 **cost is longer trial or reduced ability to estimate carryover dynamics.**

- 1223 • Three levers: longer washout, more on-drug timepoints, front-loaded drug
1224 exposure (as in OL+BDC).
- 1225 • Front-loading establishes BM-BR correlation before any off-drug observa-
1226 tions; alternative ordering (off-drug first) trades clean placebo estimation
1227 for reduced on-drug exposure.
- 1228 • Addresses root cause; cost is typically longer trial or reduced carryover
1229 estimation ability.

1230 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1231 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1232 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1233 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1234 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1235 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1236 mollit anim id est laborum.

1237 *Design-level solutions address the root cause (carryover in the data) rather*
1238 *than the symptom (power loss in the analysis). Their cost is typically a*
1239 *longer trial duration or reduced ability to estimate the carryover dynamics*
1240 *themselves.*

1241 **7.7 6.7 Comparative evaluation**

1242 The strategies described above differ in their assumptions, implementation com-
1243 plexity, and expected power recovery:

Strategy	Discards data	Complexity	Expected benefit
On-drug only	Yes	Low	High for OL/Hybrid; poor for CO
Weighted	No	Low	Moderate; depends on weight calibration
Within-subject contrast	Aggregates	Low	Moderate; robust to misspecification
Two-stage random slopes	No	Moderate	Unknown; separates estimation stages
Exclude early off-drug	Some	Low	Moderate; improves contrast clarity
Design modification	N/A	N/A	High; addresses root cause

1244 A systematic simulation comparison of these strategies under Architecture B, anal-
1245 ogous to the factorial comparison of analysis models in Section 3, would identify
1246 which approach best recovers the power lost to carryover for each trial design. We
1247 regard this as a direction for future work.

1248

1249 **8 7. Relationship to the Broader Literature**

1250 **8.1 7.1 Treatment effect heterogeneity**

1251 **The PATH Statement distinguishes risk modeling from effect modeling**
1252 **for predictive HTE analysis; this is an analysis-model distinction, not a**
1253 **DGP distinction.**

- 1254
 - Rothwell (2005) establishes the predictive/prognostic biomarker distinction;
1255 Kent et al. (2020) synthesize it into the PATH Statement.
 - PATH’s risk modeling approach stratifies on a multivariable risk score; ef-
1256 fect modeling adds explicit interaction terms.
 - PATH favors risk modeling and cautions against data-driven effect model-
1257 ing because of low power, multiplicity, and overfitting.

1260 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1261 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1262 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1263 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1264 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1265 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1266 mollit anim id est laborum.

The statistical literature on treatment effect heterogeneity is extensive. Rothwell (2005) sets out the foundational distinction between predictive biomarkers (those that modify the treatment effect) and prognostic biomarkers (those that predict outcome regardless of treatment). The current consensus statement on predictive HTE analysis is the Predictive Approaches to Treatment effect Heterogeneity (PATH) Statement of Kent et al. (2020), which distinguishes two analytic strategies for identifying predictive biomarker effects in randomized trials: a risk modeling approach, in which a multivariable risk model (developed without a treatment term) is used to stratify patients and examine variation in absolute treatment benefit across risk strata, and an effect modeling approach, in which the regression includes explicit treatment-by-covariate interaction terms. PATH expresses caution about data-driven effect modeling on the grounds of low statistical power, multiplicity, and overfitting, and favors risk modeling when the application criteria are met.

PATH's analysis-model axis (risk vs. effect modeling) is orthogonal to our DGP axis; both PATH strategies implicitly assume Architecture A.

- PATH's distinction is analysis-model; ours is DGP; they are orthogonal.
- Both PATH strategies assume heterogeneity is in the mean structure (Architecture A); Architecture B does not fit either PATH category.
- Architecture A = parametric interaction model; Architecture B = latent-class/mixture-model perspective.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

We note that the PATH risk-versus-effect-modeling distinction is an analysis-model axis. The Architecture A versus Architecture B distinction introduced in this paper is a data-generating-process axis, and the two are orthogonal. Both PATH analytic strategies implicitly assume that the underlying treatment effect heterogeneity is encoded in the mean structure of the outcome (the heterogeneity is identifiable from differences in conditional means), which corresponds to Architecture A as a DGP. Architecture B as a DGP is a separate phenomenon that does not map cleanly onto either of PATH's analysis-model categories: the interaction signal resides in the covariance structure rather than the mean, and the standard regression-based predictive HTE analysis methods endorsed by PATH may have reduced power under Architecture B precisely because they were developed under an implicit Architecture A assumption. Architecture A can accordingly be characterized as a parametric interaction model (in PATH's effect modeling sense). Architecture B can be characterized as a latent class or mixture model perspective in which the biomarker is a noisy indicator of class membership.

8.2 7.2 Prevalence of each architecture

Architecture A is the near-universal standard across the biomarker-moderated trial simulation literature; Architecture B is absent from all major threads except Hendrickson et al.

- Simon (2010), Freidlin and Korn (2014), Renfro and Sargent (2017), Haller and Ulm (2018), and the platform trial literature all use Architecture A.
- N-of-1 simulation studies (Zucker et al. 1997; Araujo et al. 2016; Duan et al. 2013) use hierarchical random-effects models, a variant of Architecture A.
- Wang and Schork (2019) and Schork (2022) share authorship with Hendrickson et al. yet adopt Architecture A, suggesting the differential-correlation choice is paper-specific.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

A survey of the simulation methodology literature reveals that Architecture A (direct mean moderation) is the near-universal standard. Simon (2010) employs it throughout the simulation frameworks he develops for biomarker-stratified designs; Freidlin and Korn (2014) adopt it for adaptive enrichment trials; Renfro and Sargent (2017) use it for basket and umbrella trials; and it underlies the platform trial literature throughout. Haller and Ulm (2018), in simulation work on biomarker-by-treatment interaction estimation in randomized trials with prognostic covariates, likewise adopt the standard regression specification, focusing on power and bias of the interaction estimator under different covariate adjustment strategies. N-of-1 simulation studies, including Zucker et al. (1997), Araujo et al. (2016), and Duan et al. (2013), use hierarchical random-effects models in which individual treatment effects are drawn from a population distribution, a variant of Architecture A in which the biomarker predicts the random treatment effect through the mean structure. Methodological work on power and design for N-of-1 trials with serial correlation and carryover, from both Wang and Schork (2019) and Schork (2022), adopts the same Architecture A specification, even though this work shares senior authorship with Hendrickson et al. (2020). It appears, then, that the differential-correlation parameterization in the latter is a paper-specific design choice rather than a methodological convention of any laboratory or research program.

Architecture B appears in latent variable / SEM / copula contexts but is unique in N-of-1 clinical trial simulation.

- Architecture B appears in latent variable / SEM / Bayesian joint modeling / copula contexts (e.g., Rizopoulos 2012).
- In N-of-1 clinical trial simulation, Hendrickson et al. (2020) appears unique in using it.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad

minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Architecture B (differential correlation via MVN) appears primarily in latent variable and structural equation modeling contexts, including the joint modeling framework discussed by Rizopoulos (2012), Bayesian joint modeling frameworks, and copula-based simulation studies. In the clinical trial simulation literature specifically, the Hendrickson et al. (2020) framework appears to be unique in using differential correlation as the mechanism for biomarker-treatment interaction in an N-of-1 trial context.

The near-universal use of Architecture A means most published power estimates are optimistic for Architecture B biomarkers under carryover.

- Virtually all published power estimates for biomarker-moderated trials assume Architecture A.
- These estimates may be optimistic for biomarkers that operate through Architecture B (covariance-structure mechanism).
- The additional carryover-induced power loss under Architecture B is not accounted for.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

This predominance of Architecture A has an important implication: the power estimates reported in the vast majority of biomarker-moderated trial simulation studies may be optimistic for biomarkers that operate through an Architecture B mechanism, because these studies do not account for the additional power loss that carryover produces when the interaction signal resides in the covariance structure rather than the mean structure.

8.3 7.3 Crossover and N-of-1 trial methodology

The crossover trial literature addresses carryover as a nuisance parameter but ignores the architectural distinction; the standard recommendation applies differently under each.

- Senn (2002), Jones and Kenward (2014), and Sturdevant and Lumley (2021) treat carryover as a nuisance parameter; none engages with DGP architecture.
- Standard recommendation (include carryover in model or use adequate washout) applies to both architectures.
- Under Architecture B, crossover designs with short washout periods may be less suitable than previously recognized.

1403 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1404 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1405 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1406 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1407 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1408 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1409 mollit anim id est laborum.

1410 *The crossover trial literature, represented most thoroughly by Senn (2002)*
1411 *and Jones and Kenward (2014), addresses carryover extensively but does not*
1412 *distinguish between architectures for biomarker-treatment interaction model-*
1413 *ing. More recent methodological work by Sturdevant and Lumley (2021) on*
1414 *testing for carryover effects via mixed-effects models likewise treats carryover*
1415 *as a nuisance parameter to be estimated and adjusted for, without engag-*
1416 *ing with how carryover interacts with biomarker-by-treatment effect signals.*
1417 *The standard recommendation, namely to always include carryover in the*
1418 *analysis model or to use adequate washout, applies to both architectures but*
1419 *has different power implications under each. The finding that Architecture*
1420 *B is more sensitive to carryover suggests that crossover designs with short*
1421 *washout periods may be less suitable for detecting correlation-based biomarker*
1422 *interactions than previously recognized.*

1423 8.4 7.4 Precision medicine trial design

1424 **Enrichment and biomarker-stratified designs assume Architecture A for**
1425 **power calculations; the architectural choice matters differently across**
1426 **three design families.**

- 1427 • Simon (2010) and Freidlin and Korn (2014) frame power calculations in
1428 the Architecture A regression idiom.
- 1429 • Three design families warrant separate consideration: adaptive enrichment,
1430 treatment-switching, and basket/umbrella trials.
- 1431 • The architectural assumption is rarely made explicit, so the practical impli-
1432 cations are easy to overlook.

1433 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1434 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1435 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1436 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1437 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1438 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1439 mollit anim id est laborum.

1440 *Enrichment and biomarker-stratified designs, as discussed by Simon (2010)*
1441 *and Freidlin and Korn (2014), use Architecture A exclusively for power cal-*
1442 *culations. Three families of design are particularly worth considering in light*
1443 *of the architectural distinction.*

1444 **Adaptive enrichment trials are vulnerable to Architecture B power short-**
1445 **falls precisely when the biomarker mechanism is most uncertain.**

- 1446 • Interim power calculations driving adaptive decisions use Architecture A;

1447 if Architecture B holds and washout occurs, realized power will be lower
1448 than predicted.

- 1449 • The mismatch biases adaptation rules toward premature futility stopping
1450 or overly narrow enrichment criteria.
- 1451 • The problem is worst when the biomarker mechanism is unknown, exactly
1452 when adaptive enrichment is most often proposed.

1453 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1454 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1455 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1456 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1457 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1458 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1459 mollit anim id est laborum.

1460 *Adaptive enrichment trials* adjust the eligible population mid-trial based
1461 on emerging evidence of a biomarker-treatment interaction. The interim
1462 power calculations that drive these adaptive decisions are universally framed
1463 in the Architecture A regression idiom. If the true biomarker mechanism is
1464 closer to Architecture B and any participants pass through washout intervals
1465 between dose-finding stages, the realized power for the interaction test will be
1466 lower than the adaptation rule expects, biasing decisions toward premature
1467 futility stopping or toward overly narrow enrichment criteria. The discrep-
1468 ancy is most acute when the biomarker mechanism is uncertain at the design
1469 stage, which is precisely the situation in which adaptive enrichment is most
1470 attractive.

1471 **Treatment-switching designs suffer Architecture B power loss through**
1472 **washout erosion; parallel-group enrichment designs do not, because there**
1473 **is no off-drug phase.**

- 1474 • Crossover, basket, platform, and N-of-1 hybrid designs all cycle participants
1475 through drug states; each washout phase erodes BM-BR correlation under
1476 Architecture B.
- 1477 • Parallel-group enrichment designs escape this loss because each participant
1478 stays on a single treatment throughout.
- 1479 • The architectural choice matters most in within-subject designs, where it
1480 is also most likely to be clinically important.

1481 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1482 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1483 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1484 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1485 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1486 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1487 mollit anim id est laborum.

1488 *Treatment-switching designs* (most crossover, basket, and platform trials,
1489 as well as N-of-1 hybrid designs) cycle individual participants through different
1490 drug states. Under Architecture B, each washout phase erodes the BM-BR
1491 correlation, so the biomarker-treatment interaction signal in the second and
1492 subsequent treatment periods is structurally weaker than in the first. The

1493 power loss documented in Section 3 applies here directly. By contrast, parallel-
1494 group enrichment designs, in which each participant remains on a single
1495 treatment for the duration of the trial, escape this loss because there is no
1496 off-drug phase during which the correlation can decay; for these designs,
1497 Architectures A and B yield similar power estimates. The architectural choice
1498 therefore matters most in exactly the designs where the biomarker mechanism
1499 is most likely to be of clinical interest, namely those that deliberately expose
1500 each participant to both treatment states in order to estimate within-subject
1501 responsiveness.

1502 **Basket and umbrella trials face an additional risk of optimistic aggregated**
1503 **power estimates when Architecture B holds in even a subset of arms.**

- 1504 • Renfro and Sargent (2017) discuss power calculations that average across
1505 biomarker-defined arms; the calculation assumes Architecture A through-
1506 out.
- 1507 • Mixed architectures across baskets produce optimistic trial-level power es-
1508 timates that are difficult to detect from aggregate results.
- 1509 • Subgroup-specific power diagnostics under both architectures are
1510 warranted when the biological mechanism is heterogeneous.

1511 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1512 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1513 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1514 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1515 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1516 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1517 mollit anim id est laborum.

1518 *Basket and umbrella trials*, as discussed by Renfro and Sargent (2017),
1519 introduce an additional layer because their power calculations average across
1520 multiple biomarker-defined arms. If even a subset of the arms operate through
1521 an Architecture B mechanism while the trial-wide power calculation assumes
1522 Architecture A throughout, the aggregated power estimates will be biased op-
1523 timistically in a way that is difficult to detect from trial-level results alone.
1524 Subgroup-specific power diagnostics under both architectures should accom-
1525 pany the primary planning analysis whenever the underlying biology is het-
1526 erogeneous across baskets.

1527 **Across all three design families, the practical recommendation is to anchor**
1528 **sample size planning on the Architecture B estimate when the biomarker**
1529 **mechanism is uncertain.**

- 1530 • Run power calculations under both architectures and treat Architecture B
1531 as the conservative anchor.
- 1532 • This recommendation parallels the advice in Section 5.3 and is consistent
1533 with the PATH Statement (Kent et al. 2020).
- 1534 • The PATH Statement's caution about data-driven interaction discovery re-
1535 inforces the case for conservative planning under architectural uncertainty.

1536 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1537 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1538 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip

ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

In all three settings the practical recommendation parallels that of Section 5.3: when the biological mechanism is uncertain, run the power calculation under both architectures and treat the Architecture B estimate as the more conservative anchor for sample size planning. This recommendation is consistent with the PATH Statement (Kent et al. 2020), which itself counsels caution about data-driven discovery of treatment-by-covariate interactions and prefers risk modeling approaches when the application criteria are met.

9 8. Conclusions

In conclusion, six points are to be emphasized.

1. The choice between direct mean moderation (Architecture A) and differential correlation (Architecture B) for generating biomarker-treatment interactions in clinical trial simulations has substantive consequences for power estimates under carryover.
2. Under Architecture A, carryover reduces power only modestly, by 1 to 3% in relative terms across the three designs examined, because the proportional biomarker-drug relationship is preserved during washout. Under Architecture B the loss is strongly design-dependent: it remains modest in the crossover and hybrid designs (3 to 5%) but reaches roughly 31% in the open-label OL+BDC design, because the differential-correlation signal erodes as the drug effect wanes.
3. The choice should be guided by the hypothesized biological mechanism: Architecture A for biomarkers that directly determine drug pharmacokinetics or pharmacodynamics; Architecture B for biomarkers that statistically predict treatment response through latent subtype membership.
4. For the prazosin-PTSD application with blood pressure as the predictive biomarker, Architecture B (MVN differential correlation) is the more appropriate model, and the resulting power sensitivity to carryover should inform trial design decisions regarding off-drug period duration.
5. The existing clinical trial simulation literature uses Architecture A almost exclusively. The Hendrickson et al. (2020) MVN approach is, to the best of our knowledge, unique in the N-of-1 trial context. This means that published power estimates for biomarker-moderated crossover and N-of-1 designs may be systematically optimistic for biomarkers that operate through

a correlation-based (Architecture B) mechanism, because they do not account for the carryover sensitivity documented here.

6. Simulation studies for biomarker-moderated treatment effects should explicitly report their DGP architecture and consider sensitivity analyses under the alternative when the biological mechanism is uncertain.

10 Conflict of interest

The authors declare no conflicts of interest.

11 Funding

This work received no specific funding.

12 Data availability statement

The data and code that support the findings of this study are openly available in the pmsimstats-ng project repository. Per-replicate Monte Carlo outputs and ADEMP-compliant cell-level summaries are versioned under `analysis/data/`; simulation drivers are at `analysis/scripts/`. Exact paths for this manuscript are listed in the Reproducibility section below.

13 Reproducibility

This manuscript was rendered with R 4.4.x inside the project's `zzcollab` Docker container (image hash recorded in `Dockerfile`; package set captured in `renv.lock`). Principal packages used by the simulation pipeline are `nlme` (continuous-time linear-mixed analysis with `corCAR1` residual correlation), `parallel` (8-core `mclapply` for parameter-grid sweeps), `data.table` (the `original-extended` simulation collection), `furrr` and `purrr` (the `tidyverse` simulation collection), and `rmarkdown` plus `bookdown` for the manuscript build.

Random-number generator. The simulation driver sets `set.seed(20260507)` at the top of the replicate loop and uses `parallel::mc.set.seed` inside `mclapply`. Per-replicate seeds derive from the master seed by `+ rep_idx`.

Data and code availability. Per-replicate simulation outputs are checkpointed at `analysis/data/quick-sim/01-dgp-replicates.rds`; Morris ADEMP

cell-level summaries are at the companion `01-dgp-summary.txt`. The simulation driver is at `analysis/scripts/quick-sim/01-dgp-prototype.R`. The manuscript source, bibliography, and SIM CSL file are at `analysis/report/01-dgp-mean-moderation-vs-mvn/`

14 References

- Araujo A, Julious S, Senn S. Understanding variation in sets of N-of-1 trials. *PLoS ONE*. 2016;11(12):e0167167.
- Duan N, Kravitz RL, Schmid CH. Single-patient (N-of-1) trials: a pragmatic clinical decision methodology. *J Clin Epidemiol*. 2013;66(8):S21-S28.
- Freidlin B, Korn EL. Biomarker enrichment strategies: matching trial design to biomarker credentials. *Nat Rev Clin Oncol*. 2014;11(2):81-90.
- Haller B, Ulm K. A simulation study on estimating biomarker-treatment interaction effects in randomized trials with prognostic variables. *Trials*. 2018;19(1):128.
- Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: statistical methods and practical considerations. *Front Digit Health*. 2020;2:13.
- Jones B, Kenward MG. *Design and Analysis of Cross-Over Trials*. 3rd ed. Boca Raton, FL: CRC Press; 2014.
- Kent DM, Paulus JK, van Klaveren D, et al. The Predictive Approaches to Treatment effect Heterogeneity (PATH) Statement. *Ann Intern Med*. 2020;172(1):35-45.
- Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Stat Med*. 2019;38(11):2074-2102.
- Raison CL, Rutherford RE, Woolwine BJ, et al. A randomized controlled trial of the tumor necrosis factor antagonist infliximab for treatment-resistant depression: the role of baseline inflammatory biomarkers. *JAMA Psychiatry*. 2013;70(1):31-41.
- Raskind MA, Peskind ER, Kanter ED, et al. Reduction of nightmares and other PTSD symptoms in combat veterans by prazosin: a placebo-controlled study. *Am J Psychiatry*. 2003;160(2):371-373.
- Raskind MA, Peskind ER, Hoff DJ, et al. A parallel group placebo-controlled study of prazosin for trauma nightmares and sleep disturbance in combat veterans with post-traumatic stress disorder. *Biol Psychiatry*. 2007;61(8):928-934.
- Raskind MA, Peterson K, Williams T, et al. A trial of prazosin for combat trauma PTSD with nightmares in active-duty soldiers returned from Iraq and Afghanistan. *Am J Psychiatry*. 2013;170(9):1003-1010.

1643 Raskind MA, Peskind ER, Chow B, et al. Trial of prazosin for post-traumatic
1644 stress disorder in military veterans. *N Engl J Med.* 2018;378(6):507-517.

1645 Renfro LA, Sargent DJ. Statistical controversies in clinical research: basket trials,
1646 umbrella trials, and other master protocols. *Ann Oncol.* 2017;28(1):34-43.

1647 Rizopoulos D. *Joint Models for Longitudinal and Time-to-Event Data: With Ap-*
1648 *plications in R.* Boca Raton, FL: CRC Press; 2012.

1649 Rothwell PM. Treating individuals 5. Subgroup analysis in randomised controlled
1650 trials: importance, indications, and interpretation. *Lancet.* 2005;365(9454):176-
1651 186.

1652 Schork NJ. Accommodating serial correlation and sequential design elements in
1653 personalized studies and aggregated personalized studies. *Harv Data Sci Rev.*
1654 2022;Special Issue 3.

1655 Senn S. *Cross-over Trials in Clinical Research.* 2nd ed. Chichester, UK: John
1656 Wiley & Sons; 2002.

1657 Senn S. Mastering variation: variance components and personalised medicine.
1658 *Stat Med.* 2016;35(7):966-977.

1659 Simon R. Clinical trial designs for evaluating the medical utility of prognostic and
1660 predictive biomarkers in oncology. *Per Med.* 2010;7(1):33-47.

1661 Sturdevant SG, Lumley T. Statistical methods for testing carryover effects: a
1662 mixed effects model approach. *Contemp Clin Trials Commun.* 2021;22:100711.

1663 Varadhan R, Segal JB, Boyd CM, Wu AW, Weiss CO. A framework for the analysis
1664 of heterogeneity of treatment effect in patient-centered outcomes research. *J Clin*
1665 *Epidemiol.* 2013;66(8):818-825.

1666 Wang Y, Schork NJ. Power and design issues in crossover-based N-of-1 clinical
1667 trials with fixed data collection periods. *Healthcare (Basel).* 2019;7(3):84.

1668 Zucker DR, Schmid CH, McIntosh MW, D'Agostino RB, Selker HP, Lau J.
1669 Combining single patient (N-of-1) trials to estimate population treatment effects
1670 and to evaluate individual patient responses to treatment. *J Clin Epidemiol.*
1671 1997;50(4):401-410.

1672 15 References

1673

1674 *Rendered on 2026-06-10 at 09:42 PDT.*

1675 *Source: ~/prj/res/36-pmsimstats-ng/pmsimstats-ng/analysis/report/01-dgp-mean-moderation-vs-mv*

2026-06-10 09:42 PDT

42

~/prj/res/36-pmsimstats-ng/pmsimstats-ng/analysis/report/
01-dgp-mean-moderation-vs-mvn/report.Rmd