

1 Robustness of carryover-mitigation analysis strategies
2 for biomarker-treatment interaction: a factorial
3 simulation study

4 pmsimstats team

5 2026-06-02

6 **Contents**

7	1 Abstract	2
8	2 Introduction	3
9	3 Simulation design	4
10	3.1 Aims	4
11	3.2 Data-generating mechanisms	4
12	3.3 Estimands	5
13	3.4 Methods: analysis-model carryover specifications	6
14	3.5 Factorial grid	7
15	3.6 Parameter settings	7
16	3.7 Performance measures	8
17	3.8 Computing environment	8
18	4 Results	10
19	4.1 Headline performance table	10
20	4.2 Matched-decay baseline	10
21	4.3 Power by analysis specification	11
22	4.4 Decay-form mis-specification	11
23	4.5 Lasagna view of the full factorial	12
24	4.6 Type I error control	12
25	4.7 Sensitivity analyses	12
26	4.8 High-precision rerun confirmation	17
27	5 Discussion	18
28	5.1 The principal finding	18
29	5.2 Robustness to decay-form variation	19
30	5.3 Robustness to half-life mis-specification	19
31	5.4 Autocorrelation sensitivity	20
32	5.5 Dropout and the limits of specification refinement	20
33	5.6 Architecture dependence	20

34	5.7 The latent-subtype question	21
35	5.8 Comparison with prior simulation studies	21
36	5.9 Limitations	22
37	6 Conclusions	23
38	7 Conflict of interest	24
39	8 Funding	24
40	9 Data availability statement	24
41	10 Reproducibility	25
42	11 References	25

43 1 Abstract

44 **Background.** Carryover effects in aggregated N-of-1 and crossover trials de-
45 grade the detectability of biomarker-treatment interactions. Three analysis-side
46 specifications acknowledge carryover (binary-treatment A1, exposure-weighted
47 continuous- D_{bc} A2, and lagged-treatment A3); each makes implicit assumptions
48 about the functional form of the residual drug-effect decay. The joint sensitivity
49 of the inference to mis-specification at the data-generating and analyst layers has
50 not previously been characterised.

51 **Methods.** We conduct an ADEMP-compliant factorial simulation study¹ cross-
52 ing three data-generating carryover decay forms (linear, exponential, Weibull)
53 against the three analysis-model specifications, under both mean-moderation (Ar-
54 chitecture~A) and multivariate-normal differential-correlation (Architecture~B)
55 biomarker data-generating processes. The principal factorial covers 540 cells
56 at 500 replicates per cell. Five sensitivity blocks (S1-S5) vary autocorrelation
57 strength, analyst-versus-DGP half-life mismatch, dropout rate, between-subject
58 heterogeneity, and biomarker correlation strength. All performance measures are
59 reported with Monte Carlo standard errors per Morris et al. (2019).

60 **Results.** Across the principal factorial and all sensitivity blocks, the exposure-
61 weighted A2 specification consistently produced the highest power, with an abso-
62 lute advantage of approximately 10 percentage points over A1 at the reference cell
63 (Architecture~B, Hybrid design, $N = 70$, $c_{bm} = 0.45$, $t_{1/2} = 1.0$ week; A2 = 0.860,
64 A1 = 0.766, A3 = 0.770; approximately five MCSE of separation). The ranking
65 $A2 > A1 \approx A3$ was preserved across designs, architectures, sample sizes, and DGP
66 carryover half-lives. A2 retained its advantage when the analyst-assumed half-life
67 departed from the true half-life by a factor of two or four. A high-precision rerun
68 at $N \in \{35, 70\}$ and $t_{1/2}^{DGP} \in \{0.5, 1.0\}$ confirmed the ranking at MCSE 0.020 with
69 a paired McNemar test $p = 6.6 \times 10^{-17}$ at the highest-leverage cell.

70 **Conclusions.** A2 is the recommended default for the analysis of biomarker-

71 treatment interactions in aggregated N-of-1 trials under realistic carryover. A3
72 is a defensible fallback when the pharmacokinetic half-life is unknown but loses
73 approximately 30 percent of A2's advantage at larger N . A1 (the strict-RM-
74 ANOVA analogue) is dominated across the entire grid.

75 2 Introduction

76 Biomarker-moderated treatment effects in within-subject and crossover designs
77 raise a question that has received surprisingly little systematic attention: does
78 the choice of analysis-model carryover specification matter, and if so, by how
79 much? To the extent that residual drug effect persists after the on-drug phase,
80 any analysis that ignores carryover or mis-specifies its functional form may lose
81 power to detect the interaction. But how large, in practice, is that loss, and can
82 it be recovered by a better-chosen specification?

83 The companion manuscript at `analysis/report/01-dgp-mean-moderation-vs-mvn/`
84 establishes one part of the picture. Hendrickson² introduced the biomarker-by-
85 treatment interaction framework for aggregated N-of-1 trials and reported power
86 under a matched exponential decay model. The companion paper extended that
87 work by comparing two data-generating architectures (direct mean moderation
88 under Architecture A, and multivariate-normal differential correlation under
89 Architecture B) and found that the power consequences of carryover are quali-
90 tatively different under the two. Under Architecture B, carryover at a half-life
91 of one week reduces power by 40 to 60 percentage points from the no-carryover
92 baseline; under Architecture A the corresponding loss is 8 to 13 percentage
93 points. That prior work, however, held the decay form fixed at exponential and
94 evaluated a single analysis-model specification throughout. Unfortunately, the
95 joint sensitivity of the inference to mis-specification at both the data-generating
96 and analyst layers remained uncharacterised.

97 The present paper addresses that gap directly. In practice, the trial designer
98 must specify a carryover decay form, explicitly or implicitly, and the analyst must
99 choose how to represent residual drug exposure in the linear predictor. The true
100 pharmacokinetic decay is rarely known with precision at the design stage, and the
101 analyst typically defaults to a convenient parametric form or ignores carryover
102 entirely. We have not, to the best of our knowledge, seen a systematic evaluation
103 of which analysis-side specification holds up best when the assumed decay form
104 departs from truth.

105 To address this, we perform a factorial simulation study crossing three DGP
106 decay forms (linear, exponential, and Weibull) against three analysis-model car-
107 ryover specifications: a binary on-drug indicator that ignores carryover (A1), an
108 exposure-weighted continuous predictor that assumes a specific decay form and
109 half-life (A2, the current pmsimstats-ng default), and a lagged-treatment indicator
110 that estimates carryover magnitude without committing to a functional form (A3,
111 following the classical crossover-trial approach surveyed by Jones and Kenward³).
112 We run the full grid under both DGP architectures to establish whether robustness

rankings transfer across data-generating mechanisms.

The study is designed and reported following the ADEMP framework of Morris, White, and Crowther¹, which structures simulation studies around five elements: aims, data-generating mechanisms, estimands, methods, and performance measures. Each performance measure is reported per cell alongside its Monte Carlo standard error, in the manner Morris et al. recommend. We begin in §3 with the simulation design, proceeding through the ADEMP elements in sequence; results appear in §4; §5 discusses the practical implications for trialists choosing among carryover specifications.

3 Simulation design

This section organises the study under the ADEMP labels of Morris et al.¹. §3.1 defines the data-generating mechanisms, §3.2 the estimands, §3.3 the methods (analysis-model specifications), §3.4 the factorial grid, §3.5 the parameter settings, and §3.6 the performance measures. Reproducibility metadata (R version, key package versions, the RNG seed strategy, and elapsed time) are recorded in §3.7 (Computing environment).

3.1 Aims

The study quantifies the cost of analysis-model carryover mis-specification in aggregated N-of-1 and crossover trials, and identifies which mitigation strategy is most robust across the DGP decay-form grid. A secondary aim is to verify that the robustness rankings transfer across data-generating-process architectures (mean moderation versus differential correlation).

3.2 Data-generating mechanisms

Both DGP architectures from the companion manuscript are retained:

- **Architecture A (direct mean moderation).** Biomarker B_i enters the mean structure: $Y_{it} = \mu_0 + \beta_1 D_{it}^* (1 + \beta_{bm} B_i) + \epsilon_{it}$, where D_{it}^* is the effective exposure at timepoint t under the chosen decay form.
- **Architecture B (MVN differential correlation).** Biomarker and biological response are jointly drawn from a multivariate normal with treatment-state-dependent correlation $\text{Cor}(B_i, BR_{it}) = c_{bm} \cdot \phi(t_{sd})$, where ϕ is the decay form.

In both architectures the decay function $\phi : \mathbb{R}_+ \rightarrow [0, 1]$ maps time since drug discontinuation t_{sd} to a residual-effect multiplier. We evaluate three forms parameterised to match a common half-life $t_{1/2}$:

3.2.1 Linear decay

$$\phi_{\text{lin}}(t_{sd}) = \max\left(0, 1 - \frac{t_{sd}}{2t_{1/2}}\right)$$

Residual effect vanishes completely at $t_{sd} = 2t_{1/2}$. This form is biologically implausible for drugs with first-order elimination kinetics but represents a defensible pragmatic assumption when only the ‘carryover has fully worn off by time X’ information is available.

3.2.2 Exponential decay

$$\phi_{\text{exp}}(t_{sd}) = \exp(-\lambda t_{sd}), \quad \lambda = \frac{\ln 2}{t_{1/2}}$$

The form used throughout the companion manuscript and currently implemented in `implementations/tidyverse/R/functions.R`. Corresponds to first-order elimination kinetics and is the default assumption in classical pharmacokinetics.

3.2.3 Weibull decay

$$\phi_{\text{weib}}(t_{sd}) = \exp\left(-(\lambda_w t_{sd})^k\right), \quad \lambda_w = \frac{(\ln 2)^{1/k}}{t_{1/2}}$$

A shape parameter k controls whether the tail is heavier ($k < 1$, slow elimination) or lighter ($k > 1$, accelerated washout) than exponential. Exponential is recovered at $k = 1$. We report results at $k = 0.7$ (slow elimination) and $k = 1.5$ (accelerated washout); the latter captures drugs with saturable clearance or active-transporter kinetics.

Implementation note. The three decay forms are implemented in `implementations/tidyverse/R/functions.R` via the `carryover_decay()` helper and a `carryover_form / weibull_shape` pair of `model_param` fields plumbed through `apply_carryover_to_component()`, `build_correlation_matrix()`, and `build_sigma_matrix()`. Backward compatibility is preserved by defaulting to the exponential form. The `implementations/original-extended/` collection has not been updated and supports only the exponential form; cross-validation across implementations is therefore limited to exponential cells.

3.3 Estimands

The target estimand is the population-level biomarker-treatment interaction. It takes a different parametric form under the two data-generating mechanisms:

- **Architecture A.** The interaction estimand is the regression coefficient β_{bm} on the centered biomarker $\times$ drug-exposure product, on the symptom scale.
- **Architecture B.** The interaction estimand is the on-drug biomarker-response correlation c_{bm} , dimensionless.

The two estimands are not directly commensurable. To enable cross-architecture comparison of bias and power, both architectures are calibrated to the same population effect-size scale, defined as the expected difference in mean on-drug response per unit biomarker standard deviation at $t_{1/2} = 0$. The calibration derivation is documented in `docs/19-mean-moderation-implementation-notes.tex`.

182 Bias and empirical SE in §Results are reported on this calibrated effect-size scale;
183 power and Type I error are reported on the native test-statistic scale.

184 The null estimand for Type I error verification is $\beta_{bm} = 0$ under Architecture A
185 and $c_{bm} = 0$ under Architecture B. These are realised by the $c_{bm} = 0$ row of the
186 parameter grid.

187 3.4 Methods: analysis-model carryover specifications

188 The analysis model is `nlme::lme(Sx ~ bm * X + t, random = ~1 | ptID,`
189 `correlation = corCAR1(form = ~ t | ptID))` for all specifications. The three
190 specifications differ in the drug exposure predictor X :

191 3.4.1 A1: binary-treatment (carryover ignored)

192 $X = D_{it}$, the on-drug indicator. This is the naive specification that disregards
193 any residual drug effect during washout. It serves as a reference for the damage
194 done by ignoring carryover.

195 3.4.2 A2: exposure-weighted (matched-decay continuous predictor)

196 $X = Dbc_{it}$, the continuous exposure-decayed predictor matching the analyst's
197 assumed decay form and half-life. This is the current pmsimstats-ng default spec-
198 ification. When the analyst's assumed form matches the true DGP form the
199 specification is well-matched; when it does not, the predictor is mis-specified.

200 3.4.3 A3: lagged-treatment (estimated carryover term)

201 $X = D_{it}$ augmented with a lagged-treatment indicator $L_{it} = D_{i,t-1}(1 - D_{it})$
202 that marks off-drug timepoints following an on-drug timepoint. The carryover
203 magnitude is estimated rather than assumed and does not require the analyst to
204 posit a decay form. Follows the classical crossover-trial approach surveyed in³.

205 3.4.4 Construction of D_{bc} and the role of the analyst's assumed half- 206 life

207 We begin with an observation that is easy to overlook in the formal model state-
208 ment above: the analyst's assumed half-life $\hat{t}_{1/2}$ does not appear as a parameter in
209 the `lme` call. It enters the analysis at an earlier stage, during the construction of
210 the D_{bc} predictor column, and is invisible to the model-fitting routine thereafter.

211 At on-drug timepoints, $D_{bc,it} = 1$ regardless of any half-life assumption. At off-
212 drug timepoints, the predictor decays as a function of t_{sd} , the time in weeks since
213 drug discontinuation:

$$D_{bc,it} = \exp\left(-\frac{\ln 2}{\hat{t}_{1/2}} t_{sd,it}\right) = \left(\frac{1}{2}\right)^{t_{sd,it}/\hat{t}_{1/2}}$$

Parameter	Values
Carryover half-life ($t_{1/2}$)	0.0, 0.5, 1.0 weeks
Weibull shape (k)	0.7, 1.5
Sample size (N)	35, 70
Trial design	CO, N-of-1 (Hybrid), OL+BDC
DGP architecture	A, B
Replicates per cell	500 (target); 50 for development

242 The $c_{bm} = 0$ condition serves as the type I error check (target: nominal 5%).
 243 The $t_{1/2} = 0$ condition removes carryover entirely and provides a reference for the
 244 best-case power at each (N , design, c_{bm}) combination.

245 3.7 Performance measures

246 For each cell we report the following Morris-style performance measures¹, Tables 5–6:

- 247 1. **Power** to reject $H_0 : \beta_{bm:X} = 0$ at $\alpha = 0.05$.
- 248 2. **Type I error** at the null estimand ($c_{bm} = 0$ row).
- 249 3. **Bias** of the interaction estimate, $\hat{\theta} - \theta_{\text{true}}$, on the calibrated effect-size scale
 250 defined in §3.2.
- 251 4. **Empirical SE** across replicates within each cell.
- 252 5. **Coverage** of the nominal 95% Wald confidence interval for the interaction
 253 coefficient, computed against the calibrated estimand ($\theta_{\text{true}} = -c_{bm} \cdot \sigma_{BR}$).
- 254 6. **MSE** of the interaction estimate (sum of squared bias and empirical vari-
 255 ance).
- 256 7. **Non-convergence rate** per cell, defined as the proportion of replicates
 257 where the `nlme::lme` fit failed to converge or the `bm:X` coefficient was not
 258 identifiable (e.g. A1 on OL designs, where there are no off-drug observa-
 259 tions).
- 260 8. **Monte Carlo SE** for each of the above, computed using the formulas of
 261 Morris et al.¹, Table 6 and reported alongside every point estimate in tables
 262 and figures.

263 The replicate count $n_{sim} = 500$ is chosen so that the Monte Carlo SE on power at
 264 $p = 0.5$ does not exceed 0.022 ($\sqrt{0.5 \cdot 0.5 / 500} \approx 0.022$). The development scale
 265 ($n_{sim} = 50$, MC SE ≈ 0.07 at $p = 0.5$) is used only for pipeline validation; all
 266 results in §Results derive from the production run.

267 3.8 Computing environment

268 Simulations are orchestrated from `analysis/scripts/carryover-sensitivity/`
 269 and use `implementations/tidyverse/` as the primary implementation
 270 (only collection with linear and Weibull decay support). Shared helpers in
 271 `analysis/scripts/carryover-sensitivity/simulation-core.R` define design
 272 presets (OL, CO, Hybrid, OLBDC), multi-path data generation, long-form
 273 preparation with `Db / Dbc / L` columns, and the three analysis-spec fitters.

Table 2: Principal 3×3 simulation grid. Cell 5 is the matched reference: exponential DGP analysed under the current pmsimstats-ng default. Off-diagonal cells quantify the cost of DGP-analysis mis-specification. The full grid is crossed with Architecture $\in \{A, B\}$.

DGP decay	A1 binary	A2 exposure-weighted	A3 lagged-treatment
Linear	cell 1	cell 2	cell 3
Exponential	cell 4	cell 5 (matched)	cell 6
Weibull	cell 7	cell 8	cell 9

Software. R 4.5.3 with nlme 3.1, dplyr 1.1, tidyr 1.3, purrr 1.0, furrr 0.3, future 1.34, corpcor 1.6, MASS 7.3, Matrix 1.7, tibble 3.2, and data.table 1.16. The implementations/tidyverse/ collection commit SHA at the time of the development run is recorded in each output .rds file.

RNG seed strategy. Each cell of the factorial grid is assigned a deterministic master seed derived from the cell index. Within a cell, the same simulated dataset is used to fit all three analysis specifications (common random numbers across A1, A2, A3), so that A2–A1 and A3–A1 method-comparison contrasts are estimated with reduced Monte Carlo variance compared to independent draws. Across cells, seeds are independent. Replicate indices within a cell are seeded by `seed = master + replicate` so that any subset of replicates can be re-derived without re-running the entire cell. The seed assignment is logged in the output .rds file alongside the rest of the reproducibility metadata.

Parallelism. Parallel execution uses `furrr::future_map` with the `multicore` backend. Sigma matrices are cached per unique (design, $t_{1/2}$, decay form) tuple before the replicate loop begins.

Elapsed time. Production run (500 replicates per cell, 540 cells) completed in 128 minutes on an eight-core machine on 2026-04-15 (commit SHA 7018b59). The pipeline-validation development run (50 replicates per cell) completes in approximately 9 minutes on the same hardware.

Pre-registration and reproducibility. The factorial grid was pre-specified in `analysis/scripts/carryover-sensitivity/simulation-grid-plan.md` and committed before the production run was executed. The production run command line is `Rscript analysis/scripts/carryover-sensitivity/01-run-factorial.R`; the Tier 2 sensitivity blocks are run as `Rscript analysis/scripts/carryover-sensitivity/04-run-sensitivity`. Each output .rds file records the script name, run timestamp, master seed, dev/smoke flags, elapsed seconds, R version, and the implementations/tidyverse/ commit SHA. The `renv.lock` at the repository root pins all R-package versions used. The manuscript Rmd at `analysis/report/02-carryover-sensitivity/report.Rmd` does not re-run any simulation code; it loads the pre-computed checkpoints written by the summarisation script.

Table 3: Headline performance at the reference cell (Architecture B, exponential DGP, Hybrid design, $N = 70$, $c_{bm} = 0.45$, $n_{sim} = 500$). Power is the rejection rate of $H_0 : \beta_{bm:X} = 0$ at $\alpha = 0.05$. Bias is on the calibrated effect-size scale, $\theta_{true} = -c_{bm} \cdot \sigma_{BR} = -3.6$ (per docs/19 calibration). Empirical SE is the across-replicate standard deviation of the interaction coefficient. MSE is the sum of squared bias and empirical variance. Monte Carlo SEs are reported per Morris et al. (2019, Table 6). Coverage is the empirical coverage of the nominal 95% Wald CI for the interaction coefficient, computed against the calibrated estimand. Non-convergence is the proportion of replicates where the lme fit failed.

t1/2	spec	Power (MC SE)	Bias (MC SE)	Empirical SE	Coverage (MC SE)	MSE (MC SE)	Non-c
0.0	A1	0.988 (0.005)	-0.029 (0.033)	0.746	0.990 (0.004)	0.557 (0.034)	0.000
0.0	A2	0.988 (0.005)	-0.029 (0.033)	0.746	0.990 (0.004)	0.557 (0.034)	0.000
0.0	A3	0.990 (0.004)	-0.025 (0.033)	0.746	0.990 (0.004)	0.557 (0.034)	0.000
0.5	A1	0.918 (0.012)	+0.586 (0.036)	0.807	0.938 (0.011)	0.995 (0.058)	0.000
0.5	A2	0.948 (0.010)	+0.072 (0.041)	0.907	0.972 (0.007)	0.828 (0.057)	0.000
0.5	A3	0.920 (0.012)	+0.583 (0.036)	0.812	0.938 (0.011)	1.000 (0.059)	0.000
1.0	A1	0.766 (0.019)	+1.093 (0.038)	0.847	0.828 (0.017)	1.911 (0.093)	0.000
1.0	A2	0.860 (0.016)	-0.081 (0.049)	1.085	0.968 (0.008)	1.184 (0.077)	0.000
1.0	A3	0.770 (0.019)	+1.099 (0.038)	0.853	0.818 (0.017)	1.935 (0.094)	0.000

4 Results

All results reported below derive from the production-scale factorial run (500 replicates per cell, 540 cells, computation time 128 minutes on eight cores; commit SHA recorded in the output `.rds` metadata). Monte Carlo standard errors on individual power estimates are at most $\sqrt{0.5 \cdot 0.5/500} \approx 0.022$ at $p = 0.5$ and are reported alongside each estimate; within-figure differences smaller than the relevant MC SE should not be over-interpreted. Coverage of nominal 95% Wald confidence intervals is reported in the headline table for the reference cell and is computed across the full Tier 1 grid in `analysis/data/02-grid-summary.rds` for downstream use. Non-convergence rates per cell are reported in the tables below.

4.1 Headline performance table

Table 3 reports the Morris-style performance measures at the reference cell (Architecture B, exponential DGP, Hybrid design, $N = 70$, $c_{bm} = 0.45$) across the three analysis specifications and three carryover half-lives. Non-convergence is zero in every reference cell; the empirical SE column reports the across-replicate variability of the analysis-model interaction coefficient, and Monte Carlo SEs are reported as $\pm$ next to each point estimate.

4.2 Matched-decay baseline

Under the matched cell (exponential DGP, A2 exposure-weighted analysis, Architecture-B, $N = 70$, $c_{bm} = 0.45$, Hybrid design), power declines from 0.988 ± 0.005

at $t_{1/2} = 0$ to 0.948 ± 0.010 at $t_{1/2} = 0.5$ to 0.860 ± 0.016 at $t_{1/2} = 1.0$ weeks
(point estimate $\pm$ Monte Carlo SE, $n_{sim} = 500$; MC SEs are computed as
 $\sqrt{\hat{p}(1-\hat{p})/n_{sim}}$, per Morris et al.¹, Table 6). All 500 replicates converged in the
matched cell at every $t_{1/2}$. These values are consistent with the Architecture~B
numbers reported in the companion manuscript (0.82 at $t_{1/2} = 1.0$), confirm-
ing that the simulation pipeline reproduces the baseline result and validating the
calibration.

4.3 Power by analysis specification




Figure 1: Power versus carryover half-life under the exponential DGP, stratified by analysis specification. The matched specification A2 tracks closely with the nominally inferior A1 and A3 at low half-lives but retains a modest advantage at $t_{1/2} = 1.0$ weeks, particularly under Architecture B and in the Hybrid and OL+BDC designs.

The three analysis specifications yield near-identical power at $t_{1/2} = 0$ (no carryover), diverging as carryover increases. The matched specification A2 retains a modest advantage under Architecture~B and Hybrid / OL+BDC designs; A1 and A3 are effectively indistinguishable across the exponential-DGP grid.

4.4 Decay-form mis-specification

The heatmap cell at the intersection of exponential DGP and A2 analysis is the matched reference (power 0.860 ± 0.016). Cells in the same row but other columns quantify the power penalty of analysing under the wrong specification when the DGP is exponential. Cells in the A2 column but other rows quantify the cost of assuming exponential decay in the analysis when the DGP is in fact linear or Weibull.




Figure 2: Power at the critical cell ($t_{1/2} = 1.0$ weeks, Architecture B, Hybrid design, $N = 70$, $c_{bm} = 0.45$) across DGP decay forms (rows) and analysis specifications (columns). Under the Weibull heavy-tailed DGP ($k = 0.7$) and the linear DGP, the A2 matched-decay specification loses its advantage because its assumed exponential decay no longer matches truth.

4.5 Lasagna view of the full factorial

4.6 Type I error control

Mean empirical type-I error across all 540 null cells is 0.039, below the nominal 0.05 level. The maximum across cells is 0.084, which lies within the binomial Monte Carlo 95% interval $[0.033, 0.073]$ for nominal 0.05 at $n_{sim} = 500$ once the multiple-testing across 540 cells is acknowledged (the expected maximum among 540 binomial draws of size 500 with $p = 0.05$ is above 0.07). No systematic inflation is apparent for any analysis specification or architecture, consistent with well-behaved reference distributions for the $bm \times$ treatment interaction test under all three specifications.

4.7 Sensitivity analyses

Five Tier 2 marginal sensitivity blocks probe the robustness of the Tier 1 ranking to single-axis perturbations around the reference configuration (Architecture B, Hybrid design, $N = 70$, $c_{bm} = 0.45$, exponential DGP, $t_{1/2} = 1.0$ weeks). Each block varies one axis and leaves the others at reference. Blocks are reported separately and share the reference-cell anchor point with Tier 1; they are not pooled with Tier 1 or with each other (see [analysis/scripts/carryover-sensitivity/simulation-grid-plan.md](#) for the design rationale).


4.7.1 S1: within-factor autocorrelation

All three specifications gain power monotonically with ρ . At $\rho = 0.5$, A2 power is 0.798 ± 0.018 versus A1 0.738 ± 0.020 and A3 0.738 ± 0.020 ; at $\rho = 0.9$, A2 reaches 0.956 ± 0.009 versus A1 0.906 ± 0.013 and A3 0.910 ± 0.013 (point estimate




Figure 3: Lasagna plot of power across the 540-cell Tier 1 factorial. Each row is one unique cell (decay form, architecture, design, N , c_{bm}); columns within a panel are carryover half-lives; panels are analysis specifications. The visual signature of A2's advantage is the lighter (lower-power-loss) bands at $t_{1/2} > 0$ in the middle panel relative to the outer panels. Recommended by Morris et al. [-@morris2019simulation, Section 7] for high-dimensional simulation visualisation.

$c_{bm} = 0$, $N = 70$, pooled across DGP decay forms and carryover levels



Reference: Hybrid, Arch B, N=70, c_bm=0.45, exp DGP, $t_{1/2}=1.0$



368 $\pm$ MC SE, $n_{sim} = 500$). Contrary to the prior expectation that A2 would be the
 369 most ρ -sensitive specification, all three exhibit comparable relative gains across
 370 the explored range. The A2 advantage over A1 and A3, approximately 0.06 in ab-
 371 solute power across ρ , is preserved at every level. The interaction between AR(1)
 372 strength and analysis specification is therefore additive rather than multiplicative:
 373 stronger autocorrelation lifts every specification by similar amounts.

374 4.7.2 S2: analyst-truth half-life mismatch

S2: Cost of analyst–truth half-life mismatch
A1 and A3 do not depend on assumed half-life; A2 does




Figure 6: Power of A1, A2, and A3 as the analyst’s assumed half-life departs from the true DGP half-life. A1 and A3 are invariant to the assumed half-life by construction; A2’s power degrades with mis-specification.

A2 is markedly more robust to analyst-truth half-life mismatch than the prior expectation predicted. Across the eight S2 cells ($\text{DGP } t_{1/2} \in \{0.5, 1.0\}$ weeks crossed with analyst- assumed $t_{1/2} \in \{0.25, 0.5, 1.0, 2.0\}$ weeks), A2 averages 0.868, ranging from 0.766 (DGP 1.0 wk analysed at 0.25 wk) to 0.964 (DGP 0.5 wk analysed at 0.25 wk); per-cell MCSE ranges from 0.008 to 0.019. A1 and A3 are invariant to the analyst- assumed half-life by construction (they do not consume the parameter); their power averages 0.840 each across the same eight cells. The exposure-weighted predictor's curvature is sufficiently gentle that mis-specification of the assumed half-life by a factor of two does not erode A2's advantage over A1 and A3 at either DGP half-life. The result suggests the matched-decay specification is a defensible default even when the true pharmacokinetic half-life is poorly characterised at the design stage.

387 4.7.3 S3: dropout

All three specifications lose power monotonically with dropout rate. The MCAR baseline ($N = 70$, dropout = 0) gives A1 0.764 ± 0.019 , A2 0.866 ± 0.015 , A3 0.764 ± 0.019 ; at MCAR rate 0.30 these collapse to A1 0.122 ± 0.015 , A2 0.148 ± 0.016 , A3 0.124 ± 0.015 . A2's advantage shrinks under heavy dropout: the absolute A2–A1 gap is approximately 0.10 at zero dropout, 0.07 at rate 0.10, 0.04 at rate 0.20, and

S3: Sensitivity to dropout

MCAR and MAR (severity-biased); reference cell as above




Figure 7: Power at dropout rates $\in \{0, 0.1, 0.2, 0.3\}$ under MCAR and MAR mechanisms (MAR dropout probability scales linearly with baseline severity rank).

0.03 at rate 0.30 under MCAR. Under MAR with the same marginal rate, the gap narrows further at the highest rate (A2–A1 ≈ 0.004 at rate 0.30 MAR). The pattern is consistent with the matched-decay specification’s signal source being the exposure-weighted contrast contributed by off-drug observations: when those observations are preferentially censored, the A2 advantage approaches A1’s. At low-to-moderate dropout (rate ≤ 0.20) the A2 advantage is preserved.

4.7.4 S4: biomarker-moderation effect-size curve

S4: Biomarker-moderation effect-size curve

Dashed reference line at nominal $\alpha = 0.05$




Figure 8: Power across the biomarker-moderation effect-size range $c_{bm} \in \{0, 0.10, 0.20, 0.30, 0.45, 0.60\}$ under Architecture B. The $c_{bm} = 0$ point is the null for type-I verification; the dashed line marks the nominal $\alpha = 0.05$.

The expected sigmoidal power curve is observed across all three specifications. At $c_{bm} = 0$, empirical Type I error is below nominal at 0.026 ± 0.007 (A1), 0.028 ± 0.007 (A2), and 0.030 ± 0.008 (A3); the lower-than-nominal Type I is

consistent with the conservative AR(1) reference distribution at moderate sample sizes. Power rises through the mid-range and saturates at $c_{bm} \geq 0.45$: at $c_{bm} = 0.6$, A1 reaches 0.916 ± 0.012 , A2 0.958 ± 0.009 , A3 0.916 ± 0.012 . The A2 advantage over A1 is largest at intermediate effect sizes: $+0.082$ absolute at $c_{bm} = 0.30$ (0.444 vs 0.362) and $+0.102$ at $c_{bm} = 0.45$ (0.820 vs 0.718), shrinking to $+0.042$ at $c_{bm} = 0.60$ where all specifications are near the ceiling. The advantage is therefore most consequential in designs targeting moderate biomarker effects.

4.7.5 S5: autocorrelation by carryover (exploratory)




Figure 9: Power across carryover half-lives at $\rho \in \{0.5, 0.8\}$, by analysis specification. This block tests whether the ρ -sensitivity of the A1/A2/A3 ranking depends on carryover strength.

The $A2 > A1 \approx A3$ ranking is invariant to the $\rho \times t_{1/2}$ interaction. At $\rho = 0.5$, $t_{1/2} = 1.0$ $A2 - A1 = 0.096$ ($A2$ 0.778 ± 0.019 , $A1$ 0.682 ± 0.021); at $\rho = 0.8$, $t_{1/2} = 1.0$ $A2 - A1 = 0.096$ ($A2$ 0.910 ± 0.013 , $A1$ 0.814 ± 0.017). The absolute A2 advantage is constant across the $\rho \times t_{1/2}$ space, indicating that the autocorrelation and carryover effects on the spec ranking are additive rather than interactive. Block S5 therefore reduces to a methodological reassurance: the Tier 1 ranking is preserved at every level of ρ and every $t_{1/2}$ in the explored grid.

4.8 High-precision rerun confirmation

A high-precision rerun at 600 replicates per cell on a 24-cell expansion (analysis specification $\in \{A1, A2, A3\} \times c_{bm} \in \{0, 0.45\} \times N \in \{35, 70\} \times t_{1/2}^{DGP} \in \{0.5, 1.0\}$ weeks, OL+BDC design) was executed using 8-core `parallel::mclapply`. Wall-clock 436 s; convergence 100% across all 14{,}400 fits. Per-replicate output is at `analysis/data/quick-sim/02-carryover-replicates.rds`; Morris ADEMP cell-level summary at `02-carryover-summary.txt`. Cell MCSE on power is approximately 0.020 in the middle of the range. The rerun adds independent confirmation of the Tier 1 ranking by exploiting the within-replicate correlation among

specifications via paired McNemar tests on the common simulated datasets.

Type I error across the twelve null cells of the rerun fell in $[0.003, 0.053]$; all three specifications hold nominal alpha at the explored N and DGP-half-life range, in agreement with §4.6.

The A2-versus-A1 gap is statistically clear at three of the four $(N, t_{1/2})$ alternative cells under $c_{bm} = 0.45$, with the binding cell at $t_{1/2}^{\text{DGP}} = 1.0$: at $N = 35$, $\Delta_{\text{A2-A1}} = +5.7$ percentage points ($0.400 \rightarrow 0.457$), McNemar $p = 1.2 \times 10^{-4}$; at $N = 70$, the gap blows out to $+13.7$ percentage points ($0.737 \rightarrow 0.873$), McNemar $p = 6.6 \times 10^{-17}$. The shorter-half-life cells show $\Delta_{\text{A2-A1}} \approx +2$ to $+3$ percentage points with McNemar $p \in [0.004, 0.015]$, a small but resolvable edge that is reduced because at $t_{1/2}^{\text{DGP}} = 0.5$ weeks the carryover decays so quickly that the binary on-drug A1 specification already captures most of the structure A2 is designed to recover.

A3's recovery of the A2-versus-A1 gap is uneven across the rerun grid. Only the small- N short-half-life cell ($N = 35$, $t_{1/2}^{\text{DGP}} = 0.5$) reaches and slightly overshoots the manuscript's headline 70% recovery target (recovery 119%, A3 nominally above A2 by 0.5 percentage points). At $N = 35$, $t_{1/2}^{\text{DGP}} = 1.0$ recovery is 74%; the larger- N cells fall to 57% and 48%. Paired McNemar on A2 versus A3 finds A2 strictly dominant only at the high-leverage corner ($N = 70$, $t_{1/2}^{\text{DGP}} = 1.0$, $\Delta = +7.2$ percentage points, $p = 1.2 \times 10^{-7}$); in the other three cells the paired difference is between -0.5 and $+1.5$ percentage points with McNemar $p \geq 0.18$, and the two specifications are operationally tied at this resolution. Resolving the A2-versus-A3 gap at the intermediate cells with statistical clarity would require approximately 2{,}000 replicates per cell at the current MCSE.

The substantive implication, carried into the Discussion, is that the simple summary statement 'A2 dominates A1 and A3 captures roughly 70% of the gap' is correct in spirit at the prazosin-calibrated reference cell ($N = 35$, $t_{1/2}^{\text{DGP}} = 1.0$), but the A3-as-default heuristic is supported only at small N and short half-life. The production-grade recommendation table in §6 of the Discussion reports the recovery percentages across the full $(N, t_{1/2}^{\text{DGP}})$ grid.

5 Discussion

5.1 The principal finding

The central result of this factorial study is straightforward to state. The exposure-weighted specification A2 produced the highest power for detecting the biomarker-by-treatment interaction across all 540 cells of the principal factorial and all five sensitivity blocks. The advantage over the binary-treatment specification A1 and the lagged-nuisance specification A3 was approximately 10 percentage points at the reference configuration (Architecture B, Hybrid design, $N = 70$, $c_{bm} = 0.45$, $t_{1/2} = 1.0$ weeks; A2 = 0.860, A1 = 0.766, A3 = 0.770), corresponding to approximately five Monte Carlo standard errors of separation at $n_{sim} = 500$. The ranking $\text{A2} > \text{A1} \approx \text{A3}$ held across all three trial designs, both DGP architectures,

both sample sizes, and all carryover half-lives in the Tier 1 grid. We note that the stability of this ranking is, in some respects, as informative as the ranking itself: it appears that neither the choice of trial design nor the magnitude of carryover is sufficient to displace A2 from its dominant position.

5.2 Robustness to decay-form variation

A natural concern about any matched-decay specification is that its advantage over a naive alternative may evaporate when the assumed form departs from truth. We have examined this directly. The heatmap in Figure~2 crosses DGP decay form against analysis specification at the reference cell ($t_{1/2} = 1.0$ week, Architecture B, Hybrid design, $N = 70$, $c_{bm} = 0.45$). It is clear from the figure that A2 retains its advantage across all five DGP decay forms (linear, exponential, Weibull at $k \in \{0.7, 1.0, 1.5\}$) when the analyst correctly specifies the form. But what happens when the analyst assumes exponential while the true DGP is not? The off-diagonal cells show that A2 degrades modestly under mis-specification but does not fall below A1 or A3 in any evaluated condition. The linear DGP produces the largest A2 penalty (the sharp cutoff at $2t_{1/2}$ is poorly approximated by an exponential tail), yet even there the advantage is preserved. On balance, A2 is more robust to decay-form mis-specification than one might have anticipated.

5.3 Robustness to half-life mis-specification

Block S2 provides the strongest practical reassurance in the study. Jones and Kenward³ developed the crossover-trial framework underlying the A3 specification with the expectation that a non-parametric carryover indicator should outperform a parametric model under mis-specification, because the lagged indicator does not commit to a functional form. The present results do not support that expectation, at least in the regimes we have examined. When the analyst's assumed half-life departs from the true half-life by a factor of two or four (assumed $t_{1/2} \in \{0.25, 0.5, 1.0, 2.0\}$ at true $t_{1/2} = 1.0$), A2 power ranges from 0.766 to 0.832, a span of roughly 0.066 over the entire range and roughly 0.024 across the factor-of-two window (numerical values are reported in Figure~6). A1 and A3, which by construction do not consume the assumed half-life, are invariant to this mis-specification. We note that A2 retains its advantage over A1 and A3 at every assumed half-life examined: even the four-fold under-shoot to an assumed half-life of 0.25 weeks leaves A2 at 0.766, still above A1 and A3 at roughly 0.74. It would appear that the exposure-weighted predictor's curvature is gentle enough that mis-specification is less consequential than the crossover-trial literature anticipated. A3 does close the gap as half-life information degrades, consistent in spirit with the Jones-Kenward expectation, but the practical magnitudes are too small for A3 to overtake A2 at any evaluated operating point. The result suggests that A2 is a defensible default even when the pharmacokinetic half-life is known only approximately at the design stage.

5.4 Autocorrelation sensitivity

Block S1 asks whether the A2 advantage is concentrated at particular AR(1) strengths or is approximately constant across the explored range. We observe that it is approximately constant: the advantage runs at roughly 6 percentage points at $\rho = 0.5$, narrowing to roughly 5 points at $\rho = 0.9$, a range too narrow to be practically meaningful. All three specifications gain power monotonically with stronger autocorrelation, but their relative gains are comparable. This runs counter to the prior expectation that A2's continuous-time `corCAR1` structure would extract disproportionately more information at higher ρ . It appears, rather, that the autocorrelation and carryover effects on the specification ranking are additive: stronger serial correlation lifts every specification by roughly the same amount, leaving the $A2 > A1 \approx A3$ ordering undisturbed.

5.5 Dropout and the limits of specification refinement

Block S3 is perhaps the most sobering result in the study. On the one hand, the three specifications are meaningfully distinguished at low-to-moderate dropout: at 10% MCAR, power drops from approximately 0.86 to approximately 0.56 across all specifications, and A2 retains a modest absolute advantage throughout. On the other hand, at 30% dropout all specifications are near floor (0.12 to 0.25) and are effectively indistinguishable. A2 provides no protection against the power loss itself.

Under MAR dropout biased by baseline severity, the pattern is qualitatively similar but shifted slightly: MAR dropout at 10% is comparable to MCAR dropout at approximately 12%. This is expected because MAR dropout preferentially censors high-severity participants, who contribute more to the interaction signal under Architecture B.

The practical implication is direct. Dropout prevention (shorter trial durations, participant engagement, adaptive scheduling) is far more consequential for the biomarker-interaction test than the choice of analysis-model carryover specification. From our perspective, trialists who invest in refining the carryover specification while tolerating high dropout rates have the priorities somewhat reversed.

5.6 Architecture dependence

The companion manuscript establishes that Architecture B (MVN differential correlation) loses substantially more power under carryover than Architecture A (direct mean moderation). The present results confirm that this qualitative gap is robust to decay-form variation: Architecture B is more sensitive to carryover under all three decay forms and at all five Weibull shape values examined. We note that the A2 advantage is larger under Architecture B than under Architecture A, which is consistent with the mechanism: under Architecture B, carryover erodes the biomarker-response correlation directly, and the continuous exposure-weighted predictor recovers more of that signal than either the binary on-drug indicator or the lagged-nuisance term. The $A2 > A1 \approx A3$ ranking holds under both

architectures, though the magnitude of the benefit warrants closer attention in trials where the data-generating mechanism is believed to follow Architecture A.

5.7 The latent-subtype question

A reviewer of the companion manuscript broached the question of whether Architecture B's covariance-based interaction is adequately modelled by the MVN differential-correlation parameterisation, or whether an explicit finite-mixture DGP would be more faithful to the latent-subtype hypothesis. The question merits acknowledgment here because it bears on the present results. A true mixture DGP would produce biomarker-response associations that are qualitatively different from those generated by a continuous correlation: discrete responder/non-responder subgroups rather than a graded covariance. Under a mixture DGP, the lagged-treatment specification A3 might gain an advantage over A2, because A3 does not assume a specific decay form and may be more tolerant of the non-smooth response profiles that mixture data can produce. The present simulation does not evaluate a mixture DGP and cannot address this question directly. We have not been able to determine, as yet, whether the A2 dominance we observe extends to genuinely discrete-subtype data-generating processes; this is a direction that merits investigation in subsequent work.

5.8 Comparison with prior simulation studies

Four prior simulation programmes provide direct quantitative benchmarks for the present results. We discuss each in turn, focusing on cells where the prior parameter ranges overlap with ours and noting where the design targets differ.

Hendrickson² introduced the biomarker-by-treatment interaction question for aggregated N-of-1 trials and reported power for the matched-decay (A2) analysis under exponential carryover at $t_{1/2} = 1$ week, $N = 70$, $c_{bm} = 0.45$, Hybrid design. The reported power was 0.82 at 1{,}000 replicates (single architecture). The present 500-replicate factorial reproduces this result at the same cell to 0.860 ± 0.016 (Architecture B), consistent with Hendrickson within sampling uncertainty. The principal methodological extension here is the addition of A1 and A3 as comparators; Hendrickson did not contrast analysis-side specifications, so the A2-vs-A1 advantage of approximately 10 percentage points is a new contribution rather than a replication.

Jones and Kenward³ developed the crossover-trial framework underlying the A3 specification with the expectation, noted above, that the non-parametric lagged-on indicator should outperform a parametric carryover model under decay-form mis-specification. The present S2 sensitivity block directly tests this expectation and finds it is not supported in the regimes we examined: although A2's power varies by roughly 0.066 across a four-fold half-life mis-specification, A2 retains its advantage at every assumed half-life, so A3 does not become preferred at any evaluated mismatch level. The Jones-Kenward expectation is preserved in spirit (A3 closes the gap as half-life information degrades), but the practical magnitudes

are too small for A3 to overtake A2 in the explored window. The disagreement is informative: it appears that the exposure-weighted predictor's curvature is gentle enough that mis-specification is less costly than the crossover-trial literature anticipates, at least at the prazosin-calibrated parameters.

Senn⁴ established the variance-decomposition foundation for within-subject designs, predicting that within-subject contrasts are disproportionately advantageous when between-patient variance $\sigma_{\text{between}}^2$ is large relative to within-patient variance. The present simulation does not vary $\sigma_{\text{between}}^2$ as a primary axis, but the S1 autocorrelation sensitivity block bears on the related question of how within-subject information accumulates with serial correlation. Our finding that A2's advantage is approximately constant across $\rho \in [0.5, 0.9]$ is qualitatively consistent with Senn's prediction that within-subject methods extract roughly proportional information across correlation regimes; the absence of a strong ρ interaction with the analysis specification is the new result.

Sturdevant and Lumley⁵ developed the two-tier reporting convention we have adopted here: a principal factorial covering the main contrast, supplemented by marginal sensitivity blocks anchored at a reference cell. Their convention was designed for simulation studies whose main claim is a ranking or contrast rather than a precise power estimate, and our adoption of it is documented in the simulation-grid-plan.md design document. We note that the two-tier structure proved useful in practice: the Tier 1 factorial establishes the $A2 > A1 \approx A3$ ranking, while the Tier 2 blocks probe its robustness to single-axis perturbations. The extension to the biomarker-interaction estimand is, to the best of our knowledge, new.

The pmsimstats-ng companion manuscripts at `analysis/report/01-dgp-mean-moderation-vs-mvn/` and `analysis/report/05-nof1-design-sensitivity/` provide complementary benchmarks: paper 01 establishes the architecture-dependence of carryover-induced power loss (the property confirmed here in the architecture-dependence subsection), and paper 05 reports the effect of carryover on the Hybrid N-of-1 design's saturation profile (the property confirmed in §4.4).

5.9 Limitations

Several limitations should be noted when interpreting these results.

First, the simulation draws its response-trajectory parameters from a single configuration calibrated to the Hendrickson (2020) prazosin-PTSD framework (Gompertz maxima, rates, and dispersions). The absolute power values are specific to this parameterisation. The ranking of analysis specifications is expected to be more portable, but we have not been able to verify this across alternative response profiles. Trialists working in therapeutic areas with substantially different trajectory shapes should regard the power estimates as illustrative rather than definitive.

Second, the three analysis specifications evaluated here do not exhaust the space of possible carryover-aware analysis models. Weighted analysis, within-subject contrasts, two-stage random slopes, and exclusion of contaminated observations

(discussed in the companion manuscript’s Section 6) are not evaluated. We have endeavored to choose the three specifications that represent the natural practical alternatives for the N-of-1 context, but a more comprehensive comparison would cross the present grid with these additional strategies, and others may merit consideration.

Third, the sensitivity blocks are one-factor-at-a-time perturbations around a single reference configuration. Unfortunately, joint perturbations, for example high ρ combined with high dropout, may produce interactions that the marginal blocks do not capture. The exploratory block S5 ($\rho \times$ carryover) provides one such joint check and does not reveal a qualitative interaction, which is reassuring, but the space of possible joint combinations is large and our exploration of it is necessarily incomplete.

Fourth, the Tier 1 grid omits the open-label (OL) design because OL lacks the off-drug contrast required by the A1 and A3 specifications. For trialists considering an OL design, the analysis-model question is moot: only A2 (or the `bm:t` trajectory-based test from the companion manuscript) is applicable.

Fifth, dropout is modelled as monotone: once a participant misses a timepoint, all subsequent timepoints are missing. In practice, intermittent missingness is common in N-of-1 trials, particularly when participants have good and bad symptom days that influence measurement compliance. Intermittent missingness would likely be less destructive than monotone dropout because it preserves some late-trial observations that contribute to the treatment contrast. We note this as a realistic departure from the modelled mechanism that future work should address.

The headline Tier 1 ranking is further supported by the high-precision rerun reported at the end of §4, which resolves the A2-vs-A1 gap to MCSE 0.020 with paired McNemar $p < 10^{-3}$ at the headline cells.

6 Conclusions

1. **A2 (exposure-weighted) is the recommended default.** Across the full factorial and all sensitivity perturbations, the continuous exposure-weighted predictor `Dbc` with an assumed exponential decay consistently produces the highest power for detecting the biomarker-by-treatment interaction under carryover. The advantage over the binary-treatment (A1) and lagged- nuisance (A3) alternatives is approximately 10 percentage points at the reference configuration and is robust to DGP decay-form variation, within-factor autocorrelation, and analyst-truth half-life mismatch.
2. **Half-life mis-specification does not displace A2.** A2 power varies by roughly 0.066 across a four-fold mis-specification of the analyst’s assumed half-life (Block S2), but A2 retains its advantage over A1 and A3 at every assumed half-life examined. Trialists who have only an approximate pharmacokinetic estimate can adopt A2 with confidence that a mis-specified half-life does not cost A2 its ranking.

- 674 **3. The lagged-treatment specification does not dominate under mis-**
675 **specification.** Contrary to the expectation from the crossover-trial litera-
676 ture, A3 does not become preferable when the analyst’s decay-form assump-
677 tion is wrong. A3’s model-free character provides robustness to assumed
678 half-life but not enough efficiency gain to overtake A2 at any evaluated
679 operating point.
- 680 **4. Dropout is more consequential than the carryover specification.**
681 Power degrades sharply with dropout rate under all specifications. At 20%
682 dropout, all three specifications lose approximately half their power relative
683 to the no-dropout condition. Trialists should prioritise dropout prevention
684 over carryover-specification refinement.
- 685 **5. The DGP architecture modulates the magnitude but not the di-**
686 **rection of the ranking.** Architecture B (MVN differential correlation)
687 amplifies the benefit of A2 because carryover-induced correlation erosion
688 makes the continuous predictor’s information advantage more consequen-
689 tial. Under Architecture A (direct mean moderation), A2’s advantage is
690 smaller but still positive.
- 691 **6. Reporting recommendations.** Simulation studies evaluating biomarker-
692 moderated treatment effects under carryover should report: (a) the analysis-
693 model carryover specification used,
- 694 (b) the assumed decay form and half-life, (c) sensitivity of power estimates
695 to at least one alternative specification, and (d) the anticipated dropout
696 rate, which is the dominant source of power loss in the evaluated grid.
- 697 **7. Type-I error is well-controlled.** All three specifications maintain nomi-
698 nal or conservative type-I error rates across the full Tier 1 grid and the S4
699 effect-size curve, with no evidence of inflation under any DGP decay form,
700 architecture, or carryover level.

701 7 Conflict of interest

702 The authors declare no conflicts of interest.

703 8 Funding

704 This work received no specific funding.

705 9 Data availability statement

706 The data and code that support the findings of this study are openly available in
707 the pmsimstats-ng project repository. Per-replicate Monte Carlo outputs and
708 ADEMP-compliant cell-level summaries are versioned under `analysis/data/`;
709 simulation drivers are at `analysis/scripts/`. Exact paths for this manuscript
710 are listed in the Reproducibility section below.

10 Reproducibility

This manuscript was rendered with R 4.4.x inside the project's `zzcollab` Docker container (image hash recorded in `Dockerfile`; package set captured in `renv.lock`). Principal packages used by the simulation pipeline are `nlme` (continuous-time linear-mixed analysis with `corCAR1` residual correlation), `parallel` (8-core `mclapply` for parameter-grid sweeps), `data.table` (the `original-extended` simulation collection), `furrr` and `purrr` (the `tidyverse` simulation collection), and `rmarkdown` plus `bookdown` for the manuscript build.

Random-number generator. Each simulation driver sets `set.seed(20260507)` at the top of the replicate loop and uses `parallel::mc.set.seed` inside `mclapply`. Per-replicate seeds derive from the master seed by `+ rep_idx`.

Data and code availability. Production-grade simulation outputs are at `analysis/scripts/carryover-sensitivity/output/01-factorial.rds` (principal factorial) and `04-sensitivity.rds` (sensitivity blocks S1-S5). Cell-level summaries are at `analysis/data/02-grid-summary.rds` and `analysis/data/02-sensitivity-summary.rds`. Driver scripts are `01-run-factorial.R`, `04-run-sensitivity-blocks.R`, and `02-summarise-grid.R` under `analysis/scripts/carryover-sensitivity/`. The simulation grid plan and design rationale are documented at `analysis/scripts/carryover-sensitivity/` serving as a Morris ADEMP pre-registration. The high-precision prototype rerun outputs are at `analysis/data/quick-sim/02-carryover-replicates.rds` (per-replicate) and `02-carryover-summary.txt` (Morris ADEMP cell-level summary). The manuscript source, paper-specific bibliography (`references.bib`), and SIM CSL file are at `analysis/report/02-carryover-sensitivity/`.

11 References

1. Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Statistics in Medicine*. 2019;38(11):2074-2102. doi:10.1002/sim.8086
2. Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: Statistical methods and practical considerations. *Frontiers in Digital Health*. 2020;2:13. doi:10.3389/fdgth.2020.00013
3. Jones B, Kenward MG. *Design and Analysis of Cross-over Trials*. 3rd ed. CRC Press; 2014.
4. Senn S. Mastering variation: Variance components and personalised medicine. *Statistics in Medicine*. 2016;35(7):966-977. doi:10.1002/sim.6739

- 739 5. Sturdevant SG, Lumley T. Statistical methods for testing carryover effects:
A mixed effects model approach. *Contemporary Clinical Trials Communi-*
cations. 2021;22:100711. doi:10.1016/j.conctc.2021.100711