

1 Carryover representation and inference for biomarker-treatment
2 interaction tests under differential-correlation moderation: a
3 simulation comparison of nine analysis specifications

4 pmsimstats team

5 August 19, 2026

6 **Contents**

7	Abstract	2
8	1 Introduction	3
9	2 Simulation design	4
10	2.1 Aims	5
11	2.2 Trial designs	5
12	2.3 Data-generating mechanism	6
13	2.4 Estimand and analysis specifications	6
14	2.5 Choice of analysis specifications	9
15	2.6 Factorial grid and parameter settings	10
16	2.7 Inference and the cluster-robust correction	11
17	2.8 Performance measures and computing environment	11
18	3 Results	12
19	3.1 Headline performance table	12
20	3.2 Tier 1 grid overview	12
21	3.3 Matched-decay baseline	14
22	3.4 Power by analysis specification	15
23	3.5 Decay-form mis-specification	15
24	3.6 Type I error control and cluster-robust recalibration	16
25	3.7 Sensitivity analyses	19
26	3.8 A higher-precision check	29
27	4 Discussion	30
28	4.1 The lagged-treatment term does no work	30
29	4.2 The exposure-weighted advantage, and where it fails	31
30	4.3 What the robustness checks add	31
31	4.4 Dropout matters more than the method	31
32	4.5 What analysis is best in which situations, in practice	32
33	4.6 Scope, related work, and limits	33

34	5 Conclusions	34
35	6 Conflict of interest	36
36	7 Funding	36
37	8 Data availability statement	36
38	9 Reproducibility	36
39	10 References	37

40 Abstract

41 **Background.** Aggregated N-of-1 trials pool repeated within-patient on/off contrasts to vali-
42 date predictive biomarkers. Under covariance moderation, a multivariate-normal differential-
43 correlation architecture in which the biomarker-treatment interaction lives in a treatment-
44 state-dependent correlation, a companion study found that carryover costs up to roughly 31%
45 of power in relative terms. The analyst must still decide how to represent residual exposure
46 in the linear predictor, and it is not established which representation is most robust to the
47 ordinarily unknown pharmacokinetic decay shape.

48 **Methods.** We compare three analysis-model specifications for the drug-exposure regressor.
49 Unadjusted is the binary on-drug indicator, ignoring carryover altogether. Lag-adjusted aug-
50 ments that indicator with a lagged just-off-drug term estimating carryover non-parametrically,
51 the classical crossover device. Exposure-weighted is a continuous indicator committing to a
52 parametric decay and half-life, the current default. We embed them in an ADEMP-structured
53 factorial simulation¹ crossing two DGP decay forms (exponential, Weibull) with the three
54 specifications, three trial designs, two sample sizes, and three interaction strengths (216 cells,
55 500 replicates each), plus four sensitivity blocks and a cluster-robust recalibration. All three
56 specifications are fitted to the same simulated datasets, so every contrast is paired, and each
57 performance measure is reported with its Monte Carlo standard error (MCSE).

58 **Results.** Two findings stand out. First, the lagged-treatment term buys nothing. Lag-
59 adjusted and Unadjusted are statistically indistinguishable across all 216 cells, differing by a
60 mean of 0.003 in power against an MCSE of 0.018, and agreeing to three decimal places on
61 bias, mean squared error, and coverage. Because the two share an estimand and differ by
62 exactly one nuisance term, this is a clean null. Second, Exposure-weighted attains the highest
63 power only in the two designs with a blinded-discontinuation phase (Hybrid and OL+BDC),
64 where its advantage is about 10 percentage points (0.860 vs 0.766 for Unadjusted at the
65 reference cell). Under the classical crossover (CO) design it is markedly inferior (0.488 vs
66 0.830). Where it leads, the lead is robust to decay-shape and half-life mis-specification and
67 to autocorrelation strength. The model-based interaction test is mildly conservative for all
68 three (mean null rejection 0.036 to 0.039), an artifact of the working covariance that a CR2
69 cluster-robust standard error removes without disturbing the ranking.

70 **Conclusions.** Three points warrant emphasis. First, under the covariance-moderation ar-

chitecture we recommend the exposure-weighted predictor, reported with a cluster-robust standard error, for the Hybrid and OL+BDC designs. Under a crossover design it should not be used, and the unadjusted binary indicator is materially better. Second, the classical lagged-treatment remedy cannot be recommended in any design examined here, since it never improved on ignoring carryover altogether. We trace this to its entering the model as a nuisance main effect, which leaves the attenuation of the interaction untouched. Third, across all conditions anticipated dropout governs power far more than the choice among specifications, and should dominate design effort accordingly. This manuscript restricts the grid to the covariance-moderation architecture and to the exponential and Weibull decay forms. The mean-moderation architecture and linear decay are evaluated in the archived companion drafts `archive/report.Rmd` and `archive/report-slim.Rmd`.

1 Introduction

In a single N-of-1 trial, one patient is crossed repeatedly between active treatment and a comparator, and the within-patient contrast between on-drug and off-drug periods estimates that patient's individual treatment effect. An aggregated N-of-1 trial pools many such series into a mixed model, trading some within-patient resolution for population-level inference at sample sizes considerably smaller than a parallel-group trial would require². The estimand is the biomarker-treatment interaction, whether a baseline covariate B_i modifies the treatment effect. We examine three trial designs that differ in how on-drug and off-drug occasions are arranged, a classical crossover (CO), an open-label design followed by blinded discontinuation (OL+BDC), and a Hybrid N-of-1 design, each defined in Section 2.2.

The inferential complication specific to this setting is carryover. A drug's pharmacological effect does not switch off abruptly at discontinuation but decays over days or weeks, so nominally off-drug observations still carry residual exposure. An exposure regressor that ignores this is a mis-measured version of what the patient actually received, and the damage is twofold. The interaction contrast is diluted, because occasions still carrying a substantial fraction of the drug effect are coded identically to occasions carrying none, biasing the biomarker-by-exposure coefficient toward zero as any covariate measured with error would. And those same occasions are heterogeneous in a way the model does not represent, so the unexplained variation inflates the standard error of the coefficient that has just been attenuated. The test loses on both counts, and power falls accordingly.

This is a loss of power, not of validity. Under the null, $c_{bm} = 0$, no biomarker-response association exists at any level of residual exposure, so a mis-coded regressor has no interaction to distort and test size is unaffected by exposure coding. The mild conservatism observed in Section 3.5 has a separate origin in the working covariance and applies equally across specifications. Carryover mis-specification is therefore a question of efficiency, and we frame the study in terms of recovered power.

Three answers are available for how to represent residual exposure in the linear predictor, corresponding to three positions an analyst can take toward the unknown decay. The analyst may deny it, coding exposure strictly on or off, may estimate it, adding a nuisance term

111 whose magnitude the data determine, or may assume it, committing to a decay curve and a
 112 half-life. A companion study³, building on Hendrickson², shows that under the covariance-
 113 moderation architecture, in which the interaction lives in the treatment-state-dependent co-
 114 variance $\text{Cor}(B_i, BR_{it}) = c_{bm} \cdot \phi(t_{sd})$, carryover erodes power directly and design-dependently.
 115 At a half-life of one week, the relative loss is roughly 31% under OL+BDC but only 3 to 5%
 116 under crossover and Hybrid, against 1 to 3% under an alternative architecture where the
 117 biomarker moderates the mean response directly. We restrict attention here to covariance
 118 moderation, where the carryover problem is largest and Hendrickson’s original results were
 119 obtained. That earlier work fixed the decay shape at exponential and evaluated only one
 120 analysis specification, leaving open which representation of residual exposure performs best
 121 when the true decay shape is not known precisely, as is ordinarily the case.

122 We compare three specifications. Unadjusted is the binary on-drug indicator, ignoring carry-
 123 over entirely and serving as the benchmark any remedy must beat. Lag-adjusted augments it
 124 with a lagged just-off-drug term that estimates carryover magnitude without assuming a func-
 125 tional form, following the classical crossover-trial device of Jones and Kenward⁴. Exposure-
 126 weighted is a continuous indicator committing to a parametric decay form and half-life, the
 127 current pmsimstats-ng default. Unadjusted and Lag-adjusted estimate the same coefficient
 128 and are fitted to the same simulated datasets, differing by exactly one nuisance term, so
 129 their contrast isolates that term’s contribution alone. Exposure-weighted differs on a sec-
 130 ond axis as well, since it changes the regressor carrying the interaction. We cross all three
 131 against exponential and Weibull data-generating decay forms and ask whether one can be
 132 recommended outright, or whether the choice depends on trial design, following the ADEMP
 133 reporting framework of Morris, White, and Crowther¹ throughout. Section 2 describes the
 134 simulation design in ADEMP order, Section 3 gives the results, and Section 4 takes up the
 135 practical implications.

136 2 Simulation design

137 The study simulates complete aggregated N-of-1 trials and analyzes each one under every
 138 specification. A single replicate proceeds as follows. Patients are allocated across the random-
 139 ization paths of the chosen trial design (Section 2.2). For each patient we draw a pre-treatment
 140 biomarker together with the three latent response components, namely the pharmacological
 141 response, the natural-history trend and the placebo-belief component, jointly across all eight
 142 measurement occasions from a multivariate normal distribution whose means follow modified
 143 Gompertz trajectories and whose covariance carries within-component AR(1) autocorrelation,
 144 cross-component correlations, and the treatment-state-dependent biomarker-response correla-
 145 tion that encodes the interaction (Section 2.3). The observed symptom score at each occasion
 146 is the patient’s baseline severity less the sum of the three components, so that a beneficial
 147 effect lowers the score. Carryover enters through a decay function of time since discontinua-
 148 tion, which keeps part of the pharmacological effect present at nominally off-drug occasions.
 149 All three analysis specifications (Section 2.4) are then fitted to that same simulated dataset,
 150 and we record the interaction coefficient, its standard error and its p -value from each. This
 151 is repeated 500 times for every parameter combination in the grid of Section 2.5.

152 Fitting all specifications to a common dataset is deliberate. It makes every contrast among
153 them paired, so that differences are estimated with substantially less Monte Carlo variance
154 than independent draws would give. Throughout, N denotes the **total** number of patients in
155 a trial, allocated as evenly as possible across the design's randomization paths, so a design
156 with P paths assigns about N/P patients to each. At $N = 70$ this is 35 per path under the
157 two-path designs and 17 or 18 per path under the four-path Hybrid design.

158 2.1 Aims

159 The aim is to quantify, under covariance moderation, the cost of mis-specifying the carryover
160 representation in the analysis model, and to determine which of the three candidate strategies
161 performs better across the grid of data-generating decay forms and trial designs. A second
162 aim, addressed in the Tier 2 sensitivity blocks of Section 3.6, is to check whether the result-
163 ing ranking survives perturbation of within-subject autocorrelation, dropout rate, half-life
164 accuracy, and interaction size.

165 2.2 Trial designs

166 The three designs differ in how on-drug and off-drug occasions are arranged, and so in how
167 much post-discontinuation information each generates. All three schedule eight measurement
168 occasions per patient across randomization paths, each a fixed on-drug/off-drug schedule,
169 and each design carries an expectancy weight (one during open-label periods, one half during
170 blinded periods) scaling the placebo-belief component.

171 **Crossover (CO).** A traditional two-period within-patient crossover, blinded throughout,
172 occasions spaced 2.5 weeks apart from 2.5 to 20 weeks. One path is active for the first
173 four occasions then comparator, and the other reverses that order and so never discontinues,
174 contributing no post-discontinuation observations.

175 **Open-label followed by blinded discontinuation (OL+BDC).** Open-label titration at
176 4, 8, 12, and 16 weeks, then weekly follow-up to week 20 with a blinded late switch to placebo.
177 The two paths discontinue after week 17 or 18, so only two or three occasions follow.

178 **Hybrid.** The Hendrickson design and this program's reference. Open-label titration at
179 weeks 4 and 8 is followed by blinded discontinuation weekly from week 9 to 12, then a brief
180 two-occasion crossover at weeks 16 and 20. The four paths vary discontinuation timing and
181 crossover-arm order. It is the richest of the three in on-off transitions and the only one allowing
182 more than one discontinuation per patient.

183 One consequence bears directly on the results. The first post-discontinuation occasion falls
184 one week after switching under OL+BDC and Hybrid, but 2.5 weeks under CO. At the half-
185 lives examined here, roughly half the pharmacological effect therefore persists into the first
186 off-drug occasion under the two non-crossover designs, against less than a fifth under CO, so
187 the designs differ substantially in how much carryover contamination they can accumulate at
188 all (Section 4.2).

2.3 Data-generating mechanism

Under covariance moderation, the biomarker and the biological response are drawn jointly from a multivariate normal distribution with a treatment-state-dependent correlation, $\text{Cor}(B_i, BR_{it}) = c_{bm} \cdot \phi(t_{sd})$, where ϕ is the decay form and t_{sd} the time since drug discontinuation. Mean moderation and the combined dual-channel architecture fall outside this manuscript's scope. The full cross-architecture comparison is in the archived full-scope drafts (`archive/report.Rmd`, `archive/report-slim.Rmd`). The decay function $\phi : \mathbb{R}_+ \rightarrow [0, 1]$ maps t_{sd} to a residual-effect multiplier, and we evaluate two forms here, each matched to a common half-life $t_{1/2}$. A third, linear decay, appears in the full-scope manuscripts but not here.

2.3.1 Exponential decay

$$\phi_{\text{exp}}(t_{sd}) = \exp(-\lambda t_{sd}), \quad \lambda = \frac{\ln 2}{t_{1/2}}$$

This is the decay form used throughout the companion manuscript and implemented in `functions.R` under `implementations/tidyverse/`. It corresponds to classical first-order elimination kinetics, the default pharmacokinetic assumption, and serves as the reference against which Weibull is compared.

2.3.2 Weibull decay

$$\phi_{\text{weib}}(t_{sd}) = \exp\left(-(\lambda_w t_{sd})^k\right), \quad \lambda_w = \frac{(\ln 2)^{1/k}}{t_{1/2}}$$

A shape parameter k governs whether the elimination tail is heavier than exponential ($k < 1$, slower elimination) or lighter ($k > 1$, faster clearance), recovering the exponential form exactly at $k = 1$. We report $k = 0.7$ (a prolonged low-concentration tail) and $k = 1.5$ (faster clearance once concentration falls below some threshold), which with $k = 1.0$ (numerically indistinguishable from exponential) gives four decay-shape levels in the grid.

2.4 Estimand and analysis specifications

The interaction estimand is the on-drug biomarker-response correlation c_{bm} , a dimensionless quantity. Bias and empirical SE are reported on a calibrated effect-size scale (the expected difference in on-drug response per SD of the biomarker at $t_{1/2} = 0$, per `docs/19-mean-moderation-implementation-notes.tex`). Power and Type I error are on the native test-statistic scale, with the null estimand $c_{bm} = 0$. The analysis model is `nlme::lme(Sx ~ bm * X + t, random = ~1 | ptID, correlation = corCAR1(form = ~ t | ptID))` for the Tier 1 three, which differ only in how X is constructed and whether a nuisance term is added. Unadjusted and Lag-adjusted both test the coefficient on $\text{bm}:D_{it}$, so every performance measure is comparable between them. Exposure-weighted tests $\text{bm}:D_{bc,it}$, a different quantity. Power and Type I error remain comparable across all three, since they are properties of the test, but bias, MSE, and coverage are not comparable across that

Table 1: All nine analysis specifications evaluated in this manuscript. G1-G3 are the Tier 1 headline three (Section 2.4); G4-G5 are Tier 2 additions replacing the retired E4/E5/E6 (Blocks S7-S9, Section 3.6); G6-G9 are CR2 cluster-robust variants of G1, G2, G4, and G3 respectively (Block S10). G10 (G5+CR2) does not apply and is omitted: G5 collapses to one row per patient before fitting, leaving no repeated-measures cluster structure for CR2 to correct.

Code	Stored	Name	Formula	Target coefficient
G1	E1	Unadjusted	Sx bm + t + Db + bm:Db	bm:Db
G2	E3	Lag-adjusted	Sx bm + t + Db + bm:Db + L	bm:Db
G3	E2	Exposure-weighted	Sx bm + t + Dbc + bm:Dbc	bm:Dbc
G4	E7	AIC-selected t1/2	G3 formula, t1/2 by AIC over {0.25,0.5,1.0,2.0}	bm:Dbc
G5	E9	Paired-difference	diff bm (patient-level OLS)	bm slope
G6	E1cr2	Unadjusted + CR2	G1 formula, CR2 SE	bm:Db
G7	E3cr2	Lag-adjusted + CR2	G2 formula, CR2 SE	bm:Db
G8	E2cr2	Exposure-weighted + CR2	G3 formula, CR2 SE	bm:Dbc
G9	E7cr2	AIC-selected t1/2 + CR2	G4 formula, CR2 SE	bm:Dbc

divide (Section 3.1). Nine specifications are evaluated in total across this manuscript: the Tier 1 three below, plus six Tier 2 additions introduced in “Choice of analysis specifications” following this section and tested in Blocks S7-S10 (Section 3.6). A short-code layer, **G1-G9**, is used for the full nine-way comparisons in Blocks S7-S10 and is defined in **spec-labels.R**; the Tier 1 three and Tier 2 discussion elsewhere in this manuscript refer to them by their E-prefixed stored codes (**E1**, **E2**, ...) or by name, per Table 1 below.

2.4.1 Unadjusted (G1): binary on-drug indicator

$X = D_{it}$, the binary on-drug indicator (one on drug, zero off). Carryover is not represented at all, so an off-drug occasion still carrying most of the pharmacological effect is coded identically to one carrying none. This is the default an analyst obtains without considering carryover, and serves as the benchmark the other two must justify their additional structure against. It consumes no half-life and so is invariant to the analyst’s pharmacokinetic assumptions.

2.4.2 Lag-adjusted (G2): binary indicator plus estimated carryover term

$X = D_{it}$ is retained and the model is augmented with a lagged-treatment indicator $L_{it} = D_{i,t-1}(1 - D_{it})$, marking off-drug timepoints immediately following an on-drug one. Carryover magnitude is estimated from the data rather than assumed, so, like Unadjusted, this specification is invariant to the assumed half-life, following the classical crossover-trial device of Jones and Kenward⁴. Critically, L_{it} enters only as a nuisance main effect, with no $\text{bm}:L_{it}$ product, so the interaction tested remains $\text{bm}:D_{it}$. Lag-adjusted can absorb mean-level carryover but cannot repair the interaction’s attenuation (Section 1). It is best understood as Unadjusted with a carryover nuisance control added, not a fundamentally different model of the interaction (Section 4.1).

2.4.3 Exposure-weighted (G3): continuous decayed predictor

$X = D_{bc,it}$, a continuous exposure-decayed predictor matched to the analyst's assumed decay form and half-life, the current pmsimstats-ng default. At on-drug timepoints $D_{bc,it} = 1$. Off drug it decays as $D_{bc,it} = (1/2)^{t_{sd,it}/\hat{t}_{1/2}}$, losing half its value every $\hat{t}_{1/2}$ weeks after discontinuation. When the assumed decay form matches the true one, the predictor is well matched. When it does not, it is mis-specified in a way probed directly in block S2. The half-life assumption enters only through the predictor values handed to `lme`, not the formula, so Exposure-weighted is in effect a different model for every assumed half-life. It is the only one of the three that consumes a half-life at all, which is what the S2 sensitivity contrast (Section 3.6) is designed to probe. A limiting relationship connects the first and third. As $\hat{t}_{1/2} \rightarrow 0$, $D_{bc,it} \rightarrow D_{it}$, so Unadjusted is exactly the boundary member of the Exposure-weighted family at zero assumed half-life, and the S2 sweep is therefore a traverse anchored at Unadjusted rather than an independent check.

2.4.4 AIC-selected half-life (G4)

Exposure-weighted's formula, $X = D_{bc,it}$, but $\hat{t}_{1/2}$ is chosen from the data rather than assumed: the model is refit at each of four candidate half-lives (0.25, 0.5, 1.0, 2.0 weeks) and the fit with the lowest AIC is kept, mirroring `characterize_carryover()` (`R/carryover_analysis.R`), the package's own real-data half-life-selection routine. Like Exposure-weighted, it targets `bm:D_{bc,it}`. Unlike Exposure-weighted, it consumes no analyst-supplied half-life at all, at the cost of a post-selection inference risk (the same data drive the half-life choice and the subsequent test) that Block S11 (Section 3.6) checks directly.

2.4.5 Paired-difference (G5)

A two-stage, correlation-structure-agnostic alternative rather than an elaboration of the mixed-model family above. For each patient with both on- and off-drug observations, compute the patient's own mean on-drug minus mean off-drug S_x ; regress that single per-patient difference on `bm` across patients (`diff ~ bm`, plain OLS). No repeated-measures model, no carryover representation, and no correlation structure of any kind. A direct biomarker-interaction extension of the paired t-test Senn, Julious and Araujo⁵ found to outperform carryover-adjusted mixed models for ordinary N-of-1 treatment-effect estimation. Patients contributing only on-drug or only off-drug observations (for instance CO's never-discontinuing path) cannot form a within-patient difference and are dropped.

2.4.6 CR2 cluster-robust variants (G6-G9)

Section 3.5 (Block S6) shows the model-based interaction test is mildly conservative and that a CR2 bias-reduced cluster-robust standard error (Bell and McCaffrey, via `clubSandwich`, paired with Satterthwaite degrees of freedom) restores nominal size without disturbing the ranking among G1-G3. That correction, `lme+CR2` in Section 2.6, was applied only to G3 there. Block S10 (Section 3.6) extends it to G1, G2, and G4 as well, replacing only the standard error, reference distribution, and *p*-value on each specification's existing point estimate: G6

(G1+CR2), G7 (G2+CR2), G9 (G4+CR2). G8 (G3+CR2) is the same specification already introduced in Section 2.6 and Block S6, carried into the nine-way G1-G9 comparison at Block S10 without being rerun. A tenth combination, G5+CR2, does not apply: G5 already collapses to one row per patient before fitting, leaving no repeated-measures cluster structure for CR2 to correct.

2.5 Choice of analysis specifications

Three considerations bounded which specifications were tested, rather than surveying the wider space of models used in the N-of-1 and crossover-trial literature.

The Tier 1 three (Unadjusted, Lag-adjusted, Exposure-weighted, above) were chosen because each represents a distinct, established position an analyst can take toward an unmeasured or assumed pharmacokinetic decay, not because they exhaust that space. Unadjusted is the default any analyst obtains without considering carryover at all, and so is the necessary benchmark for every alternative. Lag-adjusted follows Jones and Kenward's⁴ classical crossover-trial device, the shape-free remedy the pre-mixed-model crossover literature has long recommended. Exposure-weighted is the pmsimstats-ng package's own established default, so it is the specification any user of the package is actually running today, and the manuscript's practical stake in getting its behavior right is correspondingly higher than for the other two.

This left a wide range of alternative model classes untested by design. Bayesian hierarchical N-of-1 models, GEE/marginal models compared systematically rather than as a robustness check on one specification's standard error (Section 2.6), randomization/permutation tests, functional or spline-based flexible decay curves, state-space or Kalman-filter treatments of drug concentration, and two-stage "series of N-of-1 trials" meta-analytic combination are all established in the N-of-1 or crossover-trial literature and none is evaluated here. This is a scoping decision, not an oversight: the manuscript's aim is to compare carryover *representations* within a single, fixed conditional mixed-model framework, matching Hendrickson's² original analysis model, not to conduct a comprehensive methods bake-off. A wider survey belongs to future work (Section 4.6).

Blocks S7 through S9 (Section 3.6) revisit this scoping decision once, for a specific reason. Section 4.1 shows that no carryover term tested so far, however it enters, can in principle repair the interaction's attenuation if it enters only as a nuisance main effect; the natural next model interacts the carryover term with the biomarker directly. An earlier version of Block S7 tested two such specifications, E5 (a lagged term crossed with `bm`) and E6 (a continuous elapsed-time term crossed with `bm`). Both underperformed Unadjusted at every cell examined across three blocks, two trial designs, two architectures, and three carryover half-lives (Sections 3.6, 4.5), traced to a structural cause: the data-generating process examined throughout this manuscript has a single true decay parameter, so crossing a second carryover term with the biomarker adds a free parameter with no counterpart in the truth being estimated, and two collinear regressors chasing one true degree of freedom lose power to a correctly specified single-parameter model almost by construction.

Rather than test a further variant of the same interacted-term family, expected on that same

structural argument to fail similarly, E5/E6 were replaced with two specifications answering different open questions about the same underlying problem, the analyst's unknown true decay half-life. E7 asks whether estimating the half-life from the data, by AIC selection over a candidate grid (reusing the logic of `characterize_carryover()`, `R/carryover_analysis.R`, already used for real-data analysis in this package), recovers the power a correctly assumed half-life would give, addressing Block S2's question from the opposite direction: instead of asking how much a wrong assumption costs, it asks whether estimation avoids needing the assumption at all. E9 asks a more basic question raised directly by the wider literature: is a conditional mixed model even necessary. Senn, Julious and Araujo⁵ compared four analysis strategies for ordinary (no biomarker interaction) N-of-1 trials and found that a simple paired t-test on within-patient treatment differences achieved near-nominal Type I error, higher power, and comparable bias and MSE relative to a carryover-adjusted mixed model, recommending the mixed model only when carryover is specifically suspected. E9 is a direct biomarker-interaction extension of their preferred method: a two-stage, correlation-structure-agnostic regression of each patient's own mean on-drug-minus-off-drug difference on the biomarker across patients, with no repeated-measures model and no carryover representation at all.

This remains a bounded comparison. E7 and E9 were chosen because each is implementable within, or trivially adjacent to, the existing analysis pipeline (Section 2.6) and each answers a specific question this manuscript's own results raised, not because they close the wider literature gap identified above. The Bayesian, GEE-as-primary-model, randomization-test, and two-stage meta-analytic alternatives remain untested and are discussed as future work (Section 4.6).

2.6 Factorial grid and parameter settings

The principal comparison crosses DGP decay form (exponential and Weibull at $k \in \{0.7, 1.0, 1.5\}$) against the three analysis specifications, under covariance moderation only, with the exponential row analyzed under Exposure-weighted as the matched reference cell.

Parameter	Values
Biomarker moderation (c_{bm})	0.0, 0.30, 0.45
Carryover half-life ($t_{1/2}$)	0.0, 0.5, 1.0 weeks
Weibull shape (k)	0.7, 1.0, 1.5 (Weibull form only)
Sample size (N)	35, 70
Trial design	CO, N-of-1 (Hybrid), OL+BDC
DGP architecture	B (covariance moderation) only
Replicates per cell	500 (target), 50 for development

The four decay-shape levels (exponential and three Weibull shapes) against three specifications, three designs, two sample sizes, and three biomarker levels give $4 \times 3 \times 3 \times 2 \times 3 = 216$ cells. The $c_{bm} = 0$ row is the Type I error check against nominal 5%, and the $t_{1/2} = 0$ row,

with no carryover, gives a best-case power reference at each combination of N , design, and c_{bm} against which the cost of nonzero carryover can be read off directly.

2.7 Inference and the cluster-robust correction

The interaction test described so far is model-based. The standard error of $\beta_{bm:X}$ is read from the fitted `lme` object, valid only if the assumed residual covariance is correct. That assumption is not innocuous here. The working covariance combines a patient random intercept with a `corCAR1` AR(1) residual correlation, and the companion calibration study⁶ shows this does not match the covariance the data-generating mechanism actually induces, inflating the standard error by roughly six to ten percent and depressing the rejection rate correspondingly. We quantify this with the calibration ratio κ , the mean estimated SE divided by the empirical SD of the interaction estimate across replicates ($\kappa = 1$ is correctly calibrated, $\kappa > 1$ conservative). At the null cell κ runs 1.05 to 1.10 across specifications, with empirical Type I error 0.029 to 0.039 against nominal 0.05.

The remedy is a cluster-robust (sandwich) variance estimator, which derives the variance of $\hat{\beta}_{bm:X}$ from observed between-cluster variation rather than the assumed covariance, and so remains valid under mis-specification. Clusters are patients (eight occasions each, unique across randomization paths). Because the naive sandwich form is markedly downward-biased with few clusters ($N = 35$ or 70), we use the CR2 bias-reduced estimator of Bell and McCaffrey, paired with Satterthwaite degrees of freedom (both via `clubSandwich`, replacing only the standard error, reference distribution, and p -value, leaving point estimates untouched).

We apply the correction in two places. A null-cell diagnostic ($c_{bm} = 0$, exponential DGP, $t_{1/2} = 1.0$ weeks, Hybrid, $N = 70$, 3,000 replicates) establishes whether CR2 restores nominal size (Section 3.5). A five-cell block, the reference cell plus four stress cells (small N , high autocorrelation, half-life mis-specification, 30% MCAR dropout, 500 replicates each, labeled S6 in this manuscript's block numbering), then asks whether restoring size disturbs the ranking (same section). Blocks S1 through S4 are the Tier 2 sensitivity analyses of Section 3.6. Block S5, an exploratory autocorrelation-by-carryover interaction, is outside the present scope and reported only in the archived full-scope drafts.

2.8 Performance measures and computing environment

For each cell we report Power, Type I error, Bias, Empirical SE (calibrated scale), Coverage of the nominal 95% Wald CI, MSE, non-convergence rate, and the MCSE of each, following Morris et al.¹, Tables 5–6. We chose $n_{sim} = 500$ so the MCSE on power at $p = 0.5$ would not exceed 0.022. No simulation runs for this manuscript. We filter the same 540-cell production grid summary used in the archived full-scope drafts down to the 216-cell slice defined by covariance moderation and the two non-linear decay forms. All three specifications are fit to the same simulated dataset within each cell (common random numbers), so every pairwise contrast, especially Lag-adjusted minus Unadjusted, carries reduced Monte Carlo variance relative to independent draws. The Tier 2 sensitivity blocks and S6 recalibration are filtered from the same shared production outputs. Exact file paths and versions are given in the Reproducibility section.

Table 2: Headline performance at the reference cell (exponential DGP, Hybrid design, $N = 70$, $c_{bm} = 0.45$, $n_{sim} = 500$). Power is the rejection rate of $H_0 : \beta_{bm:X} = 0$ at $\alpha = 0.05$. Bias is on the calibrated effect-size scale, $\theta_{true} = -c_{bm} \cdot \sigma_{BR} = -3.6$ (per docs/19 calibration). MCSEs are reported per Morris et al. (2019, Table 6).

t1/2	Specification	Power (MCSE)	Bias (MCSE)	Empirical SE	Coverage (MCSE)	MSE (MCSE)	Non-conv
0.0	Unadjusted	0.988 (0.005)	-0.029 (0.033)	0.746	0.990 (0.004)	0.557 (0.034)	0.000
0.0	Lag-adjusted	0.990 (0.004)	-0.025 (0.033)	0.746	0.990 (0.004)	0.557 (0.034)	0.000
0.0	Exposure-weighted	0.988 (0.005)	-0.029 (0.033)	0.746	0.990 (0.004)	0.557 (0.034)	0.000
0.5	Unadjusted	0.918 (0.012)	+0.586 (0.036)	0.807	0.938 (0.011)	0.995 (0.058)	0.000
0.5	Lag-adjusted	0.920 (0.012)	+0.583 (0.036)	0.812	0.938 (0.011)	1.000 (0.059)	0.000
0.5	Exposure-weighted	0.948 (0.010)	+0.072 (0.041)	0.907	0.972 (0.007)	0.828 (0.057)	0.000
1.0	Unadjusted	0.766 (0.019)	+1.093 (0.038)	0.847	0.828 (0.017)	1.911 (0.093)	0.000
1.0	Lag-adjusted	0.770 (0.019)	+1.099 (0.038)	0.853	0.818 (0.017)	1.935 (0.094)	0.000
1.0	Exposure-weighted	0.860 (0.016)	-0.081 (0.049)	1.085	0.968 (0.008)	1.184 (0.077)	0.000

3 Results

All results reported below derive from the same 500-replicate production run used in the full-scope manuscripts, filtered to the covariance-moderation architecture and the exponential and Weibull decay forms. The Monte Carlo standard error on power does not exceed 0.022 at $p = 0.5$, and differences smaller than this should not be over-interpreted. Coverage of the nominal 95% Wald confidence interval is reported for the reference cell, and non-convergence rates are reported per cell in the tables that follow.

3.1 Headline performance table

Table 2 reports performance at the reference cell (exponential DGP, Hybrid design, $N = 70$, $c_{bm} = 0.45$), where non-convergence was zero throughout. The bias, MSE, and coverage columns are computed against Exposure-weighted's calibrated target, $\theta_{true} = -c_{bm}\sigma_{BR} = -3.6$, while Unadjusted and Lag-adjusted estimate a different coefficient, $bm:D_{it}$. Their apparent bias and under-coverage reflect that estimand mismatch, not a defect in the estimators, so those columns are comparable within the Unadjusted/Lag-adjusted pair but not across to Exposure-weighted. For the three-way comparison we rely on power and Type I error, properties of the test rather than the coefficient tested.

3.2 Tier 1 grid overview

Figures 1 through 3 give a full-grid view of the Tier 1 results as annotated heatmaps, with trial design as column facets, N as row-block facets, analysis specification as the row axis within each block, and a blue (low power) to red (high power) fill scale. Each panel fixes two of the grid's remaining dimensions and varies a third, since no single two-dimensional heatmap can display all four simultaneously.

Three features are visible at a glance that the headline table and line plots elsewhere in this section make only cell-by-cell. First, within every design x N block, the Unadjusted and Lag-adjusted rows are numerically identical to two decimal places at every biomarker level, half-life, and decay shape, the clean null of Section 3.3 displayed across the entire grid at once rather than at one reference cell. Second, the Exposure-weighted row visibly separates from




Figure 1: Tier 1 grid, biomarker-effect axis: power by trial design (columns), N (row blocks), and analysis specification (rows within block), across $c_{bm} \in \{0, 0.30, 0.45\}$ at $t_{1/2} = 1.0$ weeks, exponential DGP.




Figure 2: Tier 1 grid, carryover-axis: power by trial design (columns), N (row blocks), and analysis specification (rows within block), across $t_{1/2} \in \{0, 0.5, 1.0\}$ weeks at $c_{bm} = 0.45$, exponential DGP.




Figure 3: Tier 1 grid, decay-shape axis: power by trial design (columns), N (row blocks), and analysis specification (rows within block), across the four DGP decay-shape levels (exponential, Weibull $k \in \{0.7, 1.0, 1.5\}$) at $c_{bm} = 0.45$, $t_{1/2} = 1.0$ weeks.

the other two only in the Hybrid and OL+BDC columns, and only once carryover is present (Figure 2); in the CO column it is weaker than the other two at every half-life past zero, the design-dependence of Section 3.4 read directly off the grid. Third, Figure 3 shows the Exposure-weighted row's advantage in the Hybrid and OL+BDC columns holding across every decay shape examined, including the heaviest-tailed Weibull form ($k = 0.7$), the decay-form robustness of Section 3.5 shown across the full grid rather than at one cell.

3.3 Matched-decay baseline

At the matched cell (exponential DGP, exposure-weighted analysis, $N = 70$, $c_{bm} = 0.45$, Hybrid design), power declines from 0.988 ± 0.005 at $t_{1/2} = 0$ to 0.948 ± 0.010 at $t_{1/2} = 0.5$ to 0.860 ± 0.016 at $t_{1/2} = 1.0$ weeks (point estimate plus or minus MCSE, $n_{sim} = 500$). All 500 replicates converged at every half-life in this cell.

We had previously read the value at $t_{1/2} = 1.0$ as agreeing with the power of 0.82 reported by Hendrickson et al. for a comparable cell². Unfortunately, we must withdraw that comparison. In the reference implementation, the analysis-side and data-generating carryover parameters are set independently, and the drivers that reproduce the original figures supply a half-life to the data-generating step only, leaving the analysis-side half-life at its default of zero, at which the exposure-decayed predictor collapses onto the binary on-drug indicator. The reference analysis is therefore Unadjusted, not Exposure-weighted, so the appropriate comparator is our unadjusted power at the same cell, 0.766, and the grids differ further on interaction strength

and half-life, so “the same cell” was approximate in any case. We have not resolved what the published figure was computed under, and make no claim of agreement or disagreement here. Nothing else in this manuscript depends on that comparison, since the three specifications are compared against one another within a single simulation under common random numbers.

3.4 Power by analysis specification

Power vs carryover under matched exponential DGP

Covariance-moderation architecture, $N = 70$, $c_{bm} = 0.45$, 500 reps/cell




Figure 4: Power versus carryover half-life under the exponential DGP, by analysis specification and trial design.

Two features of Figure 4 carry the substance of this study. First, the Unadjusted and Lag-adjusted curves lie on top of one another in every panel and at every half-life. At the reference cell they give 0.766 and 0.770 at $t_{1/2} = 1.0$ weeks, and across all 216 cells the mean difference is 0.003, with mean bias, MSE, and coverage agreeing to three decimal places. Since the two share an estimand, are fitted to the same simulated datasets, and differ by exactly one nuisance term, this is as clean a null as the design can produce. The lagged-treatment term contributes nothing on any measure (Section 4.1). Second, the three specifications give nearly identical power at $t_{1/2} = 0$, where there is no carryover to model, and diverge as the half-life increases, with Exposure-weighted separating from the other two. It retains a modest but consistent advantage in the Hybrid and OL+BDC designs from $t_{1/2} = 0.5$ weeks onward, but not under the crossover design, where it is markedly inferior to both alternatives (0.488 against 0.830 for Unadjusted). This appears to be because the within-subject correlation in a crossover trial already carries most of the signal the decaying predictor targets, so it gains little while its down-weighting of off-drug observations works against it (Section 4).

3.5 Decay-form mis-specification

The heatmap cell at the intersection of the exponential DGP row and the Exposure-weighted column is the matched reference, with power 0.860 ± 0.016 . The Unadjusted and Lag-adjusted columns, invariant to the assumed decay form by construction, are near-identical throughout,




Figure 5: Power at the critical cell ($t_{1/2} = 1.0$ weeks, Hybrid, $N = 70$, $c_{bm} = 0.45$) across DGP decay forms (rows) and specifications (columns). Fill scale is stretched to the observed range, not $[0,1]$, and not comparable to other figures.

as in Figure 4, so variation down those columns is data-generating decay and Monte Carlo noise. The remaining Exposure-weighted rows give the cost of assuming exponential decay when the truth is Weibull. Even at the heaviest-tailed form examined ($k = 0.7$), its advantage narrows but does not reverse (Section 4).

3.6 Type I error control and cluster-robust recalibration




Figure 6: Empirical Type I error under the null, pooled across decay forms and half-lives, by design and specification ($N = 70$, dashed line marks nominal 0.05).

Across the null cells, empirical Type I error sits uniformly below the nominal 0.05 level for all three specifications and designs (design-level means 0.032 to 0.041). This is not sampling noise but the conservatism of the model-based test (Section 2.6), from a working covariance that does not match the data-generating one. Because the conservatism is uniform across specifications, it does not disturb the comparison among them, and the block S6 recalibration




Figure 7: Same data as the boxplot above, as an annotated heatmap of the mean Type I error per design x specification cell (diverging scale centered at nominal 0.05: blue = conservative, red = anti-conservative).

Table 3: Conservatism of the model-based interaction test and its cluster-robust correction, at the null cell ($c_{bm} = 0$, exponential DGP, $t_{1/2} = 1.0$ weeks, Hybrid design, $N = 70$, 3,000 replicates). κ is the calibration ratio defined in Section 2.6. $\kappa > 1$ indicates a conservative test. All three specifications are conservative under the model-based test, most so Exposure-weighted ($\kappa = 1.10$, Type I = 0.029). The CR2 bias-reduced cluster-robust standard error restores κ to near one and Type I to nominal (0.045 to 0.049) for all three.

Specification	κ model	Type I model	κ CR2	Type I CR2
Unadjusted	1.06	0.036	0.98	0.048
Lag-adjusted	1.05	0.039	0.98	0.049
Exposure-weighted	1.10	0.029	1.00	0.045

below confirms that a CR2 cluster-robust standard error restores nominal rejection without changing which specification leads. This also establishes that the null result of Section 3.3 is not an artifact of differential test size, since Unadjusted and Lag-adjusted are conservative to the same degree. Table 3 makes this explicit at the null cell.

Figure 8 confirms this directly. The model-based rejection rate sits below nominal across all five S6 cells and all three specifications, while CR2 restores it to within the Monte Carlo band around 0.05, and in the information-rich cells the calibration ratio κ runs from 1.07 to 1.26 under the model-based test, restored to about one under CR2, recovering two to five points of power. Table 4 reports both alternatives against the unadjusted benchmark. The Exposure-weighted advantage is held or slightly widened under CR2 at the reference cell (+0.080 to +0.092), the small- N cell (+0.066 to +0.086), and the half-life mis-specification cell (+0.010 to +0.024). It narrows under high autocorrelation (+0.052 to +0.032, though it still leads), and is small under 30% MCAR dropout (+0.010), which we read as the S3 finding restated, since dropout removes the off-drug information the advantage depends on. The Lag-adjusted gap, by contrast, is within ± 0.010 of zero at every cell under either inference, corroborating the Section 3.3 null under a second procedure and at four stress cells the main grid does not cover. This recalibration was validated only in the Hybrid, $N = 70$ cells, and should not be extended to the crossover design, where Exposure-weighted does not lead in the first place.

Table 4: Cluster-robust recalibration at the reference cell and four stress cells (500 replicates, $c_{bm} = 0.45$). EW is Exposure-weighted, LA is Lag-adjusted, U is Unadjusted. κ is the calibration ratio of Section 2.6, here for Exposure-weighted. A value above one is a conservative test. The two gap pairs give each alternative against the unadjusted benchmark under each inference. CR2 df is the mean Satterthwaite degrees of freedom for the cluster-robust test on Exposure-weighted. Computed with `clubSandwich`.

Cell	κ model	κ CR2	EW model	EW CR2	EW-U model	EW-U CR2	LA-U model	LA-U CR2	CR2 df
Reference, $N = 70$	1.15	0.99	0.829	0.863	+0.080	+0.092	-0.003	+0.002	23
Small $N = 35$	1.13	0.96	0.558	0.590	+0.066	+0.086	-0.002	+0.000	12
High autocorr, $\rho = 0.9$	1.26	0.94	0.954	0.980	+0.052	+0.032	+0.000	-0.004	23
Half-life mis-spec	1.14	1.00	0.759	0.795	+0.010	+0.024	-0.003	+0.002	23
30% MCAR dropout	0.99	0.85	0.148	0.126	+0.010	+0.010	-0.004	+0.010	5

S6: type-I error at the null ($c_{bm} = 0$), model-based vs CR2

Dashed line = nominal 0.05; grey band = binomial 95% MC interval at 500 reps




Figure 8: Empirical Type I error at the null ($c_{bm} = 0$) under model-based and CR2 cluster-robust inference, across the five S6 stress cells and the three specifications. The model-based test (open) is conservative (near 0.03), and CR2 (filled) is restored to the nominal 0.05 (dashed), within the binomial Monte Carlo band (grey). 500 null replicates per cell.




Figure 9: Same data as the dot plot above, as two annotated heatmaps (model-based, CR2) sharing one diverging scale centered at nominal 0.05.

3.7 Sensitivity analyses

Eight Tier 2 blocks (S1-S4, S7, S8, S9, S10) perturb one axis at a time around a common reference cell (Hybrid, $N = 70$, $c_{bm} = 0.45$, exponential DGP, $t_{1/2} = 1.0$), holding the remaining factors fixed, except S9, which changes the trial design itself; a ninth, S6, applies the cluster-robust correction of Section 3.5 at that same cell plus four stress cells, and a tenth, S5, is exploratory and reported only in the archived full-scope drafts. The blocks are reported separately here and are not pooled with Tier 1 or with each other.

3.7.1 S1: within-factor autocorrelation




Figure 10: Power by analysis specification across AR(1) autocorrelation $\rho \in \{0.5, 0.7, 0.8, 0.9\}$.




Figure 11: Block S1: power by AR(1) autocorrelation and analysis specification, same data as the line plot above.

All three specifications gain power monotonically with ρ , from 0.798 ± 0.018 (Exposure-weighted) and 0.738 ± 0.020 (the other two, identical to three decimals) at $\rho = 0.5$, to 0.956 ± 0.009 against 0.906 to 0.910 at $\rho = 0.9$. The Exposure-weighted advantage, roughly 0.05 to 0.08 in absolute power, is preserved at every autocorrelation level examined, suggesting autocorrelation strength and specification choice act additively rather than interacting.

504 Lag-adjusted tracks Unadjusted throughout, never differing by more than 0.008.

505 3.7.2 S2: analyst-truth half-life mismatch

S2: Cost of analyst-truth half-life mismatch

Only Exposure-weighted consumes the assumed half-life; the other two are invariant




Figure 12: Power as the analyst's assumed half-life departs from the true DGP half-life. Only Exposure-weighted consumes the assumed half-life. Unadjusted and Lag-adjusted are invariant to it by construction and appear as flat reference lines.




Figure 13: Block S2: power by analyst-assumed half-life and analysis specification, faceted by true half-life, same data as the line plot above.

506 Block S2 crosses two true half-lives ($t_{1/2} \in \{0.5, 1.0\}$ weeks) with four analyst-assumed values
 507 ($\hat{t}_{1/2} \in \{0.25, 0.5, 1.0, 2.0\}$ weeks), giving eight cells. Only Exposure-weighted consumes $\hat{t}_{1/2}$,
 508 so the cost of a wrong assumption must be read within a true half-life rather than pooled
 509 across the block.

510 The penalty is modest. At a true half-life of 1.0 weeks, Exposure-weighted attains 0.832
 511 correctly specified, versus 0.766, 0.808, and 0.826 at assumed values of 0.25, 0.5, and 2.0 weeks.
 512 The worst case, a fourfold under-assumption, costs 6.6 points and still leaves it above both
 513 alternatives (0.746, 0.742). At a true half-life of 0.5 weeks it attains 0.934 correctly specified,

against 0.964, 0.914, and 0.896 at the same four assumed values. Because Unadjusted and Lag-adjusted cannot depend on $\hat{t}_{1/2}$, their variation along that axis is pure Monte Carlo error, ranging 0.910 to 0.960 at a true half-life of 0.5 weeks. This roughly four- to five-point noise floor explains the apparent anomaly that assuming 0.25 weeks (0.964) beats correct specification (0.934), a difference well inside it. It also qualifies the one case where the direction is worth noting. At a true half-life of 0.5 weeks with an assumed value of 2.0 weeks, Exposure-weighted falls to 0.896, below both alternatives (0.926, 0.930), so a fourfold over-assumption erodes its advantage to a tie.

The practical reading is that an approximate half-life is adequate, with an asymmetry. Under-assuming is safer than over-assuming. This follows from the limiting-case relationship of Section 2.4, since shrinking $\hat{t}_{1/2}$ drives the predictor toward the binary indicator, which remains serviceable, whereas inflating it drives the predictor toward a constant, at which the on-off contrast and the estimand itself degenerate.

3.7.3 S3: dropout

S3: Sensitivity to dropout

MCAR and MAR (severity-biased); reference cell as above




Figure 14: Power at dropout rates $\in \{0, 0.1, 0.2, 0.3\}$ under MCAR and MAR mechanisms.

All three specifications lose power monotonically as dropout increases. At the MCAR baseline ($N = 70$, no dropout), Exposure-weighted gives 0.866 ± 0.015 against 0.764 ± 0.019 for the other two. At 30% MCAR dropout, all three collapse to 0.122 to 0.148. The Exposure-weighted advantage narrows steadily as dropout worsens, from about 0.10 absolute at zero dropout to 0.07 at 10%, 0.04 at 20%, and 0.03 at 30% under MCAR, narrowing further under MAR at the highest rate. This is consistent with the advantage originating in the exposure-weighted contrast from off-drug observations, which contracts as those observations are preferentially lost. The advantage is preserved at dropout of 20% or below. Lag-adjusted again tracks Unadjusted throughout, the largest discrepancy being 0.010.




Figure 15: Block S3: power by dropout rate and analysis specification, faceted by dropout mechanism, same data as the line plot above. The MAR arm has no 0% cell (Section 3.6); it shares the MCAR reference cell.




Figure 16: Power across the biomarker-moderation effect-size range $c_{bm} \in \{0, 0.10, 0.20, 0.30, 0.45, 0.60\}$. The $c_{bm} = 0$ point is the null for Type I verification.




Figure 17: Block S4: power by biomarker-moderation effect size and analysis specification, same data as the line plot above.

Table 5: Block S7: all five G1-G5 specifications at the reference cell (Hybrid, $N = 70$, $c_{bm} = 0.45$, $t_{1/2} = 1.0$, exponential DGP, $n_{sim} = 500$, all replicates converged). G4 targets bm:Dbc at the AIC-selected half-life (candidate grid 0.25, 0.5, 1.0, 2.0); G5 targets the slope of the patient-level diff bm regression. Neither is the same estimand as bm:Db (G1/G2), so only power is comparable across rows. Mean selected t1/2 is the average AIC-chosen half-life for G4 across replicates, against a true half-life of 1.0.

Specification	Power (MCSE)	Mean selected t1/2
G1: Unadjusted	0.748 (0.019)	–
G2: Lag-adjusted	0.748 (0.019)	–
G3: Exposure-weighted	0.826 (0.017)	–
G4: AIC-selected t1/2	0.832 (0.017)	1.31
G5: Paired-difference	0.804 (0.018)	–

3.7.4 S4: biomarker-moderation effect-size curve

The power curve is sigmoidal for all three specifications. At $c_{bm} = 0$, empirical Type I error sits below nominal for each (0.026 to 0.030). Power then rises through the mid-range and saturates once $c_{bm} \geq 0.45$, reaching 0.916 to 0.958 at $c_{bm} = 0.6$. The Exposure-weighted advantage over the unadjusted benchmark is largest at intermediate effect sizes, +0.082 absolute at $c_{bm} = 0.30$ and +0.102 at $c_{bm} = 0.45$, shrinking to +0.042 at $c_{bm} = 0.60$ as all three approach the ceiling, so the advantage matters most for trials targeting a moderate biomarker-treatment interaction. Lag-adjusted and Unadjusted coincide across the entire range, differing by at most 0.008.

3.7.5 S7: alternative-paradigm specifications

An earlier version of this block tested two carryover terms crossed with the biomarker, E5 ($Sx \sim bm + t + D_{it} + bm:D_{it} + L_{it} + bm:L_{it}$) and E6 ($Sx \sim bm + t + D_{it} + bm:D_{it} + t_{sd} + bm:t_{sd}$), motivated by Section 4.1's observation that a carryover term entering only as a nuisance main effect cannot repair the interaction's attenuation. Both underperformed Unadjusted at every cell examined (evaluated by a joint 2-df Wald test, since each adds a second interaction coefficient), traced to a structural cause discussed in Section 4.5: this manuscript's DGP has a single true decay parameter, so a second, DGP-unmotivated interaction coefficient mostly adds estimation noise. E5/E6 were replaced by two specifications addressing the same open problem, the analyst's unknown decay half-life, from different angles (Section 2.4): E7 (AIC-selected half-life) and E9 (paired-difference regression). Both are single-coefficient tests, unlike E5/E6.

Table 5 gives a genuinely different result from the retired E5/E6 comparison: both G4 and G5 beat Unadjusted at this cell, and G4 edges out G3 (Exposure-weighted) as well. G4 leads overall (0.832), G3 follows (0.826), then G5 (0.804, the two-stage paired-difference regression with no repeated-measures model at all), and G1/G2 tie at the back (0.748 each, the already-established Section 3.3 null replicated once more). G4's AIC selection is imperfect but usefully informative: the mean selected half-life is 1.31 weeks against a true value of 1.0, and the selection distribution concentrates on the true value and its immediate neighbor (88% of replicates choose 1.0 or 2.0 weeks), close enough to the truth that G4 slightly exceeds G3's own

S7: G1–G5 at the Hybrid reference cell




Figure 18: Block S7: power by analysis specification, same figures as the table above.

oracle-half-life power at this cell. That **G5**, which discards the repeated-measures structure, the correlation model, and any carryover representation entirely, still beats **G1/G2** is the more striking result, echoing Senn, Julious and Araujo’s⁵ finding that simplicity is not obviously costly for N-of-1 treatment-effect estimation. None of these results should be read as general, however: Block S9 tests whether the ranking survives a different trial design, and it does not for every specification equally (below).

3.7.6 S8: architecture sensitivity of the specification ranking

Every result reported to this point, Tier 1 and blocks S1-S7 alike, is restricted to Architecture B (covariance moderation), the manuscript’s stated scope (Section 2.3). Block S8 asks a narrower question than a full architecture comparison, which is the purpose of the companion manuscript on DGP architectures³: does the ranking among analysis specifications established above under Architecture B hold under Architecture A (mean moderation) as well, or is the choice of carryover representation itself architecture-dependent? We refit the five specifications of block S7’s table (**E1**, Lag-adjusted **E3**, Exposure-weighted **E2**, **E7**, **E9**) at the same reference cell under both architectures, at three half-lives ($t_{1/2} \in \{0, 0.5, 1.0\}$) rather than the single reference value, so the check covers whether the ranking is architecture-dependent uniformly or only once carryover is severe enough. Because c_{bm} is not calibrated to equal power across architectures, the two produce materially different absolute power at the same nominal value (Section 1;³ shows the mean-moderation channel requires several times the nominal effect size for a comparable power loss under carryover); only the within-architecture ranking is the object of interest here.

Table 6 and Figures 19/ 20 show the S7 result (**E7** and **E9** both beating **E1**) is architecture-specific, joining Exposure-weighted in a pattern that grows with carryover severity rather than switching on at some threshold. At $t_{1/2} = 0$ every specification collapses to nearly the same value in both architectures (no carryover to represent), so **E1**, **E2**, and **E7** agree there by construction (all 0.97-0.98 under B, 0.97 under A); **E9**’s slightly lower value even at $t_{1/2} = 0$ (0.974 under B, 0.942 under A) reflects sampling variance in its two-stage estimator rather than a carryover effect. Under Architecture B, the ranking at $t_{1/2} = 1.0$ is **E7** (0.850) > **E9** (0.842) > **E2** (0.836) > **E1/E3** (0.760): both new specifications not only beat the unadjusted benchmark, as in block S7, but slightly exceed Exposure-weighted, the manuscript’s own current recommendation. At $t_{1/2} = 0.5$ the same ordering holds but compressed (**E7** 0.936 essentially ties **E2** 0.936; **E9** 0.934 close behind; **E1/E3** 0.920-0.922 trail by about 1.5 points).

Table 6: Block S8: analysis-specification ranking by DGP architecture and carryover half-life, at the reference cell (Hybrid, $N = 70$, $c_{bm} = 0.45$, exponential DGP, $n_{sim} = 500$). Absolute power is not comparable between architectures at the same nominal c_{bm} ; only the within-architecture ranking is of interest.

Architecture	t1/2	Specification	Power (MCSE)	Non-conv
B (covariance)	0.0	Unadjusted	0.982 (0.006)	0.000
B (covariance)	0.0	Lag-adjusted	0.980 (0.006)	0.000
B (covariance)	0.0	Exposure-weighted	0.982 (0.006)	0.000
B (covariance)	0.0	AIC-selected t1/2	0.980 (0.006)	0.000
B (covariance)	0.0	Paired-difference	0.974 (0.007)	0.000
B (covariance)	0.5	Unadjusted	0.922 (0.012)	0.000
B (covariance)	0.5	Lag-adjusted	0.920 (0.012)	0.000
B (covariance)	0.5	Exposure-weighted	0.936 (0.011)	0.000
B (covariance)	0.5	AIC-selected t1/2	0.936 (0.011)	0.002
B (covariance)	0.5	Paired-difference	0.934 (0.011)	0.000
B (covariance)	1.0	Unadjusted	0.760 (0.019)	0.000
B (covariance)	1.0	Lag-adjusted	0.760 (0.019)	0.000
B (covariance)	1.0	Exposure-weighted	0.836 (0.017)	0.000
B (covariance)	1.0	AIC-selected t1/2	0.850 (0.016)	0.000
B (covariance)	1.0	Paired-difference	0.842 (0.016)	0.000
A (mean)	0.0	Unadjusted	0.972 (0.007)	0.000
A (mean)	0.0	Lag-adjusted	0.972 (0.007)	0.000
A (mean)	0.0	Exposure-weighted	0.972 (0.007)	0.000
A (mean)	0.0	AIC-selected t1/2	0.970 (0.008)	0.000
A (mean)	0.0	Paired-difference	0.942 (0.010)	0.000
A (mean)	0.5	Unadjusted	0.968 (0.008)	0.000
A (mean)	0.5	Lag-adjusted	0.968 (0.008)	0.000
A (mean)	0.5	Exposure-weighted	0.966 (0.008)	0.000
A (mean)	0.5	AIC-selected t1/2	0.970 (0.008)	0.002
A (mean)	0.5	Paired-difference	0.960 (0.009)	0.000
A (mean)	1.0	Unadjusted	0.972 (0.007)	0.000
A (mean)	1.0	Lag-adjusted	0.976 (0.007)	0.000
A (mean)	1.0	Exposure-weighted	0.946 (0.010)	0.000
A (mean)	1.0	AIC-selected t1/2	0.950 (0.010)	0.000
A (mean)	1.0	Paired-difference	0.964 (0.008)	0.000




Figure 19: Block S8, Panel A (Architecture A, mean moderation): power by carryover half-life (rows) and analysis specification (columns). Same data as the Architecture A rows of the table above.




Figure 20: Block S8, Panel B (Architecture B, covariance moderation): power by carryover half-life (rows) and analysis specification (columns). Same data as the Architecture B rows of the table above. Not directly comparable to Panel A in absolute terms (Section 3.6).

Under Architecture A the pattern reverses for E7 and E2 alike. At $t_{1/2} = 1.0$, E1/E3 lead (0.972/0.976), E9 trails only slightly (0.964, a 0.8-point gap), while E2 (0.946) and E7 (0.950) trail more substantially (about 2.2 to 2.6 points, roughly 1.5 MCSE and so a real, if modest, reversal rather than pure noise). This matches the pattern already reported from the full-scope grid in the archived drafts, that Exposure-weighted is marginally dominated in every design examined under mean moderation (Section 4.6); E7 inherits that same architecture-sensitivity since it is built on Exposure-weighted's formula. E9, built on nothing but a raw on-drug/off-drug contrast, is comparatively more robust to architecture, small under B, smaller still under A, though never reversing to a large disadvantage either way. Lag-adjusted (E3) remains statistically indistinguishable from Unadjusted (E1) under both architectures at every half-life (largest gap 0.004, well under one MCSE), extending Section 3.3's clean null to Architecture A as well. The one specification recommendation this manuscript makes, exposure weighting in the Hybrid and OL+BDC designs (Conclusions), is therefore itself conditional on Architecture B and grows more important as carryover worsens; the same qualification now extends to E7. E9 is the one specification in this table whose advantage, where it exists, does not appear to depend on which architecture generated the data, though it was tested only at this one design and reference cell (Block S9 tests one further design).

3.7.7 S9: does the reference-cell ranking survive under CO?

Blocks S7 and S8 found G4 and G5 both beat Unadjusted at the Hybrid reference cell under Architecture B, with G4 slightly exceeding Exposure-weighted (G3) itself. Block S9 asks whether that ranking is specific to Hybrid's short, compressed washout, by refitting all five G1-G5 specifications under the CO design (off-drug arm spanning 2.5 to 20 weeks post-discontinuation, Section 2.2) at two true half-lives.




Figure 21: Block S9: power under the CO design at two true half-lives, all five G1-G5 specifications.

It is not, and including G3 in this block changes the diagnosis reached from G4 alone. At $t_{1/2} = 1.0$, G1/G2 lead (0.834/0.838); G5 trails moderately (0.750); G3 and G4 both collapse, to 0.482 and 0.474 respectively, statistically indistinguishable from each other (gap 0.008, well under one MCSE). At $t_{1/2} = 0.5$ the same pattern holds (0.866/0.864 for G1/G2, 0.778 for G5, 0.546/0.552 for G3/G4, again indistinguishable). Because G3 uses the *true* half-life, not an estimated one, its near-identical collapse shows that G4's failure under CO is inherited almost entirely from Exposure-weighted's own established weakness under this design (Section 3.4:

0.488 versus 0.830 in the Tier 1 grid), not created by the AIC-selection step. The selection step does have a real, separately confirmed pathology under CO: it concentrates overwhelmingly on the largest candidate half-life in the grid (2.0 weeks, chosen in 483 of 500 replicates at $t_{1/2} = 1.0$ and 435 of 500 at $t_{1/2} = 0.5$, regardless of which true half-life generated the data), consistent with a weakly informative likelihood surface across candidates when off-drug occasions are few and widely spaced. But that mis-selection costs **G4** essentially nothing *beyond* what **G3** already loses under CO with the correct half-life in hand, so the dominant story is the continuous decayed predictor’s mismatch with CO’s design generally (Section 3.4), with AIC selection an identified but practically minor addition to it here, not the primary cause it appeared to be from **G4** in isolation. **G5**’s milder reversal has no comparable mechanical explanation: it simply reflects that CO’s own within-subject contrast (Section 3.4) already captures most of the signal a between-patient paired-difference regression is built to extract, so discarding the repeated-measures structure costs more there than it does under Hybrid. None of **G3**, **G4**, or **G5**’s Hybrid-cell behavior is design-general, then, and the Conclusions’ design-specific guidance (Section 3.6’s “what analysis is best” discussion) should be read accordingly: **G3/G4** should not be used under CO regardless of whether the half-life is assumed or estimated.

3.7.8 S10: does a CR2 correction change the ranking among G1-G5?

Section 3.5 (Block S6) showed the model-based interaction test is mildly conservative and that a CR2 cluster-robust standard error recovers two to five points of power without disturbing the ranking among **E1**, **E2**, and **E3**. That correction had not been applied to **E7** or **E9**. This block asks two questions at once, adding a compact naming layer (**G1-G9**) for the resulting nine-way comparison: does CR2 help **E7** and **E1/E3** the way it already helps **E2** (Section 3.5), and does it change which specification leads at the Hybrid, Architecture B reference cell? **G1-G5** are the five specifications introduced so far (**E1**, **E3**, **E2**, **E7**, **E9**, respectively); **G6-G9** cross the CR2 correction with **G1**, **G2**, **G4**, and **G3**. **G8** (**E2**+CR2) already exists as Block S6’s reference-cell result and is included here from that independent run (different seed, not sharing common random numbers with **G1-G7/G9**, so its comparison to the rest of this panel is not as tightly paired as the comparisons among **G1-G7/G9** are with each other). A tenth combination, **G5**+CR2, is not included: **E9**’s paired-difference regression collapses to one row per patient before fitting, so there is no repeated-measures cluster structure left for CR2 to correct.

S10: G1–G9, Hybrid reference cell




Figure 22: Block S10: power for G1 through G9 at the Hybrid, Architecture B reference cell, same figures as the table above.

Table 7: Block S10: G1-G9 at the reference cell (Hybrid, $N = 70$, $c_{bm} = 0.45$, $t_{1/2} = 1.0$, exponential DGP, Architecture B, $n_{sim} = 500$). G1-G7 and G9 share common random numbers within this block; G8 is from the independent Block S6 run (different seed) and is not paired with the rest.

Code	Method	Power (MCSE)	Non-conv
G1	Unadjusted	0.748 (0.019)	0.000
G2	Lag-adjusted	0.748 (0.019)	0.000
G3	Exposure-weighted	0.826 (0.017)	0.000
G4	AIC-selected t1/2	0.832 (0.017)	0.000
G5	Paired-difference	0.804 (0.018)	0.000
G6	Unadjusted + CR2	0.776 (0.019)	0.010
G7	Lag-adjusted + CR2	0.764 (0.019)	0.016
G8	Exposure-weighted + CR2	0.863 (source: Block S6)	–
G9	AIC-selected t1/2 + CR2	0.870 (0.015)	0.000

Table 7 and Figure 22 answer both questions together. CR2 helps every specification it was applied to, consistent with Section 3.5: G1/G2 (Unadjusted and Lag-adjusted, both 0.748 model-based) rise to G6 0.776 and G7 0.764, gains of 2.8 and 1.6 points respectively, in line with the two to five points Block S6 already documented for E2. The ranking among the five base specifications is unchanged by these particular gains, since G1/G2 remain well below G3-G5 regardless. What does change the ranking is G9. CR2 lifts G4 (AIC-selected half-life, 0.832 model-based) to 0.870, a 3.8-point gain, larger than any other CR2 correction observed here or in Block S6, and enough for G9 to overtake every other specification tested, including G8 (0.863, the Block S6 reference-cell result for Exposure-weighted+CR2). G9's gain is also notable for what did not happen: E7's AIC-selection step raised a Type I error concern earlier in this manuscript's development (post-selection inference, since the same data are used to choose the half-life and then test the resulting coefficient), and one might have expected a CR2 correction, which widens standard errors, to partly offset an anti-conservative test by coincidence. Instead every replicate converged under G9 (zero non-convergence, against 1.0% and 1.6% for G6/G7), and the power gain is the largest observed, which is at least consistent with (though does not itself establish) CR2 correcting real conservatism rather than papering over a different problem. This block did not test the null ($c_{bm} = 0$), so it does not resolve whether G4/G9's Type I error is controlled; that remains an open question raised earlier and not addressed by this comparison.

3.8 A higher-precision check

Because a number of the Tier 1 gaps are small, we reran a focused 24-cell grid (OL+BDC design, $N \in \{35, 70\}$, true half-life $\in \{0.5, 1.0\}$) at 600 trials per cell, all converging, and compared specifications with a paired McNemar test, considerably more sensitive than an independent-groups comparison since all three are fit to the same trials. Type I error across the twelve null cells fell in $[0.003, 0.053]$, so nominal alpha holds. Paired McNemar tests place Exposure-weighted ahead of the unadjusted benchmark at all four alternative cells, by +2.3

684 to +13.7 percentage points, with the largest gap at the highest-leverage cell ($N = 70$, true
685 $t_{1/2} = 1.0$ weeks, $p = 6.6 \times 10^{-17}$) and $p \leq 0.015$ at the remaining three, a sharper confirmation
686 than the Tier 1 grid alone provides.

687 One scope caveat governs the rerun's third arm. It was produced by the package's
688 `lme_analysis()` routine, whose carryover option appends time since discontinuation as a
689 linear main effect rather than the lagged indicator L_{it} used in the Tier 1 grid. It is therefore
690 not the Lag-adjusted specification of Section 2.4 and is not reported as such here. That
691 alternative remains a reasonable specification we have not evaluated (Section 4.6).

692 4 Discussion

693 4.1 The lagged-treatment term does no work

694 The clearest result of this study is a negative one. Adding the classical lagged-treatment
695 indicator to the analysis model changed nothing measurable, anywhere in the grid. Across
696 all 216 cells the mean power difference against the unadjusted benchmark was 0.003 against
697 an MCSE of 0.018, bias, MSE, and coverage agreed to three decimal places, and the equality
698 held in every sensitivity block and all five cluster-robust stress cells under both inference
699 procedures. The strength of this conclusion rests on the design rather than the grid size. The
700 two specifications estimate the same coefficient, were fitted to the same simulated datasets
701 under common random numbers, and differ by exactly one term, so there is no confounding
702 from a change of estimand, Monte Carlo variation between arms, or differential test size
703 (both are conservative to the same degree, Section 3.5). What is left is the term itself, and it
704 contributed nothing.

705 The mechanism follows from where the term sits. Carryover damages the interaction test on
706 two fronts (Section 1). It attenuates the coefficient, because contaminated occasions are coded
707 as drug-free, and it inflates the residual variance, because those occasions are heterogeneous in
708 a way the model does not represent. The lagged indicator enters only as a nuisance main effect,
709 with no $\text{bm}:L_{it}$ product term, so the tested interaction remains $\text{bm}:D_{it}$ and the attenuation
710 is untouched. It can in principle absorb some of the second effect, but evidently absorbs
711 too little to matter, since it marks a single occasion per randomization path and is constant
712 thereafter, unable to describe a graded decline across a two- to four-occasion washout.

713 A more capable, elapsed-time-indexed nuisance term might succeed where the crude lagged
714 indicator fails, so we ran a targeted check on the alternative the reference implementation
715 provides, $\text{Sx} \sim \text{bm} + \text{t} + \text{Db} + \text{bm}:\text{Db} + \text{tsd}$, called the washout-adjusted specification (E4),
716 across the three designs at true half-lives of 0.5 and 1.0 weeks ($N = 70$, $c_{bm} = 0.45$, exponential
717 DGP, 500 replicates). It does slightly better, and the improvement is instructive rather than
718 material. Type I error is at nominal for both (mean 0.050 versus 0.045), and power exceeds the
719 benchmark in all six cells by a mean of 1.3 points (sign test $p = 0.031$), reaching significance
720 under crossover at both half-lives (+1.8, $p = 0.008$; +2.2, $p = 0.003$) and OL+BDC at the
721 longer one (+2.2, $p = 0.015$), but not Hybrid (+1.0, $p = 0.27$). The mechanism is exactly the
722 one Section 2.4 predicts. Point estimates are unchanged, so nothing about the attenuation

is repaired, and the entire effect runs through the standard error, which only a long enough washout gives it something to describe. The crossover schedule places off-drug occasions 2.5 to 10 weeks after discontinuation, against one to two weeks for Hybrid, so the practical ceiling on what any carryover nuisance term can deliver is about two percentage points, and only in long-washout designs, against the nine to eleven points exposure weighting delivers where it works. This check covers only the exponential DGP at a single N and effect size, so it is a probe rather than a fourth arm, and shares every specification's structural limitation of entering carryover as a main effect only (Section 4.5). For practitioners the implication is direct. A carryover nuisance term of either form should not be expected to recover meaningful power lost to carryover, and reporting one does not constitute an effective adjustment for it.

4.2 The exposure-weighted advantage, and where it fails

The exposure-weighted predictor does help, but conditionally. At its best cell (Hybrid, $N = 70$, $c_{bm} = 0.45$, $t_{1/2} = 1.0$) it exceeds the unadjusted benchmark by about 10 points of power (0.860 versus 0.766, roughly five MCSEs and so not attributable to chance), confirmed at all four cells of the Section 3.7 high-precision rerun. Under the crossover design it performs considerably worse (0.488 versus 0.830), apparently because the within-subject correlation there already carries most of the signal the decaying predictor was designed to recover, so it gains little while its down-weighting of off-drug observations works against it. The recommendation therefore holds only in the two non-crossover designs, and should not be extended to a crossover trial, where the unadjusted binary indicator is materially better. That the one specification addressing attenuation is also the only one to gain power, while the one addressing only residual variance gains nothing, suggests attenuation is the mechanism that matters here, an interpretation consistent with the results rather than a finding the design isolates directly.

4.3 What the robustness checks add

The Section 3.6 sensitivity blocks show the exposure-weighted advantage, where it exists, is not fragile. It survives incorrect decay-shape assumptions (Figure 5), an incorrectly assumed half-life short of a large over-assumption at a short true half-life (Figure 12), and holds its five- to eight-point lead across the autocorrelation range examined (Figure 10). One caveat qualifies the Tier 1 figures throughout. Outside block S2 the analysis-side half-life is set equal to the data-generating one, granting Exposure-weighted an oracle no real analyst possesses, so its Tier 1 margins are upper bounds and block S2 gives the more realistic figure.

4.4 Dropout matters more than the method

The finding with the largest practical consequence concerns missing data, not carryover. When patients drop out (Figure 14), power collapses for all three specifications together, to between 0.12 and 0.25 at 30% dropout, essentially indistinguishable. No specification offers protection, since each can only organize whatever data remain. MAR dropout is somewhat worse than MCAR, because the sickest patients, who carry the most signal, tend to leave first. Retaining patients in the trial yields considerably more power than any choice among these three

762 specifications ever could.

763 4.5 What analysis is best in which situations, in practice

764 Collecting the findings above into a decision rule rather than a set of pairwise comparisons:

- 765 • **Covariance moderation (Architecture B), Hybrid or OL+BDC design.** Use
766 Exposure-weighted with a CR2 cluster-robust standard error. The advantage holds
767 under decay-shape mis-specification (Section 3.5), under half-life mis-specification short
768 of a large over-assumption (Section 3.6, Block S2), and across the autocorrelation range
769 examined (Block S1), and CR2 restores nominal Type I error without disturbing the
770 ranking (Block S6, validated at $N \in \{35, 70\}$ in the Hybrid design specifically; OL+BDC
771 was not itself stress-tested under CR2). This remains the recommended default: it has
772 been through the full robustness battery E7/E9 have not (next point).
- 773 • **E7 (AIC-selected half-life) and E9 (paired-difference) are promising at this**
774 **same cell, not yet validated defaults.** Both beat Exposure-weighted itself at the
775 Hybrid, Architecture B reference cell (Blocks S7-S8: E7 0.850, E9 0.842, E2 0.836, at
776 $t_{1/2} = 1.0$), and E7 in particular removes the need to assume a half-life at all. Neither
777 has been tested against decay-shape mis-specification, dropout, or under CR2 the way
778 E2 has, so this is grounds to investigate them further, not to displace the established
779 default yet.
- 780 • **Crossover (CO) design, any architecture.** Use Unadjusted. Exposure-weighted
781 is markedly worse under CO (Section 3.4), and Block S9 shows E7 and E9 fail here
782 too, E7 catastrophically. E7's AIC half-life selection systematically drifts to the largest
783 candidate in its grid under CO's sparse, widely spaced off-drug sampling (chosen in over
784 85% of replicates regardless of the true half-life), collapsing its power to roughly half of
785 Unadjusted's; it should specifically be avoided under any design with few, widely spaced
786 off-drug observations. E9 also underperforms Unadjusted under CO, though far more
787 mildly. The classical lagged-treatment remedy adds nothing here either (next point),
788 including in the one design where its washout is long enough to give it a chance (Section
789 4.1).
- 790 • **Architecture uncertain, or known to be mean moderation (Architecture A).**
791 Use Unadjusted or Lag-adjusted, not Exposure-weighted or E7. Both specifications'
792 advantage under covariance moderation reverses, marginally, under mean moderation
793 (Block S8). E9 is the one specification whose performance does not depend much on
794 which architecture generated the data, small under B, smaller under A, but it has been
795 tested at only one design and reference cell outside the Hybrid comparison (Block S9).
- 796 • **Never rely on Lag-adjusted as a carryover remedy.** It is statistically indistin-
797 guishable from Unadjusted everywhere tested in this manuscript, Tier 1 and every Tier
798 2 block alike (Section 4.1). If a reviewer or protocol expects some carryover term to be
799 present, including it costs nothing measurable, but it should not be reported as having
800 adjusted for carryover.
- 801 • **Minimize dropout before optimizing specification choice.** It is the single largest
802 lever in the entire grid (Section 4.4), and no specification, including E7/E9, has been

shown to offer any protection against it; their behavior under dropout is untested.

4.5.1 Why E7 and E9 succeed at Hybrid but fail under CO

Both new specifications' Hybrid, Architecture B advantage evaporates under the CO design (Block S9), for two different, design-specific reasons rather than a shared structural flaw of the kind that sank E5/E6 (Section 4.6). E7's failure is mechanical: with only a few, widely spaced off-drug occasions, the likelihood surface across candidate half-lives is weakly informative, and AIC selection drifts to the boundary of its candidate grid rather than locating the true value, systematically overestimating the half-life. Section 3.6's S2 block already showed over-assumption is the more damaging direction of half-life error when the half-life is *fixed* by the analyst; under CO, E7 reproduces that same damaging error through the estimation procedure itself, on nearly every replicate. This is a concrete, testable prediction: E7 should be reliable whenever a design gives several off-drug occasions across a graded range of elapsed time (as Hybrid does, despite its short overall washout), and unreliable whenever off-drug sampling is sparse, however long the washout runs in calendar time.

E9's failure has no comparable mechanical account and is milder. The most plausible explanation is that CO's own within-subject contrast (Section 3.4) already carries most of the signal a between-patient paired-difference regression is built to extract, so discarding the repeated-measures and correlation structure costs more under CO than it does under Hybrid, where that structure is doing more work. This was not tested directly and would require comparing E9 against the within-cell decomposition of E2's advantage to confirm.

Two implications follow. First, these results reinforce rather than undercut the earlier E5/E6 diagnosis (Section 4.6): the more fundamental unresolved question is not which design an alternative specification is tested under, but whether the underlying data-generating process rewards that specification's added structure at all, and neither E7 nor E9 was tested against a DGP built to favor them specifically. Second, the practical recommendation is narrower than "estimate the half-life" or "use a simpler method": both E7 and E9 are candidates worth validating further within the Hybrid/OL+BDC, Architecture B cell where this manuscript's other recommendations already concentrate, not general-purpose replacements for Exposure-weighted or Unadjusted.

4.6 Scope, related work, and limits

This manuscript deliberately narrows the full-scope grid to the covariance-moderation architecture and to the non-linear decay forms, in the interest of a self-contained evaluation of manageable length. The archived full-scope drafts (`archive/report.Rmd` and `archive/report-slim.Rmd`) show that the exposure-weighted advantage documented here does not carry over to mean moderation, where the biomarker moderates the mean response directly and Exposure-weighted is marginally dominated in every design examined. Readers who need that cross-architecture comparison should consult those drafts rather than extrapolate from this one. Establishing what the originally published power values (Section 3.2) were computed under remains outstanding work.

842 The crossover tradition surveyed by Jones and Kenward⁴ would lead one to expect a shape-
843 free lagged term to outperform a parametric predictor once the decay shape is mis-specified.
844 Our results do not support that expectation over the range examined, for two reasons. The
845 exposure-weighted penalty for an incorrect half-life is too small to be overtaken (Section 3.6),
846 and more fundamentally the lagged term does not improve on ignoring carryover at all (Section
847 4.1), leaving no advantage for decay-shape mis-specification to erode.

848 One adjacent specification bounds the negative result of Section 4.1. Every carryover term
849 considered there, lagged and washout-adjusted alike, enters as a nuisance main effect with
850 no product against the biomarker, so that finding is that a carryover *main effect* does not
851 recover the attenuated signal, not that no shape-free specification can. The natural candidate
852 is a carryover term interacted with the biomarker, either $Sx \sim \mathbf{bm} + \mathbf{t} + \mathbf{Db} + \mathbf{bm:Db} + \mathbf{L}$
853 $+ \mathbf{bm:L}$ or its washout-adjusted counterpart with $\mathbf{bm:tsd}$, letting the biomarker effect decline
854 across the washout without committing to a rate. An earlier version of Block S7 tested exactly
855 these two specifications and found a clean negative result at every cell examined, traced to a
856 structural cause (Section 4.5): this manuscript's DGP has a single true decay parameter, so a
857 second, interacted carryover coefficient adds estimation noise without a corresponding signal
858 to recover. They were replaced by two specifications addressing the same open problem, the
859 analyst's unknown decay half-life, from different angles: E7 (AIC-selected half-life) and E9 (a
860 paired-difference regression with no carryover representation at all). Both beat Unadjusted at
861 the Hybrid, Architecture B reference cell, E7 even edging out Exposure-weighted itself (Block
862 S7), but neither survives a change of trial design or architecture (Blocks S8-S9): E7's AIC
863 selection fails specifically under sparse off-drug sampling (Section 4.5), and E9's advantage
864 appears to depend on the repeated-measures structure it discards being otherwise underused,
865 which is design-dependent. Block S8 shows the Lag-adjusted null of Section 3.3 replicates
866 under Architecture A, so it is not an artifact specific to covariance moderation. Exposure-
867 weighted's advantage, and now E7's, are architecture-sensitive, reversing in direction, though
868 not substantially, under Architecture A (Section 3.6, Block S8).

869 Two further limitations bound what we can conclude. The numbers presented here come from
870 a single parameter set calibrated to a prazosin and PTSD trajectory, so absolute power values
871 should be read as illustrative and the ranking, itself conditional on trial design, as the more
872 portable result. And the sensitivity checks vary one factor at a time and can therefore miss
873 genuine interactions between factors, for instance between high autocorrelation and heavy
874 dropout. The one joint check available to us, block S5 in the full-scope manuscript, revealed
875 no such interaction.

876 5 Conclusions

- 877 1. **The lagged-treatment term does nothing.** It never improved on ignoring carryover
878 altogether, on any measure, in any design, at any half-life, and in every sensitivity block,
879 a clean null given the two specifications share an estimand and differ by a single term.
880 The term enters as a nuisance main effect only, leaving the attenuation of the interaction
881 untouched, and should not be reported as an adjustment for carryover.

- 882 2. **Exposure weighting helps, but conditionally.** It achieves the highest power, about
883 10 points over the unadjusted benchmark, only in the Hybrid and OL+BDC designs,
884 and its lead survives mis-specified decay shape and half-life there. Under crossover it is
885 markedly worse than the binary indicator (0.488 against 0.830) and should not be used.
886 Where used, we recommend a CR2 cluster-robust standard error (Block S6).
- 887 3. **An approximate half-life is adequate, but err short.** A fourfold error at a true
888 half-life of one week costs 6.6 points of power, so a rough pharmacokinetic estimate
889 suffices in practice. The error is asymmetric. Under-assuming drives the predictor
890 toward the still-serviceable binary indicator, while over-assuming erodes the advantage
891 to a tie. Outside block S2 the analyst is granted the true half-life, so the Tier 1 margins
892 are upper bounds.
- 893 4. **Dropout matters more than the choice of method.** Power falls sharply as dropout
894 increases, roughly halving at 20%, equally across all three specifications. Preventing
895 dropout is a more effective use of design effort than any choice among them.
- 896 5. **The interaction test is somewhat conservative, and correctly so.** All three
897 specifications reject slightly less often than nominal under the null. A CR2 cluster-
898 robust standard error restores the correct rate while recovering modest power, and the
899 exposure-weighted advantage is held or widened under CR2 at four of five S6 stress cells
900 (Block S6).
- 901 6. **This is a narrowed slice of a larger study.** Mean moderation and linear decay
902 fall outside this manuscript's scope. The archived drafts `archive/report.Rmd` and
903 `archive/report-slim.Rmd` are the authority on those comparisons and should be con-
904 sulted before generalizing beyond covariance moderation. We also make no claim of
905 numerical agreement with the originally published power values (Section 3.2).
- 906 7. **Carryover terms interacted with the biomarker do not help; two alternative-**
907 **paradigm specifications do, but only at the reference cell.** Crossing a lagged
908 or washout carryover term with the biomarker gave lower power than the unadjusted
909 benchmark at every cell examined (Block S7's original version), traced to a structural
910 cause: this manuscript's DGP has a single true decay parameter, so a second, interacted
911 coefficient adds estimation noise without a corresponding signal to recover (item 1's
912 mechanism, restated). They were replaced by two specifications answering the same
913 open question, the analyst's unknown decay half-life, differently: E7 estimates it by
914 AIC selection; E9 discards the decay representation, and the mixed model itself, for a
915 two-stage paired-difference regression following Senn, Julious and Araujo⁵. Both *beat*
916 Exposure-weighted at the Hybrid, Architecture B reference cell (Blocks S7-S8; E7 0.850,
917 E9 0.842, E2 0.836 at $t_{1/2} = 1.0$). Neither result generalizes. Block S8 shows both reverse,
918 like Exposure-weighted itself, under Architecture A. Block S9 shows both fail under the
919 crossover design: E9 mildly, E7 catastrophically, because AIC selection systematically
920 drifts to the boundary of its candidate half-life grid when off-drug sampling is sparse and
921 widely spaced, reproducing the S2 finding that over-assuming the half-life is the more
922 damaging error, now through the estimation procedure itself rather than a fixed wrong

assumption. Both specifications should be read as promising candidates for further validation at the one cell type where this manuscript already recommends Exposure-weighted, not as general-purpose replacements for it or for Unadjusted.

6 Conflict of interest

The authors declare no conflicts of interest.

7 Funding

This work received no specific funding.

8 Data availability statement

The data and code supporting the findings of this study are available in the pmsimstats-ng project repository. Per-replicate Monte Carlo outputs and ADEMP-compliant cell-level summaries are versioned under `analysis/data/`. Simulation drivers are at `analysis/scripts/carryover-sensitivity/`. Exact paths for this manuscript are listed in the Reproducibility section below.

9 Reproducibility

This manuscript was rendered with R 4.6.0 inside the project's zzcollab Docker container. The production simulation ran under R 4.5.3 (§3.6), with the image hash in `Dockerfile` and the package set in `renv.lock`. The simulation study uses `nlme` for the continuous-time linear-mixed analysis, `parallel` for the 8-core parameter-grid sweep, and `data.table` or `furrr/purrr` in the two simulation collections. No simulation code runs during knitting, which loads pre-computed checkpoints only. Each driver sets a master seed (20260507) and derives per-replicate seeds from it via `parallel::mc.set.seed`.

Data and code availability. All simulation outputs, cell-level summaries, and driver scripts are versioned under `analysis/scripts/carryover-sensitivity/` and `analysis/data/` in the pmsimstats-ng repository, including the principal factorial, the S1 to S5 sensitivity blocks, the S6 cluster-robust recalibration, the washout-adjusted (E4) check of Section 4.1, the high-precision prototype rerun of Section 3.7, the S7 alternative-paradigm block (09-run-sensitivity-s7.R, output/09-sensitivity-s7.rds), the S8 architecture-by-half-life block (10-run-sensitivity-s8.R, output/10-sensitivity-s8.rds), and the S9 CO-design block (14-run-sensitivity-s9.R, output/14-sensitivity-s9.rds), and the S10 CR2 comparison block (15-run-sensitivity-s10.R, output/15-sensitivity-s10-g.rds), all added 2026-08-18 (master seed 20260415, 500 replicates per cell, Hybrid reference cell except S9, which uses CO). E5/E6 (biomarker-interacted carryover terms) were tested in an earlier version of S7-S9 and are documented in Section 4.6; they were replaced by E7 (AIC-selected half-life, reusing the candidate-grid logic of `characterize_carryover()`,

R/carryover_analysis.R) and E9 (paired-difference regression). The E1-E3, E7, E9 fitters, and the E1cr2/E3cr2/E7cr2 CR2 variants used by S10 (cr2_extract(), fit_spec_s7, simulate_cell_s7), are appended to simulation-core.R alongside the Tier 1 fit_spec/fit_all_specs pair, which they do not modify. spec-labels.R additionally defines the G1-G9 compact display-code layer (g_code, g_labels) used in S10's table and figure; G8 is E2+CR2, reused directly from Block S6 rather than rerun, so it does not share common random numbers with the rest of S10's comparison. Figures 1 through 3 are produced from the same 02-grid-summary.rds Tier 1 summary as the headline table, by 12-hendrickson-heatmap.R. The corresponding Tier 2 heatmaps (S1-S4, S7, S8, S9, S10) are produced by 13-hendrickson-heatmap-tier2.R from 02-sensitivity-summary.rds and the S7/S8/S9/S10 output .rds files (plus diag-s6-cr2.rds for G8), and are additional to, not replacements for, the existing S1-S4 line plots and S7/S8 kable tables. The grid design and Morris ADEMP pre-registration are in simulation-grid-plan.md in the same directory. Superseded wider-scope drafts are retained under archive/, documented in its README.md.

Specification labels. The three specifications are stored as E1, E2, E3 and displayed here as Unadjusted, Lag-adjusted, and Exposure-weighted. The mapping is applied only at display time, by spec-labels.R, so the stored values and archived outputs are unaffected. Output archived before 2026-07-31 stores the specifications as A1, A2, A3. The migration is documented in analysis/report/NOTATION.md.

10 References

1. Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Statistics in Medicine*. 2019;38(11):2074-2102. doi:10.1002/sim.8086
2. Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: Statistical methods and practical considerations. *Frontiers in Digital Health*. 2020;2:13. doi:10.3389/fdgth.2020.00013
3. pmsimstats team. Two architectures for simulating biomarker-treatment interactions in aggregated N-of-1 trials. Published online 2026.
4. Jones B, Kenward MG. *Design and Analysis of Cross-over Trials*. 3rd ed. CRC Press; 2014.
5. Senn S, Julious SA, Araujo A. A comparison of four methods for the analysis of N-of-1 trials. *PLoS ONE*. 2014;9(2):e87752. doi:10.1371/journal.pone.0087752
6. pmsimstats team. Type-I calibration of fixed-effect interaction tests in linear mixed models: A staged diagnosis of working-covariance mis-specification and a cluster-robust remedy. Published online 2026.