

Latent-class and mixture-model formulations for biomarker-treatment interaction in N-of-1 trials

pmsimstats team

2026-05-08

Contents

1	Introduction	2
1.1	The aggregated N-of-1 trial and the biomarker-moderation question	2
1.2	Two formal answers: continuous heterogeneity and latent-class structure	3
1.3	Psychometric origins	4
1.4	Econometric and marketing-science development	5
1.5	Convergence on hybrid frameworks	7
1.6	Translation into longitudinal biostatistics	8
1.7	Why this matters for N-of-1 designs specifically	9
1.8	Aims and structure of the paper	10
2	Notation	11
3	Faithful latent-class data-generating processes	12
4	MVN approximation to mixture-induced moderation	13
5	Failure regimes of the MVN approximation	13
6	Implications for power under carryover	14
7	Identifiability at trial-relevant sample sizes	15
8	Finite-mixture taxonomy	16
8.1	Latent class and latent profile analysis	16
8.2	Growth mixture models and group-based trajectory models	17
8.3	Factor mixture models	17
8.4	Regression mixtures and mixtures of experts	18

29	9 The biomarker-moderation spectrum	18
30	9.1 Covariance-only moderation	19
31	9.2 Combined mean and covariance moderation	19
32	9.3 Location-scale moderation	20
33	9.4 Heterogeneous random-effects models	21
34	10 Implications for analysis and design	21
35	11 Software for mixture analyses	23
36	12 Discussion	23
37	12.1 Why the analysis model is simpler than the data-	
38	generating process	25
39	12.2 The latent-class DGP as a generator of apparent subaddi-	
40	tivity	26
41	12.3 Progress to date	27
42	12.4 Prototype Monte Carlo update at 240 replicates per cell .	31

1 Introduction

1.1 The aggregated N-of-1 trial and the biomarker-moderation question

The aggregated N-of-1 trial design is a class of clinical trials in which each enrolled participant cycles between active treatment and control across multiple within-person blocks, with within-person comparisons pooled across participants for population-level inference¹⁻⁴. The design preserves the within-subject contrast that gives single-participant N-of-1 trials their inferential strength while accumulating the statistical efficiency of a multi-participant trial. For chronic conditions whose treatment response is reversible on a useful timescale, the design yields more information per participant per unit time than a parallel-group randomised trial of equivalent total sample size⁵. It has been deployed in chronic pain⁶, inflammatory bowel disease, neuropsychiatric symptoms in post-traumatic stress disorder⁷, and a growing portfolio of rare and orphan disease applications where parallel-group trials are infeasible^{3,8}.

Within this design class, a question of increasing scientific weight is whether treatment response is moderated by a measured pre-treatment biomarker. A biomarker that predicts which participants respond to which treatments would let clinicians stratify treatment decisions at the

point of care, and the formal framework for such inference is laid out in the PATH statement⁹. The corresponding statistical estimand is the biomarker-by-treatment interaction term in the within-participant outcome model. Power to detect this interaction depends on the magnitude of moderation, the precision of the within-participant contrasts, the trial design (open-label, crossover, hybrid), and any time-dependent leakage between treatment phases, conventionally termed ‘carryover’, that erodes the cleanness of the within-subject contrasts.

Recent methodological work has documented substantial power losses for biomarker-by-treatment interaction inference under realistic carryover, with the loss varying sharply by how the moderation effect is encoded in the data-generating process (DGP)⁷. The present paper takes up a question that the prior work raised but did not resolve: how should biomarker moderation be represented in the DGP itself? Two formal answers, drawn from a long history of work outside biostatistics, dominate the available options.

1.2 Two formal answers: continuous heterogeneity and latent-class structure

The first answer represents biomarker moderation as a graded continuous effect: every participant has some degree of treatment response, the response is a smooth (typically linear) function of the biomarker, and population heterogeneity is captured by a single multivariate distribution with an interaction parameter. This is the dominant tradition in clinical biostatistics, where heterogeneity of treatment effect is routinely modelled via random slopes on the treatment indicator and via interaction terms in fixed-effects regression. In its most economical form, the moderation enters the conditional mean of the outcome as $\beta_{bm:D} \cdot B \cdot D$, where B is the biomarker and D is the (possibly continuous-decay) treatment indicator.

The second answer represents biomarker moderation as a discrete sub-population structure: participants belong to one of a small number of unobserved classes (most commonly ‘responder’ and ‘non-responder’), the classes have qualitatively distinct treatment-response curves, and the biomarker predicts class membership without determining it. The corresponding generative form is a finite mixture in which class membership is a Bernoulli or multinomial draw with mixing probabilities that depend

on the biomarker, and the within-class outcome model is class-specific. The marginal moderation observed in the data emerges from the mixture rather than from a smooth conditional-mean function.

These two answers correspond to genuinely different biological hypotheses. The continuous-heterogeneity view is consistent with a mechanism in which the biomarker reflects a graded physiological substrate (receptor density, enzymatic activity, baseline trait severity) that scales drug effect across the entire population. The latent-class view is consistent with a mechanism in which qualitatively distinct phenotypes exist (a metabolising versus non-metabolising allelic variant; a present versus absent target receptor) and the biomarker is an imperfect indicator of phenotype. The two views can produce identical marginal biomarker-outcome associations at a single timepoint and diverge sharply once the longitudinal structure of an N-of-1 trial is engaged, because carryover and within-person serial dependence interact differently with continuous and class-structured heterogeneity. This is the substantive point at which methodology imported from outside biostatistics becomes useful: both traditions have long histories of formalisation and identifiability analysis in their respective fields, and the corresponding tools are mature.

1.3 Psychometric origins

The continuous-trait view in formal statistical work originates with Spearman's two-factor theory of intelligence¹⁰, in which performance on cognitive tests was modelled as the sum of a single general factor and test-specific factors. Spearman's machinery was elaborated through the 1930s and 1940s by Thurstone, Hotelling, Lawley, and others into the modern factor-analytic tradition, in which observed indicator variables are written as linear combinations of unobserved continuous factors plus residual noise^{11,12}. Within this tradition, between-person heterogeneity in any latent construct is by assumption continuous, and individual differences are summarised by factor scores along one or more dimensions.

The discrete-class view emerged from a parallel but methodologically distinct literature beginning with Lazarsfeld's work on latent structure analysis in the 1950s^{13,14}. Lazarsfeld observed that in attitudinal and survey data, the joint distribution of categorical responses to multiple items frequently displayed structure inconsistent with a single contin-

uous latent trait, and instead suggested that respondents belonged to one of several qualitatively distinct latent classes within which item responses were locally independent. This ‘local independence’ formulation is the founding assumption of latent class analysis (LCA).¹⁵ placed the LCA model on rigorous statistical footing, deriving identifiability conditions and maximum-likelihood estimation procedures for the binary-indicator case in a *Biometrika* paper that remains a standard reference. Through the 1970s and 1980s, the LCA framework expanded to polytomous and continuous indicators, the latter widely known as latent profile analysis¹⁶, to multilevel and longitudinal extensions¹⁷, and to mixtures with covariate-dependent class probabilities^{18,19,20} unified much of this material into a comprehensive monograph on finite mixture models that is the standard entry point for quantitative researchers from outside the discipline.

The defining tension in psychometric methodology has been whether a given pattern of observed heterogeneity is best understood as continuous or as discrete.²¹ framed this as the choice between ‘graded’ and ‘unfolding’ models of individual differences and argued that the data themselves rarely settle the question without strong prior assumptions. The empirical literature has accumulated examples in both directions: cognitive ability is generally well-described by continuous factor models, while certain forms of psychopathology display latent-class structure consistent with qualitatively distinct subtypes. Subsequent methodological work, beginning with²² and²³, developed hybrid frameworks (factor mixture models) in which discrete classes coexist with within-class continuous factors. We return to these frameworks below.

1.4 Econometric and marketing-science development

A largely independent strand of work on finite mixtures developed in econometrics and quantitative marketing through the 1980s and 1990s. The motivating problems differed from psychometrics: rather than measurement of cognitive or psychological constructs, the econometric concern was unobserved heterogeneity in micro-econometric models (consumer choice, labour-market participation, firm productivity) where standard regression assumptions of identical agents were known to fail.²⁴, in

an Econometrica paper on duration models, demonstrated that estimates of structural parameters in survival-style models could be highly sensitive to the assumed functional form of the unobserved heterogeneity distribution, and recommended a non-parametric maximum-likelihood approach in which the heterogeneity distribution itself was estimated as a discrete mixture with unknown support points and weights. Their non-parametric maximum likelihood estimator anticipated much of what would later become standard finite-mixture methodology in biostatistics, and²⁵ placed the broader theory of mixture likelihoods on a unified mathematical footing.

In quantitative marketing, the parallel concern was customer segmentation. A population of consumers might respond to a marketing intervention differently depending on which unobserved segment they belonged to, and the firm needed both to estimate segment-specific response functions and to predict segment membership for new consumers.²⁶ formalised this as ‘clusterwise linear regression’, in which observations were partitioned into clusters with cluster-specific regression coefficients.²⁷ generalised the approach to generalised linear models.²⁸ extended it to structural equation models with unobserved segments. Within this econometric-marketing literature, the gating function (the model for the probability of segment membership as a function of covariates) became a first-class object, and methodological work focused on whether the gating model could share covariates with the within-segment regression model.²⁹, working contemporaneously in machine learning, formalised an essentially identical structure under the name ‘mixtures of experts’, which provided much of the now-standard vocabulary.

Two themes from the econometric strand are particularly relevant to the N-of-1 trial setting. First, the econometric tradition treated unobserved heterogeneity as a *nuisance* in many applications, and was correspondingly attentive to whether the heterogeneity distribution was identified, weakly identified, or over-fit. The identifiability concerns raised by²⁴ and elaborated by²⁵ translate directly into N-of-1 settings where sample sizes are small and the heterogeneity is the principal object of inference rather than a nuisance. Second, the marketing-science focus on covariate-dependent gating functions corresponds exactly to the biomarker-as-class-membership-predictor structure that motivates the present paper. The

biomarker-treatment moderation question, translated into econometric vocabulary, is whether the gating function depends on a single observed covariate (the biomarker) and whether the within-class regression coefficient on treatment differs across classes.

1.5 Convergence on hybrid frameworks

By the late 1990s, both the psychometric and the econometric strands had converged on a hybrid formulation: combine continuous within-class factor structure with discrete between-class structure.²² in Psychometrika, developed the factor mixture model as a finite mixture of confirmatory factor models, where each class had its own factor-loading and residual covariance structure.²³ placed factor mixture models in the broader Mplus-style structural equation framework.³⁰ conducted the first systematic Monte Carlo evaluation of factor mixture models and clarified the conditions under which the discrete-class component was empirically distinguishable from continuous heterogeneity.³¹ surveyed the modelling strategies and presented an applied analysis of psychiatric comorbidity.

In parallel, the longitudinal extension of these ideas yielded the growth mixture model^{32–34} and the closely related group-based trajectory model^{35,36}. Both posit that a longitudinal outcome trajectory is generated by a finite mixture of class-specific trajectories. Growth mixture models permit within-class random-effects variation around the class-specific mean trajectory; group-based trajectory models constrain within-class variance to zero, so each class has a single deterministic trajectory. Nagin’s group-based approach was developed primarily in criminology to characterise developmental trajectories of antisocial behaviour, and has been widely applied in epidemiology, public health, and clinical research where trajectory heterogeneity is the principal scientific question.

The methodological optimism around growth and factor mixture models during the early 2000s was tempered by a series of papers documenting overextraction hazards.³⁷ in Psychological Methods, demonstrated through simulation that growth mixture models fitted to data generated from a single non-gaussian distribution can spuriously identify multiple latent classes; the apparent classes are artefacts of distributional non-normality rather than evidence of substantive subpopulation structure.³⁸

elaborated on the implications for substantive psychological research and recommended caution in interpreting class structure without independent validation.³⁰ documented the analogous concerns for factor mixture models, identifying the joint conditions on sample size, class separation, and class-specific parameter differences under which class structure is recoverable.³⁹ systematised the class-enumeration question through Monte Carlo evaluation of information criteria and likelihood-ratio bootstrap procedures. The cumulative message of this literature is that mixture models can extract apparent class structure where none exists, and that distinguishing genuine from spurious class structure requires either substantial samples (typically several hundred to several thousand participants) or strong prior information.

1.6 Translation into longitudinal biostatistics

Mixture and latent-class methods entered biostatistical practice through several complementary channels.⁴⁰, in the Journal of the American Statistical Association, proposed a linear mixed-effects model in which the random-effects population was itself a finite mixture of multivariate normals. Their motivating application was longitudinal HIV-infection data in which subgroups of patients displayed qualitatively different CD4 trajectories. The model is essentially a growth mixture model in biostatistical clothing, with the standard random-intercept-and-slope linear mixed model augmented by a discrete latent class indicator.⁴¹ extended the framework to allow more flexible random-effects distributions via semi-parametric mixture specifications.

Proust-Lima and colleagues developed a more comprehensive software framework with the `lcmm` R package⁴², which implements joint latent-class mixed models for longitudinal continuous and ordinal outcomes, with optional joint modelling of survival outcomes via shared latent class structure. The `lcmm` framework has been deployed in cognitive ageing research, Alzheimer's disease progression modelling, and clinical-trial subgroup discovery. The `flexmix` package in R^{43,44} provides a more general infrastructure for fitting mixtures of regression-style models with user-supplied component specifications, and has been used in pharmacokinetic and pharmacodynamic modelling where individual-level variation in dose-response curves motivates a mixture structure.

Despite this translation activity, mixture and latent-class methods have not been systematically applied to aggregated N-of-1 trial designs. The published N-of-1 methods literature is dominated by within-participant effect estimation^{3,45}, Bayesian hierarchical models for participant-specific effects⁴⁶, and aggregated-evidence meta-analytic frameworks^{4,8}. The biomarker-moderation question has received occasional treatment under the heading of personalised-medicine power calculation^{7,47}, but in each case the moderation has been encoded as a continuous interaction term in a fixed-effects analysis, with no parallel treatment of the latent-class formulation. The mixture-formulation question is to our knowledge unaddressed in the existing N-of-1 methodology literature.

1.7 Why this matters for N-of-1 designs specifically

Three features of the aggregated N-of-1 design make the continuous-versus-discrete moderation question particularly consequential for design and analysis decisions.

First, the within-participant longitudinal structure of an N-of-1 trial provides repeated measurements at multiple drug states (on-drug, off-drug, in carryover). Under a continuous-heterogeneity DGP, the moderation signal is encoded in how the within-participant treatment contrast varies with the biomarker. Under a latent-class DGP, the moderation signal is encoded in class membership, which is fixed across all measurements within a participant and therefore in principle recoverable from the full trajectory even when individual on-drug contrasts are noisy. The two encodings imply different efficient analysis strategies and different power profiles under carryover⁷.

Second, N-of-1 trials are typically conducted at sample sizes of $N \in [30, 150]$, well below the sample sizes at which mixture- model identifiability has been validated in the psychometric literature. The Bauer-Curran-Lubke-Muthén identifiability cautions therefore apply with full force, and the analysis-strategy implications are non-trivial: even if a latent-class DGP is the biologically correct generative model, mixture analysis at trial-relevant sample sizes may be either underpowered or unidentified, and a misspecified single-component analysis may be the more reliable inferential tool despite its bias. Resolving this trade-off is in part an empirical question that the existing psychometric simulation

literature has not addressed at N-of-1- relevant sample sizes.

Third, carryover, which enters the analysis as a continuous decay in the effective drug indicator $D_{it} = \exp(-\lambda t_{sd,it})$, interacts differently with the two moderation encodings. In the continuous-heterogeneity formulation, carryover erodes the moderation signal directly because the moderation enters as a product with D_{it} . In the latent-class formulation, carryover erodes the within-class drug contrast but does not directly attenuate class membership; if class membership can be recovered from the full trajectory, some moderation signal can survive carryover that would be lost under the continuous-heterogeneity analysis. The two formulations therefore predict different power-loss profiles under carryover, and this difference is empirically large in the simulation literature accumulated by the companion carryover-sensitivity work.

1.8 Aims and structure of the paper

This paper has four aims. First, to formalise the latent-class and finite-mixture formulations of biomarker-treatment interaction in aggregated N-of-1 trials, with explicit identifiability conditions specific to the trial-relevant sample-size regime. Second, to characterise estimation strategies (expectation-maximisation, Bayesian samplers, candidate R packages) and the diagnostics required to distinguish genuine from spurious class structure at small N . Third, to report a factorial simulation study that benchmarks mixture-DGP and continuous-DGP power against the standard linear-mixed-model analysis and a class-aware analysis under realistic carryover. Fourth, to apply the developed methods to the prazosin-PTSD biomarker dataset of⁷ and to compare conclusions with those obtained under the continuous-moderation analysis reported in the original paper.

The remainder of the paper is organised as follows. A notation section fixes terminology that is used throughout. The latent- class generative form, its MVN second-moment approximation, and the regimes in which the approximation breaks down are developed in three subsequent sections. The implications for power under carryover and the identifiability constraints that bound mixture analysis at trial-relevant sample sizes are addressed next. A finite-mixture taxonomy section locates the relevant model families (LCA, GMM, FMM, regression mixtures); a biomarker-

moderation spectrum section organises the available DGP encodings moment by moment. Implications for analysis and design, and software for mixture analyses, are then drawn together. The discussion section addresses three open methodological questions and references the five-study simulation programme described in the companion simulation plan, which constitutes the empirical core of the present line of work and is the next deliverable in the programme.

2 Notation

Throughout, B_i denotes the pre-treatment biomarker value for participant i . BR_{it} denotes the biological-response component of outcome for participant i at time t , decomposed in the Hendrickson framework into three latent response factors: BR_{it} itself, a time-variant component TV_{it} , and a placebo-belief component PB_{it} . D_{it} denotes drug state at time t in the analysis-side continuous-decay form $D_{it} = \exp(-\lambda t_{sd,it})$ during off-drug periods and $D_{it} = 1$ during on-drug periods, with $t_{sd,it}$ the time since drug discontinuation and λ the decay rate. Z_i denotes an unobserved class label. The half-life of carryover is $t_{1/2} = \ln 2/\lambda$ and is allowed to differ between the DGP and the analysis model.

The principal psychometric and econometric terms used below are:

- **Latent class.** An unobserved categorical subgroup, equivalent in trial vocabulary to an unobserved responder/non-responder stratum.
- **Latent factor.** An unobserved continuous trait governing covariation among observed variables, equivalent to a frailty or to a random intercept in a linear mixed model.
- **Class-conditional distribution f_z .** The outcome distribution within a single subgroup.
- **Mixing proportion π .** The prevalence of a given subgroup in the population.
- **Measurement invariance.** Whether a model applies identically across subpopulations.
- **Growth mixture model (GMM).** A longitudinal mixed model in which the population of random effects is itself a finite mixture.
- **Factor mixture model (FMM).** A hybrid model with discrete

classes at the upper level and within-class continuous heterogeneity at the lower level.

- **Regression mixture / mixture of experts.** A finite mixture in which the gating function (class-membership probability) depends on covariates and within-class regressions are class-specific.
- **EM algorithm.** Expectation-maximisation applied to recover unobserved class labels jointly with class-conditional model parameters.

3 Faithful latent-class data-generating processes

If a candidate predictive biomarker indexes an unobserved responder/non-responder dichotomy, the mechanistically faithful data-generating process (DGP) is a finite mixture:

$$Z_i \mid B_i \sim \text{Bernoulli}(\pi(B_i)), \quad BR_{it} \mid Z_i = z \sim f_z(t, D_{it}),$$

where Z_i is the latent class indicator, $\pi(B_i) = \Pr(Z_i = 1 \mid B_i)$ is the class-membership probability (typically a logistic function of the biomarker), and the class-conditional response distributions f_0 and f_1 differ in their drug-response parameters: peak response, onset rate, or both. Within-class temporal structure follows the same modGompertz form used in single-class formulations, with class-specific parameters substituted.

Two consequences of this DGP form are central to what follows. First, the class label Z_i is fixed across all measurements within a participant; it is invariant to drug state and to time. Second, the class-conditional mean trajectory does depend on D_{it} : under full carryover washout ($D_{it} \rightarrow 0$) the two class-conditional means converge if the responder/non-responder distinction is purely in drug response. Class-label invariance therefore preserves participant-level information about the moderation signal that an analysis with access to Z_i could exploit, even where instantaneous between-class mean separation in BR is fully attenuated. The MVN-covariance encoding of biomarker moderation, in contrast, has no such latent-label information to exploit: the covariance is the entire signal, and its erosion under carryover is direct.

4 MVN approximation to mixture-induced moderation

A two-component mixture with class-dependent means and common within-class covariance produces a marginal joint distribution of (B_i, BR_{it}) whose first two moments are reproduced, to leading order, by a single multivariate normal with $\text{Cor}(B, BR) = c_{bm}$. Under within-class independence between B_i and BR_{it} ,

$$c_{bm}^2 = \frac{\pi(1-\pi)(\mu_B^1 - \mu_B^0)^2}{\text{Var}(B)} \cdot \frac{\pi(1-\pi)(\mu_{BR}^1 - \mu_{BR}^0)^2}{\text{Var}(BR)},$$

so that c_{bm}^2 is the product of the between-class variance fractions in B and in BR . The marginal correlation that an MVN can reproduce is therefore bounded by the geometric mean of the two intra-class correlations; if either class separation is small, the achievable marginal c_{bm} is small even when class structure is otherwise well-defined.

Under mild regularity conditions (adequate class overlap and mixing proportions π bounded away from 0 and 1), the mixture's second-moment structure is well-approximated by a single MVN with differential correlation across drug states: elevated correlation where class-conditional means separate (on-drug) and shrinkage towards zero where they coincide (off-drug after carryover washout). For power calculations driven by t-tests on $\beta_{bm:D}$ in a linear mixed-effects model, this approximation is adequate within the regularity conditions and captures the covariance consequences of latent-class structure without committing to a discrete generative form.

5 Failure regimes of the MVN approximation

The MVN approximation is misleading in three identifiable regimes, each with a distinct empirical signature.

- **Bimodal responses.** When f_0 and f_1 are well-separated (non-responders show near-zero drug effect and responders show a large effect), the marginal distribution of BR is bimodal: a histogram of

on-drug response would show two peaks rather than a single peak with gaussian scatter around it. No MVN matches the tails of a bimodal distribution, and power under the true mixture DGP will exceed the MVN-LMM prediction because a well-specified mixture analysis can exploit the bimodality. The pattern occurs in clinical settings where a subpopulation genuinely does not respond to a treatment, as in certain targeted oncology indications.

- **Strong class-membership gradient.** When $\pi(B_i)$ is steep, approaching a step function at a biomarker threshold, the biomarker behaves nearly deterministically as a class label. A thresholded mean-moderation specification (a thresholded Architecture-A variant) then becomes the better approximation, because the biomarker-response relationship is dominated by its mean-structure component. The corresponding clinical scenario is a diagnostic biomarker with near-perfect sensitivity and specificity for the responsive phenotype.
- **Class-varying covariance.** If autocorrelation, residual variance, or placebo-response parameters differ between classes, no single-component MVN can represent the generative structure. Analysis models assuming constant covariance are misspecified regardless of whether they encode the biomarker interaction in the mean or in the covariance. This is the latent-class analogue of heteroscedastic residual variance addressed in standard mixed-effects practice with `varIdent` or `weights` specifications, except that the stratifying variable is unobserved.

6 Implications for power under carryover

The substantial covariance-only power loss under carryover (40-60% across designs in⁷) reflects erosion of a second-moment signal. A mixture-model analysis fitted to mixture-DGP data can in principle recover additional power by estimating class membership directly, because the class label is itself carryover-invariant. The gap between the covariance-only carryover-sensitive power estimate and the mixture-faithful power estimate upper-bounds what a mixture analysis could recover at the same carryover level; whether this bound is tight in the trial-relevant parameter regime is an empirical question not addressed by

existing simulation infrastructure. The bound is asymptotic: at finite N in the trial-relevant range, finite-sample bias in mixture estimation can drive the realised gap below the asymptotic upper bound or, in adverse identifiability regimes, reverse its sign.

A focused factorial simulation isolating second-moment approximation error from analysis-model mis-specification is described in the companion simulation plan (Study 1: mixture-vs-MVN factorial). Briefly, the factorial crosses two DGPs (mixture and MVN) with two analyses (linear mixed model and joint latent-class mixed model). The difference between mixture-DGP-with-MVN-analysis and mixture-DGP-with-mixture-analysis quantifies the power recoverable by a faithful analysis model; the difference between the two MVN-analysis cells quantifies the approximation error inherent to the second-moment encoding. A null check on MVN-DGP-with-mixture-analysis verifies that the mixture model does not extract spurious power on data for which no latent class structure exists.

7 Identifiability at trial-relevant sample sizes

Mixture modelling is attractive on faithfulness grounds but carries identifiability costs that are acute in N-of-1 settings. The EM algorithm for finite mixtures of mixed-effects models requires either informative priors on class-membership probabilities or strong separation between class-conditional response curves; neither is assured in a trial designed to detect a moderation effect whose existence is itself uncertain. The standard diagnostic for mixture identifiability, the entropy of posterior class-membership probabilities, requires substantial sample sizes (N several hundred to several thousand) to stabilise at values that support confident class extraction in published psychometric applications^{37–39}. Aggregated N-of-1 trials of the size contemplated in current biomarker applications, $N \in [30, 150]$, are well below this range.

The methodological implication is that mixture analysis at trial-relevant sample sizes may be either underpowered or unidentified, even when the underlying DGP is genuinely mixture-structured. The question is therefore not whether mixture formulations are more faithful, but whether the

fidelity gain is recoverable at realistic sample sizes, or whether analysis-strategy mitigations available within standard mixed-effects practice (restricted analysis, weighting, within-subject contrasts; see⁷) deliver a better power-per-complexity trade. These mitigations have better-understood inferential properties at the relevant sample sizes. Resolving the mixture-versus-mitigation question empirically requires the simulation programme described in the companion simulation plan (Study 5 in particular addresses small- N identifiability and type I error explicitly).

8 Finite-mixture taxonomy

Four families of finite mixture model from the psychometric and econometric literatures bear directly on the biomarker-treatment interaction problem. The discussion below introduces each family, its origin, and its specific relevance to the N-of-1 trial setting.

8.1 Latent class and latent profile analysis

Latent class analysis (LCA) and latent profile analysis (LPA) are mixture models for cross-sectional data. Each participant belongs to one of a small number of unobserved classes; the observed variables have class-specific marginal distributions. LCA uses categorical indicators; LPA uses continuous indicators. A mixture fitted to baseline biomarkers alone returns posterior class-membership probabilities that can be used as covariates in subsequent outcome analysis, the two-stage analytic strategy standard in psychometric applications. In the biomarker-trial context, an LPA-style model posits two or more responder classes with class-specific marginal distributions of on-drug and off-drug response, and estimates class-membership probabilities from the observed response trajectories rather than from baseline covariates alone. Classical references include^{14, 15, 48}, and²⁰; the Vermunt and Magidson Latent GOLD software and the poLCA R package implement the standard EM estimation routines.

8.2 Growth mixture models and group-based trajectory models

Growth mixture models (GMM)^{32–34} extend LPA to longitudinal data by specifying class-specific growth (or, in the trial setting, drug-response) curves with within-class random-effects variation. A GMM is a longitudinal linear mixed model in which the population distribution of random effects is a mixture of multivariate normals, each corresponding to a distinct trajectory class. Group-based trajectory models^{35,36} are the degenerate case in which within-class random-effect variances are set to zero, so that participants within a class share an identical trajectory up to observation-level residual noise.

GMM is the closest psychometric analogue of the latent-class generative form developed above: each participant belongs to a latent class with its own drug-response trajectory, and the class label is drug-state-invariant. A GMM that uses the biomarker as a class-membership predictor is essentially the faithful generative model written down earlier in this paper.³⁷ and³⁸ document the identifiability and overextraction hazards that accompany GMM estimation at realistic sample sizes, showing in particular that GMM can extract apparent class structure from data generated by single-class non-gaussian processes. These cautions apply with full force to aggregated N-of-1 trials, where participant counts are typically well below the sample sizes at which GMM class-count recovery has been validated.

8.3 Factor mixture models

Factor mixture models (FMM)^{22,23,30} combine continuous latent factors with discrete latent classes, yielding a generative structure in which within-class variation is governed by a factor model and between-class variation is governed by class-specific means and (optionally) class-specific factor loadings or residual covariances. The FMM framework subsumes both the pure continuous (single-class factor model) and the pure discrete (LCA with no within-class structure) cases as limits.

FMMs are two-level hierarchies. At the upper level, participants are assigned to one of a small number of classes (as in LCA). At the lower level, each class has its own within-class random-effects structure. The

biological hypothesis the FMM encodes is that there are genuinely distinct responder subtypes and that, within subtype, additional continuous heterogeneity is driven by covariates or unmeasured individual differences. For the biomarker-moderation problem, FMM provides a principled way to represent a biomarker as both a noisy class indicator (via the gating function) and a continuous modulator of within-class response (via factor loadings on the response factor). The hybrid representation matches more biological scenarios than either pure LCA or pure factor analysis, at the cost of correspondingly higher estimation requirements.

8.4 Regression mixtures and mixtures of experts

Regression mixtures^{26–28} and the closely related mixture-of-experts architecture²⁹ allow class-specific regression coefficients with class-membership probabilities (the gating function) that depend on covariates. An ordinary linear model with a treatment-by-biomarker interaction posits a single regression surface in which the treatment slope varies smoothly with the biomarker. A regression mixture instead posits two or more distinct regression surfaces, one per class, with a separate covariate-dependent model governing which surface applies to which participant. The two-stage structure (classification by the gating model, then prediction by the within-class regression) parallels the two-stage structure of causal-inference methods that use propensity scores to stratify before estimating within-stratum treatment effects.

In the trial-simulation idiom, regression mixtures correspond to a DGP in which the biomarker governs both the probability of being a responder and the magnitude of the within-responder drug effect. `flexmix` in R^{43,44} implements this family with user-definable component models, including mixed-effects components suitable for repeated-measures data.

9 The biomarker-moderation spectrum

The covariance-only formulation of⁷ specifies a biomarker-dependent covariance structure with population means held constant across biomarker values. Locating this formulation within the broader spectrum of biomarker-moderated DGPs is useful because the moment-by-moment organisation of the spectrum (biomarker moderating the mean, the

covariance, both, or higher moments) clarifies which features of the carryover-induced power loss are properties of the specification itself rather than of the underlying scientific question.

9.1 Covariance-only moderation

The covariance-only specification holds $E[BR_{it}]$ fixed across biomarker strata and encodes the interaction entirely in $\text{Cov}(B_i, BR_{it} \mid D_{it})$. High- and low-biomarker participants have the same expected drug response at any given time but differ in the covariance between biomarker and response, specifically in how strongly the biomarker co-varies with response during on-drug versus off-drug periods.

This is an unusual restriction within the psychometric mixture tradition. Most standard FMM and GMM parameterisations permit class-specific means at a minimum, and the means-constant- covariance-varies case is typically encountered only as a constrained submodel used for testing measurement-invariance hypotheses^{49,50}. The corresponding biostatistical model has two subgroups with identical expected outcomes but differing residual variance or within-subject autocorrelation; such models are occasionally fitted as sensitivity analyses but rarely assumed as the primary DGP. The covariance-only encoding's appeal lies in its tractability for power simulation under multivariate normality: the entire DGP reduces to drawing from a single MVN with a constructed covariance. Its cost is that it does not correspond cleanly to any natural generative mechanism.

9.2 Combined mean and covariance moderation

²³ develop finite mixtures of conditional mean- and covariance-structure models in which both the expected-response vector and the residual covariance matrix differ between classes and may depend on observed covariates within class. This is the general psychometric form: it reduces to a mean-moderation model in the limit of a single class with covariate-dependent mean only, and to the covariance-only encoding in the limit of class-invariant means with covariate-dependent covariance only.

In the combined form, responders and non-responders differ both in typical drug response (mean moderation) and in how noisily that response is

expressed within class (covariance moderation). Both channels carry information about class membership: the mean channel via $E[BR | B]$, the covariance channel via $\text{Var}[BR | B]$. Under carryover, the mean channel is comparatively robust because class-conditional means need not converge off-drug; the covariance channel is eroded by construction because off-drug covariance is attenuated. A combined-specification simulation should therefore exhibit power loss intermediate between the pure-mean and pure-covariance extremes; this is plausibly the most biologically faithful regime for biomarker-guided psychiatric trials.⁴⁰ provide the analogous treatment for linear mixed-effects models with heterogeneity in the random-effects population, the more direct generalisation of the repeated-measures structure used in N-of-1 trial simulation.

9.3 Location-scale moderation

Generalised additive models for location, scale, and shape (GAMLSS)⁵¹ and the related distributional regression tradition⁵² allow the biomarker to enter multiple parameters of the outcome distribution simultaneously through separate link functions. In plain terms, a standard regression models the conditional mean of an outcome as a function of covariates; a GAMLSS-style distributional regression models the conditional mean, the conditional variance, and if appropriate the conditional skewness and kurtosis, each as a separate function of covariates. The familiar special case is heteroscedastic linear regression with covariate-dependent residual variance; GAMLSS generalises this to arbitrary distributional shape.

In the biomarker-trial context, a GAMLSS-based DGP would allow high-biomarker participants to have both a shifted mean drug response (a location effect, which is Architecture A) and altered response variability (a scale effect, which is analogous to Architecture B's covariance moderation but operating on marginal rather than joint moments), without requiring discrete class structure. The GAMLSS framework is thus the continuous-heterogeneity analogue of FMM and occupies an intermediate position between pure Architecture A (mean only) and pure Architecture B (covariance only). This intermediate position is attractive for sensitivity analysis because it does not require committing to a specific number of latent classes.

9.4 Heterogeneous random-effects models

⁴⁰, ⁴¹, and ⁴² develop linear mixed-effects models in which the random-effects distribution is itself a finite mixture, allowing class-specific random-intercept and random-slope distributions. The biostatistical setting this extends is the standard linear mixed model with random intercept and random treatment slope, familiar from any longitudinal trial analysis using `nlme::lme` or `lme4::lmer`. In the standard model, the random treatment slope is drawn from a single normal distribution with mean equal to the average treatment effect and variance capturing individual heterogeneity. In the heterogeneous random-effects extension, that single distribution is replaced by a mixture: a fraction of participants are drawn from a ‘responder’ distribution with a large mean treatment slope and another fraction from a ‘non-responder’ distribution with a near-zero mean treatment slope. The biomarker predicts which mixture component a participant comes from.

Applied to aggregated N-of-1 trial data, this framework represents responder heterogeneity directly at the participant-specific random treatment effect, which is arguably the most clinically transparent representation available: responders and non-responders have distinct random-slope distributions on the treatment indicator, and the biomarker predicts class membership rather than entering the analysis-model mean structure at the fixed-effect level. The `lcmm` R package⁴² implements this family and is the closest available tooling for the longitudinal aggregated N-of-1 use case: it accepts standard longitudinal mixed-model syntax, fits a mixture on the random-effects distribution, and returns posterior class probabilities for each participant. For an extension of an existing aggregated N-of-1 simulation framework, `lcmm` is the lowest-friction entry point.

10 Implications for analysis and design

Three implications follow from locating the covariance-only encoding within the broader spectrum.

First, the covariance-only parameterisation is the most restrictive operational case of a much richer family. The substantial power loss under carryover documented in⁷ is a specific property of this restrictive case and need not generalise to the mean-plus-covariance or heterogeneous-

random-effects variants that are standard in psychometric and biostatistical practice. A sensitivity analysis based solely on the covariance-only DGP risks overstating the carryover-induced power loss that a biomarker-guided aggregated N-of-1 trial should expect, because the covariance-only DGP places the entire biomarker-outcome signal in a channel that carryover erodes by construction. If a portion of the signal flows through the mean structure, as in pure mean moderation or in the combined specification of²³, that component survives carryover and overall power loss is correspondingly attenuated. A complete sensitivity analysis should report power under at least three points on the spectrum: pure mean moderation, pure covariance moderation, and a combined specification.

Second, the identifiability concerns documented in the GMM and FMM literatures^{30,37–39} bound what can be asked of mixture-based analysis models at N-of-1 trial sample sizes. Published FMM applications typically involve several hundred to several thousand participants; biomarker-trial applications with $N \in [30, 150]$ are an order of magnitude smaller. Analysis plans built around fitting `lcmm` or `flexmix` mixture models directly to trial data are unlikely to yield reliable class-membership inference at trial-relevant sample sizes and should not be relied on as the primary analysis. Mixture analysis nonetheless retains value as a sensitivity check (does class-structured residual pattern emerge from the pre-specified analysis?) and as a DGP for power simulation (does the primary analysis remain robust to class-structured departures from gaussianity?).

Third, if N-of-1 simulation infrastructure pursues a richer DGP in subsequent work, the natural target is the heterogeneous random-slopes form rather than a full FMM. The random-slopes parameterisation preserves the analysis model’s compatibility with the existing `nlme::lme` machinery (the analysis model continues to be fitted as a single linear mixed model with participant-specific random treatment slopes) while admitting the biologically faithful latent-class generative structure at the DGP level. The simulation DGP is extended in a direction internal to the standard mixed-effects framework, and the analysis model remains the pre-specified linear mixed-effects fit; this preserves the audit chain of the existing framework and permits direct comparison against the covariance-only baseline under matched effect-size calibration.

11 Software for mixture analyses

R packages relevant to estimating the models surveyed above, in approximate order of fit to the longitudinal aggregated N-of-1 use case, are as follows.

- `lcmm`⁴². Joint latent-class mixed models for longitudinal continuous and ordinal outcomes; the closest available tooling for heterogeneous random-effects specification. Syntax is closely related to `nlme::lme` with a latent-class extension.
- `flexmix`^{43,44}. General regression- mixture framework with user-definable component models; suitable for regression-mixture and mixture-of-experts specifications. More general than `lcmm` but requires more configuration.
- `mclust`⁵³. Gaussian finite mixture models with a range of class-specific covariance parameterisations; supports²³ mean-plus-covariance specifications non- longitudinally. Useful for cross-sectional latent-profile analyses of baseline biomarker panels but does not natively handle repeated-measures structure.
- `gamlss`⁵¹. Distributional regression with covariate-dependent location, scale, and shape parameters.
- `poLCA`⁵⁴. Polytomous latent class analysis for the categorical-indicator case; useful if baseline class indicators (e.g., symptom-checklist items) are incorporated into class assignment.
- `OpenMx` and `lavaan` with mixture extensions. Structural- equation mixture modelling for FMM-style specifications; more general but with steeper configuration cost than `lcmm` or `flexmix`.

The reference implementation outside R is Mplus⁵⁵, which remains the standard environment for GMM, FMM, and related mixture- SEM specifications. Much of the published psychometric simulation literature on mixture-model identifiability and overextraction hazards^{30,37,39} uses Mplus, and direct comparison with that literature occasionally requires access to it.

12 Discussion

The biomarker-treatment moderation problem in aggregated N-of-1 trials sits at the intersection of two methodological traditions that have,

until now, only weakly engaged with each other. The biostatistical tradition has developed efficient estimators for within-participant treatment contrasts and the meta-analytic machinery to combine them, but has typically encoded biomarker moderation as a continuous interaction in fixed-effects regression. The psychometric tradition has developed mature finite-mixture and latent-class machinery for representing discrete subpopulation structure but has applied it primarily to cross-sectional or observational settings with sample sizes well above the N-of-1 range. The argument of this paper is that the choice between continuous-heterogeneity and latent-class encodings is a substantive scientific question, not a parameterisation convenience, and that recognising this choice opens both a sensitivity-analysis programme (the spectrum of biomarker- moderation specifications above) and an analysis-strategy programme (class-aware mixture analyses where identifiability permits, and analysis-strategy mitigations where it does not).

Two limitations bound what the paper currently provides. First, the simulation programme that would empirically resolve the mixture-versus-mitigation efficiency question at trial-relevant sample sizes is described in the companion simulation plan but not yet executed. The five-study factorial outlined there (mixture-vs-MVN, the mean-covariance-combined triptych, heterogeneous random slopes, the three-regime breakdown grid, and the small- N identifiability and type I error pre-flight) constitutes the empirical core of the present line of work and is the next deliverable in the programme. Second, the application to real biomarker-guided trial data has not been carried out. The prazosin-PTSD analysis of⁷ provides the most immediate target; the simulation programme should run in parallel with re-analysis of that dataset under at least three biomarker- moderation encodings and a class-aware analysis where convergence permits.

Three open methodological questions remain. The first is whether class-aware analyses are usable as primary analyses at any realistic N-of-1 sample size, or only as sensitivity checks. The identifiability evidence from the broader mixture-model literature suggests the latter; only direct simulation in the trial-relevant parameter regime can settle this. The second is how analysis half-life should be specified when the DGP is a mixture rather than a single MVN. The continuous-decay drug indicator $D_{it} = \exp(-\lambda t_{sd,it})$ has a clean interpretation under a single-component

generative model; its interpretation when the underlying generative process is class-structured and class-conditional carryover dynamics differ between classes is less obvious and may require analysis-side mixture extension. The third is whether the trial design itself should respond to the moderation-encoding question. If a latent-class generative form is biologically more plausible than a continuous one, designs that maximise the per-participant contrast and minimise carryover may gain power that the standard hybrid design surrenders.

The five-study simulation programme described in the companion plan is structured to address each of these questions sequentially. The factorial designs are pre-registered through ADEMP⁵⁶ tables committed before production runs; the class-aware analyses are wrapped uniformly so that simulation results are directly comparable across studies; and the infrastructure required to add the necessary DGPs to the existing simulation framework is enumerated, scoped, and sequenced to support both this paper's revisions and the companion carryover and biomarker-interaction papers in the wider research programme.

12.1 Why the analysis model is simpler than the data-generating process

A natural reaction to the simulation programme above is that the DGP is structured (a three-component decomposition into biological response, placebo-belief response, and time-variant natural history under the standard pmsimstats parameterisation; or a latent-class mixture under the alternatives in this paper) while the analysis model is a four-term linear mixed model $\mathbf{Sx} \sim \mathbf{bm} + \mathbf{t} + \mathbf{Dbc} + \mathbf{bm:Dbc}$ with random intercept and AR(1) residual correlation. A reader newly encountering this asymmetry is reasonable to ask why the analysis does not mirror the DGP component-by-component, and whether doing so would improve inference about the moderation estimand. The short answer, with brief justification, is the following.

A component-matched analysis (a non-linear mixed-effects model fitting Gompertz BR, PB, and TV components with participant-specific random effects) is impractical for trial-relevant N for four reasons: the trial design must supply the contrasts that identify each component (open-label versus blinded for PB, on-drug versus blinded-placebo for BR, off-

drug timepoints for TV), and many designs supply only some of them; each Gompertz component costs roughly three to four parameters at the population level and additional random-effect parameters per participant, which at $N \in [30, 150]$ inflates the variance of the moderation estimand more than it reduces residual confounding; the non-linear fitting machinery has higher convergence-failure rates and is more sensitive to starting values than the linear-mixed analogue; and the moderation estimand $\beta_{bm:D}$ is robust to the BR-versus-PB attribution for the validation question, even though it is not robust to the larger questions of whether the moderation is class-structured (the present paper) or covariance-versus-mean-encoded (the `dgp-mean-moderation-vs-mvn` companion).

The smallest extension that retains tractability while improving the BR-versus-PB sensitivity of the inference is a phase indicator plus its interaction with the drug indicator, schematically $\mathbf{Sx} \sim \mathbf{bm} + \mathbf{t} + \mathbf{Dbc} + \mathbf{phase} + \mathbf{Dbc:phase} + \mathbf{bm:Dbc} + \mathbf{bm:Dbc:phase} + (1 \mid \mathbf{ptID})$. This adds three to four parameters and tests two things directly: whether the apparent drug effect shrinks under blinding (a signature that the open-label drug effect was partly PB-mediated) and whether the moderation itself shrinks under blinding (a signature that the biomarker-by-treatment interaction was partly biomarker-by-PB rather than biomarker-by-BR). For papers where pharmacological-versus-placebo disentanglement is a secondary scientific question, this extension is the right step; for papers where the primary estimand is the marginal moderation, the simple $\mathbf{Sx} \sim \mathbf{bm} + \mathbf{t} + \mathbf{Dbc} + \mathbf{bm:Dbc}$ analysis is adequate. The simulation programme fixed in the companion plan uses the simple form throughout; the phase-indicator extension is a candidate sensitivity analysis to add to Study 1 (the mixture-vs-MVN factorial) once the primary cells are run. The pedagogy doc at `docs/24-component-decomposition-pedagogy.md` develops this point in more detail with worked numerical examples.

12.2 The latent-class DGP as a generator of apparent subadditivity

A connection worth flagging explicitly: the latent-class DGP specified in Study 1 is a natural mechanism for generating **apparent subadditivity** in a downstream one-component analysis, even

when no direct $BR \times PB$ interaction is specified in the conditional generative model. The companion methods paper at [analysis/report/06-component-decomposition/](#) defines subadditivity as a non-zero coefficient γ on a $BR \times PB$ interaction term in the conditional-mean specification of the additive decomposition, and notes that $\gamma > 0$ in a fitted additive-extension analysis is ambiguous between two distinct underlying mechanisms.

The first mechanism is **direct mechanistic interaction**: shared neural substrate, ceiling effects on the symptom scale, or belief-dynamics adaptation each produce a true non-additive coupling between drug and placebo response. The second mechanism is **emergent subadditivity from latent- class structure**: if the trial population is a mixture of responder and non-responder classes (the formulation in §3 of the present paper) and class membership correlates with both BR and PB strength (responder-class participants have both higher BR and higher PB on average), then the marginal expected value of $BR \cdot PB$ over the unobserved class label has a non-zero covariance term whose magnitude depends on the gating function $\pi(B)$. From the additive-extension analysis this looks exactly like a γ -weighted interaction.

The Study 1 pre-registration (see [analysis/scripts/latent-class-mixture-application/00-ademp-pre-registration.md](#)) accordingly includes a `pb_class_correlation` axis that varies $\text{Cor}(BR, PB)$ across class labels: a PB-class-independent cell (matching the original specification) and a PB-class-correlated cell (generating apparent subadditivity). The Study 1 simulation output therefore lets the analyst distinguish the two mechanisms empirically: if the additive-extension $\hat{\gamma}$ is large in the PB-class-correlated cell but small in the PB-class-independent cell at matched marginal effect size, the mechanism is emergent rather than mechanistic. The cross-paper synthesis with the BR-PB-TV decomposition paper is therefore not just methodological convenience but a direct route to falsifying or confirming a class-aware versus mechanistic- interaction interpretation of any apparent subadditivity finding.

12.3 Progress to date

Phase 1 of the simulation programme has reached the infrastructure-and-pilot milestone. Three components are committed and exercised end-to-

end on the host R session.

Pre-registration. ADEMP tables for all five studies⁵⁶ are committed at `analysis/scripts/latent-class-mixture-application/00-ademp-pre-registration.md`, with cell counts (~2,300 cells across the programme), seeded random-number control, an effect-size calibration sub-study, and a deviation log specification. Subsequent production runs are diagnosed against these pre-registered estimands rather than against any post-hoc wrapper choice.

Class-aware analysis wrapper. An `lcmm_analysis()` wrapper around `lcmm::hlme` is committed at `01-lcmm-wrapper.R` with the same input-output contract as the existing `lme_analysis()`. The wrapper uses a two-stage fit (a single-class `hlme` is fitted first to anchor starting values for a `gridsearch`-based multi-start `hlme(ng=2)` fit) and extracts two moderation signals: the gating-coefficient z-test on the biomarker (the primary class-aware moderation test, identifying whether the biomarker predicts class membership) and the within-class $\beta_{bm:D}$ t-test (the secondary signal, retained for inferential symmetry with the linear-mixed baseline).

Heterogeneous-random-slopes DGP. A wrapper around `generateData()` that adds a Bernoulli class-membership channel with logistic gating in the standardised biomarker and class-specific multipliers on the BR factor at on-drug timepoints is committed at `02-heterogeneous-slopes-dgp.R`. On a single-replicate smoke test ($N = 70$, gating slope 1.5, responder multiplier 1.0, non-responder multiplier 0.0) the DGP produces 46% responders and a marginal $\text{Cor}(B, BR) = 0.63$ that the linear-mixed `bm:Dbc` test detects ($p = 0.002$) without class awareness; the `lcmm` fit returns a posterior class-assignment entropy of 0.86 and a BIC 170 lower than the matched single-component MVN-DGP fit, indicating that the mixture model discriminates the two DGPs at this parameter cell.

Study 5 pre-flight pilot. A 100-replicate pilot under the single-component null DGP (Architecture B with $c_{bm} = 0$, $N = 70$, hybrid OL design) is reported at `output/03-study5-pilot.rds`. Per-replicate timing is 0.05 s for the linear-mixed baseline and 1.57 s for the class-aware analysis (32:1 cost ratio, consistent with the smoke-test estimate). `lcmm` convergence rate is 96%; the four non-converging replicates encountered numerical issues during the `gridsearch` initialisation.

Type I error at $\alpha = 0.05$ (Monte Carlo standard error 0.022) varies sharply across candidate test forms, motivating the choice of the primary class-aware moderation test:

Test	Rejection rate
lme bm :Dbc Wald t-test	0.000
lcmm gating-on- bm Wald z (unconditional)	0.250
lcmm within-class bm :Dbc Wald z	0.000
BIC selects ng=2 over ng=1	0.010
LRT chi-sq(df=4), naive reference	0.146
BIC-conditional gating procedure	0.000

The unconditional Wald-on-gating test rejects at five times the nominal rate. Diagnosis of the rejecting replicates shows that the empirical standard deviation of the gating coefficient across reps ($sd = 63.9$) is roughly fifty times the median within-replicate Wald SE (0.86), and the empirical sd of the z-statistic is 1.47 against a nominal 1.0. The gating coefficient is centred at zero (median -0.10, sign balance 46/54) and shows no systematic bias; the inflation is driven by the conditional Wald SE under-estimating the true variability of $\hat{\beta}_{\text{gating}}$ at trial-relevant N . This is a finite-sample identifiability artifact predicted by the³⁷ and³⁰ overextraction literature: when the data have no class structure but **hlme** is forced to fit **ng=2**, the optimiser latches onto random patterns to define a class boundary, and the resulting Wald SE does not absorb the labelling uncertainty.

The naive LRT against a chi-squared(df=4) reference is similarly anti-conservative (rejection 0.146), because the true LRT distribution at the ng=1 boundary is a chi-bar-squared mixture rather than a chi-squared (²⁰, ch. 6). The bootstrapped LRT (BLRT) of³⁹ is the standard calibration but is computationally expensive at the per-fit budget of a class-aware simulation.

The pragmatic class-aware moderation test is therefore a **BIC-conditional gating procedure**: fit ng=1 and ng=2, prefer ng=2 only if BIC supports it ($\Delta_{\text{BIC}} > 0$ or, more conservatively, > 6 per Raftery 1995), and only then test the gating coefficient. At $N = 70$ under the null DGP, BIC selects ng=2 in 1% of replicates, and zero of those reps

also reject the gating Wald, giving the conditional procedure a Type I error of 0 at the pilot's resolution. The conditional procedure is therefore usable as a primary moderation test at trial-relevant N . Whether the same conditional procedure has adequate power to detect a true two-class DGP is the central question for the rest of Study 5; pilots at non-null cells will measure this directly.

The diagnostic distinction matters for the paper's argument. The earlier proposition that 'class-aware analyses are not usable as primary at trial-relevant N ' is too strong as stated: it applies to the unconditional Wald-on-gating test but not to the BIC-conditional procedure. The corrected proposition is that the unconditional Wald test on a single mixture-model coefficient is unreliable at small N and requires a model-comparison wrapper (BIC or BLRT) to control Type I error.

Production-budget implication. Extrapolated to the pre-registered cell counts (Study 1 has $768 \text{ cells} \times 1,000 \text{ replicates} = 768,000$ fits per analysis arm), the class-aware arm of Study 1 is approximately 340 hours of serial CPU time at the pilot's mean fit time. Parallelisation across 8-16 cores reduces this to 22-44 hours per arm. Studies 2 and 3 require similar budgets; the Study 5 production run ($5,000 \text{ replicates} \times \text{three sample sizes} \times \text{three class counts} \times \text{three DGP families}$) is the largest single block at approximately 60 hours per arm on 8 cores. The full programme is feasible on a single workstation over a week of background compute, with the calibration sub-study and Study 5 pre-flight running first to finalise effect-size axes and identifiability cuts before committing to Studies 1, 2, and 3.

Open question raised by the pilot. With Type I error of the BIC-conditional procedure controlled at the null at $N = 70$, the next-most-pressing identifiability question is whether the BIC selection step has enough sensitivity to detect a real two-class DGP at trial-relevant N . A class-aware procedure that is conservative under the null but rarely fires under true class structure is no more useful than a procedure that fires too often. Pilot 4 (planned for the next iteration) will fit the same BIC-conditional procedure to the heterogeneous- random-slopes DGP from Study 3 (with several values of the gating slope and class-separation Δ) and report the power profile of the conditional procedure alongside that of the linear-mixed baseline. The expected pattern is that the BIC-

conditional procedure outperforms the linear-mixed `bm:Dbc` test at large class separations (where `ng=2` BIC consistently wins) but underperforms it at moderate class separations (where BIC is undecided and the procedure falls back to a default that does not exploit the latent-class structure).

12.4 Prototype Monte Carlo update at 240 replicates per cell

The Study 5 prototype was rerun at 240 replicates per cell on 4-core `parallel::mclapply` with three gating-slope levels ($gs \in \{0, 1.0, 1.5\}$) and three test forms (lme `bm:Dbc` Wald, lmm gating Wald unconditional, BIC-conditional gating). Wall time 558 s; convergence 92.9% under the null and 98.3 to 98.8% at non-null cells. Per-replicate output at `analysis/data/quick-sim/03-latent-replicates.rds`; Morris summary at `03-latent-summary.txt`. Monte Carlo standard error at the working power near 0.5 is approximately 0.030.

The unconditional Wald-on-gating type I error of 0.197 (95% binomial CI [0.145, 0.250]) is now decisively distinguishable from the nominal 0.05 (more than five MCSE of separation) and is statistically consistent with the 100-replicate pilot's 0.250 point estimate. Power at $gs = 1.0$ (0.304) and $gs = 1.5$ (0.360) sits below the lme baseline (0.325 and 0.542 respectively), so the unconditional Wald buys nothing in exchange for the type I inflation. The BIC-conditional procedure does control type I (0.004 under the null, with upper CI near 0.013), but its power is 0.025 at $gs = 1.0$ and 0.004 at $gs = 1.5$: under the heterogeneous-random-slopes DGP at $N = 70$, the BIC step selects $ng = 1$ in essentially every replicate and the conditional procedure approximates a never-reject rule. The non-monotonicity ($gs = 1.5$ below $gs = 1.0$) suggests that as the heterogeneous-slopes signal strengthens, a single-class lme captures it well enough that BIC continues to favor $ng = 1$ and the conditional procedure does not fire.

The lme `bm:Dbc` test holds nominal type I (0 of 240 under the null) and reaches 0.542 power at $gs = 1.5$, dominating both lmm tests at this sample size and DGP regime. The proposition the manuscript should carry forward is that, at trial-relevant N and under the heterogeneous-random-slopes generating mechanism examined here, neither the uncondi-

tional Wald-on-gating nor the BIC-conditional gating test is competitive with the linear-mixed baseline: the unconditional test is severely anti-conservative, and the BIC-conditional test is severely under-powered. A canonical run at 1{,}000 replicates per cell with a wider gating-slope grid (and at $N \in \{35, 70, 100\}$) is required to map the regime where the BIC-conditional procedure does become competitive, likely at large class separations rather than heterogeneous random slopes.

1. Guyatt G, Sackett D, Taylor DW, Chong J, Roberts R, Pugsley S. Determining optimal therapy: Randomized trials in individual patients. *New England Journal of Medicine*. 1986;314(14):889-892.
2. Guyatt G, Keller JL, Jaeschke R, Rosenthal D, Adachi JD, Newhouse MT. The N-of-1 randomized controlled trial: Clinical usefulness. Our three-year experience. *Annals of Internal Medicine*. 1988;108(2):293-299.
3. Lillie EO, Patay B, Diamant J, Issell B, Topol EJ, Schork NJ. The n-of-1 clinical trial: The ultimate strategy for individualizing medicine? *Personalized Medicine*. 2011;8(2):161-173.
4. Senn S. Statistical pitfalls of personalized medicine. *Nature*. 2018;563(7733):619-621.
5. Mirza RD, Punja S, Vohra S, Guyatt G. The history and development of N-of-1 trials. *Journal of the Royal Society of Medicine*. 2017;110(8):330-340.
6. Yelland MJ, Schluter PJ. Defining worthwhile and desired responses to treatment of chronic low back pain: A comparison of n-of-1 trials with parallel-group trials. *Journal of Evaluation in Clinical Practice*. 2009;15(3):528-533.
7. Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: Statistical methods and practical considerations. *Frontiers in Digital Health*. 2020;2:13.

- 1063 8. Schork NJ. Personalized medicine: Time for one-person trials. *Nature*. 2015;520(7549):609-611.
- 1064 9. Kent DM, Paulus JK, Klaveren D van, et al. The Predictive Approaches to Treatment effect Heterogeneity (PATH) Statement. *Annals of Internal Medicine*. 2020;172(1):35-45.
- 1065 10. Spearman C. "General intelligence," objectively determined and measured. *American Journal of Psychology*. 1904;15(2):201-292.
- 1066 11. Thurstone LL. *Multiple-Factor Analysis*. University of Chicago Press; 1947.
- 1067 12. Lawley DN, Maxwell AE. *Factor Analysis as a Statistical Method*. 2nd ed. Butterworth; 1971.
- 1068 13. Lazarsfeld PF. The logical and mathematical foundation of latent structure analysis. In: Stouffer SA, ed. *Measurement and Prediction*. Princeton University Press; 1950:362-412.
- 1069 14. Lazarsfeld PF, Henry NW. *Latent Structure Analysis*. Houghton Mifflin; 1968.
- 1070 15. Goodman LA. Exploratory latent structure analysis using both identifiable and unidentifiable models. *Biometrika*. 1974;61(2):215-231.
- 1071 16. McCutcheon AL. *Latent Class Analysis*. Sage; 1987.
- 1072 17. Vermunt JK. Multilevel latent class models. *Sociological Methodology*. 2003;33(1):213-239.
- 1073 18. Dayton CM, Macready GB. Concomitant-variable latent-class models. *Journal of the American Statistical Association*. 1988;83(401):173-178.

- 1074 19. Bandeen-Roche K, Miglioretti DL, Zeger SL, Rathouz PJ. Latent variable regression for multiple discrete outcomes. *Journal of the American Statistical Association*. 1997;92(440):1375-1386.
- 1075 20. McLachlan GJ, Peel D. *Finite Mixture Models*. Wiley; 2000.
- 1076 21. Bartholomew DJ. *Latent Variable Models and Factor Analysis*. Griffin; 1987.
- 1077 22. Yung YF. Finite mixtures in confirmatory factor-analysis models. *Psychometrika*. 1997;62(3):297-330.
- 1078 23. Arminger G, Stein P, Wittenberg J. Mixtures of conditional mean- and covariance-structure models. *Psychometrika*. 1999;64(4):475-494.
- 1079 24. Heckman J, Singer B. A method for minimizing the impact of distributional assumptions in econometric models for duration data. *Econometrica*. 1984;52(2):271-320.
- 1080 25. Lindsay BG. *Mixture Models: Theory, Geometry, and Applications*. Vol 5. Institute of Mathematical Statistics; 1995.
- 1081 26. DeSarbo WS, Cron WL. A maximum likelihood methodology for clusterwise linear regression. *Journal of Classification*. 1988;5(2):249-282.
- 1082 27. Wedel M, DeSarbo WS. A mixture likelihood approach for generalized linear models. *Journal of Classification*. 1995;12(1):21-55.
- 1083 28. Jedidi K, Jagpal HS, DeSarbo WS. Finite-mixture structural equation models for response-based segmentation and unobserved heterogeneity. *Marketing Science*. 1997;16(1):39-59.
- 1084 29. Jacobs RA, Jordan MI, Nowlan SJ, Hinton GE. Adaptive mixtures of local experts. *Neural Computation*. 1991;3(1):79-87.

- 1085 30. Lubke GH, Muthén B. Investigating population heterogeneity with factor mixture models. *Psychological Methods*. 2005;10(1):21-39.
- 1086 31. Clark SL, Muthén B, Kaprio J, et al. Models and strategies for factor mixture analysis: An example concerning the structure underlying psychological disorders. *Structural Equation Modeling*. 2013;20(4):681-703.
- 1087 32. Muthén B, Shedden K. Finite mixture modeling with mixture outcomes using the EM algorithm. *Biometrics*. 1999;55(2):463-469.
- 1088 33. Muthén B. Latent variable mixture modeling. In: Marcoulides GA, Schumacker RE, eds. *New Developments and Techniques in Structural Equation Modeling*. Erlbaum; 2001:1-33.
- 1089 34. Muthén B. Latent variable analysis: Growth mixture modeling and related techniques for longitudinal data. In: Kaplan D, ed. *Handbook of Quantitative Methodology for the Social Sciences*. Sage; 2004:345-368.
- 1090 35. Nagin DS. Analyzing developmental trajectories: A semiparametric, group-based approach. *Psychological Methods*. 1999;4(2):139-157.
- 1091 36. Nagin DS. *Group-Based Modeling of Development*. Harvard University Press; 2005.
- 1092 37. Bauer DJ, Curran PJ. Distributional assumptions of growth mixture models: Implications for overextraction of latent trajectory classes. *Psychological Methods*. 2003;8(3):338-363.
- 1093 38. Bauer DJ. Observations on the use of growth mixture models in psychological research. *Multivariate Behavioral Research*. 2007;42(4):757-786.

- 1094 39. Nylund KL, Asparouhov T, Muthén BO. Deciding on the number of classes in latent class analysis and growth mixture modeling: A Monte Carlo simulation study. *Structural Equation Modeling*. 2007;14(4):535-569.
- 1095 40. Verbeke G, Lesaffre E. A linear mixed-effects model with heterogeneity in the random-effects population. *Journal of the American Statistical Association*. 1996;91(433):217-221.
- 1096 41. Zhang D, Davidian M. Linear mixed models with flexible distributions of random effects for longitudinal data. *Biometrics*. 2001;57(3):795-802.
- 1097 42. Proust-Lima C, Philipps V, Lique B. Estimation of extended mixed models using latent classes and latent processes: The R package lcmm. *Journal of Statistical Software*. 2017;78(2):1-56.
- 1098 43. Leisch F. FlexMix: A general framework for finite mixture models and latent class regression in R. *Journal of Statistical Software*. 2004;11(8):1-18.
- 1099 44. Grün B, Leisch F. FlexMix version 2: Finite mixtures with concomitant variables and varying and constant parameters. *Journal of Statistical Software*. 2008;28(4):1-35.
- 1100 45. Zucker DR, Schmid CH, McIntosh MW, D'Agostino RB, Selker HP, Lau J. Combining single patient (N-of-1) trials to estimate population treatment effects and to evaluate individual patient responses to treatment. *Journal of Clinical Epidemiology*. 1997;50(4):401-410.
- 1101 46. Zucker DR, Ruthazer R, Schmid CH. Individual (N-of-1) trials can be combined to give population comparative treatment effect estimates: Methodologic considerations. *Journal of Clinical Epidemiology*. 2010;63(12):1312-1323.

- 1102 47. Gabler NB, Duan N, Vohra S, Kravitz RL. N-of-1 trials in the medical literature: A systematic review. *Medical Care*. 2011;49(8):761-768.
- 1103 48. Clogg CC. Latent class models. In: Arminger G, Clogg CC, Sobel ME, eds. *Handbook of Statistical Modeling for the Social and Behavioral Sciences*. Plenum; 1995:311-359.
- 1104 49. Meredith W. Measurement invariance, factor analysis and factorial invariance. *Psychometrika*. 1993;58(4):525-543.
- 1105 50. Widaman KF, Reise SP. Exploring the measurement invariance of psychological instruments: Applications in the substance use domain. In: Bryant KJ, Windle M, West SG, eds. *The Science of Prevention*. American Psychological Association; 1997:281-324.
- 1106 51. Rigby RA, Stasinopoulos DM. Generalized additive models for location, scale and shape. *Journal of the Royal Statistical Society Series C*. 2005;54(3):507-554.
- 1107 52. Klein N, Kneib T, Klasen S, Lang S. Bayesian structured additive distributional regression for multivariate responses. *Journal of the Royal Statistical Society Series C*. 2015;64(4):569-591.
- 1108 53. Scrucca L, Fop M, Murphy TB, Raftery AE. mclust 5: Clustering, classification and density estimation using Gaussian finite mixture models. *The R Journal*. 2016;8(1):289-317.
- 1109 54. Linzer DA, Lewis JB. polCA: An R package for polytomous variable latent class analysis. *Journal of Statistical Software*. 2011;42(10):1-29.
- 1110 55. Muthén LK, Muthén BO. *Mplus User's Guide*. 8th ed. Muthén; Muthén; 1998-2017.

- 1111 56. Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Statistics in Medicine*. 2019;38(11):2074-2102.