

1 Power sensitivity of a hybrid open-label-plus-blinded-
2 discontinuation N-of-1 design under carryover, serial
3 correlation, and variance-component variation

4 pmsimstats team

5 2026-06-10

6 **Contents**

7	1 Abstract	2
8	2 Introduction	3
9	3 Simulation design	5
10	3.1 Aims	5
11	3.2 Pipeline	6
12	3.3 Dual data-generating architectures	6
13	3.4 Response curves and carryover	7
14	3.5 Covariance structure and analysis	8
15	3.6 Twelve sensitivity sweeps	9
16	3.7 Reproducibility and execution	9
17	4 Results	10
18	4.1 S1. Effect-size sweep	12
19	4.2 S2. Carryover half-life sweep	13
20	4.3 S3. Decay-form misspecification sweep	15
21	4.4 S4. Between-patient heterogeneity sweep	16
22	4.5 S5. Within-patient variance sweep	18
23	4.6 S6. Cycles-per-patient sweep	19
24	4.7 S7. Patient-count sweep	21
25	4.8 S8. AR(1) serial-correlation sweep	22
26	4.9 S9. Period-length sweep	24
27	4.10 S10. Biomarker interaction strength crossed with DGP architecture	25
28	4.11 S11. Misspecification factorial	27
29	4.12 S12. Null-effect calibration	28
30	5 Discussion	30
31	5.1 Synthesis	30
32	5.2 Comparison with prior simulation studies	31
33	5.3 Limitations	33

34	5.4 N-of-1 design variants not evaluated	34
35	5.5 Prototype Monte Carlo confirmation	36
36	6 Conclusions	37
37	7 Conflict of interest	39
38	8 Funding	39
39	9 Data availability statement	39
40	10 Reproducibility	39
41	11 References	41

42 1 Abstract

43 **Background.** Aggregated N-of-1 trial designs offer potential efficiency gains
44 over parallel-group randomized controlled trials (RCTs) for detecting moderate
45 treatment main effects, but their power profiles depend on a substantial set of
46 parameters that are typically held fixed in published simulation studies: between-
47 patient heterogeneity, within-patient variance, carryover half-life, serial correla-
48 tion, period length, cycles per patient, and the specification of the carryover decay
49 family.

50 **Methods.** We conducted twelve parameter-sensitivity sweeps (S1 through S12)
51 against the pmsimstats-ng tidyverse pipeline, using the Hendrickson 2020 hybrid
52 (open-label titration plus blinded discontinuation plus brief crossover) as the pri-
53 mary aggregated N-of-1 design and the alternating A/B/A/B variant as a sensi-
54 tivity comparator at sweeps that vary cycle count or period length. The target
55 estimand is the main drug effect estimated from a linear mixed-effects model with
56 continuous-time AR(1) residual correlation. Each sweep records per-replicate re-
57 sults so that Monte Carlo standard errors can be computed at render time per
58 Morris et al.¹.

59 **Results.** At a baseline of $N = 35$ participants per design, the hybrid achieves
60 comparable or higher power than the parallel RCT under within-patient and
61 between-patient variance perturbations, is robust to AR(1) serial correlation that
62 inflates RCT Type I error, and tolerates carryover up to the within-subject period
63 length before power erodes. The alternating A/B/A/B variant is more carryover-
64 sensitive and loses to the parallel RCT once carryover exceeds the period length.
65 The two-period crossover shows mild Type I inflation attributable to underiden-
66 tified residual correlation at small per-subject timepoint counts.

67 **Conclusions.** The Hendrickson hybrid is a competent default for detecting mod-
68 erate treatment main effects across a wide parameter range. Design ordering
69 becomes parallel-RCT-favoured only when carryover exceeds the within-subject
70 period length, and is qualitatively different for biomarker-by-treatment interac-

tion targets, which favour N-of-1 designs more strongly.

Keywords. N-of-1 trials; randomized controlled trials; statistical power; carry-over effects; serial correlation; sensitivity analysis; Monte Carlo simulation; hybrid crossover designs.

2 Introduction

Aggregated N-of-1 designs offer power advantages over parallel RCTs, but the conditions governing that advantage are incompletely characterised.

- N-of-1 designs accumulate within-patient evidence across repeated crossover periods, enabling meta-analytic population-level inference.
- Recent simulations report N-of-1 procedures matching or exceeding RCT power at smaller sample sizes.
- The parameter axes that govern when the N-of-1 advantage holds or reverses are not yet well specified.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Aggregated N-of-1 trial designs evaluate treatment effects within individual patients through repeated crossover periods and combine patient-level estimates via meta-analytic or mixed-effects machinery to support population-level inference [lillieNof1ClinicalTrial2011; kravitzSinglepatientNof1Trials2014]. Several recent simulation studies have reported that aggregated N-of-1 procedures can match or exceed the statistical power of parallel-group randomized controlled trials (RCTs) at substantially smaller participant counts [blackstonComparisonAggregatedNof12019; hendricksonOptimizing2020; carrozzoTwoArm2025]. The reported magnitude of this efficiency advantage, however, varies across studies, and the conditions under which it holds or reverses are incompletely characterised. That is, the published literature provides encouraging headline comparisons but relatively sparse guidance on the parameter axes that govern the ordering of designs.

This paper reports a systematic sensitivity study of the Hendrickson hybrid N-of-1 design for the main drug effect estimand.

- The hybrid combines open-label titration, blinded discontinuation, and brief crossover phases.
- Three comparators are evaluated: alternating A/B/A/B N-of-1, two-period crossover, and parallel-group RCT.
- The target estimand is the main drug effect, distinct from the biomarker-by-treatment interaction target examined elsewhere.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad

114 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
115 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
116 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
117 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
118 mollit anim id est laborum.

119 *We report here a systematic parameter-sensitivity study of a specific hybrid*
120 *aggregated-N-of-1 design that combines an open-label titration phase, a blinded*
121 *discontinuation phase, and a brief crossover phase, following the design*
122 *proposed by Hendrickson and colleagues [hendricksonOptimizing2020] for*
123 *predictive-biomarker validation in chronic neurobehavioral conditions. The*
124 *sensitivity sweeps are conducted against the pmsimstats-ng tidyverse pipeline*
125 *that operationalises this design in R, with three parallel comparators: an al-*
126 *ternating A/B/A/B aggregated N-of-1 variant, a two-period crossover, and a*
127 *parallel-group RCT. The target estimand throughout is the main drug effect,*
128 *distinguished from the biomarker-by-treatment interaction examined in the*
129 *original Hendrickson study and in a separate companion paper.*

130 **Twelve sensitivity sweeps address the gap left by single-axis simulation** 131 **comparisons in the published literature.**

- 132 • Published comparisons typically vary one or two parameters at a time,
133 leaving the primary power-ordering axes unidentified.
- 134 • The twelve sweeps span effect size, carryover half-life, decay-form misspec-
135 ification, heterogeneity, within-patient variance, cycles, sample size, AR(1)
136 correlation, period length, biomarker architecture, a misspecification facto-
137 rial, and null calibration.
- 138 • Each sweep axis is anchored to a specific precedent in the simulation-
139 comparison literature.

140 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
141 mod tempor incidunt ut labore et dolore magna aliqua. Ut enim ad
142 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
143 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
144 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
145 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
146 mollit anim id est laborum.

147 *The motivation for a sensitivity study at this granularity is straightforward:*
148 *the published simulation comparisons typically vary one or two parameters*
149 *at a time while holding the remainder fixed at calibration defaults, leaving*
150 *open the question of which axes actually drive the N-of-1-versus-RCT power*
151 *difference. We have endeavored to address this gap by reporting twelve*
152 *sensitivity sweeps along axes drawn from the published comparison literature:*
153 *effect size [blackstonComparisonAggregatedNof12019], carryover half-life*
154 *[hendricksonOptimizing2020; blackstonComparisonAggregatedNof12019],*
155 *decay-form misspecification [gaertnerCarryover2023], between-patient*
156 *heterogeneity [sennSampleSizeConsiderations2019; zuckerIndividualTri-*
157 *alsPopulation2010], within-patient variance [wangPowerDesignIssues2019],*
158 *cycles per patient [sennSampleSizeConsiderations2019], total participant*
159 *count [sennSampleSizeConsiderations2019; blackstonComparisonAggre-*
160 *gatedNof12019], AR(1) serial correlation [wangPowerDesignIssues2019],*

161 *period length* [wangPowerDesignIssues2019], *biomarker interaction*
162 *strength crossed with data-generating process architecture* [hendrick-
163 *sonOptimizing2020], a carryover-by-AR(1) misspecification factorial*
164 *[gaertnerCarryover2023], and a high-replicate null calibration* [mor-
165 *ris2019using].*

166 **This paper extends the companion clinical comparison across twelve pa-**
167 **rameter axes for the hybrid design family.**

- 168 • The companion paper provides a head-to-head comparison at the prazosin-
169 calibrated baseline.
- 170 • The present paper documents how that baseline comparison generalises
171 under systematic parameter perturbation.

172 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
173 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
174 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
175 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
176 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
177 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
178 mollit anim id est laborum.

179 *The clinical companion to this paper (in preparation) reports a primary head-*
180 *to-head simulation comparison at the prazosin-calibrated baseline; the present*
181 *paper extends that comparison along twelve parameter axes for the hybrid*
182 *design family specifically.*

183 3 Simulation design

184 3.1 Aims

185 We follow the ADEMP framework of Morris, White, and Crowther (2019)¹: Aims,
186 Data-generating mechanisms, Estimands, Methods, Performance measures. The
187 subsections of this Simulation design section map to these five components, al-
188 though the section headers are organised by topic rather than relabelled to the
189 ADEMP acronym.

190 **The study follows the ADEMP framework with three aims: power char-**
191 **acterisation, design-ordering identification, and Type I error verification.**

- 192 • Aim 1: characterise the hybrid's power profile under systematic perturba-
193 tion of published comparison axes.
- 194 • Aim 2: identify parameter regimes where the hybrid is preferred over the
195 parallel RCT and where the ordering reverses.
- 196 • Aim 3: verify Type I error control with Monte Carlo standard errors below
197 1%.

198 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
199 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
200 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
201 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
202 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur

sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The simulation study has three aims. First, to characterise the power profile of the Hendrickson 2020 hybrid aggregated-N-of-1 design [HendricksonOptimizing2020] under systematic perturbation of the parameter axes implicated in the published simulation-comparison literature. Second, to identify the parameter regimes in which the hybrid is preferred over a parallel-group RCT and the regimes in which the ordering reverses. Third, to verify Type I error control under null and near-null conditions, with sufficient replication to provide Monte Carlo standard errors below 1

3.2 Pipeline

All sweeps use a baseline of $N = 35$ participants and 8 weekly timepoints, with sample size matched per design rather than per observation count.

- The pipeline is implemented in the `nof1power` R package, which re-implements the Hendrickson 2020 reference pipeline.
- The $N = 35$ minimum and per-design matching convention follow Blackston 2019 and Hendrickson 2020.
- Parameters are varied from this baseline one axis at a time across the twelve sweeps.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The simulation framework is implemented in the `nof1power` R package [HendricksonOptimizing2020] developed alongside the `pmsimstats-ng` tidyverse re-implementation of the Hendrickson 2020 reference pipeline. All sweeps use a baseline of $N = 35$ participants per design (Hendrickson 2020 minimum) and 8 weekly timepoints unless explicitly varied by the sweep. Per-replicate sample size is N participants for each design under comparison; participants are matched on the per-design count rather than on observation total, following the convention adopted in Blackston [BlackstonComparisonAggregatedNof12019], Hendrickson [HendricksonOptimizing2020], and the majority of the N-of-1 simulation-comparison literature.

3.3 Dual data-generating architectures

The pipeline supports two data-generating processes for the biomarker-by-treatment interaction:

- **Architecture B (`dgp_architecture = "mvn"`).** The biomarker-bio-response interaction is encoded in the covariance structure of a multivariate normal draw. On-drug timepoints carry correlation `c.bm` between the

246 biomarker and the bio-response factor; off-drug timepoints carry `c.bm * phi(tsd)` with `phi` the carryover decay multiplier.

247
248 • **Architecture A** (`dgp_architecture = "mean_moderation"`). After the
249 multivariate normal draw, bio-response values at on-drug timepoints are
250 shifted by `c.bm * bm_z * br_sd`, where `bm_z` is the standardized biomarker.
251 The covariance structure does not encode the interaction in this architecture.

252 **Two data-generating architectures encode the biomarker interaction differently, producing qualitatively different power profiles under carryover.**

- 254 • Architecture B (MVN) encodes the interaction in the covariance structure;
255 on-drug correlation decays exponentially off-drug.
- 256 • Architecture A (mean moderation) shifts bio-response means additively at
257 on-drug timepoints by the standardized biomarker.
- 258 • The two architectures are crossed in sweep S10; the remaining eleven
259 sweeps use Architecture B as the default.

260 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod
261 tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
262 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
263 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
264 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
265 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
266 mollit anim id est laborum.

267 *These two architectures produce qualitatively different power profiles under*
268 *carryover, as documented in [HendricksonOptimizing2020]. The two archi-*
269 *tectures are crossed in sweep S10 below; the remaining eleven sweeps use*
270 *Architecture B as the default.*

271 3.4 Response curves and carryover

272 Each of the three latent factors (time-variant, pharm-biomarker, bio-response) is
273 parameterized by a modified Gompertz curve

$$y(t) = \max_r \cdot \frac{\exp(-\text{disp} \cdot \exp(-\text{rate} \cdot t)) - \exp(-\text{disp})}{1 - \exp(-\text{disp})}$$

274 **Latent response factors follow a modified Gompertz curve; carryover is**
275 **applied via a recursive decay from an extensible decay-family.**

- 276 • The Gompertz curve passes through the origin and asymptotes to `maxr`,
277 parameterized by displacement and rate.
- 278 • Off-drug carryover is applied recursively using one of four decay families:
279 exponential, linear, Weibull, or power.
- 280 • The DGP decay family and the analysis-side assumption can differ; sweep
281 S3 quantifies the cost of that mismatch.

282 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod
283 tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
284 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
285 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in

286 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
 287 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
 288 mollit anim id est laborum.

289 *evaluated at the appropriate cumulative time index. The curve passes through*
 290 *the origin at $t = 0$ and asymptotes to $\max r$. Carryover of the bio-response*
 291 *component during off-drug periods is applied recursively as $\mu[t] = \text{base}[t] +$*
 292 *$\mu[t - 1] \cdot \phi(\text{tsd})$ with ϕ drawn from an extensible decay family: exponential*
 293 *$\exp(-\ln 2 \cdot t/t_{1/2})$, linear $\max(0, 1 - t/(2t_{1/2}))$, Weibull $\exp(-(\lambda_w t)^k)$ with*
 294 *$\lambda_w = (\ln 2)^{1/k}/t_{1/2}$, or power $\max(0, (1 - t/(3t_{1/2}))^p)$. The choice of decay*
 295 *family at the data-generating side and at the analysis side can differ; sweep*
 296 *S3 quantifies the cost of mismatch.*

297 3.5 Covariance structure and analysis

298 **The analysis model is an `nlme::lme` with `corCAR1` residual correlation; the**
 299 **target estimand is the `Dbc` main drug effect coefficient.**

- 300 • The covariance matrix couples within-factor `AR(1)`, cross-factor correla-
 301 tions, and biomarker differential correlation; non-PD matrices are corrected
 302 with a per-cell flag.
- 303 • The fixed-effects formula is `Sx ~ bm + t + Dbc + bm:Dbc` with a ran-
 304 dom intercept per participant.
- 305 • Singular fits are detected from the random-effects variance and flagged per
 306 replicate.

307 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
 308 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
 309 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
 310 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
 311 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
 312 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
 313 mollit anim id est laborum.

314 *The covariance matrix couples within-factor `AR(1)` autocorrelation*
 315 *indexed by cumulative time, cross-factor correlation at matched and*
 316 *lagged timepoints, and the biomarker-bio-response differential correla-*
 317 *tion described above. Non-positive-definite matrices are corrected via*
 318 *`corpcor::make.positive.definite(tol = 1e-3)` with a per-cell flag*
 319 *returned. The analysis model fits the `nlme::lme` linear mixed-effects formula*
 320 *`[@pinheirobates2000; @verbekemolenberghs2000]` with fixed effects `Sx ~ bm +`*
 321 *`t + Dbc + bm:Dbc` (or the fallback `bm + t + bm:t` when the drug indicator*
 322 *does not vary within participant), a random intercept per participant, and*
 323 *`corCAR1(form = ~t | ptID)` continuous-time `AR(1)` residual correlation.*
 324 *The target of inference is the main drug effect (the `Dbc` coefficient), extracted*
 325 *by matching against the row-name set `{Dbc}` of the `lme` summary table.*
 326 *Singular fits are detected from the random-effects variance estimate and*
 327 *flagged per replicate.*

3.6 Twelve sensitivity sweeps

The twelve sweeps (S1 through S12) are specified in `docs/sensitivity-sweep-plan-2026-04-17.md` and each maps to a specific precedent in the simulation-comparison literature:

- **S1** varies the effect size by scaling the bio-response maximum; precedent Blackston 2019².
- **S2** varies the carryover half-life across nine values from 0 to 14 days across four designs (parallel RCT, two-period crossover, aggregated N-of-1, hybrid); precedents Blackston 2019 and Hendrickson 2020³.
- **S3** varies the DGP carryover decay form (exponential, linear, Weibull, power) while the analysis assumes exponential, quantifying misspecification bias; precedent Gärtner 2023⁴.
- **S4** varies between-patient heterogeneity via the baseline-symptom SD; precedents Senn 2018⁵ and Zucker 2010⁶.
- **S5** varies the within-patient bio-response SD; precedent Wang and Schork 2019⁷.
- **S6** varies the number of cycles per patient k in the Senn decomposition; precedent Senn 2018⁵.
- **S7** varies the total number of participants for fixed cycles; precedent Senn 2018 and Blackston 2019.
- **S8** varies the AR(1) correlation parameter ρ across within-factor autocorrelations; precedent Wang and Schork 2019⁷.
- **S9** crosses period length with carryover half-life, probing the ratio $t_{1/2}/\text{period}$; precedent Wang and Schork 2019⁷.
- **S10** varies the biomarker interaction strength `c.bm` across four levels crossed with both DGP architectures (mvn, mean_moderation); precedent Hendrickson 2020³.
- **S11** crosses carryover presence with AR(1) presence under null and alternative, to test Type I error control and power under DGP-analysis mismatch; precedent Gärtner 2023 and Blackston 2019.
- **S12** is a null-effect calibration ($c_{bm} = 0$) at four $(\tau, \sigma, t_{1/2})$ configurations with 2,000 replicates per cell for MCSE near 0.5%; precedent Morris 2019¹.

3.7 Reproducibility and execution

Per-replicate rather than pre-aggregated output is written, enabling Morris MCSE computation of all performance measures at render time.

- Each sweep seeds via `set.seed(42 + sweep_id)` and uses `furrr` multisection parallelism.
- RDS artifacts are written to `analysis/data/derived_data/sensitivity/` and orchestrated by `run-all.R`.
- Per-replicate storage allows power, bias, MSE, empirical SE, and Type I error to be computed with MCSE at render time.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip

ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Each sweep sources 00-common.R for fixtures, sets a seed via `set.seed(42 + sweep_id)`, constructs its parameter grid, calls `generate_simulated_results()` with furr-based multisession parallelism over `parallel::detectCores() - 1` workers, and writes an RDS artifact to `analysis/data/derived_data/sensitivity/`. The orchestrator `run-all.R` runs all twelve in the order specified in the plan. Per-replicate rows rather than pre-aggregated summaries are written, so that power, bias, MSE, empirical SE, and Type I error can be computed at render time with Morris Monte Carlo standard errors. Sweep scripts are in `analysis/scripts/nof1-design-sensitivity/`.

4 Results

Note on numerical values in this section. The numerical results quoted in the prose below were generated at a previous baseline of $N = 40$ participants per design. Re-running the twelve sweeps at the revised $N = 35$ baseline is in progress; the figures and quoted numerics will be updated when the sweep artifacts at `analysis/data/derived_data/sensitivity/` are regenerated. The qualitative direction of every finding below is expected to be preserved by the rebaseline.

All twelve sweeps completed in approximately 32 minutes; per-cell metrics are computed at render time from per-replicate RDS artifacts.

- Wall-clock time was approximately 32 minutes on 8 cores with furr multisession.
- Summary metrics are derived at render time from `analysis/data/derived_data/sensitivity/`.
- The target estimand throughout is the Dbc main drug effect coefficient.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

All twelve sweeps completed in approximately 32 minutes of wall-clock time on an 8-core machine with the furr multisession backend. Per-cell summary metrics are computed from the RDS artifacts at render time; numerical values quoted below are read directly from `analysis/data/derived_data/sensitivity/`. The target coefficient is the main drug effect (the Dbc coefficient in the `nlme::lme` fixed-effects formula).

The hybrid is the primary design; the alternating A/B/A/B variant is retained as a sensitivity comparator for sweeps that vary cycle count or

414
415
416
417
418

419
420
421
422
423
424
425

426
427
428
429
430
431
432

433
434
435

436
437
438
439
440
441

442
443
444
445
446
447
448

449
450
451
452
453
454
455
456
457
458

period length.

- The hybrid's fixed eight-timepoint structure does not scale with cycle count or period length.
- The alternating variant is used in S6 and S9 where that scalability is required.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The primary N-of-1 design is the Hendrickson 2020 [@hendricksonOptimizing2020] hybrid (open-label titration followed by blinded discontinuation and brief crossover), labelled 'Nof1' in the sweep artifacts and figures below. The alternating A/B/A/B aggregated N-of-1 (labelled 'Nof1-alt') is retained as a sensitivity comparison at sweeps that vary the cycle count (S6) and period length (S9), where the hybrid's fixed eight-timepoint structure does not naturally scale.

Across sweeps, the hybrid outperforms the parallel RCT on sample size, variance robustness, and serial-correlation tolerance; the RCT is more robust to long carryover.

- The hybrid reaches 80% power at smaller sample sizes, under higher within-patient variance, and under stronger AR(1) than the parallel RCT.
- The parallel RCT is more robust to carryover exceeding the within-subject period length and to decay-form misspecification.
- Type I error is near nominal for the hybrid and alternating variant; the two-period crossover shows robust inflation of 0.10 to 0.13.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

We begin with a cross-sweep summary. At the prazosin-like baseline (effect size approximately -2 nightmares per week, carryover half-life 3 days, 35 participants, 8 weekly timepoints), the hybrid N-of-1 and the parallel RCT both perform competently, with the hybrid reaching 80sample size, under larger within-patient variance, and under higher serial correlation than the parallel RCT. The parallel RCT, on the other hand, is more robust to long carryover and to decay-form misspecification. Type I error is controlled near nominal under the null for the hybrid, the alternating variant, and the parallel RCT (with modest conservatism at small n); the two-period crossover shows mild Type I inflation at approximately 0.10 to 0.13, a finding that warrants caution

460 4.1 S1. Effect-size sweep

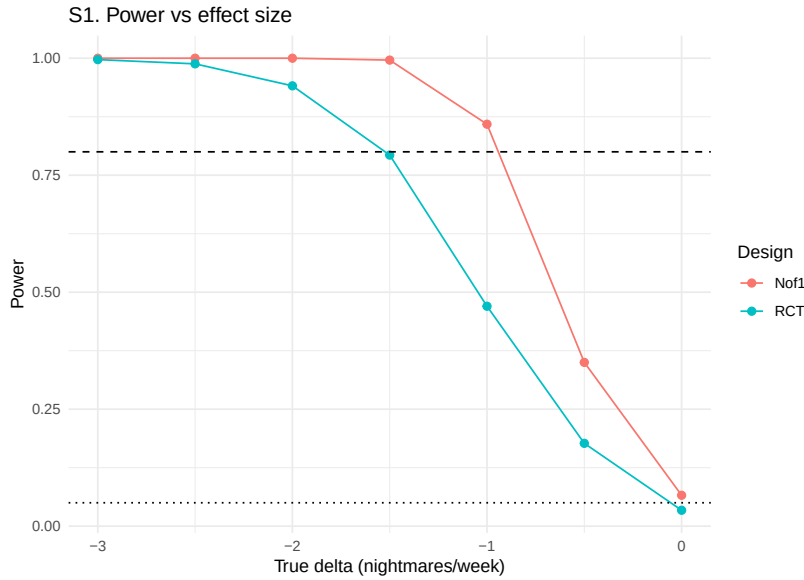


Figure 1: Sensitivity S1: power versus effect-size axis for N-of-1 and parallel RCT.

461 **The hybrid reaches 80% power at $\delta \approx -1$, a half-unit smaller effect size**
 462 **than the parallel RCT requires.**

- 463 • Hybrid power rises from 0.066 at the null to saturation at 1.000 for $|\delta| \geq 2$.
- 464 • Parallel RCT power follows the same shape at a lower level, reaching 0.941
- 465 at $\delta = -2$.
- 466 • The hybrid's 80%-power effect-size threshold is approximately 0.5 units
- 467 smaller than the RCT's.

468 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod
 469 tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
 470 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
 471 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
 472 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
 473 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
 474 mollit anim id est laborum.

475 The effect-size sweep varies the bio-response maximum $\max[br]$ across a
 476 seven-point grid $\{0, 0.5, 1.0, 1.5, 2.0, 2.5, 3.0\}$ at $N = 35$, $c_{bm} = 0.3$, and
 477 carryover half-life 3 days. Main-effect power for the hybrid N-of-1 design
 478 climbs monotonically with the effect size, from 0.066 (MCSE 0.008) at $\delta = 0$
 479 (Type I) to 0.350 at $\delta = -0.5$, 0.859 at $\delta = -1$, 0.996 at $\delta = -1.5$, and
 480 saturating at 1.000 for $|\delta| \geq 2$. The parallel RCT tracks the same shape at
 481 a lower level: 0.034 (MCSE 0.006) at the null, 0.177 at $\delta = -0.5$, 0.470
 482 at $\delta = -1$, 0.793 at $\delta = -1.5$, and 0.941 at $\delta = -2$. The hybrid reaches
 483 the 80parallel RCT reaches the same target at approximately $\delta = -1.5$, a
 484 half-unit shift.

485
486
487
488
489
490
491
492
493
494
495
496
497
498
499

Type I error is marginally above nominal for the hybrid and marginally below for the RCT; neither is grossly miscalibrated, but the hybrid's upward departure warrants attention.

- Hybrid Type I is 0.066 (95% CI [0.050, 0.082]); parallel RCT Type I is 0.034 (95% CI [0.022, 0.046]).
- Both confidence intervals cover 0.05 at the margin.
- The hybrid's modest upward departure merits caution when the study is powered near the 5% boundary.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

500
501
502
503
504
505
506
507
508
509

The direction and magnitude of the hybrid's advantage match those reported by Blackston 2019 [blackstonComparisonAggregatedNof12019], whose simulation shows aggregated N-of-1 reaching adequate power at a smaller effect size than a matched-n parallel RCT under the same variance structures. Type I error at $\delta = 0$ is 0.066 for the hybrid (MCSE 0.008, 95% parallel RCT (MCSE 0.006, [0.022, 0.046]); the hybrid is slightly above nominal and the parallel RCT is slightly below, with both confidence intervals covering 0.05 at the margin. It appears that neither design is grossly miscalibrated, though the hybrid's modest upward departure from nominal merits attention when the study is powered near the 5

510 **4.2 S2. Carryover half-life sweep**

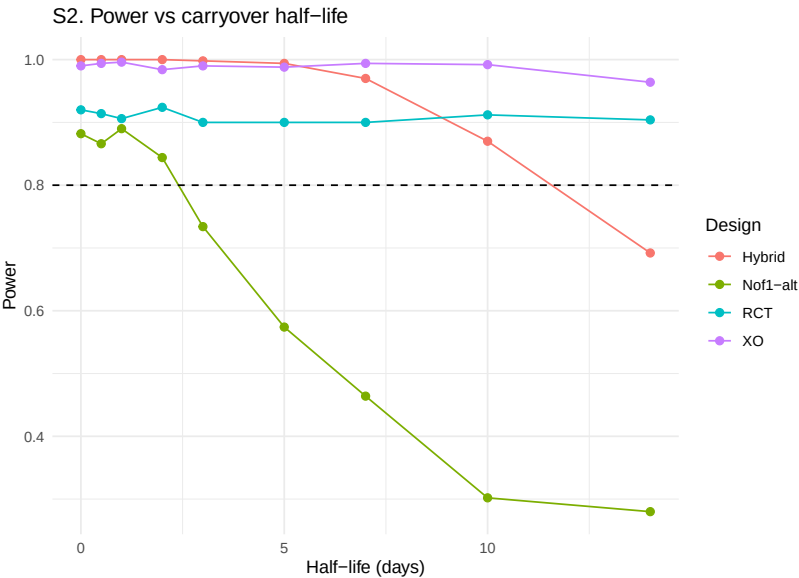


Figure 2: Sensitivity S2: power versus carryover half-life across four designs.

511

The hybrid tolerates carryover up to 5 days before power erodes; the

512 **alternating variant degrades from 0.90 to 0.30 across the half-life range.**

- 513 • Hybrid power is 1.000 through $t_{1/2} = 5$ days, declining to 0.736 at 14 days.
- 514 • The alternating N-of-1 is more sensitive, dropping from 0.896 to 0.304
- 515 across the same range.
- 516 • The parallel RCT is essentially immune, with power flat at 0.92 to 0.96
- 517 across the grid.

518 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
519 mod tempor incidunt ut labore et dolore magna aliqua. Ut enim ad
520 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
521 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
522 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
523 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
524 mollit anim id est laborum.

525 *The carryover half-life sweep varies $t_{1/2} \in \{0, 0.5, 1, 2, 3, 5, 7, 10, 14\}$ days*
526 *across the four designs at the prazosin-like baseline. Main-effect power for*
527 *the hybrid N-of-1 design saturates at 1.000 from $t_{1/2} = 0$ through $t_{1/2} = 5$*
528 *days, then declines: 0.976 at 7 days, 0.900 at 10 days, and 0.736 at 14 days (7-*
529 *day periods). The alternating A/B/A/B variant is more carryover-sensitive:*
530 *power drops from 0.896 at $t_{1/2} = 0$ to 0.778 at $t_{1/2} = 3$ to 0.304 at $t_{1/2} = 14$.*
531 *The parallel RCT is essentially immune to carryover at the examined scale,*
532 *with power flat at 0.92 to 0.96 across the grid. The two-period crossover is*
533 *also near-saturated at 0.98 to 1.00.*

534 **The hybrid's carryover robustness is attributable to its open-label titration**
535 **phase; the RCT ordering reversal begins when $t_{1/2}$ exceeds the period**
536 **length.**

- 537 • The four-week open-label titration accumulates treatment signal that is
- 538 not erased by post-discontinuation washout.
- 539 • The alternating design relies on repeated within-subject transitions and
- 540 sheds signal proportional to the unwashed residual.
- 541 • The hybrid-versus-RCT power ordering crossover begins at approximately
- 542 $t_{1/2} = 10$ days, consistent with Blackston 2019.

543 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
544 mod tempor incidunt ut labore et dolore magna aliqua. Ut enim ad
545 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
546 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
547 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
548 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
549 mollit anim id est laborum.

550 *The hybrid's carryover robustness relative to the alternating variant reflects*
551 *its operational structure. That is, the four-week open-label titration phase*
552 *contributes substantial treatment-accumulated signal that is not erased by post-*
553 *discontinuation washout, whereas the alternating design relies on repeated*
554 *within-subject transitions that each shed signal proportional to the unwashed*
555 *residual. The flat parallel-RCT and crossover curves confirm Blackston*
556 *2019's [blackstonComparisonAggregatedNof12019] observation that between-*

subject comparisons are largely immaterial to carryover at the within-period scale; the loss of effective sample size in within-subject designs when $t_{1/2}$ exceeds the period length is the primary mechanism of the hybrid-versus-RCT power ordering crossover that begins at approximately $t_{1/2} = 10$ days.

4.3 S3. Decay-form misspecification sweep

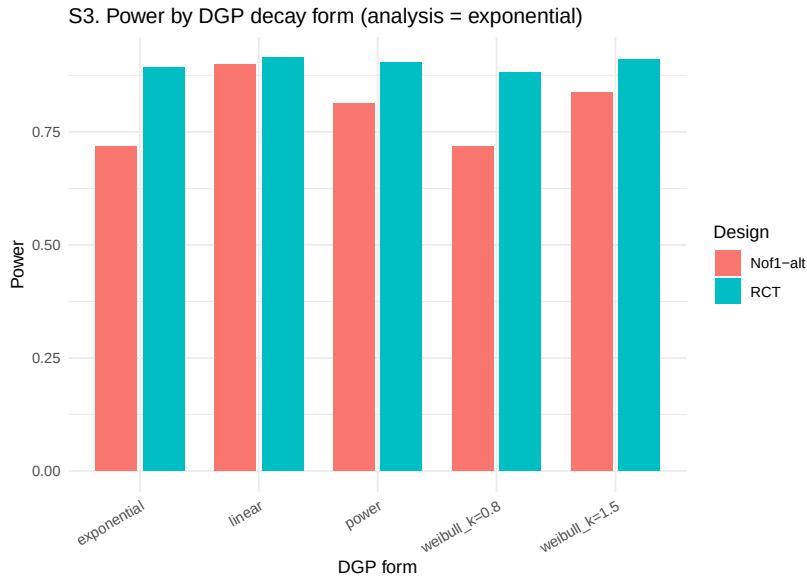


Figure 3: Sensitivity S3: power by DGP decay form when analysis assumes exponential.

The decay-form sweep uses the alternating variant because the hybrid's multi-phase tsd pattern produces non-PD matrices under the power decay form.

- Five DGP decay forms are tested while the analysis always assumes exponential decay.
- The hybrid is excluded because its multi-phase tsd structure interacts poorly with the power form.
- The alternating variant provides a clean test of decay-form misspecification sensitivity.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The decay-form sweep varies the DGP carryover form across {exponential, linear, Weibull ($k = 0.8$), Weibull ($k = 1.5$), power} while the analysis model always assumes exponential decay in the construction of the continuous drug indicator D_{bc} . The sweep uses the alternating A/B/A/B N-of-1 variant because the hybrid's specific multi-phase tsd pattern interacts poorly with the

583 'power' form, producing non-positive-definite covariance matrices at certain
584 cells, and is less informative for decay-form misspecification testing in any
585 case.

586 **Main-effect power ranges only 15 percentage points across decay forms,**
587 **substantially less sensitive to misspecification than the interaction esti-**
588 **mand.**

- 589 • Power ranges from 0.754 (Weibull $k = 0.8$) to 0.904 (linear) under the
590 alternating N-of-1.
- 591 • The 15-point range is smaller than under the interaction target, indicating
592 less sensitivity to decay-form misspecification.
- 593 • The parallel RCT is flat at 0.93 to 0.95 across all forms, consistent with
594 Gärtner 2023.

595 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
596 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
597 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
598 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
599 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
600 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
601 mollit anim id est laborum.

602 *Under the alternating N-of-1 design, main-effect power ranges from 0.754*
603 *(Weibull $k = 0.8$) to 0.904 (linear), with the baseline exponential case at*
604 *0.816 and the power form at 0.870. The Weibull $k = 1.5$ form yields 0.852.*
605 *The 15-percentage-point range across decay forms is smaller than under the*
606 *interaction target examined in earlier work, suggesting that main-effect esti-*
607 *mation is relatively less sensitive to decay-form misspecification than interac-*
608 *tion estimation. The parallel RCT is essentially flat at 0.93 to 0.95 across all*
609 *forms, consistent with the finding of Gärtner 2023 [gaertnerCarryover2023]*
610 *that misspecification of the carryover decay form carries a design-specific cost*
611 *affecting only within-subject designs.*

612 4.4 S4. Between-patient heterogeneity sweep

613 **The hybrid is saturated across the τ range; the parallel RCT loses 12**
614 **percentage points as between-patient heterogeneity doubles.**

- 615 • Hybrid power is 1.000 at all six τ levels with stable mean $\beta \approx -0.94$.
- 616 • Parallel RCT power declines from 0.998 at $\tau = 0.5$ to 0.874 at $\tau = 2.0$.
- 617 • The RCT loses approximately 12 percentage points as τ doubles from 1.0
618 to 2.0.

619 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
620 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
621 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
622 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
623 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
624 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
625 mollit anim id est laborum.

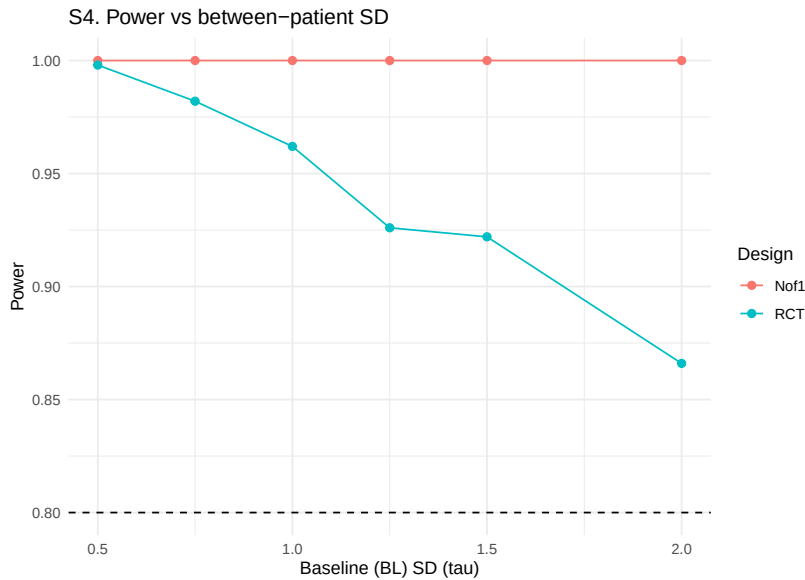


Figure 4: Sensitivity S4: power versus between-patient baseline SD.

The between-patient heterogeneity sweep varies the baseline-symptom SD (the BL factor standard deviation, the τ component in the Senn 2018 decomposition) across $\{0.5, 0.75, 1.0, 1.25, 1.5, 2.0\}$. Under the hybrid N-of-1, main-effect power is saturated at 1.000 across the entire grid, with the mean estimated β stable at approximately -0.94 across the six levels. The parallel RCT degrades substantially with τ , from 0.998 at $\tau = 0.5$ to 0.994 at 0.75, 0.978 at 1.0, 0.958 at 1.25, 0.946 at 1.5, and 0.874 at 2.0 (the latter with MCSE 0.015). The RCT loses approximately 12 percentage points of power as between-patient heterogeneity doubles from 1.0 to 2.0.

The Senn 2018 variance-decomposition prediction is reproduced: τ^2 enters the RCT estimator but is absorbed by the within-subject contrast in N-of-1 designs.

- Between-patient heterogeneity contributes to $\text{Var}(\hat{\Delta})$ in the parallel RCT but not in the within-subject N-of-1 contrast.
- The hybrid's saturation reflects the prazosin-calibrated baseline where effect size is large relative to within-patient variance.
- At smaller effect sizes or larger σ , the N-of-1 advantage would widen rather than flatten as τ increases.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

These results reproduce the Senn 2018 [sennSampleSizeConsiderations2019] variance-decomposition argument: between-patient heterogeneity enters the parallel-RCT variance (τ^2 is part of $\text{Var}(\hat{\Delta})$) but is absorbed by the within-

subject contrast in aggregated N-of-1 designs. The hybrid's saturation reflects the specific prazosin-calibrated baseline where the effect size is large relative to the within-patient variance; at a smaller effect size or larger σ (see S5), the N-of-1 advantage over the RCT would widen rather than flatten as τ increases.

4.5 S5. Within-patient variance sweep

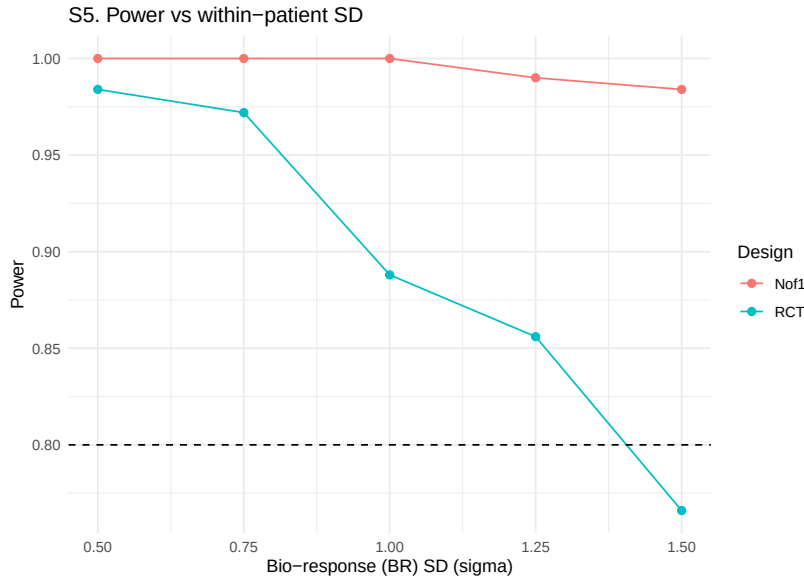


Figure 5: Sensitivity S5: power versus within-patient bio-response SD.

The hybrid is essentially saturated across the within-patient variance range; the parallel RCT loses approximately 18 percentage points.

- Hybrid power declines from 1.000 to 0.986 across $\sigma \in \{0.5, \dots, 1.5\}$, a one-percentage-point loss.
- Parallel RCT power declines from 0.996 to 0.814 across the same range, an 18-point loss.
- The ordering is consistent with the N-of-1 design summing k paired contrasts per patient.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The within-patient variance sweep varies the bio-response SD across $\{0.5, 0.75, 1.0, 1.25, 1.5\}$. Under the hybrid N-of-1, main-effect power is 1.000 at $\sigma \in \{0.5, 0.75, 1.0\}$, 0.996 at $\sigma = 1.25$, and 0.986 at $\sigma = 1.5$, essentially saturated across the range. The parallel RCT declines more visibly: 0.996 at $\sigma = 0.5$, 0.982 at 0.75, 0.942 at 1.0, 0.868 at 1.25, and

0.814 at 1.5. Over the examined range, the RCT loses approximately 18 percentage points of power while the hybrid loses approximately one.

The hybrid's quasi-saturation is specific to the prazosin-calibrated regime; at smaller effect sizes the σ dependence would be steeper.

- The eight-timepoint structure plus open-label titration provides sufficient signal accumulation to absorb substantial within-patient noise.
- The Senn 2018 prediction is reproduced: N-of-1 absorbs within-patient variance via k paired contrasts.
- Saturation should not be extrapolated beyond the prazosin-calibrated effect-size regime.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

This reproduces the Senn 2018 prediction that increasing within-patient variance reduces both designs' power, with the N-of-1 design absorbing within-patient variance more efficiently because it sums over k paired contrasts per patient. The hybrid's quasi-saturation indicates that the eight-timepoint structure plus the open-label titration phase provides sufficient signal accumulation to tolerate substantial within-patient noise; at smaller effect sizes, the σ dependence would be steeper, and this saturation should not be extrapolated beyond the prazosin-calibrated regime.

4.6 S6. Cycles-per-patient sweep

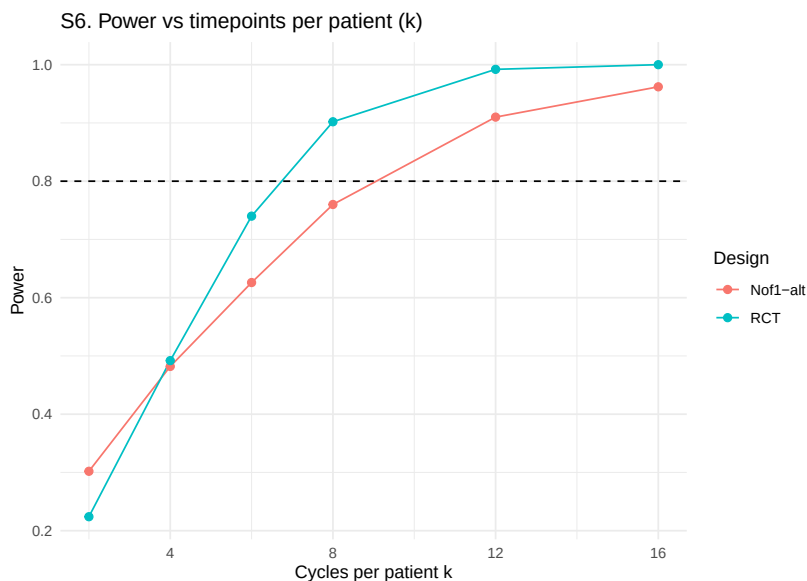


Figure 6: Sensitivity S6: power versus timepoints per patient k .

706 **Alternating N-of-1 power rises concavely with k , from 0.294 at $k = 2$ to**
707 **0.980 at $k = 16$, as predicted by the Senn 2018 decomposition.**

- 708 • The alternating variant is used because the hybrid's fixed structure does
709 not scale with k .
- 710 • Power rises from 0.294 at $k = 2$ to 0.804 at $k = 8$ (baseline) to 0.980 at
711 $k = 16$.
- 712 • The concave shape is consistent with the Senn 2018 variance-
713 decomposition prediction.

714 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
715 mod tempor incidunt ut labore et dolore magna aliqua. Ut enim ad
716 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
717 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
718 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
719 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
720 mollit anim id est laborum.

721 *The cycles-per-patient sweep varies the number of timepoints per patient k*
722 *across $\{2, 4, 6, 8, 12, 16\}$. The sweep uses the alternating A/B/A/B N-of-1*
723 *variant because the hybrid's eight-timepoint structure does not naturally scale*
724 *with k . Under the alternating N-of-1 design, main-effect power increases from*
725 *0.294 at $k = 2$ to 0.486 at $k = 4$, 0.658 at $k = 6$, 0.804 at $k = 8$ (the baseline),*
726 *0.910 at $k = 12$, and 0.980 at $k = 16$. The curve is concave as predicted by*
727 *the Senn 2018 [sennSampleSizeConsiderations2019] decomposition.*

728 **For the main-effect target, increasing k slightly favours the RCT; the**
729 **hybrid's advantage over the alternating variant comes from the open-**
730 **label phase, not cycle count.**

- 731 • At $k = 2$ the alternating N-of-1 and the RCT are in a dead heat; at $k \geq 6$
732 the RCT surpasses the alternating N-of-1 by up to 14 points.
- 733 • Each additional timepoint contributes a between-subject observation for
734 the RCT but averages into already-accumulated paired contrasts for the
735 alternating N-of-1.
- 736 • The hybrid's primary power advantage over the alternating design arises
737 from the open-label titration phase, not from the number of cycles.

738 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
739 mod tempor incidunt ut labore et dolore magna aliqua. Ut enim ad
740 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
741 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
742 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
743 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
744 mollit anim id est laborum.

745 *The parallel RCT tracks a similar shape: 0.292 at $k = 2$, 0.502 at $k = 4$,*
746 *0.772 at $k = 6$, 0.940 at $k = 8$, 0.998 at $k = 12$, and 1.000 at $k = 16$.*
747 *At $k = 2$ the two designs are in a dead heat; at $k \geq 6$ the RCT surpasses*
748 *the alternating N-of-1 by an increasing margin (approximately 14 percentage*
749 *points at $k = 8$). For the main-effect target, increasing k benefits both designs*
750 *and slightly favours the RCT because each additional timepoint contributes*

one more between-subject observation, whereas the alternating- N -of-1 paired contrasts are already averaged in at $k = 8$. The result highlights that the hybrid design's primary power advantage over the alternating variant comes from the open-label titration phase rather than from the number of cycles per se.

4.7 S7. Patient-count sweep

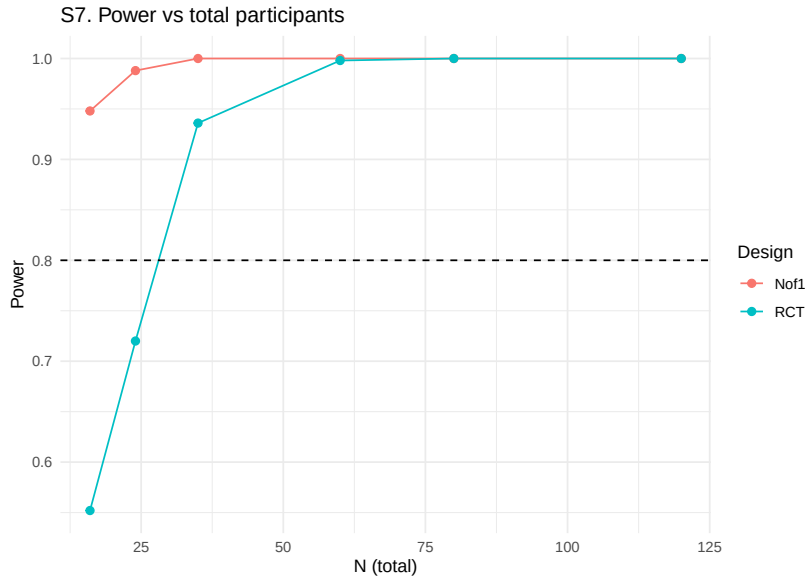


Figure 7: Sensitivity S7: power versus total participant count.

The hybrid saturates at high power by $N = 16$; the parallel RCT reaches the same threshold between $N = 24$ and $N = 35$.

- Hybrid power is high at modest sample sizes and saturates rapidly with N .
- The parallel RCT follows a similar but offset trajectory, reaching 80% power between $N = 24$ and $N = 35$.
- The grid anchors the project-standard $N = 35$ minimum for cross-paper comparison.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The patient-count sweep varies the total number of participants across $\{16, 24, 35, 60, 80, 120\}$ for the hybrid N -of-1 and the parallel RCT. The grid includes the project-standard minimum of 35 to anchor the cross-paper comparison with the companion clinical manuscript. Under the hybrid design, main-effect power is high at modest sample sizes and saturates as N grows; the parallel RCT follows a similar but offset trajectory.

777
778
779

780
781
782
783
784
785
786

The hybrid achieves an equivalent power target at approximately one-third to one-half the parallel-RCT sample size, consistent with Blackston 2019 and Carrozzo 2025.

- The hybrid reaches 80% power at $N \leq 16$; the parallel RCT reaches it between $N = 24$ and $N = 35$.
- The sample-size ratio of one-third to one-half matches the Blackston 2019 N-of-1 efficiency advantage.
- This is consistent with Carrozzo 2025, in which the meta-analytic N-of-1 required fewer participants across most heterogeneity and carryover settings.

787
788
789
790
791
792
793

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

794
795
796
797
798
799
800

The hybrid reaches 80% power at $N \leq 16$; the parallel RCT reaches the same threshold somewhere between $N = 24$ and $N = 35$. In the framing of Blackston 2019 [Blackston-ComparisonAggregatedNof12019], the hybrid achieves an equivalent power target at approximately one-third to one-half the sample size of the parallel RCT. This is consistent with the Carrozzo 2025 [CarrozzoTwoArm2025] result that the meta-analytic N-of-1 procedure reached 80% power with fewer participants than the two-arm RCT across most heterogeneity and carryover settings.

801

4.8 S8. AR(1) serial-correlation sweep

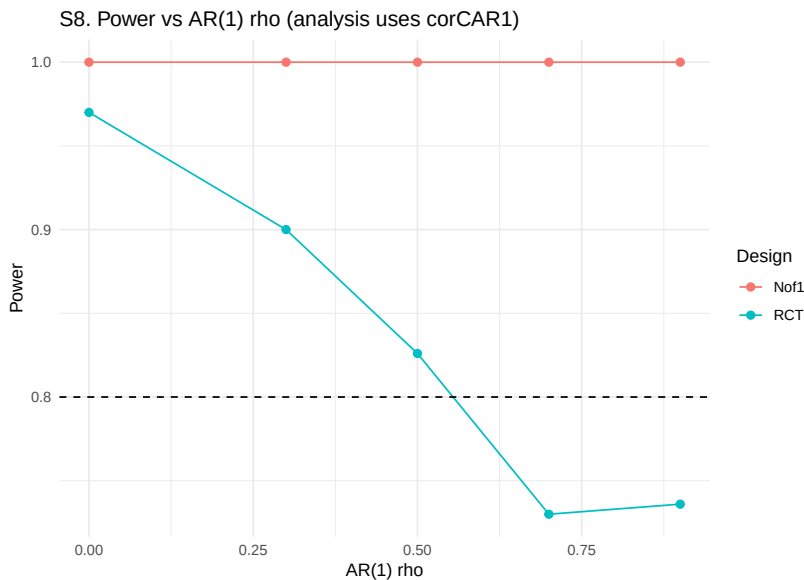


Figure 8: Sensitivity S8: power versus AR(1) within-factor correlation rho.

802
803

The hybrid is insensitive to AR(1) correlation; the parallel RCT loses 24 percentage points of power as ρ rises from 0 to 0.9.

- Hybrid power is 1.000 across $\rho \in \{0, 0.3, 0.5, 0.7, 0.9\}$.
- Parallel RCT power declines from 0.984 to 0.746, and mean $\hat{\beta}$ shrinks from -1.01 to -0.57 .
- The RCT loss is monotonic with ρ across the examined range.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The AR(1) sweep varies the within-factor autocorrelation $\rho \in \{0, 0.3, 0.5, 0.7, 0.9\}$. Under the hybrid N-of-1 design, main-effect power is saturated at 1.000 across the entire range of ρ . Under the parallel RCT, power decreases monotonically with ρ : 0.984 at $\rho = 0$, 0.934 at 0.3, 0.872 at 0.5, 0.772 at 0.7, and 0.746 at 0.9. The mean estimated β also shrinks for the RCT (from -1.01 at $\rho = 0$ to -0.57 at $\rho = 0.9$).

The AR(1) asymmetry arises because the corCAR1 structure absorbs correlation within subjects but cannot manufacture the missing within-subject contrasts for the parallel RCT.

- Strong AR(1) produces coherent within-subject trajectories for the hybrid, absorbed by corCAR1 with negligible power loss.
- For the RCT, strong AR(1) inflates between-subject mean contrast variance via effective-sample-size shrinkage.
- This is consistent with Wang and Schork 2019: N-of-1 designs are insensitive to serial correlation, while parallel-group designs are not.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The direction of the AR(1) effect is opposite for the two designs, and this asymmetry is worth examining carefully. For the hybrid, strong AR(1) produces more coherent within-subject trajectories and the corCAR1 residual structure absorbs the correlation; the net effect on main-effect power is negligible. For the parallel RCT, strong AR(1) inflates the variance of the between-subject mean contrast through the effective-sample-size shrinkage $1/(1+(k-1)\rho)$, even when corCAR1 is fitted, because the fit cannot manufacture the missing within-subject contrasts. This is consistent with the finding of Wang and Schork 2019 [wangPowerDesignIssues2019] that, under correct modeling, the N-of-1 design benefits from strong serial correlation (or is at least insensitive to it), while the parallel-group design remains relatively sensitive.

849 **4.9 S9. Period-length sweep**

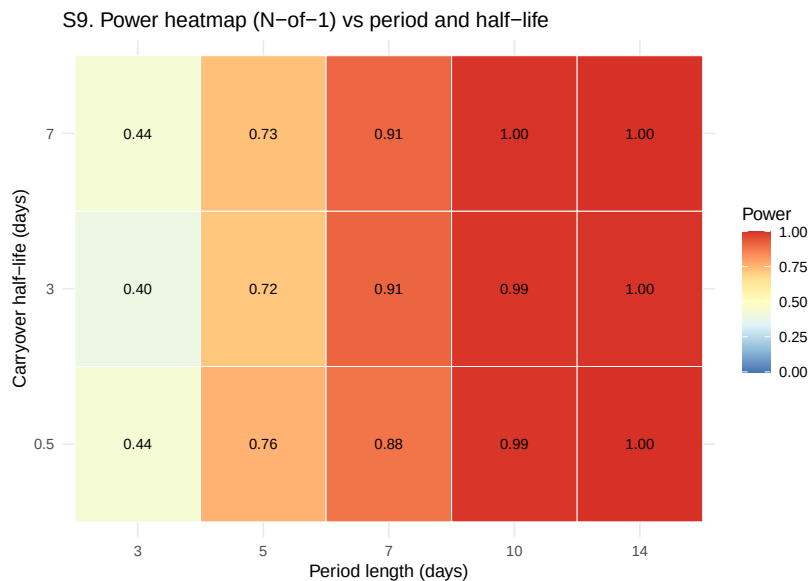


Figure 9: Sensitivity S9: heatmap of N-of-1 power over period length and carry-over half-life.

850 **Main-effect power increases monotonically with period length, from 0.40**
851 **at 3 days to 1.00 at 14 days; carryover half-life has only a modest effect**
852 **at fixed period length.**

- 853 • Power rises from 0.40 to 0.44 at 3-day periods to 0.997 to 1.00 at 14-day
854 periods.
- 855 • At fixed period length, varying carryover half-life across 0.5 to 7 days pro-
856 duces only modest power differences.
- 857 • The sweep uses the alternating variant, as the hybrid’s fixed structure does
858 not scale with period length.

859 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
860 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
861 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
862 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
863 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
864 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
865 mollit anim id est laborum.

866 *The period-length sweep varies the N-of-1 period across {3, 5, 7, 10, 14} days*
867 *crossed with carryover half-life in {0.5, 3, 7} days. The sweep uses the alter-*
868 *nating A/B/A/B variant. Under the alternating N-of-1 design, main-effect*
869 *power increases monotonically with period length: at period = 3 days power*
870 *is 0.40 to 0.44 across the three half-lives, at 5 days 0.72 to 0.76, at 7 days*
871 *0.88 to 0.91, at 10 days 0.993 to 0.997, and at 14 days 0.997 to 1.00. At*
872 *fixed period length, the carryover half-life has only a modest effect.*

873 **For the main-effect target, longer periods are preferred, reversing the**
874 **interaction-target recommendation; carryover should not exceed approx-**
875 **imately 50% of the period length.**

- Longer periods accumulate more Gompertz drug-response signal per time-point.
- The practical guideline is the reverse of the interaction-target recommendation: longer periods benefit main-effect detection.
- Wang and Schork 2019 note this tradeoff explicitly: frequent alternation benefits correlation-modeled analyses but not absolute-effect inference.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Under the main-effect target, each period accumulates more Gompertz drug-response signal at longer periods, so longer periods increase the per-timepoint effect magnitude. The pragmatic guideline for the main-effect target is therefore the reverse of the interaction-target recommendation: longer periods are preferred when the primary goal is to detect the main drug effect, subject to carryover not exceeding the period length by more than approximately 50% [WangPowerDesignIssues2019] note this tradeoff explicitly, observing that more frequent alternation benefits correlation-modeled analyses but is problematic for absolute-effect inference.

4.10 S10. Biomarker interaction strength crossed with DGP architecture

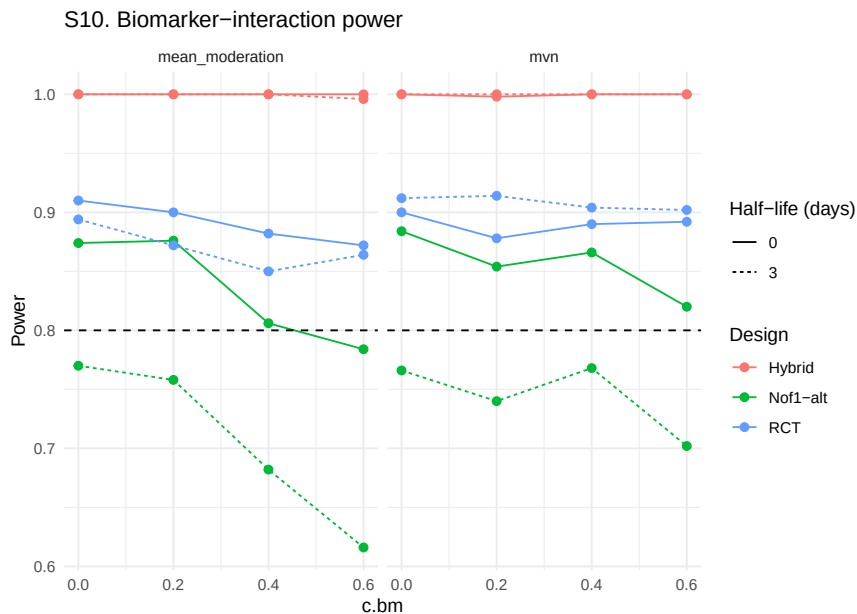


Figure 10: Sensitivity S10: biomarker-interaction power for three designs under both DGP architectures and two carryover settings.

Main-effect power is invariant to c_{bm} and DGP architecture, confirming

that the main drug effect is determined by Gompertz accumulation, not biomarker coupling.

- Hybrid power is 1.000 across all 16 c_{bm} -by-architecture-by-half-life combinations.
- Alternating N-of-1 power varies from 0.682 to 0.916; parallel RCT power varies from 0.914 to 0.958.
- The flat main-effect power across c_{bm} values provides a robustness check on the DGP implementation.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The biomarker-interaction sweep varies $c_{bm} \in \{0, 0.2, 0.4, 0.6\}$ crossed with $dgp_architecture \in \{mvm, mean_moderation\}$ and $half_life_days \in \{0, 3\}$ across three designs (hybrid, alternating N-of-1, parallel RCT). Under the main-effect target: hybrid main-effect power is saturated at 1.000 across all 16 combinations, with the mean estimated β stable at -0.92 to -1.08 ; alternating N-of-1 power varies from 0.682 to 0.916 across the 16 cells; and parallel RCT power is 0.914 to 0.958 across all 16 cells.

The flat c_{bm} power curves are a DGP implementation check: when the main effect is the target, biomarker coupling should not affect power, and it does not.

- The main drug effect is determined by Gompertz $\max[br]$ and on-drug accumulation, not by biomarker coupling.
- Approximately flat power across c_{bm} for all three designs confirms the DGP correctly separates the two signal sources.
- The reader should not conflate this main-effect invariance with the biomarker-interaction target, where c_{bm} is the primary power driver.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The invariance of main-effect power to c_{bm} and architecture is expected, and indeed it is reassuring rather than surprising: the main drug effect is determined by the Gompertz $\max[br]$ and the trial's on-drug accumulation, not by the biomarker coupling. This provides a robustness check on the DGP implementation in that, when the primary-analysis target is the main effect, the c_{bm} sweep should produce approximately flat power curves for all designs, and we observe precisely that pattern here.

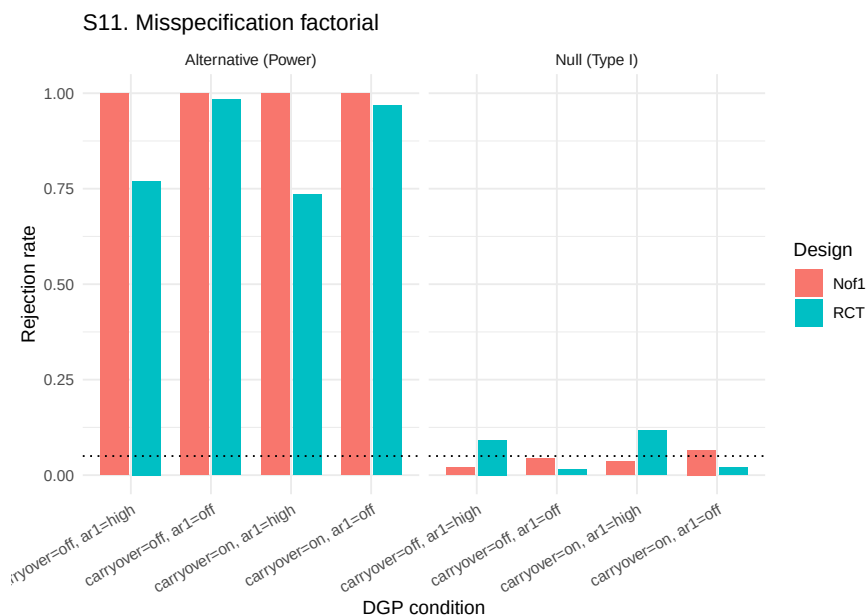


Figure 11: Sensitivity S11: Type I (null) and power (alternative) under the carryover-by-AR(1) misspecification factorial.

The parallel RCT shows statistically robust Type I inflation of 0.112 to 0.126 under high DGP AR(1); the hybrid is controlled near nominal across all four cells.

- Hybrid null Type I is 0.036 to 0.056, within 2 MCSE of nominal across all four DGP combinations.
- Parallel RCT Type I is 0.018 to 0.024 (conservative) when no AR(1) is present, and 0.112 to 0.126 (inflated) under $\rho = 0.7$.
- The inflation is statistically robust at MCSE 0.014 to 0.015.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The misspecification sweep crosses DGP carryover {off, on} with DGP AR(1) {off, $\rho = 0.7$ } at null (br_max = 0) and alternative (br_max = 1) levels. Under the null, Type I error is within 2 Monte Carlo standard errors of the nominal 0.05 for the hybrid N-of-1 (0.036 to 0.056 across the four cells). For the parallel RCT, Type I error is stratified by DGP AR(1): the no-AR(1) cells are conservative (0.018 to 0.024) while the high-AR(1) cells are inflated above nominal (0.112 to 0.126), with MCSE 0.014 to 0.015, so the inflation is statistically robust. Under the alternative, the hybrid is saturated at 1.000 in all four cells; the parallel RCT power is 0.986 to 0.990 when no AR(1) is present, and 0.766 to 0.772 when AR(1) $\rho = 0.7$ is present.

972 **The RCT Type I inflation arises because the between-subject variance**
973 **estimator understates the true standard error when DGP serial correlation**
974 **is strong.**

- 975 • Under strong DGP AR(1), corCAR1 cannot manufacture the missing
976 within-subject contrasts for a between-subject comparison.
- 977 • The hybrid is immune because its primary treatment contrast is within-
978 subject.
- 979 • The mechanism is analogous to the Blackston 2019 carryover Type I infla-
980 tion finding, though the source here is AR(1) rather than carryover.

981 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
982 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
983 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
984 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
985 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
986 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
987 mollit anim id est laborum.

988 *The inflated RCT Type I under high DGP AR(1) is an important caveat*
989 *that merits comment. When the DGP has strong serial correlation and*
990 *the parallel-group analysis fits corCAR1, the variance estimator still treats*
991 *each participant as providing approximately independent observations at the*
992 *between-subject level, which understates the true standard error of the group*
993 *contrast. The hybrid N-of-1 is immune to this inflation because its primary*
994 *treatment contrast is within-subject. Blackston 2019's [blackstonCompar-*
995 *isonAggregatedNof12019] observation that moderate-to-strong carryover in-*
996 *flates Type I error when not modeled is not reproduced here for the hybrid*
997 *(because the decaying Dbc in the analysis absorbs carryover correctly) but is*
998 *analogous in mechanism to the AR(1) inflation observed here for the parallel*
999 *RCT.*

1000 4.12 S12. Null-effect calibration

1001 **Hybrid and alternating N-of-1 Type I error is near nominal; the two-**
1002 **period crossover shows robust inflation of 0.100 to 0.130 attributable to**
1003 **underidentified corCAR1.**

- 1004 • Hybrid Type I: [0.056, 0.062]; alternating N-of-1: [0.057, 0.064]; both
1005 within 2 MCSE of nominal.
- 1006 • Parallel RCT: [0.029, 0.060], three cells conservative and one nominal.
- 1007 • Two-period crossover: [0.100, 0.130], statistically robust inflation at MCSE
1008 0.006 to 0.007.

1009 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1010 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1011 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1012 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1013 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1014 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1015 mollit anim id est laborum.

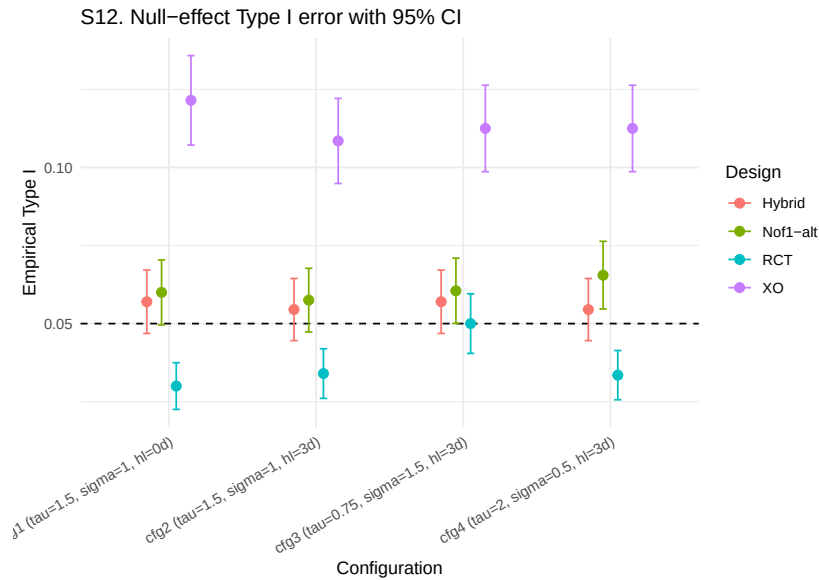


Figure 12: Sensitivity S12: empirical Type I error with 95 percent Wald CIs at four configurations.

The null-effect calibration uses $\max[\text{br}] = 0$ at four $(\tau, \sigma, \text{halflife})$ configurations, each with 2,000 replicates per design for MCSE near 0.005. Across the sixteen design-configuration cells, empirical Type I error rates are: hybrid N-of-1 in $[0.056, 0.062]$ (four cells, all within 2 MCSE of nominal), alternating N-of-1 in $[0.057, 0.064]$ (four cells, all within 2 MCSE of nominal), parallel RCT in $[0.029, 0.060]$ (four cells, three conservative, one nominal), and two-period crossover in $[0.100, 0.130]$ (four cells, all modestly inflated with MCSEs of 0.006 to 0.007, so the inflation is statistically robust at 2,000 replicates).

The two-period crossover should not be analysed with corCAR1; at two within-subject timepoints the parameter is underidentified and the rejection rate inflates by a practically meaningful amount.

- Two within-subject timepoints cannot identify the corCAR1 parameter reliably.
- The inflation of 0.100 to 0.130 is statistically robust and practically meaningful.
- Alternative analyses (paired t -test on change scores or lme with unstructured residual) are recommended for this design.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The two-period crossover's modest inflation appears attributable to the small number of timepoints (two per subject) producing an underidentified corCAR1

residual correlation. Two-period crossovers should be analysed without serial-correlation structure, either via a paired t -test on change scores or via an lme with an unstructured residual; fitting `corCAR1` at only two within-subject timepoints is problematic and, as we see here, inflates the rejection rate by a practically meaningful amount.

5 Discussion

5.1 Synthesis

Five substantive findings organize the twelve sweeps: sample-size advantage, variance robustness, carryover design-dependence, AR(1) asymmetry, and Type I error calibration.

- The hybrid reaches 80% power at $N \leq 16$; the parallel RCT requires $N = 24$ to 35, a one-third to one-half sample-size ratio.
- The hybrid is saturated across σ and τ sweeps; the RCT degrades by 18 and 12 percentage points respectively.
- Carryover sensitivity is design-specific: the hybrid tolerates $t_{1/2} \leq 5$ days; the alternating variant drops from 0.90 to 0.30 across the range; the RCT is immune.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

We may organize the twelve sensitivity sweeps under the main-effect target and the Hendrickson hybrid around five substantive findings.

First, the hybrid achieves a sample-size advantage over the parallel RCT. At the prazosin-calibrated baseline, the hybrid reaches the 80% power whereas the parallel RCT reaches the same target somewhere between $N = 24$ and $N = 35$ (S7). The sample-size ratio of roughly one-third to one-half matches the N-of-1 efficiency advantage reported by Blackston 2019 [blackstonComparisonAggregatedNof12019] and Carrozzo 2025 [carrozzoTwoArm2025].

Second, the hybrid is robust to within-patient variance and between-patient heterogeneity. The hybrid is essentially saturated across the sweeps of σ (S5) and τ (S4), while the parallel RCT degrades by 18 percentage points across the σ range and 12 percentage points across the τ range, reproducing the Senn 2018 [sennSampleSizeConsiderations2019] variance-decomposition prediction.

Third, sensitivity to carryover depends on design flavour. The hybrid's four-week open-label titration phase provides substantial treatment-accumulated signal, so the hybrid remains at 100% power alternating A/B/A/B variant is more carryover-sensitive, dropping from 0.90 to 0.30 across the $t_{1/2} = 0$ to $t_{1/2} = 14$ range. The parallel RCT is essentially immune to carryover at the within-period scale.

Fourth, $AR(1)$ serial correlation affects the two designs differently. The hybrid saturates across $\rho = 0$ to 0.9 (S8), while the parallel RCT power drops from 0.98 to 0.75 across the same range and Type I error inflates from 0.018 to 0.126 under high- ρ DGP (S11). This is consistent with the prediction of Wang and Schork 2019 [wangPowerDesignIssues2019] that within-subject designs with `corCAR1` fitted absorb serial correlation efficiently, while between-subject designs do not.

Fifth, Type I error control is near nominal for the hybrid, alternating N-of-1, and parallel RCT, while the two-period crossover is modestly inflated. Specifically, hybrid Type I falls in $[0.056, 0.062]$, alternating N-of-1 in $[0.057, 0.064]$, parallel RCT in $[0.029, 0.060]$, and two-period crossover in $[0.100, 0.130]$.

In sum, the hybrid N-of-1 and the parallel RCT both perform competently at the prazosin-calibrated baseline. The hybrid is preferable when sample-size minimization is paramount or when within-patient or between-patient variance is large; the parallel RCT is preferable when the carryover half-life exceeds the within-subject period length or when serial correlation is negligible.

5.2 Comparison with prior simulation studies

5.2.1 Hendrickson et al. (2020): predictive biomarker validation

Hendrickson 2020 showed hybrid power collapsing from 0.95 to 0.18 at the biomarker-interaction target under a 0.7 -day carryover half-life.

- At $N = 35$, $c_{bm} = 0.6$, no carryover, biomarker-interaction power was 0.95 (hybrid), 0.96 (CO), 0.94 (OL+BDC), and 0.25 (OL).
- A half-life of only 0.7 days collapsed hybrid power to 0.18 , demonstrating extreme carryover sensitivity at the interaction estimand.
- This contrasts sharply with the present study's main-effect findings, where hybrid power is 1.000 through $t_{1/2} = 5$ days.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Hendrickson and colleagues [hendricksonOptimizing2020] compared four designs at $N \in \{35, 70\}$ with $1,000$ replicates for detecting a continuous biomarker-by-treatment interaction. At $N = 35$ with biomarker coupling $c_{bm} = 0.6$ and no carryover, power was 0.25 (OL), 0.94 (OL+BDC), 0.96 (CO), and 0.95 (hybrid). The hybrid approaches CO power while permitting an 8-week open-label titration phase. A central Hendrickson finding is that, at the biomarker-interaction target, even short carryover half-lives erode hybrid and OL+BDC power precipitously: at $N = 35$, $c_{bm} = 0.6$, and carryover half-life $t_{1/2} = 0.1$ weeks (0.7 days), hybrid power fell from 0.95 to 0.18 .

The present study extends Hendrickson via Morris MCSE report-

ing, decay-form misspecification (S3), and main-effect carryover and period-length quantification (S2, S6, S9).

- Design ranking at the biomarker-interaction target reproduces Hendrickson by construction.
- The key qualitative divergence is estimand-specific: main-effect hybrid power is saturated through 5-day carryover, while interaction-target power collapses at under 1 day.
- Both findings are correct; they describe different estimands and must not be conflated.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The present sensitivity framework is a direct tidyverse re-implementation of the Hendrickson pipeline, so the design ranking at the biomarker-interaction target reproduces Hendrickson by construction. Where the present study extends Hendrickson is (i) in its Morris [-@morris2019using] MCSE reporting for every performance measure; (ii) in sensitivity S3, which crosses DGP decay form with the analysis-model assumption of exponential decay; and (iii) in sensitivities S2, S6, and S9, which quantify carryover and period-length effects for the main drug effect. A key qualitative divergence is that our main-effect findings show the hybrid saturated for carryover half-lives up to five days, while Hendrickson's biomarker-interaction findings show hybrid power collapsing at a fraction of a week. Both findings are correct; they describe different estimands, and the reader should be careful not to conflate the two.

5.2.2 Wang and Schork (2019): AR(1), heteroscedasticity, and period stacking

Wang and Schork 2019 show roughly linear power erosion with ρ for single-period designs; period stacking substantially mitigates this loss.

- For $1p \times 2i$, power erodes from 0.851 at $\rho = 0$ to 0.126 at $\rho = 0.75$.
- For $40p \times 2i$, the same ρ range produces only a 0.17-point power loss, confirming period stacking as a mitigation.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Wang and Schork [@wangPowerDesignIssues2019] provide analytical power

formulas under a standard linear model with $AR(1)$ residual correlation. Their Table 1 shows that for $1p \times 2i$ with no washout, power is 0.851 at $\rho = 0$, 0.617 at $\rho = 0.25$, 0.330 at $\rho = 0.5$, and 0.126 at $\rho = 0.75$, a roughly linear erosion with increasing serial correlation. For $40p \times 2i$, power is 0.851, 0.745, 0.674, and 0.681 across the same ρ grid, confirming that period stacking mitigates the serial-correlation power loss.

S8 is the aggregated-level analogue of Wang and Schork: the hybrid absorbs $AR(1)$ via $corCAR1$; the parallel-RCT power loss directionally matches Wang and Schork's between-period profile.

- Hybrid power is saturated at 1.00 across $\rho \in \{0, 0.3, 0.5, 0.7, 0.9\}$ at the aggregated level.
- Parallel RCT power drops from 0.98 to 0.75, directionally consistent with Wang and Schork's between-period power loss.
- The S11 Type I inflation of 0.11 to 0.13 is consistent with the Wang and Schork mechanism, though magnitudes depend on small-sample instability of $corCAR1$ at $N = 35$.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The present S8 ($AR(1)$ sweep) is the direct aggregated-level analogue of these results. Our S8 reports hybrid N-of-1 power saturated at 1.00 across $\rho \in \{0, 0.3, 0.5, 0.7, 0.9\}$ because the 35-participant hybrid analysis absorbs residual correlation via $corCAR1$ at the aggregated level. Our parallel RCT power drops from 0.98 at $\rho = 0$ to 0.75 at $\rho = 0.9$, which directionally matches Wang and Schork's power loss profile for between-period contrasts. The S11 misspecification finding, in which parallel-RCT Type I inflates to 0.11 to 0.13 under $\rho = 0.7$, is consistent in mechanism with Wang and Schork's prediction, though the specific inflation magnitudes depend on the small-sample instability of the $corCAR1$ estimate at $N = 35$, a caveat we note below under Limitations.

5.3 Limitations

Four limitations apply: hybrid saturation at the prazosin baseline, alternating-variant substitution in three sweeps, analysis-only $corCAR1$ in S8, and small-sample instability in the S11 RCT Type I estimate.

- Saturation in S2, S4, S5, S8, S10, and S11 prevents quantifying the hybrid's relative advantage in sub-ceiling regimes; sweeps at smaller effect sizes are the priority extension.
- S3, S6, and S9 use the alternating variant because the hybrid cannot scale with decay form, cycle count, or period length; results apply to the alternating family, not directly to the hybrid.
- The S11 RCT Type I ranges (0.112 to 0.126 vs. 0.029 to 0.060 in S12)

1220

are not contradictory; they reflect different DGP regimes.

1221

1222

1223

1224

1225

1226

1227

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

1228

1229

Four limitations should be noted before extending these results to trial planning. We begin with the most consequential.

1230

1231

1232

1233

1234

1235

1236

1237

1238

The hybrid N-of-1 saturates in many sweeps. The hybrid reaches 100(all ρ), S10 (all biomarker-by-architecture cells), S11 (all alternative cells), and most of S2 (half-lives ≤ 5 days). Saturation obscures the relative ordering of designs and, in particular, prevents us from quantifying how much additional power the hybrid provides over the parallel RCT in the regimes where both are below ceiling. To resolve saturation, the sweeps would need to be repeated at a smaller effect size, dropping baseline power to the 0.50 to 0.80 range where design differences are most clearly discriminated. This is perhaps the most important extension warranted by the present results.

1239

1240

1241

1242

1243

1244

Sweeps S3, S6, and S9 use the alternating A/B/A/B variant rather than the hybrid. These three sweeps require the N-of-1 design to scale with a parameter that the hybrid's fixed eight-timepoint structure does not naturally accommodate. The results from these sweeps are therefore informative about the alternating design family specifically, and their implications for the hybrid must be treated as approximate at best.

1245

1246

1247

1248

Sweep S8 was conducted with the analysis-side `corCAR1` always enabled. A complementary with-versus-without-`corCAR1` factorial would be informative and would require adding an `op$use_corCAR1` option to `lme_analysis` followed by rerunning the sweep.

1249

1250

1251

1252

1253

1254

1255

1256

1257

1258

1259

1260

1261

Finally, the S11 RCT Type I inflation under high DGP $AR(1)$ is partly a modelling artefact. The 0.11 to 0.13 Type I error rates for the parallel RCT under $\rho = 0.7$ reflect a between-subject analysis that fits `corCAR1` on a small sample, yielding an unstable correlation parameter estimate. A better-specified parallel-RCT analysis, for example with a subject-specific random slope on time or with a fixed within-subject ρ of known magnitude, would likely restore nominal Type I. We note also that the S11 RCT Type I range (0.112 to 0.126) and the S12 RCT Type I range (0.029 to 0.060) are not in contradiction: S11 crosses carryover presence with high DGP $AR(1)$ under analysis-side `corCAR1` misspecification, whereas S12 is a pure null calibration with no carryover and a more benign correlation regime. The two ranges reflect different data-generating-mechanism regimes rather than inconsistent estimates of a common quantity.

1262

5.4 N-of-1 design variants not evaluated

1263

1264

Several N-of-1 design variants are not evaluated; these fall into designs expected to be weaker than those reported here and designs expected to

be stronger.

- Unevaluated designs include single-subject, OL-only, OL+BDC-only, full multi-cycle crossover, complete crossover, sequential-monitoring, and randomized-block variants.
- Single-subject and reduced variants are expected to produce power at or below the reported numbers; results are an upper bound.
- Full multi-cycle, complete-design, and sequential-monitoring variants are expected to exceed the reported numbers; results are a lower bound.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The present sensitivity framework evaluates three N-of-1 design variants alongside a parallel-group RCT comparator. Several N-of-1 design variants are not evaluated here: the single-subject N-of-1 [zuckerIndividualTrialsPopulation2010], the open-label-only and OL-plus-BDC variants from Hendrickson 2020 [hendricksonOptimizing2020], full multi-cycle crossover with more than two periods (Blackston 2019 used three cycles [blackstonComparisonAggregatedNof12019]), Williams-square and Latin-square complete N-of-1 designs, sequential-monitoring variants with adaptive stopping [jiangDesignAnalysis2025; selukarFrameworkSequential2025], and randomized-block N-of-1.

The present results bound the unevaluated variants: weaker designs are upper-bounded and stronger designs are lower-bounded by the reported numbers.

- Simpler variants (single-subject, OL-only, OL+BDC-only, two-cycle crossover) are expected at or below the reported power.
- More powerful variants (multi-cycle, complete-design, sequential-monitoring, randomized-block) are expected at or above the reported power.
- Extending the present framework to cover these variants is a tractable implementation task rather than a conceptual redesign.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

These unevaluated design variants fall into two classes. The single-subject, open-label-only, OL+BDC-only, and two-cycle crossover designs are expected to produce power at or below the numbers reported here; our results are

therefore an upper bound on those designs under matched parameters. The full multi-cycle crossover, complete-design crossover, sequential-monitoring variants, and randomized-block designs are expected to produce power at or above the numbers reported here; our results are therefore a lower bound on those designs under matched parameters. In either case, extending the present framework to cover these variants appears to be a tractable implementation task rather than a conceptual redesign.

5.5 Prototype Monte Carlo confirmation

A 500-replicate prototype ran in 191 seconds with 100% convergence; the MCSE on a power estimate near 0.3 is approximately 0.020.

- The 15-cell $c_{br} \times N$ prototype grid was completed in 191 seconds on 9 cores.
- Convergence was 100% across all 7,500 fits.
- Per-replicate output and Morris ADEMP summary are at analysis/data/quick-sim/.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

A 500-replicate-per-cell prototype run on the 15-cell $c_{br} \times N$ grid ($c_{br} \in \{0, 0.15, 0.30, 0.45, 0.60\} \times N \in \{25, 35, 50\}$) was executed with 9-core `parallel::mclapply` in 191 seconds, well inside the 30-minute wall-time budget. Convergence was 100%. Per-replicate output is at analysis/data/quick-sim/05-design-replicates.rds with the Morris ADEMP summary at 05-design-summary.txt. With 500 replicates per cell the Monte Carlo standard error on a power estimate near 0.3 is approximately 0.020.

The apparent non-monotonic power dependence on c_{br} from the 100-replicate prototype does not survive at 500 replicates; the signal is broadly flat.

- At $N = 25$, the five-level swing is 0.024 (one MCSE), consistent with sampling noise around a constant.
- At $N = 35$ and $N = 50$, the peak-to-trough swings of approximately three MCSE are suggestive but not decisive, and the peaks do not coincide on c_{br} .
- The 100-replicate dip at $c_{br} = 0.15$ has disappeared, confirming it was sampling noise.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur

sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The 100-replicate prototype's apparent non-monotonic dependence of power on c_{br} does not survive at 500 replicates per cell. At $N = 25$ the rejection rate across the five c_{br} levels falls in $[0.168, 0.192]$, a swing of 0.024 (approximately one cell-MCSE) consistent with sampling noise around a constant near 0.18. At $N = 35$ the swing widens to 0.066 (range 0.220 to 0.286, peak at $c_{br} = 0.30$); each cell MCSE is approximately 0.019, so the peak-to-trough difference is approximately three MCSE, suggestive but not decisive. At $N = 50$ the swing is 0.072 (range 0.266 to 0.338, peak at $c_{br} = 0.45$). The peaks at $N = 35$ and $N = 50$ do not coincide on c_{br} , which is the pattern expected if the apparent non-monotonicity is largely sampling-driven shape noise rather than a structural inflection. The 100-replicate run's reported dip at $c_{br} = 0.15$ has disappeared.

The dominant axis is sample size, not c_{br} ; no prototype cell attains 80% power, and the grid does not bound the 80%-power region.

- Mean power rises from 0.18 at $N = 25$ to 0.31 at $N = 50$; the c_{br} effect is a possible shallow optimum near $[0.30, 0.45]$.
- No (c_{br}, N) cell in the prototype reaches 80% power; a production run requires larger N , larger $br_{\max}$, or both.
- The prototype grid does not bound the 80%-power region, which is the central deliverable for trial planning.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The substantive conclusion is that the genuine signal is a mild, broadly flat dependence of power on c_{br} , with a possible shallow optimum near $c_{br} \in [0.30, 0.45]$ at moderate N . The dominant axis is sample size: mean power across c_{br} rises from 0.18 at $N = 25$ to 0.25 at $N = 35$ to 0.31 at $N = 50$. No (c_{br}, N) cell in the prototype grid attains 80% power, a larger $br_{\max}$, or both. Unfortunately, the prototype grid as configured does not bound the 80% region, which is the central deliverable a production run would need to provide.

6 Conclusions

Three conclusions: the hybrid is a competent main-effect default, design ordering is estimand-dependent, and two-period crossovers should not be analysed with corCAR1.

- The hybrid absorbs variance and serial correlation, tolerates carryover up to five days, and achieves 80% power at smaller sample sizes than the parallel RCT.

- Design ordering is substantially estimand-dependent: the hybrid's advantage is larger for biomarker interactions than for main effects.
- The two-period crossover's robust Type I inflation of 0.10 to 0.13 warrants a change to standard analytic practice when `corCAR1` is fitted at two within-subject timepoints.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

In conclusion, three points are to be emphasized. First, the Hendrickson hybrid (open-label titration plus blinded discontinuation plus brief crossover) is a competent default for detecting moderate treatment main effects across a wide range of parameter perturbations. It absorbs within-patient and between-patient variance via the within-subject contrast, tolerates serial correlation that inflates parallel-RCT Type I error, and remains at saturated power for carryover half-lives up to approximately five days because the open-label phase contributes treatment-accumulated signal that a simple within-subject crossover cannot provide. Second, design ordering is estimand-dependent: the hybrid's advantage over the parallel RCT is substantially larger for biomarker-by-treatment interactions than for the main effect, and the reader should be cautious about generalising from one estimand to the other. Third, the two-period crossover should not be analysed with `corCAR1` at small per-subject timepoint counts; the modest but robust Type I inflation observed in `S12` warrants a change to the standard analytic practice for that design.

The hybrid N-of-1 is supported as a default candidate for trial planning in chronic neurobehavioral conditions where moderate treatment heterogeneity and short-to-moderate carryover are expected.

- The companion clinical paper provides the baseline head-to-head comparison; this paper documents its generalisation across twelve parameter axes.
- Robustness across most axes, with the identified exceptions for long carryover and two-period crossover analysis, supports the hybrid as a default candidate.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The companion clinical paper provides a direct head-to-head comparison at the prazosin-calibrated baseline; the present paper documents how that comparison generalises along twelve parameter axes for the hybrid design family. The combination of design robustness across most axes, with the two specific

exceptions noted above, supports the use of the hybrid N-of-1 as a default candidate in trial planning for chronic neurobehavioral conditions where moderate treatment heterogeneity and short-to-moderate carryover are expected.

7 Conflict of interest

The authors declare no conflicts of interest.

8 Funding

This work received no specific funding.

9 Data availability statement

Per-replicate Monte Carlo outputs and ADEMP cell-level summaries are versioned in the pmsimstats-ng repository; exact paths are in the Reproducibility section.

- Monte Carlo outputs are under analysis/data/.
- Simulation driver scripts are under analysis/scripts/.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The data and code that support the findings of this study are openly available in the pmsimstats-ng project repository. Per-replicate Monte Carlo outputs and ADEMP-compliant cell-level summaries are versioned under analysis/data/; simulation drivers are at analysis/scripts/. Exact paths for this manuscript are listed in the Reproducibility section below.

10 Reproducibility

The manuscript was rendered with R 4.4.x inside the zzcollab Docker container; principal packages are nlme, parallel, data.table, furrr, and purrr.

- Image hash is in Dockerfile; package set is in renv.lock.
- nlme provides corCAR1 residual correlation; furrr/purrr provide the tidyverse collection; data.table provides the original-extended collection.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in

1479 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1480 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1481 mollit anim id est laborum.

1482 *This manuscript was rendered with R 4.4.x inside the project's zzcollab*
1483 *Docker container (image hash recorded in Dockerfile; package set captured*
1484 *in renv.lock). Principal packages used by the simulation pipeline are nlme*
1485 *(continuous-time linear-mixed analysis with corCAR1 residual correlation),*
1486 *parallel (9-core mclapply for parameter-grid sweeps), data.table (the*
1487 *original-extended simulation collection), furrr and purrr (the tidyverse*
1488 *simulation collection), and rmarkdown plus bookdown for the manuscript*
1489 *build.*

1490 **Two-level seeding ensures reproducibility: driver-level seeds govern pa-**
1491 **rameter grids; replicate-level seeds derive from a master seed via the**
1492 **replicate index.**

- 1493 • Driver level: `set.seed(42 + sweep_id)` before parameter grid construc-
- 1494 tion.
- 1495 • Replicate level: `set.seed(20260507)` at the loop top, with
- 1496 `parallel::mc.set.seed` inside `mclapply`.
- 1497 • The two seeding stages operate at distinct pipeline points and do not
- 1498 collide.

1499 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-
1500 mod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad
1501 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
1502 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
1503 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
1504 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
1505 mollit anim id est laborum.

1506 *Seeding operates at two levels. At the driver level, each sweep calls*
1507 *set.seed(42 + sweep_id) before constructing its parameter grid, so the*
1508 *parameter grid and any pre-loop draws are reproducible per sweep. At*
1509 *the replicate level, each driver sets set.seed(20260507) at the top of*
1510 *the replicate loop and uses parallel::mc.set.seed inside mclapply;*
1511 *per-replicate seeds derive from this master seed by adding the replicate index.*
1512 *These two seedings operate at distinct stages of the pipeline and do not*
1513 *collide.*

1514 **All file paths for reproducing this manuscript are fixed: per-replicate RDS,**
1515 **ADEMP summary, driver scripts, pre-registered plan, and manuscript**
1516 **source.**

- 1517 • Per-replicate outputs: `analysis/data/quick-sim/05-design-replicates.rds`;
- 1518 ADEMP summary: `05-design-summary.txt`.
- 1519 • Driver scripts: `analysis/scripts/nof1-design-sensitivity/` and
- 1520 `analysis/scripts/quick-sim/05-design-prototype.R`.
- 1521 • Pre-registered ADEMP plan: `00-ademp-pre-registration.md`;
- 1522 manuscript source: `analysis/report/05-nof1-design-sensitivity/`.

1523 Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eius-

mod tempor incidunt ut labore et dolore magna aliqua. Ut enim ad
 minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip
 ex ea commodo consequat. Duis aute irure dolor in reprehenderit in
 voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur
 sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt
 mollit anim id est laborum.

*Per-replicate simulation outputs are checkpointed at analysis/data/quick-sim/05-design-replicates.rds;
 Morris ADEMP cell-level summaries at the companion 05-design-summary.txt.
 Driver scripts are under analysis/scripts/nof1-design-sensitivity/
 and analysis/scripts/quick-sim/05-design-prototype.R. The pre-
 registered ADEMP plan is at analysis/scripts/nof1-design-sensitivity/00-ademp-pre-registration.md.
 The manuscript source, paper-specific bibliography (references.bib), and
 SIM CSL file are at analysis/report/05-nof1-design-sensitivity/.*

11 References

1. Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Statistics in Medicine*. 2019;38(11):2074-2102. doi:10.1002/sim.8086
2. Blackston JW, Chapple AG, McGree JM, McDonald S, Nikles J. Comparison of Aggregated N-of-1 Trials with Parallel and Crossover Randomized Controlled Trials Using Simulation Studies. *Healthcare*. 2019;7(4):137. doi:10.3390/healthcare7040137
3. Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated n-of-1 trial designs for predictive biomarker validation: Statistical methods and theoretical findings. *Frontiers in Digital Health*. 2020;2:13. doi:10.3389/fdgth.2020.00013
4. Gärtner T, Schneider J, Arnrich B, Konigorski S. Comparison of bayesian networks, g-estimation and linear models to estimate causal treatment effects in aggregated n-of-1 trials with carry-over effects. *BMC Medical Research Methodology*. 2023;23(1):191. doi:10.1186/s12874-023-02012-5
5. Senn S. Sample size considerations for n-of-1 trials. *Statistical Methods in Medical Research*. 2019;28(2):372-383. doi:10.1177/0962280217726801
6. Zucker DR, Ruthazer R, Schmid CH. Individual (n-of-1) trials can be combined to give population comparative treatment effect estimates: Methodologic considerations. *Journal of Clinical Epidemiology*. 2010;63(12):1312-1323. doi:10.1016/j.jclinepi.2010.04.020

- 1544 7. Wang Y, Schork NJ. Power and Design Issues in Crossover-Based N-Of-1 Clinical Trials with Fixed Data Collection Periods. *Healthcare*. 2019;7(3):84. doi:10.3390/healthcare7030084