

1 Three-component decomposition of treatment response
2 in aggregated N-of-1 trials: pharmacological,
3 expectation-driven, and natural-history components

4 pmsimstats team

5 2026-06-13

6 **Contents**

7	Abstract	2
8	1 Introduction	3
9	1.1 A claim about how N-of-1 trials should be analysed	3
10	1.2 Why this matters: three failure modes of the lumped analysis . . .	4
11	1.3 The current state of the N-of-1 literature	7
12	1.4 The neurobiological and epidemiological reality of the components	11
13	1.5 What this paper contributes	12
14	2 Notation and the formal model	16
15	3 The three components	18
16	3.1 Biological response (BR)	19
17	3.2 Placebo-belief response (PB)	22
18	3.3 Time-variant natural-history component (TV)	31
19	4 Identifiability and trial-design contrasts	36
20	5 Analysis-side methodology	40
21	5.1 Four reasons not to match the DGP	40
22	5.2 A small extension that usually pays off: phase indicators	43
23	5.3 When to fit the full decomposition	45
24	6 Worked example: information loss without decomposition	47
25	6.1 What the lumped analysis estimates	47
26	6.2 A numerical illustration	49
27	7 Simulation study	54
28	7.1 Pilot validation of analysis machinery (100 replicates, 6 cells) . . .	55
29	7.2 Production-simulation specification	57
30	7.3 Study A results: bias and power at 250 replicates per cell	59

31	8 Application to predictive-biomarker validation	62
32	9 Discussion	65
33	9.1 Interpretation of the pilot diagnostics	68
34	10 Conclusions	70
35	11 Conflict of interest	72
36	12 Funding	72
37	13 Data availability statement	72
38	14 Reproducibility	73
39	15 References	74

40 Abstract

41 **Background.** Treatment response in aggregated N-of-1 clinical trials reflects
42 three causally distinct mechanisms: pharmacological action of the active com-
43 pound (BR), the patient’s belief that the treatment will work (PB), and natural-
44 history drift in the underlying condition (TV). Standard analysis models lump
45 these into a single ‘treatment effect’ plus residual error. Whether this lumping
46 biases inference is not unconditional: it depends on the estimand and on whether
47 candidate biomarkers correlate with the non-pharmacological components, a de-
48 pendence this paper makes precise.

49 **Methods.** We formalise a three-component decomposition of the within-
50 participant outcome trajectory, each component a modified Gompertz function
51 with participant-specific random effects, and characterise the trial-design con-
52 trasts (open-label, blinded discontinuation, crossover) that supply identifiability.
53 We derive the omitted-variable-bias identity for the one-component estimator,
54 which expresses both the lumped treatment effect and the lumped biomarker
55 slope as sums of component-specific terms, and we evaluate the analysis strategies
56 in a pre-registered Monte Carlo study (Study A) calibrated to the prazosin-PTSD
57 reference parameters, comparing a one-component design-contrast analysis with
58 a phase-augmented analysis at 250 replicates per cell across $m_{PB} \in \{0, 1, 3, 6, 10\}$,
59 $m_{TV} \in \{-1, 0, 1, 2\}$, and $N \in \{35, 70, 100, 150\}$.

60 **Results.** The identity shows that the lumped biomarker-by-treatment slope is dis-
61 placed from the pharmacological slope by a term proportional to the biomarker’s
62 covariance with PB and TV, not by the magnitude of those components. Study A
63 confirms this: with a biomarker coupled to BR only, the one-component estimate
64 of the biomarker-by-treatment interaction is essentially unbiased across the whole
65 grid (mean absolute bias from 0.040 at $N = 35$ to 0.012 at $N = 150$, within one
66 Monte Carlo standard error of zero), with near-nominal coverage and type I error.
67 The phase-augmented analysis provides no bias reduction and is markedly less

powerful (mean power 0.35 versus 0.59 at $N = 35$; 0.62 versus 0.91 at $N = 70$), so it is dominated under this uncontaminated biomarker. The inferential value of decomposition is therefore conditional, not universal.

Conclusions. The trial design must supply the drug-on/off and belief contrasts; given those, component-level decomposition is useful when the estimand is the attribution of response to pharmacology versus belief versus natural history, or when a candidate biomarker may correlate with PB or TV. For a biomarker-by-treatment interaction with a biomarker orthogonal to PB and TV, a single-indicator analysis that exploits the design contrast is sufficient and a parametric component model is counterproductive. A pragmatic specification for the cases that warrant it is $Sx \sim \text{bm} + \text{t} + \text{Dbc} + \text{phase} + \text{Dbc:phase} + \text{bm:Dbc} + \text{bm:Dbc:phase}$ with random intercept and AR(1) residual correlation. The decisive test of the biomarker-validation failure, a data-generating process that couples the biomarker to PB, remains to be run.

1 Introduction

1.1 A claim about how N-of-1 trials should be analysed

[1]

This paper argues that treatment response in N-of-1 trials comprises three causally distinct components that must be analysed jointly rather than lumped.

- The three components are *BR* (pharmacological), *PB* (placebo-belief), and *TV* (natural history).
- Each component is causally distinct, clinically interpretable, and identifiable from a well-designed N-of-1 trial.
- The standard single-indicator mixed-effects model conflates all three and renders biomarker validation unanswerable.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

This paper advances an opinionated methodological claim: treatment response in aggregated N-of-1 clinical trials is the sum of three causally distinct components and should routinely be analysed as such, rather than estimated as a single lumped quantity. The three components are pharmacological action (*BR*, the biological response), placebo-belief response (*PB*, the patient's belief that they are receiving treatment), and time-variant natural history (*TV*, the trajectory that would have occurred without intervention). Each is causally distinct, each is independently of clinical and regulatory interest, and each is identifiable from a properly designed N-of-1 trial. We shall argue that the standard practice of fitting a single mixed-effects model with one treatment indicator, prevalent across the published N-of-1 literature, conflates these

three components, biases the interpretation of treatment effect, and renders predictive-biomarker validation unanswerable as a scientific question.

[2]

The contribution is the consolidation and extension of an established structural-decomposition tradition, not a novel decomposition.

- The additive disease-progression, placebo, and drug-action decomposition is well established in pharmacometrics and longitudinal trial methodology.
- The present framework extends that tradition with an explicit, separately identified placebo-belief channel and brings it to the aggregated N-of-1 setting.
- The decomposition logic is not specific to N-of-1; it applies wherever a longitudinal treatment response is analysed.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

We make no claim of novelty for the decomposition itself. The structural decomposition of a longitudinal treatment response into separable disease-progression, placebo, and drug-action components is a well-established tradition in pharmacometrics¹⁻³, and the separation of drug from placebo response in longitudinal trajectories has been pursued both through growth-mixture models⁴ and through experimental decomposition of the placebo response itself⁵. Each ingredient we use (mixed-effects models with trial-phase indicators, placebo-arm contrasts, regression-to-mean adjustment, growth-curve parameterisation) has likewise appeared across decades of clinical-trial methodology. Our contribution is therefore one of consolidation and extension rather than invention: we make the three-component structure explicit for the aggregated N-of-1 setting, add a separately modelled placebo-belief channel identified by design contrasts rather than by functional form alone, and argue that the decomposition should become routine wherever a longitudinal treatment response is analysed. We document why the framing has not been routine in N-of-1 practice (parsimony, identifiability concerns, sample-size constraints, methodological inertia) and propose an analysis that is tractable at trial-relevant sample sizes through a phase-augmented linear-mixed-effects approximation to the full decomposition.

1.2 Why this matters: three failure modes of the lumped analysis

[3]

The lumped one-component analysis has three documented

149 **failure modes that grow with unmodelled component hetero-**
150 **geneity.**

- 151 • The three failure modes are placebo-attribution bias, regression-
152 to-the-mean, and biomarker-validation failure.
- 153 • Each failure mode is documented in the chronic-condition trial
154 literature, not merely hypothetical.
- 155 • The failures are present regardless of sample size.

156 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
157 *tempor incididunt ut labore et dolore magna aliqua.*

158 We begin with the failure modes of the standard approach. A one-
159 component analysis (‘the treatment effect is what we observe minus
160 baseline’) has three identifiable failure modes that arise when the pop-
161 ulation displays meaningful heterogeneity in any of the unmodelled
162 components. None of these is hypothetical: each has a documented
163 presence in the chronic-condition trial literature. Two points of preci-
164 sion, established formally in section 6.1 and borne out by the simula-
165 tion of section 7.3, should be carried through what follows. The first
166 two failure modes attach to the baseline-anchored lumped effect and
167 are removed once the design supplies a drug-on/off contrast, which
168 the one-component analysis of later sections in fact uses. The third,
169 biomarker-validation failure, is different in kind: it requires not merely
170 that PB or TV be large but that the candidate biomarker correlate
171 with them, and where that coupling is absent the lumped analysis is
172 not biased at all.

173 [4]

174 **Placebo-attribution bias is sample-size-invariant and docu-**
175 **mented across antidepressant, anxiety, and IBS literatures.**

- 176 • Lumped treatment-effect estimates overstate pharmacological ef-
177 ficacy in proportion to PB magnitude and variance.
- 178 • Kirsch et al. and Locher et al. demonstrated that antidepressant
179 apparent efficacy is largely explained by placebo response and
180 natural fluctuation.
- 181 • Placebo responsiveness has measurable biological correlates (e.g.
182 COMT genotype), confirming it is a structured phenomenon, not
183 noise.

184 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
185 *tempor incididunt ut labore et dolore magna aliqua.*

186 The first is the placebo-attribution failure. When the population con-
187 tains both high and low placebo responders, the lumped treatment
188 effect overestimates the pharmacological component. The bias scales
189 with the magnitude and variance of PB in the study population, and
190 it is invariant under sample-size increases: a larger one-component

trial converges on the same biased estimate. The most consequential example in modern clinical pharmacology is antidepressant efficacy.⁶ analysed FDA-submitted data on four widely-prescribed antidepressants and showed that, except at the most severe baseline-severity stratum, mean drug-placebo differences fell below clinical-significance thresholds, with the apparent ‘efficacy’ largely explained by placebo response and natural fluctuation;⁷ reproduced this pattern in late-life depression and explicitly invoked regression to the mean.⁸ generalised the finding to anxiety. Patient populations differ systematically in placebo responsiveness, and the difference has measurable biological correlates: COMT val158met genotype predicts placebo response in irritable bowel syndrome^{9,10}, indicating that ‘placebo’ is not a homogeneous nuisance but a structured biological phenomenon that interacts with the patient’s genome and treatment history. A trial whose population happens to be placebo-responsive will report a larger ‘drug effect’ than one whose population is not, even when the underlying pharmacology is identical.

[5]

Regression-to-the-mean contributes several points of apparent improvement in early trial weeks and is absorbed without trace by the lumped analysis.

- Enrolment at peak severity guarantees a natural-history regression component irrespective of pharmacology.
- Hrobjartsson et al. systematic reviews across 200+ trials showed much of conventional placebo response was indistinguishable from natural-history drift.
- The lumped analysis has no formal mechanism to recover the pharmacological effect once regression is absorbed.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The second failure mode is regression-to-the-mean. When patients enrol at moments of clinical engagement that coincide with peak symptom severity, natural-history regression toward typical severity contributes several points of apparent improvement in the early weeks of any trial, without any pharmacological contribution at all. The systematic reviews of^{11, 12}, and¹³, the last covering 202 trials across 60 clinical conditions, quantified this directly: in trials with both placebo and no-treatment arms, much of what is conventionally attributed to the placebo response was indistinguishable from regression to the mean and natural-history drift. The lumped analysis absorbs this regression into the treatment-effect estimate and the analyst has no formal mechanism to recover the true pharmacological effect.

[6]

Biomarker validation is invalidated when a biomarker correlates with PB or TV , rendering precision-medicine stratification wrong regardless of trial size.

- A biomarker correlated with placebo responsiveness or natural-history trajectory is indistinguishable from a true pharmacological predictor in a lumped analysis.
- The PATH statement defers mechanistic attribution to domain knowledge; the BR-PB-TV decomposition supplies that substrate.
- Where such coupling cannot be ruled out, rigorous pharmacological-biomarker validation requires the decomposition; where the biomarker is known orthogonal to PB and TV , the lumped analysis already recovers the pharmacological slope.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The third failure mode is the biomarker-validation failure. The precision-medicine question that motivates many recent N-of-1 trials is whether a baseline biomarker predicts treatment response, formalised in the PATH statement^{14,15}. A biomarker that correlates with placebo responsiveness or with a favourable natural-history trajectory will appear in a one-component analysis as a predictor of the lumped ‘treatment effect’, indistinguishable from a true pharmacological predictor. The biomarker-stratified prescribing decision that follows from such an analysis is wrong, and the trial’s regulatory utility is degraded. The PATH framework explicitly defers the mechanism of treatment-effect heterogeneity to domain knowledge^{16,17}; the BR-PB-TV decomposition supplies the mechanistic substrate that PATH leaves blank. We argue that, whenever the candidate biomarker may itself correlate with PB or TV , predictive-biomarker validation cannot be conducted rigorously in a lumped-analysis framework regardless of how large the trial is; the bias is fixed by the biomarker’s coupling to the non-pharmacological components, not by sample size. Where the biomarker can be assumed orthogonal to PB and TV , this failure mode does not arise, and the lumped analysis recovers the pharmacological slope, as the simulation of section 7.3 confirms.

1.3 The current state of the N-of-1 literature

[7]

The N-of-1 literature has matured substantially but motivates three observations that the present paper addresses.

- The design has evolved from individual-patient consultation to a framework supporting population-level inference.

- 275 • CONSORT reporting standards and quality audits document pro-
276 gressive improvement in analytic transparency.
277 • Three specific gaps in the literature motivate the decomposition
278 framework proposed here.

279 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
280 *tempor incididunt ut labore et dolore magna aliqua.*

281 The N-of-1 trial in its modern randomised form originates with¹⁸, who
282 defined the within-patient multiple-crossover RCT, demonstrated it on
283 a single airflow-limitation case, and established a clinical service to de-
284 liver it at scale. The design inherits its identifying logic from two
285 older traditions: the crossover trial of classical biostatistics, in which
286 within-subject period contrasts estimate treatment effects under ex-
287 plicit carryover assumptions, and the single-case experimental designs
288 of behavioural science, in which phase reversals (withdrawal and rein-
289 statement of treatment) identify effects within one participant. The
290 blinded-discontinuation contrast central to this paper is a direct de-
291 scendant of both. The methodological step from the individual-patient
292 trial to population-level inference was taken by¹⁹ and²⁰, who showed
293 that a series of single-patient trials can be combined through hierarchi-
294 cal meta-analysis to estimate both population treatment effects and
295 individual responses, the statistical foundation on which the aggre-
296 gated design rests. Over the last fifteen years the design has accord-
297 ingly been reframed from an individual-patient decision tool into an
298 aggregated, meta-analytic framework supporting population-level in-
299 ference: the programmatic statement of²¹, the methodological reviews
300 of²², the historical-practice audit of²³, and the more recent syntheses
301 of^{24, 25, 26}, and²⁷ trace that trajectory. Reporting standards are codi-
302 fied in the CONSORT extension²⁸ (the headline statement) and²⁹ (the
303 explanation-and-elaboration companion); the reporting-quality audits
304 of³⁰ and its successors document slow but progressive improvement in
305 analytic transparency. Three observations about this literature moti-
306 vate the present paper.

307 [8]

308 **No published N-of-1 trial employs a structural decomposition**
309 **of treatment response despite three decades of methodologi-**
310 **cal literature.**

- 311 • Audit data show the field has relied on paired t-tests, binary-
312 indicator LME, and hierarchical Bayesian regressions.
313 • Recent methodological advances (Liao 2023, Blackston 2019, Car-
314 rozzo 2025) retain the lumped-response framing.
315 • The decomposition gap persists even as carryover and identifica-
316 bility problems are increasingly recognised.

317 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*

318 *tempor incididunt ut labore et dolore magna aliqua.*

319 The first observation is that the dominant analytic models are simple. The reporting-quality audit of³⁰ documents that the published
320 N-of-1 trials between 1985 and 2013 relied predominantly on paired
321 t-tests, simple linear mixed-effects models with a binary treatment indicator, or hierarchical Bayesian regressions. On our own reading of
322 that corpus and the methodological literature that followed it, no published N-of-1 trial has employed a structural decomposition of treatment
323 response into pharmacological, expectancy, and natural-history components. The most recent methodological advances retain this
324 character:³¹ develops Bayesian distributed-lag models with autocorrelated errors as a current carryover-aware analysis, but the model
325 still treats the response as a single quantity to be estimated.³² and³³
326 compare aggregated N-of-1 trials with parallel and crossover RCT alternatives by simulation and document carryover-induced type-I-error
327 inflation as a key failure mode, but again at the level of the lumped
328 treatment effect. None of these papers decompose the response.
329
330
331
332
333
334

335 [9]

336 **Placebo-component separation is widely recognised but almost entirely unaddressed at the protocol level in N-of-1 methodology.**
337
338

- 339 • The Hendrickson hybrid blinded-discontinuation phase is among
340 the few protocol mechanisms that supply *PB*-versus-*BR* identifiability.
341
- 342 • Placebo run-in, open-label-placebo, and mediator-analysis
343 approaches have not been imported into N-of-1 designs.
- 344 • The gap between awareness of the problem and protocol-level
345 solutions motivates the present framework.

346 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
347 *tempor incididunt ut labore et dolore magna aliqua.*

348 The second observation is that the placebo-component separation
349 problem is widely known but largely unaddressed in N-of-1 specifically.
350 The blinded discontinuation phase used in the Hendrickson hybrid
351 design³⁴ is one of the few protocol-level mechanisms in routine N-of-1
352 use that supplies identifiability of *PB* versus *BR*. Other relevant
353 approaches, including placebo run-in designs, open-label-placebo
354 trials, and mediator-analysis applications, have been developed in
355 parallel but rarely imported into N-of-1 methodology.

356 [10]

357 **Kaptchuk's within-placebo decomposition is the closest**
358 **structural analogue; no published paper applies the corre-**
359 **sponding full-response decomposition to N-of-1 trials.**

- Kaptchuk et al. (2008) empirically decomposed the placebo response into observation, ritual, and relationship components.
- Open-label placebo trials confirm that clinically meaningful PB-only response is possible even when patients know they receive placebo.
- The BR-PB-TV framework is structurally analogous but operates at the level of the full treatment response.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The third observation is that the closest precedent in the literature is Kaptchuk’s three-component decomposition of the placebo response itself. In a randomised trial of acupuncture-style placebo for irritable bowel syndrome,⁵ decomposed the placebo response into observation, ritual, and patient-practitioner relationship components, demonstrating empirically that the placebo response is itself a sum of identifiable parts rather than a single ‘expectancy’ lump. The follow-up open-label placebo trials of^{35, 36}, and³⁷ in adolescent functional pain establish that even disclosed-placebo intervention can produce clinically meaningful PB-only response. The framework we propose is structurally analogous to Kaptchuk’s decomposition but applied at the level of the full treatment response (drug, placebo, natural history) rather than within the placebo response alone. To the best of our knowledge, no published paper proposes the corresponding decomposition for the aggregated N-of-1 setting, even though the methodological ingredients, including mixed-effects modelling, blinded-phase contrasts, and trajectory parameterisation, are all in routine use.

[11]

The decomposition gap is real, the tools to fill it are mature, and the failure modes are large enough to make it a methodological priority.

- The gap between available methodology and current practice is documented and not attributable to missing tools.
- Failure modes in chronic-condition pharmacology are consequential, not academic.
- The remainder of the paper develops, applies, and evaluates the decomposition framework.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The case we develop in the rest of this paper is therefore that the gap is real, that the methodological tools to fill it are mature, and that the failure modes of the lumped analysis are large enough in chronic-condition pharmacology to make the decomposition a methodological

402 priority rather than an academic refinement.

403 1.4 The neurobiological and epidemiological reality of the com- 404 ponents

405 [12]

406 **Each component of the decomposition reflects a real causal**
407 **mechanism documented in independent literatures.**

- 408 • *BR* corresponds to pharmacological action; *PB* to neurobiologi-
409 cally real belief-driven mechanisms; *TV* to natural-history drift.
- 410 • The components are not statistical constructs but distinct causal
411 pathways with separate literatures.
- 412 • The section reviews the evidence for each component before es-
413 tablishing the formal model.

414 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
415 *tempor incididunt ut labore et dolore magna aliqua.*

416 The decomposition is not a statistical convenience. Each component
417 reflects a real causal mechanism documented in independent litera-
418 tures, and we shall briefly recall each in turn before turning to the
419 formal model.

420 [13]

421 ***BR* is the pharmacological target of regulation; *PB* is neuro-**
422 **biologically real and phase-modulated, which is precisely the**
423 **design contrast that makes the decomposition identifiable.**

- 424 • *PB* is mediated by endogenous opioid/dopamine release, HPA-
425 axis regulation, and top-down attentional shifts.
- 426 • Wager et al. neurologic pain signature provides a clean fMRI dis-
427 sociation between pharmacological and expectancy contributions.
- 428 • Phase-dependent belief variation (open-label vs. blinded) is the
429 design contrast that supplies *PB* identifiability.

430 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
431 *tempor incididunt ut labore et dolore magna aliqua.*

432 The biological response component represents the pharmacological ac-
433 tion of the active compound: the chemical binding, downstream sig-
434 nalling, and physiological response that a regulatory authority cares
435 about when it grants approval. The placebo-belief response reflects
436 neurobiologically real mechanisms documented across pain, depres-
437 sion, Parkinson's, and PTSD research^{38–41}, mediated by belief-driven
438 release of endogenous opioids and dopamine, anticipatory regulation
439 of the hypothalamic-pituitary-adrenal axis, and top-down attentional
440 shifts in symptom perception. Direct neuroimaging evidence is in-
441 structive here:⁴² developed an fMRI-based 'neurologic pain signature'

(NPS) that tracks nociceptive input and is reduced by the opioid agonist remifentanyl, and the individual-participant meta-analysis of⁴³ found that placebo treatments produce moderate reductions in reported pain but only negligible reductions in the NPS, providing a neural dissociation between the pharmacological and expectancy contributions to a clinical pain response. The placebo response is heterogeneous across patients with biological correlates^{9,10} and is modulated by trial phase: when belief varies, as it does between open-label and blinded-discontinuation conditions, the placebo response varies, and the modulation is exactly the design contrast that makes the decomposition identifiable.

[14]

TV is the untreated counterfactual trajectory; its heterogeneity is the central object of precision-medicine inference and must be modelled, not absorbed into noise.

- *TV* sign depends on disease class: positive for acute recovery, near-zero for stable chronic, negative for progressive conditions.
- Selection-driven enrolment at peak severity guarantees positive *TV* in early trial weeks regardless of disease class.
- The PATH statement identifies treatment-effect heterogeneity modelling as mandatory for precision-medicine inference.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The time-variant natural-history component reflects whatever trajectory the participant would have followed without treatment. For acute conditions this trends positive (recovery on intrinsic timescale); for chronic stable conditions it is near zero on the population mean but with substantial participant-level variance; for chronic progressive conditions it trends negative; and for selection-driven enrolment processes it is positive in the early weeks of the trial regardless of underlying disease class, because of regression to the mean^{11,13}. The PATH statement^{14,15} highlights that this heterogeneity is itself the central object of precision-medicine inference: heterogeneity of treatment effect must be modelled, not absorbed into noise, if the precision-medicine question is to be answered.

1.5 What this paper contributes

[15]

The paper makes four contributions spanning formalisation, identifiability mapping, tractable methodology, and quantification of the inferential cost of lumping.

- Sections 2–3 formalise the decomposition; section 4 maps it to design contrasts; section 5 develops tractable analysis methods.
- Sections 6–7 provide worked examples and a Monte Carlo simulation study; section 8 applies the framework to biomarker validation.
- A companion pedagogical document provides extended worked examples for first-time readers.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

We make four contributions, each addressed in a section below. We begin in section 2 with the formal model and notation. Section 3 develops each component in turn, with biological motivation, parametric form, and identifiability conditions. Section 4 maps the design-contrast structure of the standard hybrid open-label-plus-blinded-discontinuation design onto the identifiability requirements of the decomposition. Section 5 develops the analysis-side methodology, including the linear-mixed-effects approximations and the phase-by-treatment indicator extension that retains BR-PB sensitivity at low parametric cost. Section 6 presents the worked examples, and section 7 presents the simulation study. Section 8 discusses applications to predictive-biomarker validation. The companion pedagogical document at docs/24-component-decomposition-pedagogy.md provides an extended worked-example treatment for readers approaching the material for the first time.

[16]

The first contribution is a formal specification of the three-component decomposition as a sum of Gompertz trajectories with participant-specific random effects.

- Each component follows a modified Gompertz function; the joint trajectory is their sum plus observation-level noise.
- Cross-component covariances are permitted, allowing participant-level heterogeneity to co-vary across components.
- The formal specification is the substrate on which identifiability and analysis-side arguments rest.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Our first contribution is to formalise the three-component decomposition for the N-of-1 trial setting (sections 2 and 3). Each component is modelled as a modified Gompertz function with participant-specific random effects; the joint trajectory is the sum plus observation-level noise; cross-component covariances are permitted. The formal specification is the substrate the framework builds on.

524

[17]

525

526

527

The second contribution maps identifiability of each component onto specific trial-design contrasts and characterises what is lost without them.

528

529

530

531

532

533

- The hybrid open-label-plus-blinded-discontinuation design supplies identifiability of all three components.
- Designs lacking specific phases lose identifiability of the corresponding component.
- The mapping clarifies which design features are mandatory versus optional for each inferential goal.

534

535

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

536

537

538

539

540

541

Our second contribution is to map identifiability onto trial-design contrasts (section 4). The standard hybrid open-label-plus-blinded-discontinuation design, together with phase-by-allocation contrasts, supplies identifiability of all three components. We characterise which designs do and do not supply this identifiability and what is lost in each case.

542

[18]

543

544

545

The third contribution is a phase-augmented linear-mixed-effects model that captures the BR-versus-PB attribution check at modest parametric cost.

546

547

548

549

550

551

552

- The full nonlinear DGP-matched fit is computationally fragile and degree-of-freedom-expensive at $N \in [30, 150]$.
- The phase-augmented LME uses `Dbc:phase` and `bm:Dbc:phase` to test whether drug and biomarker-by-drug effects attenuate under blinding.
- This model is the recommended default for trial-relevant sample sizes.

553

554

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

555

556

557

558

559

Our third contribution is to develop a tractable analysis methodology at trial-relevant sample sizes (section 5). The full nonlinear mixed-effects fit to mirror the DGP is computationally fragile at $N \in [30, 150]$ and consumes more degrees of freedom than trial-relevant sample sizes can spare. We propose a phase-augmented linear-mixed-effects model

560

561

562

```
Sx ~ bm + t + Dbc + phase + Dbc:phase
      + bm:Dbc + bm:Dbc:phase
      + (1 | ptID), correlation = corCAR1(...)
```

563

that captures the practically important part of the decomposition,

specifically the BR-versus-PB attribution check via $D_{bc} : phase$ and $bm : D_{bc} : phase$, at modest parametric cost.

[19]

The fourth contribution quantifies the inferential cost of lumping through worked examples and a calibrated Monte Carlo simulation study.

- Two participants with identical lumped responses but different (BR, PB, TV) compositions are indistinguishable under the one-component analysis.
- The simulation study quantifies bias, type I error, and power across analysis strategies and parameter regimes.
- The calibration uses prazosin-PTSD reference parameters from the Hendrickson et al. dataset.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Our fourth contribution is to quantify the inferential cost of failing to decompose, through worked numerical examples (section 6) and a Monte Carlo simulation study (section 7) calibrated to the prazosin-PTSD reference parameters. The worked example shows how two participants with identical lumped responses but different underlying (BR, PB, TV) compositions appear identical to a one-component analysis and qualitatively different under the decomposition. The simulation study quantifies bias, type I error, and power across the analysis strategies and across the parameter regimes most relevant to chronic-condition trials.

[20]

The decomposition is not optional for predictive-biomarker validation: it is the formal mechanism by which the pharmacological precision-medicine question can be asked at all.

- A biomarker predicting the lumped total is not, in general, a pharmacological predictor.
- The decomposition separates the biomarker-by- BR question from biomarker-by- PB and biomarker-by- TV confounding.
- The decomposition should be the default, not an optional sensitivity analysis, in trials with biomarker validation aims.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The framework's central application is to predictive-biomarker validation (section 8). A biomarker that predicts a lumped 'treatment effect' is not, in general, a pharmacological predictor. The decomposition is the formal mechanism that lets the analyst ask the precision-

605 medicine question (does the biomarker predict BR ?) rather than the
606 lumped surrogate question of whether the biomarker predicts the total
607 response, including placebo and natural history. We argue that the de-
608 composition should be the default, not an optional sensitivity analysis,
609 in any N-of-1 trial whose primary aim is biomarker validation.

610 2 Notation and the formal model

611 [21]

612 **This section fixes compact notation and states the formal**
613 **model; a fuller worked-example treatment is in the compan-**
614 **ion pedagogical document.**

- 615 • Definitions are precise enough to support the identifiability and
616 analysis-side arguments in later sections.
- 617 • The presentation is kept compact; extended derivations are de-
618 ferred to the companion document.
- 619 • The companion pedagogical document provides numerical worked
620 examples at each formal step.

621 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
622 *tempor incididunt ut labore et dolore magna aliqua.*

623 In this section we shall fix notation and state the formal model.
624 The three components are defined precisely enough to make
625 the identifiability and analysis-side arguments in later sections
626 rigorous, but we have endeavored to keep the presentation com-
627 pact. A fuller treatment, with worked numerical examples at
628 each step, is available in the companion pedagogical document at
629 docs/24-component-decomposition-pedagogy.md.

630 [22] where Y_{it} is the observed symptom score, BL_i is participant i 's
631 baseline score, the three components BR , PB , TV are non-negative
632 reductions in symptom severity attributable to the three causal mech-
633 anisms identified in section 1, and ε_{it} is observation-level noise. The
634 sign convention is that an increase in any component reduces Y ; all
635 three components are zero at $t = 0$ by construction (no exposure yet
636 to drug, no belief activation, no elapsed natural history) and grow over
637 the duration of the trial.

638 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
639 *tempor incididunt ut labore et dolore magna aliqua.*

640 For participant i at trial timepoint t :

$$Y_{it} = BL_i - [BR_{it} + PB_{it} + TV_{it}] + \varepsilon_{it},$$

641 [23]

Each component is driven by its own time-on-state variable, making the components structurally distinct even when they overlap in calendar time.

- BR_{it} accumulates with time-on-drug; PB_{it} with time-on-belief scaled by expectancy $\eta(\text{phase})$; TV_{it} with calendar time from trial entry.
- The distinct time indices are the structural source of identifiability from phase-contrast data.
- All components are zero at $t = 0$ by construction and grow over the trial duration.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Each component is a participant-specific function of its own time-on-state variable. BR_{it} depends on time-on-drug, the cumulative exposure to active medication since drug initiation. PB_{it} depends on time-on-belief, modulated by a phase-specific expectancy factor $\eta(\text{phase})$ that captures the patient's confidence that they are receiving active treatment. TV_{it} depends on time-since-trial-entry, calendar time elapsed regardless of treatment status.

[24] rescaled so that $C(0) = 0$ and $C(t) \rightarrow m_C(1 - e^{-d_C})$ as $t \rightarrow \infty$, for $C \in \{BR, PB, TV\}$. The three Gompertz parameters (m_C, d_C, r_C) are participant-specific random effects drawn from population distributions, and the participant-level heterogeneity in these distributions is the principal source of between-person response variability that the framework is designed to characterise.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Each component follows a modified Gompertz curve:

$$C(t) = m_C \exp(-d_C \exp(-r_C t)) - m_C \exp(-d_C),$$

[25]

The residual term has AR(1) structure and the three components are permitted to co-vary through cross-component covariances.

- Within-participant residuals follow $\text{Cov}(\varepsilon_{it}, \varepsilon_{i,t+h}) = \sigma_\varepsilon^2 \rho^{|h|}$.
- Within-time and across-time cross-component covariances permit participant-level heterogeneity in one component to co-vary with the others.
- This structure accommodates biologically plausible correlations, such as high- BR patients also tending toward high- PB .

Table 1: Principal psychometric and pharmacological terms used in the three-component decomposition.

Term	Meaning
BR	Biological response: pharmacological reduction in symptom score attributable to active drug.
PB	Placebo-belief response: reduction attributable to the patient’s belief that they are receiving treatment.
TV	Time-variant component: reduction (or increase) attributable to the natural-history trajectory of the disease.
$\eta(\text{phase})$	Expectancy multiplier for <i>PB</i> : 1 in open-label on-drug, ≈ 0.5 in blinded discontinuation, between 0 and 1 in open-label discontinuation.
BL_i	Participant i ’s baseline (pre-treatment) symptom score.
D_{it}	Continuous-decay drug indicator on the analysis side: 1 on-drug, $\exp(-\lambda t_{sd,it})$ off-drug.
$t_{sd,it}$	Time since drug discontinuation (off-drug only).
λ	Carryover decay rate, with half-life $t_{1/2} = \ln 2 / \lambda$.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

ε_{it} is a within-participant residual term with AR(1) serial-correlation structure $\text{Cov}(\varepsilon_{it}, \varepsilon_{i,t+h}) = \sigma_\varepsilon^2 \rho^{|h|}$. The three components are permitted to be jointly correlated through within-time and across-time cross-component covariances c_{cf1t} and c_{cfct} respectively, so that participant-level heterogeneity in any one component can co-vary with heterogeneity in the others.

The principal psychometric and pharmacological terms used below are summarised in Table 1.

3 The three components

[26]

This section presents each component with its biological motivation, Gompertz parameterisation, and identifiability conditions to establish the model as causally interpretable.

- The three components are presented in turn: *BR* (biological response), *PB* (placebo-belief), *TV* (natural history).
- Each component is grounded in independent biological literature before its parametric form is stated.
- The section’s purpose is to demonstrate that the model sum reflects causal structure, not mathematical convenience.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In this section we shall present the three components in turn, each with its biological motivation, the modified Gompertz parameterisation, and the design contrasts that identify it. The treatment is organised so that the reader can see, by the end of this section, that

707 the formal model is a sum of three causally interpretable trajectories
708 rather than a mathematical convenience.

709 3.1 Biological response (BR)

710 [27]

711 ***BR* is the pharmacological target of regulatory approval and**
712 **biomarker stratification, independent of belief or natural his-**
713 **tory.**

- 714 • *BR* arises from chemical action on the body and is zero for inert
715 pills regardless of patient expectation.
- 716 • It is the estimand that approval decisions, dose selection, and
717 biomarker-guided stratification require.
- 718 • Separating *BR* from *PB* and *TV* is the central precision-medicine
719 problem the framework addresses.

720 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
721 *tempor incididunt ut labore et dolore magna aliqua.*

722 *BR* is the physiological reduction in symptom score attributable to the
723 active medication's chemical action on the body. It is independent of
724 belief, suggestion, and natural history: a participant who somehow re-
725 ceived active drug without knowing it would still develop a measurable
726 *BR*, and a participant who took an inert sugar pill while believing
727 it was active drug would have $BR = 0$ no matter how strong their
728 expectation. Methodologically, *BR* is the quantity that drug regu-
729 lators and prescribers want: approval decisions, dose selection, and
730 biomarker-guided patient stratification are all framed around *BR*.

731 [28]

732 **PK/PD considerations produce a saturating *BR* trajectory**
733 **with a slow initial rise, a peak-growth phase, and a sharp**
734 **drop to zero on discontinuation.**

- 735 • Plasma steady state is reached within a few half-lives; pharma-
736 codynamic integration (e.g. receptor adaptation) adds additional
737 lag.
- 738 • For prazosin-PTSD, Raskind et al. (2013) calibrate the onset to
739 several weeks of sustained exposure.
- 740 • The saturating shape motivates the Gompertz over linear or sim-
741 ple exponential alternatives.

742 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
743 *tempor incididunt ut labore et dolore magna aliqua.*

744 The qualitative shape of the *BR* trajectory follows from pharmaco-
745 kinetic and pharmacodynamic considerations. Drug concentration in

plasma reaches steady state within a small multiple of the elimination half-life; brain concentration follows plasma with a short delay; the downstream pharmacodynamic response (reduction in symptom-relevant neural activity, in the case of prazosin's effect on noradrenergic tone during REM sleep) typically requires several days to weeks of sustained exposure to integrate fully⁴⁴. The result is a saturating curve: slow rise in the first week, steepest growth in the second to third week, plateau as receptor adaptation completes, and a decay back toward zero on discontinuation governed by the carryover half-life $t_{1/2} = \ln 2/\lambda$ rather than an instantaneous drop. The rate of this off-drug decay is the object of the carryover machinery used throughout this project; only in the limit of a short half-life relative to the measurement spacing does it approximate the sharp step assumed in the worked example of section 6.

[29]

The three Gompertz parameters (m_{BR}, r_{BR}, d_{BR}) capture ceiling, onset rate, and climb shape respectively, calibrated to prazosin-PTSD reference values.

- m_{BR} is the biological ceiling (maximum reduction at infinite exposure); r_{BR} the onset rate; d_{BR} the displacement shaping the early climb.
- Prazosin-PTSD calibration yields $m_{BR} \approx 11$, $r_{BR} \approx 0.42/\text{week}$, $d_{BR} \approx 5$.
- These are participant-specific random effects whose population distributions encode between-person response heterogeneity.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The modified Gompertz form $BR(t) = m_{BR} \exp(-d_{BR} \exp(-r_{BR}t)) - m_{BR} \exp(-d_{BR})$ captures this profile compactly. The asymptote $m_{BR}(1 - e^{-d_{BR}})$ is the biological ceiling, the largest reduction the drug can produce at infinite exposure; the scale parameter m_{BR} equals that ceiling up to the factor $(1 - e^{-d_{BR}})$, which for the calibrated $d_{BR} \approx 5$ exceeds 0.99. The onset rate r_{BR} controls how rapidly the participant approaches the ceiling: high r_{BR} means a fast onset (effect plateaus in one to two weeks), low r_{BR} a slow climb (six to eight weeks). The displacement d_{BR} shapes the climb: high d_{BR} gives an early-onset curve that rises sharply at first, low d_{BR} a more gradual climb. For PTSD prazosin, calibrated values from the³⁴ reference dataset are $m_{BR} \approx 11$ (eleven points of nightmare-score reduction at saturation), $r_{BR} \approx 0.42$ per week (half-maximum at roughly four to five weeks of sustained exposure), and $d_{BR} \approx 5$ (moderately steep early climb).

[30]

The Gompertz is preferred over linear and simple-exponential alternatives because it captures the asymmetric slow-start saturation profile of chronic-condition pharmacodynamics.

- Linear $BR = \beta t$ is biologically implausible past the first few weeks due to unbounded growth.
- Exponential $BR = m(1 - e^{-rt})$ lacks the slow-start phase from receptor adaptation and gene-expression changes.
- The Gompertz interpolates between both limits and its three-parameter form avoids overfitting; sensitivity analysis is reported separately in 07-gompertz-evaluation.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Two alternative families that do not capture the asymmetric saturation profile are worth noting for contrast. A linear function $BR = \beta t$ climbs without bound and is biologically implausible past a few weeks. An exponential decay $BR = m_{BR}(1 - \exp(-rt))$ is acceptable for fast-onset drugs but lacks the slow-start phase that characterises most chronic-condition responses, where receptor adaptation, feedback regulation, and gene-expression changes produce a measurable lag. The Gompertz interpolates between exponential decay (high d limit) and linear-in- t growth (low d limit), and its three-parameter form is flexible enough to fit a wide range of pharmacodynamic profiles without overfitting. The Gompertz is a deliberate choice over the two parametric forms most familiar from pharmacokinetic-pharmacodynamic modelling. The sigmoid $E_{\max}$ (Hill) model⁴⁵ describes the dependence of effect on concentration (a dose-response curve), whereas $BR(t)$ here is a time-course of accumulating effect at fixed dosing; and the indirect-response (turnover) models that do describe effect time-courses are specified through differential equations on a hidden response variable rather than in closed form. The modified Gompertz retains the slow-start, saturating shape those models produce while remaining a closed-form, three-parameter mean function that is tractable inside a mixed-effects likelihood. Doc 25 of the project documentation provides the full historical and biological context for the family choice and quantifies the sensitivity of pmsimstats inference to alternative families; that work is reported separately as 07-gompertz-evaluation.

[31]

BR is identified by the blinded-discontinuation contrast, where drug removal collapses BR to zero while PB remains intact.

- Within-participant on-drug versus off-drug contrasts are the primary identification mechanism.

- The blinded-discontinuation phase provides the cleanest contrast: BR drops, PB is preserved.
- This contrast does not require the analyst to model PB or TV explicitly to recover BR .

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The BR component is identified by within-participant on-drug versus off-drug contrasts, formalised in section 4 below. The most direct contrast is the blinded-discontinuation phase: when a participant is silently switched to placebo while still believing they are on active drug, BR decays toward zero as the drug clears while PB remains roughly intact, and the difference between on-drug and off-drug response in that window is a relatively clean estimate of BR that does not require the analyst to model PB or TV explicitly. The decay is not instantaneous: with the carryover half-life $t_{1/2} = \ln 2 / \lambda$ of the analysis-side decay term, residual BR persists into the early off-drug observations, so the contrast is clean only once the discontinuation window is long relative to $t_{1/2}$, and the drug state is best carried by the continuous-decay indicator D_{it} of Table 1 rather than a hard on/off indicator.

3.2 Placebo-belief response (PB)

[32]

PB is the strict placebo response: improvement that persists when the active drug is silently replaced, provided the patient’s belief is unchanged.

- PB is belief-dependent, not drug-dependent; it would survive a silent switch to placebo under sustained belief.
- It is distinct from BR in causal origin: BR requires chemical action; PB requires only belief.
- The definition is strict and operational, not vague or colloquial.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

PB is the reduction in symptom score that arises from the patient’s belief that they are receiving active treatment, independent of what they are actually receiving. That is, it is the placebo response in its strict sense: the part of the improvement that would persist if the active drug were silently replaced with an inert pill, provided the patient continued to believe they were on the active drug.

[33]

PB reflects neurobiologically real mechanisms triggered by

belief, not pharmacology; neuroimaging and genetic evidence confirm it is a structured biological phenomenon.

- Belief-driven endogenous opioid/dopamine release, HPA-axis modulation, and top-down attentional shifts are all documented mechanisms.
- Wager et al. neurologic pain signature dissociates pharmacological from expectancy contributions via fMRI.
- COMT genotype predicts placebo responsiveness, confirming biological structure rather than measurement artefact.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

It is worth being precise about what *PB* is and is not. *PB* is not ‘fake’ improvement, and it is not ‘all in the patient’s head’ in any dismissive sense; it is, moreover, a phenomenon clinicians knowingly exploit, as a national survey of US internists and rheumatologists found roughly half reporting regular use of placebo treatments⁴⁶. The placebo response is mediated by well-characterised neurobiological pathways^{38–41,47}. Belief-driven release of endogenous opioids and dopamine produces measurable analgesia and mood elevation; anticipatory regulation of the hypothalamic-pituitary-adrenal axis modulates cortisol and inflammatory cytokines; top-down attentional shifts change how symptoms are perceived and reported. Each of these mechanisms is a real biological response. What makes them ‘placebo’ rather than ‘drug’ is that they are triggered by belief rather than by pharmacology. The neuroimaging evidence is instructive:⁴² developed an fMRI-based ‘neurologic pain signature’ that is reduced by remifentanyl, and⁴³ showed across pooled participant-level data that placebo treatments reduce reported pain substantially while leaving the signature essentially unchanged, dissociating the pharmacological and expectancy contributions to a clinical pain response. Genetic correlates of placebo responsiveness^{9,10} further establish that placebo response is a structured biological phenomenon, not a measurement artefact.

[34]

***PB* is identifiable from *BR* only when belief and treatment allocation differ, as in the blinded-discontinuation phase.**

- The blinded-discontinuation phase is the canonical contrast: patients cannot tell whether they remain on active drug.
- Belief drops to a partial level ($\eta \approx 0.5$) because the patient’s expectation is hedged, not absent.
- Designs without a phase that decouples belief from treatment cannot separately identify *PB* from *BR*.

913 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
914 *tempor incididunt ut labore et dolore magna aliqua.*

915 For trial design, the relevant point is that PB is separately identi-
916 fiable from BR if and only if the trial includes a phase in which the
917 patient's belief about treatment differs from the actual treatment they
918 are receiving. The blinded-discontinuation phase is the canonical such
919 phase: the patient is told that, at some unannounced point during
920 the window, they may be silently switched to placebo, and they can-
921 not reliably tell whether they are still receiving active drug. Belief in
922 treatment correspondingly drops to a partial level (multiplier roughly
923 0.5) because the patient's expectation is hedged.

924 [35]

925 **The $\eta = 0.5$ blinded-phase expectancy is a modelling assump-**
926 **tion operationalising response expectancy, not a measured**
927 **quantity, and requires sensitivity analysis.**

- 928 • The construct is response expectancy in the sense of Kirsch
929 (1985); to our knowledge the residual expectancy under blinded
930 discontinuation has not been measured directly.
- 931 • Open-label placebo trials confirm clinically meaningful PB -only
932 responses even under explicit disclosure, bounding the residual
933 response away from zero.
- 934 • Sensitivity analyses varying η are mandatory before clinical con-
935 clusions depend on this single parameter.

936 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
937 *tempor incididunt ut labore et dolore magna aliqua.*

938 PB also follows a Gompertz curve, scaled by the expectancy factor:

$$PB(t \mid \text{phase}) = \eta(\text{phase}) \cdot m_{PB} \exp(-d_{PB} \exp(-r_{PB}t)) + (\text{offset to start at zero}).$$

939 Typical expectancy values are $\eta = 1.0$ in open-label on-drug, $\eta \approx 0.5$ in
940 blinded discontinuation, and η between 0 and 0.5 in open-label placebo
941 (the residual expectation a patient retains when told the previous ac-
942 tive phase was beneficial). The quantity η operationalises response
943 expectancy in the sense of⁴⁸, the patient's anticipation of benefit that
944 the expectancy literature treats as a determinant of experienced symp-
945 tom change. The choice of $\eta = 0.5$ for the blinded phase is a modelling
946 assumption rather than a calibrated quantity: to our knowledge the
947 residual expectancy a patient carries when uncertain about allocation
948 has not been measured directly as a fraction of the open-label level,
949 and we adopt 0.5 as a deliberately central placeholder. The avail-
950 able indirect evidence is nonetheless that the residual response is sub-

stantial. Open-label placebo trials^{35–37} demonstrate that a clinically meaningful response persists even when patients are explicitly told they are receiving an inert intervention, which bounds the residual belief-driven response away from zero. Because the value of η is not empirically pinned, sensitivity analyses varying it across its plausible range are mandatory before any conclusion is allowed to depend on this single number.

[36]

Kaptchuk et al. (2008) provide the structural precedent; the BR-PB-TV framework extends the same decomposition logic to the full treatment response.

- Kaptchuk decomposed the placebo response itself into observation, ritual, and relationship components empirically.
- The present framework applies analogous decomposition logic at the level of drug, placebo, and natural history.
- Both frameworks reject the ‘expectancy lump’ model in favour of identifiable additive parts.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The decomposition framework as proposed here has a structural analogue in the placebo-research literature.⁵ decomposed the placebo response itself into observation, ritual, and patient-practitioner relationship components in a randomised trial of acupuncture-style placebo for irritable bowel syndrome, demonstrating empirically that the placebo response is not a single ‘expectancy’ lump but a sum of identifiable parts. The present BR-PB-TV framework is analogous in spirit but operates at the level of the full treatment response (drug, placebo, natural history) rather than within the placebo response alone.

[37]

PB variance scales with η , not just its mean; ignoring this heteroscedastic structure miscalibrates standard errors and downstream biomarker tests.

- Phase-dependent variance follows $\text{SD}(PB \mid \text{phase}) = \eta(\text{phase}) \cdot \sigma_{PB}^{\text{baseline}}$.
- Open-label phases have high between-patient PB variance; blinded phases have low variance.
- Ignoring heteroscedasticity deflates standard errors in open-label and inflates them in blinded phases, with consequences for CI coverage and biomarker test calibration.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A subtle property of PB is that its variance also scales with η , not just its mean. When η is high, there is substantial variability across patients in how strongly the placebo response is expressed; high-suggestibility patients show large PB , low-suggestibility patients show essentially none. When η is low (blinded), the placebo response is uniformly attenuated and the patient-to-patient variance shrinks accordingly. The model therefore parameterises $SD(PB \mid \text{phase}) = \eta(\text{phase}) \cdot \sigma_{PB}^{\text{baseline}}$. Ignoring this heteroscedastic structure deflates standard errors in the open-label phase and inflates them in the blinded phase, with downstream consequences for confidence-interval coverage and for the calibration of any biomarker test built on top of the analysis.

3.2.1 Additivity, subadditivity, and the BR-PB interaction

[38]

Additivity is a modelling assumption, not biological fact; the section examines whether a $BR \times PB$ interaction term is warranted.

- The formal model writes total response as a sum; this section asks whether the BR - PB interaction is zero.
- The question has inferential consequences for biomarker validation and for what the phase-augmented analysis captures.
- The treatment is preparatory for the generalised interaction model introduced later in this section.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In this section we shall broach a question that the formal model leaves open: whether the three components contribute additively to symptom reduction, or whether there is a meaningful interaction between BR and PB that the additive specification misses.

[39]

The additive model may be wrong: published literature documents subadditivity between pharmacological and placebo-belief effects in some clinical contexts.

- Subadditivity means the combined $BR + PB$ effect is less than the sum of each component measured separately.
- Evidence from Kaptchuk et al. and Benedetti et al. shows the phenomenon is context-dependent, not universal.
- Formal test: $\partial Y / \partial (BR \cdot PB) > 0$ in a saturated mean-structure model.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The decomposition as written in section 2 treats the three components as additive contributors to symptom reduction:

$$Y_{it} = BL_i - (BR_{it} + PB_{it} + TV_{it}) + \varepsilon_{it}.$$

The additivity is a modelling assumption, not a fact about biology, and it is worth stating that the assumption may be wrong in directions the framework can absorb and in directions it cannot. Additivity is also scale-dependent in the sense emphasised by⁴⁹: components that combine non-additively on the raw symptom scale may combine additively on a transformed (for example logarithmic) scale, so an apparent $BR \times PB$ interaction can be partly an artefact of the chosen response metric rather than a biological fact. That belief and pharmacology interact rather than act in isolation is demonstrated directly by⁵⁰, who showed with neuroimaging that positive treatment expectation roughly doubled the analgesic benefit of the opioid remifentanyl while negative expectation abolished it, so the two channels are not in general separable by simple addition. More specifically, the published literature on placebo-drug interactions has accumulated evidence that the combined effect of receiving an active drug and believing one is receiving an active drug is, in some clinical contexts, less than the sum of the two components measured separately^{5,38}. The phenomenon is typically termed subadditivity: $\partial Y / \partial (BR \cdot PB) > 0$ when both components are present and their interaction term is included in a saturated mean-structure model.

[40]

Three candidate mechanisms for subadditivity are shared neural substrate, scale ceiling effects, and belief-dynamics adaptation.

- Shared pain-modulation circuits produce saturating occupancy when both BR and PB activate them simultaneously.
- Scale ceiling effects prevent registration of additional symptom reduction once $BR + PB$ saturates the measurement range.
- Belief-dynamics adaptation: unexpectedly large BR can recalibrate patient expectation downward, attenuating ongoing PB .

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Three mechanisms have been proposed in the placebo-research literature to explain subadditivity. First, a shared neural substrate: if the descending pain-modulation circuits that mediate placebo analgesia

overlap with the circuits modulated by the active drug, simultaneous activation by both mechanisms produces less marginal benefit than activation by either alone, in a saturating-occupancy sense. Second, ceiling effects on the symptom scale: a clinical scale that caps at zero (no symptoms) cannot register additional reduction once a participant has reached that cap, so any participant whose $BR + PB$ already saturates the scale will register no further reduction even if BR or PB would, in isolation, predict more. Third, belief-dynamics adaptation: the patient's expectation may adapt downward when the active drug produces a strong response that exceeds the expected effect, reducing the ongoing PB contribution; equivalently, unexpectedly large BR may attenuate PB via expectation recalibration.

[41] where $\gamma < 0$ encodes subadditivity, $\gamma = 0$ recovers the additive form, and $\gamma > 0$ encodes superadditivity, is the natural extension. The interaction coefficient γ is identifiable from the same trial-design contrasts that identify the marginal components: open-label phases supply $BR \cdot PB$ products at high η , blinded phases supply $BR \cdot PB$ products at attenuated η , and the γ coefficient is estimated from the difference between the two regimes once the marginal components are accounted for.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Whichever mechanism dominates in a given clinical context, the subadditivity question is empirical and design-dependent. Within the BR-PB-TV framework, subadditivity manifests as a non-zero coefficient on a $BR \times PB$ interaction term in the conditional-mean specification. The additive form treats this interaction as zero by construction; a generalised form

$$Y_{it} = \text{BL}_i - [BR_{it} + PB_{it} + TV_{it} + \gamma \cdot BR_{it} \cdot PB_{it}] + \varepsilon_{it},$$

[42]

Subadditivity biases biomarker predictive accuracy estimates and is partly absorbed by the phase-augmented analysis even without explicit γ estimation.

- Under subadditivity, the biomarker-by- BR association is partly cancelled, inflating apparent predictive accuracy in the additive model.
- Phase designs attenuating PB (blinded, crossover) push the effective biomarker-treatment estimate closer to the true biomarker-by- BR value.

- The phase-augmented analysis in section 5 absorbs subadditivity as a side benefit.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

For trial design, two consequences follow. First, a biomarker that predicts BR in an additive-decomposition analysis may overstate or understate predictive accuracy if the true mechanism is subadditive: under subadditivity, the biomarker-driven increase in BR is partly cancelled by the $\gamma \cdot BR \cdot PB$ interaction, and the marginal biomarker-treatment association is smaller than the biomarker-by- BR association. Second, trial designs that attenuate PB in some phases (blinded discontinuation, crossover) implicitly down-weight the subadditivity contribution to the response in those phases, which means that the effective biomarker-treatment association estimated from those phases is closer to the biomarker-by- BR association than is the open-label estimate. The phase-augmented analysis recommended in section 5 below therefore has the side benefit of partly absorbing subadditivity, even when γ is not estimated explicitly.

[43]

A forthcoming simulation study will quantify how bias in the biomarker-treatment interaction scales with γ ; the pre-registration predicts linear scaling.

- Bias in $\hat{\beta}_{bm:D}$ is expected to scale linearly with γ at small $|\gamma|$.
- Under subadditivity, bias sign is opposite to the true biomarker-by- BR effect; under superadditivity it is aligned.
- Empirical quantification is deferred to the production simulation programme in section 7.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A simulation study quantifying the magnitude of bias in the biomarker-treatment interaction estimate when the additive analysis is fitted to data generated under non-zero γ is forthcoming as part of the simulation programme described in section 7. The pre-registration is that bias scales linearly with γ at small $|\gamma|$, with the sign opposite to the true biomarker-by- BR effect under subadditivity and aligned with it under superadditivity.

[44]

Non-zero $\hat{\gamma}$ does not establish true subadditive interaction; latent-class structure with class-correlated components produces the same marginal pattern.

- If class membership correlates with both BR and PB strength,

the marginal covariance of $BR \cdot PB$ mimics a direct interaction term.

- This confounding is invisible from the lumped-analysis perspective but is identifiable in a class-aware analysis.
- Apparent subadditivity can arise from latent population heterogeneity with no direct BR - PB interaction in the class-conditional DGP.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A subtle but inferentially important point merits attention here: a non-zero $\hat{\gamma}$ recovered from an additive-extension analysis does not, on its own, establish that the underlying mechanism is subadditive interaction in the strict sense. The same marginal pattern can arise from latent population structure that the analyst is not modelling. If the trial population is a mixture of treatment-responders and non-responders (in the latent-class sense developed at length in [analysis/report/03-latent-class-mixture-application/](#)), and class membership correlates with both BR strength and PB strength (responder-class participants have both higher BR and higher PB on average), then the marginal expected value of $BR \cdot PB$ over the unobserved class label has a non-zero covariance term whose magnitude depends on the gating function $\pi(B)$. From the lumped-analysis perspective this looks exactly like a γ -weighted $BR \times PB$ interaction even though no such direct interaction was specified in the class-conditional DGP. That is, the latent-class structure is one mechanism that produces apparent subadditivity in a one-component analysis, regardless of whether the underlying biology has a true direct BR - PB interaction.

[45]

Distinguishing true subadditivity from latent-class confounding requires either a class-aware analysis or the phase contrasts the BR-PB-TV framework provides.

- The two mechanisms produce qualitatively similar marginal effects but are inferentially distinct.
- $\hat{\gamma}$ should be reported alongside a class-aware sensitivity analysis, not as standalone evidence for mechanistic subadditivity.
- The latent-class paper's `pb_class_correlation` simulation axis is the direct test; cross-paper synthesis is the cleanest disentanglement.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The two mechanisms, namely true subadditive interaction and latent-

class structure with class-correlated components, produce qualitatively similar marginal effects but are inferentially distinct. Distinguishing them requires either a class-aware analysis that explicitly models the latent population structure (as the latent-class paper develops at length), or a phase contrast that separates the mechanisms by exploiting the differential expectancy modulation of PB that the BR-PB-TV framework provides. Under the additive-extension analysis fitted to mixture-DGP data with class-correlated components, $\hat{\gamma}$ is identified up to confounding with the class structure, and the analyst should report $\hat{\gamma}$ alongside a class-aware sensitivity analysis rather than as standalone evidence for mechanistic subadditivity. The Study 1 simulation in the latent-class paper includes a `pb_class_correlation` axis that varies $\text{Cor}(BR, PB)$ across class labels; that axis is the direct simulation-side test of the latent-class generation of apparent subadditivity, and the cross-paper synthesis between the two papers is the cleanest way to disentangle the two mechanisms.

3.3 Time-variant natural-history component (TV)

[46]

TV is the untreated counterfactual trajectory; it is participant-specific, structured, and can be positive, negative, or near-zero depending on disease class.

- *TV* is the symptom change that would have occurred without treatment, integrated over the trial duration.
- It varies in sign across disease classes and across patients within a class.
- It is independent of treatment status and belief, distinguishing it structurally from PB .

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

TV is the reduction in symptom score that would have happened even if the patient had received no treatment at all. It is the participant's natural-history trajectory, integrated over the duration of the trial. For some participants *TV* is positive (symptoms naturally improve over the trial because of life events, therapy, time-in-condition, or regression to the mean); for others *TV* is negative (symptoms naturally worsen because of trauma anniversaries, accumulating stress, or progressive underlying disease).

[47]

TV is a structured participant-specific signal, not noise; its distinguishing feature from PB is independence from treatment status and belief.

- Unlike ε_{it} , TV is participant-level (slow trajectory) and requires explicit modelling, not absorption into noise.
- Off-drug baselines and follow-up windows are the primary data sources for estimating TV .
- TV is modulated by neither treatment assignment nor patient belief, whereas PB is modulated by both.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Like PB , TV is not noise. It is a structured participant-specific signal that the model can in principle estimate from data, given sufficient repeated measurements and a trial design that includes a long enough off-drug baseline or follow-up window. What distinguishes TV from PB is its origin: TV is independent of treatment status and of belief about treatment, whereas PB is modulated by both.

3.3.1 When does natural history actually exist?

[48]

The shape of TV depends on disease class, trial timescale, and selection process; none of these can be assumed and all must be recovered empirically.

- The assumption that subjects improve over time is true for some conditions and systematically wrong for others.
- Disease class, trial timescale, and enrolment selection jointly determine the sign and magnitude of TV .
- Recovering the empirical shape of TV is one of the decomposition's primary objectives.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A common simplifying assumption in trial analysis is that subjects will improve over time, with or without treatment. This is true for some conditions and problematic for others. The shape of TV depends on the underlying disease class, the trial timescale, and the selection process by which participants entered the trial. None of these can be assumed without thought, and the empirical shape of TV is one of the things the decomposition is designed to recover.

[49]

TV sign varies systematically by disease class: positive for acute recovery, near-zero for stable chronic, negative for progressive, and structured for cyclical conditions.

- Acute conditions (post-operative pain, acute back pain) resolve intrinsically; substantial TV -improvement appears in the placebo

- 1272 arm.
- 1273 • Progressive conditions (Alzheimer’s, ALS) require explicit
- 1274 negative-trajectory modelling; apparent treatment effect is
- 1275 slowing of progression.
- 1276 • Cyclical/triggered conditions (PTSD with trauma anniversaries,
- 1277 relapsing-remitting MS) require covariates beyond the Gompertz
- 1278 form.

1279 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
 1280 *tempor incididunt ut labore et dolore magna aliqua.*

1281 Acute conditions tend toward TV -positive: a sprained ankle, a post-
 1282 operative pain syndrome, an episode of acute back pain all resolve on a
 1283 timescale of days to weeks even without intervention, and a trial of an
 1284 analgesic in such a setting will see substantial TV -improvement in the
 1285 placebo arm. Chronic stable conditions (treated hypertension, con-
 1286 trolled chronic asthma, long-standing PTSD without active triggers)
 1287 tend toward TV -zero on the population mean but with substantial in-
 1288 dividual variance. Chronic progressive conditions (Alzheimer’s disease,
 1289 ALS, late-stage congestive heart failure, many cancers) tend toward
 1290 TV -negative: a trial of a neuroprotective agent in early Alzheimer’s
 1291 must explicitly model the negative natural-history trajectory, since the
 1292 apparent treatment effect is the slowing of progression, not absolute
 1293 improvement. Cyclical and triggered conditions (PTSD with trauma
 1294 anniversaries, seasonal-pattern major depression, relapsing-remitting
 1295 multiple sclerosis) show structured TV that is neither monotone nor
 1296 stationary, and may require explicit covariates beyond the Gompertz
 1297 form.

1298 [50]

1299 **Selection-induced TV is real and rarely zero even for stable**
 1300 **chronic conditions, because enrolment coincides with peak**
 1301 **severity triggering help-seeking.**

- 1302 • Regression to the mean alone contributes several points of appar-
- 1303 ent improvement in early trial weeks.
- 1304 • Hrobjartsson et al. systematic reviews across 200+ trials show
- 1305 much of conventional placebo response is indistinguishable from
- 1306 natural-history drift.
- 1307 • Kirsch (2008) and Locher (2015) demonstrate the clinical conse-
- 1308 quences in antidepressant and late-life depression trials.

1309 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
 1310 *tempor incididunt ut labore et dolore magna aliqua.*

1311 Selection-induced TV is real and rarely zero, even when the underlying
 1312 condition is stable. Patients enrol in trials at moments of clinical en-
 1313 gagement, which often coincide with recent worsening that triggered

help-seeking. Regression to the mean alone produces several points of apparent TV -improvement in the early weeks of a trial, even in a population whose underlying condition is stable. The systematic-review evidence of^{11, 12}, and¹³, the last covering 202 trials across 60 clinical conditions, quantifies this directly: in trials with both placebo and no-treatment arms, much of what is conventionally attributed to the placebo response was indistinguishable from regression to the mean and natural-history drift. The depression literature provides the most consequential modern example:⁶, analysing FDA-submitted antidepressant trials, showed that mean drug-placebo differences fell below clinical-significance thresholds except at the most severe baseline-severity stratum, with the apparent ‘efficacy’ largely explained by placebo response and natural fluctuation;⁷ explicitly invoked regression to the mean to explain the severity-by-treatment interaction in late-life depression.

[51]

Trial participation itself induces TV -improvement through Hawthorne, structured-visit, and diary-completion effects, independent of pharmacology or placebo.

- Hawthorne, structured-visit, and diary-completion effects each contribute symptom improvements attributable to participation, not treatment.
- These contributions are absorbed into TV in the decomposition, not into BR or PB .
- Ignoring these participation-induced effects inflates the apparent pharmacological and placebo components.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Trial participation itself also induces TV : the Hawthorne effect, the structured-visit effect, and the diary-completion effect appear in the data as TV -improvement that is independent of any pharmacological or placebo response.

3.3.2 Can we assume no TV effect?

[52]

Assuming $TV = 0$ is rarely safe in chronic psychiatric or pain trials; the conservative default is to estimate TV from the data.

- Zero- TV assumption is valid only when all four conditions hold: stable disease, short trial, non-severity-peaked recruitment, no ancillary care.

- For most chronic conditions at least one condition fails; the assumption should be tested, not asserted.
- Statistical efficiency lost by estimating TV is smaller than the bias introduced by assuming it zero.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

But how safe is it to assume TV is zero? The short answer is: rarely safe, and never safe without empirical support. The longer answer is that TV may be assumed zero only if the disease is in a stable chronic phase, the trial is short relative to the disease timescale, the recruitment process does not favour participants at peak severity, and the trial protocol does not introduce ancillary care that itself improves symptoms. For most chronic trials in psychiatric or pain conditions, at least one of these conditions fails. The conservative default is to estimate TV from the data, even at the cost of some statistical efficiency, rather than assume it is zero.

[53]

The decomposition makes the zero- TV assumption testable and shows that lumping TV into residuals produces auto-correlated, heteroscedastic residuals that bias BR and PB standard errors.

- $\hat{m}_{TV}$ with a variance estimate from the model provides an empirical test of the zero- TV assumption.
- Participant-level TV is structured and correlated with the on/off-drug schedule; it is not independent noise.
- Lumping TV into ε produces autocorrelated within-participant residuals and heteroscedastic cross-participant residuals.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The decomposition makes this assumption testable. A model that includes a TV component will return $\hat{m}_{TV}$ with a variance estimate; if both are small relative to $\hat{m}_{BR}$ and $\hat{m}_{PB}$, the analyst has empirical support for the simplifying assumption. If $\hat{m}_{TV}$ is meaningfully large, the assumption was wrong, and the trial needs the TV component to produce unbiased inference about the others. Setting TV to zero and lumping it into ε is worse than estimating it: the participant-level TV is structured (each participant has their own slow trajectory) and not independent of the on-drug-versus-off-drug schedule, so lumping it into noise produces residuals that are autocorrelated within participants and heteroscedastic across them. The resulting standard errors on the BR and PB estimates are wrong in both directions: too small for participants whose TV is small, too large for participants whose TV

is large. Estimating TV explicitly fixes both problems simultaneously.

3.3.3 Why a Gompertz, not a linear time covariate?

[54]

A linear week covariate is insufficient for TV because it enforces linearity, ignores participant heterogeneity, and cannot distinguish TV from BR without the explicit decomposition.

- A week covariate constrains each additional week to contribute a fixed natural-history change, misrepresenting curved or saturating disease trajectories.
- A single coefficient estimates a population-mean trajectory, ignoring participant-level heterogeneity in natural history.
- Without the decomposition and design contrasts, week absorbs both BR and TV , making them inseparable.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A natural objection is that a model with week as a covariate already controls for time. But how adequate is such a control? Three reasons suggest it is not enough. First, a week covariate enforces a linear trajectory, constraining each additional week to contribute a fixed amount of natural-history change; real disease trajectories curve, plateau, or accelerate, and the Gompertz form captures the S-shaped or saturating patterns linear time cannot. Second, a single week coefficient estimates a population-mean trajectory that ignores the participant-level heterogeneity in natural history. Third, BR and TV both evolve smoothly over the same trial timescale, and a week covariate cannot tell them apart from observational data alone; the design contrasts discussed in section 4 are what supply identifiability, and without the explicit decomposition the week coefficient absorbs both.

4 Identifiability and trial-design contrasts

[55]

Identifiability is supplied by trial design; each phase of the hybrid design contributes a different combination of components to the observed trajectory.

- The decomposition is mathematically specified but requires trial-design contrasts to be statistically identified.
- This section characterises the hybrid open-label-plus-blinded-discontinuation design as the primary identifiability vehicle.

- Each of the three phases supplies a distinct mixture of BR , PB , and TV .

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The mathematical decomposition above is meaningless unless the data contain enough information to pin down each component separately. Trial design is what supplies that information. In this section we shall characterise the hybrid open-label-plus-blinded-discontinuation design, used as the worked example throughout this paper, and show how each of its three phases contributes a different combination of components to the observed trajectory.

Phase	Drug?	Belief?	Components present
Open-label on drug	Yes	Full	$BR + PB(\eta = 1) + TV$
Blinded discontinuation	Hidden	Hedged	$BR \cdot \mathbf{1}_{\text{on drug}} + PB(\eta = 0.5) + TV$
Open-label crossover	Either	Knows	$BR \cdot \mathbf{1}_{\text{on drug}} + PB(\eta = 1) + TV$

[56]

Three phase-by-allocation contrasts isolate PB , BR , and TV respectively; the LME model combines them but the inference is driven by design.

- Open-label versus blinded contrast yields approximately $0.5 \cdot PB$, isolating the belief component conditional on a known blinded-phase η .
- Blinded on-drug versus blinded placebo yields BR once drug carryover has elapsed (since PB and TV match across arms within the blinded phase).
- Off-drug phases yield TV plus residual off-drug PB , with no active BR contribution beyond carryover.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Three phase-by-allocation contrasts then identify the three components. The open-label versus blinded contrast (within drug) yields $PB(\eta = 1) - PB(\eta = 0.5) \approx 0.5 \cdot PB$, which isolates the belief component, but only conditional on a known blinded-phase expectancy η : from a two-phase design just the product $\eta \cdot m_{PB}$ is identified, and separating η from m_{PB} requires the open-label placebo phase. The blinded on-drug versus blinded placebo contrast yields BR once drug carryover has elapsed, since PB and TV match across allocation

arms within the blinded phase. All phases off drug yield TV plus any residual off-drug PB , with no active BR contribution beyond carry-over and PB at the off-drug expectation level. The $\mathbf{1}_{\text{on drug}}$ indicator in the phase table above is shorthand for the continuous-decay drug state D_{it} defined in Table 1.

[57]

The LME model formally combines all three contrasts, but inferential precision is determined by the design, not model specification.

- Each contrast targets a specific component; the model integrates all three into a coherent parameter set.
- Model specification cannot compensate for missing design contrasts; the identifying information must exist in the data.
- The analyst should evaluate any proposed trial design by verifying each component has at least one identifying contrast.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Each contrast targets a different component, and the analyst combines them through the linear-mixed-effects model fitted to the full trajectory. The model is the formal mechanism, but the inference is driven by the design contrasts.

[58]

Designs lacking specific phases lose identifiability of the corresponding component; the decomposition machinery is only as useful as the design allows.

- A pure open-label design cannot separate BR from PB because both are present at full strength throughout.
- A pure blinded design cannot estimate $\eta = 1$ PB because no phase achieves full-strength belief.
- Trial designers should verify each component has at least one phase in which it is uniquely present before committing to the design.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A trial that lacks any of these phases loses the corresponding identifiability. A pure open-label design (no blinded discontinuation) cannot separate BR from PB , because both are present at full strength throughout. A pure blinded design (no open-label phase) cannot estimate $\eta = 1$ PB , because no phase has belief at full strength. The decomposition machinery is useful only to the extent that the design

supplies identifiability, and trial designers should evaluate any proposed design by checking that each component has at least one phase or contrast in which it is uniquely present.

[59]

Design-contrast identifiability is necessary but not sufficient: the three Gompertz trajectories are collinear, so joint estimation is credible only under specific design and sample-size conditions.

- The components share an onset shape and overlap in calendar time, so their parameters can be near-collinear even when the identifying contrasts are present.
- Joint recovery requires a phase in which each component varies while the others are held fixed, and an off-drug window long enough for TV to separate from a decaying BR .
- The recoverable region of the design-and-sample-size space is mapped by the pre-registered identifiability diagnostics, not asserted.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Identifiability in the design-contrast sense above is necessary but not sufficient for stable estimation at finite sample sizes. The three Gompertz trajectories share an onset shape and overlap in calendar time, so their parameters can be collinear even when the contrasts that identify them in principle are present. Two conditions make the joint estimation credible: the design must supply at least one phase in which each component varies while the others are held fixed (the blinded-placebo contrast for BR , the open-label-versus-blinded contrast for PB , and an off-drug window for TV), and that off-drug window must be long enough relative to the carryover half-life for TV to separate from a still-decaying BR . Where these conditions are weak the components are jointly identified only up to a near-collinear direction, and the symptom of that is the rank-deficiency the pilot of section 7 encountered when the full phase-augmented formula was fitted at $N = 35$. We therefore do not assert that the three components are always recoverable. The pre-registered identifiability diagnostics of section 7 (convergence rates of the component-matched fit, variance-inflation factors on the component-specific Gompertz parameters, and the sensitivity of the BR estimate to the assumed $\eta(\text{phase})$ multipliers) are the means by which the recoverable region of the design-and-sample-size space is to be mapped, and the framework's participant-level claims should be read as conditional on falling inside that region.

5 Analysis-side methodology

[60]

Four constraints make a component-matched analysis impractical at trial-relevant sample sizes; a phase-augmented LME captures the important part of the gap at minimal cost.

- The question of whether to mirror the DGP in the analysis model is reasonable; the section addresses it directly.
- The four constraints are identifiability, degrees of freedom, convergence fragility, and marginal estimand robustness.
- A phase-by-treatment indicator is the recommended modest extension that captures BR-versus-PB sensitivity without the full parametric cost.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In this section we shall take up the practical question of what analysis model to use. A reasonable reaction to the formal model in section 2 is the following: if we are simulating data with three components plus noise, why is the analysis model written as $Sx \sim \text{bm} + \text{t} + \text{Dbc}$ + $\text{bm}:\text{Dbc}$ with a single random intercept and AR(1) residual correlation? Should we not include three parametric Gompertz components in the analysis, mirroring the DGP, so that the analyst sees the same structure that generated the data? The question is reasonable. Four constraints make a component-matched analysis impractical for most aggregated N-of-1 trial sizes, and a modest extension (a phase-by-treatment indicator) captures the practically important part of the gap at minimal cost.

5.1 Four reasons not to match the DGP

[61]

The first constraint is identifiability: forcing a three-component model on a design with insufficient contrasts produces unidentified parameters, not component estimates.

- A pure open-label trial cannot distinguish *BR* from *PB*; both rise from zero together on treatment initiation.
- Fitting three components to a one-contrast design produces implausibly tight standard errors that misrepresent inferential precision.
- The analyst must audit whether the design supplies the necessary contrasts before specifying a component-matched model.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The first reason is identifiability. The trial design must supply the contrasts that identify each component, and many designs supply only part of it. The decomposition is identified by phase contrasts (open-label versus blinded for *PB*, on-drug versus blinded-placebo for *BR*, off-drug timepoints for *TV*). A pure open-label trial has none of these contrasts; in such a trial *BR* and *PB* both rise from zero on treatment initiation and are statistically indistinguishable from each other, regardless of how the analyst writes the model. Forcing an analysis to fit three components when the design supplies contrasts for only one or two of them does not recover the missing components; it produces unidentified or weakly identified parameter estimates with implausibly tight standard errors that misrepresent the analyst's actual inferential precision.

[62]

The second constraint is degrees of freedom: a component-matched analysis adds approximately a dozen non-linear parameters, inflating moderation estimand variance at $N \in [30, 150]$.

- Gompertz *BR* (3), *PB* (3 + expectancy multiplier), and *TV* (3) each require participant-specific random effects on the maxima.
- At trial-relevant N , the variance of $\hat{\beta}_{bm:D}$ is dominated by participant and timepoint counts.
- Additional parametric structure typically inflates moderation estimand variance more than it reduces bias.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The second reason is degrees of freedom. Each component term costs degrees of freedom that come out of the moderation estimand. The biomarker-moderation question (the core scientific aim of the predictive-biomarker validation programme) is asked through the *bm:D* term. A component-matched analysis adds a Gompertz *BR* (three parameters), a Gompertz *PB* (three parameters plus a phase-specific expectancy multiplier), and a Gompertz *TV* (three parameters), each with participant-specific random effects on the maxima. That is on the order of a dozen new parameters, half of them on a non-linear scale. At trial sample sizes of $N \in [30, 150]$, the variance of $\hat{\beta}_{bm:D}$ is dominated by participant and timepoint counts, and adding parametric structure to absorb residual confounding usually inflates the moderation estimand's variance by more than it reduces bias.

[63]

The third constraint is convergence fragility: component-

matched nonlinear mixed models converge below 90% at trial-relevant N , contaminating aggregated analyses.

- Simultaneous Gompertz BR , PB , and TV fitting requires `nlme::nlme`, Bayesian hierarchical, or structural-equation approximations.
- Each is more sensitive to starting values, prone to local optima, and likely to fail on individual trial datasets than the LME analogue.
- Non-converging cells at below-90% convergence rates contaminate power and bias estimates in simulation studies.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The third reason is convergence fragility. Component-matched analyses are non-linear in the parameters and their convergence is problematic in practice. Fitting Gompertz BR , PB , and TV simultaneously requires a non-linear mixed model (`nlme::nlme` with self-starting Gompertz forms), a Bayesian hierarchical fit, or a structural-equation approximation. Each is more sensitive to starting values, more prone to local optima, and more likely to fail to converge on a single trial dataset than the linear-mixed analogue. At trial-relevant N , the convergence rate of the component-matched fit is typically below 90%, and the non-converging cells contaminate any aggregated analysis.

[64]

The fourth constraint is estimand robustness: the `bm:Dbc` marginal interaction is the right estimand for the biomarker-validation question regardless of the BR-PB split.

- Regulators want to know whether the biomarker predicts treatment response under trial conditions, not a mechanistically decomposed prediction.
- The marginal `bm:Dbc` interaction captures this question whether the moderation is biological or expectation-mediated.
- Component decomposition adds mechanistic precision but is not required for the core biomarker-validation decision.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The fourth reason is that the moderation estimand is sufficiently robust to the BR-versus-PB split for the most common scientific question. For predictive-biomarker validation, the regulator wants to know whether the biomarker predicts treatment response under the trial conditions, not a clean pharmacological-versus-placebo decomposition of that prediction. The simple `bm:Dbc` interaction is the right marginal

estimand for the validation question whether the moderation is biological or expectation-mediated.

[65]

The inferential target is stated as an estimand: the population biomarker-by-treatment interaction with explicit intercurrent-event handling, while the pharmacological-component target is more ambitious and inherits the heterogeneity-identifiability caveat.

- The $bm:D_{bc}$ estimand is defined by its population, treatment contrast, endpoint, summary, and intercurrent-event strategy in the ICH E9(R1) sense.
- Dropout is handled by a treatment-policy strategy and residual carryover by the modelled exponential decay.
- Recovering a participant-level pharmacological interaction inherits Senn’s caveat that true heterogeneity is separable from within-person variation only under repeated phase contrasts.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

It is worth stating the inferential target of the analysis in the vocabulary of the estimands framework⁵¹, because the biomarker-validation question is easy to pose ambiguously. The primary estimand is the population-level biomarker-by-treatment interaction, the $bm:D_{bc}$ coefficient: its population is the trial-eligible patient population, its treatment contrast is on-drug versus off-drug exposure summarised through the continuous indicator D_{bc} , its endpoint is the symptom trajectory, its population-level summary is the interaction coefficient, and its intercurrent events are handled by a treatment-policy strategy for dropout and by the modelled exponential decay for residual carryover. The pharmacological-component target β_{bm}^{BR} is a more ambitious estimand: it is the biomarker-by- BR interaction that would obtain with the belief-driven channel held at its off-drug level, identified here through the phase contrasts of section 4 rather than by assumption. We flag explicitly that recovering a participant-level interaction of this kind inherits the difficulty Senn has emphasised for individual treatment effects^{49,52}: a genuine participant-by-component signal is separable from participant-by-period and within-person random variation only when the design supplies repeated phase contrasts, and the framework should be read as estimating these quantities under that condition rather than as dissolving the well-known problem of identifying individual effects from finite within-person data.

5.2 A small extension that usually pays off: phase indicators

[66]

The phase-augmented LME uses `Dbc:phase` and `bm:Dbc:phase` to test whether drug and biomarker-by-drug effects attenuate under blinding.

- `phase` absorbs systematic differences in *PB* expectancy across trial phases.
- `Dbc:phase` detects whether the drug effect shrinks under blinding, a signature of *PB*-mediation.
- `bm:Dbc:phase` detects whether the biomarker-by-treatment interaction shrinks under blinding, a signature that it is partly biomarker-by-*PB*.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The smallest useful step toward component-aware analysis, without paying the costs above, is to add a phase indicator and its interaction with the drug indicator. Schematically,

```
Sx ~ bm + t + Dbc + phase + Dbc:phase
      + bm:Dbc + bm:Dbc:phase
      + (1 | ptID), correlation = corCAR1(...)
```

Here `phase` is a categorical covariate (open-label, blinded, post-crossover) that absorbs the systematic differences in *PB* expectancy across phases. `Dbc:phase` tests whether the apparent drug effect shrinks under blinding, a signature that the open-label drug effect was partly *PB*-mediated. `bm:Dbc:phase` tests whether the moderation itself shrinks under blinding, a signature that the biomarker-by-treatment interaction was partly biomarker-by-*PB* rather than biomarker-by-*BR*.

[67]

The phase-augmented extension adds at most four parameters yet captures the practically relevant belief-attenuation diagnostic without requiring the Gompertz parametric form.

- Three to four parameters is well within budget at any realistic trial sample size.
- The extension does not separately estimate *BR* and *PB* components but identifies the change in apparent drug effect when belief is hedged.
- For most analyses this is the practically relevant question and is answerable without committing to the Gompertz form.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

This extension adds three to four parameters at most, well within budget at any realistic *N*. It does not separately estimate *BR* and *PB* as

components, but it does identify the change in apparent drug effect that occurs when belief is hedged. For most analyses, this is the practically relevant question, and it can be answered without committing to the Gompertz parametric form.

5.3 When to fit the full decomposition

[68]

The full decomposition merits the effort only when N is in the hundreds, the design is fully phase-balanced, and the scientific question requires the BR-PB mechanistic split.

- Feasibility conditions: several-hundred N , dense timepoints in every phase, mechanistic BR-PB split required, and informative Bayesian priors available.
- In that regime, `lcmm`, `nlme::nlme` with self-starting Gompertz forms, or `brms` can return participant-specific component estimates.
- Outside that regime the phase-augmented LME is the better tool.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A component-matched analysis merits the effort when sample size is in the several hundreds (a phase-3 trial fielded as aggregated N-of-1, or a meta-analysis pooling multiple smaller trials); the trial design provides timepoints in every phase combination, with enough density per cell to constrain the Gompertz curves; the scientific question requires the BR-PB split (a pharmacology paper claiming a specific mechanism, where ruling out a placebo contribution is the central point); and a Bayesian hierarchical fit with informative priors on the Gompertz parameters is feasible. In that regime, `lcmm`, `nlme::nlme` with self-starting Gompertz forms, or a `brms` hierarchical formulation can return participant-specific *BR*, *PB*, and *TV* component estimates that mirror the DGP. Outside that regime, the linear-mixed analysis with phase indicators is the better tool.

[69]

The DGP-analysis asymmetry is deliberate: simulation uses the rich DGP to characterise how a simpler analysis behaves; the analysis need not mirror the DGP to produce useful inference.

- Simulation controls the DGP and can specify any structure; inference must pay the identifiability cost of every parameter.
- The DGP is a characterisation tool; the analysis is a finite-data inference tool; they have different optimality criteria.

1795 • DGP-matching is not required and is often counterproductive at
1796 trial-relevant sample sizes.

1797 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
1798 *tempor incididunt ut labore et dolore magna aliqua.*

1799 The asymmetry between DGP and analysis is not a flaw of the frame-
1800 work; it is a deliberate consequence of the asymmetry between simu-
1801 lation (where the analyst controls the DGP and can specify any struc-
1802 ture, however elaborate) and inference (where the analyst must fit a
1803 model to a finite dataset and pay the identifiability cost of every addi-
1804 tional parameter). Simulation studies use the rich DGP to characterise
1805 how a simpler analysis behaves under realistic data generation; they
1806 do not require the analysis to match the DGP in order to produce
1807 useful inference.

1808 [70]

1809 **The exponential-decay carryover indicator is a deliberately**
1810 **parsimonious special case of richer carryover models, and**
1811 **the simulation should benchmark it against a distributed-lag**
1812 **alternative.**

- 1813 • The continuous indicator D_{bc} decays as $\exp(-\lambda t_{sd})$, a one-
1814 parameter form tied to a pharmacokinetic half-life.
- 1815 • Distributed-lag models, the average-period-treatment-effect esti-
1816 mand, and G-estimation relax the single-rate assumption at ad-
1817 ditional parametric cost.
- 1818 • The exponential restriction is an assumption, not a finding, and
1819 the pre-registered simulation should include a distributed-lag arm
1820 to quantify its cost.

1821 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
1822 *tempor incididunt ut labore et dolore magna aliqua.*

1823 A word on the carryover parameterisation is in order, since it is the one
1824 place where the analysis model makes a strong functional-form com-
1825 mitment. The continuous drug indicator D_{bc} decays as $\exp(-\lambda t_{sd})$ off
1826 drug, a one-parameter exponential governed by the carryover half-life.
1827 This is deliberately parsimonious and is best understood as a con-
1828 strained special case of the more flexible carryover models now avail-
1829 able for repeated-measures and N-of-1 data. Liao et al.³¹ model carry-
1830 over through a Bayesian distributed-lag structure with autocorrelated
1831 errors, relaxing the single-rate exponential assumption at the cost of
1832 additional parameters; Daza⁵³ formalises an average-period treatment-
1833 effect estimand built to accommodate autocorrelation, trend, and slow
1834 carryover jointly; and G{”a}rtner et al.⁵⁴ benchmark linear models
1835 against G-estimation and Bayesian-network alternatives specifically
1836 under carryover. We adopt the exponential form because it ties di-

rectly to a pharmacokinetic half-life and conserves degrees of freedom for the moderation estimand at trial-relevant sample sizes, but the choice is an assumption rather than a finding. The pre-registered simulation programme of section 7 should accordingly include a distributed-lag analysis arm, so that the cost of the exponential-decay restriction is quantified against the richer alternative rather than asserted.

6 Worked example: information loss without decomposition

6.1 What the lumped analysis estimates

[71]

The lumped treatment-effect estimate has a probability limit equal to the pharmacological component plus belief and natural-history contributions, so its bias does not vanish as N grows.

- Writing the reduction as $BR + PB + TV - \varepsilon$ and fitting one drug indicator gives $\text{plim } \hat{\beta}_D = m_{BR} + b_{PB} + b_{TV}$.
- b_{PB} and b_{TV} are the belief and natural-history terms that load onto the drug indicator under the chosen contrast; both are generally non-negative for an improving outcome.
- Because the displacement is in the probability limit, a larger one-component trial converges on the same biased value: the failure is identification, not precision.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Before the numerical illustration it is worth establishing algebraically what the one-component analysis actually estimates, because the qualitative failure modes of section 1 are consequences of an identity rather than empirical tendencies that require simulation to demonstrate. Write the within-participant symptom reduction as $R_{it} = \text{BL}_i - Y_{it} = BR_{it} + PB_{it} + TV_{it} - \varepsilon_{it}$. The one-component analysis fits a single drug indicator D_{it} ,

$$R_{it} = \alpha + \beta_D D_{it} + g(t) + u_i + e_{it},$$

and reports $\hat{\beta}_D$ as ‘the treatment effect’. Taking the on-drug minus off-drug contrast that β_D targets, and using that the natural-history component is by construction independent of treatment status so that its on/off difference cancels in a within-participant comparison, the probability limit of the estimator is

$$\text{plim } \hat{\beta}_D = m_{BR} + b_{PB} + b_{TV},$$

where $b_{PB} = E[PB \mid \text{on}] - E[PB \mid \text{off}]$ is the belief differential between the contrasted phases and b_{TV} collects any natural-history contribution that fails to cancel under the contrast actually used; b_{TV} is zero for a clean on/off within-participant contrast and strictly positive for a baseline-anchored pre-post analysis, in which case $b_{PB} = E[PB]$ and $b_{TV} = E[TV]$. Both terms are generally non-negative for an improving outcome, so the lumped estimate over-states the pharmacological component m_{BR} . The displacement is a property of the probability limit and is unchanged as $N \rightarrow \infty$: a larger one-component trial converges on the same biased value. This is the precise sense in which the placebo-attribution failure of section 1 is a problem of identification, not of precision.

[72]

The biomarker slope from a lumped analysis decomposes additively into component-specific slopes, so a biomarker correlated with PB or TV registers as a predictor even with no pharmacological association.

- By linearity of covariance, $\beta_{bm}^{\text{lumped}} = \beta_{bm}^{BR} + w_{PB}\beta_{bm}^{PB} + w_{TV}\beta_{bm}^{TV}$.
- A non-zero β_{bm}^{PB} or β_{bm}^{TV} produces a non-zero lumped slope even when the pharmacological slope β_{bm}^{BR} is exactly zero.
- The result is an algebraic identity under the additive decomposition, not a sample-size-dependent tendency.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The same algebra settles the biomarker-validation failure. The precision-medicine analysis regresses a participant-level response summary on a baseline biomarker B_i . If that summary is the lumped per-participant effect, which estimates $m_{BR,i} + w_{PB}m_{PB,i} + w_{TV}m_{TV,i}$ for design-determined loadings $w_{PB}, w_{TV} \geq 0$, then by the linearity of covariance the fitted biomarker slope has probability limit

$$\beta_{bm}^{\text{lumped}} = \frac{\text{Cov}(m_{BR} + w_{PB}m_{PB} + w_{TV}m_{TV}, B)}{\text{Var}(B)} = \beta_{bm}^{BR} + w_{PB}\beta_{bm}^{PB} + w_{TV}\beta_{bm}^{TV},$$

where $\beta_{bm}^C = \text{Cov}(m_C, B) / \text{Var}(B)$ is the component-specific slope. A biomarker correlated with placebo responsiveness or with the natural-history trajectory, so that β_{bm}^{PB} or β_{bm}^{TV} is non-zero, therefore registers a non-zero lumped slope even when its pharmacological association β_{bm}^{BR} is exactly zero. This is the biomarker-validation failure of section

1, and it too is an identity under the additive decomposition rather than an empirical tendency: no sample size separates the lumped slope from a genuine pharmacological one, because the two coincide in the estimand whenever $w_{PB}\beta_{bm}^{PB} + w_{TV}\beta_{bm}^{TV} \neq 0$.

[73]

The bias displaces the estimand rather than inflating variance, so a significance test of the lumped coefficient cannot detect it; the decomposition is what locates it.

- The lumped coefficient is displaced from m_{BR} by $b_{PB} + b_{TV}$ but remains a well-estimated quantity with a valid Wald test of its own nullity.
- Detecting the bias requires comparing $\hat{\beta}_D$ to m_{BR} , a comparison only the decomposition or an isolating phase contrast supplies.
- This is why a rejection-rate summary is the wrong instrument, and accounts for the flat rejection rates across the placebo-strength axis in the section 7.1 pilot.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Two consequences for how the framework should be evaluated follow. First, the quantity that carries the bias is the location of the estimand, $\text{plim } \hat{\beta}_D$, not its sampling variability: the lumped coefficient remains a well-estimated number with a valid Wald test of its own nullity, and that test rejects at the rate dictated by the displaced value. A significance test of the lumped coefficient therefore cannot, in principle, detect the bias, because the bias does not move the coefficient toward zero. Detecting it requires comparing $\hat{\beta}_D$ against m_{BR} , which is exactly the comparison the decomposition, or a phase contrast that isolates m_{BR} , makes available and the lumped analysis does not. Second, and as a corollary for the simulation evidence, a rejection-rate or power summary is the wrong instrument for this bias: the right target is the coefficient itself and its offset from m_{BR} . This is why the pilot of section 7.1 found the biomarker-by-treatment rejection rate essentially flat across the placebo-strength axis even though that axis is precisely what drives the predicted bias; the prediction lives in the mean of $\hat{\beta}_{bm:D_{bc}}$, and a production run that means to test it must read the bias off the coefficient, not off the test.

6.2 A numerical illustration

[74]

The worked example uses two participants with identical lumped responses but opposite BR-PB profiles to show that the one-component analysis cannot distinguish them.

- Participant A has high PB ($m_{PB} = 6$) and moderate BR ($m_{BR} = 4$); Participant B has low PB ($m_{PB} = 1$) and high BR ($m_{BR} = 8$).
- Both participants reach approximately 10–11 points of total improvement by week 8, making them indistinguishable to a one-component analysis.
- The blinded-discontinuation phase creates the diagnostic contrast that exposes the underlying difference.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

To illustrate the inferential cost of failing to decompose, we shall work through a simple numerical example. Consider two participants, A and B, enrolled in a sixteen-week N-of-1 trial of prazosin for PTSD. Both are otherwise similar: same age, same baseline severity, same trial design. They differ in two latent traits that the trial does not measure directly. Participant A is a high placebo responder and a moderate pharmacological responder, with true component values $m_{BR}^A = 4$, $m_{PB}^A = 6$, $m_{TV}^A = 1$. Participant B is a low placebo responder and a strong pharmacological responder, with $m_{BR}^B = 8$, $m_{PB}^B = 1$, $m_{TV}^B = 1$.

[75]

Open-label totals are clinically indistinguishable; the one-component analysis estimates a population drug effect of approximately 10 points with no differentiation between A and B.

- Both participants show approximately 10–11 points of improvement and present as clinical successes.
- A t-test or simple LME on open-label change cannot detect the heterogeneity in BR versus PB composition.
- The population drug-effect estimate conflates both pharmacological and placebo contributions.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In the open-label phase (weeks 1 to 8), both BR and PB run at full strength and TV accumulates. By week 8 both participants have reached close to their saturation values:

	BR	$PB(\eta = 1)$	TV	Total improvement
Participant A	4.0	6.0	1.0	11.0
Participant B	8.0	1.0	1.0	10.0

A clinician examining only the totals would conclude that the two

1986 participants are responding similarly: both improved by about ten
1987 points, both look like clinical successes. A simple t-test on the open-
1988 label change would estimate a population drug effect of about ten
1989 points and would not distinguish the two participants.

1990 [76]

1991 **The blinded-discontinuation phase reveals the BR-PB het-**
1992 **erogeneity: Participant A retains 4 points while B collapses**
1993 **to 1.5, exposing the diagnostic contrast.**

- 1994 • Participant A's high PB at $\eta = 0.5$ (3 points) sustains improve-
1995 ment even without active drug.
- 1996 • Participant B's low PB (0.5 points) and absent BR produce near-
1997 complete return to baseline.
- 1998 • The 4.0 versus 1.5 contrast is exactly what the decomposition
1999 uses to separate BR from PB .

2000 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
2001 *tempor incididunt ut labore et dolore magna aliqua.*

2002 In the blinded discontinuation phase (weeks 9 to 12), both partici-
2003 pants are silently switched to placebo at the start of week 9. The
2004 expectancy multiplier drops from 1.0 to 0.5. For clarity of exposition
2005 we take the carryover half-life to be short relative to the four-week
2006 blinded window, so that BR has effectively cleared by the end of week
2007 12; the more realistic case of a slowly decaying BR , modelled by the
2008 carryover machinery of section 3.1, attenuates but does not eliminate
2009 the contrast below. By the end of week 12, the components are:

	BR	$PB(\eta = 0.5)$	TV	Total improvement
Participant A	0.0	3.0	1.0	4.0
Participant B	0.0	0.5	1.0	1.5

2010 The two participants now look quite different. Participant A still
2011 shows a sizeable improvement (four points) under blinded placebo;
2012 participant B has nearly returned to baseline (one and a half points).
2013 This is the diagnostic contrast that the decomposition exploits. For
2014 simplicity the table holds TV at its week-8 value; in reality TV con-
2015 tinues to accumulate over weeks 9 to 12, which shifts both totals by a
2016 common amount and leaves the BR -versus- PB contrast on which the
2017 example turns unchanged.

2018 [77]

2019 **The one-component model estimates drug at 7.75 points with**
2020 **no principled decomposition, obscuring that A's true BR is**
2021 **4 and B's is 8.**

- The 7.75 point estimate is a conflated average of true BR values (4.0 and 8.0) plus the placebo wash-out.
- The model cannot determine whether A and B differ, nor estimate the pharmacological drug effect separately.
- Placebo and natural-history effects are not identifiable from this contrast alone.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A one-component fixed-effects regression of total symptom change on phase-by-allocation indicators sees:

- Open-label means: 11.0 (A), 10.0 (B).
- Blinded-placebo means: 4.0 (A), 1.5 (B).
- Estimated drug effect: $11.0 - 4.0 = 7.0$ (A), $10.0 - 1.5 = 8.5$ (B). Population mean: 7.75.
- Estimated placebo effect: not identifiable from this contrast.
- Estimated natural history: not identifiable from this contrast.

The one-component model has produced a single number conflated with the placebo response that washes out under blinding. It cannot say whether the drug effect is the same for participants A and B; it estimates ‘drug’ at 7.75 points on average, which is an average of the true 4.0 (A) and 8.0 (B) mixed with the placebo wash-out, with no principled way to decompose further.

[78]

The full decomposition recovers all six component estimates by exploiting phase differential expectancy and the blinded-placebo contrast.

- The LME returns $\hat{m}_{BR}^A = 4.0$, $\hat{m}_{BR}^B = 8.0$, $\hat{m}_{PB}^A = 6.0$, $\hat{m}_{PB}^B = 1.0$ at the population point estimate.
- Participant-level estimates have meaningful uncertainty at trial-relevant N ; the simulation study characterises this.
- The phase differential ($\eta = 1$ vs. $\eta = 0.5$) is the identification mechanism for the PB parameter.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A linear-mixed-effects analysis with the full BR-PB-TV decomposition, fitted to the same data, can recover the individual components by exploiting the differential expectancy between phases and the on-drug-versus-blinded-placebo contrast. After fitting, the model returns $\hat{m}_{BR}^A = 4.0$, $\hat{m}_{BR}^B = 8.0$, $\hat{m}_{PB}^A = 6.0$, $\hat{m}_{PB}^B = 1.0$, $\hat{m}_{TV}^A = \hat{m}_{TV}^B = 1.0$ (at the population point estimate; participant-level estimates have meaningful uncertainty at N in the trial-relevant range, which the

simulation study below characterises).

[79]

The one-component model produces one biased average; the three-component model produces six clinically actionable quantities, making biomarker validation possible.

- The lumped estimate of 7.75 points is a biased average of $BR_A = 4$ and $BR_B = 8$ confounded with placebo wash-out.
- The decomposition supplies individual BR , PB , and TV estimates enabling per-participant clinical decisions.
- Predictive-biomarker validation requires the decomposition because biomarker-by- BR is the target estimand.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The information loss from lumping the components is not abstract. We may compare the inferences directly:

Question	One-component answer	Three-component answer
Does the drug work for A?	~7.75 points (lumped).	$BR_A = 4.0$, distinct from her larger placebo response.
Does the drug work for B?	~7.75 points (lumped).	$BR_B = 8.0$, twice the magnitude of A.
Are A and B different responders?	No discernible difference.	B has 2x A's BR but 1/6 of A's PB .
Is a biomarker predicting BR useful here?	Cannot answer.	Yes if the biomarker correlates with $\hat{m}_{BR}$.
Should A be continued long-term?	Cannot say.	Probably yes ($BR = 4$ is real and pharmacological), but the larger placebo component will not persist if expectations decay.
Population mean drug effect?	7.75 points (biased).	$6.0 = (4.0 + 8.0)/2$ (correct mean of BR).

The one-component model has produced one biased average where the three-component model has produced six clinically actionable quantities. Predictive-biomarker validation, which is a central scientific aim

of the trial design family considered here, becomes impossible without the decomposition.

[80]

A population-level example shows that cohort differences in open-label totals driven entirely by PB produce apparent non-replicability that the decomposition resolves.

- Cohort 1 (naive, $m_{PB} = 7$) reports 13-point improvements; Cohort 2 (sceptical, $m_{PB} = 2$) reports 8 points; true $m_{BR} = 5$ in both.
- A one-component analysis would falsely conclude the drug works less well in Cohort 2 and trigger searches for confounders.
- The decomposition returns $\hat{m}_{BR} = 5$ in both cohorts, attributing the difference correctly to PB heterogeneity.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A second example at the population level illustrates the same problem. Consider two cohorts of N-of-1 participants enrolled in methodologically identical trials of the same drug. Cohort 1 (early trial, naive patients) has mean $m_{BR} = 5$, mean $m_{PB} = 7$, mean $m_{TV} = 1$. Cohort 2 (later trial, patients who have had multiple prior failed treatments and are sceptical) has mean $m_{BR} = 5$, mean $m_{PB} = 2$, mean $m_{TV} = 1$. By construction the drug works equally well in both cohorts. But the open-label totals differ substantially: cohort 1 reports thirteen-point improvements on average, cohort 2 reports eight-point improvements. A one-component analysis comparing the two cohorts would conclude that the drug works less well in cohort 2 and might trigger a search for confounders or sub-population effects. None of this is real; the difference is entirely in PB , specifically in the population baseline level of placebo responsiveness. The three-component decomposition, applied to either cohort, returns $\hat{m}_{BR} = 5$ in both. Unfortunately, this is the kind of apparent non-replicability that pharmacology has spent the last decade struggling to navigate; the decomposition is one of the formal tools that makes it tractable.

7 Simulation study

[81]

The simulation study follows ADEMP and quantifies bias and identifiability properties of three analysis strategies across a prazosin-PTSD-calibrated parameter grid.

- The Morris (2019) ADEMP framework structures the study: Aims, DGP, Estimands, Methods, Performance measures.

- 2121 • The pre-registration is at `analysis/scripts/component-decomposition/00-ademp-pre-registrat`
2122 • Three analysis strategies (one-component, phase-augmented, full
2123 decomposition) are compared across the parameter grid.

2124 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
2125 *tempor incididunt ut labore et dolore magna aliqua.*

2126 The simulation study reported and pre-registered in this sec-
2127 tion is structured under the⁵⁵ ADEMP framework, specify-
2128 ing Aims, Data-generating mechanisms, Estimands, Methods,
2129 and Performance measures. The pre-registered plan is at
2130 `analysis/scripts/component-decomposition/00-ademp-pre-registration.md`.
2131 A Monte Carlo simulation calibrated to the prazosin/PTSD reference
2132 parameters quantifies the bias and identifiability properties of the
2133 three analysis strategies described in section 5 across a parameter
2134 grid.

2135 7.1 Pilot validation of analysis machinery (100 replicates, 6 cells)

2136 [82]

2137 **The pilot is a structural check of the analysis pipeline, not a**
2138 **substantive empirical result.**

- 2139 • The production simulation in section 7.2 constitutes the empirical
2140 core of the manuscript.
2141 • The pilot confirms only that data generation, model fitting, and
2142 summary output run without error.
2143 • No bias or power claims are derived from the 100-replicate pilot.

2144 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
2145 *tempor incididunt ut labore et dolore magna aliqua.*

2146 The production simulation programme described in the following sec-
2147 tion is the empirical core of this manuscript and has not yet been
2148 executed. The pilot reported here confirms only that the analysis
2149 machinery runs end-to-end.

2150 [83]

2151 **The pilot executed a 6-cell, 100-replicate run with perfect**
2152 **convergence, confirming the pipeline is operational.**

- 2153 • The grid crossed three *pb_strength* levels against two analysis
2154 strategies.
2155 • Convergence was 1.0 across all six cells under 8-core `mclapply`.
2156 • Output is checkpointed at `analysis/data/quick-sim/`; no infer-
2157 ential claims are drawn.

2158 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
2159 *tempor incididunt ut labore et dolore magna aliqua.*

A 100-replicate-per-cell pilot was run at the reference Hybrid OL+BDC design, $N = 35$, on the 6-cell slice $pb_strength \in \{0, 6.5, 13\} \times analysis \in \{\text{one-component, phase-augmented}\}$, using 8-core `parallel::mclapply`. The convergence rate was 1.0 across all six cells. Per-replicate output is checkpointed at `analysis/data/quick-sim/06-decomp-replicates.rds` and the Morris-format cell-level summary at `analysis/data/quick-sim/06-decomp-summary.txt`. The exploratory summaries are diagnostic only. The one-component `bm:Dbc` test showed power 0.95, 0.93, and 0.83 across $pb_strength \in \{0, 6.5, 13\}$, while the phase-augmented analysis showed 0.65, 0.65, and 0.61; the mean estimated $\hat{\beta}_{bm:D_{bc}}$ was essentially flat within each analysis (about -2.25 for the one-component and -2.21 for the phase-augmented fit). At 100 replicates the Monte Carlo standard error on a power estimate near 0.7 is about 0.046, so these figures cannot by themselves separate the two analyses; their interpretation, including why the phase-augmented analysis loses its bias-detection contrast at this N , is taken up in the Discussion. No confirmatory bias or power claim is drawn from the pilot, and all substantive claims are deferred to the production programme below.

[84]

Two pilot findings revise the production specification: rank-deficiency forces a formula change, and bias must be read from coefficient means, not rejection rates.

- At $N = 35$, rank-deficiency dropped the $D_{bc}:\text{phase}$ term, eliminating the key biomarker-by-treatment-by-phase contrast.
- The predicted bias manifests in the mean of $\hat{\beta}_{bm:D_{bc}}$, not in the Wald-test rejection rate.
- The production programme addresses both by using larger N , a fuller design, and coefficient-mean reporting.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Two design-level observations from the pilot inform the production specification. First, at $N = 35$ the rank-deficiency of the full phase-augmented formula required dropping the $D_{bc}:\text{phase}$ interaction term, removing the precise contrast that would have isolated the biomarker-by-treatment-by-phase signal the framework is designed to detect. Second, the predicted bias of the one-component analysis arises in the regression coefficient $\hat{\beta}_{bm:D_{bc}}$, not in the indicator-level rejection rate, so power on the biomarker-by-treatment Wald test is not the right summary statistic to detect the framework's predicted bias. The production programme below addresses both observations through larger N , a richer trial design that admits the full phase-augmented formula, and reporting of the coefficient mean against the implied true value

rather than the rejection rate alone.

7.2 Production-simulation specification

[85]

The production simulation is organised around four deliverables under the ADEMP framework.

- Driver scripts are located at `analysis/scripts/component-decomposition/`.
- The ADEMP pre-registration is committed prior to production runs.
- The four deliverables address bias, identifiability, type I/power, and participant-level recovery.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The production programme follows the⁵⁵ ADEMP framework. The driver scripts will live at `analysis/scripts/component-decomposition/`. We shall describe the four deliverables in turn.

[86]

The first deliverable quantifies how much each analysis strategy biases the biomarker-treatment interaction estimate as placebo strength increases.

- The estimand is mean $\hat{\beta}_{bm:D_{bc}}$ with Monte Carlo SE across 1,000 replicates per cell.
- Sample sizes $N \in \{35, 70, 100, 150\}$ are crossed with a `pb_strength` grid.
- The headline measures are the bias and power of each strategy across the grid; see the realised Study A results in section 7.3.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The first deliverable is the bias of the biomarker-treatment interaction estimate: mean and Monte Carlo standard error of $\hat{\beta}_{bm:D_{bc}}$ for the one-component, phase-augmented, and full-decomposition analyses, across sample sizes $N \in \{35, 70, 100, 150\}$ and a grid of `pb_strength` values. The headline measures are the bias and power of each strategy across the grid; the realised Study A run (reported in section 7.3 below) used 250 replicates per cell.

[87]

The second deliverable assesses whether the full-decomposition analysis is identifiable at trial-relevant sample sizes.

- Identifiability is measured via convergence rates of the nonlinear-mixed-effects fit.
- Variance-inflation factors on the Gompertz parameters diagnose collinearity across components.
- Sensitivity of the BR estimate to η (phase) multipliers is reported explicitly.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The second deliverable is a set of identifiability metrics for the full-decomposition analysis: convergence rates of the nonlinear-mixed-effects fit, variance-inflation factors on the component-specific Gompertz parameters, and the sensitivity of the BR estimate to the chosen η (phase) multipliers.

[88]

The third deliverable characterises type I error and power across analysis strategies under a full sensitivity grid.

- The DGP is the three-component reference; sensitivity axes are η mismatch, TV magnitude, and dropout.
- Type I error cells use 5,000 replicates to achieve MCSE 0.003.
- All three analysis strategies (one-component, phase-augmented, full-decomposition) are compared.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The third deliverable is type I error and power for the biomarker-treatment interaction test under each analysis strategy, with the DGP held at the three-component reference and varied across a sensitivity grid (analyst-versus-truth η mismatch, TV -component magnitude, and dropout). Type I cells are run at $n_{\text{reps}} = 5,000$ for MCSE 0.003.

[89]

The fourth deliverable illustrates which analysis strategies preserve participant-level biomarker stratification.

- Participant-level interaction signals are mapped as a function of (m_{BR}, m_{PB}, m_{TV}) .
- The comparison reveals whether each strategy retains the predictive-biomarker signal.
- The lumped-total regression serves as the baseline that fails to stratify correctly.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The fourth deliverable is the participant-level interaction signal recovered by each analysis, as a function of the underlying participant-level (m_{BR}, m_{PB}, m_{TV}) values, illustrating which analysis strategies preserve the predictive-biomarker stratification that the lumped-total regression does not.

[90]

Production runs use the pmsimstats-ng tidyverse collection with checkpointed outputs and a committed pre-registration.

- The tidyverse implementation collection drives the simulation grid.
- Cell-level summaries are checkpointed under `analysis/data/`; per-replicate files under `output/`.
- The ADEMP pre-registration is committed at `00-ademp-pre-registration.md` before any production run.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The simulation will use the existing pmsimstats-ng tidyverse implementation collection. Cell-level summaries will be checkpointed under `analysis/data/`, with per-replicate output under `analysis/data/quick-sim/06-decomp-replicates.rds` (pilot) and `analysis/scripts/component-decomposition/output/` (production). The ADEMP pre-registration is committed at `analysis/scripts/component-decomposition/00-ademp-pre-registration.md` prior to the production runs.

7.3 Study A results: bias and power at 250 replicates per cell

[91]

At 250 replicates per cell the one-component estimate of $\hat{\beta}_{bm:D_{bc}}$ is essentially unbiased across the entire (m_{PB}, m_{TV}) grid, contradicting the framework's prediction of a placebo-strength-driven bias.

- The true biomarker-by-*BR* effect is -2.25 ; mean absolute bias falls from 0.040 at $N = 35$ to 0.012 at $N = 150$, in every case within one Monte Carlo SE of zero.
- Bias shows no monotone trend across $m_{PB} \in \{0, 1, 3, 6, 10\}$ or $m_{TV} \in \{-1, 0, 1, 2\}$; it behaves as sampling noise around zero.
- This is the predicted consequence of the identity in section 6.1: the DGP couples the biomarker to *BR* only, so $\beta_{bm}^{PB} = \beta_{bm}^{TV} = 0$ and the lumped slope is unbiased by construction.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The production Study A run (250 replicates per cell, 80 alternative cells crossing $m_{PB} \in \{0, 1, 3, 6, 10\}$, $m_{TV} \in \{-1, 0, 1, 2\}$, and $N \in \{35, 70, 100, 150\}$, plus four null cells) returns a result that does not support the framework's central bias prediction. The data-generating process couples the baseline biomarker to the BR component only, through the on-drug biomarker- BR covariance channel, so the true biomarker-by-treatment effect is the biomarker-by- BR slope, $\beta_{bm:D_{bc}} = -c_{bm} \sigma_{BR} / \sigma_{bm} = -2.25$, and the biomarker is by construction uncorrelated with PB and TV . Under the one-component analysis the mean estimated $\hat{\beta}_{bm:D_{bc}}$ tracks this value across the whole grid: the mean absolute bias is 0.040 at $N = 35$, 0.019 at $N = 70$, 0.011 at $N = 100$, and 0.012 at $N = 150$, in every case within one Monte Carlo standard error (respectively about 0.05, 0.04, 0.017, and 0.013) of zero, with no monotone trend across m_{PB} or m_{TV} . The predicted placebo-attribution bias of the biomarker-by-treatment estimate does not appear.

This is not a failure of the simulation; it is the behaviour the identity of section 6.1 requires. That identity writes the lumped biomarker slope as $\beta_{bm}^{BR} + w_{PB}\beta_{bm}^{PB} + w_{TV}\beta_{bm}^{TV}$. With the biomarker coupled to BR alone, $\beta_{bm}^{PB} = \beta_{bm}^{TV} = 0$, so the lumped slope equals the pharmacological slope no matter how large m_{PB} or m_{TV} grows. Study A varies the magnitudes of the unmodelled components, which the identity shows to be the innocuous axis; the axis that produces biomarker-validation bias is the biomarker's correlation with PB or TV , which this DGP holds at zero. The clean null is therefore a confirmation of the identity and a correction to the looser claim of section 1 that large PB or TV by itself biases biomarker validation. It does not; biomarker-component coupling does.

[92]

The phase-augmented analysis provides no bias reduction and incurs a substantial power penalty, so it is dominated under this uncontaminated DGP.

- Mean power at $N = 35$ is 0.35 (phase-augmented) versus 0.59 (one-component); at $N = 70$, 0.62 versus 0.91; both reach approximately 1.0 by $N = 100$.
- Phase-augmented mean absolute bias is no smaller than one-component at any N (0.051 versus 0.040 at $N = 35$).
- Convergence is approximately 100% for both analyses at every N (minimum 0.996), so the rank-deficiency that degraded the $N = 35$ pilot does not recur here.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The comparison between analyses is equally clear. The phase-

augmented analysis is not less biased than the one-component analysis at any sample size (its mean absolute bias is 0.051 versus 0.040 at $N = 35$ and is comparable elsewhere), and it is markedly less powerful for the biomarker-by-treatment test: mean power across the grid is 0.35 versus 0.59 at $N = 35$ and 0.62 versus 0.91 at $N = 70$, with both analyses reaching essentially full power by $N = 100$. Because the DGP contains no biomarker- PB contamination for the phase contrast to remove, the extra phase-by-treatment terms add estimation cost without buying any bias reduction, and the analysis is dominated. Convergence was approximately complete for both analyses at every sample size (minimum convergence rate 0.996), so the rank-deficiency that forced dropping the `Dbc:phase` term in the $N = 35$ pilot, which we had advanced as the explanation for the pilot's inverted power ranking, is not operative in the production run: the inverted ranking persists with the full formula fitted and converged. The pilot's power inversion was therefore a genuine feature of this DGP, not an artifact of rank-deficiency.

Calibration is adequate for both analyses. Nominal-95% confidence-interval coverage of $\beta_{bm:D_{bc}}$ ranges from 0.97 to 0.99 for the one-component analysis (slightly conservative) and from 0.95 to 0.98 for the phase-augmented analysis. Type I error at the four null cells is controlled, ranging from 0.032 to 0.060 for the one-component analysis and 0.020 to 0.052 for the phase-augmented analysis, in every case within Monte Carlo error of the nominal 0.05.

[93]

Study A confirms the identity but cannot exhibit the biomarker-validation failure, which requires a DGP that couples the biomarker to PB ; that corrected cell is the necessary next run.

- As pre-registered, Study A varies component magnitude with the biomarker fixed to BR , so it probes an axis the identity predicts to be null.
- Demonstrating the predicted contamination requires varying β_{bm}^{PB} , a biomarker- PB covariance, at fixed m_{PB} .
- On current evidence the phase-augmented analysis is not recommended for this design; it should be re-evaluated only under a contaminated-biomarker DGP.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Two conclusions follow for the programme. First, the result sharpens rather than supports the section 1 framing: the biomarker-validation failure is real as an identity, but it is driven by the biomarker's coupling to PB or TV , not by the magnitude of those components, and

Study A as pre-registered varies the latter while holding the former at zero. The decisive cell, which we did not run here, varies the biomarker- PB covariance β_{bm}^{PB} at fixed m_{PB} ; the identity predicts the lumped bias to scale with $w_{PB}\beta_{bm}^{PB}$, and that is the cell that would confirm or refute the practical importance of the failure mode. Second, on the present evidence the phase-augmented analysis carries a real power cost without a compensating bias reduction under an uncontaminated biomarker, so it should not be recommended as a default for this design; its value, if any, is confined to the contaminated-biomarker regime that the corrected Study A would probe. We report this negative result in full because it bounds the claims the framework can currently support.

8 Application to predictive-biomarker validation

[94]

The precision-medicine question – whether a biomarker predicts pharmacological response – is answered by the lumped total only when the biomarker is uncorrelated with PB and TV ; otherwise it requires decomposing the response.

- The question is whether the biomarker predicts the BR component, not the lumped total.
- The PATH statement formalises this as the target estimand for predictive biomarker validation.
- The BR-PB-TV decomposition is the mechanism that separates the pharmacological signal when biomarker-component coupling cannot be ruled out.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The motivating clinical question for the trial designs in this project is whether a baseline biomarker predicts treatment response, and specifically whether the biomarker predicts the pharmacological component of treatment response. This is the precision-medicine question, formalised in the PATH statement¹⁴. Whether it depends on the BR-PB-TV decomposition is itself conditional: by the identity of section 6.1, the lumped biomarker slope equals the pharmacological slope when the biomarker is uncorrelated with PB and TV , and departs from it only by the biomarker’s covariance with those components. The decomposition is therefore the mechanism the question requires precisely in the regime where that covariance is non-negligible.

[95]

A lumped-total biomarker regression conflates pharmacolog-

ical and placebo-belief prediction, producing a stratification biomarker that mis-identifies responders.

- $\hat{\beta}_{bm}$ from the lumped regression contains both pharmacological and placebo-belief prediction in unknown proportions.
- If PB covaries with the biomarker distribution, flagged 'high responders' include high- PB responders who will not benefit pharmacologically.
- Prospective application of such a biomarker would systematically mis-stratify patients for prescription decisions.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A direct regression of total response on baseline biomarker

```
total_response = beta_0 + beta_bm * biomarker + error
```

estimates the biomarker's correlation with whatever the lumped total contains. If PB varies systematically across the biomarker distribution, because the biomarker is correlated with optimism, suggestibility, or treatment-seeking history, then $\hat{\beta}_{bm}$ contains a placebo-prediction component inseparable from the pharmacological-prediction component. The resulting biomarker, applied prospectively, would mis-stratify patients: those flagged as 'high responders' would include both true high- BR responders and high- PB responders, and the latter would not benefit from prescription decisions made on a pharmacological basis.

[96]

A valid pharmacological predictor requires a non-zero $\hat{\beta}_{bm}^{BR}$; significant $\hat{\beta}_{bm}^{PB}$ or $\hat{\beta}_{bm}^{TV}$ indicate a psychological or prognostic biomarker, not a pharmacological one.

- The sole criterion for pharmacological prediction is a meaningfully non-zero BR -component regression coefficient.
- A biomarker correlating only with PB selects placebo responders, not drug responders.
- A biomarker correlating only with TV is prognostic but not predictive in the pharmacological sense.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The right approach is to regress each component on the biomarker:

$$\begin{aligned}
BR &= \beta_0^{BR} + \beta_{bm}^{BR} \cdot \text{biomarker} + \varepsilon^{BR}, \\
PB &= \beta_0^{PB} + \beta_{bm}^{PB} \cdot \text{biomarker} + \varepsilon^{PB}, \\
TV &= \beta_0^{TV} + \beta_{bm}^{TV} \cdot \text{biomarker} + \varepsilon^{TV}.
\end{aligned}$$

A biomarker is a useful pharmacological predictor if and only if $\hat{\beta}_{bm}^{BR}$ is meaningfully non-zero. Significant $\hat{\beta}_{bm}^{PB}$ or $\hat{\beta}_{bm}^{TV}$ are diagnostic findings about the biomarker, indicating that it is partly psychological or correlates with natural-history characteristics, but they are not the validation of pharmacological prediction. A biomarker that correlates only with PB would be valuable for selecting placebo responders in run-in designs but worthless for predicting pharmacological efficacy; a biomarker that correlates only with TV would be a prognostic marker, not a predictive one.

[97]

When biomarker-component coupling cannot be ruled out, the decomposition is the mechanism by which the precision-medicine question is posed; when it can, the lumped analysis already answers it.

- Where the biomarker may correlate with PB or TV , separating BR from them is what prevents validation from conflating pharmacological and non-pharmacological prediction.
- The phase-augmented `bm:Dbc:phase` interaction is a small- N proxy, but Study A shows it buys nothing under an uncontaminated biomarker, so it is warranted only when contamination is plausible.
- The full decomposition recovers component-specific regressions explicitly when sample size permits.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The decomposition is therefore the mechanism by which the precision-medicine question is posed whenever the candidate biomarker may correlate with PB or TV . Where the biomarker can be assumed orthogonal to the non-pharmacological components, the identity of section 6.1 shows the lumped analysis already targets the biomarker-by- BR estimand, and the simulation of section 7.3 confirms it recovers that estimand without bias. The phase-augmented linear-mixed analysis introduced in section 5 retains a useful subset of the decomposition machinery at small N (the `bm:Dbc:phase` interaction tests whether the biomarker-by-treatment effect attenuates under blinding, the practical proxy for biomarker-by- PB contamination), but Study A shows this

proxy earns its degrees of freedom only when such contamination is present: under an uncontaminated biomarker it reduced no bias and lost power. The full decomposition recovers the component-specific regressions explicitly when N permits and when the scientific question warrants the cost.

9 Discussion

[98]

The three-component framework reframes treatment response as a structured sum of causally distinct contributions, and the present work maps when acting on that decomposition actually changes inference.

- The framework specifies which trial-design contrasts identify which components.
- Acting on the decomposition is valuable conditionally: for attribution estimands, and for biomarkers that may correlate with PB or TV , but not for a biomarker-by-treatment interaction with a biomarker orthogonal to them.
- The target estimand for biomarker validation is the biomarker-by- BR association; whether a lumped analysis already recovers it depends on the biomarker's coupling to PB and TV .

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The three-component decomposition is the formal expression of the claim that ‘treatment response’ in an aggregated N-of-1 trial is not a single quantity but a sum of three causally distinct contributions, each with its own pharmacological, psychological, or epidemiological interpretation. What the present work establishes, analytically in section 6.1 and empirically in section 7.3, is that the inferential value of acting on this decomposition is conditional rather than universal. The identity shows that the one-component estimator is displaced from its pharmacological target only by the components’ loading onto the contrast used, and, for a biomarker slope, only by the biomarker’s covariance with PB and TV . Where those terms are zero, the design-contrast one-component analysis is already unbiased, and the production simulation confirms it: with a biomarker coupled to BR alone, the one-component biomarker-by-treatment estimate is unbiased across the whole PB - TV grid and the phase-augmented analysis is dominated. The framework’s contribution is therefore best read as a map of when decomposition changes inference, not as a blanket prescription to decompose. It makes the trial-design conditions explicit, identifies the attribution and contaminated-biomarker estimands for which decom-

position is needed, and clarifies the target estimand for biomarker validation (the biomarker-by- BR association, not the lumped total).

[99]

Four open questions bound the framework’s current scope: the simulation results are pending, the η multipliers require sensitivity analysis, non-monotone natural history is not yet handled, and participant-level recovery inherits the identifiability caveats of individual treatment effects.

- Quantitative bias-versus-power trade-offs are conjectures until the section 7 simulation is run.
- The η (phase) expectancy-modulation parameters are modelling assumptions requiring explicit sensitivity analysis.
- Cyclical or biphasic TV trajectories exceed the Gompertz form and require flexible or covariate-augmented specifications.
- Participant-level component recovery rests on separating participant-by-component signal from participant-by-period and participant-by-phase variation, the identifiability caveat Senn raises for individual treatment effects.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Four open questions remain and warrant acknowledgement. First, the simulation study described in section 7 is forthcoming; the present manuscript reports the framework rather than its empirical evaluation, and the quantitative bias-versus-power trade-offs across analysis strategies are conjectures until that work is run. Second, the η (phase) multipliers used to parameterise the expectancy modulation of PB are themselves modelling assumptions, and a sensitivity analysis varying η across plausible values is mandatory before any clinical conclusion is drawn from the framework’s outputs. Third, the framework as specified treats the three components as monotone Gompertz trajectories; cyclical or biphasic TV patterns require either flexible TV specifications or explicit covariates beyond the Gompertz form, and the boundary between ‘flexible-Gompertz’ and ‘covariate-augmented’ specifications is not yet sharply drawn. Fourth, the participant-level component estimates the framework targets inherit the identifiability caveat that⁵² and⁴⁹ raise for individual treatment effects more generally: within-participant component recovery rests on separating a true participant-by-component signal from participant-by-period and participant-by-phase variation, and this separation is credible only when the design supplies repeated phase contrasts and the cross-component covariance structure is not pathological. The framework should be read as estimating these quantities under those assumptions, not as dissolving the well-known difficulty of estimating individual ef-

fects from finite within-person data.

[100]

Two companion documents extend the manuscript: a pedagogical narrative and a Gompertz-family sensitivity evaluation.

- The pedagogy document at `docs/24-component-decomposition-pedagogy.md` provides worked examples for first-time readers.
- The Gompertz evaluation at `analysis/report/07-gompertz-evaluation/` quantifies trajectory-specification sensitivity.
- Together they address the intuitive and parametric dimensions of the framework that the present manuscript treats briefly.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The companion pedagogical document at `docs/24-component-decomposition-pedagogy.md` develops the framework with worked examples and an extended intuition narrative for readers approaching the material for the first time. The companion methodological evaluation of the Gompertz family at `analysis/report/07-gompertz-evaluation/` quantifies the sensitivity of the framework's inference to the specific parametric form chosen for the component trajectories.

[101]

The decomposition is not specific to N-of-1; the same three-component structure and identifiability logic apply to any longitudinal design with belief-varying and exposure-varying contrasts.

- The claim that a treatment response is the sum of pharmacological, belief-driven, and natural-history components is a statement about the data-generating process of any repeated-measures study.
- Parallel-group trials with repeated measures, crossover trials, and longitudinal observational cohorts all admit the required contrasts.
- The aggregated N-of-1 case is the most demanding instance because the components must be recovered from a single participant's trajectory.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A further point concerns the scope of the decomposition. Although we have developed it for the aggregated N-of-1 design, nothing in the three-component structure is specific to N-of-1 trials. The claim that

an observed longitudinal treatment response is the sum of a pharmacological component, a belief-driven component, and a natural-history component is a statement about the data-generating process of any repeated-measures treatment study, and the identifiability logic of section 4 transfers directly: wherever the design supplies a contrast in which belief varies independently of pharmacology (a blinded discontinuation, a placebo arm, an open-hidden administration) and a contrast in which drug exposure varies (on-drug versus off-drug, a crossover period), the same three components are separately identified. Parallel-group randomised trials with repeated measures, crossover trials, and longitudinal observational cohorts are all candidate settings, and the pharmacometric disease-progression models from which the present framework descends were themselves developed at the group level^{1,2}. We therefore advocate the decomposition not as an N-of-1 technique but as a default lens for any longitudinal design in which placebo and natural-history contributions are non-negligible, with the aggregated N-of-1 case treated here as the most demanding instance because it must recover the components from a single participant's trajectory.

9.1 Interpretation of the pilot diagnostics

[102]

The pilot is a 6-cell, 100-replicate run designed to probe one-component bias under a three-component DGP, with MCSE approximately 0.046 on power estimates near 0.7.

- The design crossed three *pb_strength* levels against one-component and phase-augmented analysis strategies.
- Per-replicate and Morris summary outputs are checkpointed at `analysis/data/quick-sim/`.
- At 100 replicates the Monte Carlo uncertainty is too large to draw reliable power comparisons.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The pilot of section 7.1 was run specifically to probe the predicted bias of the one-component analysis under a three-component DGP; its design, convergence, checkpointed outputs, and Monte Carlo precision are reported there. We interpret its diagnostics here.

[103]

The pilot did not confirm the framework's power prediction; rank-deficiency and a misidentified summary statistic are plausible explanations, but the pilot provides no positive evidence for the predicted effect.

- At $N = 35$ rank-deficiency dropped the `Dbc:phase` term, so the phase-augmented analysis paid fixed-effect cost without gaining the bias-detection contrast.
- Predicted bias appears in the coefficient mean $\hat{\beta}_{bm:D_{bc}}$, which is flat across *pb_strength* at this resolution, not in the Wald rejection rate.
- Coefficient means (approximately -2.25 vs -2.21) cannot be separated from sampling noise with only 100 replicates at $N = 35$.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The inverted finding reported in section 7.1, in which the phase-augmented analysis was *less* powerful than the one-component analysis across the entire *pb_strength* range, has two explanations. First, at $N = 35$ the rank-deficiency of the full phase-augmented formula required dropping the `Dbc:phase` interaction term, which removes the precise contrast that would have isolated the biomarker-by-treatment-by-phase signal the framework is designed to detect; without that term the phase-augmented analysis pays the cost of additional fixed effects without buying the bias-detection mechanism the framework promises. Second, the predicted bias of the one-component analysis arises in the regression *coefficient*, not in the indicator-level rejection rate. The mean estimated $\hat{\beta}_{bm:D_{bc}}$ is essentially flat across *pb_strength* within each analysis (roughly -2.25 for the one-component analysis and -2.21 for the phase-augmented analysis), so the rejection rate on the biomarker-by-treatment Wald test is not the right summary statistic to detect the framework's predicted bias; the bias should be read off the coefficient mean itself, which the pilot's 100-replicate resolution does not separate from sampling noise at the $N = 35$ scale. It must be stated plainly that the pilot did not confirm the framework's prediction: the coefficient mean is essentially flat across *pb_strength*, the very factor predicted to drive the bias, so the present data provide no positive evidence for the predicted effect. This question is resolved by the production run of section 7.3: at 250 replicates per cell across $N \in \{35, 70, 100, 150\}$ the one-component estimate remains unbiased across the placebo-strength axis and the phase-augmented analysis is dominated, with near-complete convergence. The pilot's inverted ranking was therefore not a rank-deficiency artifact but a genuine feature of the data-generating process, and the absence of bias is the behaviour the section 6.1 identity requires when the biomarker is coupled to *BR* alone.

[104]

The production run has now been carried out (section 7.3);

it reproduced the pilot’s inverted power ranking and found no placebo-driven bias, and the remaining empirical question is the contaminated-biomarker cell.

- The realised Study A run used 250 replicates per cell across the full $m_{PB} \times m_{TV} \times N$ grid, with near-complete convergence at every N .
- It confirmed the pilot’s qualitative finding: no bias in the one-component estimate, and the phase-augmented analysis dominated.
- The decisive outstanding cell varies the biomarker- PB covariance, the axis the identity predicts to drive the bias.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The production Study A run reported in section 7.3 supersedes the pilot as the empirical core of this manuscript. It used 250 replicates per cell at $N \in \{35, 70, 100, 150\}$ with the full phase-augmented formula fitted and converged in essentially every replicate, so the rank-deficiency that limited the $N = 35$ pilot does not arise. Its findings are unambiguous: the one-component estimator is unbiased across the placebo-strength and natural-history axes, and the phase-augmented analysis buys no bias reduction while losing power. What remains for a full empirical evaluation of the biomarker-validation failure is the contaminated-biomarker cell, a data-generating process that gives the biomarker a non-zero covariance with PB ; the section 6.1 identity predicts the lumped bias to scale with that covariance, and that cell is the natural next run.

10 Conclusions

[105]

Treatment response in aggregated N-of-1 trials is the sum of three causally distinct components, and the bias from conflating them is governed by the estimand and the biomarker’s coupling to the components, not by their raw magnitude.

- The three components are BR (pharmacological), PB (placebo-belief), and TV (natural history).
- Each has its own biological mechanism, identifiability conditions, and clinical interpretation.
- The lumped pre-post treatment effect is biased by the PB and TV contributions; the biomarker-by-treatment slope is biased only by the biomarker’s covariance with PB and TV , as Study A confirms.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In conclusion, three points are to be emphasised. First, treatment response in aggregated N-of-1 clinical trials is the sum of three causally distinct components: pharmacological action (BR), placebo-belief response (PB), and time-variant natural history (TV). Each component has its own biological mechanism, its own identifiability conditions, and its own clinical interpretation. Whether conflating them biases a given analysis is, however, governed by the omitted-variable-bias identity of section 6.1, not by the raw magnitude of PB and TV : the lumped pre-post treatment effect is displaced by the PB and TV contributions, whereas the biomarker-by-treatment slope is displaced only by the biomarker's covariance with those components. Study A bears this out: varying PB and TV magnitude with a biomarker coupled to BR alone left the one-component estimate unbiased across the entire grid.

[106]

What matters is the design contrast, not the parametric model: the trial must supply the drug-on/off and belief contrasts, but adding parametric component terms is not free and is not always worth it.

- Unique component identification requires phase contrasts built into the trial design; the one-component analysis already exploits the drug-on/off contrast.
- Under an uncontaminated biomarker, the phase-augmented analysis adds no bias reduction and loses power, so it is not recommended as a default.
- The phase-augmented and full component-matched analyses are warranted for attribution estimands and for biomarkers that may correlate with PB or TV , where the extra structure does work.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Second, the operative requirement is the design contrast, not the parametric analysis model. The trial design must include the drug-on/off and belief contrasts that identify the components, and the one-component analysis already exploits the drug-on/off contrast through its D_{bc} term. Adding parametric component structure on top is not free: under an uncontaminated biomarker the phase-augmented analysis reduced no bias and lost substantial power, so it should not be adopted as a default. Its value, and that of the full component-matched fit, is confined to the estimands where the extra structure does work, namely the attribution of response to pharmacology versus belief and the analysis of biomarkers that may themselves correlate

with PB or TV .

[107]

For predictive-biomarker validation the decomposition matters precisely when the candidate biomarker may correlate with PB or TV ; this is the case the decisive simulation still has to probe.

- A biomarker that predicts the lumped total is a pharmacological predictor only if it is uncorrelated with PB and TV ; otherwise the lumped slope is contaminated.
- Regulatory and prescribing decisions therefore require separating BR from PB and TV whenever biomarker-component coupling cannot be ruled out.
- Study A established the uncontaminated baseline; the contaminated-biomarker cell is the planned empirical test of the failure mode.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Third, for predictive-biomarker validation the decomposition earns its place precisely when the candidate biomarker may correlate with PB or TV . A biomarker that predicts the lumped total is a pharmacological predictor only when it is uncorrelated with the non-pharmacological components; where that cannot be assumed, the lumped slope mixes pharmacological and non-pharmacological prediction, and the regulatory and prescribing decisions that motivate the trial require the separation the decomposition supplies. The empirical demonstration of this failure mode, a data-generating process that couples the biomarker to PB , is the decisive cell that Study A did not include and that remains the natural next step; the present run establishes the uncontaminated baseline against which it should be compared.

11 Conflict of interest

The authors declare no conflicts of interest.

12 Funding

This work received no specific funding.

13 Data availability statement

The data and code that support the findings of this study are openly available in the pmsimstats-ng project repository. Per-replicate Monte Carlo outputs and

ADEMP-compliant cell-level summaries are versioned under `analysis/data/`;
simulation drivers are at `analysis/scripts/`. Exact paths for this manuscript
are listed in the Reproducibility section below.

14 Reproducibility

[108]

The manuscript is rendered inside a Docker container with
locked package versions; the full software environment is
pinned by `Dockerfile` and `renv.lock`.

- R 4.4.x in the `zzcollab` Docker container provides the execution environment.
- Simulation packages are `nlme`, `parallel`, `data.table`, `furrr`, and `purrr`.
- The manuscript build uses `rmarkdown` and `bookdown`.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

This manuscript was rendered with R 4.4.x inside the project's `zzcollab` Docker container (image hash recorded in `Dockerfile`; package set captured in `renv.lock`). Principal packages used by the simulation pipeline are `nlme` (continuous-time linear-mixed analysis with `corCAR1` residual correlation), `parallel` (8-core `mclapply` for parameter-grid sweeps), `data.table` (the original-extended simulation collection), `furrr` and `purrr` (the `tidyverse` simulation collection), and `rmarkdown` plus `bookdown` for the manuscript build.

[109]

Reproducibility is ensured by a master seed of 20260507 with
per-replicate seeds derived by index addition.

- `set.seed(20260507)` is called at the top of each replicate loop.
- `parallel::mc.set.seed` propagates seeds inside `mclapply`.
- Per-replicate seeds are `master_seed + rep_idx`.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Each simulation driver sets `set.seed(20260507)` at the top of the replicate loop and uses `parallel::mc.set.seed` inside `mclapply`. Per-replicate seeds derive from the master seed by `+ rep_idx`.

[110]

All simulation outputs, driver scripts, pre-registration, and
manuscript source files are versioned at documented paths
within the repository.

- Per-replicate outputs and Morris summaries are at `analysis/data/quick-sim/`.
- Driver scripts and the ADEMP pre-registration are at `analysis/scripts/component-decomposition/`.
- Manuscript source, bibliography, and CSL file are at `analysis/report/06-component-decomposition/`.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Per-replicate simulation outputs are checkpointed at `analysis/data/quick-sim/06-decomp-replicates`
 Morris ADEMP cell-level summaries at the companion `06-decomp-summary.txt`.
 Driver scripts are under `analysis/scripts/component-decomposition/`.
 The pre-registered ADEMP plan is at `analysis/scripts/component-decomposition/00-ademp-pre-reg`
 The manuscript source, paper-specific bibliography (`references.bib`),
 and SIM CSL file are at `analysis/report/06-component-decomposition/`.

15 References

1. Holford N. Clinical pharmacology = disease progression + drug action. *British Journal of Clinical Pharmacology*. 2015;79(1):18-27. doi:10.1111/bcp.12170
2. Gomeni R, Lavergne A, Merlo-Pich E. Modelling placebo response in depression trials using a longitudinal model with informative dropout. *European Journal of Pharmaceutical Sciences*. 2009;36(1):4-10. doi:10.1016/j.ejps.2008.10.025
3. Chen C, Jönsson S, Yang S, Plan EL, Karlsson MO. Detecting placebo and drug effects on Parkinson's disease symptoms by longitudinal item-score models. *CPT: Pharmacometrics & Systems Pharmacology*. 2021;10(4):309-317. doi:10.1002/psp4.12601
4. Muthén B, Brown HC. Estimating drug effects in the presence of placebo response: Causal inference using growth mixture modeling. *Statistics in Medicine*. 2009;28(27):3363-3385. doi:10.1002/sim.3721
5. Kaptchuk TJ, Kelley JM, Conboy LA, et al. Components of placebo effect: Randomised controlled trial in patients with irritable bowel syndrome. *BMJ*. 2008;336(7651):999-1003. doi:10.1136/bmj.39524.439618.25
6. Kirsch I, Deacon BJ, Huedo-Medina TB, Scoboria A, Moore TJ, Johnson BT. Initial severity and antidepressant benefits: A meta-analysis of data submitted to the Food and Drug Administration. *PLoS Medicine*. 2008;5(2):e45. doi:10.1371/journal.pmed.0050045

- 2903 7. Locher C, Kossowsky J, Gaab J, Kirsch I, et al. Moderation of antidepressant and placebo outcomes by baseline severity in late-life depression: A systematic review and meta-analysis. *Journal of Affective Disorders*. 2015;181:50-60. doi:10.1016/j.jad.2015.03.062
- 2904 8. Sugarman MA, Loree AM, Baltes BB, Grekin ER, Kirsch I. The efficacy of paroxetine and placebo in treating anxiety and depression. *PLoS ONE*. 2014;9(8):e106337. doi:10.1371/journal.pone.0106337
- 2905 9. Hall KT, Lembo AJ, Kirsch I, et al. Catechol-O-methyltransferase val158met polymorphism predicts placebo effect in irritable bowel syndrome. *PLoS ONE*. 2012;7(10):e48135. doi:10.1371/journal.pone.0048135
- 2906 10. Wang RS, Lembo AJ, Kaptchuk TJ, Hall KT, et al. Genomic effects associated with response to placebo treatment in a randomized trial of irritable bowel syndrome. *Frontiers in Pain Research*. 2022;2:775386. doi:10.3389/fpain.2021.775386
- 2907 11. Hróbjartsson A, Gøtzsche PC. Is the placebo powerless? An analysis of clinical trials comparing placebo with no treatment. *New England Journal of Medicine*. 2001;344(21):1594-1602. doi:10.1056/NEJM200105243442106
- 2908 12. Hróbjartsson A, Gøtzsche PC. Is the placebo powerless? Update of a systematic review with 52 new randomized trials. *Journal of Internal Medicine*. 2004;256(2):91-100. doi:10.1111/j.1365-2796.2004.01355.x
- 2909 13. Hróbjartsson A, Gøtzsche PC. Placebo interventions for all clinical conditions. *Cochrane Database of Systematic Reviews*. 2010;(1):CD003974. doi:10.1002/14651858.CD003974.pub3
- 2910 14. Kent DM, Paulus JK, Klaveren D van, et al. The Predictive Approaches to Treatment effect Heterogeneity (PATH) Statement. *Annals of Internal Medicine*. 2020;172(1):35-45. doi:10.7326/M18-3667
- 2911 15. Kent DM, Klaveren D van, Paulus JK, et al. The PATH statement: Explanation and elaboration. *Annals of Internal Medicine*. 2020;172(1):W1-W25. doi:10.7326/M18-3668
- 2912 16. Kent DM, Steyerberg E, Klaveren D van. Personalized evidence based medicine: Predictive approaches to heterogeneous treatment effects. *BMJ*. 2018;363:k4245. doi:10.1136/bmj.k4245
- 2913 17. Rekkas A, Paulus JK, Raman G, et al. Predictive approaches to heterogeneous treatment effects: A scoping review. *BMC Medical Research Methodology*. 2020;20(1):264. doi:10.1186/s12874-020-01145-1

- 2914 18. Guyatt G, Sackett D, Taylor DW, Chong J, Roberts R, Pugsley S. Determining optimal therapy: Randomized trials in individual patients. *New England Journal of Medicine*. 1986;314(14):889-892.
- 2915 19. Zucker DR, Schmid CH, McIntosh MW, D'Agostino RB, Selker HP, Lau J. Combining single patient (N-of-1) trials to estimate population treatment effects and to evaluate individual patient responses to treatment. *Journal of Clinical Epidemiology*. 1997;50(4):401-410. doi:10.1016/S0895-4356(96)00429-5
- 2916 20. Zucker DR, Ruthazer R, Schmid CH. Individual (N-of-1) trials can be combined to give population comparative treatment effect estimates: Methodologic considerations. *Journal of Clinical Epidemiology*. 2010;63(12):1312-1323. doi:10.1016/j.jclinepi.2010.04.020
- 2917 21. Lillie EO, Patay B, Diamant J, Issell B, Topol EJ, Schork NJ. The n-of-1 clinical trial: The ultimate strategy for individualizing medicine? *Personalized Medicine*. 2011;8(2):161-173. doi:10.2217/pme.11.7
- 2918 22. Duan N, Kravitz RL, Schmid CH. Single-patient (n-of-1) trials: A pragmatic clinical decision methodology for patient-centered comparative effectiveness research. *Journal of Clinical Epidemiology*. 2013;66(8 Suppl):S21-S28. doi:10.1016/j.jclinepi.2013.04.006
- 2919 23. Gabler NB, Duan N, Vohra S, Kravitz RL. N-of-1 trials in the medical literature: A systematic review. *Medical Care*. 2011;49(8):761-768. doi:10.1097/MLR.0b013e318215d90d
- 2920 24. Schork NJ, Goetz LH. Single-subject studies in translational nutrition research. *Annual Review of Nutrition*. 2017;37:395-422. doi:10.1146/annurev-nutr-071816-064717
- 2921 25. Schork NJ. Randomized clinical trials and personalized medicine: A commentary on Deaton and Cartwright. *Social Science and Medicine*. 2018;210:71-73. doi:10.1016/j.socscimed.2018.04.033
- 2922 26. Schork NJ, Beaulieu-Jones B, Liang WS, Smalley S, Goetz LH. Exploring human biology with N-of-1 clinical trials. *Cambridge Prisms: Precision Medicine*. 2023;1:e12. doi:10.1017/pcm.2022.15
- 2923 27. Porcino A, Vohra S. N-of-1 trials, their reporting guidelines, and the advancement of open science principles. *Harvard Data Science Review*. 2022;4(SI3). doi:10.1162/99608f92.a65a257a

- 2924 28. Vohra S, Shamseer L, Sampson M, et al. CONSORT extension for reporting N-of-1 trials (CENT) 2015 statement. *Journal of Clinical Epidemiology*. 2015;76:9-17. doi:10.1016/j.jclinepi.2015.05.004
- 2925 29. Shamseer L, Sampson M, Bukutu C, Schmid CH, et al. CONSORT extension for reporting N-of-1 trials (CENT) 2015: Explanation and elaboration. *Journal of Clinical Epidemiology*. 2015;76:18-46. doi:10.1016/j.jclinepi.2015.05.018
- 2926 30. Li J, Gao W, Punja S, et al. Reporting quality of N-of-1 trials published between 1985 and 2013: A systematic review. *Journal of Clinical Epidemiology*. 2016;76:57-64. doi:10.1016/j.jclinepi.2015.11.016
- 2927 31. Liao Z, Qian M, Kronish IM, Cheung YK. Analysis of N-of-1 trials using Bayesian distributed lag model with autocorrelated errors. *Statistics in Medicine*. 2023;42(13):2044-2060. doi:10.1002/sim.9676
- 2928 32. Blackston JW, Chapple AG, McGree JM, McDonald S, Nikles J. Comparison of aggregated N-of-1 trials with parallel and crossover randomized controlled trials using simulation studies. *Healthcare*. 2019;7(4):137. doi:10.3390/healthcare7040137
- 2929 33. Carrozzo AE, Zimmermann G, Bathke AC, et al. Two-arm crossover randomized controlled trial versus meta-analysis of N-of-1 studies: Comparison of statistical efficiency. *Biometrical Journal*. 2025;67(2):e70045. doi:10.1002/bimj.70045
- 2930 34. Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: Statistical methods and practical considerations. *Frontiers in Digital Health*. 2020;2:13. doi:10.3389/fdgth.2020.00013
- 2931 35. Kaptchuk TJ, Friedlander E, Kelley JM, et al. Placebos without deception: A randomized controlled trial in irritable bowel syndrome. *PLoS ONE*. 2010;5(12):e15591. doi:10.1371/journal.pone.0015591
- 2932 36. Lembo A, Kelley JM, Nee J, et al. Open-label placebo vs double-blind placebo for irritable bowel syndrome: A randomized clinical trial. *Pain*. 2021;162(9):2428-2435. doi:10.1097/j.pain.0000000000002234
- 2933 37. Nurko S, Saps M, Kossowsky J, others, Kelley JM. Effect of open-label placebo on children and adolescents with functional abdominal pain or irritable bowel syndrome. *JAMA Pediatrics*. 2022;176(4):349-356. doi:10.1001/jamapediatrics.2021.5750

- 2934 38. Benedetti F. Placebo effects: From the neurobiological paradigm to trans-
lational implications. *Neuron*. 2014;84(3):623-637.
- 2935 39. Frisaldi E, Shaibani A, Benedetti F. Understanding the mechanisms of
placebo and nocebo effects. *Swiss Medical Weekly*. 2020;150:w20340.
doi:10.4414/smw.2020.20340
- 2936 40. Carlino E, Frisaldi E, Benedetti F. Pain and the context. *Nature Reviews
Rheumatology*. 2014;10(6):348-355. doi:10.1038/nrrheum.2014.17
- 2937 41. Carlino E, Piedimonte A, Benedetti F. Nature of the placebo and nocebo ef-
fect in relation to functional neurologic disorders. In: *Handbook of Clinical
Neurology*. Vol 139.; 2016:597-606. doi:10.1016/B978-0-12-801772-2.00048-
5
- 2938 42. Wager TD, Atlas LY, Lindquist MA, Roy M, Woo CW, Kross E. An
fMRI-based neurologic signature of physical pain. *New England Journal
of Medicine*. 2013;368(15):1388-1397. doi:10.1056/NEJMoa1204471
- 2939 43. Zunhammer M, Bingel U, Wager TD. Placebo effects on the neurologic
pain signature: A meta-analysis of individual participant functional mag-
netic resonance imaging data. *JAMA Neurology*. 2018;75(11):1321-1330.
doi:10.1001/jamaneurol.2018.2017
- 2940 44. Raskind MA, Peterson K, Williams T, et al. A trial of prazosin for combat
trauma PTSD with nightmares in active-duty soldiers returned from Iraq
and Afghanistan. *American Journal of Psychiatry*. 2013;170(9):1003-1010.
doi:10.1176/appi.ajp.2013.12081133
- 2941 45. Holford NHG. Holford NHG and sheiner LB 'understanding the dose-effect
relationship - clinical application of pharmacokinetic-pharmacodynamic
models', *Clin Pharmacokin* 6:429-453 (1981) - the backstory. *The AAPS
Journal*. 2011;13(4):662-664. doi:10.1208/s12248-011-9306-5
- 2942 46. Tilburt JC, Emanuel EJ, Kaptchuk TJ, Curlin FA, Miller FG. Prescrib-
ing 'placebo treatments': Results of national survey of US internists and
rheumatologists. *BMJ*. 2008;337:a1938. doi:10.1136/bmj.a1938
- 2943 47. Finniss DG, Kaptchuk TJ, Miller F, Benedetti F. Biological, clinical, and
ethical advances of placebo effects. *The Lancet*. 2010;375(9715):686-695.
doi:10.1016/S0140-6736(09)61706-2
- 2944 48. Kirsch I. Response expectancy as a determinant of experience and be-
havior. *American Psychologist*. 1985;40(11):1189-1202. doi:10.1037/0003-
066X.40.11.1189

- 2945 49. Senn S. Controversies concerning randomization and additivity in clinical trials. *Statistics in Medicine*. 2004;23(24):3729-3753. doi:10.1002/sim.2074
- 2946 50. Bingel U, Wanigasekera V, Wiech K, et al. The effect of treatment expectation on drug efficacy: Imaging the analgesic benefit of the opioid remifentanyl. *Science Translational Medicine*. 2011;3(70):70ra14. doi:10.1126/scitranslmed.3001244
- 2947 51. International Council for Harmonisation of Technical Requirements for Pharmaceuticals for Human Use. ICH E9(R1) addendum on estimands and sensitivity analysis in clinical trials to the guideline on statistical principles for clinical trials. Published online 2019.
- 2948 52. Senn S. Statistical pitfalls of personalized medicine. *Nature*. 2018;563(7733):619-621.
- 2949 53. Daza EJ. Causal analysis of self-tracked time series data using a counterfactual framework for N-of-1 trials. *Methods of Information in Medicine*. 2018;57(S 01):e10-e21. doi:10.3414/ME16-02-0044
- 2950 54. Gärtner T, Schneider J, Arnrich B, Konigorski S. Comparison of Bayesian networks, G-estimation and linear models to estimate causal treatment effects in aggregated N-of-1 trials with carry-over effects. *BMC Medical Research Methodology*. 2023;23:191. doi:10.1186/s12874-023-02012-5
- 2951 55. Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Statistics in Medicine*. 2019;38(11):2074-2102.

2952

2953 *Rendered on 2026-06-13 at 17:57 PDT.*

2954 *Source: ~/prj/res/36-pmsimstats-ng/pmsimstats-ng/analysis/report/06-component-decomposition/r*

2026-06-13 17:57 PDT

79

~/prj/res/36-pmsimstats-ng/pmsimstats-ng/analysis/report/
06-component-decomposition/report.Rmd