

1 Three-component decomposition of treatment response
2 in aggregated N-of-1 trials: pharmacological,
3 expectation-driven, and natural-history components

4 pmsimstats team

5 2026-06-16

6 **Contents**

7	Abstract	2
8	1 Introduction	3
9	1.1 A claim about how N-of-1 trials should be analysed	3
10	1.2 Why this matters: three failure modes of the lumped analysis . . .	5
11	1.3 The current state of the N-of-1 literature	8
12	1.4 The neurobiological and epidemiological reality of the components	11
13	1.5 What this paper contributes	13
14	2 Notation and the formal model	14
15	3 The three components	16
16	3.1 Biological response (BR)	17
17	3.2 Placebo-belief response (PB)	21
18	3.3 Time-variant natural-history component (TV)	29
19	4 Identifiability and trial-design contrasts	34
20	5 Analysis-side methodology	38
21	5.1 Four reasons not to match the DGP	38
22	5.2 A small extension that usually pays off: phase indicators	42
23	5.3 When to fit the full decomposition	43
24	6 Worked example: information loss without decomposition	45
25	6.1 What the lumped analysis estimates	45
26	6.2 A numerical illustration	48
27	7 Simulation study	53
28	7.1 Pilot validation of analysis machinery (100 replicates, 6 cells) . . .	54
29	7.2 Production-simulation specification	55

30	7.3 Study A results: bias and power at 1,000 replicates per alternative	
31	cell	56
32	7.4 Contaminated-biomarker pilot (100 replicates per cell)	59
33	7.5 Study B: recovery under a belief-decoupling design	61
34	8 Application to predictive-biomarker validation	65
35	9 Discussion	69
36	10 Conclusions	74
37	11 Conflict of interest	77
38	12 Funding	78
39	13 Data availability statement	78
40	14 Reproducibility	78
41	15 References	80

42 Abstract

43 **Background.** Treatment response in aggregated N-of-1 clinical trials reflects
44 three causally distinct mechanisms: pharmacological action of the active com-
45 pound (BR), the patient’s belief that the treatment will work (PB), and natural-
46 history drift in the underlying condition (TV). Standard analysis models lump
47 these into a single ‘treatment effect’ plus residual error. Whether this lumping
48 biases inference is not unconditional: it depends on the estimand and on whether
49 candidate biomarkers correlate with the non-pharmacological components, a de-
50 pendence this paper makes precise.

51 **Methods.** We formalise a three-component decomposition of the within-
52 participant outcome trajectory, each component a modified Gompertz function
53 with participant-specific random effects, and characterise the trial-design con-
54 trasts (open-label, blinded discontinuation, crossover) that supply identifiability.
55 We derive the omitted-variable-bias identity for the one-component estimator,
56 which expresses both the lumped treatment effect and the lumped biomarker
57 slope as sums of component-specific terms, and we evaluate the analysis strategies
58 in a pre-registered Monte Carlo study (Study A) calibrated to the prazosin-PTSD
59 reference parameters, comparing a one-component design-contrast analysis with
60 a phase-augmented analysis at 1,000 replicates per alternative cell (5,000 per null
61 cell) across $m_{PB} \in \{0, 1, 3, 6, 10\}$, $m_{TV} \in \{-1, 0, 1, 2\}$, and $N \in \{35, 70, 100, 150\}$.

62 **Results.** The identity shows that the lumped biomarker-by-treatment slope is dis-
63 placed from the pharmacological slope by a term proportional to the biomarker’s
64 covariance with PB and TV, not by the magnitude of those components. Study A
65 confirms this: with a biomarker coupled to BR only, the one-component estimate

of the biomarker-by-treatment interaction is essentially unbiased across the whole grid (mean absolute bias 0.016 at $N = 35$ and 0.006 at $N = 150$, within one Monte Carlo standard error of zero), with near-nominal coverage and type I error. The phase-augmented analysis provides no bias reduction and is markedly less powerful (mean power 0.35 versus 0.59 at $N = 35$; 0.62 versus 0.90 at $N = 70$), so it is dominated under this uncontaminated biomarker. A contaminated-biomarker pilot then exhibits the failure mode directly: coupling the biomarker to PB drives the one-component bias up with the coupling strength, reaching +0.51 at a biomarker-PB correlation of 0.45 ($N = 150$, some 30 Monte Carlo standard errors from zero), and the operative quantity is the biomarker's correlation with the individual-level PB *variance*, not the population-mean placebo response, which does not enter the slope; the phase-augmented analysis does not remove this bias either. The inferential value of decomposition is therefore conditional, not universal.

Conclusions. The trial design must supply the drug-on/off and belief contrasts; given those, component-level decomposition is useful when the estimand is the attribution of response to pharmacology versus belief versus natural history, or when a candidate biomarker may correlate with PB or TV. For a biomarker-by-treatment interaction with a biomarker orthogonal to PB and TV, a single-indicator analysis that exploits the design contrast is sufficient and a parametric component model is counterproductive. A pragmatic specification for the cases that warrant it is $Sx \sim \text{bm} + \text{t} + \text{Dbc} + \text{phase} + \text{Dbc:phase} + \text{bm:Dbc} + \text{bm:Dbc:phase}$ with random intercept and AR(1) residual correlation. A contaminated-biomarker study demonstrates the failure mode (the lumped slope is biased in proportion to the biomarker's correlation with PB), and Study B shows that the bias is removable but only under a design that decouples drug exposure from belief: on a balanced-placebo design the belief-covariate decomposition recovers the unbiased pharmacological slope (bias within one Monte Carlo SE of zero across the feasible contamination range at $N=150$) while the one-component estimator does not, whereas on the standard hybrid design no decomposition recovers it. Recovery is therefore conditional on the trial design, not on the analysis alone.

1 Introduction

1.1 A claim about how N-of-1 trials should be analysed

[1]

This paper argues that treatment response in N-of-1 trials comprises three causally distinct components that must be analysed jointly rather than lumped.

- The three components are *BR* (pharmacological), *PB* (placebo-belief), and *TV* (natural history).
- Each component is causally distinct, clinically interpretable, and identifiable from a well-designed N-of-1 trial.

- The standard single-indicator mixed-effects model conflates all three and renders biomarker validation unanswerable.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

This paper advances an opinionated methodological claim: treatment response in aggregated N-of-1 clinical trials is the sum of three causally distinct components and should routinely be analysed as such, rather than estimated as a single lumped quantity. The three components are pharmacological action (*BR*, the biological response), placebo-belief response (*PB*, the patient’s belief that they are receiving treatment), and time-variant natural history (*TV*, the trajectory that would have occurred without intervention). Each is causally distinct, each is independently of clinical and regulatory interest, and each is identifiable from a properly designed N-of-1 trial. We shall argue that the standard practice of fitting a single mixed-effects model with one treatment indicator, prevalent across the published N-of-1 literature, conflates these three components, biases the interpretation of treatment effect, and renders predictive-biomarker validation unanswerable as a scientific question.

[2]

The contribution is the consolidation and extension of an established structural-decomposition tradition, not a novel decomposition.

- The additive disease-progression, placebo, and drug-action decomposition is well established in pharmacometrics and longitudinal trial methodology.
- The present framework extends that tradition with an explicit, separately identified placebo-belief channel and brings it to the aggregated N-of-1 setting.
- The decomposition logic is not specific to N-of-1; it applies wherever a longitudinal treatment response is analysed.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

We make no claim of novelty for the decomposition itself. The structural decomposition of a longitudinal treatment response into separable disease-progression, placebo, and drug-action components is a well-established tradition in pharmacometrics^{1–3}, and the separation of drug from placebo response in longitudinal trajectories has been pursued both through growth-mixture models⁴ and through experimental decomposition of the placebo response itself⁵. Each ingredient we use (mixed-effects models with trial-phase indicators, placebo-arm contrasts, regression-to-mean adjustment, growth-curve parameterisa-

tion) has likewise appeared across decades of clinical-trial methodology. Our contribution is therefore one of consolidation and extension rather than invention: we make the three-component structure explicit for the aggregated N-of-1 setting, add a separately modelled placebo-belief channel identified by design contrasts rather than by functional form alone, and argue that the decomposition should become routine wherever a longitudinal treatment response is analysed. We document why the framing has not been routine in N-of-1 practice (parsimony, identifiability concerns, sample-size constraints, methodological inertia) and propose an analysis that is tractable at trial-relevant sample sizes through a phase-augmented linear-mixed-effects approximation to the full decomposition.

1.2 Why this matters: three failure modes of the lumped analysis

[3]

The lumped one-component analysis has three documented failure modes that grow with unmodelled component heterogeneity.

- The three failure modes are placebo-attribution bias, regression-to-the-mean, and biomarker-validation failure.
- Each failure mode is documented in the chronic-condition trial literature, not merely hypothetical.
- The failures are present regardless of sample size.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

We begin with the failure modes of the standard approach. A one-component analysis (‘the treatment effect is what we observe minus baseline’) has three identifiable failure modes that arise when the population displays meaningful heterogeneity in any of the unmodelled components. None of these is hypothetical: each has a documented presence in the chronic-condition trial literature. Two points of precision, established formally in section 6.1 and borne out by the simulation of section 7.3, should be carried through what follows. The first two failure modes attach to the baseline-anchored lumped effect and are removed once the design supplies a drug-on/off contrast, which the one-component analysis of later sections in fact uses. The third, biomarker-validation failure, is different in kind: it requires not merely that PB or TV be large but that the candidate biomarker correlate with them, and where that coupling is absent the lumped analysis is not biased at all.

[4]

Placebo-attribution bias is sample-size-invariant and docu-

190 **mented across antidepressant, anxiety, and IBS literatures.**

- 191 • Lumped treatment-effect estimates overstate pharmacological ef-
192 ficacy in proportion to PB magnitude and variance.
- 193 • Kirsch et al. and Locher et al. demonstrated that antidepressant
194 apparent efficacy is largely explained by placebo response and
195 natural fluctuation.
- 196 • Placebo responsiveness has measurable biological correlates (e.g.
197 COMT genotype), confirming it is a structured phenomenon, not
198 noise.

199 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
200 *tempor incididunt ut labore et dolore magna aliqua.*

201 The first is the placebo-attribution failure. When the population con-
202 tains both high and low placebo responders, the lumped treatment
203 effect overestimates the pharmacological component. The bias scales
204 with the magnitude and variance of PB in the study population, and
205 it is invariant under sample-size increases: a larger one-component
206 trial converges on the same biased estimate. The most consequential
207 example in modern clinical pharmacology is antidepressant efficacy.⁶
208 analysed FDA-submitted data on four widely- prescribed antidepres-
209 sants and showed that, except at the most severe baseline-severity
210 stratum, mean drug-placebo differences fell below clinical-significance
211 thresholds, with the apparent ‘efficacy’ largely explained by placebo
212 response and natural fluctuation;⁷ reproduced this pattern in late-life
213 depression and explicitly invoked regression to the mean.⁸ generalised
214 the finding to anxiety. Patient populations differ systematically in
215 placebo responsiveness, and the difference has measurable biological
216 correlates: COMT val158met genotype predicts placebo response in
217 irritable bowel syndrome^{9,10}, indicating that ‘placebo’ is not a homoge-
218 neous nuisance but a structured biological phenomenon that interacts
219 with the patient’s genome and treatment history. A trial whose pop-
220 ulation happens to be placebo-responsive will report a larger ‘drug
221 effect’ than one whose population is not, even when the underlying
222 pharmacology is identical.

223 [5]

224 **Regression-to-the-mean contributes several points of appar-**
225 **ent improvement in early trial weeks and is absorbed without**
226 **trace by the lumped analysis.**

- 227 • Enrolment at peak severity guarantees a natural-history regres-
228 sion component irrespective of pharmacology.
- 229 • Hrobjartsson et al. systematic reviews across 200+ trials showed
230 much of conventional placebo response was indistinguishable
231 from natural-history drift.

- 232 • The lumped analysis has no formal mechanism to recover the
233 pharmacological effect once regression is absorbed.

234 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
235 *tempor incididunt ut labore et dolore magna aliqua.*

236 The second failure mode is regression-to-the-mean. When patients en-
237 rol at moments of clinical engagement that coincide with peak symp-
238 tom severity, natural-history regression toward typical severity con-
239 tributes several points of apparent improvement in the early weeks
240 of any trial, without any pharmacological contribution at all. The
241 systematic reviews of^{11, 12}, and¹³, the last covering 202 trials across 60
242 clinical conditions, quantified this directly: in trials with both placebo
243 and no-treatment arms, much of what is conventionally attributed to
244 the placebo response was indistinguishable from regression to the mean
245 and natural-history drift. The lumped analysis absorbs this regression
246 into the treatment-effect estimate and the analyst has no formal mech-
247 anism to recover the true pharmacological effect.

248 [6]

249 **Biomarker validation is invalidated when a biomarker corre-**
250 **lates with *PB* or *TV*, rendering precision-medicine stratifica-**
251 **tion wrong regardless of trial size.**

- 252 • A biomarker correlated with placebo responsiveness or natural-
253 history trajectory is indistinguishable from a true pharmacologi-
254 cal predictor in a lumped analysis.
- 255 • The PATH statement defers mechanistic attribution to do-
256 main knowledge; the BR-PB-TV decomposition supplies that
257 substrate.
- 258 • Where such coupling cannot be ruled out, rigorous pharmacological-
259 biomarker validation requires the decomposition; where the
260 biomarker is known orthogonal to *PB* and *TV*, the lumped
261 analysis already recovers the pharmacological slope.

262 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
263 *tempor incididunt ut labore et dolore magna aliqua.*

264 The third failure mode is the biomarker-validation failure. The
265 precision-medicine question that motivates many recent N-of-1 trials
266 is whether a baseline biomarker predicts treatment response, for-
267 malised in the PATH statement^{14,15}. A biomarker that correlates with
268 placebo responsiveness or with a favourable natural-history trajectory
269 will appear in a one-component analysis as a predictor of the lumped
270 ‘treatment effect’, indistinguishable from a true pharmacological
271 predictor. The biomarker-stratified prescribing decision that follows
272 from such an analysis is wrong, and the trial’s regulatory utility is
273 degraded. The PATH framework explicitly defers the mechanism

of treatment-effect heterogeneity to domain knowledge^{16,17}; the BR-PB-TV decomposition supplies the mechanistic substrate that PATH leaves blank. We argue that, whenever the candidate biomarker may itself correlate with PB or TV , predictive-biomarker validation cannot be conducted rigorously in a lumped-analysis framework regardless of how large the trial is; the bias is fixed by the biomarker's coupling to the non-pharmacological components, not by sample size. Where the biomarker can be assumed orthogonal to PB and TV , this failure mode does not arise, and the lumped analysis recovers the pharmacological slope, as the simulation of section 7.3 confirms. We stress at the outset the precise form this failure mode takes, because the paper's own simulation history clarifies it. Study A (section 7.3) varies the *magnitude* of PB and TV with the biomarker coupled to BR alone, which the section 6.1 identity predicts to be innocuous, and it finds no bias. The decisive experiment is the contaminated-biomarker pilot (section 7.4), which couples the biomarker to PB and exhibits the failure directly: the lumped bias grows with the biomarker- PB correlation. The operative quantity is the biomarker's correlation with the individual-level PB variance ($c_{bm,PB} \sigma_{PB}$), not the population-mean placebo response m_{PB} , which does not enter the slope; a trial population can have a large average placebo response and yet pose no validation hazard, while one with zero average but a biomarker-correlated spread of placebo responsiveness is at risk. The failure-mode claims of this section are therefore consequences of the identity, demonstrated in simulation (section 7.4), and the bias is recoverable: Study B (section 7.5) shows that a belief-covariate decomposition recovers the unbiased pharmacological slope under contamination, but only on a design that decouples drug exposure from belief, not on the standard hybrid design.

1.3 The current state of the N-of-1 literature

[7]

The N-of-1 literature has matured substantially but motivates three observations that the present paper addresses.

- The design has evolved from individual-patient consultation to a framework supporting population-level inference.
- CONSORT reporting standards and quality audits document progressive improvement in analytic transparency.
- Three specific gaps in the literature motivate the decomposition framework proposed here.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The N-of-1 trial in its modern randomised form originates with¹⁸, who

defined the within-patient multiple-crossover RCT, demonstrated it on a single airflow-limitation case, and established a clinical service to deliver it at scale. The design inherits its identifying logic from two older traditions: the crossover trial of classical biostatistics, in which within-subject period contrasts estimate treatment effects under explicit carryover assumptions, and the single-case experimental designs of behavioural science, in which phase reversals (withdrawal and reinstatement of treatment) identify effects within one participant. The blinded-discontinuation contrast central to this paper is a direct descendant of both. The methodological step from the individual-patient trial to population-level inference was taken by¹⁹ and²⁰, who showed that a series of single-patient trials can be combined through hierarchical meta-analysis to estimate both population treatment effects and individual responses, the statistical foundation on which the aggregated design rests. Over the last fifteen years the design has accordingly been reframed from an individual-patient decision tool into an aggregated, meta-analytic framework supporting population-level inference: the programmatic statement of²¹, the methodological reviews of²², the historical-practice audit of²³, and the more recent syntheses of²⁴,²⁵,²⁶, and²⁷ trace that trajectory. Reporting standards are codified in the CONSORT extension²⁸ (the headline statement) and²⁹ (the explanation-and-elaboration companion); the reporting-quality audits of³⁰ and its successors document slow but progressive improvement in analytic transparency. Three observations about this literature motivate the present paper.

[8]

No published N-of-1 trial employs a structural decomposition of treatment response despite three decades of methodological literature.

- Audit data show the field has relied on paired t-tests, binary-indicator LME, and hierarchical Bayesian regressions.
- Recent methodological advances (Liao 2023, Blackston 2019, Carrozzo 2025) retain the lumped-response framing.
- The decomposition gap persists even as carryover and identifiability problems are increasingly recognised.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The first observation is that the dominant analytic models are simple. The reporting-quality audit of³⁰ documents that the published N-of-1 trials between 1985 and 2013 relied predominantly on paired t-tests, simple linear mixed-effects models with a binary treatment indicator, or hierarchical Bayesian regressions. On our own reading of that corpus and the methodological literature that followed it, no pub-

lished N-of-1 trial has employed a structural decomposition of treatment response into pharmacological, expectancy, and natural-history components. The most recent methodological advances retain this character:³¹ develops Bayesian distributed-lag models with autocorrelated errors as a current carryover-aware analysis, but the model still treats the response as a single quantity to be estimated.³² and³³ compare aggregated N-of-1 trials with parallel and crossover RCT alternatives by simulation and document carryover-induced type-I-error inflation as a key failure mode, but again at the level of the lumped treatment effect. None of these papers decompose the response.

[9]

Placebo-component separation is widely recognised but almost entirely unaddressed at the protocol level in N-of-1 methodology.

- The Hendrickson hybrid blinded-discontinuation phase is among the few protocol mechanisms that supply *PB*-versus-*BR* identifiability.
- Placebo run-in, open-label-placebo, and mediator-analysis approaches have not been imported into N-of-1 designs.
- The gap between awareness of the problem and protocol-level solutions motivates the present framework.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The second observation is that the placebo-component separation problem is widely known but largely unaddressed in N-of-1 specifically. The blinded discontinuation phase used in the Hendrickson hybrid design³⁴ is one of the few protocol-level mechanisms in routine N-of-1 use that supplies identifiability of *PB* versus *BR*. Other relevant approaches, including placebo run-in designs, open-label-placebo trials, and mediator-analysis applications, have been developed in parallel but rarely imported into N-of-1 methodology.

[10]

Kaptchuk's within-placebo decomposition is the closest structural analogue; no published paper applies the corresponding full-response decomposition to N-of-1 trials.

- Kaptchuk et al. (2008) empirically decomposed the placebo response into observation, ritual, and relationship components.
- Open-label placebo trials confirm that clinically meaningful PB-only response is possible even when patients know they receive placebo.
- The BR-PB-TV framework is structurally analogous but operates

400 at the level of the full treatment response.

401 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
402 *tempor incididunt ut labore et dolore magna aliqua.*

403 The third observation is that the closest precedent in the literature is
404 Kaptchuk’s three-component decomposition of the placebo response
405 itself. In a randomised trial of acupuncture-style placebo for irritable
406 bowel syndrome,⁵ decomposed the placebo response into observation,
407 ritual, and patient-practitioner relationship components, demonstrat-
408 ing empirically that the placebo response is itself a sum of identifiable
409 parts rather than a single ‘expectancy’ lump. The follow-up open-label
410 placebo trials of^{35,36}, and³⁷ in adolescent functional pain establish that
411 even disclosed-placebo intervention can produce clinically meaningful
412 PB-only response. The framework we propose is structurally anal-
413 ogous to Kaptchuk’s decomposition but applied at the level of the
414 full treatment response (drug, placebo, natural history) rather than
415 within the placebo response alone. To the best of our knowledge, no
416 published paper proposes the corresponding decomposition for the ag-
417 gregated N-of-1 setting, even though the methodological ingredients,
418 including mixed-effects modelling, blinded-phase contrasts, and trajec-
419 tory parameterisation, are all in routine use.

420 [11]

421 **The decomposition gap is real, the tools to fill it are ma-**
422 **ture, and the failure modes are large enough to make it a**
423 **methodological priority.**

- 424 • The gap between available methodology and current practice is
425 documented and not attributable to missing tools.
- 426 • Failure modes in chronic-condition pharmacology are consequen-
427 tial, not academic.
- 428 • The remainder of the paper develops, applies, and evaluates the
429 decomposition framework.

430 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
431 *tempor incididunt ut labore et dolore magna aliqua.*

432 The case we develop in the rest of this paper is therefore that the gap
433 is real, that the methodological tools to fill it are mature, and that
434 the failure modes of the lumped analysis are large enough in chronic-
435 condition pharmacology to make the decomposition a methodological
436 priority rather than an academic refinement.

437 1.4 The neurobiological and epidemiological reality of the com- 438 ponents

439 [12]

440 **Each component of the decomposition reflects a real causal**
441 **mechanism documented in independent literatures.**

- 442 • *BR* corresponds to pharmacological action; *PB* to neurobiologi-
443 cally real belief-driven mechanisms; *TV* to natural-history drift.
- 444 • The components are not statistical constructs but distinct causal
445 pathways with separate literatures.
- 446 • The section reviews the evidence for each component before es-
447 tablishing the formal model.

448 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
449 *tempor incididunt ut labore et dolore magna aliqua.*

450 The decomposition is not a statistical convenience. Each component
451 reflects a real causal mechanism documented in independent litera-
452 tures, and we shall briefly recall each in turn before turning to the
453 formal model.

454 [13]

455 ***BR* is the pharmacological target of regulation; *PB* is neuro-**
456 **biologically real and phase-modulated, which is precisely the**
457 **design contrast that makes the decomposition identifiable.**

- 458 • *PB* is mediated by endogenous opioid/dopamine release, HPA-
459 axis regulation, and top-down attentional shifts.
- 460 • Wager et al. neurologic pain signature provides a clean fMRI dis-
461 sociation between pharmacological and expectancy contributions.
- 462 • Phase-dependent belief variation (open-label vs. blinded) is the
463 design contrast that supplies *PB* identifiability.

464 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
465 *tempor incididunt ut labore et dolore magna aliqua.*

466 The biological response component represents the pharmacological ac-
467 tion of the active compound: the chemical binding, downstream sig-
468 nalling, and physiological response that a regulatory authority cares
469 about when it grants approval. The placebo-belief response reflects
470 neurobiologically real mechanisms documented across pain, depres-
471 sion, Parkinson's, and PTSD research^{38–41}, mediated by belief-driven
472 release of endogenous opioids and dopamine, anticipatory regulation
473 of the hypothalamic-pituitary-adrenal axis, and top-down attentional
474 shifts in symptom perception. Direct neuroimaging evidence is in-
475 structive here:⁴² developed an fMRI-based 'neurologic pain signature'
476 (NPS) that tracks nociceptive input and is reduced by the opioid ag-
477 onist remifentanyl, and the individual-participant meta-analysis of⁴³
478 found that placebo treatments produce moderate reductions in re-
479 ported pain but only negligible reductions in the NPS, providing a
480 neural dissociation between the pharmacological and expectancy con-
481 tributions to a clinical pain response. The placebo response is hetero-

482 geneous across patients with biological correlates^{9,10} and is modulated
483 by trial phase: when belief varies, as it does between open-label and
484 blinded-discontinuation conditions, the placebo response varies, and
485 the modulation is exactly the design contrast that makes the decom-
486 position identifiable.

487 [14]

488 ***TV* is the untreated counterfactual trajectory; its heterogene-**
489 **ity is the central object of precision-medicine inference and**
490 **must be modelled, not absorbed into noise.**

- 491 • *TV* sign depends on disease class: positive for acute recovery,
492 near-zero for stable chronic, negative for progressive conditions.
- 493 • Selection-driven enrolment at peak severity guarantees positive
494 *TV* in early trial weeks regardless of disease class.
- 495 • The PATH statement identifies treatment-effect heterogeneity
496 modelling as mandatory for precision-medicine inference.

497 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
498 *tempor incididunt ut labore et dolore magna aliqua.*

499 The time-variant natural-history component reflects whatever trajec-
500 tory the participant would have followed without treatment. For acute
501 conditions this trends positive (recovery on intrinsic timescale); for
502 chronic stable conditions it is near zero on the population mean but
503 with substantial participant-level variance; for chronic progressive con-
504 ditions it trends negative; and for selection-driven enrolment processes
505 it is positive in the early weeks of the trial regardless of underly-
506 ing disease class, because of regression to the mean^{11,13}. The PATH
507 statement^{14,15} highlights that this heterogeneity is itself the central ob-
508 ject of precision-medicine inference: heterogeneity of treatment effect
509 must be modelled, not absorbed into noise, if the precision-medicine
510 question is to be answered.

511 1.5 What this paper contributes

512 [15]

513 **The paper makes four contributions spanning formalisation,**
514 **identifiability mapping, tractable methodology, and quantifi-**
515 **cation of the inferential cost of lumping.**

- 516 • Sections 2–3 formalise the decomposition; section 4 maps it to
517 design contrasts; section 5 develops tractable analysis methods.
- 518 • Sections 6–7 provide worked examples and a Monte Carlo simu-
519 lation study; section 8 applies the framework to biomarker vali-
520 dation.

- A companion pedagogical document provides extended worked examples for first-time readers.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

We make four contributions, each addressed in a section below. We begin in section 2 with the formal model and notation. Section 3 develops each component in turn, with biological motivation, parametric form, and identifiability conditions. Section 4 maps the design-contrast structure of the standard hybrid open-label-plus-blinded-discontinuation design onto the identifiability requirements of the decomposition. Section 5 develops the analysis-side methodology, including the linear-mixed-effects approximations and the phase-by-treatment indicator extension that retains BR-PB sensitivity at low parametric cost. Section 6 presents the worked examples, and section 7 presents the simulation study. Section 8 discusses applications to predictive-biomarker validation. The companion pedagogical document at docs/24-component-decomposition-pedagogy.md provides an extended worked-example treatment for readers approaching the material for the first time.

2 Notation and the formal model

[16]

This section fixes compact notation and states the formal model; a fuller worked-example treatment is in the companion pedagogical document.

- Definitions are precise enough to support the identifiability and analysis-side arguments in later sections.
- The presentation is kept compact; extended derivations are deferred to the companion document.
- The companion pedagogical document provides numerical worked examples at each formal step.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In this section we shall fix notation and state the formal model. The three components are defined precisely enough to make the identifiability and analysis-side arguments in later sections rigorous, but we have endeavored to keep the presentation compact. A fuller treatment, with worked numerical examples at each step, is available in the companion pedagogical document at docs/24-component-decomposition-pedagogy.md.

[17] where Y_{it} is the observed symptom score, BL_i is participant i 's

baseline score, the three components BR , PB , TV are non-negative reductions in symptom severity attributable to the three causal mechanisms identified in section 1, and ε_{it} is observation-level noise. The sign convention is that an increase in any component reduces Y ; all three components are zero at $t = 0$ by construction (no exposure yet to drug, no belief activation, no elapsed natural history) and grow over the duration of the trial.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

For participant i at trial timepoint t :

$$Y_{it} = \text{BL}_i - [BR_{it} + PB_{it} + TV_{it}] + \varepsilon_{it},$$

[18]

Each component is driven by its own time-on-state variable, making the components structurally distinct even when they overlap in calendar time.

- BR_{it} accumulates with time-on-drug; PB_{it} with time-on-belief scaled by expectancy $\eta(\text{phase})$; TV_{it} with calendar time from trial entry.
- The distinct time indices are the structural source of identifiability from phase-contrast data.
- All components are zero at $t = 0$ by construction and grow over the trial duration.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Each component is a participant-specific function of its own time-on-state variable. BR_{it} depends on time-on-drug, the cumulative exposure to active medication since drug initiation. PB_{it} depends on time-on-belief, modulated by a phase-specific expectancy factor $\eta(\text{phase})$ that captures the patient's confidence that they are receiving active treatment. TV_{it} depends on time-since-trial-entry, calendar time elapsed regardless of treatment status.

[19] rescaled so that $C(0) = 0$ and $C(t) \rightarrow m_C(1 - e^{-d_C t})$ as $t \rightarrow \infty$, for $C \in \{BR, PB, TV\}$. The three Gompertz parameters (m_C, d_C, r_C) are participant-specific random effects drawn from population distributions, and the participant-level heterogeneity in these distributions is the principal source of between-person response variability that the framework is designed to characterise.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Table 1: Principal psychometric and pharmacological terms used in the three-component decomposition.

Term	Meaning
BR	Biological response: pharmacological reduction in symptom score attributable to active drug.
PB	Placebo-belief response: reduction attributable to the patient’s belief that they are receiving treatment.
TV	Time-variant component: reduction (or increase) attributable to the natural-history trajectory of the disease.
$\eta(\text{phase})$	Expectancy multiplier for <i>PB</i> : 1 in open-label on-drug, ≈ 0.5 in blinded discontinuation, between 0 and 1 in open-label discontinuation.
BL_i	Participant i ’s baseline (pre-treatment) symptom score.
D_{it}	Continuous-decay drug indicator on the analysis side: 1 on-drug, $\exp(-\lambda t_{sd,it})$ off-drug.
$t_{sd,it}$	Time since drug discontinuation (off-drug only).
λ	Carryover decay rate, with half-life $t_{1/2} = \ln 2 / \lambda$.

Each component follows a modified Gompertz curve:

$$C(t) = m_C \exp(-d_C \exp(-r_C t)) - m_C \exp(-d_C),$$

[20]

The residual term has AR(1) structure and the three components are permitted to co-vary through cross-component covariances.

- Within-participant residuals follow $\text{Cov}(\varepsilon_{it}, \varepsilon_{i,t+h}) = \sigma_\varepsilon^2 \rho^{|h|}$.
- Within-time and across-time cross-component covariances permit participant-level heterogeneity in one component to co-vary with the others.
- This structure accommodates biologically plausible correlations, such as high-*BR* patients also tending toward high-*PB*.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

ε_{it} is a within-participant residual term with AR(1) serial-correlation structure $\text{Cov}(\varepsilon_{it}, \varepsilon_{i,t+h}) = \sigma_\varepsilon^2 \rho^{|h|}$. The three components are permitted to be jointly correlated through within-time and across-time cross-component covariances c_{cf1t} and c_{cfct} respectively, so that participant-level heterogeneity in any one component can co-vary with heterogeneity in the others.

The principal psychometric and pharmacological terms used below are summarised in Table 1.

3 The three components

[21]

622 **This section presents each component with its biological mo-**
623 **tivation, Gompertz parameterisation, and identifiability con-**
624 **ditions to establish the model as causally interpretable.**

- 625 • The three components are presented in turn: BR (biological re-
626 response), PB (placebo-belief), TV (natural history).
- 627 • Each component is grounded in independent biological literature
628 before its parametric form is stated.
- 629 • The section’s purpose is to demonstrate that the model sum re-
630 flects causal structure, not mathematical convenience.

631 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
632 *tempor incididunt ut labore et dolore magna aliqua.*

633 In this section we shall present the three components in turn, each
634 with its biological motivation, the modified Gompertz parameterisa-
635 tion, and the design contrasts that identify it. The treatment is or-
636 ganised so that the reader can see, by the end of this section, that
637 the formal model is a sum of three causally interpretable trajectories
638 rather than a mathematical convenience.

639 **3.1 Biological response (BR)**

640 [22]

641 **BR is the pharmacological target of regulatory approval and**
642 **biomarker stratification, independent of belief or natural his-**
643 **tory.**

- 644 • BR arises from chemical action on the body and is zero for inert
645 pills regardless of patient expectation.
- 646 • It is the estimand that approval decisions, dose selection, and
647 biomarker-guided stratification require.
- 648 • Separating BR from PB and TV is the central precision-medicine
649 problem the framework addresses.

650 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
651 *tempor incididunt ut labore et dolore magna aliqua.*

652 BR is the physiological reduction in symptom score attributable to the
653 active medication’s chemical action on the body. It is independent of
654 belief, suggestion, and natural history: a participant who somehow re-
655 ceived active drug without knowing it would still develop a measurable
656 BR , and a participant who took an inert sugar pill while believing
657 it was active drug would have $BR = 0$ no matter how strong their
658 expectation. Methodologically, BR is the quantity that drug regu-
659 lators and prescribers want: approval decisions, dose selection, and
660 biomarker-guided patient stratification are all framed around BR .

661 [23]

PK/PD considerations produce a saturating BR trajectory with a slow initial rise, a peak-growth phase, and a sharp drop to zero on discontinuation.

- Plasma steady state is reached within a few half-lives; pharmacodynamic integration (e.g. receptor adaptation) adds additional lag.
- For prazosin-PTSD, Raskind et al. (2013) calibrate the onset to several weeks of sustained exposure.
- The saturating shape motivates the Gompertz over linear or simple exponential alternatives.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The qualitative shape of the BR trajectory follows from pharmacokinetic and pharmacodynamic considerations. Drug concentration in plasma reaches steady state within a small multiple of the elimination half-life; brain concentration follows plasma with a short delay; the downstream pharmacodynamic response (reduction in symptom-relevant neural activity, in the case of prazosin's effect on noradrenergic tone during REM sleep) typically requires several days to weeks of sustained exposure to integrate fully⁴⁴. The result is a saturating curve: slow rise in the first week, steepest growth in the second to third week, plateau as receptor adaptation completes, and a decay back toward zero on discontinuation governed by the carryover half-life $t_{1/2} = \ln 2/\lambda$ rather than an instantaneous drop. The rate of this off-drug decay is the object of the carryover machinery used throughout this project; only in the limit of a short half-life relative to the measurement spacing does it approximate the sharp step assumed in the worked example of section 6.

[24]

The three Gompertz parameters (m_{BR}, r_{BR}, d_{BR}) capture ceiling, onset rate, and climb shape respectively, calibrated to prazosin-PTSD reference values.

- m_{BR} is the biological ceiling (maximum reduction at infinite exposure); r_{BR} the onset rate; d_{BR} the displacement shaping the early climb.
- Prazosin-PTSD calibration yields $m_{BR} \approx 11$, $r_{BR} \approx 0.42/\text{week}$, $d_{BR} \approx 5$.
- These are participant-specific random effects whose population distributions encode between-person response heterogeneity.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The modified Gompertz form $BR(t) = m_{BR} \exp(-d_{BR} \exp(-r_{BR}t)) -$

$m_{BR} \exp(-d_{BR})$ captures this profile compactly. The asymptote $m_{BR}(1 - e^{-d_{BR}})$ is the biological ceiling, the largest reduction the drug can produce at infinite exposure; the scale parameter m_{BR} equals that ceiling up to the factor $(1 - e^{-d_{BR}})$, which for the calibrated $d_{BR} \approx 5$ exceeds 0.99. The onset rate r_{BR} controls how rapidly the participant approaches the ceiling: high r_{BR} means a fast onset (effect plateaus in one to two weeks), low r_{BR} a slow climb (six to eight weeks). The displacement d_{BR} shapes the climb: high d_{BR} gives an early-onset curve that rises sharply at first, low d_{BR} a more gradual climb. For PTSD prazosin, calibrated values from the³⁴ reference dataset are $m_{BR} \approx 11$ (eleven points of nightmare-score reduction at saturation), $r_{BR} \approx 0.42$ per week (half-maximum at roughly four to five weeks of sustained exposure), and $d_{BR} \approx 5$ (moderately steep early climb).

[25]

The Gompertz is preferred over linear and simple-exponential alternatives because it captures the asymmetric slow-start saturation profile of chronic-condition pharmacodynamics.

- Linear $BR = \beta t$ is biologically implausible past the first few weeks due to unbounded growth.
- Exponential $BR = m(1 - e^{-rt})$ lacks the slow-start phase from receptor adaptation and gene-expression changes.
- The Gompertz interpolates between both limits and its three-parameter form avoids overfitting; sensitivity analysis is reported separately in 07-gompertz-evaluation.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Two alternative families that do not capture the asymmetric saturation profile are worth noting for contrast. A linear function $BR = \beta t$ climbs without bound and is biologically implausible past a few weeks. An exponential decay $BR = m_{BR}(1 - \exp(-rt))$ is acceptable for fast-onset drugs but lacks the slow-start phase that characterises most chronic-condition responses, where receptor adaptation, feedback regulation, and gene-expression changes produce a measurable lag. The Gompertz interpolates between exponential decay (high d limit) and linear-in- t growth (low d limit), and its three-parameter form is flexible enough to fit a wide range of pharmacodynamic profiles without overfitting. The Gompertz is a deliberate choice over the two parametric forms most familiar from pharmacokinetic-pharmacodynamic modelling. The sigmoid $E_{\max}$ (Hill) model⁴⁵ describes the dependence of effect on concentration (a dose-response curve), whereas $BR(t)$ here is a time-course of accumulating effect at fixed dosing; and the indirect-response (turnover) models that do describe effect time-courses are

specified through differential equations on a hidden response variable rather than in closed form. The modified Gompertz retains the slow-start, saturating shape those models produce while remaining a closed-form, three-parameter mean function that is tractable inside a mixed-effects likelihood. Doc 25 of the project documentation provides the full historical and biological context for the family choice and quantifies the sensitivity of pmsimstats inference to alternative families; that work is reported separately as 07-gompertz-evaluation. The headline result of that evaluation is reassuring for the present framework: in a 1{,}000-replicate-per-cell simulation that generated data under the logistic, hyperbolic-tangent, and piecewise-linear breakpoint families and analysed them under the standard linear-mixed model, the trajectory family was essentially irrelevant to the power of the biomarker-treatment interaction test (the family-to-family spread fell within two Monte Carlo standard errors of sampling noise) and to type I error (0.025 to 0.036 across null cells, with no family-specific inflation). The conclusions drawn here are therefore not artefacts of the Gompertz form, though the regimes the family cannot represent (cyclical or biphasic trajectories) remain genuine exclusions.

[26]

***BR* is identified by the blinded-discontinuation contrast, where drug removal collapses *BR* to zero while *PB* remains intact.**

- Within-participant on-drug versus off-drug contrasts are the primary identification mechanism.
- The blinded-discontinuation phase provides the cleanest contrast: *BR* drops, *PB* is preserved.
- This contrast does not require the analyst to model *PB* or *TV* explicitly to recover *BR*.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The *BR* component is identified by within-participant on-drug versus off-drug contrasts, formalised in section 4 below. The most direct contrast is the blinded-discontinuation phase: when a participant is silently switched to placebo while still believing they are on active drug, *BR* decays toward zero as the drug clears while *PB* remains roughly intact, and the difference between on-drug and off-drug response in that window is a relatively clean estimate of *BR* that does not require the analyst to model *PB* or *TV* explicitly. The decay is not instantaneous: with the carryover half-life $t_{1/2} = \ln 2/\lambda$ of the analysis-side decay term, residual *BR* persists into the early off-drug observations, so the contrast is clean only once the discontinuation window is long relative to $t_{1/2}$, and the drug state is best carried by

790 the continuous-decay indicator D_{it} of Table 1 rather than a hard on/off
791 indicator.

792 3.2 Placebo-belief response (PB)

793 [27]

794 ***PB* is the strict placebo response: improvement that per-**
795 **sists when the active drug is silently replaced, provided the**
796 **patient’s belief is unchanged.**

- 797 • *PB* is belief-dependent, not drug-dependent; it would survive a
798 silent switch to placebo under sustained belief.
- 799 • It is distinct from *BR* in causal origin: *BR* requires chemical
800 action; *PB* requires only belief.
- 801 • The definition is strict and operational, not vague or colloquial.

802 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
803 *tempor incididunt ut labore et dolore magna aliqua.*

804 *PB* is the reduction in symptom score that arises from the patient’s
805 belief that they are receiving active treatment, independent of what
806 they are actually receiving. That is, it is the placebo response in
807 its strict sense: the part of the improvement that would persist if
808 the active drug were silently replaced with an inert pill, provided the
809 patient continued to believe they were on the active drug.

810 [28]

811 ***PB* reflects neurobiologically real mechanisms triggered by**
812 **belief, not pharmacology; neuroimaging and genetic evidence**
813 **confirm it is a structured biological phenomenon.**

- 814 • Belief-driven endogenous opioid/dopamine release, HPA-axis
815 modulation, and top-down attentional shifts are all documented
816 mechanisms.
- 817 • Wager et al. neurologic pain signature dissociates pharmacologi-
818 cal from expectancy contributions via fMRI.
- 819 • COMT genotype predicts placebo responsiveness, confirming bi-
820 ological structure rather than measurement artefact.

821 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
822 *tempor incididunt ut labore et dolore magna aliqua.*

823 It is worth being precise about what *PB* is and is not. *PB* is not
824 ‘fake’ improvement; it is a phenomenon clinicians knowingly exploit, a
825 national survey of US internists and rheumatologists finding roughly
826 half report regular use of placebo treatments⁴⁶. As the introduction
827 sets out, the placebo response is mediated by well-characterised neu-
828robiological pathways^{38–41,47}: belief-driven release of endogenous opi-
829oids and dopamine, anticipatory hypothalamic-pituitary-adrenal reg-

830 ulation, and top-down attentional shifts, dissociated from pharmacol-
831 ogy by the⁴² neurologic pain signature⁴³ and shown to have genetic
832 correlates^{9,10}. What makes these mechanisms ‘placebo’ rather than
833 ‘drug’ is that they are triggered by belief rather than by pharmacol-
834 ogy.

835 [29]

836 ***PB* is identifiable from *BR* only when belief and treatment**
837 **allocation differ, as in the blinded-discontinuation phase.**

- 838 • The blinded-discontinuation phase is the canonical contrast: pa-
839 tients cannot tell whether they remain on active drug.
- 840 • Belief drops to a partial level ($\eta \approx 0.5$) because the patient’s
841 expectation is hedged, not absent.
- 842 • Designs without a phase that decouples belief from treatment
843 cannot separately identify *PB* from *BR*.

844 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
845 *tempor incididunt ut labore et dolore magna aliqua.*

846 For trial design, the relevant point is that *PB* is separately identi-
847 fiable from *BR* if and only if the trial includes a phase in which the
848 patient’s belief about treatment differs from the actual treatment they
849 are receiving. The blinded-discontinuation phase is the canonical such
850 phase: the patient is told that, at some unannounced point during
851 the window, they may be silently switched to placebo, and they can-
852 not reliably tell whether they are still receiving active drug. Belief in
853 treatment correspondingly drops to a partial level (multiplier roughly
854 0.5) because the patient’s expectation is hedged.

855 [30]

856 **The $\eta = 0.5$ blinded-phase expectancy is a modelling assump-**
857 **tion operationalising response expectancy, not a measured**
858 **quantity, and requires sensitivity analysis.**

- 859 • The construct is response expectancy in the sense of Kirsch
860 (1985); to our knowledge the residual expectancy under blinded
861 discontinuation has not been measured directly.
- 862 • Open-label placebo trials confirm clinically meaningful *PB*-only
863 responses even under explicit disclosure, bounding the residual
864 response away from zero.
- 865 • Sensitivity analyses varying η are mandatory before clinical con-
866 clusions depend on this single parameter.

867 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
868 *tempor incididunt ut labore et dolore magna aliqua.*

869 *PB* also follows a Gompertz curve, scaled by the expectancy factor:

$$PB(t \mid \text{phase}) = \eta(\text{phase}) \cdot m_{PB} \exp(-d_{PB} \exp(-r_{PB}t)) + (\text{offset to start at zero}).$$

Typical expectancy values are $\eta = 1.0$ in open-label on-drug, $\eta \approx 0.5$ in blinded discontinuation, and η between 0 and 0.5 in open-label placebo (the residual expectation a patient retains when told the previous active phase was beneficial). The quantity η operationalises response expectancy in the sense of⁴⁸, the patient's anticipation of benefit that the expectancy literature treats as a determinant of experienced symptom change. The choice of $\eta = 0.5$ for the blinded phase is a modelling assumption rather than a calibrated quantity: to our knowledge the residual expectancy a patient carries when uncertain about allocation has not been measured directly as a fraction of the open-label level, and we adopt 0.5 as a deliberately central placeholder. The available indirect evidence is nonetheless that the residual response is substantial. Open-label placebo trials^{35–37} demonstrate that a clinically meaningful response persists even when patients are explicitly told they are receiving an inert intervention, which bounds the residual belief-driven response away from zero. Because the value of η is not empirically pinned, sensitivity analyses varying it across its plausible range are mandatory before any conclusion is allowed to depend on this single number.

[31]

Kaptchuk et al. (2008) provide the structural precedent; the BR-PB-TV framework extends the same decomposition logic to the full treatment response.

- Kaptchuk decomposed the placebo response itself into observation, ritual, and relationship components empirically.
- The present framework applies analogous decomposition logic at the level of drug, placebo, and natural history.
- Both frameworks reject the 'expectancy lump' model in favour of identifiable additive parts.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The decomposition framework as proposed here has a structural analogue in the placebo-research literature.⁵ decomposed the placebo response itself into observation, ritual, and patient-practitioner relationship components in a randomised trial of acupuncture-style placebo for irritable bowel syndrome, demonstrating empirically that the placebo response is not a single 'expectancy' lump but a sum of identifiable parts. The present BR-PB-TV framework is analogous in spirit but operates at the level of the full treatment response (drug, placebo,

909 natural history) rather than within the placebo response alone.

910 [32]

911 ***PB* variance scales with η , not just its mean; ignoring this**
912 **heteroscedastic structure miscalibrates standard errors and**
913 **downstream biomarker tests.**

- 914 • Phase-dependent variance follows $\text{SD}(PB \mid \text{phase}) = \eta(\text{phase}) \cdot$
915 $\sigma_{PB}^{\text{baseline}}$.
- 916 • Open-label phases have high between-patient *PB* variance;
917 blinded phases have low variance.
- 918 • Ignoring heteroscedasticity deflates standard errors in open-label
919 and inflates them in blinded phases, with consequences for CI
920 coverage and biomarker test calibration.

921 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
922 *tempor incididunt ut labore et dolore magna aliqua.*

923 A subtle property of *PB* is that its variance also scales with η ,
924 not just its mean. When η is high, there is substantial variability
925 across patients in how strongly the placebo response is expressed;
926 high-suggestibility patients show large *PB*, low-suggestibility pa-
927 tients show essentially none. When η is low (blinded), the placebo
928 response is uniformly attenuated and the patient-to-patient vari-
929 ance shrinks accordingly. The model therefore parameterises
930 $\text{SD}(PB \mid \text{phase}) = \eta(\text{phase}) \cdot \sigma_{PB}^{\text{baseline}}$. Ignoring this heteroscedastic
931 structure deflates standard errors in the open-label phase and inflates
932 them in the blinded phase, with downstream consequences for
933 confidence-interval coverage and for the calibration of any biomarker
934 test built on top of the analysis.

935 3.2.1 Additivity, subadditivity, and the BR-PB interaction

936 [33]

937 **Additivity is a modelling assumption, not biological fact; the**
938 **section examines whether a $BR \times PB$ interaction term is war-**
939 **ranted.**

- 940 • The formal model writes total response as a sum; this section
941 asks whether the *BR-PB* interaction is zero.
- 942 • The question has inferential consequences for biomarker valida-
943 tion and for what the phase-augmented analysis captures.
- 944 • The treatment is preparatory for the generalised interaction
945 model introduced later in this section.

946 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
947 *tempor incididunt ut labore et dolore magna aliqua.*

948 In this section we shall broach a question that the formal model leaves
949 open: whether the three components contribute additively to symptom
950 reduction, or whether there is a meaningful interaction between BR
951 and PB that the additive specification misses.

952 [34]

953 **The additive model may be wrong: published literature doc-**
954 **uments subadditivity between pharmacological and placebo-**
955 **belief effects in some clinical contexts.**

- 956 • Subadditivity means the combined $BR + PB$ effect is less than
957 the sum of each component measured separately.
- 958 • Evidence from Kaptchuk et al. and Benedetti et al. shows the
959 phenomenon is context-dependent, not universal.
- 960 • Formal test: $\partial Y / \partial (BR \cdot PB) > 0$ in a saturated mean-structure
961 model.

962 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
963 *tempor incididunt ut labore et dolore magna aliqua.*

964 The decomposition as written in section 2 treats the three components
965 as additive contributors to symptom reduction:

$$Y_{it} = BL_i - (BR_{it} + PB_{it} + TV_{it}) + \varepsilon_{it}.$$

966 The additivity is a modelling assumption, not a fact about biology,
967 and it is worth stating that the assumption may be wrong in direc-
968 tions the framework can absorb and in directions it cannot. Additiv-
969 ity is also scale-dependent in the sense emphasised by⁴⁹: components
970 that combine non-additively on the raw symptom scale may combine
971 additively on a transformed (for example logarithmic) scale, so an ap-
972 parent $BR \times PB$ interaction can be partly an artefact of the chosen
973 response metric rather than a biological fact. That belief and pharma-
974 cology interact rather than act in isolation is demonstrated directly
975 by⁵⁰, who showed with neuroimaging that positive treatment expecta-
976 tion roughly doubled the analgesic benefit of the opioid remifentanyl
977 while negative expectation abolished it, so the two channels are not in
978 general separable by simple addition. More specifically, the published
979 literature on placebo-drug interactions has accumulated evidence that
980 the combined effect of receiving an active drug and believing one is
981 receiving an active drug is, in some clinical contexts, less than the sum
982 of the two components measured separately^{5,38}. The phenomenon is
983 typically termed subadditivity: $\partial Y / \partial (BR \cdot PB) > 0$ when both compo-
984 nents are present and their interaction term is included in a saturated
985 mean-structure model.

986 [35]

987 **Three candidate mechanisms for subadditivity are shared**
988 **neural substrate, scale ceiling effects, and belief-dynamics**
989 **adaptation.**

- 990 • Shared pain-modulation circuits produce saturating occupancy
991 when both BR and PB activate them simultaneously.
- 992 • Scale ceiling effects prevent registration of additional symptom
993 reduction once $BR + PB$ saturates the measurement range.
- 994 • Belief-dynamics adaptation: unexpectedly large BR can recal-
995ibrate patient expectation downward, attenuating ongoing PB .

996 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
997 *tempor incididunt ut labore et dolore magna aliqua.*

998 Three mechanisms have been proposed in the placebo-research litera-
999 ture to explain subadditivity. First, a shared neural substrate: if the
1000 descending pain-modulation circuits that mediate placebo analgesia
1001 overlap with the circuits modulated by the active drug, simultaneous
1002 activation by both mechanisms produces less marginal benefit than
1003 activation by either alone, in a saturating-occupancy sense. Second,
1004 ceiling effects on the symptom scale: a clinical scale that caps at zero
1005 (no symptoms) cannot register additional reduction once a participant
1006 has reached that cap, so any participant whose $BR + PB$ already sat-
1007 urates the scale will register no further reduction even if BR or PB
1008 would, in isolation, predict more. Third, belief-dynamics adaptation:
1009 the patient's expectation may adapt downward when the active drug
1010 produces a strong response that exceeds the expected effect, reducing
1011 the ongoing PB contribution; equivalently, unexpectedly large BR
1012 may attenuate PB via expectation recalibration.

1013 [36] where $\gamma < 0$ encodes subadditivity, $\gamma = 0$ recovers the additive
1014 form, and $\gamma > 0$ encodes superadditivity, is the natural extension.
1015 The interaction coefficient γ is identifiable from the same trial-design
1016 contrasts that identify the marginal components: open-label phases
1017 supply $BR \cdot PB$ products at high η , blinded phases supply $BR \cdot PB$
1018 products at attenuated η , and the γ coefficient is estimated from the
1019 difference between the two regimes once the marginal components are
1020 accounted for.

1021 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
1022 *tempor incididunt ut labore et dolore magna aliqua.*

1023 Whichever mechanism dominates in a given clinical context, the subad-
1024 ditivity question is empirical and design-dependent. Within the BR-
1025 PB-TV framework, subadditivity manifests as a non-zero coefficient
1026 on a $BR \times PB$ interaction term in the conditional-mean specification.
1027 The additive form treats this interaction as zero by construction; a
1028 generalised form

$$Y_{it} = \text{BL}_i - [BR_{it} + PB_{it} + TV_{it} + \gamma \cdot BR_{it} \cdot PB_{it}] + \varepsilon_{it},$$

[37]

Subadditivity biases biomarker predictive accuracy estimates and is partly absorbed by the phase-augmented analysis even without explicit γ estimation.

- Under subadditivity, the biomarker-by- BR association is partly cancelled, inflating apparent predictive accuracy in the additive model.
- Phase designs attenuating PB (blinded, crossover) push the effective biomarker-treatment estimate closer to the true biomarker-by- BR value.
- The phase-augmented analysis in section 5 absorbs subadditivity as a side benefit.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

For trial design, two consequences follow. First, a biomarker that predicts BR in an additive-decomposition analysis may overstate or understate predictive accuracy if the true mechanism is subadditive: under subadditivity, the biomarker-driven increase in BR is partly cancelled by the $\gamma \cdot BR \cdot PB$ interaction, and the marginal biomarker-treatment association is smaller than the biomarker-by- BR association. Second, trial designs that attenuate PB in some phases (blinded discontinuation, crossover) implicitly down-weight the subadditivity contribution to the response in those phases, which means that the effective biomarker-treatment association estimated from those phases is closer to the biomarker-by- BR association than is the open-label estimate. The phase-augmented analysis recommended in section 5 below therefore has the side benefit of partly absorbing subadditivity, even when γ is not estimated explicitly.

[38]

A forthcoming simulation study will quantify how bias in the biomarker-treatment interaction scales with γ ; the pre-registration predicts linear scaling.

- Bias in $\hat{\beta}_{bm:D}$ is expected to scale linearly with γ at small $|\gamma|$.
- Under subadditivity, bias sign is opposite to the true biomarker-by- BR effect; under superadditivity it is aligned.
- Empirical quantification is deferred to the production simulation programme in section 7.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A simulation study quantifying the magnitude of bias in the biomarker-treatment interaction estimate when the additive analysis is fitted to data generated under non-zero γ is forthcoming as part of the simulation programme described in section 7. The pre-registration is that bias scales linearly with γ at small $|\gamma|$, with the sign opposite to the true biomarker-by- BR effect under subadditivity and aligned with it under superadditivity.

[39]

Non-zero $\hat{\gamma}$ does not establish true subadditive interaction; latent-class structure with class-correlated components produces the same marginal pattern.

- If class membership correlates with both BR and PB strength, the marginal covariance of $BR \cdot PB$ mimics a direct interaction term.
- This confounding is invisible from the lumped-analysis perspective but is identifiable in a class-aware analysis.
- Apparent subadditivity can arise from latent population heterogeneity with no direct BR - PB interaction in the class-conditional DGP.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A subtle but inferentially important point merits attention here: a non-zero $\hat{\gamma}$ recovered from an additive-extension analysis does not, on its own, establish that the underlying mechanism is subadditive interaction in the strict sense. The same marginal pattern can arise from latent population structure that the analyst is not modelling. If the trial population is a mixture of treatment-responders and non-responders (in the latent-class sense developed at length in [analysis/report/03-latent-class-mixture-application/](#)), and class membership correlates with both BR strength and PB strength (responder-class participants have both higher BR and higher PB on average), then the marginal expected value of $BR \cdot PB$ over the unobserved class label has a non-zero covariance term whose magnitude depends on the gating function $\pi(B)$. From the lumped-analysis perspective this looks exactly like a γ -weighted $BR \times PB$ interaction even though no such direct interaction was specified in the class-conditional DGP. That is, the latent-class structure is one mechanism that produces apparent subadditivity in a one-component analysis, regardless of whether the underlying biology has a true direct BR - PB interaction.

[40]

Distinguishing true subadditivity from latent-class confounding requires either a class-aware analysis or the phase contrasts the BR-PB-TV framework provides.

- The two mechanisms produce qualitatively similar marginal effects but are inferentially distinct.
- $\hat{\gamma}$ should be reported alongside a class-aware sensitivity analysis, not as standalone evidence for mechanistic subadditivity.
- The latent-class paper's `pb_class_correlation` simulation axis is the direct test; cross-paper synthesis is the cleanest disentanglement.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The two mechanisms, namely true subadditive interaction and latent-class structure with class-correlated components, produce qualitatively similar marginal effects but are inferentially distinct. Distinguishing them requires either a class-aware analysis that explicitly models the latent population structure (as the latent-class paper develops at length), or a phase contrast that separates the mechanisms by exploiting the differential expectancy modulation of *PB* that the BR-PB-TV framework provides. Under the additive-extension analysis fitted to mixture-DGP data with class-correlated components, $\hat{\gamma}$ is identified up to confounding with the class structure, and the analyst should report $\hat{\gamma}$ alongside a class-aware sensitivity analysis rather than as standalone evidence for mechanistic subadditivity. The Study 1 simulation in the latent-class paper includes a `pb_class_correlation` axis that varies $\text{Cor}(BR, PB)$ across class labels; that axis is the direct simulation-side test of the latent-class generation of apparent subadditivity, and the cross-paper synthesis between the two papers is the cleanest way to disentangle the two mechanisms.

3.3 Time-variant natural-history component (TV)

[41]

***TV* is the untreated counterfactual trajectory; it is participant-specific, structured, and can be positive, negative, or near-zero depending on disease class.**

- *TV* is the symptom change that would have occurred without treatment, integrated over the trial duration.
- It varies in sign across disease classes and across patients within a class.
- It is independent of treatment status and belief, distinguishing it structurally from *PB*.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

TV is the reduction in symptom score that would have happened even if the patient had received no treatment at all. It is the participant's natural-history trajectory, integrated over the duration of the trial. For some participants TV is positive (symptoms naturally improve over the trial because of life events, therapy, time-in-condition, or regression to the mean); for others TV is negative (symptoms naturally worsen because of trauma anniversaries, accumulating stress, or progressive underlying disease).

[42]

TV is a structured participant-specific signal, not noise; its distinguishing feature from PB is independence from treatment status and belief.

- Unlike ε_{it} , TV is participant-level (slow trajectory) and requires explicit modelling, not absorption into noise.
- Off-drug baselines and follow-up windows are the primary data sources for estimating TV .
- TV is modulated by neither treatment assignment nor patient belief, whereas PB is modulated by both.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Like PB , TV is not noise. It is a structured participant-specific signal that the model can in principle estimate from data, given sufficient repeated measurements and a trial design that includes a long enough off-drug baseline or follow-up window. What distinguishes TV from PB is its origin: TV is independent of treatment status and of belief about treatment, whereas PB is modulated by both.

3.3.1 When does natural history actually exist?

[43]

The shape of TV depends on disease class, trial timescale, and selection process; none of these can be assumed and all must be recovered empirically.

- The assumption that subjects improve over time is true for some conditions and systematically wrong for others.
- Disease class, trial timescale, and enrolment selection jointly determine the sign and magnitude of TV .
- Recovering the empirical shape of TV is one of the decomposition's primary objectives.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A common simplifying assumption in trial analysis is that subjects will improve over time, with or without treatment. This is true for some conditions and problematic for others. The shape of TV depends on the underlying disease class, the trial timescale, and the selection process by which participants entered the trial. None of these can be assumed without thought, and the empirical shape of TV is one of the things the decomposition is designed to recover.

[44]

TV sign varies systematically by disease class: positive for acute recovery, near-zero for stable chronic, negative for progressive, and structured for cyclical conditions.

- Acute conditions (post-operative pain, acute back pain) resolve intrinsically; substantial TV -improvement appears in the placebo arm.
- Progressive conditions (Alzheimer's, ALS) require explicit negative-trajectory modelling; apparent treatment effect is slowing of progression.
- Cyclical/triggered conditions (PTSD with trauma anniversaries, relapsing-remitting MS) require covariates beyond the Gompertz form.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Acute conditions tend toward TV -positive: a sprained ankle, a post-operative pain syndrome, an episode of acute back pain all resolve on a timescale of days to weeks even without intervention, and a trial of an analgesic in such a setting will see substantial TV -improvement in the placebo arm. Chronic stable conditions (treated hypertension, controlled chronic asthma, long-standing PTSD without active triggers) tend toward TV -zero on the population mean but with substantial individual variance. Chronic progressive conditions (Alzheimer's disease, ALS, late-stage congestive heart failure, many cancers) tend toward TV -negative: a trial of a neuroprotective agent in early Alzheimer's must explicitly model the negative natural-history trajectory, since the apparent treatment effect is the slowing of progression, not absolute improvement. Cyclical and triggered conditions (PTSD with trauma anniversaries, seasonal-pattern major depression, relapsing-remitting multiple sclerosis) show structured TV that is neither monotone nor stationary, and may require explicit covariates beyond the Gompertz form.

[45]

Selection-induced TV is real and rarely zero even for stable chronic conditions, because enrolment coincides with peak severity triggering help-seeking.

- Regression to the mean alone contributes several points of apparent improvement in early trial weeks.
- Hrobjartsson et al. systematic reviews across 200+ trials show much of conventional placebo response is indistinguishable from natural-history drift.
- Kirsch (2008) and Locher (2015) demonstrate the clinical consequences in antidepressant and late-life depression trials.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Selection-induced TV is real and rarely zero, even when the underlying condition is stable. Patients enrol at moments of clinical engagement that often coincide with recent worsening, so regression to the mean alone produces several points of apparent TV -improvement in the early weeks of a trial. As the introduction documents, the systematic-review evidence of¹¹, ¹², and¹³ (the last covering 202 trials across 60 conditions) and the antidepressant analyses of⁶ and⁷ show that much of what is conventionally attributed to placebo response is indistinguishable from regression to the mean and natural-history drift.

[46]

Trial participation itself induces TV -improvement through Hawthorne, structured-visit, and diary-completion effects, independent of pharmacology or placebo.

- Hawthorne, structured-visit, and diary-completion effects each contribute symptom improvements attributable to participation, not treatment.
- These contributions are absorbed into TV in the decomposition, not into BR or PB .
- Ignoring these participation-induced effects inflates the apparent pharmacological and placebo components.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Trial participation itself also induces TV : the Hawthorne effect, the structured-visit effect, and the diary-completion effect appear in the data as TV -improvement that is independent of any pharmacological or placebo response.

3.3.2 Can we assume no TV effect?

[47]

Assuming $TV = 0$ is rarely safe in chronic psychiatric or pain trials; the conservative default is to estimate TV from the data.

- Zero- TV assumption is valid only when all four conditions hold: stable disease, short trial, non-severity-peaked recruitment, no ancillary care.
- For most chronic conditions at least one condition fails; the assumption should be tested, not asserted.
- Statistical efficiency lost by estimating TV is smaller than the bias introduced by assuming it zero.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

But how safe is it to assume TV is zero? The short answer is: rarely safe, and never safe without empirical support. The longer answer is that TV may be assumed zero only if the disease is in a stable chronic phase, the trial is short relative to the disease timescale, the recruitment process does not favour participants at peak severity, and the trial protocol does not introduce ancillary care that itself improves symptoms. For most chronic trials in psychiatric or pain conditions, at least one of these conditions fails. The conservative default is to estimate TV from the data, even at the cost of some statistical efficiency, rather than assume it is zero.

[48]

The decomposition makes the zero- TV assumption testable and shows that lumping TV into residuals produces autocorrelated, heteroscedastic residuals that bias BR and PB standard errors.

- $\hat{m}_{TV}$ with a variance estimate from the model provides an empirical test of the zero- TV assumption.
- Participant-level TV is structured and correlated with the on/off-drug schedule; it is not independent noise.
- Lumping TV into ε produces autocorrelated within-participant residuals and heteroscedastic cross-participant residuals.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The decomposition makes this assumption testable. A model that includes a TV component will return $\hat{m}_{TV}$ with a variance estimate; if both are small relative to $\hat{m}_{BR}$ and $\hat{m}_{PB}$, the analyst has empirical support for the simplifying assumption. If $\hat{m}_{TV}$ is meaningfully large, the assumption was wrong, and the trial needs the TV component to produce unbiased inference about the others. Setting TV to zero and lumping it into ε is worse than estimating it: the participant-level TV

is structured (each participant has their own slow trajectory) and not independent of the on-drug-versus-off-drug schedule, so lumping it into noise produces residuals that are autocorrelated within participants and heteroscedastic across them. The resulting standard errors on the BR and PB estimates are wrong in both directions: too small for participants whose TV is small, too large for participants whose TV is large. Estimating TV explicitly fixes both problems simultaneously.

3.3.3 Why a Gompertz, not a linear time covariate?

[49]

A linear week covariate is insufficient for TV because it enforces linearity, ignores participant heterogeneity, and cannot distinguish TV from BR without the explicit decomposition.

- A **week** covariate constrains each additional week to contribute a fixed natural-history change, misrepresenting curved or saturating disease trajectories.
- A single coefficient estimates a population-mean trajectory, ignoring participant-level heterogeneity in natural history.
- Without the decomposition and design contrasts, **week** absorbs both BR and TV , making them inseparable.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A natural objection is that a model with **week** as a covariate already controls for time. But how adequate is such a control? Three reasons suggest it is not enough. First, a **week** covariate enforces a linear trajectory, constraining each additional week to contribute a fixed amount of natural-history change; real disease trajectories curve, plateau, or accelerate, and the Gompertz form captures the S-shaped or saturating patterns linear time cannot. Second, a single **week** coefficient estimates a population-mean trajectory that ignores the participant-level heterogeneity in natural history. Third, BR and TV both evolve smoothly over the same trial timescale, and a **week** covariate cannot tell them apart from observational data alone; the design contrasts discussed in section 4 are what supply identifiability, and without the explicit decomposition the **week** coefficient absorbs both.

4 Identifiability and trial-design contrasts

[50]

Identifiability is supplied by trial design; each phase of the hybrid design contributes a different combination of compo-

1351
1352
1353
1354
1355
1356
1357
1358
1359
1360
1361
1362
1363
1364
1365
1366
1367
1368
1369
1370
1371
1372
1373
1374
1375
1376
1377
1378
1379
1380
1381
1382

nents to the observed trajectory.

- The decomposition is mathematically specified but requires trial-design contrasts to be statistically identified.
- This section characterises the hybrid open-label-plus-blinded-discontinuation design as the primary identifiability vehicle.
- Each of the three phases supplies a distinct mixture of BR , PB , and TV .

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The mathematical decomposition above is meaningless unless the data contain enough information to pin down each component separately. Trial design is what supplies that information. In this section we shall characterise the hybrid open-label-plus-blinded-discontinuation design, used as the worked example throughout this paper, and show how each of its three phases contributes a different combination of components to the observed trajectory.

Phase	Drug?	Belief?	Components present
Open-label on drug	Yes	Full	$BR + PB(\eta = 1) + TV$
Blinded discontinuation	Hidden	Hedged	$BR \cdot \mathbf{1}_{\text{on drug}} + PB(\eta = 0.5) + TV$
Open-label crossover	Either	Knows	$BR \cdot \mathbf{1}_{\text{on drug}} + PB(\eta = 1) + TV$

[51]

Three phase-by-allocation contrasts isolate PB , BR , and TV respectively; the LME model combines them but the inference is driven by design.

- Open-label versus blinded contrast yields approximately $0.5 \cdot PB$, isolating the belief component conditional on a known blinded-phase η .
- Blinded on-drug versus blinded placebo yields BR once drug carryover has elapsed (since PB and TV match across arms within the blinded phase).
- Off-drug phases yield TV plus residual off-drug PB , with no active BR contribution beyond carryover.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Three phase-by-allocation contrasts then identify the three components. The open-label versus blinded contrast (within drug) yields

$PB(\eta = 1) - PB(\eta = 0.5) \approx 0.5 \cdot PB$, which isolates the belief component, but only conditional on a known blinded-phase expectancy η : from a two-phase design just the product $\eta \cdot m_{PB}$ is identified, and separating η from m_{PB} requires the open-label placebo phase. The blinded on-drug versus blinded placebo contrast yields BR once drug carryover has elapsed, since PB and TV match across allocation arms within the blinded phase. All phases off drug yield TV plus any residual off-drug PB , with no active BR contribution beyond carryover and PB at the off-drug expectation level. The $\mathbf{1}_{\text{on drug}}$ indicator in the phase table above is shorthand for the continuous-decay drug state D_{it} defined in Table 1.

This conditioning on η is the framework's principal identification limitation and deserves emphasis rather than a parenthetical. With a two-phase (open-label and blinded) design the data identify only the product $\eta \cdot m_{PB}$, so PB is recovered up to the assumed blinded-phase expectancy; the working value $\eta \approx 0.5$ is a modelling assumption, not an estimate, and absent a belief or expectancy measurement it is not falsifiable from the outcome data alone. The identified set for m_{PB} is accordingly the interval $\{\eta \widehat{m_{PB}} / \eta : \eta \in (0, 1]\}$, which collapses to a point only when η is pinned down. Two design features close the gap: an open-label placebo phase, which supplies a second equation separating η from m_{PB} , and direct measurement of belief, as in the open-hidden and balanced-placebo designs^{51,52} that manipulate or record the patient's treatment expectation. We therefore recommend that any trial relying on the decomposition for a PB -versus- BR attribution either include such a phase or collect an expectancy measure, and that analyses report sensitivity of the attribution to η across its plausible range.

[52]

Designs lacking specific phases lose identifiability of the corresponding component; the decomposition machinery is only as useful as the design allows.

- A pure open-label design cannot separate BR from PB because both are present at full strength throughout.
- A pure blinded design cannot estimate $\eta = 1$ PB because no phase achieves full-strength belief.
- Trial designers should verify each component has at least one phase in which it is uniquely present before committing to the design.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A trial that lacks any of these phases loses the corresponding identifiability. A pure open-label design (no blinded discontinuation) can-

not separate BR from PB , because both are present at full strength throughout. A pure blinded design (no open-label phase) cannot estimate $\eta = 1$ PB , because no phase has belief at full strength. The decomposition machinery is useful only to the extent that the design supplies identifiability, and trial designers should evaluate any proposed design by checking that each component has at least one phase or contrast in which it is uniquely present.

[53]

Design-contrast identifiability is necessary but not sufficient: the three Gompertz trajectories are collinear, so joint estimation is credible only under specific design and sample-size conditions.

- The components share an onset shape and overlap in calendar time, so their parameters can be near-collinear even when the identifying contrasts are present.
- Joint recovery requires a phase in which each component varies while the others are held fixed, and an off-drug window long enough for TV to separate from a decaying BR .
- The recoverable region of the design-and-sample-size space is mapped by the pre-registered identifiability diagnostics, not asserted.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Identifiability in the design-contrast sense above is necessary but not sufficient for stable estimation at finite sample sizes. The three Gompertz trajectories share an onset shape and overlap in calendar time, so their parameters can be collinear even when the contrasts that identify them in principle are present. Two conditions make the joint estimation credible: the design must supply at least one phase in which each component varies while the others are held fixed (the blinded-placebo contrast for BR , the open-label-versus-blinded contrast for PB , and an off-drug window for TV), and that off-drug window must be long enough relative to the carryover half-life for TV to separate from a still-decaying BR . Where these conditions are weak the components are jointly identified only up to a near-collinear direction, and the symptom of that is the rank-deficiency the pilot of section 7 encountered when the full phase-augmented formula was fitted at $N = 35$. We therefore do not assert that the three components are always recoverable. The pre-registered identifiability diagnostics of section 7 (convergence rates of the component-matched fit, variance-inflation factors on the component-specific Gompertz parameters, and the sensitivity of the BR estimate to the assumed $\eta(\text{phase})$ multipliers) are the means by which the recoverable region of the design-and-sample-size space is

to be mapped, and the framework’s participant-level claims should be read as conditional on falling inside that region.

5 Analysis-side methodology

[54]

Four constraints make a component-matched analysis impractical at trial-relevant sample sizes; a phase-augmented LME captures the important part of the gap at minimal cost.

- The question of whether to mirror the DGP in the analysis model is reasonable; the section addresses it directly.
- The four constraints are identifiability, degrees of freedom, convergence fragility, and marginal estimand robustness.
- A phase-by-treatment indicator is the recommended modest extension that captures BR-versus-PB sensitivity without the full parametric cost.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In this section we shall take up the practical question of what analysis model to use. A reasonable reaction to the formal model in section 2 is the following: if we are simulating data with three components plus noise, why is the analysis model written as $Sx \sim \text{bm} + \text{t} + \text{Dbc} + \text{bm}:\text{Dbc}$ with a single random intercept and AR(1) residual correlation? Should we not include three parametric Gompertz components in the analysis, mirroring the DGP, so that the analyst sees the same structure that generated the data? The question is reasonable. Four constraints make a component-matched analysis impractical for most aggregated N-of-1 trial sizes, and a modest extension (a phase-by-treatment indicator) captures the practically important part of the gap at minimal cost.

5.1 Four reasons not to match the DGP

[55]

The first constraint is identifiability: forcing a three-component model on a design with insufficient contrasts produces unidentified parameters, not component estimates.

- A pure open-label trial cannot distinguish *BR* from *PB*; both rise from zero together on treatment initiation.
- Fitting three components to a one-contrast design produces implausibly tight standard errors that misrepresent inferential precision.

- The analyst must audit whether the design supplies the necessary contrasts before specifying a component-matched model.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The first reason is identifiability. The trial design must supply the contrasts that identify each component, and many designs supply only part of it. The decomposition is identified by phase contrasts (open-label versus blinded for PB , on-drug versus blinded-placebo for BR , off-drug timepoints for TV). A pure open-label trial has none of these contrasts; in such a trial BR and PB both rise from zero on treatment initiation and are statistically indistinguishable from each other, regardless of how the analyst writes the model. Forcing an analysis to fit three components when the design supplies contrasts for only one or two of them does not recover the missing components; it produces unidentified or weakly identified parameter estimates with implausibly tight standard errors that misrepresent the analyst's actual inferential precision.

[56]

The second constraint is degrees of freedom: a component-matched analysis adds approximately a dozen non-linear parameters, inflating moderation estimand variance at $N \in [30, 150]$.

- Gompertz BR (3), PB (3 + expectancy multiplier), and TV (3) each require participant-specific random effects on the maxima.
- At trial-relevant N , the variance of $\hat{\beta}_{bm:D}$ is dominated by participant and timepoint counts.
- Additional parametric structure typically inflates moderation estimand variance more than it reduces bias.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The second reason is degrees of freedom. Each component term costs degrees of freedom that come out of the moderation estimand. The biomarker-moderation question (the core scientific aim of the predictive-biomarker validation programme) is asked through the $bm:D_{bc}$ term. A component-matched analysis adds a Gompertz BR (three parameters), a Gompertz PB (three parameters plus a phase-specific expectancy multiplier), and a Gompertz TV (three parameters), each with participant-specific random effects on the maxima. That is on the order of a dozen new parameters, half of them on a non-linear scale. At trial sample sizes of $N \in [30, 150]$, the variance of $\hat{\beta}_{bm:D}$ is dominated by participant and timepoint counts, and adding parametric structure to absorb residual confounding

usually inflates the moderation estimand's variance by more than it reduces bias.

[57]

The third constraint is convergence fragility: component-matched nonlinear mixed models converge below 90% at trial-relevant N , contaminating aggregated analyses.

- Simultaneous Gompertz BR , PB , and TV fitting requires `nlme::nlme`, Bayesian hierarchical, or structural-equation approximations.
- Each is more sensitive to starting values, prone to local optima, and likely to fail on individual trial datasets than the LME analogue.
- Non-converging cells at below-90% convergence rates contaminate power and bias estimates in simulation studies.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The third reason is convergence fragility. Component-matched analyses are non-linear in the parameters and their convergence is problematic in practice. Fitting Gompertz BR , PB , and TV simultaneously requires a non-linear mixed model (`nlme::nlme` with self-starting Gompertz forms), a Bayesian hierarchical fit, or a structural-equation approximation. Each is more sensitive to starting values, more prone to local optima, and more likely to fail to converge on a single trial dataset than the linear-mixed analogue. At trial-relevant N , the convergence rate of the component-matched fit is typically below 90%, and the non-converging cells contaminate any aggregated analysis.

[58]

The fourth constraint is estimand robustness: the `bm:Dbc` marginal interaction is the right estimand for the biomarker-validation question regardless of the BR - PB split.

- Regulators want to know whether the biomarker predicts treatment response under trial conditions, not a mechanistically decomposed prediction.
- The marginal `bm:Dbc` interaction captures this question whether the moderation is biological or expectation-mediated.
- Component decomposition adds mechanistic precision but is not required for the core biomarker-validation decision.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The fourth reason is that the moderation estimand is sufficiently robust to the BR-versus-PB split for the most common scientific question. For predictive-biomarker validation, the regulator wants to know whether the biomarker predicts treatment response under the trial conditions, not a clean pharmacological-versus-placebo decomposition of that prediction. The simple $bm:D_{bc}$ interaction is the right marginal estimand for the validation question whether the moderation is biological or expectation-mediated.

[59]

The inferential target is stated as an estimand: the population biomarker-by-treatment interaction with explicit intercurrent-event handling, while the pharmacological-component target is more ambitious and inherits the heterogeneity-identifiability caveat.

- The $bm:D_{bc}$ estimand is defined by its population, treatment contrast, endpoint, summary, and intercurrent-event strategy in the ICH E9(R1) sense.
- Dropout is handled by a treatment-policy strategy and residual carryover by the modelled exponential decay.
- Recovering a participant-level pharmacological interaction inherits Senn’s caveat that true heterogeneity is separable from within-person variation only under repeated phase contrasts.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

It is worth stating the inferential target of the analysis in the vocabulary of the estimands framework⁵³, because the biomarker-validation question is easy to pose ambiguously. The primary estimand is the population-level biomarker-by-treatment interaction, the $bm:D_{bc}$ coefficient: its population is the trial-eligible patient population, its treatment contrast is on-drug versus off-drug exposure summarised through the continuous indicator D_{bc} , its endpoint is the symptom trajectory, its population-level summary is the interaction coefficient, and its intercurrent events are handled by a treatment-policy strategy for dropout and by the modelled exponential decay for residual carryover. The pharmacological-component target β_{bm}^{BR} is a more ambitious estimand: it is the biomarker-by- BR interaction that would obtain with the belief-driven channel held at its off-drug level, identified here through the phase contrasts of section 4 rather than by assumption. We flag explicitly that recovering a participant-level interaction of this kind inherits the difficulty Senn has emphasised for individual treatment effects^{49,54}: a genuine participant-by-component signal is separable from participant-by-period and within-person random variation only when the design supplies repeated phase contrasts,

and the framework should be read as estimating these quantities under that condition rather than as dissolving the well-known problem of identifying individual effects from finite within-person data.

5.2 A small extension that usually pays off: phase indicators

[60]

The phase-augmented LME uses `Dbc:phase` and `bm:Dbc:phase` to test whether drug and biomarker-by-drug effects attenuate under blinding.

- `phase` absorbs systematic differences in *PB* expectancy across trial phases.
- `Dbc:phase` detects whether the drug effect shrinks under blinding, a signature of *PB*-mediation.
- `bm:Dbc:phase` detects whether the biomarker-by-treatment interaction shrinks under blinding, a signature that it is partly biomarker-by-*PB*.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The smallest useful step toward component-aware analysis, without paying the costs above, is to add a phase indicator and its interaction with the drug indicator. Schematically,

```
Sx ~ bm + t + Dbc + phase + Dbc:phase
      + bm:Dbc + bm:Dbc:phase
      + (1 | ptID), correlation = corCAR1(...)
```

Here `phase` is a categorical covariate (open-label, blinded, post-crossover) that absorbs the systematic differences in *PB* expectancy across phases. `Dbc:phase` tests whether the apparent drug effect shrinks under blinding, a signature that the open-label drug effect was partly *PB*-mediated. `bm:Dbc:phase` tests whether the moderation itself shrinks under blinding, a signature that the biomarker-by-treatment interaction was partly biomarker-by-*PB* rather than biomarker-by-*BR*.

[61]

The phase-augmented extension adds at most four parameters yet captures the practically relevant belief-attenuation diagnostic without requiring the Gompertz parametric form.

- Three to four parameters is well within budget at any realistic trial sample size.
- The extension does not separately estimate *BR* and *PB* components but identifies the change in apparent drug effect when belief is hedged.

- For most analyses this is the practically relevant question and is answerable without committing to the Gompertz form.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

This extension adds three to four parameters at most, well within budget at any realistic N . It does not separately estimate BR and PB as components, but it does identify the change in apparent drug effect that occurs when belief is hedged. For most analyses, this is the practically relevant question, and it can be answered without committing to the Gompertz parametric form. Where both the lumped and the phase-augmented analyses are reported, the additional `Dbc:phase` and `bm:Dbc:phase` terms introduce a small family of belief-attenuation tests, and the attribution inference they support should be treated as a pre-specified secondary analysis with its multiplicity acknowledged rather than as a second bite at the primary interaction test.

5.3 When to fit the full decomposition

[62]

The full decomposition merits the effort only when N is in the hundreds, the design is fully phase-balanced, and the scientific question requires the BR-PB mechanistic split.

- Feasibility conditions: several-hundred N , dense timepoints in every phase, mechanistic BR-PB split required, and informative Bayesian priors available.
- In that regime, `lcmm`, `nlme::nlme` with self-starting Gompertz forms, or `brms` can return participant-specific component estimates.
- Outside that regime the phase-augmented LME is the better tool.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A component-matched analysis merits the effort when sample size is in the several hundreds (a phase-3 trial fielded as aggregated N-of-1, or a meta-analysis pooling multiple smaller trials); the trial design provides timepoints in every phase combination, with enough density per cell to constrain the Gompertz curves; the scientific question requires the BR-PB split (a pharmacology paper claiming a specific mechanism, where ruling out a placebo contribution is the central point); and a Bayesian hierarchical fit with informative priors on the Gompertz parameters is feasible. In that regime, `lcmm`, `nlme::nlme` with self-starting Gompertz forms, or a `brms` hierarchical formulation can return participant-specific BR , PB , and TV component estimates that mirror the DGP. Outside that regime, the linear-mixed analysis

with phase indicators is the better tool.

[63]

The DGP-analysis asymmetry is deliberate: simulation uses the rich DGP to characterise how a simpler analysis behaves; the analysis need not mirror the DGP to produce useful inference.

- Simulation controls the DGP and can specify any structure; inference must pay the identifiability cost of every parameter.
- The DGP is a characterisation tool; the analysis is a finite-data inference tool; they have different optimality criteria.
- DGP-matching is not required and is often counterproductive at trial-relevant sample sizes.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The asymmetry between DGP and analysis is not a flaw of the framework; it is a deliberate consequence of the asymmetry between simulation (where the analyst controls the DGP and can specify any structure, however elaborate) and inference (where the analyst must fit a model to a finite dataset and pay the identifiability cost of every additional parameter). Simulation studies use the rich DGP to characterise how a simpler analysis behaves under realistic data generation; they do not require the analysis to match the DGP in order to produce useful inference.

[64]

The exponential-decay carryover indicator is a deliberately parsimonious special case of richer carryover models, and the simulation should benchmark it against a distributed-lag alternative.

- The continuous indicator D_{bc} decays as $\exp(-\lambda t_{sd})$, a one-parameter form tied to a pharmacokinetic half-life.
- Distributed-lag models, the average-period-treatment-effect estimand, and G-estimation relax the single-rate assumption at additional parametric cost.
- The exponential restriction is an assumption, not a finding, and the pre-registered simulation should include a distributed-lag arm to quantify its cost.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A word on the carryover parameterisation is in order, since it is the one place where the analysis model makes a strong functional-form commitment. The continuous drug indicator D_{bc} decays as $\exp(-\lambda t_{sd})$ off

drug, a one-parameter exponential governed by the carryover half-life. This is deliberately parsimonious and is best understood as a constrained special case of the more flexible carryover models now available for repeated-measures and N-of-1 data. Liao et al.³¹ model carryover through a Bayesian distributed-lag structure with autocorrelated errors, relaxing the single-rate exponential assumption at the cost of additional parameters; Daza⁵⁵ formalises an average-period treatment-effect estimand built to accommodate autocorrelation, trend, and slow carryover jointly; and G{”a}rtner et al.⁵⁶ benchmark linear models against G-estimation and Bayesian-network alternatives specifically under carryover. We adopt the exponential form because it ties directly to a pharmacokinetic half-life and conserves degrees of freedom for the moderation estimand at trial-relevant sample sizes, but the choice is an assumption rather than a finding. The pre-registered simulation programme of section 7 should accordingly include a distributed-lag analysis arm, so that the cost of the exponential-decay restriction is quantified against the richer alternative rather than asserted.

6 Worked example: information loss without decomposition

6.1 What the lumped analysis estimates

[65]

The lumped treatment-effect estimate has a probability limit equal to the pharmacological component plus belief and natural-history contributions, so its bias does not vanish as N grows.

- Writing the reduction as $BR + PB + TV - \varepsilon$ and fitting one drug indicator gives $\text{plim } \hat{\beta}_D = m_{BR} + b_{PB} + b_{TV}$.
- b_{PB} and b_{TV} are the belief and natural-history terms that load onto the drug indicator under the chosen contrast; both are generally non-negative for an improving outcome.
- Because the displacement is in the probability limit, a larger one-component trial converges on the same biased value: the failure is identification, not precision.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Before the numerical illustration it is worth establishing algebraically what the one-component analysis actually estimates, because the qualitative failure modes of section 1 are consequences of an identity rather than empirical tendencies that require simulation to demonstrate. Write the within-participant symptom reduction as

$R_{it} = \text{BL}_i - Y_{it} = BR_{it} + PB_{it} + TV_{it} - \varepsilon_{it}$. The one-component analysis fits a single drug indicator D_{it} ,

$$R_{it} = \alpha + \beta_D D_{it} + g(t) + u_i + e_{it},$$

and reports $\hat{\beta}_D$ as ‘the treatment effect’. Taking the on-drug minus off-drug contrast that β_D targets, and using that the natural-history component is by construction independent of treatment status so that its on/off difference cancels in a within-participant comparison, the probability limit of the estimator is

$$\text{plim } \hat{\beta}_D = m_{BR} + b_{PB} + b_{TV},$$

where $b_{PB} = E[PB \mid \text{on}] - E[PB \mid \text{off}]$ is the belief differential between the contrasted phases and b_{TV} collects any natural-history contribution that fails to cancel under the contrast actually used; b_{TV} is zero for a clean on/off within-participant contrast and strictly positive for a baseline-anchored pre-post analysis, in which case $b_{PB} = E[PB]$ and $b_{TV} = E[TV]$. To fix which case is in play below: the one-component analysis used in the worked example and the simulations of this paper is the design-contrast (drug-on/off) form, so $b_{TV} \approx 0$ there and the operative displacement is the belief term b_{PB} ; the baseline-anchored pre-post form, which additionally carries $b_{TV} = E[TV]$, is the worse case retained only for comparison. Both terms are generally non-negative for an improving outcome, so the lumped estimate over-states the pharmacological component m_{BR} . The displacement is a property of the probability limit and is unchanged as $N \rightarrow \infty$: a larger one-component trial converges on the same biased value. This is the precise sense in which the placebo-attribution failure of section 1 is a problem of identification, not of precision.

[66]

The biomarker slope from a lumped analysis decomposes additively into component-specific slopes, so a biomarker correlated with PB or TV registers as a predictor even with no pharmacological association.

- By linearity of covariance, $\beta_{bm}^{\text{lumped}} = \beta_{bm}^{BR} + w_{PB}\beta_{bm}^{PB} + w_{TV}\beta_{bm}^{TV}$.
- A non-zero β_{bm}^{PB} or β_{bm}^{TV} produces a non-zero lumped slope even when the pharmacological slope β_{bm}^{BR} is exactly zero.
- The result is an algebraic identity under the additive decomposition, not a sample-size-dependent tendency.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The same algebra settles the biomarker-validation failure. The precision-medicine analysis regresses a participant-level response summary on a baseline biomarker B_i . If that summary is the lumped per-participant effect, which estimates $m_{BR,i} + w_{PB} m_{PB,i} + w_{TV} m_{TV,i}$ for design-determined loadings $w_{PB}, w_{TV} \geq 0$, then by the linearity of covariance the fitted biomarker slope has probability limit

$$\beta_{bm}^{\text{lumped}} = \frac{\text{Cov}(m_{BR} + w_{PB} m_{PB} + w_{TV} m_{TV}, B)}{\text{Var}(B)} = \beta_{bm}^{BR} + w_{PB} \beta_{bm}^{PB} + w_{TV} \beta_{bm}^{TV},$$

where $\beta_{bm}^C = \text{Cov}(m_C, B) / \text{Var}(B)$ is the component-specific slope. A biomarker correlated with placebo responsiveness or with the natural-history trajectory, so that β_{bm}^{PB} or β_{bm}^{TV} is non-zero, therefore registers a non-zero lumped slope even when its pharmacological association β_{bm}^{BR} is exactly zero. This is the biomarker-validation failure of section 1, and it too is an identity under the additive decomposition rather than an empirical tendency: no sample size separates the lumped slope from a genuine pharmacological one, because the two coincide in the estimand whenever $w_{PB} \beta_{bm}^{PB} + w_{TV} \beta_{bm}^{TV} \neq 0$.

We make two qualifications explicit, because they bound how much weight the framework can carry. First, both displays are immediate corollaries of the additive decomposition and the linearity of the projection that defines the estimator; they are not a derived result in any deeper sense, and their value is in naming the displacement terms, not in the algebra. Second, they assume the components combine additively. Under the subadditive interaction documented in section 3.3, where the joint effect of BR and PB falls below their sum, a first-order correction proportional to $\text{Cov}(m_{BR} \cdot m_{PB}, B)$ enters the biomarker slope, and the identities above should be read as the leading behaviour in the additive case rather than as exact under interaction. The empirical magnitude of that correction term is itself an open question that the framework's γ -interaction simulation (section 3.3) is designed to address.

[67]

The bias displaces the estimand rather than inflating variance, so a significance test of the lumped coefficient cannot detect it; the decomposition is what locates it.

- The lumped coefficient is displaced from m_{BR} by $b_{PB} + b_{TV}$ but remains a well-estimated quantity with a valid Wald test of its own nullity.
- Detecting the bias requires comparing $\hat{\beta}_D$ to m_{BR} , a comparison only the decomposition or an isolating phase contrast supplies.

- This is why a rejection-rate summary is the wrong instrument, and accounts for the flat rejection rates across the placebo-strength axis in the section 7.1 pilot.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Two consequences for how the framework should be evaluated follow. First, the quantity that carries the bias is the location of the estimand, $\text{plim } \hat{\beta}_D$, not its sampling variability: the lumped coefficient remains a well-estimated number with a valid Wald test of its own nullity, and that test rejects at the rate dictated by the displaced value. A significance test of the lumped coefficient therefore cannot, in principle, detect the bias, because the bias does not move the coefficient toward zero. Detecting it requires comparing $\hat{\beta}_D$ against m_{BR} , which is exactly the comparison the decomposition, or a phase contrast that isolates m_{BR} , makes available and the lumped analysis does not. Second, and as a corollary for the simulation evidence, a rejection-rate or power summary is the wrong instrument for this bias: the right target is the coefficient itself and its offset from m_{BR} . This is why the pilot of section 7.1 found the biomarker-by-treatment rejection rate essentially flat across the placebo-strength axis even though that axis is precisely what drives the predicted bias; the prediction lives in the mean of $\hat{\beta}_{bm:D_{bc}}$, and a production run that means to test it must read the bias off the coefficient, not off the test.

6.2 A numerical illustration

[68]

The worked example uses two participants with identical lumped responses but opposite BR-PB profiles to show that the one-component analysis cannot distinguish them.

- Participant A has high PB ($m_{PB} = 6$) and moderate BR ($m_{BR} = 4$); Participant B has low PB ($m_{PB} = 1$) and high BR ($m_{BR} = 8$).
- Both participants reach approximately 10–11 points of total improvement by week 8, making them indistinguishable to a one-component analysis.
- The blinded-discontinuation phase creates the diagnostic contrast that exposes the underlying difference.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

To illustrate the inferential cost of failing to decompose, we shall work through a simple numerical example. Consider two participants, A and B, enrolled in a sixteen-week N-of-1 trial of prazosin for PTSD. Both

are otherwise similar: same age, same baseline severity, same trial design. They differ in two latent traits that the trial does not measure directly. Participant A is a high placebo responder and a moderate pharmacological responder, with true component values $m_{BR}^A = 4$, $m_{PB}^A = 6$, $m_{TV}^A = 1$. Participant B is a low placebo responder and a strong pharmacological responder, with $m_{BR}^B = 8$, $m_{PB}^B = 1$, $m_{TV}^B = 1$.

[69]

Open-label totals are clinically indistinguishable; the one-component analysis estimates a population drug effect of approximately 10 points with no differentiation between A and B.

- Both participants show approximately 10–11 points of improvement and present as clinical successes.
- A t-test or simple LME on open-label change cannot detect the heterogeneity in BR versus PB composition.
- The population drug-effect estimate conflates both pharmacological and placebo contributions.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In the open-label phase (weeks 1 to 8), both BR and PB run at full strength and TV accumulates. By week 8 both participants have reached close to their saturation values:

	BR	$PB(\eta = 1)$	TV	Total improvement
Participant A	4.0	6.0	1.0	11.0
Participant B	8.0	1.0	1.0	10.0

A clinician examining only the totals would conclude that the two participants are responding similarly: both improved by about ten points, both look like clinical successes. A simple t-test on the open-label change would estimate a population drug effect of about ten points and would not distinguish the two participants.

[70]

The blinded-discontinuation phase reveals the BR-PB heterogeneity: Participant A retains 4 points while B collapses to 1.5, exposing the diagnostic contrast.

- Participant A's high PB at $\eta = 0.5$ (3 points) sustains improvement even without active drug.
- Participant B's low PB (0.5 points) and absent BR produce near-complete return to baseline.

- The 4.0 versus 1.5 contrast is exactly what the decomposition uses to separate *BR* from *PB*.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In the blinded discontinuation phase (weeks 9 to 12), both participants are silently switched to placebo at the start of week 9. The expectancy multiplier drops from 1.0 to 0.5. For clarity of exposition we take the carryover half-life to be short relative to the four-week blinded window, so that *BR* has effectively cleared by the end of week 12; the more realistic case of a slowly decaying *BR*, modelled by the carryover machinery of section 3.1, attenuates but does not eliminate the contrast below. By the end of week 12, the components are:

	<i>BR</i>	<i>PB</i> ($\eta = 0.5$)	<i>TV</i>	Total improvement
Participant A	0.0	3.0	1.0	4.0
Participant B	0.0	0.5	1.0	1.5

The two participants now look quite different. Participant A still shows a sizeable improvement (four points) under blinded placebo; participant B has nearly returned to baseline (one and a half points). This is the diagnostic contrast that the decomposition exploits. For simplicity the table holds *TV* at its week-8 value; in reality *TV* continues to accumulate over weeks 9 to 12, which shifts both totals by a common amount and leaves the *BR*-versus-*PB* contrast on which the example turns unchanged.

[71]

The one-component model estimates drug at 7.75 points with no principled decomposition, obscuring that A's true *BR* is 4 and B's is 8.

- The 7.75 point estimate is a conflated average of true *BR* values (4.0 and 8.0) plus the placebo wash-out.
- The model cannot determine whether A and B differ, nor estimate the pharmacological drug effect separately.
- Placebo and natural-history effects are not identifiable from this contrast alone.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A one-component fixed-effects regression of total symptom change on phase-by-allocation indicators sees:

- Open-label means: 11.0 (A), 10.0 (B).
- Blinded-placebo means: 4.0 (A), 1.5 (B).

- Estimated drug effect: $11.0 - 4.0 = 7.0$ (A), $10.0 - 1.5 = 8.5$ (B). Population mean: 7.75.
- Estimated placebo effect: not identifiable from this contrast.
- Estimated natural history: not identifiable from this contrast.

The one-component model has produced a single number conflated with the placebo response that washes out under blinding. It cannot say whether the drug effect is the same for participants A and B; it estimates ‘drug’ at 7.75 points on average, which is an average of the true 4.0 (A) and 8.0 (B) mixed with the placebo wash-out, with no principled way to decompose further.

[72]

The full decomposition recovers all six component estimates by exploiting phase differential expectancy and the blinded-placebo contrast.

- The LME returns $\hat{m}_{BR}^A = 4.0$, $\hat{m}_{BR}^B = 8.0$, $\hat{m}_{PB}^A = 6.0$, $\hat{m}_{PB}^B = 1.0$ at the population point estimate.
- Participant-level estimates have meaningful uncertainty at trial-relevant N ; the simulation study characterises this.
- The phase differential ($\eta = 1$ vs. $\eta = 0.5$) is the identification mechanism for the PB parameter.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A linear-mixed-effects analysis with the full BR-PB-TV decomposition, fitted to the same data, can recover the individual components by exploiting the differential expectancy between phases and the on-drug-versus-blinded-placebo contrast. After fitting, the model returns $\hat{m}_{BR}^A = 4.0$, $\hat{m}_{BR}^B = 8.0$, $\hat{m}_{PB}^A = 6.0$, $\hat{m}_{PB}^B = 1.0$, $\hat{m}_{TV}^A = \hat{m}_{TV}^B = 1.0$ (at the population point estimate; participant-level estimates have meaningful uncertainty at N in the trial-relevant range, which the simulation study below characterises).

[73]

The one-component model produces one biased average; the three-component model produces six clinically actionable quantities, making biomarker validation possible.

- The lumped estimate of 7.75 points is a biased average of $BR_A = 4$ and $BR_B = 8$ confounded with placebo wash-out.
- The decomposition supplies individual BR , PB , and TV estimates enabling per-participant clinical decisions.
- Predictive-biomarker validation requires the decomposition because biomarker-by- BR is the target estimand.

2021 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
2022 *tempor incididunt ut labore et dolore magna aliqua.*

2023 The information loss from lumping the components is not abstract.
2024 We may compare the inferences directly:

Question	One-component answer	Three-component answer
Does the drug work for A?	~7.75 points (lumped).	$BR_A = 4.0$, distinct from her larger placebo response.
Does the drug work for B?	~7.75 points (lumped).	$BR_B = 8.0$, twice the magnitude of A.
Are A and B different responders?	No discernible difference.	B has 2x A's BR but 1/6 of A's PB .
Is a biomarker predicting BR useful here?	Cannot answer.	Yes if the biomarker correlates with $\hat{m}_{BR}$.
Should A be continued long-term?	Cannot say.	Probably yes ($BR = 4$ is real and pharmacological), but the larger placebo component will not persist if expectations decay.
Population mean drug effect?	7.75 points (biased).	$6.0 = (4.0 + 8.0)/2$ (correct mean of BR).

2025 The one-component model has produced one biased average where the
2026 three-component model has produced six clinically actionable quanti-
2027 ties. Predictive-biomarker validation, which is a central scientific aim
2028 of the trial design family considered here, becomes impossible without
2029 the decomposition.

2030 [74]

2031 **A population-level example shows that cohort differences in**
2032 **open-label totals driven entirely by PB produce apparent**
2033 **non-replicability that the decomposition resolves.**

- 2034 • Cohort 1 (naive, $m_{PB} = 7$) reports 13-point improvements; Co-
2035 hort 2 (sceptical, $m_{PB} = 2$) reports 8 points; true $m_{BR} = 5$ in
2036 both.
- 2037 • A one-component analysis would falsely conclude the drug works
2038 less well in Cohort 2 and trigger searches for confounders.

- The decomposition returns $\hat{m}_{BR} = 5$ in both cohorts, attributing the difference correctly to PB heterogeneity.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A second example at the population level illustrates the same problem. Consider two cohorts of N-of-1 participants enrolled in methodologically identical trials of the same drug. Cohort 1 (early trial, naive patients) has mean $m_{BR} = 5$, mean $m_{PB} = 7$, mean $m_{TV} = 1$. Cohort 2 (later trial, patients who have had multiple prior failed treatments and are sceptical) has mean $m_{BR} = 5$, mean $m_{PB} = 2$, mean $m_{TV} = 1$. By construction the drug works equally well in both cohorts. But the open-label totals differ substantially: cohort 1 reports thirteen-point improvements on average, cohort 2 reports eight-point improvements. A one-component analysis comparing the two cohorts would conclude that the drug works less well in cohort 2 and might trigger a search for confounders or sub-population effects. None of this is real; the difference is entirely in PB , specifically in the population baseline level of placebo responsiveness. The three-component decomposition, applied to either cohort, returns $\hat{m}_{BR} = 5$ in both. Unfortunately, this is the kind of apparent non-replicability that pharmacology has spent the last decade struggling to navigate; the decomposition is one of the formal tools that makes it tractable.

7 Simulation study

[75]

The simulation study follows ADEMP and quantifies bias and identifiability properties of three analysis strategies across a prazosin-PTSD-calibrated parameter grid.

- The Morris (2019) ADEMP framework structures the study: Aims, DGP, Estimands, Methods, Performance measures.
- The pre-registration is at [analysis/scripts/component-decomposition/00-ademp-pre-registration](#).
- Three analysis strategies (one-component, phase-augmented, full decomposition) are compared across the parameter grid.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The simulation study reported and pre-registered in this section is structured under the⁵⁷ ADEMP framework, specifying Aims, Data-generating mechanisms, Estimands, Methods, and Performance measures. The pre-registered plan is at [analysis/scripts/component-decomposition/00-ademp-pre-registration.md](#). A Monte Carlo simulation calibrated to the prazosin/PTSD reference parameters quantifies the bias and identifiability properties of the

three analysis strategies described in section 5 across a parameter grid.

7.1 Pilot validation of analysis machinery (100 replicates, 6 cells)

[76]

The pilot is a structural check of the analysis pipeline, not a substantive empirical result.

- The realised Study A run (section 7.3) is the empirical core of the manuscript; this pilot precedes it.
- The pilot confirms only that data generation, model fitting, and summary output run without error.
- No bias or power claims are derived from the 100-replicate pilot.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The production simulation programme is the empirical core of this manuscript; its first realised run, Study A, is reported in section 7.3, while the fuller pre-registered grid (section 7.2) remains partly outstanding. The pilot reported here precedes both and confirms only that the analysis machinery runs end-to-end.

[77]

The pilot executed a 6-cell, 100-replicate run with perfect convergence, confirming the pipeline is operational.

- The grid crossed three *pb_strength* levels against two analysis strategies.
- Convergence was 1.0 across all six cells under 8-core `mclapply`.
- Output is checkpointed at `analysis/data/quick-sim/`; no inferential claims are drawn.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A 100-replicate-per-cell pilot was run at the reference Hybrid OL+BDC design, $N = 35$, on the 6-cell slice $pb_strength \in \{0, 6.5, 13\} \times analysis \in \{\text{one-component, phase-augmented}\}$, using 8-core `parallel::mclapply`. The convergence rate was 1.0 across all six cells. Per-replicate output is checkpointed at `analysis/data/quick-sim/06-decomp-replicates.rds` and the Morris-format cell-level summary at `analysis/data/quick-sim/06-decomp-summary.txt`. The exploratory summaries are diagnostic only. The one-component `bm:Dbc` test showed power 0.95, 0.93, and 0.83 across $pb_strength \in \{0, 6.5, 13\}$, while the phase-augmented analysis showed 0.65, 0.65, and 0.61; the mean estimated $\hat{\beta}_{bm:Dbc}$ was essentially flat within each analysis (about -2.25 for the one-component

and -2.21 for the phase-augmented fit). At 100 replicates the Monte Carlo standard error on a power estimate near 0.7 is about 0.046, so these figures cannot by themselves separate the two analyses; their interpretation, including why the phase-augmented analysis loses its bias-detection contrast at this N , is resolved by the production run of section 7.3. No confirmatory bias or power claim is drawn from the pilot, and all substantive claims are deferred to the production programme below.

[78]

Two pilot findings revise the production specification: rank-deficiency forces a formula change, and bias must be read from coefficient means, not rejection rates.

- At $N = 35$, rank-deficiency dropped the D_{bc} :phase term, eliminating the key biomarker-by-treatment-by-phase contrast.
- The predicted bias manifests in the mean of $\hat{\beta}_{bm:D_{bc}}$, not in the Wald-test rejection rate.
- The production programme addresses both by using larger N , a fuller design, and coefficient-mean reporting.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Two design-level observations from the pilot inform the production specification. First, at $N = 35$ the rank-deficiency of the full phase-augmented formula required dropping the D_{bc} :phase interaction term, removing the precise contrast that would have isolated the biomarker-by-treatment-by-phase signal the framework is designed to detect. Second, the predicted bias of the one-component analysis arises in the regression coefficient $\hat{\beta}_{bm:D_{bc}}$, not in the indicator-level rejection rate, so power on the biomarker-by-treatment Wald test is not the right summary statistic to detect the framework's predicted bias. The production programme below addresses both observations through larger N , a richer trial design that admits the full phase-augmented formula, and reporting of the coefficient mean against the implied true value rather than the rejection rate alone.

7.2 Production-simulation specification

[79]

The production programme follows the ADEMP framework with four deliverables; the realised Study A run in section 7.3 is the first of these.

- The deliverables are interaction-estimate bias, full-decomposition identifiability, type I error and power, and participant-level signal recovery.

- Sample sizes $N \in \{35, 70, 100, 150\}$ cross a *pb_strength* grid; type I cells target MCSE 0.003 at 5,000 replicates.
- Driver scripts, checkpointed outputs, and the ADEMP pre-registration reside under `analysis/scripts/component-decomposition/`.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The production programme follows the⁵⁷ ADEMP framework and targets four deliverables: the bias of $\hat{\beta}_{bm:D_{bc}}$ for the one-component, phase-augmented, and full-decomposition analyses across $N \in \{35, 70, 100, 150\}$ and a *pb_strength* grid; identifiability metrics for the full-decomposition fit (convergence, variance-inflation factors, and sensitivity of the *BR* estimate to the η (phase) multipliers); type I error and power for the interaction test under an η -mismatch, *TV*-magnitude, and dropout sensitivity grid, with null cells at $n_{\text{reps}} = 5,000$ for MCSE 0.003; and the participant-level interaction signal recovered as a function of (m_{BR}, m_{PB}, m_{TV}) , which the lumped-total regression does not preserve. The `tidyverse` implementation collection drives the grid; driver scripts, checkpointed cell-level summaries, per-replicate output, and the committed ADEMP pre-registration all reside under `analysis/scripts/component-decomposition/`. The realised Study A run reported next used 1,000 replicates per alternative cell.

7.3 Study A results: bias and power at 1,000 replicates per alternative cell

[80]

At 1,000 replicates per alternative cell the one-component estimate of $\hat{\beta}_{bm:D_{bc}}$ is essentially unbiased across the entire (m_{PB}, m_{TV}) grid, contradicting the framework's prediction of a placebo-strength-driven bias.

- The true biomarker-by-*BR* effect is -2.25 ; mean absolute bias is 0.016 at $N = 35$, 0.016 at $N = 70$, 0.006 at $N = 100$, and 0.006 at $N = 150$, in every case within one Monte Carlo SE of zero.
- Bias shows no monotone trend across $m_{PB} \in \{0, 1, 3, 6, 10\}$ or $m_{TV} \in \{-1, 0, 1, 2\}$; it behaves as sampling noise around zero.
- This is the predicted consequence of the identity in section 6.1: the DGP couples the biomarker to *BR* only, so $\beta_{bm}^{PB} = \beta_{bm}^{TV} = 0$ and the lumped slope is unbiased by construction.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The production Study A run (1,000 replicates per alternative cell, 5,000 per null cell, 80 alternative cells crossing $m_{PB} \in \{0, 1, 3, 6, 10\}$,

$m_{TV} \in \{-1, 0, 1, 2\}$, and $N \in \{35, 70, 100, 150\}$, plus four null cells) returns a result that does not support the framework's central bias prediction. The data-generating process couples the baseline biomarker to the BR component only, through the on-drug biomarker- BR covariance channel, so the true biomarker-by-treatment effect is the biomarker-by- BR slope, $\beta_{bm:D_{bc}} = -c_{bm} \sigma_{BR} / \sigma_{bm} = -2.25$, and the biomarker is by construction uncorrelated with PB and TV . Under the one-component analysis the mean estimated $\hat{\beta}_{bm:D_{bc}}$ tracks this value across the whole grid: the mean absolute bias is 0.016 at $N = 35$, 0.016 at $N = 70$, 0.006 at $N = 100$, and 0.006 at $N = 150$, in every case within one Monte Carlo standard error (respectively about 0.028, 0.019, 0.009, and 0.007) of zero, with no monotone trend across m_{PB} or m_{TV} . The predicted placebo-attribution bias of the biomarker-by-treatment estimate does not appear.

This is not a failure of the simulation; it is the behaviour the identity of section 6.1 requires. That identity writes the lumped biomarker slope as $\beta_{bm}^{BR} + w_{PB} \beta_{bm}^{PB} + w_{TV} \beta_{bm}^{TV}$. With the biomarker coupled to BR alone, $\beta_{bm}^{PB} = \beta_{bm}^{TV} = 0$, so the lumped slope equals the pharmacological slope no matter how large m_{PB} or m_{TV} grows. Study A varies the magnitudes of the unmodelled components, which the identity shows to be the innocuous axis; the axis that produces biomarker-validation bias is the biomarker's correlation with PB or TV , which this DGP holds at zero. The clean null is therefore a confirmation of the identity and a correction to the looser claim of section 1 that large PB or TV by itself biases biomarker validation. It does not; biomarker-component coupling does.

[81]

The phase-augmented analysis provides no bias reduction and incurs a substantial power penalty, so it is dominated under this uncontaminated DGP.

- Mean power at $N = 35$ is 0.35 (phase-augmented) versus 0.59 (one-component); at $N = 70$, 0.62 versus 0.90; both reach approximately 1.0 by $N = 100$.
- Phase-augmented mean absolute bias is no smaller than one-component at any N (0.028 versus 0.016 at $N = 35$).
- Convergence is approximately 100% for both analyses at every N (minimum 0.998), so the rank-deficiency that degraded the $N = 35$ pilot does not recur here.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The comparison between analyses is equally clear. The phase-augmented analysis is not less biased than the one-component analysis at any sample size (its mean absolute bias is 0.028 versus

0.016 at $N = 35$ and is comparable elsewhere), and it is markedly less powerful for the biomarker-by-treatment test: mean power across the grid is 0.35 versus 0.59 at $N = 35$ and 0.62 versus 0.90 at $N = 70$, with both analyses reaching essentially full power by $N = 100$. Because the DGP contains no biomarker- PB contamination for the phase contrast to remove, the extra phase-by-treatment terms add estimation cost without buying any bias reduction, and the analysis is dominated. Convergence was approximately complete for both analyses at every sample size (minimum convergence rate 0.998), so the rank-deficiency that forced dropping the `Dbc:phase` term in the $N = 35$ pilot, which we had advanced as the explanation for the pilot's inverted power ranking, is not operative in the production run: the inverted ranking persists with the full formula fitted and converged. The pilot's power inversion was therefore a genuine feature of this DGP, not an artifact of rank-deficiency.

Calibration is adequate for both analyses. Nominal-95% confidence-interval coverage of $\beta_{bm:D_{bc}}$ ranges from 0.96 to 0.99 for the one-component analysis (slightly conservative) and from 0.94 to 0.98 for the phase-augmented analysis. Type I error at the four null cells is controlled, ranging from 0.031 to 0.056 for the one-component analysis and 0.038 to 0.049 for the phase-augmented analysis; the one-component estimate at $N = 100$ (0.056, MCSE 0.003) sits 1.9 standard errors above nominal but is within the simulation's Monte Carlo uncertainty.

[82]

Study A confirms the identity but cannot exhibit the biomarker-validation failure; a contaminated-biomarker pilot reported in the next subsection provides the decisive evidence.

- As pre-registered, Study A varies component magnitude with the biomarker fixed to BR , so it probes an axis the identity predicts to be null.
- The decisive axis is $c_{bm,PB}$, the biomarker- PB correlation; this is varied in the contaminated-biomarker pilot below.
- On current evidence the phase-augmented analysis is not recommended for the uncontaminated-biomarker regime; the pilot below evaluates it under contamination.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Two conclusions follow for the Study A programme. First, the result sharpens rather than supports the section 1 framing: the biomarker-validation failure is real as an identity, but it is driven by the biomarker's coupling to PB or TV , not by the magnitude of those

components, and Study A as pre-registered varies the latter while holding the former at zero. The decisive axis is the biomarker- PB correlation $c_{bm,PB}$, which the contaminated-biomarker pilot in the following subsection varies directly. Second, on the present evidence the phase-augmented analysis carries a real power cost without a compensating bias reduction under an uncontaminated biomarker, so it should not be recommended as a default for this design; the pilot below evaluates whether contamination changes this verdict.

7.4 Contaminated-biomarker pilot (100 replicates per cell)

[83]

A pilot varying the biomarker- PB correlation $c_{bm,PB} \in \{0, 0.15, 0.30, 0.45\}$ confirms that bias in the one-component estimator scales with coupling strength, not component magnitude, and that the phase-augmented analysis provides no protection.

- At $N = 150$, mean $\hat{\beta}_{bm:D_{bc}}$ shifts from -2.25 (true) to -1.74 at $c_{bm,PB} = 0.45$ (bias $+0.51$, 30 Monte Carlo SEs from zero), confirming the identity's prediction.
- The bias is driven by $c_{bm,PB} \times \sigma_{PB}$, not by the population-mean m_{PB} : cells with $m_{PB} = 0$ and $m_{PB} = 6$ show identical bias at every $c_{bm,PB}$ level, because the individual-level PB variance ($\sigma_{PB} = 2$) provides the contamination channel even when the mean is zero.
- The phase-augmented analysis shows the same bias as the one-component analysis at every contamination level, offering no protection against biomarker- PB coupling.
- This pilot uses 100 replicates per cell; a production run at 1,000 replicates per cell is required to confirm the pattern and assess type I error under contamination.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A supplementary pilot varying the biomarker- PB correlation $c_{bm,PB} \in \{0, 0.15, 0.30, 0.45\}$ at two PB magnitude levels ($m_{PB} \in \{0, 6\}$) and two sample sizes ($N \in \{70, 150\}$), using 100 replicates per cell, confirms the identity-based prediction and reveals an unexpected structural result.

At $N = 150$ the one-component bias scales monotonically with $c_{bm,PB}$: mean $\hat{\beta}_{bm:D_{bc}}$ moves from -2.31 ($c_{bm,PB} = 0$, bias -0.06) to -2.37 ($c_{bm,PB} = 0.15$, bias -0.12) to -2.08 ($c_{bm,PB} = 0.30$, bias $+0.17$) to -1.74 ($c_{bm,PB} = 0.45$, bias $+0.51$). The last value sits 30 Monte Carlo standard errors from the true -2.25 , establishing the effect unambigu-

ously. The direction of bias is positive (toward zero): high- $c_{bm,PB}$ participants experience stronger PB , which in the open-label phase adds to the apparent drug effect, but this elevation is partially absorbed by the bm main effect, leaving the $bm : D_{bc}$ interaction coefficient attenuated relative to the pure pharmacological value.

The unexpected result is that m_{PB} does not modulate the bias: at the contaminated cells ($c_{bm,PB} \geq 0.30$, $N = 150$) the $m_{PB} = 0$ and $m_{PB} = 6$ biases are identical to within 0.006, well inside Monte Carlo error, even as the bias itself grows to +0.51; the only larger discrepancy is at $c_{bm,PB} = 0$, where both biases are near zero and differ by 0.067 at the pilot's 100-replicate resolution. The driver is therefore $c_{bm,PB} \times \sigma_{PB}$, not $c_{bm,PB} \times m_{PB}$. Even when the population-mean PB amplitude is zero, individual participants have nonzero PB draws (standard deviation $\sigma_{PB} = 2$), and coupling the biomarker to that variance is sufficient to contaminate the estimator. A true no- PB control requires $\sigma_{PB} = 0$, not $m_{PB} = 0$. This refines the Study A design: varying m_{PB} at $c_{bm,PB} = 0$ correctly produced zero bias, but because the relevant null was $c_{bm,PB} = 0$, not $m_{PB} = 0$.

The phase-augmented analysis provides no protection: its bias at every $(c_{bm,PB}, m_{PB}, N)$ cell is indistinguishable from the one-component bias. Adding phase-by-treatment interactions to the model does not decompose the contamination because both analyses see the biomarker- PB covariance through the same open-label versus blinded-discontinuation contrast; the contrast changes the weight on PB (via expectancy) but does not isolate the PB component's contribution to the biomarker slope. Separating that contribution requires both a design that decouples drug exposure from belief and a decomposition that exploits the decoupling; Study B (section 7.5) shows that the standard hybrid design cannot achieve this separation, whereas a balanced-placebo design can.

The $N = 70$ results are uninformative at 100 replicates (Monte Carlo standard error ≈ 0.06 , comparable to or larger than the bias at moderate $c_{bm,PB}$). The production run will use 1,000 replicates per cell to establish the $N = 70$ pattern and 5,000 null replicates to assess type I error under contamination. Until the production run is complete, the quantitative bias estimates from this pilot should be treated as directionally informative but imprecise.

[84]

The contaminated-biomarker pilot reframes the paper's practical guidance: the risk of biomarker-validation failure arises from biomarker- PB correlation in the study population, not from the presence of strong placebo response per se.

- A trial in a population where the biomarker is uncorrelated with

placebo responsiveness can use the one-component estimator without bias, regardless of how large the placebo response is.

- A trial where the biomarker covaries with optimism, suggestibility, or treatment-seeking history is at material risk of contamination; bias of +0.51 (23% of the true effect) at $c_{bm,PB} = 0.45$ illustrates the practical scale.
- Neither the one-component nor the phase-augmented analysis removes this bias; full component decomposition is required.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The pilot results reframe the paper’s practical guidance. The risk of biomarker-validation failure in an aggregated N-of-1 study arises not from the presence of a strong placebo response in the population but from the biomarker’s correlation with individual-level placebo responsiveness. In a population where the baseline biomarker is uncorrelated with each participant’s placebo-belief response amplitude, the one-component estimator recovers the pharmacological interaction without bias, and the full decomposition machinery is not required. In a population where the biomarker covaries with psychological characteristics that predict placebo response – optimism, prior treatment experience, suggestibility, expectancy at enrolment – the one-component analysis conflates pharmacological and placebo-belief prediction, and the scale of that conflation grows with the coupling strength. At $c_{bm,PB} = 0.45$, a value that equals the pharmacological coupling used throughout this paper, the one-component bias is approximately 0.51 points on the interaction scale, representing 23% of the true pharmacological effect. This is a practically meaningful distortion for any precision-medicine stratification decision.

The production run of the contaminated-biomarker study, at 1,000 alternative-cell replicates and 5,000 null-cell replicates, is required to confirm these estimates, establish the $N = 70$ pattern, and characterise type I error under contamination. Results are pending.

7.5 Study B: recovery under a belief-decoupling design

[85]

Study B shows the contamination is removable, but only under a design that decouples drug exposure from belief: on a balanced-placebo design a belief-covariate decomposition recovers the unbiased pharmacological slope across the contamination range, while the one-component estimator does not.

- On the standard hybrid design the drug state D_{bc} and the expectancy η are collinear, so no decomposition can separate the

channels; we confirmed that several tractable decompositions fail there, no better than the one-component analysis.

- A balanced-placebo design adds a covert on-drug phase (drug given, belief withheld, $\eta = 0$) and an open-placebo phase (drug withheld, belief sustained, $\eta = 1$), which decouples D_{bc} from η .
- The decomposition $Sx \sim bm + t + D_{bc} + \eta + bm:D_{bc} + bm:\eta$ then isolates β_{bm}^{BR} on the $bm:D_{bc}$ term while $bm:\eta$ absorbs the placebo channel.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Study B asks whether the contamination demonstrated in section 7.4 can be removed by decomposition. The answer is conditional on the trial design. The mechanism is that, in the MVN data-generating process, the biomarker-*BR* coupling acts only on drug (and so loads on the drug-state contrast D_{bc}), whereas the biomarker-*PB* coupling scales with the expectancy η (the placebo component's amplitude is proportional to η , so its biomarker-coupled part loads on η). A decomposition that interacts the biomarker with both D_{bc} and η can therefore isolate β_{bm}^{BR} on $bm:D_{bc}$, but only if D_{bc} and η vary independently in the design. In the standard hybrid open-label-plus-blinded-discontinuation design they do not: on-drug time and belief accumulate together (their analysis-stage correlation is about +0.6), and we verified that the belief-covariate decomposition, a blinded-stratum drug contrast, and a Gompertz-basis component decomposition all fail to recover β_{bm}^{BR} on that design, performing no better than, and sometimes worse than, the one-component analysis. The remedy is a design that breaks the collinearity. A balanced-placebo (open-hidden) design^{51,52} interposes a covert on-drug phase, in which the drug is administered without the participant's knowledge ($\eta = 0$), and an open-placebo phase, in which the participant is off drug but believes treatment continues ($\eta = 1$); the first supplies drug variation at fixed belief and the second belief variation at fixed drug exposure, reducing the D_{bc} - η correlation to about -0.66 .

[86]

At 1,000 replicates per cell on the balanced-placebo design, the decomposition is essentially unbiased across the contamination range while the one-component estimator degrades monotonically.

- At $N = 150$, the decomposition's bias of $\hat{\beta}_{bm:D_{bc}}$ stays within one Monte Carlo SE of zero across $c_{bm,PB} \in \{0, 0.1, 0.2, 0.3\}$ ($-0.022, +0.008, +0.021, -0.026$; MCSE ≈ 0.02), while the one-component bias rises monotonically to $+0.16, +0.34, +0.48$.

- Nominal-95% coverage is near-nominal for the decomposition (0.94–0.96) and falls to 0.86 for the one-component analysis at the highest contamination; convergence is 100 percent for both.
- The contamination grid is capped at $c_{bm,PB} = 0.30$ because beyond approximately 0.45 the implied covariance is non-positive-definite (minimum eigenvalue -0.35 at 0.45) and the contaminated DGP is no longer faithfully representable.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The production run (1{,}000 replicates per cell, balanced-placebo design with three open-label, six covert on-drug, and six open-placebo timepoints, $N \in \{70, 150\}$, $c_{bm,PB} \in \{0, 0.1, 0.2, 0.3\}$) confirms the recovery at full resolution. At $N = 150$ the mean estimated $\hat{\beta}_{bm:D_{bc}}$ from the decomposition tracks the pharmacological target -2.25 across the entire contamination range, with bias -0.022 , $+0.008$, $+0.021$, and -0.026 at $c_{bm,PB} = 0, 0.1, 0.2, 0.3$ respectively, every value within one Monte Carlo standard error (about 0.02) of zero. The one-component estimator on the same data is unbiased only at $c_{bm,PB} = 0$ and develops a monotone bias of $+0.16$, $+0.34$, and $+0.48$ as the biomarker- PB correlation grows. Calibration tracks the bias: the decomposition holds near-nominal interval coverage (0.94 to 0.96), whereas one-component coverage degrades to 0.86 at $c_{bm,PB} = 0.3$. The $N = 70$ cells show the same pattern with the decomposition unbiased (absolute bias ≤ 0.08) and lower power (about 0.66 versus 0.93 at $N = 150$), the precision cost of the smaller sample and the extra belief-covariate terms. Convergence was complete in every cell. Per-replicate output and the cell-level summary are at `analysis/data/derived_data/component-decomposition/study-b-balanced/`. The contamination range is bounded above by positive-definiteness: at $c_{bm,PB} = 0.45$ the biomarker would have to correlate 0.45 with both BR and PB , which the covariance cannot accommodate (minimum eigenvalue -0.35), so we restrict the production grid to the feasible range $c_{bm,PB} \leq 0.30$. \end{orig}

[87]

Recovery is therefore conditional on the design, not on the analysis alone: the decomposition is necessary but the belief-decoupling design is what makes the pharmacological

slope identifiable.

- The practical implication is that a biomarker-validation trial exposed to possible biomarker-*PB* coupling should adopt a balanced-placebo or open-hidden design, not the standard hybrid design, if the pharmacological interaction is to be recovered.
- This sharpens the framework’s central claim: decomposition is the mechanism, but a design that varies belief independently of drug exposure is the enabling condition.
- Open questions remaining: recovery beyond the positive-definite range, the participant-level component recovery of the original Study B identifiability aims, and replication of the pattern under Architecture A (mean-moderation) coupling.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The result sharpens the paper’s central claim in an important way. The inferential failure of section 7.4 is real and the decomposition does remove it, but the decomposition is not self-sufficient: on the standard hybrid design it does not recover the pharmacological slope, because the design does not supply independent variation in drug exposure and belief. The enabling condition is a design that decouples the two, of which the balanced-placebo and open-hidden paradigms^{51,52} are the established instances. The practical recommendation that follows is specific: a biomarker-validation trial in which the candidate biomarker may correlate with placebo responsiveness should be fielded as a balanced-placebo or open-hidden design rather than the standard hybrid design, because only the former makes the pharmacological biomarker-by-treatment interaction separately identifiable. Three questions remain open. Recovery beyond the positive-definite contamination range is not characterised, since the MVN architecture cannot represent it; the participant-level component recovery that motivated the original identifiability aims of Study B is not addressed by the slope-level

analysis here; and whether the same recovery holds under Architecture A (mean-moderation) coupling, where the biomarker shifts component means rather than correlating with component draws, remains to be confirmed. \end{orig}

8 Application to predictive-biomarker validation

[88]

The precision-medicine question – whether a biomarker predicts pharmacological response – is answered by the lumped total only when the biomarker is uncorrelated with PB and TV ; otherwise it requires decomposing the response.

- The question is whether the biomarker predicts the BR component, not the lumped total.
- The PATH statement formalises this as the target estimand for predictive biomarker validation.
- The BR-PB-TV decomposition is the mechanism that separates the pharmacological signal when biomarker-component coupling cannot be ruled out.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The motivating clinical question for the trial designs in this project is whether a baseline biomarker predicts treatment response, and specifically whether the biomarker predicts the pharmacological component of treatment response. This is the precision-medicine question, formalised in the PATH statement¹⁴. Whether it depends on the BR-PB-TV decomposition is itself conditional: by the identity of section 6.1, the lumped biomarker slope equals the pharmacological slope when the biomarker is uncorrelated with PB

and TV , and departs from it only by the biomarker's covariance with those components. The decomposition is therefore the mechanism the question requires precisely in the regime where that covariance is non-negligible.

[89]

A lumped-total biomarker regression conflates pharmacological and placebo-belief prediction, producing a stratification biomarker that mis-identifies responders.

- $\hat{\beta}_{bm}$ from the lumped regression contains both pharmacological and placebo-belief prediction in unknown proportions.
- If PB covaries with the biomarker distribution, flagged 'high responders' include high- PB responders who will not benefit pharmacologically.
- Prospective application of such a biomarker would systematically mis-stratify patients for prescription decisions.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A direct regression of total response on baseline biomarker

```
total_response = beta_0 + beta_bm * biomarker + error
```

estimates the biomarker's correlation with whatever the lumped total contains. If PB varies systematically across the biomarker distribution, because the biomarker is correlated with optimism, suggestibility, or treatment-seeking history, then $\hat{\beta}_{bm}$ contains a placebo-prediction component inseparable from the pharmacological-prediction component. The resulting biomarker, applied prospectively, would mis-stratify patients: those flagged as 'high responders' would include both true high- BR responders and high- PB respon-

ders, and the latter would not benefit from prescription decisions made on a pharmacological basis.

[90]

A valid pharmacological predictor requires a non-zero $\hat{\beta}_{bm}^{BR}$; significant $\hat{\beta}_{bm}^{PB}$ or $\hat{\beta}_{bm}^{TV}$ indicate a psychological or prognostic biomarker, not a pharmacological one.

- The sole criterion for pharmacological prediction is a meaningfully non-zero BR -component regression coefficient.
- A biomarker correlating only with PB selects placebo responders, not drug responders.
- A biomarker correlating only with TV is prognostic but not predictive in the pharmacological sense.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The right approach is to regress each component on the biomarker:

$$\begin{aligned} BR &= \beta_0^{BR} + \beta_{bm}^{BR} \cdot \text{biomarker} + \varepsilon^{BR}, \\ PB &= \beta_0^{PB} + \beta_{bm}^{PB} \cdot \text{biomarker} + \varepsilon^{PB}, \\ TV &= \beta_0^{TV} + \beta_{bm}^{TV} \cdot \text{biomarker} + \varepsilon^{TV}. \end{aligned}$$

A biomarker is a useful pharmacological predictor if and only if $\hat{\beta}_{bm}^{BR}$ is meaningfully non-zero. Significant $\hat{\beta}_{bm}^{PB}$ or $\hat{\beta}_{bm}^{TV}$ are diagnostic findings about the biomarker, indicating that it is partly psychological or correlates with natural-history characteristics, but they are not the validation of pharmacological prediction. A biomarker that correlates only with PB would be valuable for selecting placebo responders in run-in designs but worthless for predicting pharmacological efficacy; a biomarker that correlates only with TV would be a prognostic

marker, not a predictive one.

[91]

When biomarker-component coupling cannot be ruled out, the decomposition is the mechanism by which the precision-medicine question is posed; when it can, the lumped analysis already answers it.

- Where the biomarker may correlate with PB or TV , separating BR from them is what prevents validation from conflating pharmacological and non-pharmacological prediction.
- The phase-augmented `bm:Dbc:phase` interaction is a small- N proxy, but Study A shows it buys nothing under an uncontaminated biomarker, so it is warranted only when contamination is plausible.
- The full decomposition recovers component-specific regressions explicitly when sample size permits.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The decomposition is therefore the mechanism by which the precision-medicine question is posed whenever the candidate biomarker may correlate with PB or TV . Where the biomarker can be assumed orthogonal to the non-pharmacological components, the identity of section 6.1 shows the lumped analysis already targets the biomarker-by- BR estimand, and the simulation of section 7.3 confirms it recovers that estimand without bias. The phase-augmented linear-mixed analysis introduced in section 5 retains a useful subset of the decomposition machinery at small N (the `bm:Dbc:phase` interaction tests whether the biomarker-by-treatment effect attenuates under blinding, the practical proxy for biomarker-by- PB contamination), but Study A shows this

proxy earns its degrees of freedom only when such contamination is present: under an uncontaminated biomarker it reduced no bias and lost power. The full decomposition recovers the component-specific regressions explicitly when N permits and when the scientific question warrants the cost.

9 Discussion

[92]

The three-component framework reframes treatment response as a structured sum of causally distinct contributions, and the present work maps when acting on that decomposition actually changes inference.

- The framework specifies which trial-design contrasts identify which components.
- Acting on the decomposition is valuable conditionally: for attribution estimands, and for biomarkers that may correlate with PB or TV , but not for a biomarker-by-treatment interaction with a biomarker orthogonal to them.
- The target estimand for biomarker validation is the biomarker-by- BR association; whether a lumped analysis already recovers it depends on the biomarker's coupling to PB and TV .

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The three-component decomposition is the formal expression of the claim that ‘treatment response’ in an aggregated N-of-1 trial is not a single quantity but a sum of three causally distinct contributions, each with its own pharmacological, psychological, or epidemiological interpretation. What the present work establishes,

analytically in section 6.1 and empirically in section 7.3, is that the inferential value of acting on this decomposition is conditional rather than universal. The identity shows that the one-component estimator is displaced from its pharmacological target only by the components' loading onto the contrast used, and, for a biomarker slope, only by the biomarker's covariance with PB and TV . Where those terms are zero, the design-contrast one-component analysis is already unbiased, and the production simulation confirms it: with a biomarker coupled to BR alone, the one-component biomarker-by-treatment estimate is unbiased across the whole $PB-TV$ grid and the phase-augmented analysis is dominated. The framework's contribution is therefore best read as a map of when decomposition changes inference, not as a blanket prescription to decompose. It makes the trial-design conditions explicit, identifies the attribution and contaminated-biomarker estimands for which decomposition is needed, and clarifies the target estimand for biomarker validation (the biomarker-by- BR association, not the lumped total).

[93]

Four open questions bound the framework's current scope: the simulation results are pending, the η multipliers require sensitivity analysis, non-monotone natural history is not yet handled, and participant-level recovery inherits the identifiability caveats of individual treatment effects.

- Quantitative bias-versus-power trade-offs are conjectures until the section 7 simulation is run.
- The η (phase) expectancy-modulation parameters are modelling assumptions requiring explicit sensitivity analysis.
- Cyclical or biphasic TV trajectories

exceed the Gompertz form and require flexible or covariate-augmented specifications.

- Participant-level component recovery rests on separating participant-by-component signal from participant-by-period and participant-by-phase variation, the identifiability caveat Senn raises for individual treatment effects.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Four open questions remain and warrant acknowledgement. First, the simulation study described in section 7 is forthcoming; the present manuscript reports the framework rather than its empirical evaluation, and the quantitative bias-versus-power trade-offs across analysis strategies are conjectures until that work is run. Second, the η (phase) multipliers used to parameterise the expectancy modulation of PB are themselves modelling assumptions, and a sensitivity analysis varying η across plausible values is mandatory before any clinical conclusion is drawn from the framework's outputs. Third, the framework as specified treats the three components as monotone Gompertz trajectories; cyclical or biphasic TV patterns require either flexible TV specifications or explicit covariates beyond the Gompertz form, and the boundary between 'flexible-Gompertz' and 'covariate-augmented' specifications is not yet sharply drawn. Fourth, the participant-level component estimates the framework targets inherit the identifiability caveat that⁵⁴ and⁴⁹ raise for individual treatment effects more generally: within-participant component recovery rests on separating a true participant-by-component signal from participant-by-period and participant-by-phase variation, and this separation is credible only when the design supplies repeated phase contrasts and the

cross-component covariance structure is not pathological. The framework should be read as estimating these quantities under those assumptions, not as dissolving the well-known difficulty of estimating individual effects from finite within-person data.

[94]

Two companion documents extend the manuscript: a pedagogical narrative and a Gompertz-family sensitivity evaluation.

- The pedagogy document at `docs/24-component-decomposition-pedagogy.md` provides worked examples for first-time readers.
- The Gompertz evaluation at `analysis/report/07-gompertz-evaluation/` quantifies trajectory-specification sensitivity.
- Together they address the intuitive and parametric dimensions of the framework that the present manuscript treats briefly.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The companion pedagogical document at `docs/24-component-decomposition-pedagogy.md` develops the framework with worked examples and an extended intuition narrative for readers approaching the material for the first time. The companion methodological evaluation of the Gompertz family at `analysis/report/07-gompertz-evaluation/` quantifies the sensitivity of the framework's inference to the specific parametric form chosen for the component trajectories.

[95]

The components exist in any longitudinal treatment response, but participant-level identifiability requires a within-subject belief-by-exposure contrast that not every

placebo-controlled design supplies.

- The three-component structure is a property of the data-generating process of any repeated-measures treatment study; its identifiability is design-specific.
- Parallel-group placebo control identifies the population-mean PB but not a participant-level PB , and supplies no within-subject belief contrast; only designs with such a contrast (blinded or open-hidden discontinuation, randomised withdrawal, crossover, hybrid N-of-1) identify the participant-level decomposition the biomarker question needs.
- The broad claim is therefore scoped to designs supplying a within-subject belief-by-exposure contrast, not to placebo-controlled longitudinal trials in general.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A final point concerns the scope of the decomposition, where it is essential to distinguish the existence of the components from their identifiability. The claim that an observed longitudinal treatment response is the sum of a pharmacological, a belief-driven, and a natural-history component is a statement about the data-generating process of any repeated-measures treatment study, and in that sense the decomposition is not specific to N-of-1 trials. Identifiability, by contrast, is design-specific, and the distinction between population-level and participant-level recovery is decisive for the biomarker application. A standard parallel-group placebo-controlled trial supplies a between-arm placebo contrast that identifies the population-mean belief response but not a participant-level PB , and it contains no within-subject contrast in which belief

varies independently of pharmacology; the participant-level biomarker-by- BR slope that the precision-medicine question targets is therefore not identified from such a design, no matter how large it is. The designs that do supply the required within-subject belief contrast are the blinded or open-hidden discontinuation and randomised-withdrawal designs^{51,52}, the crossover with on/off periods, and the hybrid N-of-1 design developed here; the pharmacometric disease-progression-plus-placebo models from which the framework descends¹⁻³ operate at the group level and recover population-mean components only. We therefore advocate the decomposition as a default lens for longitudinal designs that supply a within-subject belief-by-exposure contrast, with the aggregated N-of-1 case the most demanding instance because it must recover the components from a single participant's trajectory. We do not claim participant-level identifiability from parallel-group placebo control alone, and where only a between-arm placebo contrast is available the framework's claims are restricted to the population mean.

10 Conclusions

[96]

Treatment response in aggregated N-of-1 trials is the sum of three causally distinct components, and the bias from conflating them is governed by the estimand and the biomarker's coupling to the components, not by their raw magnitude.

- The three components are BR (pharmacological), PB (placebo-belief), and TV (natural history).
- Each has its own biological mechanism, identifiability conditions, and clinical interpretation.

- The lumped pre-post treatment effect is biased by the PB and TV contributions; the biomarker-by-treatment slope is biased only by the biomarker's covariance with PB and TV , as Study A confirms.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In conclusion, three points are to be emphasised. First, treatment response in aggregated N-of-1 clinical trials is the sum of three causally distinct components: pharmacological action (BR), placebo-belief response (PB), and time-variant natural history (TV). Each component has its own biological mechanism, its own identifiability conditions, and its own clinical interpretation. Whether conflating them biases a given analysis is, however, governed by the omitted-variable-bias identity of section 6.1, not by the raw magnitude of PB and TV : the lumped pre-post treatment effect is displaced by the PB and TV contributions, whereas the biomarker-by-treatment slope is displaced only by the biomarker's covariance with those components. Study A bears this out: varying PB and TV magnitude with a biomarker coupled to BR alone left the one-component estimate unbiased across the entire grid.

[97]

What matters is the design contrast, not the parametric model: the trial must supply the drug-on/off and belief contrasts, but adding parametric component terms is not free and is not always worth it.

- Unique component identification requires phase contrasts built into the trial design; the one-component analysis already exploits the drug-on/off contrast.
- Under an uncontaminated biomarker,

the phase-augmented analysis adds no bias reduction and loses power, so it is not recommended as a default.

- The phase-augmented and full component-matched analyses are warranted for attribution estimands and for biomarkers that may correlate with PB or TV , where the extra structure does work.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Second, the operative requirement is the design contrast, not the parametric analysis model. The trial design must include the drug-on/off and belief contrasts that identify the components, and the one-component analysis already exploits the drug-on/off contrast through its D_{bc} term. Adding parametric component structure on top is not free: under an uncontaminated biomarker the phase-augmented analysis reduced no bias and lost substantial power, so it should not be adopted as a default. Its value, and that of the full component-matched fit, is confined to the estimands where the extra structure does work, namely the attribution of response to pharmacology versus belief and the analysis of biomarkers that may themselves correlate with PB or TV .

[98]

For predictive-biomarker validation the decomposition matters precisely when the candidate biomarker may correlate with PB or TV ; the contaminated-biomarker study demonstrates this failure and Study B shows it is recoverable under a belief-decoupling design.

- A biomarker that predicts the lumped total is a pharmacological predictor only if it is uncorrelated with PB and TV ; otherwise the lumped slope is contaminated.

- Regulatory and prescribing decisions therefore require separating BR from PB and TV whenever biomarker-component coupling cannot be ruled out.
- Study A established the uncontaminated baseline, the contaminated-biomarker study (section 7.4) demonstrated the failure, and Study B (section 7.5) showed the bias is recoverable on a balanced-placebo design but not on the standard hybrid design.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Third, for predictive-biomarker validation the decomposition earns its place precisely when the candidate biomarker may correlate with PB or TV . A biomarker that predicts the lumped total is a pharmacological predictor only when it is uncorrelated with the non-pharmacological components; where that cannot be assumed, the lumped slope mixes pharmacological and non-pharmacological prediction, and the regulatory and prescribing decisions that motivate the trial require the separation the decomposition supplies. The contaminated-biomarker study (section 7.4) demonstrates this failure directly, and Study B (section 7.5) establishes that the bias is recoverable: a belief-covariate decomposition returns the unbiased pharmacological slope across the feasible contamination range, but only on a balanced-placebo design that decouples drug exposure from belief, not on the standard hybrid design. The decisive practical lesson is therefore as much about design as about analysis.

11 Conflict of interest

The authors declare no conflicts of interest.

12 Funding

This work received no specific funding.

13 Data availability statement

The data and code that support the findings of this study are openly available in the pmsimstats-ng project repository. Per-replicate Monte Carlo outputs and ADEMP-compliant cell-level summaries are versioned under `analysis/data/`; simulation drivers are at `analysis/scripts/`. Exact paths for this manuscript are listed in the Reproducibility section below.

14 Reproducibility

[99]

The manuscript is rendered inside a Docker container with locked package versions; the full software environment is pinned by Dockerfile and `renv.lock`.

- R 4.4.x in the zzcollab Docker container provides the execution environment.
- Simulation packages are `nlme`, `parallel`, `data.table`, `furrr`, and `purrr`.
- The manuscript build uses `rmarkdown` and `bookdown`.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

This manuscript was rendered with R 4.4.x inside the project’s zzcollab Docker container (image hash recorded in `Dockerfile`; package set captured in `renv.lock`). Principal packages used by the simulation pipeline are `nlme` (continuous-time linear-mixed analysis with `corCAR1` residual correlation), `parallel` (8-core `mclapply` for parameter-grid sweeps), `data.table` (the original-extended simulation collection), `furrr` and `purrr` (the `tidyverse`

simulation collection), and `rmarkdown` plus `bookdown` for the manuscript build.

[100]

Reproducibility is ensured by a master seed of 20260507 with per-replicate seeds derived by index addition.

- `set.seed(20260507)` is called at the top of each replicate loop.
- `parallel::mc.set.seed` propagates seeds inside `mclapply`.
- Per-replicate seeds are `master_seed + rep_idx`.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Each simulation driver sets `set.seed(20260507)` at the top of the replicate loop and uses `parallel::mc.set.seed` inside `mclapply`. Per-replicate seeds derive from the master seed by `+ rep_idx`.

[101]

All simulation outputs, driver scripts, pre-registration, and manuscript source files are versioned at documented paths within the repository.

- Per-replicate outputs and Morris summaries are at `analysis/data/quick-sim/`.
- Driver scripts and the ADEMP pre-registration are at `analysis/scripts/component-decomposition/`.
- Manuscript source, bibliography, and CSL file are at `analysis/report/06-component-decomposition/`.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Per-replicate simulation outputs are checkpointed at `analysis/data/quick-sim/06-decomp-replicates.rds`; Morris ADEMP cell-level summaries at the companion `06-decomp-summary.txt`. Driver scripts are under `analysis/scripts/component-decomposition/`. The pre-registered ADEMP plan is at

analysis/scripts/component-decomposition/00-ademp-pre-registration.md.
 The manuscript source, paper-specific bibliography (references.bib), and SIM CSL
 file are at analysis/report/06-component-decomposition/.

15 References

1. Holford N. Clinical pharmacology = disease progression + drug action. *British Journal of Clinical Pharmacology*. 2015;79(1):18-27. doi:10.1111/bcp.12170
2. Gomeni R, Lavergne A, Merlo-Pich E. Modelling placebo response in depression trials using a longitudinal model with informative dropout. *European Journal of Pharmaceutical Sciences*. 2009;36(1):4-10. doi:10.1016/j.ejps.2008.10.025
3. Chen C, Jönsson S, Yang S, Plan EL, Karlsson MO. Detecting placebo and drug effects on Parkinson's disease symptoms by longitudinal item-score models. *CPT: Pharmacometrics & Systems Pharmacology*. 2021;10(4):309-317. doi:10.1002/psp4.12601
4. Muthén B, Brown HC. Estimating drug effects in the presence of placebo response: Causal inference using growth mixture modeling. *Statistics in Medicine*. 2009;28(27):3363-3385. doi:10.1002/sim.3721
5. Kaptchuk TJ, Kelley JM, Conboy LA, et al. Components of placebo effect: Randomised controlled trial in patients with irritable bowel syndrome. *BMJ*. 2008;336(7651):999-1003. doi:10.1136/bmj.39524.439618.25

- 3180 6. Kirsch I, Deacon BJ, Huedo-Medina TB, Scoboria A, Moore TJ, Johnson BT. Initial severity and antidepressant benefits: A meta-analysis of data submitted to the Food and Drug Administration. *PLoS Medicine*. 2008;5(2):e45. doi:10.1371/journal.pmed.0050045
- 3181 7. Locher C, Kossowsky J, Gaab J, Kirsch I, et al. Moderation of antidepressant and placebo outcomes by baseline severity in late-life depression: A systematic review and meta-analysis. *Journal of Affective Disorders*. 2015;181:50-60. doi:10.1016/j.jad.2015.03.062
- 3182 8. Sugarman MA, Loree AM, Baltes BB, Grekin ER, Kirsch I. The efficacy of paroxetine and placebo in treating anxiety and depression. *PLoS ONE*. 2014;9(8):e106337. doi:10.1371/journal.pone.0106337
- 3183 9. Hall KT, Lembo AJ, Kirsch I, et al. Catechol-O-methyltransferase val158met polymorphism predicts placebo effect in irritable bowel syndrome. *PLoS ONE*. 2012;7(10):e48135. doi:10.1371/journal.pone.0048135
- 3184 10. Wang RS, Lembo AJ, Kaptchuk TJ, Hall KT, et al. Genomic effects associated with response to placebo treatment in a randomized trial of irritable bowel syndrome. *Frontiers in Pain Research*. 2022;2:775386. doi:10.3389/fpain.2021.775386
- 3185 11. Hróbjartsson A, Gøtzsche PC. Is the placebo powerless? An analysis of clinical trials comparing placebo with no treatment. *New England Journal of Medicine*. 2001;344(21):1594-1602. doi:10.1056/NEJM200105243442106

- 3186 12. Hróbjartsson A, Gøtzsche PC. Is the placebo powerless? Update of a systematic review with 52 new randomized trials. *Journal of Internal Medicine*. 2004;256(2):91-100. doi:10.1111/j.1365-2796.2004.01355.x
- 3187 13. Hróbjartsson A, Gøtzsche PC. Placebo interventions for all clinical conditions. *Cochrane Database of Systematic Reviews*. 2010;(1):CD003974. doi:10.1002/14651858.CD003974.pub3
- 3188 14. Kent DM, Paulus JK, Klaveren D van, et al. The Predictive Approaches to Treatment effect Heterogeneity (PATH) Statement. *Annals of Internal Medicine*. 2020;172(1):35-45. doi:10.7326/M18-3667
- 3189 15. Kent DM, Klaveren D van, Paulus JK, et al. The PATH statement: Explanation and elaboration. *Annals of Internal Medicine*. 2020;172(1):W1-W25. doi:10.7326/M18-3668
- 3190 16. Kent DM, Steyerberg E, Klaveren D van. Personalized evidence based medicine: Predictive approaches to heterogeneous treatment effects. *BMJ*. 2018;363:k4245. doi:10.1136/bmj.k4245
- 3191 17. Rekkas A, Paulus JK, Raman G, et al. Predictive approaches to heterogeneous treatment effects: A scoping review. *BMC Medical Research Methodology*. 2020;20(1):264. doi:10.1186/s12874-020-01145-1
- 3192 18. Guyatt G, Sackett D, Taylor DW, Chong J, Roberts R, Pugsley S. Determining optimal therapy: Randomized trials in individual patients. *New England Journal of Medicine*. 1986;314(14):889-892.

- 3193 19. Zucker DR, Schmid CH, McIntosh MW, D'Agostino RB, Selker HP, Lau J. Combining single patient (N-of-1) trials to estimate population treatment effects and to evaluate individual patient responses to treatment. *Journal of Clinical Epidemiology*. 1997;50(4):401-410. doi:10.1016/S0895-4356(96)00429-5
- 3194 20. Zucker DR, Ruthazer R, Schmid CH. Individual (N-of-1) trials can be combined to give population comparative treatment effect estimates: Methodologic considerations. *Journal of Clinical Epidemiology*. 2010;63(12):1312-1323. doi:10.1016/j.jclinepi.2010.04.020
- 3195 21. Lillie EO, Patay B, Diamant J, Issell B, Topol EJ, Schork NJ. The n-of-1 clinical trial: The ultimate strategy for individualizing medicine? *Personalized Medicine*. 2011;8(2):161-173. doi:10.2217/pme.11.7
- 3196 22. Duan N, Kravitz RL, Schmid CH. Single-patient (n-of-1) trials: A pragmatic clinical decision methodology for patient-centered comparative effectiveness research. *Journal of Clinical Epidemiology*. 2013;66(8 Suppl):S21-S28. doi:10.1016/j.jclinepi.2013.04.006
- 3197 23. Gabler NB, Duan N, Vohra S, Kravitz RL. N-of-1 trials in the medical literature: A systematic review. *Medical Care*. 2011;49(8):761-768. doi:10.1097/MLR.0b013e318215d90d
- 3198 24. Schork NJ, Goetz LH. Single-subject studies in translational nutrition research. *Annual Review of Nutrition*. 2017;37:395-422. doi:10.1146/annurev-nutr-071816-064717

- 3199 25. Schork NJ. Randomized clinical trials and personalized medicine: A commentary on Deaton and Cartwright. *Social Science and Medicine*. 2018;210:71-73. doi:10.1016/j.socscimed.2018.04.033
- 3200 26. Schork NJ, Beaulieu-Jones B, Liang WS, Smalley S, Goetz LH. Exploring human biology with N-of-1 clinical trials. *Cambridge Prisms: Precision Medicine*. 2023;1:e12. doi:10.1017/pcm.2022.15
- 3201 27. Porcino A, Vohra S. N-of-1 trials, their reporting guidelines, and the advancement of open science principles. *Harvard Data Science Review*. 2022;4(SI3). doi:10.1162/99608f92.a65a257a
- 3202 28. Vohra S, Shamseer L, Sampson M, et al. CONSORT extension for reporting N-of-1 trials (CENT) 2015 statement. *Journal of Clinical Epidemiology*. 2015;76:9-17. doi:10.1016/j.jclinepi.2015.05.004
- 3203 29. Shamseer L, Sampson M, Bukutu C, Schmid CH, et al. CONSORT extension for reporting N-of-1 trials (CENT) 2015: Explanation and elaboration. *Journal of Clinical Epidemiology*. 2015;76:18-46. doi:10.1016/j.jclinepi.2015.05.018
- 3204 30. Li J, Gao W, Punja S, et al. Reporting quality of N-of-1 trials published between 1985 and 2013: A systematic review. *Journal of Clinical Epidemiology*. 2016;76:57-64. doi:10.1016/j.jclinepi.2015.11.016
- 3205 31. Liao Z, Qian M, Kronish IM, Cheung YK. Analysis of N-of-1 trials using Bayesian distributed lag model with autocorrelated errors. *Statistics in Medicine*. 2023;42(13):2044-2060. doi:10.1002/sim.9676

- 3206 32. Blackston JW, Chapple AG, McGree JM, McDonald S, Nikles J. Comparison of aggregated N-of-1 trials with parallel and crossover randomized controlled trials using simulation studies. *Healthcare*. 2019;7(4):137. doi:10.3390/healthcare7040137
- 3207 33. Carrozzo AE, Zimmermann G, Bathke AC, et al. Two-arm crossover randomized controlled trial versus meta-analysis of N-of-1 studies: Comparison of statistical efficiency. *Biometrical Journal*. 2025;67(2):e70045. doi:10.1002/bimj.70045
- 3208 34. Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: Statistical methods and practical considerations. *Frontiers in Digital Health*. 2020;2:13. doi:10.3389/fdgth.2020.00013
- 3209 35. Kaptchuk TJ, Friedlander E, Kelley JM, et al. Placebos without deception: A randomized controlled trial in irritable bowel syndrome. *PLoS ONE*. 2010;5(12):e15591. doi:10.1371/journal.pone.0015591
- 3210 36. Lembo A, Kelley JM, Nee J, et al. Open-label placebo vs double-blind placebo for irritable bowel syndrome: A randomized clinical trial. *Pain*. 2021;162(9):2428-2435. doi:10.1097/j.pain.0000000000002234
- 3211 37. Nurko S, Saps M, Kossowsky J, others, Kelley JM. Effect of open-label placebo on children and adolescents with functional abdominal pain or irritable bowel syndrome. *JAMA Pediatrics*. 2022;176(4):349-356. doi:10.1001/jamapediatrics.2021.5750
- 3212 38. Benedetti F. Placebo effects: From the neurobiological paradigm to translational implications. *Neuron*. 2014;84(3):623-637.

- 3213 39. Frisaldi E, Shaibani A, Benedetti F. Understanding the mechanisms of placebo and nocebo effects. *Swiss Medical Weekly*. 2020;150:w20340. doi:10.4414/smw.2020.20340
- 3214 40. Carlino E, Frisaldi E, Benedetti F. Pain and the context. *Nature Reviews Rheumatology*. 2014;10(6):348-355. doi:10.1038/nrrheum.2014.17
- 3215 41. Carlino E, Piedimonte A, Benedetti F. Nature of the placebo and nocebo effect in relation to functional neurologic disorders. In: *Handbook of Clinical Neurology*. Vol 139.; 2016:597-606. doi:10.1016/B978-0-12-801772-2.00048-5
- 3216 42. Wager TD, Atlas LY, Lindquist MA, Roy M, Woo CW, Kross E. An fMRI-based neurologic signature of physical pain. *New England Journal of Medicine*. 2013;368(15):1388-1397. doi:10.1056/NEJMoa1204471
- 3217 43. Zunhammer M, Bingel U, Wager TD. Placebo effects on the neurologic pain signature: A meta-analysis of individual participant functional magnetic resonance imaging data. *JAMA Neurology*. 2018;75(11):1321-1330. doi:10.1001/jamaneurol.2018.2017
- 3218 44. Raskind MA, Peterson K, Williams T, et al. A trial of prazosin for combat trauma PTSD with nightmares in active-duty soldiers returned from Iraq and Afghanistan. *American Journal of Psychiatry*. 2013;170(9):1003-1010. doi:10.1176/appi.ajp.2013.12081133

- 3219 45. Holford NHG. Holford NHG and sheiner LB 'understanding the dose-effect relationship - clinical application of pharmacokinetic-pharmacodynamic models', *Clin Pharmacokin* 6:429-453 (1981) - the backstory. *The AAPS Journal*. 2011;13(4):662-664. doi:10.1208/s12248-011-9306-5
- 3220 46. Tilburt JC, Emanuel EJ, Kaptchuk TJ, Curlin FA, Miller FG. Prescribing 'placebo treatments': Results of national survey of US internists and rheumatologists. *BMJ*. 2008;337:a1938. doi:10.1136/bmj.a1938
- 3221 47. Finniss DG, Kaptchuk TJ, Miller F, Benedetti F. Biological, clinical, and ethical advances of placebo effects. *The Lancet*. 2010;375(9715):686-695. doi:10.1016/S0140-6736(09)61706-2
- 3222 48. Kirsch I. Response expectancy as a determinant of experience and behavior. *American Psychologist*. 1985;40(11):1189-1202. doi:10.1037/0003-066X.40.11.1189
- 3223 49. Senn S. Controversies concerning randomization and additivity in clinical trials. *Statistics in Medicine*. 2004;23(24):3729-3753. doi:10.1002/sim.2074
- 3224 50. Bingel U, Wanigasekera V, Wiech K, et al. The effect of treatment expectation on drug efficacy: Imaging the analgesic benefit of the opioid remifentanyl. *Science Translational Medicine*. 2011;3(70):70ra14. doi:10.1126/scitranslmed.3001244
- 3225 51. Colloca L, Lopiano L, Lanotte M, Benedetti F. Overt versus covert treatment for pain, anxiety, and Parkinson's disease. *The Lancet Neurology*. 2004;3(11):679-684.

- 3226 52. Rohsenow DJ, Marlatt GA. The balanced placebo design: Methodological considerations. *Addictive Behaviors*. 1981;6(2):107-122.
- 3227 53. International Council for Harmonisation of Technical Requirements for Pharmaceuticals for Human Use. ICH E9(R1) addendum on estimands and sensitivity analysis in clinical trials to the guideline on statistical principles for clinical trials. Published online 2019.
- 3228 54. Senn S. Statistical pitfalls of personalized medicine. *Nature*. 2018;563(7733):619-621.
- 3229 55. Daza EJ. Causal analysis of self-tracked time series data using a counterfactual framework for N-of-1 trials. *Methods of Information in Medicine*. 2018;57(S 01):e10-e21. doi:10.3414/ME16-02-0044
- 3230 56. Gärtner T, Schneider J, Arnrich B, Konigorski S. Comparison of Bayesian networks, G-estimation and linear models to estimate causal treatment effects in aggregated N-of-1 trials with carry-over effects. *BMC Medical Research Methodology*. 2023;23:191. doi:10.1186/s12874-023-02012-5
- 3231 57. Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Statistics in Medicine*. 2019;38(11):2074-2102.

3232

3233

3234

Rendered on 2026-06-16 at 13:34 PDT.

Source: ~/prj/res/36-pmsimstats-ng/pmsimstats-ng/analysis/report/06-component-dec