

1 Three-component decomposition of treatment response
2 in aggregated N-of-1 trials: pharmacological,
3 expectation-driven, and natural-history components

4 pmsimstats team

5 2026-06-17

6 **Contents**

7	Abstract	2
8	1 Introduction	3
9	2 Notation and the formal model	6
10	3 The three components	8
11	3.1 Biological response (BR)	9
12	3.2 Placebo-belief response (PB)	13
13	3.3 Time-variant natural-history component (TV)	17
14	4 Identifiability and trial-design contrasts	19
15	5 Analysis-side methodology	22
16	5.1 Four reasons not to match the DGP	23
17	5.2 A small extension that usually pays off: phase indicators	26
18	5.3 When to fit the full decomposition	27
19	6 Worked example: information loss without decomposition	30
20	6.1 What the lumped analysis estimates	30
21	6.2 A numerical illustration	33
22	7 Simulation study	33
23	7.1 Pilot validation	34
24	7.2 Production-simulation specification	35
25	7.3 Study A: orthogonal-biomarker arm (1,000 replicates per alterna-	
26	tive cell)	35
27	7.4 Study B pilot: coupling-arm simulation (100 replicates per cell) . .	38
28	7.5 Study C: recovery under a belief-decoupling design	41
29	8 Application to predictive-biomarker validation	43

30	9 Discussion	46
31	10 Conclusions	49
32	11 Conflict of interest	52
33	12 Funding	52
34	13 Data availability statement	52
35	14 Reproducibility	52
36	15 References	53

37 Abstract

38 **Background.** Treatment response in aggregated N-of-1 clinical trials reflects
39 three causally distinct mechanisms: pharmacological action of the active com-
40 pound (BR), the patient’s belief that the treatment will work (PB), and natural-
41 history drift in the underlying condition (TV). Standard analysis models lump
42 these into a single ‘treatment effect’ plus residual error. Whether this lump-
43 ing biases inference, and whether a component decomposition repays its cost in
44 degrees of freedom and power, is not unconditional: it depends jointly on the
45 estimand, on the trial design, and on whether a candidate biomarker correlates
46 with the non-pharmacological components. This paper makes that dependence
47 precise and resolves it into a design-by-correlation-structure rule.

48 **Methods.** We formalise a three-component decomposition of the within-
49 participant outcome trajectory, each component a modified Gompertz function
50 with participant-specific random effects, and characterise the trial-design con-
51 trasts (open-label, blinded discontinuation, crossover, balanced-placebo) that
52 supply identifiability. We derive the omitted-variable-bias identity for the one-
53 component estimator, which expresses both the lumped treatment effect and the
54 lumped biomarker slope as sums of component-specific terms. We then evaluate
55 analysis strategies in three Monte Carlo studies calibrated to the prazosin-PTSD
56 reference parameters: Study A couples the biomarker to BR only ($c_{bm,PB} = 0$)
57 and varies the unmodelled-component magnitudes (1,000 replicates per alter-
58 native cell, 5,000 per null, across $m_{PB} \in \{0, 1, 3, 6, 10\}$, $m_{TV} \in \{-1, 0, 1, 2\}$,
59 $N \in \{35, 70, 100, 150\}$); Study B (100-replicate pilot) couples the biomarker to
60 PB; and Study C (1,000 replicates) contrasts the standard hybrid design with a
61 belief-decoupling balanced-placebo design across $c_{bm,PB} \in \{0, 0.1, 0.2, 0.3\}$.

62 **Results.** The identity shows that the lumped biomarker-by-treatment slope is dis-
63 placed from the pharmacological slope by a term proportional to the biomarker’s
64 covariance with PB and TV, not by the magnitude of those components; three
65 regimes follow. (i) *Biomarker orthogonal to PB and TV, hybrid design:* the one-
66 component estimator is essentially unbiased across the whole grid (mean absolute
67 bias 0.016 at $N = 35$, 0.006 at $N = 150$, within one Monte Carlo standard error of

zero) and strictly more powerful than the phase-augmented decomposition (mean power 0.59 versus 0.35 at $N = 35$; 0.90 versus 0.62 at $N = 70$); decomposition removes no bias and only sacrifices power, so it is dominated. (ii) *Biomarker coupled to PB, standard hybrid design*: the one-component slope is biased in proportion to $c_{bm,PB}$ (one-component bias rising to +0.48, coverage falling to 0.86 at $N = 150$), but because the drug-state and belief contrasts are collinear under this design (analysis-stage correlation $\approx +0.6$), no decomposition recovers the pharmacological slope either. (iii) *Biomarker coupled to PB, balanced-placebo design*: adding covert-drug and open-placebo phases decouples the contrasts (≈ -0.66), and the belief-covariate decomposition then returns the unbiased pharmacological slope across the feasible contamination range (bias within one Monte Carlo SE of zero, coverage 0.94–0.96) where the one-component estimator does not, at the cost of some power.

Conclusions. The benefit of decomposition is jointly determined by the biomarker’s correlation structure and the trial design, not by either alone. When the candidate biomarker is orthogonal to PB and TV, a single-indicator analysis that exploits the design contrast is unbiased and most powerful, and a parametric component model is counterproductive. When the biomarker may correlate with PB or TV, the lumped slope is biased and decomposition is necessary, but the bias is recoverable only under a design that varies belief independently of drug exposure (balanced-placebo or open-hidden); on the standard hybrid design no analysis recovers it, so a biomarker-validation trial at risk of placebo-belief contamination must be fielded with the decoupling design. Component decomposition is likewise warranted whenever the estimand is the attribution of response to pharmacology versus belief versus natural history. A pragmatic specification for the cases that warrant it is $Sx \sim bm + t + Dbc + phase + Dbc:phase + bm:Dbc + bm:Dbc:phase$ with random intercept and AR(1) residual correlation.

1 Introduction

[1]

Partitioning treatment response into pharmacological (*BR*), placebo-belief (*PB*), and natural-history (*TV*) channels has biological precedent and clinical appeal; this paper rigorously evaluates when the partition is worth making.

- Precedent runs from pharmacometric disease-progression models [holford2015; gomeni2009] through Kaptchuk’s empirical decomposition of the placebo response itself [kaptchuk2008components] to neuroimaging dissociations of drug from expectancy effects [wager2013nps; zunhammer2018nps]; each motivates the partition on independent biological grounds.
- Clinical appeal is direct: *BR* is what persists after expectancy decays; *PB* drives apparent response in high-placebo-

responsiveness populations and misleads precision-medicine stratification when mistaken for pharmacological prediction; TV is what would have changed without any treatment at all.

- Partitioning costs degrees of freedom, requires specific trial-design contrasts for identifiability, and adds analytic complexity; whether the cost is warranted is an empirical and design-dependent question this paper answers.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Decomposing treatment response into pharmacological, expectancy, and natural-history contributions is not a new idea. Pharmacometrics has long modelled disease progression and drug effect as separable longitudinal processes^{1,2}; Kaptchuk and colleagues decomposed the placebo response itself into observation, ritual, and practitioner-relationship components in a three-arm IBS trial³; neuroimaging has directly dissociated pharmacological from expectancy-driven analgesia using objective neural signatures^{4,5}. The partition has intuitive clinical appeal: a patient whose symptom improvement is pharmacological will benefit from long-term treatment even after novelty fades, whereas one whose improvement is primarily belief-driven may not sustain it as expectations normalise. The question this paper addresses is not whether the partition is biologically coherent – it is – but when it is necessary for valid inference.

[2]

The optimal choice is conditional: decompose when the candidate biomarker may correlate with PB or TV ; use the lumped analysis when domain knowledge supports orthogonality.

- Decompose: COMT val158met genotype as a biomarker in IBS or PTSD trials – COMT regulates dopaminergic expectancy processing and is a documented predictor of placebo-response magnitude [hall2012comt; wang2022comt], making it a PB moderator whose lumped biomarker-treatment slope conflates pharmacological and belief-driven prediction. Baseline symptom severity is similarly confounded: it predicts pharmacological response but also couples to regression-to-mean TV [hrobjartsson2001; hrobjartsson2010] and heightened treatment expectancy [kirsch2008; locher2015].
- Lumped sufficient: CYP metabolizer phenotype as a pharmacokinetic biomarker – enzymatic drug metabolism cannot plausibly couple to belief state or untreated disease trajectory, so the lumped pharmacological slope is already the target estimand. ADRB2 Arg16Gly in asthma bronchodilator trials provides an

analogous case: the variant affects receptor downregulation via a direct pharmacological mechanism, lung function is objectively measured by spirometry, and no expectancy pathway links ADRB2 genotype to PB .

- In the PTSD/prazosin context of this paper: baseline PCL-5 severity belongs in the decompose column (a psychological measure with documented coupling to both PB and TV); an adrenergic receptor variant affecting prazosin binding affinity belongs in the lumped-sufficient column (pharmacological mechanism, no PB coupling).

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The decision criterion follows from an algebraic identity established in section 6: the lumped biomarker-treatment slope decomposes as $\beta_{bm}^{lumped} = \beta_{bm}^{BR} + w_{PB}\beta_{bm}^{PB} + w_{TV}\beta_{bm}^{TV}$, where the loadings w_{PB} and w_{TV} are determined by the trial design and the β_{bm}^C terms are component-specific slopes. Bias is zero if and only if $\beta_{bm}^{PB} = \beta_{bm}^{TV} = 0$, that is, if and only if the biomarker is uncorrelated with the non-pharmacological components. Whether that condition holds is determined by biology and domain knowledge, not by the magnitude of PB or TV in the study population. A population with a large placebo response but a biomarker orthogonal to it yields an unbiased lumped estimate; a population with a small placebo response but a biomarker that correlates with it yields a biased one.

COMT genotype exemplifies the first category. COMT modulates catechol-O-methyltransferase activity in prefrontal cortex, regulates dopaminergic signalling in expectancy circuits, and predicts placebo-response magnitude in IBS and chronic pain populations^{6,7}. A trial using COMT as a predictive biomarker for prazosin response would face a lumped slope that conflates pharmacological prediction with expectancy moderation. CYP2D6 metabolizer phenotype exemplifies the second: the genotype determines hepatic drug metabolism and thereby plasma drug concentration, with no plausible pathway to expectancy or natural-history effects. ADRB2 Arg16Gly in asthma bronchodilator trials provides a further example: spirometric outcomes are objective and receptor binding affinity does not modulate belief. For both, the lumped pharmacological slope is already the target estimand.

[3]

An algebraic identity establishes the conditional rule exactly; Study A confirms the orthogonal arm; Study B (pre-registered, pilot complete) addresses the coupling arm; together they bound the inferential cost of the choice.

- The identity (section 6) is exact: bias in the lumped slope =

196 $w_{PB}\beta_{bm}^{PB} + w_{TV}\beta_{bm}^{TV}$; it is zero if and only if the biomarker is
197 orthogonal to PB and TV , regardless of the magnitude of those
198 components.

- 199 • Study A (section 7.3; 1,000 replicates per alternative cell, 5,000
200 per null) confirms the orthogonal arm: mean absolute bias in the
201 lumped $\hat{\beta}_{bm:D_{bc}}$ is 0.016 against a true value of -2.25 at $N = 70$,
202 with no trend across $m_{PB} \in \{0, 1, 3, 6, 10\}$ or $m_{TV} \in \{-1, 0, 1, 2\}$.
- 203 • Study B (section 7.4; 100-replicate pilot reported here; produc-
204 tion run pre-registered) addresses the coupling arm: pilot data
205 show bias scaling with $c_{bm,PB}$, consistent with the algebraic pre-
206 diction and motivating the production run.

207 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
208 *tempor incididunt ut labore et dolore magna aliqua.*

209 The paper proceeds as follows. Section 2 fixes compact notation and
210 states the formal model. Section 3 develops each component with
211 biological motivation, parametric form, and identifiability conditions.
212 Section 4 maps the hybrid open-label-plus-blinded-discontinuation de-
213 sign’s phase structure onto the decomposition’s identifiability require-
214 ments. Section 5 develops the analysis-side methodology, including the
215 phase-augmented LME approximation. Section 6 works through the
216 numerical illustration and derives the algebraic identity that anchors
217 the conditional rule. Section 7 presents Study A (orthogonal arm) and
218 the Study B pilot (coupling arm). Section 8 translates the conditional
219 rule into a practical biomarker-validation decision criterion under the
220 PATH framework⁸. No published N-of-1 trial has applied this evalua-
221 tion framework; the paper’s primary contribution is determining when
222 decomposition changes the inferential answer rather than advocating
223 for it as a universal default.

224 2 Notation and the formal model

225 [4]

226 **This section fixes compact notation and states the formal**
227 **model; a fuller worked-example treatment is in the compan-**
228 **ion pedagogical document.**

- 229 • Definitions are precise enough to support the identifiability and
230 analysis-side arguments in later sections.
- 231 • The presentation is kept compact; extended derivations are de-
232 ferred to the companion document.
- 233 • The companion pedagogical document provides numerical worked
234 examples at each formal step.

235 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
236 *tempor incididunt ut labore et dolore magna aliqua.*

In this section we shall fix notation and state the formal model. The three components are defined precisely enough to make the identifiability and analysis-side arguments in later sections rigorous, but we have endeavored to keep the presentation compact. A fuller treatment, with worked numerical examples at each step, is available in the companion pedagogical document at docs/24-component-decomposition-pedagogy.md.

[5] where Y_{it} is the observed symptom score, BL_i is participant i 's baseline score, the three components BR , PB , TV are non-negative reductions in symptom severity attributable to the three causal mechanisms identified in section 1, and ε_{it} is observation-level noise. The sign convention is that an increase in any component reduces Y ; all three components are zero at $t = 0$ by construction (no exposure yet to drug, no belief activation, no elapsed natural history) and grow over the duration of the trial.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

For participant i at trial timepoint t :

$$Y_{it} = BL_i - [BR_{it} + PB_{it} + TV_{it}] + \varepsilon_{it},$$

[6]

Each component is driven by its own time-on-state variable, making the components structurally distinct even when they overlap in calendar time.

- BR_{it} accumulates with time-on-drug; PB_{it} with time-on-belief scaled by expectancy $\eta(\text{phase})$; TV_{it} with calendar time from trial entry.
- The distinct time indices are the structural source of identifiability from phase-contrast data.
- All components are zero at $t = 0$ by construction and grow over the trial duration.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Each component is a participant-specific function of its own time-on-state variable. BR_{it} depends on time-on-drug, the cumulative exposure to active medication since drug initiation. PB_{it} depends on time-on-belief, modulated by a phase-specific expectancy factor $\eta(\text{phase})$ that captures the patient's confidence that they are receiving active treatment. TV_{it} depends on time-since-trial-entry, calendar time elapsed regardless of treatment status.

[7] rescaled so that $C(0) = 0$ and $C(t) \rightarrow m_C(1 - e^{-d_C})$ as $t \rightarrow \infty$, for $C \in \{BR, PB, TV\}$. The three Gompertz parameters (m_C, d_C, r_C) are participant-specific random effects drawn from population distributions, and the participant-level heterogeneity in these distributions is the principal source of between-person response variability that the framework is designed to characterise.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Each component follows a modified Gompertz curve:

$$C(t) = m_C \exp(-d_C \exp(-r_C t)) - m_C \exp(-d_C),$$

[8]

The residual term has AR(1) structure and the three components are permitted to co-vary through cross-component covariances.

- Within-participant residuals follow $\text{Cov}(\varepsilon_{it}, \varepsilon_{i,t+h}) = \sigma_\varepsilon^2 \rho^{|h|}$.
- Within-time and across-time cross-component covariances permit participant-level heterogeneity in one component to co-vary with the others.
- This structure accommodates biologically plausible correlations, such as high-*BR* patients also tending toward high-*PB*.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

ε_{it} is a within-participant residual term with AR(1) serial-correlation structure $\text{Cov}(\varepsilon_{it}, \varepsilon_{i,t+h}) = \sigma_\varepsilon^2 \rho^{|h|}$. The three components are permitted to be jointly correlated through within-time and across-time cross-component covariances c_{cf1t} and c_{cfct} respectively, so that participant-level heterogeneity in any one component can co-vary with heterogeneity in the others.

The principal psychometric and pharmacological terms used below are summarised in Table 1.

3 The three components

[9]

This section presents each component with its biological motivation, Gompertz parameterisation, and identifiability conditions to establish the model as causally interpretable.

- The three components are presented in turn: *BR* (biological response), *PB* (placebo-belief), *TV* (natural history).

Table 1: Principal psychometric and pharmacological terms used in the three-component decomposition.

Term	Meaning
BR	Biological response: pharmacological reduction in symptom score attributable to active drug.
PB	Placebo-belief response: reduction attributable to the patient’s belief that they are receiving treatment.
TV	Time-variant component: reduction (or increase) attributable to the natural-history trajectory of the disease.
$\eta(\text{phase})$	Expectancy multiplier for <i>PB</i> : 1 in open-label on-drug, ≈ 0.5 in blinded discontinuation, between on-drug and off-drug.
BL_i	Participant i ’s baseline (pre-treatment) symptom score.
D_{it}	Continuous-decay drug indicator on the analysis side: 1 on-drug, $\exp(-\lambda t_{sd,it})$ off-drug.
$t_{sd,it}$	Time since drug discontinuation (off-drug only).
λ	Carryover decay rate, with half-life $t_{1/2} = \ln 2 / \lambda$.

- Each component is grounded in independent biological literature before its parametric form is stated.
- The section’s purpose is to demonstrate that the model sum reflects causal structure, not mathematical convenience.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In this section we shall present the three components in turn, each with its biological motivation, the modified Gompertz parameterisation, and the design contrasts that identify it. The treatment is organised so that the reader can see, by the end of this section, that the formal model is a sum of three causally interpretable trajectories rather than a mathematical convenience.

3.1 Biological response (BR)

[10]

***BR* is the pharmacological target of regulatory approval and biomarker stratification, independent of belief or natural history.**

- *BR* arises from chemical action on the body and is zero for inert pills regardless of patient expectation.
- It is the estimand that approval decisions, dose selection, and biomarker-guided stratification require.
- Separating *BR* from *PB* and *TV* is the central precision-medicine problem the framework addresses.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

BR is the physiological reduction in symptom score attributable to the active medication’s chemical action on the body. It is independent of

belief, suggestion, and natural history: a participant who somehow received active drug without knowing it would still develop a measurable BR , and a participant who took an inert sugar pill while believing it was active drug would have $BR = 0$ no matter how strong their expectation. Methodologically, BR is the quantity that drug regulators and prescribers want: approval decisions, dose selection, and biomarker-guided patient stratification are all framed around BR .

[11]

PK/PD considerations produce a saturating BR trajectory with a slow initial rise, a peak-growth phase, and a sharp drop to zero on discontinuation.

- Plasma steady state is reached within a few half-lives; pharmacodynamic integration (e.g. receptor adaptation) adds additional lag.
- For prazosin-PTSD, Raskind et al. (2013) calibrate the onset to several weeks of sustained exposure.
- The saturating shape motivates the Gompertz over linear or simple exponential alternatives.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The qualitative shape of the BR trajectory follows from pharmacokinetic and pharmacodynamic considerations. Drug concentration in plasma reaches steady state within a small multiple of the elimination half-life; brain concentration follows plasma with a short delay; the downstream pharmacodynamic response (reduction in symptom-relevant neural activity, in the case of prazosin's effect on noradrenergic tone during REM sleep) typically requires several days to weeks of sustained exposure to integrate fully⁹. The result is a saturating curve: slow rise in the first week, steepest growth in the second to third week, plateau as receptor adaptation completes, and a decay back toward zero on discontinuation governed by the carryover half-life $t_{1/2} = \ln 2/\lambda$ rather than an instantaneous drop. The rate of this off-drug decay is the object of the carryover machinery used throughout this project; only in the limit of a short half-life relative to the measurement spacing does it approximate the sharp step assumed in the worked example of section 6.

[12]

The three Gompertz parameters (m_{BR}, r_{BR}, d_{BR}) capture ceiling, onset rate, and climb shape respectively, calibrated to prazosin-PTSD reference values.

- m_{BR} is the biological ceiling (maximum reduction at infinite exposure); r_{BR} the onset rate; d_{BR} the displacement shaping the

- early climb.
- Prazosin-PTSD calibration yields $m_{BR} \approx 11$, $r_{BR} \approx 0.42/\text{week}$, $d_{BR} \approx 5$.
 - These are participant-specific random effects whose population distributions encode between-person response heterogeneity.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The modified Gompertz form $BR(t) = m_{BR} \exp(-d_{BR} \exp(-r_{BR}t)) - m_{BR} \exp(-d_{BR})$ captures this profile compactly. The asymptote $m_{BR}(1 - e^{-d_{BR}})$ is the biological ceiling, the largest reduction the drug can produce at infinite exposure; the scale parameter m_{BR} equals that ceiling up to the factor $(1 - e^{-d_{BR}})$, which for the calibrated $d_{BR} \approx 5$ exceeds 0.99. The onset rate r_{BR} controls how rapidly the participant approaches the ceiling: high r_{BR} means a fast onset (effect plateaus in one to two weeks), low r_{BR} a slow climb (six to eight weeks). The displacement d_{BR} shapes the climb: high d_{BR} gives an early-onset curve that rises sharply at first, low d_{BR} a more gradual climb. For PTSD prazosin, calibrated values from the¹⁰ reference dataset are $m_{BR} \approx 11$ (eleven points of nightmare-score reduction at saturation), $r_{BR} \approx 0.42$ per week (half-maximum at roughly four to five weeks of sustained exposure), and $d_{BR} \approx 5$ (moderately steep early climb).

[13]

The Gompertz is preferred over linear and simple-exponential alternatives because it captures the asymmetric slow-start saturation profile of chronic-condition pharmacodynamics.

- Linear $BR = \beta t$ is biologically implausible past the first few weeks due to unbounded growth.
- Exponential $BR = m(1 - e^{-rt})$ lacks the slow-start phase from receptor adaptation and gene-expression changes.
- The Gompertz interpolates between both limits and its three-parameter form avoids overfitting; sensitivity analysis is reported separately in 07-gompertz-evaluation.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Two alternative families that do not capture the asymmetric saturation profile are worth noting for contrast. A linear function $BR = \beta t$ climbs without bound and is biologically implausible past a few weeks. An exponential decay $BR = m_{BR}(1 - \exp(-rt))$ is acceptable for fast-onset drugs but lacks the slow-start phase that characterises most chronic-condition responses, where receptor adaptation, feedback regulation, and gene-expression changes produce a measurable lag. The

Gompertz interpolates between exponential decay (high d limit) and linear-in- t growth (low d limit), and its three-parameter form is flexible enough to fit a wide range of pharmacodynamic profiles without overfitting. The Gompertz is a deliberate choice over the two parametric forms most familiar from pharmacokinetic-pharmacodynamic modelling. The sigmoid $E_{\max}$ (Hill) model¹¹ describes the dependence of effect on concentration (a dose-response curve), whereas $BR(t)$ here is a time-course of accumulating effect at fixed dosing; and the indirect-response (turnover) models that do describe effect time-courses are specified through differential equations on a hidden response variable rather than in closed form. The modified Gompertz retains the slow-start, saturating shape those models produce while remaining a closed-form, three-parameter mean function that is tractable inside a mixed-effects likelihood. Doc 25 of the project documentation provides the full historical and biological context for the family choice and quantifies the sensitivity of pmsimstats inference to alternative families; that work is reported separately as 07-gompertz-evaluation. The headline result of that evaluation is reassuring for the present framework: in a 1{,}000-replicate-per-cell simulation that generated data under the logistic, hyperbolic-tangent, and piecewise-linear breakpoint families and analysed them under the standard linear-mixed model, the trajectory family was essentially irrelevant to the power of the biomarker-treatment interaction test (the family-to-family spread fell within two Monte Carlo standard errors of sampling noise) and to type I error (0.025 to 0.036 across null cells, with no family-specific inflation). The conclusions drawn here are therefore not artefacts of the Gompertz form, though the regimes the family cannot represent (cyclical or biphasic trajectories) remain genuine exclusions.

[14]

BR is identified by the blinded-discontinuation contrast, where drug removal collapses BR to zero while PB remains intact.

- Within-participant on-drug versus off-drug contrasts are the primary identification mechanism.
- The blinded-discontinuation phase provides the cleanest contrast: BR drops, PB is preserved.
- This contrast does not require the analyst to model PB or TV explicitly to recover BR .

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The BR component is identified by within-participant on-drug versus off-drug contrasts, formalised in section 4 below. The most direct contrast is the blinded-discontinuation phase: when a participant is

silently switched to placebo while still believing they are on active drug, BR decays toward zero as the drug clears while PB remains roughly intact, and the difference between on-drug and off-drug response in that window is a relatively clean estimate of BR that does not require the analyst to model PB or TV explicitly. The decay is not instantaneous: with the carryover half-life $t_{1/2} = \ln 2 / \lambda$ of the analysis-side decay term, residual BR persists into the early off-drug observations, so the contrast is clean only once the discontinuation window is long relative to $t_{1/2}$, and the drug state is best carried by the continuous-decay indicator D_{it} of Table 1 rather than a hard on/off indicator.

3.2 Placebo-belief response (PB)

[15]

PB is the strict placebo response: improvement that persists when the active drug is silently replaced, provided the patient's belief is unchanged.

- PB is belief-dependent, not drug-dependent; it would survive a silent switch to placebo under sustained belief.
- It is distinct from BR in causal origin: BR requires chemical action; PB requires only belief.
- The definition is strict and operational, not vague or colloquial.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

PB is the reduction in symptom score that arises from the patient's belief that they are receiving active treatment, independent of what they are actually receiving. That is, it is the placebo response in its strict sense: the part of the improvement that would persist if the active drug were silently replaced with an inert pill, provided the patient continued to believe they were on the active drug.

[16]

PB reflects neurobiologically real mechanisms triggered by belief, not pharmacology; neuroimaging and genetic evidence confirm it is a structured biological phenomenon.

- Belief-driven endogenous opioid/dopamine release, HPA-axis modulation, and top-down attentional shifts are all documented mechanisms.
- Wager et al. neurologic pain signature dissociates pharmacological from expectancy contributions via fMRI.
- COMT genotype predicts placebo responsiveness, confirming biological structure rather than measurement artefact.

505 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
506 *tempor incididunt ut labore et dolore magna aliqua.*

507 It is worth being precise about what PB is and is not. PB is not
508 ‘fake’ improvement; it is a phenomenon clinicians knowingly exploit, a
509 national survey of US internists and rheumatologists finding roughly
510 half report regular use of placebo treatments¹². As the introduc-
511 tion sets out, the placebo response is mediated by well-characterised
512 neurobiological pathways^{13–17}: belief-driven release of endogenous opi-
513 oids and dopamine, anticipatory hypothalamic-pituitary-adrenal reg-
514 ulation, and top-down attentional shifts, dissociated from pharmacol-
515 ogy by the⁴ neurologic pain signature⁵ and shown to have genetic
516 correlates^{6,7}. What makes these mechanisms ‘placebo’ rather than
517 ‘drug’ is that they are triggered by belief rather than by pharmacol-
518 ogy.

519 [17]

520 **PB is identifiable from BR only when belief and treatment**
521 **allocation differ, as in the blinded-discontinuation phase.**

- 522 • The blinded-discontinuation phase is the canonical contrast: pa-
523 tients cannot tell whether they remain on active drug.
- 524 • Belief drops to a partial level ($\eta \approx 0.5$) because the patient’s
525 expectation is hedged, not absent.
- 526 • Designs without a phase that decouples belief from treatment
527 cannot separately identify PB from BR .

528 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
529 *tempor incididunt ut labore et dolore magna aliqua.*

530 For trial design, the relevant point is that PB is separately identi-
531 fiable from BR if and only if the trial includes a phase in which the
532 patient’s belief about treatment differs from the actual treatment they
533 are receiving. The blinded-discontinuation phase is the canonical such
534 phase: the patient is told that, at some unannounced point during
535 the window, they may be silently switched to placebo, and they can-
536 not reliably tell whether they are still receiving active drug. Belief in
537 treatment correspondingly drops to a partial level (multiplier roughly
538 0.5) because the patient’s expectation is hedged.

539 [18]

540 **The $\eta = 0.5$ blinded-phase expectancy is a modelling assump-**
541 **tion operationalising response expectancy, not a measured**
542 **quantity, and requires sensitivity analysis.**

- 543 • The construct is response expectancy in the sense of Kirsch
544 (1985); to our knowledge the residual expectancy under blinded
545 discontinuation has not been measured directly.

- Open-label placebo trials confirm clinically meaningful PB -only responses even under explicit disclosure, bounding the residual response away from zero.
- Sensitivity analyses varying η are mandatory before clinical conclusions depend on this single parameter.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

PB also follows a Gompertz curve, scaled by the expectancy factor:

$$PB(t \mid \text{phase}) = \eta(\text{phase}) \cdot m_{PB} \exp(-d_{PB} \exp(-r_{PB}t)) + (\text{offset to start at zero}).$$

Typical expectancy values are $\eta = 1.0$ in open-label on-drug, $\eta \approx 0.5$ in blinded discontinuation, and η between 0 and 0.5 in open-label placebo (the residual expectation a patient retains when told the previous active phase was beneficial). The quantity η operationalises response expectancy in the sense of¹⁸, the patient's anticipation of benefit that the expectancy literature treats as a determinant of experienced symptom change. The choice of $\eta = 0.5$ for the blinded phase is a modelling assumption rather than a calibrated quantity: to our knowledge the residual expectancy a patient carries when uncertain about allocation has not been measured directly as a fraction of the open-label level, and we adopt 0.5 as a deliberately central placeholder. The available indirect evidence is nonetheless that the residual response is substantial. Open-label placebo trials^{19–21} demonstrate that a clinically meaningful response persists even when patients are explicitly told they are receiving an inert intervention, which bounds the residual belief-driven response away from zero. Because the value of η is not empirically pinned, sensitivity analyses varying it across its plausible range are mandatory before any conclusion is allowed to depend on this single number.

[19]

Kaptchuk et al. (2008) provide the structural precedent; the BR-PB-TV framework extends the same decomposition logic to the full treatment response.

- Kaptchuk decomposed the placebo response itself into observation, ritual, and relationship components empirically.
- The present framework applies analogous decomposition logic at the level of drug, placebo, and natural history.
- Both frameworks reject the 'expectancy lump' model in favour of identifiable additive parts.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod

584 *tempor incididunt ut labore et dolore magna aliqua.*

585 The decomposition framework as proposed here has a structural ana-
586 logue in the placebo-research literature.³ decomposed the placebo re-
587 sponse itself into observation, ritual, and patient-practitioner relation-
588 ship components in a randomised trial of acupuncture-style placebo for
589 irritable bowel syndrome, demonstrating empirically that the placebo
590 response is not a single ‘expectancy’ lump but a sum of identifiable
591 parts. The present BR-PB-TV framework is analogous in spirit but
592 operates at the level of the full treatment response (drug, placebo,
593 natural history) rather than within the placebo response alone.

594 [20]

595 ***PB* variance scales with η , not just its mean; ignoring this**
596 **heteroscedastic structure miscalibrates standard errors and**
597 **downstream biomarker tests.**

- 598 • Phase-dependent variance follows $\text{SD}(PB \mid \text{phase}) = \eta(\text{phase}) \cdot$
599 $\sigma_{PB}^{\text{baseline}}$.
- 600 • Open-label phases have high between-patient *PB* variance;
601 blinded phases have low variance.
- 602 • Ignoring heteroscedasticity deflates standard errors in open-label
603 and inflates them in blinded phases, with consequences for CI
604 coverage and biomarker test calibration.

605 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
606 *tempor incididunt ut labore et dolore magna aliqua.*

607 A subtle property of *PB* is that its variance also scales with η ,
608 not just its mean. When η is high, there is substantial variability
609 across patients in how strongly the placebo response is expressed;
610 high-suggestibility patients show large *PB*, low-suggestibility pa-
611 tients show essentially none. When η is low (blinded), the placebo
612 response is uniformly attenuated and the patient-to-patient vari-
613 ance shrinks accordingly. The model therefore parameterises
614 $\text{SD}(PB \mid \text{phase}) = \eta(\text{phase}) \cdot \sigma_{PB}^{\text{baseline}}$. Ignoring this heteroscedastic
615 structure deflates standard errors in the open-label phase and inflates
616 them in the blinded phase, with downstream consequences for
617 confidence-interval coverage and for the calibration of any biomarker
618 test built on top of the analysis.

619 3.2.1 Additivity, subadditivity, and the BR-PB interaction

620 *[Extended treatment in supplement S2.]*

621 [21]

622 **Additivity is a modelling assumption; published literature**
623 **documents subadditivity between pharmacological and belief**

624 **effects in some contexts; the phase-augmented analysis partly**
625 **absorbs it as a side benefit.**

- 626 • The generalised model adds $\gamma \cdot BR_{it} \cdot PB_{it}$; $\gamma < 0$ encodes subad-
627 ditivity, $\gamma = 0$ recovers the additive form.
- 628 • Apparent subadditivity can arise from latent-class structure with
629 class-correlated BR and PB ; distinguishing true interaction from
630 confounding requires a class-aware analysis or the phase contrasts
631 of section 4.
- 632 • Simulation quantifying bias scaling with γ is deferred to the pro-
633 duction programme (section 7).

634 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
635 *tempor incididunt ut labore et dolore magna aliqua.*

636 The additive decomposition of section 2 treats the three components as
637 independent contributors. Subadditivity – where the combined $BR +$
638 PB effect falls below the sum of the components measured separately
639 – is documented in some clinical contexts^{3,14,22} and can arise from
640 shared neural substrate, ceiling effects, or belief-dynamics adaptation.
641 A generalised model with $\gamma \cdot BR_{it} \cdot PB_{it}$ is identifiable from the same
642 phase contrasts. Apparent subadditivity can also arise from latent-
643 class structure with class-correlated components rather than a direct
644 BR - PB interaction; disentanglement requires a class-aware analysis
645 or the phase contrasts provided here. The phase-augmented analysis
646 of section 5 absorbs subadditivity partly as a side benefit.

647 3.3 Time-variant natural-history component (TV)

648 [22]

649 **TV is the untreated counterfactual trajectory; it is**
650 **participant-specific, structured, and can be positive,**
651 **negative, or near-zero depending on disease class.**

- 652 • TV is the symptom change that would have occurred without
653 treatment, integrated over the trial duration.
- 654 • It varies in sign across disease classes and across patients within
655 a class.
- 656 • It is independent of treatment status and belief, distinguishing it
657 structurally from PB .

658 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
659 *tempor incididunt ut labore et dolore magna aliqua.*

660 TV is the reduction in symptom score that would have happened even
661 if the patient had received no treatment at all. It is the participant's
662 natural-history trajectory, integrated over the duration of the trial.
663 For some participants TV is positive (symptoms naturally improve

over the trial because of life events, therapy, time-in-condition, or regression to the mean); for others TV is negative (symptoms naturally worsen because of trauma anniversaries, accumulating stress, or progressive underlying disease).

[23]

TV is a structured participant-specific signal, not noise; its distinguishing feature from PB is independence from treatment status and belief.

- Unlike ε_{it} , TV is participant-level (slow trajectory) and requires explicit modelling, not absorption into noise.
- Off-drug baselines and follow-up windows are the primary data sources for estimating TV .
- TV is modulated by neither treatment assignment nor patient belief, whereas PB is modulated by both.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Like PB , TV is not noise. It is a structured participant-specific signal that the model can in principle estimate from data, given sufficient repeated measurements and a trial design that includes a long enough off-drug baseline or follow-up window. What distinguishes TV from PB is its origin: TV is independent of treatment status and of belief about treatment, whereas PB is modulated by both.

3.3.1 Natural history: existence, testability, and functional form

[24]

TV sign varies by disease class; the zero- TV assumption is rarely safe in chronic trials; a linear week covariate is insufficient because it enforces population-mean linearity and cannot separate TV from BR without the design contrasts of section 4.

- Disease class determines TV sign: positive for acute recovery, near-zero for stable chronic, negative for progressive conditions; selection-driven enrolment adds a positive TV component regardless of class due to regression to the mean [hrobjartsson2001; hrobjartsson2010].
- Zero- TV is safe only when all four conditions hold: stable disease, short trial, non-severity-peaked recruitment, no ancillary care; in most chronic trials at least one fails.
- A week covariate enforces linearity, ignores participant-level heterogeneity, and conflates BR with TV ; the Gompertz form and the design contrasts of section 4 supply identifiability.

704 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
705 *tempor incididunt ut labore et dolore magna aliqua.*

706 The shape of TV depends on disease class, trial timescale, and enrol-
707 ment selection. Acute conditions trend TV -positive (intrinsic recovery); chronic stable conditions trend near $TV = 0$ on the population
708 mean but with substantial participant-level variance; progressive conditions trend TV -negative. Selection-driven enrolment adds a positive
709 TV component regardless of class, because patients enrol at moments
710 of peak severity and regression to the mean then contributes apparent
711 improvement in early trial weeks^{23,24}. The zero- TV assumption
712 is testable: a model including TV returns $\hat{m}_{TV}$ with a variance estimate; lumping TV into ε is worse, because the participant-level
713 natural-history trajectory is structured and correlated with the on/off-
714 drug schedule, producing autocorrelated within-participant residuals
715 and heteroscedastic between-participant residuals that bias the standard
716 errors on BR and PB . A simple `week` covariate is insufficient:
717 it enforces a linear, population-mean trajectory and, without the design
718 contrasts of section 4, absorbs both BR and TV into a single
719 coefficient that identifies neither.
720
721
722

723 4 Identifiability and trial-design contrasts

724 [25]

725 **Identifiability is supplied by trial design; each phase of the**
726 **hybrid design contributes a different combination of compo-**
727 **nents to the observed trajectory.**

- 728 • The decomposition is mathematically specified but requires trial-
729 design contrasts to be statistically identified.
- 730 • This section characterises the hybrid open-label-plus-blinded-
731 discontinuation design as the primary identifiability vehicle.
- 732 • Each of the three phases supplies a distinct mixture of BR , PB ,
733 and TV .

734 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
735 *tempor incididunt ut labore et dolore magna aliqua.*

736 The mathematical decomposition above is meaningless unless the data
737 contain enough information to pin down each component separately.
738 Trial design is what supplies that information. In this section we
739 shall characterise the hybrid open-label-plus-blinded-discontinuation
740 design, used as the worked example throughout this paper, and show
741 how each of its three phases contributes a different combination of
742 components to the observed trajectory.

Phase	Drug?	Belief?	Components present
Open-label on drug	Yes	Full	$BR + PB(\eta = 1) + TV$
Blinded discontinuation	Hidden	Hedged	$BR \cdot \mathbf{1}_{\text{on drug}} + PB(\eta = 0.5) + TV$
Open-label crossover	Either	Knows	$BR \cdot \mathbf{1}_{\text{on drug}} + PB(\eta = 1) + TV$

[26]

Three phase-by-allocation contrasts isolate PB , BR , and TV respectively; the LME model combines them but the inference is driven by design.

- Open-label versus blinded contrast yields approximately $0.5 \cdot PB$, isolating the belief component conditional on a known blinded-phase η .
- Blinded on-drug versus blinded placebo yields BR once drug carryover has elapsed (since PB and TV match across arms within the blinded phase).
- Off-drug phases yield TV plus residual off-drug PB , with no active BR contribution beyond carryover.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Three phase-by-allocation contrasts then identify the three components. The open-label versus blinded contrast (within drug) yields $PB(\eta = 1) - PB(\eta = 0.5) \approx 0.5 \cdot PB$, which isolates the belief component, but only conditional on a known blinded-phase expectancy η : from a two-phase design just the product $\eta \cdot m_{PB}$ is identified, and separating η from m_{PB} requires the open-label placebo phase. The blinded on-drug versus blinded placebo contrast yields BR once drug carryover has elapsed, since PB and TV match across allocation arms within the blinded phase. All phases off drug yield TV plus any residual off-drug PB , with no active BR contribution beyond carryover and PB at the off-drug expectation level. The $\mathbf{1}_{\text{on drug}}$ indicator in the phase table above is shorthand for the continuous-decay drug state D_{it} defined in Table 1.

This conditioning on η is the framework’s principal identification limitation and deserves emphasis rather than a parenthetical. With a two-phase (open-label and blinded) design the data identify only the product $\eta \cdot m_{PB}$, so PB is recovered up to the assumed blinded-phase expectancy; the working value $\eta \approx 0.5$ is a modelling assumption, not an estimate, and absent a belief or expectancy measurement it is not falsifiable from the outcome data alone. The identified set for m_{PB}

is accordingly the interval $\{\widehat{\eta m_{PB}}/\eta : \eta \in (0,1]\}$, which collapses to a point only when η is pinned down. Two design features close the gap: an open-label placebo phase, which supplies a second equation separating η from m_{PB} , and direct measurement of belief, as in the open-hidden and balanced-placebo designs^{25,26} that manipulate or record the patient's treatment expectation. We therefore recommend that any trial relying on the decomposition for a PB -versus- BR attribution either include such a phase or collect an expectancy measure, and that analyses report sensitivity of the attribution to η across its plausible range.

[27]

Designs lacking specific phases lose identifiability of the corresponding component; the decomposition machinery is only as useful as the design allows.

- A pure open-label design cannot separate BR from PB because both are present at full strength throughout.
- A pure blinded design cannot estimate $\eta = 1$ PB because no phase achieves full-strength belief.
- Trial designers should verify each component has at least one phase in which it is uniquely present before committing to the design.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A trial that lacks any of these phases loses the corresponding identifiability. A pure open-label design (no blinded discontinuation) cannot separate BR from PB , because both are present at full strength throughout. A pure blinded design (no open-label phase) cannot estimate $\eta = 1$ PB , because no phase has belief at full strength. The decomposition machinery is useful only to the extent that the design supplies identifiability, and trial designers should evaluate any proposed design by checking that each component has at least one phase or contrast in which it is uniquely present.

[28]

Design-contrast identifiability is necessary but not sufficient: the three Gompertz trajectories are collinear, so joint estimation is credible only under specific design and sample-size conditions.

- The components share an onset shape and overlap in calendar time, so their parameters can be near-collinear even when the identifying contrasts are present.
- Joint recovery requires a phase in which each component varies

818 while the others are held fixed, and an off-drug window long
819 enough for TV to separate from a decaying BR .
820 • The recoverable region of the design-and-sample-size space is
821 mapped by the pre-registered identifiability diagnostics, not as-
822 serted.

823 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
824 *tempor incididunt ut labore et dolore magna aliqua.*

825 Identifiability in the design-contrast sense above is necessary but not
826 sufficient for stable estimation at finite sample sizes. The three Gompertz trajectories share an onset shape and overlap in calendar time, so
827 their parameters can be collinear even when the contrasts that identify
828 them in principle are present. Two conditions make the joint estimation credible: the design must supply at least one phase in which each
829 component varies while the others are held fixed (the blinded-placebo
830 contrast for BR , the open-label-versus-blinded contrast for PB , and
831 an off-drug window for TV), and that off-drug window must be long
832 enough relative to the carryover half-life for TV to separate from a
833 still-decaying BR . Where these conditions are weak the components
834 are jointly identified only up to a near-collinear direction, and the
835 symptom of that is the rank-deficiency the pilot of section 7 encountered
836 when the full phase-augmented formula was fitted at $N = 35$.
837 We therefore do not assert that the three components are always recoverable.
838 The pre-registered identifiability diagnostics of section 7 (convergence rates of the component-matched fit, variance-inflation factors
839 on the component-specific Gompertz parameters, and the sensitivity
840 of the BR estimate to the assumed $\eta(\text{phase})$ multipliers) are the means
841 by which the recoverable region of the design-and-sample-size space is
842 to be mapped, and the framework’s participant-level claims should be
843 read as conditional on falling inside that region.

847 5 Analysis-side methodology

848 [29]

849 **Four constraints make a component-matched analysis im-**
850 **practical at trial-relevant sample sizes; a phase-augmented**
851 **LME captures the important part of the gap at minimal cost.**

- 852 • The question of whether to mirror the DGP in the analysis model
853 is reasonable; the section addresses it directly.
- 854 • The four constraints are identifiability, degrees of freedom, con-
855 vergence fragility, and marginal estimand robustness.
- 856 • A phase-by-treatment indicator is the recommended modest ex-
857 tension that captures BR-versus-PB sensitivity without the full
858 parametric cost.

859 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
860 *tempor incididunt ut labore et dolore magna aliqua.*

861 In this section we shall take up the practical question of what analysis
862 model to use. A reasonable reaction to the formal model in section
863 2 is the following: if we are simulating data with three components
864 plus noise, why is the analysis model written as $Sx \sim \text{bm} + \text{t} + \text{Dbc}$
865 $+ \text{bm}:\text{Dbc}$ with a single random intercept and AR(1) residual corre-
866 lation? Should we not include three parametric Gompertz compo-
867 nents in the analysis, mirroring the DGP, so that the analyst sees the
868 same structure that generated the data? The question is reasonable.
869 Four constraints make a component-matched analysis impractical for
870 most aggregated N-of-1 trial sizes, and a modest extension (a phase-
871 by-treatment indicator) captures the practically important part of the
872 gap at minimal cost.

873 5.1 Four reasons not to match the DGP

874 [30]

875 **The first constraint is identifiability: forcing a three-**
876 **component model on a design with insufficient contrasts**
877 **produces unidentified parameters, not component estimates.**

- 878 • A pure open-label trial cannot distinguish *BR* from *PB*; both
879 rise from zero together on treatment initiation.
- 880 • Fitting three components to a one-contrast design produces im-
881 plausibly tight standard errors that misrepresent inferential pre-
882 cision.
- 883 • The analyst must audit whether the design supplies the necessary
884 contrasts before specifying a component-matched model.

885 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
886 *tempor incididunt ut labore et dolore magna aliqua.*

887 The first reason is identifiability. The trial design must supply the
888 contrasts that identify each component, and many designs supply only
889 part of it. The decomposition is identified by phase contrasts (open-
890 label versus blinded for *PB*, on-drug versus blinded-placebo for *BR*,
891 off-drug timepoints for *TV*). A pure open-label trial has none of these
892 contrasts; in such a trial *BR* and *PB* both rise from zero on treat-
893 ment initiation and are statistically indistinguishable from each other,
894 regardless of how the analyst writes the model. Forcing an analysis to
895 fit three components when the design supplies contrasts for only one
896 or two of them does not recover the missing components; it produces
897 unidentified or weakly identified parameter estimates with implausibly
898 tight standard errors that misrepresent the analyst's actual inferential
899 precision.

900

[31]

901

902

903

904

The second constraint is degrees of freedom: a component-matched analysis adds approximately a dozen non-linear parameters, inflating moderation estimand variance at $N \in [30, 150]$.

905

906

907

908

909

910

- Gompertz *BR* (3), *PB* (3 + expectancy multiplier), and *TV* (3) each require participant-specific random effects on the maxima.
- At trial-relevant N , the variance of $\hat{\beta}_{bm:D}$ is dominated by participant and timepoint counts.
- Additional parametric structure typically inflates moderation estimand variance more than it reduces bias.

911

912

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

913

914

915

916

917

918

919

920

921

922

923

924

925

926

The second reason is degrees of freedom. Each component term costs degrees of freedom that come out of the moderation estimand. The biomarker-moderation question (the core scientific aim of the predictive-biomarker validation programme) is asked through the `bm:Dbc` term. A component-matched analysis adds a Gompertz *BR* (three parameters), a Gompertz *PB* (three parameters plus a phase-specific expectancy multiplier), and a Gompertz *TV* (three parameters), each with participant-specific random effects on the maxima. That is on the order of a dozen new parameters, half of them on a non-linear scale. At trial sample sizes of $N \in [30, 150]$, the variance of $\hat{\beta}_{bm:D}$ is dominated by participant and timepoint counts, and adding parametric structure to absorb residual confounding usually inflates the moderation estimand's variance by more than it reduces bias.

927

[32]

928

929

930

The third constraint is convergence fragility: component-matched nonlinear mixed models converge below 90% at trial-relevant N , contaminating aggregated analyses.

931

932

933

934

935

936

937

938

- Simultaneous Gompertz *BR*, *PB*, and *TV* fitting requires `nlme::nlme`, Bayesian hierarchical, or structural-equation approximations.
- Each is more sensitive to starting values, prone to local optima, and likely to fail on individual trial datasets than the LME analogue.
- Non-converging cells at below-90% convergence rates contaminate power and bias estimates in simulation studies.

939

940

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

941 The third reason is convergence fragility. Component-matched anal-
942 yses are non-linear in the parameters and their convergence is prob-
943 lematic in practice. Fitting Gompertz *BR*, *PB*, and *TV* simulta-
944 neously requires a non-linear mixed model (`nlme::nlme` with self-
945 starting Gompertz forms), a Bayesian hierarchical fit, or a structural-
946 equation approximation. Each is more sensitive to starting values,
947 more prone to local optima, and more likely to fail to converge on a
948 single trial dataset than the linear-mixed analogue. At trial-relevant
949 N , the convergence rate of the component-matched fit is typically be-
950 low 90%, and the non-converging cells contaminate any aggregated
951 analysis.

952 [33]

953 **The fourth constraint is estimand robustness: the `bm:Dbc`**
954 **marginal interaction is the right estimand for the biomarker-**
955 **validation question regardless of the BR-PB split.**

- 956 • Regulators want to know whether the biomarker predicts treat-
957 ment response under trial conditions, not a mechanistically de-
958 composed prediction.
- 959 • The marginal `bm:Dbc` interaction captures this question whether
960 the moderation is biological or expectation-mediated.
- 961 • Component decomposition adds mechanistic precision but is not
962 required for the core biomarker-validation decision.

963 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
964 *tempor incididunt ut labore et dolore magna aliqua.*

965 The fourth reason is that the moderation estimand is sufficiently ro-
966 bust to the BR-versus-PB split for the most common scientific ques-
967 tion. For predictive-biomarker validation, the regulator wants to know
968 whether the biomarker predicts treatment response under the trial con-
969 ditions, not a clean pharmacological-versus-placebo decomposition of
970 that prediction. The simple `bm:Dbc` interaction is the right marginal
971 estimand for the validation question whether the moderation is biolog-
972 ical or expectation-mediated.

973 [34]

974 **The inferential target is stated as an estimand: the pop-**
975 **ulation biomarker-by-treatment interaction with explicit**
976 **intercurrent-event handling, while the pharmacological-**
977 **component target is more ambitious and inherits the**
978 **heterogeneity-identifiability caveat.**

- 979 • The $bm:D_{bc}$ estimand is defined by its population, treatment con-
980 trast, endpoint, summary, and intercurrent-event strategy in the
981 ICH E9(R1) sense.

- Dropout is handled by a treatment-policy strategy and residual carryover by the modelled exponential decay.
- Recovering a participant-level pharmacological interaction inherits Senn’s caveat that true heterogeneity is separable from within-person variation only under repeated phase contrasts.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

It is worth stating the inferential target of the analysis in the vocabulary of the estimands framework²⁷, because the biomarker-validation question is easy to pose ambiguously. The primary estimand is the population-level biomarker-by-treatment interaction, the $bm:D_{bc}$ coefficient: its population is the trial-eligible patient population, its treatment contrast is on-drug versus off-drug exposure summarised through the continuous indicator D_{bc} , its endpoint is the symptom trajectory, its population-level summary is the interaction coefficient, and its intercurrent events are handled by a treatment-policy strategy for dropout and by the modelled exponential decay for residual carryover. The pharmacological-component target β_{bm}^{BR} is a more ambitious estimand: it is the biomarker-by- BR interaction that would obtain with the belief-driven channel held at its off-drug level, identified here through the phase contrasts of section 4 rather than by assumption. We flag explicitly that recovering a participant-level interaction of this kind inherits the difficulty Senn has emphasised for individual treatment effects^{28,29}: a genuine participant-by-component signal is separable from participant-by-period and within-person random variation only when the design supplies repeated phase contrasts, and the framework should be read as estimating these quantities under that condition rather than as dissolving the well-known problem of identifying individual effects from finite within-person data.

5.2 A small extension that usually pays off: phase indicators

[35]

The phase-augmented LME uses `Dbc:phase` and `bm:Dbc:phase` to test whether drug and biomarker-by-drug effects attenuate under blinding.

- `phase` absorbs systematic differences in PB expectancy across trial phases.
- `Dbc:phase` detects whether the drug effect shrinks under blinding, a signature of PB -mediation.
- `bm:Dbc:phase` detects whether the biomarker-by-treatment interaction shrinks under blinding, a signature that it is partly biomarker-by- PB .

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The smallest useful step toward component-aware analysis, without paying the costs above, is to add a phase indicator and its interaction with the drug indicator. Schematically,

```
Sx ~ bm + t + Dbc + phase + Dbc:phase
      + bm:Dbc + bm:Dbc:phase
      + (1 | ptID), correlation = corCAR1(...)
```

Here `phase` is a categorical covariate (open-label, blinded, post-crossover) that absorbs the systematic differences in *PB* expectancy across phases. `Dbc:phase` tests whether the apparent drug effect shrinks under blinding, a signature that the open-label drug effect was partly *PB*-mediated. `bm:Dbc:phase` tests whether the moderation itself shrinks under blinding, a signature that the biomarker-by-treatment interaction was partly biomarker-by-*PB* rather than biomarker-by-*BR*.

[36]

The phase-augmented extension adds at most four parameters yet captures the practically relevant belief-attenuation diagnostic without requiring the Gompertz parametric form.

- Three to four parameters is well within budget at any realistic trial sample size.
- The extension does not separately estimate *BR* and *PB* components but identifies the change in apparent drug effect when belief is hedged.
- For most analyses this is the practically relevant question and is answerable without committing to the Gompertz form.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

This extension adds three to four parameters at most, well within budget at any realistic *N*. It does not separately estimate *BR* and *PB* as components, but it does identify the change in apparent drug effect that occurs when belief is hedged. For most analyses, this is the practically relevant question, and it can be answered without committing to the Gompertz parametric form. Where both the lumped and the phase-augmented analyses are reported, the additional `Dbc:phase` and `bm:Dbc:phase` terms introduce a small family of belief-attenuation tests, and the attribution inference they support should be treated as a pre-specified secondary analysis with its multiplicity acknowledged rather than as a second bite at the primary interaction test.

5.3 When to fit the full decomposition

[37]

The full decomposition merits the effort only when N is in the hundreds, the design is fully phase-balanced, and the scientific question requires the BR-PB mechanistic split.

- Feasibility conditions: several-hundred N , dense timepoints in every phase, mechanistic BR-PB split required, and informative Bayesian priors available.
- In that regime, `lcmm`, `nlme::nlme` with self-starting Gompertz forms, or `brms` can return participant-specific component estimates.
- Outside that regime the phase-augmented LME is the better tool.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A component-matched analysis merits the effort when sample size is in the several hundreds (a phase-3 trial fielded as aggregated N-of-1, or a meta-analysis pooling multiple smaller trials); the trial design provides timepoints in every phase combination, with enough density per cell to constrain the Gompertz curves; the scientific question requires the BR-PB split (a pharmacology paper claiming a specific mechanism, where ruling out a placebo contribution is the central point); and a Bayesian hierarchical fit with informative priors on the Gompertz parameters is feasible. In that regime, `lcmm`, `nlme::nlme` with self-starting Gompertz forms, or a `brms` hierarchical formulation can return participant-specific BR , PB , and TV component estimates that mirror the DGP. Outside that regime, the linear-mixed analysis with phase indicators is the better tool.

[38]

The DGP-analysis asymmetry is deliberate: simulation uses the rich DGP to characterise how a simpler analysis behaves; the analysis need not mirror the DGP to produce useful inference.

- Simulation controls the DGP and can specify any structure; inference must pay the identifiability cost of every parameter.
- The DGP is a characterisation tool; the analysis is a finite-data inference tool; they have different optimality criteria.
- DGP-matching is not required and is often counterproductive at trial-relevant sample sizes.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The asymmetry between DGP and analysis is not a flaw of the framework; it is a deliberate consequence of the asymmetry between simu-

lation (where the analyst controls the DGP and can specify any structure, however elaborate) and inference (where the analyst must fit a model to a finite dataset and pay the identifiability cost of every additional parameter). Simulation studies use the rich DGP to characterise how a simpler analysis behaves under realistic data generation; they do not require the analysis to match the DGP in order to produce useful inference.

[39]

The exponential-decay carryover indicator is a deliberately parsimonious special case of richer carryover models, and the simulation should benchmark it against a distributed-lag alternative.

- The continuous indicator D_{bc} decays as $\exp(-\lambda t_{sd})$, a one-parameter form tied to a pharmacokinetic half-life.
- Distributed-lag models, the average-period-treatment-effect estimand, and G-estimation relax the single-rate assumption at additional parametric cost.
- The exponential restriction is an assumption, not a finding, and the pre-registered simulation should include a distributed-lag arm to quantify its cost.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A word on the carryover parameterisation is in order, since it is the one place where the analysis model makes a strong functional-form commitment. The continuous drug indicator D_{bc} decays as $\exp(-\lambda t_{sd})$ off drug, a one-parameter exponential governed by the carryover half-life. This is deliberately parsimonious and is best understood as a constrained special case of the more flexible carryover models now available for repeated-measures and N-of-1 data. Liao et al.³⁰ model carryover through a Bayesian distributed-lag structure with autocorrelated errors, relaxing the single-rate exponential assumption at the cost of additional parameters; Daza³¹ formalises an average-period treatment-effect estimand built to accommodate autocorrelation, trend, and slow carryover jointly; and G{a}rtner et al.³² benchmark linear models against G-estimation and Bayesian-network alternatives specifically under carryover. We adopt the exponential form because it ties directly to a pharmacokinetic half-life and conserves degrees of freedom for the moderation estimand at trial-relevant sample sizes, but the choice is an assumption rather than a finding. The pre-registered simulation programme of section 7 should accordingly include a distributed-lag analysis arm, so that the cost of the exponential-decay restriction is quantified against the richer alternative rather than asserted.

6 Worked example: information loss without decomposition

6.1 What the lumped analysis estimates

[40]

The lumped treatment-effect estimate has a probability limit equal to the pharmacological component plus belief and natural-history contributions, so its bias does not vanish as N grows.

- Writing the reduction as $BR + PB + TV - \varepsilon$ and fitting one drug indicator gives $\text{plim } \hat{\beta}_D = m_{BR} + b_{PB} + b_{TV}$.
- b_{PB} and b_{TV} are the belief and natural-history terms that load onto the drug indicator under the chosen contrast; both are generally non-negative for an improving outcome.
- Because the displacement is in the probability limit, a larger one-component trial converges on the same biased value: the failure is identification, not precision.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Before the numerical illustration it is worth establishing algebraically what the one-component analysis actually estimates, because the qualitative failure modes of section 1 are consequences of an identity rather than empirical tendencies that require simulation to demonstrate. Write the within-participant symptom reduction as $R_{it} = BL_i - Y_{it} = BR_{it} + PB_{it} + TV_{it} - \varepsilon_{it}$. The one-component analysis fits a single drug indicator D_{it} ,

$$R_{it} = \alpha + \beta_D D_{it} + g(t) + u_i + e_{it},$$

and reports $\hat{\beta}_D$ as ‘the treatment effect’. Taking the on-drug minus off-drug contrast that β_D targets, and using that the natural-history component is by construction independent of treatment status so that its on/off difference cancels in a within-participant comparison, the probability limit of the estimator is

$$\text{plim } \hat{\beta}_D = m_{BR} + b_{PB} + b_{TV},$$

where $b_{PB} = E[PB \mid \text{on}] - E[PB \mid \text{off}]$ is the belief differential between the contrasted phases and b_{TV} collects any natural-history contribution that fails to cancel under the contrast actually used; b_{TV} is zero for a clean on/off within-participant contrast and strictly positive for a baseline-anchored pre-post analysis, in which case $b_{PB} = E[PB]$

and $b_{TV} = E[TV]$. To fix which case is in play below: the one-component analysis used in the worked example and the simulations of this paper is the design-contrast (drug-on/off) form, so $b_{TV} \approx 0$ there and the operative displacement is the belief term b_{PB} ; the baseline-anchored pre-post form, which additionally carries $b_{TV} = E[TV]$, is the worse case retained only for comparison. Both terms are generally non-negative for an improving outcome, so the lumped estimate over-states the pharmacological component m_{BR} . The displacement is a property of the probability limit and is unchanged as $N \rightarrow \infty$: a larger one-component trial converges on the same biased value. This is the precise sense in which the placebo-attribution failure of section 1 is a problem of identification, not of precision.

[41]

The biomarker slope from a lumped analysis decomposes additively into component-specific slopes, so a biomarker correlated with PB or TV registers as a predictor even with no pharmacological association.

- By linearity of covariance, $\beta_{bm}^{\text{lumped}} = \beta_{bm}^{BR} + w_{PB}\beta_{bm}^{PB} + w_{TV}\beta_{bm}^{TV}$.
- A non-zero β_{bm}^{PB} or β_{bm}^{TV} produces a non-zero lumped slope even when the pharmacological slope β_{bm}^{BR} is exactly zero.
- The result is an algebraic identity under the additive decomposition, not a sample-size-dependent tendency.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The same algebra settles the biomarker-validation failure. The precision-medicine analysis regresses a participant-level response summary on a baseline biomarker B_i . If that summary is the lumped per-participant effect, which estimates $m_{BR,i} + w_{PB}m_{PB,i} + w_{TV}m_{TV,i}$ for design-determined loadings $w_{PB}, w_{TV} \geq 0$, then by the linearity of covariance the fitted biomarker slope has probability limit

$$\beta_{bm}^{\text{lumped}} = \frac{\text{Cov}(m_{BR} + w_{PB}m_{PB} + w_{TV}m_{TV}, B)}{\text{Var}(B)} = \beta_{bm}^{BR} + w_{PB}\beta_{bm}^{PB} + w_{TV}\beta_{bm}^{TV},$$

where $\beta_{bm}^C = \text{Cov}(m_C, B)/\text{Var}(B)$ is the component-specific slope. A biomarker correlated with placebo responsiveness or with the natural-history trajectory, so that β_{bm}^{PB} or β_{bm}^{TV} is non-zero, therefore registers a non-zero lumped slope even when its pharmacological association β_{bm}^{BR} is exactly zero. This is the biomarker-validation failure of section 1, and it too is an identity under the additive decomposition rather than an empirical tendency: no sample size separates the lumped slope from a genuine pharmacological one, because the two coincide in the

estimand whenever $w_{PB}\beta_{bm}^{PB} + w_{TV}\beta_{bm}^{TV} \neq 0$.

We make two qualifications explicit, because they bound how much weight the framework can carry. First, both displays are immediate corollaries of the additive decomposition and the linearity of the projection that defines the estimator; they are not a derived result in any deeper sense, and their value is in naming the displacement terms, not in the algebra. Second, they assume the components combine additively. Under the subadditive interaction documented in section 3.3, where the joint effect of BR and PB falls below their sum, a first-order correction proportional to $\text{Cov}(m_{BR} \cdot m_{PB}, B)$ enters the biomarker slope, and the identities above should be read as the leading behaviour in the additive case rather than as exact under interaction. The empirical magnitude of that correction term is itself an open question that the framework's γ -interaction simulation (section 3.3) is designed to address.

[42]

The bias displaces the estimand rather than inflating variance, so a significance test of the lumped coefficient cannot detect it; the decomposition is what locates it.

- The lumped coefficient is displaced from m_{BR} by $b_{PB} + b_{TV}$ but remains a well-estimated quantity with a valid Wald test of its own nullity.
- Detecting the bias requires comparing $\hat{\beta}_D$ to m_{BR} , a comparison only the decomposition or an isolating phase contrast supplies.
- This is why a rejection-rate summary is the wrong instrument, and accounts for the flat rejection rates across the placebo-strength axis in the section 7.1 pilot.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Two consequences for how the framework should be evaluated follow. First, the quantity that carries the bias is the location of the estimand, $\text{plim } \hat{\beta}_D$, not its sampling variability: the lumped coefficient remains a well-estimated number with a valid Wald test of its own nullity, and that test rejects at the rate dictated by the displaced value. A significance test of the lumped coefficient therefore cannot, in principle, detect the bias, because the bias does not move the coefficient toward zero. Detecting it requires comparing $\hat{\beta}_D$ against m_{BR} , which is exactly the comparison the decomposition, or a phase contrast that isolates m_{BR} , makes available and the lumped analysis does not. Second, and as a corollary for the simulation evidence, a rejection-rate or power summary is the wrong instrument for this bias: the right target is the coefficient itself and its offset from m_{BR} . This is why the pilot of section 7.1 found the biomarker-by-treatment rejection rate

essentially flat across the placebo-strength axis even though that axis is precisely what drives the predicted bias; the prediction lives in the mean of $\hat{\beta}_{bm:D_{bc}}$, and a production run that means to test it must read the bias off the coefficient, not off the test.

6.2 A numerical illustration

[43]

Figure 1 illustrates the diagnostic contrast: two participants with identical open-label responses but opposite BR - PB profiles become distinguishable only in the blinded-discontinuation phase.

- Participant A ($m_{BR} = 4$, $m_{PB} = 6$) and Participant B ($m_{BR} = 8$, $m_{PB} = 1$) both show approximately ten points of open-label improvement, indistinguishable to a one-component analysis that estimates drug at 7.75 points.
- In the blinded-discontinuation phase A retains four points (PB persists at $\eta = 0.5$) while B collapses to 1.5 points (drug cleared, minimal PB): a seven-point diagnostic contrast.
- The three-component decomposition recovers $\hat{m}_{BR}^A = 4$, $\hat{m}_{BR}^B = 8$, $\hat{m}_{PB}^A = 6$, $\hat{m}_{PB}^B = 1$ – six clinically actionable quantities versus one biased average; predictive-biomarker validation requires these.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

[Figure 1 here: two-panel trajectory plot. Left panel shows open-label phase (weeks 1–8) with overlapping response trajectories for Participants A and B. Right panel shows blinded-discontinuation phase (weeks 9–12) with the diagnostic divergence: A retains four points while B collapses to 1.5. Component contributions are labelled in each panel.]

Two participants in a sixteen-week N-of-1 trial of prazosin for PTSD share the same baseline but differ in latent traits: Participant A has $m_{BR}^A = 4$, $m_{PB}^A = 6$; Participant B has $m_{BR}^B = 8$, $m_{PB}^B = 1$. Both reach approximately ten points of open-label improvement by week eight, making them indistinguishable to a one-component analysis that estimates a population drug effect of 7.75 points. The blinded-discontinuation phase supplies the identifying contrast: A retains four points of improvement (belief persists at $\eta = 0.5$) while B returns nearly to baseline (drug cleared, minimal belief). The three-component decomposition recovers all six component estimates by exploiting this phase differential; the one-component model cannot.

7 Simulation study

[44]

The simulation study follows ADEMP and quantifies bias and identifiability properties of three analysis strategies across a prazosin-PTSD-calibrated parameter grid.

- The Morris (2019) ADEMP framework structures the study: Aims, DGP, Estimands, Methods, Performance measures.
- The pre-registration is at `analysis/scripts/component-decomposition/00-ademp-pre-registration`.
- Three analysis strategies (one-component, phase-augmented, full decomposition) are compared across the parameter grid.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The simulation study reported and pre-registered in this section is structured under the³³ ADEMP framework, specifying Aims, Data-generating mechanisms, Estimands, Methods, and Performance measures. The pre-registered plan is at `analysis/scripts/component-decomposition/00-ademp-pre-registration.md`. A Monte Carlo simulation calibrated to the prazosin/PTSD reference parameters quantifies the bias and identifiability properties of the three analysis strategies described in section 5 across a parameter grid.

7.1 Pilot validation

[Details in supplement S1.]

[45]

A 100-replicate pilot confirmed the analysis pipeline is operational; two design-level findings inform the production specification.

- Convergence was 1.0 across six cells at $N = 35$; per-replicate output is checkpointed at `analysis/data/quick-sim/`.
- At $N = 35$, rank-deficiency dropped the D_{bc} :phase interaction term, requiring larger N and a richer design in the production programme.
- The predicted bias manifests in the mean of $\hat{\beta}_{bm:D_{bc}}$, not the Wald-test rejection rate; the production programme reports coefficient means against true values.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A 100-replicate pilot at $N = 35$ on a 6-cell grid confirmed the end-to-end pipeline. Two observations inform the production specification: at $N = 35$, rank-deficiency dropped the D_{bc} :phase interaction term,

removing the key contrast; and the predicted bias appears in the coefficient mean, not the rejection rate, so the production programme reports $\hat{\beta}_{bm:D_{bc}}$ against its implied true value rather than power alone. Full details and exploratory summaries are in supplement S1.

7.2 Production-simulation specification

[46]

The production programme follows the ADEMP framework with four deliverables; the realised Study A run in section 7.3 is the first of these.

- The deliverables are interaction-estimate bias, full-decomposition identifiability, type I error and power, and participant-level signal recovery.
- Sample sizes $N \in \{35, 70, 100, 150\}$ cross a *pb_strength* grid; type I cells target MCSE 0.003 at 5,000 replicates.
- Driver scripts, checkpointed outputs, and the ADEMP pre-registration reside under `analysis/scripts/component-decomposition/`.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The production programme follows the³³ ADEMP framework and targets four deliverables: the bias of $\hat{\beta}_{bm:D_{bc}}$ for the one-component, phase-augmented, and full-decomposition analyses across $N \in \{35, 70, 100, 150\}$ and a *pb_strength* grid; identifiability metrics for the full-decomposition fit (convergence, variance-inflation factors, and sensitivity of the *BR* estimate to the η (phase) multipliers); type I error and power for the interaction test under an η -mismatch, *TV*-magnitude, and dropout sensitivity grid, with null cells at $n_{\text{reps}} = 5,000$ for MCSE 0.003; and the participant-level interaction signal recovered as a function of (m_{BR}, m_{PB}, m_{TV}) , which the lumped-total regression does not preserve. The `tidyverse` implementation collection drives the grid; driver scripts, checkpointed cell-level summaries, per-replicate output, and the committed ADEMP pre-registration all reside under `analysis/scripts/component-decomposition/`. The realised Study A run reported next used 1,000 replicates per alternative cell.

7.3 Study A: orthogonal-biomarker arm (1,000 replicates per alternative cell)

[47]

Study A is the orthogonal-biomarker arm of the conditional: with $c_{bm,PB} = 0$, the one-component estimate is essentially

unbiased across the entire (m_{PB}, m_{TV}) grid, exactly as the algebraic identity requires.

- The true biomarker-by- BR effect is -2.25 ; mean absolute bias is 0.016 at $N = 35$, 0.016 at $N = 70$, 0.006 at $N = 100$, and 0.006 at $N = 150$, in every case within one Monte Carlo SE of zero.
- Bias shows no monotone trend across $m_{PB} \in \{0, 1, 3, 6, 10\}$ or $m_{TV} \in \{-1, 0, 1, 2\}$; it behaves as sampling noise around zero.
- This is the predicted consequence of the identity in section 6.1: the DGP couples the biomarker to BR only, so $\beta_{bm}^{PB} = \beta_{bm}^{TV} = 0$ and the lumped slope is unbiased by construction.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The production Study A run (1,000 replicates per alternative cell, 5,000 per null cell, 80 alternative cells crossing $m_{PB} \in \{0, 1, 3, 6, 10\}$, $m_{TV} \in \{-1, 0, 1, 2\}$, and $N \in \{35, 70, 100, 150\}$, plus four null cells) is the orthogonal-biomarker arm of the conditional evaluation. The data-generating process couples the baseline biomarker to the BR component only, through the on-drug biomarker- BR covariance channel, so the true biomarker-by-treatment effect is the biomarker-by- BR slope, $\beta_{bm:D_{bc}} = -c_{bm} \sigma_{BR} / \sigma_{bm} = -2.25$, and the biomarker is by construction uncorrelated with PB and TV . Under the one-component analysis the mean estimated $\hat{\beta}_{bm:D_{bc}}$ tracks this value across the whole grid: the mean absolute bias is 0.016 at $N = 35$, 0.016 at $N = 70$, 0.006 at $N = 100$, and 0.006 at $N = 150$, in every case within one Monte Carlo standard error (respectively about 0.028 , 0.019 , 0.009 , and 0.007) of zero, with no monotone trend across m_{PB} or m_{TV} . The predicted placebo-attribution bias of the biomarker-by-treatment estimate does not appear.

This is not a failure of the simulation; it is the behaviour the identity of section 6.1 requires. That identity writes the lumped biomarker slope as $\beta_{bm}^{BR} + w_{PB} \beta_{bm}^{PB} + w_{TV} \beta_{bm}^{TV}$. With the biomarker coupled to BR alone, $\beta_{bm}^{PB} = \beta_{bm}^{TV} = 0$, so the lumped slope equals the pharmacological slope no matter how large m_{PB} or m_{TV} grows. Study A varies the magnitudes of the unmodelled components, which the identity shows to be the innocuous axis; the axis that produces biomarker-validation bias is the biomarker's correlation with PB or TV , which this DGP holds at zero. The clean null is therefore a confirmation of the identity and a correction to the looser claim of section 1 that large PB or TV by itself biases biomarker validation. It does not; biomarker-component coupling does.

[48]

The phase-augmented analysis provides no bias reduction and incurs a substantial power penalty, so it is dominated

under this uncontaminated DGP.

- Mean power at $N = 35$ is 0.35 (phase-augmented) versus 0.59 (one-component); at $N = 70$, 0.62 versus 0.90; both reach approximately 1.0 by $N = 100$.
- Phase-augmented mean absolute bias is no smaller than one-component at any N (0.028 versus 0.016 at $N = 35$).
- Convergence is approximately 100% for both analyses at every N (minimum 0.998), so the rank-deficiency that degraded the $N = 35$ pilot does not recur here.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The comparison between analyses is equally clear. The phase-augmented analysis is not less biased than the one-component analysis at any sample size (its mean absolute bias is 0.028 versus 0.016 at $N = 35$ and is comparable elsewhere), and it is markedly less powerful for the biomarker-by-treatment test: mean power across the grid is 0.35 versus 0.59 at $N = 35$ and 0.62 versus 0.90 at $N = 70$, with both analyses reaching essentially full power by $N = 100$. Because the DGP contains no biomarker- PB contamination for the phase contrast to remove, the extra phase-by-treatment terms add estimation cost without buying any bias reduction, and the analysis is dominated. Convergence was approximately complete for both analyses at every sample size (minimum convergence rate 0.998), so the rank-deficiency that forced dropping the `Dbc:phase` term in the $N = 35$ pilot, which we had advanced as the explanation for the pilot's inverted power ranking, is not operative in the production run: the inverted ranking persists with the full formula fitted and converged. The pilot's power inversion was therefore a genuine feature of this DGP, not an artifact of rank-deficiency.

Calibration is adequate for both analyses. Nominal-95% confidence-interval coverage of $\beta_{bm:D_{bc}}$ ranges from 0.96 to 0.99 for the one-component analysis (slightly conservative) and from 0.94 to 0.98 for the phase-augmented analysis. Type I error at the four null cells is controlled, ranging from 0.031 to 0.056 for the one-component analysis and 0.038 to 0.049 for the phase-augmented analysis; the one-component estimate at $N = 100$ (0.056, MCSE 0.003) sits 1.9 standard errors above nominal but is within the simulation's Monte Carlo uncertainty.

[49]

Study A confirms the identity but cannot exhibit the biomarker-validation failure; a contaminated-biomarker pilot reported in the next subsection provides the decisive evidence.

- As pre-registered, Study A varies component magnitude with the biomarker fixed to BR , so it probes an axis the identity predicts to be null.
- The decisive axis is $c_{bm,PB}$, the biomarker- PB correlation; this is varied in the contaminated-biomarker pilot below.
- On current evidence the phase-augmented analysis is not recommended for the uncontaminated-biomarker regime; the pilot below evaluates it under contamination.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Two conclusions follow for the Study A programme. First, the result sharpens rather than supports the section 1 framing: the biomarker-validation failure is real as an identity, but it is driven by the biomarker's coupling to PB or TV , not by the magnitude of those components, and Study A as pre-registered varies the latter while holding the former at zero. The decisive axis is the biomarker- PB correlation $c_{bm,PB}$, which the contaminated-biomarker pilot in the following subsection varies directly. Second, on the present evidence the phase-augmented analysis carries a real power cost without a compensating bias reduction under an uncontaminated biomarker, so it should not be recommended as a default for this design; the pilot below evaluates whether contamination changes this verdict.

7.4 Study B pilot: coupling-arm simulation (100 replicates per cell)

[50]

A pilot varying the biomarker- PB correlation $c_{bm,PB} \in \{0, 0.15, 0.30, 0.45\}$ confirms that bias in the one-component estimator scales with coupling strength, not component magnitude, and that the phase-augmented analysis provides no protection.

- At $N = 150$, mean $\hat{\beta}_{bm:D_{bc}}$ shifts from -2.25 (true) to -1.74 at $c_{bm,PB} = 0.45$ (bias $+0.51$, 30 Monte Carlo SEs from zero), confirming the identity's prediction.
- The bias is driven by $c_{bm,PB} \times \sigma_{PB}$, not by the population-mean m_{PB} : cells with $m_{PB} = 0$ and $m_{PB} = 6$ show identical bias at every $c_{bm,PB}$ level, because the individual-level PB variance ($\sigma_{PB} = 2$) provides the contamination channel even when the mean is zero.
- The phase-augmented analysis shows the same bias as the one-component analysis at every contamination level, offering no protection against biomarker- PB coupling.

- This pilot uses 100 replicates per cell; a production run at 1,000 replicates per cell is required to confirm the pattern and assess type I error under contamination.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A supplementary pilot varying the biomarker- PB correlation $c_{bm,PB} \in \{0, 0.15, 0.30, 0.45\}$ at two PB magnitude levels ($m_{PB} \in \{0, 6\}$) and two sample sizes ($N \in \{70, 150\}$), using 100 replicates per cell, confirms the identity-based prediction and reveals an unexpected structural result.

At $N = 150$ the one-component bias scales monotonically with $c_{bm,PB}$: mean $\hat{\beta}_{bm:D_{bc}}$ moves from -2.31 ($c_{bm,PB} = 0$, bias -0.06) to -2.37 ($c_{bm,PB} = 0.15$, bias -0.12) to -2.08 ($c_{bm,PB} = 0.30$, bias $+0.17$) to -1.74 ($c_{bm,PB} = 0.45$, bias $+0.51$). The last value sits 30 Monte Carlo standard errors from the true -2.25 , establishing the effect unambiguously. The direction of bias is positive (toward zero): high- $c_{bm,PB}$ participants experience stronger PB , which in the open-label phase adds to the apparent drug effect, but this elevation is partially absorbed by the bm main effect, leaving the $bm : D_{bc}$ interaction coefficient attenuated relative to the pure pharmacological value.

The unexpected result is that m_{PB} does not modulate the bias: cells with $m_{PB} = 0$ and $m_{PB} = 6$ produce identical bias at every contamination level (differences ≤ 0.006 at $N = 150$, well within Monte Carlo error). The driver is $c_{bm,PB} \times \sigma_{PB}$, not $c_{bm,PB} \times m_{PB}$. Even when the population-mean PB amplitude is zero, individual participants have nonzero PB draws (standard deviation $\sigma_{PB} = 2$), and coupling the biomarker to that variance is sufficient to contaminate the estimator. A true no- PB control requires $\sigma_{PB} = 0$, not $m_{PB} = 0$. This refines the Study A design: varying m_{PB} at $c_{bm,PB} = 0$ correctly produced zero bias, but because the relevant null was $c_{bm,PB} = 0$, not $m_{PB} = 0$.

The phase-augmented analysis provides no protection: its bias at every $(c_{bm,PB}, m_{PB}, N)$ cell is indistinguishable from the one-component bias. Adding phase-by-treatment interactions to the model does not decompose the contamination because both analyses see the biomarker- PB covariance through the same open-label versus blinded-discontinuation contrast; the contrast changes the weight on PB (via expectancy) but does not isolate the PB component's contribution to the biomarker slope. Separating that contribution requires both a design that decouples drug exposure from belief and a decomposition that exploits it; Study C shows that the standard hybrid design cannot achieve this separation, whereas a balanced-placebo design can.

The $N = 70$ results are uninformative at 100 replicates (Monte Carlo

standard error ≈ 0.06 , comparable to or larger than the bias at moderate $c_{bm,PB}$). The production run will use 1,000 replicates per cell to establish the $N = 70$ pattern and 5,000 null replicates to assess type I error under contamination. Until the production run is complete, the quantitative bias estimates from this pilot should be treated as directionally informative but imprecise.

[51]

The contaminated-biomarker pilot reframes the paper’s practical guidance: the risk of biomarker-validation failure arises from biomarker- PB correlation in the study population, not from the presence of strong placebo response per se.

- A trial in a population where the biomarker is uncorrelated with placebo responsiveness can use the one-component estimator without bias, regardless of how large the placebo response is.
- A trial where the biomarker covaries with optimism, suggestibility, or treatment-seeking history is at material risk of contamination; bias of $+0.51$ (23% of the true effect) at $c_{bm,PB} = 0.45$ illustrates the practical scale.
- Neither the one-component nor the phase-augmented analysis removes this bias; full component decomposition is required.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The pilot results reframe the paper’s practical guidance. The risk of biomarker-validation failure in an aggregated N-of-1 study arises not from the presence of a strong placebo response in the population but from the biomarker’s correlation with individual-level placebo responsiveness. In a population where the baseline biomarker is uncorrelated with each participant’s placebo-belief response amplitude, the one-component estimator recovers the pharmacological interaction without bias, and the full decomposition machinery is not required. In a population where the biomarker covaries with psychological characteristics that predict placebo response – optimism, prior treatment experience, suggestibility, expectancy at enrolment – the one-component analysis conflates pharmacological and placebo-belief prediction, and the scale of that conflation grows with the coupling strength. At $c_{bm,PB} = 0.45$, a value that equals the pharmacological coupling used throughout this paper, the one-component bias is approximately 0.51 points on the interaction scale, representing 23% of the true pharmacological effect. This is a practically meaningful distortion for any precision-medicine stratification decision.

The production run of the contaminated-biomarker study, at 1,000 alternative-cell replicates and 5,000 null-cell replicates, is required to

confirm these estimates, establish the $N = 70$ pattern, and characterise type I error under contamination. Results are pending.

7.5 Study C: recovery under a belief-decoupling design

[52]

Study C shows the contamination is recoverable, but only under a design that decouples drug exposure from belief: a belief-covariate decomposition on a balanced-placebo design returns the unbiased pharmacological slope across the feasible contamination range, while the one-component estimator does not.

- In the MVN process the biomarker- BR coupling loads on the drug state D_{bc} and the biomarker- PB coupling scales with the expectancy η ; the decomposition $Sx \sim bm + t + D_{bc} + \eta + bm:D_{bc} + bm:\eta$ isolates β_{bm}^{BR} on $bm:D_{bc}$ only if D_{bc} and η vary independently.
- On the standard hybrid design they are collinear (analysis-stage correlation $\approx +0.6$), and the belief-covariate, blinded-stratum, and Gompertz-basis decompositions all fail there, no better than the one-component analysis.
- A balanced-placebo design adds a covert on-drug phase ($\eta = 0$) and an open-placebo phase ($\eta = 1$), decoupling the contrasts (≈ -0.66) and enabling recovery.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Study C asks whether the contamination demonstrated in the coupling-arm study can be removed by decomposition, and finds that the answer depends on the trial design. In the MVN process the biomarker- BR coupling acts only on drug and so loads on the drug-state contrast D_{bc} , whereas the biomarker- PB coupling scales with the expectancy η ; interacting the biomarker with both isolates β_{bm}^{BR} on $bm:D_{bc}$, but only where D_{bc} and η vary independently. On the standard hybrid design they do not (on-drug time and belief accumulate together, with an analysis-stage correlation of about $+0.6$), and we verified that the belief-covariate decomposition, a blinded-stratum drug contrast, and a Gompertz-basis component decomposition all fail to recover β_{bm}^{BR} there. The remedy is a design that breaks the collinearity: a balanced-placebo (open-hidden) design^{25,26} interposes a covert on-drug phase, in which the drug is given without the participant's knowledge ($\eta = 0$), and an open-placebo phase, in which the participant is off drug but believes treatment continues ($\eta = 1$), reducing the D_{bc} - η correlation to about -0.66 .

[53]

At 1,000 replicates per cell the decomposition is unbiased across the contamination range while the one-component estimator degrades monotonically; recovery is conditional on the design and bounded by positive-definiteness.

- At $N = 150$ the decomposition bias of $\hat{\beta}_{bm:D_{bc}}$ stays within one Monte Carlo SE of zero across $c_{bm,PB} \in \{0, 0.1, 0.2, 0.3\}$ ($-0.022, +0.008, +0.021, -0.026$; MCSE ≈ 0.02), while the one-component bias rises to $+0.16, +0.34, +0.48$.
- Interval coverage is near-nominal for the decomposition (0.94–0.96) and falls to 0.86 for the one-component analysis at the highest contamination; convergence is 100 percent.
- The one-component power advantage is confined to low contamination: it is $+0.19$ (power 0.85 versus 0.66) at $c_{bm,PB} = 0$, $N = 70$, erodes monotonically as contamination grows, and reverses by $c_{bm,PB} = 0.3$ (0.64 versus 0.67), because the biased one-component estimate is dragged toward the null.
- The grid is capped at $c_{bm,PB} = 0.30$ because beyond approximately 0.45 the implied covariance is non-positive-definite and the contaminated DGP is no longer faithfully representable.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The production run (1{,}000 replicates per cell, $N \in \{70, 150\}$, $c_{bm,PB} \in \{0, 0.1, 0.2, 0.3\}$) confirms the recovery. At $N = 150$ the decomposition tracks the pharmacological target -2.25 across the contamination range, with bias $-0.022, +0.008, +0.021$, and -0.026 , every value within one Monte Carlo standard error (about 0.02) of zero, while the one-component estimator develops a monotone bias of $+0.16, +0.34$, and $+0.48$ as $c_{bm,PB}$ grows. Interval coverage holds near nominal for the decomposition (0.94 to 0.96) and degrades to 0.86 for the one-component analysis. The power comparison (Table 2) is more nuanced than a uniform decomposition penalty: at zero contamination the one-component analysis is more powerful, by 0.19 at $N = 70$ (0.85 versus 0.66) and 0.07 at $N = 150$, reflecting the degrees of freedom the decomposition spends, but this advantage erodes monotonically as $c_{bm,PB}$ grows and reverses by $c_{bm,PB} = 0.3$ (0.64 versus 0.67 at $N = 70$). The reason is that contamination in this direction drags the one-component point estimate toward the null and so costs it power, whereas the decomposition holds the estimate at the pharmacological target; the one-component power advantage is accordingly confined to the low-contamination regime in which decomposition is unnecessary. Recovery is therefore conditional on the design, not on the analysis alone: the decomposition is necessary but a design that varies belief independently of drug exposure is the enabling condition. The practical recommendation is that a biomarker-validation trial in

1679 which the candidate biomarker may correlate with placebo responsive-
1680 ness should be fielded as a balanced-placebo or open-hidden design
1681 rather than the standard hybrid design. Per-replicate output is at
1682 `analysis/data/derived_data/component-decomposition/study-b-balanced/`.
1683 Recovery beyond the positive-definite range, participant-level compo-
1684 nent recovery, and replication under Architecture A coupling remain
1685 open.

Table 2: Rejection rate (power) for the biomarker-by-treatment interaction at $\alpha = 0.05$ under the balanced-placebo design, one-component versus belief-covariate decomposition, by contamination $c_{bm,PB}$ and sample size (1,000 replicates per cell; paired standard error of the gap ≤ 0.014). The one-component advantage at $c_{bm,PB} = 0$ erodes and reverses as contamination grows, because contamination drags the biased one-component estimate toward the null.

$c_{bm,PB}$	N	One-component	Decomposition	Gap
0.0	70	0.85	0.66	+0.19
0.1	70	0.78	0.67	+0.11
0.2	70	0.74	0.69	+0.05
0.3	70	0.64	0.67	-0.03
0.0	150	0.99	0.93	+0.07
0.1	150	0.98	0.93	+0.05
0.2	150	0.97	0.93	+0.04
0.3	150	0.93	0.93	-0.00

1686 8 Application to predictive-biomarker validation

1687 [54]

1688 **The precision-medicine question – whether a biomarker pre-**
1689 **dicts pharmacological response – is answered by the lumped**
1690 **total only when the biomarker is uncorrelated with PB and**
1691 **TV ; otherwise it requires decomposing the response.**

- 1692 • The question is whether the biomarker predicts the BR compo-
1693 nent, not the lumped total.
- 1694 • The PATH statement formalises this as the target estimand for
1695 predictive biomarker validation.
- 1696 • The BR-PB-TV decomposition is the mechanism that separates
1697 the pharmacological signal when biomarker-component coupling
1698 cannot be ruled out.

1699 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod*
1700 *tempor incididunt ut labore et dolore magna aliqua.*

The motivating clinical question for the trial designs in this project is whether a baseline biomarker predicts treatment response, and specifically whether the biomarker predicts the pharmacological component of treatment response. This is the precision-medicine question, formalised in the PATH statement⁸. Whether it depends on the BR-PB-TV decomposition is itself conditional: by the identity of section 6.1, the lumped biomarker slope equals the pharmacological slope when the biomarker is uncorrelated with PB and TV , and departs from it only by the biomarker's covariance with those components. The decomposition is therefore the mechanism the question requires precisely in the regime where that covariance is non-negligible.

[55]

A lumped-total biomarker regression conflates pharmacological and placebo-belief prediction, producing a stratification biomarker that mis-identifies responders.

- $\hat{\beta}_{bm}$ from the lumped regression contains both pharmacological and placebo-belief prediction in unknown proportions.
- If PB covaries with the biomarker distribution, flagged 'high responders' include high- PB responders who will not benefit pharmacologically.
- Prospective application of such a biomarker would systematically mis-stratify patients for prescription decisions.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A direct regression of total response on baseline biomarker

```
total_response = beta_0 + beta_bm * biomarker + error
```

estimates the biomarker's correlation with whatever the lumped total contains. If PB varies systematically across the biomarker distribution, because the biomarker is correlated with optimism, suggestibility, or treatment-seeking history, then $\hat{\beta}_{bm}$ contains a placebo-prediction component inseparable from the pharmacological-prediction component. The resulting biomarker, applied prospectively, would mis-stratify patients: those flagged as 'high responders' would include both true high- BR responders and high- PB responders, and the latter would not benefit from prescription decisions made on a pharmacological basis.

[56]

A valid pharmacological predictor requires a non-zero $\hat{\beta}_{bm}^{BR}$; significant $\hat{\beta}_{bm}^{PB}$ or $\hat{\beta}_{bm}^{TV}$ indicate a psychological or prognostic biomarker, not a pharmacological one.

- The sole criterion for pharmacological prediction is a meaningfully non-zero BR -component regression coefficient.
- A biomarker correlating only with PB selects placebo responders, not drug responders.
- A biomarker correlating only with TV is prognostic but not predictive in the pharmacological sense.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The right approach is to regress each component on the biomarker:

$$\begin{aligned} BR &= \beta_0^{BR} + \beta_{bm}^{BR} \cdot \text{biomarker} + \varepsilon^{BR}, \\ PB &= \beta_0^{PB} + \beta_{bm}^{PB} \cdot \text{biomarker} + \varepsilon^{PB}, \\ TV &= \beta_0^{TV} + \beta_{bm}^{TV} \cdot \text{biomarker} + \varepsilon^{TV}. \end{aligned}$$

A biomarker is a useful pharmacological predictor if and only if $\hat{\beta}_{bm}^{BR}$ is meaningfully non-zero. Significant $\hat{\beta}_{bm}^{PB}$ or $\hat{\beta}_{bm}^{TV}$ are diagnostic findings about the biomarker, indicating that it is partly psychological or correlates with natural-history characteristics, but they are not the validation of pharmacological prediction. A biomarker that correlates only with PB would be valuable for selecting placebo responders in run-in designs but worthless for predicting pharmacological efficacy; a biomarker that correlates only with TV would be a prognostic marker, not a predictive one.

[57]

When biomarker-component coupling cannot be ruled out, the decomposition is the mechanism by which the precision-medicine question is posed; when it can, the lumped analysis already answers it.

- Where the biomarker may correlate with PB or TV , separating BR from them is what prevents validation from conflating pharmacological and non-pharmacological prediction.
- The phase-augmented `bm:Dbc:phase` interaction is a small- N proxy, but Study A shows it buys nothing under an uncontaminated biomarker, so it is warranted only when contamination is plausible.
- The full decomposition recovers component-specific regressions explicitly when sample size permits.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The decomposition is therefore the mechanism by which the precision-

medicine question is posed whenever the candidate biomarker may correlate with PB or TV . Where the biomarker can be assumed orthogonal to the non-pharmacological components, the identity of section 6.1 shows the lumped analysis already targets the biomarker-by- BR estimand, and the simulation of section 7.3 confirms it recovers that estimand without bias. The phase-augmented linear-mixed analysis introduced in section 5 retains a useful subset of the decomposition machinery at small N (the `bm:Dbc:phase` interaction tests whether the biomarker-by-treatment effect attenuates under blinding, the practical proxy for biomarker-by- PB contamination), but Study A shows this proxy earns its degrees of freedom only when such contamination is present: under an uncontaminated biomarker it reduced no bias and lost power. The full decomposition recovers the component-specific regressions explicitly when N permits and when the scientific question warrants the cost.

9 Discussion

[58]

The three-component framework reframes treatment response as a structured sum of causally distinct contributions, and the present work maps when acting on that decomposition actually changes inference.

- The framework specifies which trial-design contrasts identify which components.
- Acting on the decomposition is valuable conditionally: for attribution estimands, and for biomarkers that may correlate with PB or TV , but not for a biomarker-by-treatment interaction with a biomarker orthogonal to them.
- The target estimand for biomarker validation is the biomarker-by- BR association; whether a lumped analysis already recovers it depends on the biomarker's coupling to PB and TV .

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The three-component decomposition is the formal expression of the claim that ‘treatment response’ in an aggregated N-of-1 trial is not a single quantity but a sum of three causally distinct contributions, each with its own pharmacological, psychological, or epidemiological interpretation. What the present work establishes, analytically in section 6.1 and empirically in section 7.3, is that the inferential value of acting on this decomposition is conditional rather than universal. The identity shows that the one-component estimator is displaced from its pharmacological target only by the components’ loading onto the contrast

used, and, for a biomarker slope, only by the biomarker's covariance with PB and TV . Where those terms are zero, the design-contrast one-component analysis is already unbiased, and the production simulation confirms it: with a biomarker coupled to BR alone, the one-component biomarker-by-treatment estimate is unbiased across the whole $PB-TV$ grid and the phase-augmented analysis is dominated. The framework's contribution is therefore best read as a map of when decomposition changes inference, not as a blanket prescription to decompose. It makes the trial-design conditions explicit, identifies the attribution and contaminated-biomarker estimands for which decomposition is needed, and clarifies the target estimand for biomarker validation (the biomarker-by- BR association, not the lumped total).

[59]

Four open questions bound the framework's current scope: the simulation results are pending, the η multipliers require sensitivity analysis, non-monotone natural history is not yet handled, and participant-level recovery inherits the identifiability caveats of individual treatment effects.

- Quantitative bias-versus-power trade-offs are conjectures until the section 7 simulation is run.
- The $\eta(\text{phase})$ expectancy-modulation parameters are modelling assumptions requiring explicit sensitivity analysis.
- Cyclical or biphasic TV trajectories exceed the Gompertz form and require flexible or covariate-augmented specifications.
- Participant-level component recovery rests on separating participant-by-component signal from participant-by-period and participant-by-phase variation, the identifiability caveat Senn raises for individual treatment effects.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Four open questions remain and warrant acknowledgement. First, the simulation study described in section 7 is forthcoming; the present manuscript reports the framework rather than its empirical evaluation, and the quantitative bias-versus-power trade-offs across analysis strategies are conjectures until that work is run. Second, the $\eta(\text{phase})$ multipliers used to parameterise the expectancy modulation of PB are themselves modelling assumptions, and a sensitivity analysis varying η across plausible values is mandatory before any clinical conclusion is drawn from the framework's outputs. Third, the framework as specified treats the three components as monotone Gompertz trajectories; cyclical or biphasic TV patterns require either flexible TV specifications or explicit covariates beyond the Gompertz form, and the boundary between 'flexible-Gompertz' and 'covariate-augmented'

specifications is not yet sharply drawn. Fourth, the participant-level component estimates the framework targets inherit the identifiability caveat that²⁹ and²⁸ raise for individual treatment effects more generally: within-participant component recovery rests on separating a true participant-by-component signal from participant-by-period and participant-by-phase variation, and this separation is credible only when the design supplies repeated phase contrasts and the cross-component covariance structure is not pathological. The framework should be read as estimating these quantities under those assumptions, not as dissolving the well-known difficulty of estimating individual effects from finite within-person data.

[60]

Two companion documents extend the manuscript: a pedagogical narrative and a Gompertz-family sensitivity evaluation.

- The pedagogy document at `docs/24-component-decomposition-pedagogy.md` provides worked examples for first-time readers.
- The Gompertz evaluation at `analysis/report/07-gompertz-evaluation/` quantifies trajectory-specification sensitivity.
- Together they address the intuitive and parametric dimensions of the framework that the present manuscript treats briefly.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The companion pedagogical document at `docs/24-component-decomposition-pedagogy.md` develops the framework with worked examples and an extended intuition narrative for readers approaching the material for the first time. The companion methodological evaluation of the Gompertz family at `analysis/report/07-gompertz-evaluation/` quantifies the sensitivity of the framework's inference to the specific parametric form chosen for the component trajectories.

[61]

The components exist in any longitudinal treatment response, but participant-level identifiability requires a within-subject belief-by-exposure contrast that not every placebo-controlled design supplies.

- The three-component structure is a property of the data-generating process of any repeated-measures treatment study; its identifiability is design-specific.
- Parallel-group placebo control identifies the population-mean PB but not a participant-level PB , and supplies no within-subject belief contrast; only designs with such a contrast (blinded or

open-hidden discontinuation, randomised withdrawal, crossover, hybrid N-of-1) identify the participant-level decomposition the biomarker question needs.

- The broad claim is therefore scoped to designs supplying a within-subject belief-by-exposure contrast, not to placebo-controlled longitudinal trials in general.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A final point concerns the scope of the decomposition, where it is essential to distinguish the existence of the components from their identifiability. The claim that an observed longitudinal treatment response is the sum of a pharmacological, a belief-driven, and a natural-history component is a statement about the data-generating process of any repeated-measures treatment study, and in that sense the decomposition is not specific to N-of-1 trials. Identifiability, by contrast, is design-specific, and the distinction between population-level and participant-level recovery is decisive for the biomarker application. A standard parallel-group placebo-controlled trial supplies a between-arm placebo contrast that identifies the population-mean belief response but not a participant-level PB , and it contains no within-subject contrast in which belief varies independently of pharmacology; the participant-level biomarker-by- BR slope that the precision-medicine question targets is therefore not identified from such a design, no matter how large it is. The designs that do supply the required within-subject belief contrast are the blinded or open-hidden discontinuation and randomised-withdrawal designs^{25,26}, the crossover with on/off periods, and the hybrid N-of-1 design developed here; the pharmacometric disease-progression-plus-placebo models from which the framework descends^{1,2,34} operate at the group level and recover population-mean components only. We therefore advocate the decomposition as a default lens for longitudinal designs that supply a within-subject belief-by-exposure contrast, with the aggregated N-of-1 case the most demanding instance because it must recover the components from a single participant's trajectory. We do not claim participant-level identifiability from parallel-group placebo control alone, and where only a between-arm placebo contrast is available the framework's claims are restricted to the population mean.

10 Conclusions

[62]

Treatment response in aggregated N-of-1 trials is the sum of three causally distinct components, and the bias from con-

flating them is governed by the estimand and the biomarker's coupling to the components, not by their raw magnitude.

- The three components are BR (pharmacological), PB (placebo-belief), and TV (natural history).
- Each has its own biological mechanism, identifiability conditions, and clinical interpretation.
- The lumped pre-post treatment effect is biased by the PB and TV contributions; the biomarker-by-treatment slope is biased only by the biomarker's covariance with PB and TV , as Study A confirms.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In conclusion, three points are to be emphasised. First, treatment response in aggregated N-of-1 clinical trials is the sum of three causally distinct components: pharmacological action (BR), placebo-belief response (PB), and time-variant natural history (TV). Each component has its own biological mechanism, its own identifiability conditions, and its own clinical interpretation. Whether conflating them biases a given analysis is, however, governed by the omitted-variable-bias identity of section 6.1, not by the raw magnitude of PB and TV : the lumped pre-post treatment effect is displaced by the PB and TV contributions, whereas the biomarker-by-treatment slope is displaced only by the biomarker's covariance with those components. Study A bears this out: varying PB and TV magnitude with a biomarker coupled to BR alone left the one-component estimate unbiased across the entire grid.

[63]

What matters is the design contrast, not the parametric model: the trial must supply the drug-on/off and belief contrasts, but adding parametric component terms is not free and is not always worth it.

- Unique component identification requires phase contrasts built into the trial design; the one-component analysis already exploits the drug-on/off contrast.
- Under an uncontaminated biomarker, the phase-augmented analysis adds no bias reduction and loses power, so it is not recommended as a default.
- The phase-augmented and full component-matched analyses are warranted for attribution estimands and for biomarkers that may correlate with PB or TV , where the extra structure does work.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Second, the operative requirement is the design contrast, not the parametric analysis model. The trial design must include the drug-on/off and belief contrasts that identify the components, and the one-component analysis already exploits the drug-on/off contrast through its D_{bc} term. Adding parametric component structure on top is not free: under an uncontaminated biomarker the phase-augmented analysis reduced no bias and lost substantial power, so it should not be adopted as a default. Its value, and that of the full component-matched fit, is confined to the estimands where the extra structure does work, namely the attribution of response to pharmacology versus belief and the analysis of biomarkers that may themselves correlate with PB or TV .

[64]

For predictive-biomarker validation the decomposition matters precisely when the candidate biomarker may correlate with PB or TV ; this is the case the decisive simulation still has to probe.

- A biomarker that predicts the lumped total is a pharmacological predictor only if it is uncorrelated with PB and TV ; otherwise the lumped slope is contaminated.
- Regulatory and prescribing decisions therefore require separating BR from PB and TV whenever biomarker-component coupling cannot be ruled out.
- Study A established the uncontaminated baseline; the contaminated-biomarker cell is the planned empirical test of the failure mode.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Third, for predictive-biomarker validation the decomposition earns its place precisely when the candidate biomarker may correlate with PB or TV . A biomarker that predicts the lumped total is a pharmacological predictor only when it is uncorrelated with the non-pharmacological components; where that cannot be assumed, the lumped slope mixes pharmacological and non-pharmacological prediction, and the regulatory and prescribing decisions that motivate the trial require the separation the decomposition supplies. The coupling-arm study demonstrates this failure directly, and Study C establishes that the bias is recoverable: a belief-covariate decomposition returns the unbiased pharmacological slope across the feasible contamination range, but only on a balanced-placebo design that decouples drug exposure from belief, not on the standard hybrid design. The decisive practical lesson is therefore as much about design as about analysis.

11 Conflict of interest

The authors declare no conflicts of interest.

12 Funding

This work received no specific funding.

13 Data availability statement

The data and code that support the findings of this study are openly available in the pmsimstats-ng project repository. Per-replicate Monte Carlo outputs and ADEMP-compliant cell-level summaries are versioned under `analysis/data/`; simulation drivers are at `analysis/scripts/`. Exact paths for this manuscript are listed in the Reproducibility section below.

14 Reproducibility

[65]

The manuscript is rendered inside a Docker container with locked package versions; the full software environment is pinned by Dockerfile and renv.lock.

- R 4.4.x in the zzcollab Docker container provides the execution environment.
- Simulation packages are `nlme`, `parallel`, `data.table`, `furrr`, and `purrr`.
- The manuscript build uses `rmarkdown` and `bookdown`.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

This manuscript was rendered with R 4.4.x inside the project's zzcollab Docker container (image hash recorded in `Dockerfile`; package set captured in `renv.lock`). Principal packages used by the simulation pipeline are `nlme` (continuous-time linear-mixed analysis with `corCAR1` residual correlation), `parallel` (8-core `mclapply` for parameter-grid sweeps), `data.table` (the original-extended simulation collection), `furrr` and `purrr` (the tidyverse simulation collection), and `rmarkdown` plus `bookdown` for the manuscript build.

[66]

Reproducibility is ensured by a master seed of 20260507 with per-replicate seeds derived by index addition.

- `set.seed(20260507)` is called at the top of each replicate loop.
- `parallel::mc.set.seed` propagates seeds inside `mclapply`.

- Per-replicate seeds are `master_seed + rep_idx`.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Each simulation driver sets `set.seed(20260507)` at the top of the replicate loop and uses `parallel::mc.set.seed` inside `mclapply`. Per-replicate seeds derive from the master seed by `+ rep_idx`.

[67]

All simulation outputs, driver scripts, pre-registration, and manuscript source files are versioned at documented paths within the repository.

- Per-replicate outputs and Morris summaries are at `analysis/data/quick-sim/`.
- Driver scripts and the ADEMP pre-registration are at `analysis/scripts/component-decomposition/`.
- Manuscript source, bibliography, and CSL file are at `analysis/report/06-component-decomposition/`.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Per-replicate simulation outputs are checkpointed at `analysis/data/quick-sim/06-decomp-replicates`.
Morris ADEMP cell-level summaries at the companion `06-decomp-summary.txt`.

Driver scripts are under `analysis/scripts/component-decomposition/`.

The pre-registered ADEMP plan is at `analysis/scripts/component-decomposition/00-ademp-pre-reg`.

The manuscript source, paper-specific bibliography (`references.bib`),
and SIM CSL file are at `analysis/report/06-component-decomposition/`.

15 References

1. Holford N. Clinical pharmacology = disease progression + drug action. *British Journal of Clinical Pharmacology*. 2015;79(1):18-27. doi:10.1111/bcp.12170
2. Gomeni R, Lavergne A, Merlo-Pich E. Modelling placebo response in depression trials using a longitudinal model with informative dropout. *European Journal of Pharmaceutical Sciences*. 2009;36(1):4-10. doi:10.1016/j.ejps.2008.10.025
3. Kaptchuk TJ, Kelley JM, Conboy LA, et al. Components of placebo effect: Randomised controlled trial in patients with irritable bowel syndrome. *BMJ*. 2008;336(7651):999-1003. doi:10.1136/bmj.39524.439618.25
4. Wager TD, Atlas LY, Lindquist MA, Roy M, Woo CW, Kross E. An fMRI-based neurologic signature of physical pain. *New England Journal of Medicine*. 2013;368(15):1388-1397. doi:10.1056/NEJMoa1204471

- 2091 5. Zunhammer M, Bingel U, Wager TD. Placebo effects on the neurologic pain signature: A meta-analysis of individual participant functional magnetic resonance imaging data. *JAMA Neurology*. 2018;75(11):1321-1330. doi:10.1001/jamaneurol.2018.2017
- 2092 6. Hall KT, Lembo AJ, Kirsch I, et al. Catechol-O-methyltransferase val158met polymorphism predicts placebo effect in irritable bowel syndrome. *PLoS ONE*. 2012;7(10):e48135. doi:10.1371/journal.pone.0048135
- 2093 7. Wang RS, Lembo AJ, Kaptchuk TJ, Hall KT, et al. Genomic effects associated with response to placebo treatment in a randomized trial of irritable bowel syndrome. *Frontiers in Pain Research*. 2022;2:775386. doi:10.3389/fpain.2021.775386
- 2094 8. Kent DM, Paulus JK, Klaveren D van, et al. The Predictive Approaches to Treatment effect Heterogeneity (PATH) Statement. *Annals of Internal Medicine*. 2020;172(1):35-45. doi:10.7326/M18-3667
- 2095 9. Raskind MA, Peterson K, Williams T, et al. A trial of prazosin for combat trauma PTSD with nightmares in active-duty soldiers returned from Iraq and Afghanistan. *American Journal of Psychiatry*. 2013;170(9):1003-1010. doi:10.1176/appi.ajp.2013.12081133
- 2096 10. Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: Statistical methods and practical considerations. *Frontiers in Digital Health*. 2020;2:13. doi:10.3389/fdgth.2020.00013
- 2097 11. Holford NHG. Holford NHG and sheiner LB 'understanding the dose-effect relationship - clinical application of pharmacokinetic-pharmacodynamic models', Clin Pharmacokin 6:429-453 (1981) - the backstory. *The AAPS Journal*. 2011;13(4):662-664. doi:10.1208/s12248-011-9306-5
- 2098 12. Tilburt JC, Emanuel EJ, Kaptchuk TJ, Curlin FA, Miller FG. Prescribing 'placebo treatments': Results of national survey of US internists and rheumatologists. *BMJ*. 2008;337:a1938. doi:10.1136/bmj.a1938
- 2099 13. Finniss DG, Kaptchuk TJ, Miller F, Benedetti F. Biological, clinical, and ethical advances of placebo effects. *The Lancet*. 2010;375(9715):686-695. doi:10.1016/S0140-6736(09)61706-2
- 2100 14. Benedetti F. Placebo effects: From the neurobiological paradigm to translational implications. *Neuron*. 2014;84(3):623-637.

- 2101 15. Frisaldi E, Shaibani A, Benedetti F. Understanding the mechanisms of placebo and nocebo effects. *Swiss Medical Weekly*. 2020;150:w20340. doi:10.4414/smw.2020.20340
- 2102 16. Carlino E, Frisaldi E, Benedetti F. Pain and the context. *Nature Reviews Rheumatology*. 2014;10(6):348-355. doi:10.1038/nrrheum.2014.17
- 2103 17. Carlino E, Piedimonte A, Benedetti F. Nature of the placebo and nocebo effect in relation to functional neurologic disorders. In: *Handbook of Clinical Neurology*. Vol 139.; 2016:597-606. doi:10.1016/B978-0-12-801772-2.00048-5
- 2104 18. Kirsch I. Response expectancy as a determinant of experience and behavior. *American Psychologist*. 1985;40(11):1189-1202. doi:10.1037/0003-066X.40.11.1189
- 2105 19. Kaptchuk TJ, Friedlander E, Kelley JM, et al. Placebos without deception: A randomized controlled trial in irritable bowel syndrome. *PLoS ONE*. 2010;5(12):e15591. doi:10.1371/journal.pone.0015591
- 2106 20. Lembo A, Kelley JM, Nee J, et al. Open-label placebo vs double-blind placebo for irritable bowel syndrome: A randomized clinical trial. *Pain*. 2021;162(9):2428-2435. doi:10.1097/j.pain.0000000000002234
- 2107 21. Nurko S, Saps M, Kossowsky J, others, Kelley JM. Effect of open-label placebo on children and adolescents with functional abdominal pain or irritable bowel syndrome. *JAMA Pediatrics*. 2022;176(4):349-356. doi:10.1001/jamapediatrics.2021.5750
- 2108 22. Bingel U, Wanigasekera V, Wiech K, et al. The effect of treatment expectation on drug efficacy: Imaging the analgesic benefit of the opioid remifentanyl. *Science Translational Medicine*. 2011;3(70):70ra14. doi:10.1126/scitranslmed.3001244
- 2109 23. Hróbjartsson A, Gøtzsche PC. Is the placebo powerless? An analysis of clinical trials comparing placebo with no treatment. *New England Journal of Medicine*. 2001;344(21):1594-1602. doi:10.1056/NEJM200105243442106
- 2110 24. Hróbjartsson A, Gøtzsche PC. Placebo interventions for all clinical conditions. *Cochrane Database of Systematic Reviews*. 2010;(1):CD003974. doi:10.1002/14651858.CD003974.pub3
- 2111 25. Colloca L, Lopiano L, Lanotte M, Benedetti F. Overt versus covert treatment for pain, anxiety, and Parkinson's disease. *The Lancet Neurology*. 2004;3(11):679-684.

- 2112 26. Rohsenow DJ, Marlatt GA. The balanced placebo design: Methodological considerations. *Addictive Behaviors*. 1981;6(2):107-122.
- 2113 27. International Council for Harmonisation of Technical Requirements for Pharmaceuticals for Human Use. ICH E9(R1) addendum on estimands and sensitivity analysis in clinical trials to the guideline on statistical principles for clinical trials. Published online 2019.
- 2114 28. Senn S. Controversies concerning randomization and additivity in clinical trials. *Statistics in Medicine*. 2004;23(24):3729-3753. doi:10.1002/sim.2074
- 2115 29. Senn S. Statistical pitfalls of personalized medicine. *Nature*. 2018;563(7733):619-621.
- 2116 30. Liao Z, Qian M, Kronish IM, Cheung YK. Analysis of N-of-1 trials using Bayesian distributed lag model with autocorrelated errors. *Statistics in Medicine*. 2023;42(13):2044-2060. doi:10.1002/sim.9676
- 2117 31. Daza EJ. Causal analysis of self-tracked time series data using a counterfactual framework for N-of-1 trials. *Methods of Information in Medicine*. 2018;57(S 01):e10-e21. doi:10.3414/ME16-02-0044
- 2118 32. Gärtner T, Schneider J, Arnrich B, Konigorski S. Comparison of Bayesian networks, G-estimation and linear models to estimate causal treatment effects in aggregated N-of-1 trials with carry-over effects. *BMC Medical Research Methodology*. 2023;23:191. doi:10.1186/s12874-023-02012-5
- 2119 33. Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Statistics in Medicine*. 2019;38(11):2074-2102.
- 2120 34. Chen C, Jönsson S, Yang S, Plan EL, Karlsson MO. Detecting placebo and drug effects on Parkinson's disease symptoms by longitudinal item-score models. *CPT: Pharmacometrics & Systems Pharmacology*. 2021;10(4):309-317. doi:10.1002/psp4.12601