

1 Three-component decomposition of treatment response in
2 aggregated N-of-1 trials: pharmacological, expectation-driven,
3 and natural-history components

4 pmsimstats team

5 August 15, 2026

6 **Contents**

7	Abstract	2
8	1 Introduction	3
9	2 Notation and the formal model	4
10	3 The three components	5
11	3.1 Biological response (BR)	6
12	3.2 Placebo-belief response (PB)	8
13	3.3 Time-variant natural-history component (TV)	10
14	4 Identifiability and trial-design contrasts	11
15	5 Analysis-side methodology	13
16	5.1 Four reasons not to match the DGP	13
17	5.2 A small extension that usually pays off: phase indicators	15
18	5.3 When to fit the full decomposition	16
19	6 Worked example: information loss without decomposition	17
20	6.1 What the one-component analysis estimates	17
21	6.2 A numerical illustration	19
22	7 Simulation study	19
23	7.1 Pilot validation	20
24	7.2 Production-simulation specification	20
25	7.3 Study A: orthogonal-biomarker arm (1,000 replicates per alternative cell) . . .	20
26	7.4 Study B pilot: coupling-arm simulation (100 replicates per cell)	22
27	7.5 Study C: recovery under a belief-decoupling design	23
28	8 Application to predictive-biomarker validation	25
29	9 Discussion	26

30	10 Conclusions	28
31	11 Conflict of interest	29
32	12 Funding	30
33	13 Data availability statement	30
34	14 Reproducibility	30
35	15 References	30

36 Abstract

37 **Background.** Treatment response in aggregated N-of-1 clinical trials reflects three causally
38 distinct mechanisms: pharmacological action of the active compound (BR), the patient's
39 belief that the treatment will work (PB), and natural-history drift in the underlying condition
40 (TV). Standard analysis models combine these into a single 'treatment effect' plus residual
41 error. Whether this one-component approach biases inference, and whether a component
42 decomposition repays its cost in degrees of freedom and power, is not unconditional: it depends
43 jointly on the estimand, on the trial design, and on whether a candidate biomarker correlates
44 with the non-pharmacological components. This paper makes that dependence precise and
45 resolves it into a design-by-correlation-structure rule.

46 **Methods.** We formalize a three-component decomposition of the within-participant out-
47 come trajectory, each component a modified Gompertz function with participant-specific ran-
48 dom effects, and characterize the trial-design contrasts (open-label, blinded discontinuation,
49 crossover, balanced-placebo) that supply identifiability. We derive the omitted-variable-bias
50 identity for the one-component estimator, which expresses both the one-component treatment
51 effect and the one-component biomarker slope as sums of component-specific terms. We then
52 evaluate analysis strategies in three Monte Carlo studies calibrated to the prazosin-PTSD
53 reference parameters: Study A couples the biomarker to BR only ($c_{bm,PB} = 0$) and varies the
54 unmodeled-component magnitudes (1,000 replicates per alternative cell, 5,000 per null, across
55 $m_{PB} \in \{0, 1, 3, 6, 10\}$, $m_{TV} \in \{-1, 0, 1, 2\}$, $N \in \{35, 70, 100, 150\}$); Study B (100-replicate pi-
56 lot) couples the biomarker to PB; and Study C (1,000 replicates) contrasts the standard hybrid
57 design with a belief-decoupling balanced-placebo design across $c_{bm,PB} \in \{0, 0.1, 0.2, 0.3\}$.

58 **Results.** The identity shows that the one-component biomarker-treatment slope is displaced
59 from the pharmacological slope by a term proportional to the biomarker's covariance with
60 PB and TV, not by the magnitude of those components; three regimes follow. (i) *Biomarker*
61 *orthogonal to PB and TV, hybrid design:* the one-component estimator is essentially unbi-
62 ased across the whole grid (mean absolute bias 0.016 at $N = 35$, 0.006 at $N = 150$, within
63 one Monte Carlo standard error (MCSE) of zero) and strictly more powerful than the phase-
64 augmented decomposition (mean power 0.59 versus 0.35 at $N = 35$; 0.90 versus 0.62 at
65 $N = 70$); decomposition removes no bias and only sacrifices power, so it is dominated. (ii)
66 *Biomarker coupled to PB, standard hybrid design:* the one-component slope is biased in propor-

tion to $c_{bm,PB}$ (one-component bias rising to +0.48, coverage falling to 0.86 at $N = 150$), but because the drug-state and belief contrasts are collinear under this design (analysis-stage correlation $\approx +0.6$), no decomposition recovers the pharmacological slope either. (iii) *Biomarker coupled to PB, balanced-placebo design*: adding covert-drug and open-placebo phases decouples the contrasts (≈ -0.66), and the belief-covariate decomposition then returns the unbiased pharmacological slope across the feasible contamination range (bias within one MCSE of zero, coverage 0.94–0.96) where the one-component estimator does not, at the cost of some power.

Conclusions. The benefit of decomposition is jointly determined by the biomarker’s correlation structure and the trial design, not by either alone. When the candidate biomarker is orthogonal to PB and TV, a single-indicator analysis that exploits the design contrast is unbiased and most powerful, and a parametric component model is counterproductive. When the biomarker may correlate with PB or TV, the one-component slope is biased and decomposition is necessary, but the bias is recoverable only under a design that varies belief independently of drug exposure (balanced-placebo or open-hidden); on the standard hybrid design no analysis recovers it, so a biomarker-validation trial at risk of placebo-belief contamination must be fielded with the decoupling design. Component decomposition is likewise warranted whenever the estimand is the attribution of response to pharmacology versus belief versus natural history. A pragmatic specification for the cases that warrant it is $Sx \sim \mathbf{bm} + \mathbf{t} + \mathbf{D}bc + \mathbf{phase} + \mathbf{D}bc:\mathbf{phase} + \mathbf{bm}:\mathbf{D}bc + \mathbf{bm}:\mathbf{D}bc:\mathbf{phase}$ with random intercept and AR(1) residual correlation.

1 Introduction

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Decomposing treatment response into pharmacological, expectancy, and natural-history contributions is not a new idea. Pharmacometrics has long modeled disease progression and drug effect as separable longitudinal processes^{1,2}; Kaptchuk and colleagues decomposed the placebo response itself into observation, ritual, and practitioner-relationship components in a three-arm IBS trial³; neuroimaging has directly dissociated pharmacological from expectancy-driven analgesia using objective neural signatures^{4,5}. The partition has intuitive clinical appeal: a patient whose symptom improvement is pharmacological will benefit from long-term treatment even after novelty fades, whereas one whose improvement is primarily belief-driven may not sustain it as expectations normalize. The question this paper addresses is not whether the partition is biologically coherent – it is – but when it is necessary for valid inference.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The decision criterion follows from an algebraic identity established in section 6: the one-component biomarker-treatment slope decomposes as $\beta_{bm}^{\text{lumped}} = \beta_{bm}^{BR} + w_{PB}\beta_{bm}^{PB} + w_{TV}\beta_{bm}^{TV}$, where the loadings w_{PB} and w_{TV} are determined by the trial

design and the β_{bm}^C terms are component-specific slopes. Bias is zero if and only if $\beta_{bm}^{PB} = \beta_{bm}^{TV} = 0$, that is, if and only if the biomarker is uncorrelated with the non-pharmacological components. Whether that condition holds is determined by biology and domain knowledge, not by the magnitude of PB or TV in the study population. A population with a large placebo response but a biomarker orthogonal to it yields an unbiased one-component estimate; a population with a small placebo response but a biomarker that correlates with it yields a biased one.

COMT genotype exemplifies the first category. COMT modulates catechol-O-methyltransferase activity in prefrontal cortex, regulates dopaminergic signaling in expectancy circuits, and predicts placebo-response magnitude in IBS and chronic pain populations^{6,7}. A trial using COMT as a predictive biomarker for prazosin response would face a one-component slope that conflates pharmacological prediction with expectancy moderation. CYP2D6 metabolizer phenotype exemplifies the second: the genotype determines hepatic drug metabolism and thereby plasma drug concentration, with no plausible pathway to expectancy or natural-history effects. ADRB2 Arg16Gly in asthma bronchodilator trials provides a further example: spirometric outcomes are objective and receptor binding affinity does not modulate belief. For both, the one-component pharmacological slope is already the target estimand.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The paper proceeds as follows. Section 2 fixes compact notation and states the formal model. Section 3 develops each component with biological motivation, parametric form, and identifiability conditions. Section 4 maps the hybrid open-label-plus-blinded-discontinuation design’s phase structure onto the decomposition’s identifiability requirements. Section 5 develops the analysis-side methodology, including the phase-augmented linear mixed-effects model (LME) approximation. Section 6 works through the numerical illustration and derives the algebraic identity that anchors the conditional rule. Section 7 presents Study A (orthogonal arm) and the Study B pilot (coupling arm). Section 8 translates the conditional rule into a practical biomarker-validation decision criterion under the PATH framework⁸. No published N-of-1 trial has applied this evaluation framework; the paper’s primary contribution is determining when decomposition changes the inferential answer rather than advocating for it as a universal default.

2 Notation and the formal model

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In this section we shall fix notation and state the formal model. The three components are defined precisely enough to make the identifiability and analysis-side arguments in later sections rigorous, but we have endeavored to

keep the presentation compact. A fuller treatment, with worked numerical examples at each step, is available in the companion pedagogical document at docs/24-component-decomposition-pedagogy.md.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

For participant i at trial timepoint t :

$$Y_{it} = \text{BL}_i - [BR_{it} + PB_{it} + TV_{it}] + \varepsilon_{it},$$

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Each component is a participant-specific function of its own time-on-state variable. BR_{it} depends on time-on-drug, the cumulative exposure to active medication since drug initiation. PB_{it} depends on time-on-belief, modulated by a phase-specific expectancy factor $\eta(\text{phase})$ that captures the patient's confidence that they are receiving active treatment. TV_{it} depends on time-since-trial-entry, calendar time elapsed regardless of treatment status.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Each component follows a modified Gompertz curve:

$$C(t) = m_C \exp(-d_C \exp(-r_C t)) - m_C \exp(-d_C),$$

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

ε_{it} is a within-participant residual term with AR(1) serial-correlation structure $\text{Cov}(\varepsilon_{it}, \varepsilon_{i,t+h}) = \sigma_\varepsilon^2 \rho^{|h|}$. The three components are permitted to be jointly correlated through within-time and across-time cross-component covariances c_{cf1t} and c_{cfct} respectively, so that participant-level heterogeneity in any one component can co-vary with heterogeneity in the others.

The principal psychometric and pharmacological terms used below are summarized in Table 1.

3 The three components

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In this section we shall present the three components in turn, each with its biological motivation, the modified Gompertz parameterization, and the design contrasts that identify it. The treatment is organized so that the reader can see, by the end

Table 1: Principal psychometric and pharmacological terms used in the three-component decomposition.

Term	Meaning
BR	Biological response: pharmacological reduction in symptom score attributable to active drug.
PB	Placebo-belief response: reduction attributable to the patient’s belief that they are receiving treatment.
TV	Time-variant component: reduction (or increase) attributable to the natural-history trajectory of the symptom.
$\eta(\text{phase})$	Expectancy multiplier for <i>PB</i> : 1 in open-label on-drug, ≈ 0.5 in blinded discontinuation, between 0 and 1 in the interim.
BL_i	Participant <i>i</i> ’s baseline (pre-treatment) symptom score.
D_{it}	Continuous-decay drug indicator on the analysis side: 1 on-drug, $\exp(-\lambda t_{sd,it})$ off-drug.
$t_{sd,it}$	Time since drug discontinuation (off-drug only).
λ	Carryover decay rate, with half-life $t_{1/2} = \ln 2/\lambda$.

of this section, that the formal model is a sum of three causally interpretable trajectories rather than a mathematical convenience.

3.1 Biological response (BR)

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

BR is the physiological reduction in symptom score attributable to the active medication’s chemical action on the body. It is independent of belief, expectancy, and natural history: a participant who somehow received active drug without knowing it would still develop a measurable *BR*, and a participant who took an inert sugar pill while believing it was active drug would have $BR = 0$ no matter how strong their expectation. Methodologically, *BR* is the quantity that drug regulators and prescribers want: approval decisions, dose selection, and biomarker-guided patient stratification are all framed around *BR*.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The qualitative shape of the *BR* trajectory follows from pharmacokinetic and pharmacodynamic considerations. Drug concentration in plasma reaches steady state within a small multiple of the elimination half-life; brain concentration follows plasma with a short delay; the downstream pharmacodynamic response (reduction in symptom-relevant neural activity, in the case of prazosin’s effect on noradrenergic tone during REM sleep) typically requires several days to weeks of sustained exposure to integrate fully⁹. The result is a saturating curve: slow rise in the first week, steepest growth in the second to third week, plateau as receptor adaptation completes, and a decay back toward zero on discontinuation governed by the carryover half-life $t_{1/2} = \ln 2/\lambda$ rather than an instantaneous drop. The rate of this off-drug decay is the object of the carryover machinery used throughout this project; only in the limit of a short half-life relative to the measurement spacing does it approximate the sharp step assumed in the worked example of section 6.

209 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
210 *ut labore et dolore magna aliqua.*

211 The modified Gompertz form $BR(t) = m_{BR} \exp(-d_{BR} \exp(-r_{BR}t)) -$
212 $m_{BR} \exp(-d_{BR})$ captures this profile compactly. The asymptote $m_{BR}(1 - e^{-d_{BR}})$
213 is the biological ceiling, the largest reduction the drug can produce at infinite
214 exposure; the scale parameter m_{BR} equals that ceiling up to the factor $(1 - e^{-d_{BR}})$,
215 which for the calibrated $d_{BR} \approx 5$ exceeds 0.99. The onset rate r_{BR} controls
216 how rapidly the participant approaches the ceiling: high r_{BR} means a fast onset
217 (effect plateaus in one to two weeks), low r_{BR} a slow climb (six to eight weeks).
218 The displacement d_{BR} shapes the climb: high d_{BR} gives an early-onset curve
219 that rises sharply at first, low d_{BR} a more gradual climb. For PTSD prazosin,
220 calibrated values from the¹⁰ reference dataset are $m_{BR} \approx 11$ (eleven points of
221 nightmare-score reduction at saturation), $r_{BR} \approx 0.42$ per week (half-maximum at
222 roughly four to five weeks of sustained exposure), and $d_{BR} \approx 5$ (moderately steep
223 early climb).

224 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
225 *ut labore et dolore magna aliqua.*

226 Two alternative families that do not capture the asymmetric saturation profile
227 are worth noting for contrast. A linear function $BR = \beta t$ climbs without bound
228 and is biologically implausible past a few weeks. An exponential decay $BR =$
229 $m_{BR}(1 - \exp(-rt))$ is acceptable for fast-onset drugs but lacks the slow-start
230 phase that characterizes most chronic-condition responses, where receptor adapta-
231 tion, feedback regulation, and gene-expression changes produce a measurable lag.
232 The Gompertz interpolates between exponential decay (high d limit) and linear-in-
233 t growth (low d limit), and its three-parameter form is flexible enough to fit a wide
234 range of pharmacodynamic profiles without overfitting. The Gompertz is a deliber-
235 ate choice over the two parametric forms most familiar from pharmacokinetic-
236 pharmacodynamic modeling. The sigmoid $E_{\max}$ (Hill) model¹¹ describes the de-
237 pendence of effect on concentration (a dose-response curve), whereas $BR(t)$ here
238 is a time-course of accumulating effect at fixed dosing; and the indirect-response
239 (turnover) models that do describe effect time-courses are specified through differ-
240 ential equations on a hidden response variable rather than in closed form. The
241 modified Gompertz retains the slow-start, saturating shape those models produce
242 while remaining a closed-form, three-parameter mean function that is tractable
243 inside a mixed-effects likelihood. Doc 25 of the project documentation provides
244 the full historical and biological context for the family choice and quantifies the
245 sensitivity of pmsimstats inference to alternative families; that work is reported
246 separately as 07-gompertz-evaluation. The headline result of that evaluation
247 is reassuring for the present framework: in a 1{,}000-replicate-per-cell simulation
248 that generated data under the logistic, hyperbolic-tangent, and piecewise-linear
249 breakpoint families and analyzed them under the standard linear-mixed model,
250 the trajectory family was essentially irrelevant to the power of the biomarker-

251 treatment interaction test (the family-to-family spread fell within two MCSEs of
252 sampling noise) and to Type I error (0.025 to 0.036 across null cells, with no
253 family-specific inflation). The conclusions drawn here are therefore not artifacts
254 of the Gompertz form, though the regimes the family cannot represent (cyclical
255 or biphasic trajectories) remain genuine exclusions.

256 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
257 *ut labore et dolore magna aliqua.*

258 The BR component is identified by within-participant on-drug versus off-drug
259 contrasts, formalized in section 4 below. The most direct contrast is the blinded-
260 discontinuation phase: when a participant is silently switched to placebo while
261 still believing they are on active drug, BR decays toward zero as the drug clears
262 while PB remains roughly intact, and the difference between on-drug and off-drug
263 response in that window is a relatively clean estimate of BR that does not require
264 the analyst to model PB or TV explicitly. The decay is not instantaneous: with
265 the carryover half-life $t_{1/2} = \ln 2 / \lambda$ of the analysis-side decay term, residual BR
266 persists into the early off-drug observations, so the contrast is clean only once the
267 discontinuation window is long relative to $t_{1/2}$, and the drug state is best carried by
268 the continuous-decay indicator D_{it} of Table 1 rather than a hard on/off indicator.

269 3.2 Placebo-belief response (PB)

270 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
271 *ut labore et dolore magna aliqua.*

272 PB is the reduction in symptom score that arises from the patient’s belief that they
273 are receiving active treatment, independent of what they are actually receiving.
274 That is, it is the placebo response in its strict sense: the part of the improvement
275 that would persist if the active drug were silently replaced with an inert pill,
276 provided the patient continued to believe they were on the active drug.

277 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
278 *ut labore et dolore magna aliqua.*

279 It is worth being precise about what PB is and is not. PB is not ‘fake’ im-
280 provement; it is a phenomenon clinicians knowingly exploit, a national survey
281 of US internists and rheumatologists finding roughly half report regular use of
282 placebo treatments¹². As the introduction sets out, the placebo response is medi-
283 ated by well-characterized neurobiological pathways^{13–17}: belief-driven release of
284 endogenous opioids and dopamine, anticipatory hypothalamic-pituitary-adrenal
285 regulation, and top-down attentional shifts, dissociated from pharmacology by
286 the⁴ neurologic pain signature⁵ and shown to have genetic correlates^{6,7}. What
287 makes these mechanisms ‘placebo’ rather than ‘drug’ is that they are triggered by
288 belief rather than by pharmacology.

289 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
290 *ut labore et dolore magna aliqua.*

For trial design, the relevant point is that PB is separately identifiable from BR if and only if the trial includes a phase in which the patient's belief about treatment differs from the actual treatment they are receiving. The blinded-discontinuation phase is the canonical such phase: the patient is told that, at some unannounced point during the window, they may be silently switched to placebo, and they cannot reliably tell whether they are still receiving active drug. Belief in treatment correspondingly drops to a partial level (multiplier roughly 0.5) because the patient's expectation is hedged.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

PB also follows a Gompertz curve, scaled by the expectancy factor:

$$PB(t \mid \text{phase}) = \eta(\text{phase}) \cdot m_{PB} \exp(-d_{PB} \exp(-r_{PB}t)) + (\text{offset to start at zero}).$$

Typical expectancy values are $\eta = 1.0$ in open-label on-drug, $\eta \approx 0.5$ in blinded discontinuation, and η between 0 and 0.5 in open-label placebo (the residual expectation a patient retains when told the previous active phase was beneficial). The quantity η operationalises response expectancy in the sense of¹⁸, the patient's anticipation of benefit that the expectancy literature treats as a determinant of experienced symptom change. The choice of $\eta = 0.5$ for the blinded phase is a modeling assumption rather than a calibrated quantity: to our knowledge the residual expectancy a patient carries when uncertain about allocation has not been measured directly as a fraction of the open-label level, and we adopt 0.5 as a deliberately central placeholder. The available indirect evidence is nonetheless that the residual response is substantial. Open-label placebo trials^{19–21} demonstrate that a clinically meaningful response persists even when patients are explicitly told they are receiving an inert intervention, which bounds the residual belief-driven response away from zero. Because the value of η is not empirically pinned, sensitivity analyses varying it across its plausible range are mandatory before any conclusion is allowed to depend on this single number.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The decomposition framework as proposed here has a structural analogue in the placebo-research literature.³ decomposed the placebo response itself into observation, ritual, and patient-practitioner relationship components in a randomized trial of acupuncture-style placebo for irritable bowel syndrome, demonstrating empirically that the placebo response is not a single 'expectancy' lump but a sum of identifiable parts. The present BR-PB-TV framework is analogous in spirit but operates at the level of the full treatment response (drug, placebo, natural history) rather than within the placebo response alone.

329 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
330 *ut labore et dolore magna aliqua.*

331 A subtle property of PB is that its variance also scales with η , not just its mean.
332 When η is high, there is substantial variability across patients in how strongly
333 the placebo response is expressed; high-expectancy patients show large PB , low-
334 expectancy patients show essentially none. When η is low (blinded), the placebo
335 response is uniformly attenuated and the patient-to-patient variance shrinks ac-
336 cordingly. The model therefore parameterizes $SD(PB \mid \text{phase}) = \eta(\text{phase}) \cdot \sigma_{PB}^{\text{baseline}}$.
337 Ignoring this heteroscedastic structure deflates standard errors in the open-label
338 phase and inflates them in the blinded phase, with downstream consequences for
339 confidence-interval coverage and for the calibration of any biomarker test built on
340 top of the analysis.

341 3.2.1 Additivity, subadditivity, and the BR-PB interaction

342 *[Extended treatment in supplement S2.]*

343 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
344 *ut labore et dolore magna aliqua.*

345 The additive decomposition of section 2 treats the three components as indepen-
346 dent contributors. Subadditivity – where the combined $BR + PB$ effect falls below
347 the sum of the components measured separately – is documented in some clinical
348 contexts^{3,14,22} and can arise from shared neural substrate, ceiling effects, or belief-
349 dynamics adaptation. A generalized model with $\gamma \cdot BR_{it} \cdot PB_{it}$ is identifiable
350 from the same phase contrasts. Apparent subadditivity can also arise from latent-
351 class structure with class-correlated components rather than a direct BR - PB in-
352 teraction; disentanglement requires a class-aware analysis or the phase contrasts
353 provided here. The phase-augmented analysis of section 5 absorbs subadditivity
354 partly as a side benefit.

355 3.3 Time-variant natural-history component (TV)

356 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
357 *ut labore et dolore magna aliqua.*

358 TV is the reduction in symptom score that would have happened even if the
359 patient had received no treatment at all. It is the participant’s natural-history
360 trajectory, integrated over the duration of the trial. For some participants TV is
361 positive (symptoms naturally improve over the trial because of life events, therapy,
362 time-in-condition, or regression to the mean); for others TV is negative (symptoms
363 naturally worsen because of trauma anniversaries, accumulating stress, or progres-
364 sive underlying disease).

365 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
366 *ut labore et dolore magna aliqua.*

Like PB , TV is not noise. It is a structured participant-specific signal that the model can in principle estimate from data, given sufficient repeated measurements and a trial design that includes a long enough off-drug baseline or follow-up window. What distinguishes TV from PB is its origin: TV is independent of treatment status and of belief about treatment, whereas PB is modulated by both.

3.3.1 Natural history: existence, testability, and functional form

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The shape of TV depends on disease class, trial timescale, and enrollment selection. Acute conditions trend TV -positive (intrinsic recovery); chronic stable conditions trend near $TV = 0$ on the population mean but with substantial participant-level variance; progressive conditions trend TV -negative. Selection-driven enrollment adds a positive TV component regardless of class, because patients enrol at moments of peak severity and regression to the mean then contributes apparent improvement in early trial weeks^{23,24}. The zero- TV assumption is testable: a model including TV returns $\hat{m}_{TV}$ with a variance estimate; folding TV into ε is worse, because the participant-level natural-history trajectory is structured and correlated with the on/off-drug schedule, producing autocorrelated within-participant residuals and heteroscedastic between-participant residuals that bias the standard errors on BR and PB . A simple week covariate is insufficient: it enforces a linear, population-mean trajectory and, without the design contrasts of section 4, absorbs both BR and TV into a single coefficient that identifies neither.

4 Identifiability and trial-design contrasts

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The mathematical decomposition above is meaningless unless the data contain enough information to pin down each component separately. Trial design is what supplies that information. In this section we shall characterize the hybrid open-label-plus-blinded-discontinuation design, used as the worked example throughout this paper, and show how each of its three phases contributes a different combination of components to the observed trajectory.

Phase	Drug?	Belief?	Components present
Open-label on drug	Yes	Full	$BR + PB(\eta = 1) + TV$
Blinded discontinuation	Hidden	Hedged	$BR \cdot \mathbf{1}_{\text{on drug}} + PB(\eta = 0.5) + TV$
Open-label crossover	Either	Knows	$BR \cdot \mathbf{1}_{\text{on drug}} + PB(\eta = 1) + TV$

398 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
399 *ut labore et dolore magna aliqua.*

400 Three phase-by-allocation contrasts then identify the three components. The open-
401 label versus blinded contrast (within drug) yields $PB(\eta = 1) - PB(\eta = 0.5) \approx$
402 $0.5 \cdot PB$, which isolates the belief component, but only conditional on a known
403 blinded-phase expectancy η : from a two-phase design just the product $\eta \cdot m_{PB}$ is
404 identified, and separating η from m_{PB} requires the open-label placebo phase. The
405 blinded on-drug versus blinded placebo contrast yields BR once drug carryover
406 has elapsed, since PB and TV match across allocation arms within the blinded
407 phase. All phases off drug yield TV plus any residual off-drug PB , with no active
408 BR contribution beyond carryover and PB at the off-drug expectation level. The
409 $\mathbf{1}_{\text{on drug}}$ indicator in the phase table above is shorthand for the continuous-decay
410 drug state D_{it} defined in Table 1.

411 This conditioning on η is the framework's principal identification limitation and
412 deserves emphasis rather than a parenthetical. With a two-phase (open-label and
413 blinded) design the data identify only the product $\eta \cdot m_{PB}$, so PB is recovered up
414 to the assumed blinded-phase expectancy; the working value $\eta \approx 0.5$ is a modeling
415 assumption, not an estimate, and absent a belief or expectancy measurement it is
416 not falsifiable from the outcome data alone. The identified set for m_{PB} is accord-
417 ingly the interval $\{\eta \widehat{m_{PB}} / \eta : \eta \in (0, 1]\}$, which collapses to a point only when η
418 is pinned down. Two design features close the gap: an open-label placebo phase,
419 which supplies a second equation separating η from m_{PB} , and direct measurement
420 of belief, as in the open-hidden and balanced-placebo designs^{25,26} that manipulate
421 or record the patient's treatment expectation. We therefore recommend that any
422 trial relying on the decomposition for a PB -versus- BR attribution either include
423 such a phase or collect an expectancy measure, and that analyses report sensitivity
424 of the attribution to η across its plausible range.

425 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
426 *ut labore et dolore magna aliqua.*

427 A trial that lacks any of these phases loses the corresponding identifiability. A
428 pure open-label design (no blinded discontinuation) cannot separate BR from
429 PB , because both are present at full strength throughout. A pure blinded design
430 (no open-label phase) cannot estimate $\eta = 1 \cdot PB$, because no phase has belief at
431 full strength. The decomposition machinery is useful only to the extent that the
432 design supplies identifiability, and trial designers should evaluate any proposed
433 design by checking that each component has at least one phase or contrast in
434 which it is uniquely present.

435 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
436 *ut labore et dolore magna aliqua.*

437 Identifiability in the design-contrast sense above is necessary but not sufficient for
438 stable estimation at finite sample sizes. The three Gompertz trajectories share

an onset shape and overlap in calendar time, so their parameters can be collinear even when the contrasts that identify them in principle are present. Two conditions make the joint estimation credible: the design must supply at least one phase in which each component varies while the others are held fixed (the blinded-placebo contrast for BR , the open-label-versus-blinded contrast for PB , and an off-drug window for TV), and that off-drug window must be long enough relative to the carryover half-life for TV to separate from a still-decaying BR . Where these conditions are weak the components are jointly identified only up to a near-collinear direction, and the symptom of that is the rank-deficiency the pilot of section 7 encountered when the full phase-augmented formula was fitted at $N = 35$. We therefore do not assert that the three components are always recoverable. The pre-registered identifiability diagnostics of section 7 (convergence rates of the component-matched fit, variance-inflation factors on the component-specific Gompertz parameters, and the sensitivity of the BR estimate to the assumed $\eta(\text{phase})$ multipliers) are the means by which the recoverable region of the design-and-sample-size space is to be mapped, and the framework's participant-level claims should be read as conditional on falling inside that region.

5 Analysis-side methodology

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In this section we shall take up the practical question of what analysis model to use. A reasonable reaction to the formal model in section 2 is the following: if we are simulating data with three components plus noise, why is the analysis model written as $Sx \sim \text{bm} + \text{t} + \text{Dbc} + \text{bm:Dbc}$ with a single random intercept and AR(1) residual correlation? Should we not include three parametric Gompertz components in the analysis, mirroring the DGP, so that the analyst sees the same structure that generated the data? The question is reasonable. Four constraints make a component-matched analysis impractical for most aggregated N-of-1 trial sizes, and a modest extension (a phase-by-treatment indicator) captures the practically important part of the gap at minimal cost.

5.1 Four reasons not to match the DGP

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The first reason is identifiability. The trial design must supply the contrasts that identify each component, and many designs supply only part of it. The decomposition is identified by phase contrasts (open-label versus blinded for PB , on-drug versus blinded-placebo for BR , off-drug timepoints for TV). A pure open-label trial has none of these contrasts; in such a trial BR and PB both rise from zero on treatment initiation and are statistically indistinguishable from each other, re-

478 gardless of how the analyst writes the model. Forcing an analysis to fit three
479 components when the design supplies contrasts for only one or two of them does
480 not recover the missing components; it produces unidentified or weakly identified
481 parameter estimates with implausibly tight standard errors that misrepresent the
482 analyst's actual inferential precision.

483 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
484 *ut labore et dolore magna aliqua.*

485 The second reason is degrees of freedom. Each component term costs degrees of
486 freedom that come out of the moderation estimand. The biomarker-moderation
487 question (the core scientific aim of the predictive-biomarker validation program)
488 is asked through the `bm:Dcb` term. A component-matched analysis adds a Gompertz *BR* (three parameters), a Gompertz *PB* (three parameters plus a phase-specific expectancy multiplier), and a Gompertz *TV* (three parameters), each with participant-specific random effects on the maxima. That is on the order of
492 a dozen new parameters, half of them on a non-linear scale. At trial sample sizes
493 of $N \in [30, 150]$, the variance of $\hat{\beta}_{bm:D}$ is dominated by participant and timepoint
494 counts, and adding parametric structure to absorb residual confounding usually
495 inflates the moderation estimand's variance by more than it reduces bias.

496 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
497 *ut labore et dolore magna aliqua.*

498 The third reason is convergence fragility. Component-matched analyses are non-
499 linear in the parameters and their convergence is problematic in practice. Fitting
500 Gompertz *BR*, *PB*, and *TV* simultaneously requires a non-linear mixed model
501 (`nlme::nlme` with self-starting Gompertz forms), a Bayesian hierarchical fit, or
502 a structural-equation approximation. Each is more sensitive to starting values,
503 more prone to local optima, and more likely to fail to converge on a single trial
504 dataset than the linear-mixed analogue. At trial-relevant N , the convergence rate
505 of the component-matched fit is typically below 90%, and the non-converging cells
506 contaminate any aggregated analysis.

507 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
508 *ut labore et dolore magna aliqua.*

509 The fourth reason is that the moderation estimand is sufficiently robust to the BR-
510 versus-PB split for the most common scientific question. For predictive-biomarker
511 validation, the regulator wants to know whether the biomarker predicts treatment
512 response under the trial conditions, not a clean pharmacological-versus-placebo
513 decomposition of that prediction. The simple `bm:Dcb` interaction is the right
514 marginal estimand for the validation question whether the moderation is biological
515 or expectation-mediated.

516 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
517 *ut labore et dolore magna aliqua.*

518 It is worth stating the inferential target of the analysis in the vocabulary of the

estimands framework²⁷, because the biomarker-validation question is easy to pose ambiguously. The primary estimand is the population-level biomarker-treatment interaction, the $bm:D_{bc}$ coefficient: its population is the trial-eligible patient population, its treatment contrast is on-drug versus off-drug exposure summarized through the continuous indicator D_{bc} , its endpoint is the symptom trajectory, its population-level summary is the interaction coefficient, and its intercurrent events are handled by a treatment-policy strategy for dropout and by the modeled exponential decay for residual carryover. The pharmacological-component target β_{bm}^{BR} is a more ambitious estimand: it is the biomarker-by- BR interaction that would obtain with the belief-driven channel held at its off-drug level, identified here through the phase contrasts of section 4 rather than by assumption. We flag explicitly that recovering a participant-level interaction of this kind inherits the difficulty Senn has emphasized for individual treatment effects^{28,29}: a genuine participant-by-component signal is separable from participant-by-period and within-person random variation only when the design supplies repeated phase contrasts, and the framework should be read as estimating these quantities under that condition rather than as dissolving the well-known problem of identifying individual effects from finite within-person data.

5.2 A small extension that usually pays off: phase indicators

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The smallest useful step toward component-aware analysis, without paying the costs above, is to add a phase indicator and its interaction with the drug indicator. Schematically,

```
Sx ~ bm + t + Dbc + phase + Dbc:phase
      + bm:Dbc + bm:Dbc:phase
      + (1 | ptID), correlation = corCAR1(...)
```

Here `phase` is a categorical covariate (open-label, blinded, post-crossover) that absorbs the systematic differences in PB expectancy across phases. `Dbc:phase` tests whether the apparent drug effect shrinks under blinding, a signature that the open-label drug effect was partly PB -mediated. `bm:Dbc:phase` tests whether the moderation itself shrinks under blinding, a signature that the biomarker-treatment interaction was partly biomarker-by- PB rather than biomarker-by- BR .

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

This extension adds three to four parameters at most, well within budget at any realistic N . It does not separately estimate BR and PB as components, but it does identify the change in apparent drug effect that occurs when belief is hedged. For most analyses, this is the practically relevant question, and it can be answered without committing to the Gompertz parametric form. Where both the one-component

559 and the phase-augmented analyses are reported, the additional `Dbc:phase` and
560 `bm:Dbc:phase` terms introduce a small family of belief-attenuation tests, and the
561 attribution inference they support should be treated as a pre-specified secondary
562 analysis with its multiplicity acknowledged rather than as a second bite at the
563 primary interaction test.

564 5.3 When to fit the full decomposition

565 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
566 *ut labore et dolore magna aliqua.*

567 A component-matched analysis merits the effort when sample size is in the several
568 hundreds (a phase-3 trial fielded as aggregated N-of-1, or a meta-analysis pool-
569 ing multiple smaller trials); the trial design provides timepoints in every phase
570 combination, with enough density per cell to constrain the Gompertz curves; the
571 scientific question requires the BR-PB split (a pharmacology paper claiming a spe-
572 cific mechanism, where ruling out a placebo contribution is the central point); and
573 a Bayesian hierarchical fit with informative priors on the Gompertz parameters is
574 feasible. In that regime, `lcmm`, `nlme::nlme` with self-starting Gompertz forms, or
575 a `brms` hierarchical formulation can return participant-specific *BR*, *PB*, and *TV*
576 component estimates that mirror the DGP. Outside that regime, the linear-mixed
577 analysis with phase indicators is the better tool.

578 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
579 *ut labore et dolore magna aliqua.*

580 The asymmetry between DGP and analysis is not a flaw of the framework; it is a
581 deliberate consequence of the asymmetry between simulation (where the analyst
582 controls the DGP and can specify any structure, however elaborate) and inference
583 (where the analyst must fit a model to a finite dataset and pay the identifiability
584 cost of every additional parameter). Simulation studies use the rich DGP to
585 characterize how a simpler analysis behaves under realistic data generation; they
586 do not require the analysis to match the DGP in order to produce useful inference.

587 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
588 *ut labore et dolore magna aliqua.*

589 A word on the carryover parameterization is in order, since it is the one place
590 where the analysis model makes a strong functional-form commitment. The con-
591 tinuous drug indicator D_{bc} decays as $\exp(-\lambda t_{sd})$ off drug, a one-parameter ex-
592 ponential governed by the carryover half-life. This is deliberately parsimonious
593 and is best understood as a constrained special case of the more flexible carry-
594 over models now available for repeated-measures and N-of-1 data. Liao et al.³⁰
595 model carryover through a Bayesian distributed-lag structure with autocorrelated
596 errors, relaxing the single-rate exponential assumption at the cost of additional pa-
597 rameters; Daza³¹ formalizes an average-period treatment-effect estimand built to
598 accommodate autocorrelation, trend, and slow carryover jointly; and G{”a}rtner

et al.³² benchmark linear models against G-estimation and Bayesian-network alternatives specifically under carryover. We adopt the exponential form because it ties directly to a pharmacokinetic half-life and conserves degrees of freedom for the moderation estimand at trial-relevant sample sizes, but the choice is an assumption rather than a finding. The pre-registered simulation program of section 7 should accordingly include a distributed-lag analysis arm, so that the cost of the exponential-decay restriction is quantified against the richer alternative rather than asserted.

6 Worked example: information loss without decomposition

6.1 What the one-component analysis estimates

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Before the numerical illustration it is worth establishing algebraically what the one-component analysis actually estimates, because the qualitative failure modes of section 1 are consequences of an identity rather than empirical tendencies that require simulation to demonstrate. Write the within-participant symptom reduction as $R_{it} = BL_i - Y_{it} = BR_{it} + PB_{it} + TV_{it} - \varepsilon_{it}$. The one-component analysis fits a single drug indicator D_{it} ,

$$R_{it} = \alpha + \beta_D D_{it} + g(t) + u_i + e_{it},$$

and reports $\hat{\beta}_D$ as ‘the treatment effect’. Taking the on-drug minus off-drug contrast that β_D targets, and using that the natural-history component is by construction independent of treatment status so that its on/off difference cancels in a within-participant comparison, the probability limit of the estimator is

$$\text{plim } \hat{\beta}_D = m_{BR} + b_{PB} + b_{TV},$$

where $b_{PB} = E[PB \mid \text{on}] - E[PB \mid \text{off}]$ is the belief differential between the contrasted phases and b_{TV} collects any natural-history contribution that fails to cancel under the contrast actually used; b_{TV} is zero for a clean on/off within-participant contrast and strictly positive for a baseline-anchored pre-post analysis, in which case $b_{PB} = E[PB]$ and $b_{TV} = E[TV]$. To fix which case is in play below: the one-component analysis used in the worked example and the simulations of this paper is the design-contrast (drug-on/off) form, so $b_{TV} \approx 0$ there and the operative displacement is the belief term b_{PB} ; the baseline-anchored pre-post form, which additionally carries $b_{TV} = E[TV]$, is the worse case retained only for comparison. Both terms are generally non-negative for an improving outcome, so the one-component estimate over-states the pharmacological component m_{BR} . The displacement is a property of the probability limit and is unchanged as $N \rightarrow \infty$:

a larger one-component trial converges on the same biased value. This is the precise sense in which the placebo-attribution failure of section 1 is a problem of identification, not of precision.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The same algebra settles the biomarker-validation failure. The precision-medicine analysis regresses a participant-level response summary on a baseline biomarker B_i . If that summary is the one-component per-participant effect, which estimates $m_{BR,i} + w_{PB} m_{PB,i} + w_{TV} m_{TV,i}$ for design-determined loadings $w_{PB}, w_{TV} \geq 0$, then by the linearity of covariance the fitted biomarker slope has probability limit

$$\beta_{bm}^{\text{lumped}} = \frac{\text{Cov}(m_{BR} + w_{PB} m_{PB} + w_{TV} m_{TV}, B)}{\text{Var}(B)} = \beta_{bm}^{BR} + w_{PB} \beta_{bm}^{PB} + w_{TV} \beta_{bm}^{TV},$$

where $\beta_{bm}^C = \text{Cov}(m_C, B) / \text{Var}(B)$ is the component-specific slope. A biomarker correlated with placebo responsiveness or with the natural-history trajectory, so that β_{bm}^{PB} or β_{bm}^{TV} is non-zero, therefore registers a non-zero one-component slope even when its pharmacological association β_{bm}^{BR} is exactly zero. This is the biomarker-validation failure of section 1, and it too is an identity under the additive decomposition rather than an empirical tendency: no sample size separates the one-component slope from a genuine pharmacological one, because the two coincide in the estimand whenever $w_{PB} \beta_{bm}^{PB} + w_{TV} \beta_{bm}^{TV} \neq 0$.

We make two qualifications explicit, because they bound how much weight the framework can carry. First, both displays are immediate corollaries of the additive decomposition and the linearity of the projection that defines the estimator; they are not a derived result in any deeper sense, and their value is in naming the displacement terms, not in the algebra. Second, they assume the components combine additively. Under the subadditive interaction documented in section 3.3, where the joint effect of BR and PB falls below their sum, a first-order correction proportional to $\text{Cov}(m_{BR} \cdot m_{PB}, B)$ enters the biomarker slope, and the identities above should be read as the leading behavior in the additive case rather than as exact under interaction. The empirical magnitude of that correction term is itself an open question that the framework's γ -interaction simulation (section 3.3) is designed to address.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Two consequences for how the framework should be evaluated follow. First, the quantity that carries the bias is the location of the estimand, $\text{plim } \hat{\beta}_D$, not its sampling variability: the one-component coefficient remains a well-estimated number with a valid Wald test of its own nullity, and that test rejects at the rate dictated by the displaced value. A significance test of the one-component coefficient

therefore cannot, in principle, detect the bias, because the bias does not move the coefficient toward zero. Detecting it requires comparing $\hat{\beta}_D$ against m_{BR} , which is exactly the comparison the decomposition, or a phase contrast that isolates m_{BR} , makes available and the one-component analysis does not. Second, and as a corollary for the simulation evidence, a rejection-rate or power summary is the wrong instrument for this bias: the right target is the coefficient itself and its offset from m_{BR} . This is why the pilot of section 7.1 found the biomarker-treatment rejection rate essentially flat across the placebo-strength axis even though that axis is precisely what drives the predicted bias; the prediction lives in the mean of $\hat{\beta}_{bm:D_{bc}}$, and a production run that means to test it must read the bias off the coefficient, not off the test.

6.2 A numerical illustration

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

[Figure 1 here: two-panel trajectory plot. Left panel shows open-label phase (weeks 1–8) with overlapping response trajectories for Participants A and B. Right panel shows blinded-discontinuation phase (weeks 9–12) with the diagnostic divergence: A retains four points while B collapses to 1.5. Component contributions are labeled in each panel.]

Two participants in a sixteen-week N-of-1 trial of prazosin for PTSD share the same baseline but differ in latent traits: Participant A has $m_{BR}^A = 4$, $m_{PB}^A = 6$; Participant B has $m_{BR}^B = 8$, $m_{PB}^B = 1$. Both reach approximately ten points of open-label improvement by week eight, making them indistinguishable to a one-component analysis that estimates a population drug effect of 7.75 points. The blinded-discontinuation phase supplies the identifying contrast: A retains four points of improvement (belief persists at $\eta = 0.5$) while B returns nearly to baseline (drug cleared, minimal belief). The three-component decomposition recovers all six component estimates by exploiting this phase differential; the one-component model cannot.

7 Simulation study

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The simulation study reported and pre-registered in this section is structured under the³³ ADEMP framework, specifying Aims, Data-generating mechanisms, Estimands, Methods, and Performance measures. The pre-registered plan is at `analysis/scripts/component-decomposition/00-ademp-pre-registration.md`. A Monte Carlo simulation calibrated to the prazosin/PTSD reference parameters quantifies the bias and identifiability properties of the three analysis strategies described in section 5 across a parameter grid.

7.1 Pilot validation

[Details in supplement S1.]

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A 100-replicate pilot at $N = 35$ on a 6-cell grid confirmed the end-to-end pipeline. Two observations inform the production specification: at $N = 35$, rank-deficiency dropped the D_{bc} :phase interaction term, removing the key contrast; and the predicted bias appears in the coefficient mean, not the rejection rate, so the production program reports $\hat{\beta}_{bm:D_{bc}}$ against its implied true value rather than power alone. Full details and exploratory summaries are in supplement S1.

7.2 Production-simulation specification

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The production program follows the³³ ADEMP framework and targets four deliverables: the bias of $\hat{\beta}_{bm:D_{bc}}$ for the one-component, phase-augmented, and full-decomposition analyses across $N \in \{35, 70, 100, 150\}$ and a *pb_strength* grid; identifiability metrics for the full-decomposition fit (convergence, variance-inflation factors, and sensitivity of the *BR* estimate to the η (phase) multipliers); Type I error and power for the interaction test under an η -mismatch, *TV*-magnitude, and dropout sensitivity grid, with null cells at $n_{\text{reps}} = 5,000$ for MCSE 0.003; and the participant-level interaction signal recovered as a function of (m_{BR}, m_{PB}, m_{TV}) , which the one-component-total regression does not preserve. The *tidyverse* implementation collection drives the grid; driver scripts, checkpointed cell-level summaries, per-replicate output, and the committed ADEMP pre-registration all reside under `analysis/scripts/component-decomposition/`. The realised Study A run reported next used 1,000 replicates per alternative cell.

7.3 Study A: orthogonal-biomarker arm (1,000 replicates per alternative cell)

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The production Study A run (1,000 replicates per alternative cell, 5,000 per null cell, 80 alternative cells crossing $m_{PB} \in \{0, 1, 3, 6, 10\}$, $m_{TV} \in \{-1, 0, 1, 2\}$, and $N \in \{35, 70, 100, 150\}$, plus four null cells) is the orthogonal-biomarker arm of the conditional evaluation. The data-generating process couples the baseline biomarker to the *BR* component only, through the on-drug biomarker-*BR* covariance channel, so the true biomarker-treatment effect is the biomarker-by-*BR* slope, $\beta_{bm:D_{bc}} = -c_{bm} \sigma_{BR} / \sigma_{bm} = -2.25$, and the biomarker is by construction uncorrelated with *PB* and *TV*. Under the one-component analysis the mean estimated $\hat{\beta}_{bm:D_{bc}}$ tracks this value across the whole grid: the mean absolute bias is

0.016 at $N = 35$, 0.016 at $N = 70$, 0.006 at $N = 100$, and 0.006 at $N = 150$, in every case within one MCSE (respectively about 0.028, 0.019, 0.009, and 0.007) of zero, with no monotone trend across m_{PB} or m_{TV} . The predicted placebo-attribution bias of the biomarker-treatment estimate does not appear.

This is not a failure of the simulation; it is the behavior the identity of section 6.1 requires. That identity writes the one-component biomarker slope as $\beta_{bm}^{BR} + w_{PB}\beta_{bm}^{PB} + w_{TV}\beta_{bm}^{TV}$. With the biomarker coupled to BR alone, $\beta_{bm}^{PB} = \beta_{bm}^{TV} = 0$, so the one-component slope equals the pharmacological slope no matter how large m_{PB} or m_{TV} grows. Study A varies the magnitudes of the unmodeled components, which the identity shows to be the innocuous axis; the axis that produces biomarker-validation bias is the biomarker's correlation with PB or TV , which this DGP holds at zero. The clean null is therefore a confirmation of the identity and a correction to the looser claim of section 1 that large PB or TV by itself biases biomarker validation. It does not; biomarker-component coupling does.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The comparison between analyses is equally clear. The phase-augmented analysis is not less biased than the one-component analysis at any sample size (its mean absolute bias is 0.028 versus 0.016 at $N = 35$ and is comparable elsewhere), and it is markedly less powerful for the biomarker-treatment test: mean power across the grid is 0.35 versus 0.59 at $N = 35$ and 0.62 versus 0.90 at $N = 70$, with both analyses reaching essentially full power by $N = 100$. Because the DGP contains no biomarker- PB contamination for the phase contrast to remove, the extra phase-by-treatment terms add estimation cost without buying any bias reduction, and the analysis is dominated. Convergence was approximately complete for both analyses at every sample size (minimum convergence rate 0.998), so the rank-deficiency that forced dropping the `Dcb:phase` term in the $N = 35$ pilot, which we had advanced as the explanation for the pilot's inverted power ranking, is not operative in the production run: the inverted ranking persists with the full formula fitted and converged. The pilot's power inversion was therefore a genuine feature of this DGP, not an artifact of rank-deficiency.

Calibration is adequate for both analyses. Nominal-95% confidence-interval coverage of $\beta_{bm:D_{bc}}$ ranges from 0.96 to 0.99 for the one-component analysis (slightly conservative) and from 0.94 to 0.98 for the phase-augmented analysis. Type I error at the four null cells is controlled, ranging from 0.031 to 0.056 for the one-component analysis and 0.038 to 0.049 for the phase-augmented analysis; the one-component estimate at $N = 100$ (0.056, MCSE 0.003) sits 1.9 standard errors above nominal but is within the simulation's Monte Carlo uncertainty.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Two conclusions follow for the Study A program. First, the result sharpens rather than supports the section 1 framing: the biomarker-validation failure is real as an identity, but it is driven by the biomarker's coupling to PB or TV , not by the magnitude of those components, and Study A as pre-registered varies the latter while holding the former at zero. The decisive axis is the biomarker- PB correlation $c_{bm,PB}$, which the contaminated-biomarker pilot in the following subsection varies directly. Second, on the present evidence the phase-augmented analysis carries a real power cost without a compensating bias reduction under an uncontaminated biomarker, so it should not be recommended as a default for this design; the pilot below evaluates whether contamination changes this verdict.

7.4 Study B pilot: coupling-arm simulation (100 replicates per cell)

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A supplementary pilot varying the biomarker- PB correlation $c_{bm,PB} \in \{0, 0.15, 0.30, 0.45\}$ at two PB magnitude levels ($m_{PB} \in \{0, 6\}$) and two sample sizes ($N \in \{70, 150\}$), using 100 replicates per cell, confirms the identity-based prediction and reveals an unexpected structural result.

At $N = 150$ the one-component bias scales monotonically with $c_{bm,PB}$: mean $\hat{\beta}_{bm:D_{bc}}$ moves from -2.31 ($c_{bm,PB} = 0$, bias -0.06) to -2.37 ($c_{bm,PB} = 0.15$, bias -0.12) to -2.08 ($c_{bm,PB} = 0.30$, bias $+0.17$) to -1.74 ($c_{bm,PB} = 0.45$, bias $+0.51$). The last value sits 30 MCSEs from the true -2.25 , establishing the effect unambiguously. The direction of bias is positive (toward zero): high- $c_{bm,PB}$ participants experience stronger PB , which in the open-label phase adds to the apparent drug effect, but this elevation is partially absorbed by the bm main effect, leaving the $bm : D_{bc}$ interaction coefficient attenuated relative to the pure pharmacological value.

The unexpected result is that m_{PB} does not modulate the bias: cells with $m_{PB} = 0$ and $m_{PB} = 6$ produce identical bias at every contamination level (differences ≤ 0.006 at $N = 150$, well within Monte Carlo error). The driver is $c_{bm,PB} \times \sigma_{PB}$, not $c_{bm,PB} \times m_{PB}$. Even when the population-mean PB amplitude is zero, individual participants have nonzero PB draws (standard deviation $\sigma_{PB} = 2$), and coupling the biomarker to that variance is sufficient to contaminate the estimator. A true no- PB control requires $\sigma_{PB} = 0$, not $m_{PB} = 0$. This refines the Study A design: varying m_{PB} at $c_{bm,PB} = 0$ correctly produced zero bias, but because the relevant null was $c_{bm,PB} = 0$, not $m_{PB} = 0$.

The phase-augmented analysis provides no protection: its bias at every ($c_{bm,PB}, m_{PB}, N$) cell is indistinguishable from the one-component bias. Adding phase-by-treatment interactions to the model does not decompose the contamination because both analyses see the biomarker- PB covariance through the same open-label versus blinded-discontinuation contrast; the contrast changes

the weight on PB (via expectancy) but does not isolate the PB component's contribution to the biomarker slope. Separating that contribution requires both a design that decouples drug exposure from belief and a decomposition that exploits it; Study C shows that the standard hybrid design cannot achieve this separation, whereas a balanced-placebo design can.

The $N = 70$ results are uninformative at 100 replicates (Monte Carlo standard error ≈ 0.06 , comparable to or larger than the bias at moderate $c_{bm,PB}$). The production run will use 1,000 replicates per cell to establish the $N = 70$ pattern and 5,000 null replicates to assess Type I error under contamination. Until the production run is complete, the quantitative bias estimates from this pilot should be treated as directionally informative but imprecise.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The pilot results reframe the paper's practical guidance. The risk of biomarker-validation failure in an aggregated N-of-1 study arises not from the presence of a strong placebo response in the population but from the biomarker's correlation with individual-level placebo responsiveness. In a population where the baseline biomarker is uncorrelated with each participant's placebo-belief response amplitude, the one-component estimator recovers the pharmacological interaction without bias, and the full decomposition machinery is not required. In a population where the biomarker covaries with psychological characteristics that predict placebo response – optimism, prior treatment experience, suggestibility, expectancy at enrollment – the one-component analysis conflates pharmacological and placebo-belief prediction, and the scale of that conflation grows with the coupling strength. At $c_{bm,PB} = 0.45$, a value that equals the pharmacological coupling used throughout this paper, the one-component bias is approximately 0.51 points on the interaction scale, representing 23% of the true pharmacological effect. This is a practically meaningful distortion for any precision-medicine stratification decision.

The production run of the contaminated-biomarker study, at 1,000 alternative-cell replicates and 5,000 null-cell replicates, is required to confirm these estimates, establish the $N = 70$ pattern, and characterize Type I error under contamination. Results are pending.

7.5 Study C: recovery under a belief-decoupling design

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Study C asks whether the contamination demonstrated in the coupling-arm study can be removed by decomposition, and finds that the answer depends on the trial design. In the MVN process the biomarker- BR coupling acts only on drug and so loads on the drug-state contrast D_{bc} , whereas the biomarker- PB coupling

scales with the expectancy η ; interacting the biomarker with both isolates β_{bm}^{BR} on $bm:D_{bc}$, but only where D_{bc} and η vary independently. On the standard hybrid design they do not (on-drug time and belief accumulate together, with an analysis-stage correlation of about +0.6), and we verified that the belief-covariate decomposition, a blinded-stratum drug contrast, and a Gompertz-basis component decomposition all fail to recover β_{bm}^{BR} there. The remedy is a design that breaks the collinearity: a balanced-placebo (open-hidden) design^{25,26} interposes a covert on-drug phase, in which the drug is given without the participant's knowledge ($\eta = 0$), and an open-placebo phase, in which the participant is off drug but believes treatment continues ($\eta = 1$), reducing the D_{bc} - η correlation to about -0.66.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The production run (1{,}000 replicates per cell, $N \in \{70, 150\}$, $c_{bm, PB} \in \{0, 0.1, 0.2, 0.3\}$) confirms the recovery. At $N = 150$ the decomposition tracks the pharmacological target -2.25 across the contamination range, with bias -0.022, +0.008, +0.021, and -0.026, every value within one MCSE (about 0.02) of zero, while the one-component estimator develops a monotone bias of +0.16, +0.34, and +0.48 as $c_{bm, PB}$ grows. Interval coverage holds near nominal for the decomposition (0.94 to 0.96) and degrades to 0.86 for the one-component analysis. The power comparison (Table 2) is more nuanced than a uniform decomposition penalty: at zero contamination the one-component analysis is more powerful, by 0.19 at $N = 70$ (0.85 versus 0.66) and 0.07 at $N = 150$, reflecting the degrees of freedom the decomposition spends, but this advantage erodes monotonically as $c_{bm, PB}$ grows and reverses by $c_{bm, PB} = 0.3$ (0.64 versus 0.67 at $N = 70$). The reason is that contamination in this direction drags the one-component point estimate toward the null and so costs it power, whereas the decomposition holds the estimate at the pharmacological target; the one-component power advantage is accordingly confined to the low-contamination regime in which decomposition is unnecessary. Recovery is therefore conditional on the design, not on the analysis alone: the decomposition is necessary but a design that varies belief independently of drug exposure is the enabling condition. The practical recommendation is that a biomarker-validation trial in which the candidate biomarker may correlate with placebo responsiveness should be fielded as a balanced-placebo or open-hidden design rather than the standard hybrid design. Per-replicate output is at `analysis/data/derived_data/component-decomposition/study-b-balanced/`. Recovery beyond the positive-definite range, participant-level component recovery, and replication under the mean-moderation architecture coupling remain open.

Table 2: Rejection rate (power) for the biomarker-treatment interaction at $\alpha = 0.05$ under the balanced-placebo design, one-component versus belief-covariate decomposition, by contamination $c_{bm,PB}$ and sample size (1,000 replicates per cell; paired standard error of the gap ≤ 0.014). The one-component advantage at $c_{bm,PB} = 0$ erodes and reverses as contamination grows, because contamination drags the biased one-component estimate toward the null.

$c_{bm,PB}$	N	One-component	Decomposition	Gap
0.0	70	0.85	0.66	+0.19
0.1	70	0.78	0.67	+0.11
0.2	70	0.74	0.69	+0.05
0.3	70	0.64	0.67	−0.03
0.0	150	0.99	0.93	+0.07
0.1	150	0.98	0.93	+0.05
0.2	150	0.97	0.93	+0.04
0.3	150	0.93	0.93	−0.00

8 Application to predictive-biomarker validation

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The motivating clinical question for the trial designs in this project is whether a baseline biomarker predicts treatment response, and specifically whether the biomarker predicts the pharmacological component of treatment response. This is the precision-medicine question, formalized in the PATH statement⁸. Whether it depends on the BR-PB-TV decomposition is itself conditional: by the identity of section 6.1, the one-component biomarker slope equals the pharmacological slope when the biomarker is uncorrelated with PB and TV , and departs from it only by the biomarker’s covariance with those components. The decomposition is therefore the mechanism the question requires precisely in the regime where that covariance is non-negligible.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A direct regression of total response on baseline biomarker

```
total_response = beta_0 + beta_bm * biomarker + error
```

estimates the biomarker’s correlation with whatever the one-component total contains. If PB varies systematically across the biomarker distribution, because the biomarker is correlated with optimism, suggestibility, or treatment-seeking history, then $\hat{\beta}_{bm}$ contains a placebo-prediction component inseparable from the pharmacological-prediction component. The resulting biomarker, applied prospectively, would mis-stratify patients: those flagged as ‘high responders’ would include both true high- BR responders and high- PB responders, and the latter would not

benefit from prescription decisions made on a pharmacological basis.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The right approach is to regress each component on the biomarker:

$$\begin{aligned} BR &= \beta_0^{BR} + \beta_{bm}^{BR} \cdot \text{biomarker} + \varepsilon^{BR}, \\ PB &= \beta_0^{PB} + \beta_{bm}^{PB} \cdot \text{biomarker} + \varepsilon^{PB}, \\ TV &= \beta_0^{TV} + \beta_{bm}^{TV} \cdot \text{biomarker} + \varepsilon^{TV}. \end{aligned}$$

A biomarker is a useful pharmacological predictor if and only if $\hat{\beta}_{bm}^{BR}$ is meaningfully non-zero. Significant $\hat{\beta}_{bm}^{PB}$ or $\hat{\beta}_{bm}^{TV}$ are diagnostic findings about the biomarker, indicating that it is partly psychological or correlates with natural-history characteristics, but they are not the validation of pharmacological prediction. A biomarker that correlates only with PB would be valuable for selecting placebo responders in run-in designs but worthless for predicting pharmacological efficacy; a biomarker that correlates only with TV would be a prognostic marker, not a predictive one.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The decomposition is therefore the mechanism by which the precision-medicine question is posed whenever the candidate biomarker may correlate with PB or TV . Where the biomarker can be assumed orthogonal to the non-pharmacological components, the identity of section 6.1 shows the one-component analysis already targets the biomarker-by- BR estimand, and the simulation of section 7.3 confirms it recovers that estimand without bias. The phase-augmented linear-mixed analysis introduced in section 5 retains a useful subset of the decomposition machinery at small N (the `bm:Dbc:phase` interaction tests whether the biomarker-treatment effect attenuates under blinding, the practical proxy for biomarker-by- PB contamination), but Study A shows this proxy earns its degrees of freedom only when such contamination is present: under an uncontaminated biomarker it reduced no bias and lost power. The full decomposition recovers the component-specific regressions explicitly when N permits and when the scientific question warrants the cost.

9 Discussion

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

The three-component decomposition is the formal expression of the claim that ‘treatment response’ in an aggregated N-of-1 trial is not a single quantity but a sum of three causally distinct contributions, each with its own pharmacological,

971 psychological, or epidemiological interpretation. What the present work estab-
972 lishes, analytically in section 6.1 and empirically in section 7.3, is that the inferen-
973 tial value of acting on this decomposition is conditional rather than universal. The
974 identity shows that the one-component estimator is displaced from its pharmaco-
975 logical target only by the components' loading onto the contrast used, and, for
976 a biomarker slope, only by the biomarker's covariance with PB and TV . Where
977 those terms are zero, the design-contrast one-component analysis is already un-
978 biased, and the production simulation confirms it: with a biomarker coupled to
979 BR alone, the one-component biomarker-treatment estimate is unbiased across
980 the whole $PB-TV$ grid and the phase-augmented analysis is dominated. The
981 framework's contribution is therefore best read as a map of when decomposition
982 changes inference, not as a blanket prescription to decompose. It makes the trial-
983 design conditions explicit, identifies the attribution and contaminated-biomarker
984 estimands for which decomposition is needed, and clarifies the target estimand for
985 biomarker validation (the biomarker-by- BR association, not the one-component
986 total).

987 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
988 *ut labore et dolore magna aliqua.*

989 Four open questions remain and warrant acknowledgement. First, the simula-
990 tion study described in section 7 is forthcoming; the present manuscript reports
991 the framework rather than its empirical evaluation, and the quantitative bias-
992 versus-power trade-offs across analysis strategies are conjectures until that work is
993 run. Second, the η (phase) multipliers used to parameterize the expectancy mod-
994 ulation of PB are themselves modeling assumptions, and a sensitivity analysis
995 varying η across plausible values is mandatory before any clinical conclusion is
996 drawn from the framework's outputs. Third, the framework as specified treats
997 the three components as monotone Gompertz trajectories; cyclical or biphasic TV
998 patterns require either flexible TV specifications or explicit covariates beyond the
999 Gompertz form, and the boundary between 'flexible-Gompertz' and 'covariate-
1000 augmented' specifications is not yet sharply drawn. Fourth, the participant-level
1001 component estimates the framework targets inherit the identifiability caveat that²⁹
1002 and²⁸ raise for individual treatment effects more generally: within-participant com-
1003 ponent recovery rests on separating a true participant-by-component signal from
1004 participant-by-period and participant-by-phase variation, and this separation is
1005 credible only when the design supplies repeated phase contrasts and the cross-
1006 component covariance structure is not pathological. The framework should be
1007 read as estimating these quantities under those assumptions, not as dissolving
1008 the well-known difficulty of estimating individual effects from finite within-person
1009 data.

1010 *Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt*
1011 *ut labore et dolore magna aliqua.*

1012 The companion pedagogical document at docs/24-component-decomposition-pedagogy.md

develops the framework with worked examples and an extended intuition narrative for readers approaching the material for the first time. The companion methodological evaluation of the Gompertz family at [analysis/report/07-gompertz-evaluation/](#) quantifies the sensitivity of the framework's inference to the specific parametric form chosen for the component trajectories.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

A final point concerns the scope of the decomposition, where it is essential to distinguish the existence of the components from their identifiability. The claim that an observed longitudinal treatment response is the sum of a pharmacological, a belief-driven, and a natural-history component is a statement about the data-generating process of any repeated-measures treatment study, and in that sense the decomposition is not specific to N-of-1 trials. Identifiability, by contrast, is design-specific, and the distinction between population-level and participant-level recovery is decisive for the biomarker application. A standard parallel-group placebo-controlled trial supplies a between-arm placebo contrast that identifies the population-mean belief response but not a participant-level *PB*, and it contains no within-subject contrast in which belief varies independently of pharmacology; the participant-level biomarker-by-*BR* slope that the precision-medicine question targets is therefore not identified from such a design, no matter how large it is. The designs that do supply the required within-subject belief contrast are the blinded or open-hidden discontinuation and randomized-withdrawal designs^{25,26}, the crossover with on/off periods, and the hybrid N-of-1 design developed here; the pharmacometric disease-progression-plus-placebo models from which the framework descends^{1,2,34} operate at the group level and recover population-mean components only. We therefore advocate the decomposition as a default lens for longitudinal designs that supply a within-subject belief-by-exposure contrast, with the aggregated N-of-1 case the most demanding instance because it must recover the components from a single participant's trajectory. We do not claim participant-level identifiability from parallel-group placebo control alone, and where only a between-arm placebo contrast is available the framework's claims are restricted to the population mean.

10 Conclusions

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

In conclusion, three points are to be emphasized. First, treatment response in aggregated N-of-1 clinical trials is the sum of three causally distinct components: pharmacological action (*BR*), placebo-belief response (*PB*), and time-variant natural history (*TV*). Each component has its own biological mechanism, its own iden-

tifiability conditions, and its own clinical interpretation. Whether conflating them biases a given analysis is, however, governed by the omitted-variable-bias identity of section 6.1, not by the raw magnitude of PB and TV : the one-component pre-post treatment effect is displaced by the PB and TV contributions, whereas the biomarker-treatment slope is displaced only by the biomarker's covariance with those components. Study A bears this out: varying PB and TV magnitude with a biomarker coupled to BR alone left the one-component estimate unbiased across the entire grid.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Second, the operative requirement is the design contrast, not the parametric analysis model. The trial design must include the drug-on/off and belief contrasts that identify the components, and the one-component analysis already exploits the drug-on/off contrast through its D_{bc} term. Adding parametric component structure on top is not free: under an uncontaminated biomarker the phase-augmented analysis reduced no bias and lost substantial power, so it should not be adopted as a default. Its value, and that of the full component-matched fit, is confined to the estimands where the extra structure does work, namely the attribution of response to pharmacology versus belief and the analysis of biomarkers that may themselves correlate with PB or TV .

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Third, for predictive-biomarker validation the decomposition earns its place precisely when the candidate biomarker may correlate with PB or TV . A biomarker that predicts the one-component total is a pharmacological predictor only when it is uncorrelated with the non-pharmacological components; where that cannot be assumed, the one-component slope mixes pharmacological and non-pharmacological prediction, and the regulatory and prescribing decisions that motivate the trial require the separation the decomposition supplies. The coupling-arm study demonstrates this failure directly, and Study C establishes that the bias is recoverable: a belief-covariate decomposition returns the unbiased pharmacological slope across the feasible contamination range, but only on a balanced-placebo design that decouples drug exposure from belief, not on the standard hybrid design. The decisive practical lesson is therefore as much about design as about analysis.

11 Conflict of interest

The authors declare no conflicts of interest.

12 Funding

This work received no specific funding.

13 Data availability statement

The data and code that support the findings of this study are openly available in the pmsimstats-ng project repository. Per-replicate Monte Carlo outputs and ADEMP-compliant cell-level summaries are versioned under **analysis/data/**; simulation drivers are at **analysis/scripts/**. Exact paths for this manuscript are listed in the Reproducibility section below.

14 Reproducibility

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

This manuscript was rendered with R 4.4.x inside the project’s zzcollab Docker container (image hash recorded in `Dockerfile`; package set captured in `renv.lock`). Principal packages used by the simulation pipeline are `nlme` (continuous-time linear-mixed analysis with `corCAR1` residual correlation), `parallel` (8-core `mclapply` for parameter-grid sweeps), `data.table` (the original-extended simulation collection), `furrr` and `purrr` (the tidyverse simulation collection), and `rmarkdown` plus `bookdown` for the manuscript build.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Each simulation driver sets `set.seed(20260507)` at the top of the replicate loop and uses `parallel::mc.set.seed` inside `mclapply`. Per-replicate seeds derive from the master seed by `+ rep_idx`.

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Per-replicate simulation outputs are checkpointed at `analysis/data/quick-sim/06-decomp-replicates`. Morris ADEMP cell-level summaries at the companion `06-decomp-summary.txt`. Driver scripts are under `analysis/scripts/component-decomposition/`. The pre-registered ADEMP plan is at `analysis/scripts/component-decomposition/00-ademp-pre-regist`. The manuscript source, paper-specific bibliography (`references.bib`), and SIM CSL file are at `analysis/report/06-component-decomposition/`.

15 References

1. Holford N. Clinical pharmacology = disease progression + drug action. *British Journal of Clinical Pharmacology*. 2015;79(1):18-27. doi:10.1111/bcp.12170

- 1123 2. Gomeni R, Lavergne A, Merlo-Pich E. Modelling placebo response in depression trials using a longitudinal model with informative dropout. *European Journal of Pharmaceutical Sciences*. 2009;36(1):4-10. doi:10.1016/j.ejps.2008.10.025
- 1124 3. Kaptchuk TJ, Kelley JM, Conboy LA, et al. Components of placebo effect: Randomised controlled trial in patients with irritable bowel syndrome. *BMJ*. 2008;336(7651):999-1003. doi:10.1136/bmj.39524.439618.25
- 1125 4. Wager TD, Atlas LY, Lindquist MA, Roy M, Woo CW, Kross E. An fMRI-based neurologic signature of physical pain. *New England Journal of Medicine*. 2013;368(15):1388-1397. doi:10.1056/NEJMoa1204471
- 1126 5. Zunhammer M, Bingel U, Wager TD. Placebo effects on the neurologic pain signature: A meta-analysis of individual participant functional magnetic resonance imaging data. *JAMA Neurology*. 2018;75(11):1321-1330. doi:10.1001/jamaneurol.2018.2017
- 1127 6. Hall KT, Lembo AJ, Kirsch I, et al. Catechol-O-methyltransferase val158met polymorphism predicts placebo effect in irritable bowel syndrome. *PLoS ONE*. 2012;7(10):e48135. doi:10.1371/journal.pone.0048135
- 1128 7. Wang RS, Lembo AJ, Kaptchuk TJ, Hall KT, et al. Genomic effects associated with response to placebo treatment in a randomized trial of irritable bowel syndrome. *Frontiers in Pain Research*. 2022;2:775386. doi:10.3389/fpain.2021.775386
- 1129 8. Kent DM, Paulus JK, Klaveren D van, et al. The Predictive Approaches to Treatment effect Heterogeneity (PATH) Statement. *Annals of Internal Medicine*. 2020;172(1):35-45. doi:10.7326/M18-3667
- 1130 9. Raskind MA, Peterson K, Williams T, et al. A trial of prazosin for combat trauma PTSD with nightmares in active-duty soldiers returned from Iraq and Afghanistan. *American Journal of Psychiatry*. 2013;170(9):1003-1010. doi:10.1176/appi.ajp.2013.12081133
- 1131 10. Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: Statistical methods and practical considerations. *Frontiers in Digital Health*. 2020;2:13. doi:10.3389/fdgth.2020.00013
- 1132 11. Holford NHG. Holford NHG and sheiner LB 'understanding the dose-effect relationship - clinical application of pharmacokinetic-pharmacodynamic models', Clin Pharmacokin 6:429-453 (1981) - the backstory. *The AAPS Journal*. 2011;13(4):662-664. doi:10.1208/s12248-011-9306-5

- 1133 12. Tilburt JC, Emanuel EJ, Kaptchuk TJ, Curlin FA, Miller FG. Prescribing 'placebo treatments': Results of national survey of US internists and rheumatologists. *BMJ*. 2008;337:a1938. doi:10.1136/bmj.a1938
- 1134 13. Finniss DG, Kaptchuk TJ, Miller F, Benedetti F. Biological, clinical, and ethical advances of placebo effects. *The Lancet*. 2010;375(9715):686-695. doi:10.1016/S0140-6736(09)61706-2
- 1135 14. Benedetti F. Placebo effects: From the neurobiological paradigm to translational implications. *Neuron*. 2014;84(3):623-637.
- 1136 15. Frisaldi E, Shaibani A, Benedetti F. Understanding the mechanisms of placebo and nocebo effects. *Swiss Medical Weekly*. 2020;150:w20340. doi:10.4414/smw.2020.20340
- 1137 16. Carlino E, Frisaldi E, Benedetti F. Pain and the context. *Nature Reviews Rheumatology*. 2014;10(6):348-355. doi:10.1038/nrrheum.2014.17
- 1138 17. Carlino E, Piedimonte A, Benedetti F. Nature of the placebo and nocebo effect in relation to functional neurologic disorders. In: *Handbook of Clinical Neurology*. Vol 139.; 2016:597-606. doi:10.1016/B978-0-12-801772-2.00048-5
- 1139 18. Kirsch I. Response expectancy as a determinant of experience and behavior. *American Psychologist*. 1985;40(11):1189-1202. doi:10.1037/0003-066X.40.11.1189
- 1140 19. Kaptchuk TJ, Friedlander E, Kelley JM, et al. Placebos without deception: A randomized controlled trial in irritable bowel syndrome. *PLoS ONE*. 2010;5(12):e15591. doi:10.1371/journal.pone.0015591
- 1141 20. Lembo A, Kelley JM, Nee J, et al. Open-label placebo vs double-blind placebo for irritable bowel syndrome: A randomized clinical trial. *Pain*. 2021;162(9):2428-2435. doi:10.1097/j.pain.0000000000002234
- 1142 21. Nurko S, Saps M, Kossowsky J, others, Kelley JM. Effect of open-label placebo on children and adolescents with functional abdominal pain or irritable bowel syndrome. *JAMA Pediatrics*. 2022;176(4):349-356. doi:10.1001/jamapediatrics.2021.5750
- 1143 22. Bingel U, Wanigasekera V, Wiech K, et al. The effect of treatment expectation on drug efficacy: Imaging the analgesic benefit of the opioid remifentanyl. *Science Translational Medicine*. 2011;3(70):70ra14. doi:10.1126/scitranslmed.3001244
- 1144 23. Hróbjartsson A, Gøtzsche PC. Is the placebo powerless? An analysis of clinical trials comparing placebo with no treatment. *New England Journal of Medicine*. 2001;344(21):1594-1602. doi:10.1056/NEJM200105243442106

- 1145 24. Hróbjartsson A, Gøtzsche PC. Placebo interventions for all clinical conditions. *Cochrane Database of Systematic Reviews*. 2010;(1):CD003974. doi:10.1002/14651858.CD003974.pub3
- 1146 25. Colloca L, Lopiano L, Lanotte M, Benedetti F. Overt versus covert treatment for pain, anxiety, and Parkinson's disease. *The Lancet Neurology*. 2004;3(11):679-684.
- 1147 26. Rohsenow DJ, Marlatt GA. The balanced placebo design: Methodological considerations. *Addictive Behaviors*. 1981;6(2):107-122.
- 1148 27. International Council for Harmonisation of Technical Requirements for Pharmaceuticals for Human Use. ICH E9(R1) addendum on estimands and sensitivity analysis in clinical trials to the guideline on statistical principles for clinical trials. Published online 2019.
- 1149 28. Senn S. Controversies concerning randomization and additivity in clinical trials. *Statistics in Medicine*. 2004;23(24):3729-3753. doi:10.1002/sim.2074
- 1150 29. Senn S. Statistical pitfalls of personalized medicine. *Nature*. 2018;563(7733):619-621.
- 1151 30. Liao Z, Qian M, Kronish IM, Cheung YK. Analysis of N-of-1 trials using Bayesian distributed lag model with autocorrelated errors. *Statistics in Medicine*. 2023;42(13):2044-2060. doi:10.1002/sim.9676
- 1152 31. Daza EJ. Causal analysis of self-tracked time series data using a counterfactual framework for N-of-1 trials. *Methods of Information in Medicine*. 2018;57(S 01):e10-e21. doi:10.3414/ME16-02-0044
- 1153 32. Gärtner T, Schneider J, Arnrich B, Königorski S. Comparison of Bayesian networks, G-estimation and linear models to estimate causal treatment effects in aggregated N-of-1 trials with carry-over effects. *BMC Medical Research Methodology*. 2023;23:191. doi:10.1186/s12874-023-02012-5
- 1154 33. Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Statistics in Medicine*. 2019;38(11):2074-2102.
- 1155 34. Chen C, Jönsson S, Yang S, Plan EL, Karlsson MO. Detecting placebo and drug effects on Parkinson's disease symptoms by longitudinal item-score models. *CPT: Pharmacometrics & Systems Pharmacology*. 2021;10(4):309-317. doi:10.1002/psp4.12601