

1 Test-procedure choice for the biomarker-treatment interaction
2 in aggregated N-of-1 trials: dichotomization, classical
3 RM-ANOVA, mixed-effects, and generalized estimating
4 equations

5 pmsimstats team

6 August 15, 2026

7 **Contents**

8	1 Abstract	2
9	2 Introduction	3
10	2.1 A claim about how design and test choices interact	3
11	2.2 Why the question matters for chronic-condition pharmacology	3
12	2.3 What the published N-of-1 methodology literature does and does not do	4
13	2.4 Run-in design and the placebo-selection controversy	4
14	2.5 What this paper contributes	4
15	3 Closed-form strict RM-ANOVA test of the biomarker-treatment interaction	5
16	3.1 Notation and design	5
17	3.2 Sums of squares and the interaction F -statistic	6
18	3.3 Relation to the linear-mixed Wald test	8
19	4 Cycle-by-period design grid	10
20	5 Simulation design	10
21	6 Results	12
22	6.1 Test-procedure comparison at fixed design	12
23	7 Design-and-test recommendation	14
24	8 Discussion	15
25	8.1 What the Monte Carlo study settles	15
26	9 Conclusions	16
27	10 Conflict of interest	17
28	11 Funding	17

29	12 Data availability statement	17
30	13 Reproducibility	17
31	14 References	18

32 1 Abstract

33 **Background.** Inference about the biomarker-treatment interaction in aggregated N-of-1
34 trials depends on the analyst’s test procedure: classical repeated-measures ANOVA (RM-
35 ANOVA), the linear mixed-effects model (LME), or generalized estimating equations (GEE). A
36 common but rarely examined feature of the classical analysis is that it requires the continuous
37 biomarker to be dichotomized. We quantify the consequences of test-procedure choice for the
38 interaction, separating the cost of dichotomization from the cost of the test procedure proper.

39 **Methods.** We derive the closed-form strict RM-ANOVA F -statistic for the biomarker-
40 stratified treatment-by-time interaction in a balanced split-plot design with categorical pre-
41 dictors and the sphericity assumption, and relate it to the Wald t -test for the corresponding
42 fixed-effects coefficient in a linear mixed model with continuous-time AR(1) residual correla-
43 tion. We then conduct a Monte Carlo simulation study, designed and reported under the¹
44 ADEMP framework, that crosses five test procedures (strict RM-ANOVA on the dichotomized
45 biomarker; a continuous-biomarker within-subject contrast; the standard pmsimstats linear-
46 mixed analysis with `corCAR1`; and GEE with naive and with Mancl-DeRouen small-sample
47 sandwich variance) against the biomarker-moderation coefficient $\beta_{bm} \in \{0, 0.45\}$ and sample
48 size $N \in \{25, 35, 70, 100\}$, a 40-cell factorial at the prazosin-PTSD reference design with 1000
49 replicates per power cell and 2000 per Type I cell. The continuous-biomarker contrast dif-
50 fers from strict RM-ANOVA only in retaining the continuous biomarker, isolating the cost of
51 dichotomization. The cycle-by-period design grid is specified here but its empirical sweep is
52 deferred to a companion paper.

53 **Results.** The closed-form strict RM-ANOVA test agrees with the linear-mixed Wald test
54 under sphericity and balanced design. In the simulation, most of the power gap between
55 strict RM-ANOVA and the mixed model is the cost of dichotomization rather than of the test
56 procedure: at $N = 25$, replacing the dichotomized biomarker with the continuous one raises
57 power by 0.180, while the further step to the full mixed model adds only 0.027, a pattern
58 stable across N (dichotomization 0.08–0.21, test procedure at most 0.04). Among the four
59 continuous-biomarker procedures, the linear-mixed analysis attains the highest power within
60 nominal Type I error; GEE with the naive sandwich inflates Type I error to between $1.5\times$ and
61 $1.9\times$ nominal across all four sample sizes, which the Mancl-DeRouen correction removes at
62 a moderate (about 10-percentage-point) power cost while restoring 93–95% interval coverage.
63 The inline Mancl-DeRouen standard error agrees with the `geesmv` reference to within 2.5%.

64 **Conclusions.** The analyst’s most consequential decision is to keep the biomarker contin-
65 uous rather than dichotomize it; among continuous-biomarker procedures the linear-mixed
66 analysis is the preferred default and a dichotomization-free within-subject contrast is a close,
67 transparent alternative. Naive GEE should not be used at the sample sizes examined without

a small-sample correction. This paper resolves the test-procedure axis at a fixed reference design; the joint cycle-by-period design grid is specified here and deferred to a companion paper.

2 Introduction

2.1 A claim about how design and test choices interact

The present paper develops the methodological case for treating the analyst's two principal choices in an aggregated N-of-1 trial jointly rather than separately. The two choices are the test procedure for the biomarker-treatment interaction (strict repeated-measures ANOVA, the linear mixed-effects model with continuous-time AR(1) residual correlation, the mixed-effects model for repeated measures, generalized estimating equations) and the trial-design parameters that govern within-subject treatment cycling (number of treatment cycles, on-drug and off-drug period lengths, on/off symmetry, optional run-in and run-out phases). The validity and power of any given test depend on what the design assumes about carryover, autocorrelation, and within-subject variance structure; conversely, the optimal design parameters depend on which test the analyst plans to use. Treating the two as independent optimizations, as the published methodological literature typically does, produces locally optimal but jointly suboptimal combinations, and we shall argue that this gap is consequential at the sample sizes characteristic of precision-medicine pilot studies.

2.2 Why the question matters for chronic-condition pharmacology

In this section we shall set out three reasons why the joint treatment of test procedure and design parameters is warranted specifically for the aggregated N-of-1 design class.

We begin with the physiological motivation. Within-participant residuals in chronic-condition trials are autocorrelated not as an analytic convenience but because most chronic-condition response trajectories exhibit serial correlation on a timescale set by drug pharmacokinetics, condition-specific natural-history drift, and measurement reactivity (Hawthorne, structured-visit, and diary-completion effects). Continuous-time AR(1) residual correlation is, in our view, the appropriate structural prior for this physiology, but it does not satisfy the sphericity assumption required by strict RM-ANOVA.² show through simulation that uncorrected RM-ANOVA inflates Type I error under sphericity violations; they also report, importantly, that the Greenhouse-Geisser-corrected RM-ANOVA and the multilevel linear model perform comparably once the correction is applied, and that the multilevel model is not uniformly safer (it can be the more liberal of the two in some conditions). The degrees-of-freedom corrections^{3,4} recover nominal levels only conditionally.⁵ extend this evidence with finer epsilon resolution, with the conventional rule favoring Greenhouse-Geisser when $\epsilon < 0.75$. The advantage we claim for the mixed-effects model with explicit `corCAR1` residual structure is thus not that it is uniformly more powerful than a well-corrected RM-ANOVA, but that it models the AR(1) structure directly, accommodates the continuous biomarker and unequal time spacing, and so avoids both the dichotomization and the degrees-of-freedom penalty that the classical route incurs.

2.3 What the published N-of-1 methodology literature does and does not do

We begin with the state of the analytic-model literature. The systematic review of⁶ documents that 83.8% of the published N-of-1 studies they catalogue ignore autocorrelation entirely;⁷ and⁸ reach similar conclusions in neurology and behavioral-medicine subdomains. CONSORT extension reporting standards^{9,10} codify reporting quality but do not address the analytic-model decision in any detail, and¹¹ identifies the analytic-method gap as the highest-priority methodological problem in the field. It appears, then, that the majority of N-of-1 analyses are conducted under a misspecified residual covariance assumption, and that the field has not yet converged on remedies.

Test-procedure comparisons for longitudinal data are not new. The generic contrast of RM-ANOVA against GEE and mixed-effects models is well established:¹² for instance, report through simulation that RM-ANOVA is appreciably less powerful than GEE and mixed models for continuous longitudinal endpoints, the same ordering we recover. What that literature does not do is treat the biomarker-treatment interaction estimand, the aggregated N-of-1 / cross-over design, or the continuous-time AR(1) structure together, nor does it isolate the dichotomization cost that the classical route imposes specifically when the moderator is continuous. The MMRM consensus established by Mallinckrodt and colleagues^{13–16} applies to parallel-group longitudinal trials with monotone missingness, not to aggregated N-of-1 designs. The GEE methodology of¹⁷ with sandwich variance estimation has been adapted to small-cluster settings via bias-corrected estimators^{18,19}, but the small-sample concerns are acute at trial-relevant N-of-1 cluster counts:²⁰ report that in stepped-wedge cluster trials with fewer than six independent units only 6.6% of analyses use appropriate small-sample corrections, illustrating the gap between method availability and practice. None of these papers address the biomarker-interaction estimand in the cross-over or N-of-1 setting directly.

The gap, then, is this. Across the recent systematic reviews^{6,7,21} and the methodological literature surveyed above, no published study crosses test procedure (strict RM-ANOVA, MMRM, GEE, mixed-effects with `corCAR1`) against a design-parameter grid (cycles, period lengths, asymmetry, run-in, run-out) under simulation for the aggregated N-of-1 design class. The present paper endeavors to address that gap, at least on the test-procedure axis.

2.4 Run-in design and the placebo-selection controversy

2.5 What this paper contributes

We make four contributions, each addressed in a section below.

First, we derive the closed-form strict RM-ANOVA F -statistic for the biomarker-stratified treatment-by-time interaction in a balanced split-plot design with categorical predictors and the sphericity assumption (§2). We relate the strict-RM-ANOVA test to the Wald t -test on the biomarker-treatment fixed-effects coefficient in a linear mixed model with continuous-time AR(1) residual correlation, identifying the conditions on residual covariance under which the two tests agree and the regimes in which they diverge. The derivation extends standard

146 cross-over methodology^{22,23} to the biomarker-stratified case and clarifies the connection to
147 the contemporary sphericity-violation literature^{2,5}.

148 Third, we conduct a Monte Carlo simulation study (§4-§5) that crosses five test pro-
149 cedures (strict RM-ANOVA on the dichotomized biomarker; a continuous-biomarker
150 within-subject contrast; the standard pmsimstats linear-mixed analysis with `corCAR1`;
151 and GEE with naive and with Mancl-DeRouen²⁴ small-sample sandwich variance)
152 against the biomarker-moderation coefficient and sample size at the fixed reference
153 design. The continuous-biomarker contrast is the dichotomization-free analogue
154 of the strict RM-ANOVA test and lets us separate the cost of dichotomizing the
155 biomarker from the cost of the test procedure proper. The simulation runs on the
156 pmsimstats package data-generating process under the mean-moderation architec-
157 ture via `generateData()` and `lme_analysis()`, holding fixed the response-curve and
158 covariance-structure axes addressed in `analysis/report/07-gompertz-evaluation/`
159 and `analysis/report/01-dgp-mean-moderation-vs-mvn/`. Results are reported in the¹
160 ADEMP framework adopted as the project-wide simulation-reporting standard.

161 3 Closed-form strict RM-ANOVA test of the biomarker- 162 treatment interaction

163 In this section we shall derive the closed-form strict RM-ANOVA F -statistic for the biomarker-
164 treatment interaction in a balanced split-plot design with categorical predictors and the
165 sphericity assumption, and relate it to the Wald t -statistic on the corresponding linear-
166 mixed coefficient. The derivation specializes classical split-plot machinery²⁵ to the biomarker-
167 stratified case in the simplest design that supports the test: two biomarker strata, two treat-
168 ment phases, and T within-phase timepoints. The presentation is deliberately confined to
169 strict RM-ANOVA. Modern extensions — the mixed-effects model for repeated measures²⁶,
170 generalized estimating equations¹⁷, and ANCOVA with random subject effects — appear only
171 for context; the closed-form expression we report is the textbook compound-symmetric form.
172 Three findings anchor the rest of the paper:

- 173 1. Under sphericity and balance, the strict RM-ANOVA test of the biomarker-treatment
174 interaction is identical to a two-sample t -test on the within-subject treatment differences,
175 with degrees of freedom $2(n - 1)$.
- 176 2. The closed form yields a clean non-central F power formula that we use as the analytic
177 anchor for the simulation study in §4-§5.
- 178 3. Five assumptions of the closed form are violated by the pmsimstats data-generating
179 process (DGP; continuous biomarker, AR(1) within-subject covariance, exponential car-
180 ryover, missing data, unequal time spacing). Each violation favors the linear-mixed
181 analysis over the strict RM-ANOVA test in operating characteristics, motivating the
182 primary-and-sensitivity reporting strategy of §6.

183 3.1 Notation and design

184 Specifically, we consider the following structural elements:

- Two biomarker strata $i \in \{1, 2\}$, formed by dichotomizing the continuous biomarker at a pre-specified cut-point. Strict RM-ANOVA requires categorical predictors; the dichotomization is the price of admission to the framework.
- n subjects per stratum, balanced. Index $l \in \{1, \dots, n\}$.
- Two treatment phases $j \in \{1, 2\}$, active drug and placebo. In an N-of-1 design these are two phases through which every subject passes; in a parallel-group design they are between-subject and the design is not a split-plot.
- T within-phase timepoints, indexed $k \in \{1, \dots, T\}$. In the simplest case $T = 1$, the analyst pre-averages within-phase observations to a single mean per subject per phase.

We write Y_{ijkl} for the symptom outcome at biomarker stratum i , treatment phase j , subject l within stratum i , and timepoint k . The strict classical split-plot model is

$$Y_{ijkl} = \mu + \alpha_i + \beta_j + (\alpha\beta)_{ij} + \pi_{l(i)} + \tau_k + (\alpha\tau)_{ik} + (\beta\tau)_{jk} + (\alpha\beta\tau)_{ijk} + e_{ijkl}, \quad (1)$$

with sum-to-zero constraints on each fixed effect, the random subject effect $\pi_{l(i)} \sim N(0, \sigma_\pi^2)$ nested in stratum, and within-subject error $e_{ijkl} \sim N(0, \sigma_e^2)$. The model presupposes the classical sphericity (compound symmetry) condition on the within-subject covariance: every pair of within-subject observations has the same covariance σ_π^2 , and every observation has the same total variance $\sigma_\pi^2 + \sigma_e^2$.

3.2 Sums of squares and the interaction F -statistic

We begin the derivation by partitioning the total sum of squares. The total sum of squares partitions into between-subjects and within-subjects components:

$$SS_{\text{tot}} = SS_{\text{between}} + SS_{\text{within}}. \quad (2)$$

The between-subjects component decomposes further into SS_A (biomarker stratum) and $SS_{S(A)}$ (subject within stratum):

$$SS_A = bnT \sum_{i=1}^a (\bar{Y}_{i...} - \bar{Y}_{...})^2, \quad SS_{S(A)} = bT \sum_{i=1}^a \sum_{l=1}^n (\bar{Y}_{i.l} - \bar{Y}_{i...})^2. \quad (3)$$

The within-subjects component decomposes into SS_B (treatment phase), SS_{AB} (the biomarker-treatment interaction, the primary target of the test), and a within-subjects error term $SS_{B \times S(A)}$:

$$SS_{AB} = nT \sum_{i,j} (\bar{Y}_{ij..} - \bar{Y}_{i...} - \bar{Y}_{.j..} + \bar{Y}_{...})^2, \quad (4)$$

$$SS_{B \times S(A)} = T \sum_{i,j,l} (\bar{Y}_{ijl} - \bar{Y}_{i..} - \bar{Y}_{.jl} + \bar{Y}_{...})^2. \quad (5)$$

Corresponding degrees of freedom are $df_A = a - 1$, $df_{S(A)} = a(n - 1)$, $df_B = b - 1$, $df_{AB} = (a - 1)(b - 1)$, and $df_{B \times S(A)} = a(n - 1)(b - 1)$, the last being the within-subjects error degrees of freedom appropriate for the treatment effect and its interaction with the between-subjects factor. The biomarker-treatment interaction is then tested by the F -ratio

$$F_{AB} = \frac{MS_{AB}}{MS_{B \times S(A)}} = \frac{SS_{AB}/df_{AB}}{SS_{B \times S(A)}/df_{B \times S(A)}}. \quad (6)$$

Under the null $H_0 : (\alpha\beta)_{ij} \equiv 0$ and the sphericity assumption, $F_{AB} \sim F(df_{AB}, df_{B \times S(A)})$.

3.2.1 The 2×2 split-plot case

We now specialize to $a = b = 2$ (two biomarker strata, two phases), setting $T = 1$ for cleanest exposition; that is, the analyst pre-averages within-phase observations to a single value per subject per phase before computing the interaction. The sample interaction contrast is

$$\hat{L} = (\bar{Y}_{11.} - \bar{Y}_{12.}) - (\bar{Y}_{21.} - \bar{Y}_{22.}), \quad (7)$$

That is, $\hat{L}$ is the difference of within-subject treatment differences across biomarker strata, and is the natural biomarker-treatment interaction statistic. The interaction parameter under the sum-to-zero ANOVA parameterization satisfies $\widehat{(\alpha\beta)}_{11} = \hat{L}/4$ (and the other three cells take values $\pm\hat{L}/4$), so the interaction sum of squares is $SS_{AB} = nT\hat{L}^2/4$, and with $df_{AB} = 1$ this equals the interaction mean square. Under sphericity and balance, $MS_{B \times S(A)} = \hat{\sigma}_e^2$. Substituting into equation (6) yields

$$F_{AB} = \frac{nT\hat{L}^2}{4\hat{\sigma}_e^2}, \quad df_1 = 1, \quad df_2 = 2(n - 1). \quad (8)$$

3.2.2 Equivalence to a within-subject t -test

We note that equation (8) is equivalent to a two-sample t -test on the within-subject differences. Define $\Delta_{il} = \bar{Y}_{i1l} - \bar{Y}_{i2l}$ as the within-subject treatment difference for subject l in biomarker stratum i , averaged over T within-phase timepoints. Then

$$\hat{L} = \bar{\Delta}_1 - \bar{\Delta}_2, \quad \bar{\Delta}_i = \frac{1}{n} \sum_{l=1}^n \Delta_{il}. \quad (9)$$

Under sphericity, Δ_{il} are i.i.d. within stratum with $\text{Var}(\Delta_{il}) = 2\sigma_e^2/T$, since the random subject effect cancels in the within-subject difference. The variance of $\hat{L}$ is therefore

$$\text{Var}(\hat{L}) = \frac{2 \cdot 2\sigma_e^2}{nT} = \frac{4\sigma_e^2}{nT}, \quad (10)$$

and the t -statistic for the contrast is

$$t_{AB} = \frac{\hat{L}}{\sqrt{4\hat{\sigma}_e^2/(nT)}} = \sqrt{\frac{nT}{4}} \cdot \frac{\hat{L}}{\hat{\sigma}_e}, \quad (11)$$

3.3 Relation to the linear-mixed Wald test

We shall now relate the strict-RM-ANOVA F -statistic of equation (8) to the `pmsimstats` `nlme::lme` Wald t -statistic on the `bm:Dbc` coefficient. The two are asymptotically equivalent under three conditions: (i) the biomarker is dichotomized at the cut-point that defines the strata, (ii) the within-subject covariance is exactly compound-symmetric, and (iii) `Dbc` is binary, the E1 specification of the carryover-sensitivity manuscript rather than the matched-decay E2 specification. When all three conditions hold, the two tests reach the same coefficient and the same standard error up to small-sample correction factors. When any one fails, the mixed-model test extracts more information than the strict RM-ANOVA: the continuous biomarker contributes within-stratum variation, the AR(1) covariance is modeled directly, and the continuous exposure-decay predictor draws on the off-drug trajectory.

The `pmsimstats` DGP imposes AR(1) within-factor correlation, which is not compound-symmetric. Under AR(1), Box's ϵ is strictly below 1, and treating the full multi-timepoint series as repeated within-subject levels under classical RM-ANOVA inflates Type I error unless a degrees-of-freedom correction is applied. The relevant magnitude is recorded in the project's own audit (`docs/06-ar1-residual-correlation.tex`): at $T = 8$ timepoints, fitting a mixed model with no residual-correlation structure (the naive analogue of treating AR(1) data as exchangeable) produces Type I error of 13 to 17 percent for the crossover (CO) and open-label (OL) designs, which switching to `nlme::lme` with `corCAR1` returns to nominal. We stress that this 13-to-17-percent figure is for the unstructured/naive fit, not for the strict RM-ANOVA arm evaluated in this paper: that arm pre-averages each phase to a single value (the 2×2 form of §2.2), so its within-subject covariance is 2×2 , sphericity holds by construction, and its Type I error is at nominal (§5). The sphericity penalty is therefore a property of the un-pre-averaged classical analysis; for the pre-averaged strict RM-ANOVA the operative cost is dichotomization, not sphericity. Either way the linear-mixed analysis with `corCAR1` avoids both costs, modeling the AR(1) structure directly rather than absorbing it via a degrees-of-freedom penalty.

3.3.1 Power calculation under the closed form

For the 2×2 case, equation (8) yields a closed-form power formula that we shall use as the analytic anchor for the simulation study in §4–§5. Under the alternative $H_1 : \hat{L} = \delta$, the statistic $F_{AB} = nT \hat{L}^2 / (4\hat{\sigma}_e^2)$ follows a non-central F distribution with non-centrality parameter

$$\lambda = \frac{nT \delta^2}{4\sigma_e^2}, \quad (12)$$

and power at level α is

$$1 - \beta = \Pr \left[F_{AB} > F_{1,2(n-1),1-\alpha} \mid \lambda \right], \quad (13)$$

3.3.2 Limitations of the strict classical formulation

Unfortunately, the closed-form simplicity of equation (8) is bought at the price of five strong assumptions, each of which is empirically problematic in the pmsimstats setting. We present them in approximate order of severity.

1. *Categorical biomarker.* The pmsimstats biomarker (resting SBP in the prazosin-PTSD application) is continuous and has no defensible cut-point. Dichotomizing at the median or at a clinical threshold loses the within-stratum biomarker variation, attenuates the interaction estimator, and inflates Type II error. In the carryover-sensitivity production data the median split gives a $\hat{L}$ standard error roughly 1.4 to 1.7 times that of the continuous- biomarker linear-mixed coefficient, translating to a 30 to 50 percent power penalty.
2. *Sphericity.* Compound-symmetric within-subject covariance is empirically rejected for time-series symptom data: nearby timepoints are more correlated than distant ones. The AR(1) structure pmsimstats implements explicitly violates sphericity. After Greenhouse-Geisser correction, the test recovers nominal Type I error but loses degrees of freedom and power.
3. *No carryover model.* Strict RM-ANOVA has no provision for residual drug effect during the off-drug period; the model treats off-drug observations as if no drug had been administered. Under the pmsimstats DGP with $t_{1/2} = 1$ week, this misspecification biases $\hat{L}$ toward zero (attenuation by 30 to 35 percent in the production data, comparable to the E1 specification evaluated in the carryover-sensitivity manuscript).
4. *Balanced design required.* Equations (6) and (8) assume equal stratum sizes and equal numbers of within-phase timepoints. Real trials have missing data; classical RM-ANOVA handles missingness only through complete-case deletion. The MMRM framework's likelihood-based handling of missing-at-random (MAR) data²⁶ is genuinely better-behaved.
5. *Equal time spacing.* Strict RM-ANOVA treats time as a categorical factor with implicitly equal-spaced levels. Real N-of-1 trials have unequally-spaced visits; pmsimstats implements continuous-time AR(1) (`corCAR1`), which the strict-RM-ANOVA framework cannot accommodate.

These limitations motivate retaining the linear-mixed analysis with continuous-time AR(1) residual correlation as primary, and reporting the strict-RM-ANOVA closed-form F -statistic

as a sensitivity comparator. The five violations are not equally consequential: the categorical-biomarker and no-carryover assumptions are the most problematic in the pmsimstats setting, together accounting for a 30 to 50 percent attenuation of the interaction estimator relative to the linear-mixed analysis. The recommendation operationalized in §6 follows this strategy.

4 Cycle-by-period design grid

The cycle-by-period design grid that varies the number of treatment cycles, on-drug and off-drug period lengths, and on/off symmetry around the²⁷ hybrid reference is specified in `docs/27-cycle-period-design-sweep-plan.md` and constitutes the second axis of the joint design-by-test sensitivity question framed in §1. The empirical sweep across this grid is deferred to a follow-up paper. We shall note here, in the interest of transparency, that the present paper narrows its scope to the test-procedure axis at the prazosin/PTSD reference design (Hendrickson hybrid open-label plus blinded discontinuation, OL+BDC; $N \in \{25, 35, 70, 100\}$, $\beta_{bm} \in \{0, 0.45\}$), which the simulation infrastructure does support at production rep counts. The §6 recommendation is correspondingly scoped to test-procedure choice rather than the joint (cycle count, period length, test procedure) optimum that the design-by-test sensitivity question requires. The follow-up paper executing the cycle-by-period sweep is the natural extension and is recorded as a portfolio item.

5 Simulation design

The simulation program follows the ADEMP framework of¹. The primary estimand is the biomarker-treatment interaction inference under each of the five test procedures. Four operate on the continuous biomarker and share the same interaction estimand $\beta_{bm:D}$ on the symptom scale: a within-subject treatment-effect contrast regressed on the continuous biomarker (the dichotomization-free analogue of the two-sample t -test of equation (11)); the linear-mixed-effects Wald t on $\hat{\beta}_{bm:D}$; the GEE Wald z on $\hat{\beta}_{bm:D}$ with naive sandwich variance; and the GEE Wald z with Mancl-DeRouen small-sample sandwich variance per²⁴. The fifth, the strict classical RM-ANOVA F -statistic, operates on the median-dichotomized biomarker and so targets a dichotomized estimand on a different scale. The continuous-biomarker contrast is included specifically to separate the cost of dichotomization (strict RM-ANOVA versus the contrast, both pre-averaged to one value per phase) from the cost of the test procedure proper (the contrast versus the longitudinal linear-mixed and GEE analyses). Primary inferential outputs are power $\pi = \Pr(p < 0.05)$ for $H_0 : \beta_{bm:D} = 0$ at $\alpha = 0.05$, and Type I error under $\beta_{bm:D} = 0$. Secondary estimands, reported for the four coefficient-bearing arms, are bias of $\hat{\beta}_{bm:D}$ relative to a high-precision Monte-Carlo reference value of the estimand, the empirical and the mean model-based standard error, and nominal-95% CI coverage. Strict RM-ANOVA is excluded from these coefficient-based measures because its estimand is on the dichotomized scale.

The full Aims, Data-generating mechanisms, Estimands, Methods, and Performance measures are pre-registered at `analysis/scripts/test-procedure-design-sensitivity/00-ademp-pre-registration` which fixes the test-procedure times design-grid factorial across three studies (Study 1: test-

348 procedure comparison at fixed design; Study 2: cycle-by-period grid at fixed test, 240 cells;
 349 Study 3: joint design-by-test heatmap, 240 cells) before the production runs commit. The
 350 present paper executes a revised Study 1; the executed design departs from the pre-registered
 351 Study 1 on several factor levels, and we document every departure in Table 1 below in the
 352 interest of full ADEMP transparency.

Table 1: Pre-registered versus executed Study 1. Departures are recorded for transparency; the executed superset of sample sizes (which restores the pre-registered N levels and adds $N = 25$) and the added continuous-biomarker contrast arm respond to referee comments on the first internal draft.

Factor	Pre-registered	Executed (this paper)
Test procedures	3 (RM-ANOVA, LME, GEE-MD)	5 (adds naive GEE and the continuous-biomarker contrast)
Biomarker effect β_{bm}	$\{0, 0.3, 0.6\}$	$\{0, 0.45\}$ (the project-wide reference effect size)
Sample size N	$\{35, 70, 100\}$	$\{25, 35, 70, 100\}$ (superset; adds $N = 25$)
Replicates per cell	1000 (5000 for Type I cells)	1000 power, 2000 Type I
Carryover half-life	$t_{1/2} = 3$ days	$t_{1/2} = 0$ (no carryover at the fixed reference)

353 **Units.** Carryover half-lives $t_{1/2}$ are quoted in days in this table, following the pre-registration
 354 and the main-effect manuscripts of this series. The interaction-focused manuscripts quote $t_{1/2}$
 355 in weeks on the canonical grid $\{0, 0.5, 1.0\}$; the conversion is 1 week = 7 days. Timepoints t
 356 are weekly in both conventions.

357 In the scope of the present paper (test-procedure axis only at the reference design), the sim-
 358 ulation runs on the pmsimstats package data-generating process under the mean-moderation
 359 architecture (mean moderation), exercised through `generateData()` and the standard
 360 `lme_analysis()` fit, rather than on the reduced `implementations/simple` substrate. The
 361 factorial comprises $n_{\text{reps}} = 1000$ per power cell and 2000 per Type I cell, across the 40
 362 cells defined by $\text{procedure} \in \{\text{strict RM-ANOVA, continuous-biomarker contrast, linear-}$
 363 $\text{mixed with corCAR1, GEE with naive sandwich, GEE with Mancl-DeRouen small-sample}$
 364 $\text{correction}\} \times \beta_{bm} \in \{0, 0.45\} \times N \in \{25, 35, 70, 100\}$. Execution uses `parallel::mclapply`
 365 with `RNGkind("L'Ecuyer-CMRG")` and `mc.set.seed = TRUE`, so each replicate draws from
 366 an independent, reproducible random-number stream. The Mancl-DeRouen correction
 367 was implemented inline (no `geesmv`-style external dependency in the production path) by
 368 inflating each cluster's residual via $(I - H_{ii})^{-1} r_i$ in the sandwich meat. It produces identical
 369 point estimates and larger standard errors than the naive sandwich, and its SE for the
 370 interaction coefficient agrees with `geesmv::GEE.var.md` to within approximately 2.5% (mean
 371 inline/`geesmv` ratio 1.02, standard deviation 0.02, over 25 probe replicates at $N = 35$,
 372 $\beta_{bm} = 0.45$), so the correction magnitude is externally validated rather than merely plausible.

373 6 Results

374 The simulation program was executed across the 40 cells of the five-procedure-by-effect-by-
 375 N factorial in 587 s using 7-core `parallel::mclapply` with `RNGkind("L'Ecuyer-CMRG")`
 376 per-stream seeding. Convergence was 100% across all 60{,}000 fits. Per-replicate output
 377 is at `analysis/data/quick-sim/08-test-replicates.rds`; the Morris ADEMP cell-level
 378 summary is at `analysis/data/quick-sim/08-test-summary.txt`. With 1000 replicates per
 379 power cell the MCSE on a power estimate near 0.5 is approximately 0.016, and with 2000
 380 replicates per Type I cell the MCSE on a rate near 0.05 is approximately 0.005. We begin
 381 the results with the test-procedure comparison at the fixed reference design.

382 6.1 Test-procedure comparison at fixed design

383 We note at the outset that strict RM-ANOVA targets a different estimand from the other
 384 four procedures. It operates on a median-dichotomized biomarker, pre-averaged to one value
 385 per phase, and so estimates a dichotomized-stratum interaction; the continuous-biomarker
 386 contrast, the linear-mixed model, and both GEE variants estimate the interaction coefficient
 387 $\beta_{bm:D}$ on the continuous symptom scale. An earlier version of this comparison contrasted
 388 strict RM-ANOVA directly with the mixed model and attributed the resulting power gap to
 389 the test procedure; that attribution conflated two distinct costs. We therefore introduce the
 390 continuous-biomarker contrast of §4, which differs from strict RM-ANOVA only in retaining
 391 the continuous biomarker (it shares the pre-averaging and the absence of any AR(1) machin-
 392 ery) and differs from the linear-mixed model only in discarding the within-phase longitudinal
 393 structure. The RM-ANOVA $\rightarrow$ contrast step therefore measures the cost of dichotomization in
 394 isolation, and the contrast $\rightarrow$ linear-mixed step measures the cost of the test procedure proper.
 395 Reading the two steps separately is the honest decomposition the earlier direct comparison
 396 lacked.

397 Type I error at the null cell ($\beta_{bm} = 0$), percent $\pm$ MCSE, is shown in Table 2; power at
 398 $\beta_{bm} = 0.45$ in Table 3.

Table 2: Type I error rate (percent $\pm$ MCSE) under $H_0 : \beta_{bm:D} = 0$, 2000 replicates per cell. Nominal level 5%.

Procedure	$N = 25$	$N = 35$	$N = 70$	$N = 100$
Strict RM-ANOVA (dich.)	4.55 $\pm$ 0.47	4.15 $\pm$ 0.45	5.95 $\pm$ 0.53	5.20 $\pm$ 0.50
Contrast (continuous bm)	5.25 $\pm$ 0.50	4.50 $\pm$ 0.46	6.45 $\pm$ 0.55	4.60 $\pm$ 0.47
Linear-mixed (<code>corCAR1</code>)	2.75 $\pm$ 0.37	2.05 $\pm$ 0.32	3.35 $\pm$ 0.40	2.90 $\pm$ 0.38
GEE (naive sandwich)	9.70 $\pm$ 0.66	8.55 $\pm$ 0.63	7.90 $\pm$ 0.60	7.40 $\pm$ 0.59
GEE (Mancl-DeRouen)	6.35 $\pm$ 0.55	5.30 $\pm$ 0.50	6.25 $\pm$ 0.54	5.90 $\pm$ 0.53

Table 3: Power (percent $\pm$ MCSE) at $\beta_{bm} = 0.45$, 1000 replicates per cell.

Procedure	$N = 25$	$N = 35$	$N = 70$	$N = 100$
Strict RM-ANOVA (dich.)	35.8 ± 1.52	46.8 ± 1.58	75.3 ± 1.36	91.3 ± 0.89
Contrast (continuous bm)	53.8 ± 1.58	67.7 ± 1.48	91.5 ± 0.88	98.8 ± 0.34
Linear-mixed (corCAR1)	56.5 ± 1.57	71.2 ± 1.43	94.5 ± 0.72	99.2 ± 0.28
GEE (naive sandwich)	67.2 ± 1.48	75.9 ± 1.35	94.5 ± 0.72	99.2 ± 0.28
GEE (Mancl-DeRouen)	57.3 ± 1.56	69.1 ± 1.46	92.7 ± 0.82	99.1 ± 0.30

Four substantive findings emerge from the tables. First, and most consequentially for how the comparison should be read, the power gap between strict RM-ANOVA and the linear-mixed model is predominantly a cost of dichotomizing the biomarker rather than a property of the test procedure. At $N = 25$ the strict RM-ANOVA power is 0.358; the continuous-biomarker contrast, which differs only in retaining the continuous biomarker, reaches 0.538, so dichotomization alone accounts for 0.180 of lost power. The further step to the full linear-mixed model raises power only to 0.565, an increment of 0.027 attributable to the test procedure proper. The decomposition is stable across the grid: the dichotomization component is 0.209 at $N = 35$ ($0.468 \rightarrow 0.677$), 0.162 at $N = 70$ ($0.753 \rightarrow 0.915$), and 0.075 at $N = 100$ ($0.913 \rightarrow 0.988$), while the contrast-to-linear-mixed increment never exceeds 0.04. The earlier framing, which compared strict RM-ANOVA against the mixed model directly and read the difference as a test-procedure effect, therefore overstated the role of the test procedure by roughly a factor of five; the honest statement is that a classical, sphericity-free, AR(1)-free within-subject contrast on the continuous biomarker captures most of the available power, and the mixed model adds a small further increment.

Second, GEE with the naive sandwich is anti-conservative across the entire sample-size range examined. Its null rejection rate is 0.097 at $N = 25$ and declines only slowly with N to 0.086, 0.079, and 0.074 at $N = 35, 70, 100$ – still roughly $1.5\times$ nominal at $N = 100$. The inflation is therefore not a small-cluster curiosity that disappears at trial-relevant sizes; it persists. The Mancl-DeRouen 2001 small-sample sandwich correction brings the rate within or close to MCSE of nominal at every N (0.064, 0.053, 0.063, 0.059), at a power cost of about ten percentage points at $N = 25$ ($0.672 \rightarrow 0.573$) and seven at $N = 35$ ($0.759 \rightarrow 0.691$). The corrected procedure’s apparent power advantage over the linear-mixed model at small N (Table 3) is the residue of its remaining slight over-rejection, not a genuine efficiency gain.

Third, the secondary estimands (Table 4) locate the procedures’ behavior in the standard-error calibration that the rejection rates only summarize. At $\beta_{bm} = 0.45$, the continuous-biomarker contrast and the Mancl-DeRouen GEE are well-calibrated: their nominal-95% CI coverage is 93–95% and their mean model-based SE is within a few percent of the empirical SE. Naive GEE under-covers at small N (90.0% at $N = 25$, 91.5% at $N = 35$) because its model-based

SE is 9 to 12% smaller than the empirical SE, the calibration counterpart of its Type I inflation. The linear-mixed model over-covers (97.0 to 97.7%) because its model-based SE exceeds the empirical SE by 11 to 16%; this is the same conservatism visible in its sub-nominal Type I error.

Table 4: Secondary estimands at $\beta_{bm} = 0.45$ for the four coefficient-bearing arms (the interaction estimand is $\beta_{bm:D} = -2.21$, the Monte-Carlo reference value). Coverage is of the nominal-95% Wald interval; SE ratio is mean model-based SE divided by empirical SE.

Arm	Coverage $N = 25$	Coverage $N = 100$	SE ratio $N = 25$	SE ratio $N = 100$	Max rel. bias
Contrast (continuous bm)	94.9	94.8	0.99	0.98	3.0%
Linear-mixed (corCAR1)	97.7	97.0	1.16	1.11	2.6%
GEE (naive sandwich)	90.0	93.3	0.91	0.98	3.8%
GEE (Mancl-DeRouen)	94.3	94.0	1.07	1.02	3.8%

Fourth, the comparison is not contaminated by point-estimate bias. The relative bias of $\hat{\beta}_{bm:D}$ is below 4% of the estimand $|\beta_{bm:D}| = 2.21$ for the contrast, the linear-mixed model, and both GEE variants in every cell, so the differences in power and coverage reflect variance estimation and information use rather than systematic mis-estimation. The linear-mixed procedure's Type I error sits below nominal throughout (0.021 to 0.034), the conservatism already visible in its over-coverage and its model/empirical SE ratio above one; it is consistent with corCAR1 degrees-of-freedom accounting under the short ($T = 8$) serial series. A Kenward-Roger²⁸ or Satterthwaite²⁹ small-sample degrees-of-freedom correction would address the conservatism and recover the small power increment it forgoes; this extension is recorded as a follow-up item.

7 Design-and-test recommendation

The recommendation table below maps test procedure choice to typical chronic-condition N-of-1 trial scenarios at the prazosin/PTSD reference design (Hendrickson hybrid OL+BDC, biomarker-treatment interaction estimand). The cycle count and period length are held at the reference values defined in §3 because the joint optimization over those axes is deferred to the follow-up paper. We present the recommendation as a table of preferred (procedure, scenario) combinations rather than as a single ‘best’ procedure, because the preferred choice depends on features of the application that the analyst must determine. The overriding recommendation, however, is design-agnostic: retain the biomarker as a continuous predictor, since dichotomization forfeits the large majority of the available power (§5, Table 3).

Scenario	Preferred procedure	Rationale
Default (continuous biomarker, any N)	linear-mixed (<code>corCAR1</code>)	Highest power within nominal Type I; <code>corCAR1</code> matches AR(1) physiology; well-behaved across N
Lightweight / transparent classical analysis	continuous-biomarker contrast	Within 0.03 of the mixed model on power, nominal Type I, no covariance modeling required
$N \leq 35$, concern about LME conservatism	linear-mixed (<code>corCAR1</code>) with Kenward-Roger	LME sub-nominal and over-covers; KR addresses this without inflating Type I
Robust-SE-required setting (e.g. unmodeled clustering)	GEE Mancl-DeRouen	Restores nominal Type I and 93–95% coverage; modest power cost
Reviewer-mandated ‘classical’ analysis	strict RM-ANOVA	Dichotomized sensitivity comparator; report alongside a continuous-biomarker analysis with the power-loss decomposition
Naive GEE (no small-sample correction)	not recommended	Type I error 1.5–1.9 $\times$ nominal across all N ; under-covers

452 The dominant message is that the biomarker should be analyzed as a continuous predic-
 453 tor: dichotomizing it for a strict RM-ANOVA forfeits the large majority of the power gap
 454 that the present study isolates, while modeling choices among the continuous-biomarker
 455 procedures move power only at the margin. Within that class, the linear-mixed analysis
 456 with continuous-time AR(1) residual correlation is the preferred default for the biomarker-
 457 treatment interaction in aggregated N-of-1 trials at the parameter scale evaluated here, with
 458 the dichotomization-free within-subject contrast a close and transparent alternative that re-
 459 quires no covariance modeling. Strict RM-ANOVA remains useful as a dichotomized sensi-
 460 tivity comparator, as the §2 closed-form derivation makes explicit, but it should be reported
 461 alongside a continuous-biomarker analysis and the power-loss decomposition rather than on
 462 its own. Naive GEE should not be used at the sample sizes examined without a small-sample
 463 correction; the Mancl-DeRouen correction restores nominal Type I error and 93–95% coverage
 464 at a moderate power cost.

465 8 Discussion

466 8.1 What the Monte Carlo study settles

467 We begin the discussion by summarizing what the Monte Carlo study settles at the resolution
 468 it was run. The run on the 40-cell five-procedure factorial (procedure $\in$ {strict RM-ANOVA,
 469 continuous-biomarker contrast, linear-mixed with `corCAR1`, GEE with naive sandwich, GEE
 470 with Mancl-DeRouen small-sample correction} $\times \beta_{bm} \in \{0, 0.45\} \times N \in \{25, 35, 70, 100\}$,

with 1000 replicates per power cell and 2000 per Type I cell) was executed in 587 s with 100% convergence across all 60,000 fits. Per-replicate output is at `analysis/data/quick-sim/08-test-replicates.rds`; Morris ADEMP summary at `08-test-summary.txt`. With 1000 power replicates the MCSE on a power estimate near 0.5 is approximately 0.016 and with 2000 Type I replicates the MCSE on a rate near 0.05 is approximately 0.005. The Mancl-DeRouen 2001²⁴ correction was implemented inline by inflating each cluster's residual via $(I - H_{ii})^{-1}r_i$ in the sandwich meat; it produces identical point estimates and larger standard errors than the naive sandwich, and its SE for the interaction coefficient agrees with `geesmv::GEE.var.md` to within approximately 2.5% (mean inline/`geesmv` ratio 1.02 over 25 probe replicates), so the correction magnitude is externally validated.

9 Conclusions

In conclusion, three points are to be emphasized. First, the analyst's most consequential choice for the biomarker-treatment interaction in aggregated N-of-1 trials is to keep the biomarker continuous rather than to dichotomize it for a classical RM-ANOVA. The continuous-biomarker contrast recovers about four-fifths of the apparent RM-ANOVA-to-mixed-model power gap (dichotomization costs 0.08 to 0.21 of power across the sample sizes evaluated, while the test procedure proper adds at most 0.04); the linear-mixed analysis with continuous-time AR(1) residual correlation has the highest power within nominal Type I among the continuous-biomarker procedures, with the dichotomization-free within-subject contrast a close alternative. Naive GEE with the standard sandwich estimator inflates Type I error to between $1.5\times$ and $1.9\times$ nominal across the whole $N \in \{25, 35, 70, 100\}$ range, disqualifying its rejection rate as a power summary; the Mancl-DeRouen 2001 small-sample sandwich correction restores nominal Type I control and 93–95% coverage at a moderate (about 10-percentage-point) power cost, making it the operational GEE comparator rather than the naive sandwich.

Second, the closed-form strict-RM-ANOVA test of the biomarker-treatment interaction (§2) reduces to a two-sample t -test on the within-subject treatment differences under sphericity and balance. The reduction makes the strict-RM-ANOVA result analytically transparent and connects the simulation findings to the classical split-plot literature. It also pins down what does and does not limit the test here. The strict RM-ANOVA arm we evaluate pre-averages each phase to a single value, so its within-subject covariance is 2×2 and sphericity holds by construction; consistent with this, its Type I error is at nominal (Table 2), and its power deficit is wholly attributable to the median dichotomization of the biomarker, as the continuous-biomarker contrast demonstrates. The sphericity penalty that motivates much of §2 applies instead to the alternative classical analysis that treats the full $T = 8$ series as repeated within-subject levels: there the AR(1) residual correlation violates sphericity and a Greenhouse-Geisser correction recovers nominal Type I only at a degrees-of-freedom cost, which the linear-mixed analysis with `corCAR1` avoids by modeling the AR(1) structure directly. Both considerations point the same way – away from a dichotomized classical analysis – but through different mechanisms, and we are careful to keep them distinct.

512 Third, the cycle-by-period design grid that constitutes the second axis of the joint design-by-
513 test sensitivity question framed in §1 is deferred to a follow-up paper. The infrastructure for
514 the sweep is in place at `analysis/scripts/test-procedure-design-sensitivity/`, and the
515 design specification is committed at `docs/27-cycle-period-design-sweep-plan.md`. The
516 follow-up paper will deliver the joint (cycle count, period length, test procedure) recommen-
517 dation table that the present paper's §6 narrows to the test-procedure-only axis.

518 10 Conflict of interest

519 The authors declare no conflicts of interest.

520 11 Funding

521 This work received no specific funding.

522 12 Data availability statement

523 The data and code that support the findings of this study are openly available in the
524 `pmsimstats-ng` project repository. Per-replicate Monte Carlo outputs and ADEMP-
525 compliant cell-level summaries are versioned under `analysis/data/`; simulation drivers
526 are at `analysis/scripts/`. Exact paths relevant to this manuscript are listed in the
527 Reproducibility section below.

528 13 Reproducibility

529 This manuscript's simulation was executed with R 4.6.x; the project also ships a `zzcollab`
530 Docker container (image hash recorded in `Dockerfile`; package set captured in `renv.lock`)
531 for container-based reproduction. Principal packages used by the Study 1 pipeline are
532 `nlme` (continuous-time linear-mixed analysis with `corCAR1` residual correlation), `geepack`
533 (GEE fitting) plus an inline Mancl-DeRouen 2001 small-sample sandwich correction, `geesmv`
534 (independent benchmark of that correction), `data.table`, and `parallel` (`mclapply` for
535 the parameter-grid sweep), with `rmarkdown` plus `bookdown` for the manuscript build. The
536 Study 1 simulation runs on the `pmsimstats` package data-generating process under the mean-
537 moderation architecture via `generateData()` and `lme_analysis()`, not on the reduced
538 `implementations/simple` substrate.

539 **Random-number generator.** The driver sets `RNGkind("L'Ecuyer-CMRG")` and
540 `set.seed(20260507)`, then runs the cell-by-replicate jobs through `parallel::mclapply`
541 with `mc.set.seed = TRUE`, so each replicate draws from an independent, reproducible
542 random-number stream.

543 **Data and code availability.** Per-replicate simulation outputs are checkpointed at
544 `analysis/data/quick-sim/08-test-replicates.rds`; the extended Morris ADEMP
545 cell-level summary (rejection rate, MCSE, convergence, bias, empirical and model

SE, coverage) at the companion `08-test-summary.txt`, with the Monte-Carlo reference estimand at `08-test-betatruer.txt` and the Mancl-DeRouen benchmark at `08-test-mancl-benchmark.txt`. The driver is `analysis/scripts/test-procedure-design-sensitivity/01` (superseding the earlier `analysis/scripts/quick-sim/08-test-procedure-prototype.R`). The pre-registered ADEMP plan, and the deviations recorded in Table 1, are at `analysis/scripts/test-procedure-design-sensitivity/00-ademp-pre-registration.md`. The manuscript source, paper-specific bibliography (`references.bib`), and SIM CSL file are at `analysis/report/08-test-procedure-design-sensitivity/`.

14 References

1. Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Statistics in Medicine*. 2019;38(11):2074-2102.
2. Haverkamp N, Beauducel A. Violation of the sphericity assumption and its effect on type-I error rates in repeated measures ANOVA and multi-level linear models. *Frontiers in Psychology*. 2017;8:1841. doi:10.3389/fpsyg.2017.01841
3. Huynh H, Feldt LS. Estimation of the Box correction for degrees of freedom from sample data in randomized block and split-plot designs. *Journal of Educational Statistics*. 1976;1(1):69-82. doi:10.3102/10769986001001069
4. Greenhouse SW, Geisser S. On methods in the analysis of profile data. *Psychometrika*. 1959;24(2):95-112. doi:10.1007/BF02289823
5. Blanca MJ, Arnau J, García-Castro FJ, Alarcón R, Bono R. Repeated measures ANOVA and adjusted tests when sphericity is violated: Which procedure is best? *Frontiers in Psychology*. 2023;14:1192453. doi:10.3389/fpsyg.2023.1192453
6. Natesan Batley P, McClure EB, Brewer B, et al. Evidence and reporting standards in N-of-1 medical studies: A systematic review. *Translational Psychiatry*. 2023;13(1):263. doi:10.1038/s41398-023-02562-8
7. Stunnenberg BC, Berends J, Griggs RC, et al. N-of-1 trials in neurology: A systematic review. *Neurology*. 2022;98(2):e174-e185. doi:10.1212/WNL.00000000000012998
8. Shaffer JA, Kronish IM, Falzon L, Cheung YK, Davidson KW. N-of-1 randomized intervention trials in health psychology: A systematic review and methodology critique. *Annals of Behavioral Medicine*. 2018;52(9):731-742. doi:10.1093/abm/kax026
9. Vohra S, Shamseer L, Sampson M, et al. CONSORT extension for reporting N-of-1 trials (CENT) 2015 statement. *Journal of Clinical Epidemiology*. 2015;76:9-17. doi:10.1016/j.jclinepi.2015.05.004

- 564 10. Shamseer L, Sampson M, Bukutu C, Schmid CH, et al. CENT 2015: Explanation and elaboration. *Journal of Clinical Epidemiology*. 2015;76:18-46. doi:10.1016/j.jclinepi.2015.05.018
- 565 11. Kwasnicka D, Inauen J, Nieuwenboom W, et al. Challenges and solutions for N-of-1 design studies in health psychology. *Health Psychology Review*. 2019;13(2):163-178. doi:10.1080/17437199.2018.1564627
- 566 12. Ma Y, Mazumdar M, Memtsoudis SG. Beyond repeated-measures analysis of variance: Advanced statistical methods for the analysis of longitudinal data in anesthesia research. *Regional Anesthesia and Pain Medicine*. 2012;37(1):99-105. doi:10.1097/AAP.0b013e31823ebc74
- 567 13. Mallinckrodt CH, Clark SWS, Carroll RJ, Molenberghs G. Assessing response profiles from incomplete longitudinal clinical trial data under regulatory considerations. *Journal of Biopharmaceutical Statistics*. 2003;13(2):179-190. doi:10.1081/BIP-120019265
- 568 14. Mallinckrodt CH, Clark WS, David SR. Accounting for dropout bias using mixed-effects models. *Journal of Biopharmaceutical Statistics*. 2001;11(1-2):9-21. doi:10.1081/BIP-100104194
- 569 15. Mallinckrodt CH, Raskin J, Wohlreich MM, Watkin JG, Detke MJ. The efficacy of duloxetine: A comprehensive summary of results from MMRM and LOCF_ANCOVA in eight clinical trials. *BMC Psychiatry*. 2004;4:26. doi:10.1186/1471-244X-4-26
- 570 16. Prakash A, Risser RC, Mallinckrodt CH. The impact of analytic method on interpretation of outcomes in longitudinal clinical trials. *International Journal of Clinical Practice*. 2008;62(8):1147-1158. doi:10.1111/j.1742-1241.2008.01808.x
- 571 17. Liang KY, Zeger SL. Longitudinal data analysis using generalized linear models. *Biometrika*. 1986;73(1):13-22. doi:10.1093/biomet/73.1.13
- 572 18. Wang M, Kong L, Li Z, Zhang L. Covariance estimators for generalized estimating equations (GEE) in longitudinal analysis with small samples. *Statistics in Medicine*. 2016;35(10):1706-1721. doi:10.1002/sim.6817
- 573 19. Pan W. Sample size and power calculations with correlated binary data. *Controlled Clinical Trials*. 2001;22(3):211-227. doi:10.1016/s0197-2456(01)00131-3
- 574 20. Tong G, Nevins P, Ryan M, others, Taljaard M. A review of current practice in the design and analysis of extremely small stepped-wedge cluster randomized trials. *Clinical Trials*. 2024;22(1):45-56. doi:10.1177/17407745241276137

- 575 21. Defelippe VM, Thiel GJMW van, Otte WM, Schutgens REG, Stunnenberg B, et al. Toward responsible clinical n-of-1 strategies for rare diseases. *Drug Discovery Today*. 2023;28(10):103688. doi:10.1016/j.drudis.2023.103688
- 576 22. Senn S. The AB/BA crossover: Past, present and future? *Statistical Methods in Medical Research*. 1994;3(4):303-324. doi:10.1177/096228029400300402
- 577 23. Senn S. *Cross-over Trials in Clinical Research*. 2nd ed. Wiley; 2002.
- 578 24. Mancl LA, DeRouen TA. A covariance estimator for GEE with improved small-sample properties. *Biometrics*. 2001;57(1):126-134.
- 579 25. Maxwell SE, Delaney HD, Kelley K. *Designing Experiments and Analyzing Data: A Model Comparison Perspective*. 3rd ed. Routledge; 2018.
- 580 26. Mallinckrodt CH, Lane PW, Schnell D, Peng Y, Mancuso JP. Recommendations for the primary analysis of continuous endpoints in longitudinal clinical trials. *Drug Information Journal*. 2008;42(4):303-319. doi:10.1177/009286150804200402
- 581 27. Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: Statistical methods and practical considerations. *Frontiers in Digital Health*. 2020;2:13. doi:10.3389/fdgth.2020.00013
- 582 28. Kenward MG, Roger JH. Small sample inference for fixed effects from restricted maximum likelihood. *Biometrics*. 1997;53(3):983-997. doi:10.2307/2533558
- 583 29. Satterthwaite FE. An approximate distribution of estimates of variance components. *Biometrics Bulletin*. 1946;2(6):110-114. doi:10.2307/3002019