

1 A combined data-generating architecture for
2 biomarker-treatment interactions: channel super-additivity and
3 mixing-weight sensitivity in aggregated N-of-1 trials

4 pmsimstats team

5 August 15, 2026

6 **Contents**

7	1 Abstract	1
8	2 1. Introduction	2
9	3 2. The Combined Architecture	3
10	3.1 2.1 Definition	3
11	3.2 2.2 Boundary reductions and identifiability	3
12	3.3 2.3 Biological rationale	4
13	4 3. Channel Super-additivity: an Analytical Account	5
14	5 4. Simulation Study	6
15	5.1 4.1 Design	6
16	5.2 4.2 Boundary validation against the companion baselines	6
17	5.3 4.3 Empirical power over the channel-weight grid	6
18	6 5. When a Combined Architecture is Appropriate	7
19	7 6. Discussion	8
20	8 Conflict of interest	9
21	9 Funding	9
22	10 Reproducibility	9
23	11 References	9

24 **1 Abstract**

25 **Background.** A companion paper¹ distinguishes two data-generating architectures for
26 biomarker-treatment interactions in aggregated N-of-1 trials: direct mean moderation (the
27 mean-moderation architecture), in which the biomarker scales the treatment effect in the

conditional mean, and differential correlation (the covariance-moderation architecture, the multivariate-normal approach), in which the interaction is encoded as treatment-state-dependent correlation in the joint distribution. That paper treats A and B as alternatives. Here we consider the case in which both channels are active at once.

Methods. We define the dual-channel architecture, which carries independent parameters $c_{bm,A}$ (mean channel) and $c_{bm,B}$ (covariance channel), and which reduces algebraically to the pure mean-moderation architecture when $c_{bm,B} = 0$ and to pure the covariance-moderation architecture when $c_{bm,A} = 0$. We show analytically that the biomarker-treatment signal delivered to the fixed analysis model is additive across the two channels, and that, because the Monte-Carlo power function is convex in the noncentrality over the sub-ceiling regime, the two channels combine super-additively in power. We then quantify the combined architecture by Monte Carlo simulation on a 3×3 grid of channel weights crossed with three within-subject designs and three carryover half-lives.

Results. [To be finalized after the common-driver boundary re-run; see Section 4.2.] Empirical power increases monotonically in each channel weight, and interior cells exceed the probability-scale additive benchmark, confirming the analytical prediction of super-additivity. The covariance and mean channels contribute comparably to carryover sensitivity, and the design ranking established for the pure architectures is preserved throughout the combined grid.

Conclusions. A combined architecture is the appropriate data-generating model when a biomarker is hypothesized to act through both a pharmacodynamic-scaling pathway and a co-regulatory coupling pathway. It should not be a default: it carries a second free parameter and requires independent evidence for each channel. Its principal use is as a sensitivity framework spanning the two pure-architecture extremes.

52

2 1. Introduction

To be completed by rgt.

A companion paper¹ formalizes two data-generating process (DGP) architectures for simulation studies of predictive biomarker-treatment interactions in aggregated N-of-1 trials, and shows that the choice between them materially changes the estimated power loss under carryover. Under the mean-moderation architecture the biomarker scales the treatment effect in the conditional mean; under the covariance-moderation architecture the interaction is carried by treatment-state-dependent correlation between biomarker and response in the joint distribution. That paper treats the two as competing choices. The present paper takes up the case, deferred there, in which both mechanisms operate at once. We call the resulting data-generating model the dual-channel architecture. It is not a third

alternative on the same footing as A and B; it is their common generalization, with A and B recovered exactly as its single-channel limits.

To be completed by rgt.

The paper makes two contributions. First, we give an analytical account (Section~3) of why the two channels, when active together, combine super-additively in power rather than merely adding. The argument identifies the additivity of the biomarker-treatment signal delivered to the analysis model as the mechanism, and the convexity of the Monte-Carlo power function over the sub-ceiling regime as the reason additivity of signal becomes super-additivity of power. This addresses a gap noted for the descriptive combined results reported in the companion paper, where the super-additivity was observed but not explained. Second, we characterize the combined architecture empirically over a 3×3 grid of channel weights crossed with three within-subject designs and three carryover half-lives, with the boundary cells validated numerically against the pure-architecture baselines of the companion paper.

3 2. The Combined Architecture

3.1 2.1 Definition

To be completed by rgt.

the dual-channel architecture applies both interaction mechanisms in sequence. First, the multivariate-normal covariance matrix is constructed with biomarker-to-biological-response (BM-BR) differential correlation weighted by $c_{bm,B}$, exactly as under the covariance-moderation architecture: the on-drug correlation is $c_{bm,B}$, decaying to $c_{bm,B} \cdot e^{-\lambda t_{sd}}$ during off-drug periods, where t_{sd} is time since discontinuation. Second, after the MVN draw, the BR component is additively shifted at on-drug timepoints by

$$\Delta BR_{it} = c_{bm,A} \cdot z_{B,i} \cdot \sigma_{BR} \cdot \mathbf{1}[\text{on-drug}],$$

where $z_{B,i} = (B_i - \mu_B) / \sigma_B$ is the standardized biomarker for participant i , exactly as under the mean-moderation architecture.

3.2 2.2 Boundary reductions and identifiability

To be completed by rgt.

The two parameters $c_{bm,A}$ and $c_{bm,B}$ are independent. Setting $c_{bm,B} = 0$ removes the covariance channel and recovers the pure mean-moderation architecture; setting $c_{bm,A} = 0$ removes the mean-shift channel and recovers the

pure covariance-moderation architecture. This combined form is the mean-and-covariance-structure model of Arminger et al.², the general psychometric parameterization within which the mean-moderation and covariance-moderation architectures are constrained submodels; a companion paper³ situates it as one node of a biomarker-moderation spectrum bridging continuous-heterogeneity and latent-class encodings. The boundary cases are enforced algebraically rather than by convention, so they serve as internal validation gates for any simulation study using the dual-channel architecture: a correctly implemented combined driver must reproduce the pure-architecture power values at the boundaries when run under matched design, sample-size, and calibration settings. Section~4.2 uses exactly this property to validate the simulation against the companion paper's baselines.

We note that when both parameters are set to their full values ($c_{bm,A} = c_{bm,B} = 0.45$), the total interaction effect on the BR mean is approximately double that of either pure architecture at $c_{bm} = 0.45$. A researcher wishing to hold total effect size constant across architectures may instead parameterize as $c_{bm,A} = (1 - \alpha) \cdot c$ and $c_{bm,B} = \alpha \cdot c$ for a mixing weight $\alpha \in [0, 1]$ and total weight c . We do not adopt this constraint in the simulations below, preferring the factorial parameterization that spans the full $(c_{bm,A}, c_{bm,B})$ space.

3.3 2.3 Biological rationale

To be completed by rgt.

The two channels are mechanistically distinct and need not scale in proportion. The mean-shift channel (the mean-moderation architecture) reflects a pharmacodynamic relationship: the drug's biological effect on BR is amplified in participants whose biomarker is above the population mean, because the biomarker governs the effective dose or receptor occupancy. This mechanism operates at the level of the mean trajectory for each participant. The covariance channel (the covariance-moderation architecture) reflects a different process: on-drug periods impose a co-regulatory coupling between biomarker expression and treatment response, perhaps because the drug activates a pathway whose output is gated by the biomarker signal. This coupling is absent during off-drug periods, producing the differential correlation structure that the covariance-moderation architecture encodes in the covariance matrix. A given drug could activate one, both, or neither channel. Cast in latent-class terms, both channels carry information about an unobserved responder class: the mean channel through the class-conditional expectation $E[BR | B]$ and the covariance channel through the class-conditional dispersion $\text{Var}[BR | B]$ ³. The linear-mixed analogue, in which the heterogeneity resides in the random-effects population, is developed by Verbeke and Lesaffre⁴.

4 3. Channel Super-additivity: an Analytical Account

To be completed by rgt.

The fixed analysis model detects the interaction through the coefficient on the biomarker-by-exposure term, whose estimate is driven by the on-drug covariance between the biomarker B and the biological response BR . Consider that covariance under each channel. Under the covariance channel alone (the covariance-moderation architecture), the on-drug joint distribution sets

$$\text{Cov}(B, BR \mid \text{on-drug}) = c_{bm,B} \sigma_B \sigma_{BR}.$$

Under the mean channel alone (the mean-moderation architecture), the on-drug shift $\Delta BR = c_{bm,A} z_B \sigma_{BR} = c_{bm,A} (B - \mu_B) \sigma_{BR} / \sigma_B$ contributes

$$\text{Cov}(B, \Delta BR) = \frac{c_{bm,A} \sigma_{BR}}{\sigma_B} \text{Var}(B) = c_{bm,A} \sigma_B \sigma_{BR}.$$

Under the combined architecture the on-drug response is the sum of the correlated draw and the mean shift, so the two covariance contributions add:

$$\text{Cov}(B, BR_C \mid \text{on-drug}) = (c_{bm,A} + c_{bm,B}) \sigma_B \sigma_{BR}.$$

The biomarker-treatment signal delivered to the analysis model is therefore exactly additive in the two channel weights. To leading order the residual variance is unchanged by the channels (the covariance channel preserves the marginal BR variance, and the mean channel inflates it only by the second-order term $c_{bm,A}^2 \sigma_{BR}^2$), so the noncentrality parameter of the interaction test is, to first order, additive across channels: $\eta_C \approx \eta_A + \eta_B$.

To be completed by rgt.

Additivity of the noncentrality does not imply additivity of power, because power is a nonlinear function of the noncentrality. Write $\pi(\eta)$ for the rejection probability of the interaction test as a function of its noncentrality $\eta \geq 0$. On the range over which π is below approximately one half, π is convex in η : each additional unit of noncentrality buys more power than the last, until the function inflects and saturates toward one. For η_A, η_B in this convex regime, convexity gives

$$\pi(\eta_A + \eta_B) - \pi(0) > [\pi(\eta_A) - \pi(0)] + [\pi(\eta_B) - \pi(0)],$$

which is precisely super-additivity relative to the probability-scale additive benchmark $p(c_{bm,A}, 0) + p(0, c_{bm,B}) - p(0, 0)$. The combined power exceeds the benchmark exactly when the two single-channel cells sit in the convex, sub-ceiling part of the power curve, and the excess shrinks, and can reverse to sub-additivity, once cells approach the ceiling and enter the concave regime. The empirical grid of Section~4 lies in the sub-ceiling regime (single-channel powers between roughly 0.03 and 0.30), so super-additivity is the predicted and observed pattern there.

174 5 4. Simulation Study

175 5.1 4.1 Design

176 *To be completed by rgt.*

177 We ran a factorial simulation on the dual-channel architecture parameter grid.
 178 The design factors are the crossover (CO), Hybrid, and open-label-with-blinded-
 179 discontinuation (OL+BDC) designs; $t_{1/2} \in \{0, 0.5, 1.0\}$ weeks; 1000 replicates
 180 per cell; and the E2 analysis specification. The channel-weight panel is $c_{bm,A} \times$
 181 $c_{bm,B} \in \{0, 0.22, 0.45\}^2$, nine cells per design-carryover combination. The analysis
 182 model is held fixed throughout and is identical to that of the companion paper:
 183 an `nlme::lme` fit with `corCAR1` continuous-time residual correlation and a single
 184 biomarker-by-exposure interaction term.

185 5.2 4.2 Boundary validation against the companion baselines

186 *To be completed by rgt.*

187 **Status: pending re-run.** The combined-architecture results must be produced
 188 by a driver whose sample-size convention and Gompertz and biomarker calibra-
 189 tion match the companion paper's production driver exactly, so that the boundary
 190 cells reproduce the pure-architecture power values numerically and not merely al-
 191 gebraically. An earlier combined driver used a different sample-size convention
 192 ($N = 35$ as a total rather than per randomization path), which made its boundary
 193 cells uniformly lower than, and not directly comparable to, the companion base-
 194 lines. This section will report a boundary-validation table placing the combined
 195 driver's $c_{bm,B} = 0$ column and $c_{bm,A} = 0$ row alongside the companion paper's the
 196 mean-moderation architecture and the covariance-moderation architecture base-
 197 lines, and will confirm agreement to within Monte Carlo standard error (approx-
 198 imately 0.013 to 0.016 at these cell sizes). The results in Section~4.3 are reported
 199 at the total-sample convention of the earlier driver and will be re-stamped once
 200 the matched re-run is complete.

201 5.3 4.3 Empirical power over the channel-weight grid

202 *To be completed by rgt.*

203 Table~1 reports empirical power across the full 3×3 channel-weight panel at the
 204 most demanding carryover condition ($t_{1/2} = 1.0$ week). Three features merit com-
 205 ment. First, the monotone structure is exact: power increases strictly with each
 206 channel weight, holding the other fixed, across all three designs, confirming that

Table 1: the dual-channel architecture empirical power: 3×3 channel-weight panel at carryover $t_{1/2} = 1.0$ week, E2 analysis specification, 1000 replicates per cell. Rows index $c_{bm,A}$ (mean channel); columns index $c_{bm,B}$ (covariance channel). Boundary column $c_{bm,B} = 0$ corresponds to the pure mean-moderation architecture; boundary row $c_{bm,A} = 0$ corresponds to the pure covariance-moderation architecture. Values are at the earlier driver’s total-sample convention and are pending the matched re-run of Section 4.2; Monte Carlo standard error columns will be added at that time.

$c_{bm,A}$ (mean channel)	$c_{bm,B}$ (covariance channel)		
	0.00	0.22	0.45
<i>Crossover (CO)</i>			
0.00	0.046	0.064	0.206
0.22	0.070	0.179	0.361
0.45	0.172	0.354	0.528
<i>Hybrid</i>			
0.00	0.034	0.083	0.273
0.22	0.108	0.299	0.621
0.45	0.389	0.642	0.870
<i>OL+BDC</i>			
0.00	0.017	0.049	0.148
0.22	0.086	0.164	0.362
0.45	0.276	0.473	0.688

the channels contribute independently to detectability. Second, interior cells exceed the probability-scale additive benchmark. At $(c_{bm,A} = 0.22, c_{bm,B} = 0.22)$ in the Hybrid design, power is 0.299; the additive baseline $p(0.22, 0) + p(0, 0.22) - p(0, 0) = 0.108 + 0.083 - 0.034 = 0.157$ is exceeded by 0.142, the super-additivity predicted in Section~3. Third, the covariance and mean channels contribute comparably to carryover sensitivity, and design choice remains consequential: at the full-signal, one-week-carryover cell, Hybrid (0.870) substantially outperforms OL+BDC (0.688) and CO (0.528).

6 5. When a Combined Architecture is Appropriate

To be completed by rgt.

the dual-channel architecture is appropriate when there is independent evidence that both the pharmacodynamic-scaling pathway and the co-regulatory coupling pathway are biologically active for the same biomarker-drug pair. Such evidence might include mechanistic assay data (for example, the drug enhances biomarker-outcome correlation in vitro, consistent with the B channel, while receptor occupancy scales with biomarker level, consistent with the A channel) or replication

224 across independent trials individually underpowered to distinguish the channels.
225 The dual-channel architecture should not be the default choice for simulation
226 studies despite its generality. It carries two free parameters where either pure
227 architecture requires one, and calibrating both independently demands more prior
228 information than is typically available at the trial-design stage. This estimation
229 caution is not a claim that the combined mechanism is biologically unusual: a
230 companion paper³ argues that the combined mean-and-covariance form is plausi-
231 bly the most biologically faithful regime for biomarker-guided psychiatric trials,
232 because a biomarker that indexes a responder subtype will in general shift both
233 the class-conditional mean and the class-conditional dispersion of response. The
234 recommendation against the dual-channel architecture as a default is therefore
235 about identifiability and prior information at the design stage, not about bio-
236 logical plausibility: the most faithful generative model is not always the most
237 estimable one, and when independent evidence on the two channels is unavailable
238 the single-parameter pure architectures remain the more disciplined default. Its
239 appropriate use cases are formal sensitivity analyses that ask how power estimates
240 change across the $(c_{bm,A}, c_{bm,B})$ grid when the channel structure is genuinely un-
241 certain, and evaluation studies designed to determine which channel is dominant
242 for a given biomarker-drug pair.

243 In the prazosin-PTSD context, the dual-channel architecture captures the possibil-
244 ity that systolic blood pressure simultaneously governs the expected prazosin effect
245 at the mean level (consistent with the pharmacological argument that noradren-
246 ergic tone determines the therapeutic window) and alters the coupling between
247 blood pressure and symptom change during active drug exposure (consistent with
248 a signaling-pathway interpretation). We regard the dual-channel architecture as
249 the appropriate framework for the prazosin-blood-pressure interaction if and when
250 mechanistic data become available to calibrate both channels separately.

251

252 7 6. Discussion

253 *To be completed by rgt.*

254 The combined architecture places the two architectures of the companion paper
255 within a single model, with each recovered as a single-channel limit. This reframes
256 the A-versus-B contrast as a question about which channel weight is nonzero
257 rather than a choice between unrelated data-generating conventions. The super-
258 additivity result of Section~3 explains, rather than merely reports, the compound-
259 ing of the two channels: it follows from the additivity of the biomarker-treatment
260 signal and the convexity of the power function, and is therefore a property of the
261 estimand and the test, not of the particular calibration. As in the companion

paper, the analysis model is held fixed; the conservatism of the standard `corCAR1` Wald test, and the cluster-robust remedy for it, are treated in the companion calibration study⁵ and apply here as a common-mode effect across the channel-weight grid. The combined form itself is not new to this series: a companion paper³ introduces it conceptually as the general node of a biomarker-moderation spectrum and supplies its psychometric lineage^{2,4}. The contribution of the present paper is complementary and specific: the analytical super-additivity result of Section~3 and the calibrated channel-weight simulation of Section~4, neither of which the conceptual treatment provides.

271

272 8 Conflict of interest

273 The pmsimstats team declares no conflicts of interest.

274 9 Funding

275 This work received no specific funding.

276 10 Reproducibility

277 This manuscript was rendered with R~4.6.0 inside the project's zzzcollab Docker container
278 (`rocker/tidyverse:4.6.0`; package set captured in `renv.lock`). The combined-architecture
279 driver is at `analysis/scripts/dgp-combined/01-run-combined-factorial.R` and
280 uses `furrr::future_map_dfr` with `plan(multicore)`; it sets `set.seed(20260610)` and
281 `furrr_options(seed = 20260610L)`. Per-replicate outputs are at `analysis/data/dgp-combined/01-combined`
282 and cell-level summaries at `01-combined-summary.rds`. The matched boundary re-
283 run described in Section~4.2 will be added as a separate driver that adopts the per-
284 randomization-path sample-size convention of the companion paper's production driver
285 (`analysis/scripts/quick-sim/01-dgp-prototype.R`). The manuscript source, bibliogra-
286 phy, and CSL file are at `analysis/report/11-combined-dgp-architecture/`.

287 11 References

- 288 1. pmsimstats team. Two architectures for simulating biomarker-treatment interactions: Implications for statistical power under carryover. Published online 2026.
- 289 2. Arminger G, Stein P, Wittenberg J. Mixtures of conditional mean- and covariance-structure models. *Psychometrika*. 1999;64(4):475-494.

- 290 3. pmsimstats team. Latent-class and mixture-model formulations for biomarker-treatment interaction in N-of-1 trials. Published online 2026.
- 291 4. Verbeke G, Lesaffre E. A linear mixed-effects model with heterogeneity in the random-effects population. *Journal of the American Statistical Association*. 1996;91(433):217-221.
- 292 5. pmsimstats team. Type I calibration of fixed-effect interaction tests in linear mixed models: A staged diagnosis of working-covariance misspecification and a cluster-robust remedy. Published online 2026.