

Carryover Half-Life and the Power of the Biomarker-Treatment Interaction Test: Closed-Form Results for N-of-1 and Crossover Designs

2026-10-02 16:33 PDT

Author. pmsimstats team

1. Summary

This paper isolates problem 2 of docs/36, the collapse of power under small carryover reported by Hendrickson et al. (2020) for the Hybrid (N-of-1) design, and studies it analytically. The core setting is the smallest case that still shows the problem: the Hybrid design, $N = 70$, the published compound-symmetry (CS) correlation construct, and the paired-difference interaction statistic (E9) of the compendium. For that statistic every moment that matters has a closed form. Sections 5.6 to 5.8 extend the framework to a fourth data-generating rule, to the OL+BDC and crossover (CO) designs, and to two AR(1) correlation structures, with the variance written in factor-by-time (separable) form. The results:

1. **The half-life enters in one place only.** The variance of each patient's on-minus-off contrast Δ does not depend on the carryover half-life $t_{1/2}$ at all: carryover acts on the response means, which the contrast's intercept absorbs, and on the biomarker coupling. The slope of Δ on the biomarker is

$$\gamma = -c_{bm} \frac{\sigma_{BR}}{\sigma_B} (1 - \bar{g}_{\text{off}}),$$

where $\bar{g}_{\text{off}}$ is the mean biomarker coupling over the patient's off-drug visits: the average fraction of the biomarker effect present in the biomarker's correlation with the drug response when the patient is off drug (defined in Section 2; derived in Section 4). Power is a strictly decreasing function of $\bar{g}_{\text{off}}$, so the entire effect of carryover on power runs through this one number (Section 5).

2. **Under the published coupling rule, power does not decline with the half-life; it switches off.** The published rule sets the coupling to its full value wherever any residual drug effect remains, so $\bar{g}_{\text{off}} = 1$ for every $t_{1/2} > 0$, however small. The slope is then exactly zero, and power equals the test size, for every positive half-life (Section 5.1). The published "precipitous decline" is a step at $t_{1/2} = 0$, not a decline. The residual power in the simulated published analysis (0.08-0.11 at $c_{bm} = 0.25$) comes from the published analysis counting the baseline visit as off drug, which leaves a slope of $1/(n_{\text{off}} + 1)$ of its no-carryover value; the closed form reproduces the published mean estimate to within 0.003.

3. **Under graded coupling, power declines monotonically, and the decline is provable.** With the coupling decaying as the drug effect does, $\bar{g}_{\text{off}}$ increases strictly with $t_{1/2}$, so the slope shrinks, the residual variance grows and power falls, strictly, for every $t_{1/2}$ (Section 5.2). The loss is negligible at the half-lives Hendrickson examined (power 0.721 at $t_{1/2} = 0.1$, the same as without carryover, and 0.712 at 0.2) and substantial only when carryover is long (0.473 at one week, 0.278 at two). The published step rule is the $t_{1/2} \rightarrow \infty$ limit of graded coupling: the published construct behaves, whenever carryover is nonzero, as if it were infinitely long.
4. **Under mean moderation, power does not depend on the half-life** (Section 5.3).
5. **Decayed mean moderation behaves like graded coupling.** If the biomarker-dependent shift of the drug response fades off drug at the rate of the drug effect, the slope has the same form as under graded coupling and power declines monotonically, with no ceiling on the effect size (Section 5.6).
6. **The results carry over to other designs, with one structural difference.** OL+BDC behaves like Hybrid. In CO the placebo-first path's off-drug visits come before any exposure, keep zero coupling under every rule, and so keep their full contrast: under the step rule CO loses only the active-first path's slope, and its power falls to 0.20 rather than to the test size. In CO the E9 statistic's common intercept also absorbs a large period effect, which costs power (Section 5.7).
7. **AR(1) structures cut power but do not change its shape.** Under the AR(1) covar and separable structures, power at $c_{bm} = 0.25$ is between a third (CO) and a half (Hybrid) of its CS value, with the same step, monotone decline and flat behavior across the rules. The factor-by-time decomposition of the contrast variance shows why: under CS the contrast cancels the large person-level term, under AR(1) widely spaced visits keep nearly the whole factor variance (Section 5.8).
8. **The closed form reproduces the simulations.** It agrees with a direct Monte Carlo of the statistic in 283 cells (three designs, three structures, four rules, effect sizes 0.1 to 0.3), and its baseline-inclusive variant reproduces the mean estimates of the published linear mixed-model (`lmer`) analysis in all three designs and its power in the Hybrid design (Section 6).

2. Setting

Design. The Hybrid design of Hendrickson et al. (2020): eight post-baseline visits at weeks 4, 8, 9, 10, 11, 12, 16 and 20, with expectancy 1 at the two open-label visits and 0.5 afterwards, and four paths. $N = 70$ patients are allocated 18, 18, 17, 17 across the paths, as in the published code. The time since discontinuation t_{sd} at each visit, from the published `buildtrialdesign()` semantics (verified):

Path	On drug (visits 1-8)	t_{sd} (weeks)	n_{on}	n_{off}
1	1 1 1 1 0 0 1 0	0 0 0 0 1 2 0 4	5	3
2	1 1 1 1 0 0 0 1	0 0 0 0 1 2 6 0	5	3
3	1 1 1 0 0 0 1 0	0 0 0 1 2 3 0 4	4	4
4	1 1 1 0 0 0 0 1	0 0 0 1 2 3 7 0	4	4

Every off-drug visit follows an exposure, so every off-drug visit has $t_{sd} \geq 1$ week. This matters below.

Data-generating process. The published construct (commit 58b32a9, configuration A of docs/36): patient-level baseline severity BL, biomarker B , and three response factors TV_t , PB_t , BR_t at each visit, jointly normal; outcome $Y_t = BL - (TV_t + PB_t + BR_t)$, and $Y_0 = BL$ at baseline. Standard deviations $\sigma_{TV} = 10$, $\sigma_{PB,t} = 10e_t$, $\sigma_{BR} = 8$, $\sigma_B = 15.36$. Correlations: within a factor, $\rho = 0.8$ between any two visits (compound symmetry); between factors, $c_1 = 0.2$ at the same visit and $c_\times = 0.1$ at different visits. BL is uncorrelated with everything else. The biomarker is correlated with BR_t only,

$$\text{Cor}(B, BR_t) = c_{bm} g_t,$$

Coupling. We call g_t the *coupling* at visit t : the fraction of the biomarker effect c_{bm} that is present in the biomarker's correlation with the drug response at that visit, with $0 \leq g_t \leq 1$. The term is ours, not Hendrickson et al.'s; docs/36 writes the same quantity as the coupling vector r , with $r_t = c_{bm}g_t$. In the covariance construct the biomarker does not enter the outcome equation; it moderates the drug effect only through this correlation, so a patient with a high biomarker tends to have a larger drug response at the visits where the coupling is present. The interaction the analysis is meant to detect is the difference between on-drug and off-drug coupling: if the biomarker predicts the response as strongly off drug as on drug, there is no biomarker-by-treatment interaction to find. Both covariance rules below set $g_t = 1$ on drug; they differ in the off-drug coupling, and that difference is the subject of this paper. Under the two mean-moderation rules the interaction is not carried by a correlation, $g_t = 0$ throughout, and the corresponding quantity is the *moderation profile* h_t : the fraction of the full biomarker-dependent shift of the drug response present at visit t .

The interaction is encoded by one of four rules:

- **Step (published).** $g_t = 1[\mu_{BR,t} \neq 0]$: full coupling wherever the mean drug response is nonzero, which off drug is the case whenever $t_{1/2} > 0$.

- **Graded.** $g_t = 1$ on drug and $g_t = \phi_t = 2^{-t_{sd}/t_{1/2}}$ off drug: coupling decays with the drug effect (adjustment 2 of `docs/36`).
- **Mean moderation.** $g_t = 0$; instead the on-drug BR_t is shifted by $c_{bm}\sigma_{BR}z$, z the standardized biomarker (adjustment 3 of `docs/36`): $h_t = 1$ on drug and 0 off drug.
- **Decayed mean moderation.** The same shift with $h_t = 1$ on drug and $h_t = \phi_t = 2^{-t_{sd}/t_{1/2}}$ off drug: the biomarker-dependent part of the drug response fades after discontinuation at the rate of the drug effect (`docs/38`, Section 4, gives the pharmacological argument for a proportional profile; this anchored form is the version directly comparable with graded coupling).

In every rule, an off-drug visit before any exposure ($t_{sd} = 0$, as in the placebo-first crossover path) has coupling and moderation 0.

For example, path 1's off-drug visits fall 1, 2 and 4 weeks after discontinuation. At $t_{1/2} = 0.5$ weeks their coupling is 1, 1 and 1 under the step rule, and 0.25, 0.06 and 0.004 under the graded rule.

Unless stated otherwise, $c_{bm} = 0.25$. Under CS the joint distribution exists only for c_{bm} below a ceiling, 0.256 in the Hybrid design without carryover (`docs/36`, Section 6.5); 0.25 is the largest round value feasible for every rule, design and structure at every half-life, so no matrix is repaired anywhere in this paper. The published value 0.3 is infeasible without carryover and 0.6 in most cells; the published results there come from repaired matrices whose coupling the formulas below would have to take from the repaired matrix.

3. The statistic

The compendium's paired-difference specification E9 (`analysis/report/compendium-mathematics.tex`, Section "Analysis specifications for carryover") is a two-stage estimator. For patient i ,

$$\Delta_i = \bar{Y}_i^{\text{on}} - \bar{Y}_i^{\text{off}} = \sum_t a_t Y_{it}, \quad a_t = \begin{cases} 1/n_{\text{on}} & \text{on drug} \\ -1/n_{\text{off}} & \text{off drug} \end{cases},$$

followed by the ordinary least squares (OLS) regression $\Delta_i = \beta_0 + \beta_1 B_i + \varepsilon_i$ over all N patients and a t -test of β_1 with $N - 2$ degrees of freedom.

The slope β_1 is the biomarker-by-treatment interaction: the change in a patient's on-minus-off outcome difference per unit of biomarker, in outcome points per biomarker unit. Since $Y = \text{BL} - (\dots)$ is a symptom score that falls as patients improve, a biomarker that predicts a larger drug response gives a negative slope. Its expected value is written γ below; without carryover it equals $-c_{bm}\sigma_{BR}/\sigma_B$, the drug-response gain in standard-deviation units per standard deviation of

biomarker, rescaled to the outcome and biomarker scales. It is the same quantity the published mixed model estimates as the `bm:Db` coefficient.

Two variants are treated:

- **E9**, the specification as written, on the eight post-baseline visits.
- **E9b**, the same with the baseline visit counted as one more off-drug observation ($n_{\text{off}} + 1$ off-drug rows). This is how the published analysis codes it: `lme_analysis()` adds a baseline row with $t = 0$ and $D_b = 0$ before fitting $Sx \sim \text{bm} + \text{Db} + \text{t} + \text{bm}*\text{Db} + (1|\text{ptID})$. E9b is not that mixed model, but Section 6 shows that its moments and power match the mixed model's closely in this setting, so it serves as a closed-form proxy for the published analysis.

In both variants the weights over all rows sum to zero, so BL, which enters every row including baseline, cancels:

$$\Delta = - \sum_{t=1}^8 a_t S_t, \quad S_t = TV_t + PB_t + BR_t.$$

For E9, $\sum_{t \geq 1} a_t = 0$. For E9b the post-baseline weights sum to $1/(n_{\text{off}} + 1)$, the baseline row carrying the remaining $-1/(n_{\text{off}} + 1)$ against an outcome of BL alone.

4. Closed-form moments

Δ and B are jointly normal within a path, so the regression of Δ on B is exactly linear,

$$E(\Delta | B) = \mu + \gamma(B - \mu_B), \quad \text{Var}(\Delta | B) = \tau^2,$$

and three quantities per path determine everything: the variance of Δ , its covariance with B , and its mean.

4.1 The variance of the contrast

With $A_1 = \sum a_t$, $A_2 = \sum a_t^2$, $A_e = \sum a_t e_t$, $A_{ee} = \sum a_t^2 e_t^2$ and $A_{2e} = \sum a_t^2 e_t$ (sums over the eight post-baseline visits), expanding $\text{Var}(a^\top S)$ over the CS structure gives

$$\begin{aligned} \text{Var}(\Delta) = & (\sigma_{TV}^2 + \sigma_{BR}^2)[\rho A_1^2 + (1 - \rho)A_2] + \sigma_{PB}^2[\rho A_e^2 + (1 - \rho)A_{ee}] \\ & + 2\sigma_{TV}\sigma_{BR}[c_\times A_1^2 + (c_1 - c_\times)A_2] + 2(\sigma_{TV} + \sigma_{BR})\sigma_{PB}[c_\times A_1 A_e + (c_1 - c_\times)A_{2e}]. \end{aligned}$$

For E9, $A_1 = 0$, and the person-level part of every factor (the ρ terms) cancels except where the placebo factor's visit-dependent scale e_t keeps A_e away from zero. Evaluated at the published values (verified, by hand for path 1 and by the script for all paths):

Variant	Paths 1, 2	Paths 3, 4
E9	44.13	45.03
E9b	51.04	51.77

Proposition 1. $\text{Var}(\Delta)$ *does not depend on $t_{1/2}$* . The formula contains the response standard deviations and correlations only; the half-life enters the published construct through the response means ($\mu_{BR,t}$, with its recursive carry-over) and through the coupling g_t , and neither appears. The script confirms the value is identical at every half-life to the last digit (verified). The argument is not specific to CS: it holds for any construct in which $t_{1/2}$ does not enter the response block, which includes every configuration of `docs/36`.

Counting the baseline as off drug (E9b) increases the variance by about 15%, because the baseline row adds BL-free weight that no longer cancels the person-level parts ($A_1 \neq 0$).

4.2 The slope

Only BR_t is correlated with B , with covariance $c_{bm}g_t\sigma_B\sigma_{BR}$, so $\text{Cov}(\Delta, B) = -c_{bm}\sigma_B\sigma_{BR}\sum_t a_t g_t$ and

$$\gamma = -c_{bm} \frac{\sigma_{BR}}{\sigma_B} \sum_{t=1}^8 a_t g_t.$$

On-drug visits have $g_t = 1$ and weights summing to 1. Writing $\bar{g}_{\text{off}}$ for the mean coupling over the n_{off} post-baseline off-drug visits, and recalling that the baseline has $g_0 = 0$:

Proposition 2.

$$\gamma_{\text{E9}} = -c_{bm} \frac{\sigma_{BR}}{\sigma_B} (1 - \bar{g}_{\text{off}}), \quad \gamma_{\text{E9b}} = -c_{bm} \frac{\sigma_{BR}}{\sigma_B} \left(1 - \frac{n_{\text{off}}}{n_{\text{off}} + 1} \bar{g}_{\text{off}}\right).$$

Without carryover $\bar{g}_{\text{off}} = 0$ and both equal the expected interaction slope $-c_{bm}\sigma_{BR}/\sigma_B = -0.1302$. Under the mean-moderation rules the shift $c_{bm}\sigma_{BR}z h_t$ contributes $-c_{bm}\sigma_{BR}z \sum_t a_t h_t$ to Δ , so

$$\gamma = -c_{bm} \frac{\sigma_{BR}}{\sigma_B} \sum_t a_t h_t = -c_{bm} \frac{\sigma_{BR}}{\sigma_B} (1 - \bar{h}_{\text{off}}) \quad (\text{E9}),$$

with $\bar{h}_{\text{off}} = 0$ under mean moderation, so $\gamma = -0.1302$ whatever the half-life, and $\bar{h}_{\text{off}} = \bar{\phi}_{\text{off}}$ under decayed mean moderation, the same form as graded coupling. In every rule the off-drug average includes the zero coupling of any off-drug visit before first exposure.

4.3 The residual variance and the correlation

For the covariance rules (step and graded), B explains part of the existing variance of Δ , so

$$\tau^2 = \text{Var}(\Delta) - \gamma^2 \sigma_B^2.$$

Under mean moderation the shift adds variance rather than explaining it, so $\text{Var}(\Delta)$ grows by $\gamma^2 \sigma_B^2$ and τ^2 is the covariance-rule $\text{Var}(\Delta)$ itself. This is why, at the same slope, mean moderation has slightly less power than the covariance rules without carryover (Section 5).

Power depends on these only through the correlation between the contrast and the biomarker,

$$\rho_{\Delta B} = \frac{\gamma \sigma_B}{\sqrt{\text{Var}(\Delta)}} = -c_{bm} \frac{\sigma_{BR}(1 - \bar{g}_{\text{off}})}{\text{SD}(\Delta)} \quad (\text{E9, covariance rules}),$$

0.30 in every path without carryover.

4.4 The pooled slope and its power

For a single path of n patients, write $S_B = \sum_i (B_i - \bar{B})^2$. Conditional on the biomarker values the OLS slope is exactly normal, $\hat{\beta}_1 | B \sim N(\gamma, \tau^2/S_B)$, and its t -statistic is exactly noncentral t with $n - 2$ degrees of freedom and noncentrality $\gamma\sqrt{S_B}/\tau$. Since $S_B \sim \sigma_B^2 \chi_{n-1}^2$, power is the noncentral- t tail probability averaged over that χ^2 law, and $\text{Var}(\hat{\beta}_1) = \tau^2 E(1/S_B) = \tau^2 / \{\sigma_B^2(n - 3)\}$. At $S_B = \sigma_B^2(n - 1)$ the noncentrality is

$$\delta = \sqrt{n - 1} \frac{\rho_{\Delta B}}{\sqrt{1 - \rho_{\Delta B}^2}},$$

so power is governed by the contrast-biomarker correlation alone.

E9 pools the four paths with a common intercept. With $\pi_p = n_p/N$ and $\mu_p = E(\Delta)$ in path p , the pooled slope has mean $\bar{\gamma} = \sum_p \pi_p \gamma_p$, and the differences between paths act as additional residual variance,

$$\sigma_{\text{res}}^2 = \sum_p \pi_p [\tau_p^2 + (\mu_p - \bar{\mu})^2 + \sigma_B^2 (\gamma_p - \bar{\gamma})^2],$$

so that the t -statistic is, to first order in the path differences, noncentral t with $N - 2$ degrees of freedom and noncentrality $\bar{\gamma}\sqrt{S_B}/\sigma_{\text{res}}$, $S_B \sim \sigma_B^2 \chi_{N-1}^2$, and

$$\text{Var}(\hat{\beta}_1) \approx \frac{\sum_p \pi_p [\tau_p^2 + (\mu_p - \bar{\mu})^2]}{(N-3)\sigma_B^2} + \frac{2}{N} \sum_p \pi_p (\gamma_p - \bar{\gamma})^2.$$

Power is computed by averaging the noncentral- t tail probability over 400 quantiles of χ_{N-1}^2 . A first version of the calculation used N in place of $N - 3$ and ignored the variation of S_B ; it overstated power by about 0.01-0.02, which the direct simulation of Section 6 detected. In the Hybrid and OL+BDC designs the path-mean term is small: in Hybrid at most 0.30 against residual variances of 41-49, under 0.7% of the total, decreasing with $t_{1/2}$; in OL+BDC at most 0.6% (verified). In CO it is not small (Section 5.7).

5. The effect of the half-life

By Propositions 1 and 2, $t_{1/2}$ affects power only through $\bar{g}_{\text{off}}$, and $|\rho_{\Delta B}|$, the noncentrality and power are all strictly decreasing in $\bar{g}_{\text{off}}$ (or $\bar{h}_{\text{off}}$ for the mean rules). The rules differ only in how that average depends on $t_{1/2}$. Sections 5.1 to 5.5 treat the Hybrid design under CS; Sections 5.6 to 5.8 extend the results.

Closed-form power and mean slope (verified):

Rule	Variant	$t_{1/2} =$					
		0	0.1	0.2	0.5	1	2
Step (published)	E9	0.721	0.050	0.050	0.050	0.050	0.050
	E9b	0.654	0.081	0.081	0.081	0.081	0.081
Graded	E9	0.721	0.721	0.712	0.631	0.473	0.278
	E9b	0.654	0.653	0.647	0.586	0.473	0.328
Mean moderation	E9	0.681	0.681	0.682	0.682	0.682	0.682
	E9b	0.619	0.619	0.620	0.620	0.621	0.621

Rule	Variant	Mean slope, $t_{1/2} =$					
		0	0.1	0.2	0.5	1	2
Step (pub- lished)	E9	-0.130	0	0	0	0	0
	E9b	-0.130	-0.029	-0.029	-0.029	-0.029	-0.029
Graded	E9	-0.130	-0.130	-0.129	-0.118	-0.098	-0.072
	E9b	-0.130	-0.130	-0.129	-0.121	-0.106	-0.085
Mean moder- ation	both	-0.130	-0.130	-0.130	-0.130	-0.130	-0.130




Figure 1: Power of the paired-difference statistic against carryover half-life

Figure 1. Closed-form power (lines; filled points at $t_{1/2} = 0$) against carryover half-life on a logarithmic axis, with direct Monte Carlo estimates (open points), for the four rules (decayed mean moderation dashed; Section 5.6). Hybrid design, $N = 70$, $c_{bm} = 0.25$, CS construct.

5.1 Step coupling: a switch, not a decline

Under the published rule $g_t = \mathbf{1}[\mu_{BR,t} \neq 0]$. With $t_{1/2} > 0$ the off-drug mean is $\mu_{BR,t} = \mu_{BR,t-1} 2^{-t_{sd}/t_{1/2}} > 0$ at every off-drug visit, because every off-drug visit in the Hybrid design follows an exposure. Hence $\bar{g}_{\text{off}} = 1$ for every $t_{1/2} > 0$, and by Proposition 2

$$\gamma_{E9} = 0, \quad \gamma_{E9b} = -c_{bm} \frac{\sigma_{BR}}{\sigma_B} \cdot \frac{1}{n_{\text{off}} + 1} \quad \text{for all } t_{1/2} > 0.$$

Power is therefore a step function of the half-life: its no-carryover value at $t_{1/2} = 0$, and a constant at every positive half-life, whether 0.01 weeks or 10. For E9 that constant is the nominal size, 0.050, because the slope is exactly zero; for E9b it is 0.081, carried by the baseline row alone, which is the only observation whose coupling stays at zero. The magnitude of the half-life plays no role. What the published figure shows as sensitivity to small carryover is the coupling rule redefining the estimand the moment any carryover is present: with constant coupling the biomarker moderates on-drug and off-drug response equally, and there is no interaction left to detect.

One numerical footnote. Below about 0.012 weeks the recursive off-drug mean at the latest off-drug visits underflows to exactly zero in double precision (it is about 2^{-1300} at the last off-drug visit of path 4 when $t_{1/2} = 0.01$), so the step rule, as coded both here and in the published code, switches those visits' coupling off. Figure 1 therefore starts at 0.02 weeks. This is a floating-point effect, not a property of the model.

The residual published power is explained exactly. E9b predicts mean slopes of $-0.1302/4 = -0.0325$ in paths 1 and 2 and $-0.1302/5 = -0.0260$ in paths 3 and 4, pooled -0.029 ; the published analysis, simulated with the published construct, gives -0.026 to -0.032 (Section 6).

5.2 Graded coupling: a provable monotone decline

Under graded coupling,

$$\bar{g}_{\text{off}}(t_{1/2}) = \frac{1}{n_{\text{off}}} \sum_{t_{\text{off}}} 2^{-t_{sd,t}/t_{1/2}}, \quad \frac{d\bar{g}_{\text{off}}}{dt_{1/2}} = \frac{\ln 2}{n_{\text{off}} t_{1/2}^2} \sum_{t_{\text{off}}} t_{sd,t} 2^{-t_{sd,t}/t_{1/2}} > 0,$$

since every off-drug visit has $t_{sd,t} \geq 1$.

Proposition 3. *Under graded coupling, in every path, $|\gamma_p|$ is strictly decreasing and τ_p^2 strictly increasing in $t_{1/2}$; hence the noncentrality, and the power of the path-stratified test, are strictly decreasing in $t_{1/2}$ on $(0, \infty)$. With $\bar{g}_{\text{off}}$ strictly increasing, Proposition 2 gives the slope, Section 4.3 gives $\tau_p^2 = \text{Var}(\Delta) - \gamma_p^2 \sigma_B^2$ with $\text{Var}(\Delta)$ fixed (Proposition 1), and the noncentral t tail probability is increasing in $|\delta|$. For the pooled E9 statistic as specified, with its common intercept, the path-mean term of Section 4.4 also moves with $t_{1/2}$; it is under 0.7% of the variance and decreases with $t_{1/2}$, and power is non-increasing at every step of a 120-point grid from 0.02 to 4 weeks for both E9 and E9b (verified), with steps too small to resolve only at the shortest half-lives, where the loss is $O(2^{-1/t_{1/2}})$.*

The limits make the relation to the published rule precise:

- $t_{1/2} \rightarrow 0^+$. $\bar{g}_{\text{off}} \leq 2^{-1/t_{1/2}}$, so the loss vanishes faster than any power of $t_{1/2}$: at $t_{1/2} = 0.1$ week, $\bar{g}_{\text{off}} \leq 2^{-10} \approx 0.001$ and power is unchanged to within 0.001. Power is continuous at zero.
- $t_{1/2} \rightarrow \infty$. $\bar{g}_{\text{off}} \rightarrow 1$, so graded coupling converges to the step rule. The published construct is the infinite-half-life limit of graded coupling, applied at every positive half-life.

Two features are worth noting. The decline is slow at first and faster later because $\bar{g}_{\text{off}}$ is a sum of $2^{-t_{sd}/t_{1/2}}$ terms, each negligible until $t_{1/2}$ is a substantial fraction of t_{sd} , which is at least one week here. And E9 and E9b cross: E9 has more power while carryover is short, because its variance is lower; E9b has more beyond about one week, because its baseline row keeps zero coupling however long the carryover and so retains part of the slope.

5.3 Mean moderation: no dependence

Under mean moderation the interaction is a shift of the on-drug mean only, so $\gamma = -c_{bm}\sigma_{BR}/\sigma_B$ and τ^2 are both free of $t_{1/2}$, and power is constant (0.681-0.682 for E9). The recursive carryover of the mean response off drug changes μ_p but not γ_p or τ_p^2 ; it enters only through the negligible path-mean term.

5.4 Power as a function of the coupling gap

Sections 5.1 to 5.3 are three special cases of one relation. With on-drug coupling 1, the slope is proportional to the *coupling gap* $\bar{g}_{\text{on}} - \bar{g}_{\text{off}}$, and power depends on the coupling profile only through it: neither which off-drug visits are coupled nor how the coupling is spread across them matters under E9, which weights every off-drug visit equally. Holding the response structure fixed and setting the off-drug coupling to $1 - \text{gap}$ at every off-drug visit gives power as a function of the gap alone (closed form, verified):

Coupling gap	0	0.25	0.5	0.75	1
Power, E9	0.050	0.094	0.236	0.468	0.721
Power, E9b	0.081	0.161	0.294	0.468	0.654

The coupling rules are points on this curve. Without carryover the gap is 1. Under graded coupling it is 0.991 at $t_{1/2} = 0.2$, 0.906 at 0.5, 0.754 at 1 and 0.551 at 2 weeks (patient-weighted over paths). Under the published step rule it is 0 at every positive half-life.




Figure 2: Power against the coupling gap

Figure 2. Closed-form power against the coupling gap, the mean on-drug minus mean off-drug coupling over the post-baseline visits, with the positions of the coupling rules marked (E9). Under E9b the baseline row keeps zero coupling, so its power does not fall to the test size at gap 0. Hybrid design, $N = 70$, $c_{bm} = 0.25$.

Three things follow. First, power falls steeply with the gap: the noncentrality is proportional to it, so the sample size needed for a given power scales roughly as $1/\text{gap}^2$, and halving the gap calls for about four times as many patients. Second, the curve is steepest near gap 1 and flattest near 0: the first 0.1 of gap lost costs 0.097 of power (0.721 to 0.624), the last 0.1 only 0.007 (0.057 to 0.050). Short carryover nevertheless costs almost nothing under graded coupling, because the gap itself barely moves (0.991 at $t_{1/2} = 0.2$), not because the curve is flat there. Third, the published step rule sits at the bottom of the curve for every positive half-life. The difference between the published result and a smooth decline is the difference between jumping from gap 1 to gap 0 and moving gradually along the curve. The curve is computed for the covariance construct; under mean moderation the gap is always 1 and the residual variance slightly larger (Section 4.3).

5.5 Power by effect size and half-life

Figure 3 extends the results from $c_{bm} = 0.25$ to a grid of effect sizes, $c_{bm} \in \{0, 0.1, 0.2, 0.3\}$ against $t_{1/2} \in \{0, 0.1, 0.2, 0.5, 1, 2\}$ weeks, one panel per rule (including decayed mean moderation, Section 5.6), in the layout of Hendrickson et al.'s Figure 4B. Each cell is the closed form of Section 4; every feasible cell at $c_{bm} = 0.1, 0.2$ and 0.3 was also simulated, and agrees (Section 6). Cells where c_{bm} is at or above the CS ceiling are not evaluated: there the correlation matrix is not positive definite, no joint distribution exists, and the formulas would return a number without meaning. The ceiling depends on the rule and the half-life (verified):

Rule	$t_{1/2} = 0$	0.1	0.2	0.5	1	2
Step (pub- lished)	0.256	0.893	0.893	0.893	0.893	0.893
Graded	0.256	0.256	0.257	0.274	0.307	0.369
Mean	none	none	none	none	none	none
moderation, decayed or not						




Figure 3: Power by biomarker effect and carryover half-life

Figure 3. Closed-form power by biomarker effect c_{bm} and carryover half-life, for each of the four rules (columns) and each variant of the statistic (rows). Gray cells lie at or above the CS ceiling. Hybrid design, $N = 70$.

The figure shows the three behaviors of Section 5 at every effect size:

- **Step rule: a wall at $t_{1/2} = 0$.** Every column after the first is at the test size for E9 (0.05) and barely above it for E9b (0.05-0.10), however large c_{bm} . The rule's ceiling jumps to 0.893 as soon as $t_{1/2} > 0$, because constant coupling costs almost nothing (docs/36, Section 6.5), so $c_{bm} = 0.3$ is feasible there but carries no signal. E9b at $c_{bm} = 0.3$ gives 0.09-0.10 at every positive half-life, close to the 0.12 and 0.10 that the published results report for the Hybrid design at $N = 70$, $c_{bm} = 0.3$, $t_{1/2} = 0.1$ and 0.2 (docs/36, Section 5.1); those published cells were not repaired, since 0.3 is below the 0.893 ceiling.
- **Graded rule: a gradual decline, and a rising ceiling.** At each feasible effect size power falls smoothly with the half-life. The ceiling rises with the half-life (0.256 to 0.369), because smoother coupling pays a smaller switch cost, so $c_{bm} = 0.3$ becomes feasible from $t_{1/2} = 1$ week. The result that power at $c_{bm} = 0.3$, $t_{1/2} = 1$ (0.63) exceeds power at $c_{bm} = 0.2$ without carryover (0.52 for E9) is a comparison across effect sizes, not a gain from carryover.
- **Mean moderation: flat in the half-life.** Each row is constant, and every cell is feasible, including $c_{bm} = 0.3$ (0.83 for E9).

In every rule, power at $c_{bm} = 0$ is the test size, 0.05, as it must be.

5.6 Decayed mean moderation

Under decayed mean moderation the slope is that of graded coupling (Section 4.2): the per-path slopes are identical, path by path and half-life by half-life (verified). The residual variance is not: under mean moderation the biomarker-dependent shift adds variance instead of explaining part of the existing variance (Section 4.3), so at every half-life power is a little lower than under graded coupling (Hybrid, CS, E9, closed form, verified):

Rule	$t_{1/2} = 0$	0.1	0.2	0.5	1	2
Graded coupling	0.721	0.721	0.712	0.631	0.473	0.278
Decayed mean moderation	0.681	0.681	0.674	0.598	0.454	0.271
Mean moderation	0.681	0.681	0.682	0.682	0.682	0.682

The decline is strictly monotone for the same reason as under graded coupling (Proposition 3 applies with h_t in place of g_t), and it is verified on the half-life

grid in every design and structure of this paper. Decayed mean moderation has no ceiling, so its decline can be traced at effect sizes where the covariance rules cannot be simulated: at $c_{bm} = 0.3$, E9 power is 0.83, 0.83, 0.82, 0.75, 0.60 and 0.37 across the half-lives (Figure 3).

5.7 Other designs: OL+BDC and crossover

The framework needs no change for other designs: each path has its own contrast weights a_t and its own coupling pattern, and pooling is as in Section 4.4. The two other published designs, with $N = 70$ allocated 35 and 35 (verified):

Design	Path	On drug (visits 1-8)	t_{sd} (weeks)
OL+BDC (weeks 4, 8, 12, 16, 17, 18, 19, 20)	1	1 1 1 1 1 1 0 0	0 0 0 0 0 0 1 2
	2	1 1 1 1 1 0 0 0	0 0 0 0 0 1 2 3
CO (every 2.5 weeks)	1	1 1 1 1 0 0 0 0	0 0 0 0 2.5 5 7.5 10
	2	0 0 0 0 1 1 1 1	0 0 0 0 0 0 0 0

OL+BDC has the structure of Hybrid: every off-drug visit follows an exposure, one to three weeks after it. Its results are those of Sections 5.1 to 5.5 in every respect (CS, E9, closed form, verified): the step rule falls to the test size at every positive half-life, and graded coupling declines monotonically (0.598, 0.597, 0.586, 0.477, 0.305, 0.156), faster than in Hybrid because its off-drug visits come sooner after discontinuation.

CO differs in one structural respect: path 2 receives placebo first, so its four off-drug visits precede any exposure and have $t_{sd} = 0$. Their coupling and moderation are 0 under every rule, whatever the half-life. Path 2 therefore keeps the full slope -0.1302 under every rule, and the half-life acts only on path 1. Two consequences follow.

- **Under the step rule CO keeps half its slope.** Path 1 loses its slope entirely; path 2 keeps it. The pooled E9 slope is -0.065 and E9b's is -0.078 at every positive half-life (path 1 keeps $1/(n_{\text{off}} + 1) = 1/5$ through the baseline row). The simulated published analysis gives -0.075 to -0.082 (Section 6). This is why the published crossover results lose less power than the N-of-1 results: at $N = 70$, $c_{bm} = 0.3$, from 0.85 to 0.44 and 0.46, against 0.74 to 0.12 and 0.10 for Hybrid (docs/36, Section 5.1).
- **Under graded coupling CO is robust to short carryover,** because path 1's off-drug visits are 2.5 to 10 weeks after discontinuation: power is 0.622 at every half-life up to 0.2 weeks and 0.581 at one week.

The period effect in CO. E9’s common intercept assumes that the paths share the same mean contrast. In CO they do not: path 1 is on drug in the first half of the trial and path 2 in the second, and the time-related and expectancy factors trend over the trial, so the mean contrasts differ (CS, mean moderation: -2.3 in path 1 against -11.1 in path 2 without carryover). The path-mean term of Section 4.4 is then 19.5 against a residual variance of 35.9 , 54% of it, and it grows with the half-life (to 25.5 at two weeks) because carryover narrows path 1’s contrast (verified). This term, not the interaction, is why E9 power under mean moderation falls in CO (0.590 to 0.548 over two weeks), and it is the main source of E9’s inefficiency there relative to a mixed model with a time term. Setting the term to zero, as an analysis with path-specific intercepts would, gives (closed form, not simulated):

CO, CS, E9	$t_{1/2} = 0$	0.1	1	2
Step, as specified	0.622	0.199	0.193	0.184
Step, path term removed	0.817	0.283	0.283	0.283
Graded, as specified	0.622	0.622	0.581	0.497
Graded, path term removed	0.817	0.817	0.794	0.737
Mean moderation, as specified	0.590	0.590	0.575	0.548
Mean moderation, path term removed	0.771	0.771	0.771	0.771

With the period effect removed, the three behaviors are exactly those of the Hybrid design. In Hybrid and OL+BDC removing the term changes power by at most 0.003 .

5.8 AR(1) correlation structures and the factor-by-time decomposition

The framework applies to any response correlation structure, because the slope (Proposition 2) does not depend on it and the variance (Proposition 1) does not depend on the half-life under any of them. Two AR(1) structures from `docs/36` were evaluated: `covar` (AR(1) within a factor, cross-factor correlation $c_{\times}\rho^{\text{gap}}$, configuration F) and the separable form (cross-factor $c_1\rho^{\text{gap}}$, configuration C), with $\rho = 0.8$ per week of visit separation.

Factorization. Writing $s_t = (\sigma_{TV}, \sigma_{PB,t}, \sigma_{BR})$ for the factor standard deviations at visit t , each structure’s response correlation is a sum of factor-by-time Kronecker terms (`docs/36`, Appendix A.7), and the contrast variance factors accordingly:

$$\text{Var}(\Delta) = \sum_k \sum_{t,s} a_t a_s T_{ts}^{(k)} (s_t^\top K^{(k)} s_s),$$

Structure	Terms $K^{(k)} \otimes T^{(k)}$
Separable AR(1)	$K \otimes A, K = (1 - c_1)I_3 + c_1J_3$
AR(1) covar	$K_a \otimes A + K_b \otimes I,$ $K_a = (1 - c_\times)I_3 + c_\times J_3,$ $K_b = (c_1 - c_\times)(J_3 - I_3)$
CS (published)	$K_o \otimes I + K_p \otimes J,$ $K_o = (1 - \rho)I_3 + (c_1 - c_\times)(J_3 - I_3),$ $K_p = \rho I_3 + c_\times(J_3 - I_3)$

with $A_{ts} = \rho^{|w_t - w_s|}$ the AR(1) time kernel, I the occasion kernel and J the all-ones (person-level) kernel. When the placebo scale e_t is constant over the contrast, as in CO, the separable variance is a single product, $\text{Var}(\Delta) = (s^\top K s)(a^\top A a)$. The separable ceiling factors as well, $c_{bm}^* = ([K^{-1}]_{33} \max_p g_p^\top A^{-1} g_p)^{-1/2}$. The factored variance agrees with the full-matrix value to 10^{-13} in all 48 design-structure-statistic-path cases, and the factored ceiling with the direct ceiling to 10^{-16} in all 36 cases (verified).

Why AR(1) costs power. The decomposition makes the comparison transparent (E9 contrast variance, mean over paths, verified):

Design	CS	AR(1) covar	Separable AR(1)
Hybrid	44.6	121.5	117.9
OL+BDC	58.4	186.0	188.3
CO	35.9	201.7	212.9

Under CS the person-level term $K_p \otimes J$ is constant across visits, so a contrast whose weights sum to zero cancels it, leaving only the occasion-level term, whose factor variances are $1 - \rho = 0.2$. Under AR(1) there is no constant component to cancel: the correlation between visits decays with their separation (0.8 at one week, 0.57 at the 2.5-week CO spacing, 0.33 at five weeks), so a contrast between visits weeks apart retains most of each factor’s variance. The contrast variance is 2.6 times larger in Hybrid and up to 5.9 times larger in CO, and since the slope is the same under every structure, the comparison is at an identical interaction and only the noise differs. Power at $c_{bm} = 0.25$ falls accordingly (E9, closed form, verified):

Design	Structure	Step, 0 / any $t_{1/2} > 0$	Graded, $t_{1/2}$ $= 0 / 0.5 / 1$ $/ 2$	Mean moderation
Hybrid	CS	0.721 / 0.050	0.721 / 0.631 / 0.473 / 0.278	0.681

Design	Structure	Step, 0 / any $t_{1/2} > 0$	Graded, $t_{1/2}$ = 0 / 0.5 / 1 / 2	Mean moderation
Hybrid	AR(1) covar	0.326 / 0.050	0.326 / 0.276 / 0.205 / 0.130	0.317
Hybrid	Separable	0.335 / 0.050	0.335 / 0.283 / 0.210 / 0.133	0.325
OL+BDC	CS	0.598 / 0.050	0.598 / 0.477 / 0.305 / 0.156	0.568
OL+BDC	AR(1) covar	0.229 / 0.050	0.229 / 0.183 / 0.127 / 0.082	0.225
OL+BDC	Separable	0.226 / 0.050	0.226 / 0.181 / 0.126 / 0.082	0.222
CO	CS	0.622 / 0.199	0.622 / 0.615 / 0.581 / 0.497	0.590-0.548
CO	AR(1) covar	0.199 / 0.086	0.199 / 0.198 / 0.190 / 0.170	0.196-0.192
CO	Separable	0.192 / 0.084	0.192 / 0.190 / 0.183 / 0.164	0.189-0.186

The shape of the dependence on the half-life is the same under every structure: a switch under the step rule, a monotone decline under graded coupling and decayed mean moderation, and no dependence under mean moderation except through CO's period effect. Separability changes power little relative to **covar** (at most 0.009 in this table); its effect is on feasibility. The ceilings (verified):

Design	Structure	Step, 0 / $t_{1/2} > 0$	Graded, $t_{1/2} = 0$ / 1 / 2
Hybrid	CS	0.256 / 0.893	0.256 / 0.307 / 0.369
Hybrid	AR(1) covar	0.291 / 0.595	0.291 / 0.486 / 0.559
Hybrid	Separable	0.440 / 0.588	0.440 / 0.544 / 0.581
OL+BDC	CS	0.262 / 0.893	0.262 / 0.348 / 0.469

Design	Structure	Step, 0 / $t_{1/2} > 0$	Graded, $t_{1/2} = 0$ / 1 / 2
OL+BDC	AR(1) covar	0.303 / 0.594	0.303 / 0.517 / 0.592
OL+BDC	Separable	0.468 / 0.588	0.468 / 0.578 / 0.609
CO	CS	0.256 / 0.256	0.256 / 0.256 / 0.256
CO	AR(1) covar	0.627 / 0.576	0.627 / 0.627 / 0.627
CO	Separable	0.637 / 0.567	0.637 / 0.637 / 0.637

In CO under CS the ceiling does not move with the half-life, because the placebo-first path's off-drug visits keep zero coupling and preserve the contrast that sets it.

Figure 4. Closed-form power (lines) and Monte Carlo estimates (open points) against carryover half-life for the three response correlation structures, Hybrid design, $N = 70$, $c_{bm} = 0.25$.

Figure 5. E9 power against carryover half-life by design (columns) and response structure (rows), $N = 70$, $c_{bm} = 0.25$. The CO step-rule curve stays above the test size because the placebo-first path keeps its contrast.

6. Validation

Direct Monte Carlo. The statistic was simulated directly from the published construct for each rule at $t_{1/2} \in \{0, 0.1, 0.2, 0.5, 1, 2\}$.

Each cell used 5000 simulated trials of 70 patients, drawn from the full joint distribution (not from the formulas), giving Monte Carlo standard errors of about 0.007 on power, 0.0007 on the mean slope and 0.0005 on its SD. E9 (verified):

Rule	$t_{1/2}$	Mean, closed	Mean, MC	SD, closed	SD, MC	Power, closed	Power, MC
Step	0	-0.1302	-0.1299	0.0507	0.0508	0.721	0.715
Step	0.1	0	0.0010	0.0531	0.0536	0.050	0.054
Step	0.5	0	-0.0010	0.0531	0.0525	0.050	0.050
Step	2	0	-0.0008	0.0531	0.0533	0.050	0.052
Graded	0	-0.1302	-0.1301	0.0507	0.0498	0.721	0.723
Graded	0.2	-0.1290	-0.1296	0.0508	0.0496	0.712	0.726
Graded	0.5	-0.1180	-0.1190	0.0511	0.0516	0.631	0.634

Rule	$t_{1/2}$	Mean, closed	Mean, MC	SD, closed	SD, MC	Power, closed	Power, MC
Graded	1	-0.0982	-0.0977	0.0517	0.0510	0.473	0.479
Graded	2	-0.0717	-0.0724	0.0524	0.0516	0.278	0.273
Mean	0	-0.1302	-0.1307	0.0531	0.0521	0.681	0.688
moder- ation							
Mean	1	-0.1302	-0.1314	0.0531	0.0538	0.682	0.681
moder- ation							
Mean	2	-0.1302	-0.1284	0.0531	0.0529	0.682	0.677
moder- ation							

The full validation covers 283 cells, each simulated with 5000 trials and analyzed with both statistics, 566 cell-variants in all (full table in `e9-closed-form-vs-mc.csv`; verified):

Block	Cell- variants	Largest $ z $	Mean z	Beyond 2 SE	Largest mean difference
Hybrid, CS, $c_{bm} =$ 0.25, four rules	48	2.66	0.21	4	0.0017
Hybrid, CS, $c_{bm} =$ 0.1, 0.2, 0.3 (Figure 3)	134	2.45	-0.21	2	0.0018
Hybrid, AR(1) structures	96	2.18	-0.01	2	0.0027
OL+BDC, three structures	144	2.83	0.05	5	0.0038
CO, three structures	144	2.63	-0.09	4	0.0044

Here z is the difference between simulated and closed-form power in binomial standard errors. Seventeen of 566 cell-variants (3.0%) lie beyond 2 standard errors, against about 5% expected by chance, with deviations of both signs. The heatmap rows of Figure 3 include cells close to the ceiling (graded coupling




Figure 4: Power against carryover half-life by response structure




Figure 5: Power against carryover half-life by design and structure

at $c_{bm} = 0.3$ and $t_{1/2} = 1$, ceiling 0.307), where any approximation error would be largest; there the deviations are under 0.6 standard errors. The qualitative results of Section 5, the exact zero slope under the step rule, the floor at the test size, the CO floor of Section 5.7 and the monotone declines, are reproduced without exception.

Against the published mixed-model analysis. docs/36 simulated the published construct (configuration A, step rule) and mean moderation (configuration D) with the published `lmer` analysis, 500 replicates per cell, at $c_{bm} = 0.25$. E9b reproduces it (verified):

Rule	$t_{1/2}$	<code>lmer</code> mean	E9b mean	<code>lmer</code> SD	E9b SD	<code>lmer</code> power	E9b power
Step	0	-0.126	-0.130	0.055	0.055	0.632	0.654
Step	0.1	-0.028	-0.029	0.055	0.057	0.080	0.081
Step	0.2	-0.026	-0.029	0.058	0.057	0.096	0.081
Step	0.5	-0.029	-0.029	0.058	0.057	0.108	0.081
Step	1	-0.032	-0.029	0.059	0.057	0.108	0.081
Mean	0	-0.131	-0.130	0.059	0.057	0.628	0.619
moder- ation							
Mean	0.1	-0.132	-0.130	0.058	0.057	0.672	0.619
moder- ation							
Mean	0.2	-0.127	-0.130	0.058	0.057	0.636	0.620
moder- ation							
Mean	0.5	-0.131	-0.130	0.059	0.057	0.646	0.620
moder- ation							
Mean	1	-0.129	-0.130	0.058	0.057	0.652	0.621
moder- ation							

The Monte Carlo standard error of the `lmer` power is about 0.02 and of its mean estimate about 0.0026. The mean estimates agree to within 0.003 and the standard deviations to within 0.002. Power agrees to within 0.03 in nine of the ten cells; the exception is mean moderation at $t_{1/2} = 0.1$ (0.672 against 0.619, about 2.6 Monte Carlo standard errors of the `lmer` estimate). Under mean moderation the mixed model is ahead in every cell (0.63-0.67 against 0.62), plausibly because it weights the visits by their covariance where E9b uses a simple on-off contrast; this was not tested separately. The structure of the published curve, a level of about 0.63 without carryover and a flat floor near 0.08-0.11 at every positive half-life, is exactly the structure Section 5.1 derives.

In the other two designs (`e9-vs-lmer.csv`, verified), E9b reproduces the `lmer` mean estimates in every cell, including the partial floors of the step rule: -0.078 against -0.075 to -0.082 in CO, and -0.038 against -0.034 to -0.040 in OL+BDC. Its power is lower than the mixed model's, for two different reasons. In CO the mixed model is more efficient: its estimates vary less (SD 0.048-0.053 against 0.058), consistent with its time term absorbing the period effect of Section 5.7 (inferred). In OL+BDC the SDs agree (0.063-0.070 against 0.065-0.067), yet `lmer` power is higher (0.56-0.63 against 0.49 under mean moderation, 0.12-0.13 against 0.087 under the step rule with carryover). That pattern, equal variability with more rejections, points to an anticonservative mixed-model test, and it is consistent with the elevated Type I error of the published analysis in OL+BDC under compound symmetry, 0.06-0.11, that `docs/36` (Section 6.8) reports but does not explain (inferred; not tested here). The closed form is therefore a proxy for the published analysis's mean behavior in every design, and for its power in the Hybrid design.

7. Discussion

What problem 2 is. In the simplest setting that shows it, the power collapse has a complete analytic account. Carryover does not degrade the information in the data: the variance of the contrast is fixed, and the half-life enters only through the biomarker coupling at off-drug visits. The published rule makes that coupling an indicator of whether any drug effect remains, so the interaction the simulation encodes vanishes at the first instant of carryover. The analysis then correctly finds nothing to detect. The problem is in the data-generating rule, not in the trial design or the analysis.

What a decline should look like. If the biomarker's association with the drug response fades as the drug effect fades, power falls monotonically and smoothly with the half-life, by Proposition 3, starting from no loss at all for half-lives short relative to the off-drug intervals. At the published half-lives (0.1 and 0.2 weeks, against off-drug intervals of at least one week) the loss is under 0.01. This is the behavior the published discussion appears to have expected, and it is what the corrected coupling of `docs/36` produces in the full simulation.

One number to watch. Because power depends on carryover only through $\bar{g}_{\text{off}}$, a design can be judged by how much coupling it leaves at its off-drug visits: off-drug visits soon after discontinuation (small t_{sd}) are the ones that lose their contrast first as carryover lengthens. In the Hybrid design these are the first blinded-discontinuation visits, one week after stopping.

Covariance and mean moderation. With the same slope, the covariance rules give slightly more power than mean moderation without carryover (0.721 against 0.681 for E9), because under covariance moderation the biomarker explains part of the variance the contrast already has, whereas under mean moderation it adds variance on top of it. The same holds between graded coupling and decayed mean moderation, which share a slope at every half-life.

Design and structure. The design determines whether carryover can remove the contrast at all: a design with off-drug observations before any exposure, like the placebo-first crossover path, keeps them free of carryover under every rule. The correlation structure determines how much noise the contrast carries, not how the half-life acts on it: under AR(1) the same interaction is detected with a third to a half of the power, but the step, the monotone decline and the flat profile appear exactly as under CS.

8. Limitations

- **Three structures, three designs.** The results cover CS and two AR(1) structures, and the Hybrid, OL+BDC and CO designs; the OL design has no off-drug observation and E9 is not defined for it. Proposition 1 holds for any construct whose response block does not depend on $t_{1/2}$.
- **One sample size.** $N = 70$ throughout, and $c_{bm} = 0.25$ except in Figure 3. The structure of the results does not depend on these choices; the numbers do.
- **Below the ceiling only.** The formulas assume the joint distribution exists. At the published $c_{bm} = 0.3$ without carryover, and at 0.6 in most cells, it does not; the published results there use repaired matrices, whose couplings differ from $c_{bm}g_t$.
- **E9 is not the published mixed model.** E9b reproduces the mixed model’s mean estimates in all three designs and its power in Hybrid; in CO it is less efficient, and in OL+BDC the mixed model’s test appears anticonservative (Section 6). The closed form is exact for the paired-difference statistic, not for `lmer`.
- **The CO period effect.** E9’s common intercept absorbs the period effect of the crossover poorly. The path-stratified figures of Section 5.7 are closed form and were not simulated.
- **Large-sample pooling.** The pooled-slope variance is first order in $1/N$ and power uses a normal approximation to the slope; the single-path results are exact.
- **Binary on/off contrast.** E9 contrasts on-drug against off-drug means. A contrast weighted by the expected exposure (the compendium’s D_{bc}) would change the weights a_t and is not treated.

9. Reproducibility

```
Rscript analysis/scripts/quick-sim/carryover-closed-form/01-e9-closed-form.R \
--mc-reps 5000 --cores 6
```

Run from the repository root. The script reads the published construct from `analysis/scripts/quick-sim/hendrickson-problems/04-power-simulation.R` (which reproduces the matrices of the 58b32a9 code exactly; docs/36, Section

6.5) and the `lmer` results from `analysis/data/quick-sim/hendrickson-problems/power-sim/power-summary`. It writes the per-path moments, the closed-form curves for all designs, structures and rules, the ceilings, the factorization check, the power heatmap and both validation tables to `analysis/data/quick-sim/carryover-closed-form/`, and Figures 1 to 5 to `docs/figures/37-fig*.png`. Each Monte Carlo cell is cached under `mc-cells-reps5000/`, so a rerun simulates only missing cells.

10. References

1. Hendrickson RC, Thomas RG, Schork NJ, Raskind MA. Optimizing aggregated N-of-1 trial designs for predictive biomarker validation: statistical methods and theoretical findings. *Frontiers in Digital Health* 2020; 2:13. Code: github.com/rchendrickson/pmsimstats, commit 58b32a9.
2. pmsimstats team. *Compendium mathematics*, `analysis/report/compendium-mathematics.tex` (specification E9; Paper 08's closed-form RM-ANOVA power result).
3. pmsimstats team. *Positive definiteness and carryover sensitivity in the Hendrickson et al. (2020) simulation framework*, `docs/36-hendrickson-pd-and-carryover.md`.