

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29

Data Generation Machinery for N-of-1 Clinical Trial

Simulations: Appendix B

pmsimstats team

September 30, 2026

Contents

Appendix B. Correlation structure of the covariance-moderation and published data-generating processes	1
B.1 Variable ordering and block partition	1
B.2 Within-factor blocks A_c	2
B.3 Cross-factor blocks $C_{cc'}$	3
B.4 Biomarker coupling vector r	3
B.5 A worked example	4
B.6 The three data-generating processes	6
B.7 Summary of differences	9
B.8 The mean-covariance contrast	11

This document reproduces Appendix B of the manuscript `report.Rmd`. Cross-references to numbered sections and to Appendix A refer to that manuscript.

Appendix B. Correlation structure of the covariance-moderation and published data-generating processes

This appendix sets out the correlation matrix that each covariance-based data-generating process assembles, so that the differences discussed in Section 2.2.4 can be read off directly. Throughout, `covar` denotes Architecture B as defined in Section 2.2.2 and `orig` denotes the published implementation at commit `58b32a9`. Symbols follow Section 2.1, with the participant subscript i suppressed because every participant’s trajectory is drawn from the same matrix: B is the biomarker B_i , and BR_t is the biological response component BR_{it} at occasion t . The vector of biomarker-response correlations is written r , with entries r_t , so that b_i keeps its meaning of the standardized biomarker. Mean moderation does not appear below, because it writes no biomarker-response correlation into the covariance at all; its vector r is identically zero and its interaction is carried in the mean instead.

B.1 Variable ordering and block partition

Both implementations draw each participant's trajectory from a single multivariate normal over the same ordered variable list. With n measurement occasions at cumulative weeks $w_1 < w_2 < \dots < w_n$, the vector is

$$X = (B, BL, TV_1, \dots, TV_n, PB_1, \dots, PB_n, BR_1, \dots, BR_n)^\top, \quad \dim X = 2 + 3n.$$

The ordering is factor-major: all n occasions of one factor are contiguous, rather than all three factors at one occasion. Writing A_c for the $n \times n$ within-factor block of factor $c \in \{tv, pb, br\}$ and $C_{cc'}$ for the $n \times n$ cross-factor block, the correlation matrix partitions as

$$R = \begin{pmatrix} 1 & 0 & 0^\top & 0^\top & r^\top \\ 0 & 1 & 0^\top & 0^\top & 0^\top \\ 0 & 0 & A_{tv} & C_{tv,pb} & C_{tv,br} \\ 0 & 0 & C_{pb,tv} & A_{pb} & C_{pb,br} \\ r & 0 & C_{br,tv} & C_{br,pb} & A_{br} \end{pmatrix}.$$

Two features hold in both implementations. The baseline term BL is uncorrelated with every other variable, so its row and column are those of the identity. The biomarker B couples only to the response factor BR , through the vector $r \in \mathbb{R}^n$; its correlations with TV and PB are zero. The two implementations differ in the interior of A_c , in the interior of $C_{cc'}$, and in r .

B.2 Within-factor blocks A_c

Let ρ_c denote the autocorrelation parameter of factor c (`c.tv`, `c.pb`, `c.br` in the code).

orig: compound symmetry. A single constant is assigned to every pair of occasions, irrespective of their separation:

$$[A_c]_{ij} = \begin{cases} 1, & i = j \\ \rho_c, & i \neq j \end{cases} \iff A_c = (1 - \rho_c)I_n + \rho_c J_n,$$

with J_n the $n \times n$ matrix of ones. For $n = 4$,

$$A_c^{\text{orig}} = \begin{pmatrix} 1 & \rho_c & \rho_c & \rho_c \\ \rho_c & 1 & \rho_c & \rho_c \\ \rho_c & \rho_c & 1 & \rho_c \\ \rho_c & \rho_c & \rho_c & 1 \end{pmatrix}.$$

The block has two distinct eigenvalues, $1 + (n - 1)\rho_c$ once and $1 - \rho_c$ with multiplicity $n - 1$, so positive definiteness of the block requires $-1/(n - 1) < \rho_c < 1$. We note that it is the

lower bound that moves with n , and it binds only for negative ρ_c . At the positive values used throughout this paper the admissible range of ρ_c is not affected by the number of occasions. Compound symmetry does restrict the moderation parameter c_{bm} more tightly than AR(1) does, but not by this route. The mechanism is given in B.7.

covar: AR(1) in calendar time. The correlation decays with the elapsed time between occasions, using the cumulative week gap rather than the index difference:

$$[A_c]_{ij} = \rho_c^{|w_i - w_j|}.$$

For $n = 4$, writing $d_{ij} = |w_i - w_j|$,

$$A_c^{\text{covar}} = \begin{pmatrix} 1 & \rho_c^{d_{12}} & \rho_c^{d_{13}} & \rho_c^{d_{14}} \\ \rho_c^{d_{12}} & 1 & \rho_c^{d_{23}} & \rho_c^{d_{24}} \\ \rho_c^{d_{13}} & \rho_c^{d_{23}} & 1 & \rho_c^{d_{34}} \\ \rho_c^{d_{14}} & \rho_c^{d_{24}} & \rho_c^{d_{34}} & 1 \end{pmatrix}.$$

Because the exponent uses calendar time, unequally spaced designs produce unequal decay across consecutive occasions. In the Hybrid design, whose occasions fall at weeks 4, 8, 9, 10, 11, 12, 16, 20, consecutive gaps of one week during the discontinuation phase yield much higher adjacent correlations than the four-week gaps at either end.

B.3 Cross-factor blocks $C_{cc'}$

Let c_1 denote the same-occasion cross-factor correlation (`c.cf1t`) and $c_\times$ the different-occasion cross-factor correlation (`c.cfct`).

orig. Both are constants:

$$[C_{cc'}]_{ij} = \begin{cases} c_1, & i = j \\ c_\times, & i \neq j \end{cases} \iff C_{cc'} = c_\times J_n + (c_1 - c_\times) I_n.$$

covar. The same-occasion entry is unchanged, but the different-occasion entries inherit the AR(1) decay:

$$[C_{cc'}]_{ij} = \begin{cases} c_1, & i = j \\ c_\times \rho^{|w_i - w_j|}, & i \neq j. \end{cases}$$

73 B.4 Biomarker coupling vector r

74 The vector r carries the biomarker-treatment interaction. Both implementations write it in
75 terms of the same quantity, the residual fraction of the on-drug response at occasion t ,

$$76 \quad \phi_t = \begin{cases} (1/2)^{t_{sd,t}/t_{1/2}}, & \text{occasion } t \text{ off drug, after first exposure, } t_{1/2} > 0 \\ 0, & \text{otherwise,} \end{cases}$$

77 where $t_{sd,t}$ is the time since discontinuation. ϕ_t is zero throughout when $t_{1/2} = 0$ and at
78 off-drug occasions before first exposure, and it is the factor by which carryover multiplies the
79 on-drug response mean.

80 **covar.** The entry equals c_{bm} on drug and is scaled by the residual fraction off drug:

$$81 \quad r_t = \begin{cases} c_{bm}, & \text{occasion } t \text{ on drug} \\ c_{bm} \phi_t, & \text{occasion } t \text{ off drug.} \end{cases}$$

82 Since $\phi_t = e^{-\lambda t_{sd,t}}$ with $\lambda = \ln 2/t_{1/2}$, this is the exponential decay of Section 2.2.2.

83 **orig.** The entry equals c_{bm} wherever the response mean $\mu_{BR,t}$ is nonzero and zero elsewhere.
84 Because $\mu_{BR,t}$ is nonzero off drug exactly when $\phi_t > 0$, this is

$$85 \quad r_t = \begin{cases} c_{bm}, & \text{occasion } t \text{ on drug} \\ c_{bm} \mathbf{1}[\phi_t > 0], & \text{occasion } t \text{ off drug.} \end{cases}$$

86 The two vectors are therefore nonzero at the same occasions. The published code's condition
87 on $\mu_{BR,t}$ and **covar**'s branch on the drug indicator switch the channel on in the same places,
88 and they differ only in how large the off-drug entry is: graded in proportion to the residual
89 effect under **covar**, and a step to the full on-drug value under **orig**. At $t_{1/2} = 0$, $\phi_t = 0$
90 at every off-drug occasion and the two vectors are equal entry for entry. For $t_{1/2} > 0$ both
91 sets of off-drug entries rise above zero, but the **covar** entries rise only as far as the residual
92 fraction allows, while the **orig** entries jump to c_{bm} . In a design whose every path begins on
93 drug, such as the Hybrid design, $r^{\text{orig}} = c_{bm} \mathbf{1}_n$ at any positive half-life, so the vector carries
94 no on-drug versus off-drug contrast at all.

95 B.5 A worked example

96 To make the three components concrete, take four occasions drawn from the Hybrid design at
97 cumulative weeks $w = (10, 11, 12, 16)$, with the first and fourth on drug and the second and
98 third off it, so that $t_{sd} = (0, 1, 2, 0)$. Set $\rho_{br} = 0.6$, $\rho_{pb} = 0.4$, $c_1 = 0.3$, $c_x = 0.1$, $c_{bm} = 0.45$,

99 and a carryover half-life of one week, giving $\lambda = \ln 2 \approx 0.693$. The matrix of week gaps is

$$100 \quad (|w_i - w_j|) = \begin{pmatrix} 0 & 1 & 2 & 6 \\ 1 & 0 & 1 & 5 \\ 2 & 1 & 0 & 4 \\ 6 & 5 & 4 & 0 \end{pmatrix}.$$

101 The unequal spacing matters: occasions one to three are a week apart, while the fourth sits
102 four weeks after the third.

103 **B.5.1 Within-factor block A_{br}**

$$104 \quad A_{br}^{\text{orig}} = \begin{pmatrix} 1.0000 & 0.6000 & 0.6000 & 0.6000 \\ 0.6000 & 1.0000 & 0.6000 & 0.6000 \\ 0.6000 & 0.6000 & 1.0000 & 0.6000 \\ 0.6000 & 0.6000 & 0.6000 & 1.0000 \end{pmatrix}, \quad A_{br}^{\text{covar}} = \begin{pmatrix} 1.0000 & 0.6000 & 0.3600 & 0.0467 \\ 0.6000 & 1.0000 & 0.6000 & 0.0778 \\ 0.3600 & 0.6000 & 1.0000 & 0.1296 \\ 0.0467 & 0.0778 & 0.1296 & 1.0000 \end{pmatrix}.$$

105 Under **orig** every pair of occasions carries 0.6, so the week-apart pair (1, 2) and the six-weeks-
106 apart pair (1, 4) are treated as equally associated. Under **covar** the same two pairs carry 0.6
107 and 0.0467, a factor of thirteen apart. The consequence for the fourth occasion is the striking
108 one: it is nearly independent of the first three under AR(1), and as strongly tied to them as
109 they are to each other under compound symmetry.

110 **B.5.2 Cross-factor block $C_{tv,pb}$**

$$111 \quad C_{tv,pb}^{\text{orig}} = \begin{pmatrix} 0.3000 & 0.1000 & 0.1000 & 0.1000 \\ 0.1000 & 0.3000 & 0.1000 & 0.1000 \\ 0.1000 & 0.1000 & 0.3000 & 0.1000 \\ 0.1000 & 0.1000 & 0.1000 & 0.3000 \end{pmatrix}, \quad C_{tv,pb}^{\text{covar}} = \begin{pmatrix} 0.3000 & 0.0400 & 0.0160 & 0.0004 \\ 0.0400 & 0.3000 & 0.0400 & 0.0010 \\ 0.0160 & 0.0400 & 0.3000 & 0.0026 \\ 0.0004 & 0.0010 & 0.0026 & 0.3000 \end{pmatrix}.$$

112 The diagonal, which carries the same-occasion cross-factor correlation c_1 , is identical in both.
113 The off-diagonal entries decay under **covar** and do not under **orig**, and by the fourth occasion
114 the difference is two orders of magnitude.

115 A caution attaches to the **covar** block. In the implementation the decay rate applied to a
116 cross-factor pair is taken from whichever factor the outer loop is currently visiting, and because
117 each unordered pair is visited in both orders, the later write survives. The rate appearing
118 in $C_{tv,pb}$ is therefore determined by loop order rather than by any modeling decision. The
119 values above use ρ_{pb} on that basis. This is an implementation artifact and not a property of
120 the AR(1) specification.

121 B.5.3 Biomarker vector r

122 At a carryover half-life of one week,

$$123 \quad r^{\text{covar}} = (0.4500, 0.2250, 0.1125, 0.4500)^\top, \quad r^{\text{orig}} = (0.4500, 0.4500, 0.4500, 0.4500)^\top.$$

124 The two off-drug entries halve and quarter under **covar**, tracking $\phi_t = (1/2, 1/4)$, whereas
 125 under **orig** they take the full c_{bm} , because carryover has left the response mean nonzero and
 126 the step therefore returns the full on-drug value.

127 Removing carryover collapses the difference entirely:

$$128 \quad r^{\text{covar}} = r^{\text{orig}} = (0.4500, 0.0000, 0.0000, 0.4500)^\top \quad (t_{1/2} = 0).$$

129 This pair of displays is the essential comparison. The vector r is the only channel through
 130 which the biomarker-treatment interaction enters the covariance, and an interaction is identi-
 131 fied by r varying between on-drug and off-drug occasions. At $t_{1/2} = 0$ the two implementations
 132 produce the same vector. At any positive half-life r^{orig} is constant, so it carries no on-drug
 133 versus off-drug contrast, while r^{covar} retains one. Across the grid $t_{1/2} \in \{0, 0.5, 1.0\}$ used in
 134 this paper, **orig** therefore supplies one cell identical to **covar** and two in which its interaction
 135 channel is degenerate.

136 B.6 The three data-generating processes

137 The table below lists only the entries in which the three processes differ, grouped by what
 138 each group governs. Rows in which all three agree are recorded in prose beneath it, so that
 139 the table can be read as a list of differences rather than a full specification.

	mean	covar	orig
<i>Where the interaction is written</i>			
Moment used	μ	Σ	Σ
<i>Correlation structure</i>			
A_c within-factor	$\rho_c^{ w_t - w_s }$	$\rho_c^{ w_t - w_s }$	ρ_c (constant)
$C_{cc'}$ off-diagonal	$c_{\times} \rho^{ w_t - w_s }$	$c_{\times} \rho^{ w_t - w_s }$	$c_{\times}$ (constant)
<i>Interaction channel</i>			
r_t , on drug	0	c_{bm}	c_{bm}
r_t , off drug	0	$c_{bm} \phi_t$	$c_{bm} \mathbf{1}[\phi_t > 0]$
μ_{BR} shift, on drug	$\beta_{bm} b_i \sigma_{BR}$	none	none
<i>Carryover on the BR mean</i>			
Decay anchored at	μ_{BR,t^*}	μ_{BR,t^*}	$\mu_{BR,t-1}$

Notation. μ and Σ are the mean vector and covariance matrix of the multivariate normal from which each participant's trajectory is drawn. c indexes the three response factors, $c \in \{tv, pb, br\}$, and t indexes measurement occasions, which fall at cumulative weeks $w_1 < \dots < w_n$.

A_c is the $n \times n$ within-factor block of factor c , giving the correlation between that factor at two occasions. $C_{cc'}$ is the $n \times n$ cross-factor block between two distinct factors. Its diagonal holds same-occasion entries and its off-diagonal different-occasion entries.

ρ_c is the within-factor autocorrelation parameter of factor c (`c.tv`, `c.pb`, `c.br` in the code). ρ without a subscript is the corresponding decay rate applied to cross-factor entries. c_1 is the same-occasion cross-factor correlation (`c.cf1t`) and $c_\times$ the different-occasion one (`c.cfct`). The exponent $|w_t - w_s|$ is the gap in cumulative weeks between occasions t and s , not the difference in occasion index.

r_t is the correlation between the biomarker and the response factor at occasion t , the vector through which covariance moderation carries the interaction. c_{bm} is the moderation parameter, $t_{sd,t}$ the time since drug discontinuation, and $t_{1/2}$ the carryover half-life. $\phi_t = (1/2)^{t_{sd,t}/t_{1/2}}$ is the residual fraction of the on-drug response at an off-drug occasion that follows drug exposure, and $\phi_t = 0$ otherwise, including throughout when $t_{1/2} = 0$ and at off-drug occasions before first exposure. `covar` and `orig` are nonzero at the same occasions; they differ in whether r_t scales with the residual fraction or steps to its full on-drug value.

μ_{BR} is the response factor's mean trajectory and $b_i = (B_i - \mu_{bm})/\sigma_{bm}$ the participant's standardized biomarker, so the mean-moderation shift $\beta_{bm} b_i \sigma_{BR}$ is expressed per standard deviation of the biomarker. t^* denotes the most recent on-drug occasion and $t - 1$ the immediately preceding occasion, which is the whole of the difference in the final row: all three processes apply the same factor $(1/2)^{t_{sd,t}/t_{1/2}}$ and differ only in what they apply it to.

BL is baseline symptom severity, B the biomarker, and TV , PB , BR the natural-history, placebo-belief and pharmacological response factors respectively.

Five entries agree across all three processes and are omitted above. The same-occasion cross-factor correlation is c_1 throughout. The baseline row and column are those of the identity. The biomarker correlates with neither the natural-history nor the placebo-belief factor. The off-drug μ_{BR} shift is absent everywhere, since mean moderation applies its shift only at on-drug occasions. And all three apply the same exponential decay factor $(1/2)^{t_{sd,t}/t_{1/2}}$, differing only in the value it multiplies, which is the last row of the table.

Three features of this comparison are worth drawing out.

First, `mean` and `covar` share their entire covariance matrix apart from the vector r . Setting r to zero and adding the mean shift converts one into the other, with every within-factor and cross-factor block untouched. This is the decomposition of Appendix A made concrete: the two architectures are not different models but the two places a single joint distribution can carry an interaction.

Second, `orig` differs from `covar` on three rows simultaneously, in A_c , in the off-diagonal of $C_{cc'}$, and in r . No single row can therefore be held responsible for any difference in results obtained under it, and a comparison against `orig` bounds the joint effect of the three differences rather than decomposing them. Separating them would require running the correlation structure and the off-drug coupling rule as crossed factors. The carryover row is a fourth difference, but it is of our making rather than Hendrickson's, and Section B.7 explains why it should be held fixed across arms rather than counted among the architectural differences.

Third, the off-drug r_t row is the consequential one, and it is consequential because of how r_t is sized, not where it is nonzero. The two covariance-based processes switch the interaction channel on at exactly the same occasions: on drug, and off drug after exposure whenever

164 $t_{1/2} > 0$. We confirmed this against the package's own matrix construction on every path
 165 of the three designs at each half-life in the grid. They differ in what happens there. Under
 166 **covar** the correlation is scaled by the residual fraction ϕ_t , so that it falls in proportion to the
 167 residual drug effect in the response mean. Under **orig** it steps to the full c_{bm} as soon as any
 168 residual effect is present, however small. With carryover, the published coupling is therefore
 169 constant at every post-exposure occasion, and its interaction channel loses the on-drug versus
 170 off-drug contrast on which identification depends.

171 The carryover row requires explanation, since the three processes agreed on it until recently
 172 and a reader comparing code will want to identify the difference quickly.

173 All three adjust the response mean at an off-drug occasion by adding a decayed fraction of
 174 the previous response. They differ in what is decayed. Writing t^* for the most recent on-drug
 175 occasion,

$$176 \quad \underbrace{\mu_{BR,t} + \mu_{BR,t^*} (1/2)^{t_{sd,t}/t_{1/2}}}_{\text{mean, covar}} \quad \text{against} \quad \underbrace{\mu_{BR,t} + \mu_{BR,t-1} (1/2)^{t_{sd,t}/t_{1/2}}}_{\text{orig}}.$$

177 The published form multiplies the *already adjusted* mean at $t - 1$ while $t_{sd,t}$ is measured
 178 cumulatively from discontinuation, so elapsed time is counted twice. The corrected form
 179 applies the cumulative factor once, to the mean at discontinuation.

180 Two consequences make the difference easy to spot in output. The first off-drug occasion of
 181 any run agrees exactly under both forms, since there is nothing yet to inherit. Divergence
 182 begins at the second. In the worked example of Section B.5, at $t_{1/2} = 1$ week, occasion 3
 183 carries $0.25 \mu_{BR,1}$ under the corrected form and $0.125 \mu_{BR,1}$ under the published one, a factor
 184 of two. The gap widens with each further off-drug occasion, so designs with long uninterrupted
 185 washouts show it most.

186 We corrected **mean** and **covar** in September 2026 and leave **orig** as published, since the object
 187 of that column is to describe the published code rather than to improve it.

188 The consequence for the re-run flagged in Section 2.2.4 should be settled before that run
 189 is designed. A comparison of **covar** against **orig** now carries one difference beyond the
 190 four enumerated there, and it is of our making. Running **covar** as corrected asks what the
 191 architecture delivers today and credits the correction to **covar**. Running **covar** with the
 192 recursion restored isolates the four architectural differences and answers the question Section
 193 2.2.4 poses. We prefer the second for the headline comparison, with the first reported alongside
 194 it, so that the cost of the correction remains visible.

195 B.7 Summary of differences

	Component	orig	covar
196	A_c interior	ρ_c	$\rho_c^{ w_t-w_s }$
	$C_{cc'}, t \neq s$	$c_\times$	$c_\times \rho^{ w_t-w_s }$
	r_t , off drug	$c_{bm} \mathbf{1}[\phi_t > 0]$	$c_{bm} \phi_t$
	Carryover on μ_{BR}	recursive	anchored

197 The table lists only the entries that differ. Everything else in the correlation matrix agrees:
 198 the same-occasion cross-factor entry c_1 , the on-drug value c_{bm} of the coupling vector, the BL
 199 row and column, which are those of the identity in both processes, and the $B-TV$ and $B-PB$
 200 entries, which are zero in both. That last agreement is worth stating rather than leaving
 201 implicit, since it is the assumption that the biomarker couples to no factor other than the
 202 response, and a good deal rests on it.

203 The four differences are independent decisions rather than consequences of one another, and
 204 they are not of equal standing. Only one of them repairs a defect. We set out what each is
 205 for, since a reader is entitled to ask whether a departure from published code was necessary
 206 or merely preferred.

207 **The off-drug coupling (r_t off drug) corrects a defect.** This is the one substantive
 208 repair. Both processes switch the coupling on at the same occasions, so the defect is not
 209 where the published code opens the channel but how far. Under the published rule the off-
 210 drug correlation takes the full on-drug value c_{bm} at any occasion carrying residual drug effect,
 211 however small that residual is. The coupling is therefore a step function of the residual effect
 212 rather than a function of its size: an off-drug occasion at which the response mean has fallen
 213 to a few percent of its on-drug value is assigned the same biomarker-response correlation
 214 as an on-drug occasion. In a design whose paths all begin on drug, the coupling vector
 215 is then constant at every positive half-life, and the interaction channel carries no on-drug
 216 versus off-drug contrast at all. Since the interaction is identified by precisely that contrast,
 217 the introduction of any carryover at all removes it, which inverts the graded effect residual
 218 exposure should have. Architecture B instead scales the coupling by the residual fraction ϕ_t ,
 219 so that the biomarker-response correlation falls in proportion to the residual drug effect in
 220 the response mean. The ceiling discussion below shows that the published rule admits a wider
 221 nominal c_{bm} only because its vector has stopped carrying an interaction, which is a further
 222 symptom of the same defect rather than a separate finding.

223 The repository was itself revised in 2024, replacing the step with a graded coupling. We do
 224 not offer that revision as independent corroboration. The commit in question (8609f12, 6
 225 May 2024) is authored by Hendrickson but its message describes it as the present author's
 226 edits curated into working code, so the two lines of reasoning are not independent. The
 227 argument for the graded form rests on the identification point above, not on the repository's
 228 later history.

229 **The within-factor form (A_c) is a modeling choice.** It is adopted for clinical plausibility,
 230 since nearby measurements should be more strongly correlated than distant ones, and for the
 231 wider feasible range discussed below. It repairs nothing. Compound symmetry is a defensible
 232 model of within-participant correlation, and it is the standard implication of a participant
 233 random intercept.

234 **The cross-factor form ($C_{cc'}$, $i \neq j$) follows from the previous row.** It is not listed
 235 among the four differences of Section 2.2.4, and on its own it would be difficult to defend.
 236 Taken without the AR(1) change it lowers the feasible range of c_{bm} substantially, which is
 237 the numerical signature of an incoherent structure: a constant cross-factor term alongside a
 238 decaying within-factor block asserts that a factor's correlation with a *different* factor at eight
 239 occasions' remove equals its correlation with that factor at the adjacent occasion, while its
 240 correlation with *itself* has decayed away. `covar` takes the two rows together for consistency,
 241 not because either entails the other.

242 **The carryover row is ours, not Hendrickson's.** All three processes originally adjusted
 243 the response mean by the recursive form, and we corrected `mean` and `covar` in September
 244 2026 for the reason given in Section B.6. The row is listed here because a reader comparing
 245 code will encounter it, but it should not be counted among the architectural differences that
 246 Section 2.2.4 enumerates.

247 **The feasible range of c_{bm} .** The two forms differ in how much moderation the covariance
 248 can carry. The mechanism is worth setting out, because the obvious explanation is wrong and
 249 the direction of the difference is not constant.

250 The binding constraint is not positive definiteness of the within-factor block. It is a Schur
 251 complement on the biomarker row, and we set the argument out in full because it replaces a
 252 grid search with a formula.

253 Drop the BL row and column, which are those of the identity and therefore factor out of every
 254 determinant. Order the remainder with the biomarker first and partition accordingly:

$$255 \quad R = \begin{pmatrix} 1 & \tilde{r}^\top \\ \tilde{r} & M \end{pmatrix},$$

256 where M is the $3n \times 3n$ block of the three response factors and $\tilde{r}$ is the biomarker coupling
 257 vector extended by zeros, equal to r on the BR positions and vanishing on the others. The
 258 Schur complement of M in R is the scalar $1 - \tilde{r}^\top M^{-1} \tilde{r}$, and R is positive definite if and only
 259 if M is positive definite and that scalar is positive.

260 The scalar has a statistical reading which makes the constraint intuitive rather than merely
 261 algebraic. On the correlation scale the biomarker has unit variance. The quantity $\tilde{r}^\top M^{-1} \tilde{r}$
 262 is the coefficient of determination from regressing the biomarker on the whole set of response
 263 variables, and $1 - \tilde{r}^\top M^{-1} \tilde{r}$ is the residual variance of the biomarker once that regression

is removed. The condition therefore says only this: the biomarker may not be more than perfectly predictable from the response factors. An analyst who specifies correlations strong enough to imply an R^2 above one has not specified a covariance matrix at all, having asked for a negative variance.

Writing $\tilde{r} = c_{bm}v$, where v carries the on-drug and off-drug pattern of Section B.4 with the moderation parameter factored out, the condition becomes $c_{bm}^2 v^\top M^{-1}v < 1$ and the exact ceiling follows:

$$c_{bm}^* = (v^\top M^{-1}v)^{-1/2}.$$

The bound is attained rather than approached. At $c_{bm} = c_{bm}^*$ the implied R^2 is exactly one, the residual variance is exactly zero, and the smallest eigenvalue of R reaches zero at the same point. We have verified this correspondence directly on small cases, where the three quantities can be watched together as c_{bm} is increased.

Two features of this expression matter. The ceiling depends on the correlation structure only through how the coupling pattern v sits relative to M^{-1} , and it depends on v , which differs between the two processes and, under `orig`, changes with the carryover half-life.

That second point reverses the comparison. At $t_{1/2} = 0$ the `orig` coupling is the on-drug indicator and its ceiling is the lower of the two. At any positive half-life the `orig` coupling becomes constant, as Section B.4 shows, and a constant v aligns with the direction in which M is best conditioned. Its ceiling is then the higher of the two. The apparent gain is worthless: it arises precisely because the vector has lost the on-drug versus off-drug contrast on which identification depends. A process that can carry a larger nominal c_{bm} while carrying no interaction is not thereby preferable.

We note that the number of occasions is not what separates the two forms. Both ceilings fall as occasions are added, in similar proportion. The smallest eigenvalue of the compound-symmetric block M is $(1 - \rho_c) - (c_1 - c_\times)$, which contains neither n nor the number of factors. An earlier version of this appendix attributed the difference to a constraint on ρ_c that tightens with n ; that attribution was incorrect.

[TO RE-VERIFY: the relations above were derived analytically and confirmed on an evenly spaced schedule. Design-specific thresholds for the CO, Hybrid and OL+BDC schedules of Section 2.3 have not been computed, and none is quoted here. Evaluate c_{bm}^* on those three designs, under 58b32a9 for `orig`, before restoring numeric thresholds.]

B.8 The mean-covariance contrast

Sections B.6 and B.7 compare Architecture B against the published implementation. We give the corresponding comparison for the contrast this paper actually reports, mean moderation against covariance moderation, because it is much the cleaner of the two.

300 **Where the interaction is written.** The two processes place the interaction in different
 301 objects, and the first table records only that.

	Component	mean	covar
302	r_t , on drug	0	c_{bm}
	r_t , off drug	0	$c_{bm} \phi_t$
	μ_{BR} shift, on drug	$\beta_{bm} b_i \sigma_{BR}$	none

303 Everything else is shared. Both processes use the AR(1) within-factor form, the same cross-
 304 factor blocks, the same anchored carryover adjustment, the same baseline row and column,
 305 and the same zero correlations between the biomarker and the two non-response factors.
 306 Setting r to zero and adding the mean shift converts one process into the other, and nothing
 307 else moves.

308 **A caution about the zero.** The entry $r_t = 0$ describes the covariance matrix from which
 309 the data are drawn. It does not say that the biomarker and the response are uncorrelated
 310 under mean moderation. The shift is applied after the draw, and it induces an association
 311 of its own. Reading the zero as an absence of biomarker-response dependence would be a
 312 mistake, and the second table is included to forestall it.

313 **What each process realizes.** Writing σ for σ_{BR} , taking an on-drug occasion, and setting
 314 the two parameters to their matched value $\beta_{bm} = c_{bm}$, written c_{bm} in the table, the two
 315 processes differ in every second moment while agreeing exactly on the first.

	Quantity	mean	covar
	$E[BR_t B_i]$	$\mu_{BR} + c_{bm} \frac{\sigma}{\sigma_{bm}} (B_i - \mu_{bm})$	same
	Slope of BR_t on B_i	$c_{bm} \sigma / \sigma_{bm}$	same
316	$SD[BR_t B_i]$	σ	$\sigma \sqrt{1 - c_{bm}^2}$
	Marginal $SD[BR]$	$\sigma \sqrt{1 + c_{bm}^2}$	σ
	$Cor(B_i, BR_t)$	$c_{bm} / \sqrt{1 + c_{bm}^2}$	c_{bm}

317 At $c_{bm} = 0.45$ and $\sigma = 8$ these give a realized correlation of 0.410 under mean moderation
 318 against 0.450 under covariance moderation, conditional standard deviations of 8.00 against
 319 7.15, and marginal standard deviations of 8.77 against 8.00. A two-million-draw simulation
 320 reproduces each to three decimal places.

321 Three things follow, and together they are the reason we prefer this comparison to the one
 322 against **orig**.

323 First, the conditional means agree exactly. The mean shift is calibrated to reproduce the
 324 conditional expectation that Architecture B's covariance implies, so the two processes present

the same regression slope to the analysis model. Whatever difference Section 3 reports is not a difference in the quantity being estimated.

Second, they reach that common slope by opposite routes, and this is the substantive content of the architecture distinction. Covariance moderation holds the marginal variance fixed and shrinks the conditional variance, so the biomarker explains variation already present in the population. Mean moderation holds the conditional variance fixed and inflates the marginal, so the biomarker introduces new between-patient spread. Appendix A.2 puts the same point qualitatively: a biomarker that indexes an unobserved responder status against one that fixes the drug effect. The dilution of the correlation from 0.450 to 0.410 is the arithmetic residue of the added variance.

Third, the comparison is clean in a way the `orig` comparison is not. Section B.7 lists four differences, and a power difference observed against `orig` cannot be attributed to any one of them without further work. Here there is a single difference, expressed in two complementary rows, and its consequences are the five moments tabulated above.

We note one practical implication. Because the marginal variances differ, absolute power is not comparable between the two architectures at a common nominal c_{bm} , and Section 3 accordingly compares the profile across carryover within each architecture rather than the levels between them.

The two processes do not, however, carry the interaction at the same occasions. Mean moderation applies its shift only where the drug indicator is 1, so its interaction is absent at every off-drug occasion whatever the carryover half-life. Covariance moderation sets $r_t = c_{bm}$ at those same on-drug occasions and also carries a graded interaction, $c_{bm}\phi_t$, at off-drug occasions that follow exposure whenever $t_{1/2} > 0$. The two therefore agree exactly when $t_{1/2} = 0$ and diverge off drug as soon as carryover is present. This is the data-generating counterpart of the point made in Section 2.2.1: the analysis regressor $D_{bc,it}$ has the same graded off-drug profile as ϕ_t , so it matches covariance moderation's off-drug interaction and finds none under mean moderation. The published implementation shares covariance moderation's support but not its scaling, which is the difference Section B.7 identifies as consequential.